


Extra Class

Introduction to Transformer

Nguyen Quoc Thai



CONTENT

(1) – Attention

(2) – Transformer-Encoder


(3) – Text Classification

(4) – Vision Transformer

1 – Attention



RNNs Model

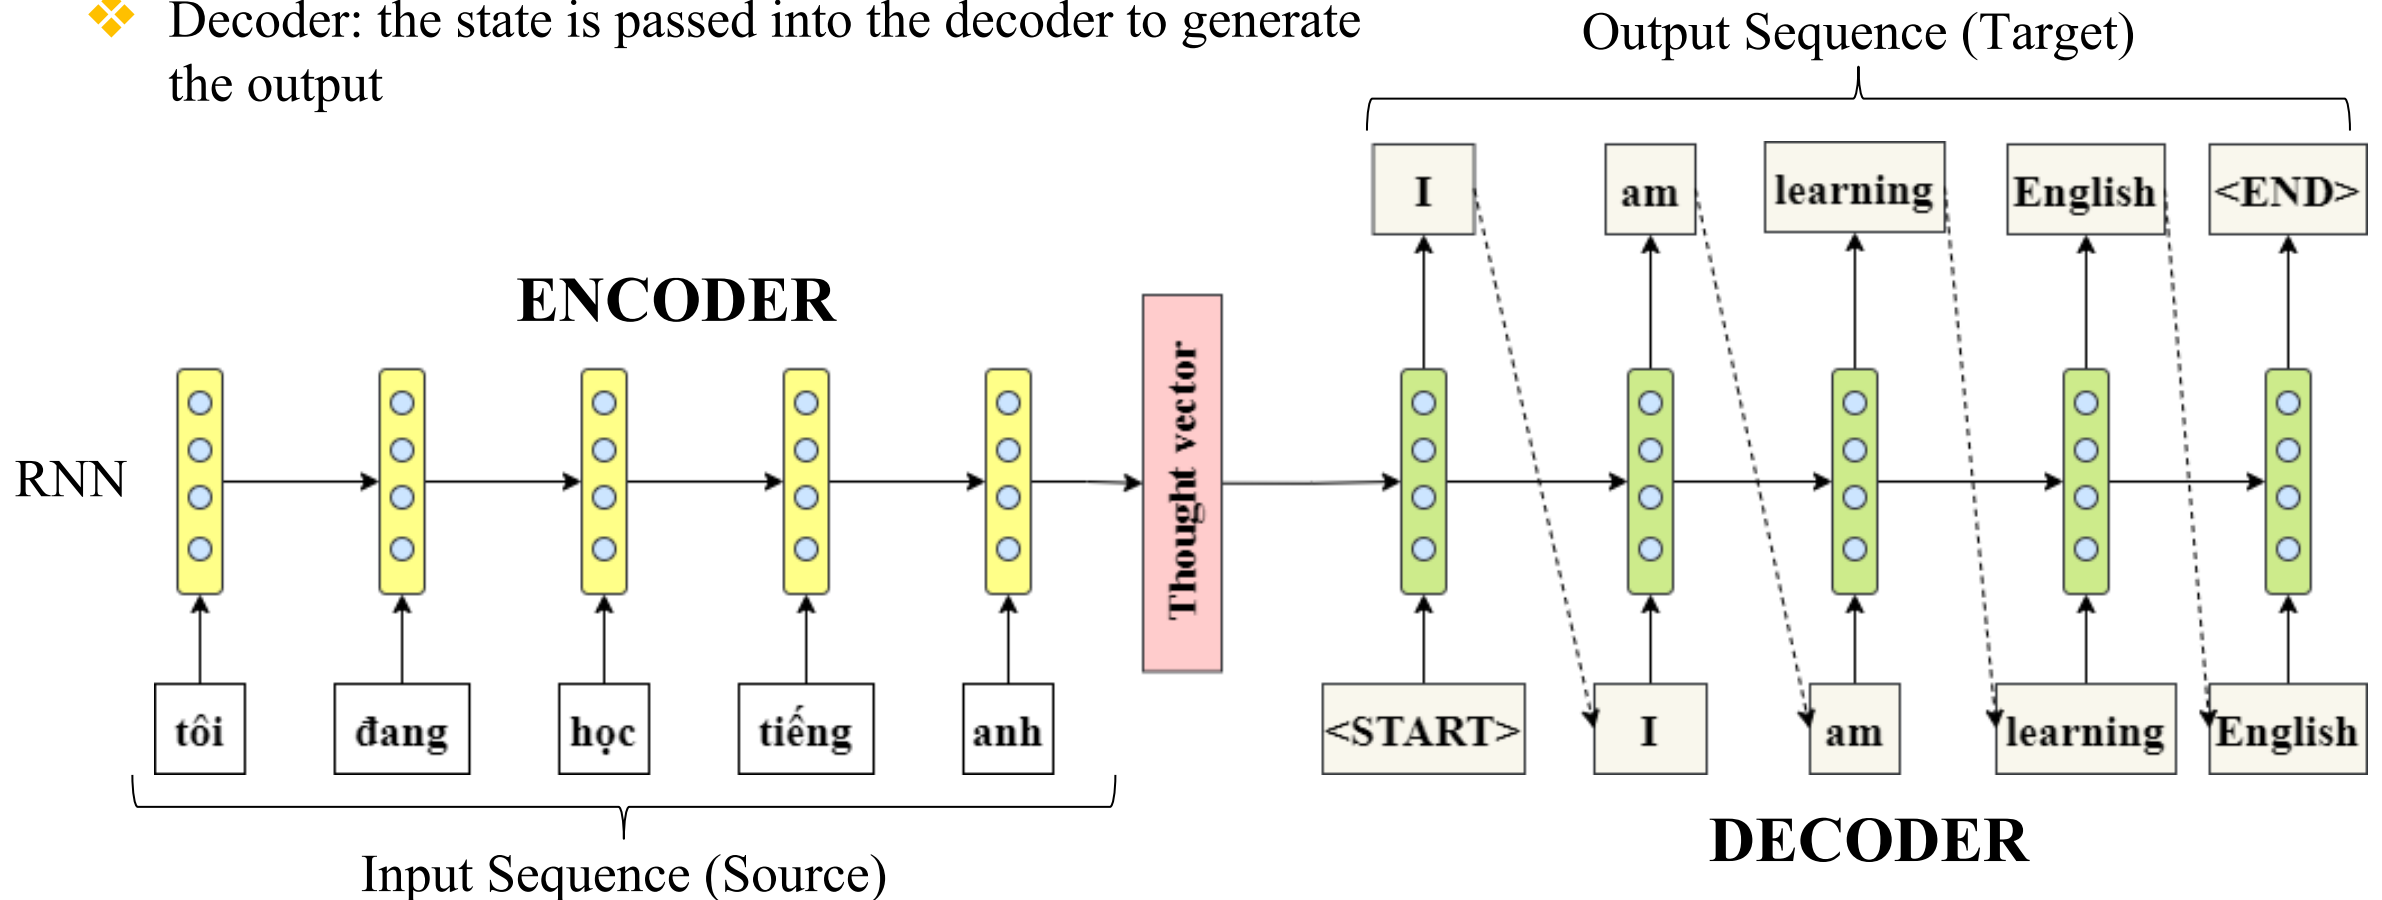


1 – Attention



Sequence-to-Sequence Architecture


- ❖ Encoder: encoding the inputs into state (thought vector)
- ❖ Decoder: the state is passed into the decoder to generate the output



1 – Attention




Scaled Dot-Product Attention



1 – Attention




Scaled Dot-Product Attention



1 – Attention




Scaled Dot-Product Attention



1 – Attention



Scaled Dot-Product Attention

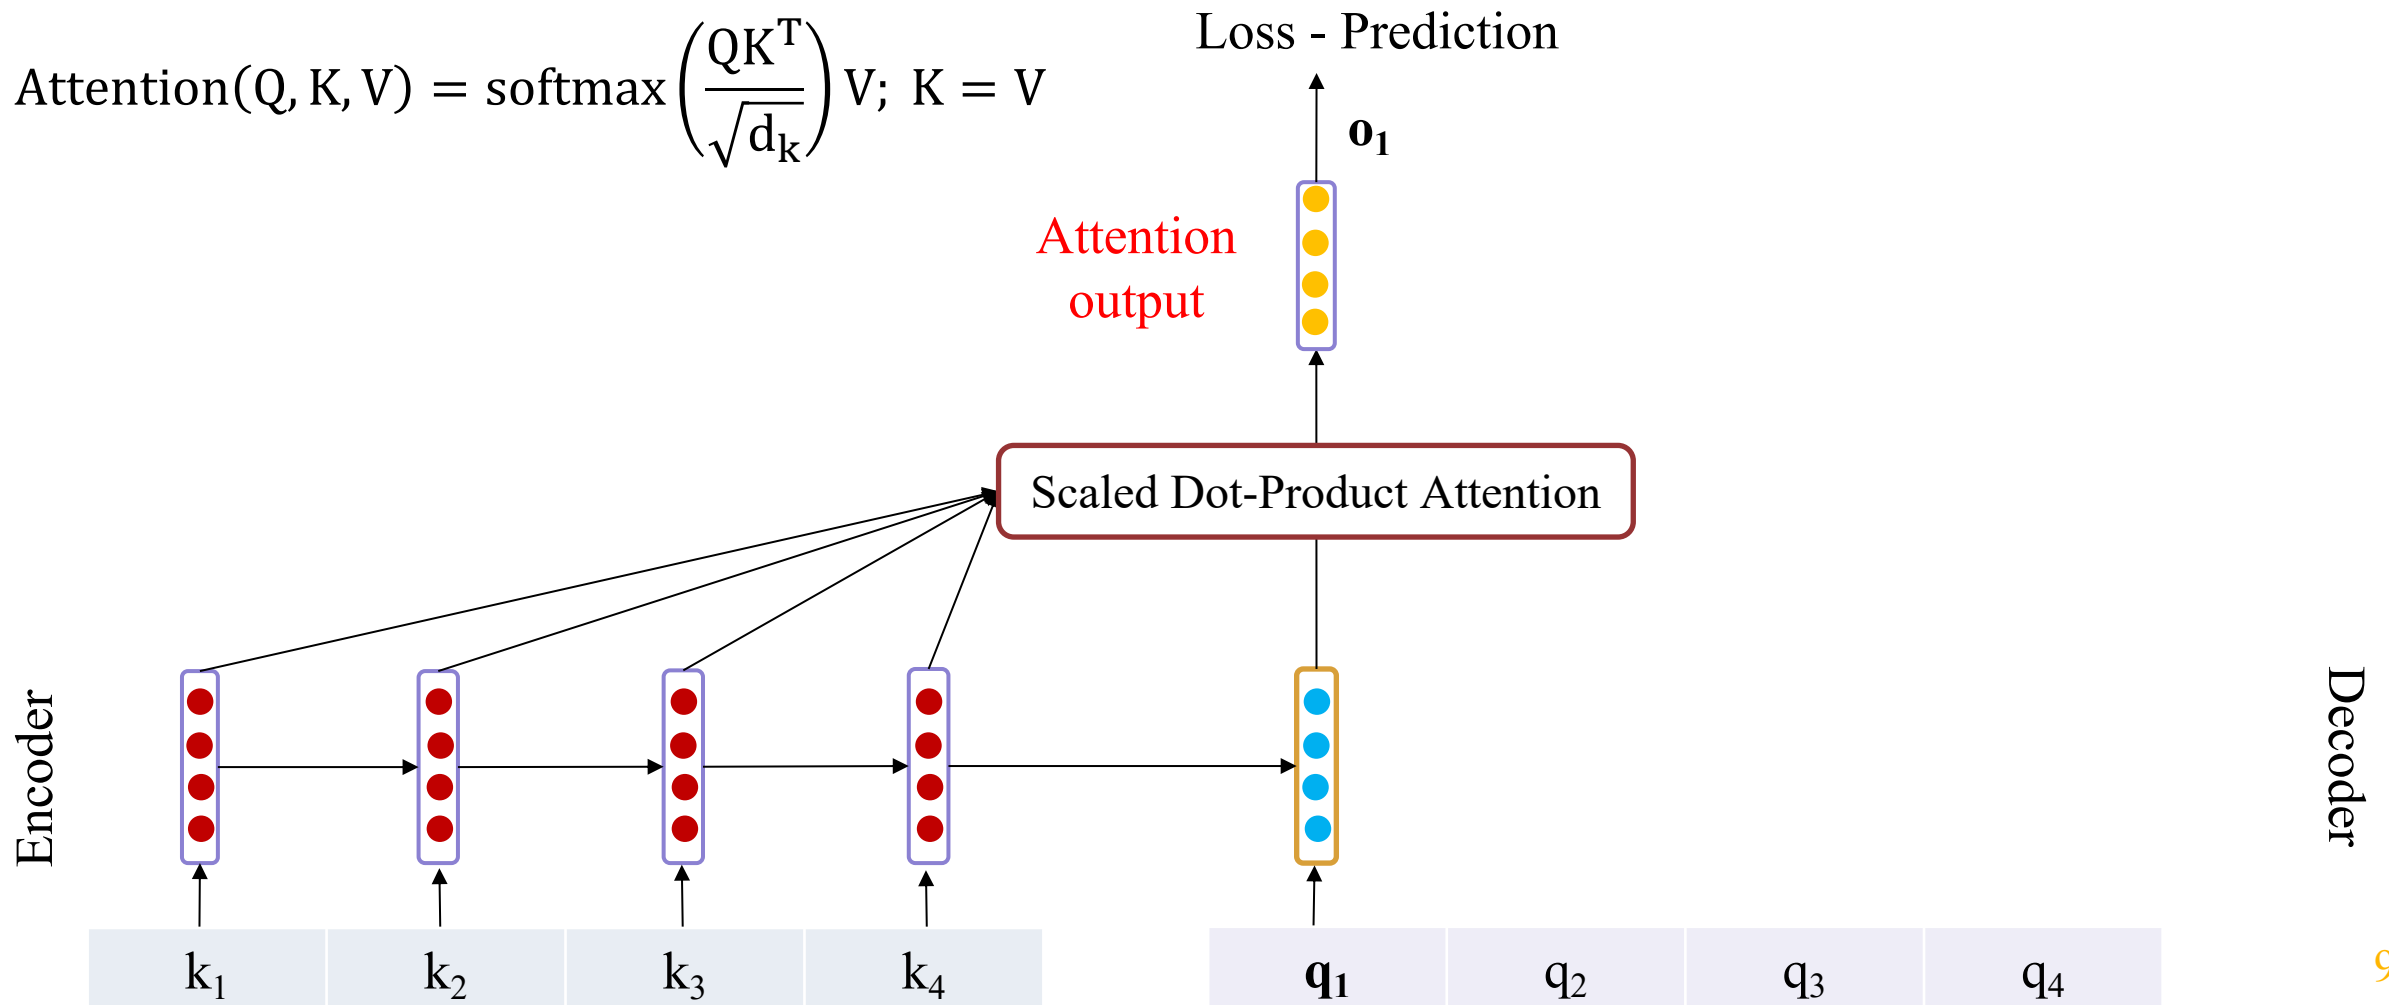


1 – Attention



Scaled Dot-Product Attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V; K = V$$





1 – Attention



Scaled Dot-Product Attention

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V; K = V$$



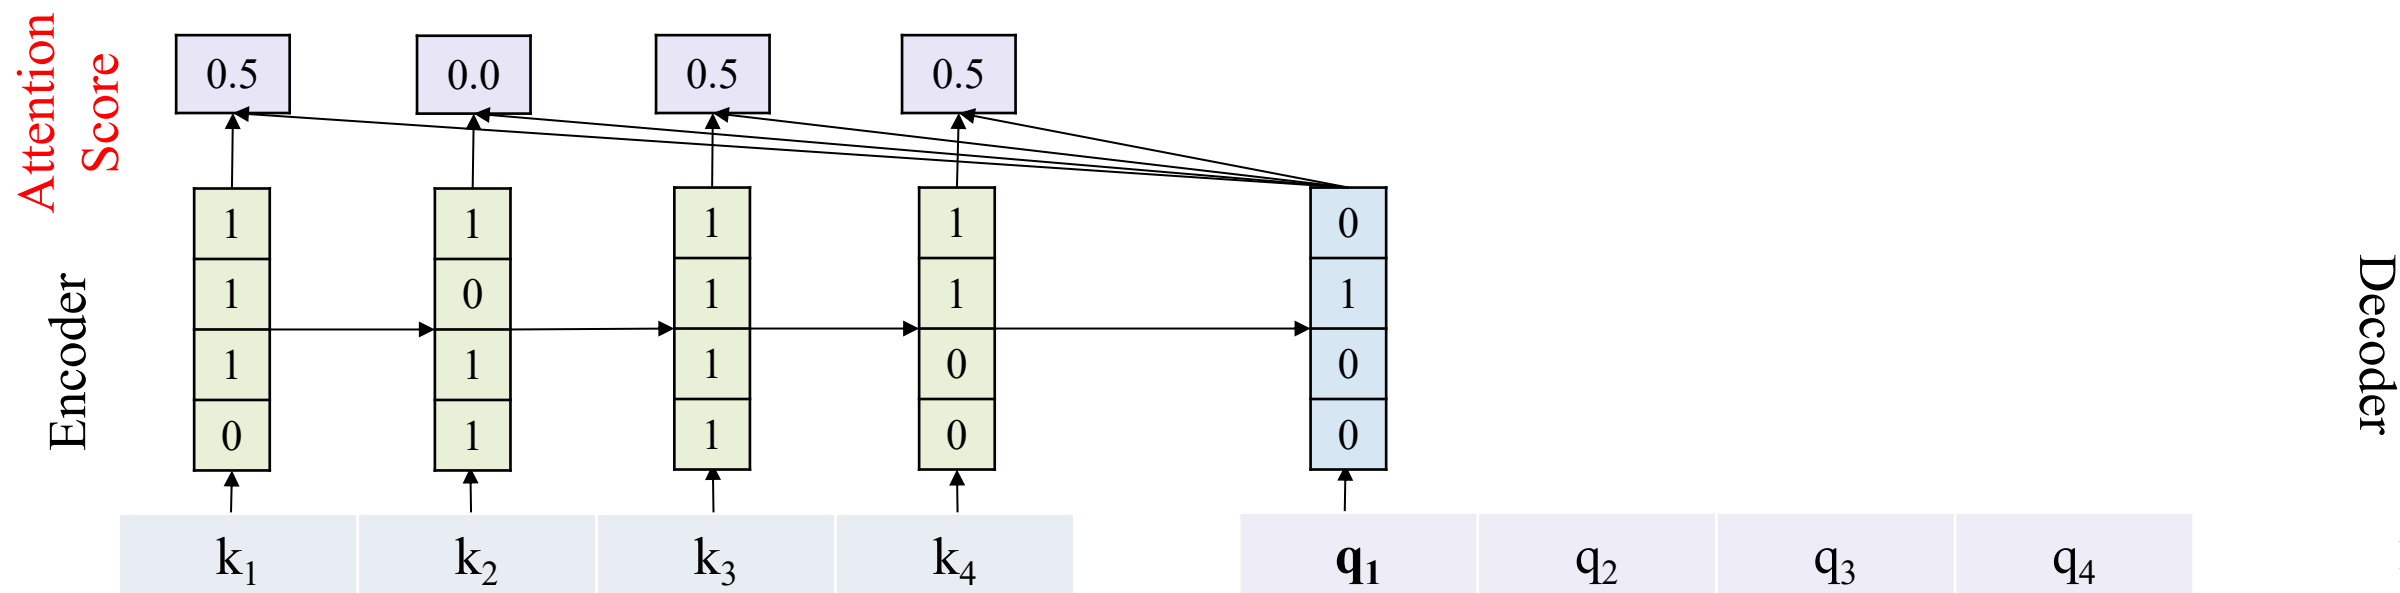


1 – Attention



Scaled Dot-Product Attention - Example

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V; K = V$$

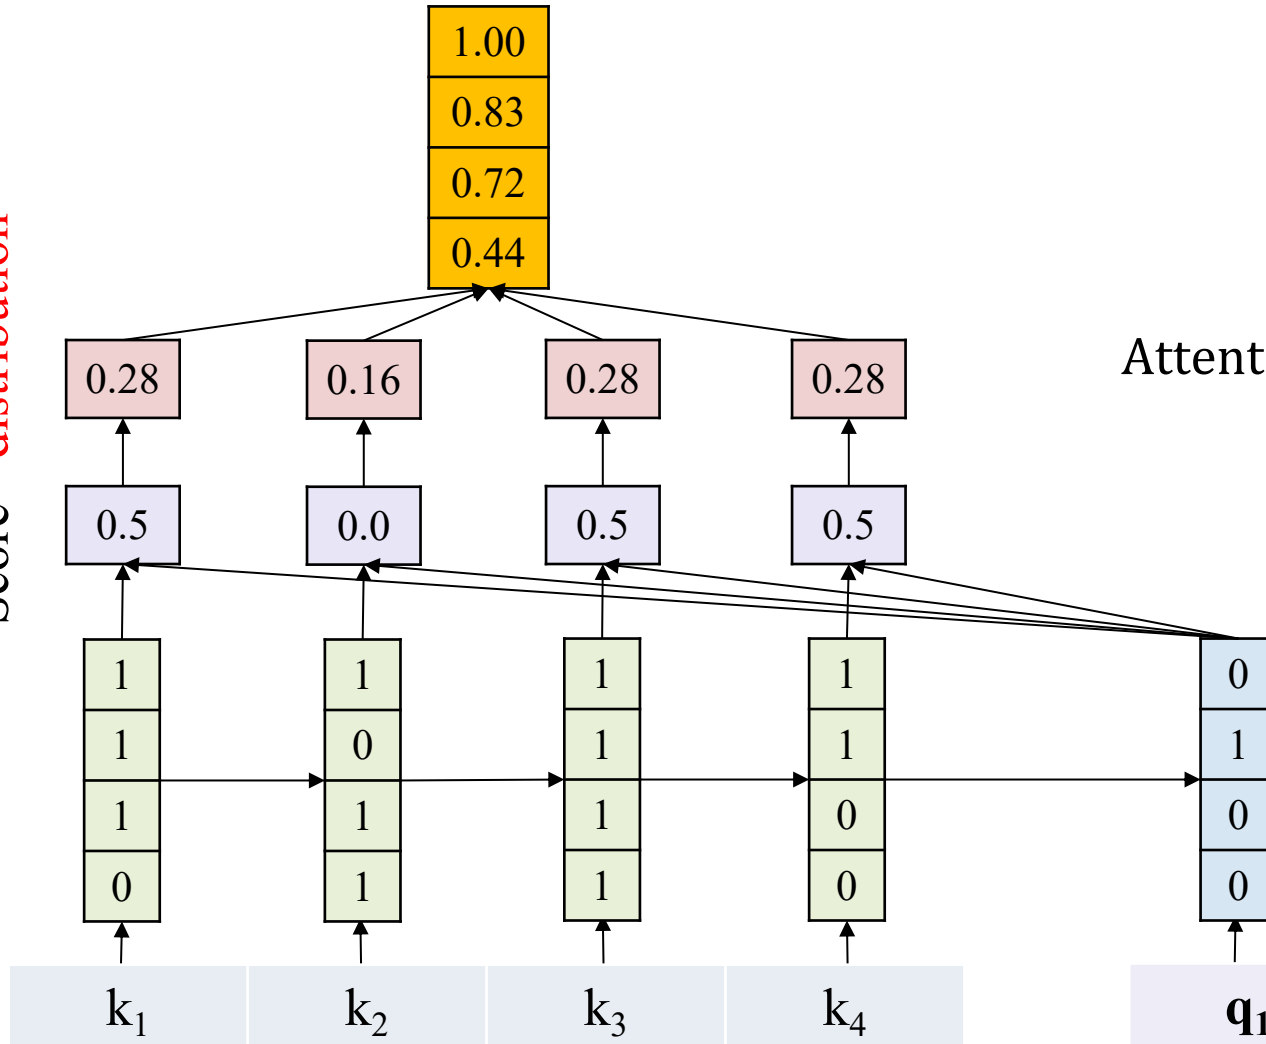


1 – Attention



Scaled Dot-Product Attention - Example

Attention
Score
Encoder



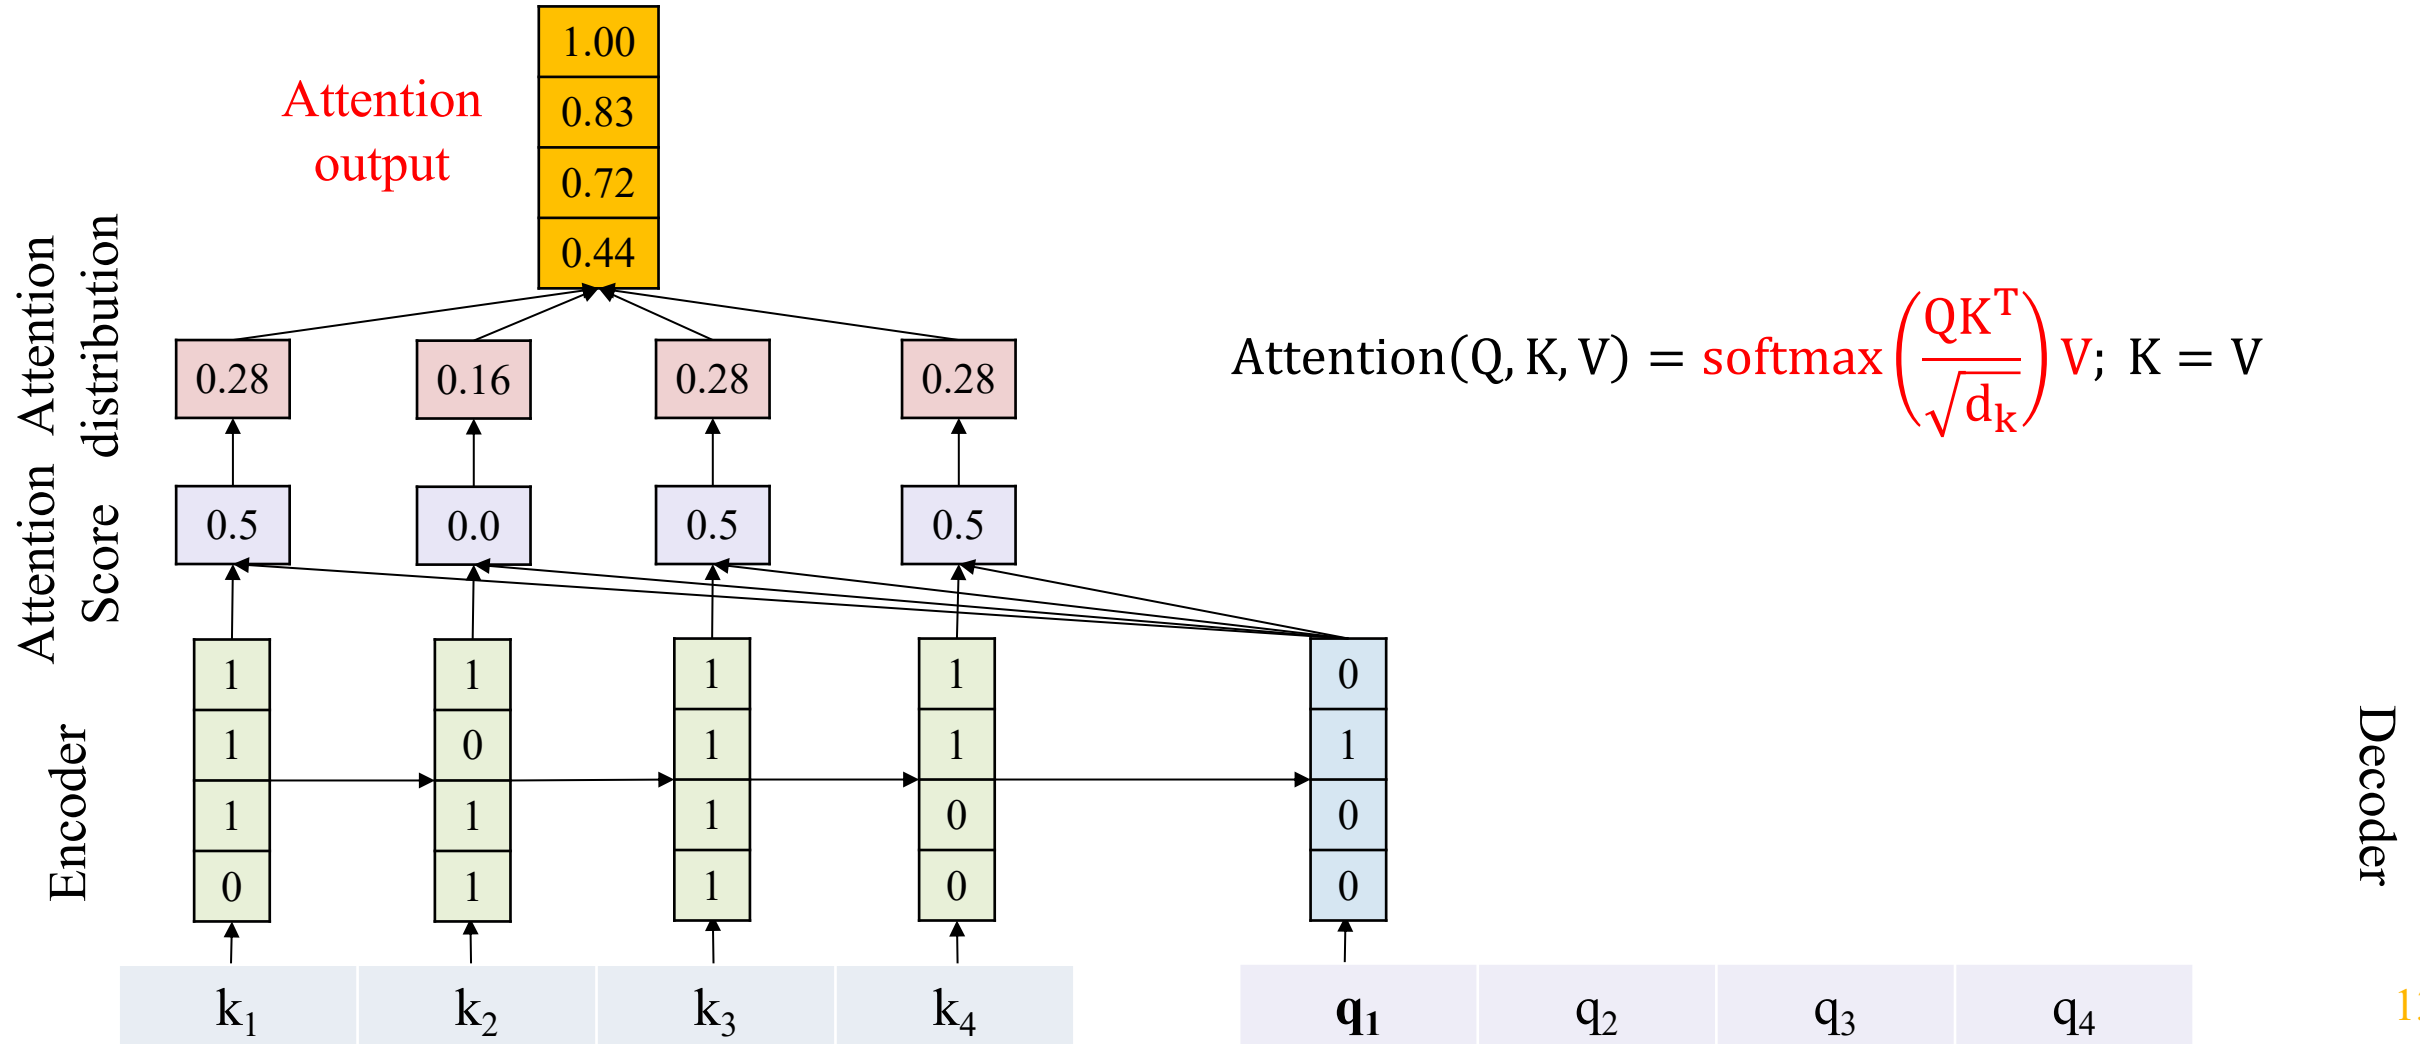
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V; K = V$$

Decoder

1 – Attention




Scaled Dot-Product Attention - Example



1 – Attention

Scaled Dot-Product Attention - Example

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V; K = V$$



1 – Attention



Scaled Dot-Product Attention - Demo

```
query = torch.randint(  
    high=2,  
    size=(1, 4, 4), # batch_size x seq_len x embedding_dim  
    dtype=torch.float32  
)  
query
```

```
tensor([[[0., 1., 0., 0.],  
         [0., 1., 0., 0.],  
         [0., 1., 0., 0.],  
         [0., 0., 1., 0.]]])
```

```
key = torch.randint(  
    high=2,  
    size=(1, 4, 4),  
    dtype=torch.float32  
)  
key
```

```
tensor([[[1., 1., 1., 0.],  
         [1., 0., 1., 1.],  
         [1., 1., 1., 1.],  
         [1., 1., 0., 0.]]])
```

```
value = key  
value
```

```
tensor([[[1., 1., 1., 0.],  
         [1., 0., 1., 1.],  
         [1., 1., 1., 1.],  
         [1., 1., 0., 0.]]])
```

```
attentionn_weight = F.scaled_dot_product_attention(  
    query=query,  
    key=key,  
    value=value  
)  
attentionn_weight
```

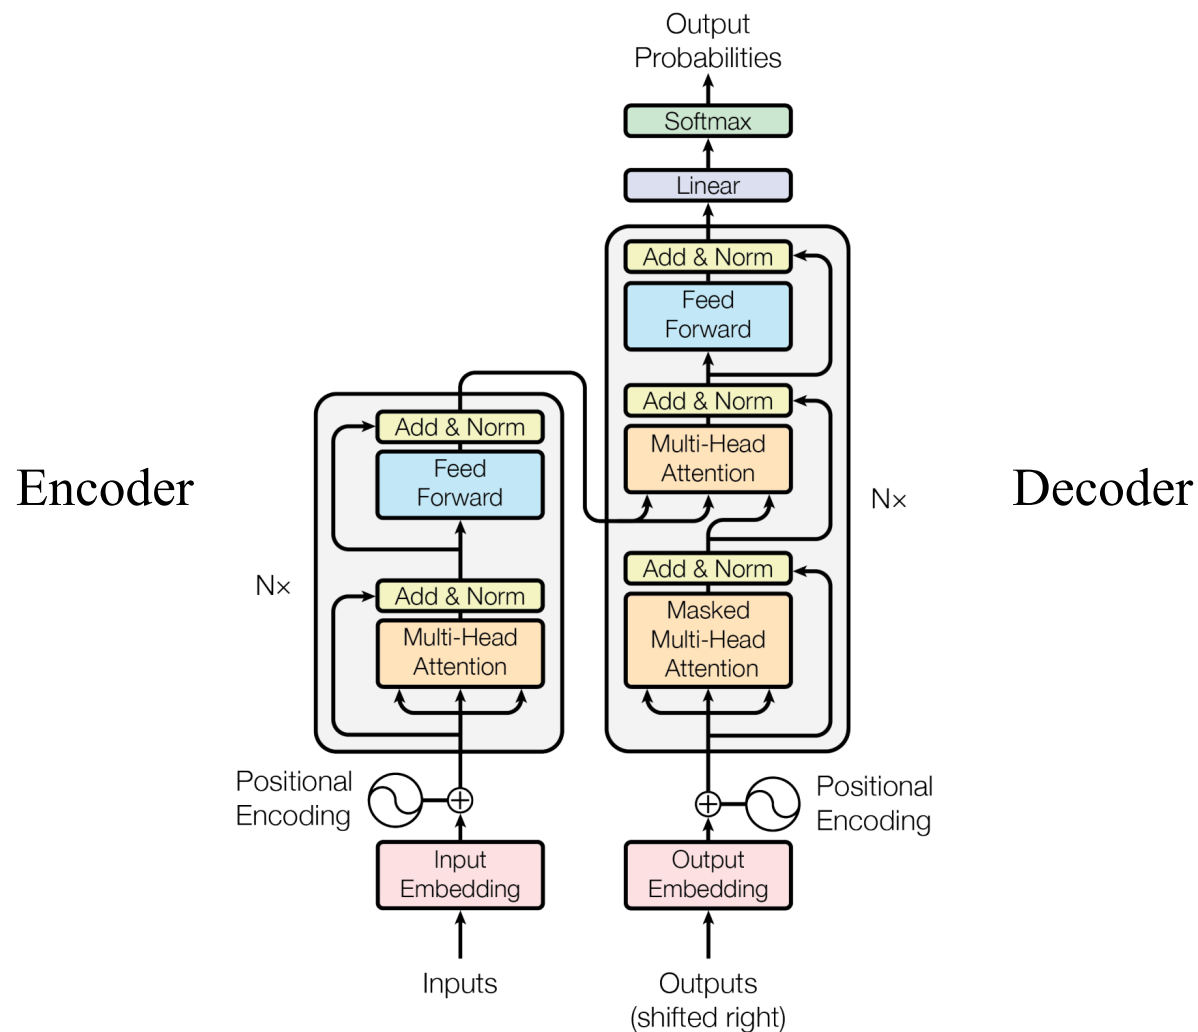
```
tensor([[[1.0000, 0.8318, 0.7227, 0.4455],  
         [1.0000, 0.8318, 0.7227, 0.4455],  
         [1.0000, 0.8318, 0.7227, 0.4455],  
         [1.0000, 0.7227, 0.8318, 0.5545]])
```

2 – Transformer-Encoder



Transformer

- ❖ Architecture:
 - N Encoder Layer
 - N Decoder Layer
- ❖ Core technique: attention
- ❖ Loss function: cross-entropy

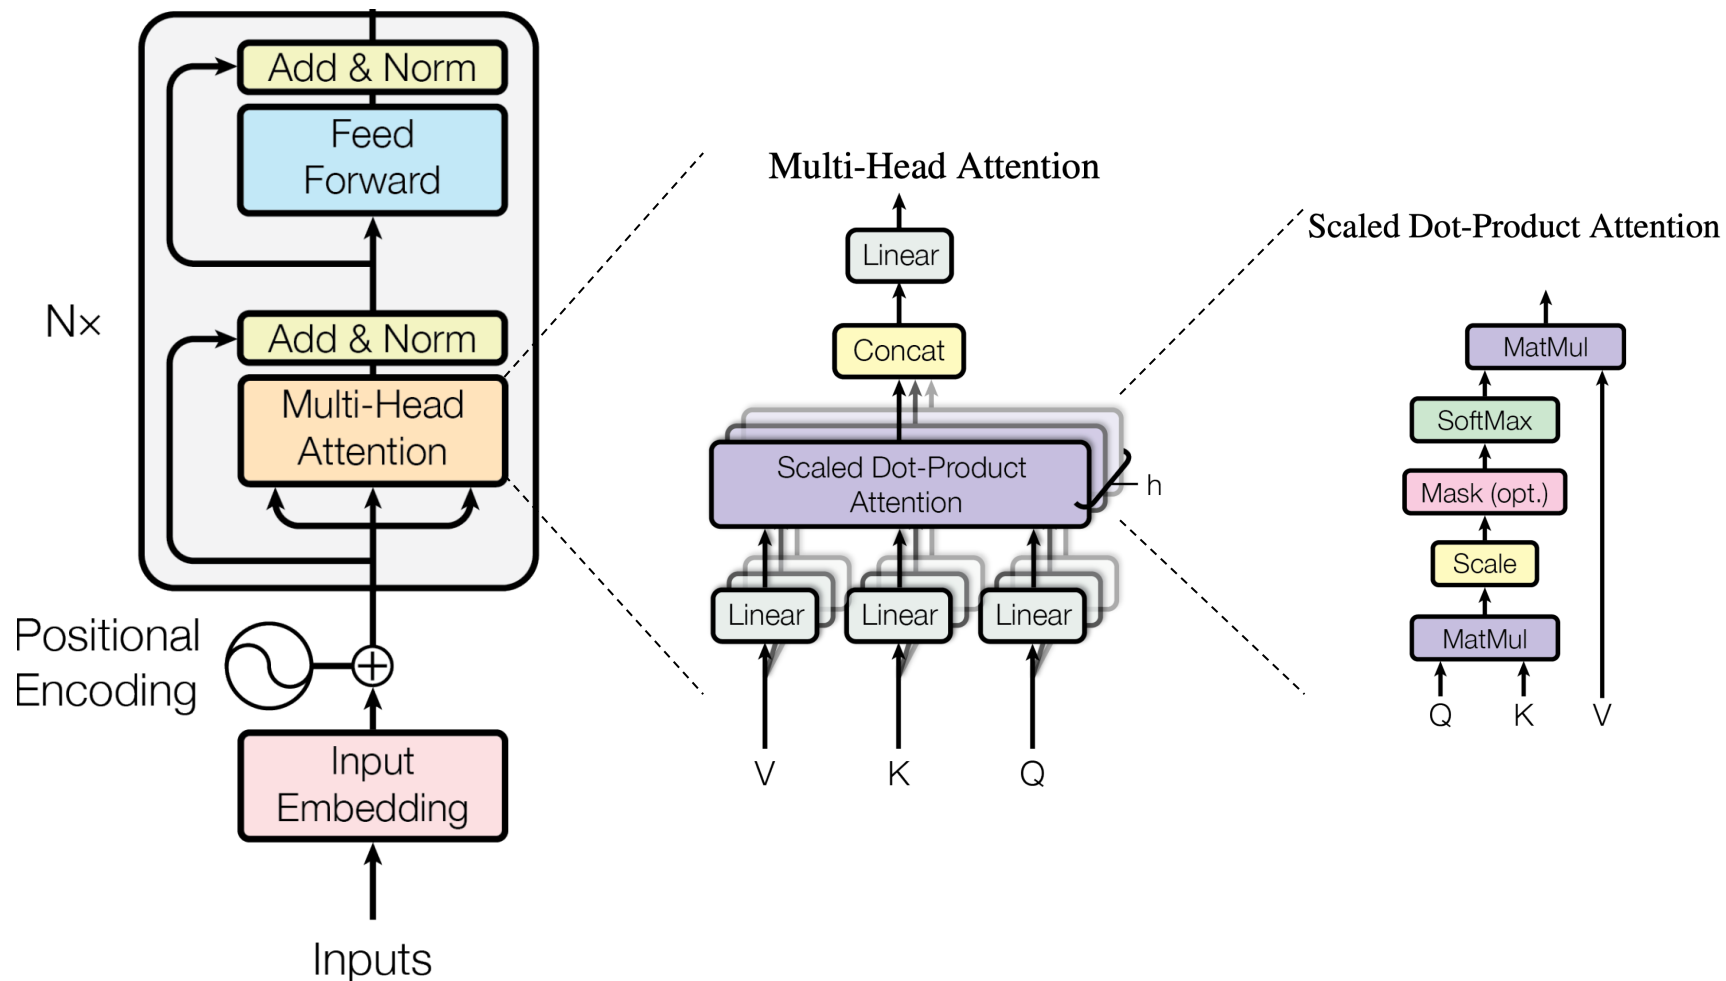


2 – Transformer-Encoder



Transformer-Encoder

- ❖ Input Embedding
- ❖ Positional Encoding
- ❖ Multi-Head Attention
- ❖ Feed Forward
- ❖ Add & Norm



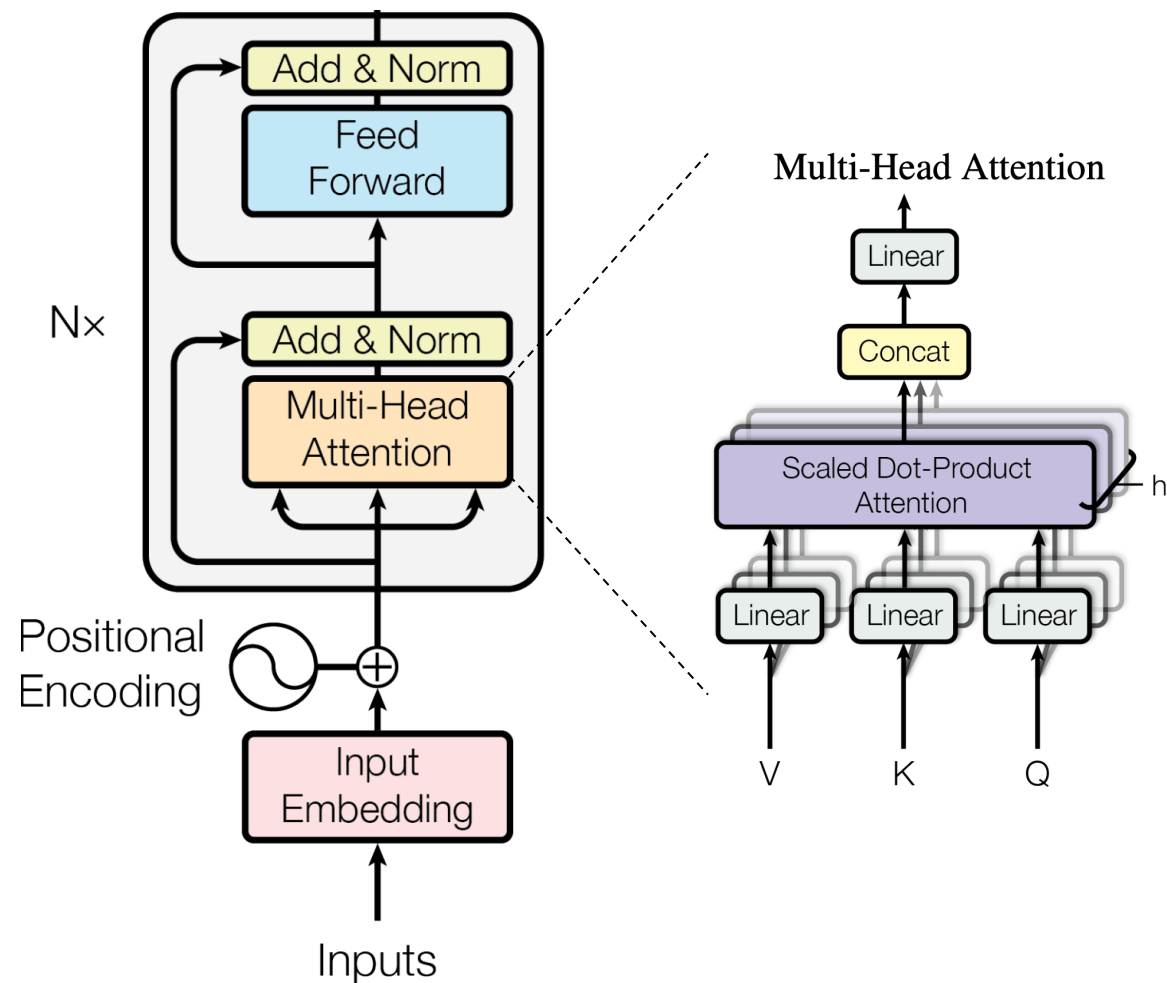
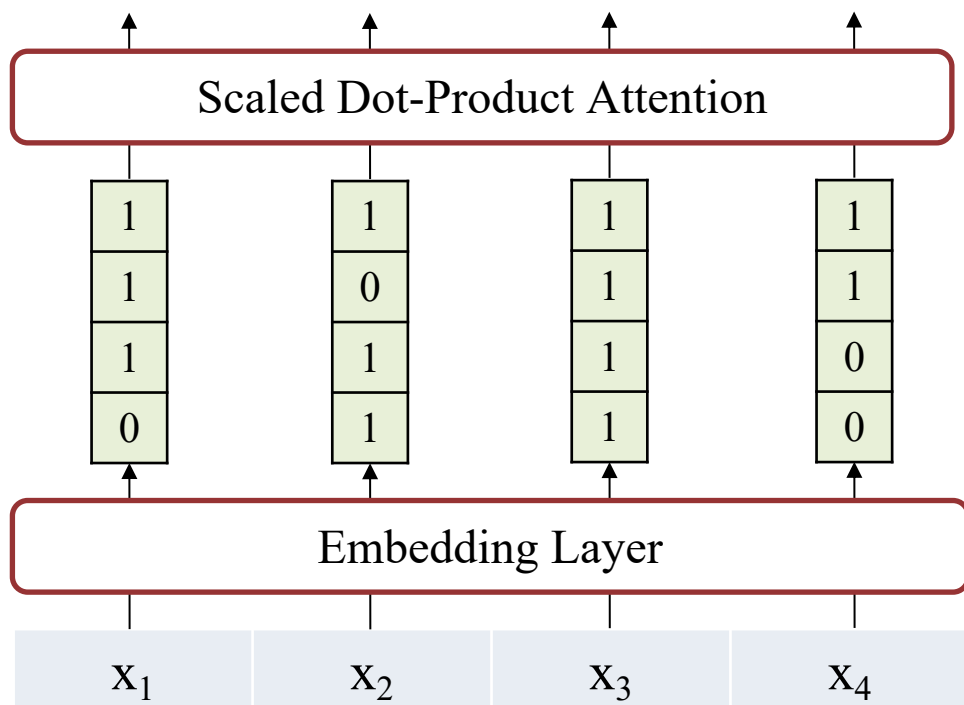
2 – Transformer-Encoder



Input Embedding

❖ Input Embedding: Embedding Layer

QUERY – KEY – VALUE ?

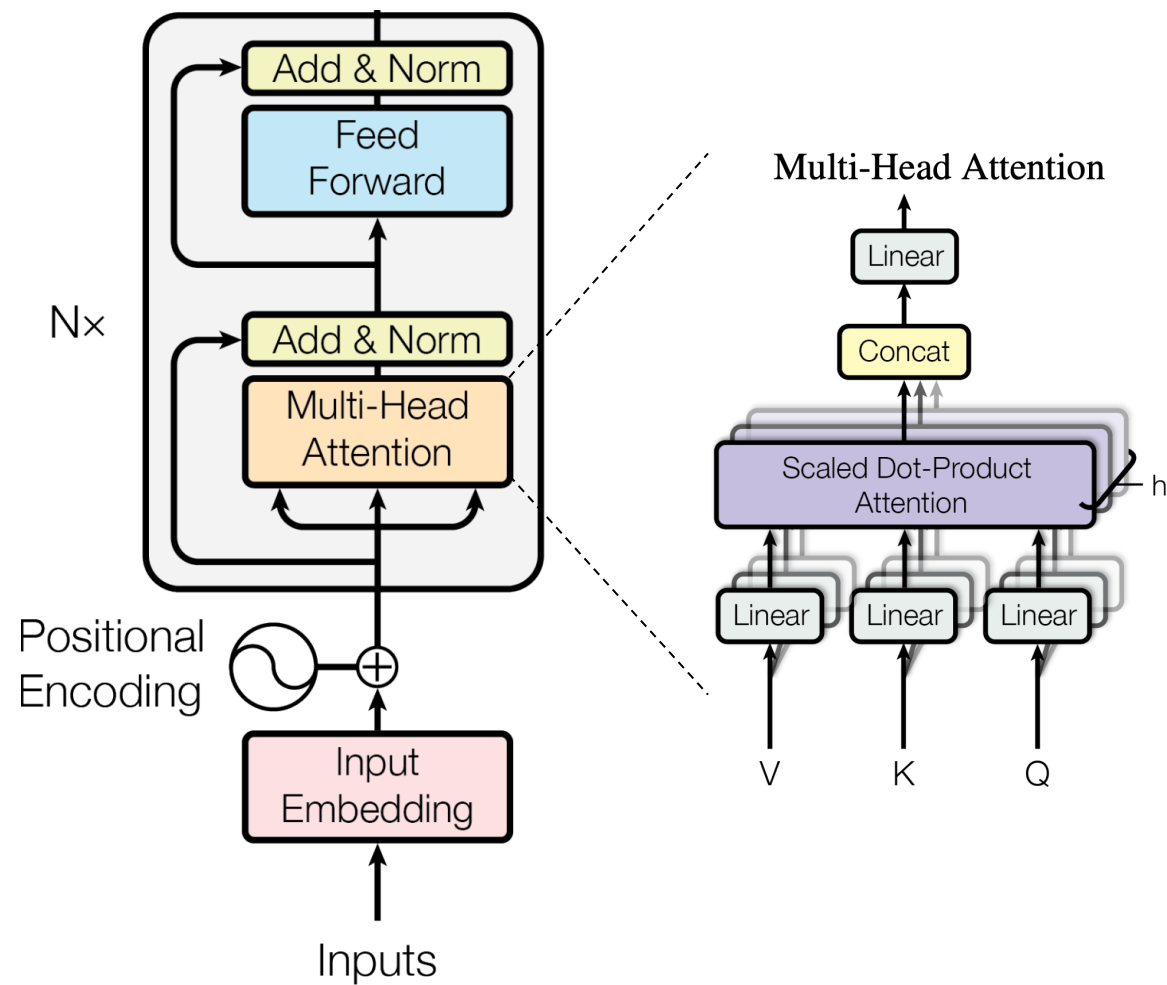
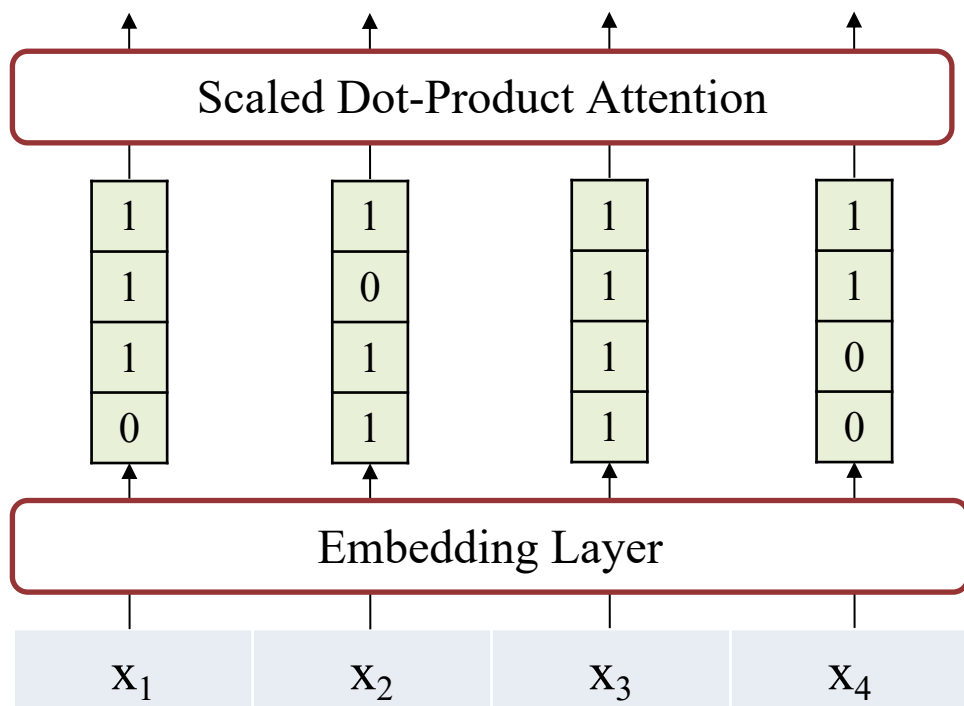


2 – Transformer-Encoder

!

Self-Attention




QUERY = KEY = VALUE = EMBEDDED



2 – Transformer-Encoder




Self-Attention



2 – Transformer-Encoder



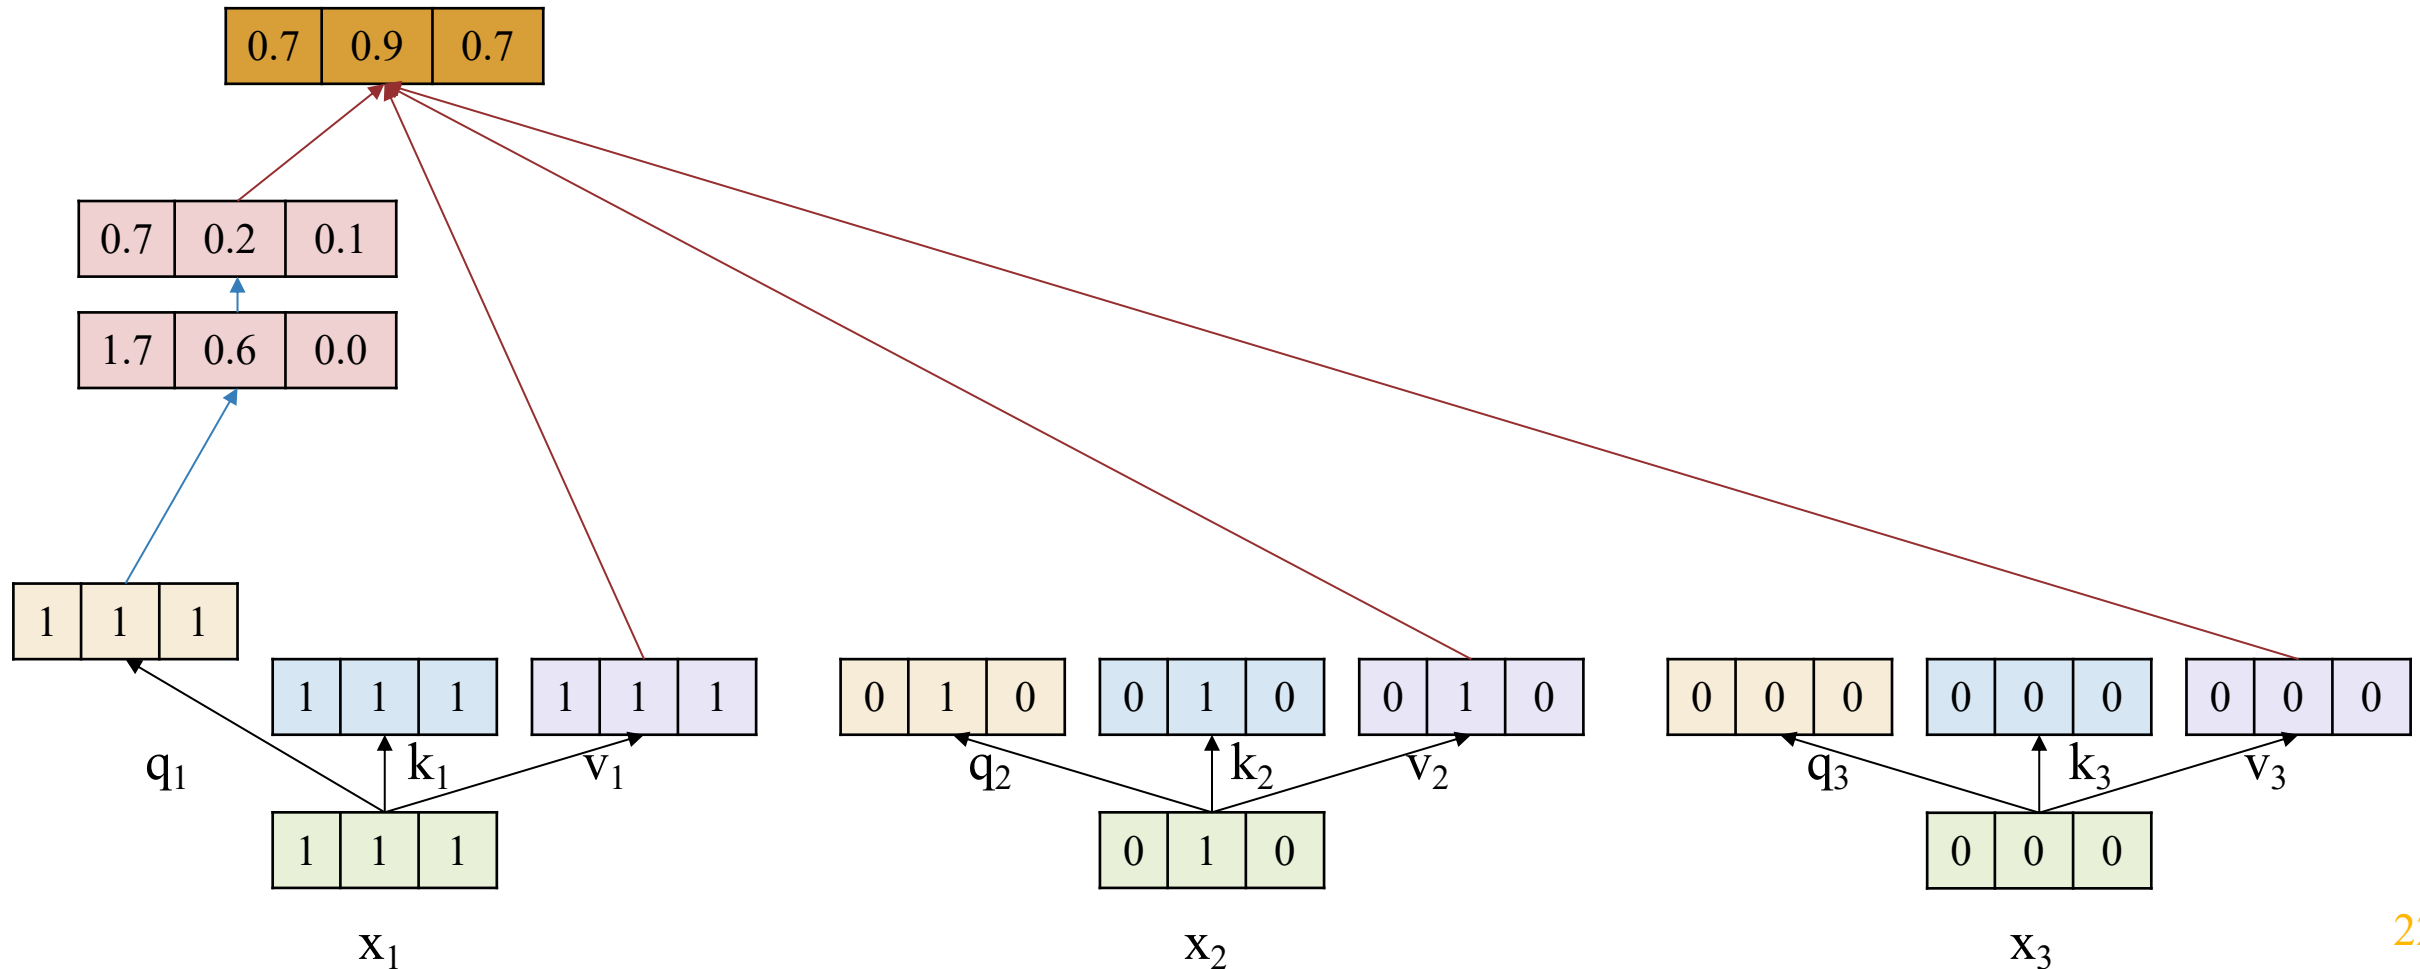
Self-Attention



2 – Transformer-Encoder



Self-Attention

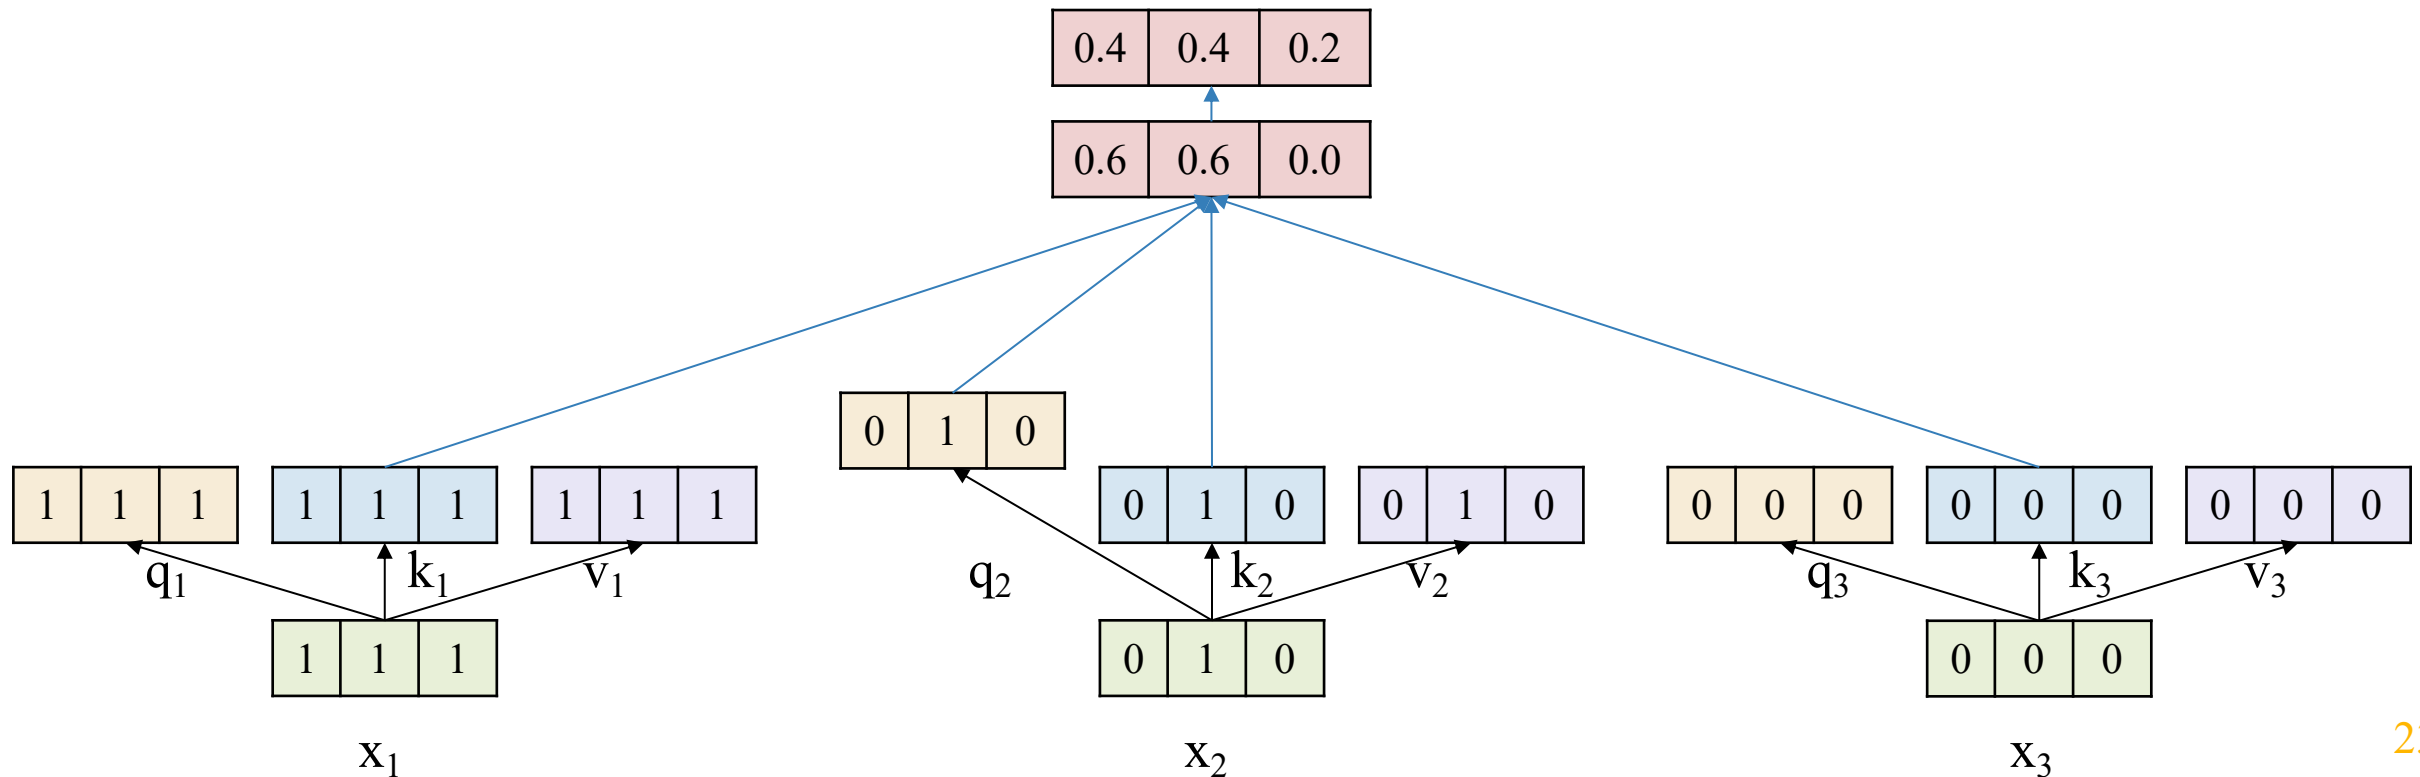


2 – Transformer-Encoder

!

Self-Attention

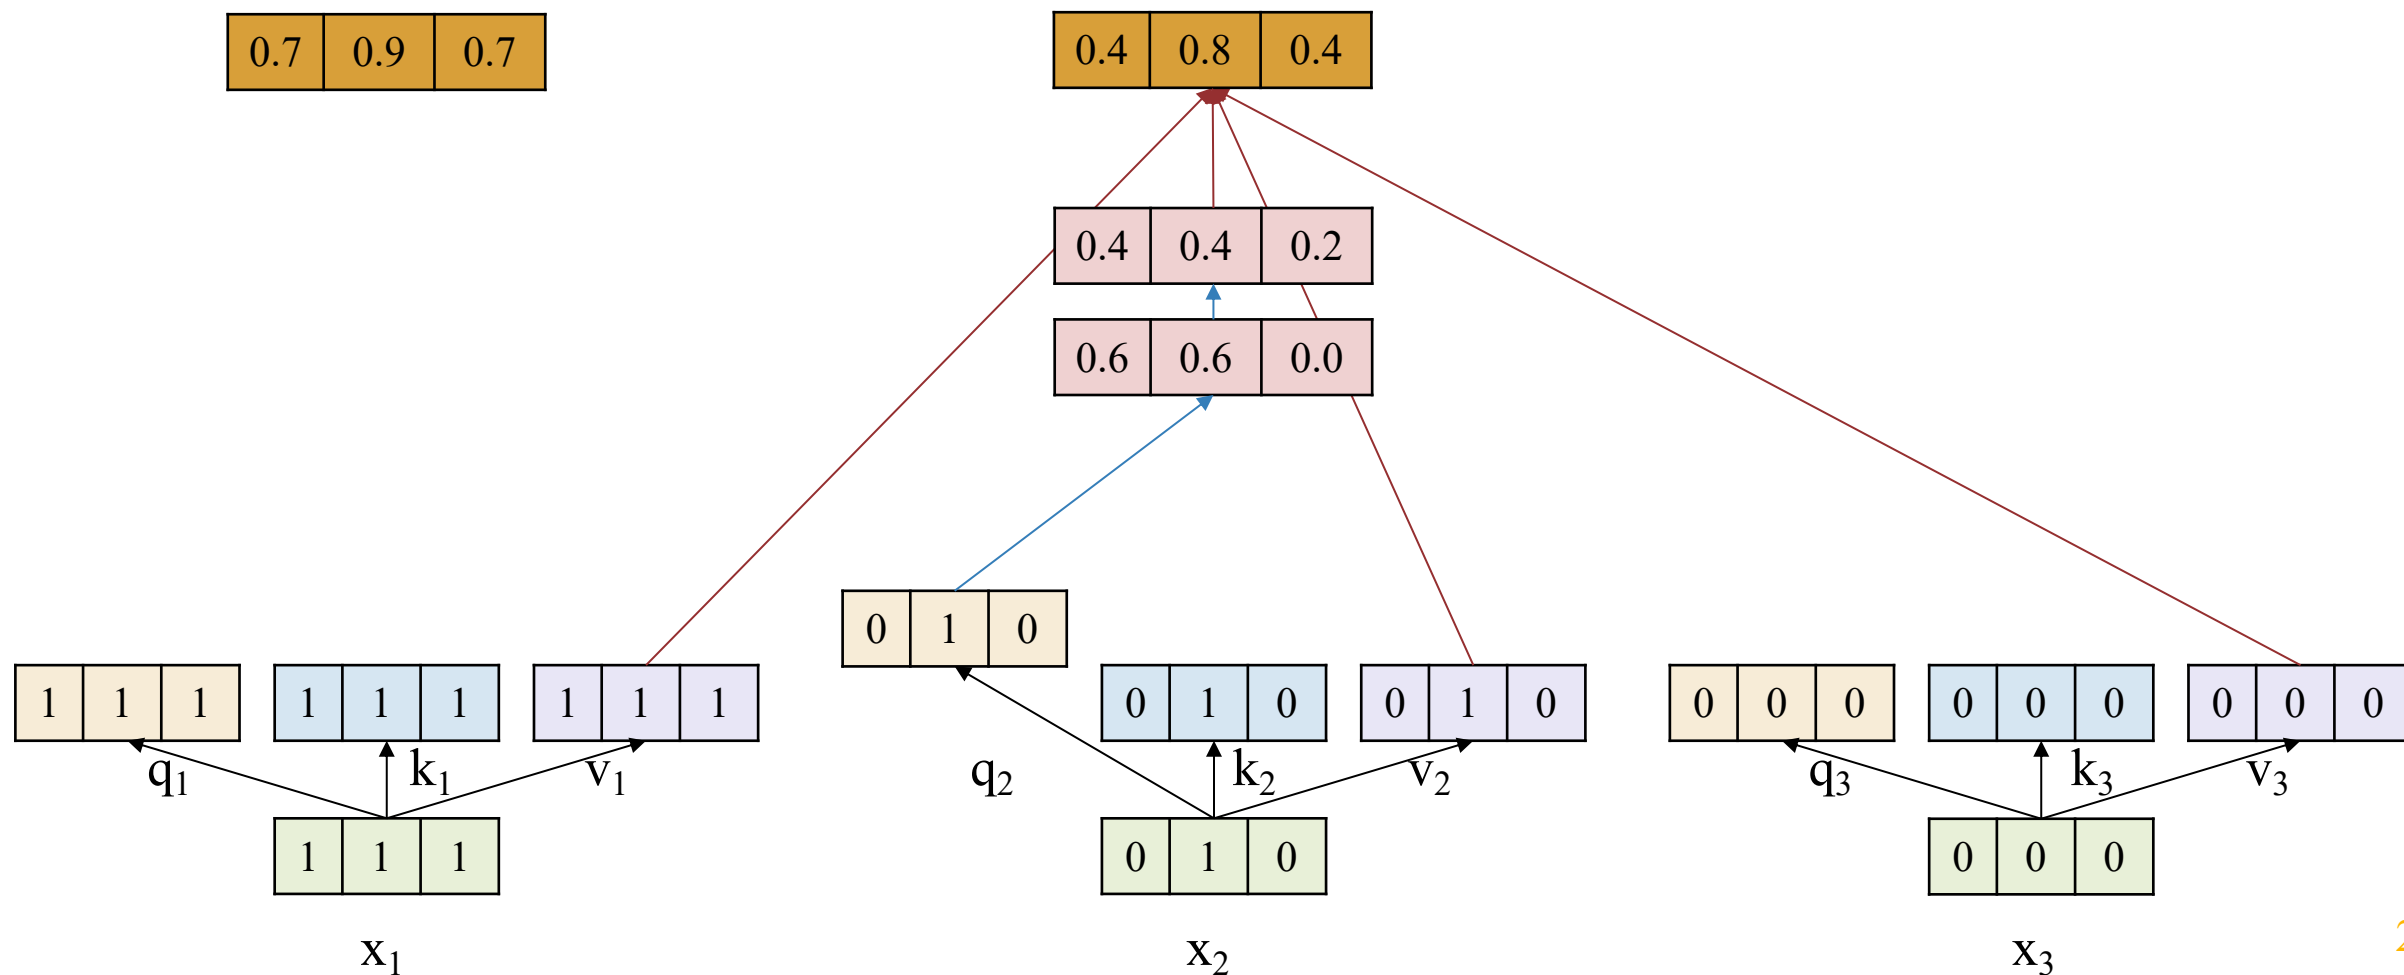
| | | |
|-----|-----|-----|
| 0.7 | 0.9 | 0.7 |
|-----|-----|-----|



2 – Transformer-Encoder



Self-Attention




2 – Transformer-Encoder



Self-Attention

| | | |
|-----|-----|-----|
| 0.7 | 0.9 | 0.7 |
|-----|-----|-----|

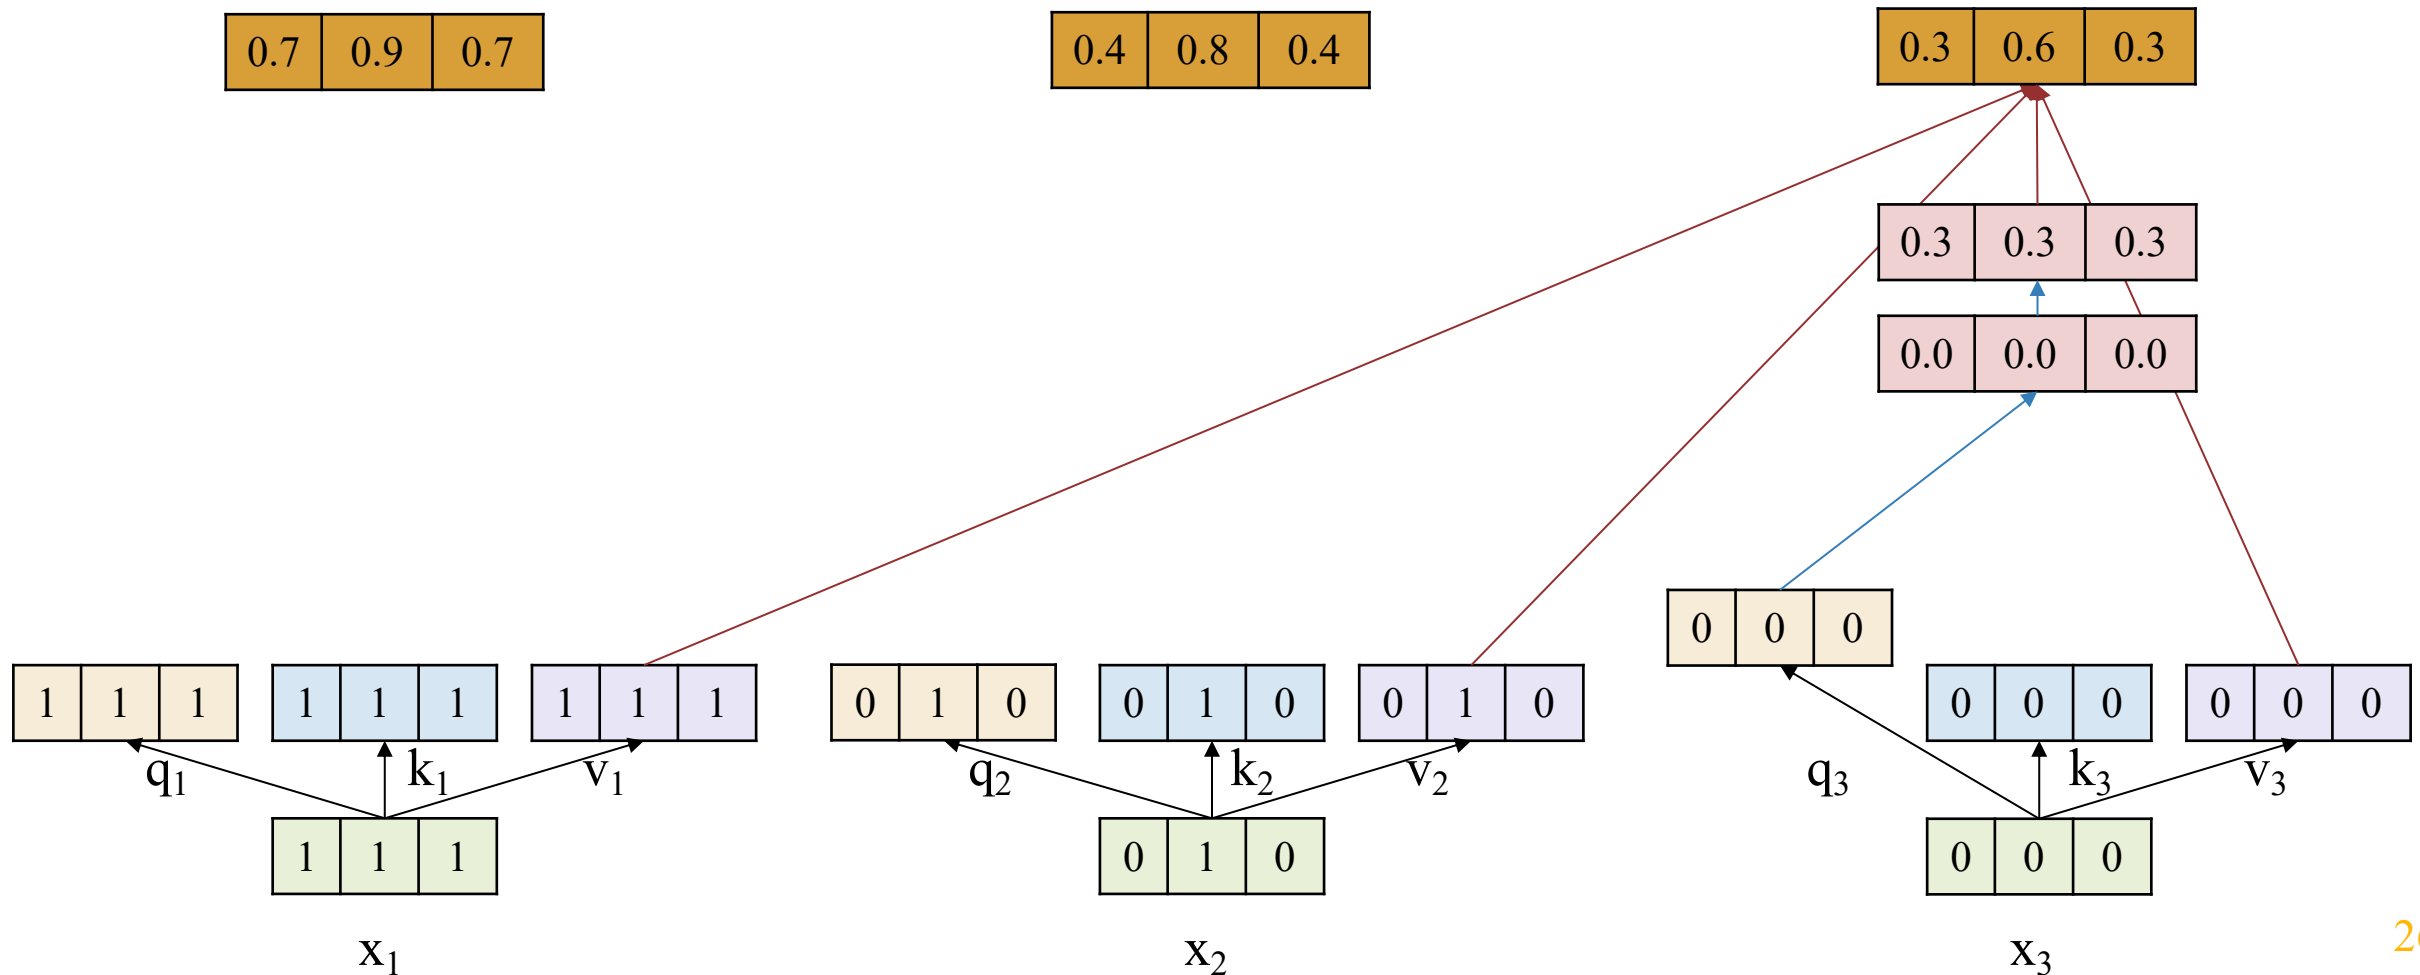
| | | |
|-----|-----|-----|
| 0.4 | 0.8 | 0.4 |
|-----|-----|-----|



2 – Transformer-Encoder

!

Self-Attention



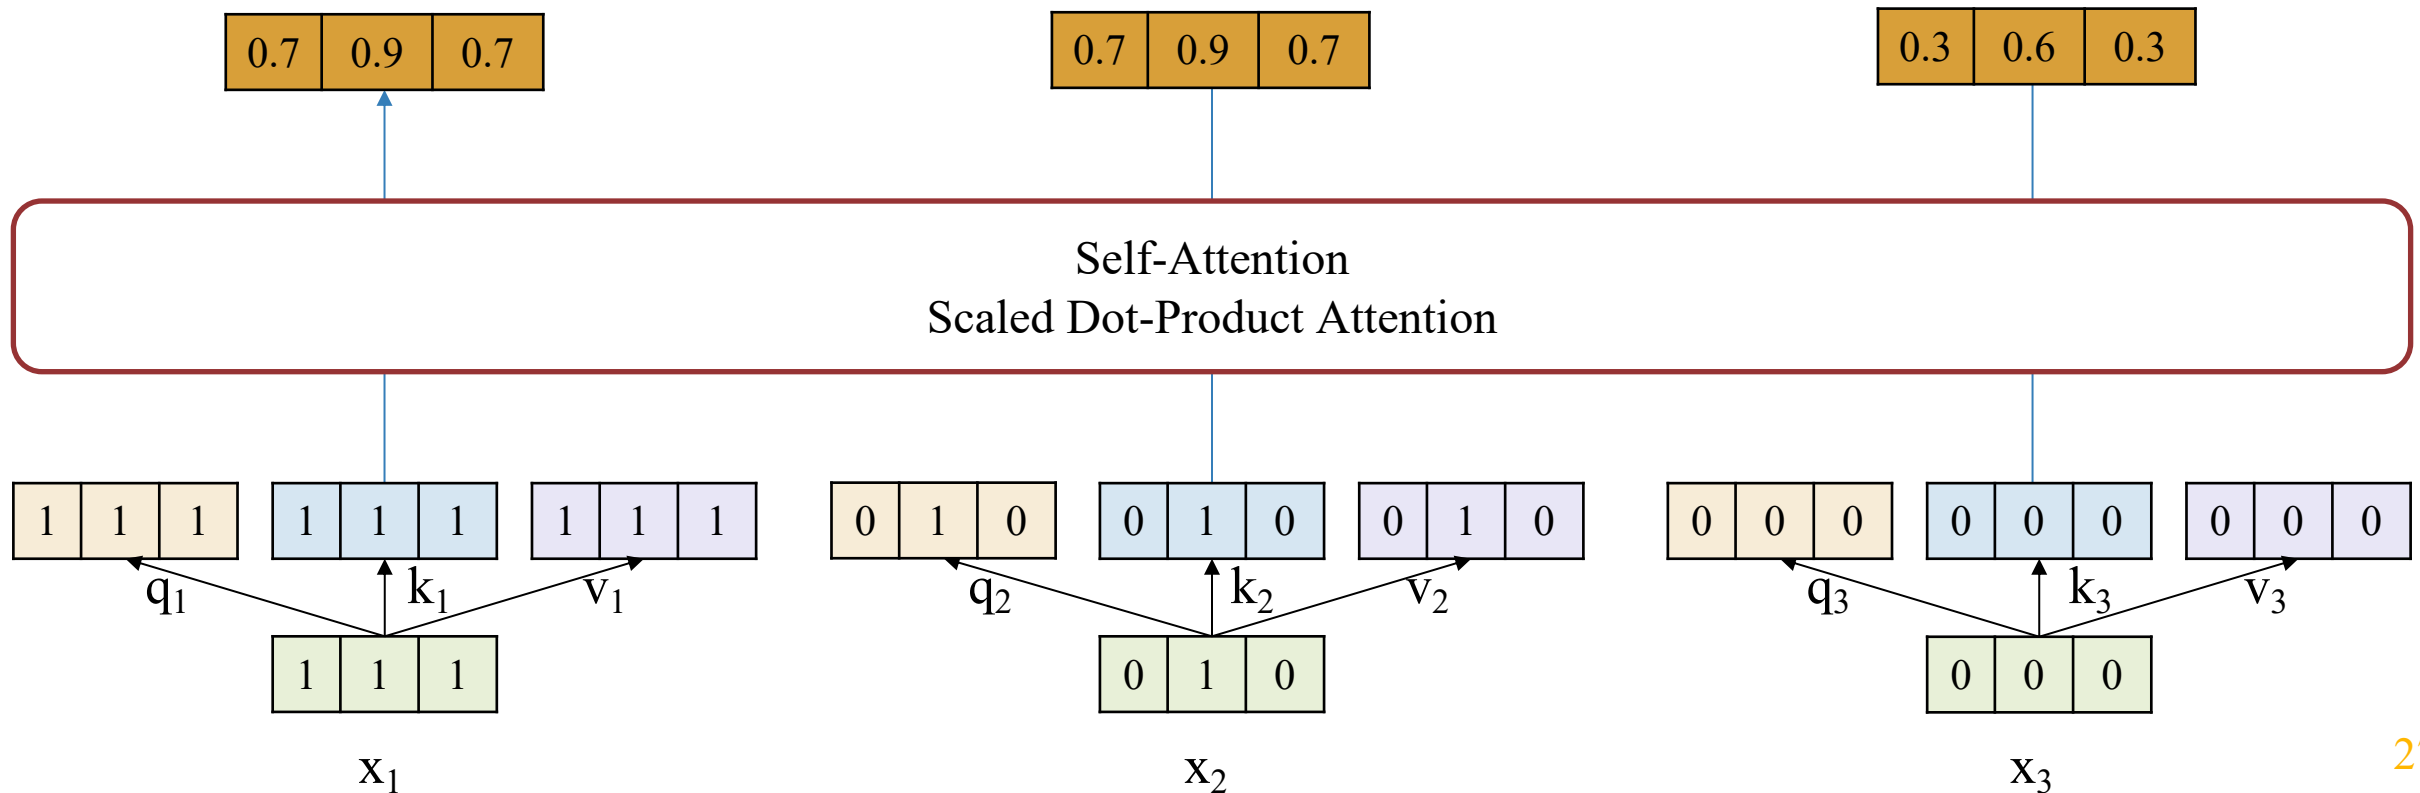
2 – Transformer-Encoder



Self-Attention

- ❖ To learn the relationship between word in the sentence

Ignore the order of words in the sentence ?

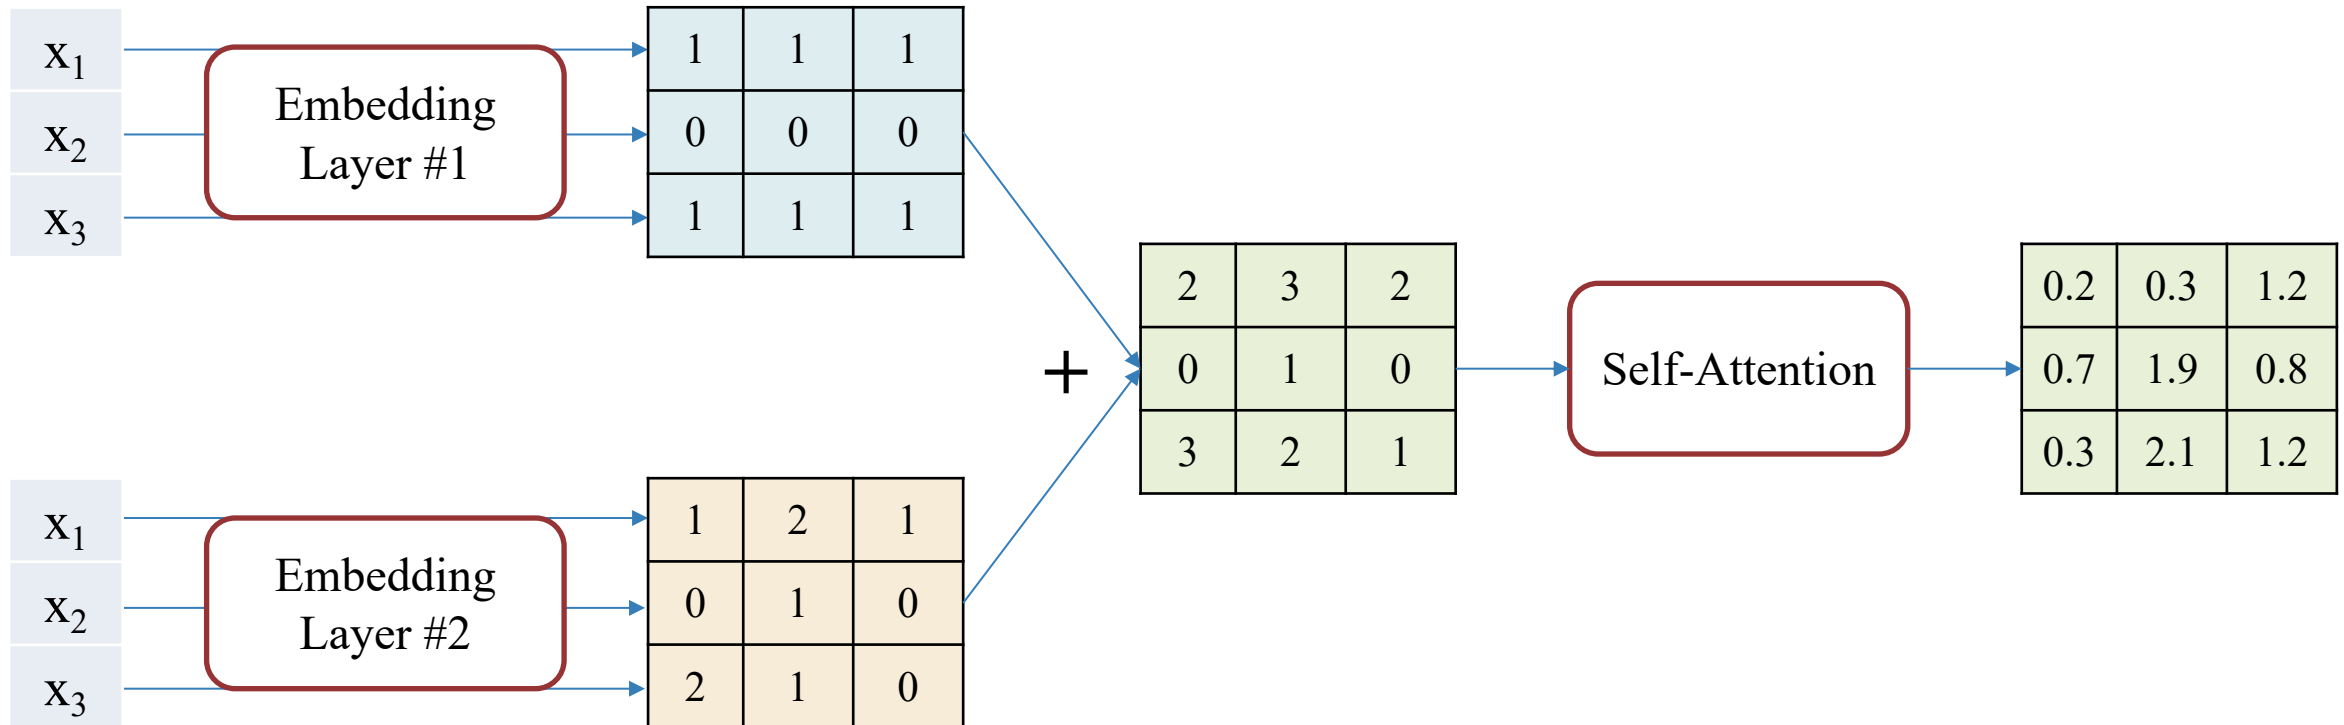


2 – Transformer-Encoder



Positional Encoding

- ❖ The position of a token in a sentence as unique representation – each position is mapped to a vector
- ❖ Methods: Sinusoid; Learned positional embedding (as learned input embedding)



2 – Transformer-Encoder



Positional Encoding – Demo

```
class TokenAndPositionEmbedding(nn.Module):
    def __init__(self, vocab_size, embed_dim, max_length, device='cpu'):
        super().__init__()
        self.device = device
        self.word_emb = nn.Embedding(
            num_embeddings=vocab_size,
            embedding_dim=embed_dim
        )
        self.pos_emb = nn.Embedding(
            num_embeddings=max_length,
            embedding_dim=embed_dim
        )

    def forward(self, x):
        N, seq_len = x.size()
        positions = torch.arange(0, seq_len).expand(N, seq_len).to(self.device)
        output1 = self.word_emb(x)
        output2 = self.pos_emb(positions)
        output = output1 + output2
        return output
```

```
vocab_size = 10000
embed_dim = 200
max_length = 50
embedding = TokenAndPositionEmbedding(
    vocab_size,
    embed_dim,
    max_length
)
```

```
batch_size = 32
```

```
input = torch.randint(
    high=2,
    size=(batch_size, max_length),
    dtype=torch.int64
)
```

```
embedded = embedding(input)
```

```
embedded.shape
```

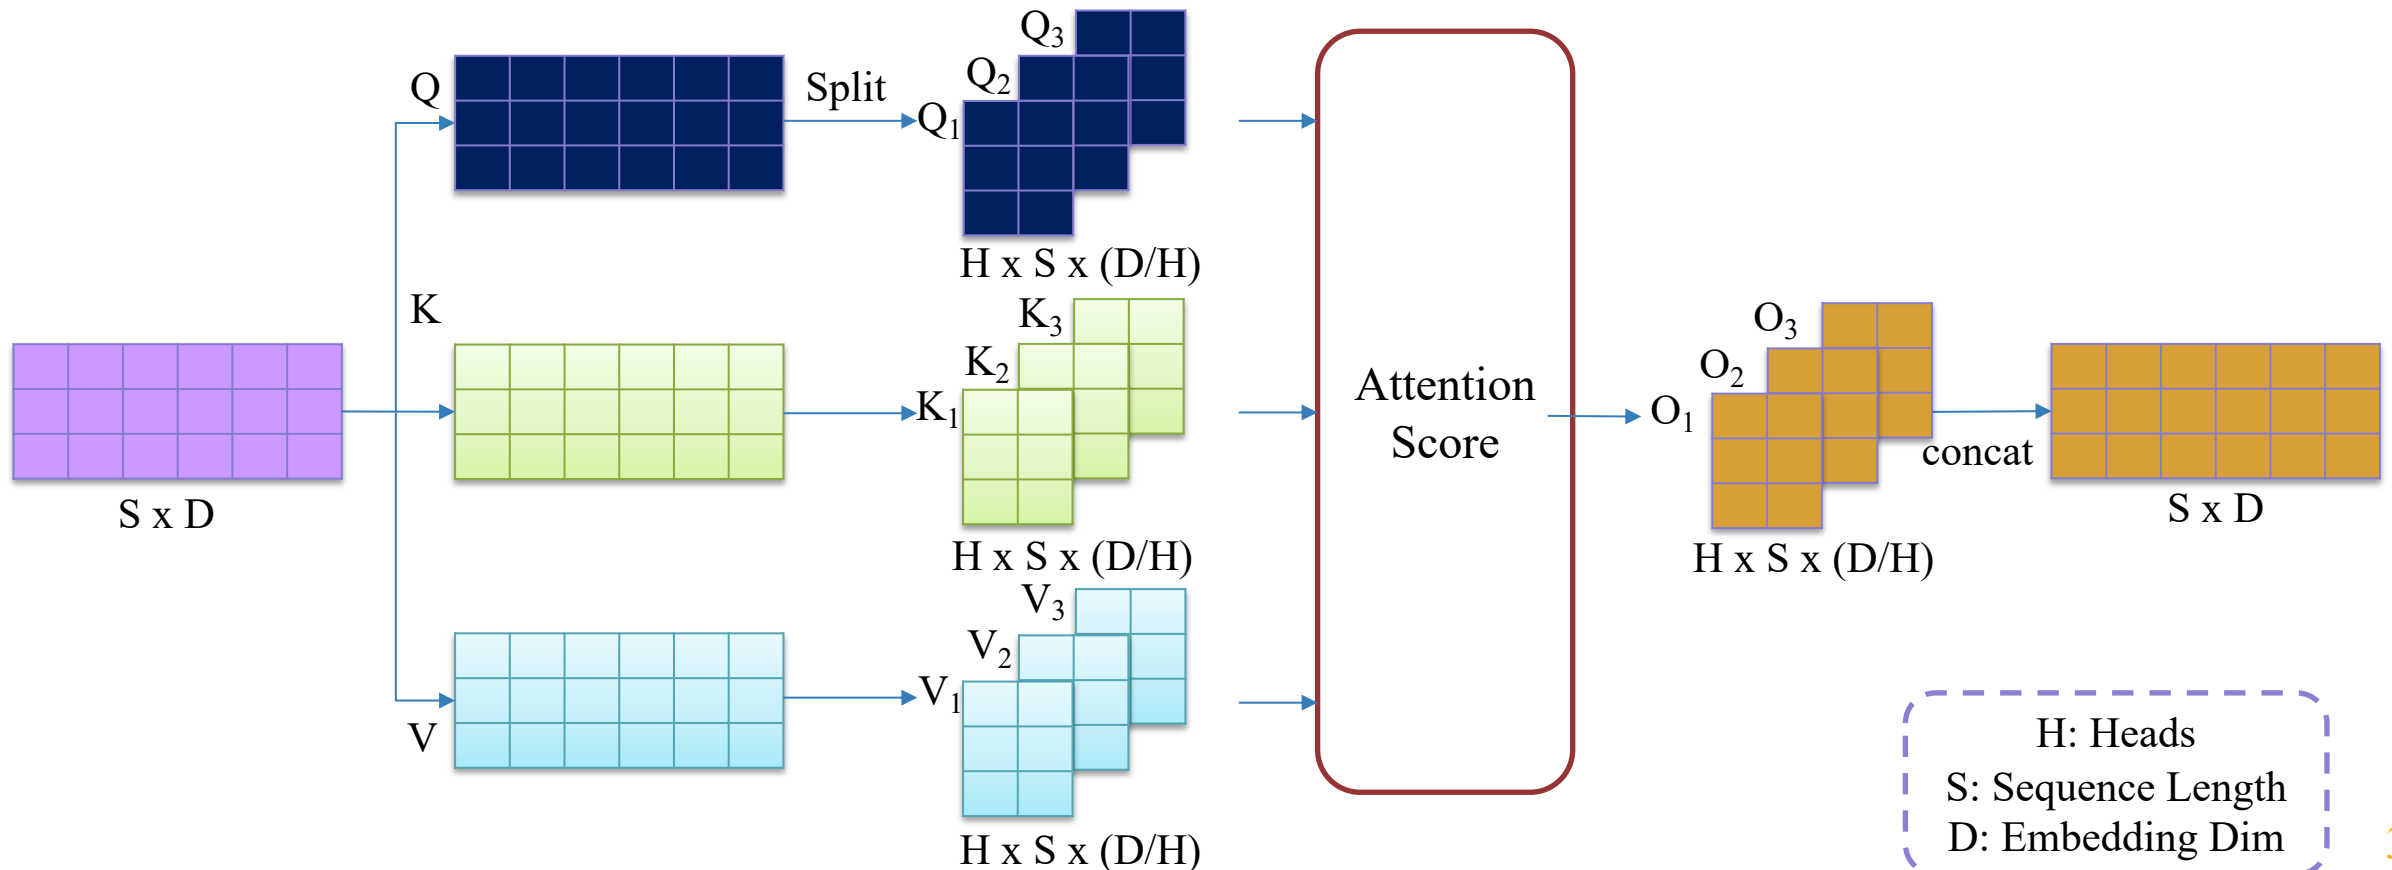
```
torch.Size([32, 50, 200])
```

2 – Transformer-Encoder



Multi-Head Attention

- ❖ Split into the multiple attention heads (process independently) \Rightarrow self-attention \Rightarrow concat



2 – Transformer-Encoder



Multi-Head Attention – Demo

```
batch_size = 1
seq_len = 50
embedding_dim = 200

input = torch.randint(
    high=2,
    size=(batch_size, seq_len, embedding_dim),
    dtype=torch.float32
)
input

tensor([[[[0., 1., 1., ..., 0., 1., 1.],
          [0., 1., 0., ..., 0., 0., 0.],
          [1., 0., 1., ..., 1., 1., 1.],
          ...,
          [0., 0., 0., ..., 1., 0., 0.],
          [1., 0., 1., ..., 0., 1., 1.],
          [0., 1., 1., ..., 1., 1., 1.]]]]])
```

```
embedding_dim = 200
num_heads = 5

att_layer = nn.MultiheadAttention(
    embed_dim=embedding_dim,
    num_heads=num_heads,
    batch_first=True
)
```

```
attn_output, attn_output_weights = att_layer(
    query=input,
    key=input,
    value=input
)
```

```
attn_output.shape
```

```
torch.Size([1, 50, 200])
```

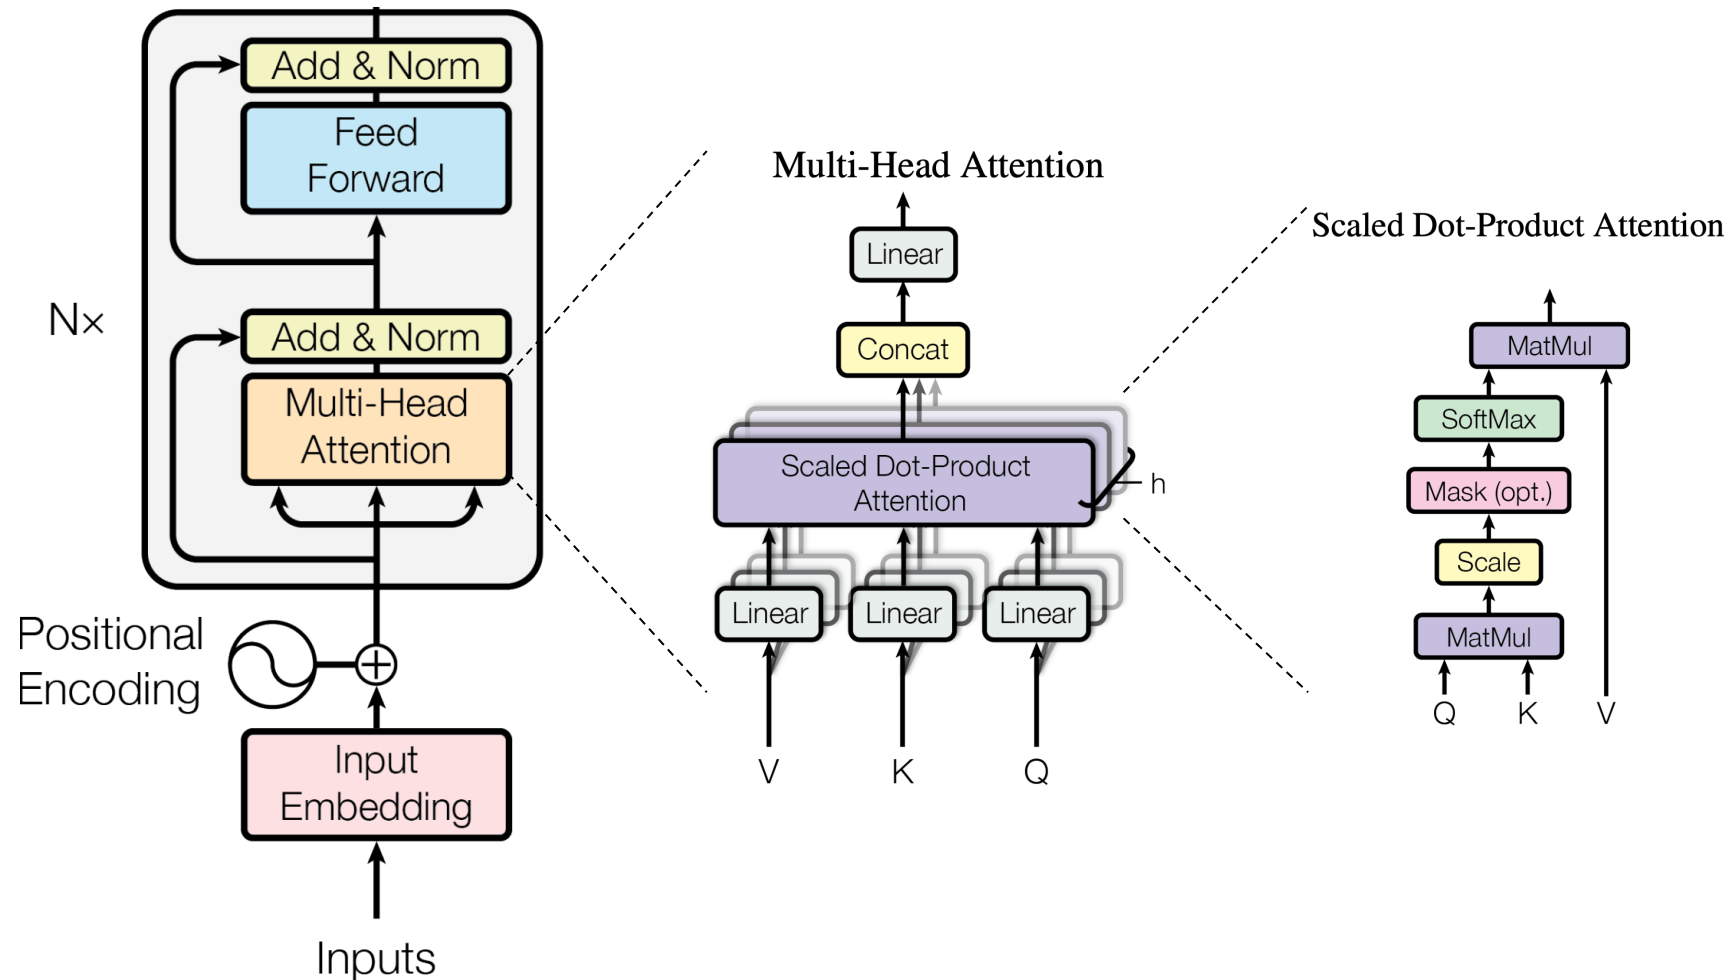
```
attn_output_weights.shape
```

```
torch.Size([1, 50, 50])
```

2 – Transformer-Encoder



Transformer-Encoder



2 – Transformer-Encoder

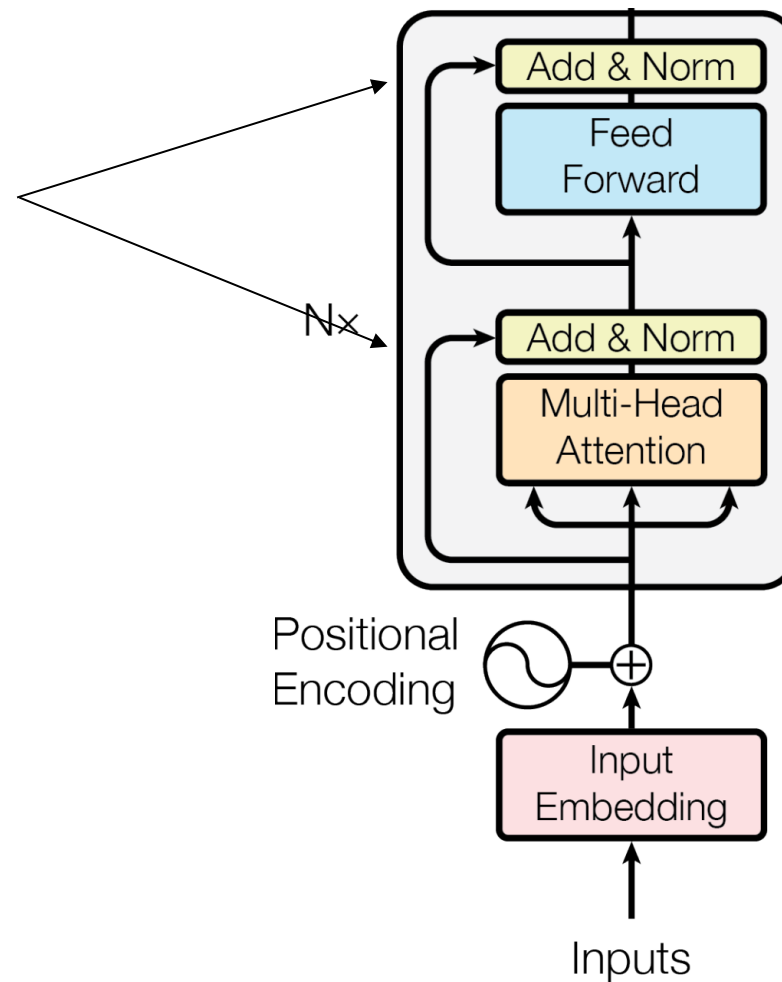


Layer Normalization

$$\mu_i = \frac{1}{m} \sum_{j=1}^m x_{ij}$$

$$\sigma_i^2 = \frac{1}{m} \sum_{j=1}^m (x_{ij} - \mu_i)^2$$

$$\hat{x}_{ij} = \frac{(x_{ij} - \mu_i)}{\sqrt{\sigma_i^2 + \epsilon}}$$

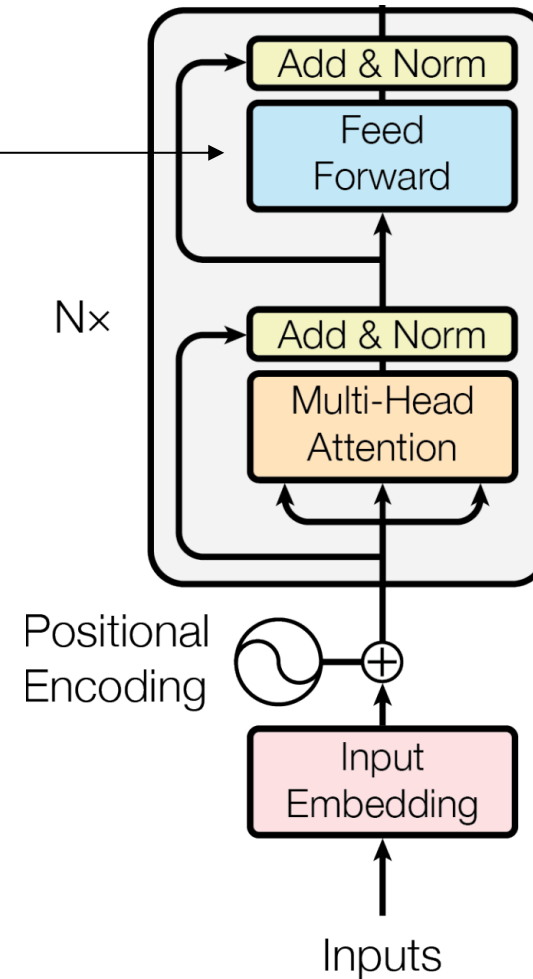



2 – Transformer-Encoder



Feed Forward

❖ 2 FC Layer



2 – Transformer-Encoder



Transformer-Encoder – Demo

```
class TransformerEncoder(nn.Module):
    def __init__(self, embed_dim, num_heads, ff_dim, dropout=0.1):
        super().__init__()
        self.attn = nn.MultiheadAttention(
            embed_dim=embed_dim,
            num_heads=num_heads,
            batch_first=True
        )
        self.ffn = nn.Sequential(
            nn.Linear(in_features=embed_dim, out_features=ff_dim, bias=True),
            nn.ReLU(),
            nn.Linear(in_features=ff_dim, out_features=embed_dim, bias=True)
        )
        self.layernorm_1 = nn.LayerNorm(normalized_shape=embed_dim, eps=1e-6)
        self.layernorm_2 = nn.LayerNorm(normalized_shape=embed_dim, eps=1e-6)
        self.dropout_1 = nn.Dropout(p=dropout)
        self.dropout_2 = nn.Dropout(p=dropout)

    def forward(self, query, key, value):
        attn_output, _ = self.attn(query, key, value)
        attn_output = self.dropout_1(attn_output)
        out_1 = self.layernorm_1(query + attn_output)
        ffn_output = self.ffn(out_1)
        ffn_output = self.dropout_2(ffn_output)
        out_2 = self.layernorm_2(out_1 + ffn_output)
        return out_2
```

```
encoder_layer = TransformerEncoder(
    embed_dim=200,
    num_heads=5,
    ff_dim=1024
)
```

```
embedded.shape
```

```
torch.Size([32, 50, 200])
```

```
encoded = encoder_layer(embedded, embedded, embedded)
```

```
encoded.shape
```

```
torch.Size([32, 50, 200])
```

3 – Text Classification



NTC-SCV Dataset

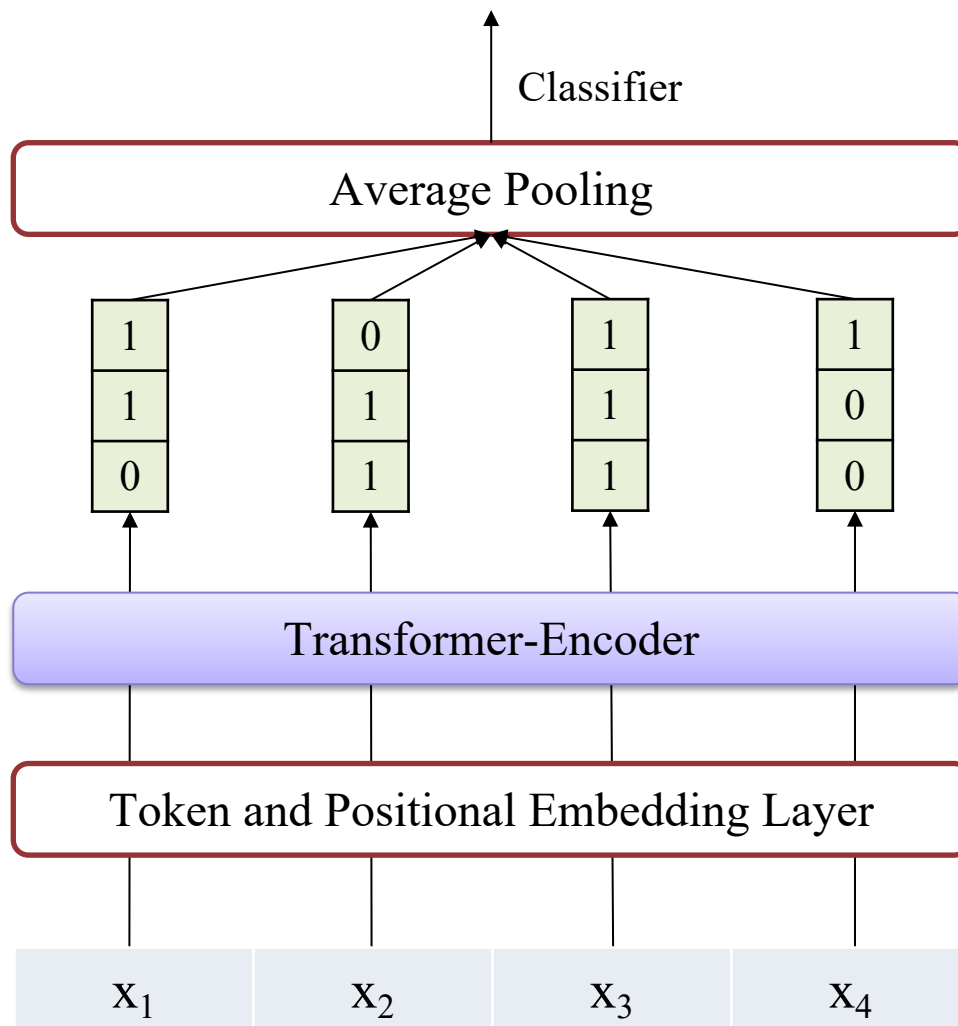
❖ Sentiment Analysis

| Positive Example | Negative Example |
|--|--|
| Mình được 1 cô bạn giới_thiệu đến đây , tìm địa_chỉ khá dễ . Menu nước uống chất khỏi nói . Mình muốn cũng đc 8 loại nước ở đây , món nào cũng ngon và bổ_dưỡng cả . | Quán chế_biến đồ_ăn lâu , Cá_Sapa nướng ướp rất dở , sò Lông ko tươi , nước_chấm ko ngon\nTôm_lại sẽ ko bao_giờ ghé nữa , ăn_dở mà uống tiền |
| Mỗi lần thèm trà sữa là làm 1 ly . Quán dễ kiếm , không_gian lại rộng_rãi . Nhân_viên thì dễ_thương gần_gũi . Nói_chung thèm trà sữa là mình ghé Quán ở đây vì gần nhà . | Quán này thấy khá nhiều người bảo mình nên mình đã đi ăn thử , nhưng thực_sự ăn xong thấy không được như mong_đợi lắm . |

3 – Text Classification



Modeling



3 – Text Classification



Modeling – Demo

```
class TransformerEncoder(nn.Module):
    def __init__(self, embed_dim, num_heads, ff_dim, dropout=0.1):
        super().__init__()
        self.attn = nn.MultiheadAttention(
            embed_dim=embed_dim,
            num_heads=num_heads,
            batch_first=True
        )
        self.ffn = nn.Sequential(
            nn.Linear(in_features=embed_dim, out_features=ff_dim, bias=True),
            nn.ReLU(),
            nn.Linear(in_features=ff_dim, out_features=embed_dim, bias=True)
        )
        self.layernorm_1 = nn.LayerNorm(normalized_shape=embed_dim, eps=1e-6)
        self.layernorm_2 = nn.LayerNorm(normalized_shape=embed_dim, eps=1e-6)
        self.dropout_1 = nn.Dropout(p=dropout)
        self.dropout_2 = nn.Dropout(p=dropout)



    def forward(self, query, key, value):
        attn_output, _ = self.attn(query, key, value)
        attn_output = self.dropout_1(attn_output)
        out_1 = self.layernorm_1(query + attn_output)
        ffn_output = self.ffn(out_1)
        ffn_output = self.dropout_2(ffn_output)
        out_2 = self.layernorm_2(out_1 + ffn_output)
        return out_2
```

3 – Text Classification



Training

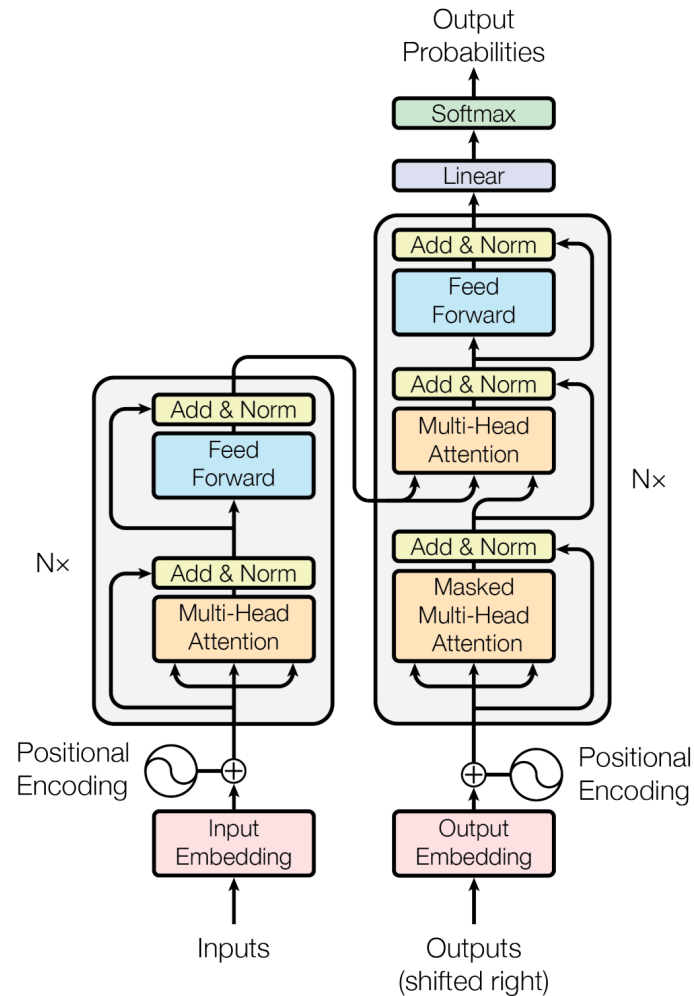
❖ Testing: 83.66%



4 – Vision Transformer



ViT



Transformers are so successful in NLP,
Can we use them for images?


4 – Vision Transformer



From text to image

“I Love AI VN”

Tokenization



4 – Vision Transformer



From text to image

Can we tokenize an image?





4 – Vision Transformer



From text to image

Can we tokenize an image?






4 – Vision Transformer



From text to image

Can we tokenize an image?

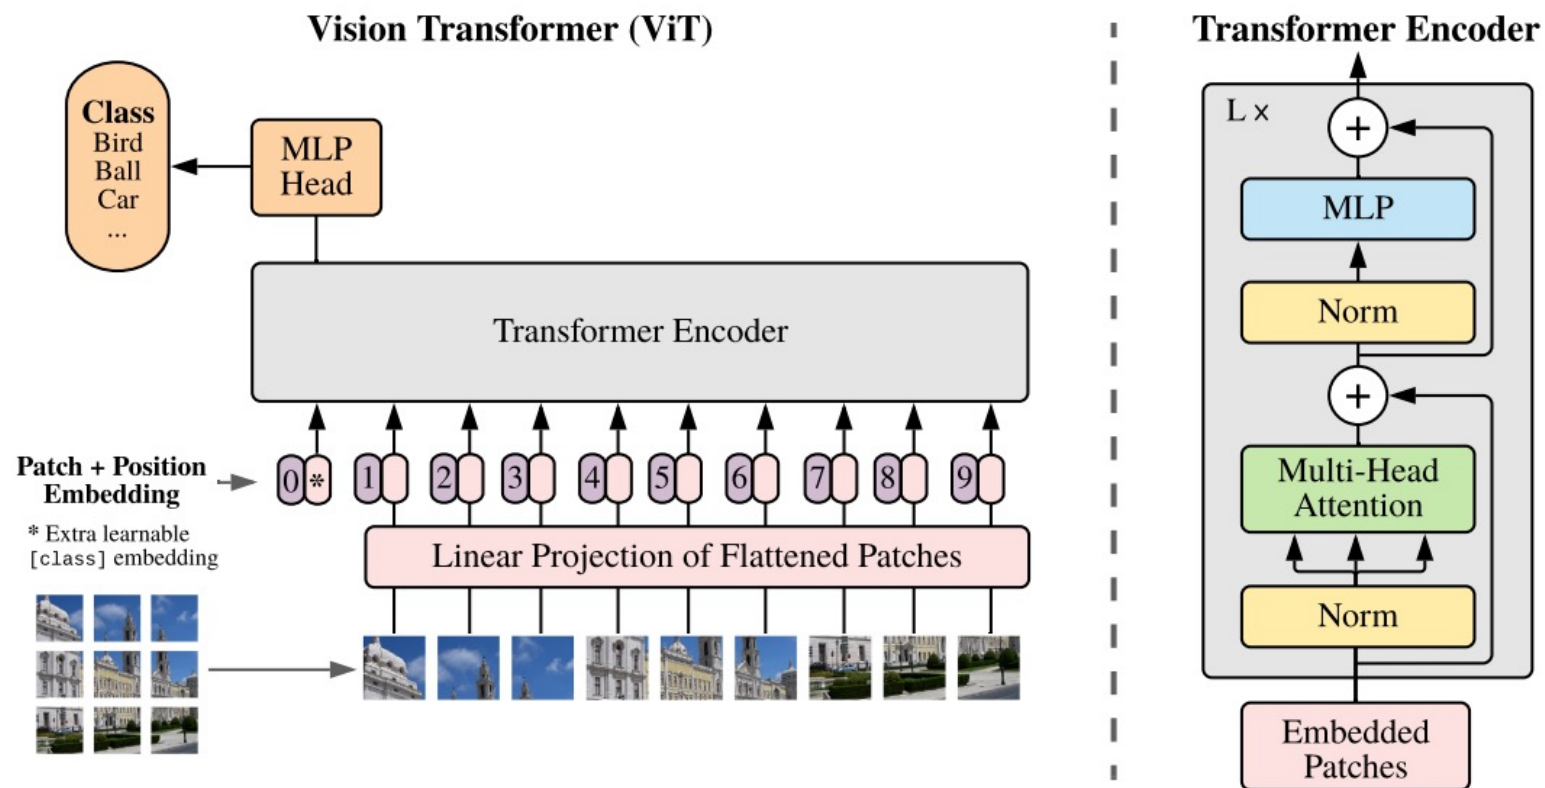


Flattening

4 – Vision Transformer

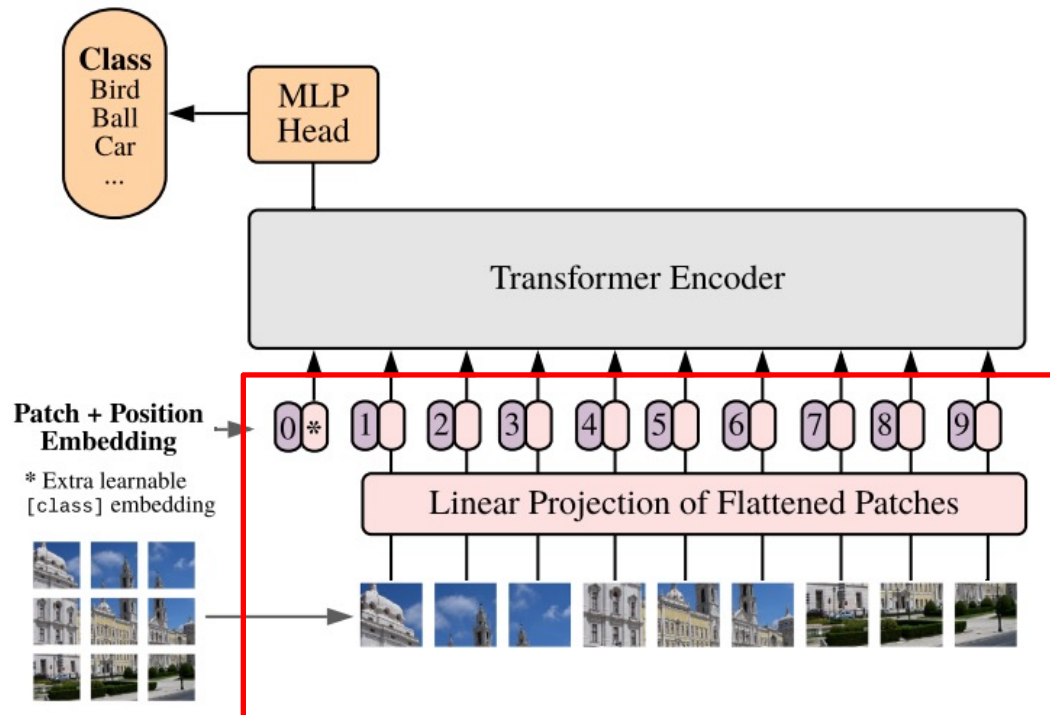


ViT Architecture

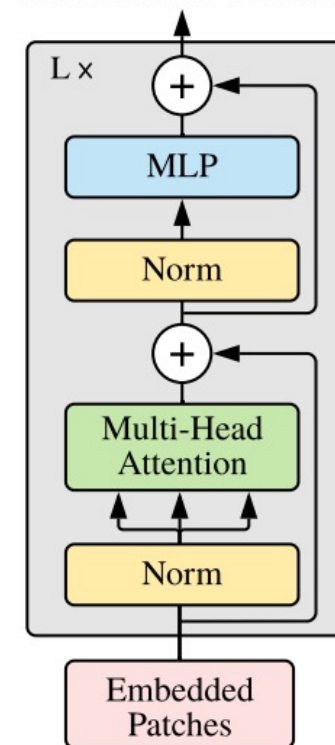


!

Vision Transformer (ViT)



Transformer Encoder



4 – Vision Transformer




Patch embedding

```
1 class PatchEmbedding(nn.Module):
2     def __init__(self, embed_dim=512, patch_size=16, image_size=224):
3         super().__init__()
4         self.conv1 = nn.Conv2d(in_channels=3, out_channels=embed_dim, kernel_size=patch_size, stride=patch_size, bias=False)
5
6     def forward(self, x):
7         x = self.conv1(x) # shape = [* , width, grid, grid]
8         x = x.reshape(x.shape[0], x.shape[1], -1) # shape = [* , width, grid ** 2]
9         x = x.permute(0, 2, 1) # shape = [* , grid ** 2, width]
10        return x

1 patch_embedding = PatchEmbedding()
2 x = torch.randn(1, 3, 224, 224)
3
4 out = patch_embedding(x)
5 print(out.shape)

torch.Size([1, 196, 512])
```



4 – Vision Transformer



Patch embedding

| | | | | | | | |
|---|---|---|---|---|---|---|---|
| 1 | 2 | 4 | 2 | 2 | 3 | 3 | 2 |
| 1 | 0 | 2 | 1 | 2 | 1 | 1 | 1 |
| 2 | 2 | 3 | 4 | 3 | 4 | 1 | 3 |
| 2 | 1 | 3 | 0 | 0 | 2 | 3 | 0 |
| 3 | 3 | 4 | 0 | 2 | 0 | 2 | 2 |
| 1 | 4 | 4 | 3 | 4 | 0 | 4 | 0 |
| 1 | 2 | 0 | 0 | 0 | 3 | 2 | 3 |
| 4 | 1 | 4 | 1 | 0 | 0 | 0 | 0 |

Patch size: 4

| | | | |
|---|---|---|---|
| 1 | 0 | 0 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 0 | 0 | 1 |
| 1 | 1 | 0 | 1 |

| | | | |
|---|---|---|---|
| 1 | 1 | 0 | 0 |
| 0 | 0 | 1 | 0 |
| 0 | 1 | 0 | 0 |
| 1 | 1 | 0 | 0 |



| | |
|----|----|
| 14 | 10 |
|----|----|

| | |
|----|----|
| 14 | 12 |
|----|----|


| | |
|----|----|
| 18 | 17 |
|----|----|

| | |
|----|---|
| 11 | 9 |
|----|---|


4 – Vision Transformer

!

Positional embedding



Embed + add



4 – Vision Transformer




Positional embedding

```
1 class PatchPositionEmbedding(nn.Module):
2     def __init__(self, embed_dim=512, patch_size=16, image_size=224):
3         super().__init__()
4         self.conv1 = nn.Conv2d(in_channels=3, out_channels=embed_dim, kernel_size=patch_size, stride=patch_size, bias=False)
5
6         scale = embed_dim ** -0.5
7         self.positional_embedding = nn.Parameter(scale * torch.randn((image_size // patch_size) ** 2, embed_dim))
8
9     def forward(self, x):
10         x = self.conv1(x) # shape = [* , width, grid, grid]
11         x = x.reshape(x.shape[0], x.shape[1], -1) # shape = [* , width, grid ** 2]
12         x = x.permute(0, 2, 1) # shape = [* , grid ** 2, width]
13
14         x = x + self.positional_embedding.to(x.dtype)
15         return x
```

```
[36] 1 patchpos_embedding = PatchPositionEmbedding()
      2 x = torch.randn(1, 3, 224, 224)
      3
      4 out = patchpos_embedding(x)
      5 print(out.shape)
```

```
torch.Size([1, 196, 512])
```



4 – Vision Transformer



Positional embedding

14 | 10

0 | 0

14 | 10

14 | 12

1 | 1

15 | 13

+

=

18 | 17

1 | 1

19 | 18

11 | 9

0 | 0

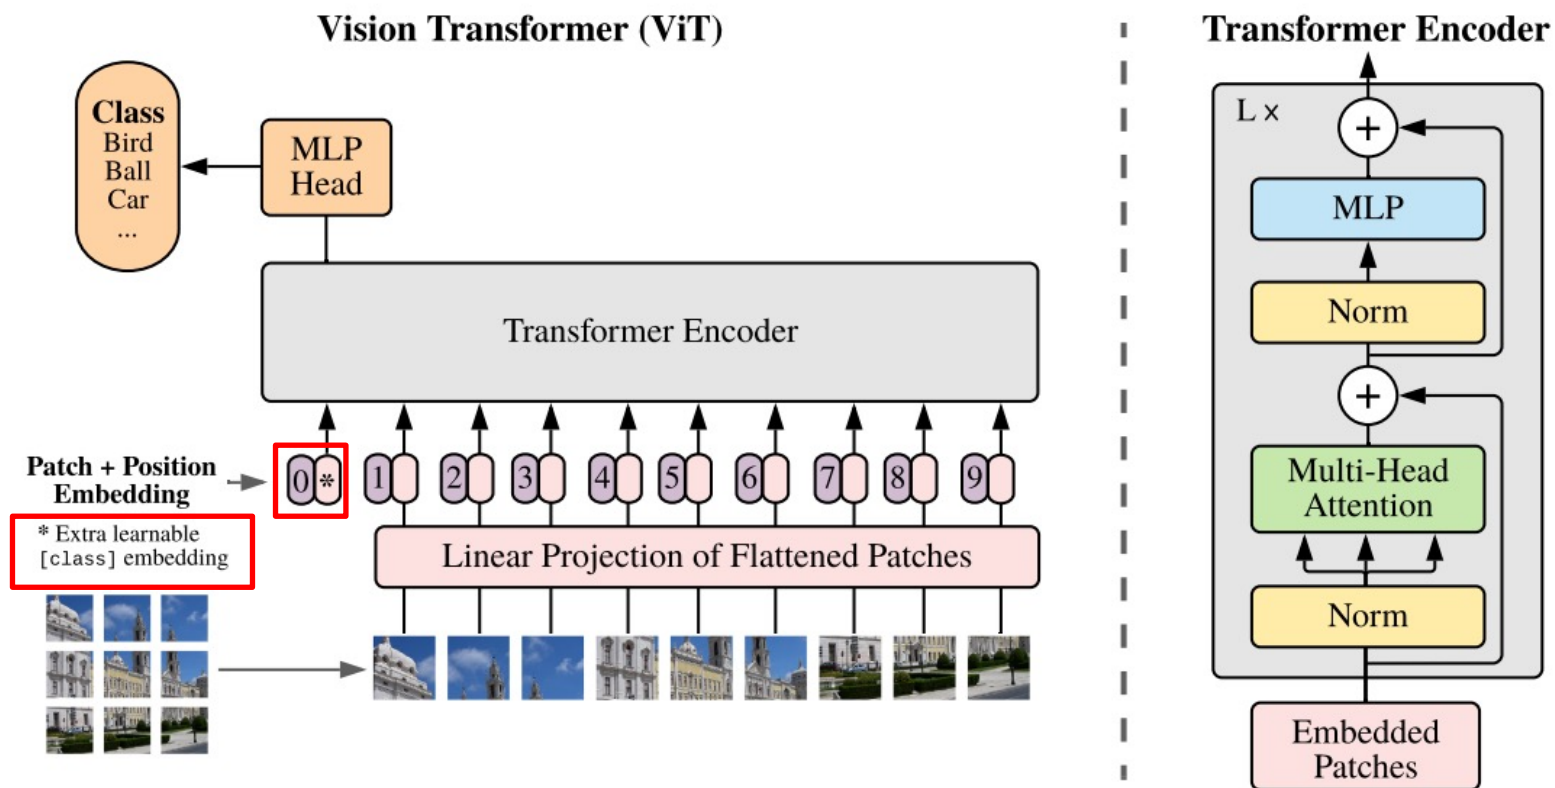
11 | 9

Positional
Embedding

4 – Vision Transformer



[CLS] Token



4 – Vision Transformer




[CLS] Token – Why?

Alternatives?

- Global Average Pooling
- Max Pooling
- ...


Have info of
other tokens



Attention n

...

Average Pooling




t1

t2

t3

Attention 1



CLS

t1

t2

t3

4 – Vision Transformer



[CLS] Token – Demo

```
1 class PatchPositionEmbedding(nn.Module):
2     def __init__(self, embed_dim=512, patch_size=16, image_size=224):
3         super().__init__()
4         self.conv1 = nn.Conv2d(in_channels=3, out_channels=embed_dim, kernel_size=patch_size, stride=patch_size, bias=False)
5
6         scale = embed_dim ** -0.5
7         self.class_embedding = nn.Parameter(scale * torch.randn(embed_dim))
8         self.positional_embedding = nn.Parameter(scale * torch.randn((image_size // patch_size) ** 2 + 1, embed_dim))
9
10    def forward(self, x):
11        x = self.conv1(x) # shape = [*, width, grid, grid]
12        x = x.reshape(x.shape[0], x.shape[1], -1) # shape = [*, width, grid ** 2]
13        x = x.permute(0, 2, 1) # shape = [*, grid ** 2, width]
14        # expanding the CLS embedding
15        cls_embs = self.class_embedding.to(x.dtype) + torch.zeros(x.shape[0], 1, x.shape[-1], dtype=x.dtype, device=x.device)
16        x = torch.cat([cls_embs, x], dim=1) # shape = [*, grid ** 2 + 1, width]
17
18        x = x + self.positional_embedding.to(x.dtype)
19        return x
```

4 – Vision Transformer



Modeling

Change Token
Embedding with
Patch Embedding



[CLS] token instead
of pooling (can still
use pooling)

```
1 class VisionTransformerCls(nn.Module):
2     def __init__(self,
3         image_size, embed_dim, num_heads, ff_dim,
4         dropout=0.1, device='cpu', num_classes = 10, patch_size=16
5     ):
6         super().init ()
7         self.embd_layer = PatchPositionEmbedding(
8             image_size=image_size, embed_dim=embed_dim, patch_size=patch_size
9         )
10        self.transformer_layer = TransformerEncoder(
11            embed_dim, num_heads, ff_dim, dropout
12        )
13        # self.pooling = nn.AvgPool1d(kernel_size=max_length)
14        self.fc1 = nn.Linear(in_features=embed_dim, out_features=20)
15        self.fc2 = nn.Linear(in_features=20, out_features=num_classes)
16        self.dropout = nn.Dropout(p=dropout)
17        self.relu = nn.ReLU()
18    def forward(self, x):
19        output = self.embd_layer(x)
20        output = self.transformer_layer(output, output, output)
21        output = output[:, 0, :]
22        output = self.dropout(output)
23        output = self.fc1(output)
24        output = self.dropout(output)
25        output = self.fc2(output)
26        return output
```

4 – Vision Transformer



Training





AI VIET NAM

@aivietnam.edu.vn

Thanks!

Any questions?