


Extra Class

UNet

Nguyen Quoc Thai



CONTENT

(1) – Image Segmentation


(2) – Transposed Convolution

(3) – UNet Model

1 – Image Segmentation



Image Classification



MODEL
(LeNet, ResNet,...)

Class: CAT

Pretrained Model

1 – Image Segmentation



Multiple Objects



DOG

CAT



MODEL
(LeNet, ResNet,...)

Class: ?

1 – Image Segmentation




Image Segmentation





DOG

CAT



MODEL



1 – Image Segmentation



Image Segmentation




DOG

CAT

MODEL


UNet



1 – Image Segmentation



Unet Model



2 – Transposed Convolution



Review: Convolution

0	3	1	1
3	1	2	0
3	4	2	3
3	0	0	2

Input: M x N

Padding: (P, Q)

0	0	0	0	0	0
0	0	3	1	1	0
0	3	1	2	0	0
0	3	4	2	3	0
0	3	0	0	2	0
0	0	0	0	0	0

Shape: (M+2P) x (N+2Q)

Stride: (S, T)

1	1	1
1	1	1
0	1	0

Kernel: K x O

*


1

Bias

=

7	8
15	13

$$\left\lfloor \frac{M + 2P - K}{S} + 1 \right\rfloor \times \left\lfloor \frac{N + 2Q - O}{T} + 1 \right\rfloor$$



2 – Transposed Convolution



Review: Pooling

❖ Max Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: 2
Stride: 2

3	3	3
4	4	4
4	4	4

Output: 3 x 3

❖ Average Pooling

3	2	1	0	0	3
0	3	3	1	1	0
3	1	4	1	1	0
2	4	1	1	0	4
1	0	3	0	3	0
3	4	4	3	3	4

Input: 6 x 6

Kernel Size: (3, 2)
Stride: 2

2.0	1.7	0.8
1.8	1.6	1.3

Output: 2 x 3

2 – Transposed Convolution



Transposed Convolution

0	3	1	1
3	1	2	0
3	4	2	3
3	0	0	2

Convolution

7	8
15	13

7	8
15	13

Transposed

Convolution

0	3	1	1
3	1	2	0
3	4	2	3
3	0	0	2

2 – Transposed Convolution



Transposed Convolution

Input

2	2
1	4

Kernel

1	1
1	1

2	2	
2	2	

+

	2	2
	2	2

+

1	1	
1	1	

+

	4	4
	4	4

2	4	2
3	9	6
1	5	4

2 – Transposed Convolution



Transposed Convolution – Demo

```
input = torch.randint(  
    5, (1, 2, 2),  
    dtype=torch.float32  
)  
input
```

```
tensor([[[2., 2.],  
         [1., 4.]]])
```

```
conv_layer = nn.ConvTranspose2d(  
    in_channels=1,  
    out_channels=1,  
    kernel_size=2,  
    bias=False  
)
```

```
conv_layer.weight.data = torch.ones(  
    conv_layer.weight.data.shape  
)  
conv_layer.weight
```

```
Parameter containing:  
tensor([[[[1., 1.],  
          [1., 1.]]]], requires_grad=True)
```

```
output = conv_layer(input)  
output
```

```
tensor([[[2., 4., 2.],  
         [3., 9., 6.],  
         [1., 5., 4.] ]], grad_fn=<SqueezeBackward1>)
```

2 – Transposed Convolution



Transposed Convolution

2	2
1	4

Input

1	1
1	1

Kernel



2	4	2
3	9	6
1	5	4



1

Bias

3	5	3
4	10	7
2	6	5

2 – Transposed Convolution



Transposed Convolution – Demo

```
input = torch.randint(  
    5, (1, 2, 2),  
    dtype=torch.float32  
)  
input
```

```
tensor([[[2., 2.],  
         [1., 4.]]])
```

```
conv_layer = nn.ConvTranspose2d(  
    in_channels=1,  
    out_channels=1,  
    kernel_size=2,  
    bias=True  
)
```

```
conv_layer.weight.data = torch.ones(  
    conv_layer.weight.data.shape  
)  
conv_layer.weight
```

Parameter containing:
tensor([[[[1., 1.],
 [1., 1.]]]], requires_grad=True)

```
conv_layer.bias = nn.Parameter(  
    torch.tensor([1], dtype=torch.float32)  
)  
conv_layer.bias
```

Parameter containing:
tensor([1.], requires_grad=True)

```
output = conv_layer(input)  
output
```

```
tensor([[[ 3.,  5.,  3.],  
         [ 4., 10.,  7.],  
         [ 2.,  6.,  5.] ]], grad_fn=<SqueezeBackward1>)
```

2 – Transposed Convolution



Padding

2	2
1	4

Input

1	1
1	1

Kernel



2	4	2
3	9	6
1	5	4

Padding (1, 1)

9

2	4	2
3	9	6
1	5	4

Padding (1, 0)

3 9 6

2 – Transposed Convolution



Padding – Demo

```
input = torch.randint(
    5, (1, 2, 2),
    dtype=torch.float32
)
input
tensor([[[2., 2.],
         [1., 4.]])]
```

```
conv_layer = nn.ConvTranspose2d(
    in_channels=1,
    out_channels=1,
    kernel_size=2,
    padding=1,
    bias=False
```

```
conv_layer.weight.data = torch.ones(
    conv_layer.weight.data.shape
)
conv_layer.weight
Parameter containing:
tensor([[[[1., 1.],
         [1., 1.]]]], requires_grad=True)
```

```
output = conv_layer(input)
output
```

```
tensor([[[[9.]]], grad_fn=<SqueezeBackward1>)
```

```
conv_layer = nn.ConvTranspose2d(
    in_channels=1,
    out_channels=1,
    kernel_size=2,
    padding=(1, 0),
    bias=False
)
```

```
conv_layer.weight.data = torch.ones(
    conv_layer.weight.data.shape
)
conv_layer.weight
Parameter containing:
tensor([[[[1., 1.],
         [1., 1.]]]], requires_grad=True)
```

```
output = conv_layer(input)
output
```

```
tensor([[[[3., 9., 6.]]], grad_fn=<SqueezeBackward1>)
```


2 – Transposed Convolution



Stride

2	2
1	4

Input

1	1
1	1

Kernel



Stride (1, 1)

2	4	2
3	9	6
1	5	4



Stride (2, 2)

2	2	2	2
2	2	2	2
1	1	4	4
1	1	4	4

2 – Transposed Convolution



Stride – Demo

```
input = torch.randint(  
    5, (1, 2, 2),  
    dtype=torch.float32  
)  
input  
  
tensor([[[2., 2.],  
         [1., 4.]])])
```

```
conv_layer = nn.ConvTranspose2d(  
    in_channels=1,  
    out_channels=1,  
    kernel_size=2,  
    stride=2,  
    bias=False  
)
```

```
conv_layer.weight.data = torch.ones(  
    conv_layer.weight.data.shape  
)  
conv_layer.weight
```

```
Parameter containing:  
tensor([[[[1., 1.],  
         [1., 1.]]]], requires_grad=True)
```

```
output = conv_layer(input)  
output
```

```
tensor([[[2., 2., 2., 2.],  
         [2., 2., 2., 2.],  
         [1., 1., 4., 4.],  
         [1., 1., 4., 4.]]], grad_fn=<SqueezeBackward1>)
```

2 – Transposed Convolution



Padding & Stride

2	2
1	4

Input

1	1
1	1

Kernel

Stride (2, 2)

2	2	2	2
2	2	2	2
1	1	4	4
1	1	4	4

Padding (1, 1)

2	2
1	4

2 – Transposed Convolution



Padding & Stride – Demo

```
input = torch.randint(  
    5, (1, 2, 2),  
    dtype=torch.float32  
)  
input
```

```
tensor([[[2., 2.],  
         [1., 4.]])
```

```
conv_layer = nn.ConvTranspose2d(  
    in_channels=1,  
    out_channels=1,  
    kernel_size=2,  
    stride=2,  
    padding=1,  
    bias=False  
)
```

```
conv_layer.weight.data = torch.ones(  
    conv_layer.weight.data.shape  
)  
conv_layer.weight
```

```
Parameter containing:  
tensor([[[[1., 1.],  
         [1., 1.]])], requires_grad=True)
```

```
output = conv_layer(input)  
output
```

```
tensor([[[2., 2.],  
         [1., 4.]])], grad_fn=<SqueezeBackward1>)
```

2 – Transposed Convolution



Padding & Stride


2	2
1	4

Input

 $M \times N$

1	1
1	1


Kernel

 $K \times O$ 

Stride (2, 2)

 $S \times T$

2	2	2	2
2	2	2	2
1	1	4	4
1	1	4	4



Padding (1, 1)

 $P \times Q$

2	2
1	4

$$H_{\text{out}} = (M - 1) * S - 2P + (K - 1) + 1$$

$$W_{\text{out}} = (N - 1) * T - 2Q + (O - 1) + 1$$

2 – Transposed Convolution



Multiple Channels

0	1
3	0

Input #1

1	1
1	1

Kernel #1

Stride (2, 2)

0	0	1	1
0	0	1	1
3	3	0	0
3	3	0	0

1	1
2	4

Input #2

1	1
1	1

Kernel #2

Stride (2, 2)

1	1	1	1
1	1	1	1
2	2	4	4
2	2	4	4

1	1	2	2
1	1	2	2
5	5	4	4
5	5	4	4

2 – Transposed Convolution



Multiple Channels – Demo

```
input = torch.randint(  
    5, (2, 2, 2),  
    dtype=torch.float32  
)  
input
```

```
tensor([[[0., 1.],  
         [3., 0.]],  
  
        [[1., 1.],  
         [2., 4.]]])
```

```
conv_layer = nn.ConvTranspose2d(  
    in_channels=2,  
    out_channels=1,  
    kernel_size=2,  
    stride=2,  
    bias=False  
)
```

```
conv_layer.weight.data = torch.ones(  
    conv_layer.weight.data.shape  
)  
conv_layer.weight
```

Parameter containing:

```
tensor([[[[1., 1.],  
         [1., 1.]]],
```

```
        [[1., 1.],  
         [1., 1.]]]], requires_grad=True)
```


```
output = conv_layer(input)  
output
```

```
tensor([[[1., 1., 2., 2.],  
         [1., 1., 2., 2.],  
         [5., 5., 4., 4.],  
         [5., 5., 4., 4.]]], grad_fn=<SqueezeBackward1>)
```

3 – UNet Model



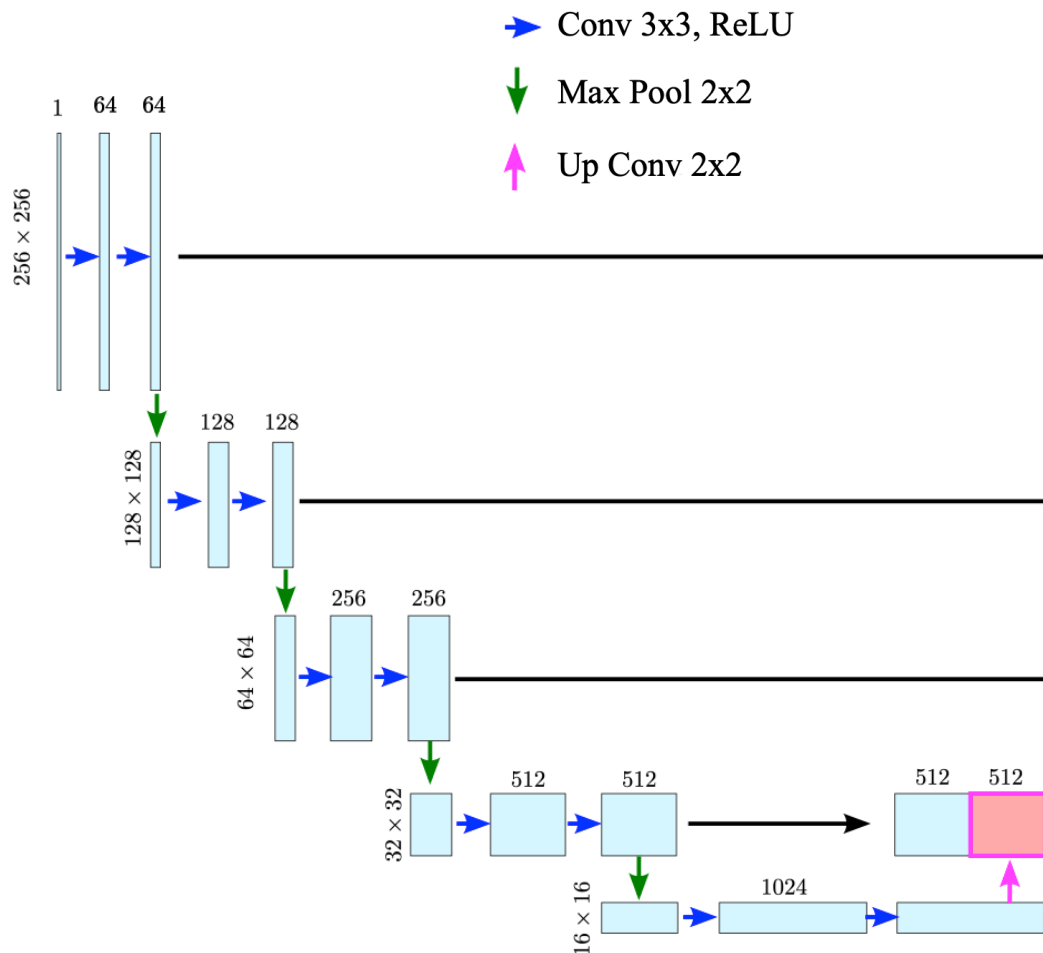
UNet



3 – UNet Model



ConvBlock – Demo



```
class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels) -> None:
        super().__init__()
        self.conv_block = nn.Sequential(
            nn.Conv2d(
                in_channels, out_channels, kernel_size=3, padding=1, bias=False
            ),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True),
            nn.Conv2d(
                out_channels, out_channels, kernel_size=3, padding=1, bias=False
            ),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True),
        )

    def forward(self, x):
        x = self.conv_block(x)
        return x
```

```
input = torch.randint(5, (1, 1, 256, 256), dtype=torch.float32)
```

```
in_conv = ConvBlock(in_channels=1, out_channels=64)
```


```
output = in_conv(input)
output.shape
```

```
torch.Size([1, 64, 256, 256])
```

3 – UNet Model



Encoder – Demo



```
# Down-Sample
class Encoder(nn.Module):
    def __init__(self, in_channels, out_channels) -> None:
        super().__init__()
        self.encoder = nn.Sequential(
            nn.MaxPool2d(2),
            ConvBlock(in_channels, out_channels)
        )

    def forward(self, x):
        x = self.encoder(x)
        return x
```

```
encoder = Encoder(in_channels=64, out_channels=128)
```

```
input = torch.randint(5, (1, 64, 256, 256), dtype=torch.float32)
```

```
output = encoder(input)
output.shape
```

```
torch.Size([1, 128, 128, 128])
```

3 – UNet Model



Decoder – Demo

Up-Sample

```
class Decoder(nn.Module):
    def __init__(self, in_channels, out_channels) -> None:
        super().__init__()
        self.conv_trans = nn.ConvTranspose2d(
            in_channels, out_channels, kernel_size=4, stride=2, padding=1
        )
        self.conv_block = ConvBlock(in_channels, out_channels)

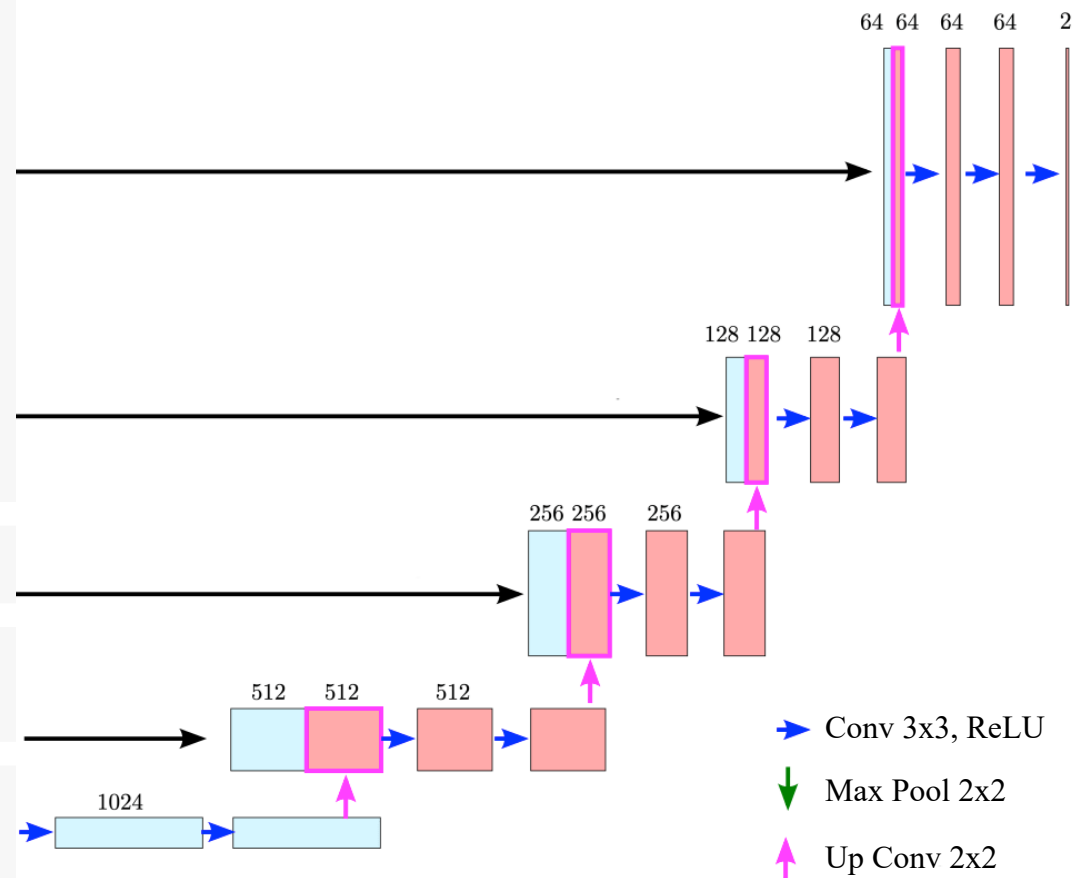
    def forward(self, x1, x2):
        x1 = self.conv_trans(x1)
        x = torch.cat([x2, x1], dim=1)
        x = self.conv_block(x)
        return x
```

```
decoder = Decoder(in_channels=128, out_channels=64)
```

```
x1 = torch.randint(5, (1, 128, 128, 128), dtype=torch.float32)
x2 = torch.randint(5, (1, 64, 256, 256), dtype=torch.float32)
```

```
output = decoder(x1, x2)
output.shape
```


```
torch.Size([1, 64, 256, 256])
```



3 – UNet Model



UNet – Demo



3 – UNet Model



UNet – Demo

```
class UNet(nn.Module):
    def __init__(self, n_channels, n_classes):
        super().__init__()
        self.n_channels = n_channels
        self.n_classes = n_classes

        self.in_conv = ConvBlock(n_channels, 64)

        self.enc_1 = Encoder(64, 128)
        self.enc_2 = Encoder(128, 256)
        self.enc_3 = Encoder(256, 512)
        self.enc_4 = Encoder(512, 1024)

        self.dec_1 = Decoder(1024, 512)
        self.dec_2 = Decoder(512, 256)
        self.dec_3 = Decoder(256, 128)
        self.dec_4 = Decoder(128, 64)

        self.out_conv = nn.Conv2d(
            64, n_classes, kernel_size=1
        )
```

```
def forward(self, x):
    x1 = self.in_conv(x)

    x2 = self.enc_1(x1)
    x3 = self.enc_2(x2)
    x4 = self.enc_3(x3)
    x5 = self.enc_4(x4)

    x = self.dec_1(x5, x4)
    x = self.dec_2(x, x3)
    x = self.dec_3(x, x2)
    x = self.dec_4(x, x1)

    x = self.out_conv(x)

    return x
```

```
model = UNet(n_channels=1, n_classes=2)

input = torch.randint(
    5, (4, 1, 256, 256), dtype=torch.float32
)

predictions = model(input)

predictions.shape

torch.Size([4, 2, 256, 256])
```



AI VIET NAM

@aivietnam.edu.vn

Thanks!

Any questions?