

**AI VIETNAM**  
**All-in-One Course**  
**(TA Session)**

# Traffic Sign Detection Project



**AI VIET NAM**  
[@aivietnam.edu.vn](http://aivietnam.edu.vn)

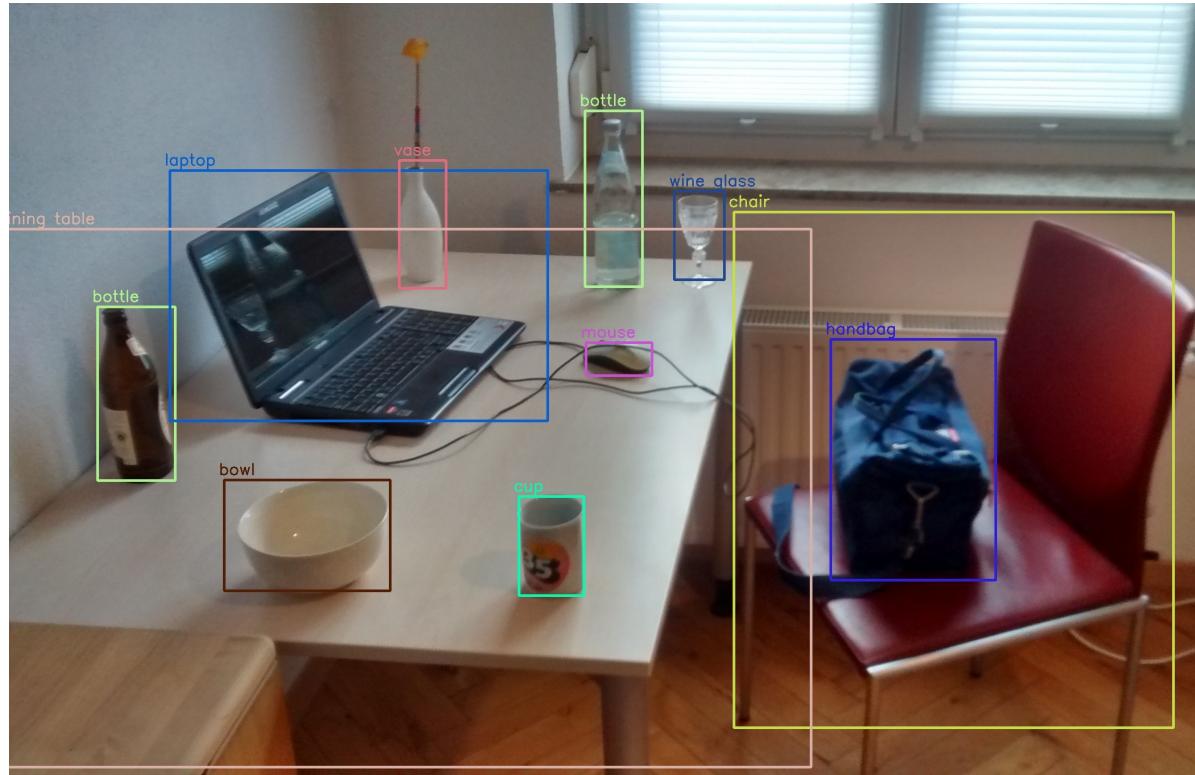
**Dinh-Thang Duong – TA**

# Outline

- Introduction
- Project Pipeline
- Traffic Sign Classification
- Traffic Sign Localization
- Question

# Introduction

## ❖ Object Detection



**Object Detection:** A computer vision task that involves identifying and locating objects of interest in an image or video.

# Introduction

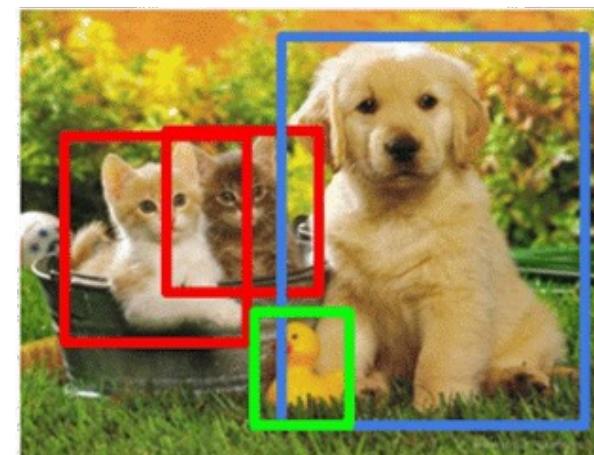
## ❖ Image Classification vs Object Detection

Classification



CAT

Object Detection



CAT, DOG, DUCK

Unlike image classification, where the goal is to assign a label to an entire image, object detection aims to **identify multiple objects** and **provide bounding boxes** around them, specifying where each of object is in the image.

# Introduction

## ❖ Two-stage Object Detection

**Classification**



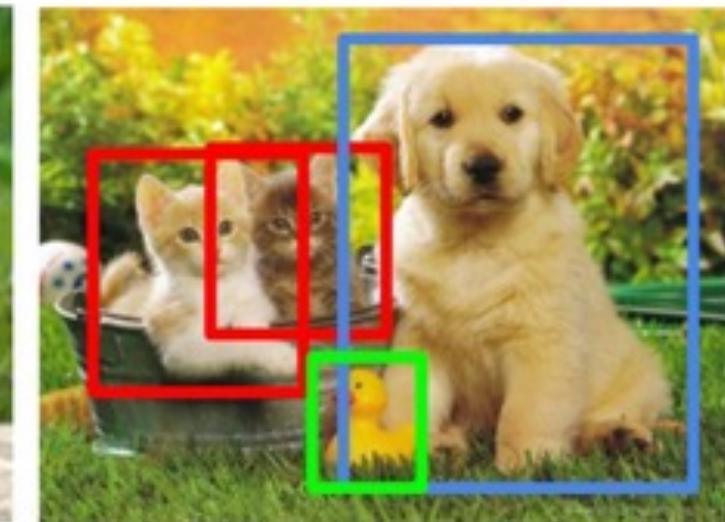
CAT

**Classification + Localization**



CAT

**Object Detection**

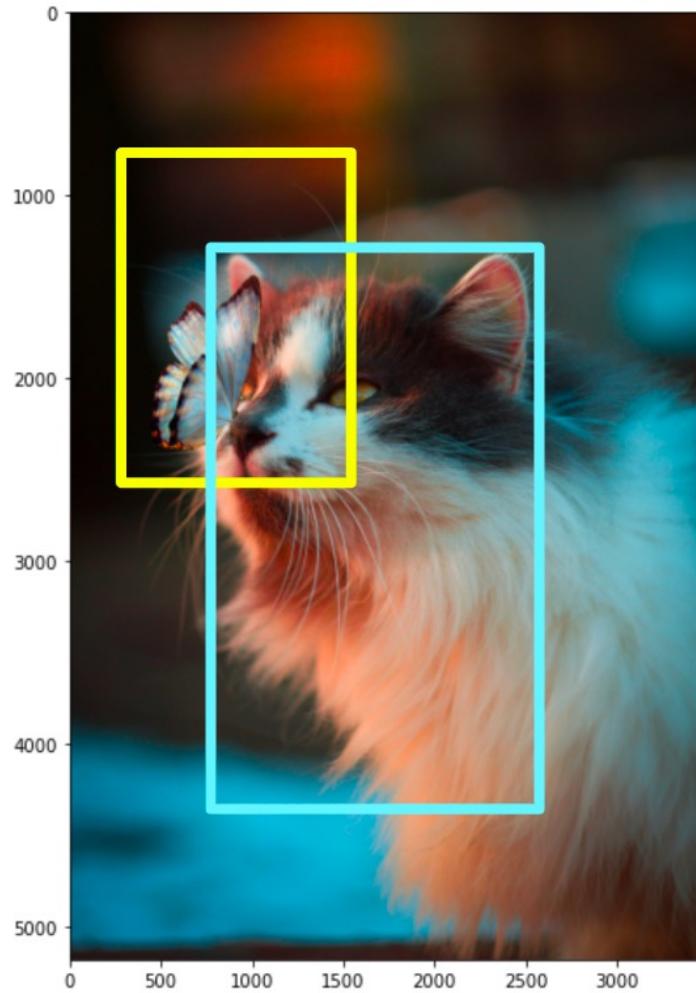


CAT, DOG, DUCK

**Two-stage Object Detection:** A combination of Object Localization and Object Classification.

# Introduction

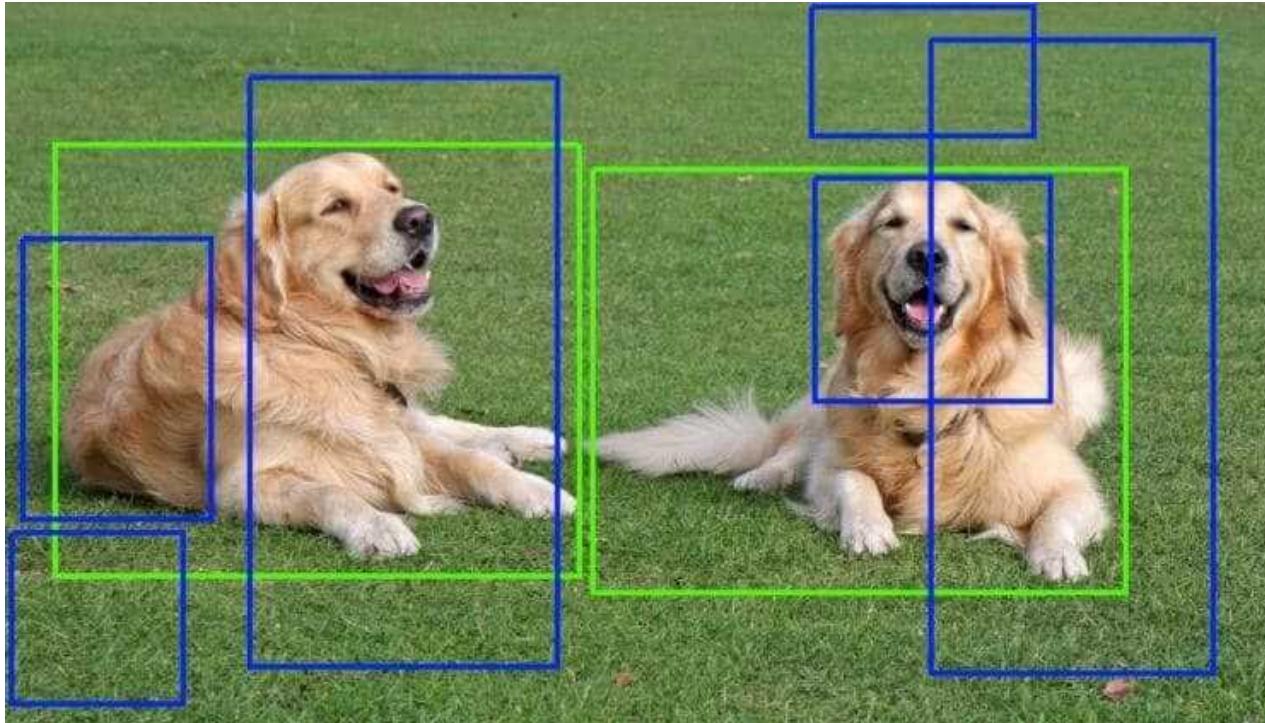
## ❖ Object Localization



**Object Localization:** A computer vision task that seeks to pinpoint the location of a specific object in an image. The primary output for object localization is usually a bounding box that encompasses the objects of interest.

# Introduction

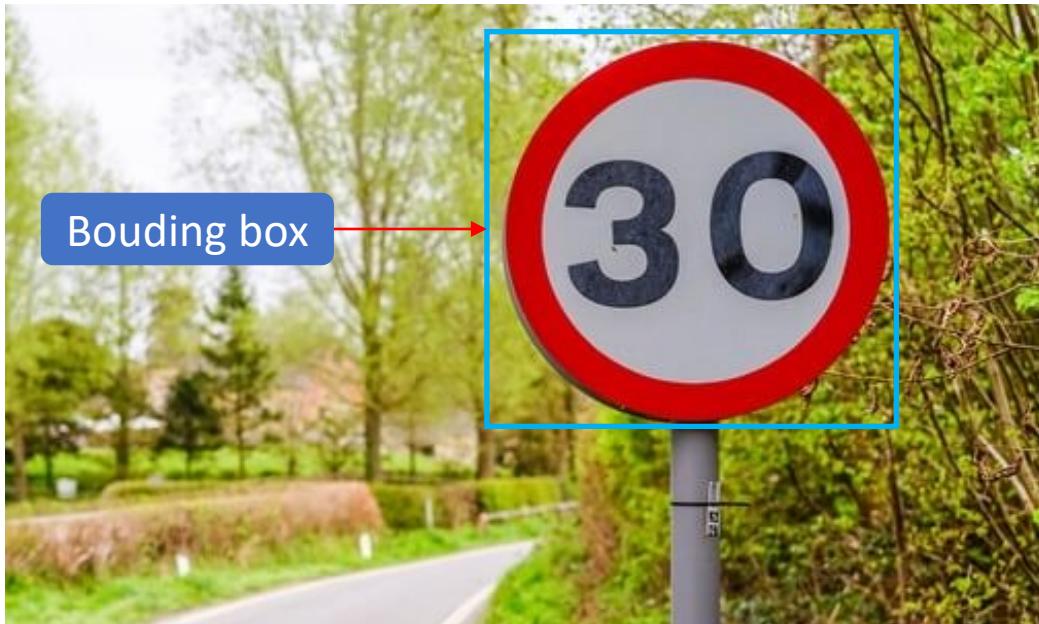
## ❖ Object Localization



**Example:** Multiple bounding boxes are presented to identify the position of two dogs. Green seems more accurate than Blue.

# Introduction

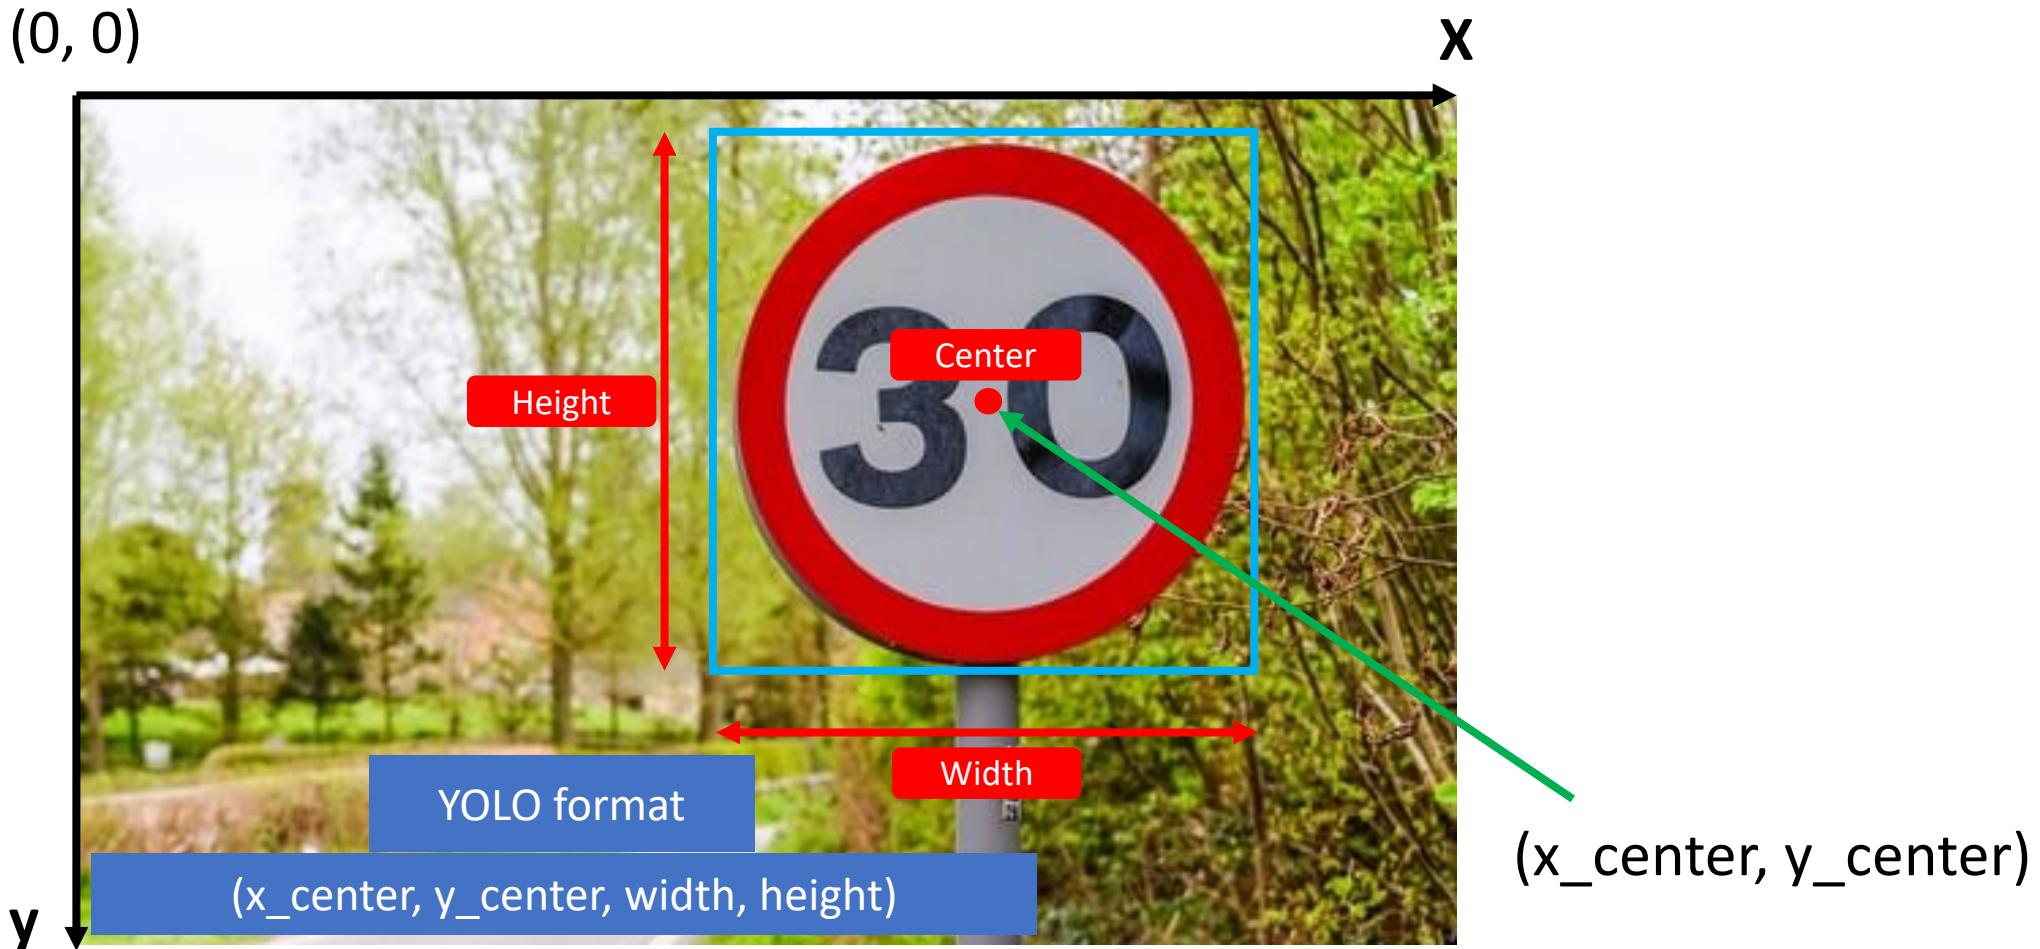
## ❖ Object Localization: Bounding box



**Bounding box:** A rectangular box used to describe the location of an object within an image. Bounding box is often represented as a tuple of (x\_center, y\_center, width, height) or (xmin, ymin, xmax, ymax).

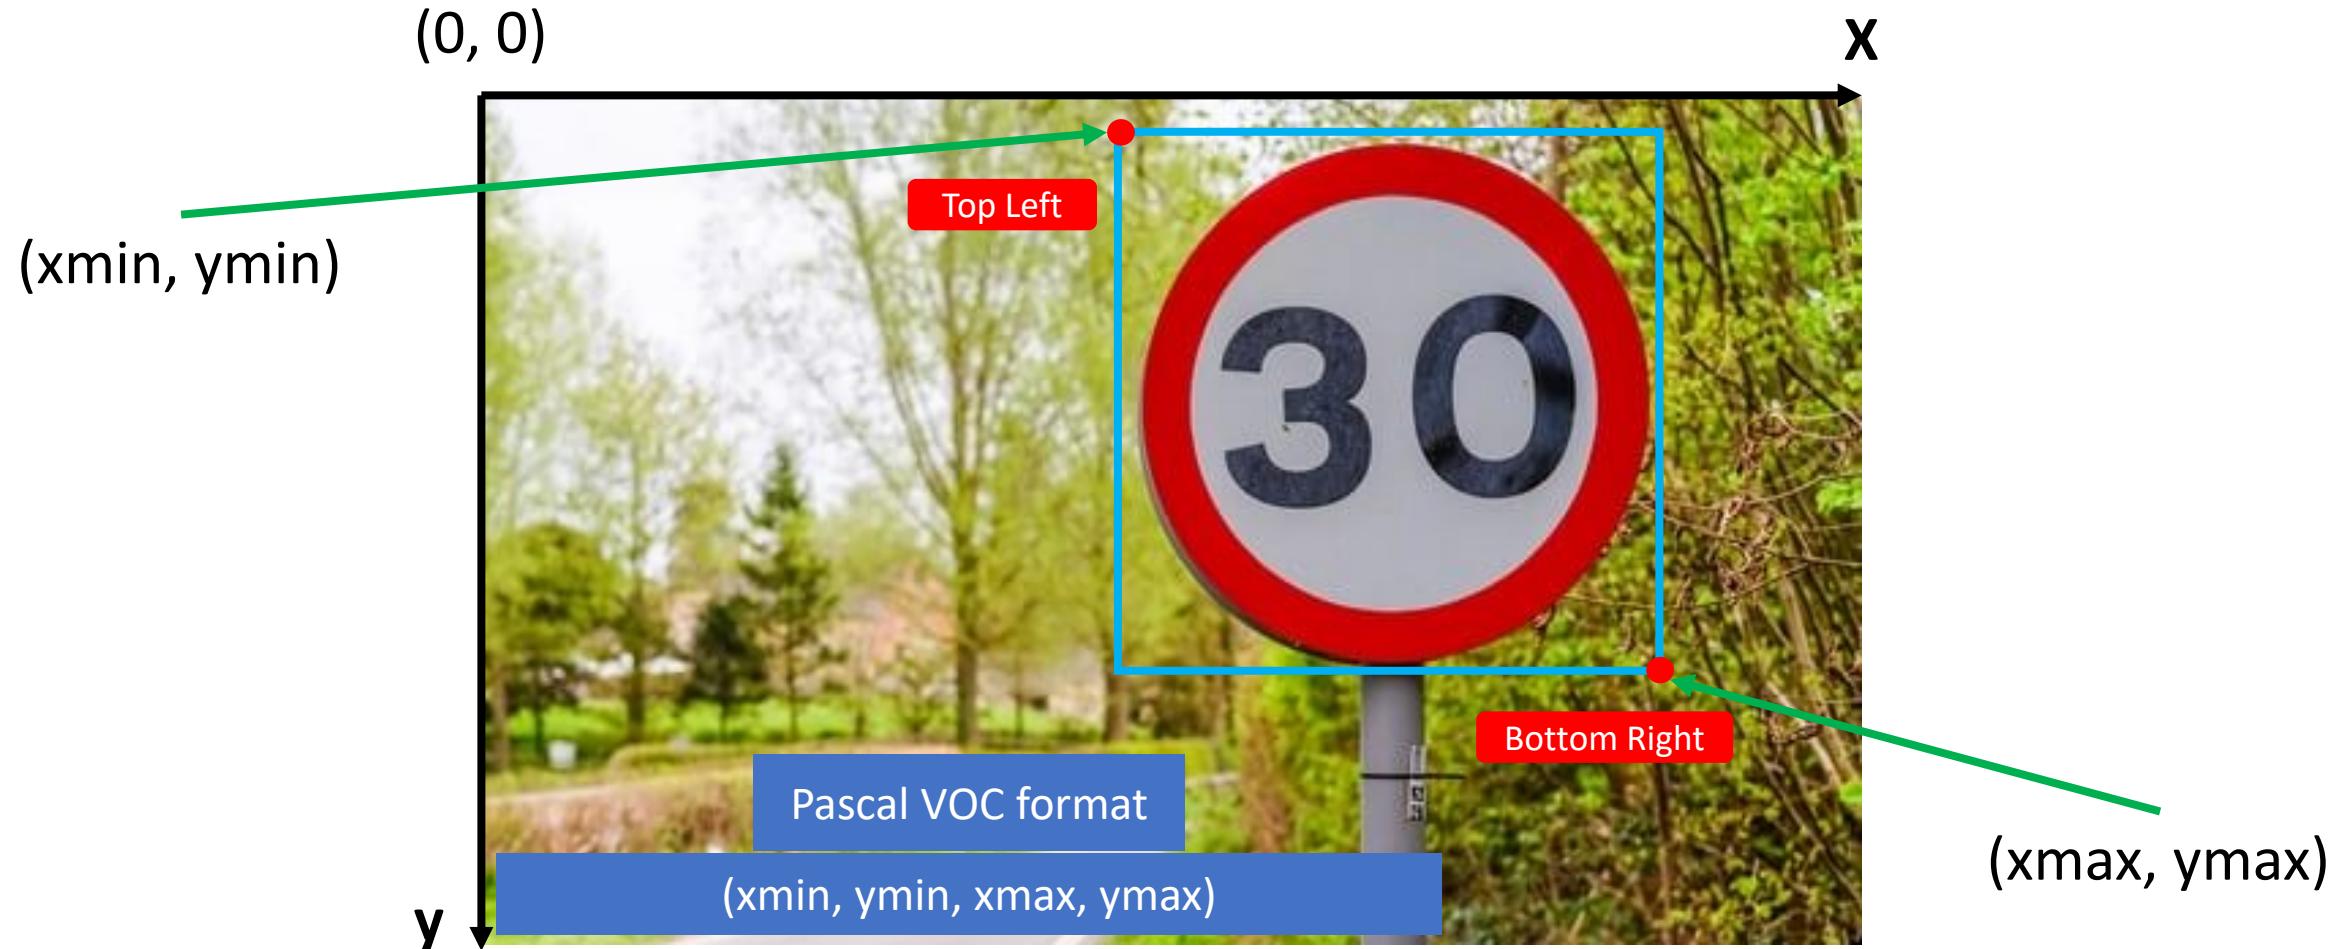
# Introduction

## ❖ Object Localization: Bounding box



# Introduction

## ❖ Object Localization: Bounding box



# Introduction

## ❖ Object Localization: Draw bounding box to an image

Syntax: `cv2.rectangle(image, start_point, end_point, color, thickness)`

Parameters:

**image:** It is the image on which rectangle is to be drawn.

**start\_point:** It is the starting coordinates of rectangle. The coordinates are represented as tuples of two values i.e. (X coordinate value, Y coordinate value).

**end\_point:** It is the ending coordinates of rectangle. The coordinates are represented as tuples of two values i.e. (X coordinate value, Y coordinate value).

**color:** It is the color of border line of rectangle to be drawn. For **BGR**, we pass a tuple. eg: (255, 0, 0) for blue color.

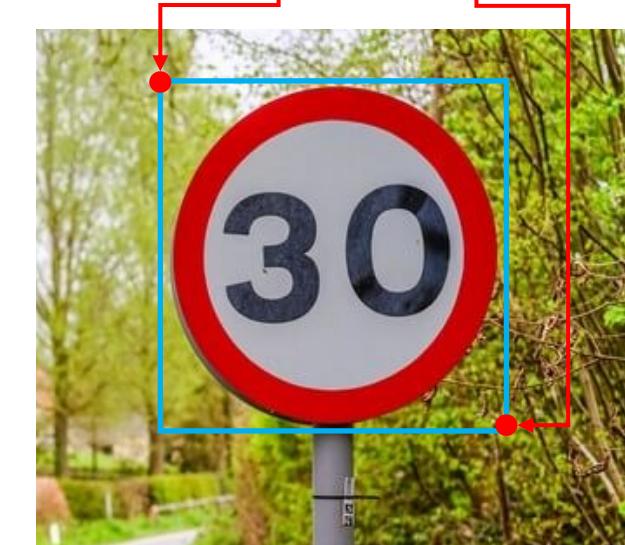
**thickness:** It is the thickness of the rectangle border line in **px**. Thickness of **-1 px** will fill the rectangle shape by the specified color.

**Return Value:** It returns an image.

We can use `cv2.rectangle()` to draw a bounding box

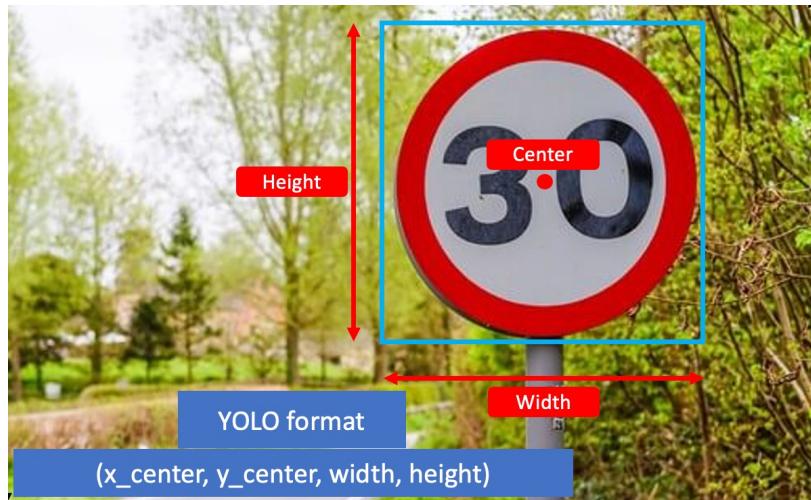
Draw bounding box into image using `cv2.rectangle`

```
1 image = cv2.rectangle(image, top_left, bottom_right, color, thickness)
```

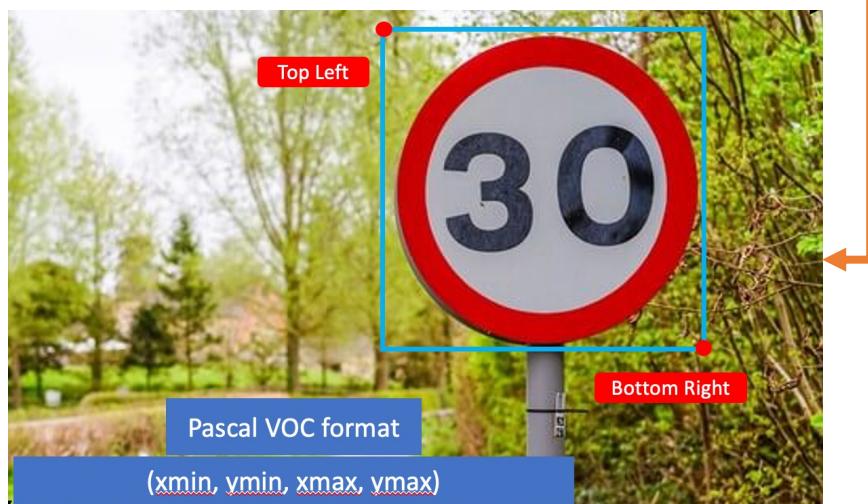


# Introduction

## ❖ Object Localization: Convert bounding box format



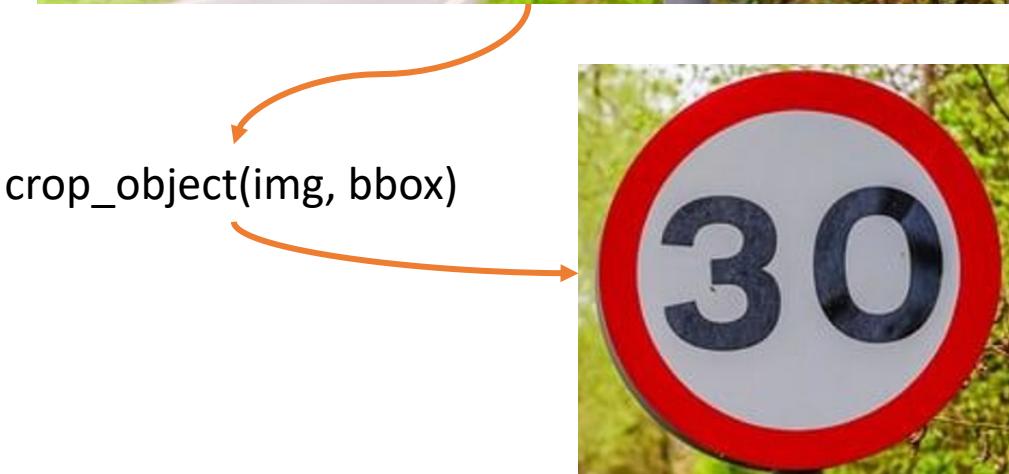
Given a tuple bbox of (x\_center, y\_center, width, height):



```
1 def xywh2xyxy(bbox):
2     """
3     Input:
4         - bbox: Tuple(x_center, y_center, width, height)
5     Output:
6         - Tuple(x1, y1, x2, y2) Where:
7             - x1, y1: Top left point
8             - x2, y2: Bottom right point
9     """
10    x1 = bbox[0] - bbox[2] / 2 # top left x
11    y1 = bbox[1] - bbox[3] / 2 # top left y
12    x2 = bbox[0] + bbox[2] / 2 # bottom right x
13    y2 = bbox[1] + bbox[3] / 2 # bottom right y
14    return (x1, y1, x2, y2)
```

# Introduction

## ❖ Object Localization: Convert bounding box format



Given an image and a tuple bbox of (xmin, ymin, xmax, ymax):

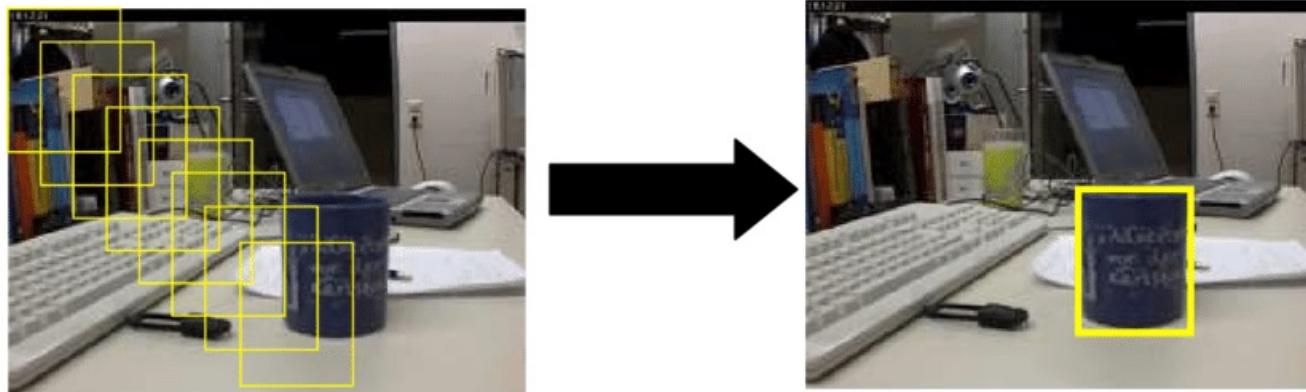
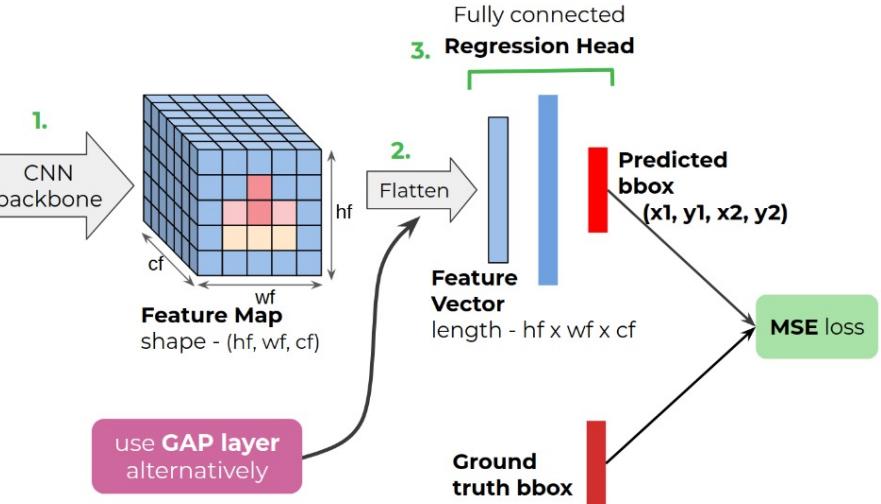
```
1 def crop_object(img, bbox):  
2     xmin, ymin, xmax, ymax = bbox  
3     object_img = img[ymin:ymax, xmin:xmax]  
4  
5     return object_img
```

# Introduction

## ❖ Object Localization: Methods



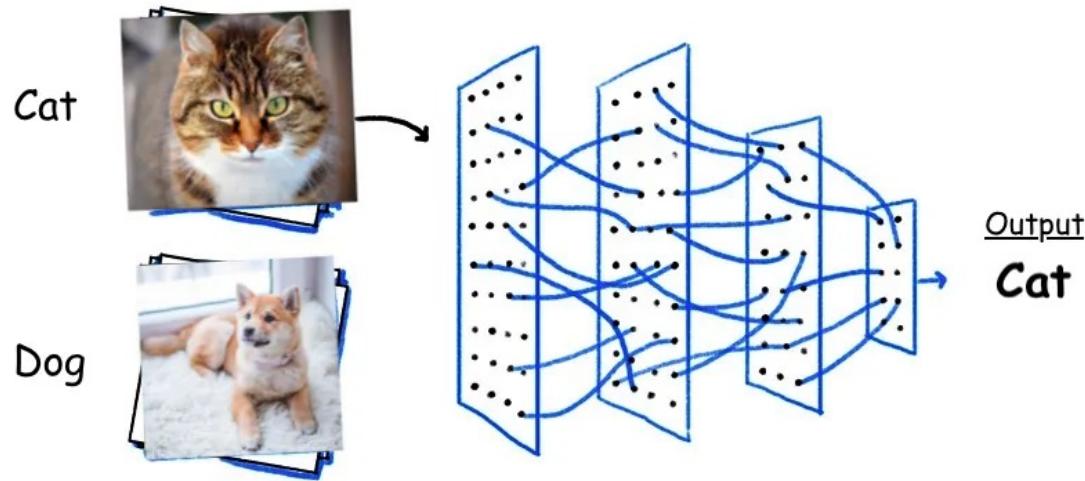
Input Image



There are several methods to solve the object localization task. In this project, we will use **Sliding Window** technique.

# Introduction

## ❖ Object Classification

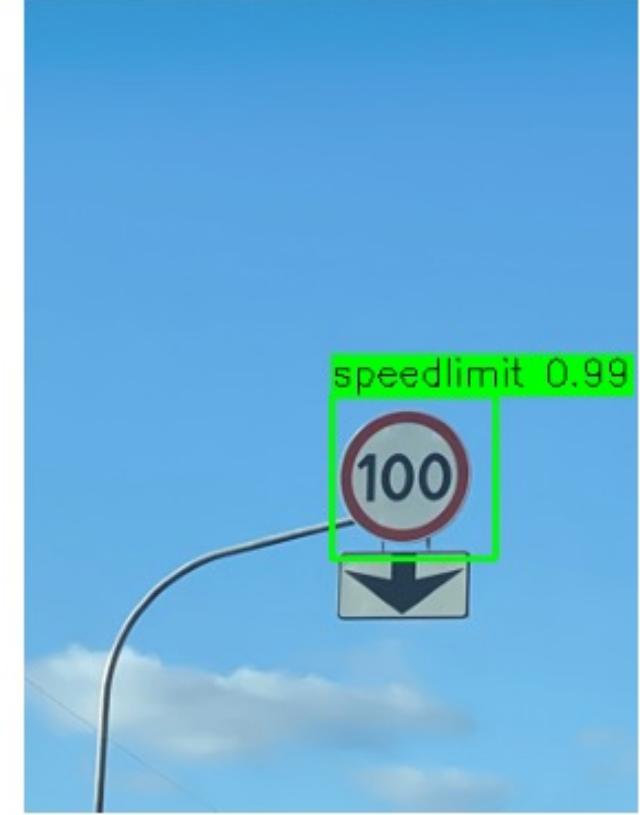
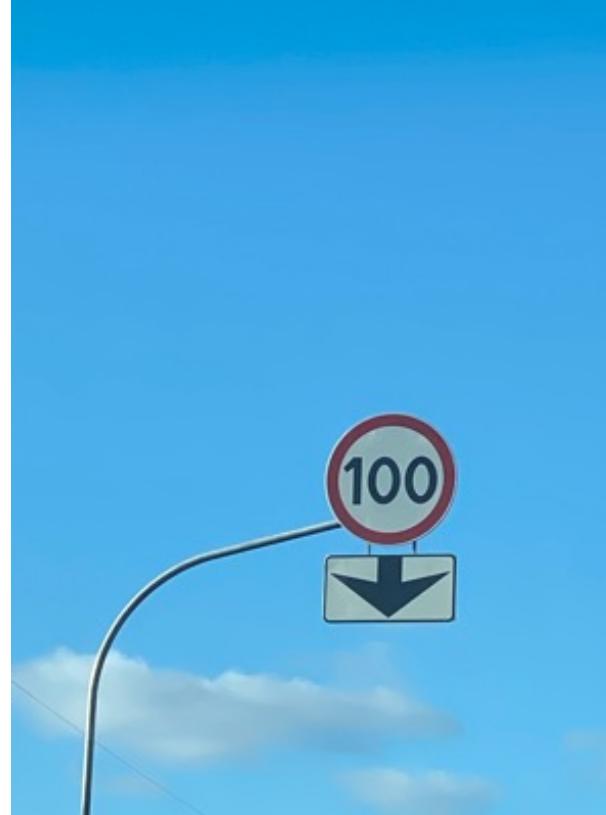


**Object Classification (Image Classification):** A computer vision task that involves assigning a predefined category or label to an image. In this project, we will use SVM to build an image classifier.

# Introduction

## ❖ Object Detection Output

**Output:** A list of predicted bounding boxes in the form of (xmin, ymin, xmax, ymax, confidence\_score)



Output Visualization

# Introduction

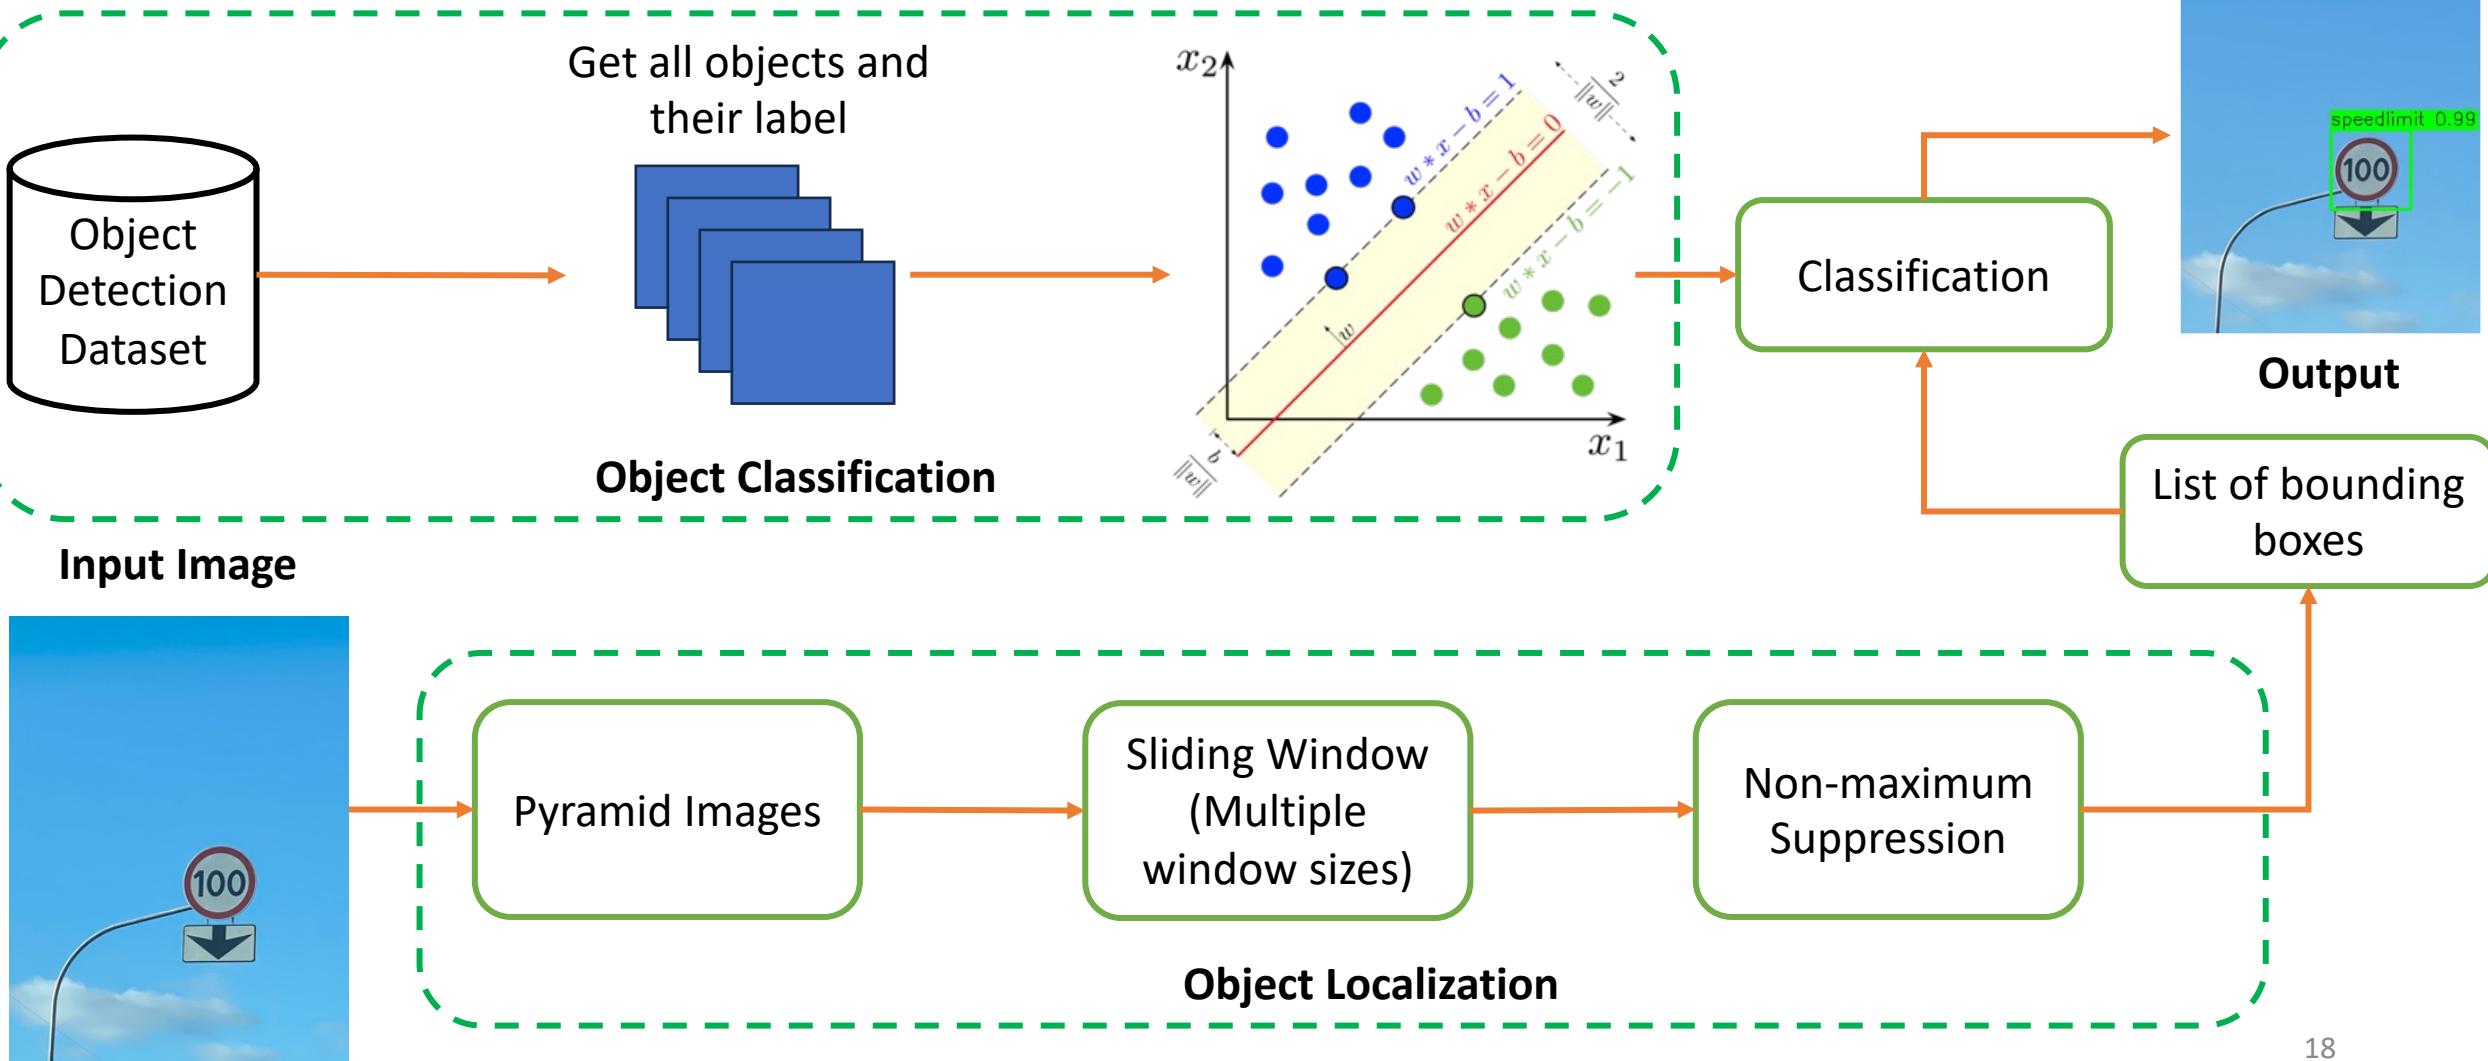
## ❖ Object Detection Confidence Score



**Confidence Score:** A value represents how ‘confident’ the model is to assign a class to the corresponding object. Confidence Score is often **taken from the output of Object Classifier**.

# Project Pipeline

## ❖ Introduction



# Project Pipeline

## ❖ Import essential libraries

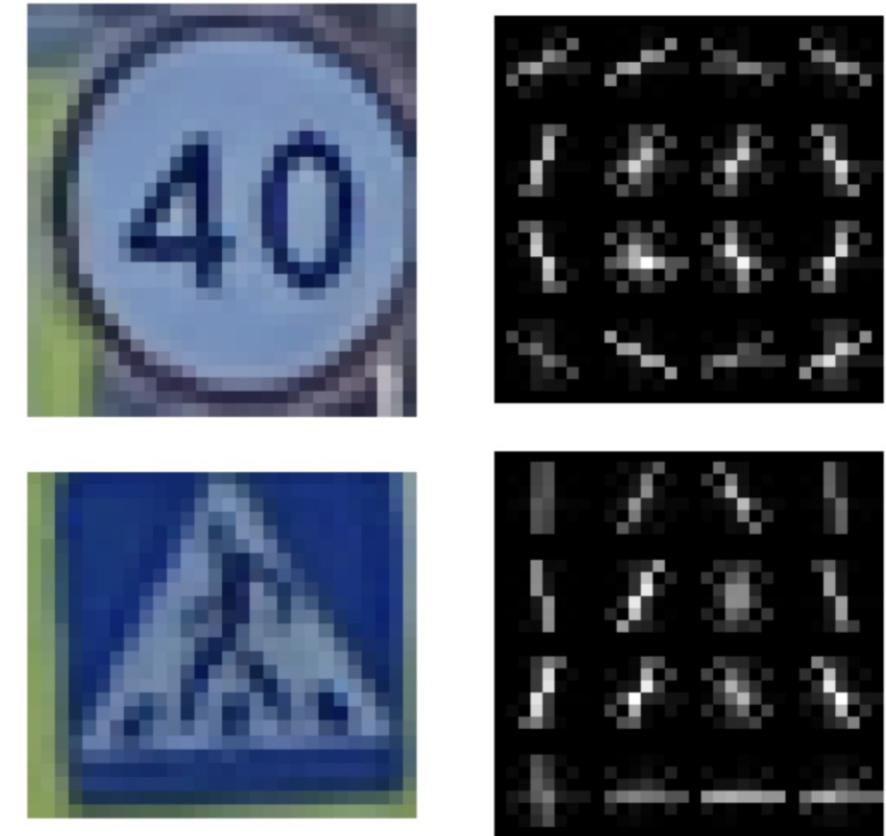
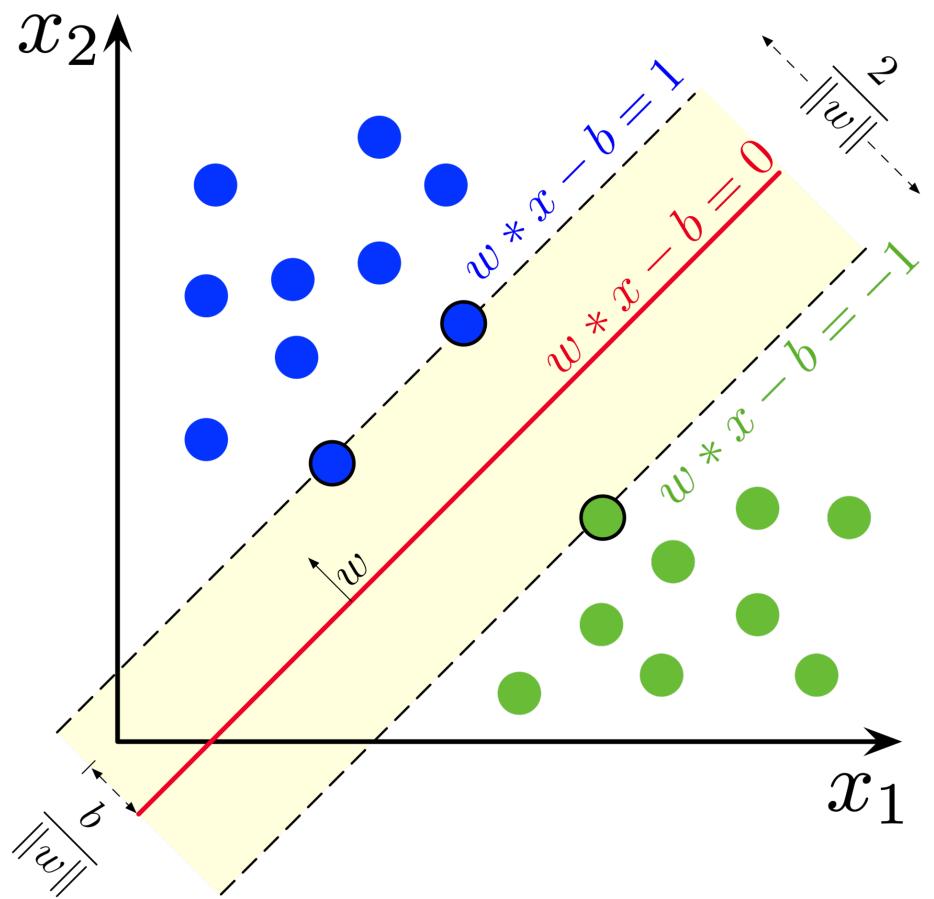
```
1 import time
2 import os
3 import cv2
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 import xml.etree.ElementTree as ET
8
9 from skimage.transform import resize
10 from skimage import feature
11 from sklearn.svm import SVC
12 from sklearn.preprocessing import LabelEncoder
13 from sklearn.preprocessing import StandardScaler
14 from sklearn.model_selection import train_test_split
15 from sklearn.metrics import accuracy_score
```



# Traffic Sign Classification

## ❖ Introduction

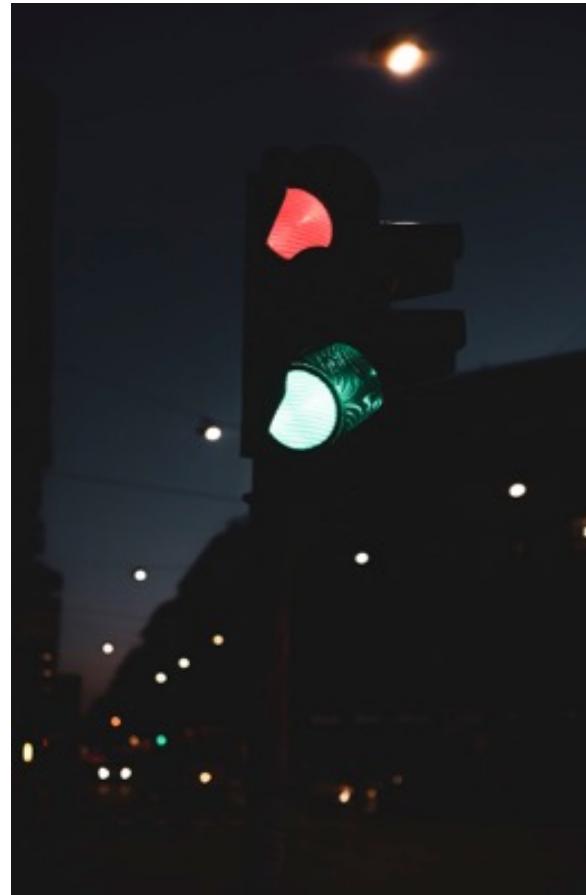
**Description:** Given a dataset of [Traffic Sign Detection](#). Build a Traffic Sign Image Classifier using SVM model.



# Traffic Sign Classification

## ❖ Step 1: Read dataset

images	annotations
road0.png	road0.xml
road1.png	road1.xml
road10.png	road10.xml
road100.png	road100.xml
road101.png	road101.xml
road102.png	road102.xml
road103.png	road103.xml
road104.png	road104.xml
road105.png	road105.xml



```
1 <annotation>
2   <folder>images</folder>
3   <filename>road0.png</filename>
4   <size>
5     <width>267</width>
6     <height>400</height>
7     <depth>3</depth>
8   </size>
9   <segmented>0</segmented>
10  <object>
11    <name>trafficlight</name>
12    <pose>Unspecified</pose>
13    <truncated>0</truncated>
14    <occluded>0</occluded>
15    <difficult>0</difficult>
16    <bndbox>
17      <xmin>98</xmin>
18      <ymin>62</ymin>
19      <xmax>208</xmax>
20      <ymax>232</ymax>
21    </bndbox>
22  </object>
23 </annotation>
```

# Traffic Sign Classification

## ❖ Step 1: Read dataset

```
1 <annotation>
2   <folder>images</folder> ← Folder contains the image
3   <filename>road0.png</filename> ← Filename of the image
4   <size>
5     <width>267</width>
6     <height>400</height> } ← Shape of the image
7     <depth>3</depth>
8   </size>
9   <segmented>0</segmented>
10  <object>
11    <name>trafficlight</name> ← Label of the object (classname)
12    <pose>Unspecified</pose>
13    <truncated>0</truncated>
14    <occluded>0</occluded>
15    <difficult>0</difficult>
16    <bndbox>
17      <xmin>98</xmin>
18      <ymin>62</ymin>
19      <xmax>208</xmax>
20      <ymax>232</ymax>
21    </bndbox> } ← Bounding box (Pascal VOC format) of the
22  </object> object
23 </annotation>
```

# Traffic Sign Classification

## ❖ Step 1: Read dataset

```
1 <annotation>
2     <folder>images</folder>
3     <filename>road0.png</filename>
4     <size>
5         <width>267</width>
6         <height>400</height>
7         <depth>3</depth>
8     </size>
9     <segmented>0</segmented>
10    <object>
11        <name>trafficlight</name>
12        <pose>Unspecified</pose>
13        <truncated>0</truncated>
14        <occluded>0</occluded>
15        <difficult>0</difficult>
16        <bndbox>
17            <xmin>98</xmin>
18            <ymin>62</ymin>
19            <xmax>208</xmax>
20            <ymax>232</ymax>
21        </bndbox>
22    </object>
23 </annotation>
```

## xml.etree.ElementTree — The ElementTree XML API

[Source code: Lib/xml/etree/ElementTree.py](#)

The `xml.etree.ElementTree` module implements a simple and efficient API for parsing and creating XML data.

*Changed in version 3.3:* This module will use a fast implementation whenever available.

*Deprecated since version 3.3:* The `xml.etree.cElementTree` module is deprecated.

```
1 import xml.etree.ElementTree as ET
2
3 xml_file = 'annotations/road0.xml'
4 tree = ET.parse(xml_file)
5 root = tree.getroot()
6 folder_name = root.find('folder').text
7 filename = root.find('filename').text
8
9 print(f'Folder: {folder_name}')
10 print(f'Filename: {filename}')
```

Folder: images  
Filename: road0.png

# Traffic Sign Classification

## ❖ Step 1: Read dataset

images	annotations
road0.png	road0.xml
road1.png	road1.xml
road10.png	road10.xml
road100.png	road100.xml
road101.png	road101.xml
road102.png	road102.xml
road103.png	road103.xml
road104.png	road104.xml
road105.png	road105.xml

We want to extract all objects as images using bounding box and their corresponding label as string.

# Traffic Sign Classification

## ❖ Step 1: Read dataset

```
1 <annotation>
2   <folder>images</folder>
3   <filename>road0.png</filename>
4   <size>
5     <width>267</width>
6     <height>400</height>
7     <depth>3</depth>
8   </size>
9   <segmented>0</segmented>
10  <object>
11    <name>trafficlight</name>
12    <pose>Unspecified</pose>
13    <truncated>0</truncated>
14    <occluded>0</occluded>
15    <difficult>0</difficult>
16    <bndbox>
17      <xmin>98</xmin>
18      <ymin>62</ymin>
19      <xmax>208</xmax>
20      <ymax>232</ymax>
21    </bndbox>
22  </object>
23 </annotation>
```

```
1 import xml.etree.ElementTree as ET
2
3 xml_file = 'annotations/road0.xml'
4 tree = ET.parse(xml_file)
5 root = tree.getroot()
6
7 for obj in root.findall('object'):
8     classname = obj.find('name').text
9
10    xmin = int(obj.find('bndbox/xmin').text)
11    ymin = int(obj.find('bndbox/ymin').text)
12    xmax = int(obj.find('bndbox/xmax').text)
13    ymax = int(obj.find('bndbox/ymax').text)
14
15    print(f'Classname: {classname}')
16    print(f'Bounding box: {xmin, ymin, xmax, ymax}')
```

Classname: trafficlight  
Bounding box: (98, 62, 208, 232)

# Traffic Sign Classification

## ❖ Step 1: Read dataset



```
1 print('Number of objects: ', len(img_lst))
2 print('Class names: ', list(set(label_lst)))
```

Number of objects: 1074

Class names: ['stop', 'crosswalk', 'speedlimit']

```
1 annotations_dir = 'annotations'
2 img_dir = 'images'
3
4 img_lst = []
5 label_lst = []
6 for xml_file in os.listdir(annotations_dir):
7     xml_filepath = os.path.join(annotations_dir, xml_file)
8     tree = ET.parse(xml_filepath)
9     root = tree.getroot()
10
11    folder = root.find('folder').text
12    img_filename = root.find('filename').text
13    img_filepath = os.path.join(img_dir, img_filename)
14    img = cv2.imread(img_filepath)
15
16    for obj in root.findall('object'):
17        classname = obj.find('name').text
18        if classname == 'trafficlight':
19            continue
20
21        xmin = int(obj.find('bndbox/xmin').text)
22        ymin = int(obj.find('bndbox/ymin').text)
23        xmax = int(obj.find('bndbox/xmax').text)
24        ymax = int(obj.find('bndbox/ymax').text)
25
26        object_img = img[ymin:ymax, xmin:xmax]
27        img_lst.append(object_img)
28        label_lst.append(classname)
```

We ignore 'trafficlight' class since we do Traffic Sign Classification

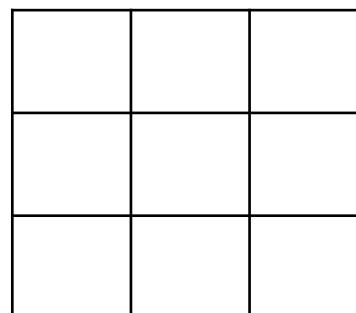
# Traffic Sign Classification

## ❖ Step 2: Preprocess image



```
1 print('Image shape: ')
2 print(img_lst[1].shape)
```

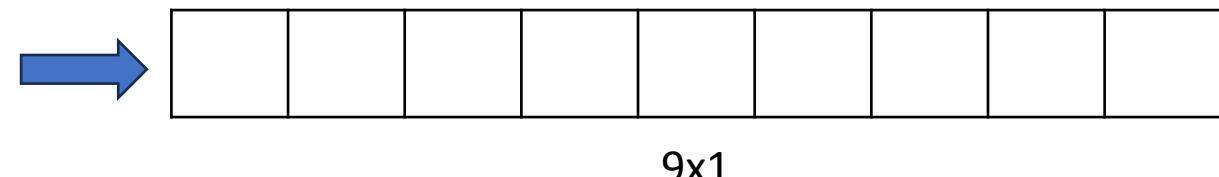
Image shape:  
(53, 55, 3)



3x3

We can flatten  
image to have the  
 $\text{len}(\text{shape}) == 1$ .

Currently, we cannot feed the entire image ( $\text{len}(\text{shape}) == 3$ ) as X to train  
an SVM model.

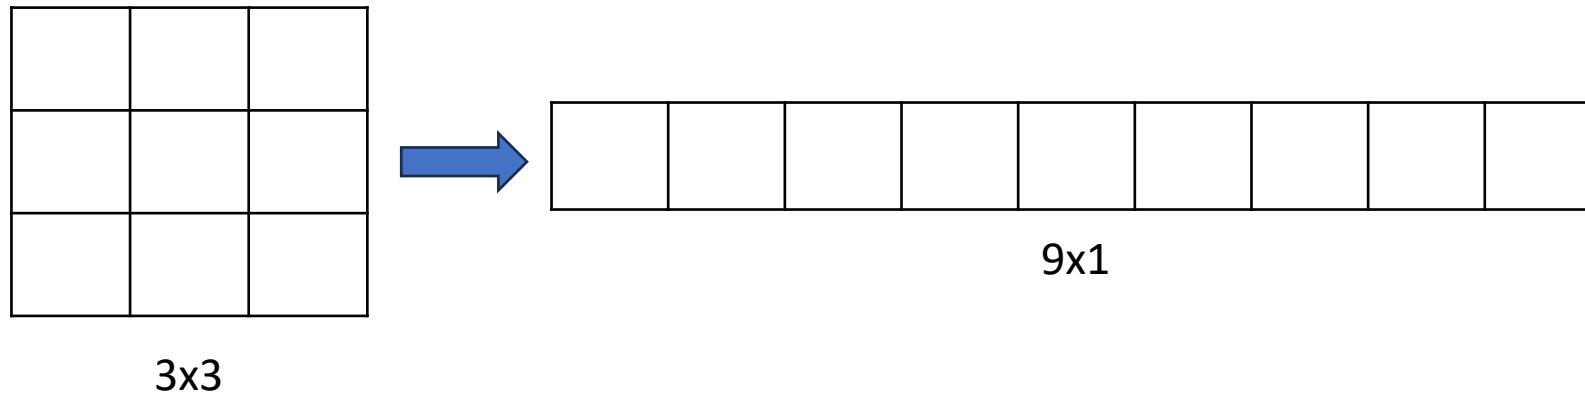


```
1 print('Image shape: ')
2 print(img_lst[1].shape)
3 flattened_img = img_lst[1].flatten()
4 print('Image shape after flattened: ')
5 print(flattened_img.shape)
```

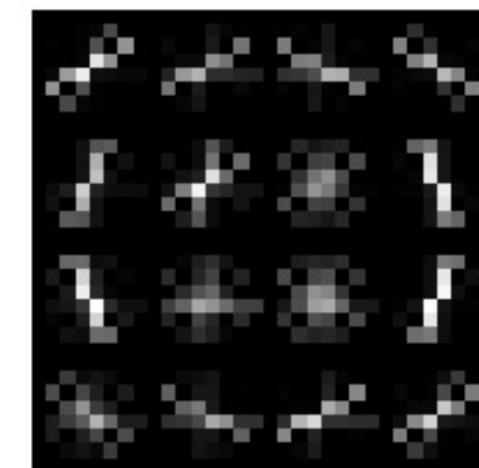
Image shape:  
(53, 55, 3)  
Image shape after flattened:  
(8745, )

# Traffic Sign Classification

## ❖ Step 2: Preprocess image



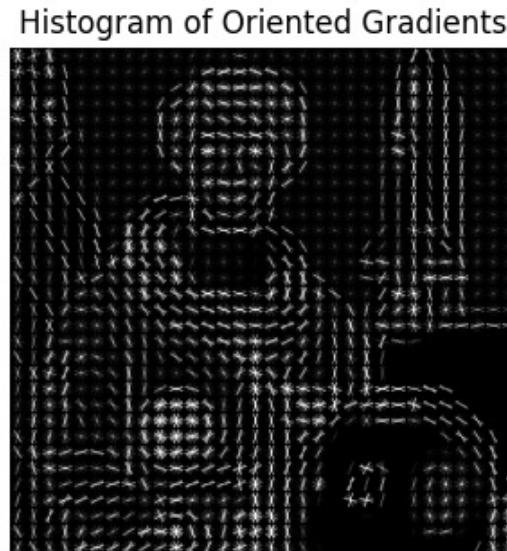
**Flattening on raw image** does not seem good enough since the vector does not represent the image well.



We need  
to use a  
better  
representa-  
tion

# Traffic Sign Classification

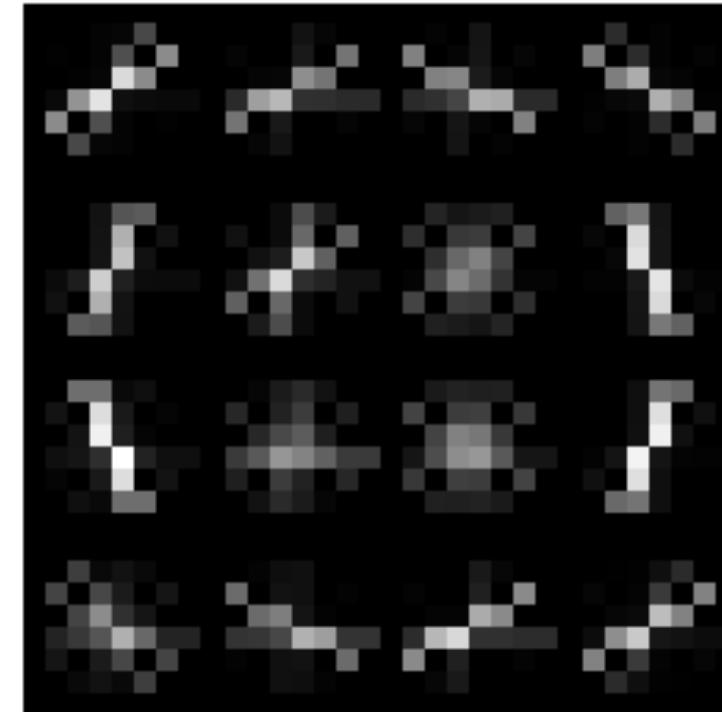
## ❖ HoG



**Histogram of Oriented Gradients (HoG):** A feature descriptor used primarily in object detection. It was introduced in a 2005 paper by Dalal and Triggs, where it was applied to pedestrian detection and achieved significant improvements over then-existing algorithms.

# Traffic Sign Classification

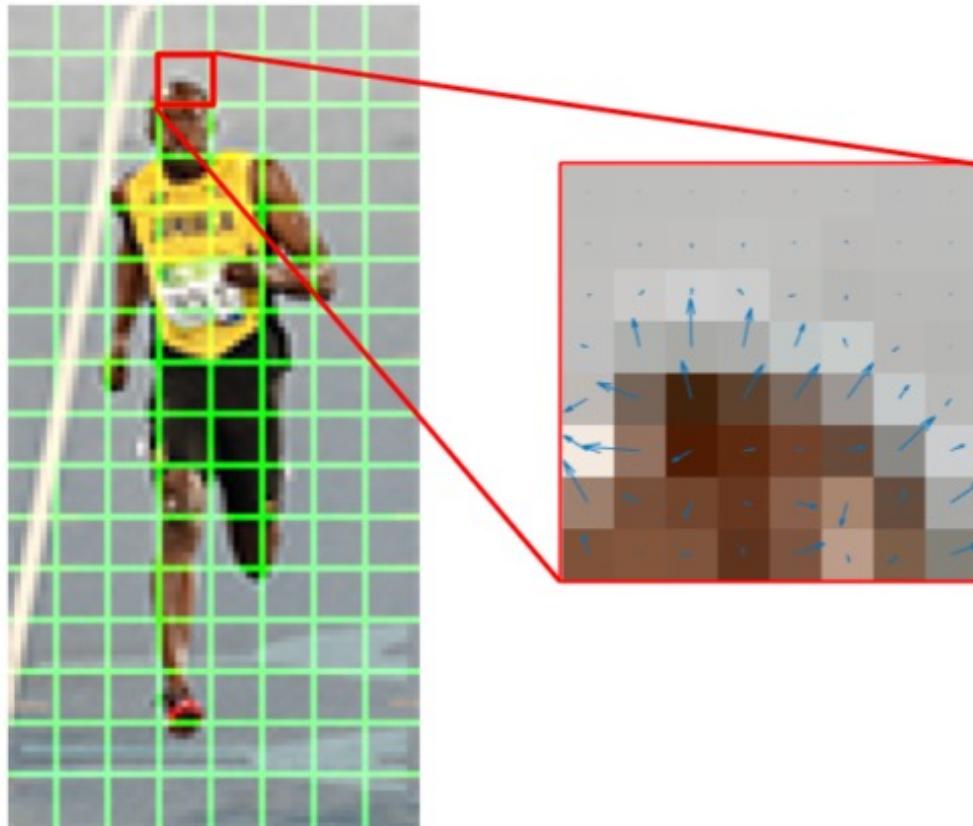
## ❖ HoG



**Objective:** Convert raw image to HoG representation, then do flattening on HoG.

# Traffic Sign Classification

## ❖ HoG



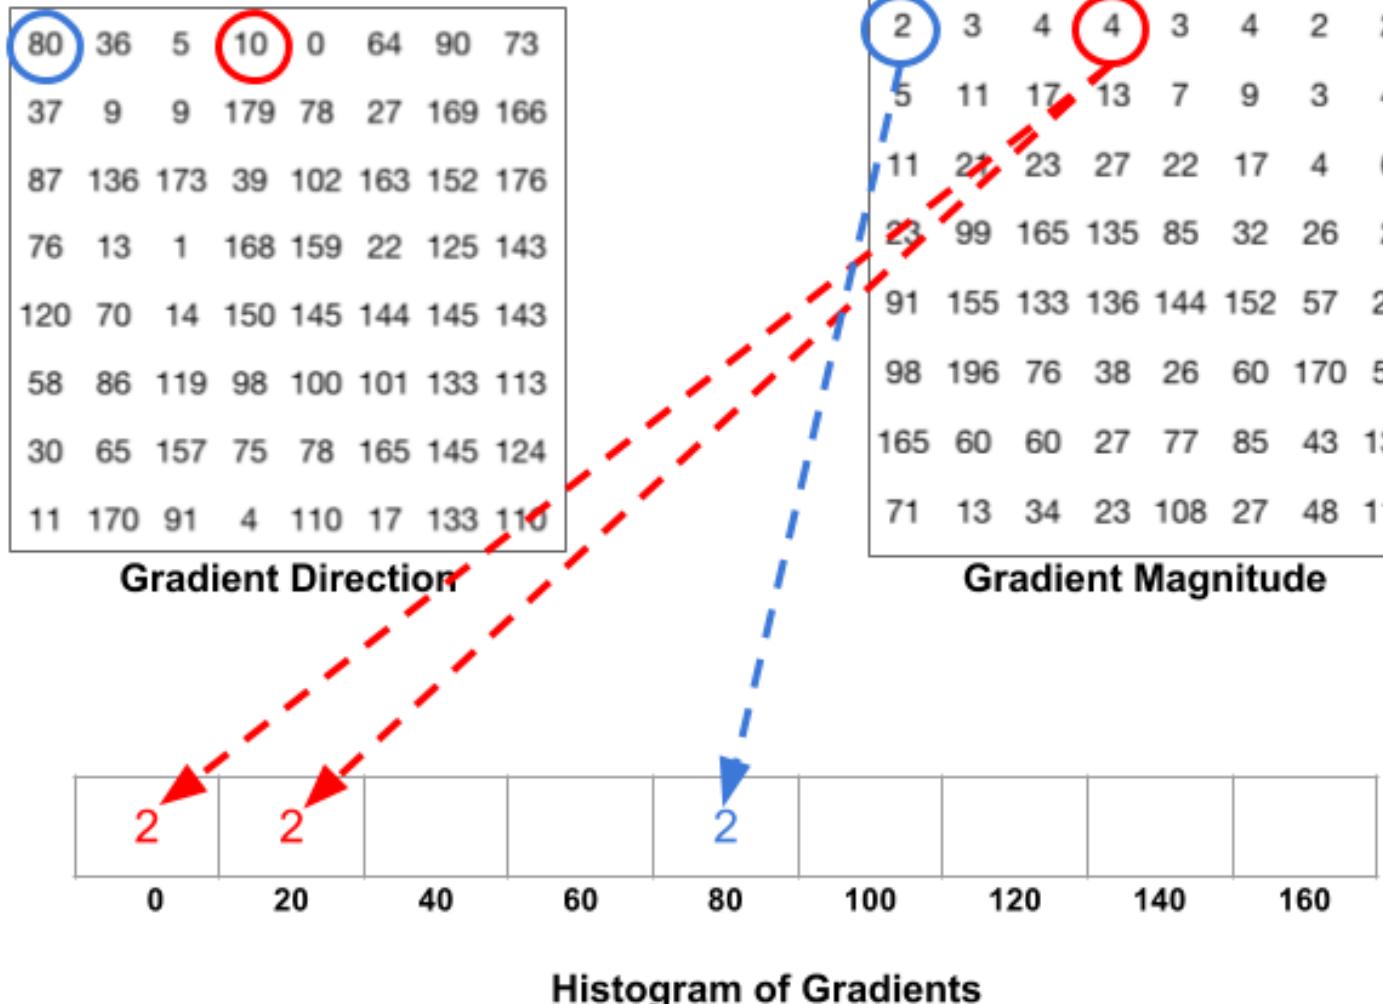
1. Split the images into cells. Each cell contains 8x8 pixels.
2. For each cell, compute the Gradient Direction and Gradient Magnitude.

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} * I \quad G = \sqrt{G_x^2 + G_y^2}$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} * I \quad \theta = \text{arctan}\left(\frac{G_y}{G_x}\right)$$

# Traffic Sign Classification

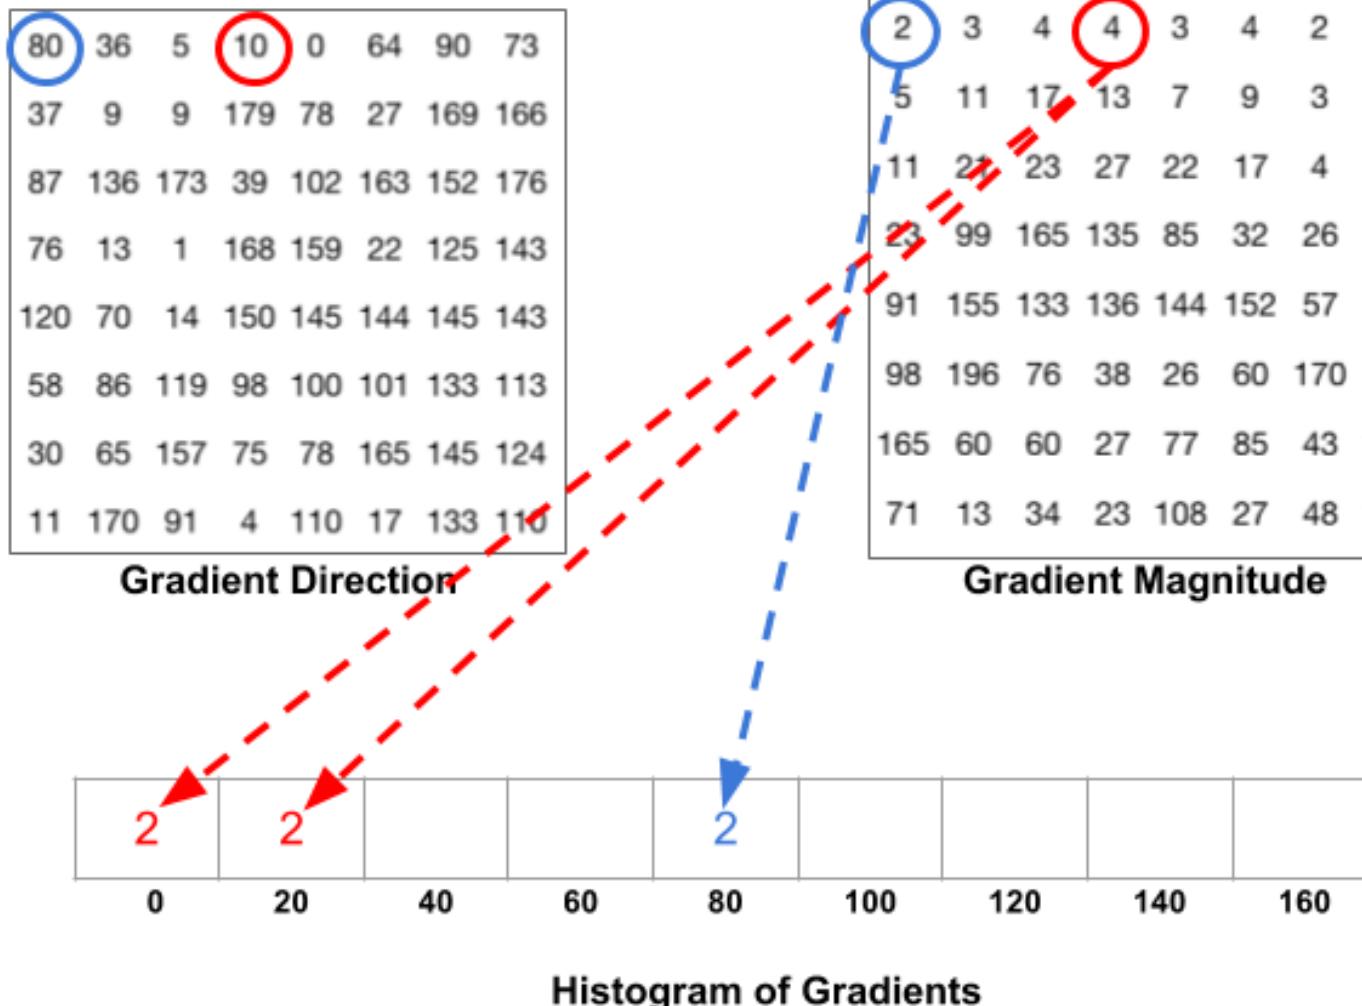
## ❖ HoG: Compute Histogram of Gradients



We create a vector of 9 elements (9 bins), representing the value range of Gradient Direction [0, 180]. If direction is in a bin, assign magnitude to the bin.

# Traffic Sign Classification

## ❖ HoG: Compute Histogram of Gradients



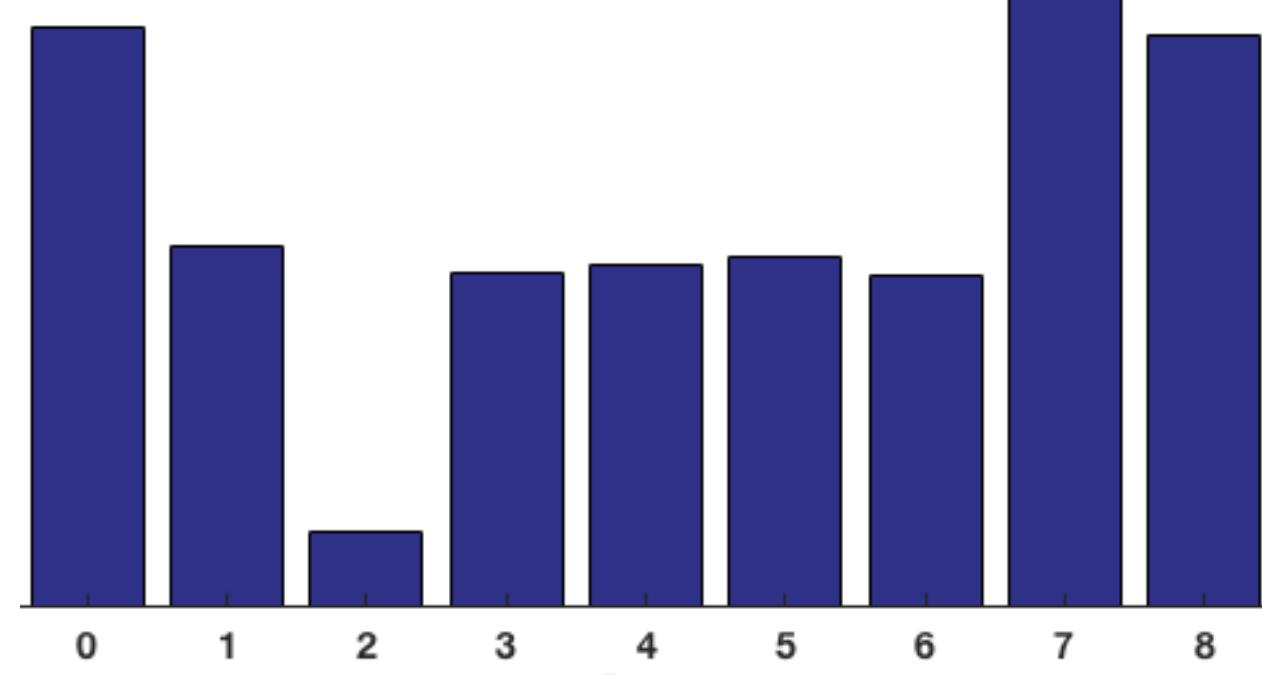
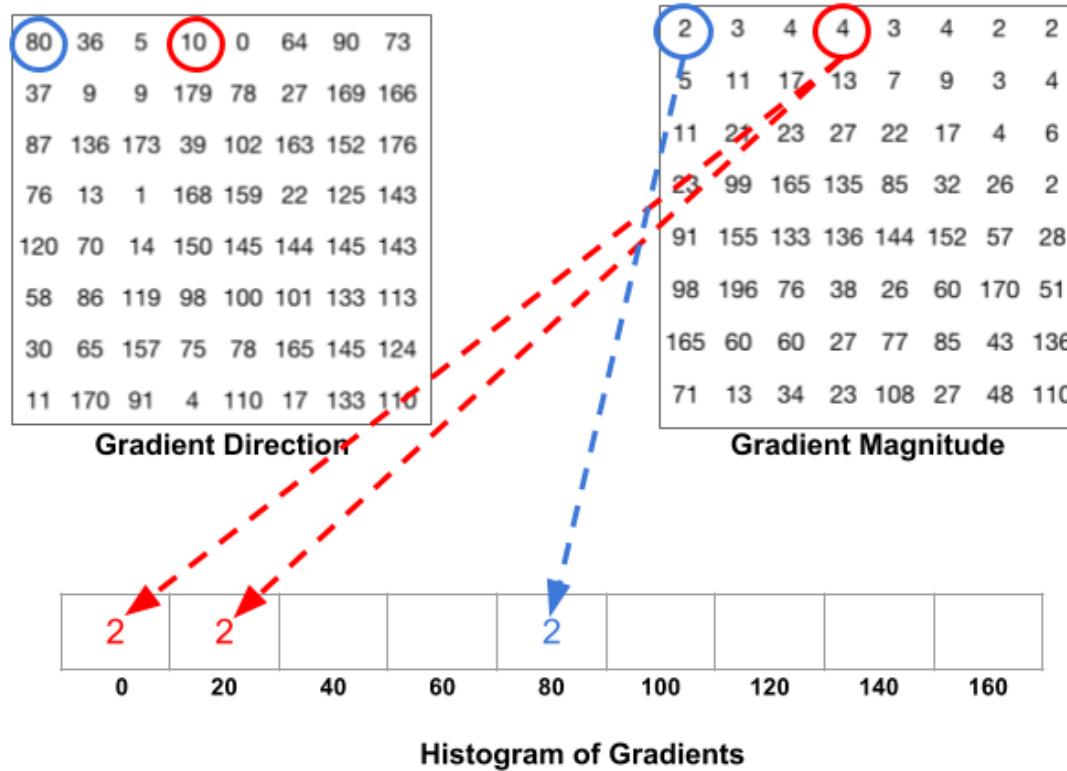
If not, assign magnitude to two adjacent bins  $x \in [x_0, x_1]$  using the following formulas:

$$x_{l-1} = \frac{(x_1 - x)}{x_1 - x_0} * magnitude$$

$$x_l = \frac{(x - x_0)}{x_1 - x_0} * magnitude$$

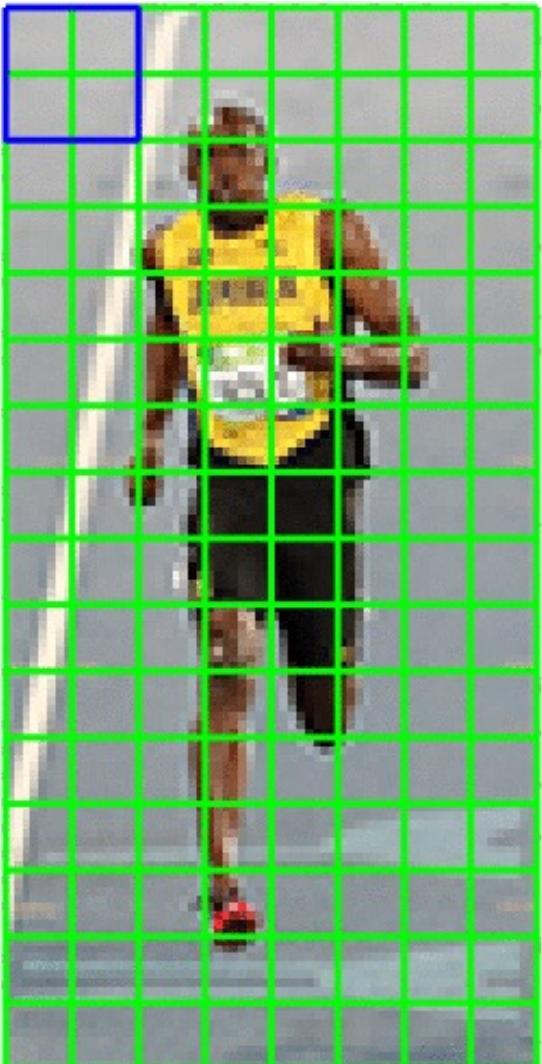
# Traffic Sign Classification

## ❖ HoG: Compute Histogram of Gradients

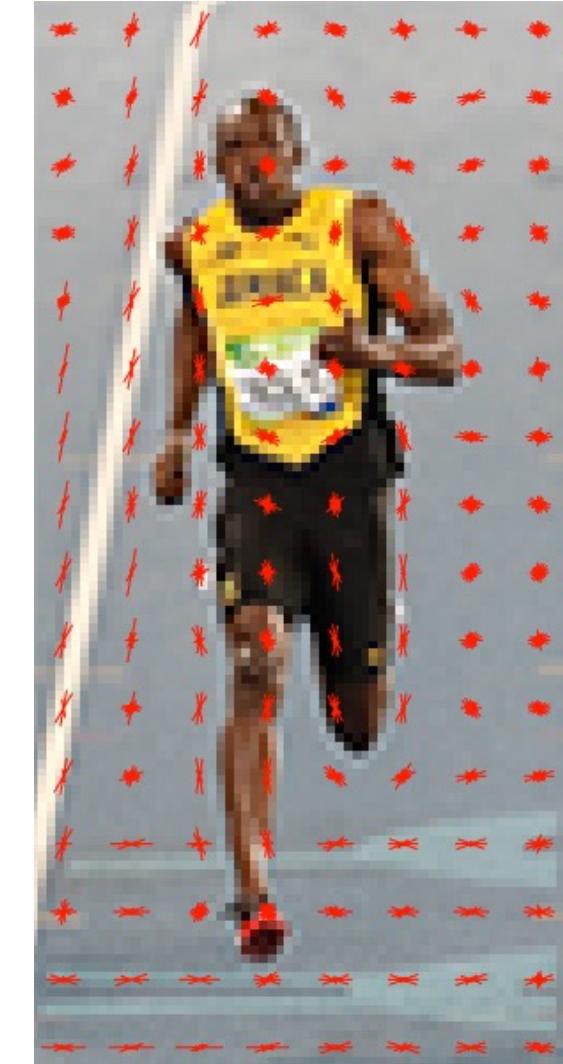
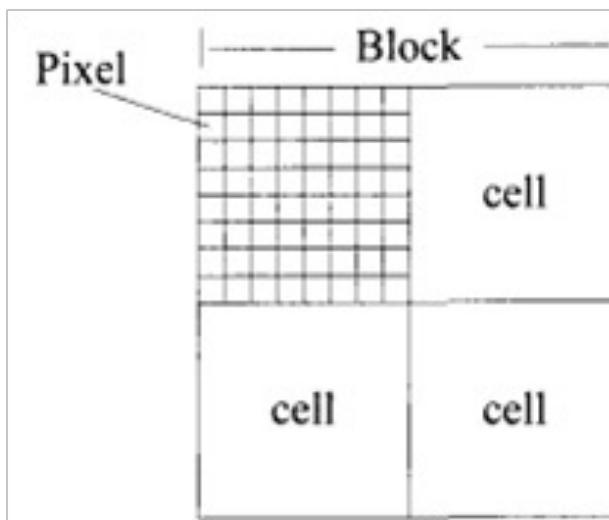


# Traffic Sign Classification

## ❖ HoG: Normalize block of cells



Do normalization on each 2x2 cells block =  
normalization on 1x36 vector.

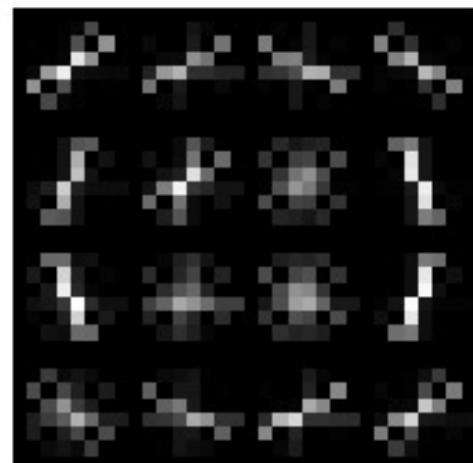


# Traffic Sign Classification

## ❖ Step 2: Preprocess image



Convert all raw images to HOG representation



```
1 def preprocess_img(img):
2     if len(img.shape) > 2:
3         img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
4     img = img.astype(np.float32)
5
6     resized_img = resize(
7         img,
8         output_shape=(32, 32),
9         anti_aliasing=True
10    )
11
12     hog_feature = feature.hog(
13         resized_img,
14         orientations=9,
15         pixels_per_cell=(8, 8),
16         cells_per_block=(2, 2),
17         transform_sqrt=True,
18         block_norm="L2",
19         feature_vector=True
20    )
21
22     return hog_feature
```

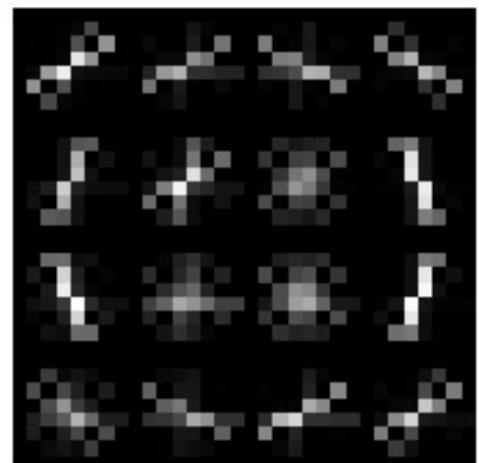
We resize all images to the same size to ensure HOG vectors have the same shape.

# Traffic Sign Classification

## ❖ Step 2: Preprocess image



Convert all raw  
images to HOG  
representation



```
1 img_features_lst = []
2 for img in img_lst:
3     hog_feature = preprocess_img(img)
4     img_features_lst.append(hog_feature)
5
6 img_features = np.array(img_features_lst)
7 print('X shape:')
8 print(img_features.shape)
```

X shape:  
(1074, 324)

# Traffic Sign Classification

## ❖ Step 3: Encode label

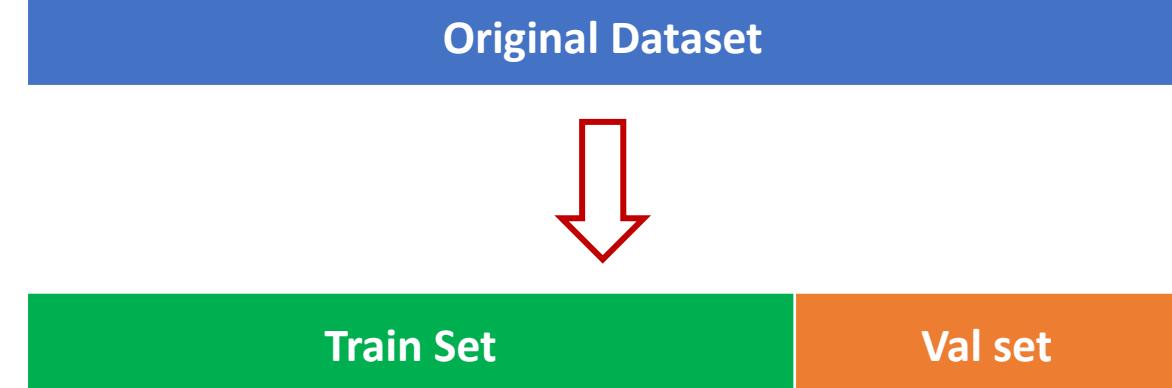
Label Encoder	
'crosswalk'	0
'speedlimit'	1
'stop'	2

```
1 label_encoder = LabelEncoder()  
2 encoded_labels = label_encoder.fit_transform(label_lst)  
  
1 label_encoder.classes_  
array(['crosswalk', 'speedlimit', 'stop'], dtype='<U10')  
  
1 encoded_labels  
array([1, 1, 2, ..., 0, 1, 0])
```

# Traffic Sign Classification

## ❖ Step 4: Create train, val set

```
1 random_state = 0
2 test_size = 0.3
3 is_shuffle = True
4
5 X_train, X_val, y_train, y_val = train_test_split(
6     img_features, encoded_labels,
7     test_size=test_size,
8     random_state=random_state,
9     shuffle=is_shuffle
10 )
```



# Traffic Sign Classification

## ❖ Step 5: Normalization

Using `sklearn.preprocessing.StandardScaler()`  
to scale all values in dataset.

```
1 scaler = StandardScaler()  
2 X_train = scaler.fit_transform(X_train)  
3 X_val = scaler.transform(X_val)
```

$$z = \frac{x_i - \mu}{\sigma}$$

```
1 X_train
```

```
array([[ 0.5161459 ,  1.4854372 , -0.5399236 , ...,  0.03029941,  
       -0.29923776, -0.12044104],  
       [ 0.86671036,  0.45618102, -0.6950155 , ..., -0.2180403 ,  
       -0.1757149 ,  0.2960105 ],  
       [-0.78160155, -0.0928254 ,  0.02434517, ..., -0.6275371 ,  
       -0.52018636, -0.6298612 ],  
       ...,  
       [ 0.24302678, -0.7746629 , -1.5431894 , ..., -0.6275371 ,  
       -0.4982059 ,  0.39656585],  
       [-0.8555847 , -0.21795394, -0.52630377, ..., -0.6275371 ,  
       -0.4973315 , -0.5843945 ],  
       [-0.74369645, -0.38316953,  0.10891949, ..., -0.503356 ,  
       -0.5255429 , -0.49774203]], dtype=float32)
```

# Traffic Sign Classification

## ❖ Step 6, 7: Train and evaluate SVM model

```
1 clf = SVC(  
2     kernel='rbf',  
3     random_state=random_state,  
4     probability=True,  
5     C=0.5  
6 )  
7 clf.fit(X_train, y_train)
```

SVC

```
SVC(C=0.5, probability=True, random_state=0)
```

```
1 y_pred = clf.predict(X_val)  
2 score = accuracy_score(y_pred, y_val)  
3  
4 print('Evaluation results on val set')  
5 print('Accuracy: ', score)
```

Evaluation results on val set  
Accuracy: 0.9659442724458205

# Traffic Sign Classification

## ❖ Step 6, 7: Train and evaluate SVM model



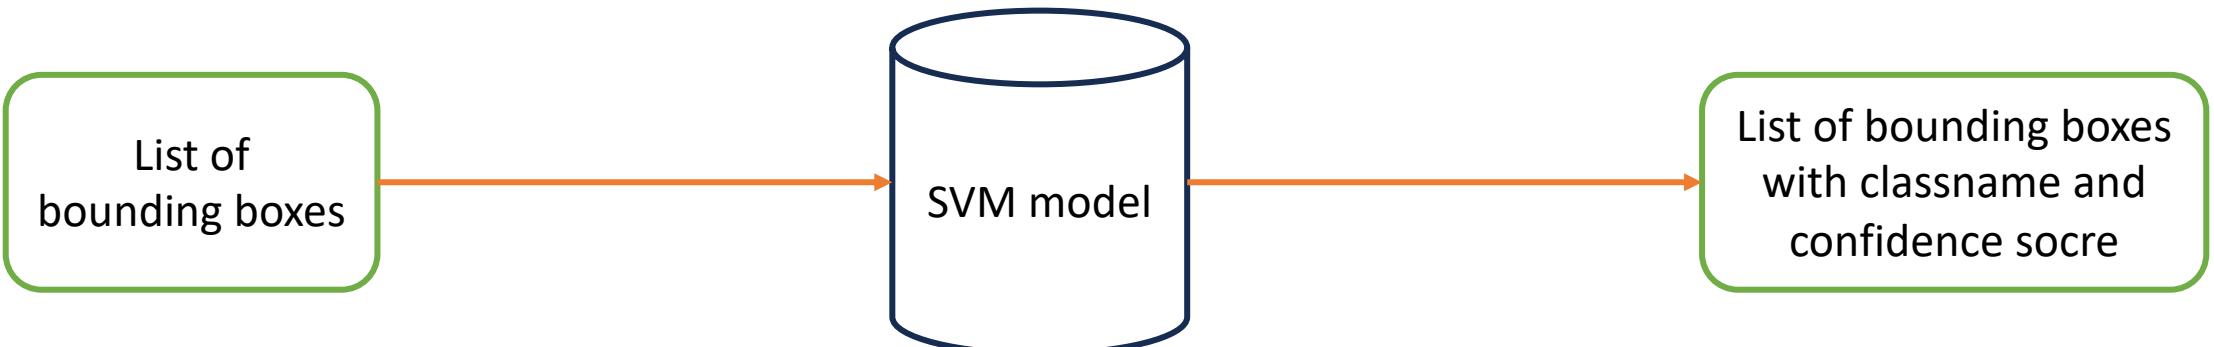
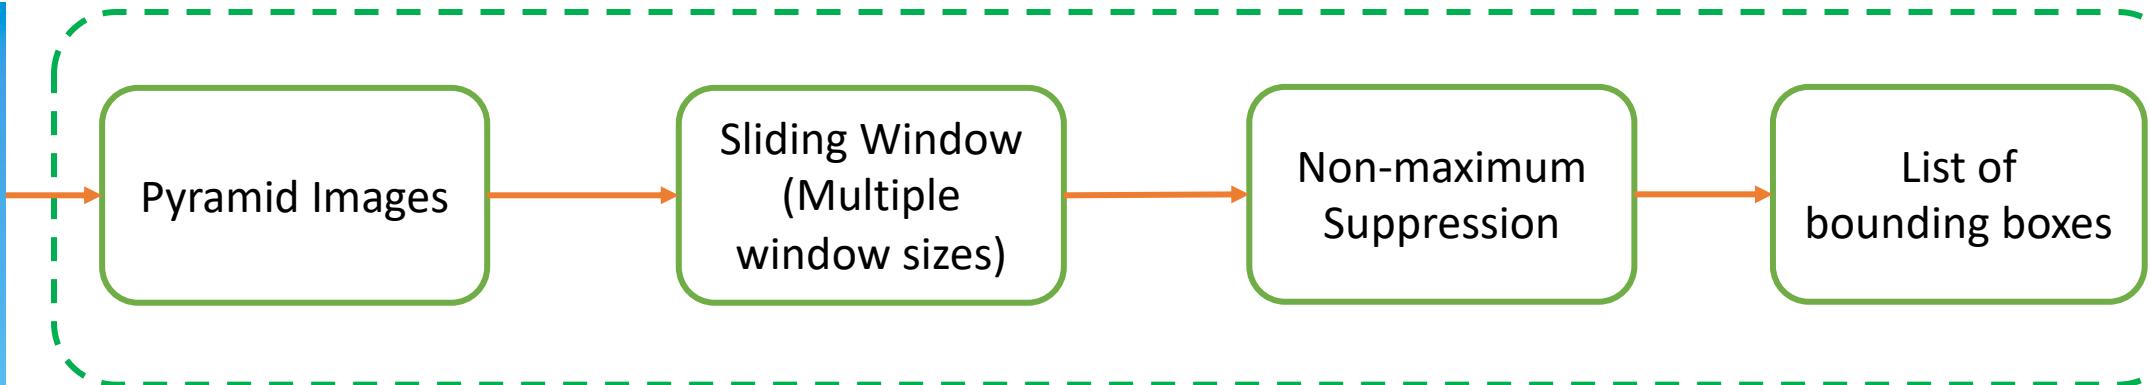
```
1 input_img = img_lst[1]
2 normalized_img = preprocess_img(input_img)
3 y_pred = clf.predict([normalized_img])[0]
4 print(f'Normal prediction: {y_pred}')
5 y_pred_prob = clf.predict_proba([normalized_img])
6 prediction = np.argmax(y_pred_prob)
7 y_pred_prob = [f'{p:.10f}' for p in y_pred_prob[0]]
8 print(f'Probability of each class: {y_pred_prob}')
9 print(f'Class with highest probability: {prediction}')
```

Normal prediction: 1  
Probability of each class: ['0.0000000040', '0.9999306581', '0.0000693379']  
Class with highest probability: 1

# Traffic Sign Localization

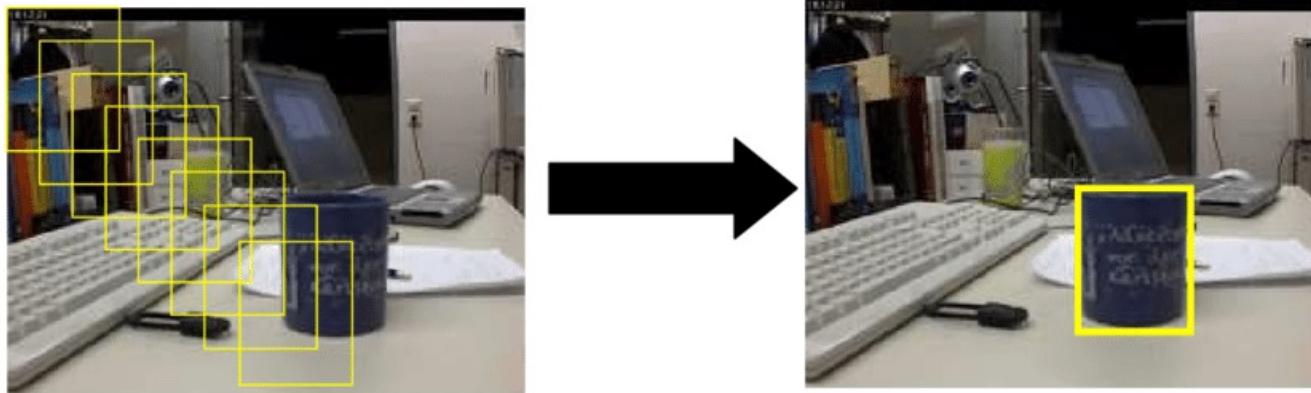
## ❖ Introduction

### Input Image



# Traffic Sign Localization

## ❖ Sliding Window

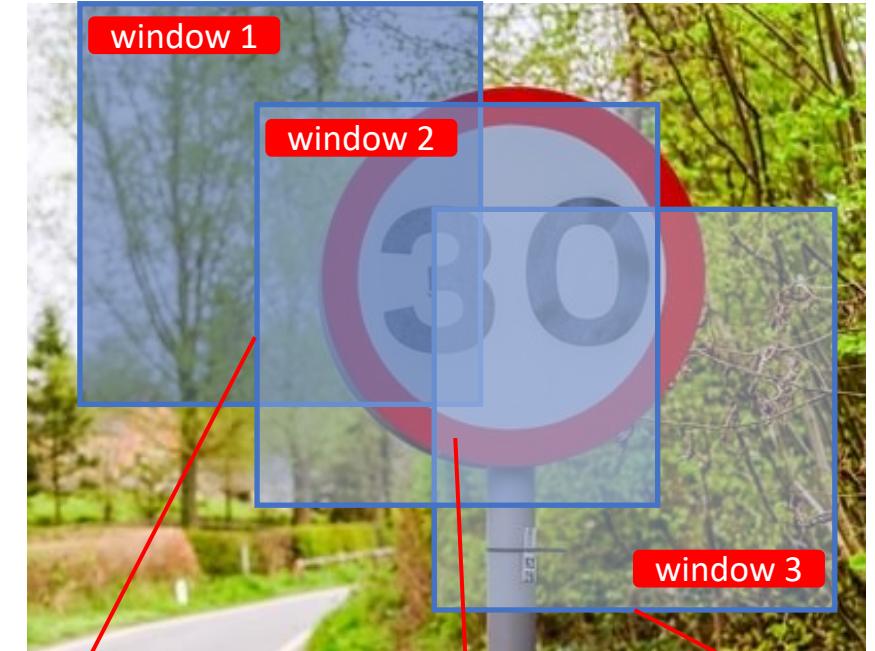


**Sliding Window:** A technique used in computer vision and signal processing where a fixed-size “window” (or a sub-rectangle) is moved across an image (or a signal) to analyze a local region at a time.

**Example:** Using sliding window to find a cup within an image.

# Traffic Sign Localization

## ❖ Sliding Window

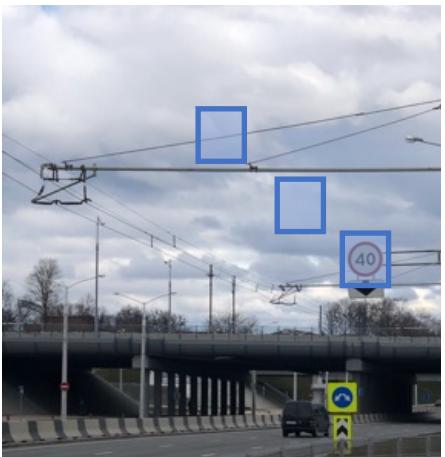


# Traffic Sign Localization

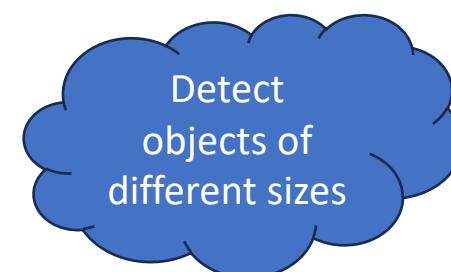
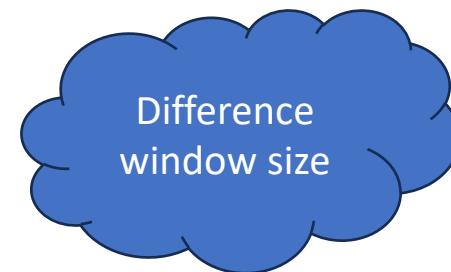
## ❖ Sliding Window



Window size 1



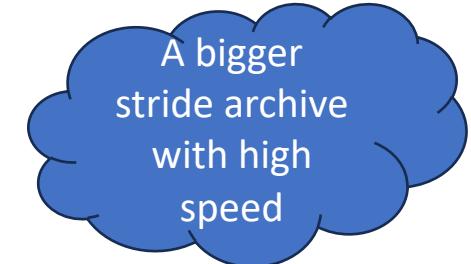
Window size 2



Stride = 10



Stride = 20



# Traffic Sign Localization

## ❖ Sliding Window

```
1 def sliding_window(img, window_sizes, stride, scale_factor):
2     img_height, img_width = img.shape[:2]
3     windows = []
4     for window_size in window_sizes:
5         window_width, window_height = window_size
6         for ymin in range(0, img_height - window_height + 1, stride):
7             for xmin in range(0, img_width - window_width + 1, stride):
8                 xmax = xmin + window_width
9                 ymax = ymin + window_height
10
11             windows.append([xmin, ymin, xmax, ymax])
12
13 return windows
```

# Traffic Sign Localization

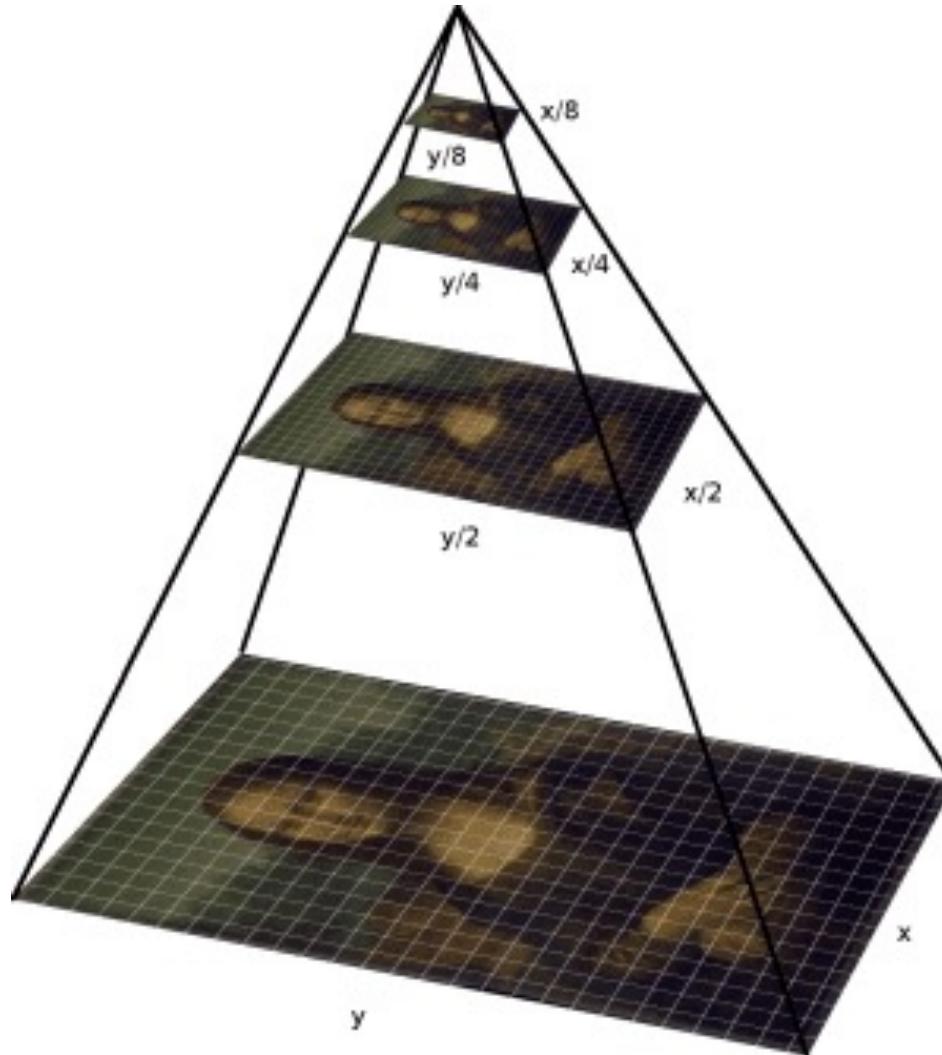
## ❖ Small objects problem



In many situations, **the objects of interest are often presented very small in the image**. Directly apply a fixed-size window might not be effectively handle this.

# Traffic Sign Localization

## ❖ Pyramid Images



**Pyramid Image:** A multi-scale representation of an input image where the image is processed at various scales (resolutions) to detect objects at different sizes.

# Traffic Sign Localization

## ❖ Pyramid Images

Scaled at: 1.0



Scaled at: 0.5120000000000001



Scaled at: 0.32768000000000014



# Traffic Sign Localization

## ❖ Pyramid Images

Scaled at: 1.0



Scaled at: 0.5120000000000001



Apply sliding window on each scaled image to find the objects of interest.

# Traffic Sign Localization

## ❖ Pyramid Images

```
1 def pyramid(img, scale=0.8, min_size=(30, 30)):
2     acc_scale = 1.0
3     pyramid_imgs = [(img, acc_scale)]
4
5     i = 0
6     while True:
7         acc_scale = acc_scale * scale
8         h = int(img.shape[0] * acc_scale)
9         w = int(img.shape[1] * acc_scale)
10        if h < min_size[1] or w < min_size[0]:
11            break
12        img = cv2.resize(img, (w, h))
13        pyramid_imgs.append((img, acc_scale * (scale ** i)))
14        i += 1
15
16    return pyramid_imgs
```

# Traffic Sign Localization

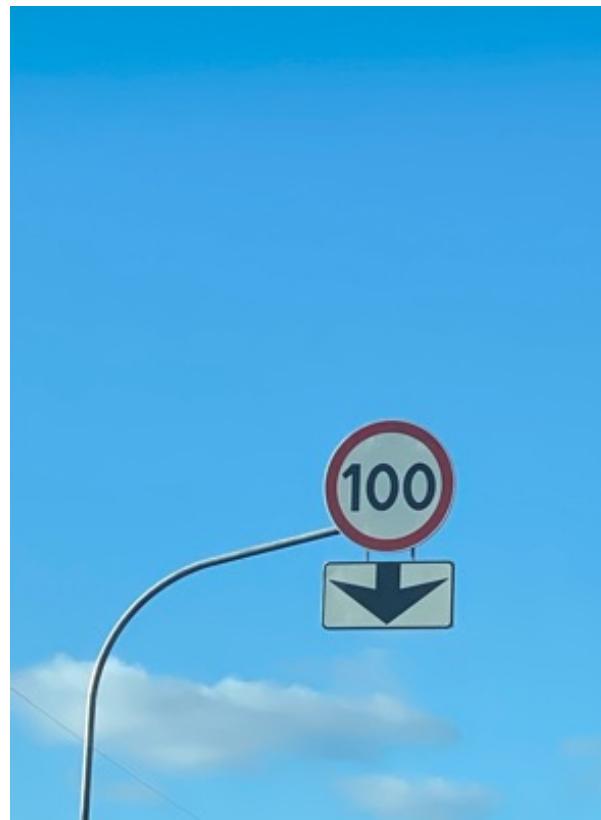
## ❖ Apply Sliding Window on Pyramid Images



# Traffic Sign Localization

## ❖ Visualize bounding box

With this list of bounding box, we need to put them on the image for a better understanding of the final results.



# Traffic Sign Localization

## ❖ Visualize bounding box

```
1 def visualize_bbox(img, bboxes, label_encoder):
2     img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
3
4     for box in bboxes:
5         xmin, ymin, xmax, ymax, predict_id, conf_score = box
6
7         cv2.rectangle(img, (xmin, ymin), (xmax, ymax), (0, 255, 0), 2)
8
9         classname = label_encoder.inverse_transform([predict_id])[0]
10        label = f'{classname} {conf_score:.2f}'
11
12        (w, h), _ = cv2.getTextSize(label, cv2.FONT_HERSHEY_SIMPLEX, 0.6, 1)
13
14        cv2.rectangle(img, (xmin, ymin - 20), (xmin + w, ymin), (0, 255, 0), -1)
15
16        cv2.putText(img, label, (xmin, ymin - 5), cv2.FONT_HERSHEY_SIMPLEX, 0.6, (0, 0, 0), 1)
17
18    plt.imshow(img)
19    plt.axis('off')
20    plt.show()
```

# Traffic Sign Localization

# Prediction

```
for pyramid_img_info in pyramid_imgs:
    pyramid_img, scale_factor = pyramid_img_info
    window_lst = sliding_window(
        pyramid_img,
        window_sizes=window_sizes,
        stride=stride,
        scale_factor=scale_factor
    )
    for window in window_lst:
        xmin, ymin, xmax, ymax = window
        object_img = pyramid_img[ymin:ymax, xmin:xmax]
        preprocessed_img = preprocess_img(object_img)
        normalized_img = scaler.transform([preprocessed_img])[0]
        decision = clf.predict_proba([normalized_img])[0]
        predict_id = np.argmax(decision)
        conf_score = decision[predict_id]
        xmin = int(xmin / scale_factor)
        ymin = int(ymin / scale_factor)
        xmax = int(xmax / scale_factor)
        ymax = int(ymax / scale_factor)
        bboxes.append([
            [xmin, ymin, xmax, ymax, predict_id, conf_score]
        ])
```

# Traffic Sign Localization

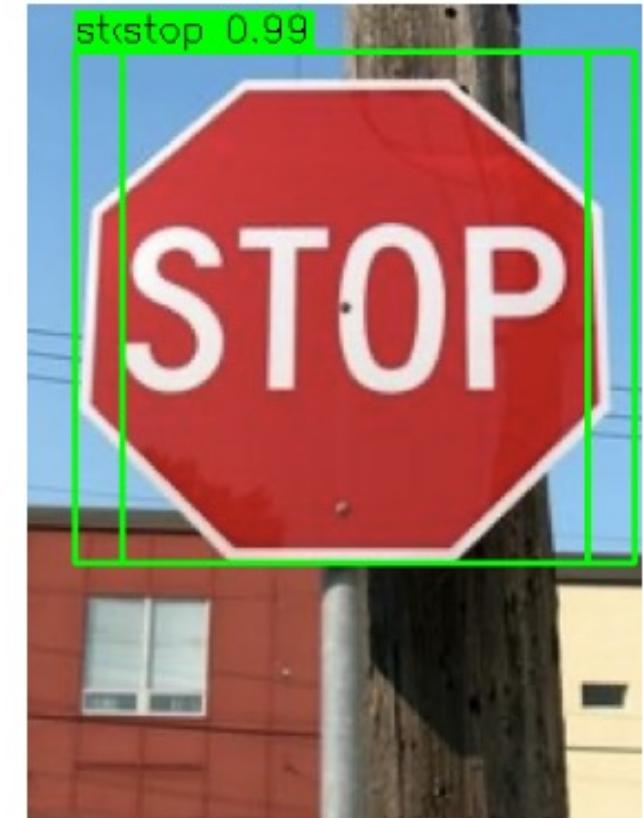
## Prediction

We want to select certain bounding boxes (e.g: bounding boxes which have a higher or equal than the minimum confidence score => Threshold the confidence score.

```
for window in window_lst:
    xmin, ymin, xmax, ymax = window
    object_img = pyramid_img[ymin:ymax, xmin:xmax]
    preprocessed_img = preprocess_img(object_img)
    normalized_img = scaler.transform([preprocessed_img])[0]
    decision = clf.predict_proba([normalized_img])[0]
    if np.all(decision < conf_threshold):
        continue
    else:
        predict_id = np.argmax(decision)
        conf_score = decision[predict_id]
        xmin = int(xmin / scale_factor)
        ymin = int(ymin / scale_factor)
        xmax = int(xmax / scale_factor)
        ymax = int(ymax / scale_factor)
        bboxes.append(
            [xmin, ymin, xmax, ymax, predict_id, conf_score]
        )
```

# Traffic Sign Localization

## ❖ Prediction



Results with a confidence score threshold = 0.95

# Traffic Sign Localization

## ❖ Prediction



# Traffic Sign Localization

## ❖ Overlapped bounding boxes problem



Usually, the raw outcome of object detection may have a lot of bounding boxes overlapping each other, with almost the same confidence score. We need to compress into the best one.

# Traffic Sign Localization

## ❖ Non-maximum Suppression

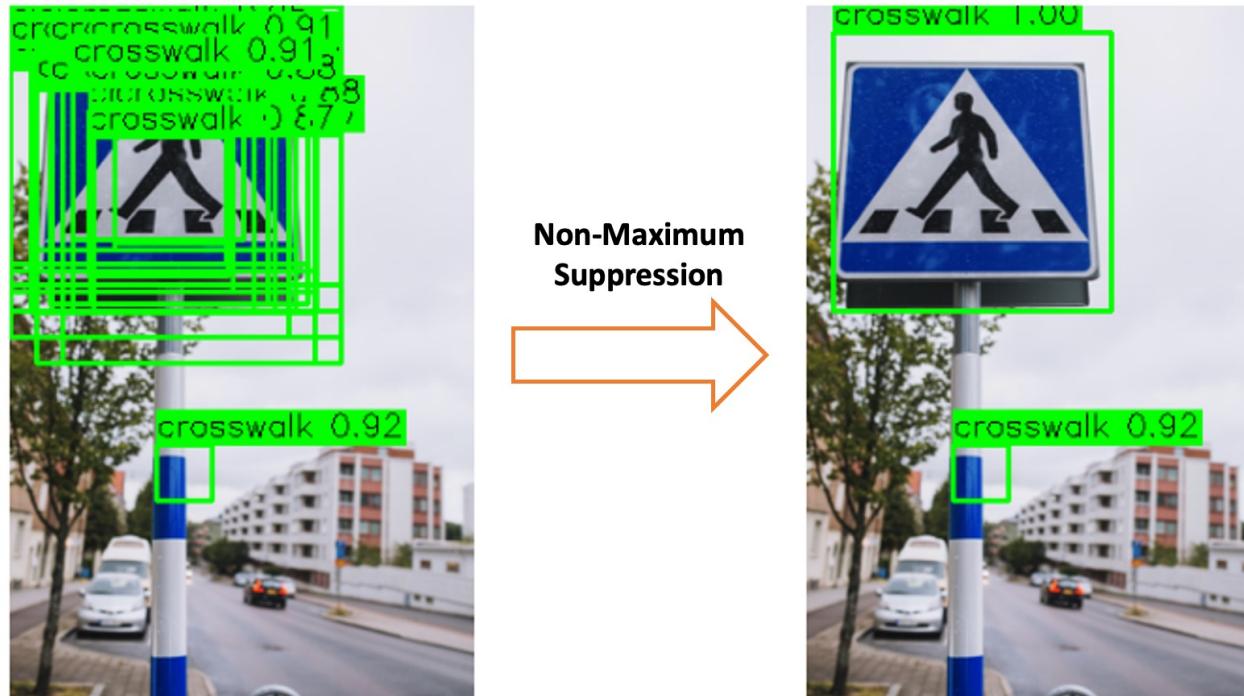


Non-Maximum  
Suppression



# Traffic Sign Localization

## ❖ Non-maximum Suppression



**Non-maximum suppression (NMS):** A technique used in object detection tasks to prune multiple bounding boxes that are overlapping and refer to the same object.

# Traffic Sign Localization

## ❖ Non-maximum suppression



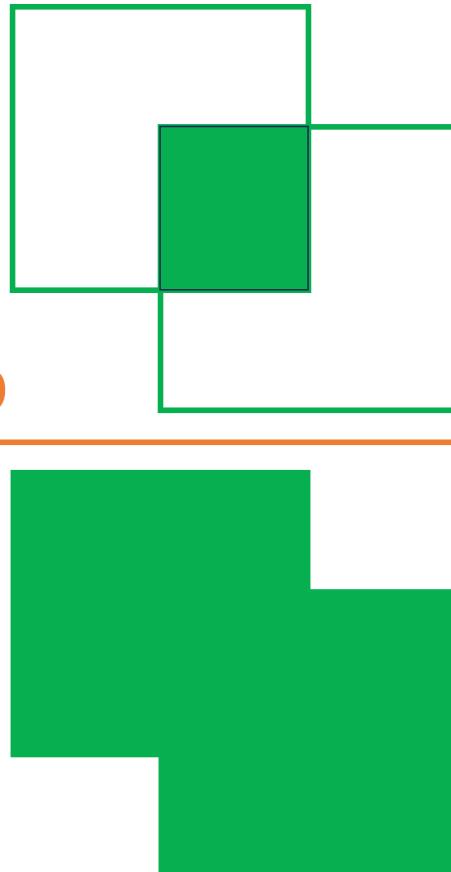
### Non-maximum suppression:

1. Select the bounding box with highest confidence score (A).
2. Compute IoU between (A) with other images.
3. Only retain images that have IoU less than a threshold ( $iou\_threshold$ ).  
Otherwise, remove all.

# Traffic Sign Localization

## ❖ Intersection over Union

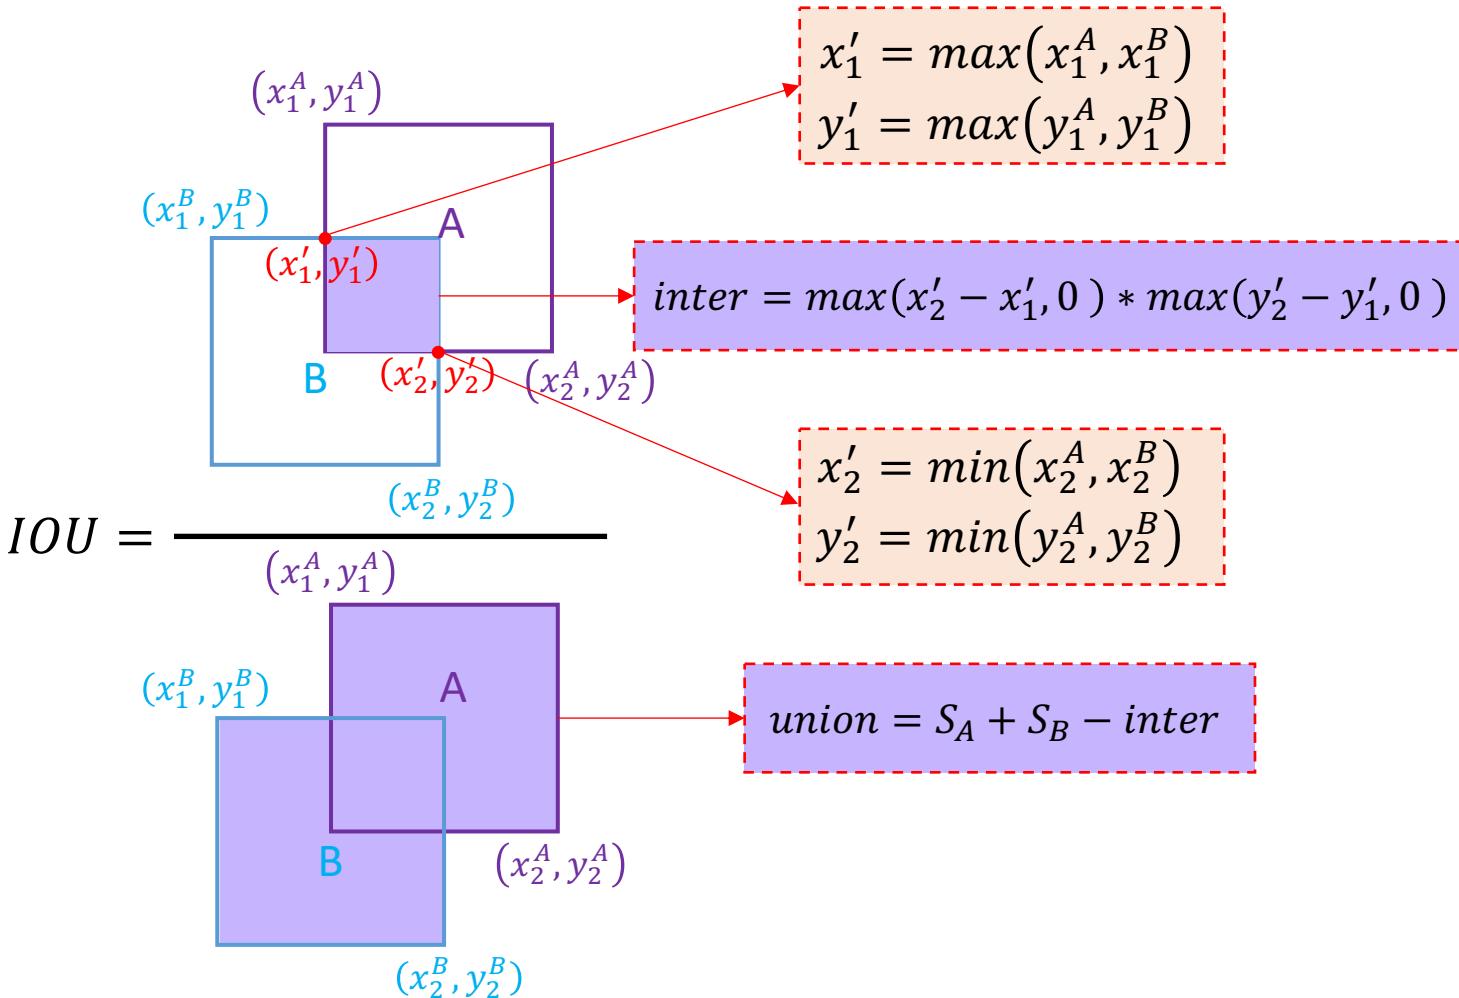
$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$



**Intersection over Union (IoU):** A metric used to measure the overlap between two bounding boxes.

# Traffic Sign Localization

## ❖ Intersection over Union



```

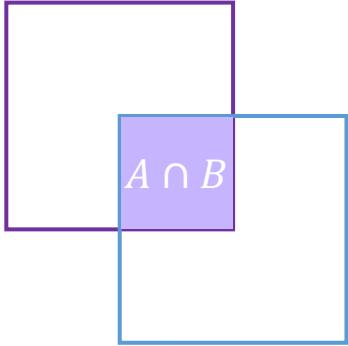
1  def iou_bbox(box1, box2):
2      """
3          - Input:
4              - box1 (x1, y1, x2, y2)
5              - box2 (x1, y1, x2, y2)
6          - Output:
7              - iou of box1 and box2
8
9      """
10     b1_x1, b1_y1, b1_x2, b1_y2 = box1[0], box1[1], box1[2], box1[3]
11     b2_x1, b2_y1, b2_x2, b2_y2 = box2[0], box2[1], box2[2], box2[3]
12
13     # Step 1: Compute inter area
14     x1 = max(b1_x1, b2_x1)
15     y1 = max(b1_y1, b2_y1)
16     x2 = min(b1_x2, b2_x2)
17     y2 = min(b1_y2, b2_y2)
18
19     inter = max((x2 - x1), 0) * max((y2 - y1), 0)
20
21     # Step 2: Compute union area
22     box1Area = abs((b1_x2 - b1_x1) * (b1_y2 - b1_y1))
23     box2Area = abs((b2_x2 - b2_x1) * (b2_y2 - b2_y1))
24
25     union = float(box1Area + box2Area - inter)
26
27     iou = inter / union
28     return iou

```

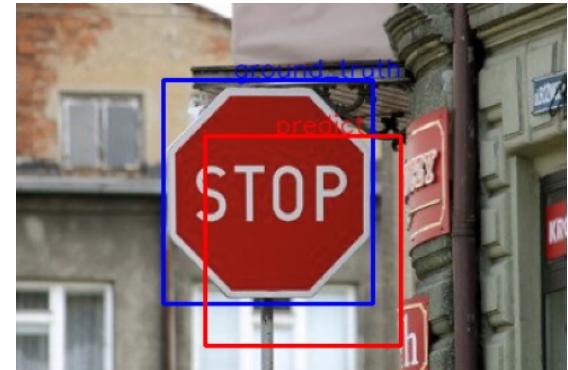
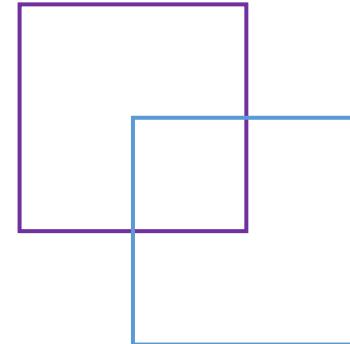
# Traffic Sign Localization

## ❖ Intersection over Union

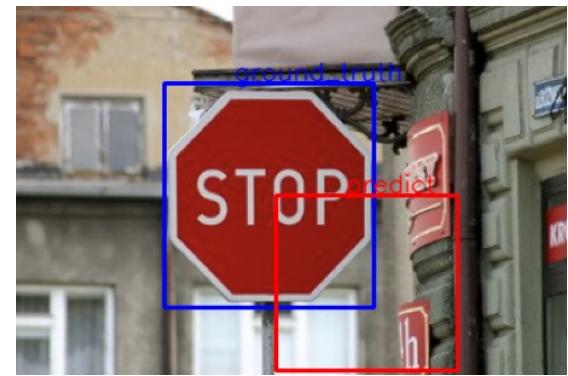
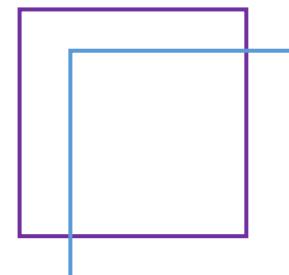
$$IoU = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$



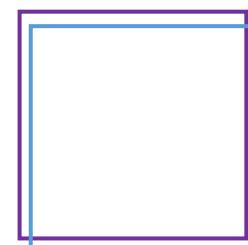
Poor



Good



Excellent



# Traffic Sign Localization

## ❖ Non-maximum suppression

### Non-maximum suppression:

1. Select the bounding box with highest confidence score (A).
2. Compute IoU between (A) with other images.
3. Only retain images that have IoU less than a threshold (iou\_threshold). Otherwise, remove all.

```
1 def nms(bboxes, iou_threshold):
2     if not bboxes:
3         return []
4     scores = np.array([bbox[5] for bbox in bboxes])
5     sorted_indices = np.argsort(scores)[::-1]
6
7     xmin = np.array([bbox[0] for bbox in bboxes])
8     ymin = np.array([bbox[1] for bbox in bboxes])
9     xmax = np.array([bbox[2] for bbox in bboxes])
10    ymax = np.array([bbox[3] for bbox in bboxes])
11
12    areas = (xmax - xmin + 1) * (ymax - ymin + 1)
13
14    keep = []
15    while sorted_indices.size > 0:
16        i = sorted_indices[0]
17        keep.append(i)
18
19        iou = compute_iou(
20            [xmin[i], ymin[i], xmax[i], ymax[i]],
21            np.array(
22                [
23                    xmin[sorted_indices[1:]],
24                    ymin[sorted_indices[1:]],
25                    xmax[sorted_indices[1:]],
26                    ymax[sorted_indices[1:]]
27                ].T,
28                areas[i],
29                areas[sorted_indices[1:]]
30            )
31
32        idx_to_keep = np.where(iou <= iou_threshold)[0]
33        sorted_indices = sorted_indices[idx_to_keep + 1]
34
35    return [bboxes[i] for i in keep]
```

# Traffic Sign Localization

## ❖ Final Results



# Question

