

AI VIETNAM  
All-in-One Course  
(TA Session)

# Convolutional Neural Networks

## Exercise



AI VIET NAM  
[@aivietnam.edu.vn](http://aivietnam.edu.vn)

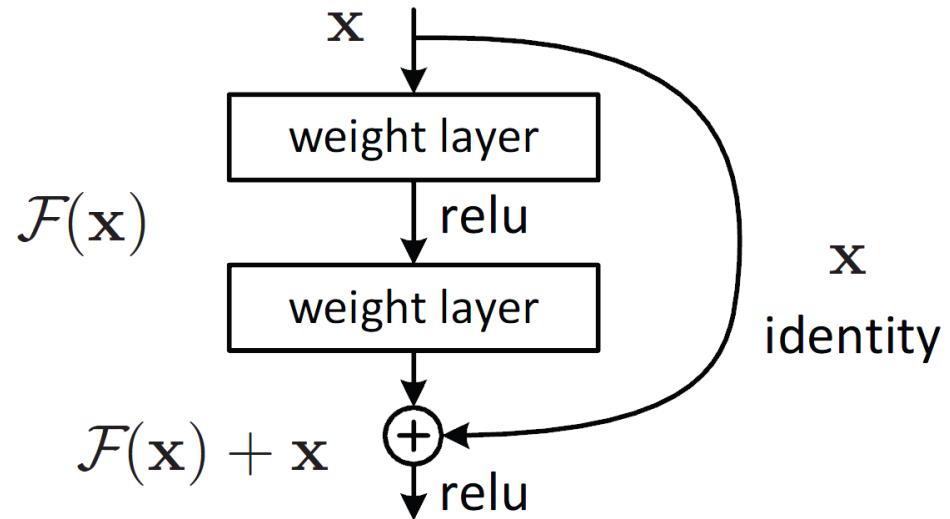
Dinh-Thang Duong – TA

# Outline

- ResNet
- DenseNet
- Question

# ResNet

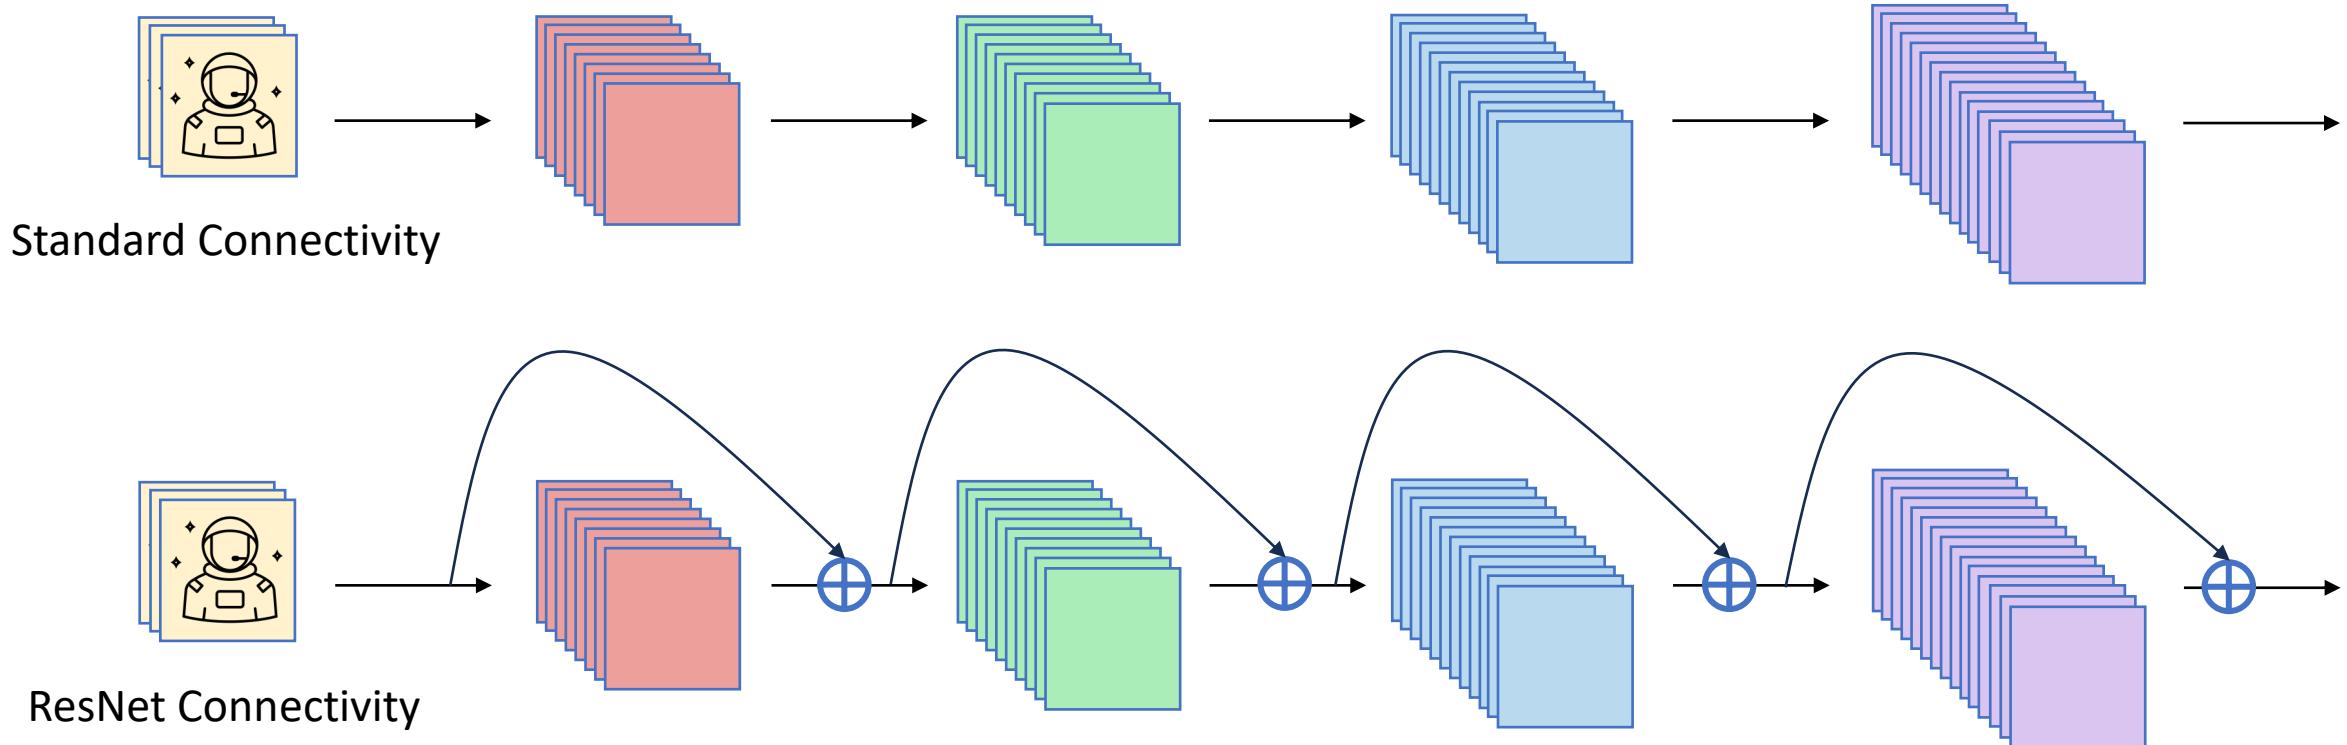
## ❖ ResNet Architecture



**Residual Networks (ResNet):** A type of CNNs first introduced in 2015. Its innovative approach involves using “residual blocks”, which allow the network to efficiently train extremely deep models by learning residual (difference) between the input and desired output.

# ResNet

## ❖ ResNet Architecture



In **ResNet**, the output of a layer is combined with the input to create a residual connection or shortcut connection.

# ResNet

## ❖ Coding Exercise

**Problem Statement:** Given the weather image dataset (download [here](#)). Train a model to determine which type of weather is in a image using ResNet.



# ResNet

## ❖ Step 1: Import libraries

```
1 import torch
2 import torch.nn as nn
3 import os
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7
8 from PIL import Image
9 from torch.utils.data import Dataset, DataLoader
10 from sklearn.model_selection import train_test_split
```



# ResNet

## ❖ Step 2: Read dataset

- weather-dataset
  - dataset
    - ▶ dew
    - ▶ fogsmog
    - ▶ frost
    - ▶ glaze
    - ▶ hail
    - ▶ lightning
    - ▶ rain
    - ▶ rainbow
    - ▶ rime
    - ▶ sandstorm
    - ▶ snow

```
1 root_dir = 'dataset/weather-dataset/dataset'  
2 classes = {  
3     label_idx: class_name \  
4         for label_idx, class_name in enumerate(  
5             sorted(os.listdir(root_dir))  
6         )  
7 }
```

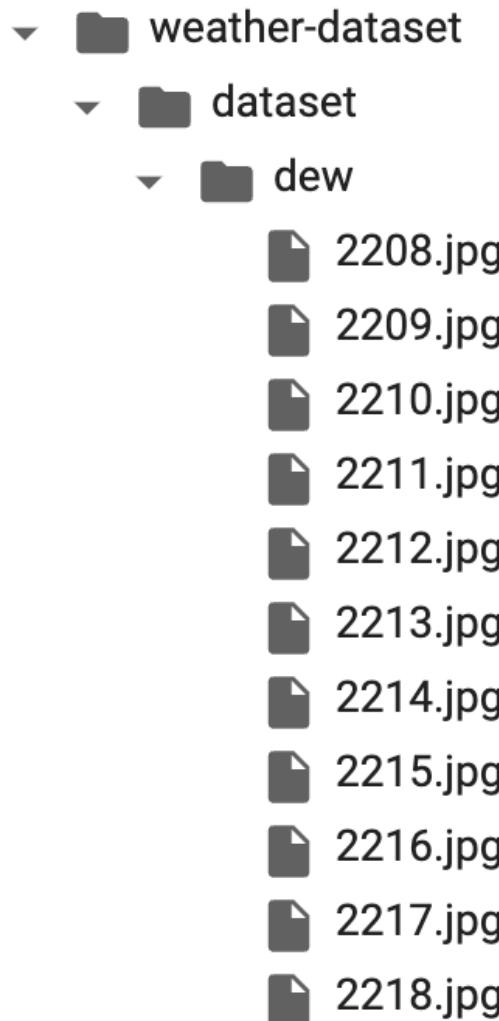
1 classes

```
{0: 'dew',  
 1: 'fogsmog',  
 2: 'frost',  
 3: 'glaze',  
 4: 'hail',  
 5: 'lightning',  
 6: 'rain',  
 7: 'rainbow',  
 8: 'rime',  
 9: 'sandstorm',  
10: 'snow'}
```

Each folder name is also a class name, thus we could get list of classnames by reading over them.

# ResNet

## ❖ Step 2: Read dataset



```
9 img_paths = []
10 labels = []
11 for label_idx, class_name in classes.items():
12     class_dir = os.path.join(root_dir, class_name)
13     for img_filename in os.listdir(class_dir):
14         img_path = os.path.join(class_dir, img_filename)
15         img_paths.append(img_path)
16         labels.append(label_idx)
```

1 img\_paths

```
['weather-dataset/dataset/dew/2405.jpg',
 'weather-dataset/dataset/dew/2340.jpg',
 'weather-dataset/dataset/dew/2772.jpg',
 'weather-dataset/dataset/dew/2794.jpg',
 'weather-dataset/dataset/dew/2827.jpg',
 'weather-dataset/dataset/dew/2235.jpg',
 'weather-dataset/dataset/dew/2396.jpg',
 'weather-dataset/dataset/dew/2349.jpg',
 'weather-dataset/dataset/dew/2879.jpg',
 'weather-dataset/dataset/dew/2734.jpg',
```

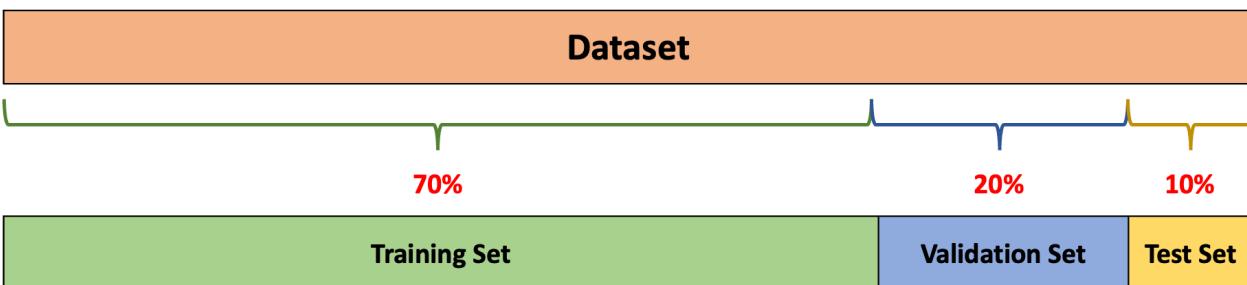
1 labels

```
[0,
 0,
 0,
 0,
 0,
 0,
 0,
 0,
 0,
 0,
```

We get all image paths  
as well as their  
corresponding label.

# ResNet

## ❖ Step 3: Train, val, test split



```
1 seed = 0
2 val_size = 0.2
3 test_size = 0.125
4 is_shuffle = True
5
6 X_train, X_val, y_train, y_val = train_test_split(
7     img_paths, labels,
8     test_size=val_size,
9     random_state=seed,
10    shuffle=is_shuffle
11 )
12
13 X_train, X_test, y_train, y_test = train_test_split(
14     X_train, y_train,
15     test_size=val_size,
16     random_state=seed,
17     shuffle=is_shuffle
18 )
```

# ResNet

## ❖ Step 4: Create pytorch datasets

```
1 class WeatherDataset(Dataset):
2     def __init__(
3         self,
4         X, y,
5         transform=None
6     ):
7         self.transform = transform
8         self.img_paths = X
9         self.labels = y
10
11    def __len__(self):
12        return len(self.img_paths)
13
14    def __getitem__(self, idx):
15        img_path = self.img_paths[idx]
16        img = Image.open(img_path).convert("RGB")
17
18        if self.transform:
19            img = self.transform(img)
20
21        return img, self.labels[idx]
```

```
1 def transform(img, img_size=(224, 224)):
2     img = img.resize(img_size)
3     img = np.array(img)[..., :3]
4     img = torch.tensor(img).permute(2, 0, 1).float()
5     normalized_img = img / 255.0
6
7     return normalized_img
```

```
1 train_dataset = WeatherDataset(
2     X_train, y_train,
3     transform=transform
4 )
5 val_dataset = WeatherDataset(
6     X_val, y_val,
7     transform=transform
8 )
9 test_dataset = WeatherDataset(
10    X_test, y_test,
11    transform=transform
12 )
```

Declare train, val, test datasets object

# ResNet

## ❖ Step 5: Create dataloader

```
1 train_batch_size = 512    }
2 test_batch_size = 8
3
4 train_loader = DataLoader(
5     train_dataset,
6     batch_size=train_batch_size,
7     shuffle=True
8 )
9 val_loader = DataLoader(
10    val_dataset,
11    batch_size=test_batch_size,
12    shuffle=False
13 )
14 test_loader = DataLoader(
15    test_dataset,
16    batch_size=test_batch_size,
17    shuffle=False
18 )
```

We must also declare the value of batch size for training and testing.

DataLoader is an iterator, so we can get a sample (image\_path, label) and visualize it.

```
1 train_features, train_labels = next(iter(train_loader))
2 print(f'Feature batch shape: {train_features.size()}')
3 print(f'Labels batch shape: {train_labels.size()}')
4 img = train_features[0].permute(1, 2, 0)
5 label = train_labels[0].item()
6 plt.imshow(img)
7 plt.axis('off')
8 plt.title(f'Label: {classes[label]}')
9 plt.show()
```

Feature batch shape: torch.Size([512, 3, 224, 224])  
Labels batch shape: torch.Size([512])

Label: rainbow



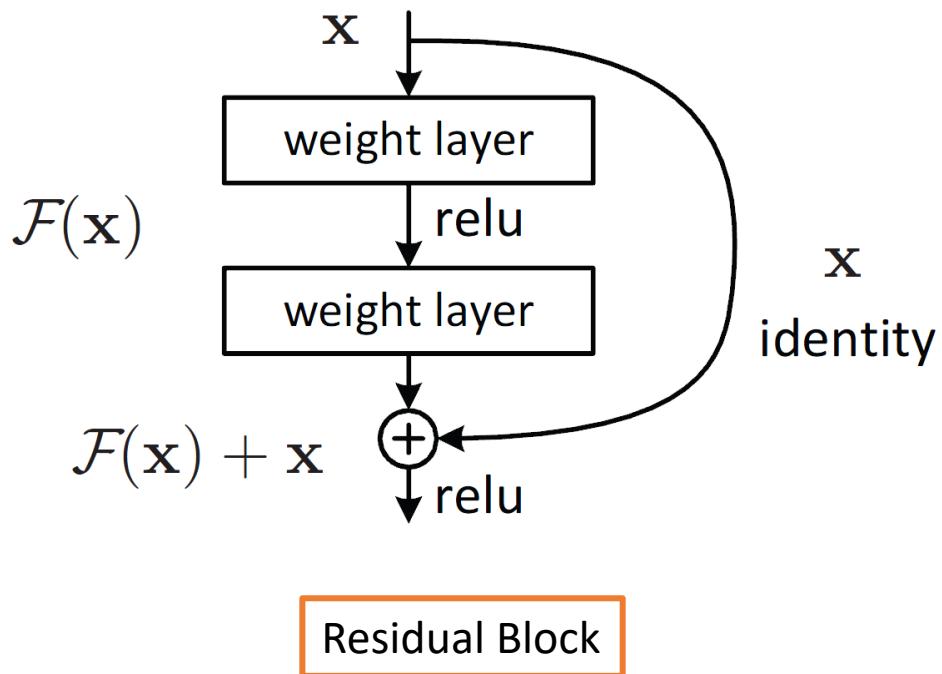
# ResNet

## ❖ Step 6: ResNet Architecture

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
conv2_x	56×56			3×3 max pool, stride 2		
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1			average pool, 1000-d fc, softmax		
FLOPs		$1.8 \times 10^9$	$3.6 \times 10^9$	$3.8 \times 10^9$	$7.6 \times 10^9$	$11.3 \times 10^9$

# ResNet

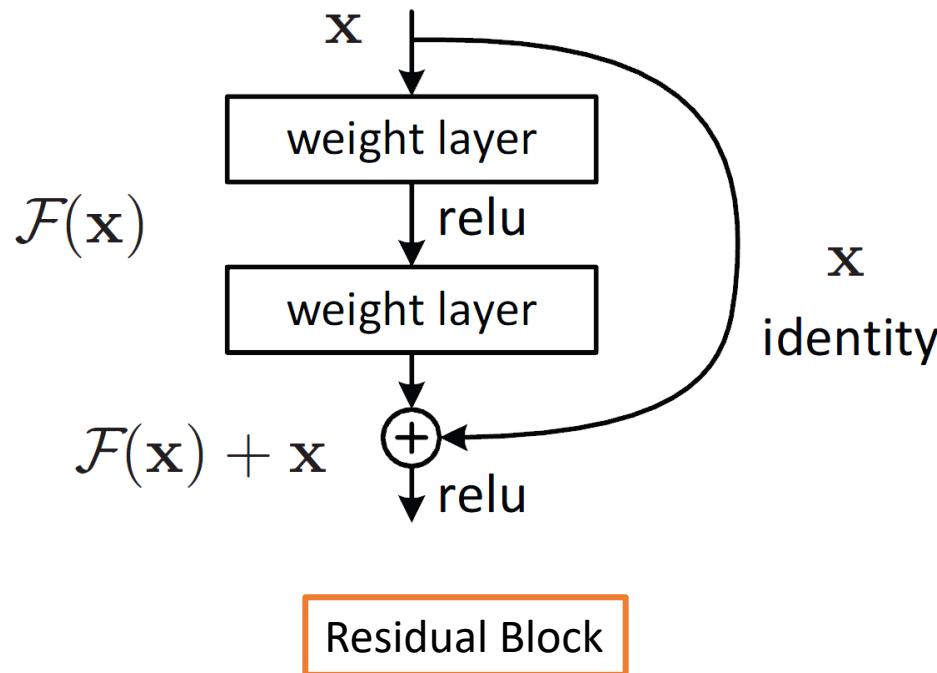
## ❖ Step 6: Implement Residual Block



```
1 class ResidualBlock(nn.Module):
2     def __init__(self, in_channels, out_channels, stride=1):
3         super(ResidualBlock, self).__init__()
4         self.conv1 = nn.Conv2d(
5             in_channels, out_channels,
6             kernel_size=3, stride=stride, padding=1
7         )
8         self.batch_norm1 = nn.BatchNorm2d(out_channels)
9         self.conv2 = nn.Conv2d(
10            out_channels, out_channels,
11            kernel_size=3, stride=1, padding=1
12        )
13         self.batch_norm2 = nn.BatchNorm2d(out_channels)
14
15         self.downsample = nn.Sequential()
16         if stride != 1 or in_channels != out_channels:
17             self.downsample = nn.Sequential(
18                 nn.Conv2d(
19                     in_channels, out_channels,
20                     kernel_size=1, stride=stride
21                 ),
22                 nn.BatchNorm2d(out_channels)
23             )
24         self.relu = nn.ReLU()
```

# ResNet

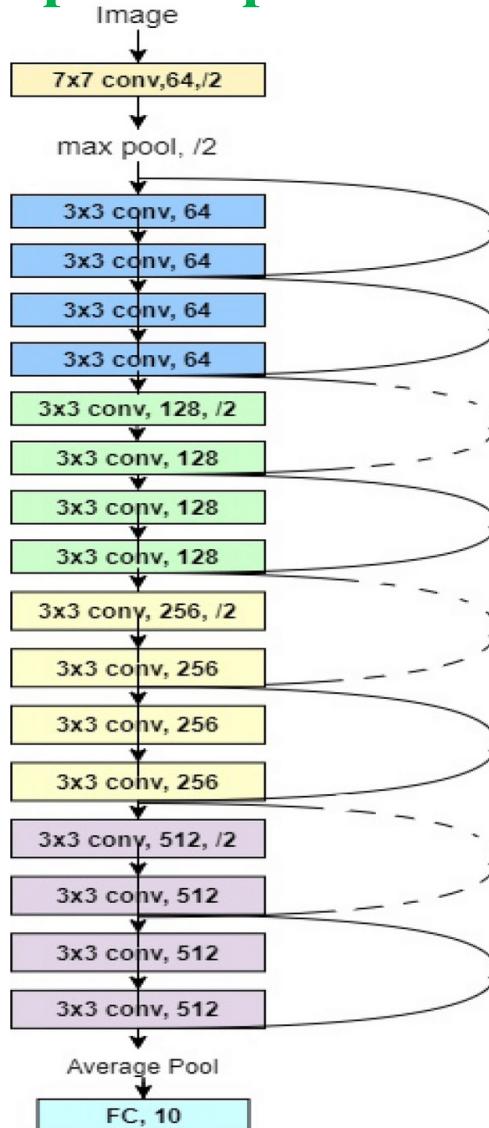
## ❖ Step 6: Implement Residual Block



```
26 def forward(self, x):  
27     shortcut = x.clone()  
28     x = self.conv1(x)  
29     x = self.batch_norm1(x)  
30     x = self.relu(x)  
31     x = self.conv2(x)  
32     x = self.batch_norm2(x)  
33     x += self.downsample(shortcut)  
34     x = self.relu(x)  
35  
36     return x
```

# ResNet

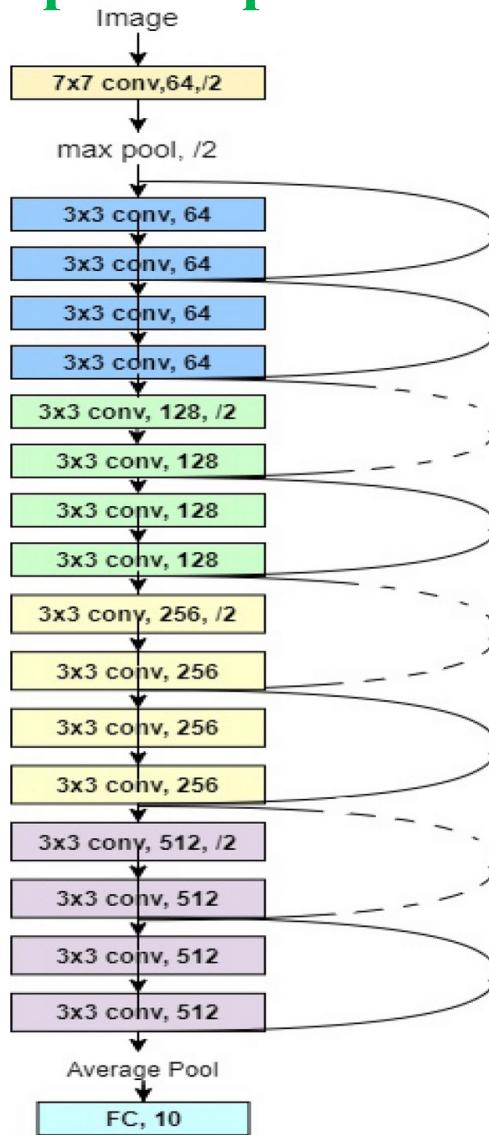
## ❖ Step 7: Implement ResNet (ResNet-18)



```
16 def create_layer(self, residual_block, in_channels, out_channels, n_blocks, stride):  
17     blocks = []  
18     first_block = residual_block(in_channels, out_channels, stride)  
19     blocks.append(first_block)  
20  
21     for idx in range(1, n_blocks):  
22         block = residual_block(out_channels, out_channels, stride)  
23         blocks.append(block)  
24  
25     block_sequential = nn.Sequential(*blocks)  
26  
27     return block_sequential  
1 class ResNet(nn.Module):  
2     def __init__(self, residual_block, n_blocks_lst, n_classes):  
3         super(ResNet, self).__init__()  
4         self.conv1 = nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3)  
5         self.batch_norm1 = nn.BatchNorm2d(64)  
6         self.relu = nn.ReLU()  
7         self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)  
8         self.conv2 = self.create_layer(residual_block, 64, 64, n_blocks_lst[0], 1)  
9         self.conv3 = self.create_layer(residual_block, 64, 128, n_blocks_lst[1], 2)  
10        self.conv4 = self.create_layer(residual_block, 128, 256, n_blocks_lst[2], 2)  
11        self.conv5 = self.create_layer(residual_block, 256, 512, n_blocks_lst[3], 2)  
12        self.avgpool = nn.AdaptiveAvgPool2d(1)  
13        self.flatten = nn.Flatten()  
14        self.fc1 = nn.Linear(512, n_classes)
```

# ResNet

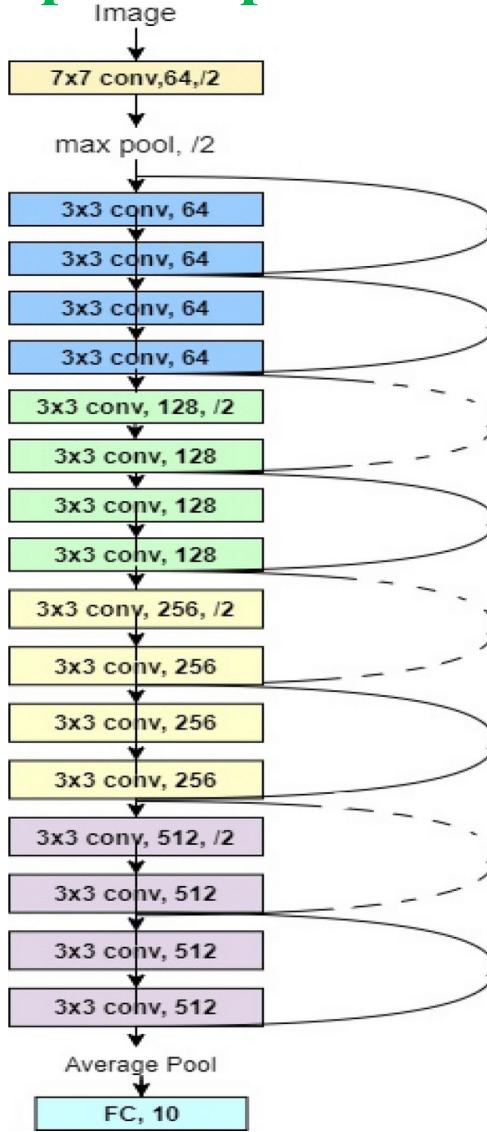
## ❖ Step 7: Implement ResNet (ResNet-18)



```
1 class ResNet(nn.Module):
2     def __init__(self, residual_block, n_blocks_lst, n_classes):
3         super(ResNet, self).__init__()
4         self.conv1 = nn.Conv2d(3, 64, kernel_size=7, stride=2, padding=3)
5         self.batch_norm1 = nn.BatchNorm2d(64)
6         self.relu = nn.ReLU()
7         self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)
8         self.conv2 = self.create_layer(residual_block, 64, 64, n_blocks_lst[0], 1)
9         self.conv3 = self.create_layer(residual_block, 64, 128, n_blocks_lst[1], 2)
10        self.conv4 = self.create_layer(residual_block, 128, 256, n_blocks_lst[2], 2)
11        self.conv5 = self.create_layer(residual_block, 256, 512, n_blocks_lst[3], 2)
12        self.avgpool = nn.AdaptiveAvgPool2d(1)
13        self.flatten = nn.Flatten()
14        self.fc1 = nn.Linear(512, n_classes)
```

# ResNet

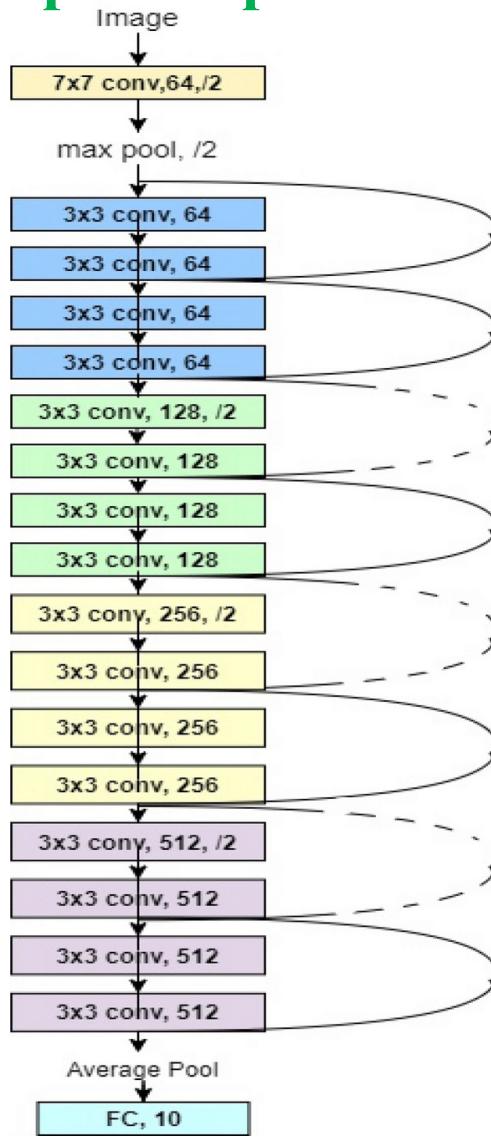
## ❖ Step 7: Implement ResNet (ResNet-18)



```
29 def forward(self, x):  
30     x = self.conv1(x)  
31     x = self.batch_norm1(x)  
32     x = self.maxpool(x)  
33     x = self.relu(x)  
34     x = self.conv2(x)  
35     x = self.conv3(x)  
36     x = self.conv4(x)  
37     x = self.conv5(x)  
38     x = self.avgpool(x)  
39     x = self.flatten(x)  
40     x = self.fc1(x)  
41  
42     return x
```

# ResNet

## ❖ Step 7: Implement ResNet (ResNet-18)



```
1 n_classes = len(list(classes.keys()))
2 device = 'cuda' if torch.cuda.is_available() else 'cpu'
3
4 model = ResNet(
5     residual_block=ResidualBlock,
6     n_blocks_lst=[2, 2, 2, 2],
7     n_classes=n_classes
8 ).to(device)
```

```
1 model.eval()
2
3 dummy_tensor = torch.randn(1, 3, 224, 224).to(device)
4
5 with torch.no_grad():
6     output = model(dummy_tensor)
7
8 print('Output shape:', output.shape)
```

Output shape: torch.Size([1, 11])

After initialize the model,  
we should whether it  
works correctly or not

# ResNet

## ❖ Step 8: Create fit function

**From line 1 to line 9:** Declare fit() function with the following parameters:

- model: the pytorch model to use.
- train\_loader: the dataloader of train set.
- val\_loader: the dataloader of val set.
- criterion: the pytorch loss function.
- optimizer: the pytorch learning algorithm.
- device: ‘cuda’ or ‘cpu’.
- epochs: number of epochs.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'EPOCH {epoch + 1}: \tTrain loss: {train_loss:.4f} \tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 8: Create fit function

**Line 10, 11:** Declare two empty lists for storing losses on train set and val set over epochs. These will be the return values of our function.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'EPOCH {epoch + 1}:\tTrain loss: {train_loss:.4f}\tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 8: Create fit function

**Line 13, 14, 15, 16:** Iterate over each epoch, create an empty list for storing loss on each iteration and turn the training mode of the model on.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'EPOCH {epoch + 1}: \tTrain loss: {train_loss:.4f} \tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 8: Create fit function

**Line 17, 18:** Iterate over each batch of samples in train dataloader, convert the input tensor into correspoding device (CPU or GPU).

**Line 20, 21:** Clear the previous gradient and execute the prediction on inputs.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'EPOCH {epoch + 1}: \tTrain loss: {train_loss:.4f} \tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 8: Create fit function

**Line 22, 23, 24, 26:** Compute the loss, gradient and update the computed gradient to the optimizer.  
Then, append the computed loss to empty list declared on line 14.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'EPOCH {epoch + 1}: \tTrain loss: {train_loss:.4f} \tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 8: Create fit function

**Line 28, 29:** With a list of loss value of each iteration, we average them to get a single representative value for the current epoch.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28    train_loss = sum(batch_train_losses) / len(batch_train_losses)
29    train_losses.append(train_loss)
30
31    val_loss, val_acc = evaluate(
32        model, val_loader,
33        criterion, device
34    )
35    val_losses.append(val_loss)
36
37    print(f'EPOCH {epoch + 1}:\tTrain loss: {train_loss:.4f}\tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 8: Create fit function

**From line 31 to 39:** Evaluate the model on val set, append the val loss to empty list declared in line 11, print the training result on the current epoch. If the iteration over epoch ends, return the train\_losses and val\_losses.

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'EPOCH {epoch + 1}: \tTrain loss: {train_loss:.4f} \tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# ResNet

## ❖ Step 9: Training

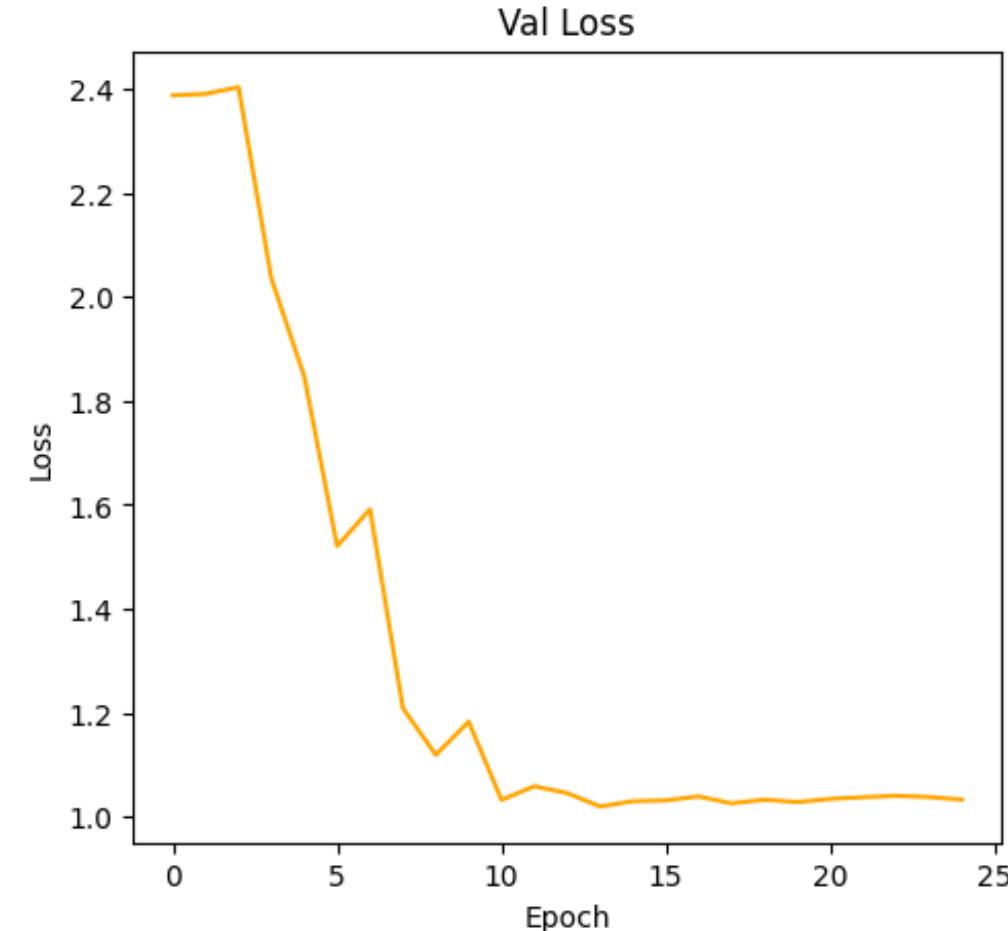
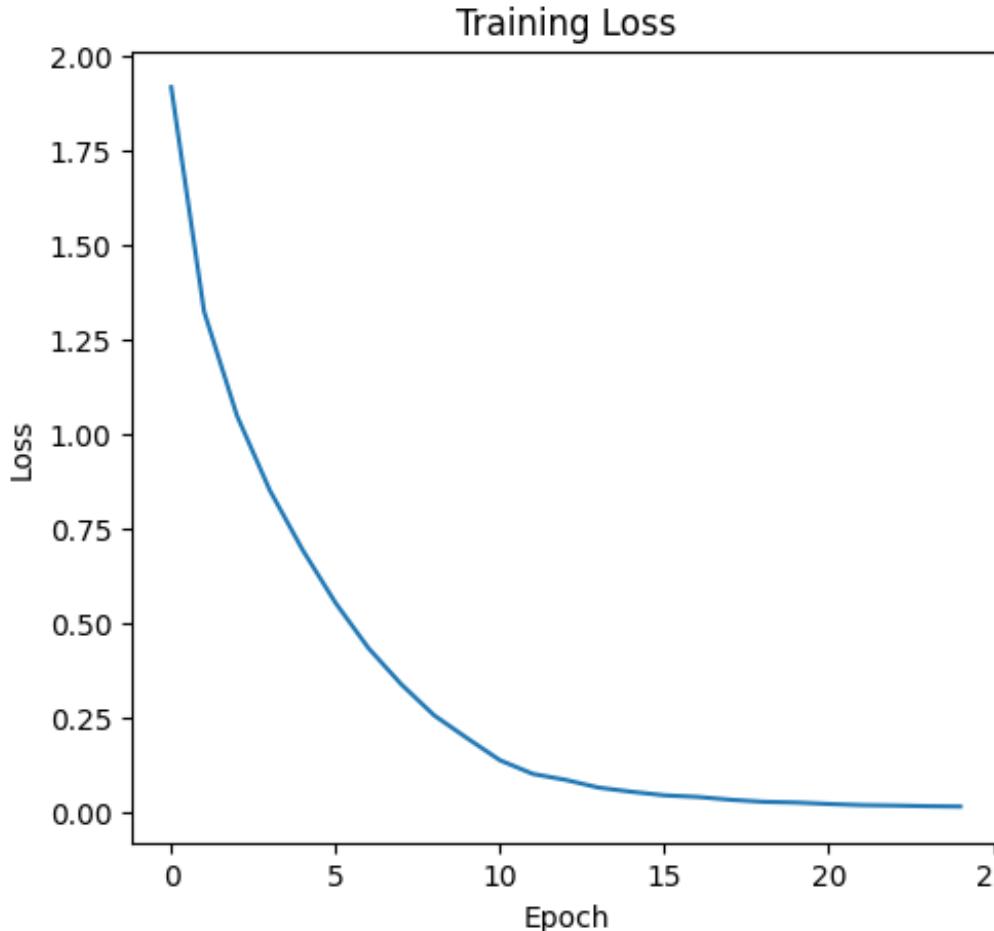
```
1 lr = 1e-2
2 epochs = 25
3
4 criterion = nn.CrossEntropyLoss()
5 optimizer = torch.optim.SGD(
6     model.parameters(),
7     lr=lr
8 )
```

```
1 train_losses, val_losses = fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 )
```

EPOCH 1:	Train loss: 1.9155	Val loss: 2.3871
EPOCH 2:	Train loss: 1.3232	Val loss: 2.3897
EPOCH 3:	Train loss: 1.0485	Val loss: 2.4028
EPOCH 4:	Train loss: 0.8511	Val loss: 2.0351
EPOCH 5:	Train loss: 0.6936	Val loss: 1.8467
EPOCH 6:	Train loss: 0.5541	Val loss: 1.5207
EPOCH 7:	Train loss: 0.4349	Val loss: 1.5914
EPOCH 8:	Train loss: 0.3392	Val loss: 1.2092
EPOCH 9:	Train loss: 0.2574	Val loss: 1.1194
EPOCH 10:	Train loss: 0.1973	Val loss: 1.1830
EPOCH 11:	Train loss: 0.1389	Val loss: 1.0327
EPOCH 12:	Train loss: 0.1029	Val loss: 1.0589
EPOCH 13:	Train loss: 0.0871	Val loss: 1.0458
EPOCH 14:	Train loss: 0.0667	Val loss: 1.0200
EPOCH 15:	Train loss: 0.0558	Val loss: 1.0302
EPOCH 16:	Train loss: 0.0464	Val loss: 1.0318
EPOCH 17:	Train loss: 0.0422	Val loss: 1.0393
EPOCH 18:	Train loss: 0.0348	Val loss: 1.0261
EPOCH 19:	Train loss: 0.0294	Val loss: 1.0330
EPOCH 20:	Train loss: 0.0274	Val loss: 1.0285
EPOCH 21:	Train loss: 0.0236	Val loss: 1.0346
EPOCH 22:	Train loss: 0.0206	Val loss: 1.0378
EPOCH 23:	Train loss: 0.0194	Val loss: 1.0406
EPOCH 24:	Train loss: 0.0178	Val loss: 1.0382
EPOCH 25:	Train loss: 0.0169	Val loss: 1.0332

# ResNet

## ❖ Step 9: Training



# ResNet

## ❖ Step 10: Create evaluation function

```
1 def evaluate(model, dataloader, criterion, device):
2     model.eval()
3     correct = 0
4     total = 0
5     losses = []
6     with torch.no_grad():
7         for inputs, labels in dataloader:
8             inputs, labels = inputs.to(device), labels.to(device)
9             outputs = model(inputs)
10            loss = criterion(outputs, labels)
11            losses.append(loss.item())
12            _, predicted = torch.max(outputs.data, 1)
13            total += labels.size(0)
14            correct += (predicted == labels).sum().item()
15
16            loss = sum(losses) / len(losses)
17            acc = correct / total
18
19    return loss, acc
```

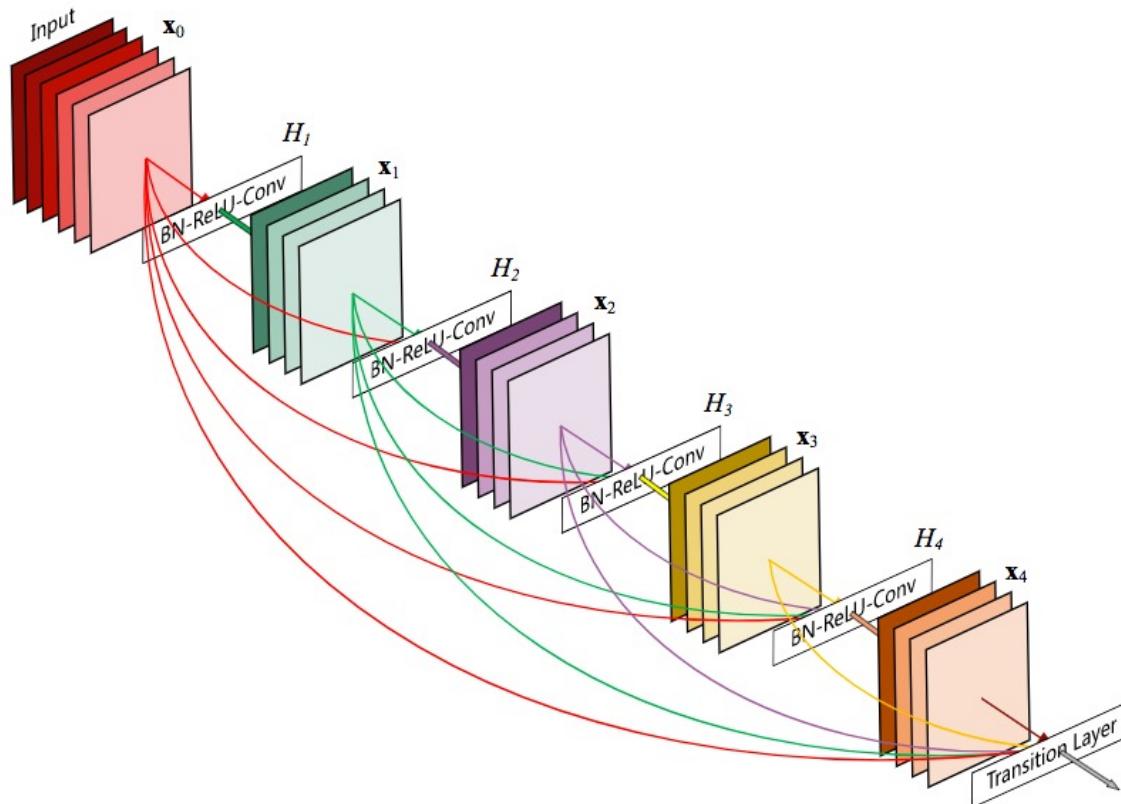
```
val_loss, val_acc = evaluate(
    model,
    val_loader,
    criterion,
    device
)
test_loss, test_acc = evaluate(
    model,
    test_loader,
    criterion,
    device
)

print('Evaluation on val/test dataset')
print('Val accuracy: ', val_acc)
print('Test accuracy: ', test_acc)
```

```
Evaluation on val/test dataset
Val accuracy: 0.6758922068463219
Test accuracy: 0.7040072859744991
```

# DenseNet

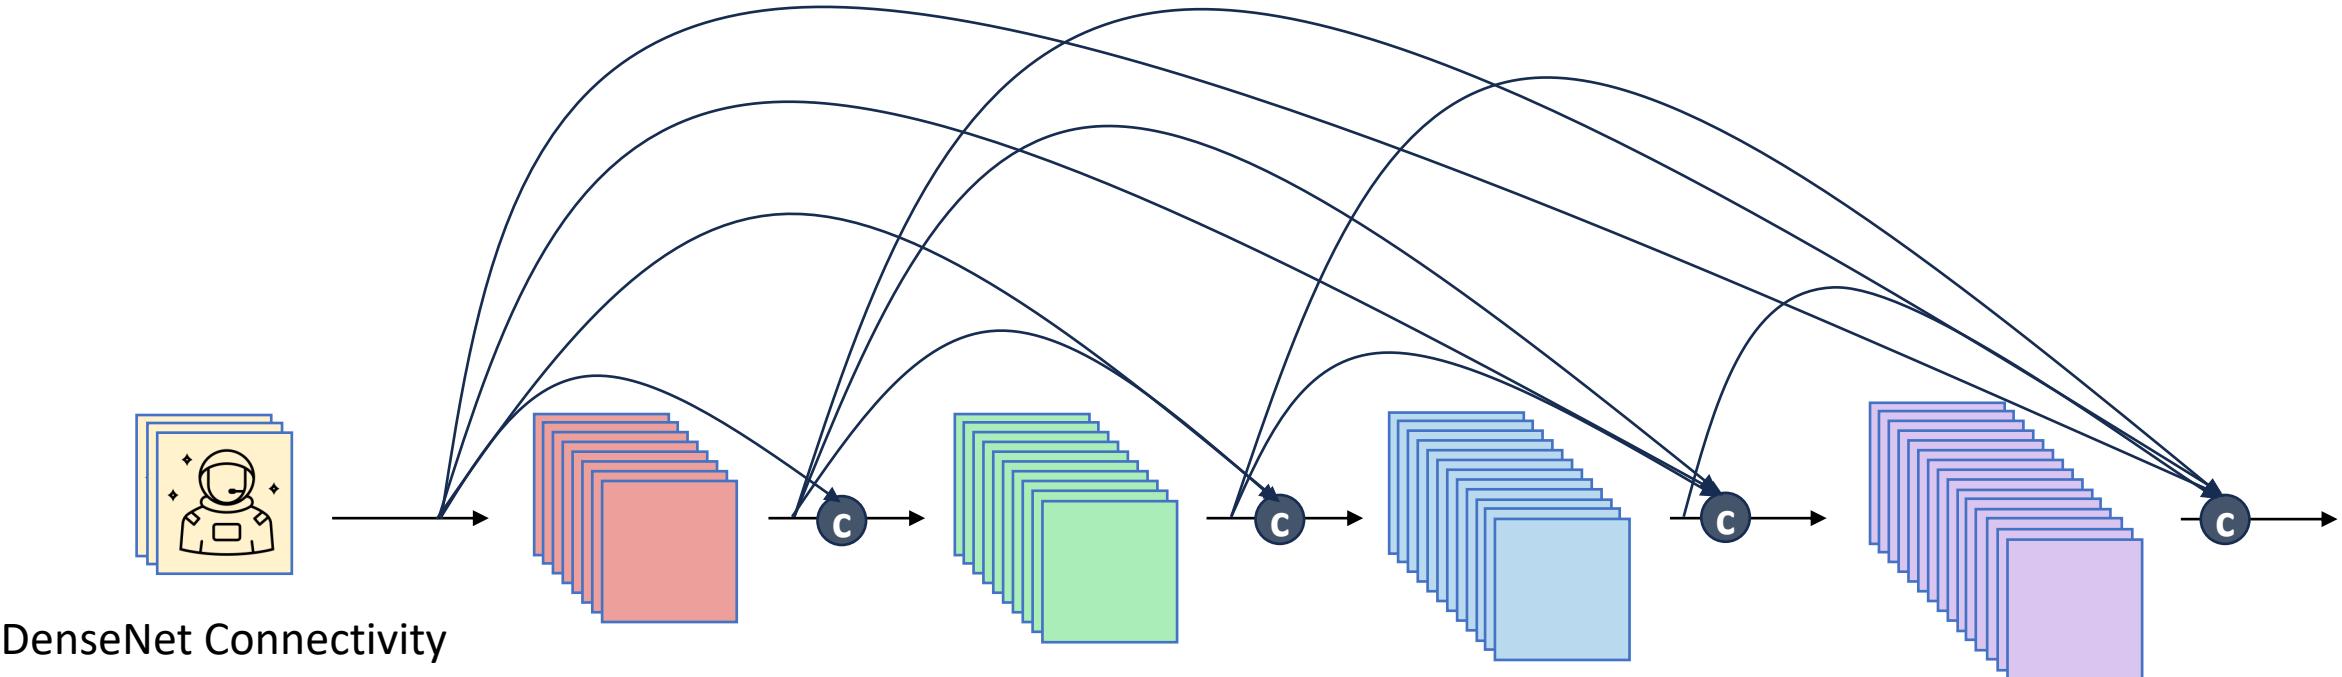
## ❖ DenseNet Architecture



**Densely Connected Convolutional Networks (DenseNet):** A type of CNNs introduced in 2017. It revolutionized CNNs by introducing dense connectivity (DenseBlock), facilitating information flow and feature reuse.

# DenseNet

## ❖ DenseNet architecture

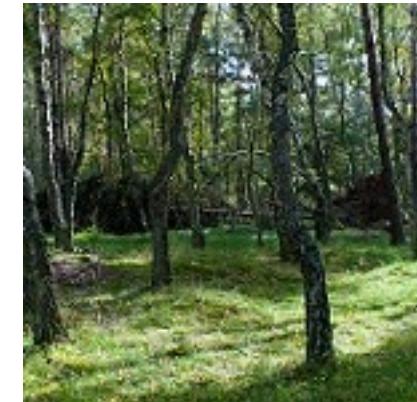


In **DenseNet**, each layer is not only connected to the previous layer but also to all subsequent layers in the network.

# DenseNet

## ❖ Coding Exercise

**Problem Statement:** Given the scenes image dataset (download [here](#)), build a scenes image classifier using DenseNet architecture.



# DenseNet

## ❖ Step 1: Import libraries

```
1 import torch
2 import torch.nn as nn
3 import os
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7
8 from PIL import Image
9 from torch.utils.data import Dataset, DataLoader
10 from sklearn.model_selection import train_test_split
```



# DenseNet

## ❖ Step 2: Read dataset

```
scenes_classification
  train
    buildings
    forest
    glacier
    mountain
    sea
    street
  val
    buildings
    forest
    glacier
    mountain
    sea
    street
```

```
1 root_dir = 'dataset/scenes_classification'
2 train_dir = os.path.join(root_dir, 'train')
3 test_dir = os.path.join(root_dir, 'val')
4
5 classes = {
6     label_idx: class_name \
7         for label_idx, class_name in enumerate(
8             sorted(os.listdir(train_dir))
9         )
10 }
```

1 classes

```
{0: 'buildings',
1: 'forest',
2: 'glacier',
3: 'mountain',
4: 'sea',
5: 'street'}
```

Each folder name is also a class name, thus we could get list of classnames by reading over them.

# DenseNet

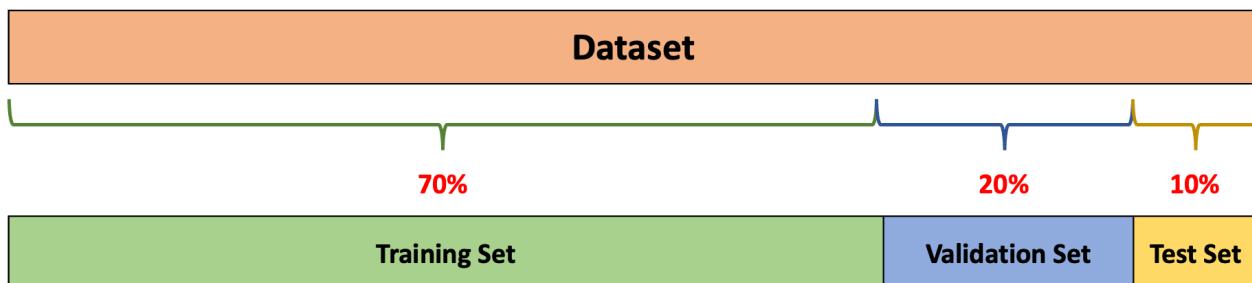
## ❖ Step 2: Read dataset



```
12 X_train = []
13 y_train = []
14 X_test = []
15 y_test = []
16 for dataset_path in [train_dir, test_dir]:
17     for label_idx, class_name in classes.items():
18         class_dir = os.path.join(dataset_path, class_name)
19         for img_filename in os.listdir(class_dir):
20             img_path = os.path.join(class_dir, img_filename)
21             if 'train' in dataset_path:
22                 X_train.append(img_path)
23                 y_train.append(label_idx)
24             else:
25                 X_test.append(img_path)
26                 y_test.append(label_idx)
```

# DenseNet

## ❖ Step 3: Train, val, test split



```
1 seed = 0
2 val_size = 0.2
3 is_shuffle = True
4
5 X_train, X_val, y_train, y_val = train_test_split(
6     X_train, y_train,
7     test_size=val_size,
8     random_state=seed,
9     shuffle=is_shuffle
10 )
```

We only need to split once since we already have test set

# DenseNet

## ❖ Step 4: Create pytorch dataset

```
1 class ScenesDataset(Dataset):
2     def __init__(self, X, y, transform=None):
3         self.transform = transform
4         self.img_paths = X
5         self.labels = y
6
7     def __len__(self):
8         return len(self.img_paths)
9
10    def __getitem__(self, idx):
11        img_path = self.img_paths[idx]
12        img = Image.open(img_path).convert("RGB")
13
14        if self.transform:
15            img = self.transform(img)
16
17        return img, self.labels[idx]
```

```
1 def transform(img, img_size=(224, 224)):
2     img = img.resize(img_size)
3     img = np.array(img)[..., :3]
4     img = torch.tensor(img).permute(2, 0, 1).float()
5     normalized_img = img / 255.0
6
7     return normalized_img
```

```
1 train_dataset = ScenesDataset(
2     X_train, y_train,
3     transform=transform
4 )
5 val_dataset = ScenesDataset(
6     X_val, y_val,
7     transform=transform
8 )
9 test_dataset = ScenesDataset(
10    X_test, y_test,
11    transform=transform
12 )
```

Declare train, val, test datasets object

# DenseNet

## ❖ Step 5: Create dataloader

```
1 train_batch_size = 64
2 test_batch_size = 8
3
4 train_loader = DataLoader(
5     train_dataset,
6     batch_size=train_batch_size,
7     shuffle=True
8 )
9 val_loader = DataLoader(
10    val_dataset,
11    batch_size=test_batch_size,
12    shuffle=False
13 )
14 test_loader = DataLoader(
15    test_dataset,
16    batch_size=test_batch_size,
17    shuffle=False
18 )
```

We must also declare the value of batch size for training and testing.

```
1 train_features, train_labels = next(iter(train_loader))
2 print(f'Feature batch shape: {train_features.size()}')
3 print(f'Labels batch shape: {train_labels.size()}')
4 img = train_features[0].permute(1, 2, 0)
5 label = train_labels[0].item()
6 plt.imshow(img)
7 plt.axis('off')
8 plt.title(f'Label: {classes[label]}')
9 plt.show()
```

Feature batch shape: torch.Size([64, 3, 224, 224])  
Labels batch shape: torch.Size([64])

Label: buildings



DataLoader is an iterator, so we can get a sample (image\_path, label) and visualize it.

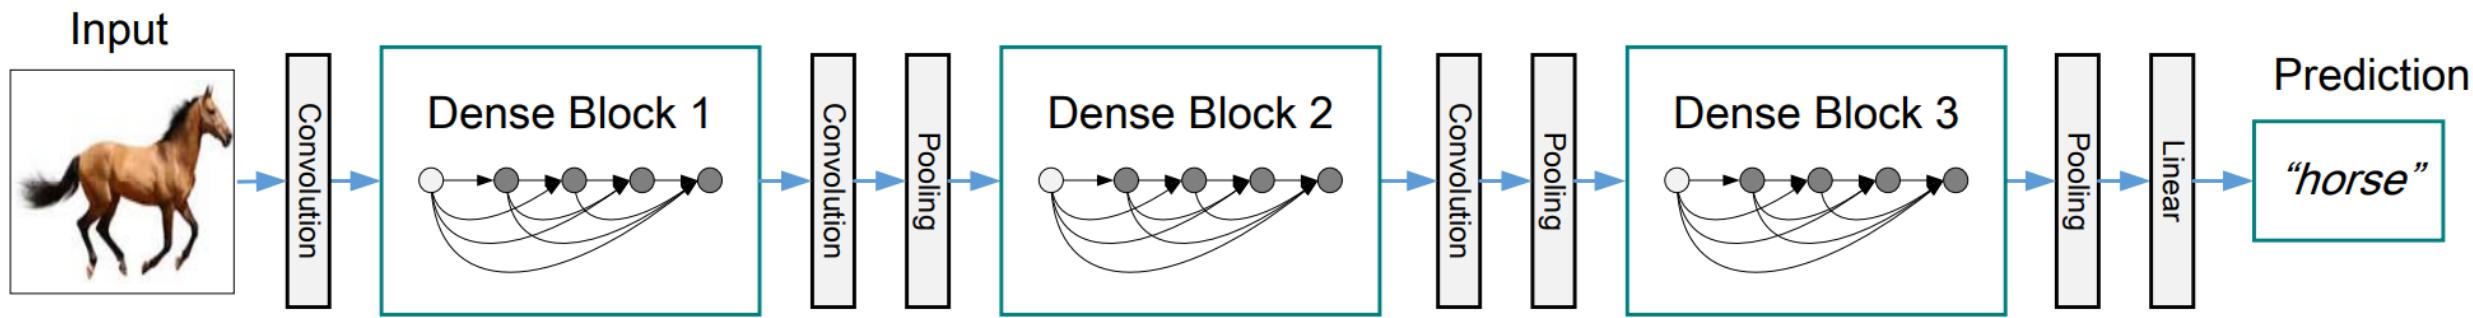
# DenseNet

## ❖ Step 6: DenseNet Architecture

Layers	Output Size	DenseNet-121	DenseNet-169	DenseNet-201	DenseNet-264
Convolution	$112 \times 112$		$7 \times 7$ conv, stride 2		
Pooling	$56 \times 56$		$3 \times 3$ max pool, stride 2		
Dense Block (1)	$56 \times 56$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 6$
Transition Layer (1)	$56 \times 56$			$1 \times 1$ conv	
	$28 \times 28$			$2 \times 2$ average pool, stride 2	
Dense Block (2)	$28 \times 28$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 12$
Transition Layer (2)	$28 \times 28$			$1 \times 1$ conv	
	$14 \times 14$			$2 \times 2$ average pool, stride 2	
Dense Block (3)	$14 \times 14$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 24$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 48$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 64$
Transition Layer (3)	$14 \times 14$			$1 \times 1$ conv	
	$7 \times 7$			$2 \times 2$ average pool, stride 2	
Dense Block (4)	$7 \times 7$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 16$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 32$	$\begin{bmatrix} 1 \times 1 \text{ conv} \\ 3 \times 3 \text{ conv} \end{bmatrix} \times 48$
Classification Layer	$1 \times 1$			$7 \times 7$ global average pool	
				1000D fully-connected, softmax	

# DenseNet

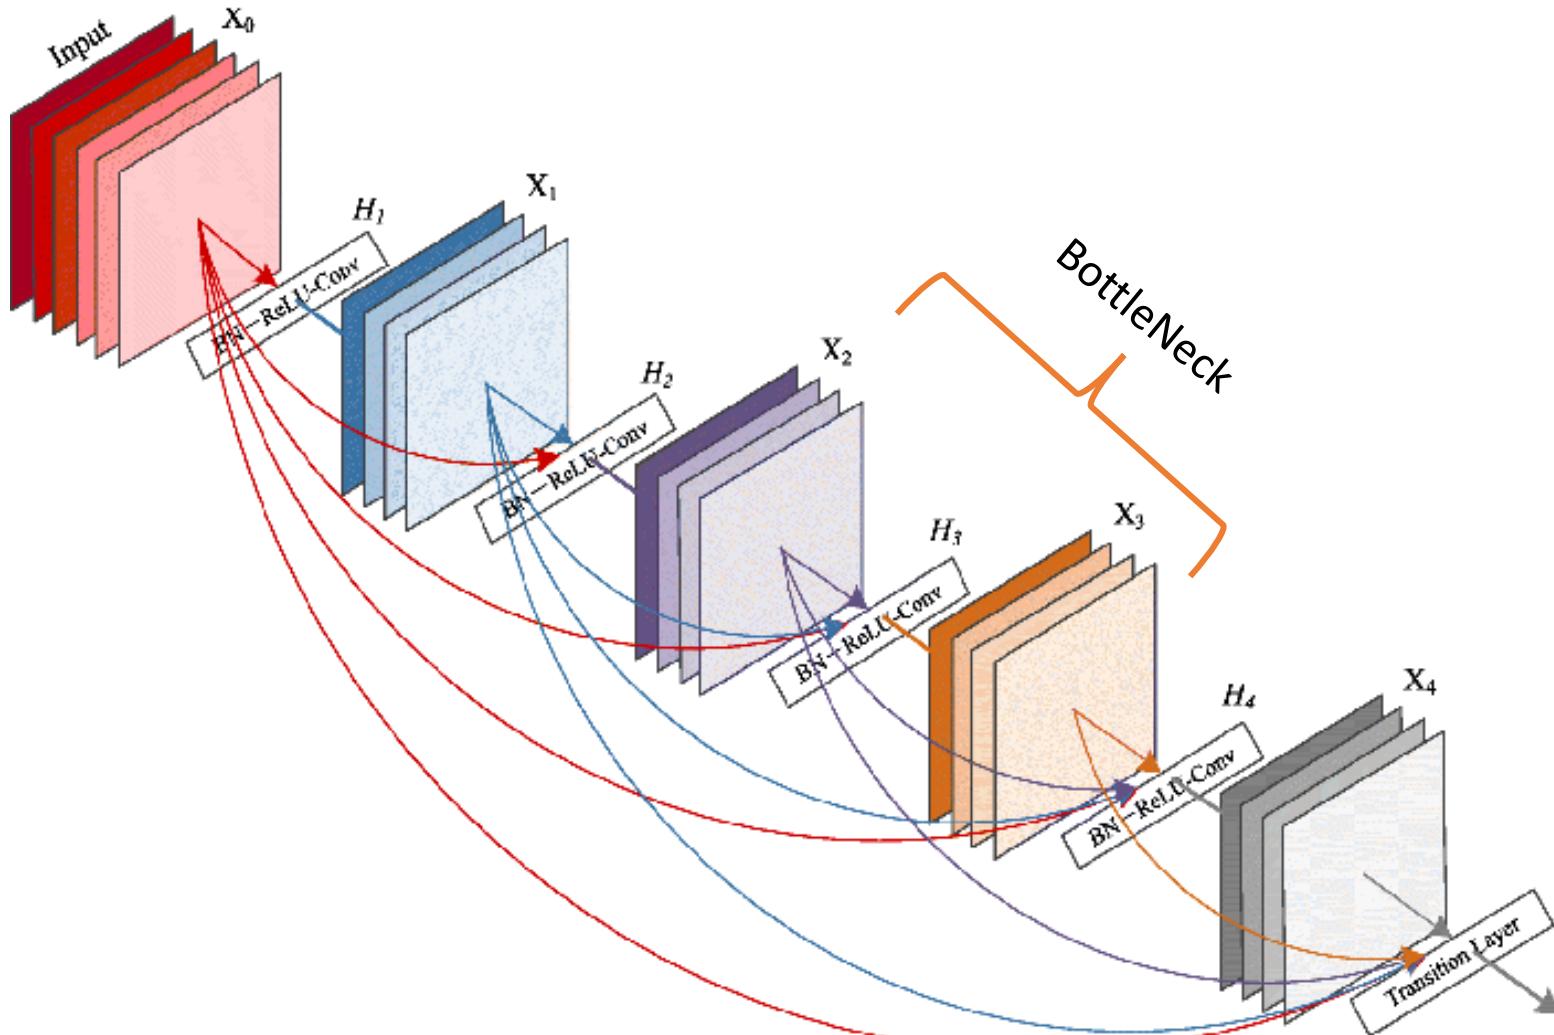
## ❖ Step 6: DenseNet Architecture Visualization



**Figure 2:** A deep DenseNet with three dense blocks. The layers between two adjacent blocks are referred to as transition layers and change feature-map sizes via convolution and pooling.

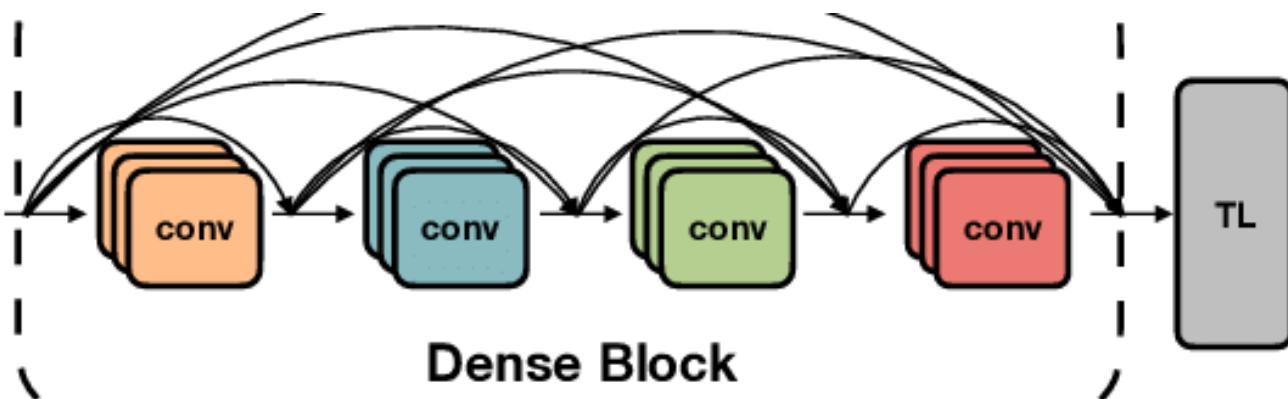
# DenseNet

## ❖ Step 6: DenseBlock Architecture



# DenseNet

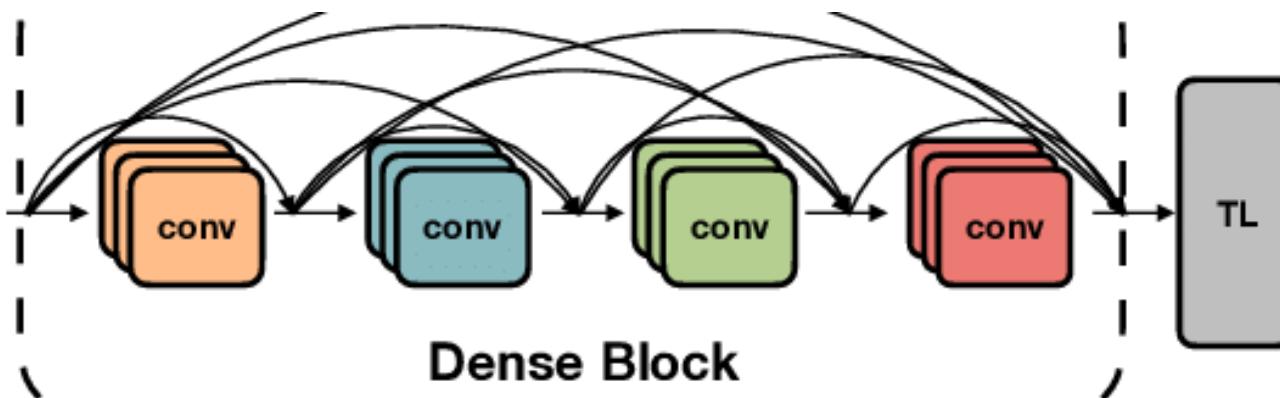
## ❖ Step 6: Implement DenseBlock



```
1 class BottleneckBlock(nn.Module):
2     def __init__(self, in_channels, growth_rate):
3         super(BottleneckBlock, self).__init__()
4         self.bn1 = nn.BatchNorm2d(in_channels)
5         self.conv1 = nn.Conv2d(
6             in_channels, 4 * growth_rate,
7             kernel_size=1,
8             bias=False
9         )
10        self.bn2 = nn.BatchNorm2d(4 * growth_rate)
11        self.conv2 = nn.Conv2d(
12            4 * growth_rate,
13            growth_rate,
14            kernel_size=3,
15            padding=1,
16            bias=False
17        )
18        self.relu = nn.ReLU()
19
20    def forward(self, x):
21        res = x.clone().detach()
22        x = self.bn1(x)
23        x = self.relu(x)
24        x = self.conv1(x)
25        x = self.bn2(x)
26        x = self.relu(x)
27        x = self.conv2(x)
28        x = torch.cat([res, x], 1)
29
30        return x
```

# DenseNet

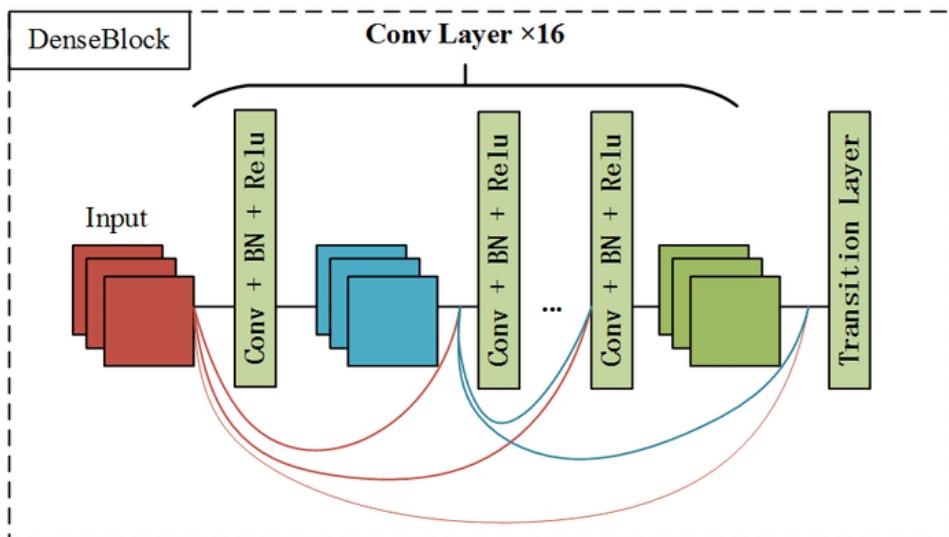
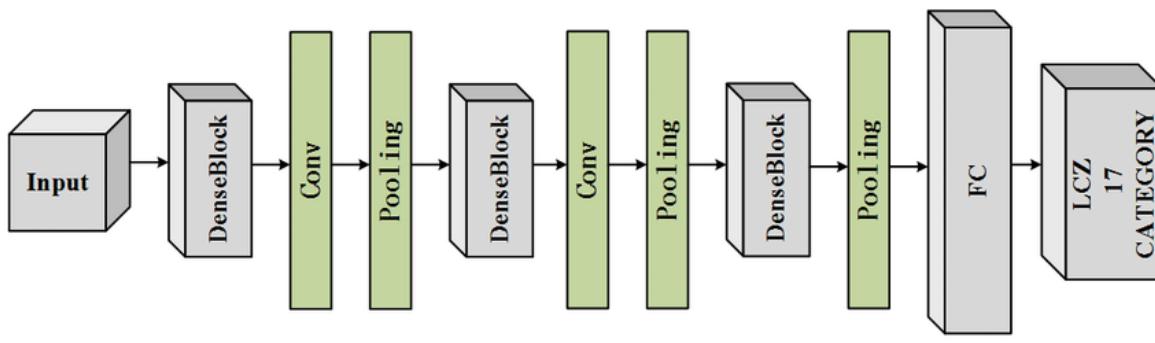
## ❖ Step 6: Implement DenseBlock



```
32 class DenseBlock(nn.Module):
33     def __init__(self, num_layers, in_channels, growth_rate):
34         super(DenseBlock, self).__init__()
35         layers = []
36         for i in range(num_layers):
37             layers.append(
38                 BottleneckBlock(
39                     in_channels + i * growth_rate,
40                     growth_rate
41                 )
42             )
43         self.block = nn.Sequential(*layers)
44
45     def forward(self, x):
46         return self.block(x)
```

# DenseNet

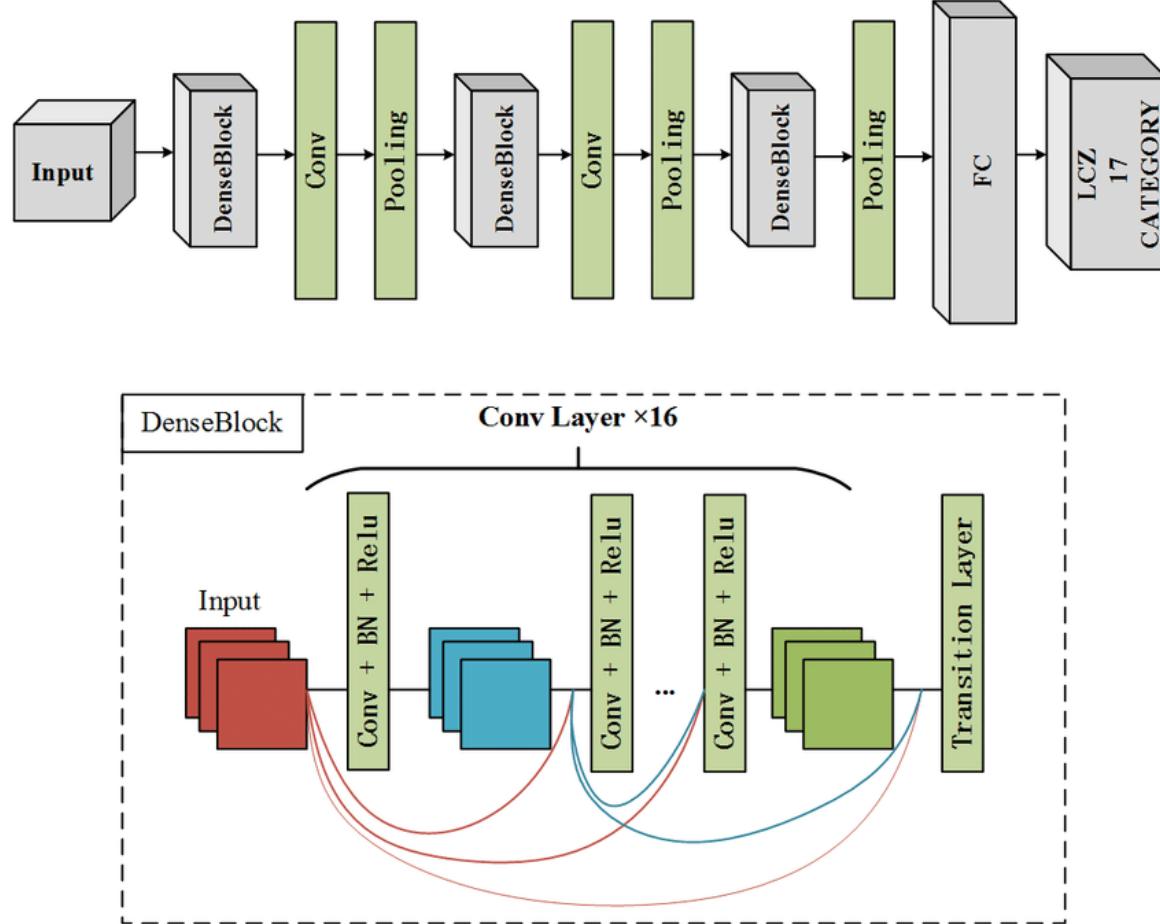
## ❖ Step 6: Implement DenseNet



```
1 class DenseNet(nn.Module):
2     def __init__(self, num_blocks, growth_rate, num_classes):
3         super(DenseNet, self).__init__()
4         self.conv1 = nn.Conv2d(
5             3, 2 * growth_rate,
6             kernel_size=7, padding=3,
7             stride=2, bias=False
8         )
9         self.bn1 = nn.BatchNorm2d(2 * growth_rate)
10        self.pool1 = nn.MaxPool2d(
11            kernel_size=3, stride=2, padding=1
12        )
13
14        self.dense_blocks = nn.ModuleList()
15        in_channels = 2 * growth_rate
16        for i, num_layers in enumerate(num_blocks):
17            self.dense_blocks.append(
18                DenseBlock(
19                    num_layers, in_channels, growth_rate
20                )
21            )
22            in_channels += num_layers * growth_rate
23            if i != len(num_blocks) - 1:
24                out_channels = in_channels // 2
25                self.dense_blocks.append(nn.Sequential(
26                    nn.BatchNorm2d(in_channels),
27                    nn.Conv2d(
28                        in_channels, out_channels,
29                        kernel_size=1, bias=False
30                    ),
31                    nn.AvgPool2d(kernel_size=2, stride=2)
32                ))
33                in_channels = out_channels
```

# DenseNet

## ❖ Step 6: Implement DenseNet



```
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
```

```
def forward(self, x):
    x = self.conv1(x)
    x = self.bn1(x)
    x = self.relu(x)
    x = self.pool1(x)

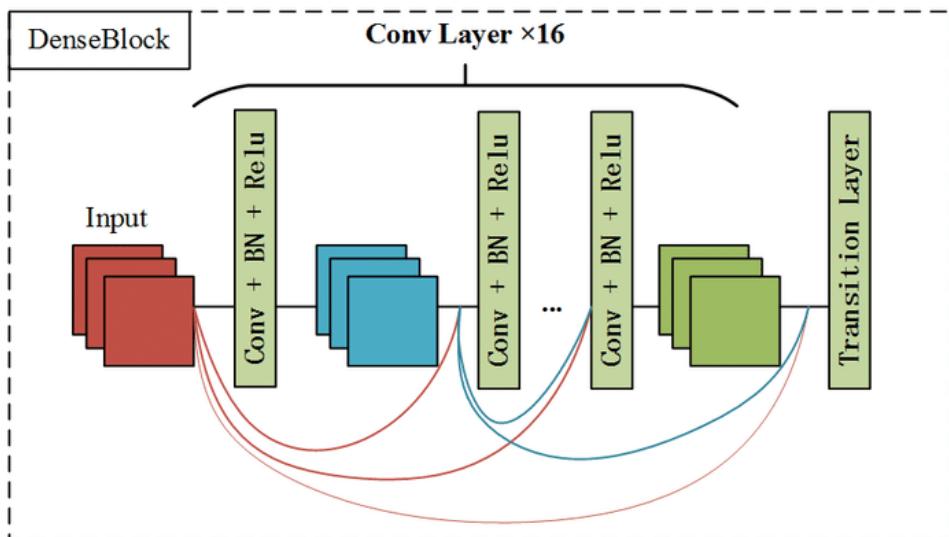
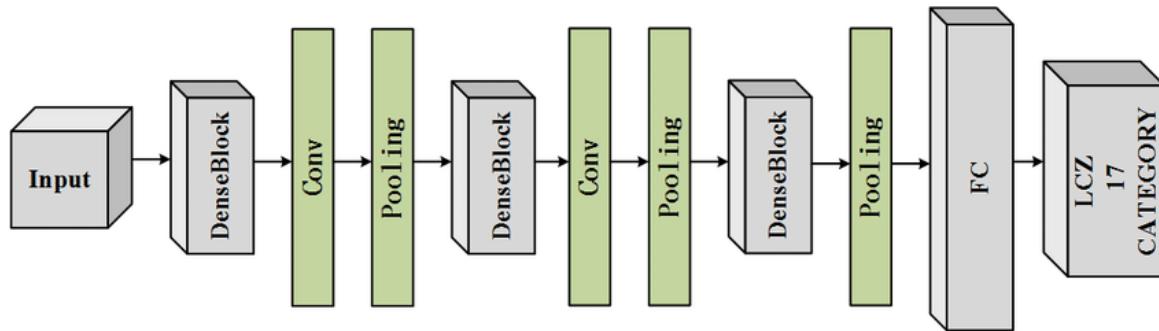
    for block in self.dense_blocks:
        x = block(x)

        x = self.bn2(x)
        x = self.relu(x)
        x = self.pool2(x)
        x = x.view(x.size(0), -1)
        x = self.fc(x)

    return x
```

# DenseNet

## ❖ Step 6: Implement DenseNet



```
1 n_classes = len(list(classes.keys()))
2 device = 'cuda' if torch.cuda.is_available() else 'cpu'
3
4 model = DenseNet(
5     [6, 12, 24, 16],
6     growth_rate=32,
7     num_classes=n_classes
8 ).to(device)
```

```
1 model.eval()
2
3 dummy_tensor = torch.randn(1, 3, 224, 224).to(device)
4
5 with torch.no_grad():
6     output = model(dummy_tensor)
7
8 print('Output shape:', output.shape)
```

Output shape: torch.Size([1, 6])

# DenseNet

## ❖ Step 8, 9: Create fit and evaluate function

```
1 def evaluate(model, dataloader, criterion, device):
2     model.eval()
3     correct = 0
4     total = 0
5     losses = []
6     with torch.no_grad():
7         for inputs, labels in dataloader:
8             inputs, labels = inputs.to(device), labels.to(device)
9             outputs = model(inputs)
10            loss = criterion(outputs, labels)
11            losses.append(loss.item())
12            _, predicted = torch.max(outputs.data, 1)
13            total += labels.size(0)
14            correct += (predicted == labels).sum().item()
15
16    loss = sum(losses) / len(losses)
17    acc = correct / total
18
19    return loss, acc
```

```
1 def fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 ):
10    train_losses = []
11    val_losses = []
12
13    for epoch in range(epochs):
14        batch_train_losses = []
15
16        model.train()
17        for idx, (inputs, labels) in enumerate(train_loader):
18            inputs, labels = inputs.to(device), labels.to(device)
19
20            optimizer.zero_grad()
21            outputs = model(inputs)
22            loss = criterion(outputs, labels)
23            loss.backward()
24            optimizer.step()
25
26            batch_train_losses.append(loss.item())
27
28        train_loss = sum(batch_train_losses) / len(batch_train_losses)
29        train_losses.append(train_loss)
30
31        val_loss, val_acc = evaluate(
32            model, val_loader,
33            criterion, device
34        )
35        val_losses.append(val_loss)
36
37        print(f'Epoch {epoch + 1}: \tTrain loss: {train_loss:.4f} \tVal loss: {val_loss:.4f}')
38
39    return train_losses, val_losses
```

# DenseNet

## ❖ Step 10: Training

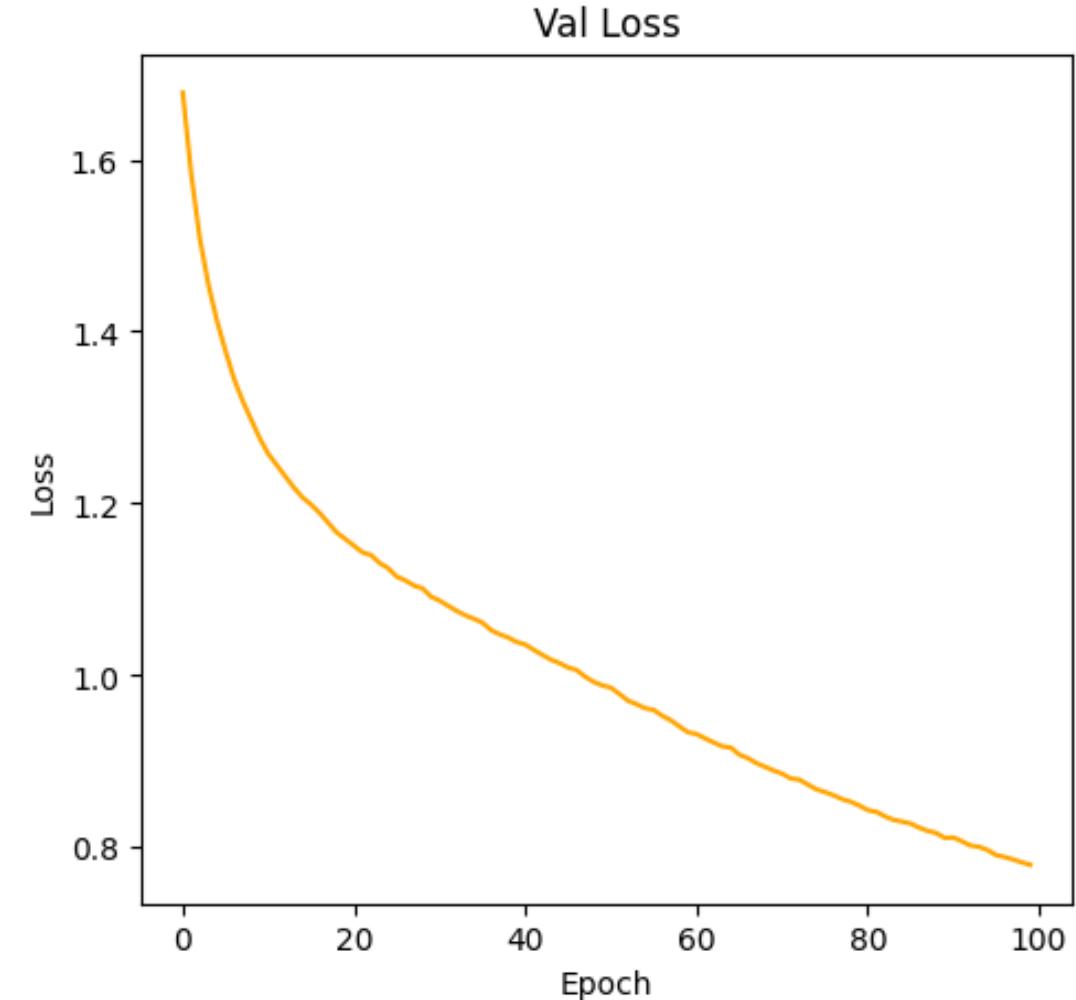
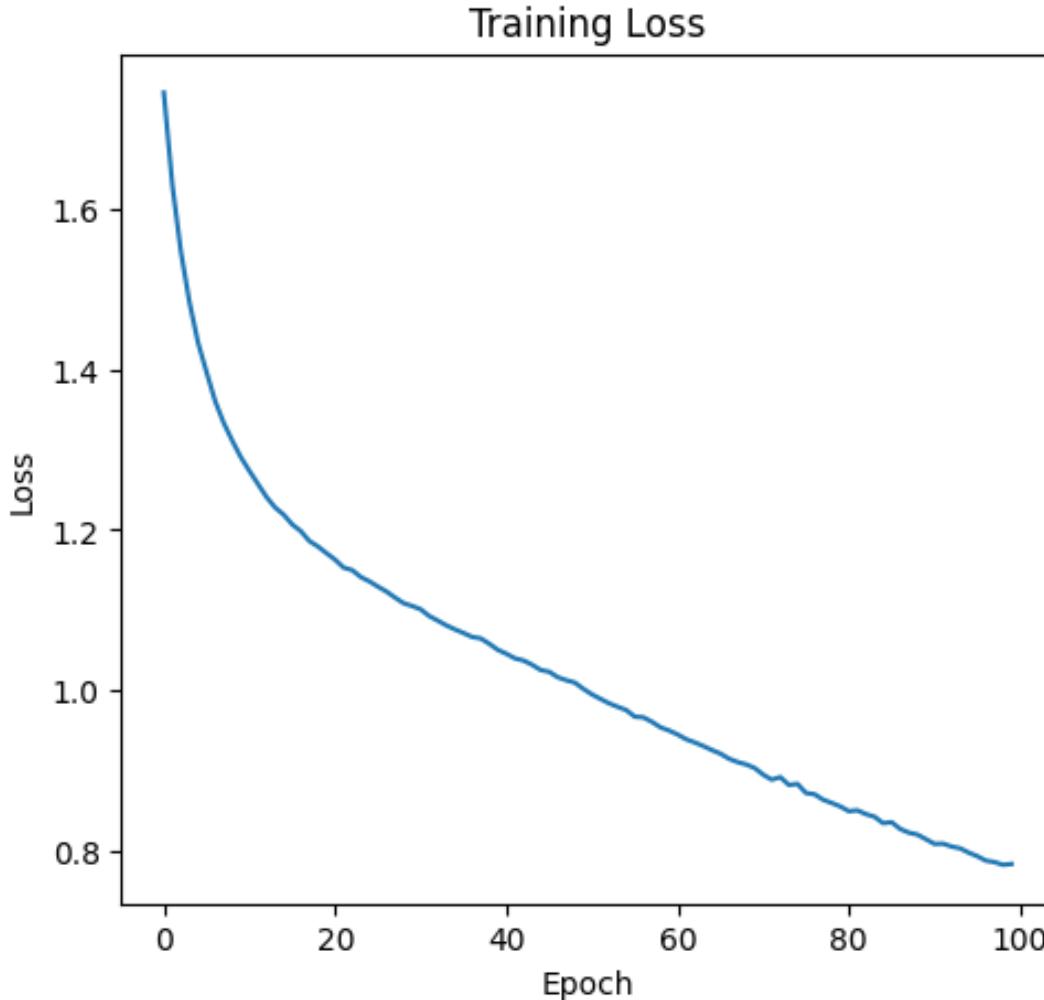
```
1 lr = 1e-3
2 epochs = 100
3
4 criterion = nn.CrossEntropyLoss()
5 optimizer = torch.optim.SGD(
6     model.parameters(),
7     lr=lr
8 )
```

```
1 train_losses, val_losses = fit(
2     model,
3     train_loader,
4     val_loader,
5     criterion,
6     optimizer,
7     device,
8     epochs
9 )
```

EPOCH 1:	Train loss: 1.7440	Val loss: 1.6777
EPOCH 2:	Train loss: 1.6288	Val loss: 1.5821
EPOCH 3:	Train loss: 1.5459	Val loss: 1.5079
EPOCH 4:	Train loss: 1.4831	Val loss: 1.4544
EPOCH 5:	Train loss: 1.4330	Val loss: 1.4122
EPOCH 6:	Train loss: 1.3960	Val loss: 1.3770
EPOCH 7:	Train loss: 1.3602	Val loss: 1.3445
EPOCH 8:	Train loss: 1.3331	Val loss: 1.3194
EPOCH 9:	Train loss: 1.3110	Val loss: 1.2973
EPOCH 10:	Train loss: 1.2904	Val loss: 1.2756
EPOCH 11:	Train loss: 1.2734	Val loss: 1.2575
EPOCH 12:	Train loss: 1.2571	Val loss: 1.2445
EPOCH 13:	Train loss: 1.2409	Val loss: 1.2312
EPOCH 14:	Train loss: 1.2279	Val loss: 1.2182
EPOCH 15:	Train loss: 1.2190	Val loss: 1.2069
EPOCH 16:	Train loss: 1.2069	Val loss: 1.1987
EPOCH 17:	Train loss: 1.1985	Val loss: 1.1888
EPOCH 18:	Train loss: 1.1863	Val loss: 1.1775
EPOCH 19:	Train loss: 1.1795	Val loss: 1.1662

# DenseNet

## ❖ Step 10: Training



# DenseNet

## ❖ Step 10: Training

```
1 def evaluate(model, dataloader, criterion, device):
2     model.eval()
3     correct = 0
4     total = 0
5     losses = []
6     with torch.no_grad():
7         for inputs, labels in dataloader:
8             inputs, labels = inputs.to(device), labels.to(device)
9             outputs = model(inputs)
10            loss = criterion(outputs, labels)
11            losses.append(loss.item())
12            _, predicted = torch.max(outputs.data, 1)
13            total += labels.size(0)
14            correct += (predicted == labels).sum().item()
15
16    loss = sum(losses) / len(losses)
17    acc = correct / total
18
19    return loss, acc
```

```
val_loss, val_acc = evaluate(
    model,
    val_loader,
    criterion,
    device
)
test_loss, test_acc = evaluate(
    model,
    test_loader,
    criterion,
    device
)

print('Evaluation on val/test dataset')
print('Val accuracy: ', val_acc)
print('Test accuracy: ', test_acc)
```

```
Evaluation on val/test dataset
Val accuracy: 0.7267545422158889
Test accuracy: 0.7133333333333334
```

# Question

