

# Generative Adversarial Networks

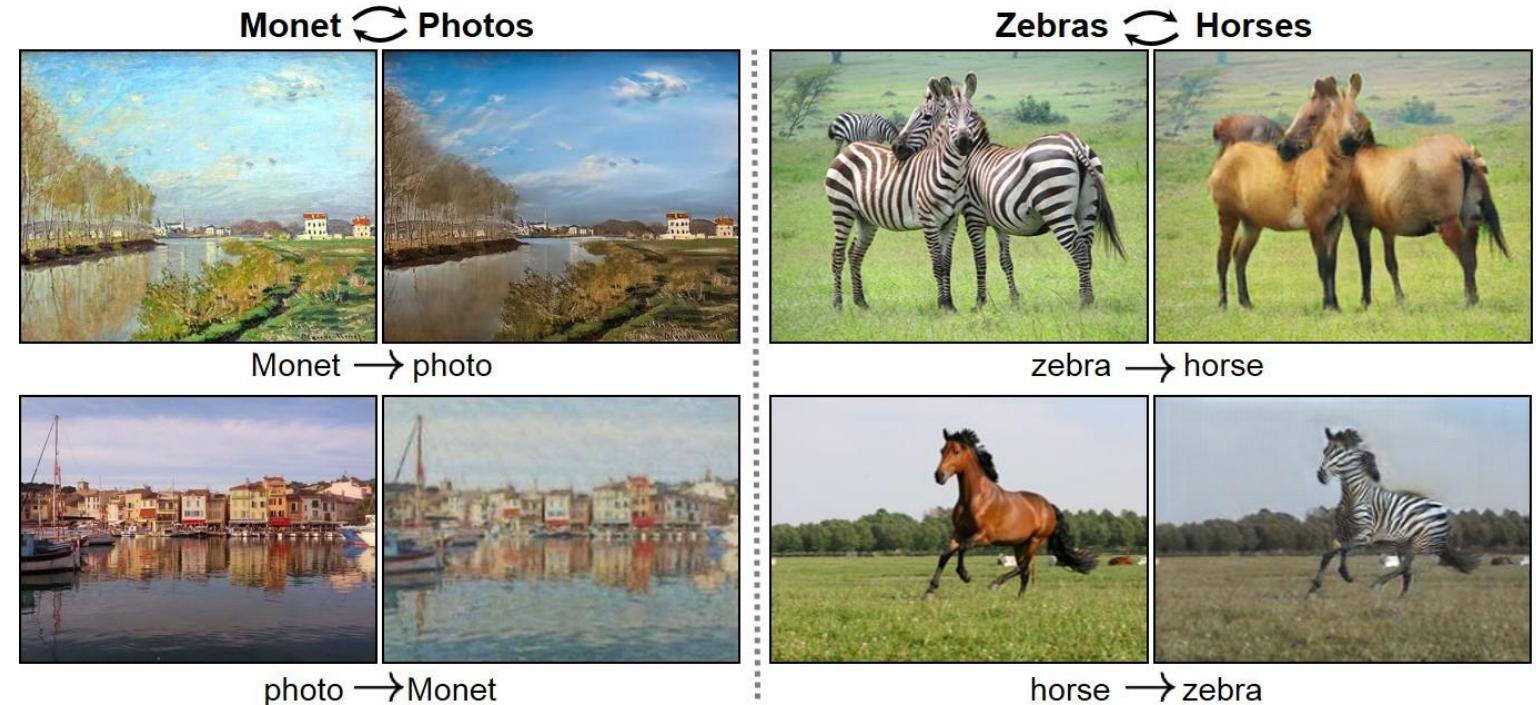
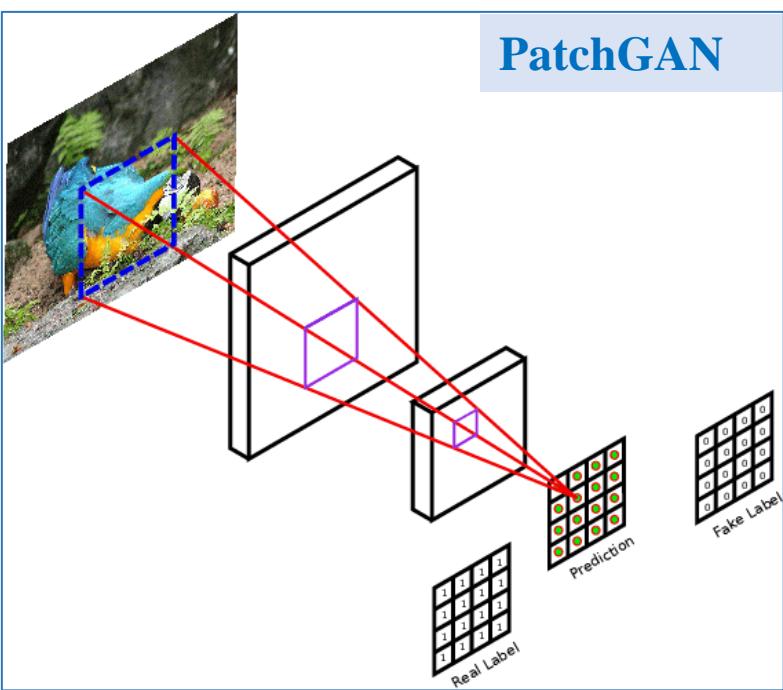
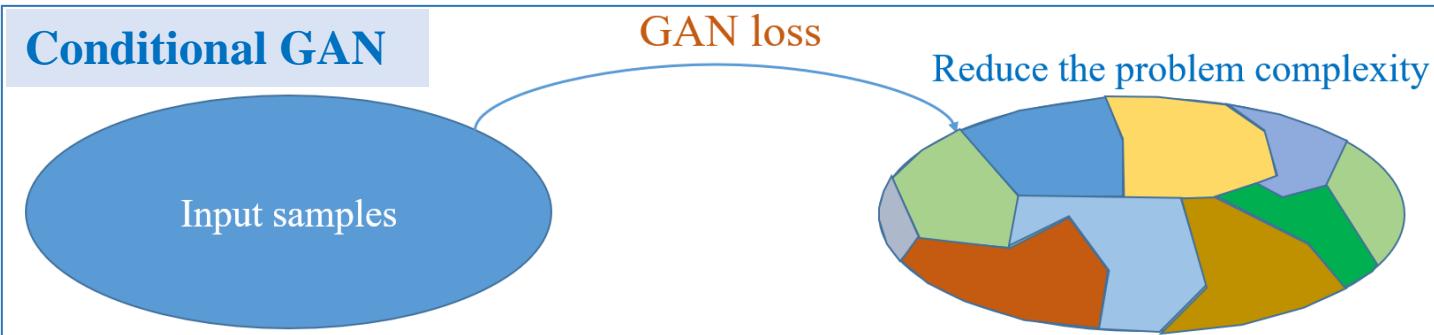
## Toward to CycleGAN

Quang-Vinh Dinh  
Ph.D. in Computer Science

# Objectives

- ✓ Study conditional GAN
- ✓ Study Pix2pix

- ✓ Study PatchGAN
- ✓ Study CycleGAN

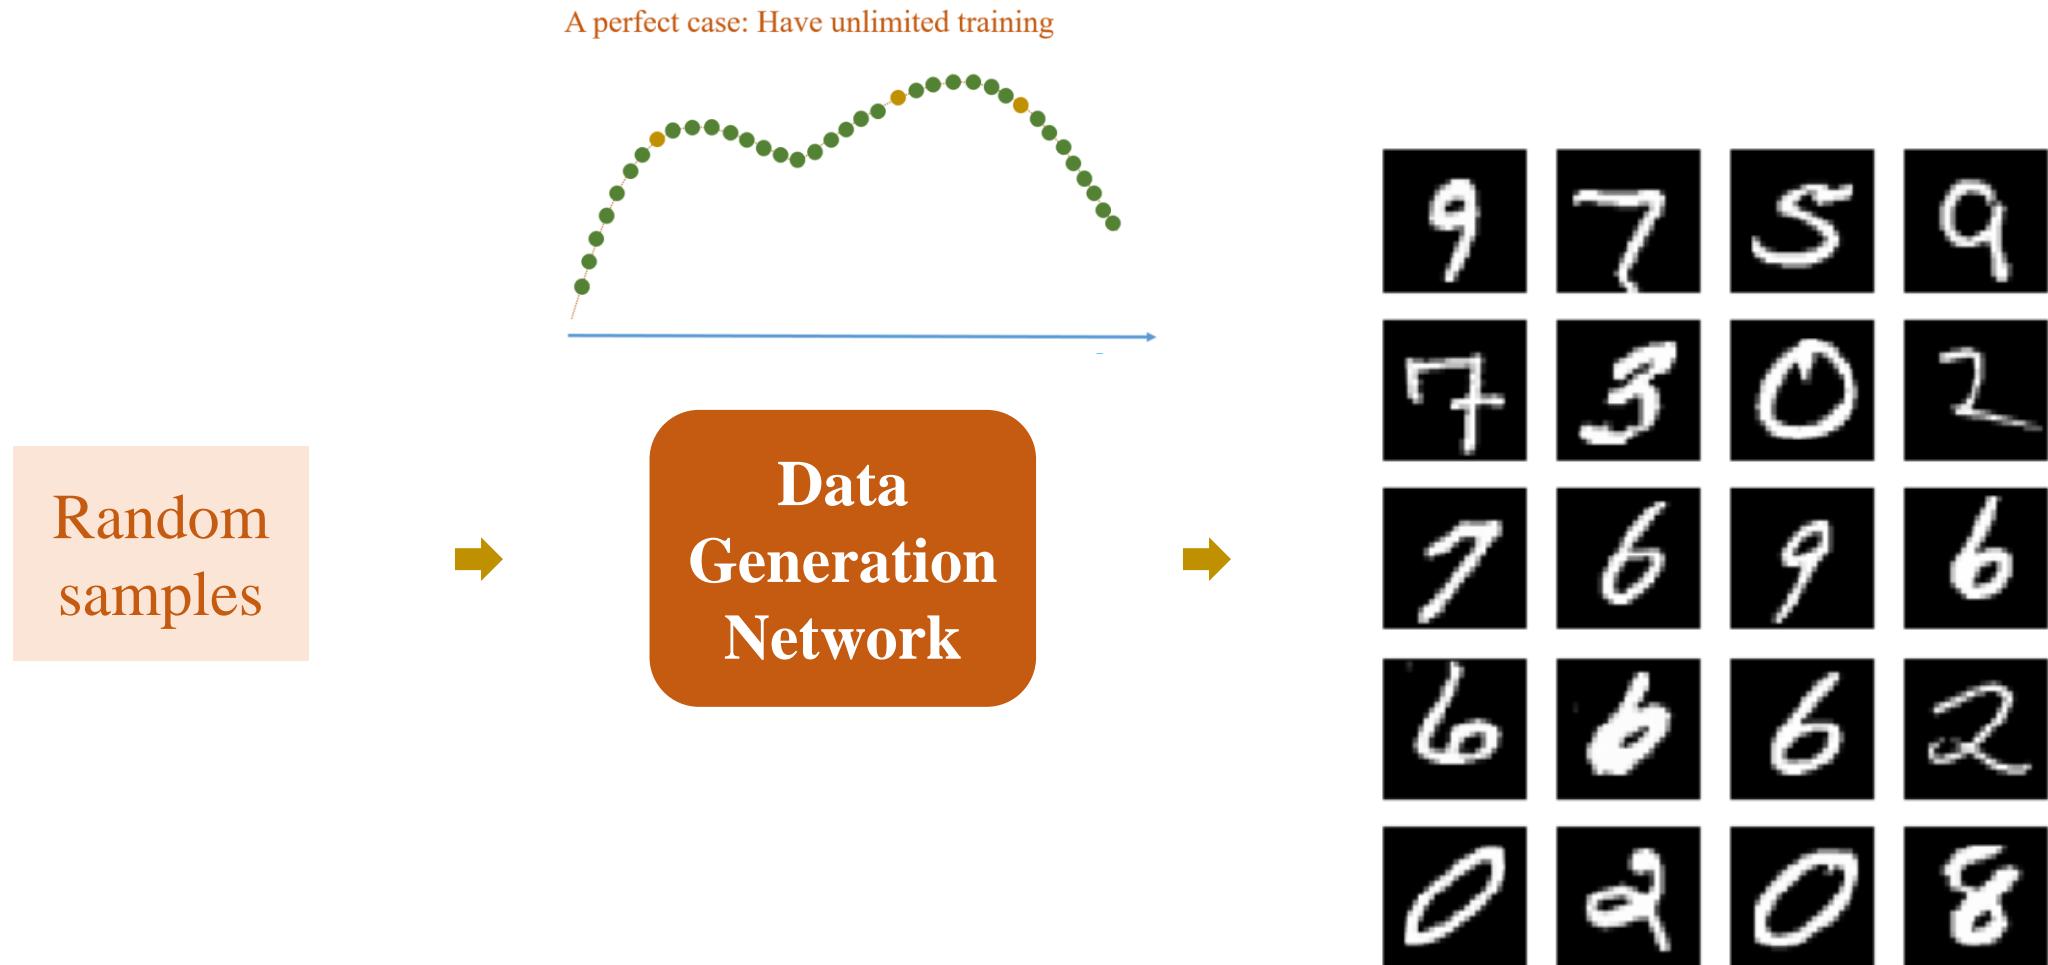


# Outline

- Quick Review on GAN/DCGAN
- Conditional GAN
- PatchGAN
- Pix2Pix
- CycleGAN

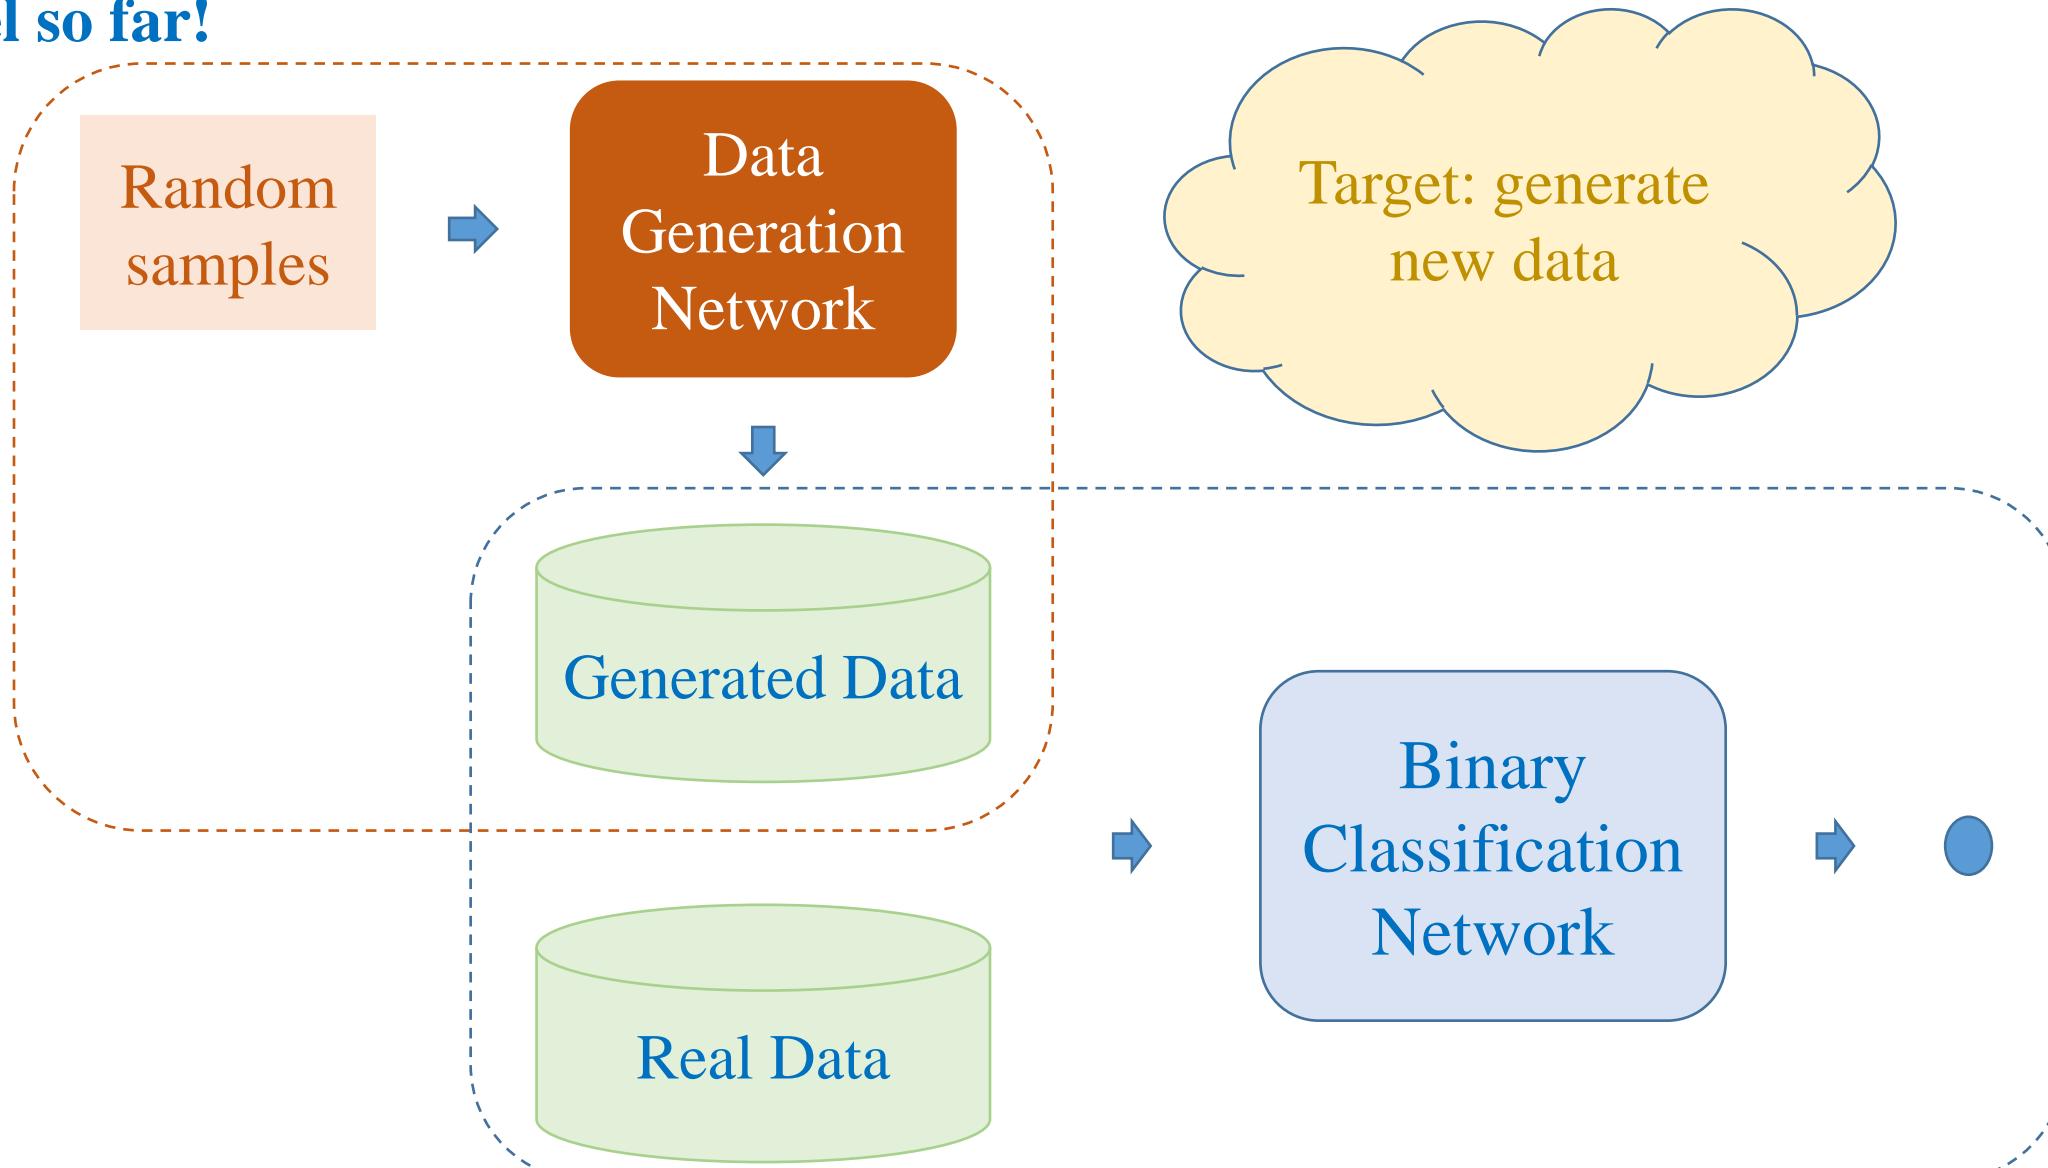
# Generative Adversarial Networks

## ❖ Input?



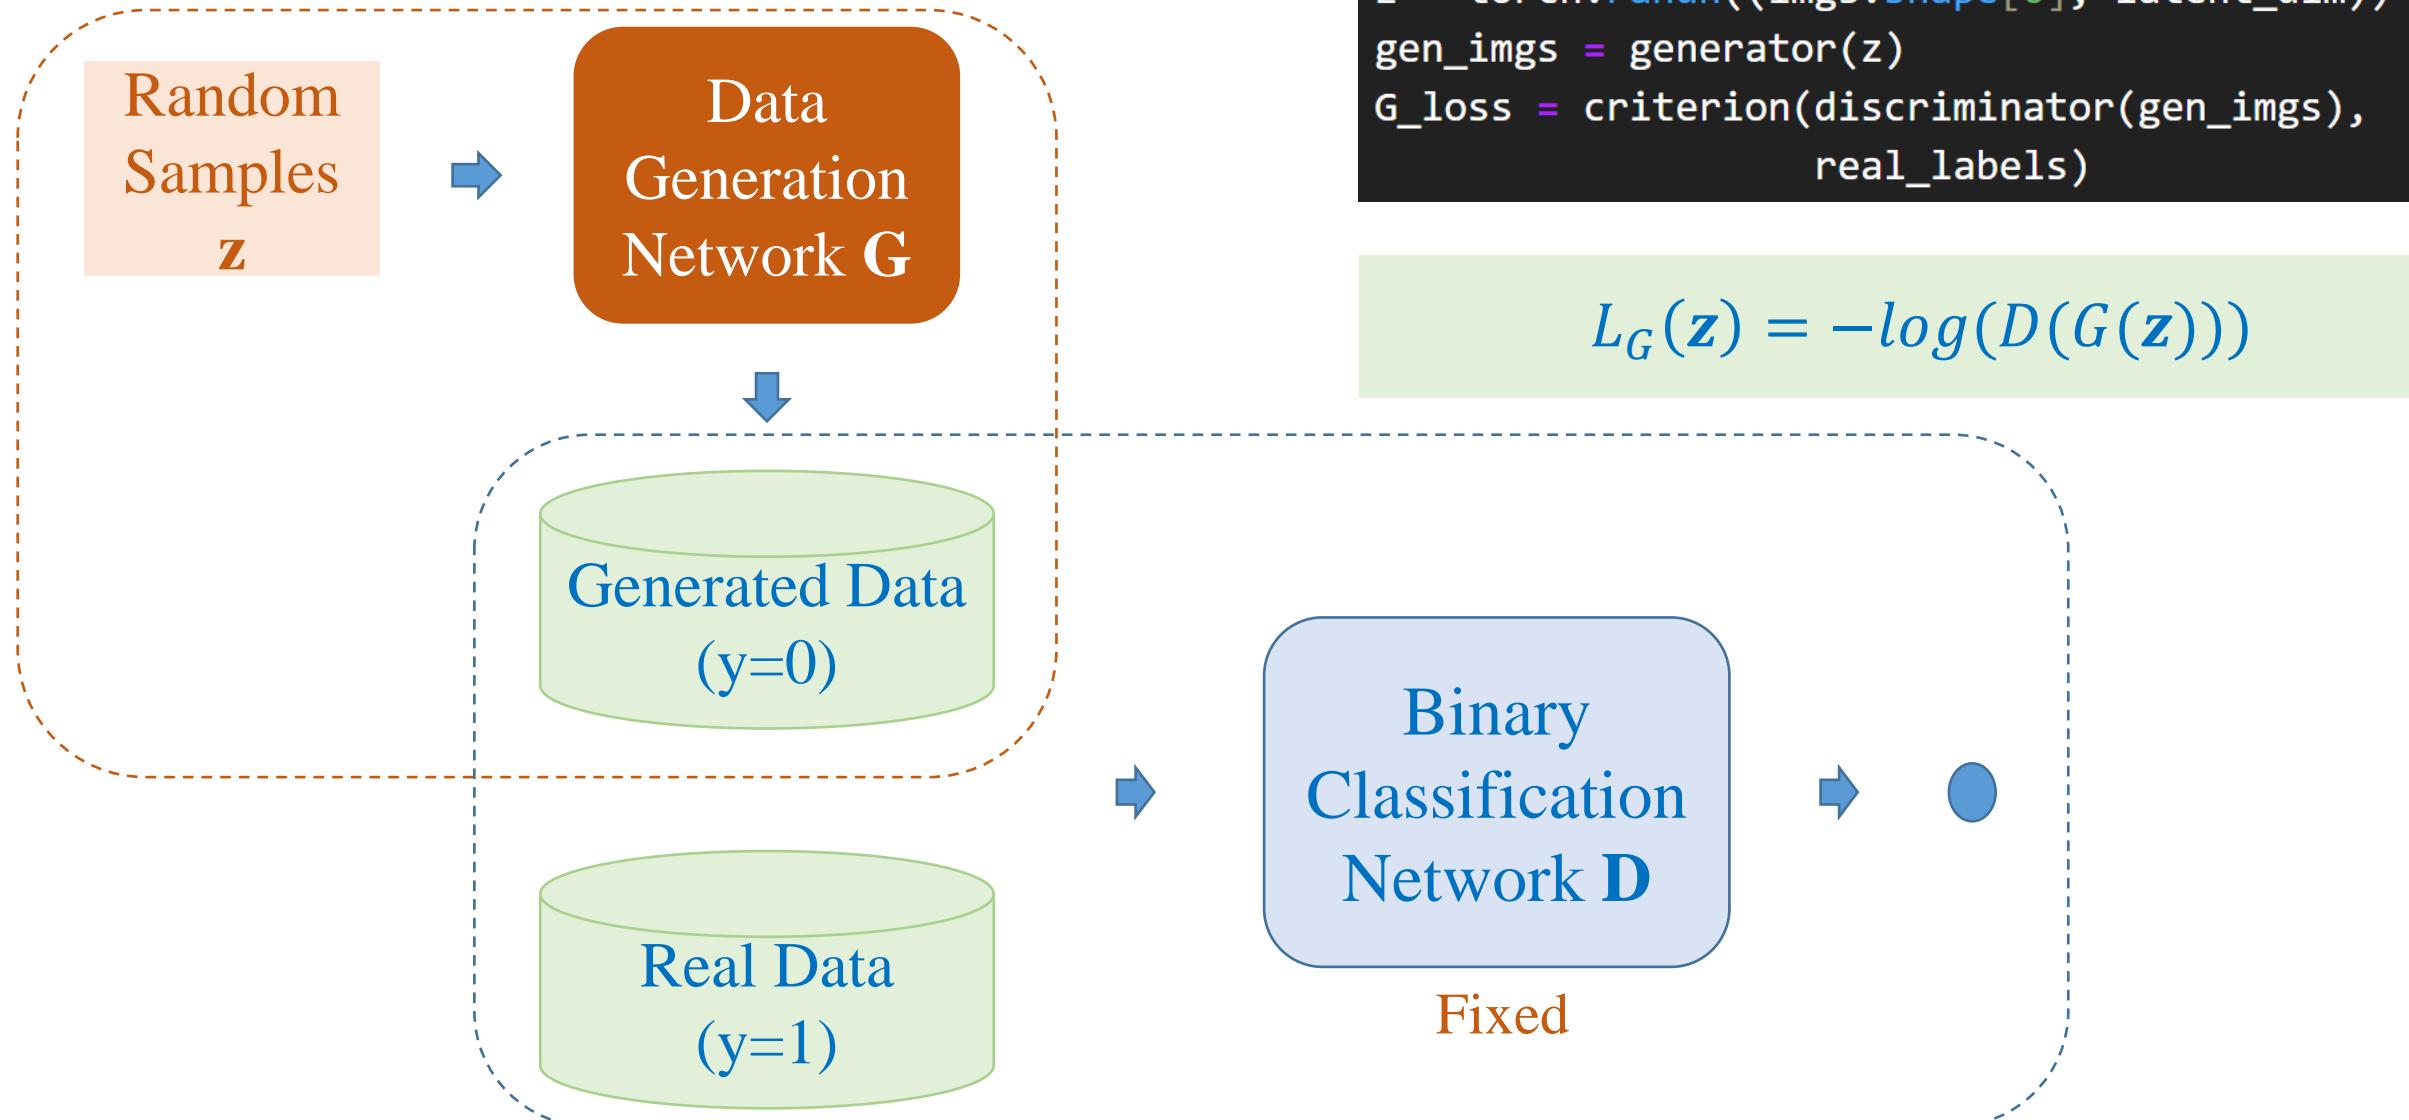
# Generative Adversarial Networks

Model so far!



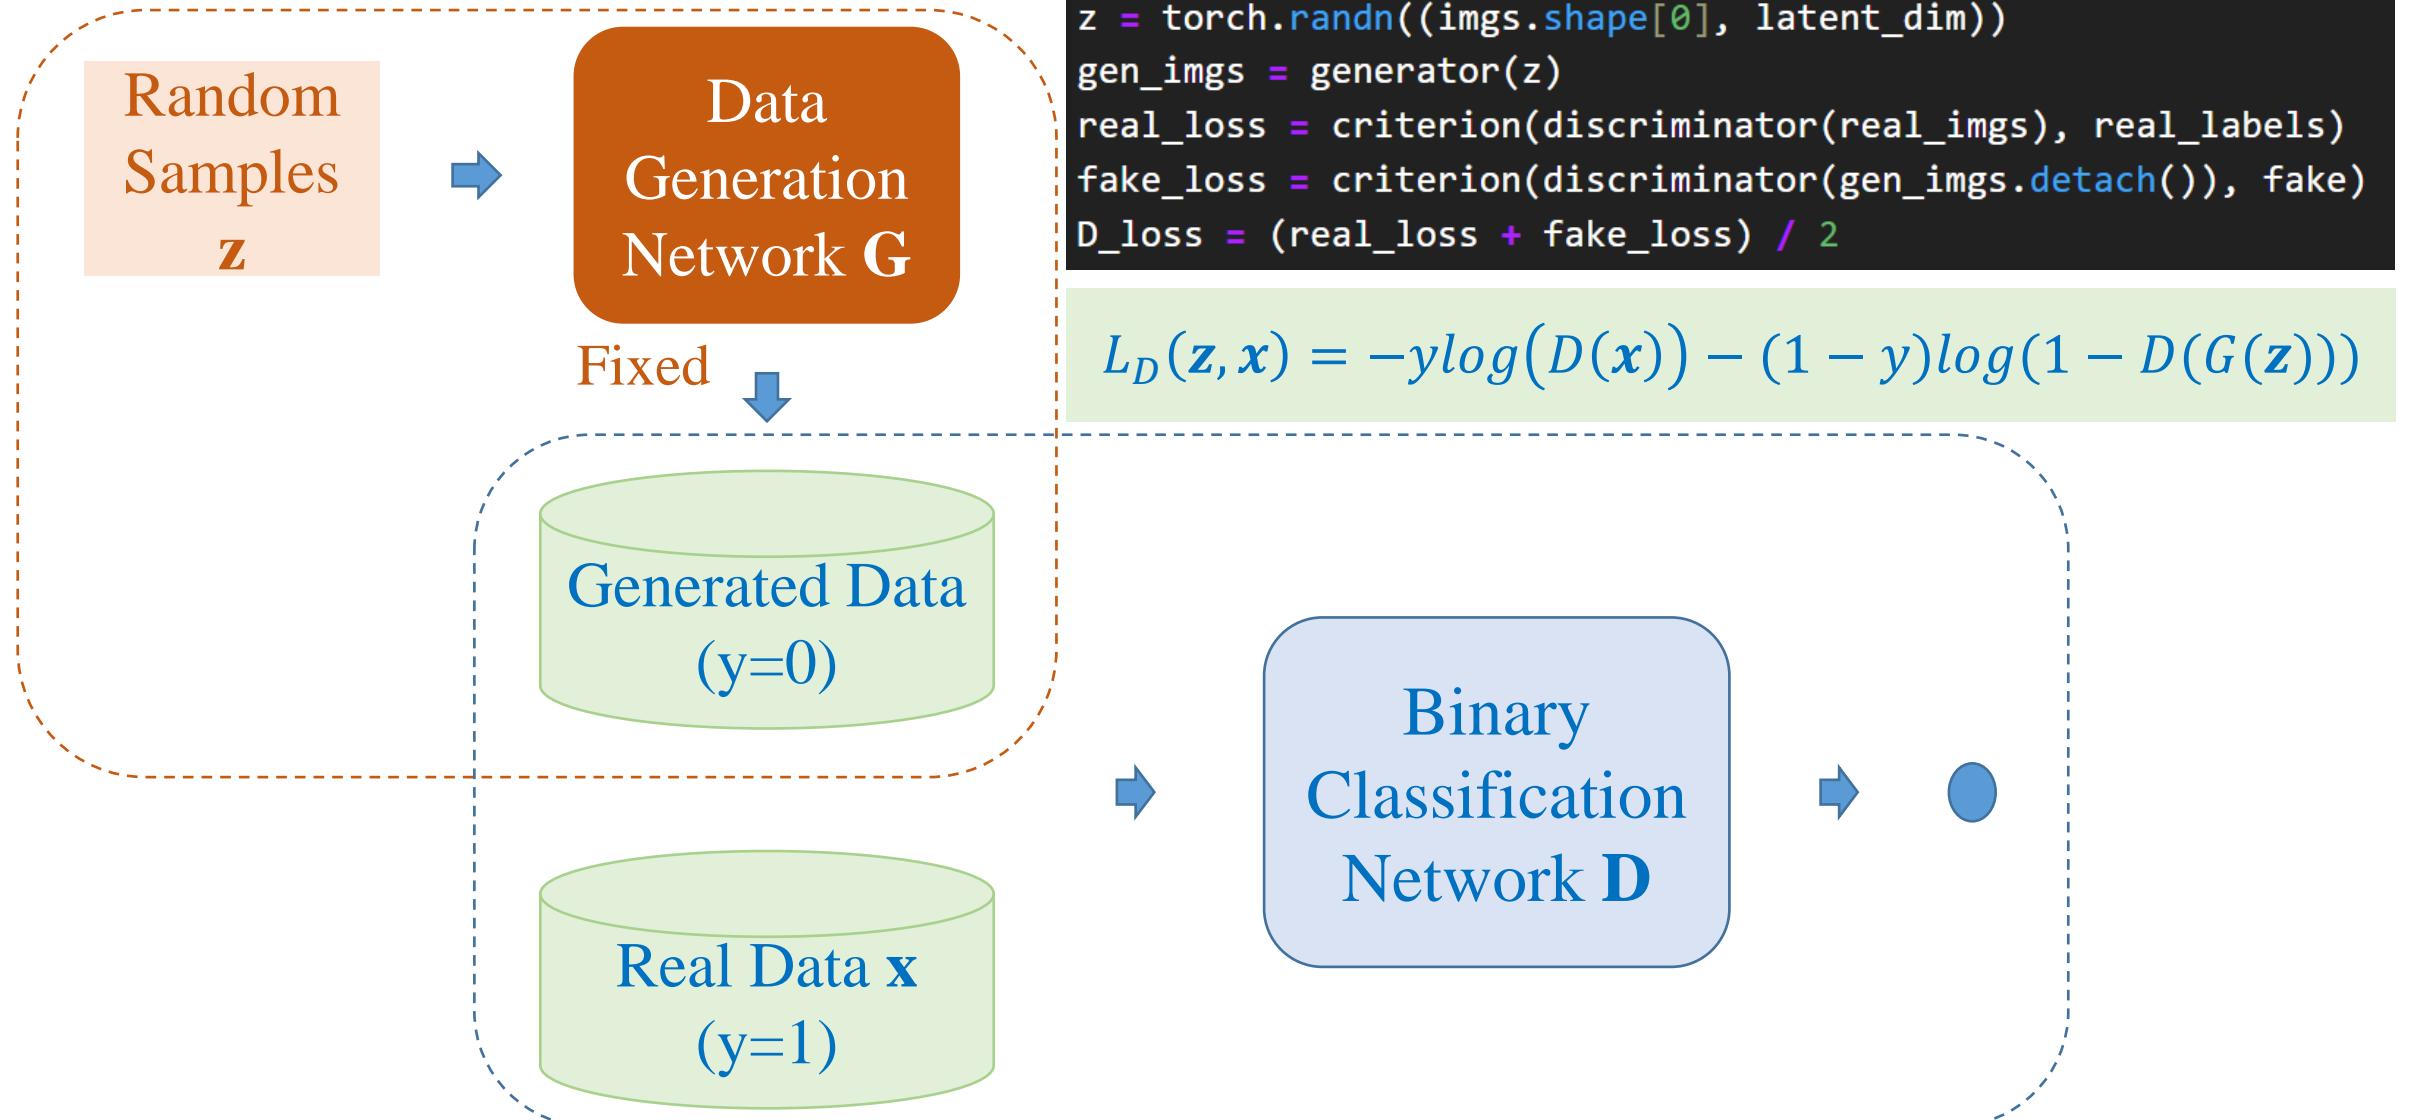
# Generator Loss

## Model so far!



# Discriminator Loss

Model so far!

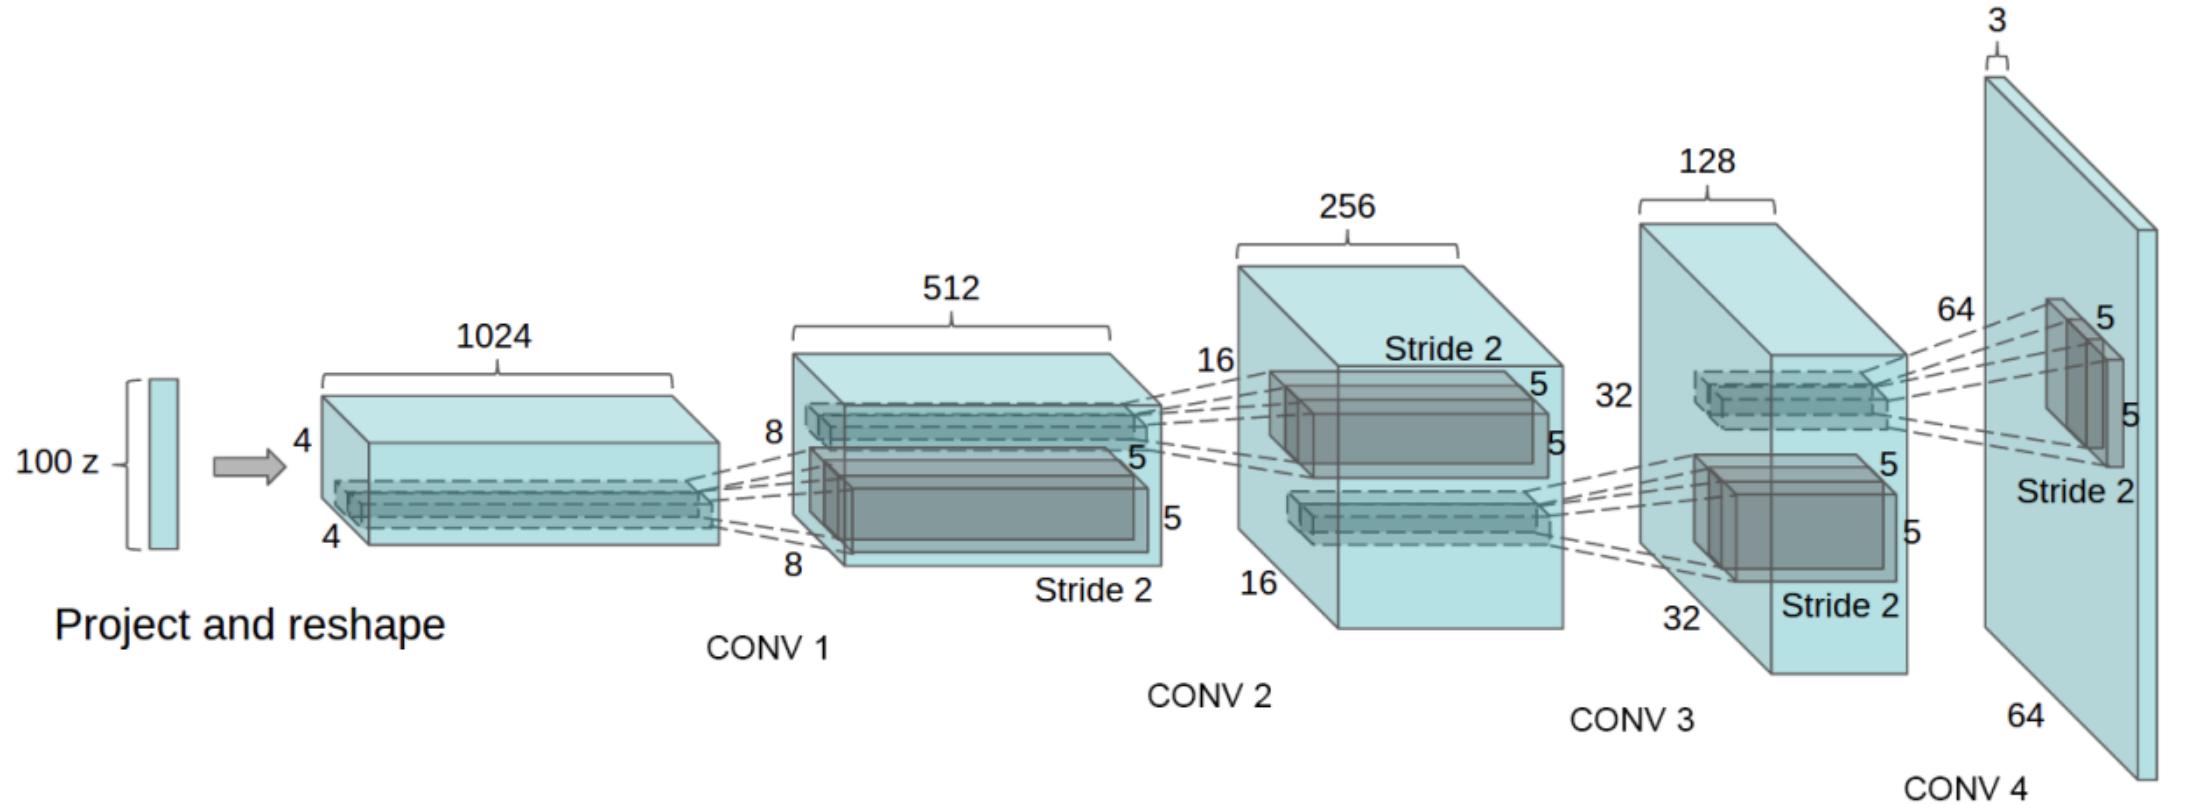


# Review

## ❖ DCGANs (2015)

### Architecture guidelines for stable Deep Convolutional GANs

- Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- Use batchnorm in both the generator and the discriminator.
- Remove fully connected hidden layers for deeper architectures.
- Use ReLU activation in generator for all layers except for the output, which uses Tanh.
- Use LeakyReLU activation in the discriminator for all layers.



# Review

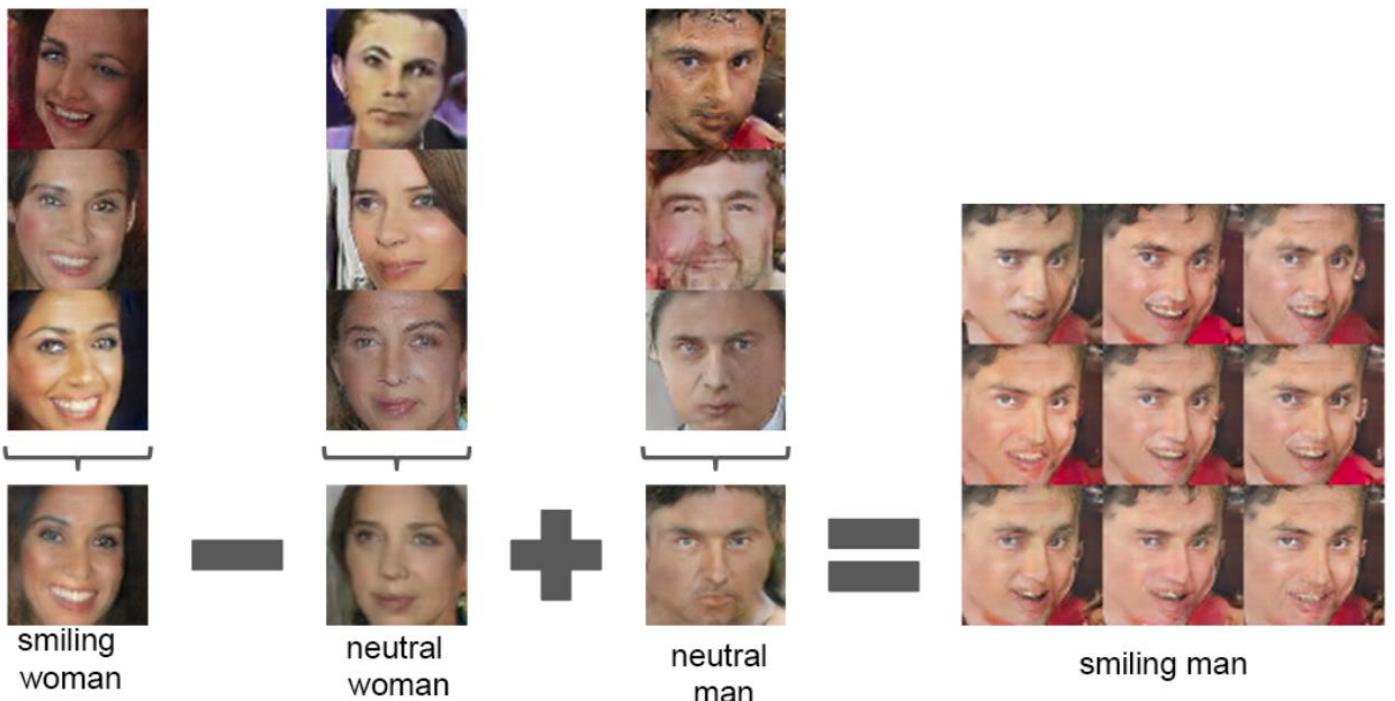
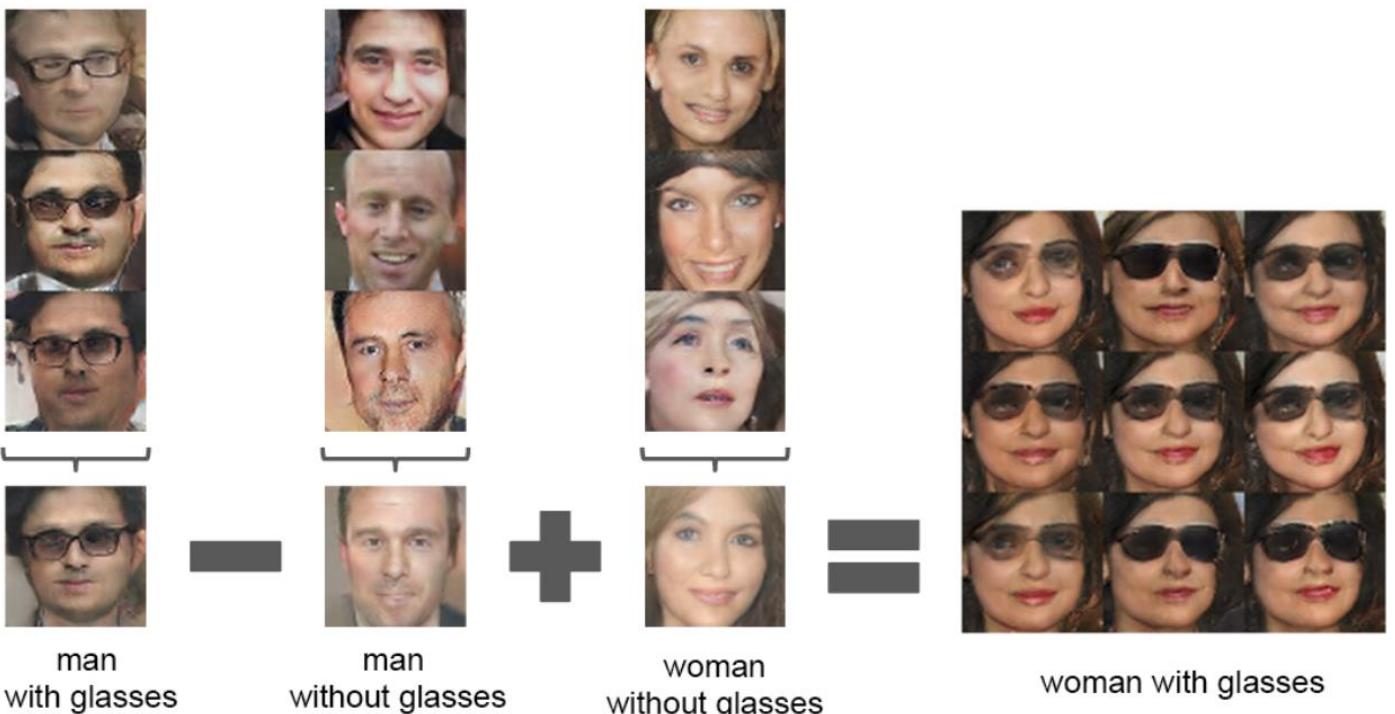
## ❖ DCGANs (2015)



# Review

## ❖ DCGANs (2015)

Unsupervised Representation  
Learning with Deep Convolutional  
Generative Adversarial Networks

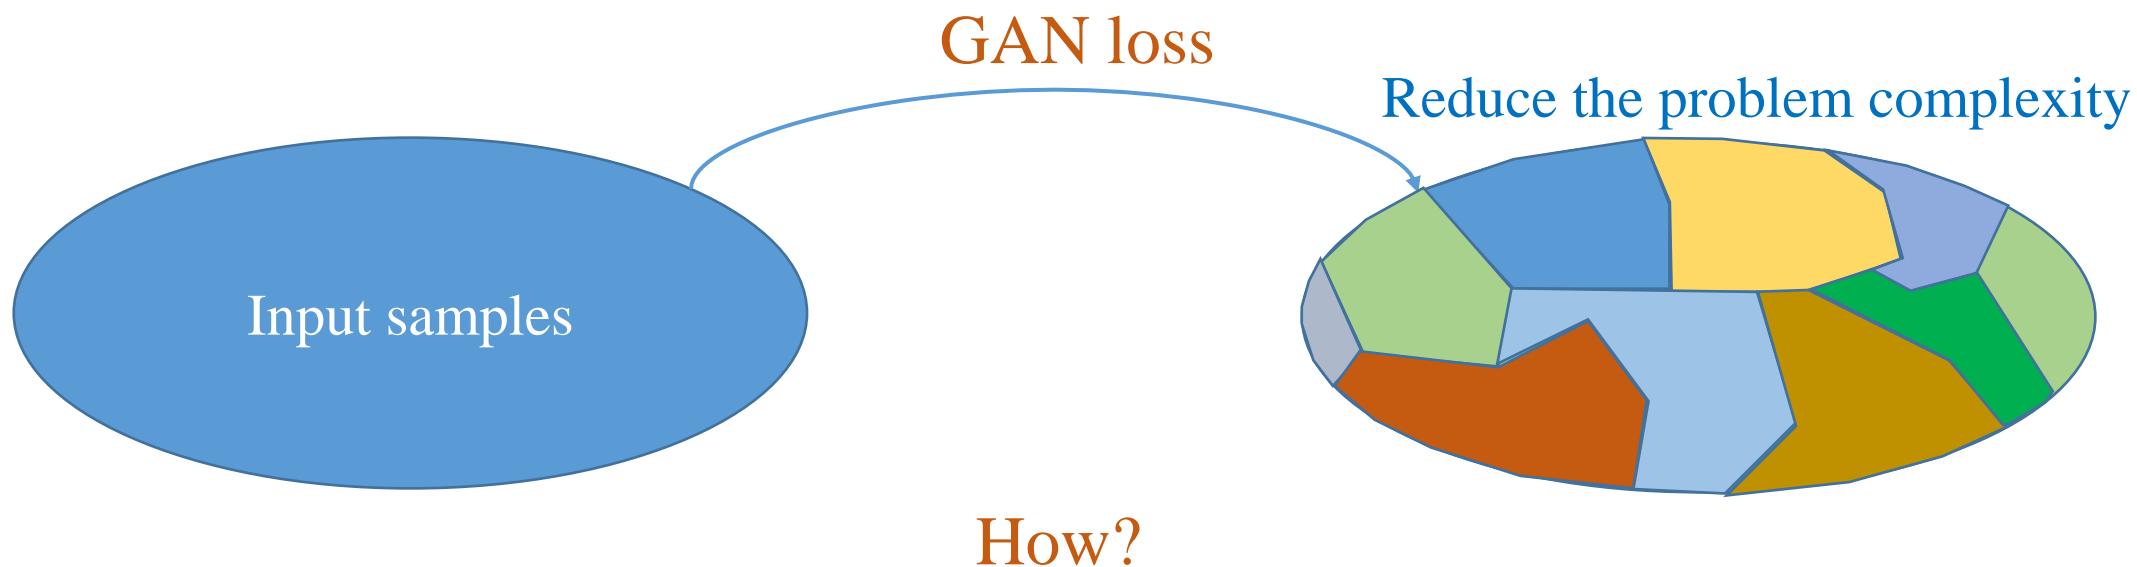
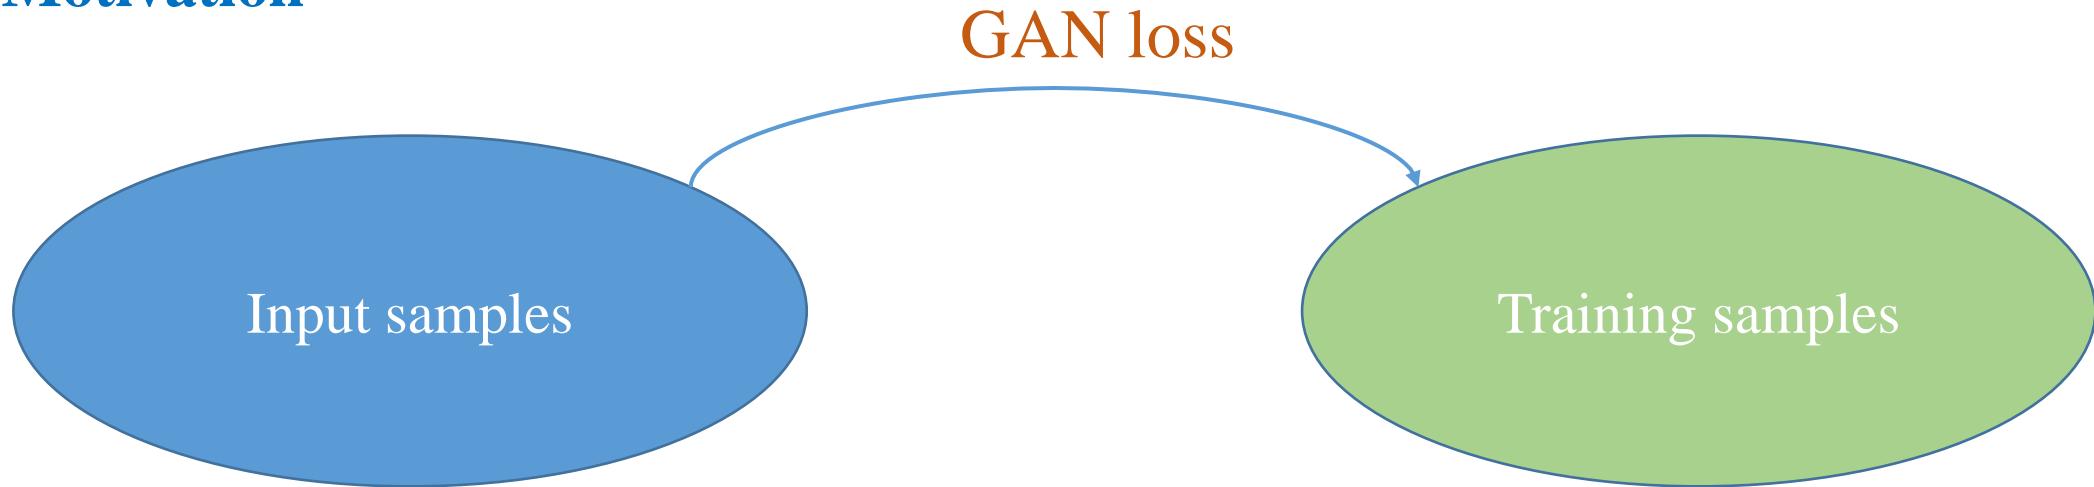


# Outline

- Quick Review on GAN/DCGAN
- Conditional GAN
- PatchGAN
- Pix2Pix
- CycleGAN

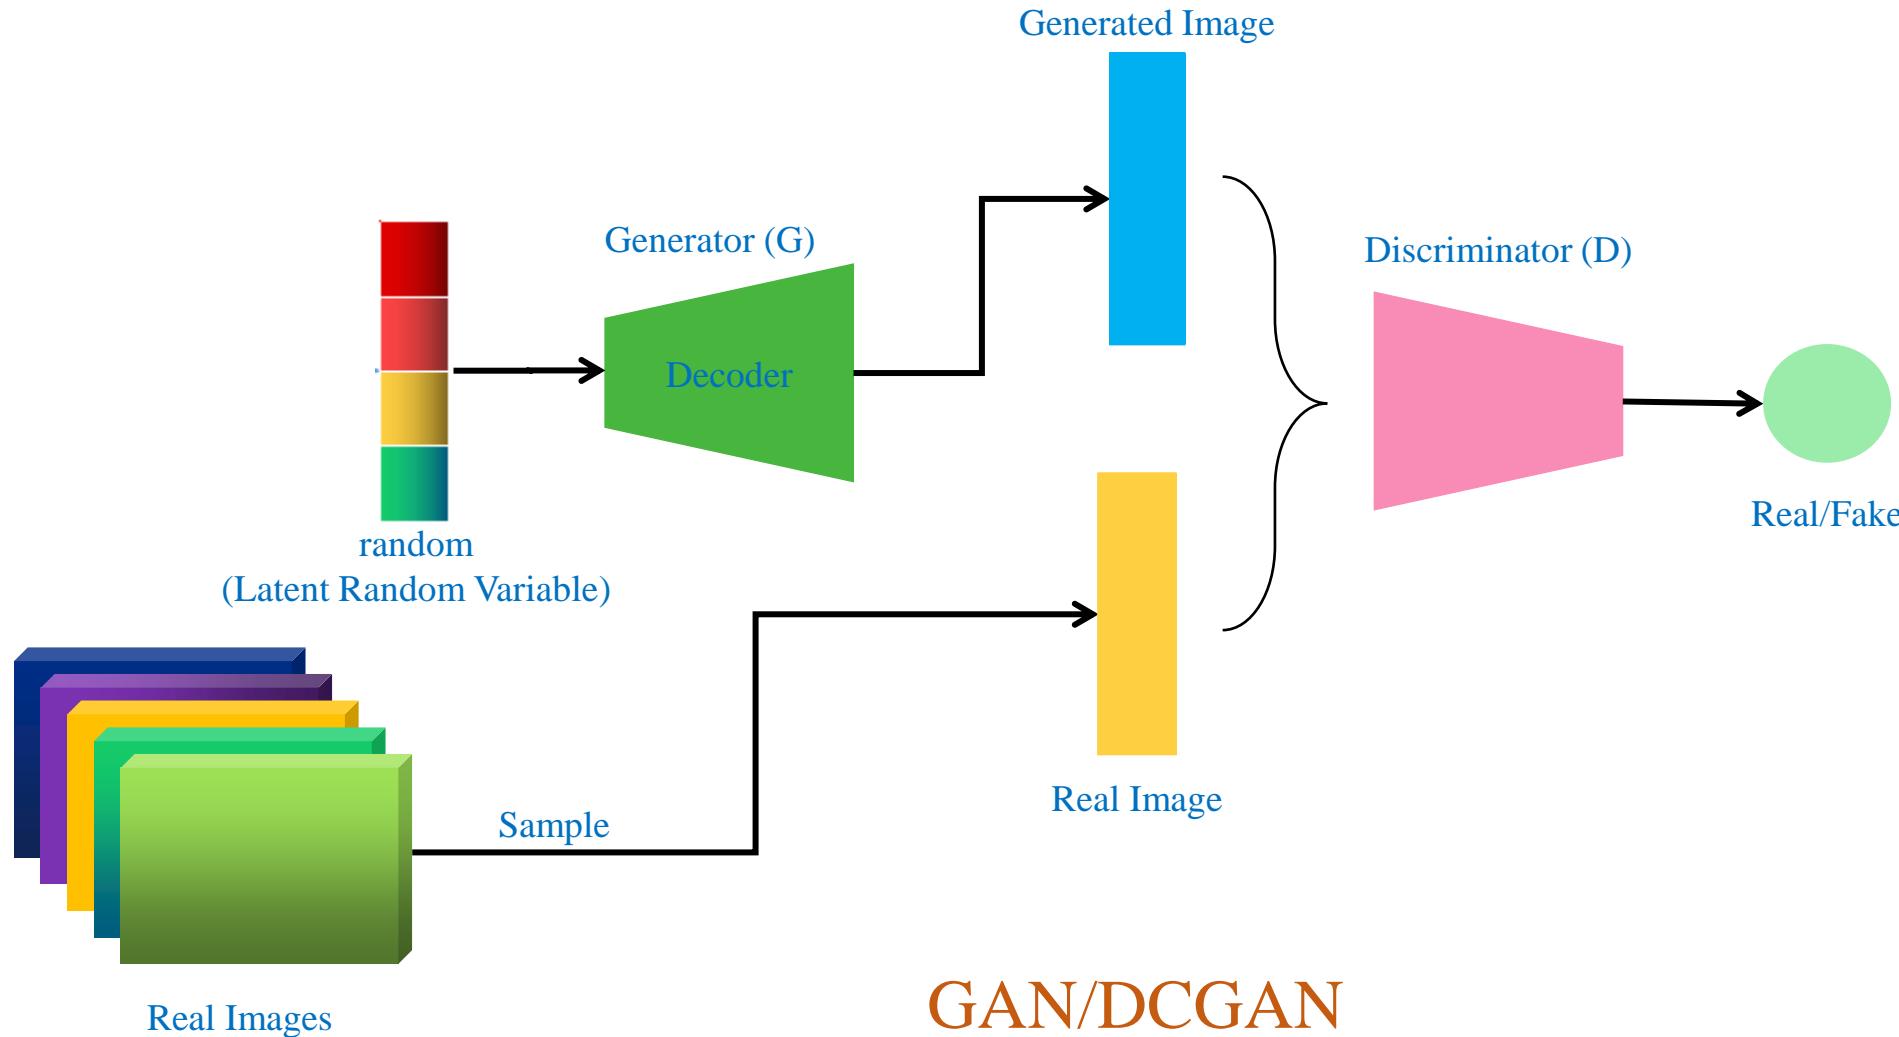
# Conditional GAN (cGAN)

## ❖ Motivation



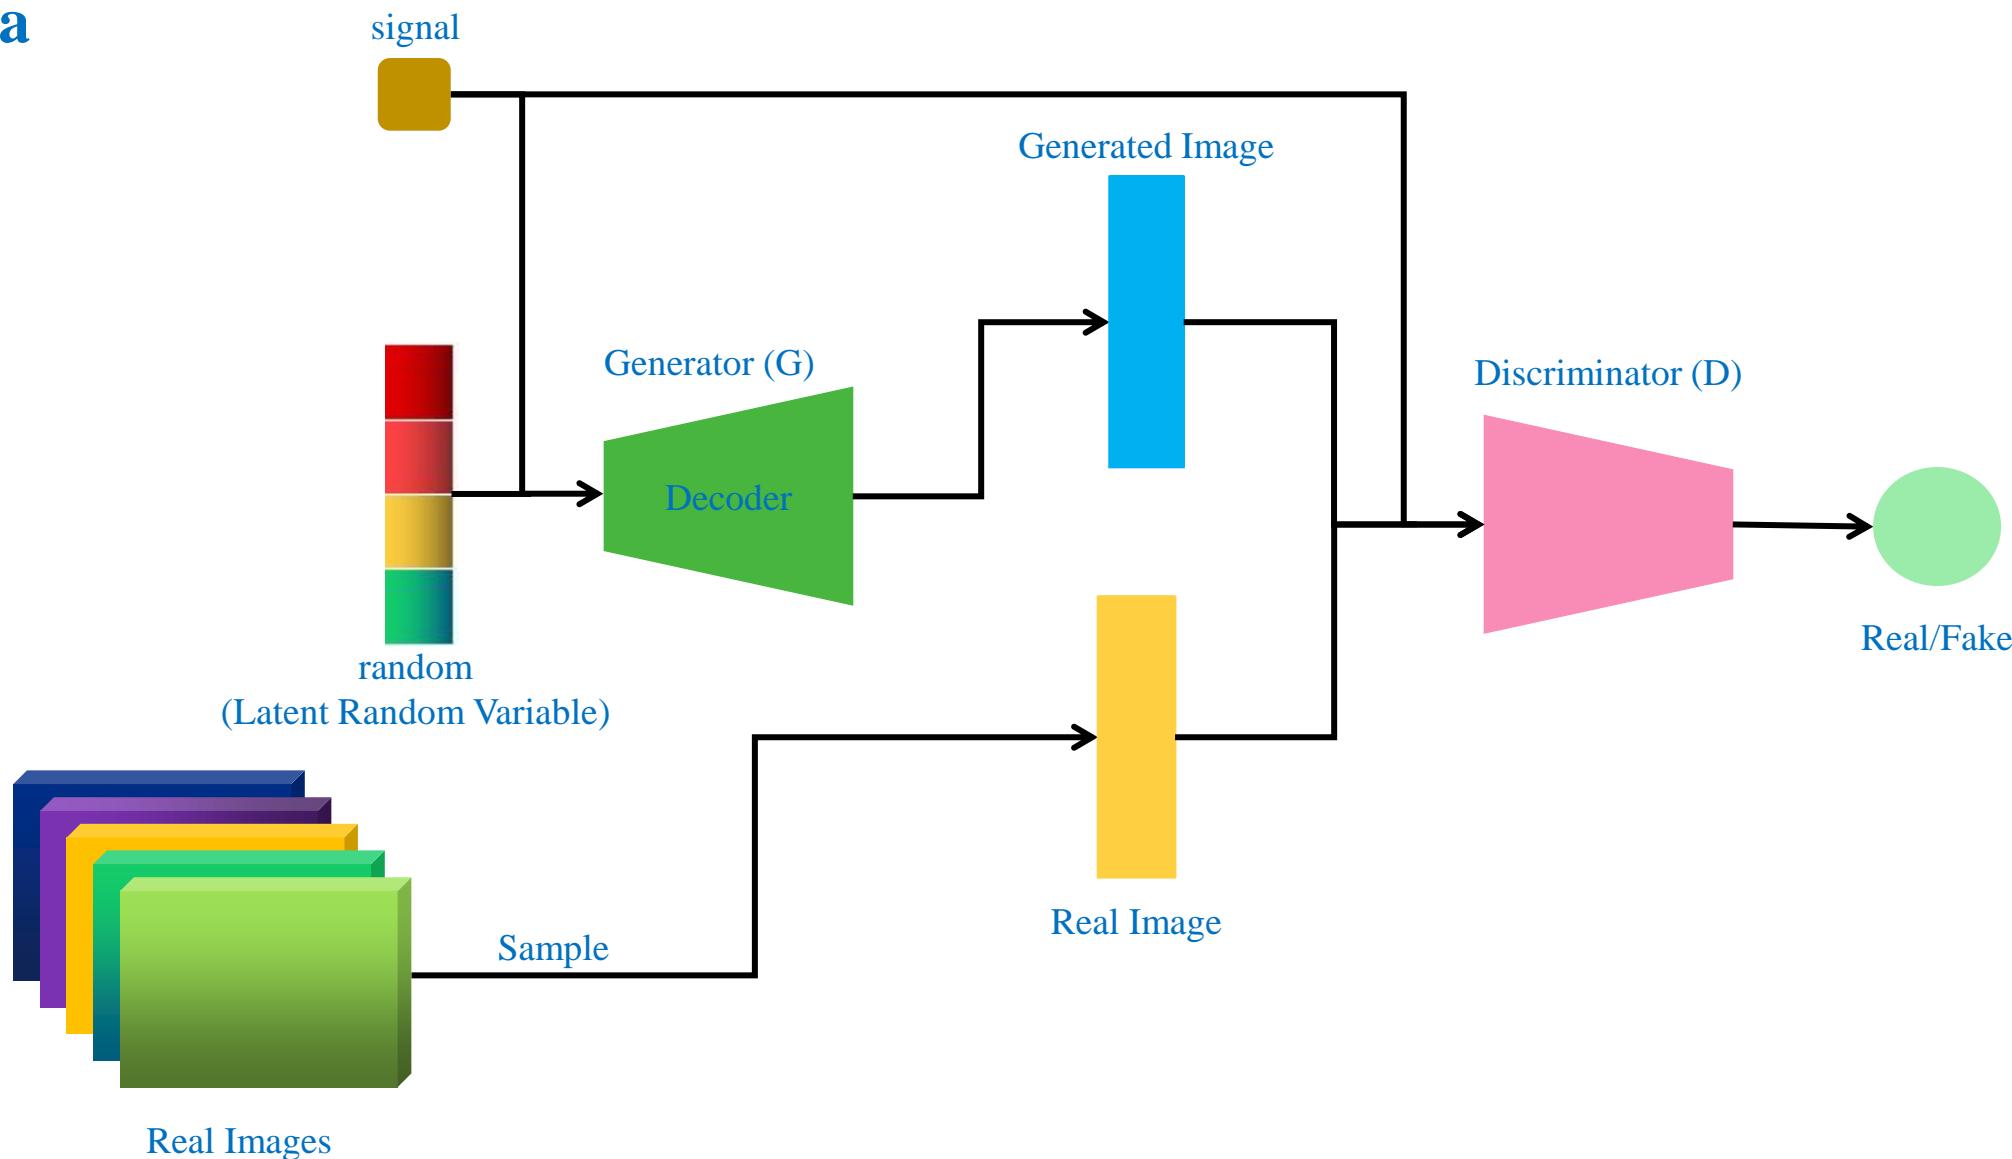
# Conditional GAN (cGAN)

❖ Idea?



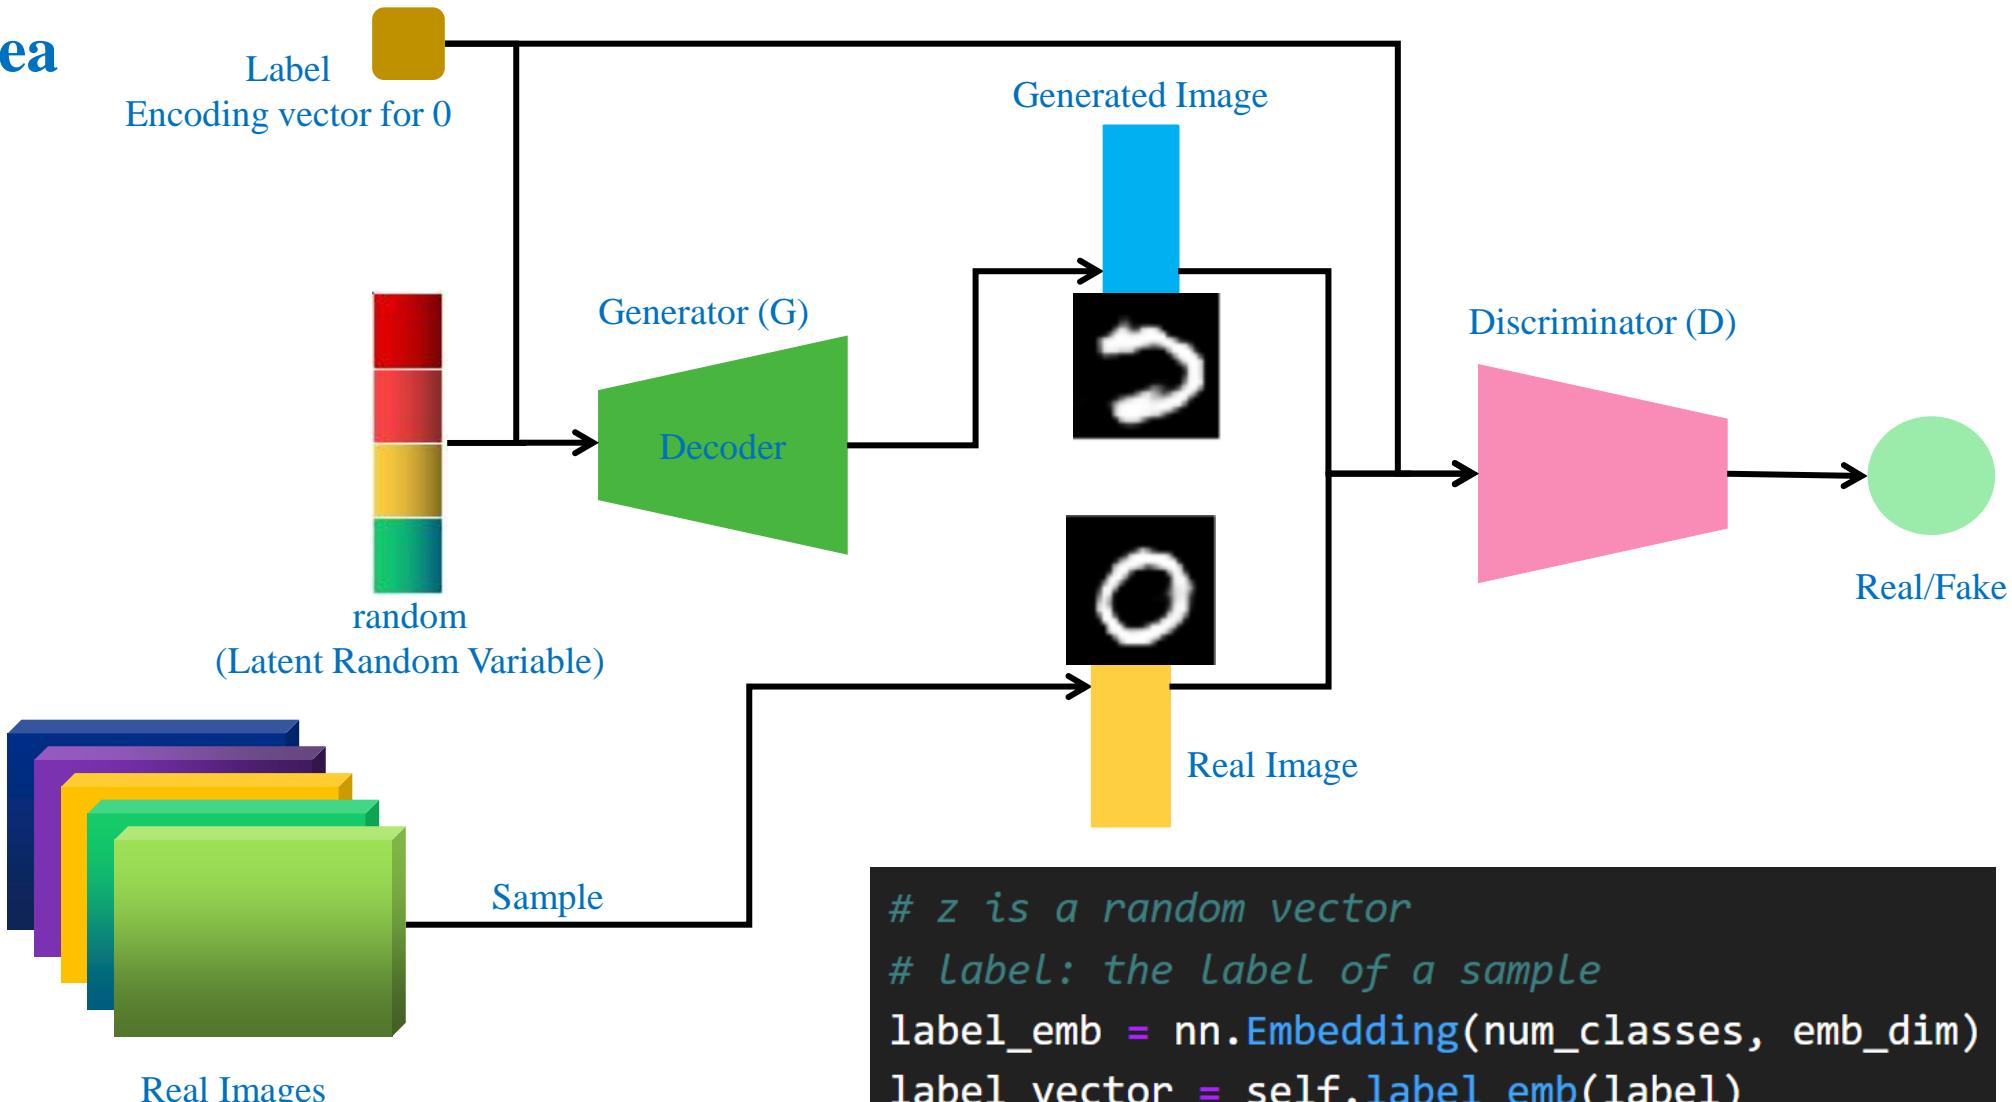
# Conditional GAN (cGAN)

## ❖ Idea



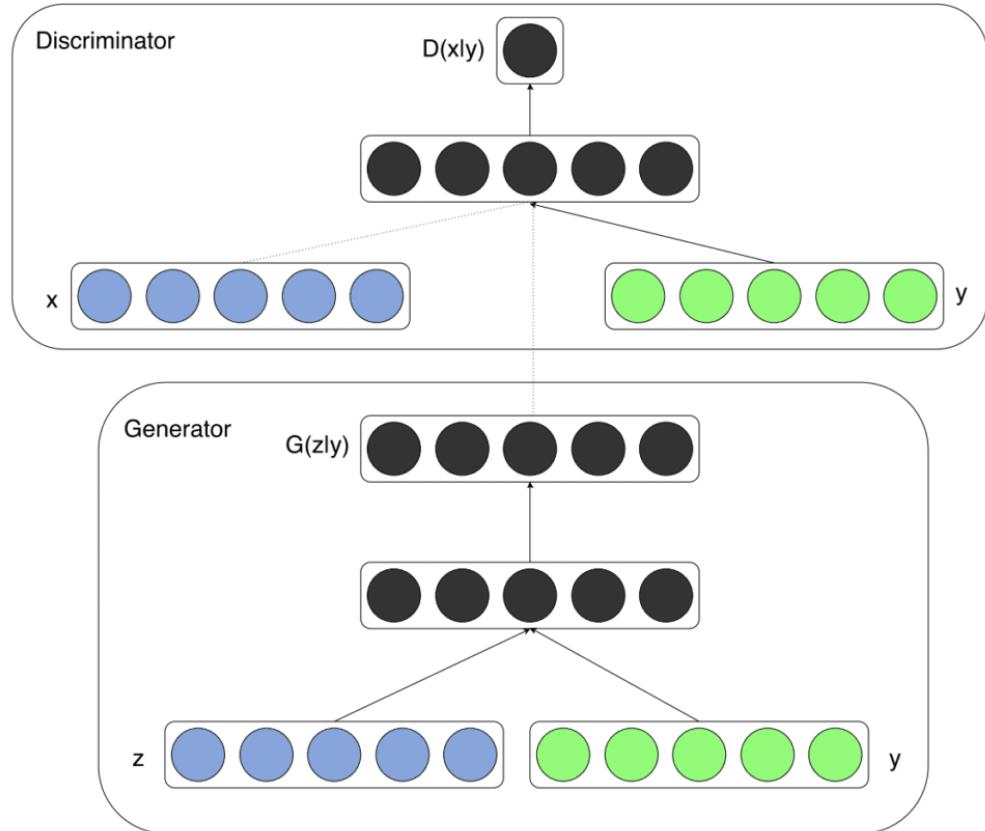
# Conditional GAN (cGAN)

## ❖ Idea



# Conditional GAN

## ❖ Model: Generator



$$L_G(z, y) = -\log(D(G(z|y)))$$

```
# cGAN - Generator Loss
# real_imgs
# Labels

# Loss and optimizer
criterion = nn.BCELoss()
optimizer_G = torch.optim.Adam(generator.parameters(),
                               lr=0.0001)

# prepare objective data
real_labels = torch.ones((imgs.shape[0], 1))

optimizer_G.zero_grad()

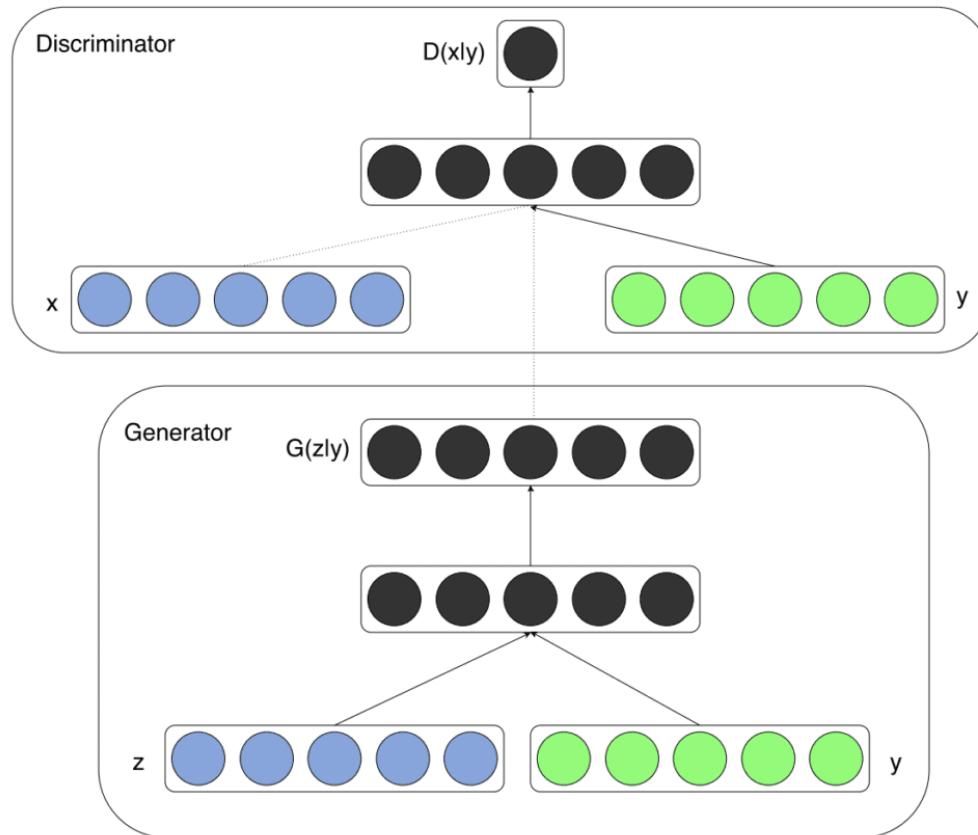
# random input for Generator
z = torch.randn((imgs.shape[0], latent_dim))

# train
gen_imgs = generator(z, labels)
validity = discriminator(gen_imgs, labels)
G_loss = criterion(validity, real_labels)

# optimize
G_loss.backward()
optimizer_G.step()
```

# Conditional GAN

## ❖ Model: Discriminator



$$L_D(z, x, y) = -y \log(D(x|y)) - (1 - y) \log(1 - D(G(z|y)))$$

```
# cGAN - Discriminator Loss
# real_imgs
# Labels

# Loss and optimizer
criterion = nn.BCELoss()
optimizer_D = torch.optim.Adam(discriminator.parameters(),
                               lr=0.0002)

# prepare objective data
real_labels = torch.ones((imgs.shape[0], 1))
fake_labels = torch.zeros((imgs.shape[0], 1))
optimizer_D.zero_grad()

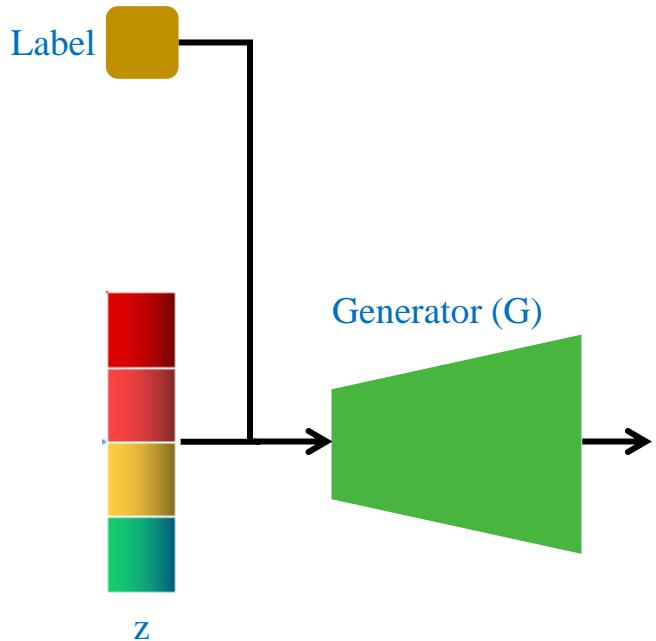
# random input for Generator
z = torch.randn((imgs.shape[0], latent_dim))

# train
real_validity = discriminator(real_imgs, labels)
real_loss = criterion(real_validity, real_labels)
fake_validity = discriminator(gen_imgs.detach(), labels)
fake_loss = criterion(fake_validity, fake_labels)
D_loss = (real_loss + fake_loss) / 2

# optimize
D_loss.backward()
optimizer_D.step()
```

# Conditional GAN (cGAN)

## ❖ Implementation



```
class Generator(nn.Module):
    def __init__(self, num_classes, emb_dim):
        super().__init__()
        self.label_emb = nn.Embedding(num_classes, emb_dim)

        self.model = nn.Sequential(
            nn.Linear(100+emb_dim, 256),
            nn.BatchNorm1d(256),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Linear(256, 512),
            nn.BatchNorm1d(512),
            nn.LeakyReLU(0.2, inplace=True),

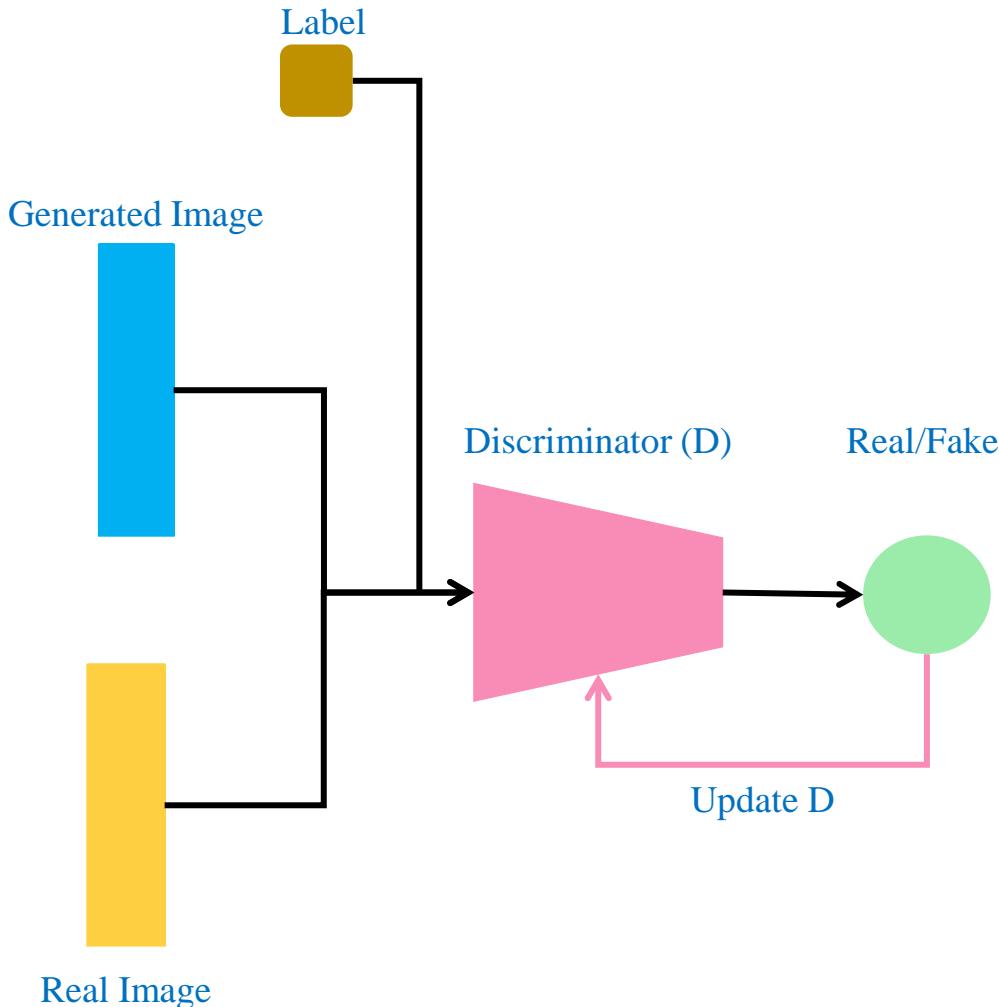
            nn.Linear(512, 1024),
            nn.BatchNorm1d(1024),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Linear(1024, 32*32),
            nn.Tanh())

    def forward(self, z, label):
        cond = self.label_emb(label)
        x = torch.cat([z, cond], 1)
        img = self.model(x)
        img = img.view(img.size(0), *img_shape)
        return img
```

# Conditional GAN (cGAN)

## ❖ Implementation



```
class Descriminator(nn.Module):
    def __init__(self, num_classes, emb_dim):
        super().__init__()
        self.label_emb = nn.Embedding(num_classes, emb_dim)

        self.model = nn.Sequential(
            nn.Linear(32*32 + emb_dim, 512),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Linear(512, 256),
            nn.LeakyReLU(0.2, inplace=True),

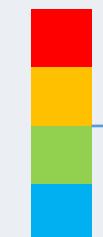
            nn.Linear(256, 1),
            nn.Sigmoid()
        )

    def forward(self, img, label):
        img_flat = img.view(img.size(0), -1)
        cond = self.label_emb(label)
        x = torch.cat([img_flat, cond], 1)
        validity = self.model(x)
        return validity
```

Train Generator

$$L_G(\mathbf{z}, \mathbf{y}) = -\log(D(G(\mathbf{z}|\mathbf{y})))$$

Random noise



Label



Generator (G)

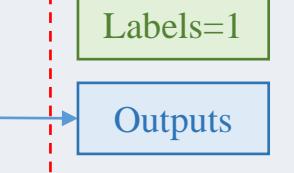
Generated Image

Update

G\_Optimizer

Loss Function

Discriminator (D)



$$L_D(\mathbf{z}, \mathbf{x}, \mathbf{y}) = -y\log(D(\mathbf{x}|\mathbf{y})) - (1-y)\log(1-D(G(\mathbf{z}|\mathbf{y})))$$

Random noise

Label



Generator (G)

Generated Image

Discriminator (D)

Labels=0

Outputs

Loss Function

Outputs

Labels=1

Loss Function

Update

D\_Optimizer

Fake\_label: 0

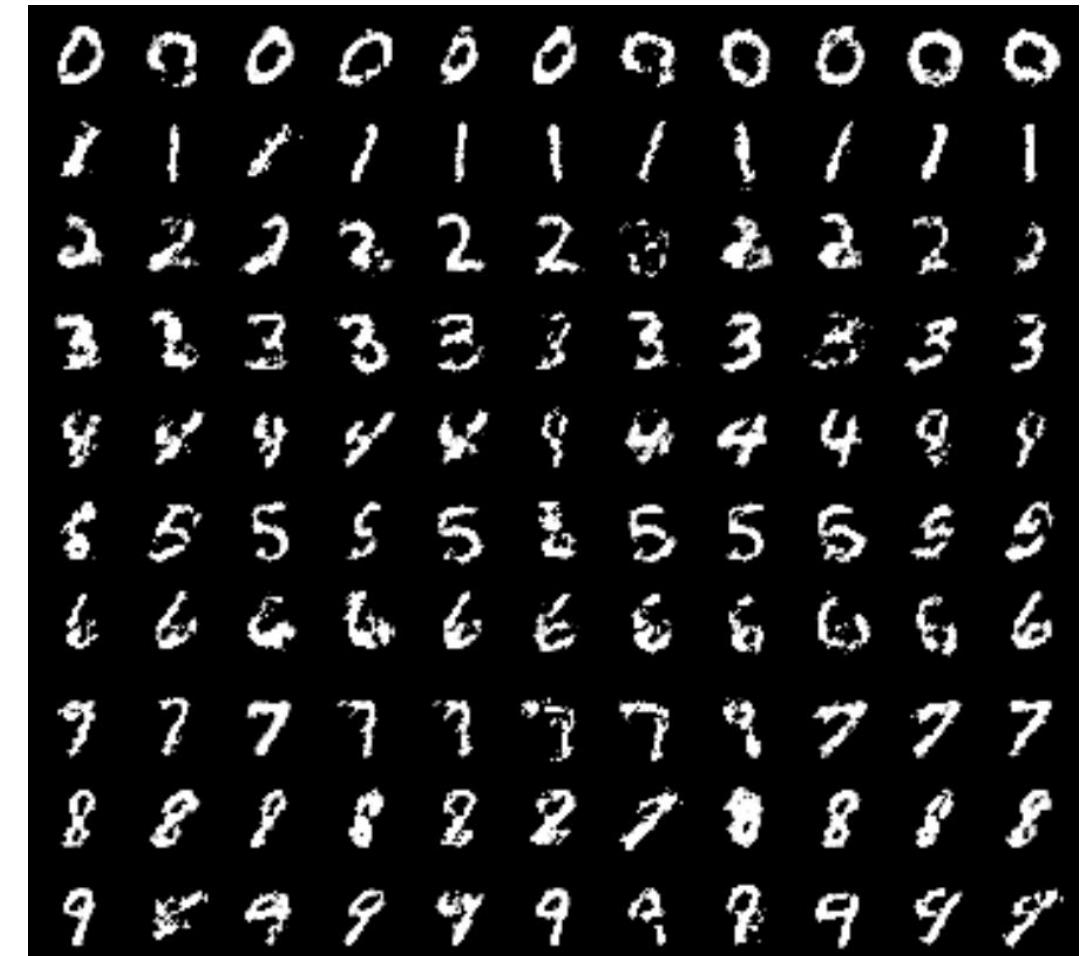
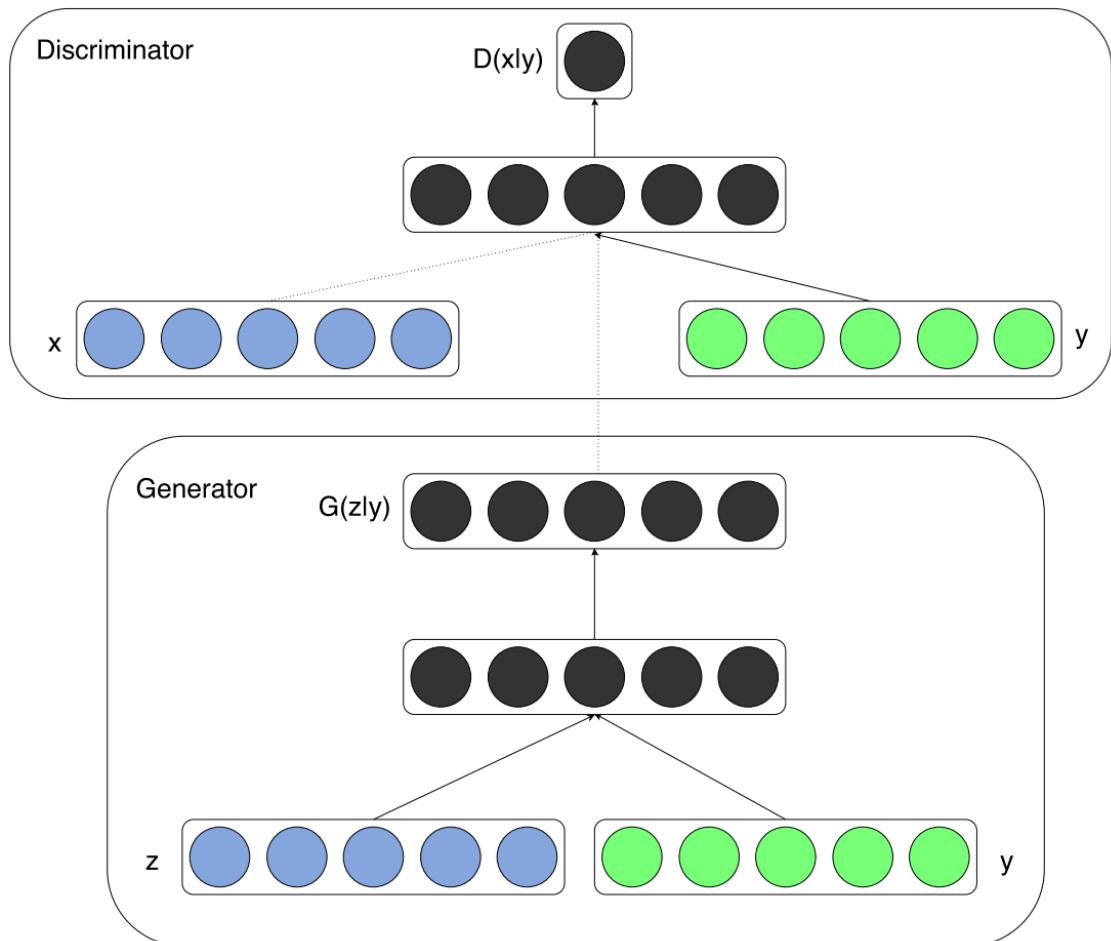
Real\_label: 1

Real Image

Train Discriminator

# GANs Survey

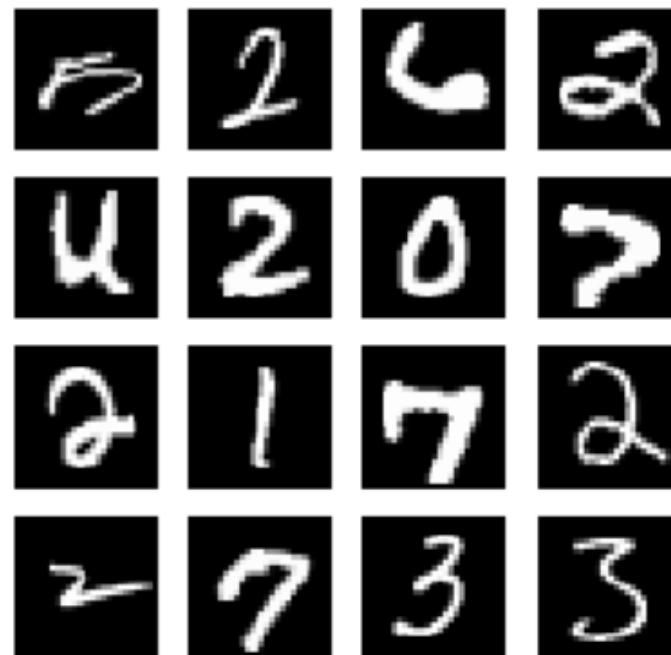
## ❖ Conditional Generative Adversarial Nets



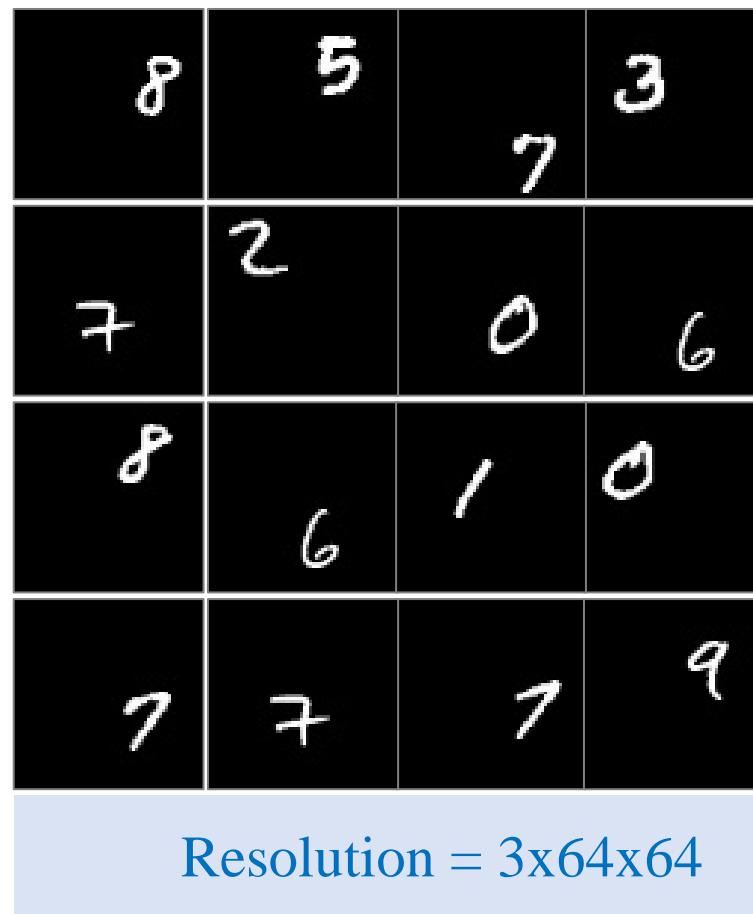
# Outline

- Quick Review on GAN/DCGAN
- Conditional GAN
- PatchGAN
- Pix2Pix
- CycleGAN

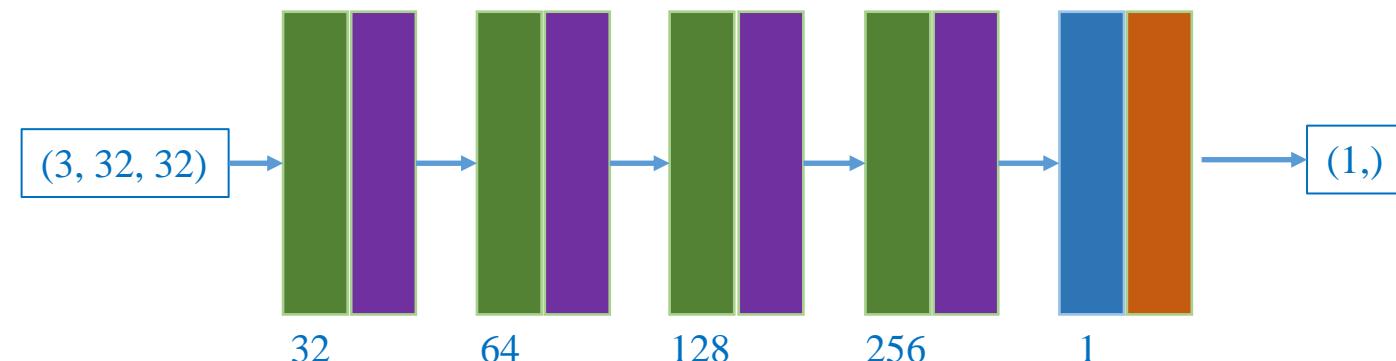
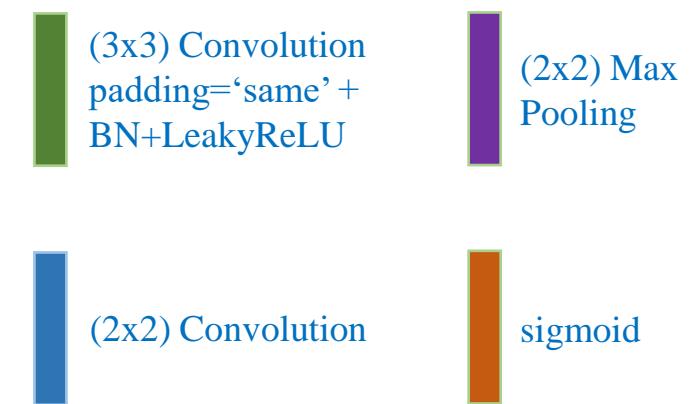
## ❖ From detection problem



Resolution = 3x32x32

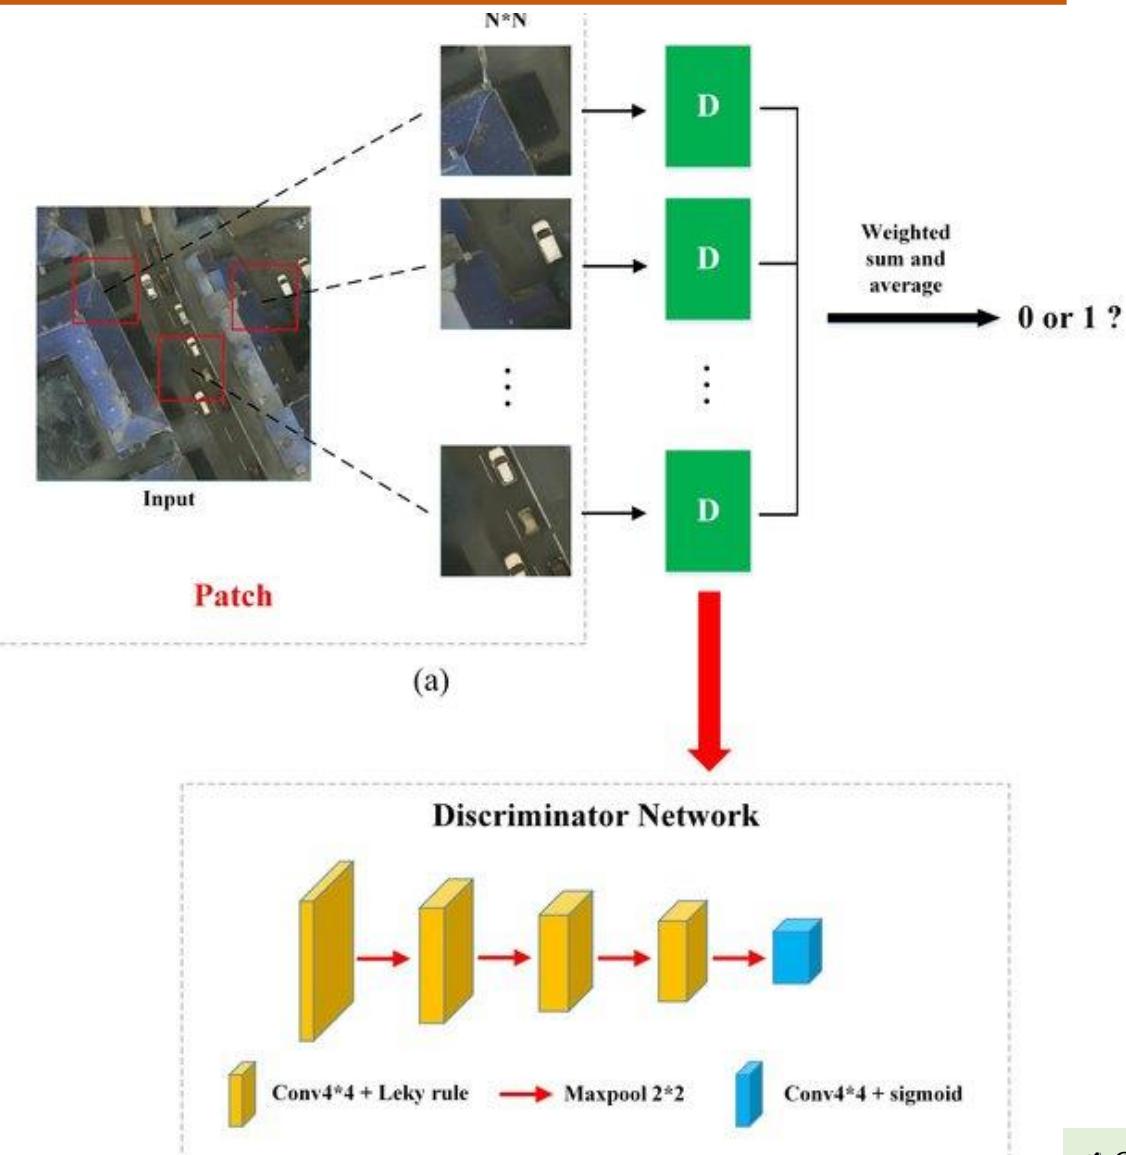
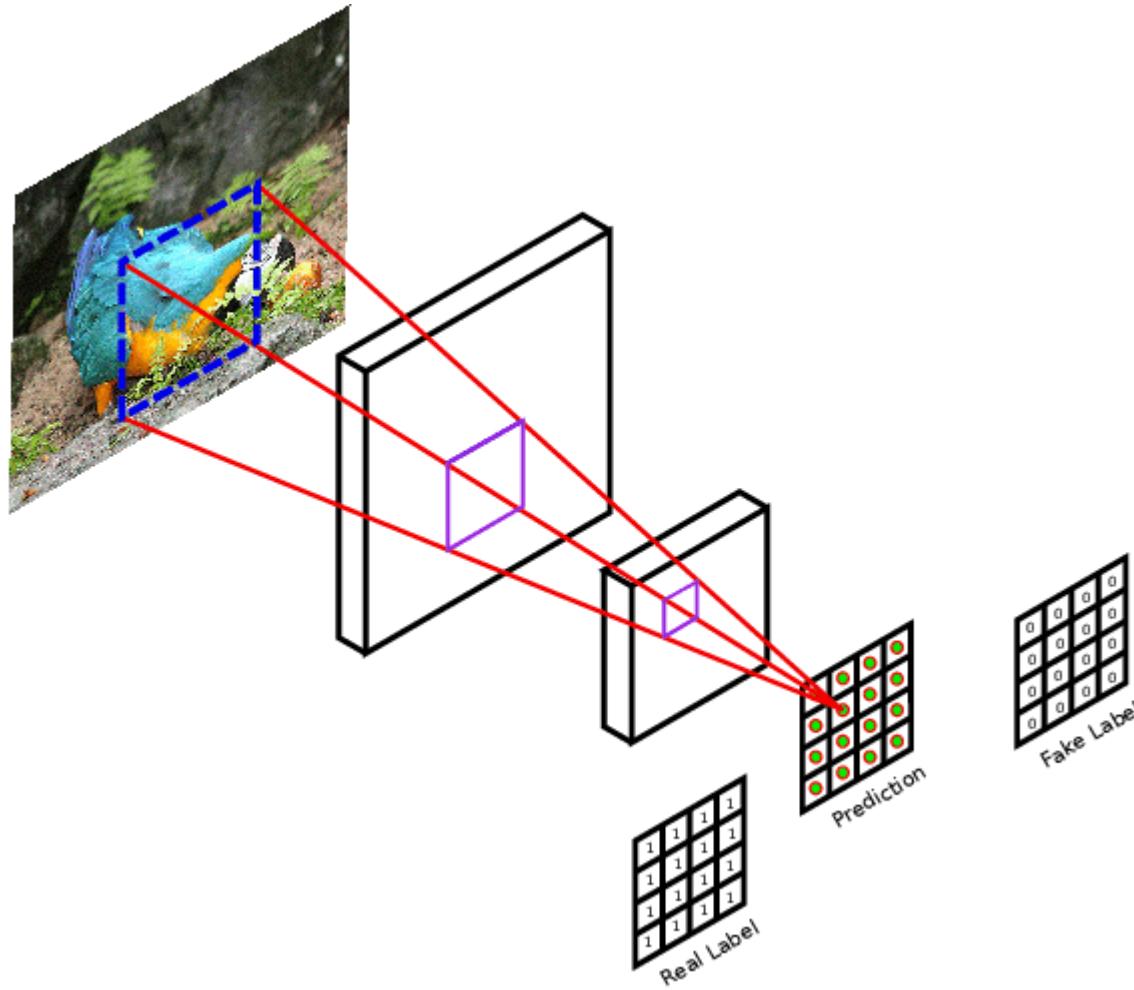


Resolution = 3x64x64



# PatchGAN

## ❖ Idea

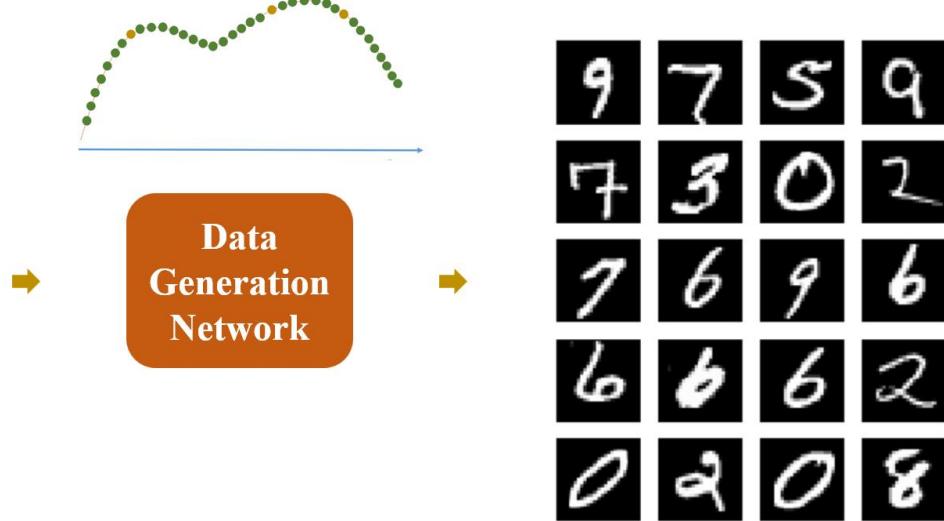


# Outline

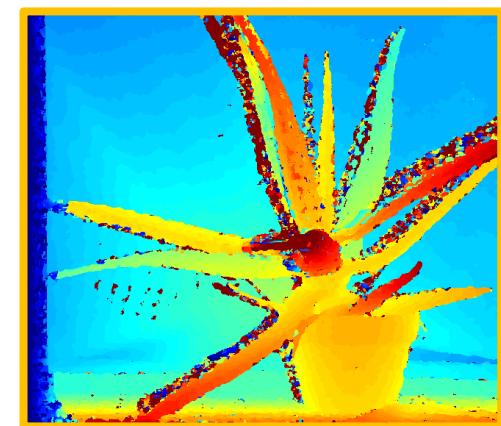
- Quick Review on GAN/DCGAN
- Conditional GAN
- PatchGAN
- Pix2Pix
- CycleGAN

# Pix2Pix

## ❖ Motivation



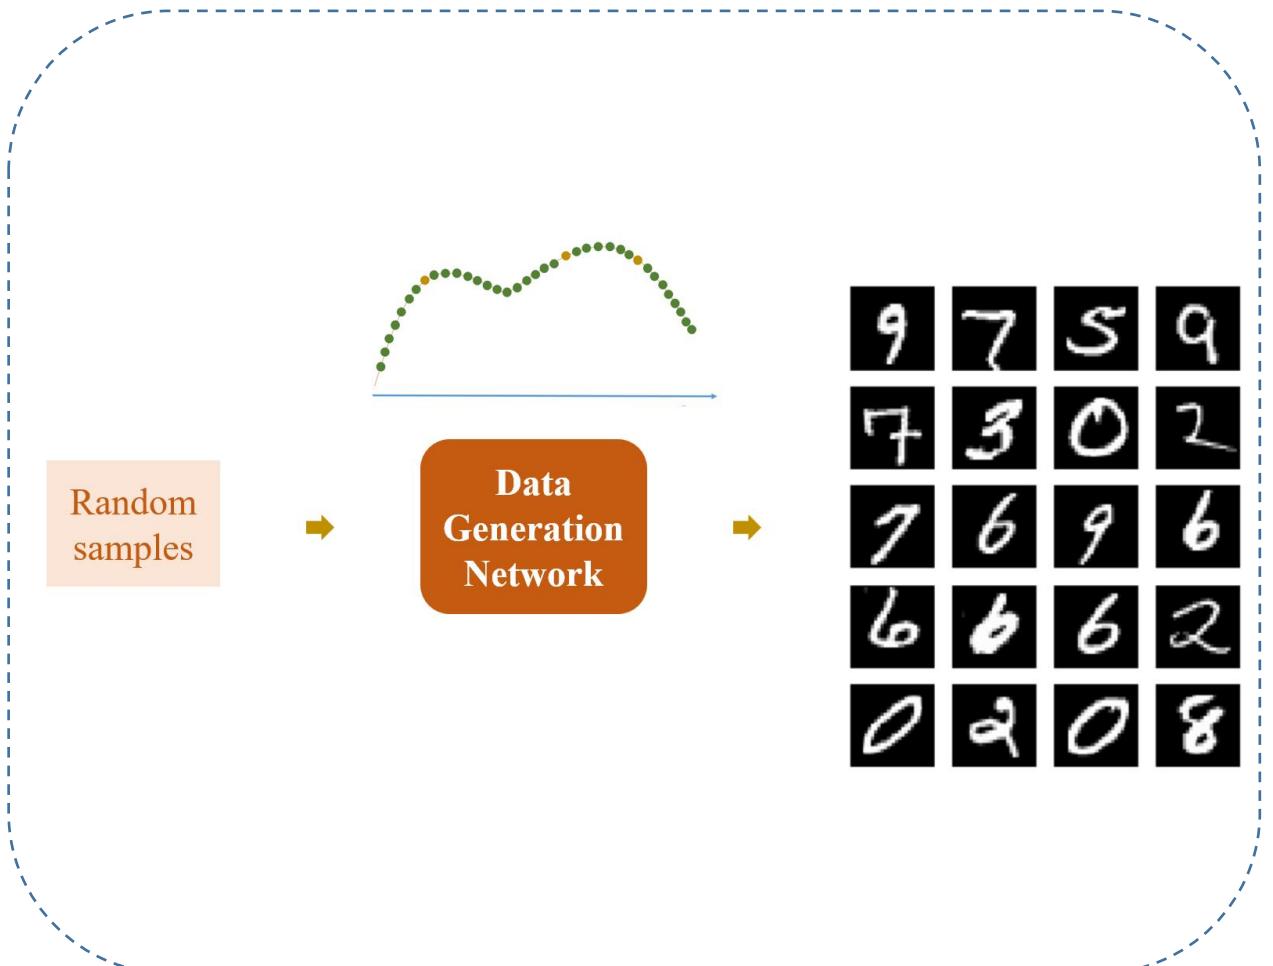
Convert from a noise space to an image space



Convert from an image space to another image space

# Pix2Pix

## ❖ Problems of GAN/DCGAN



Convert from a noise space to an image space

```
class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.init_size = 8
        self.fc = nn.Linear(latent_dim, 128*8*8)
        self.conv_blocks = nn.Sequential(
            nn.BatchNorm2d(128),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(128, 128, 3, padding=1),

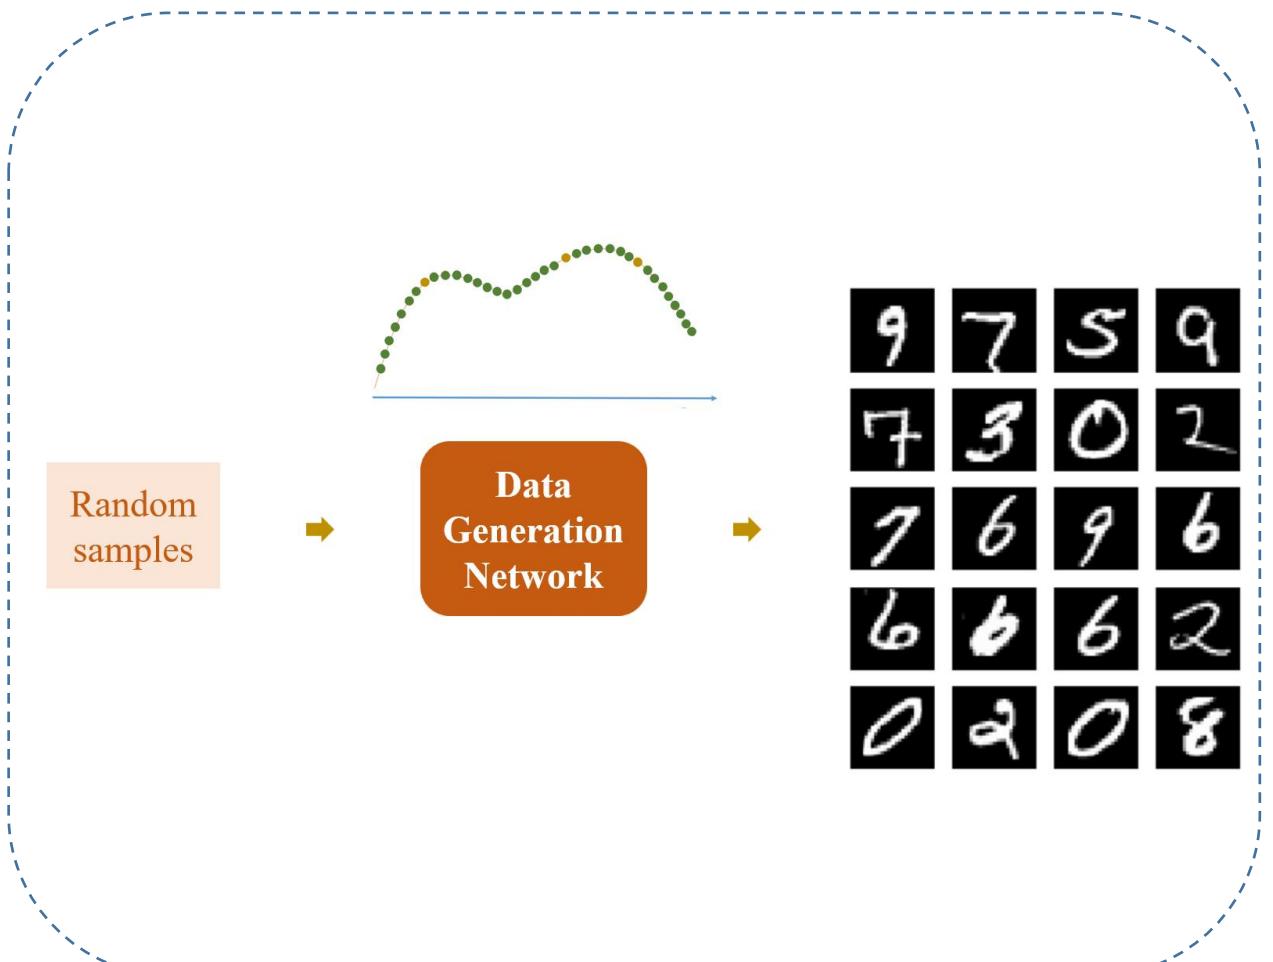
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Upsample(scale_factor=2),
            nn.Conv2d(128, 64, 3, padding=1),

            nn.BatchNorm2d(64),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(64, channels,
                     kernel_size=3, padding=1),
            nn.Tanh()
        )

    def forward(self, z):
        x = self.fc(z)
        x = x.view(x.shape[0], 128,
                   self.init_size, self.init_size)
        img = self.conv_blocks(x)
        return img
```

# Pix2Pix

## ❖ Problems of GAN/DCGAN



```
class Descriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Conv2d(channels, 16, kernel_size=3,
                     stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True)

            nn.Conv2d(16, 32, kernel_size=3,
                     stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True),

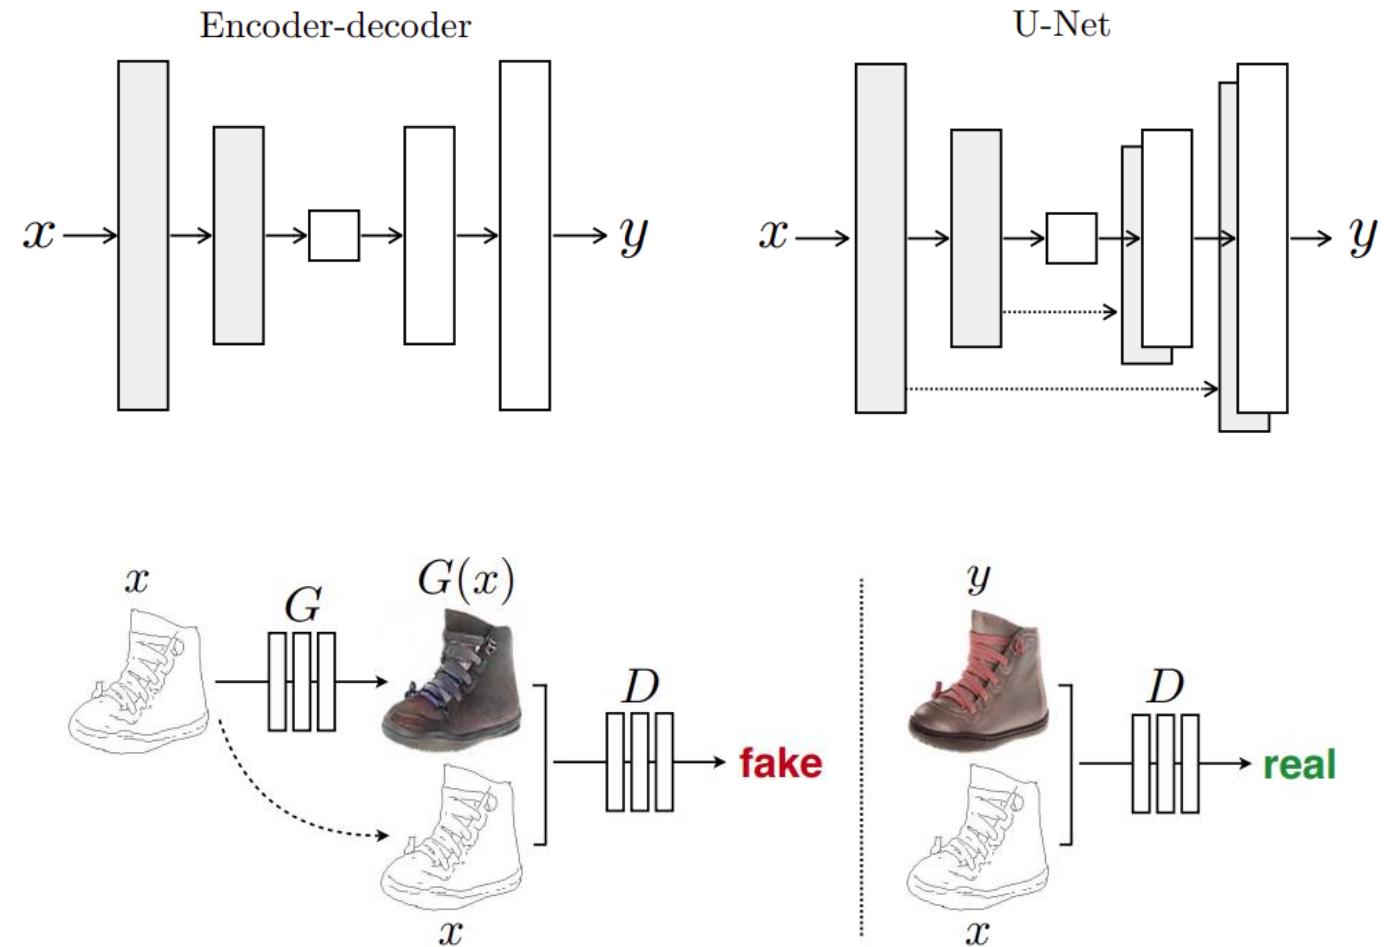
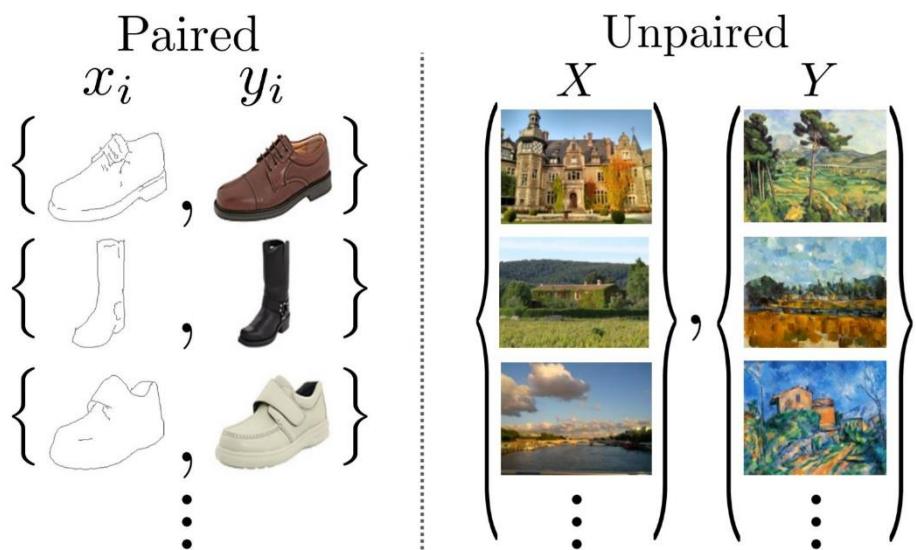
            nn.Conv2d(32, 64, kernel_size=3,
                     stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Conv2d(64, 128, kernel_size=3,
                     stride=2, padding=1),
            nn.LeakyReLU(0.2, inplace=True))
        self.adv_layer = nn.Sequential(
            nn.Linear(128*2*2, 1),
            nn.Sigmoid())

    def forward(self, img):
        x = self.model(img)
        x = x.view(x.shape[0], -1)
        validity = self.adv_layer(x)
        return validity
```

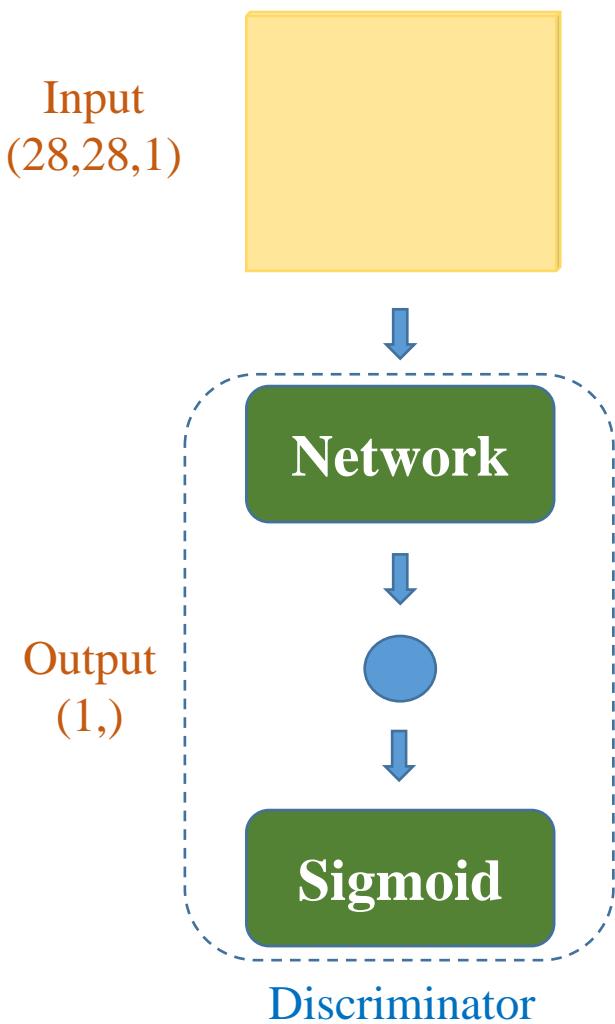
# Pix2Pix

## ❖ Network



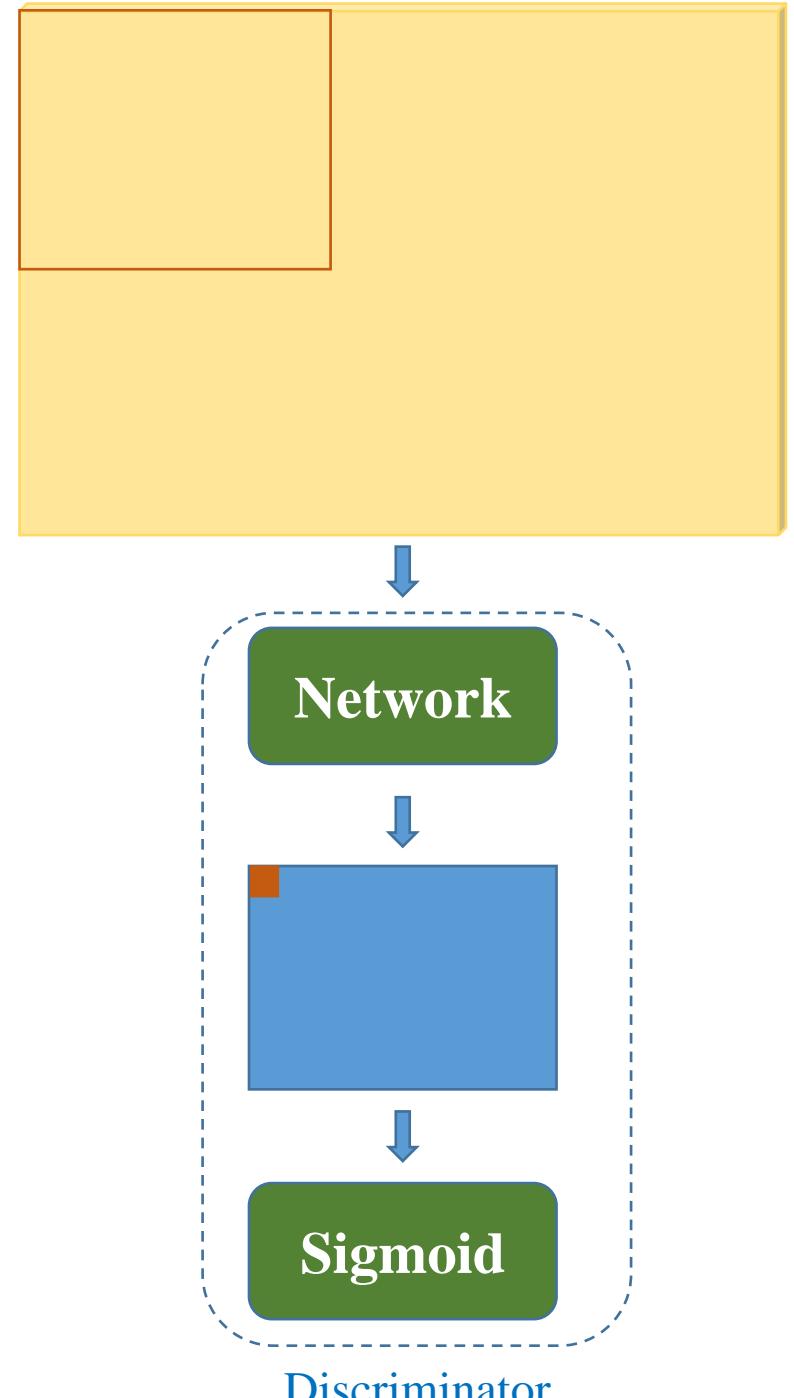
# Pix2Pix

## ❖ Loss function



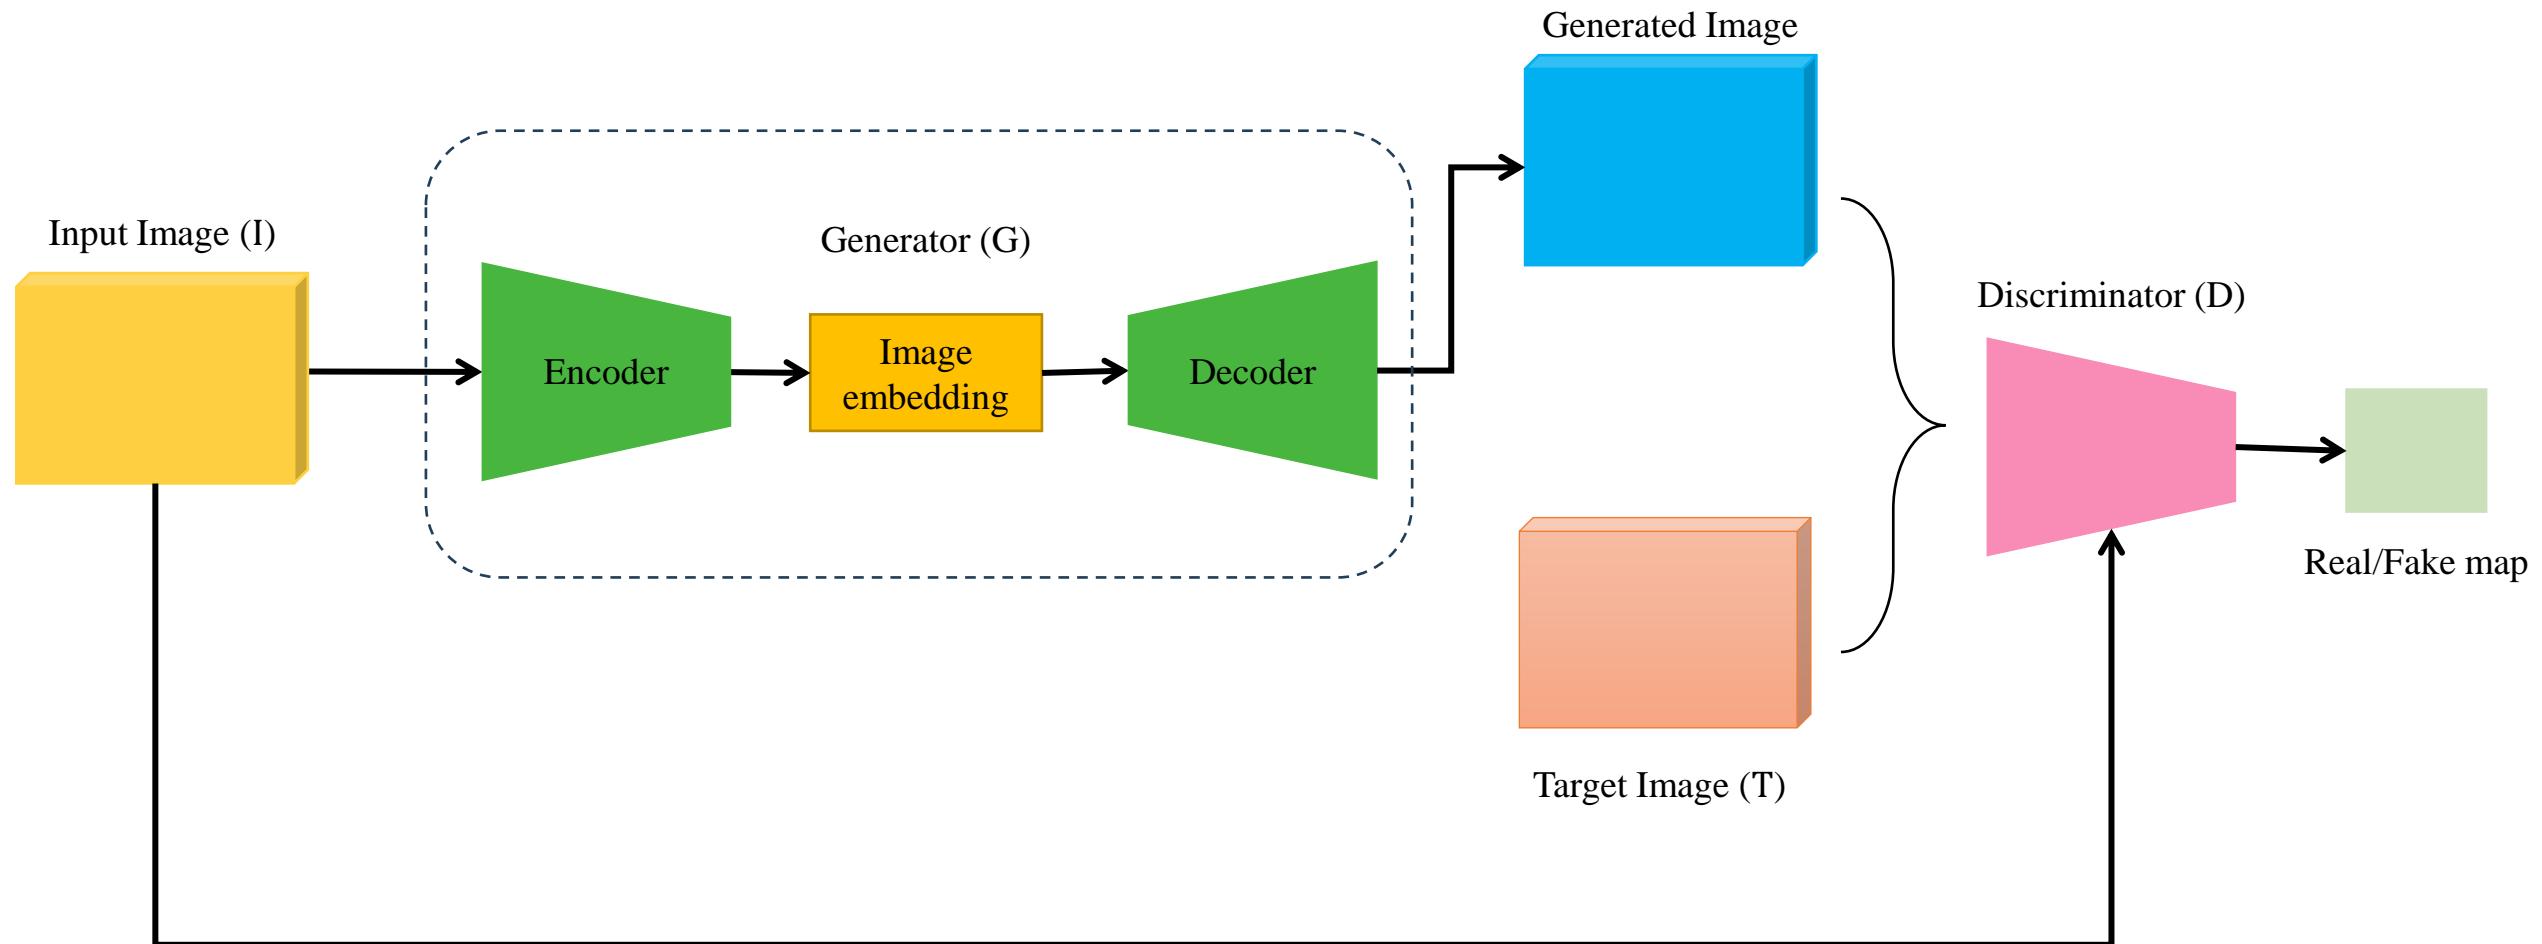
Input  
 $(256, 256, 3)$

Output  
 $(30, 30, 1)$

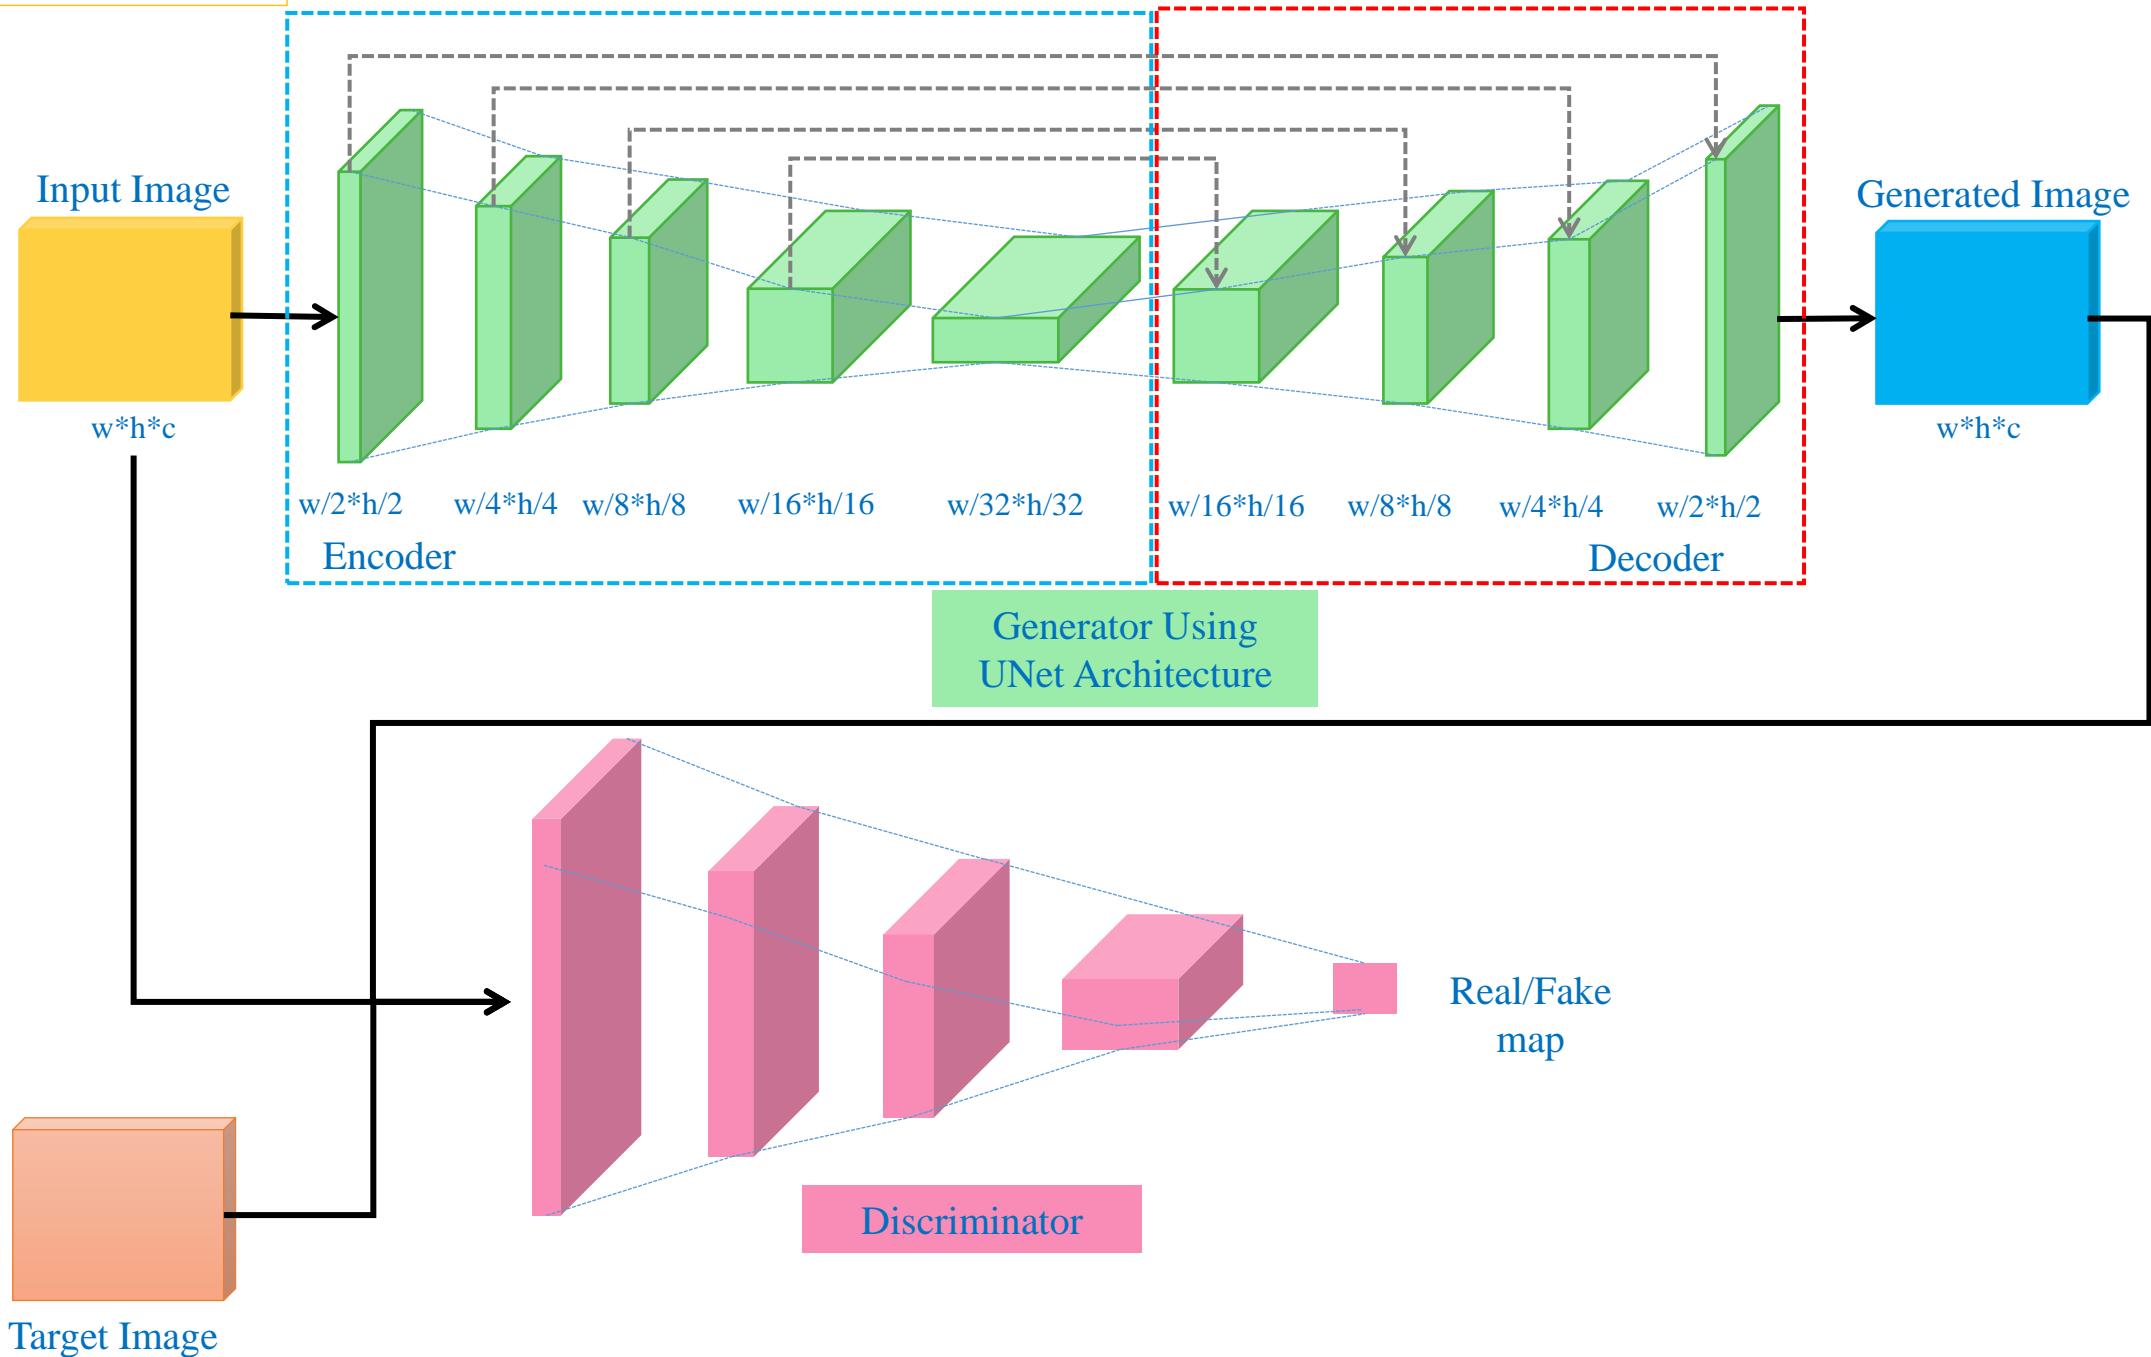


# Pix2Pix

## ❖ Model

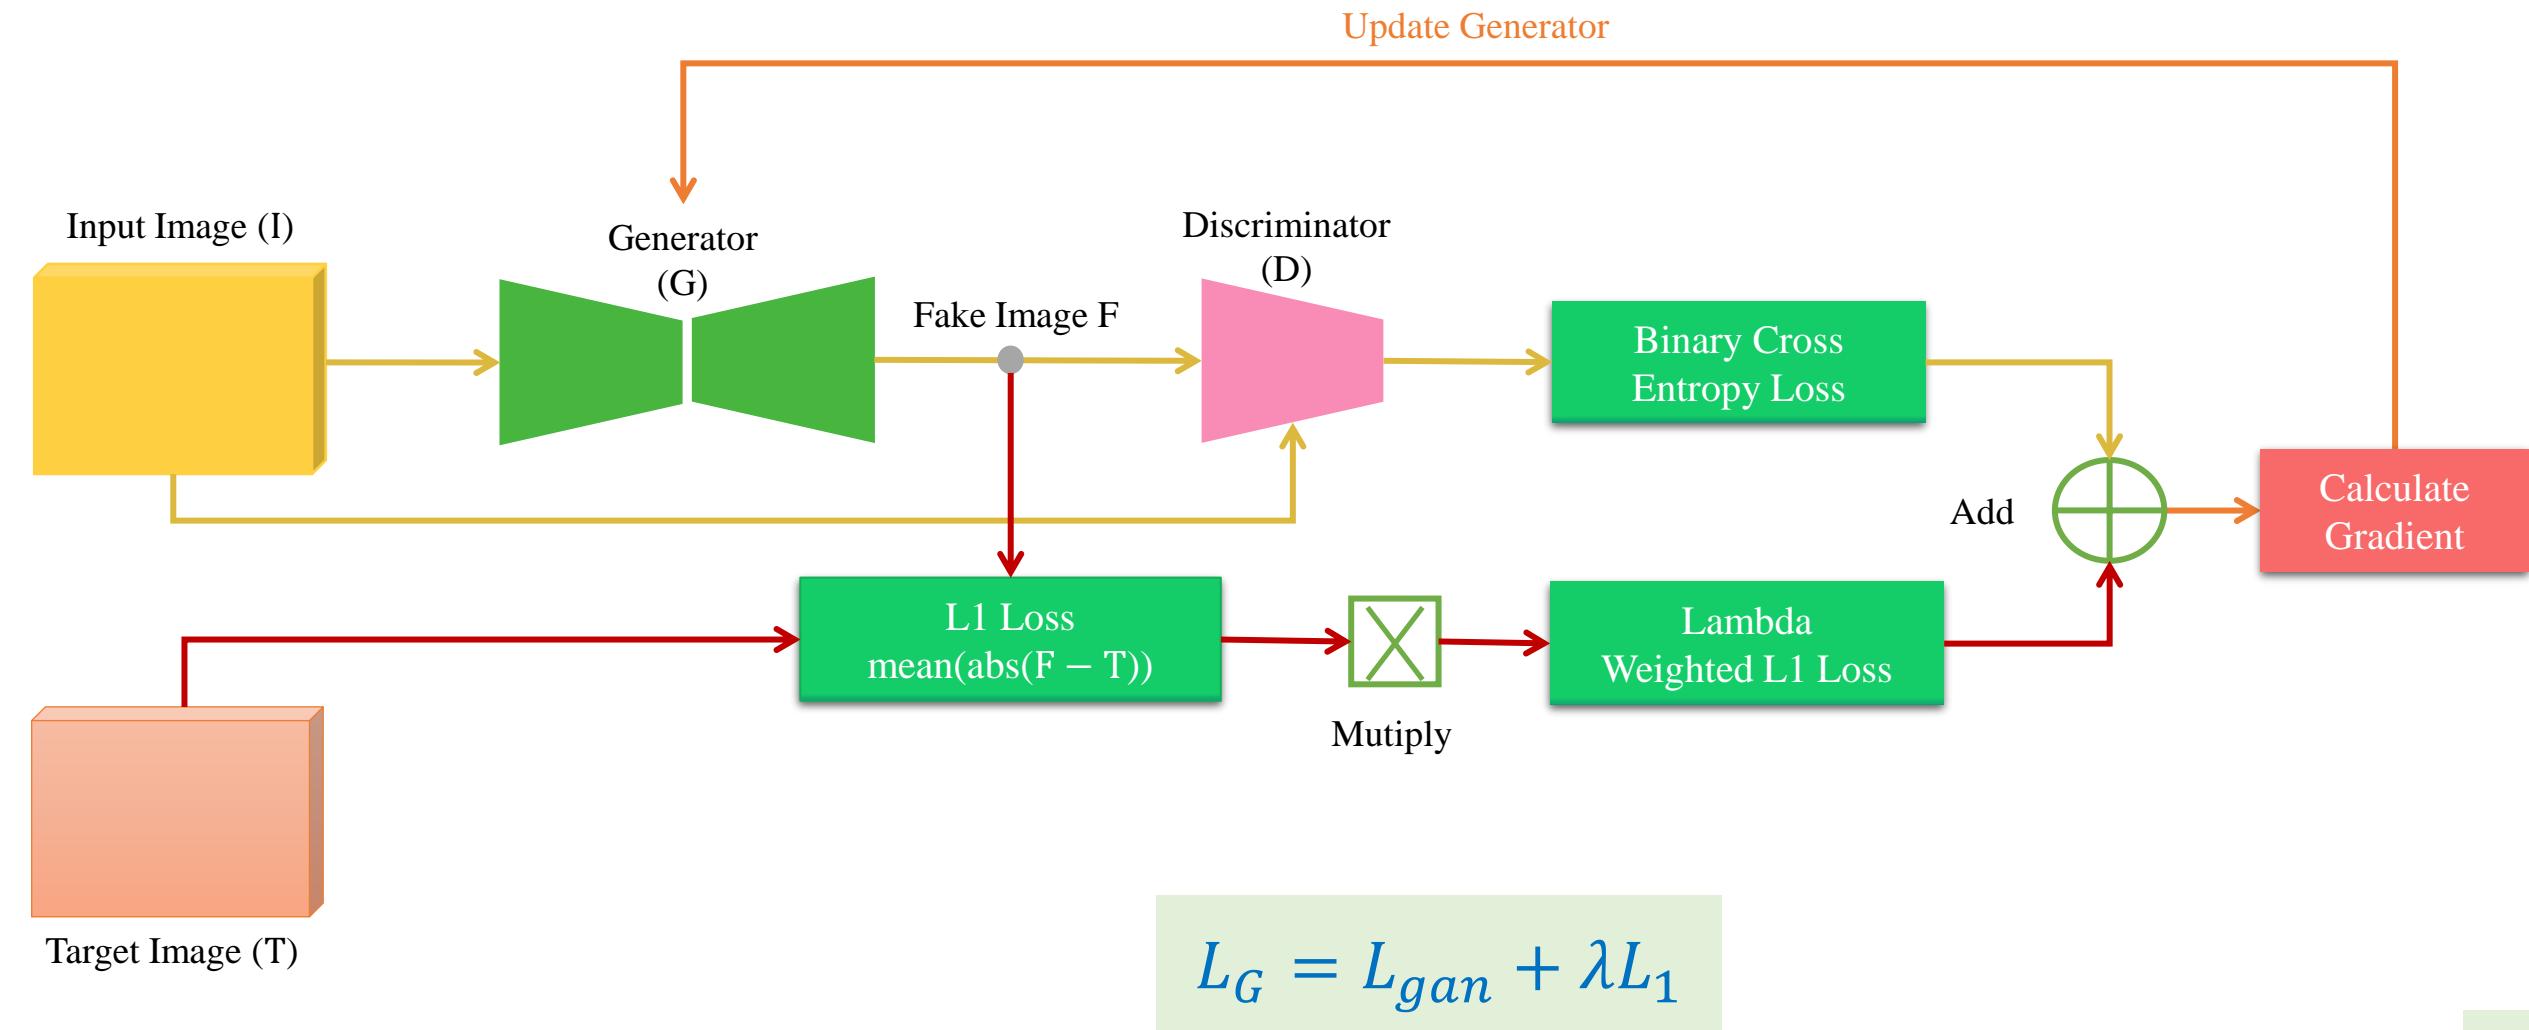


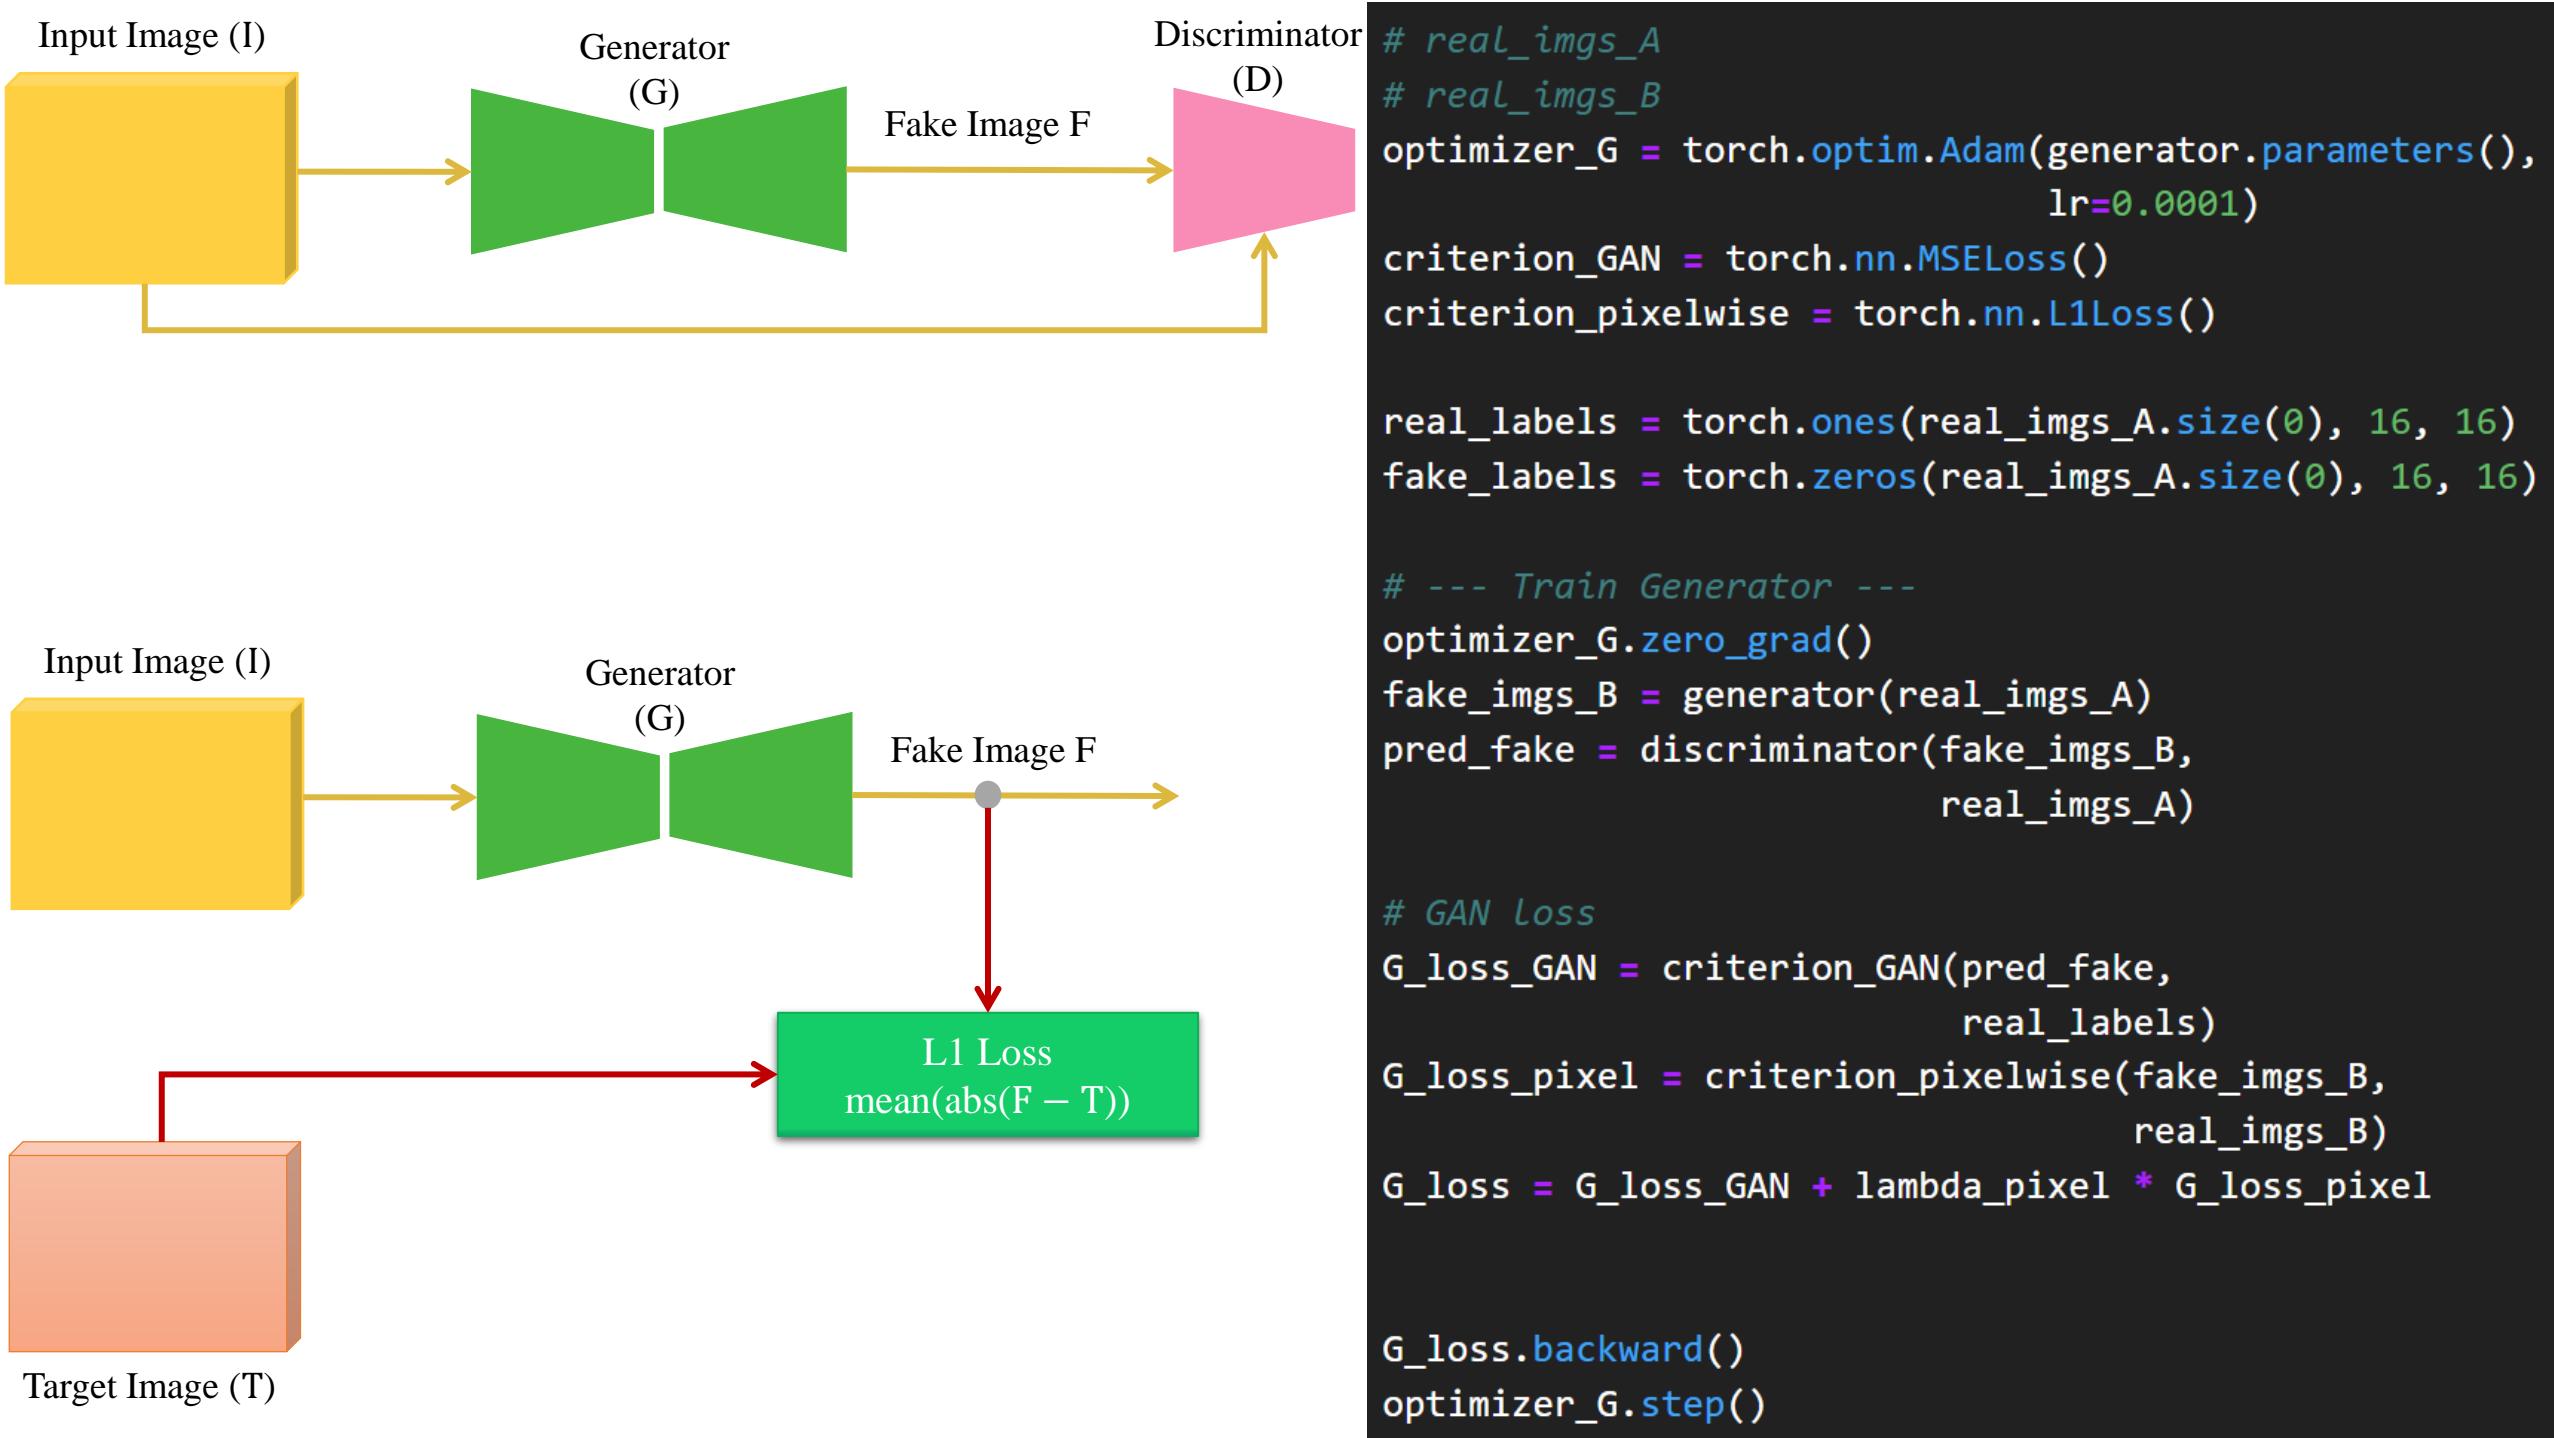
## Pix2Pix Architecture



# Pix2Pix

## Training the Generator

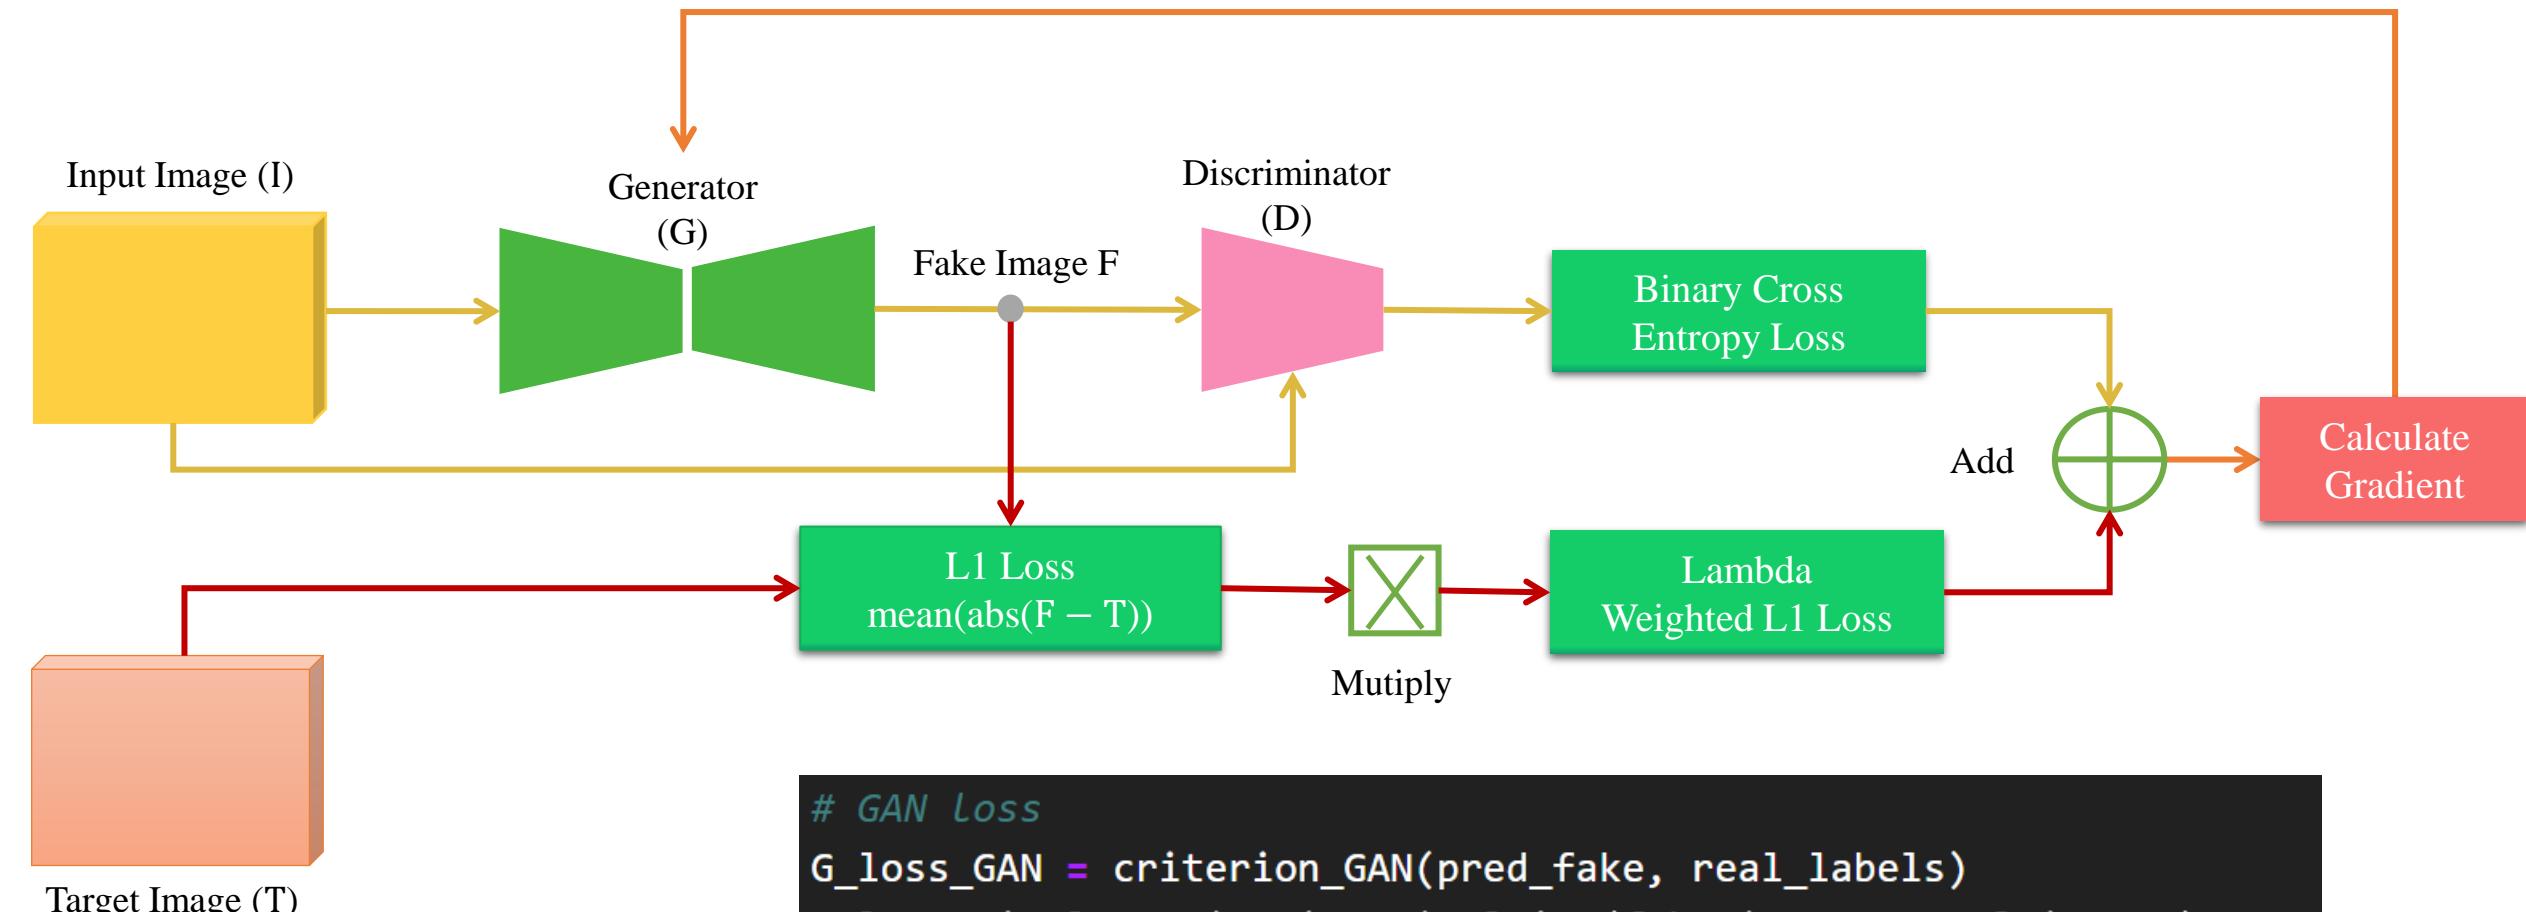




# Pix2Pix

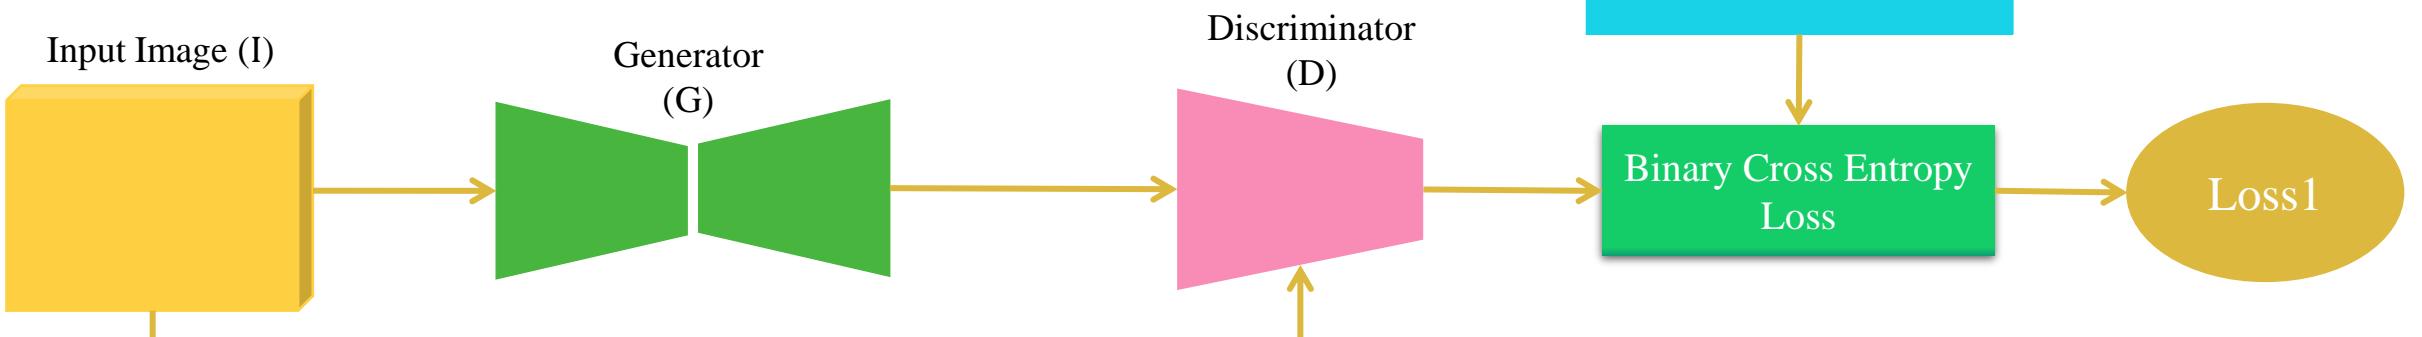
$$L_G = L_{gan} + \lambda L_1$$

## Training the Generator

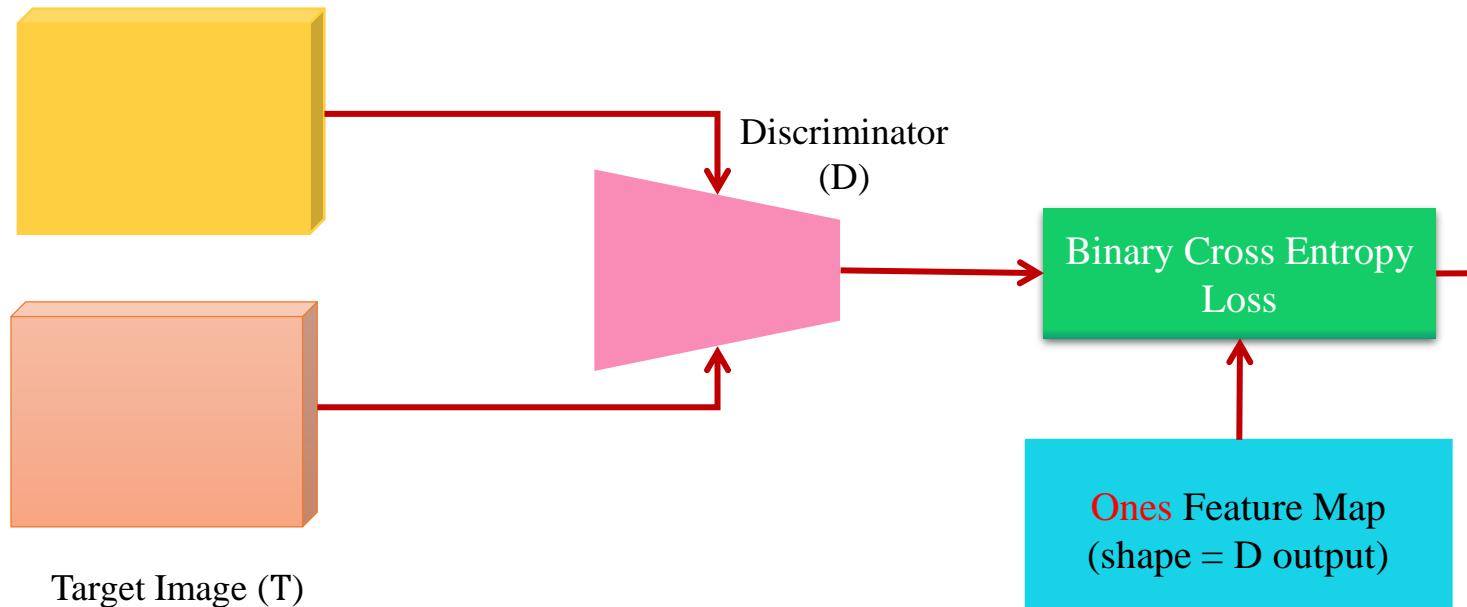


```
# GAN Loss
G_loss_GAN = criterion_GAN(pred_fake, real_labels)
G_loss_pixel = criterion_pixelwise(fake_imgs_B, real_imgs_B)
G_loss = G_loss_GAN + lambda_pixel * G_loss_pixel
```

## STAGE 1



## STAGE 2



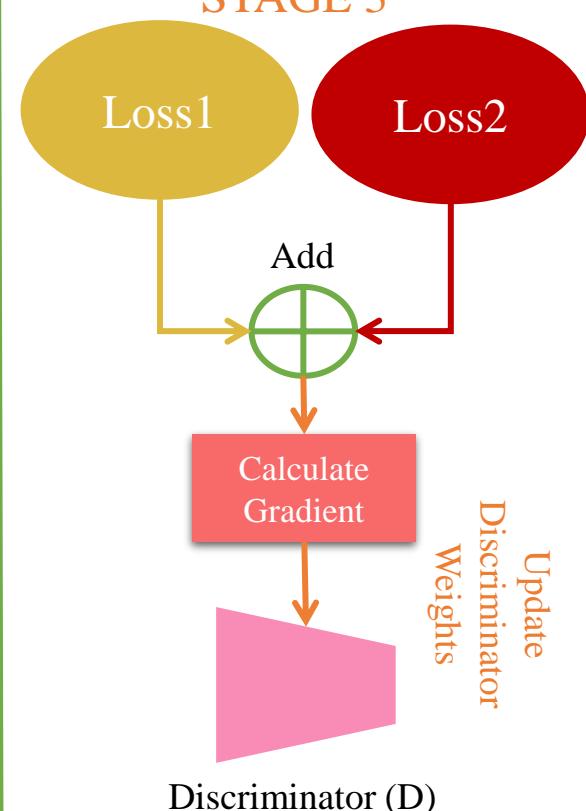
Training Procedure for the Discriminator

Zeros Feature Map  
(shape = D output)

Binary Cross Entropy  
Loss

Loss1

## STAGE 3



Discriminator (D)

Update  
Discriminator  
Weights

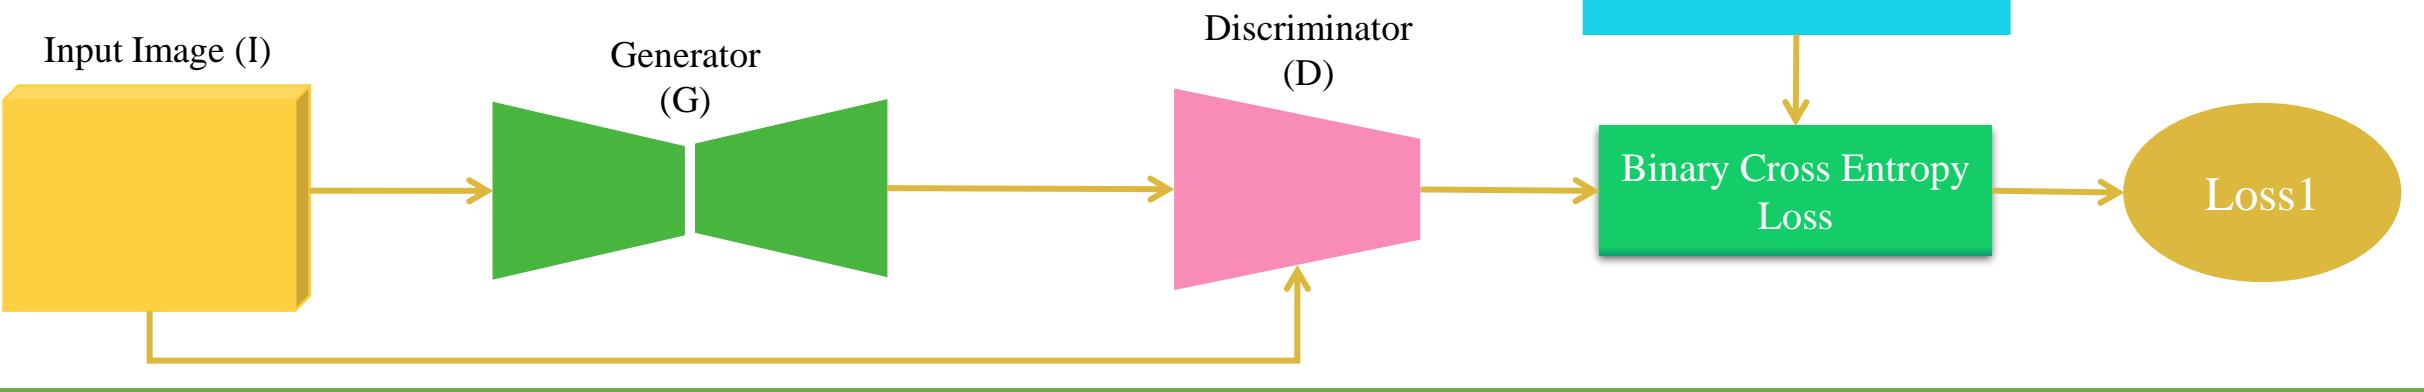
Calculate  
Gradient

Loss1

Loss2

Add

## STAGE 1



# Pix2Pix

## Training the Discriminator

```
# real_imgs_A
# real_imgs_B

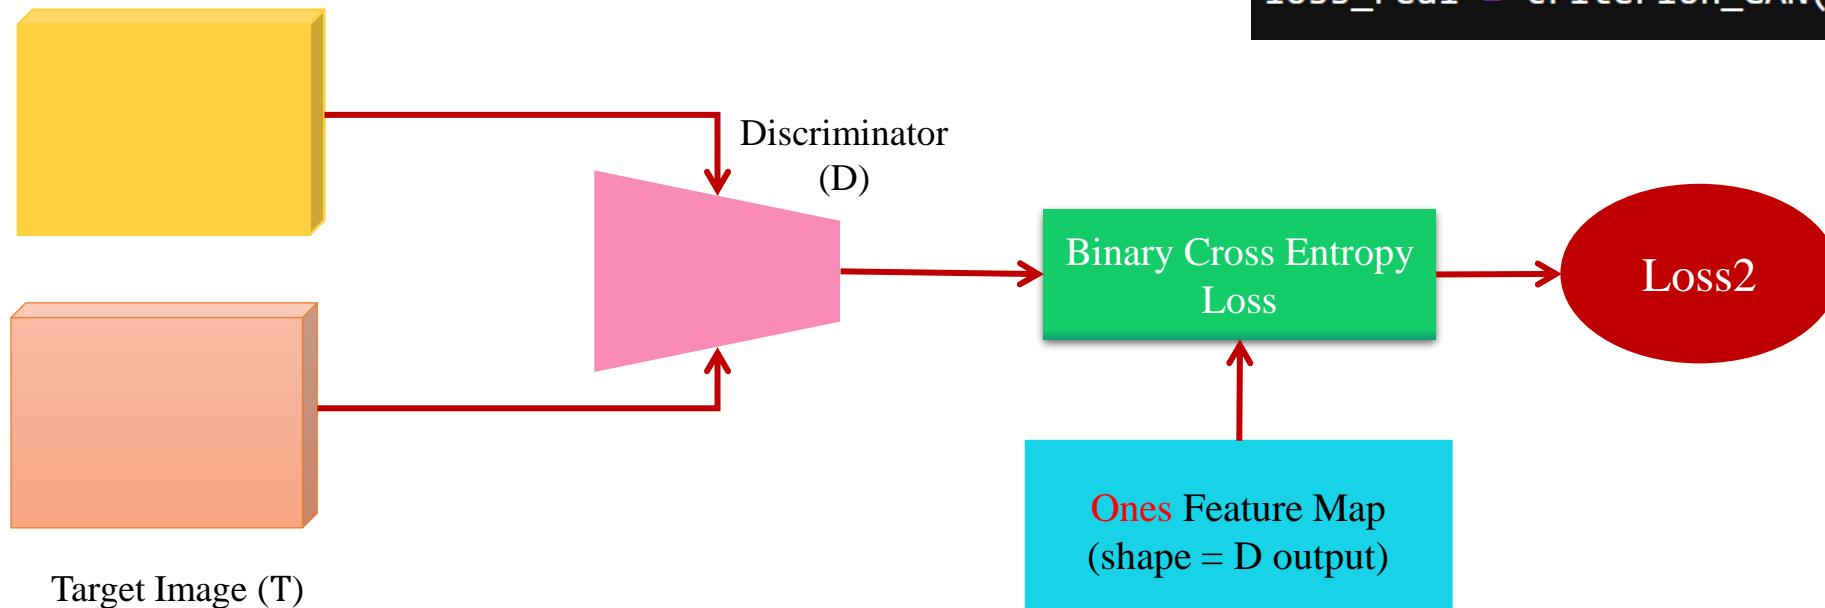
criterion = torch.nn.MSELoss()

# prepare Labels
fake_labels = torch.zeros(real_imgs_A.size(0), *patch)

# Train
pred_fake = discriminator(fake_imgs_B.detach(), real_imgs_A)
loss_fake = criterion(pred_fake, fake_labels)
```

# Pix2Pix

## Training the Discriminator



```
# real_imgs_A  
# real_imgs_B  
  
criterion = torch.nn.MSELoss()  
  
# prepare labels  
real_labels = torch.ones(real_imgs_A.size(0), *patch)  
  
# Train  
pred_real = discriminator(real_imgs_B, real_imgs_A)  
loss_real = criterion_GAN(pred_real, real_labels)
```

Training Procedure for the Discriminator

```

optimizer_D = Adam(discriminator.parameters(),
                   lr=0.0002)
criterion = torch.nn.MSELoss()

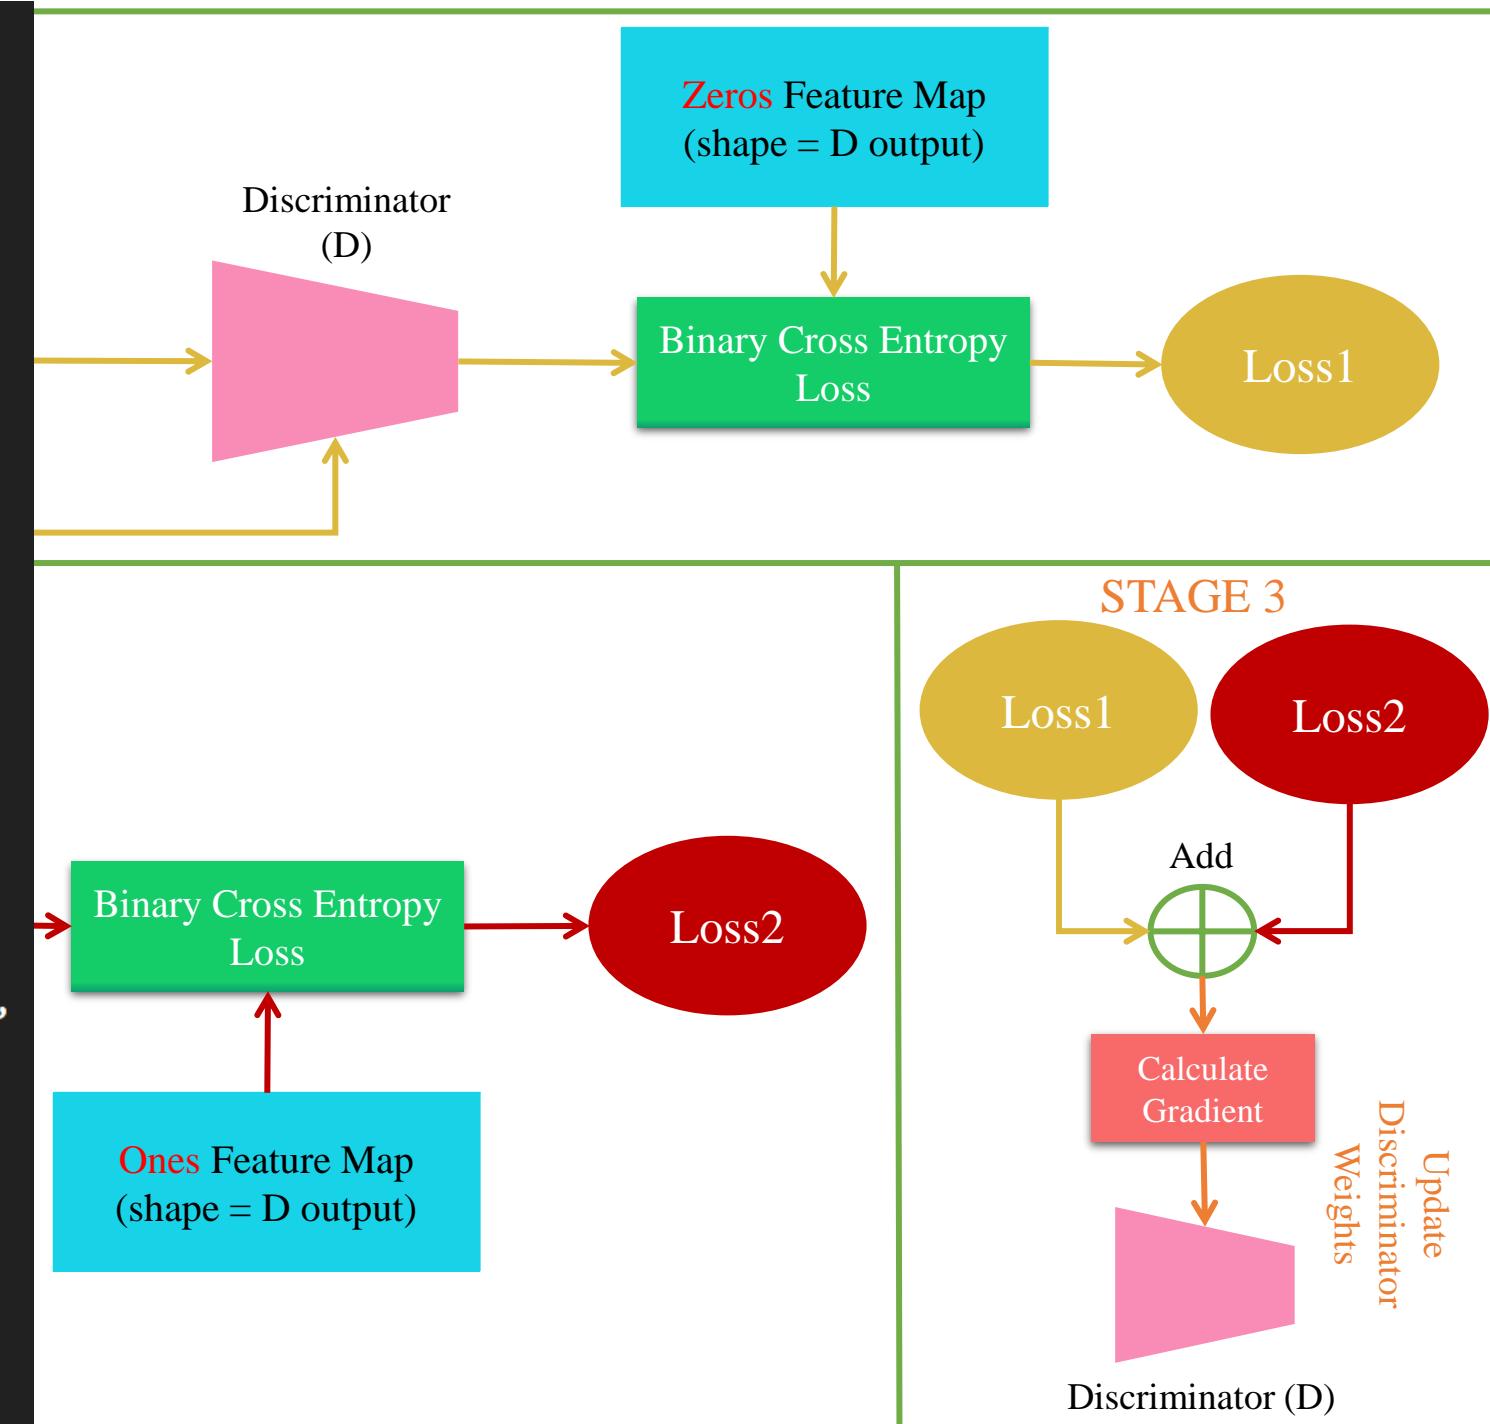
# prepare labels
real_labels = torch.ones(real_imgs_A.size(0),
                         *patch)
fake_labels = torch.zeros(real_imgs_A.size(0),
                         *patch)

# --- Train Discriminator
optimizer_D.zero_grad()
pred_real = discriminator(real_imgs_B,
                           real_imgs_A)
loss_real = criterion_GAN(pred_real,
                          real_labels)

pred_fake = discriminator(fake_imgs_B.detach(),
                           real_imgs_A)
loss_fake = criterion(pred_fake, fake_labels)
D_loss = (loss_real + loss_fake) / 2

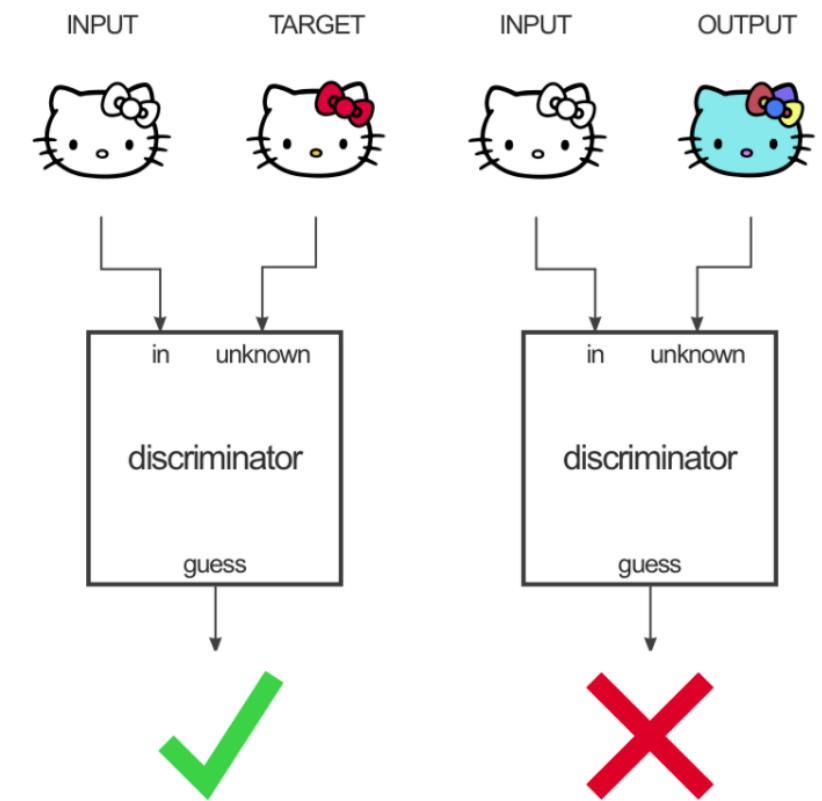
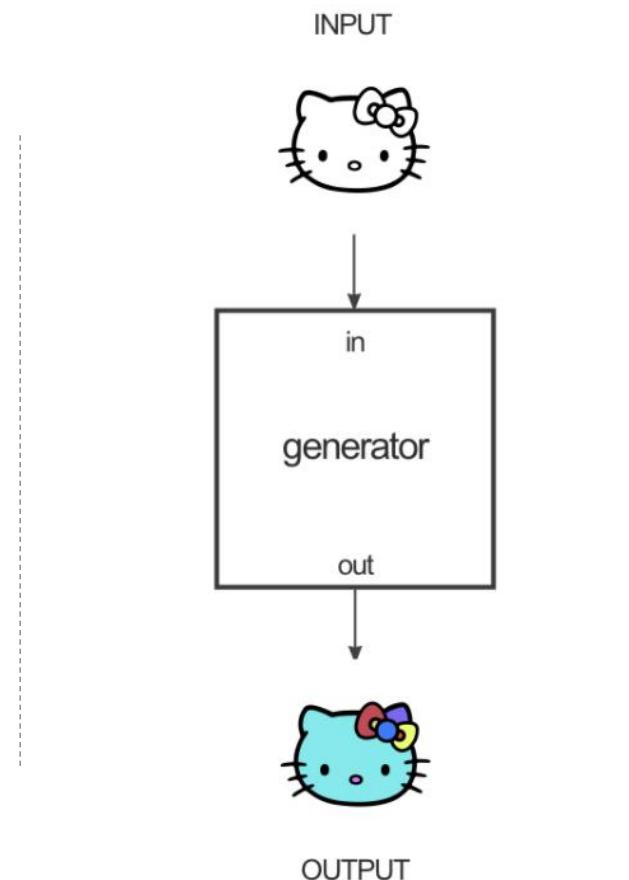
# update
D_loss.backward()
optimizer_D.step()

```



# Pix2Pix

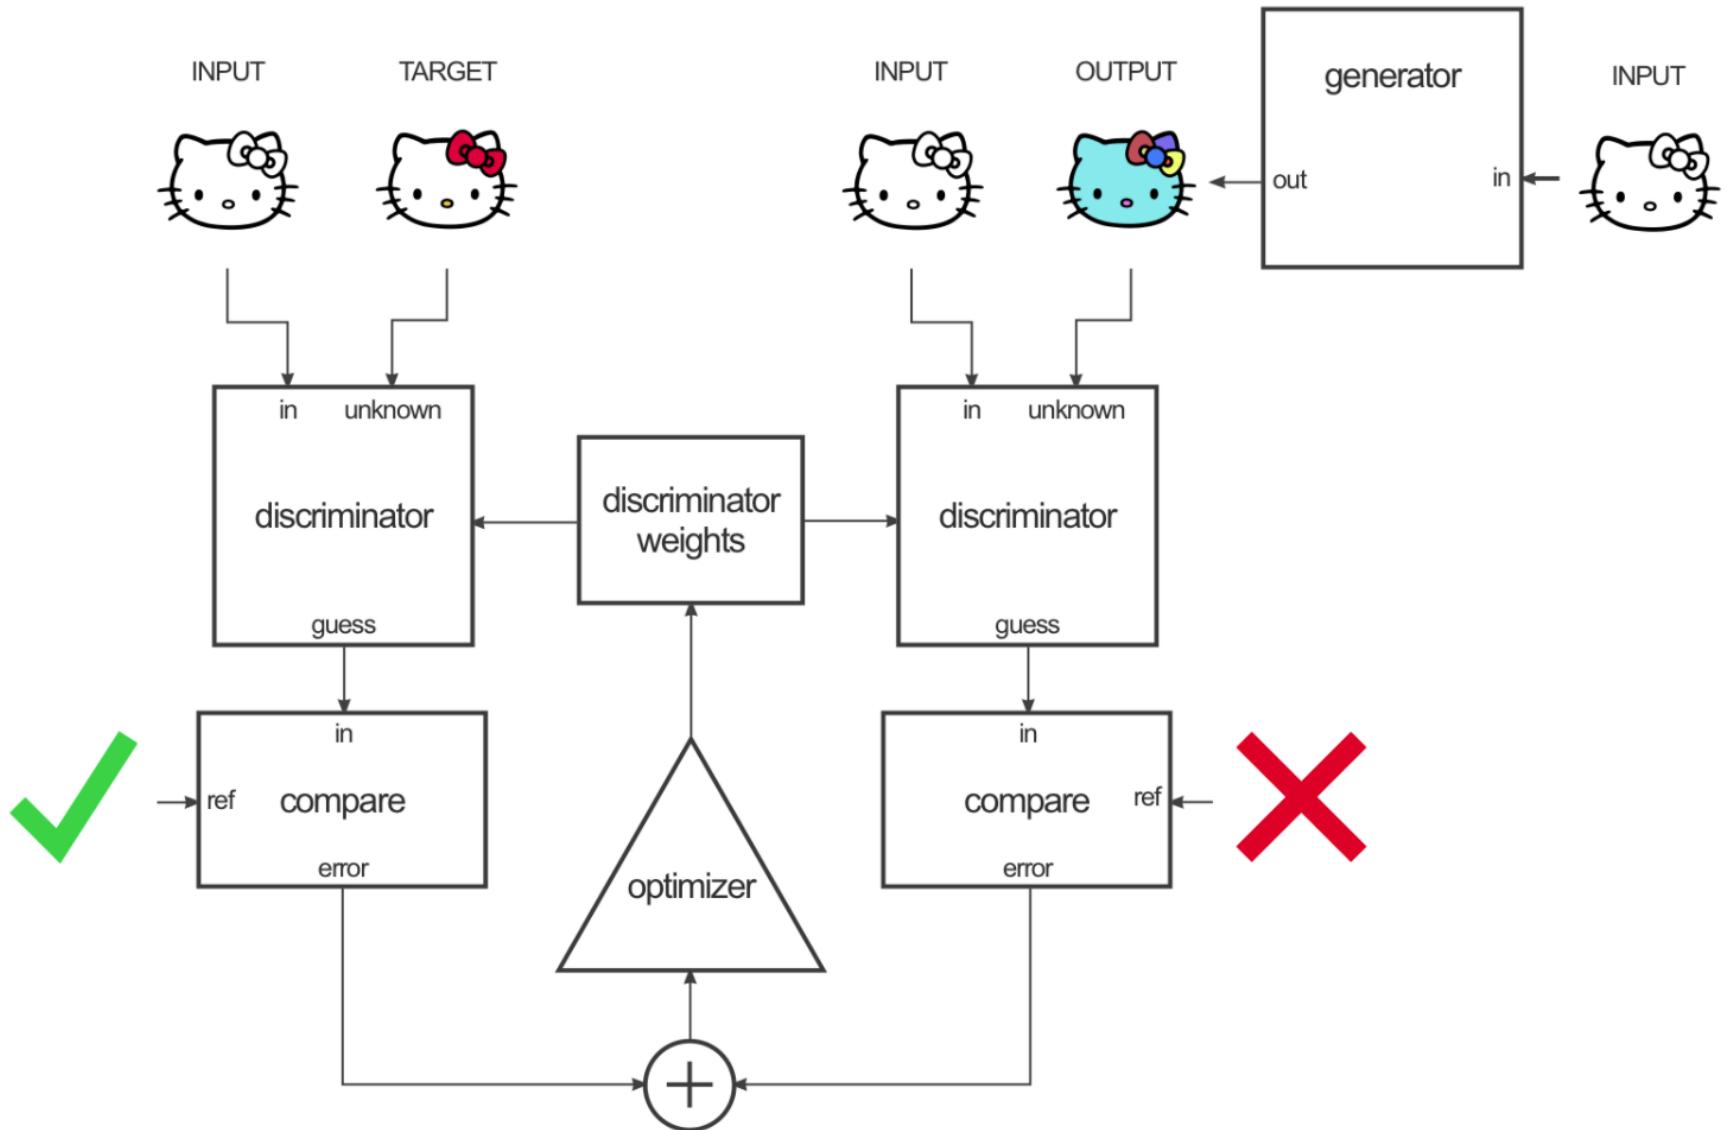
## ❖ Train the model



# Pix2Pix

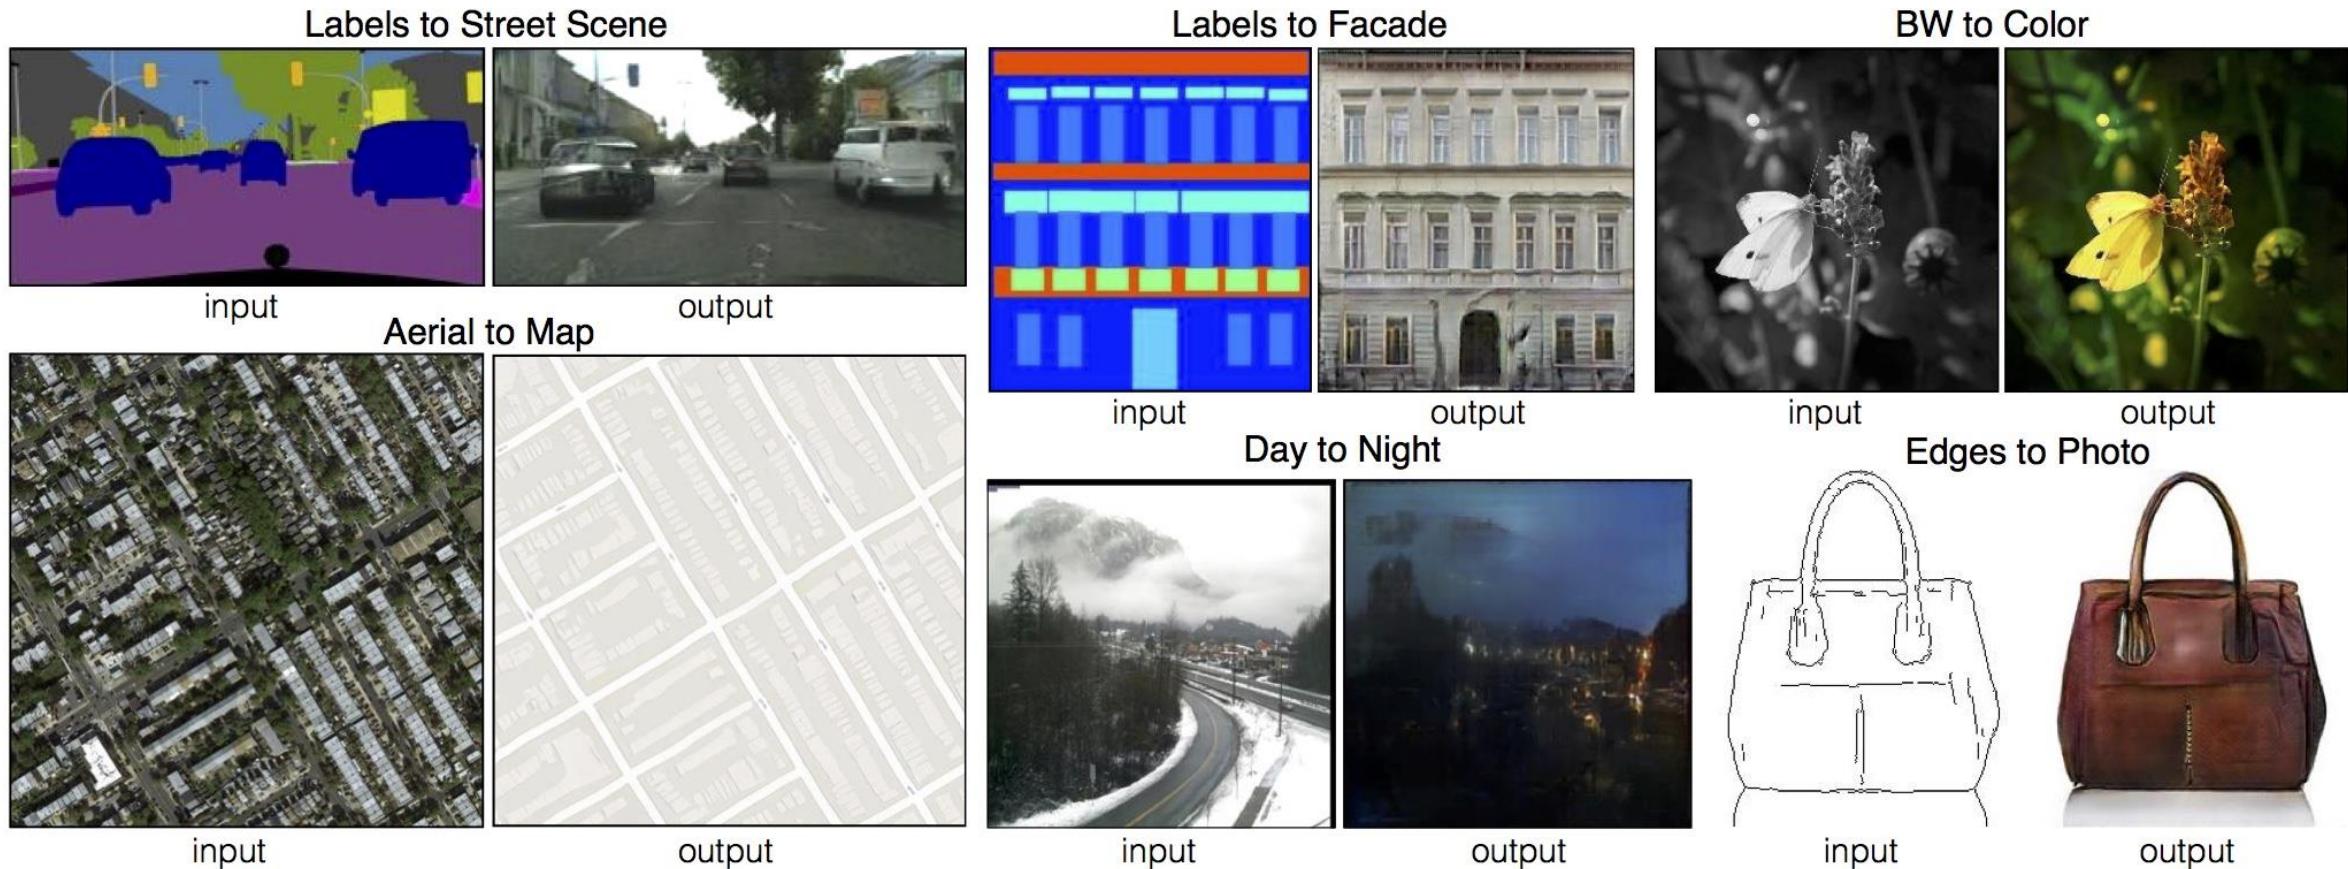
## ❖ Train the model

Train the  
discriminator



# Pix2Pix

## ❖ Applications



# Pix2Pix

## ❖ Applications



GTA → Cityscapes

# Pix2Pix

## ❖ Applications



winter Yosemite → summer Yosemite



summer Yosemite → winter Yosemite

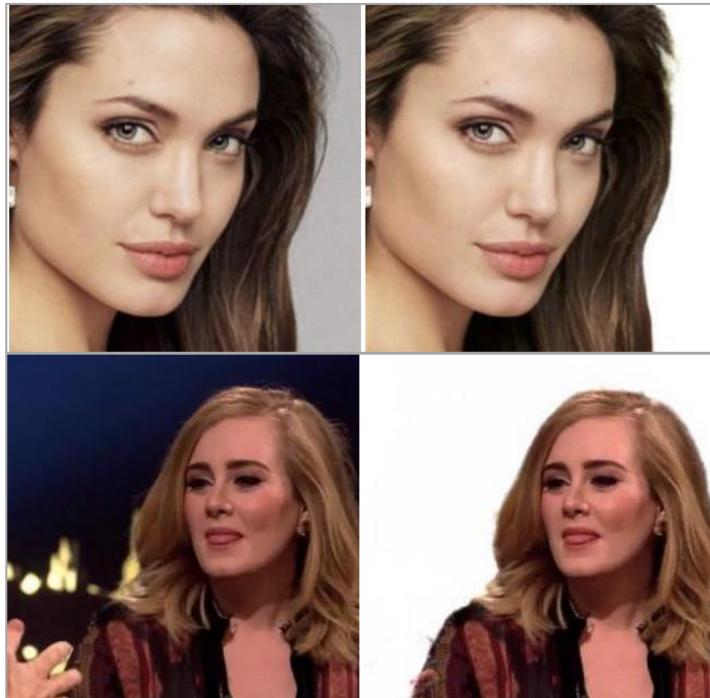
# Pix2Pix

## ❖ Applications



# Pix2Pix

## ❖ Applications



Background masking

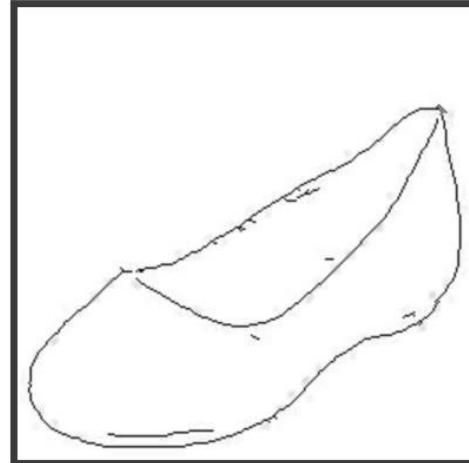
<https://affinelayer.com/pixsrv/>

### edges2shoes

#### TOOL

line  
eraser

#### INPUT



pix2pix  
process

undo clear random

#### OUTPUT



save

### edges2handbags

#### TOOL

line  
eraser

#### INPUT



pix2pix  
process

undo clear random

#### OUTPUT



save

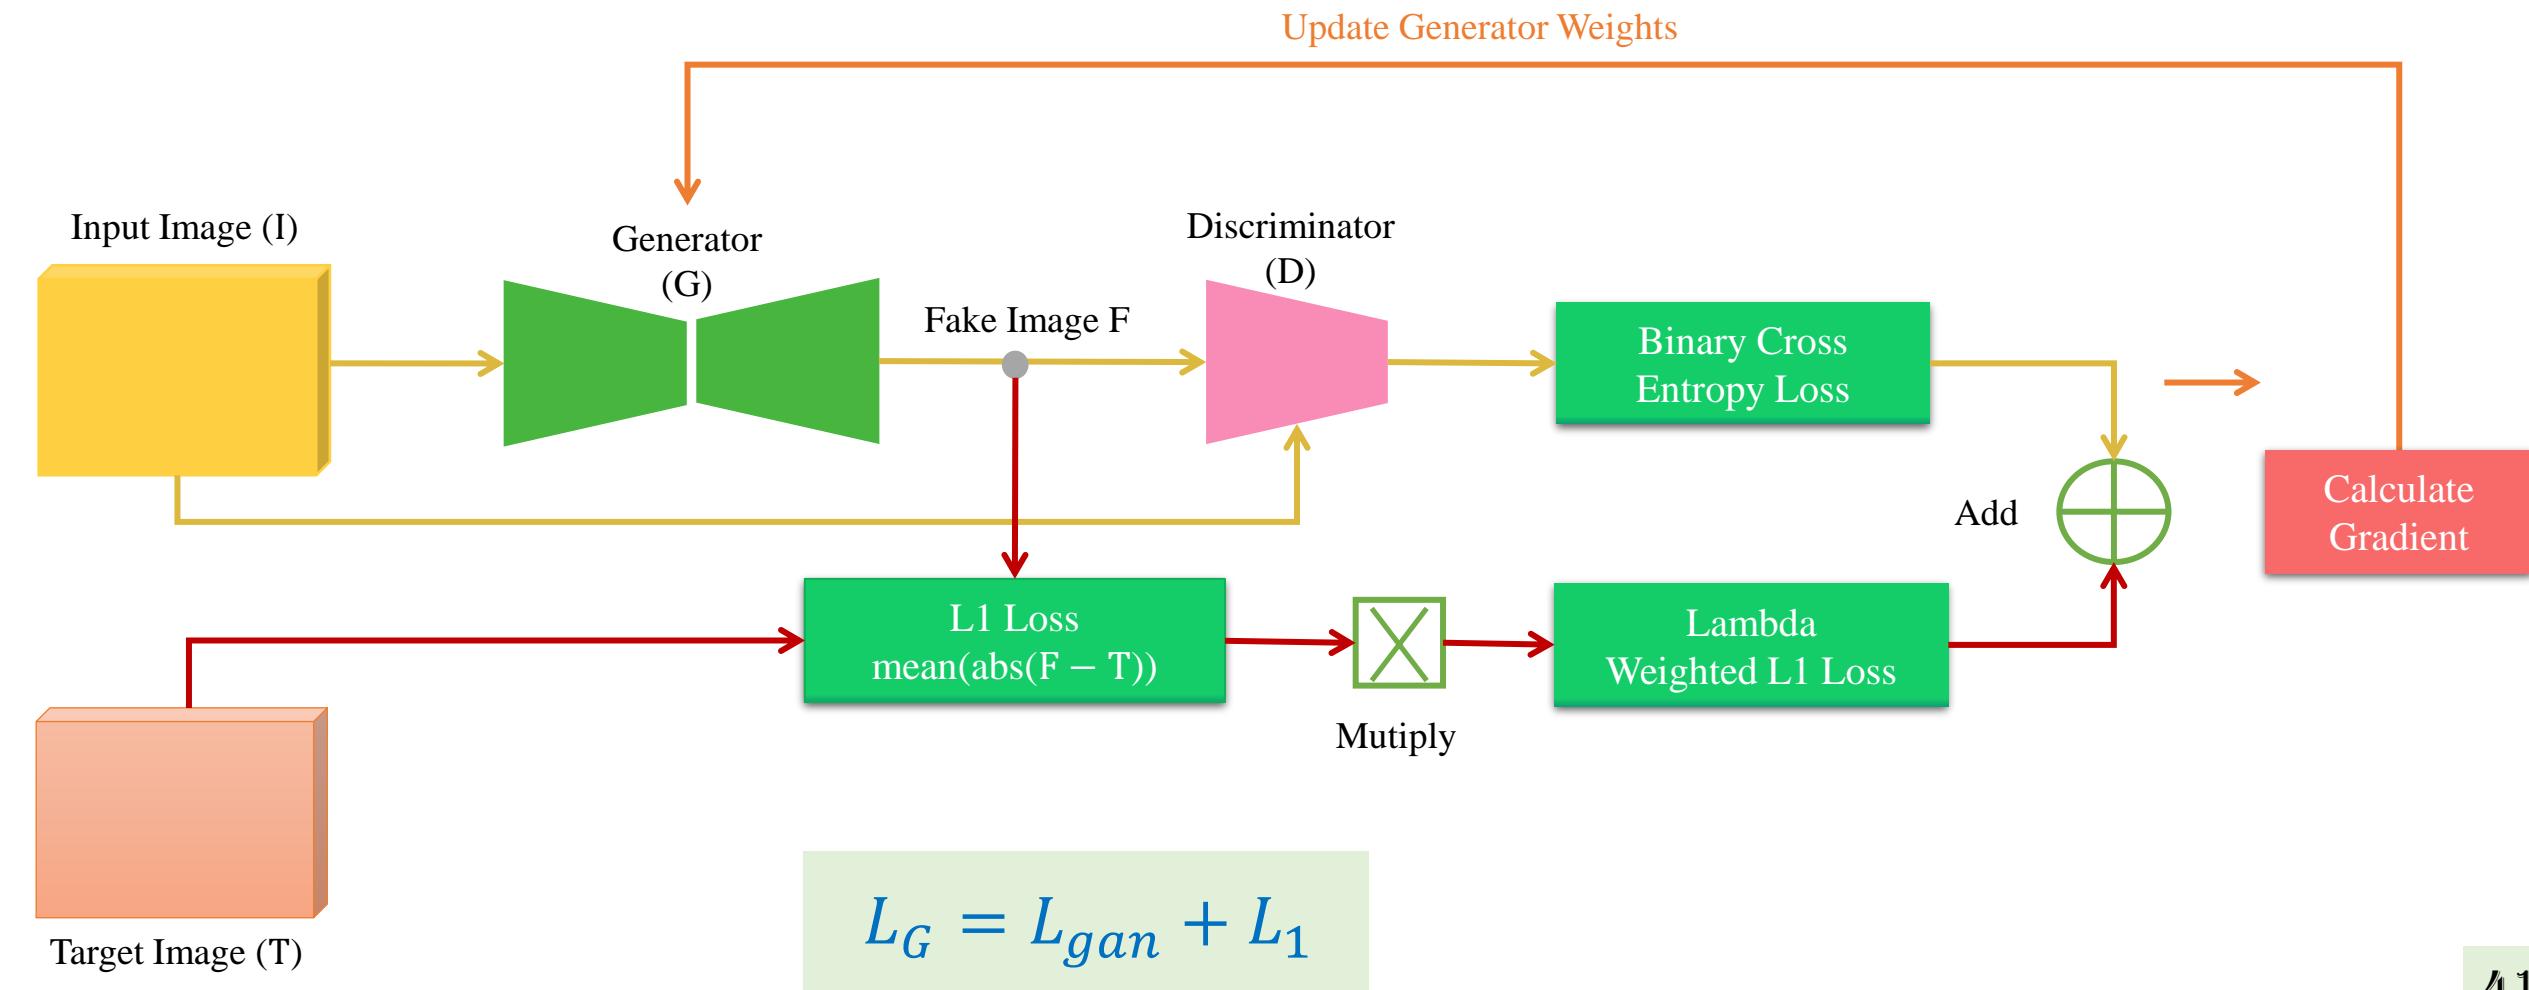
**QUIZ TIME**

# Outline

- Quick Review on GAN/DCGAN
- Conditional GAN
- PatchGAN
- Pix2Pix
- CycleGAN

# CycleGAN

## ❖ Pix2Pix - Problem?



# CycleGAN

## ❖ Motivation

Horses2Zebra Dataset



# CycleGAN

## ❖ Motivation

Apple2Orange Dataset

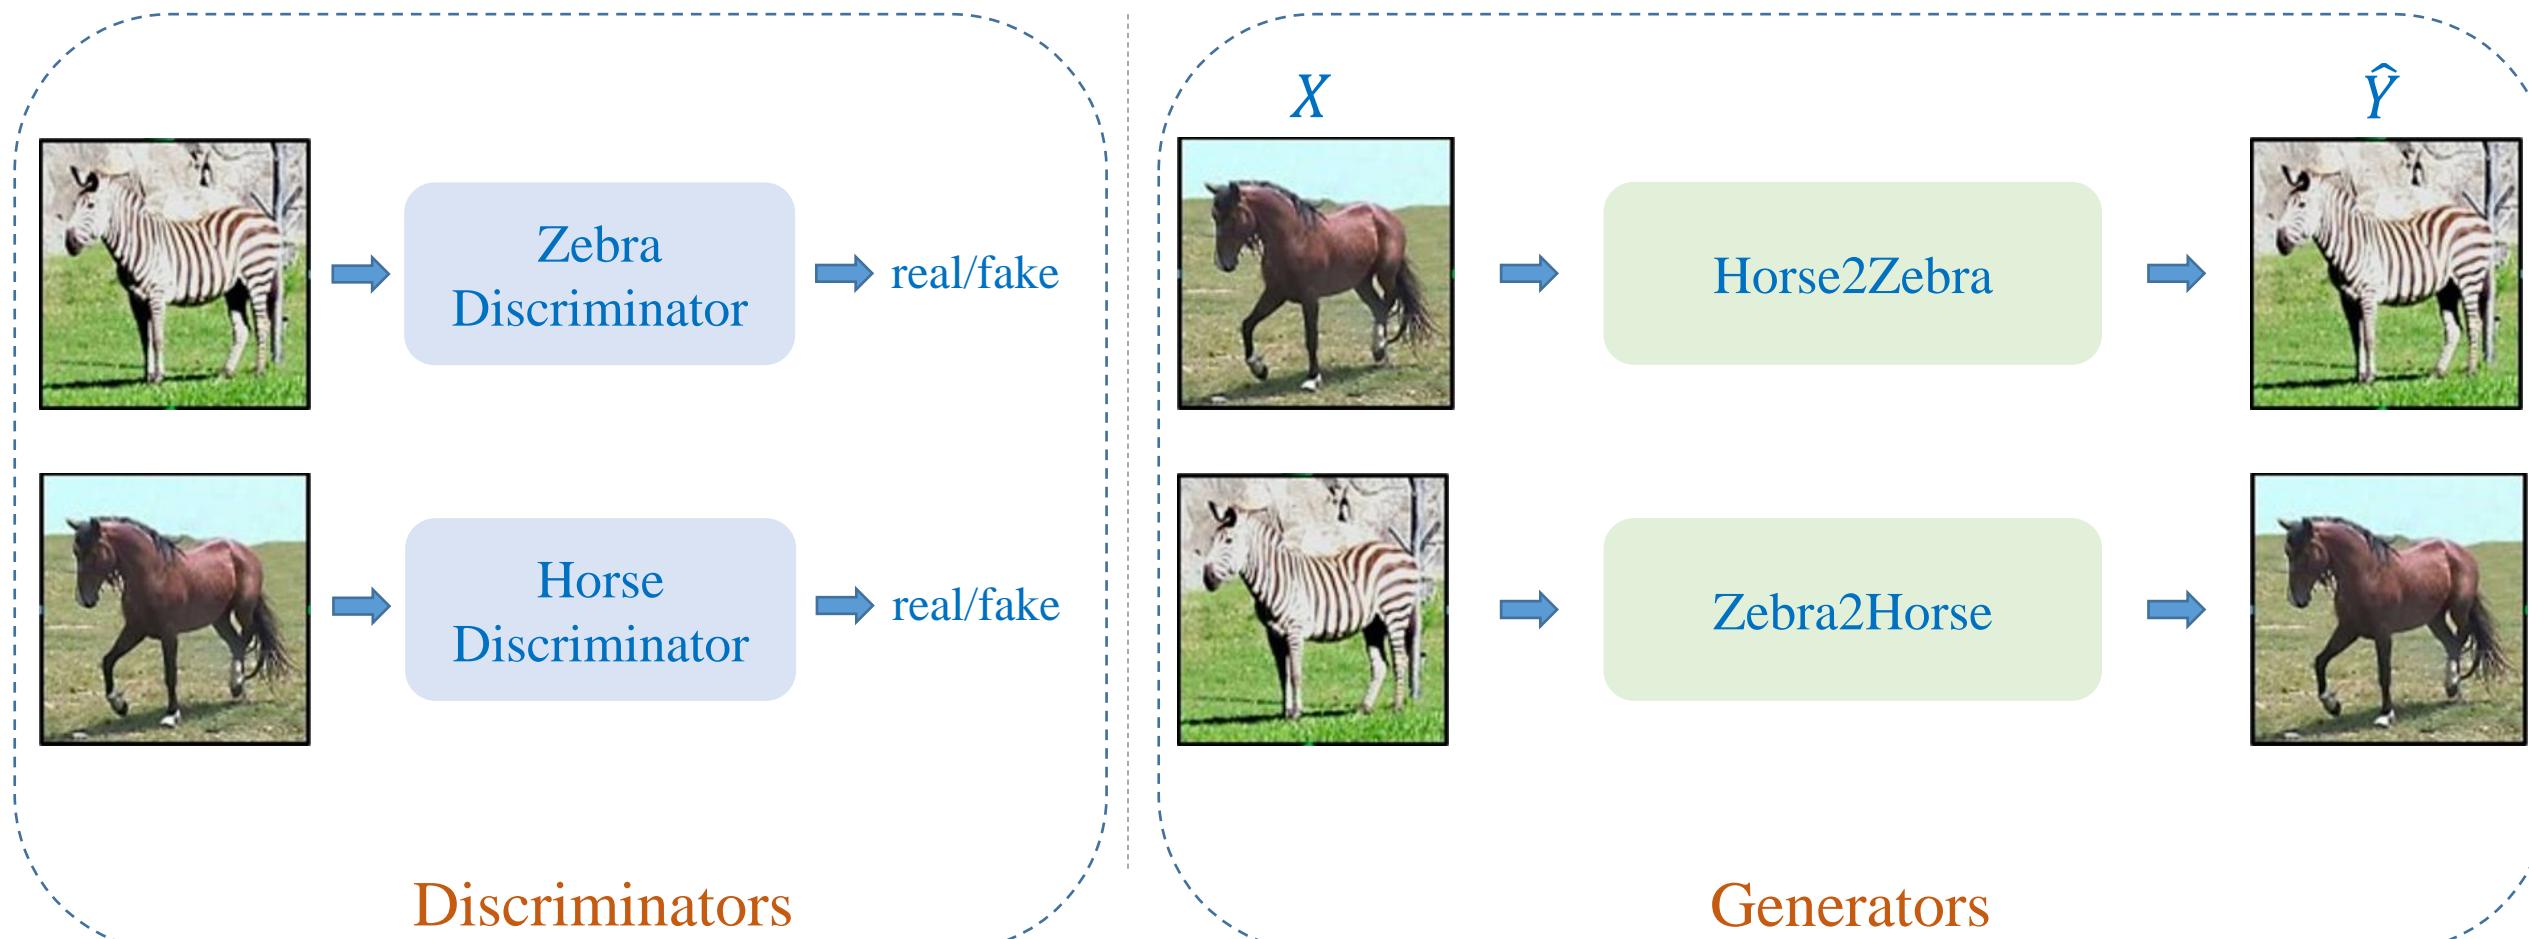


# CycleGAN

## ❖ Idea

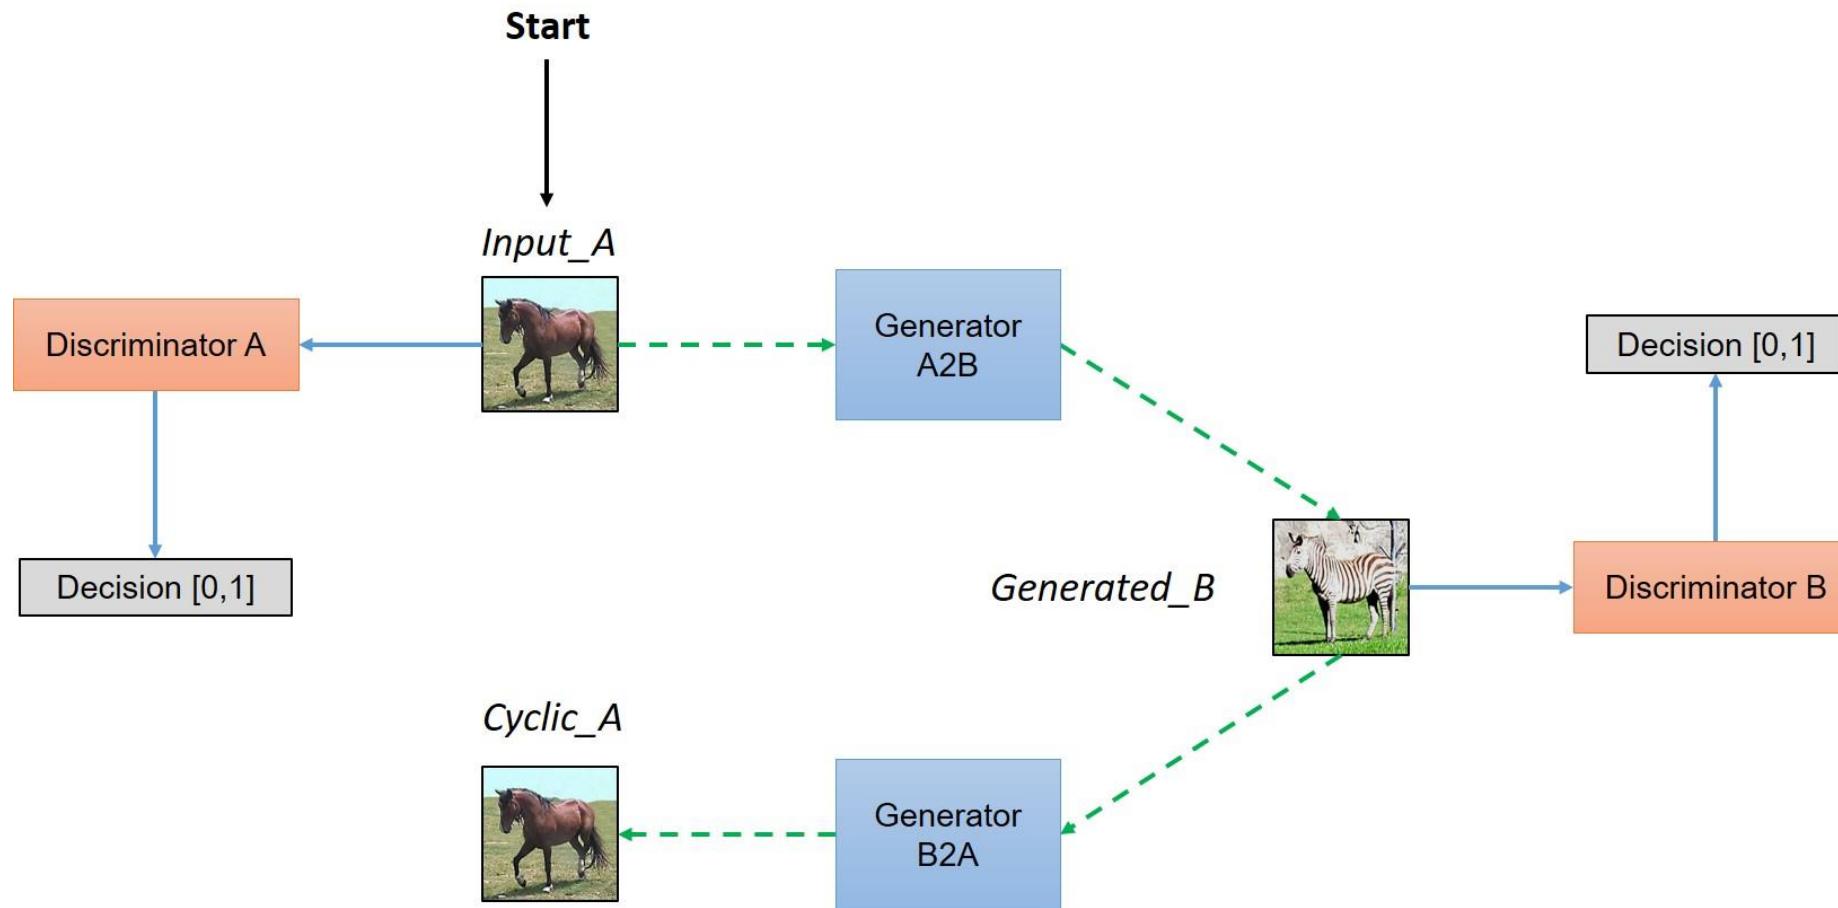
Using L1 loss?

Using GAN loss?



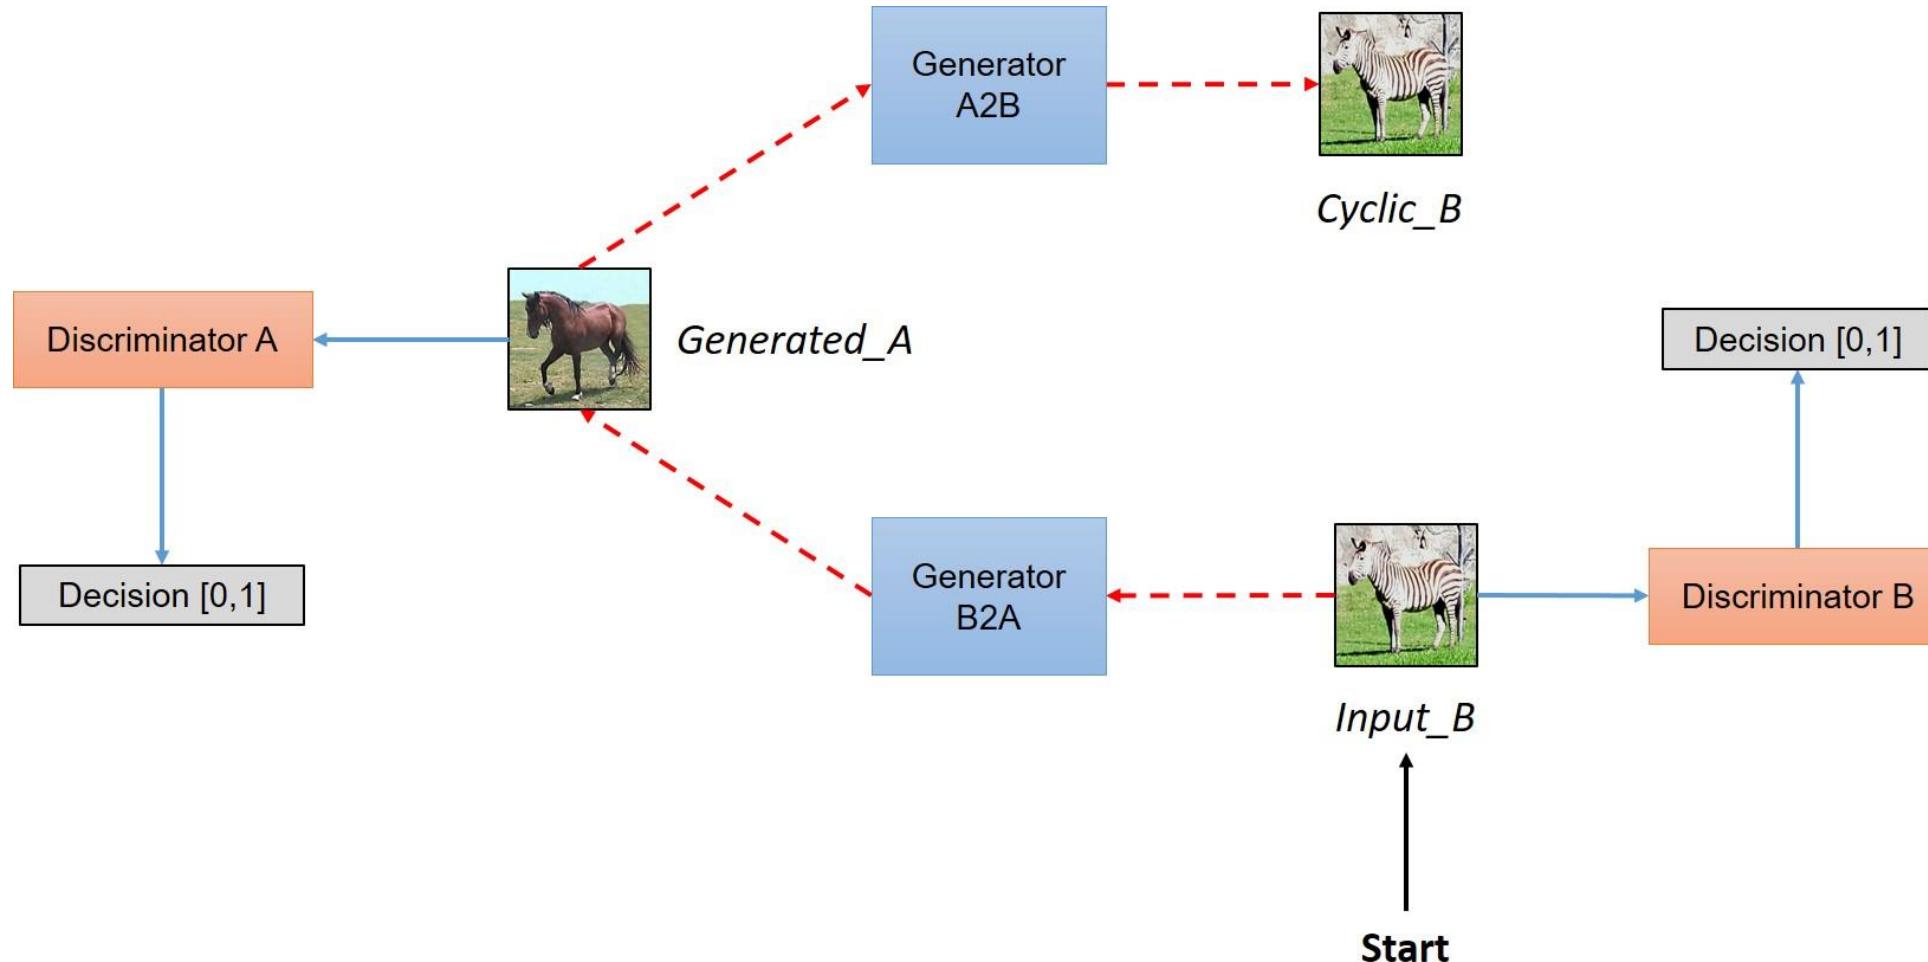
# CycleGAN

## ❖ Model



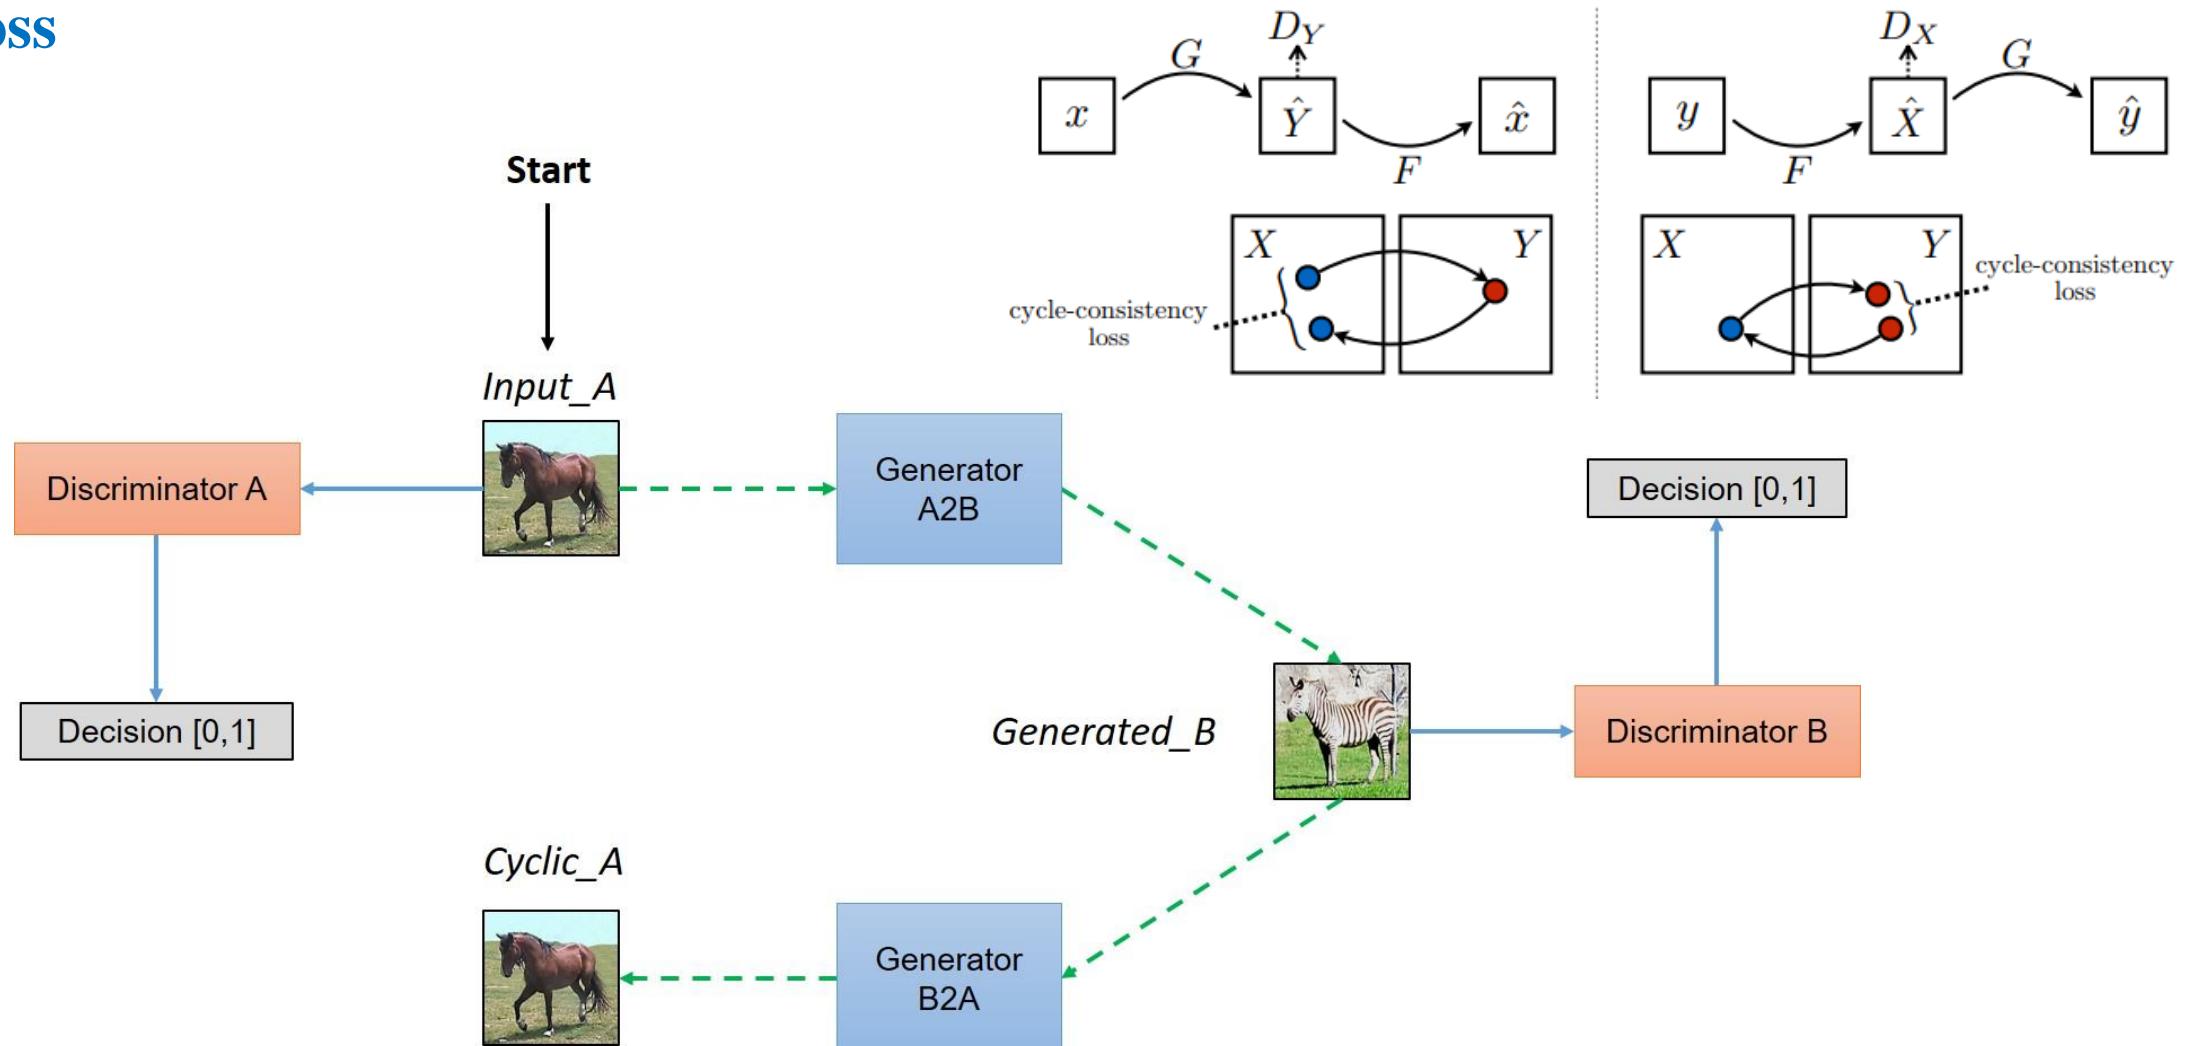
# CycleGAN

## ❖ Model



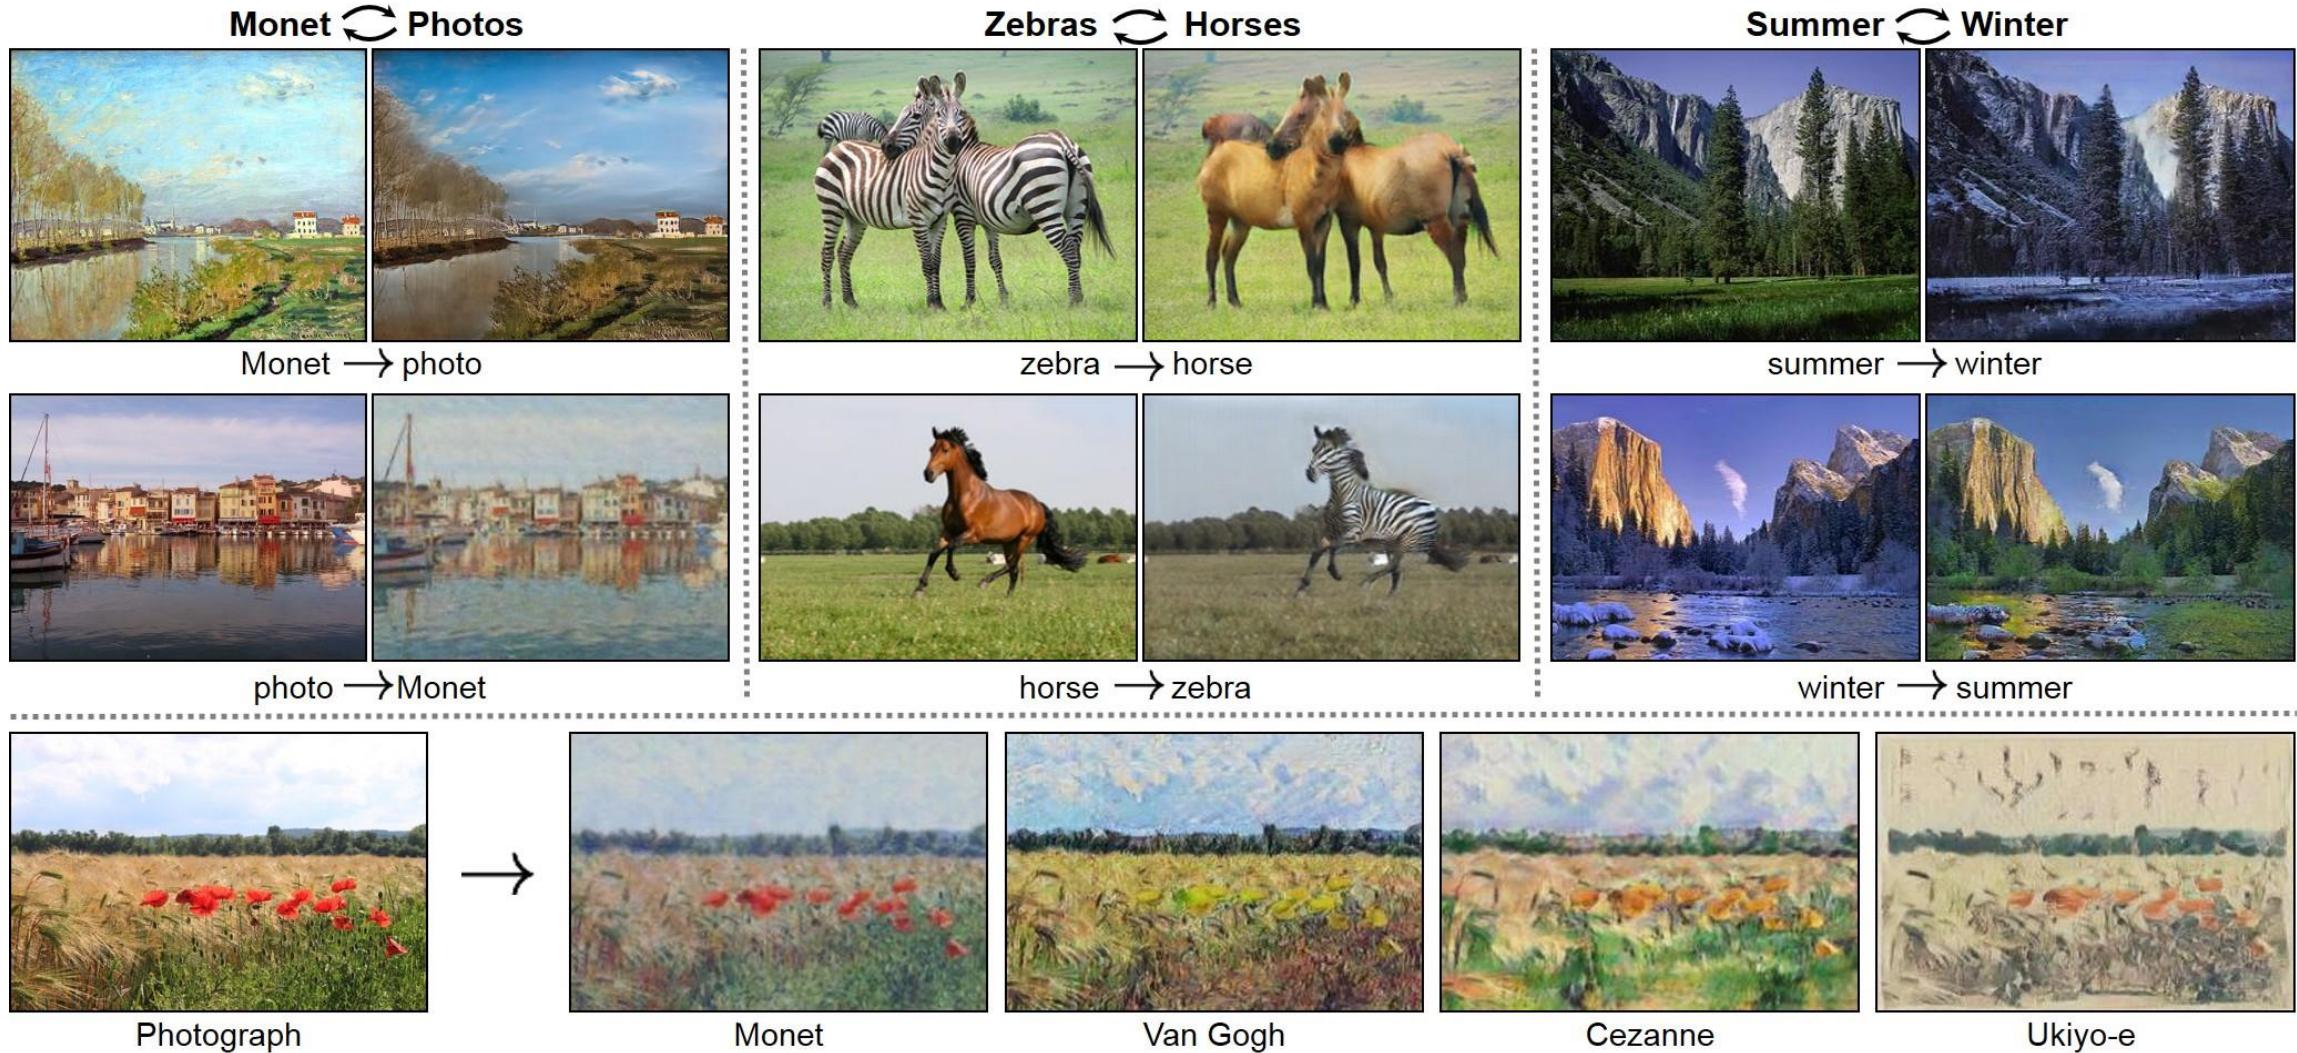
# CycleGAN

## ❖ Loss



# CycleGAN

## ❖ Applications



# Summary

- Quick Review on GAN/DCGAN
- Conditional GAN
- PatchGAN
- Pix2Pix
- CycleGAN

# Objectives

✓ GAN/DCGAN

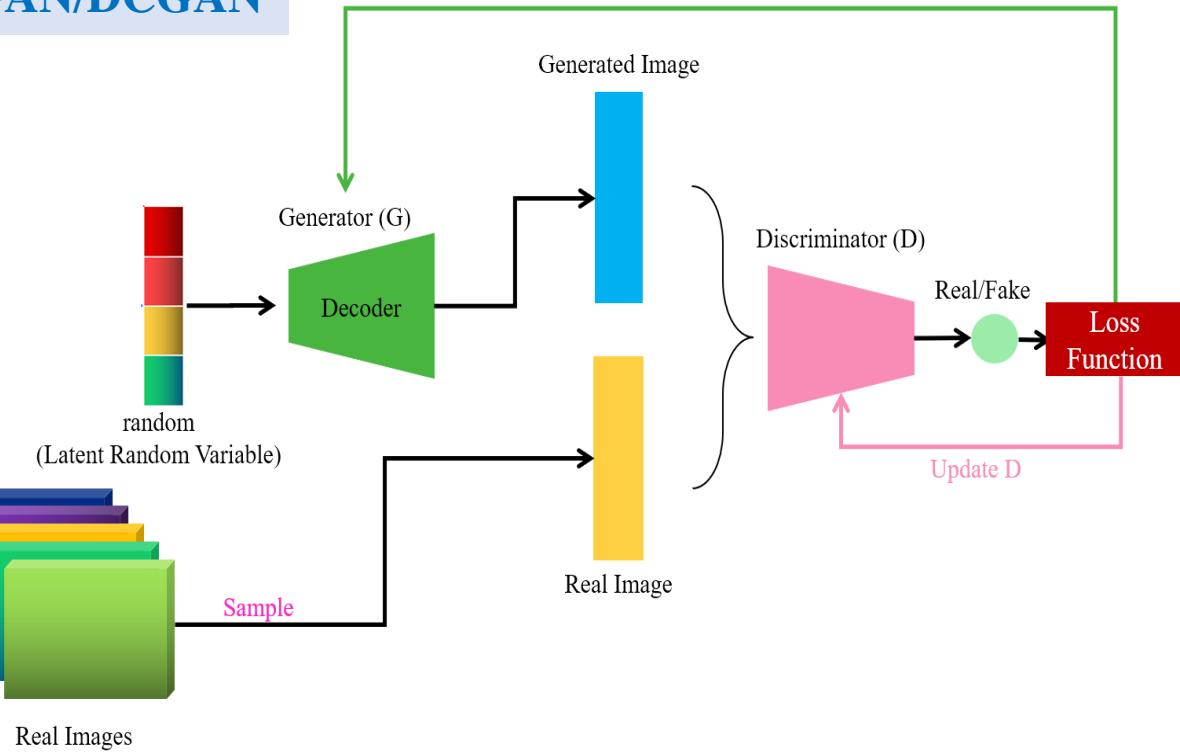
✓ Studied conditional GAN

✓ Studied PatchGAN

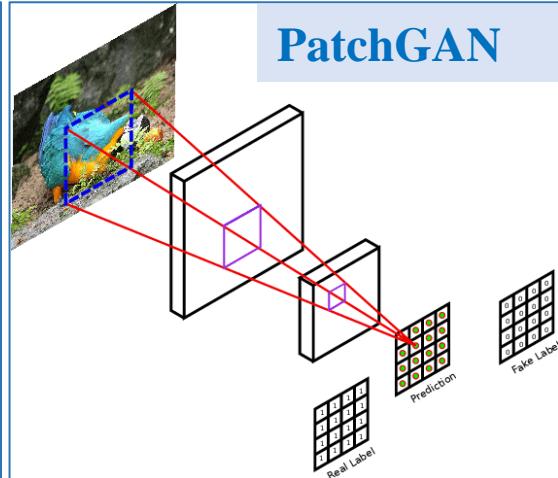
✓ Studied Pix2pix

✓ Studied CycleGAN

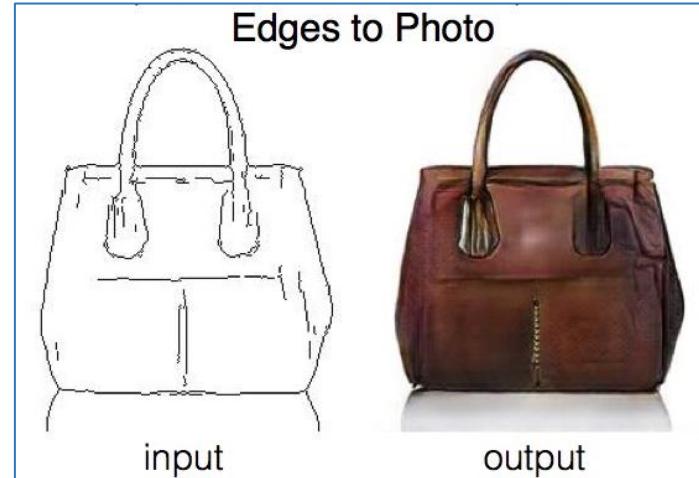
## GAN/DCGAN



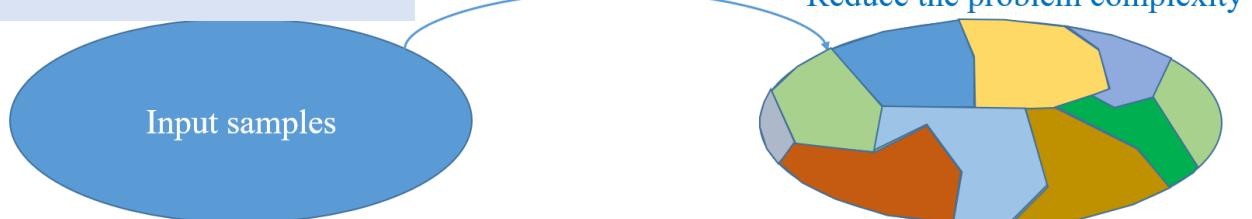
## PatchGAN



## Edges to Photo



## Conditional GAN



## GAN loss

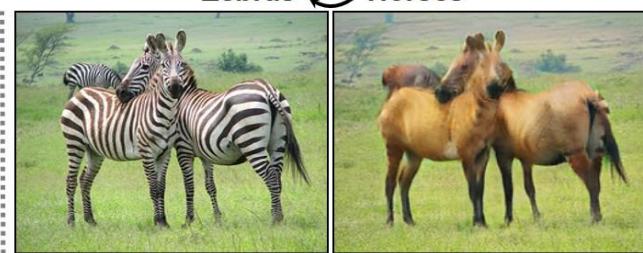
Reduce the problem complexity

## Monet ↤ Photos



Monet → photo

## Zebras ↤ Horses



zebra → horse



photo → Monet



horse → zebra

# On Sunday

## ❖ Text to Image

### Generative Adversarial Text to Image Synthesis

Scott Reed, Zeynep Akata, Xinchen Yan, Lajanugen Logeswaran  
Bernt Schiele, Honglak Lee

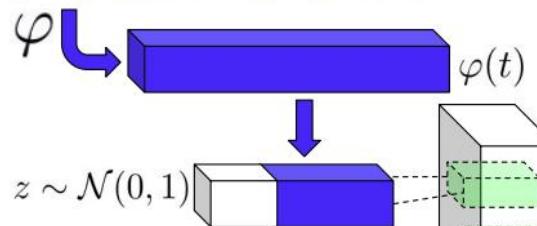
REEDSCOT<sup>1</sup>, AKATA<sup>2</sup>, XCYAN<sup>1</sup>, LLAJAN<sup>1</sup>  
SCHIELE<sup>2</sup>, HONGLAK<sup>1</sup>

<sup>1</sup> University of Michigan, Ann Arbor, MI, USA (UMICH.EDU)

<sup>2</sup> Max Planck Institute for Informatics, Saarbrücken, Germany (MPI-INF.MPG.DE)

### Generative Adversarial Text to Image Synthesis

*This flower has small, round violet petals with a dark purple center*



$$\hat{x} := G(z, \varphi(t))$$

Generator Network

this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



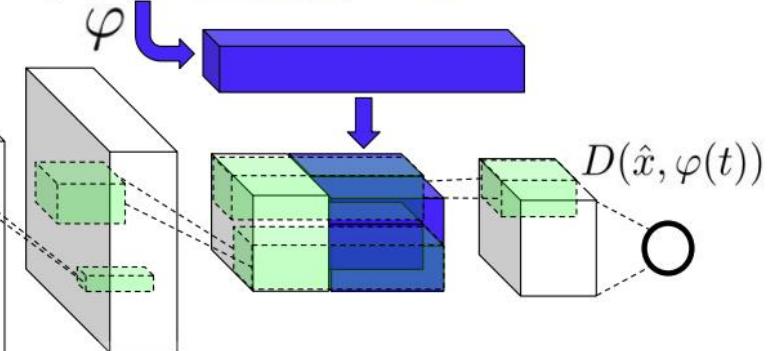
the flower has petals that are bright pinkish purple with white stigma



this white and yellow flower have thin white petals and a round yellow stamen



*This flower has small, round violet petals with a dark purple center*



Discriminator Network

