

AI VIETNAM
All-in-One Course

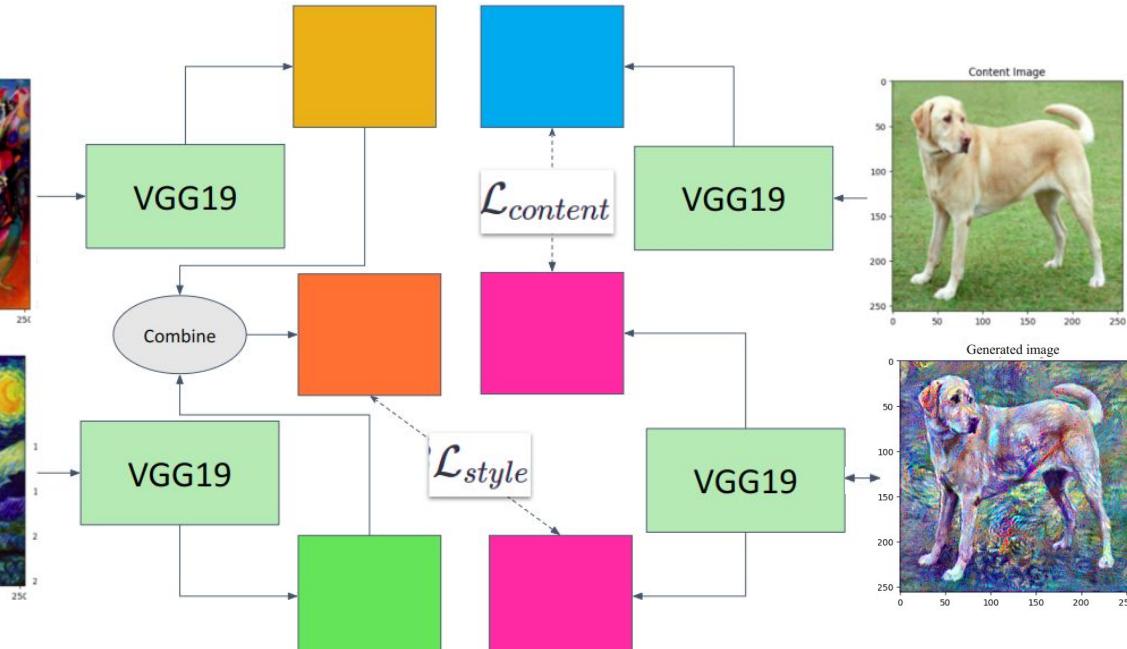
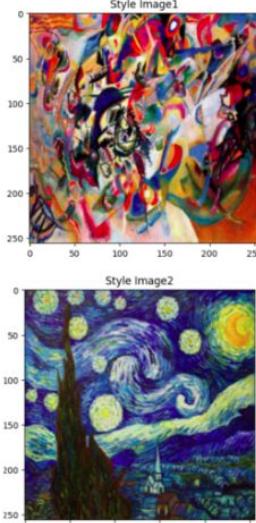
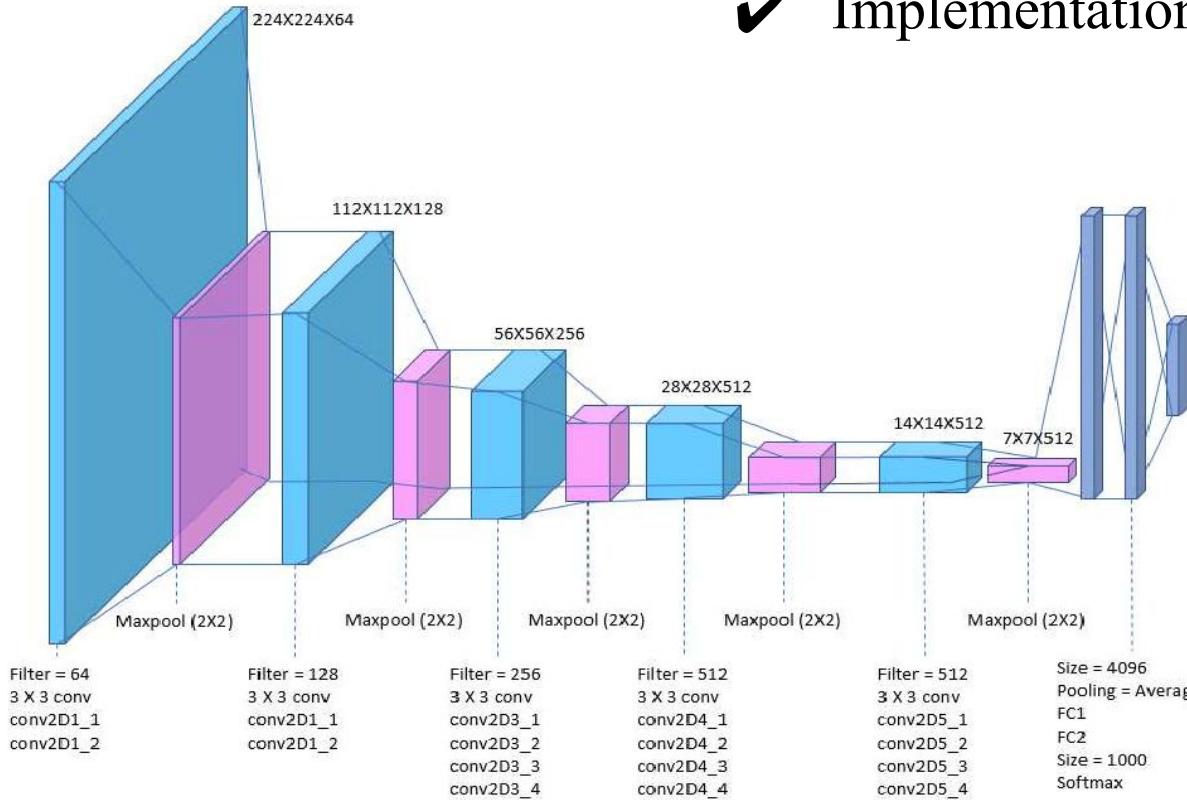
Neural Style Transfer

(TA Session)

Khoa Nguyen
Ph.D. Student

Objectives

- ✓ Review Style Transfer (Definition, ...)
- ✓ Review requisite knowledge
- ✓ Introduce Style Transfer Examples
- ✓ Introduce Milestone Papers
- ✓ Implementation of Style Transfer with 2 Style Images
- ✓ Implementation of Style Transfer with Rotation + DE



Outline

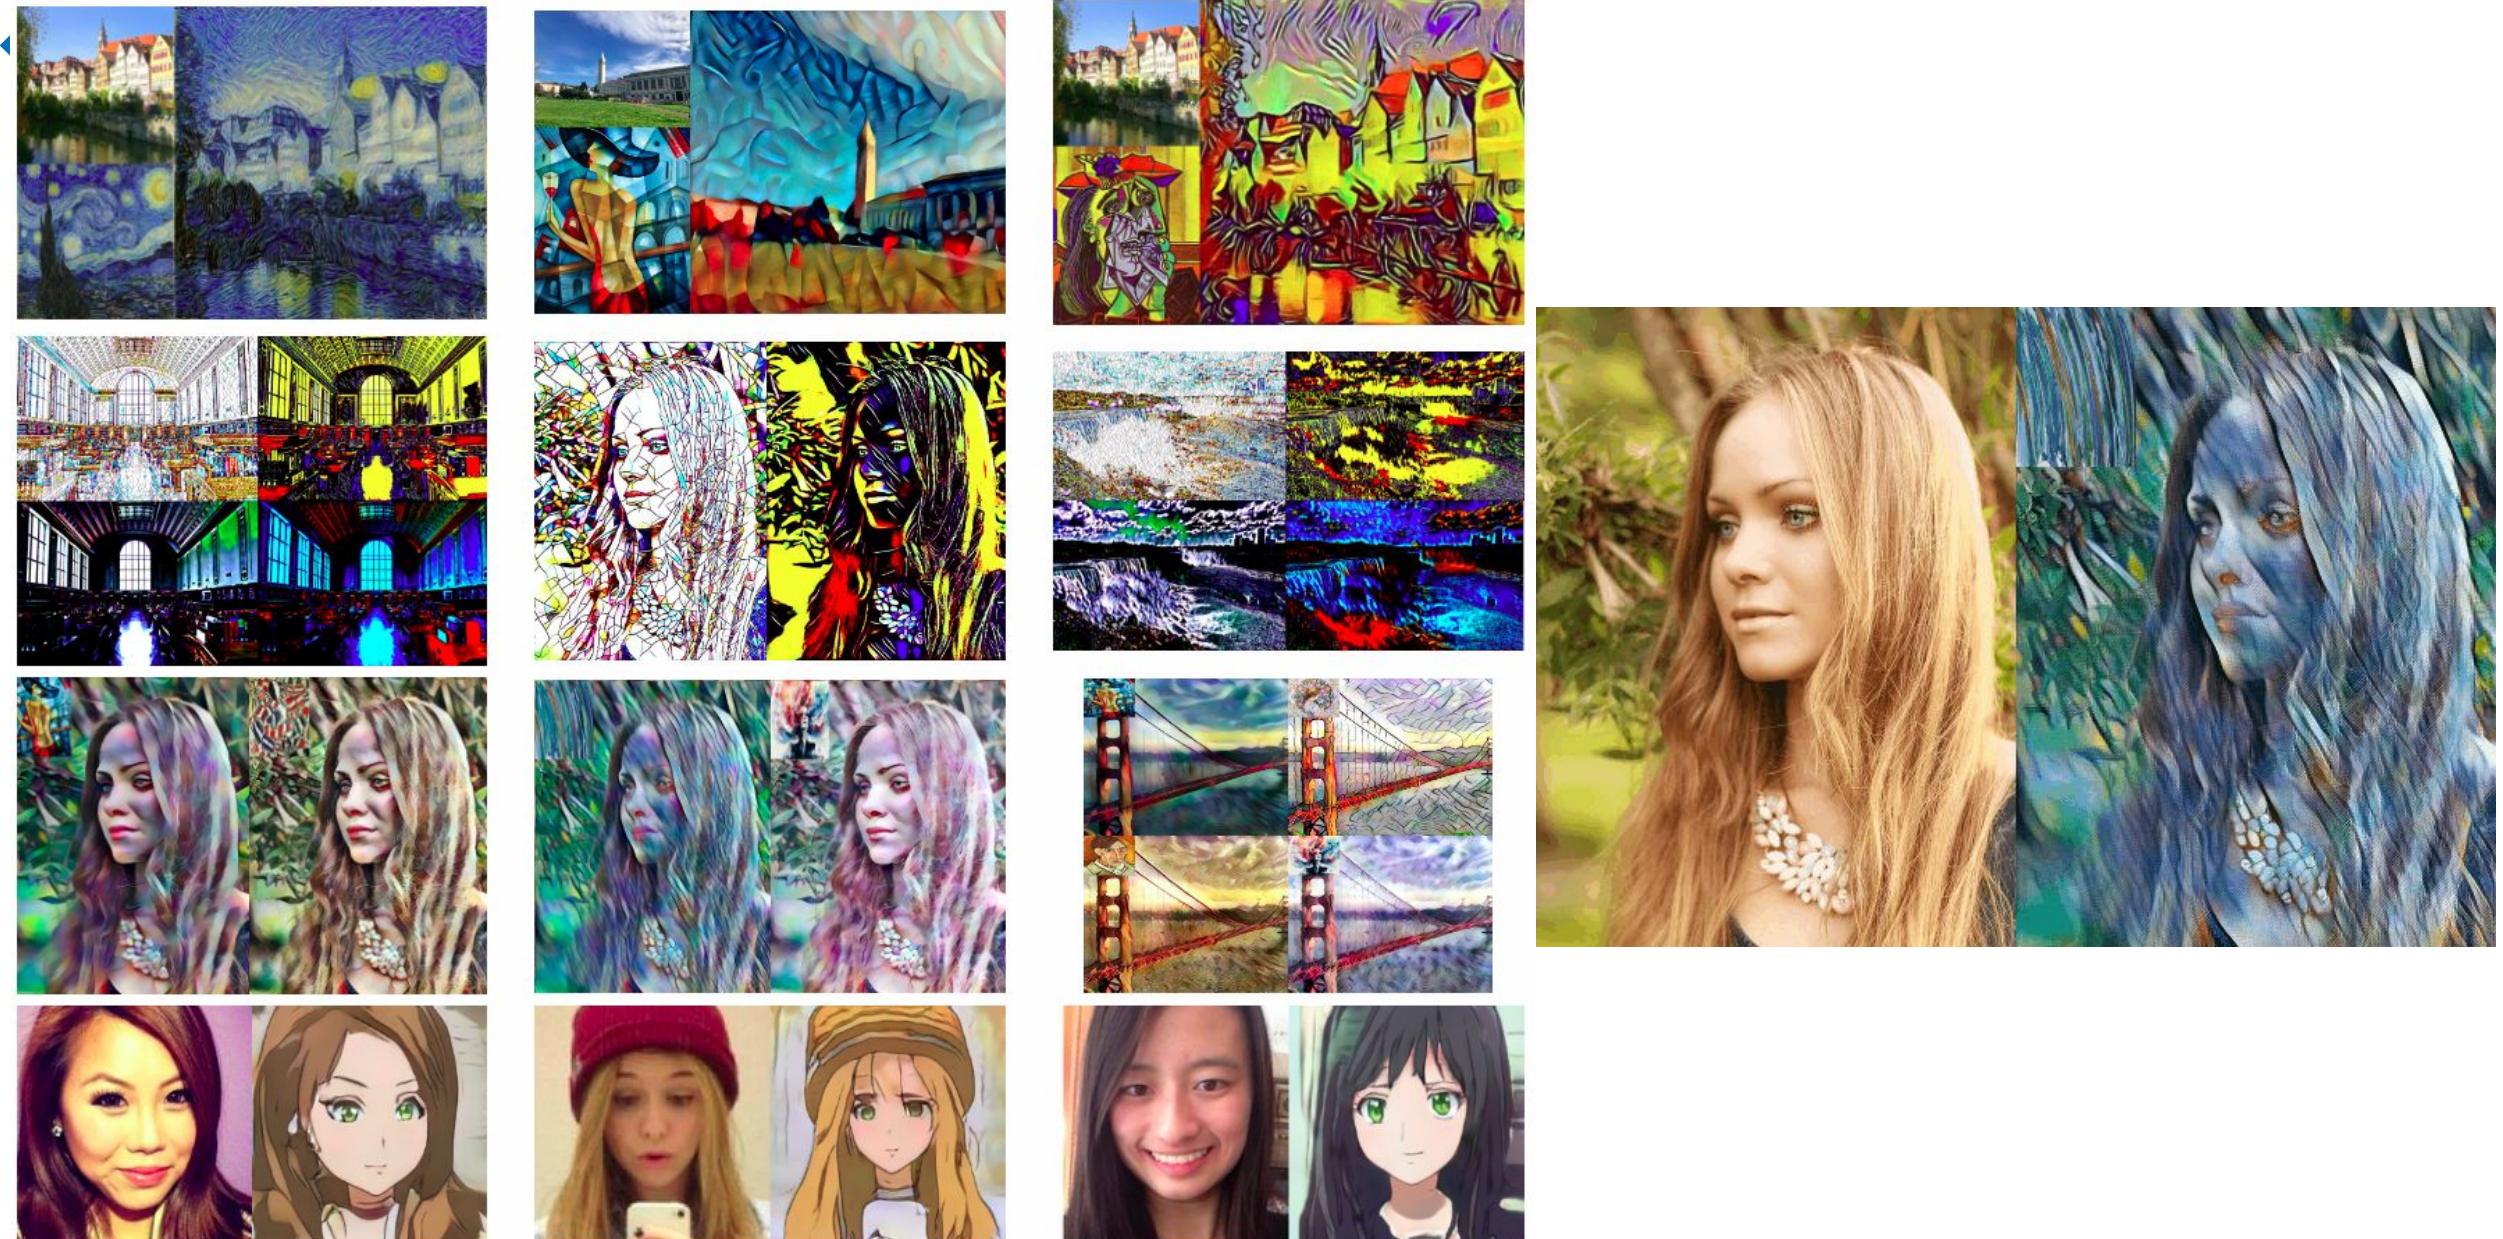
- Introduction
- Basic Knowledge
- Style Transfer Example
- Milestone Papers
- Style Transfer with 2 Style Images
- Style Transfer with Rotation + DE

Introduction

❖ Definition



Introduction



Introduction

◆ Definition

Neural Style Transfer is a fascinating technique in the field of **computer vision** and **artificial intelligence** that allows the **style of one image** to be **applied** to the **content of another image**, effectively blending the two. This is achieved through the use of Convolutional Neural Networks (CNNs), which are a class of deep neural networks commonly used in analyzing visual imagery.

Content Image: This is the primary image to which you want to apply a new style. It provides the structure and content.

Style Image: This image provides the artistic style you wish to apply to the content image. It could be a famous painting or any other image with a distinctive artistic style.

Generated Image: This is the output image that starts as a copy of the content image and iteratively transforms to match the style of the style image while retaining the content of the original content image.

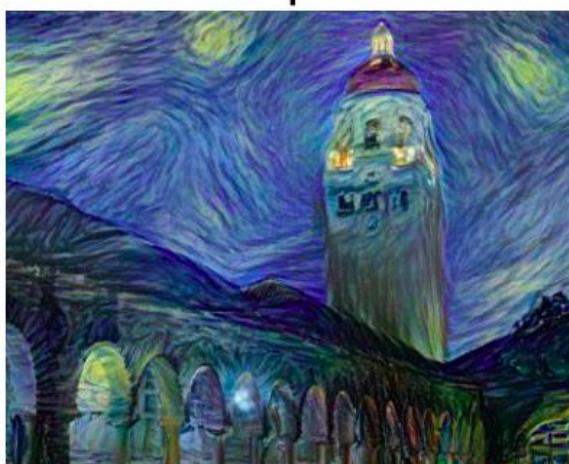
Introduction

◆ Definition



Content

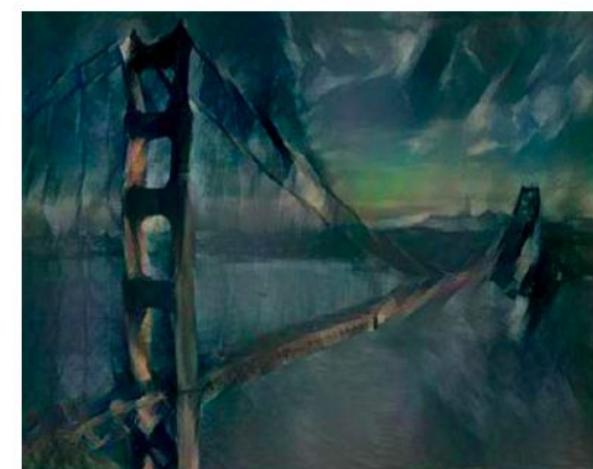
Style



Generated image

Content

Style



Generated image

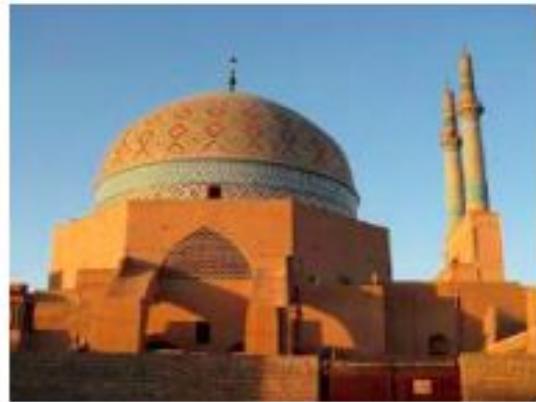
AI VIETNAM
All-in-One Course

Basic Knowledge

Basic Knowledge

❖ Generate Image

Content



Style



Initialize Random Image
- $I : C \times H \times W$

Use gradient descent to minimize the loss function

Basic Knowledge

❖ Neural Style Transfer

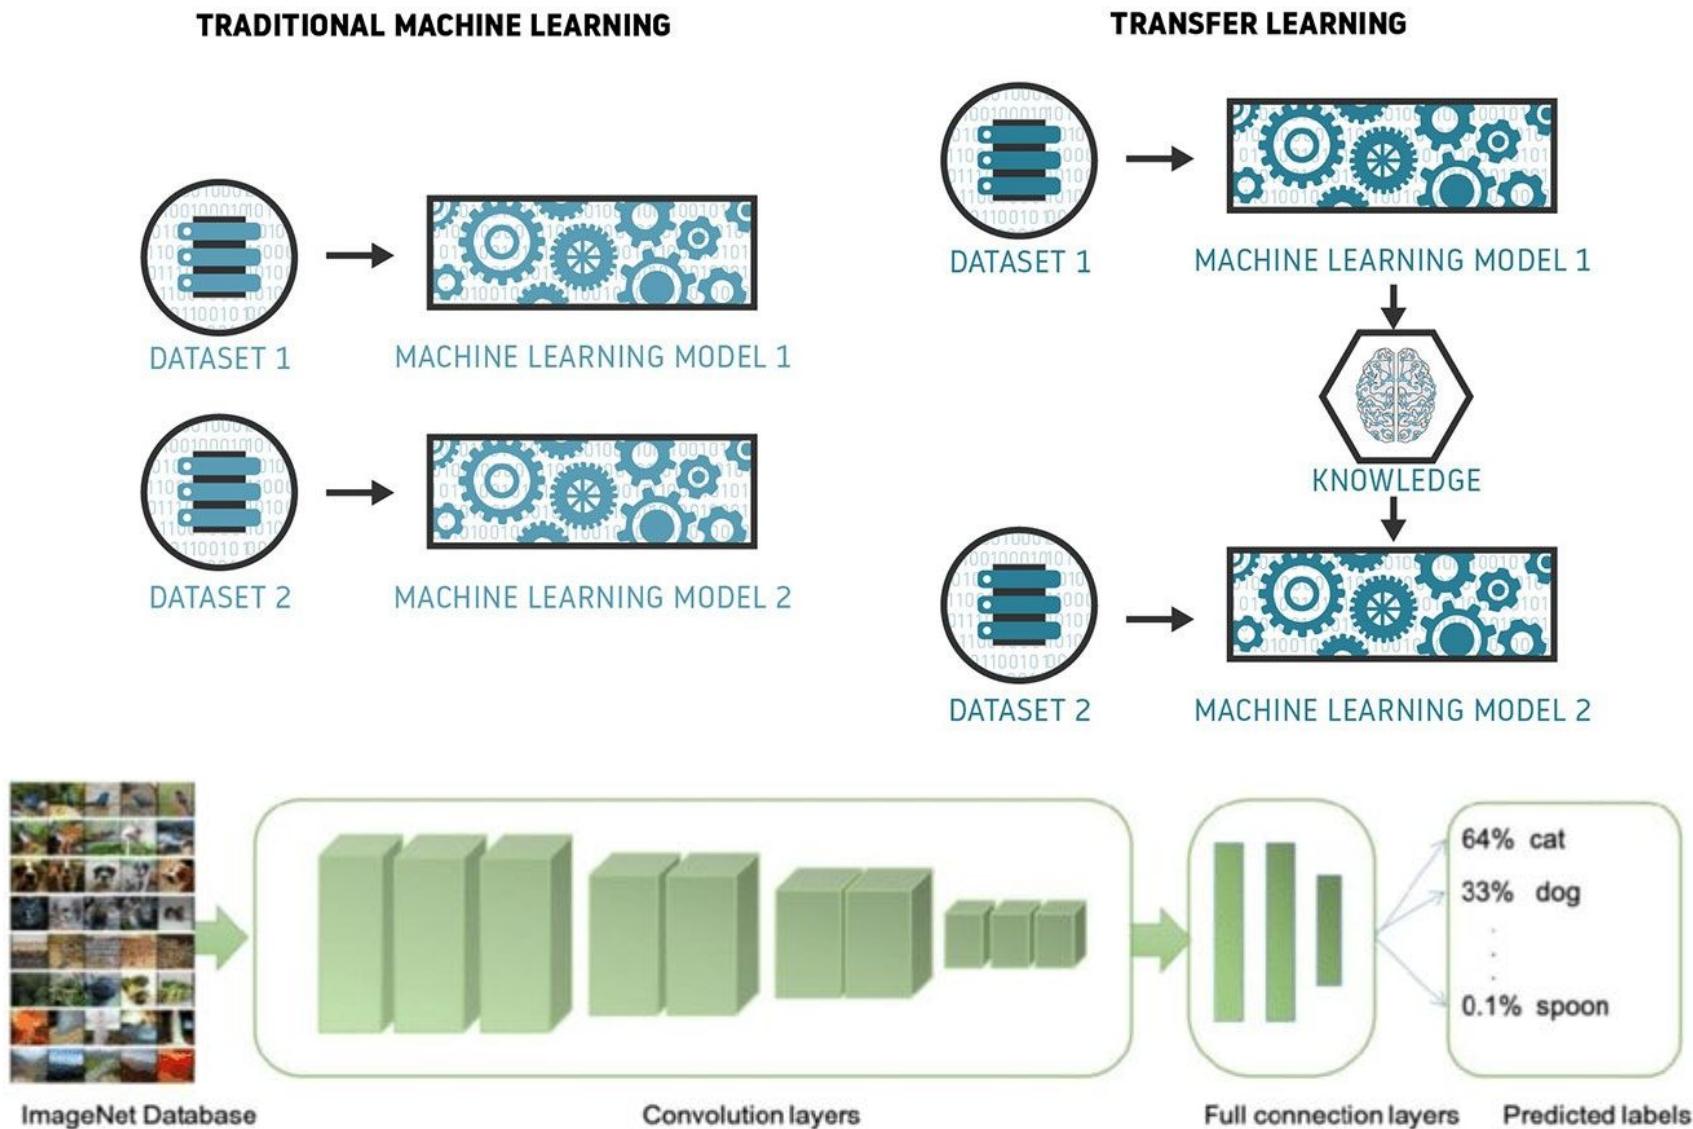


$$\mathcal{L}_{total} = \alpha \mathcal{L}_{content} + \beta \mathcal{L}_{style}$$

using CNN representation (VGG)

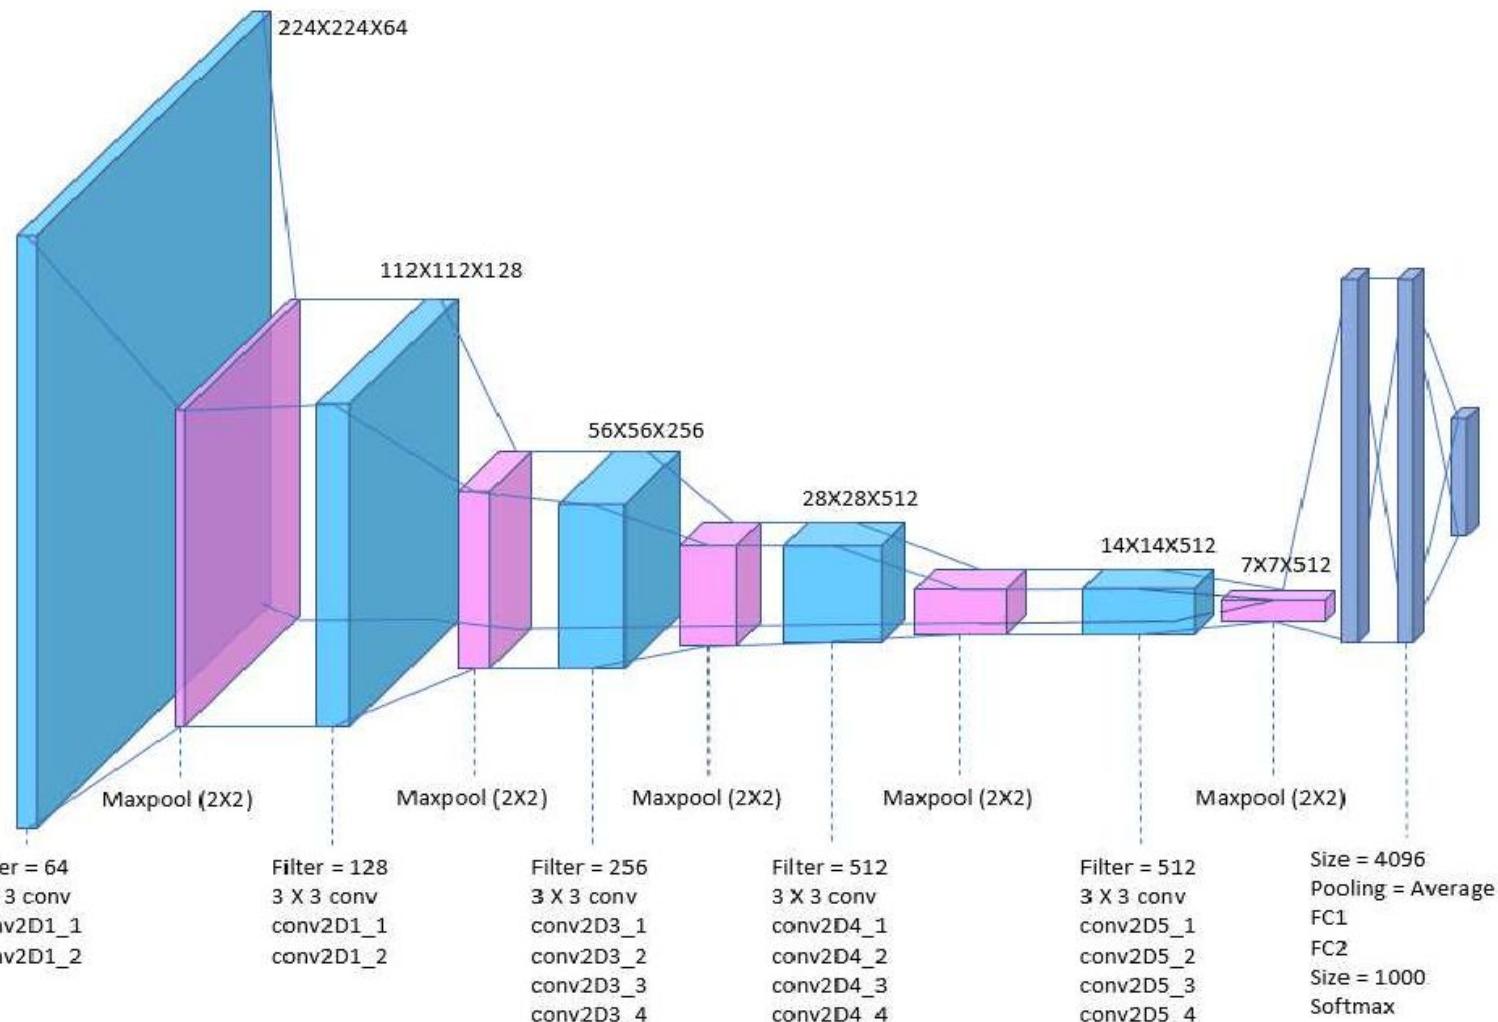
Basic Knowledge

◆ Pre-trained model



Basic Knowledge

❖ VGG



Sofmax	
FC 1000	
FC 4096	
FC 4096	
Pool	
3x3 conv,512	
3x3 conv,128	
3x3 conv,128	
Pool	
3x3 conv,512	
Input	
VGG16	
VGG19	

Basic Knowledge

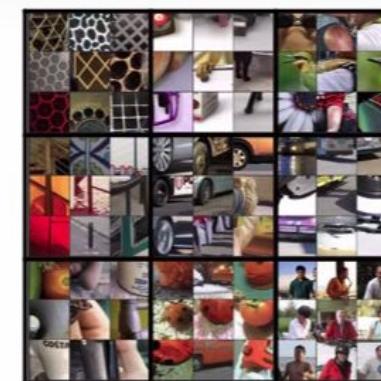
◆ Visualizing deep layers: Layer 1



Layer 1



Layer 2



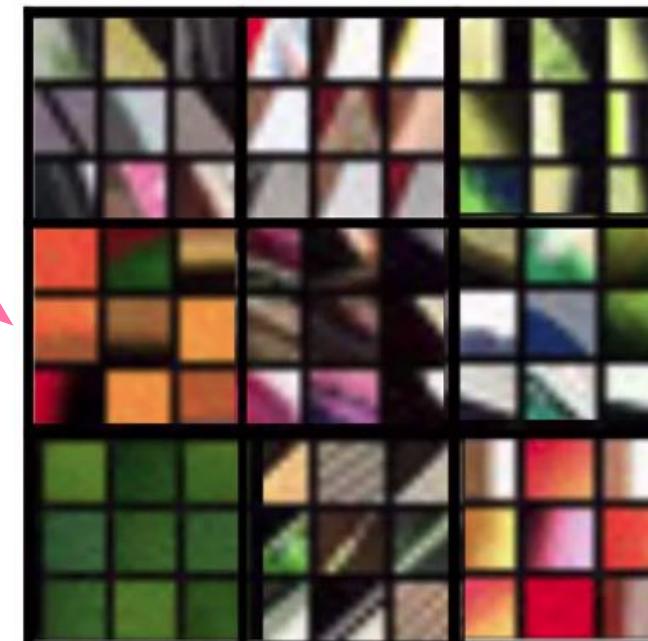
Layer 3



Layer 4



Layer 5



Basic Knowledge

❖ Visualizing deep layers: Layer 2



Layer 1



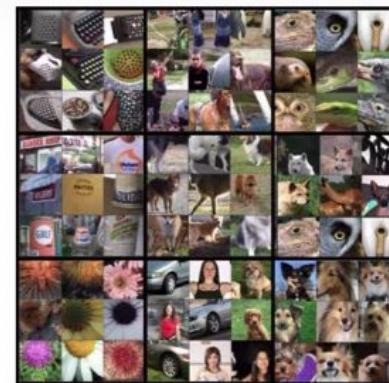
Layer 2



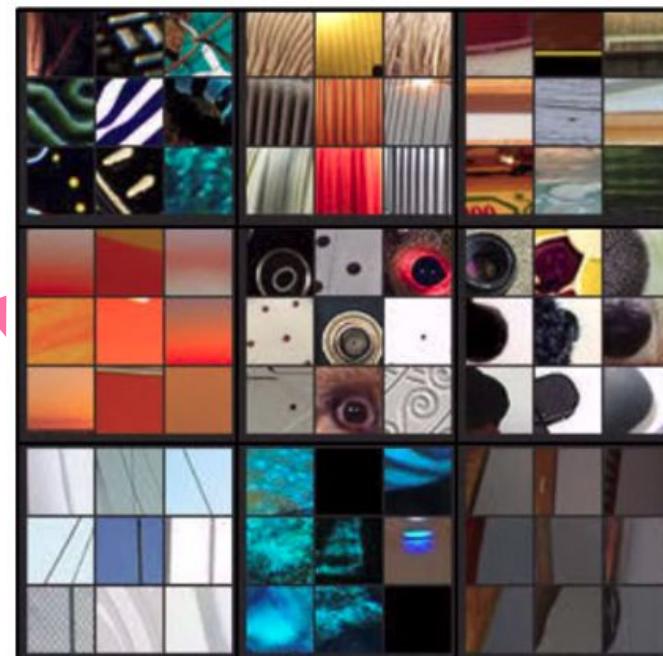
Layer 3



Layer 4



Layer 5



Basic Knowledge

❖ Visualizing deep layers: Layer 3



Layer 1



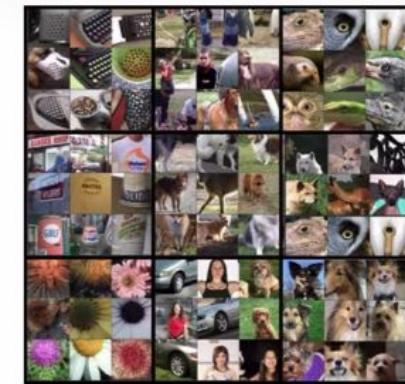
Layer 2



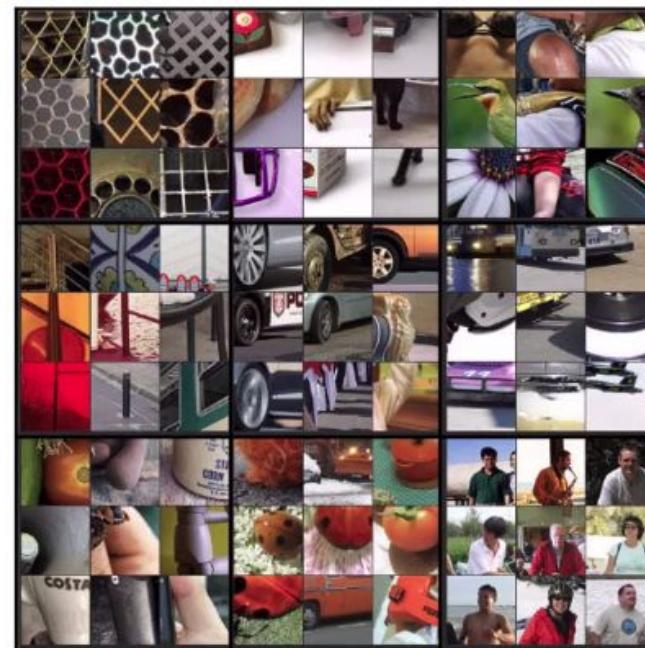
Layer 3



Layer 4



Layer 5



Basic Knowledge

❖ Visualizing deep layers: Layer 4



Layer 1



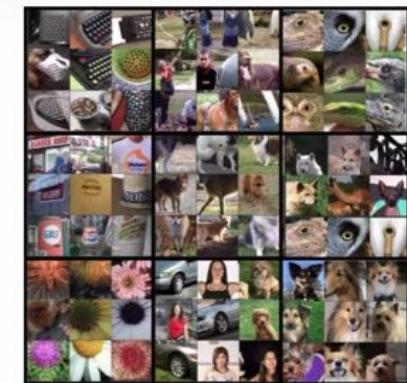
Layer 2



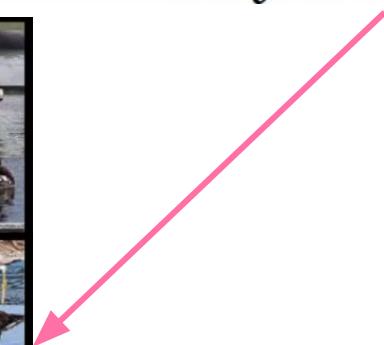
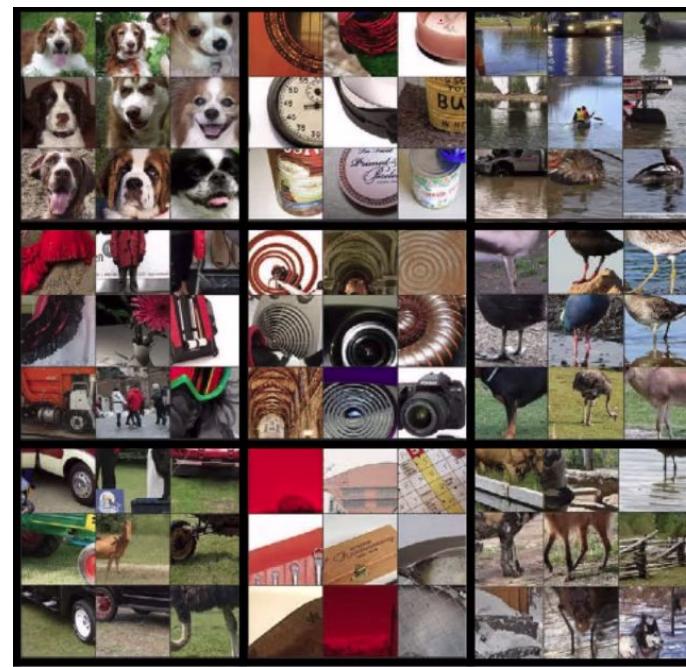
Layer 3



Layer 4



Layer 5



Basic Knowledge

❖ Visualizing deep layers: Layer 5



Layer 1



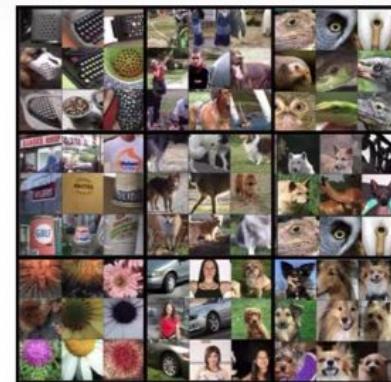
Layer 2



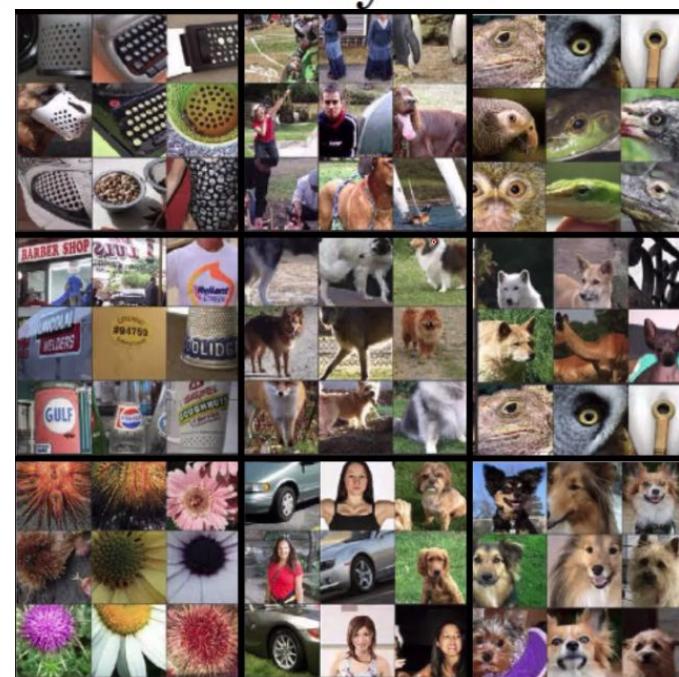
Layer 3



Layer 4

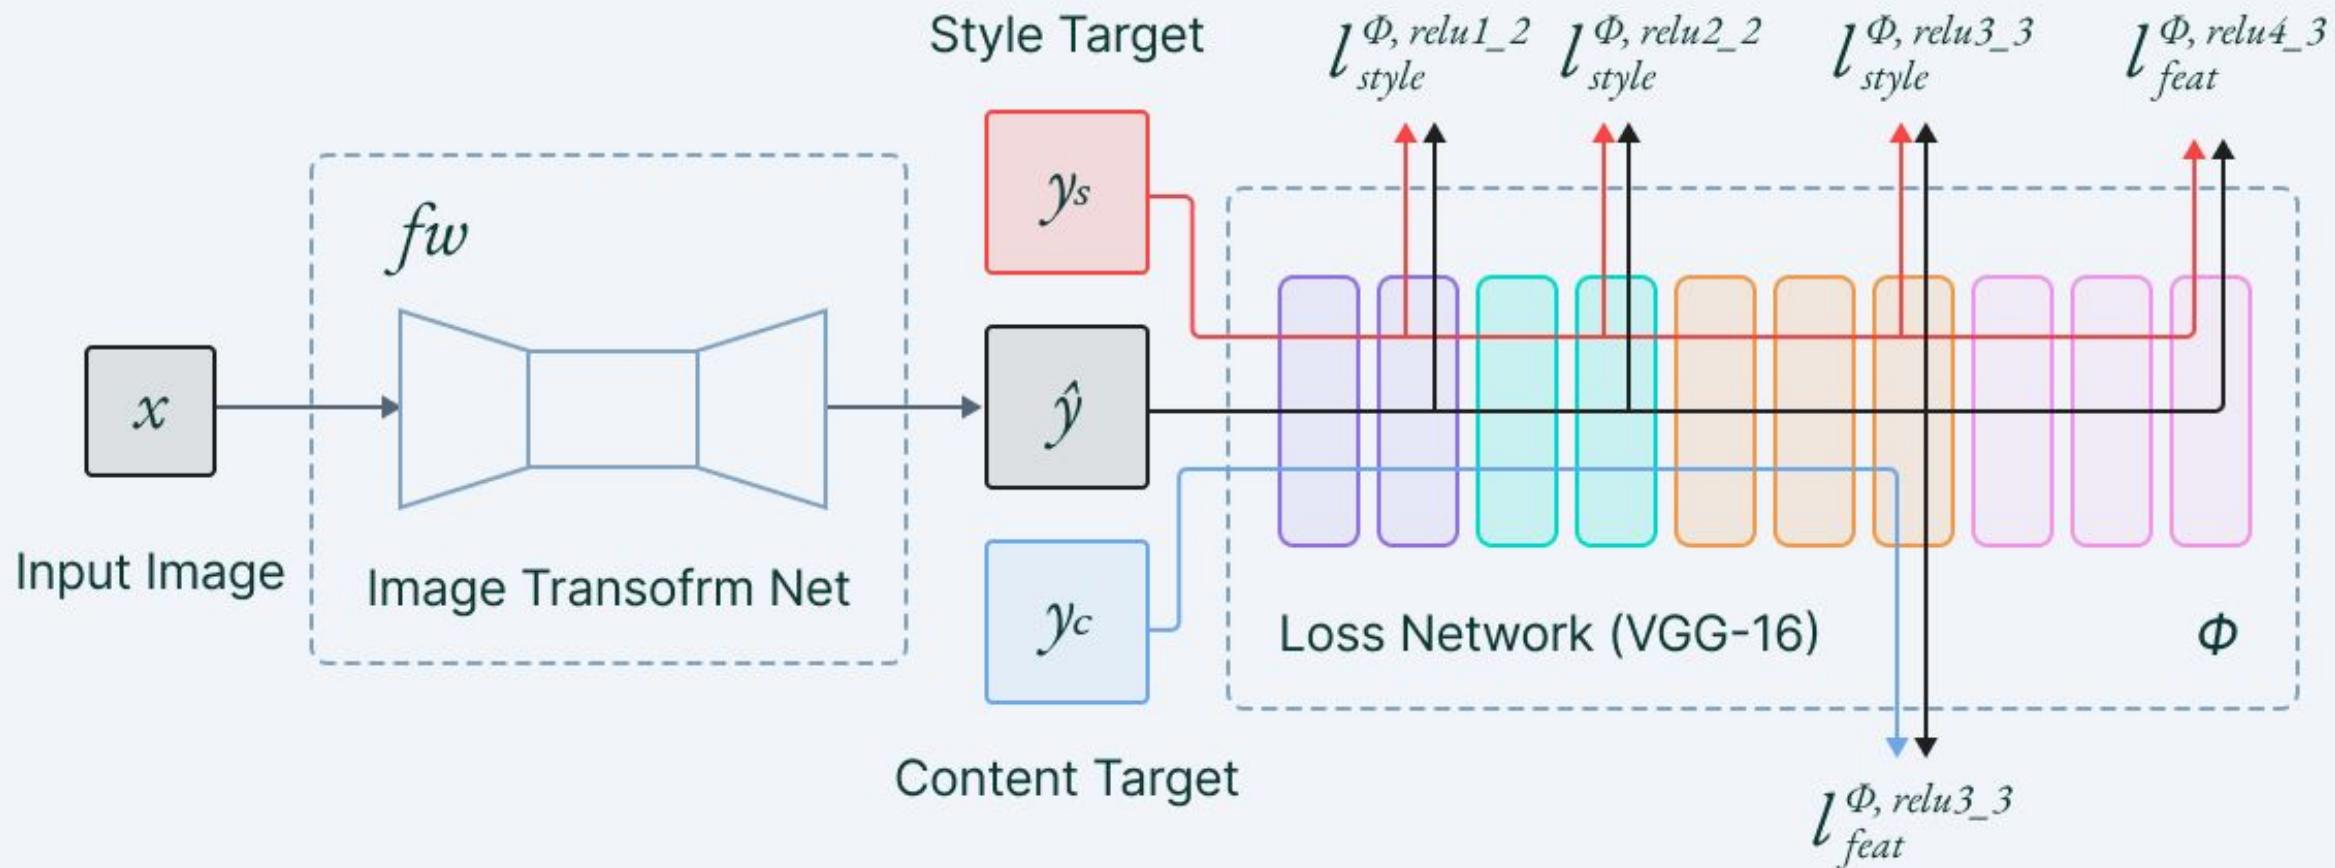


Layer 5



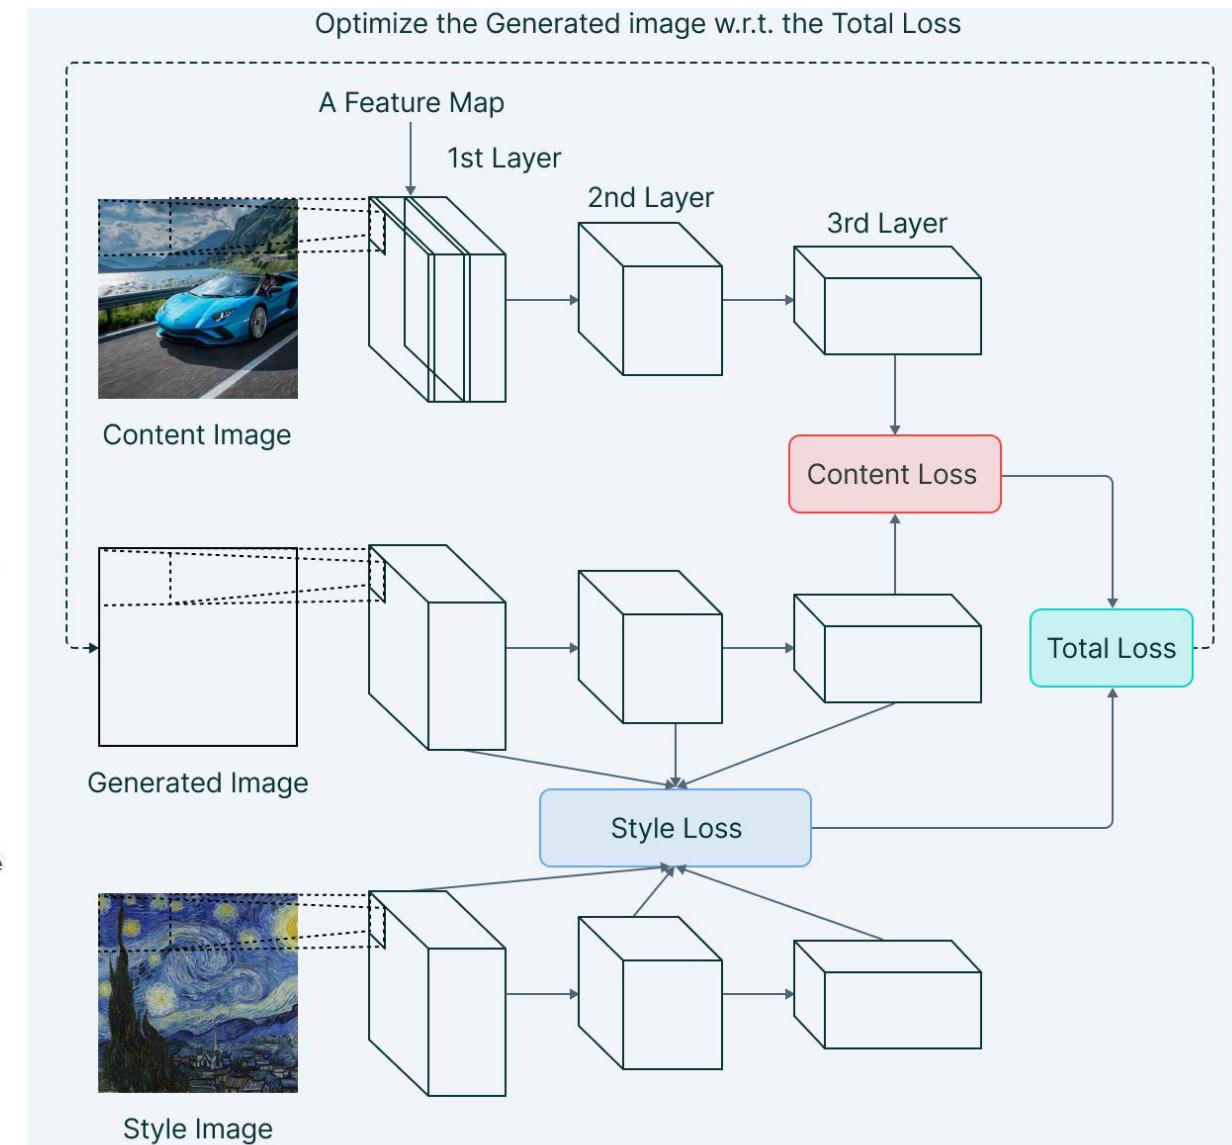
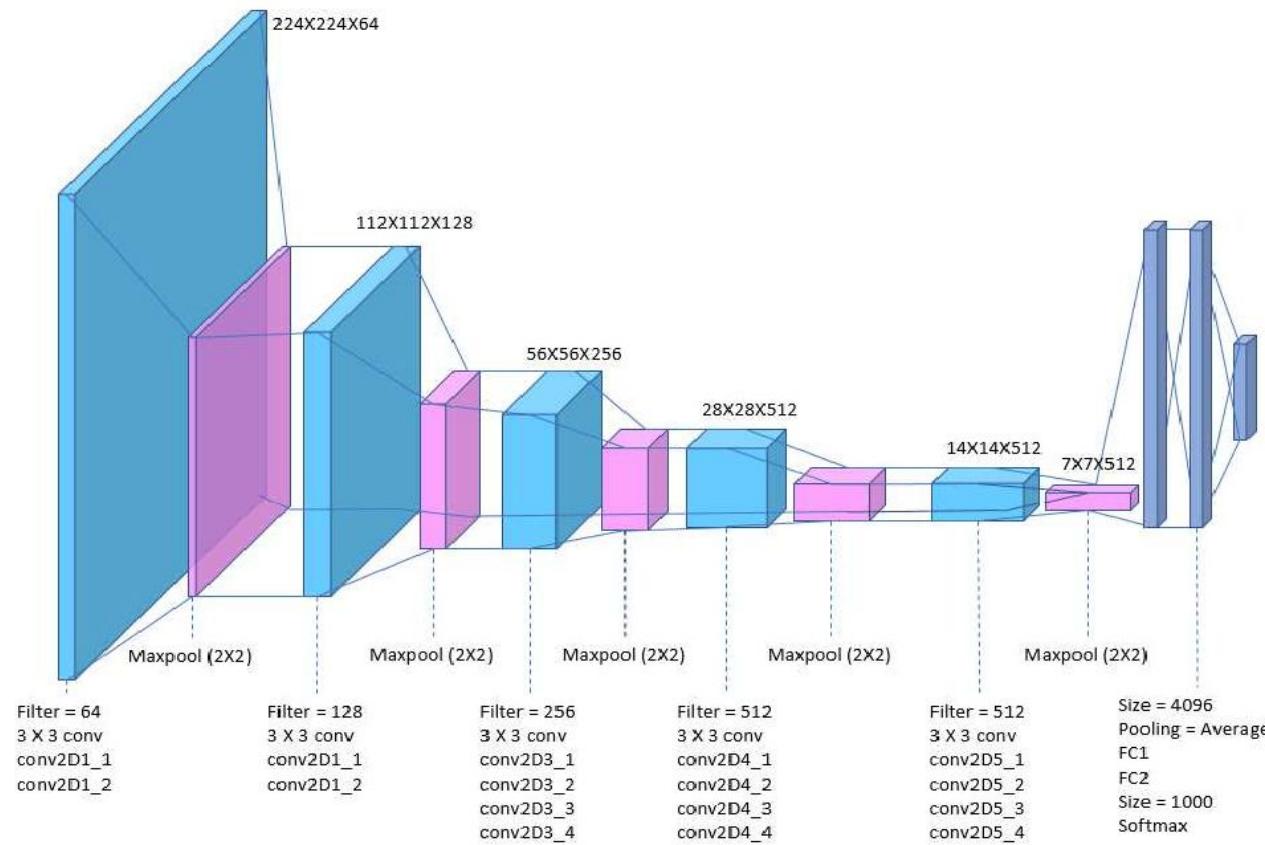
Basic Knowledge

❖ Neural Style Transfer basic structure



Basic Knowledge

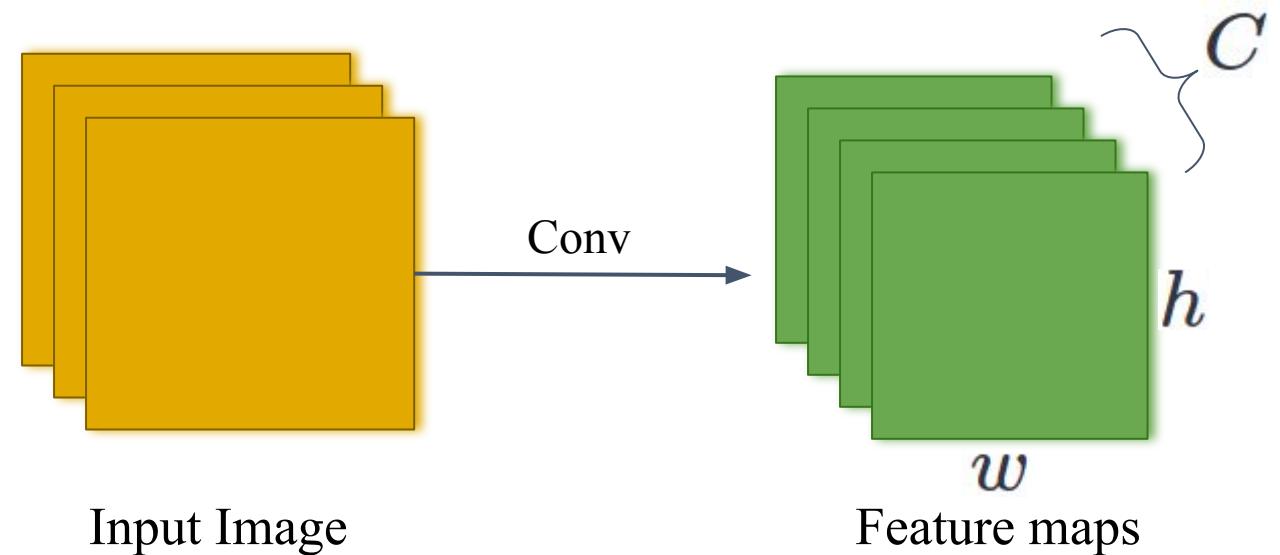
❖ Neural Style Transfer basic structure



Basic Knowledge

◆ Content Loss

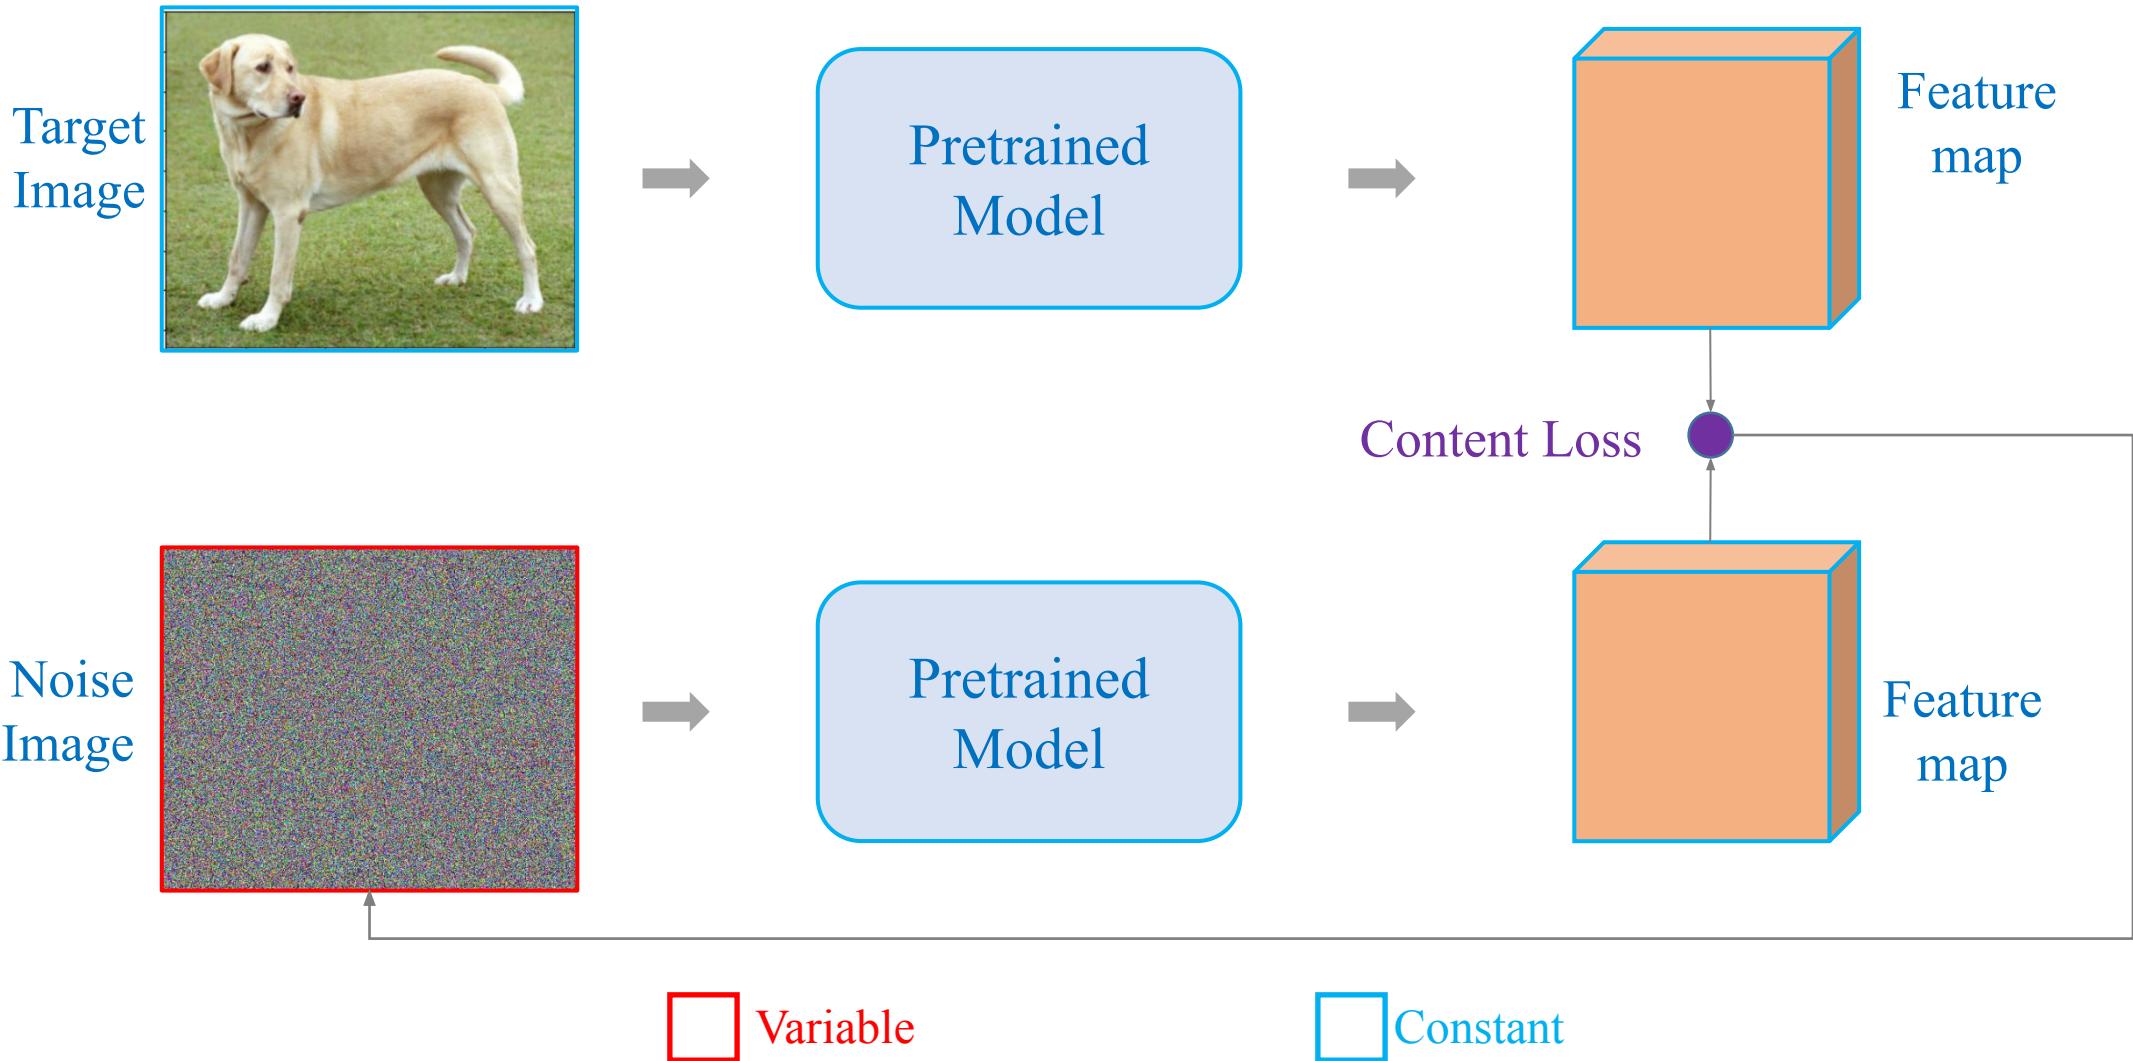
- Layer l
- Number of filters C
- Generated image x
- Generated feature map F
- Content image p
- Content feature map P



$$\mathcal{L}_{content}(p, x, l) = \frac{1}{2} \sum_{i,j,k} (F_{ijk}^l - P_{ijk}^l)^2$$

Basic Knowledge

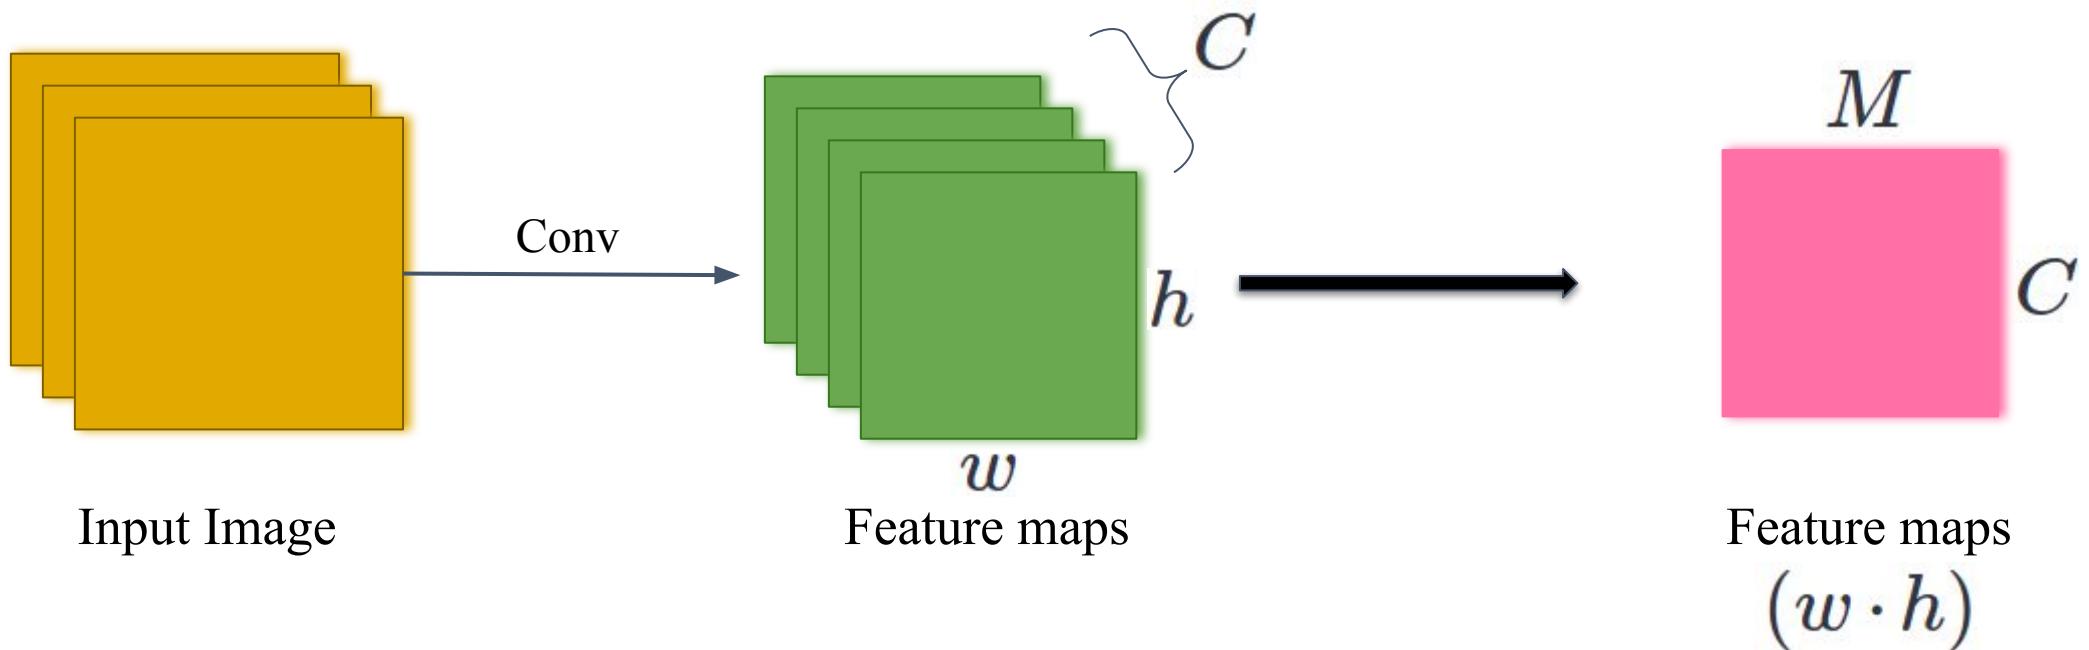
◆ Content Loss



Basic Knowledge

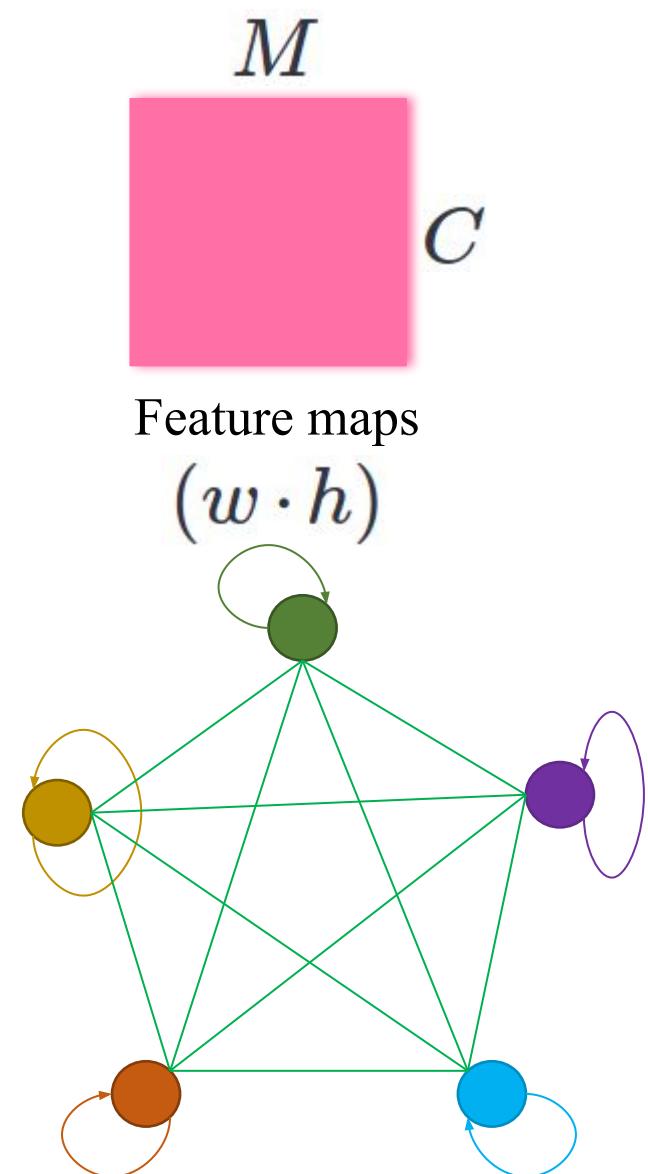
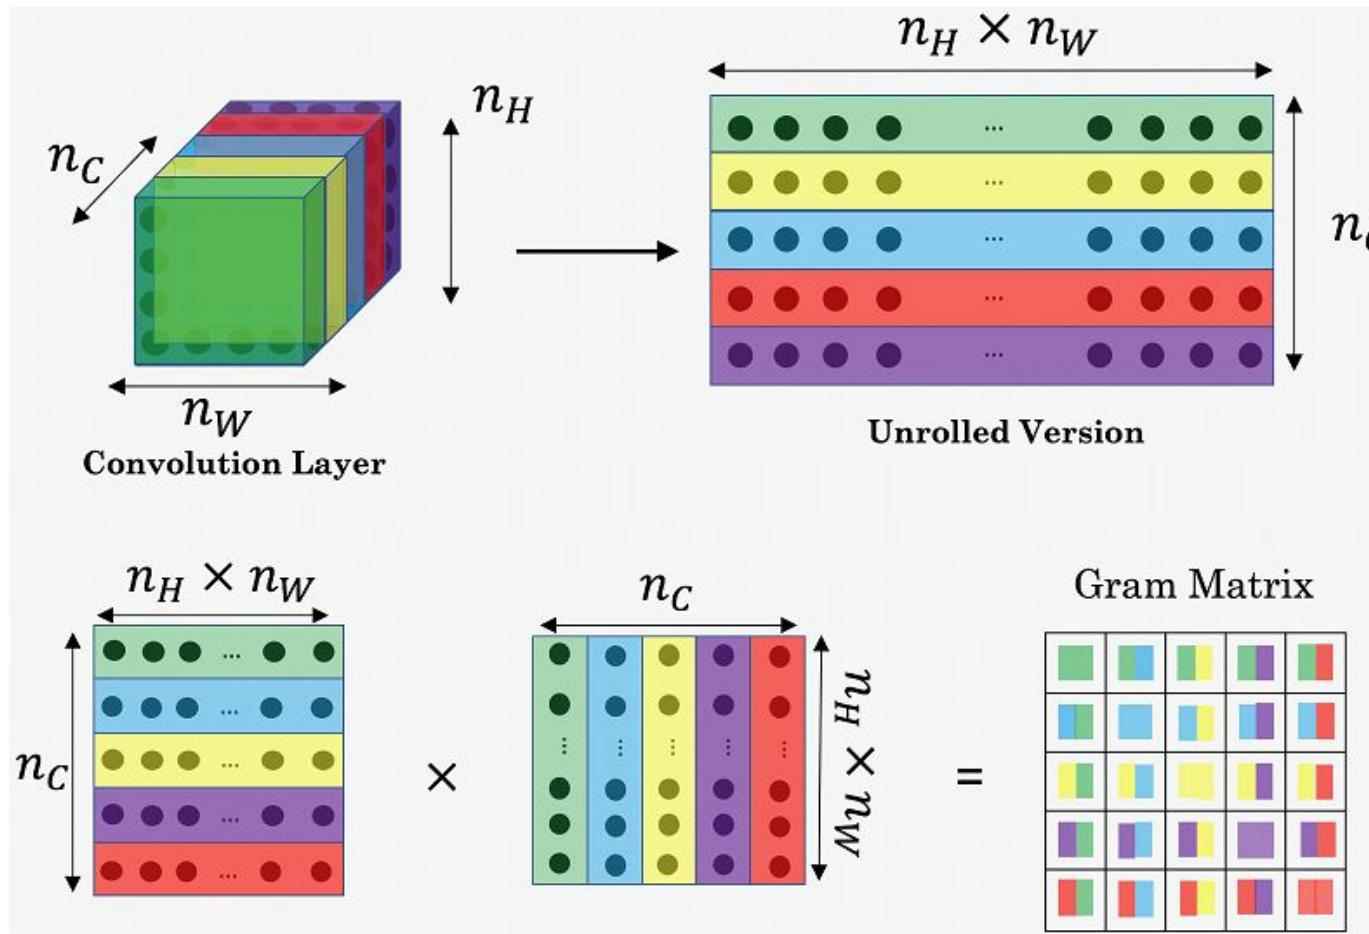
◆ Style Loss

- Layer l
- Number of filters C
- Feature map size $M (w \cdot h)$
- Gram Matrix G



Basic Knowledge

◆ Gram Matrix

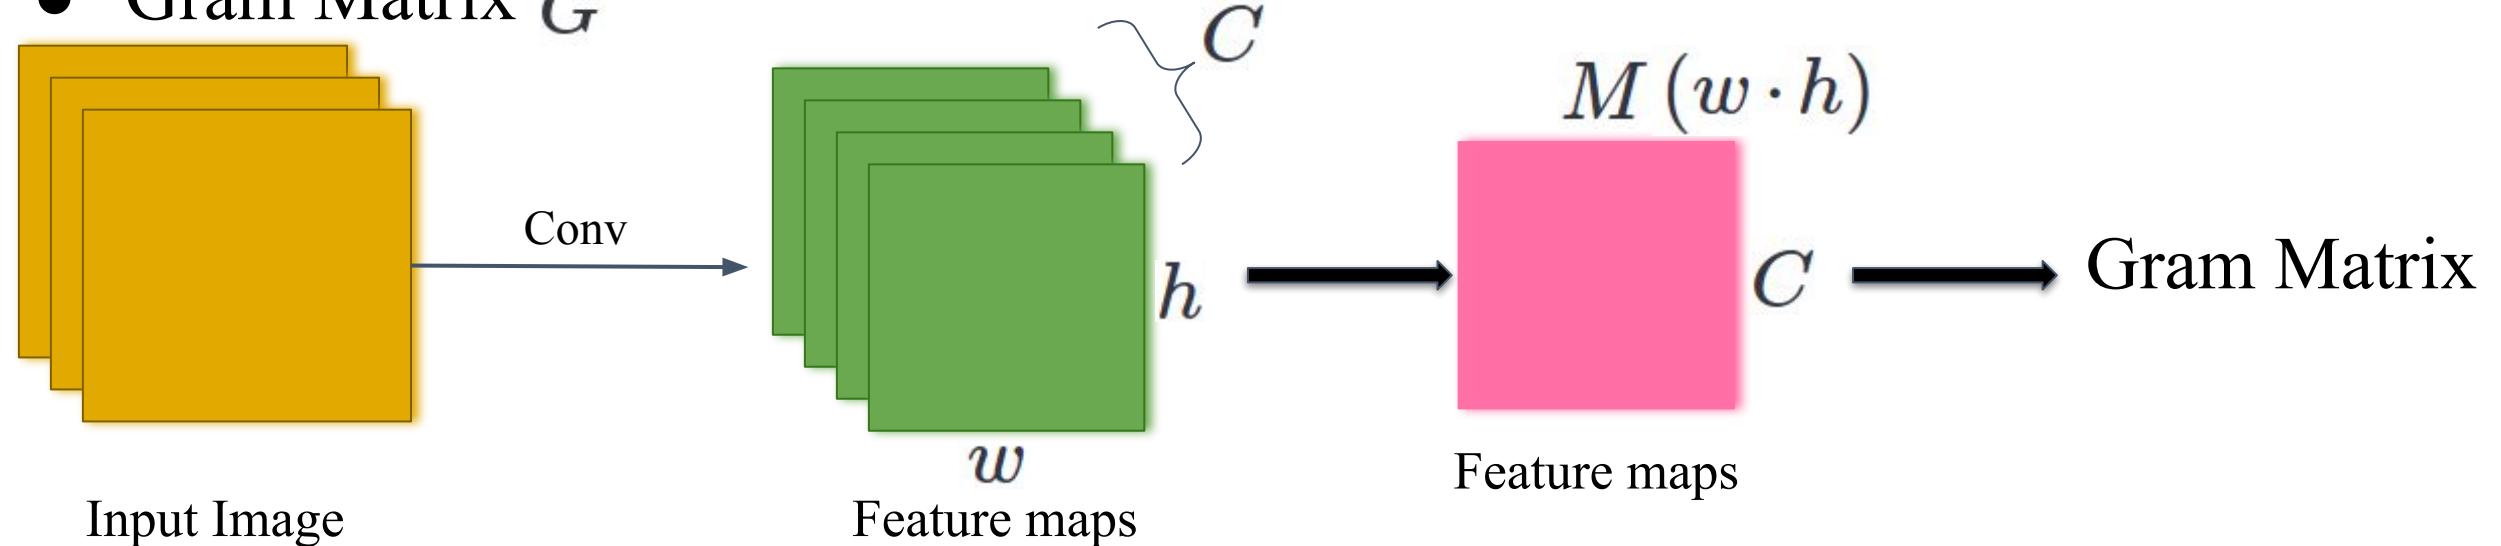


Basic Knowledge

◆ Style Loss

- Layer l
- Number of filters C
- Feature map size $M (w \cdot h)$
- Gram Matrix G

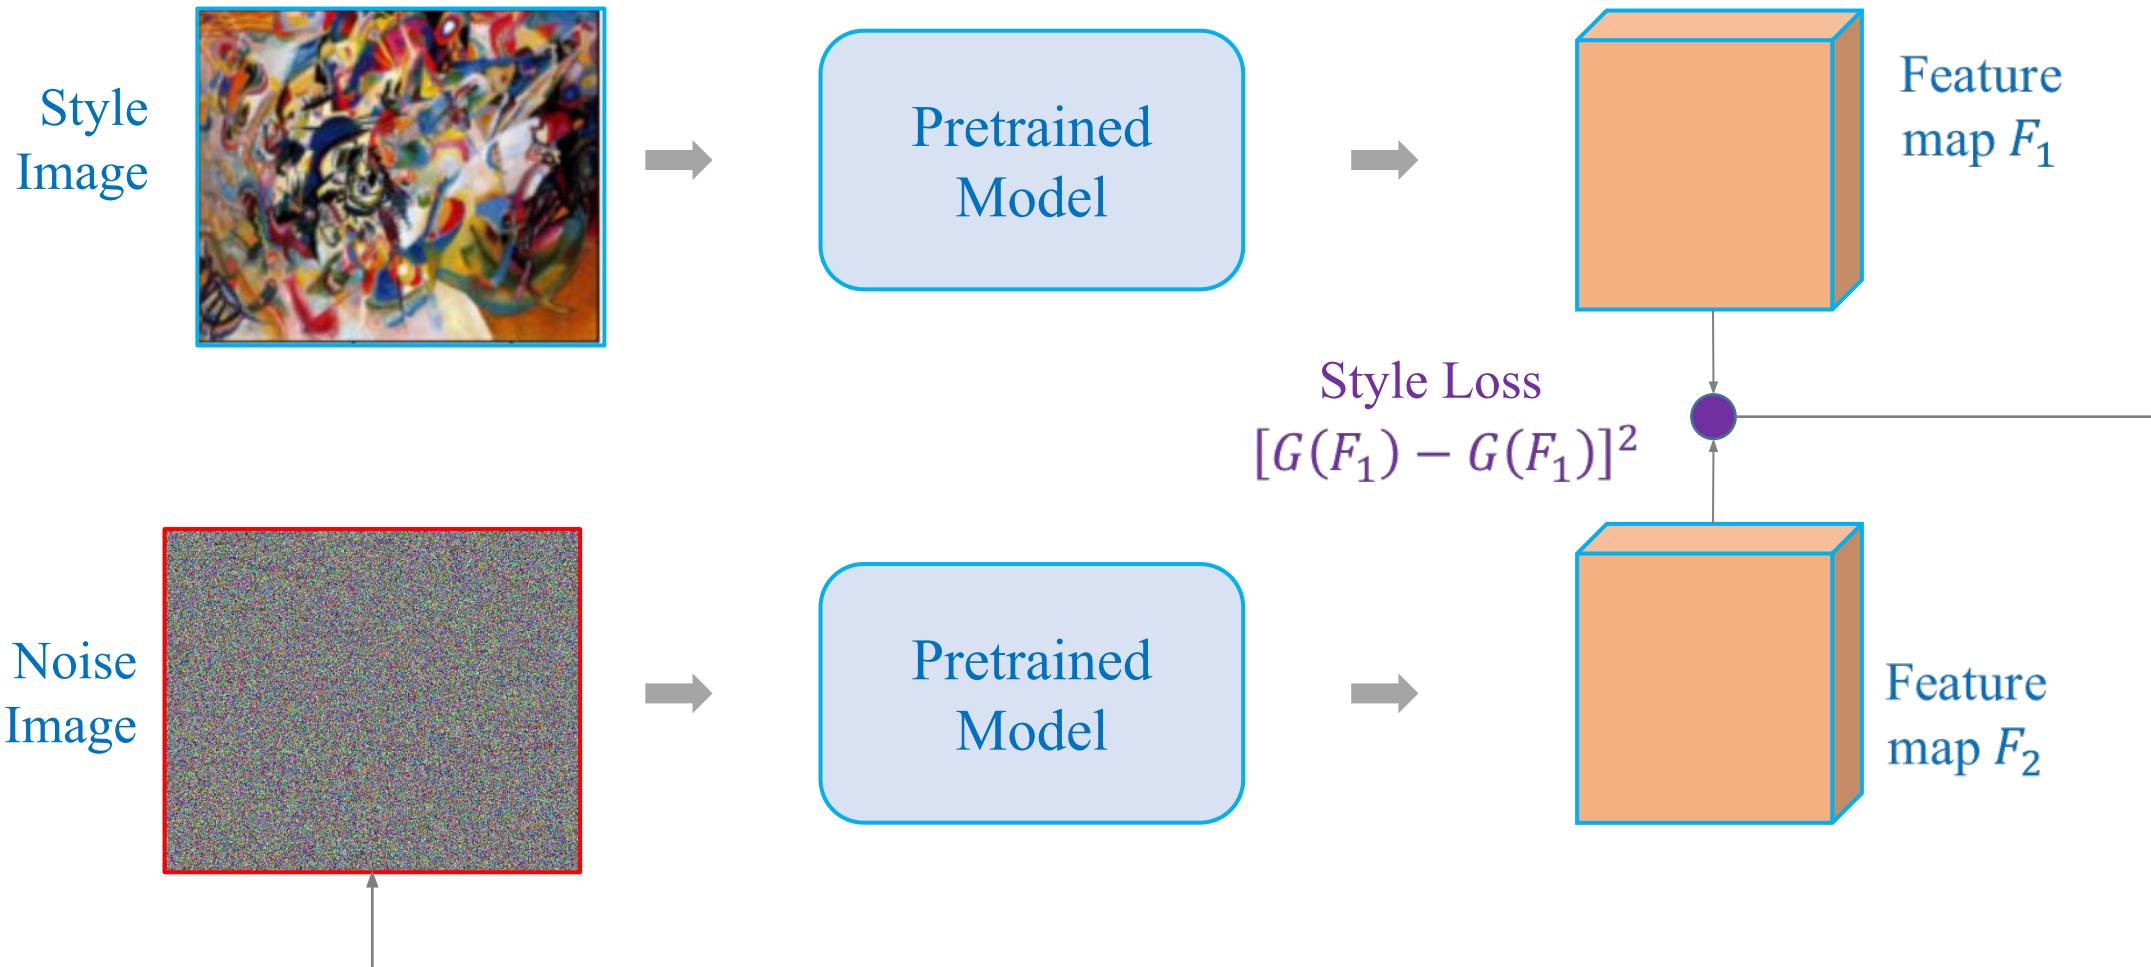
- Generated image x
- Generated feature map F
- Style image s
- Style feature map S



$$\mathcal{L}_{style}(s, x, l) = \frac{1}{2} \sum_{m,n} (G(F)_m^l - G(S)_m^l)^2$$

Basic Knowledge

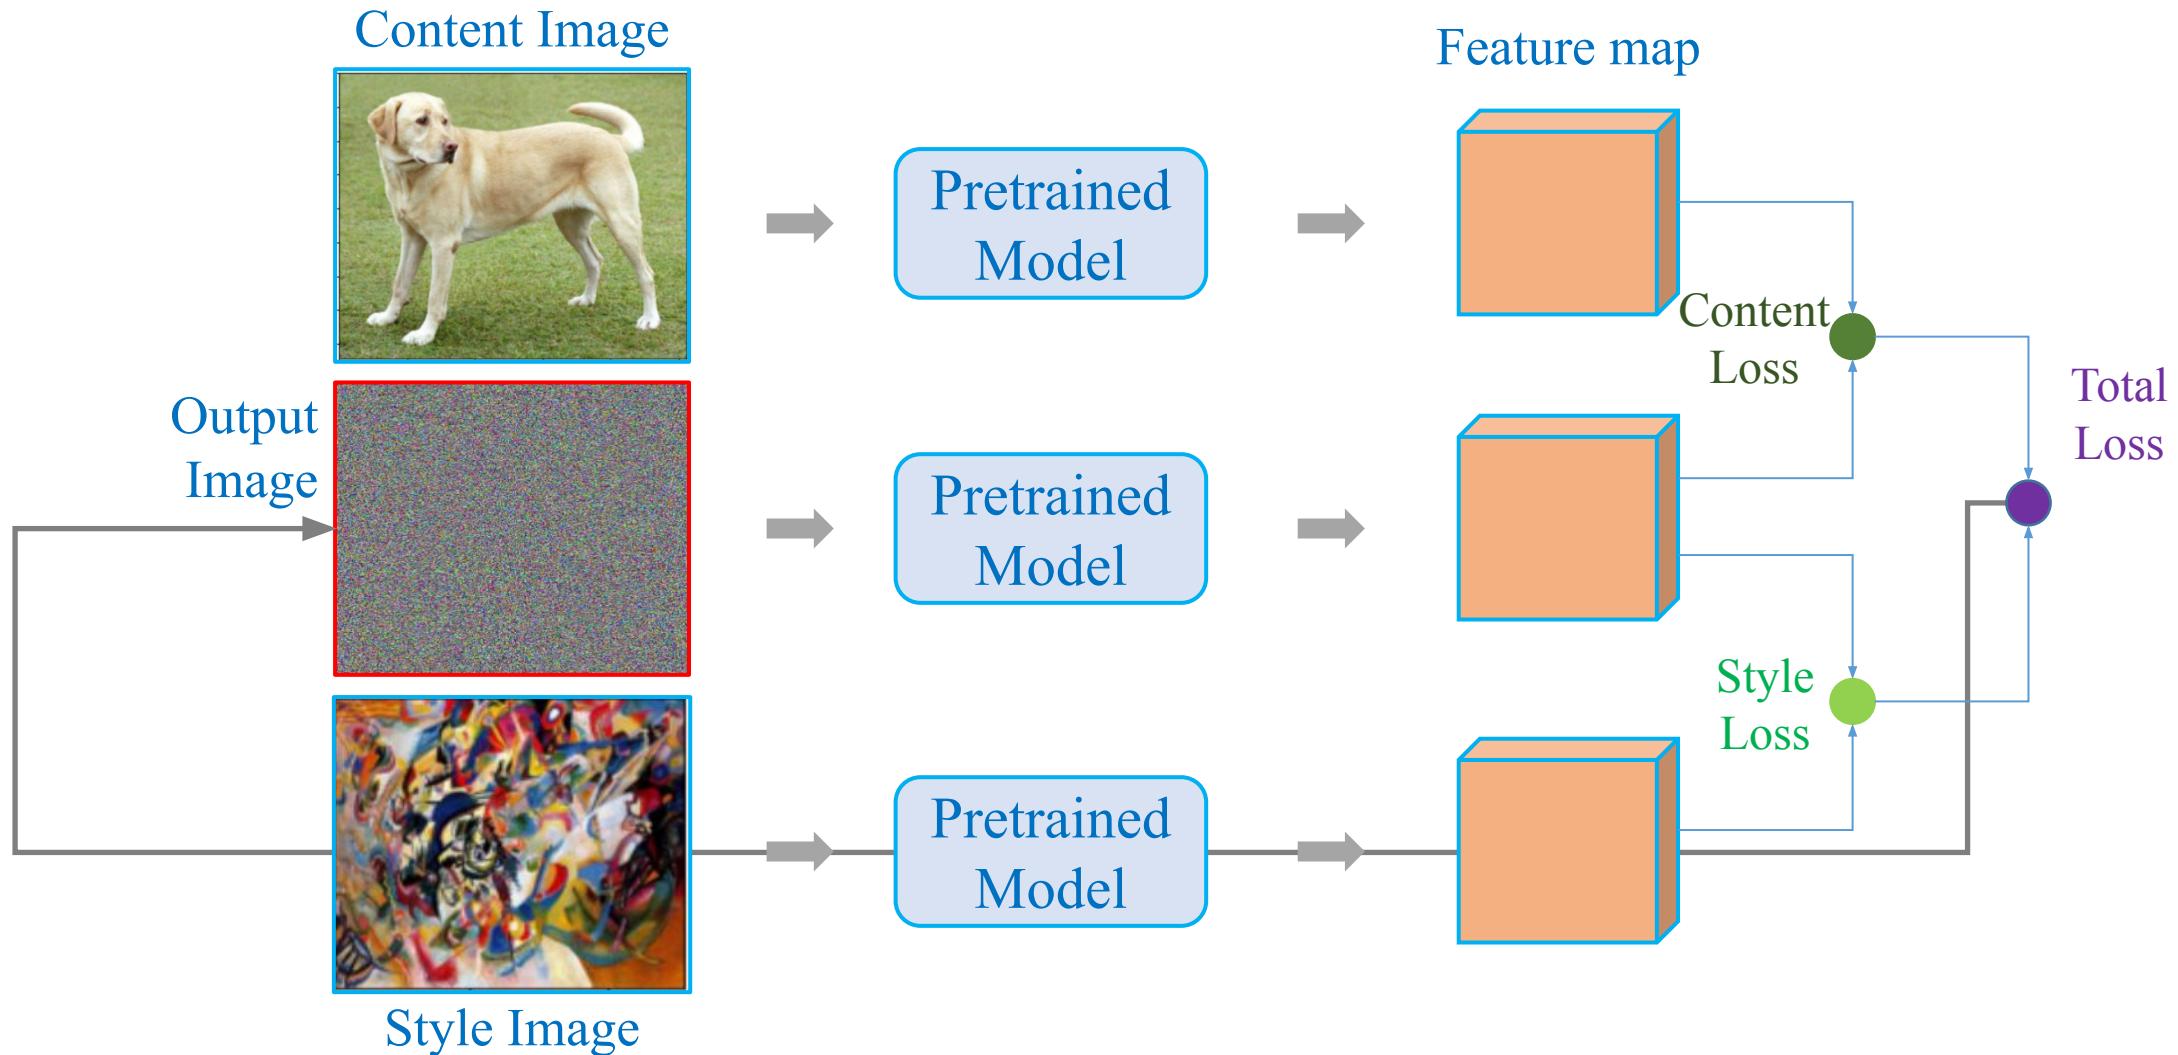
◆ Style Loss



Basic Knowledge

◆ Total Loss

$$\mathcal{L}_{total}(p, s, x) = \alpha \mathcal{L}_{content}(p, x) + \beta \mathcal{L}_{style}(s, x)$$



Style Transfer Example

Style Transfer Example

❖ Style Transfer

Algorithm 1 Neural Style Transfer

Start with a copy of the content reference image, call this *input* image

Run content and style reference images through VGG-19 model

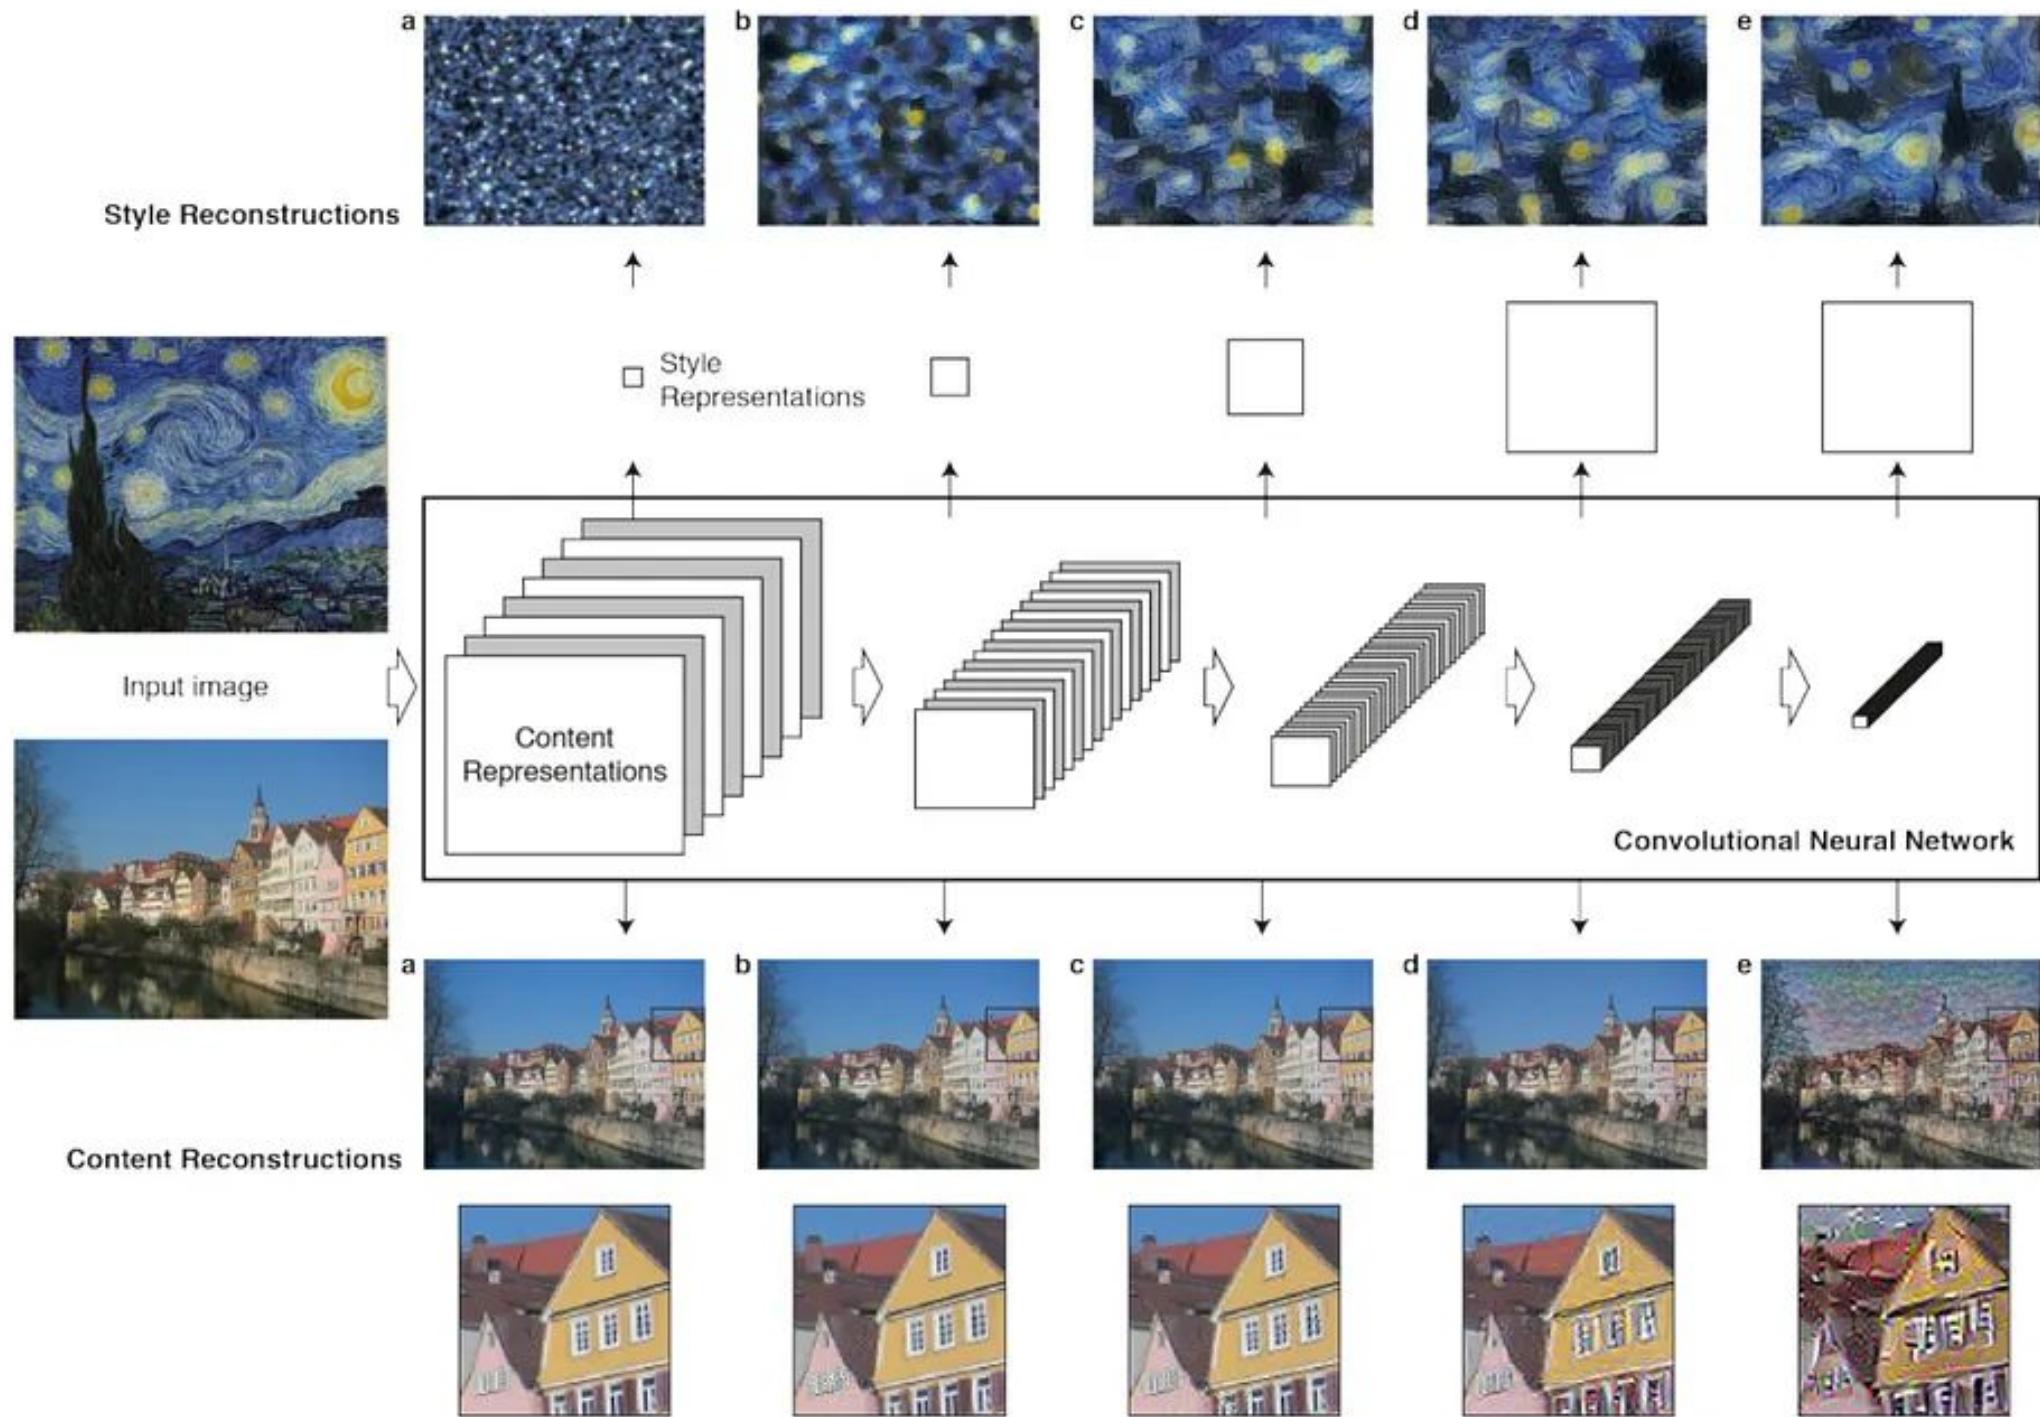
for num_epochs **do**

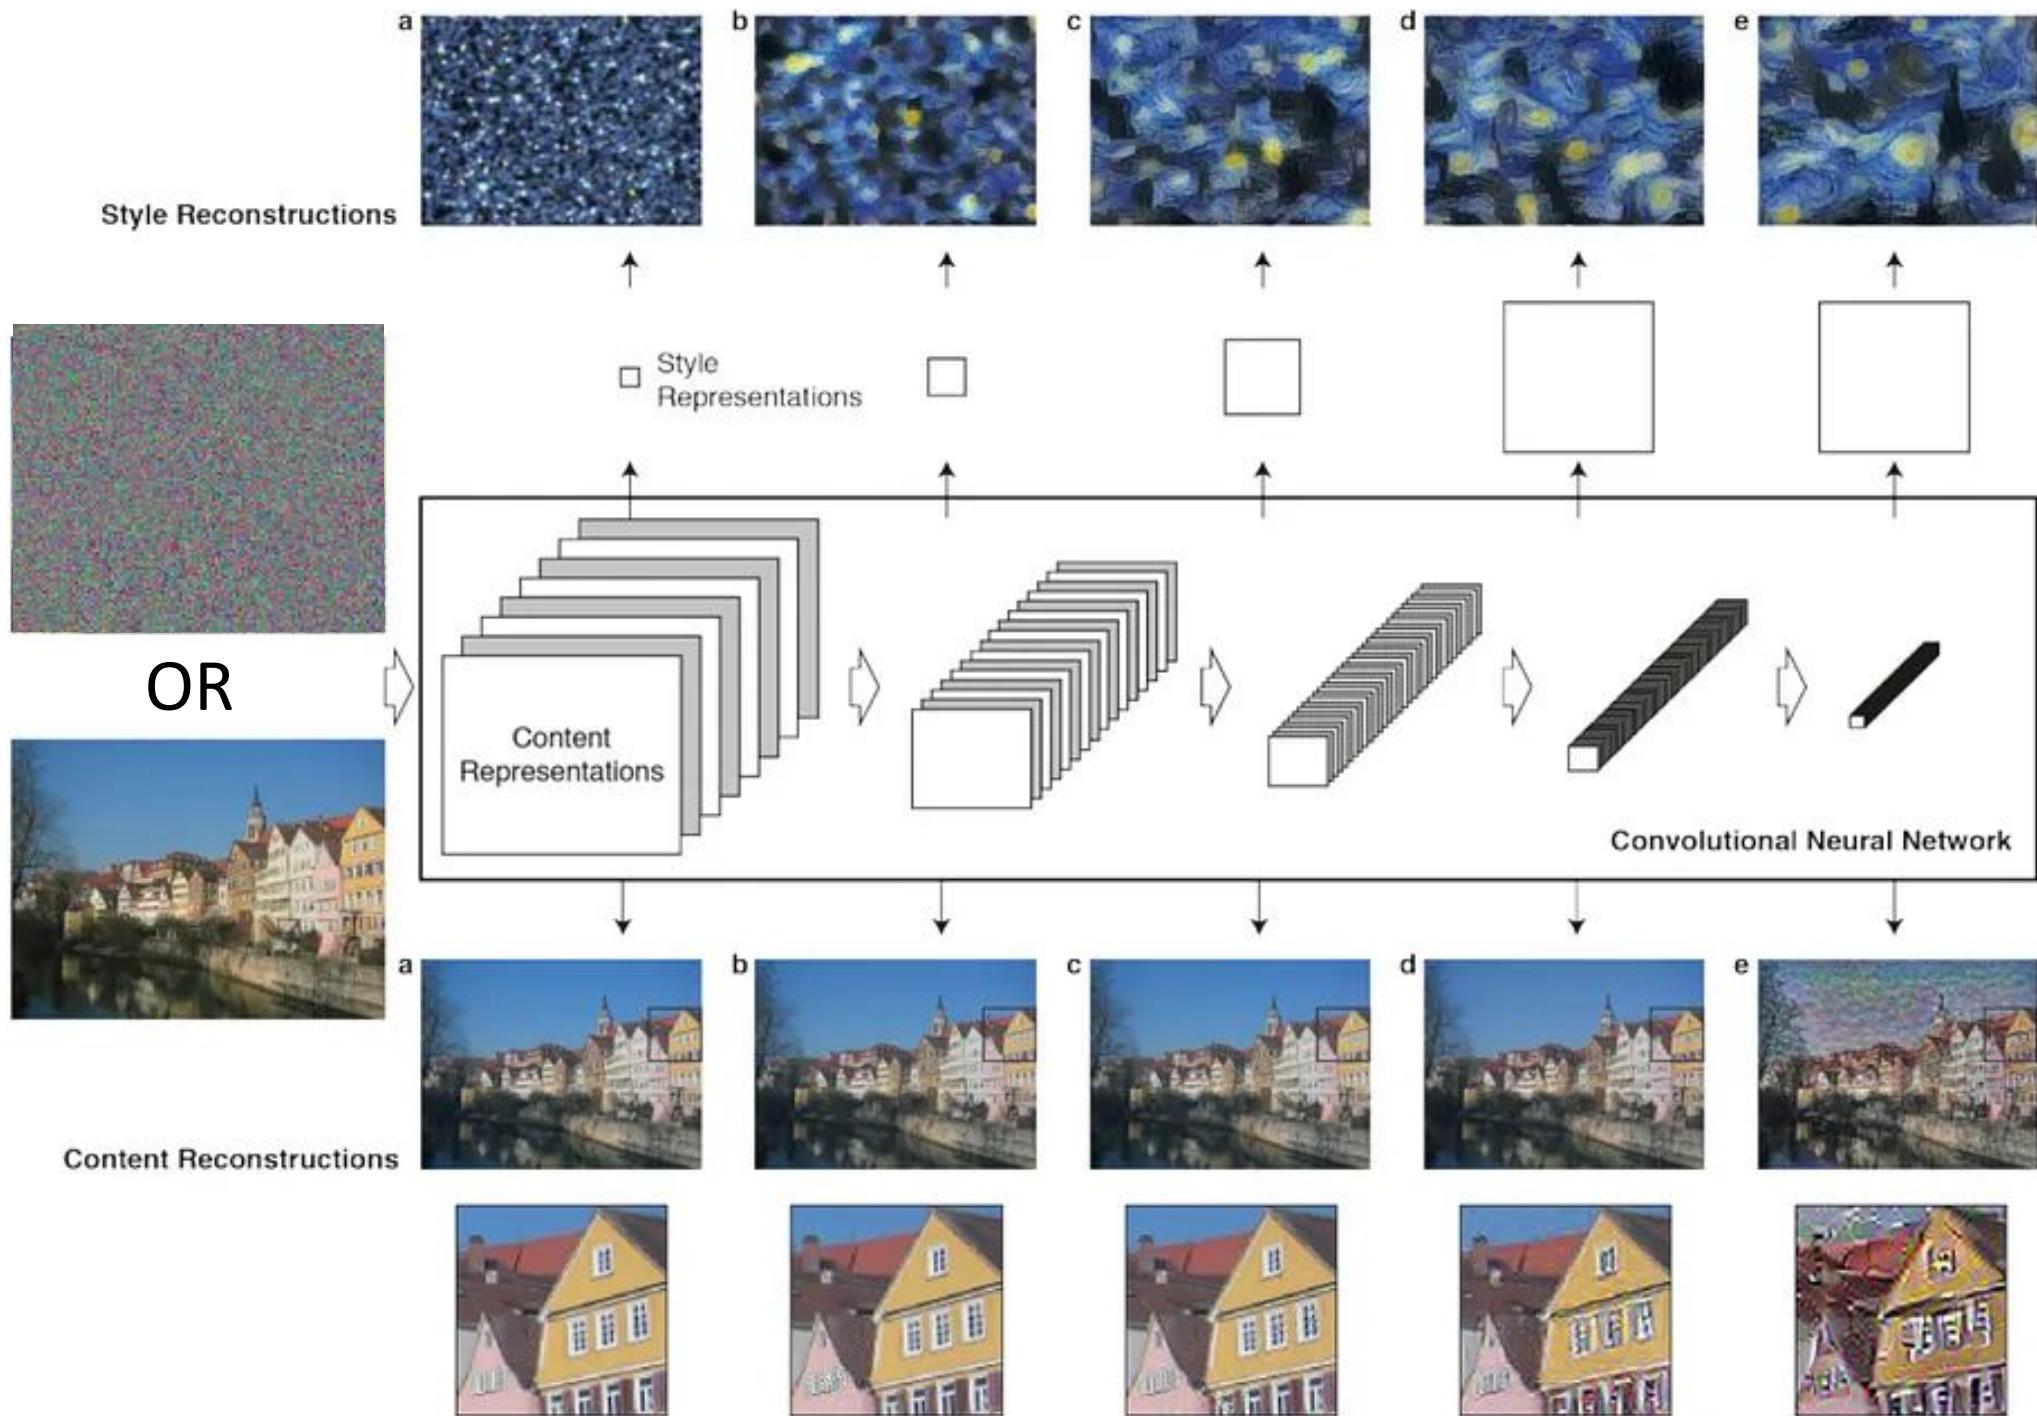
 Run input image through VGG-19

 Calculate L using intermediate feature maps

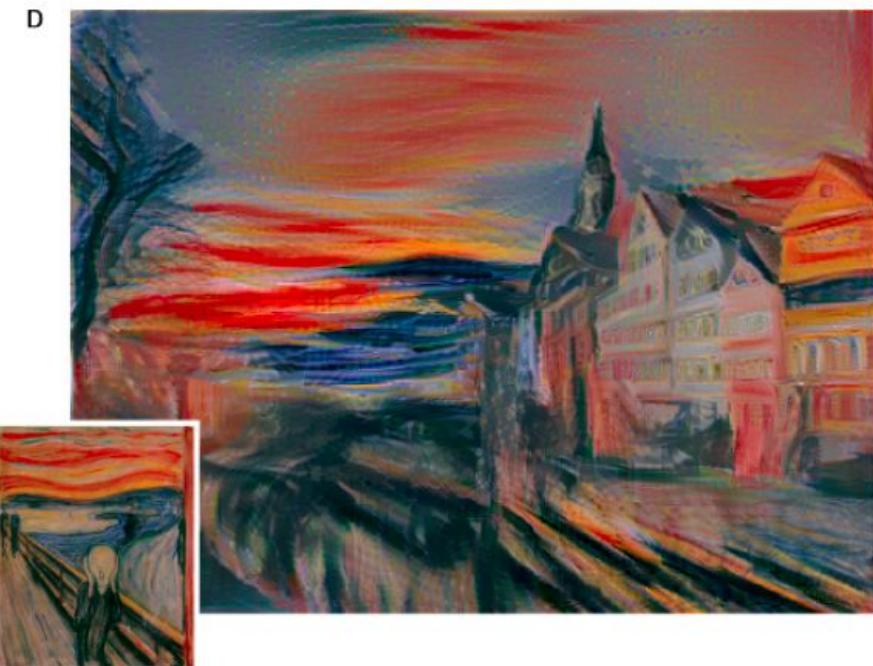
 Backpropagate error through VGG-19 and adjust pixel values of input image

end for





◆ Style

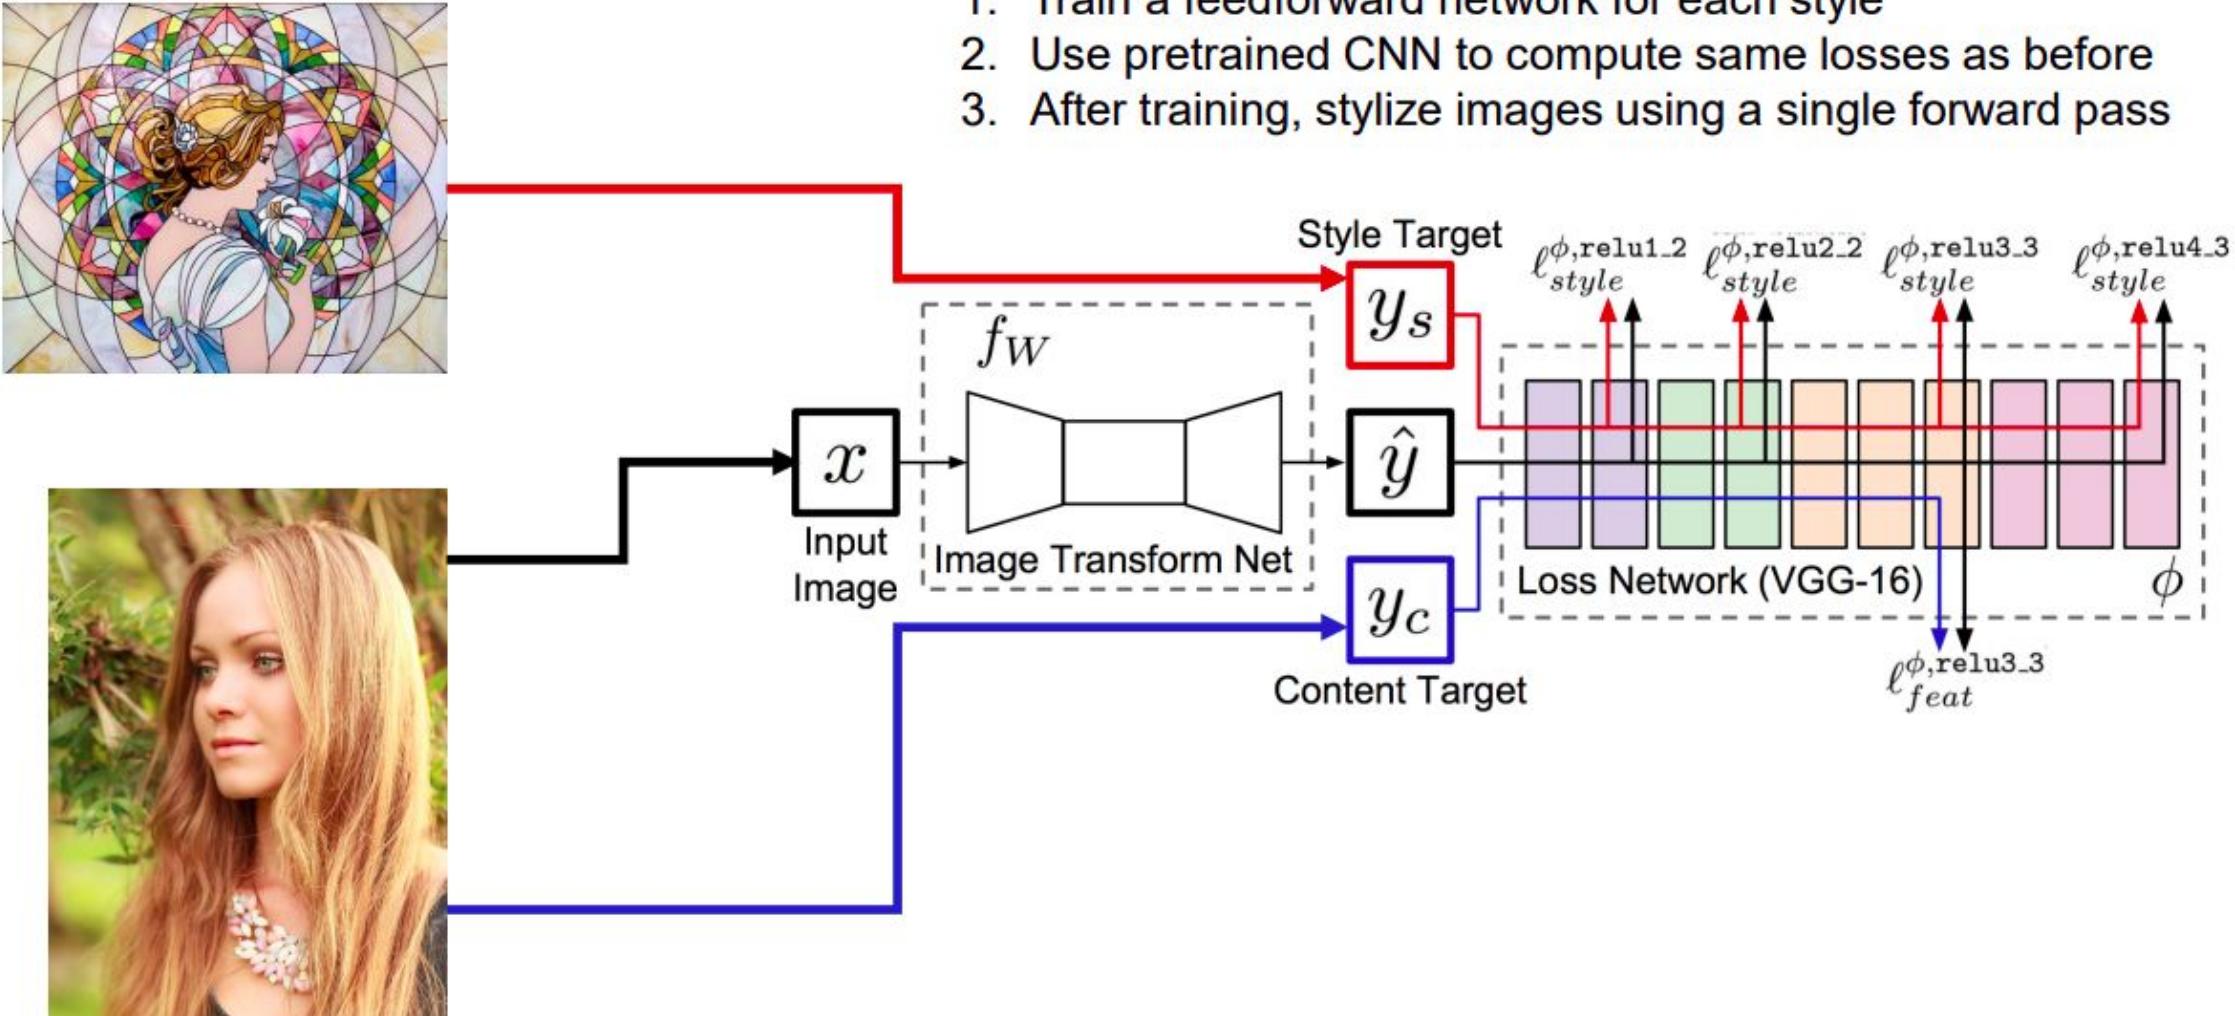


AI VIETNAM
All-in-One Course

Milestone Papers

Milestone Papers

❖ Fast Neural Style Transfer

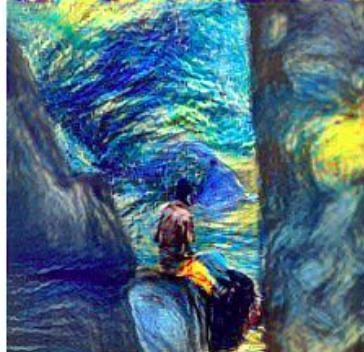
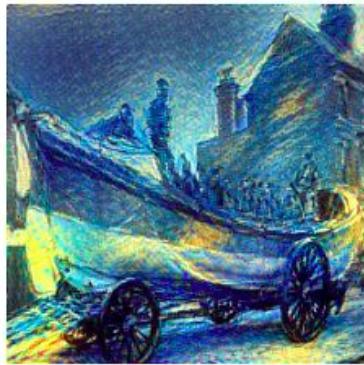


Milestone Papers

❖ Fast Neural Style Transfer

Style

The Starry Night,
Vincent van Gogh,
1889



Style

The Muse,
Pablo Picasso,
1935



Milestone Papers

❖ Fast Neural Style Transfer

Style
Sketch



Style
The Simpsons



Content

[10]

Ours

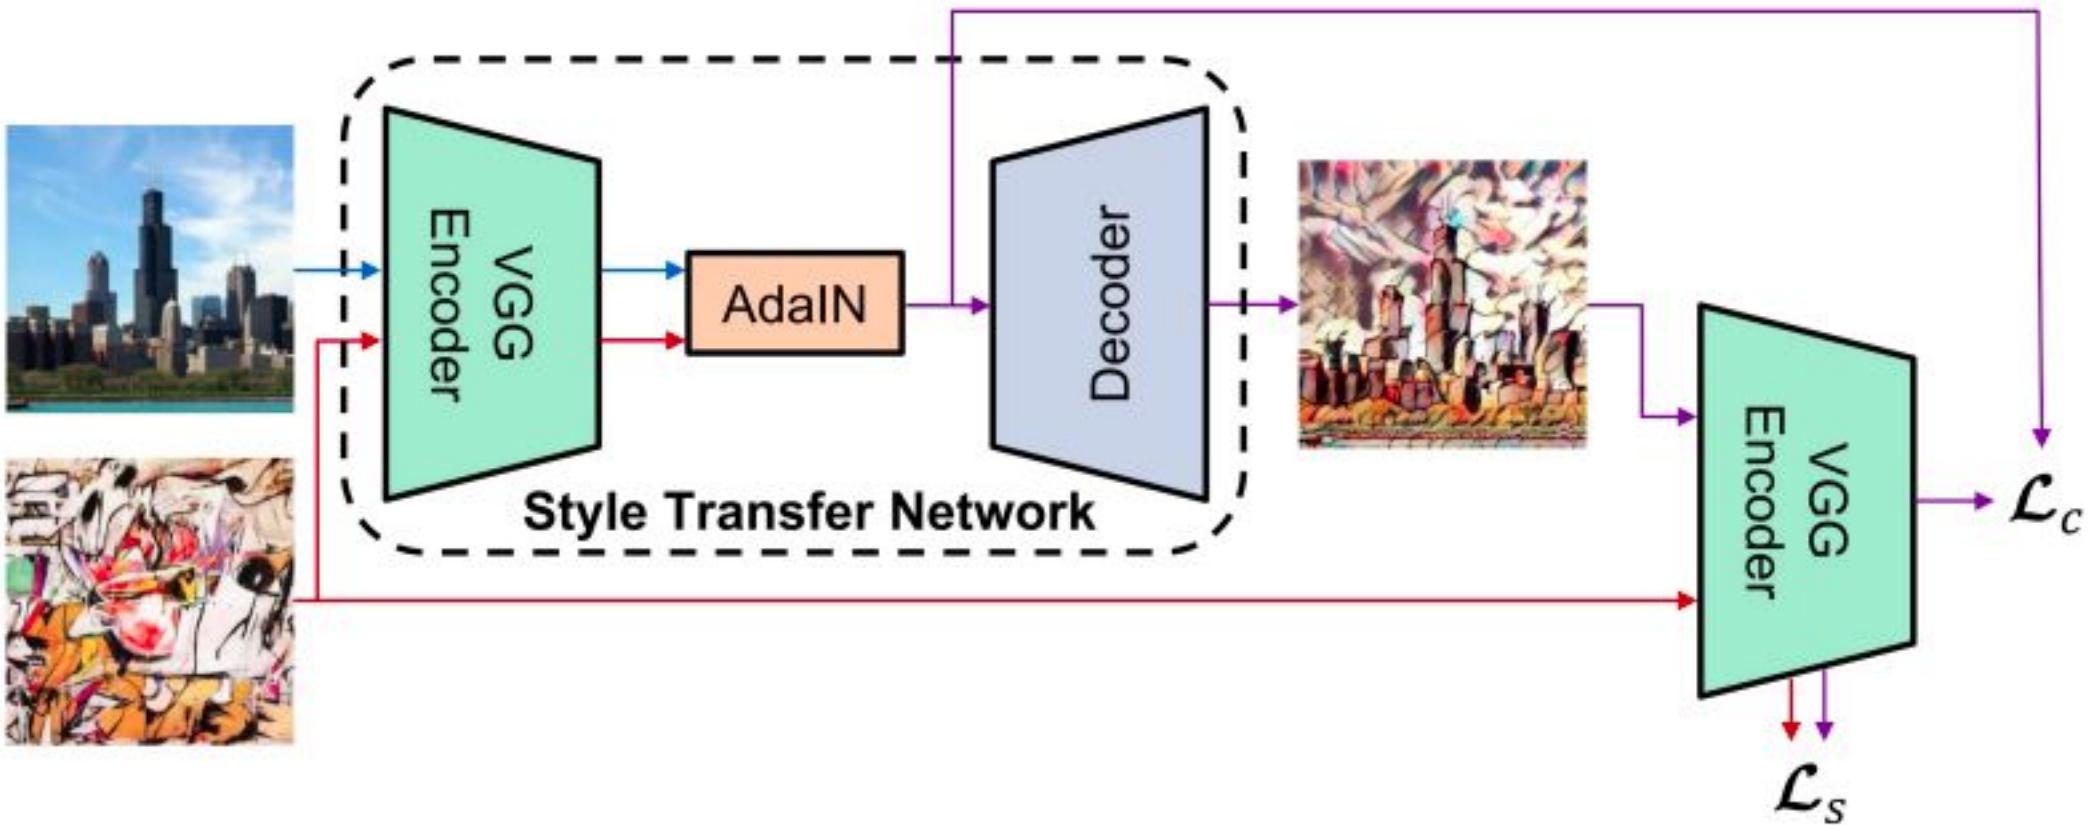
Content

[10]

Ours

Milestone Papers

❖ Arbitrary Style Transfer



Milestone Papers

❖ Arbitrary Style Transfer



(a) $\alpha = 0$



(b) $\alpha = 0.25$



(c) $\alpha = 0.5$



(d) $\alpha = 0.75$



(e) $\alpha = 1$



(f) Style Image

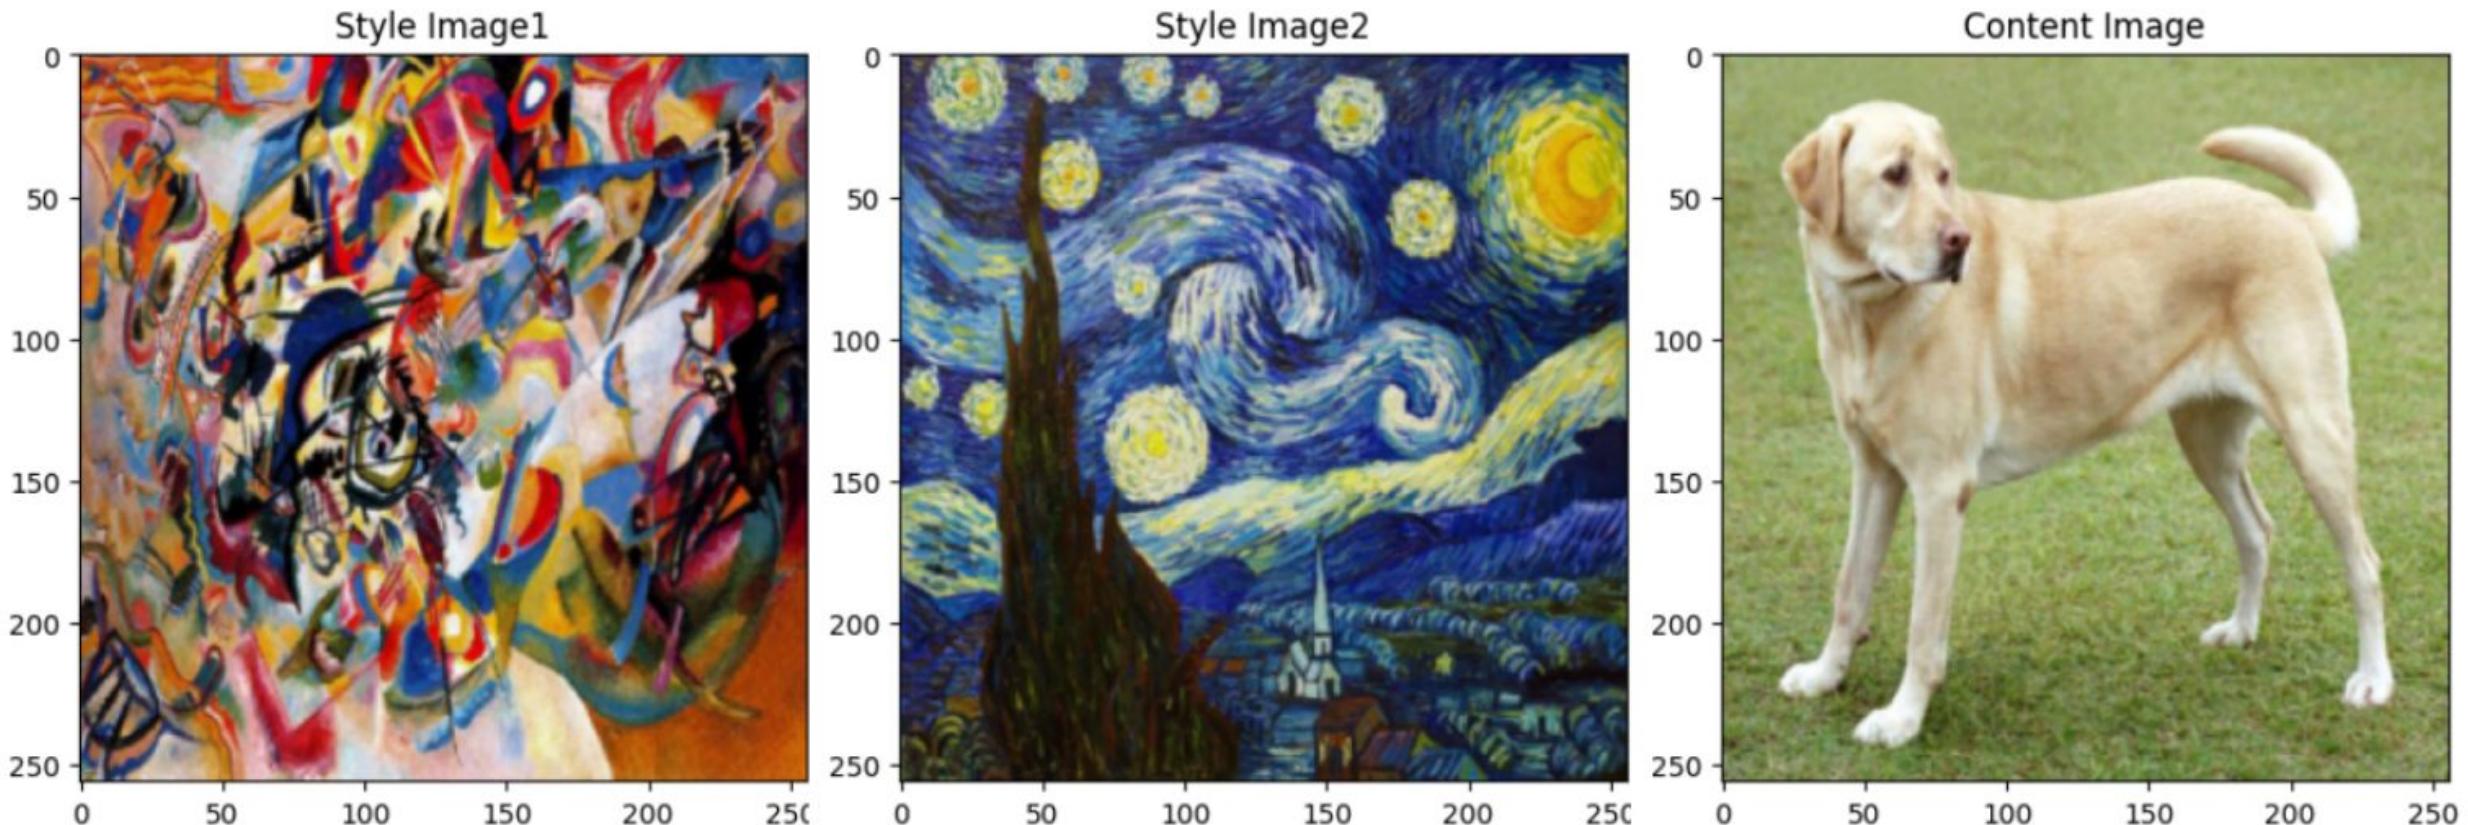
QUIZ



AI VIETNAM
All-in-One Course

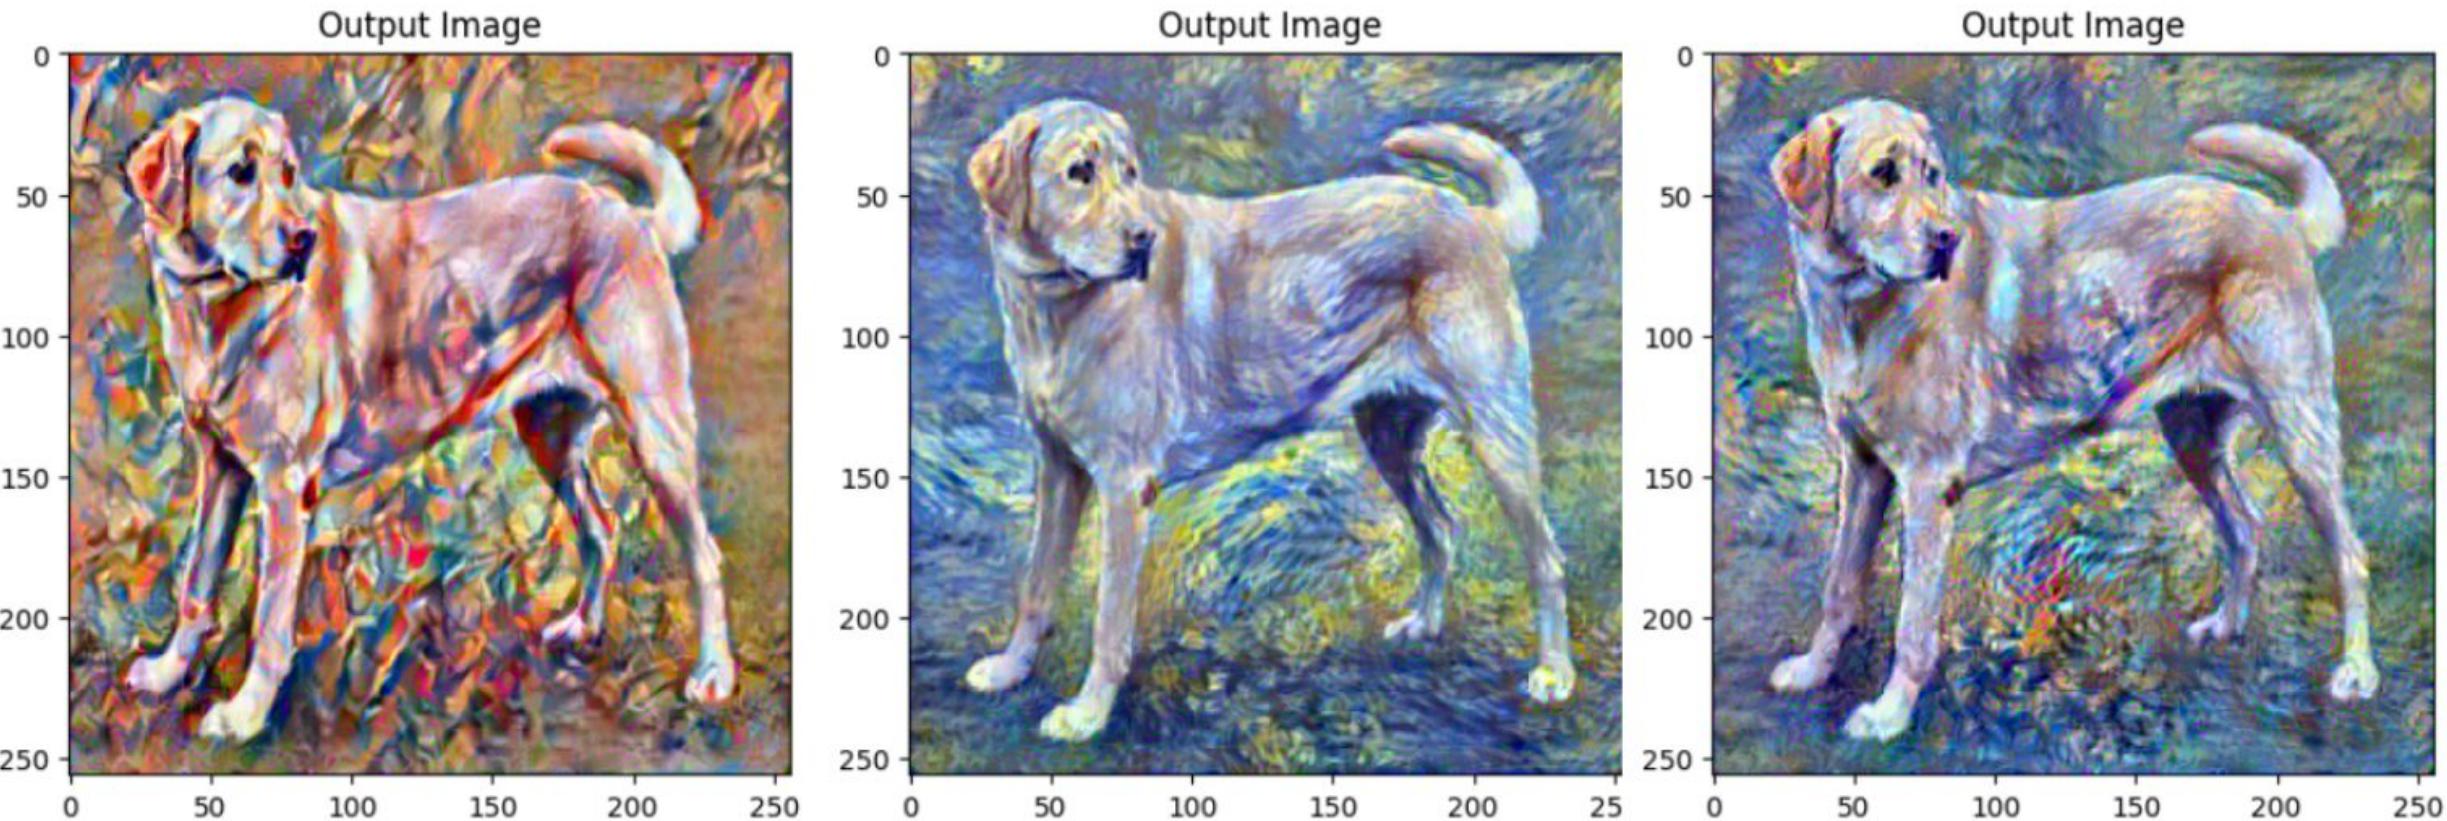
Style Transfer with 2 Style Images

Style Transfer with 2 Style Images



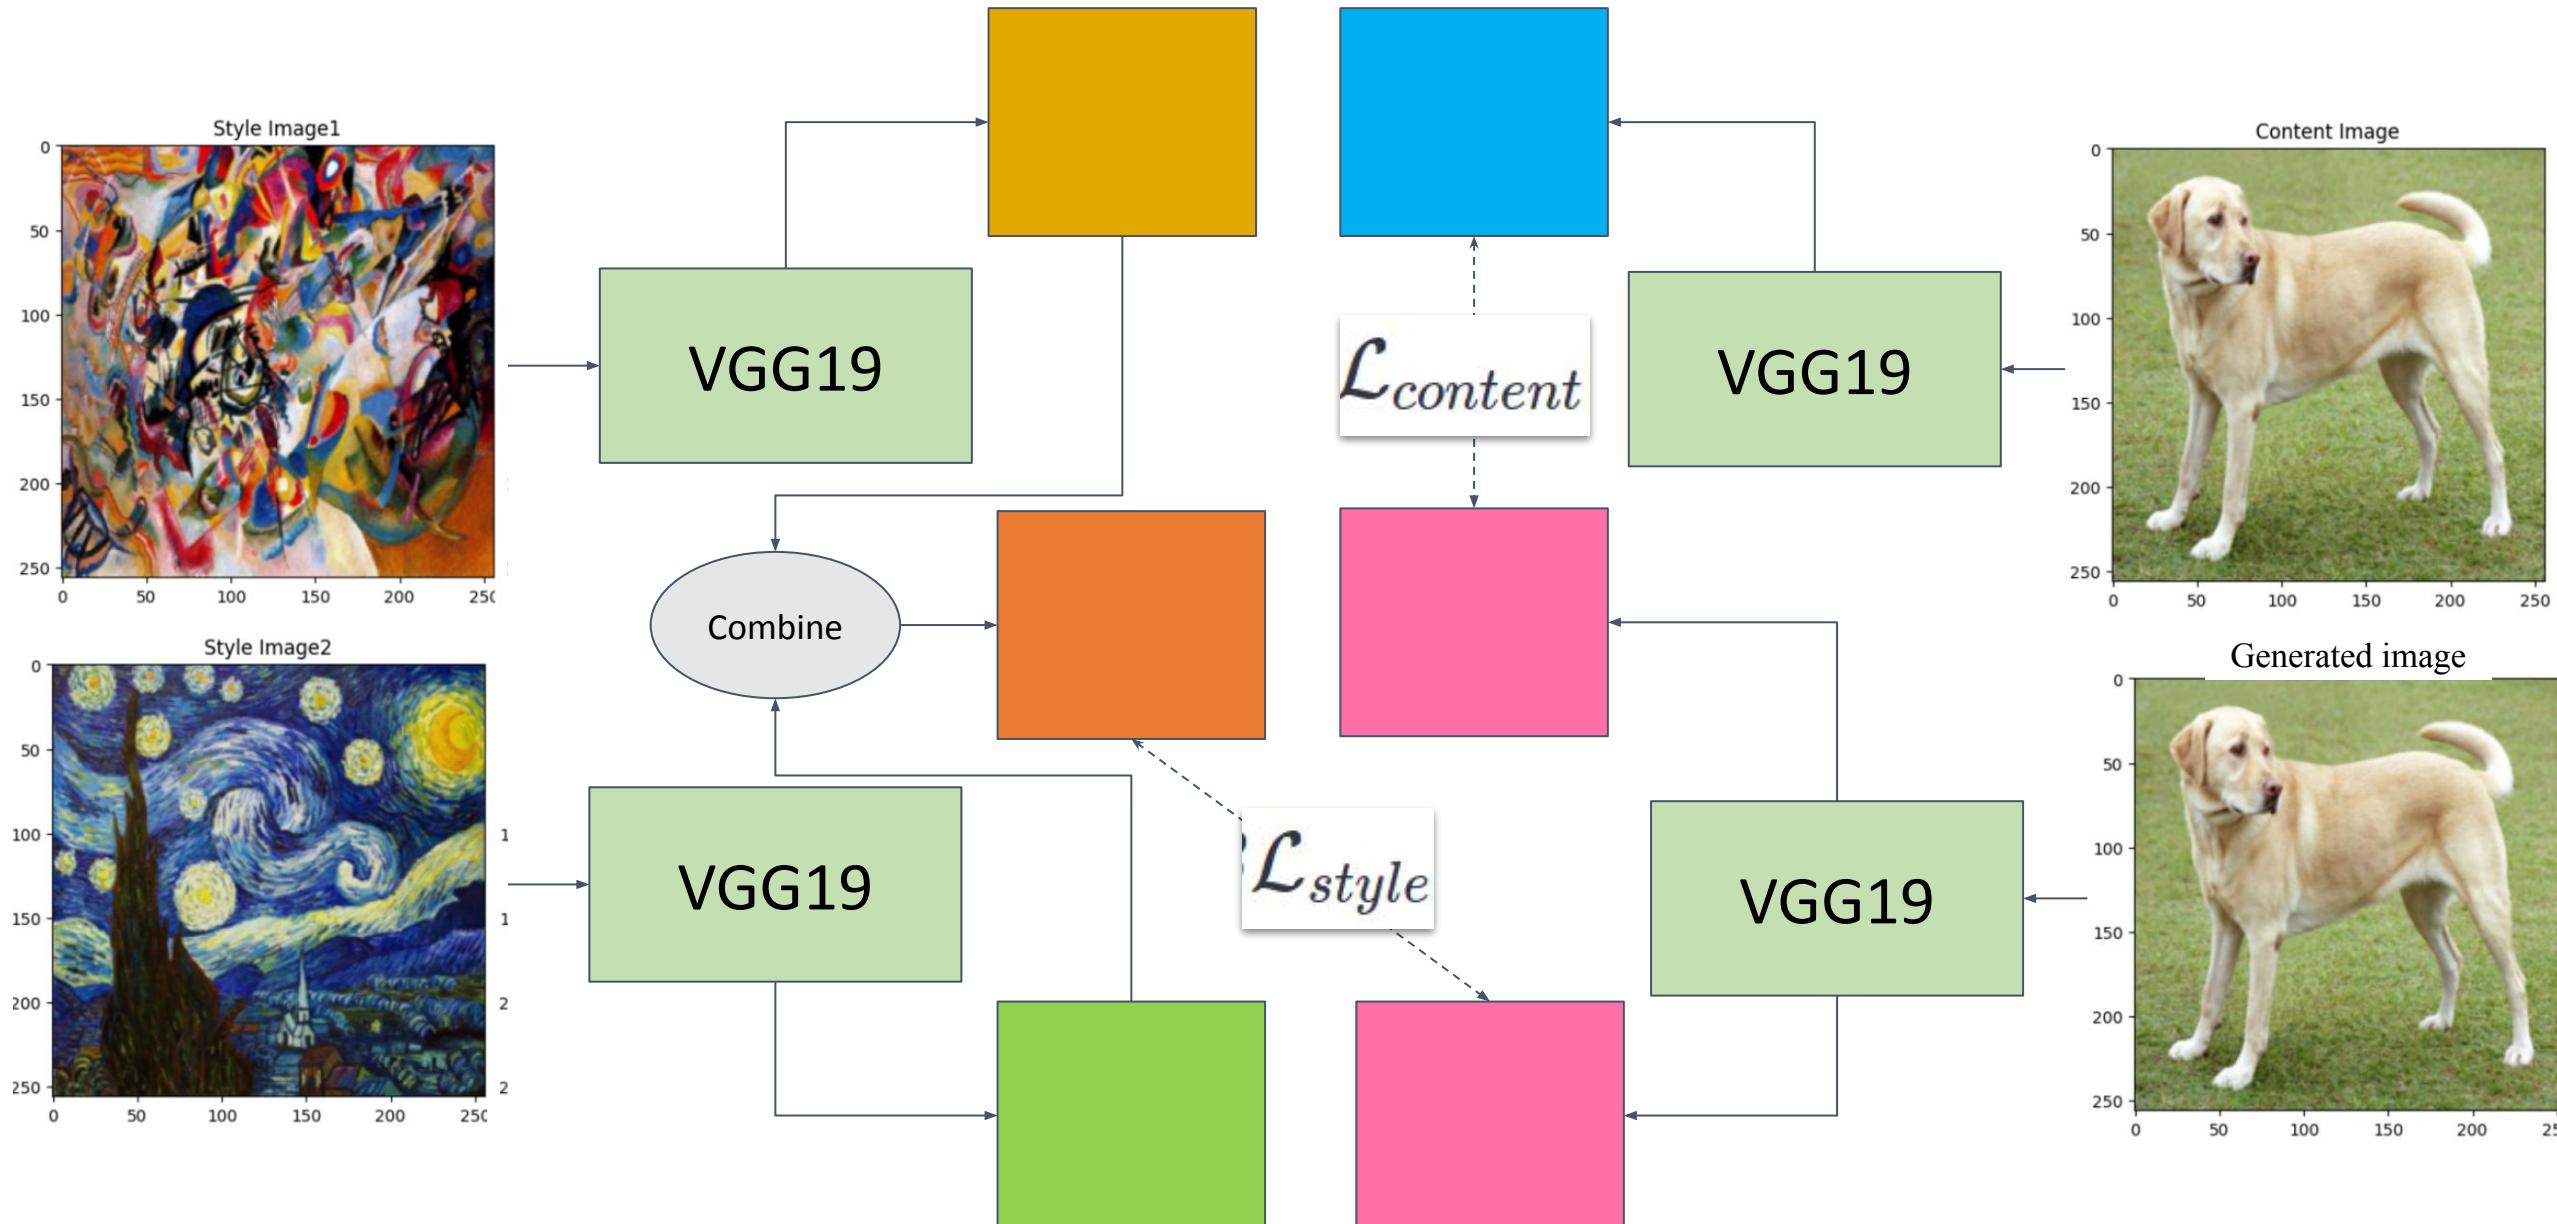
(a) Ảnh đầu vào của bài tập 1, ảnh trái là ảnh style 1, ảnh giữa là ảnh style 2, và ảnh phải là ảnh content

Style Transfer with 2 Style Images

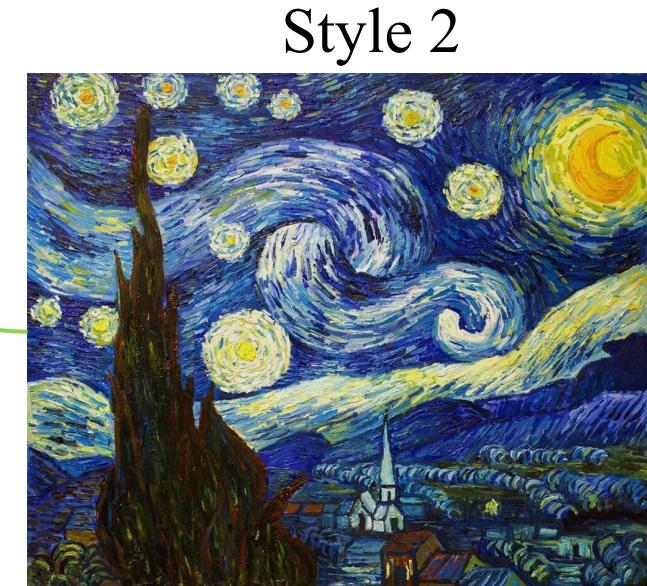
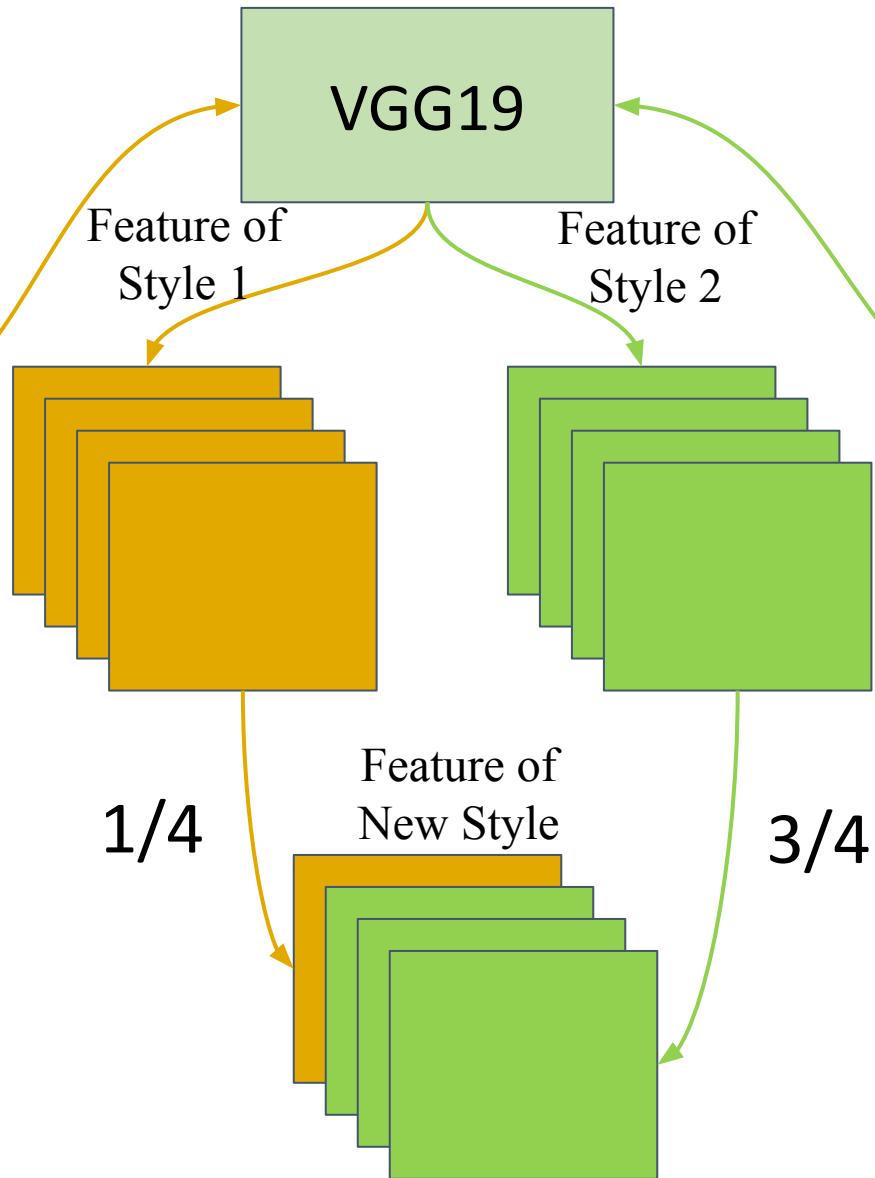


(b) Ảnh trái là ảnh khi transfer theo style 1, ảnh giữa là ảnh khi transfer theo style 2, và ảnh phải là ảnh output của bài tập 1 khi transfer theo $1/4$ style1 + $3/4$ style 2

Style Transfer with 2 Style Images



Style Transfer with 2 Style Images



Style Transfer with 2 Style Images

◆ Code Overview

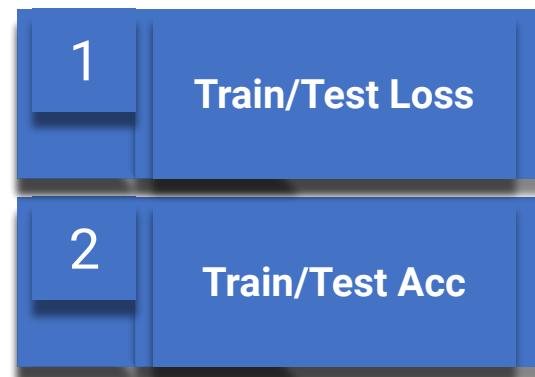
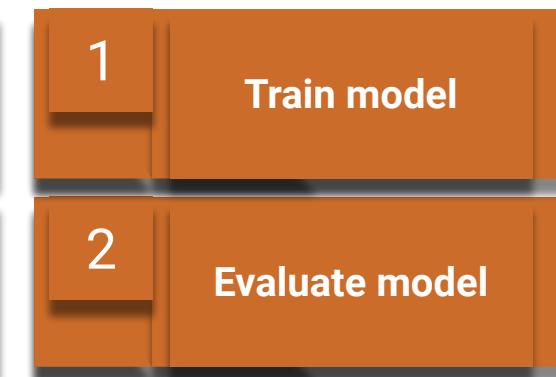
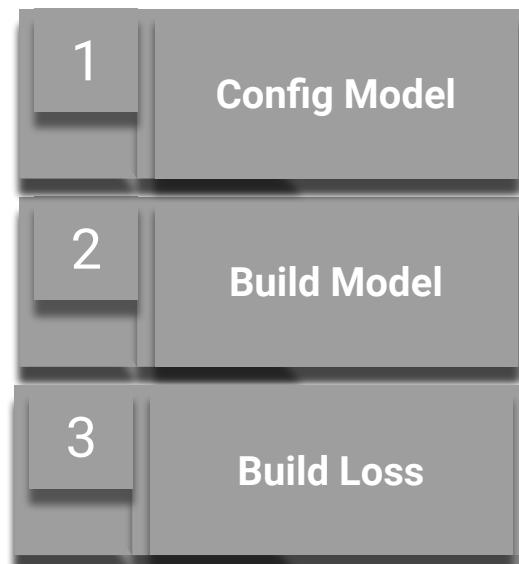
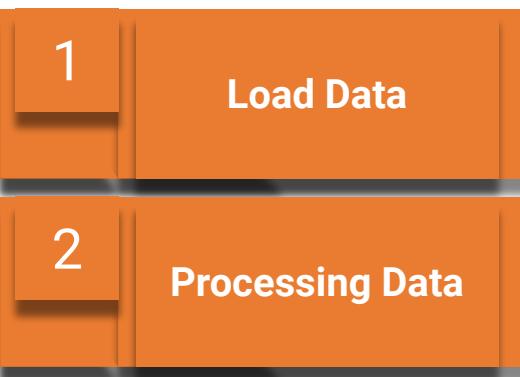
0: Import Library

1: Data Preparation

2: Modeling

3: Train and Evaluate

4: Plot Results



Style Transfer with 2 Style Images

❖ Data

```
1 from PIL import Image
2 import torchvision.transforms as transforms
3
4 imsize = 256
5
6 img_transforms = transforms.Compose([
7     transforms.Resize((imsize, imsize)),
8     transforms.ToTensor(),
9 ])
10
```

```
1 def image_loader(image_name):
2     image = Image.open(image_name)
3     image = img_transforms(image).unsqueeze(0)
4     return image.to(device, torch.float)
5
6 style_img1 = image_loader("style_img.jpg")
7 style_img2 = image_loader("style_img2.jpg")
8 content_img = image_loader("content_img.jpg")
```

Đầu tiên chúng ta cần load ảnh content và style bằng thư viện **PIL**, rồi resize về cùng kích thước **256x256** và convert thành tensor để sử dụng với Pytorch. Đối với bài 1 chúng ta cần load 3 ảnh (2 style và 1 content), còn đối với bài 2 chỉ cần 1 style và 1 content

Style Transfer with 2 Style Images

❖ Loss

▼ 3.1 Content Loss

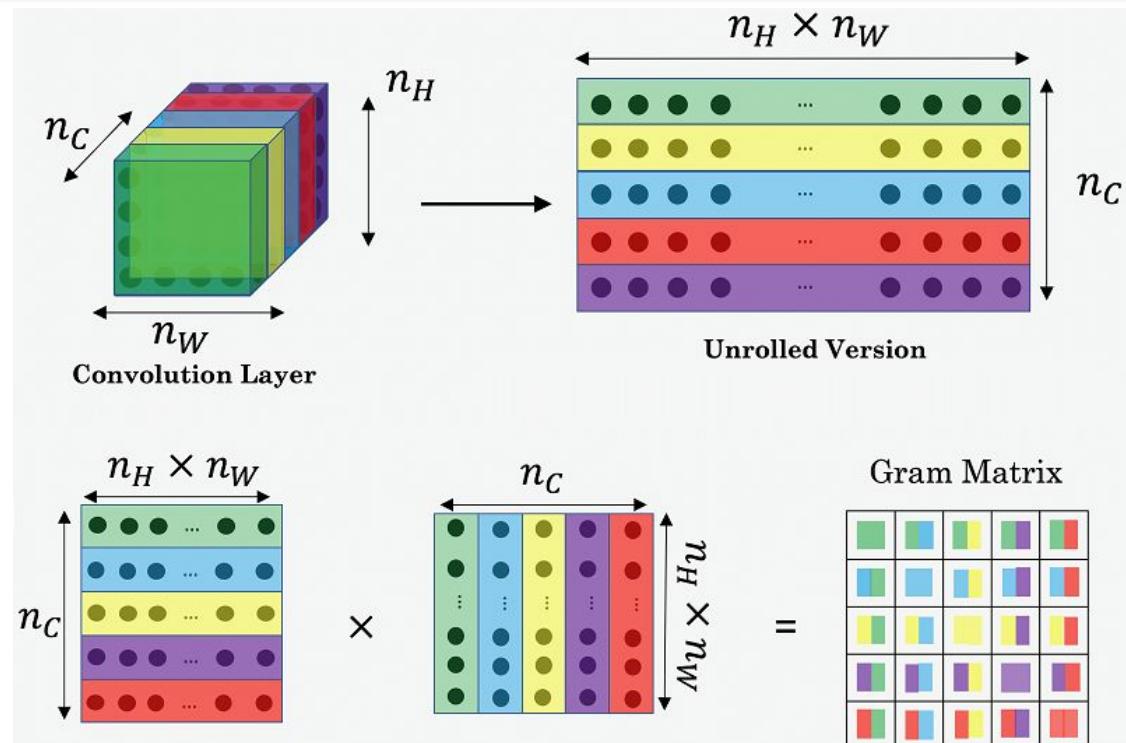
```
[ ] 1 content_weight = 1
2 ContentLoss = nn.MSELoss()
```

▼ 3.2 Style Loss

```
[ ] 1 def gram_matrix(tensor):
2     a, b, c, d = tensor.size()
3     tensor = tensor.view(a * b, c * d)
4     G = torch.mm(tensor, tensor.t())
5     return G.div(a * b * c * d)
6
7 style_weight = 1e6
8 StyleLoss = nn.MSELoss()
```

$$\mathcal{L}_{content}(p, x, l) = \frac{1}{2} \sum_{i,j,k} (F_{ijk}^l - P_{ijk}^l)^2$$

$$\mathcal{L}_{style}(s, x, l) = \frac{1}{2} \sum_{m,n} (G(F)_mn^l - G(S)_mn^l)^2$$



Softmax

FC 1000 se

FC 4096

FC 4096

Pool

3x3 conv,512

3x3 conv,512

3x3 conv,512

3x3 conv,512

Pool

3x3 conv,512

3x3 conv,512

3x3 conv,512

Pool

3x3 conv,512

3x3 conv,512

3x3 conv,512

3x3 conv,512

Pool

3x3 conv,128

3x3 conv,128

Pool

3x3 conv,64

3x3 conv,64

Input

Style Transfer with 2 Style Images

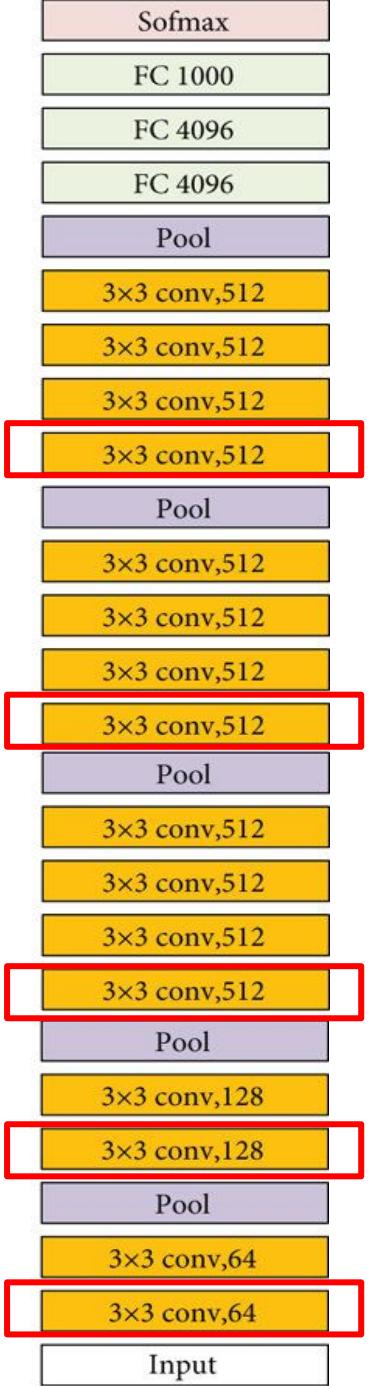
❖ Model

```

1 from torchvision.models import vgg19, VGG19_Weights
2
3 VGG19_pretrained = vgg19(weights=VGG19_Weights.DEFAULT).features.eval()
4 VGG19_pretrained.to(device)

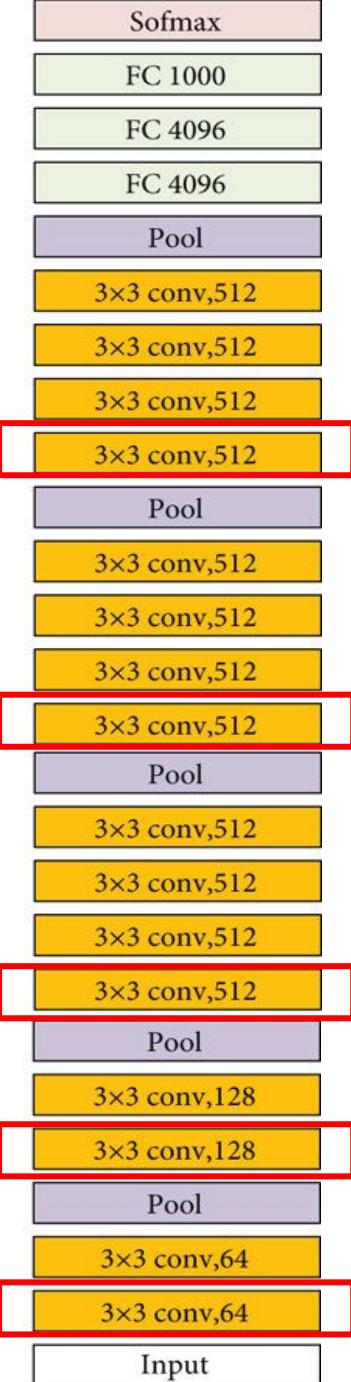
(0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1))
(1): ReLU(inplace=True)
(2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1))
(3): ReLU(inplace=True)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1)
(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1))
(6): ReLU(inplace=True)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1))
(8): ReLU(inplace=True)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1))
(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1))
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1))
(15): ReLU(inplace=True)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1))
(17): ReLU(inplace=True)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(22): ReLU(inplace=True)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(24): ReLU(inplace=True)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(26): ReLU(inplace=True)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(31): ReLU(inplace=True)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(33): ReLU(inplace=True)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1))
(35): ReLU(inplace=True)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1)

```



```
1 def get_features(pretrained_model, image):
2     layers = {
3         '0': 'conv_1',
4         '5': 'conv_2',
5         '10': 'conv_3',
6         '19': 'conv_4',
7         '28': 'conv_5'
8     }
```

```
(0): Conv2d(3, 64, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(1): ReLU(inplace=True)
(2): Conv2d(64, 64, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(3): ReLU(inplace=True)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(5): Conv2d(64, 128, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(6): ReLU(inplace=True)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(8): ReLU(inplace=True)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(15): ReLU(inplace=True)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(17): ReLU(inplace=True)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(22): ReLU(inplace=True)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(24): ReLU(inplace=True)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(26): ReLU(inplace=True)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(31): ReLU(inplace=True)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(33): ReLU(inplace=True)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=1, padding=1, bias=False)
(35): ReLU(inplace=True)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

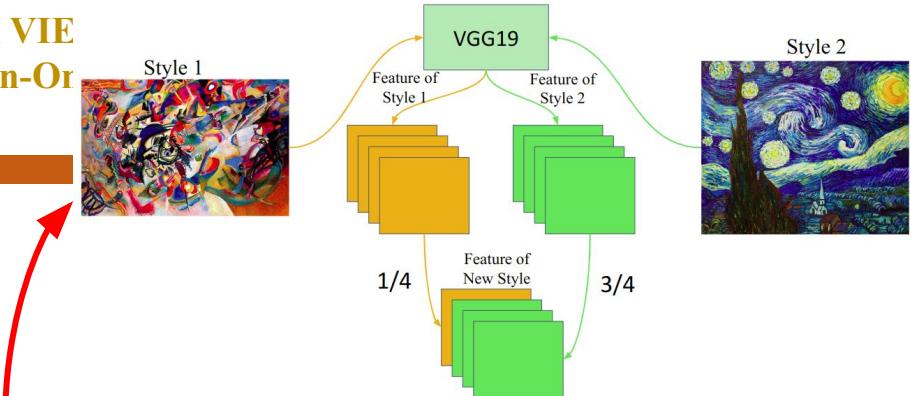


```

9     features = {}
10    x = image
11    for name, pretrained_layer in pretrained_model._modules.items():
12        x = pretrained_layer(x)
13        if name in layers:
14            features[layers[name]] = x
15    return features
16

(0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(1): ReLU(inplace=True)
(2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(3): ReLU(inplace=True)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(6): ReLU(inplace=True)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(8): ReLU(inplace=True)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(15): ReLU(inplace=True)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(17): ReLU(inplace=True)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(22): ReLU(inplace=True)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(24): ReLU(inplace=True)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(26): ReLU(inplace=True)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(31): ReLU(inplace=True)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(33): ReLU(inplace=True)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(35): ReLU(inplace=True)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

```



```
content_layers = ['conv_4']
style_layers = ['conv_1', 'conv_2',
                'conv_3', 'conv_4',
                'conv_5']
```

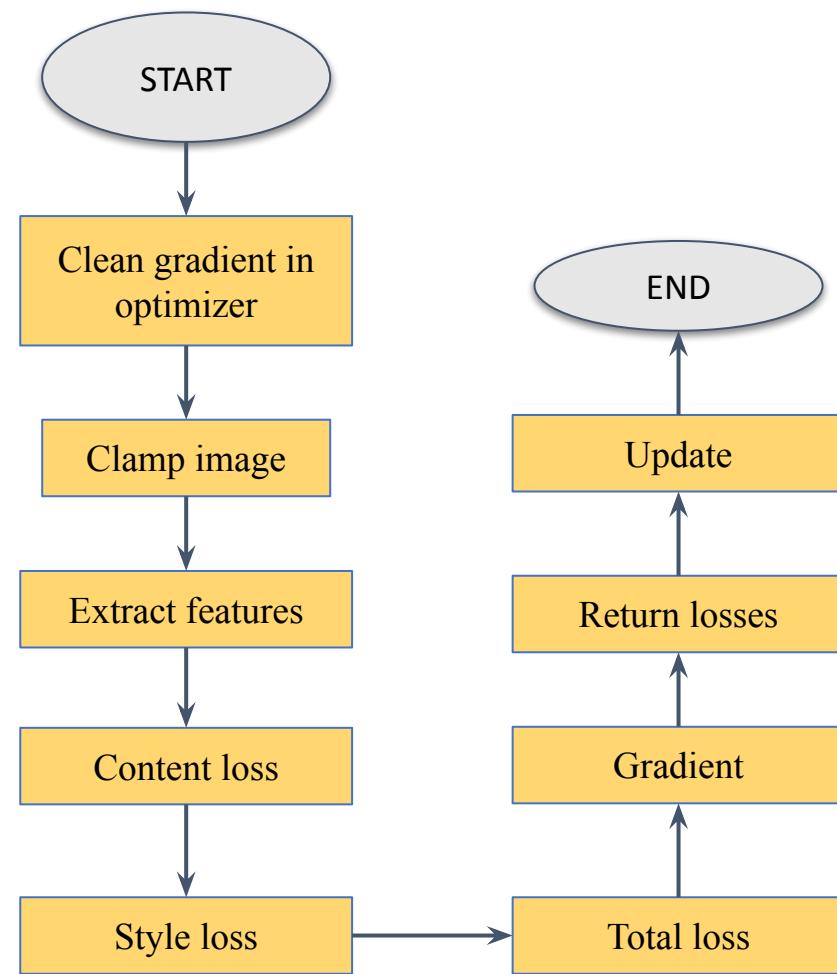
```
1 def get_dual_style(style_features1, style_features2, style_layers):
2     final_style_features = {}
3     for layer in style_layers:
4         sf1 = style_features1[layer]
5         sf2 = style_features2[layer]
6         sf1_size = int(sf1.size()[1] / 4)
7         final_style_features[layer] = torch.concatenate([
8             sf1[:, :sf1_size, :, :], 
9             sf2[:, sf1_size:, :, :]
10            ], dim=1)
11 return final_style_features
```

```
1 content_features = get_features(VGG19_pretrained, content_img)
2 style_features1 = get_features(VGG19_pretrained, style_img1)
3 style_features2 = get_features(VGG19_pretrained, style_img2)
4 final_style_features = get_dual_style(style_features1, style_features2, style_layers)
```

Style Transfer with 2 Style Images

❖ Style Transfer

```
1 def style_transfer_(model, optimizer, target_img,
2                     content_features, style_features,
3                     style_layers, content_weight, style_weight):
4
5     optimizer.zero_grad()
6     with torch.no_grad():
7         target_img.clamp_(0, 1)
8     target_features = get_features(model, target_img)
9
10    content_loss = ContentLoss(content_features['conv_4'],
11                               target_features['conv_4'])
12
13    style_loss = 0
14    for layer in style_layers:
15        target_gram = gram_matrix(target_features[layer])
16        style_gram = gram_matrix(style_features[layer])
17        style_loss += StyleLoss(style_gram, target_gram)
18
19    total_loss = content_loss*content_weight + style_loss*style_weight
20    total_loss.backward(retain_graph=True)
21    optimizer.step()
22    return total_loss, content_loss, style_loss
```



Style Transfer with 2 Style Images

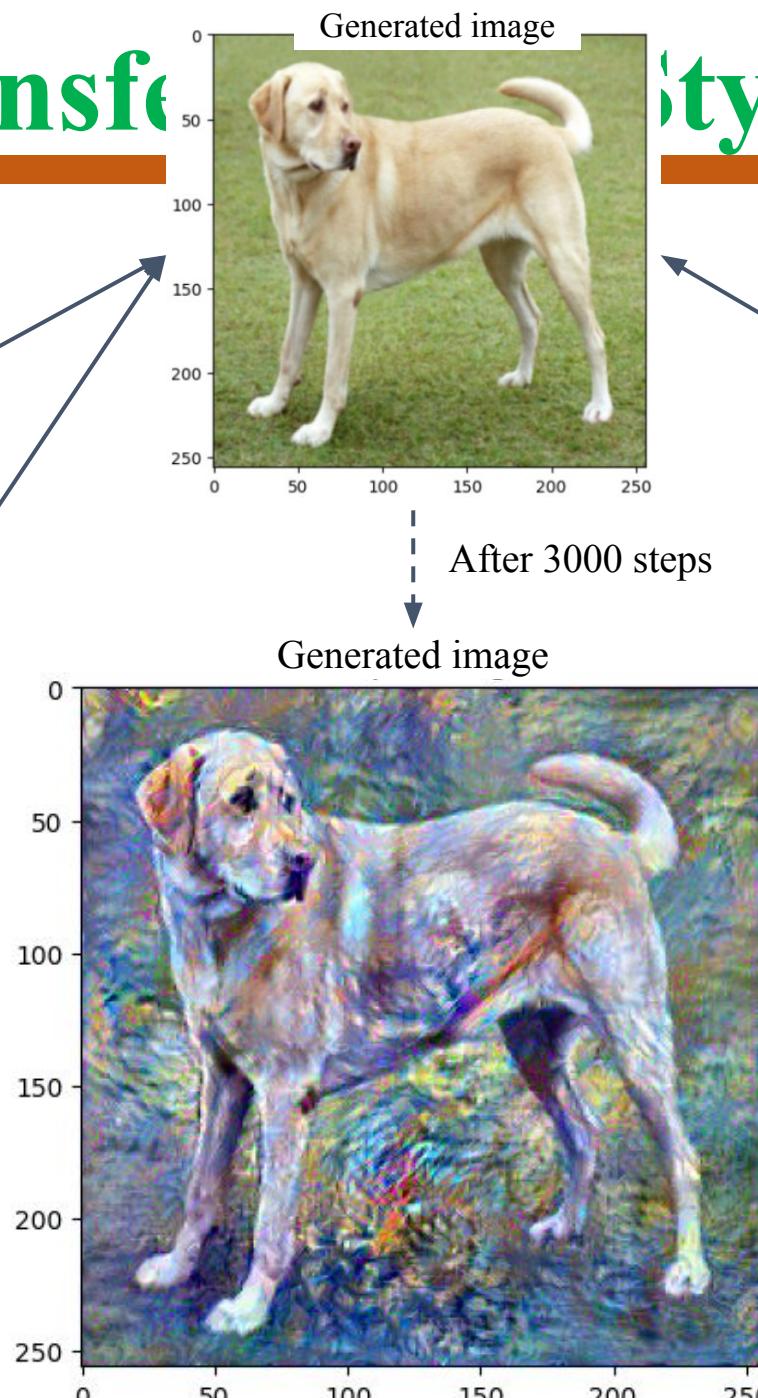
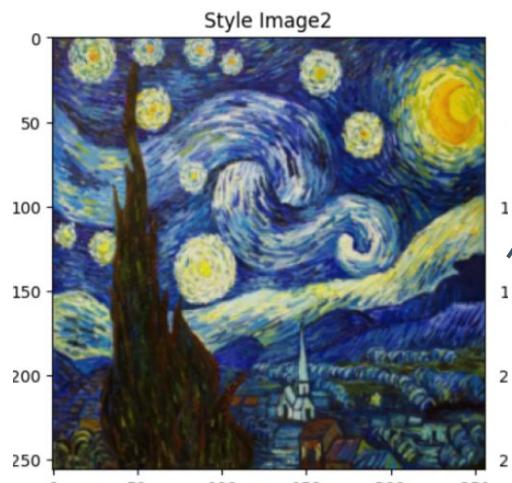
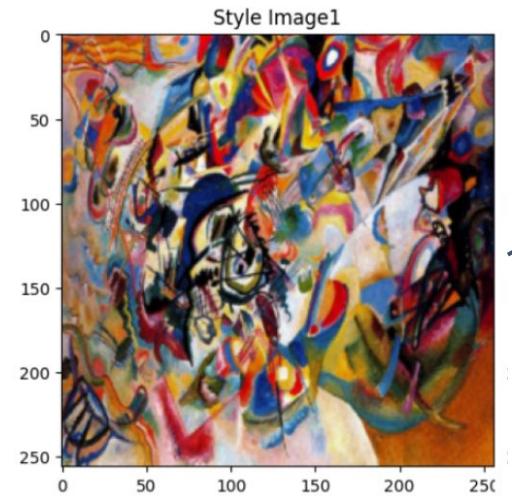
❖ Style Transfer

```
3 target_img = content_img.clone().requires_grad_(True).to(device)
4 optimizer = optim.Adam([target_img], lr=0.02)
```

```
1 STEPS = 3000
2
3 for step in range(STEPS):
4     optimizer.zero_grad()
5     with torch.no_grad():
6         target_img.clamp_(0, 1)
7
8     total_loss, content_loss, style_loss = style_transfer_(VGG19_pretrained, optimizer, target_img,
9                                         content_features, final_style_features,
10                                        style_layers, content_weight, style_weight)
11
12    if step % 100 == 99:
13        print(f"Epoch [{step+1}/{STEPS}] Total loss: {total_loss.item():.6f} - \
14          Content loss: {content_loss.item():.6f} - Style loss: {style_loss.item():.6f}")
15
16    with torch.no_grad():
17        target_img.clamp_(0, 1)
```

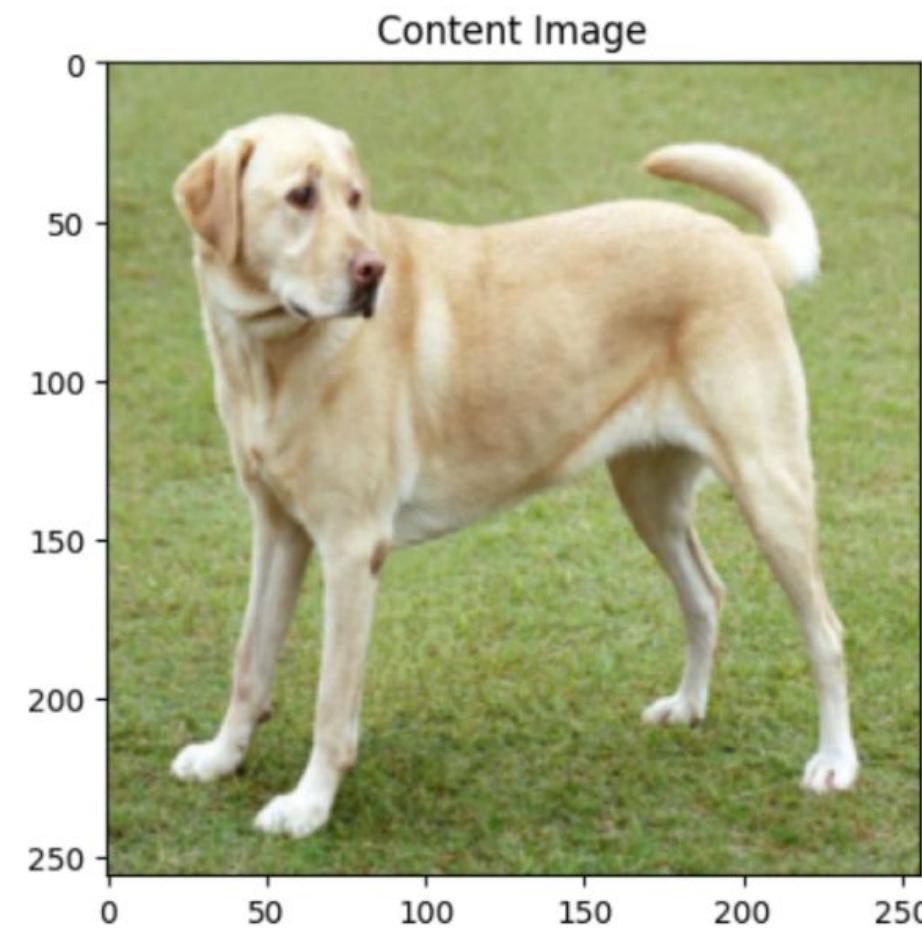
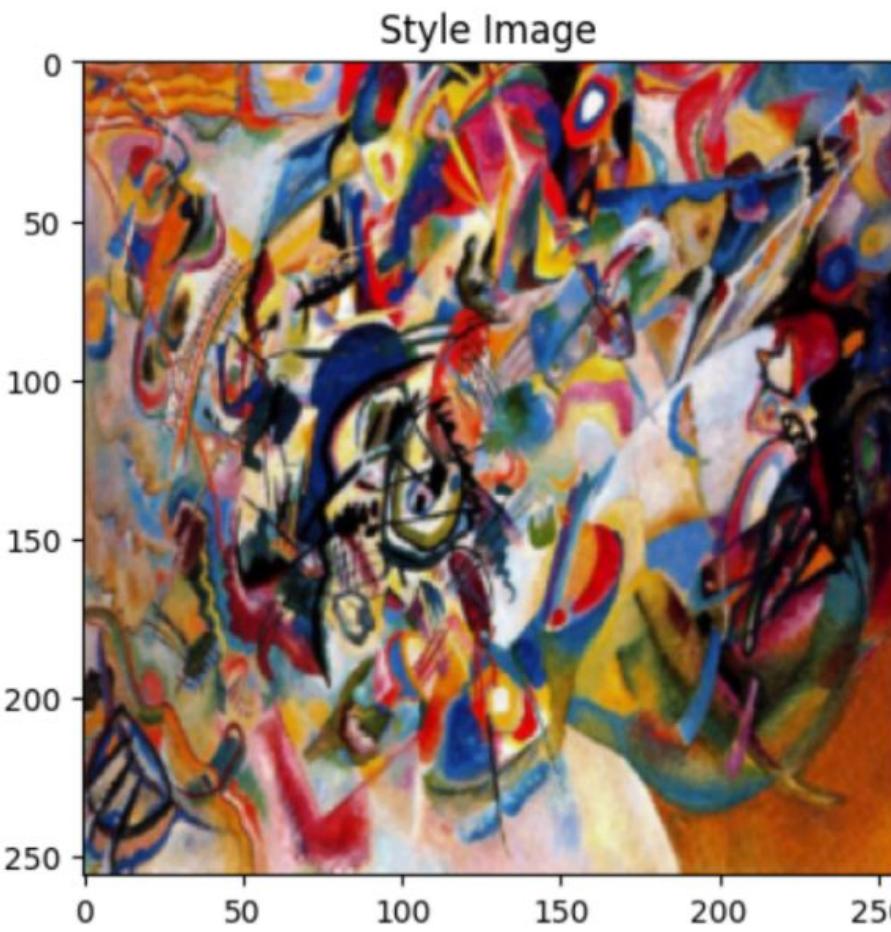
Style Transfer

Style Images



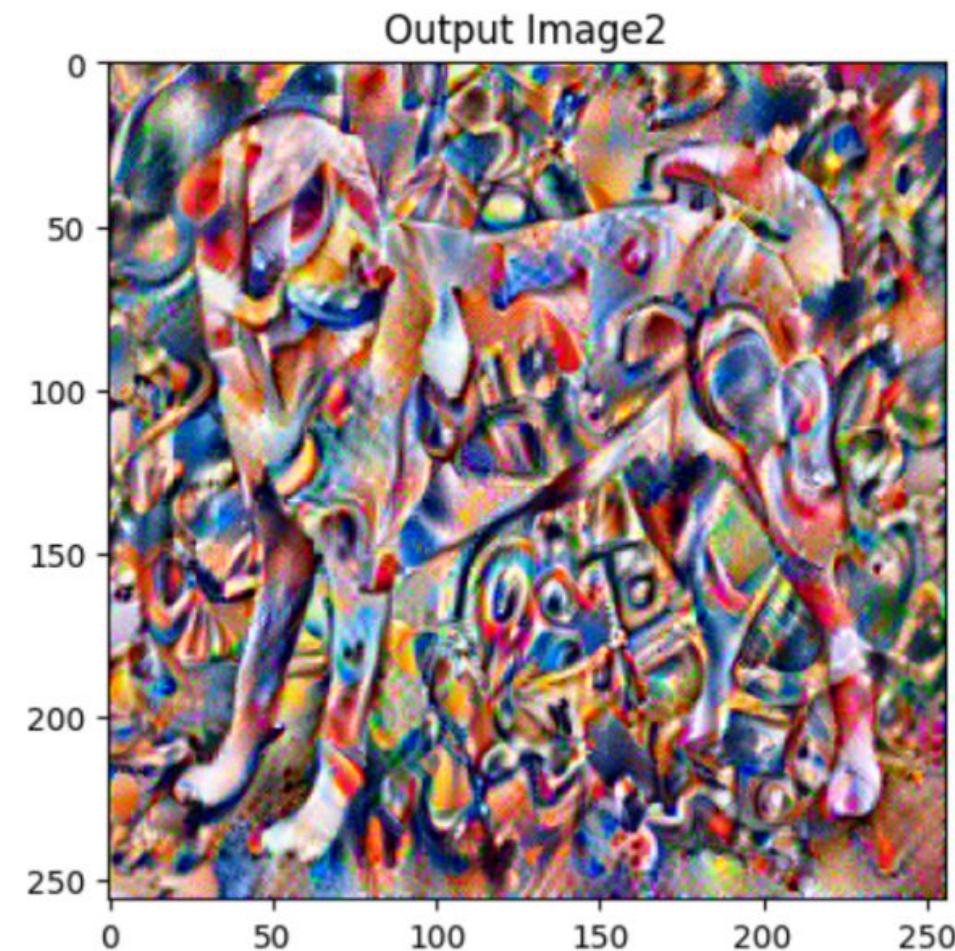
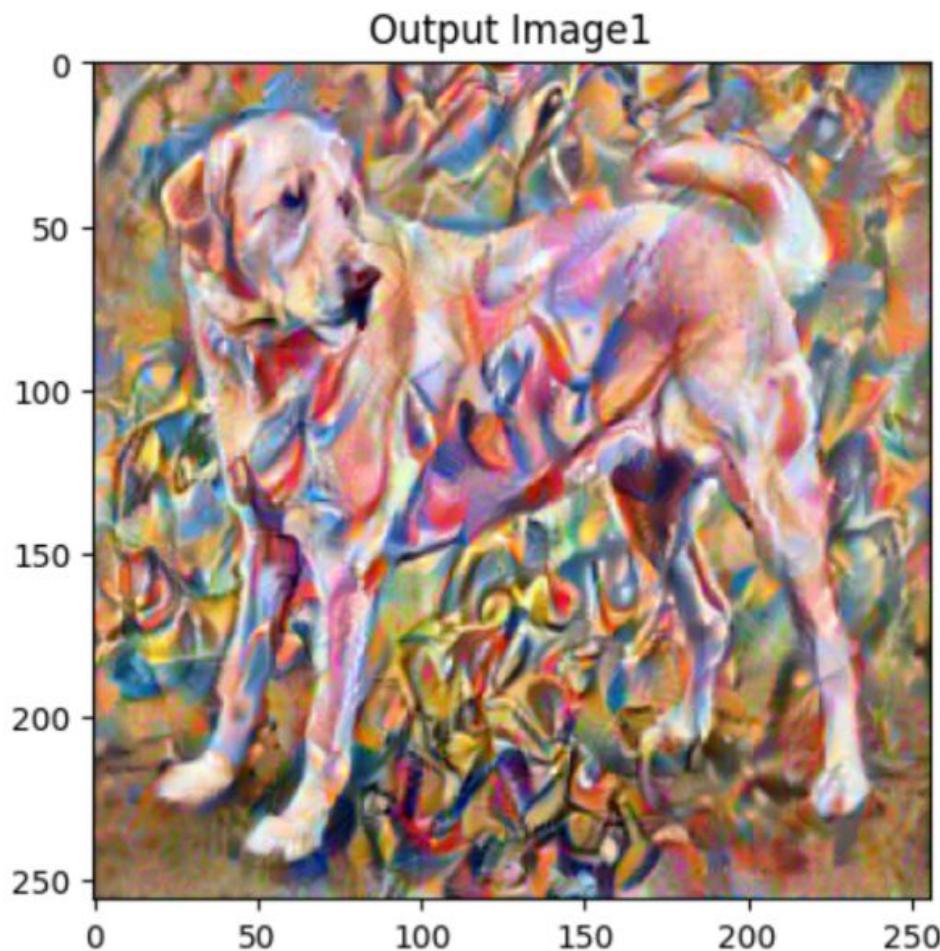
Style Transfer with Rotation + DE

Style Transfer with Rotation + DE

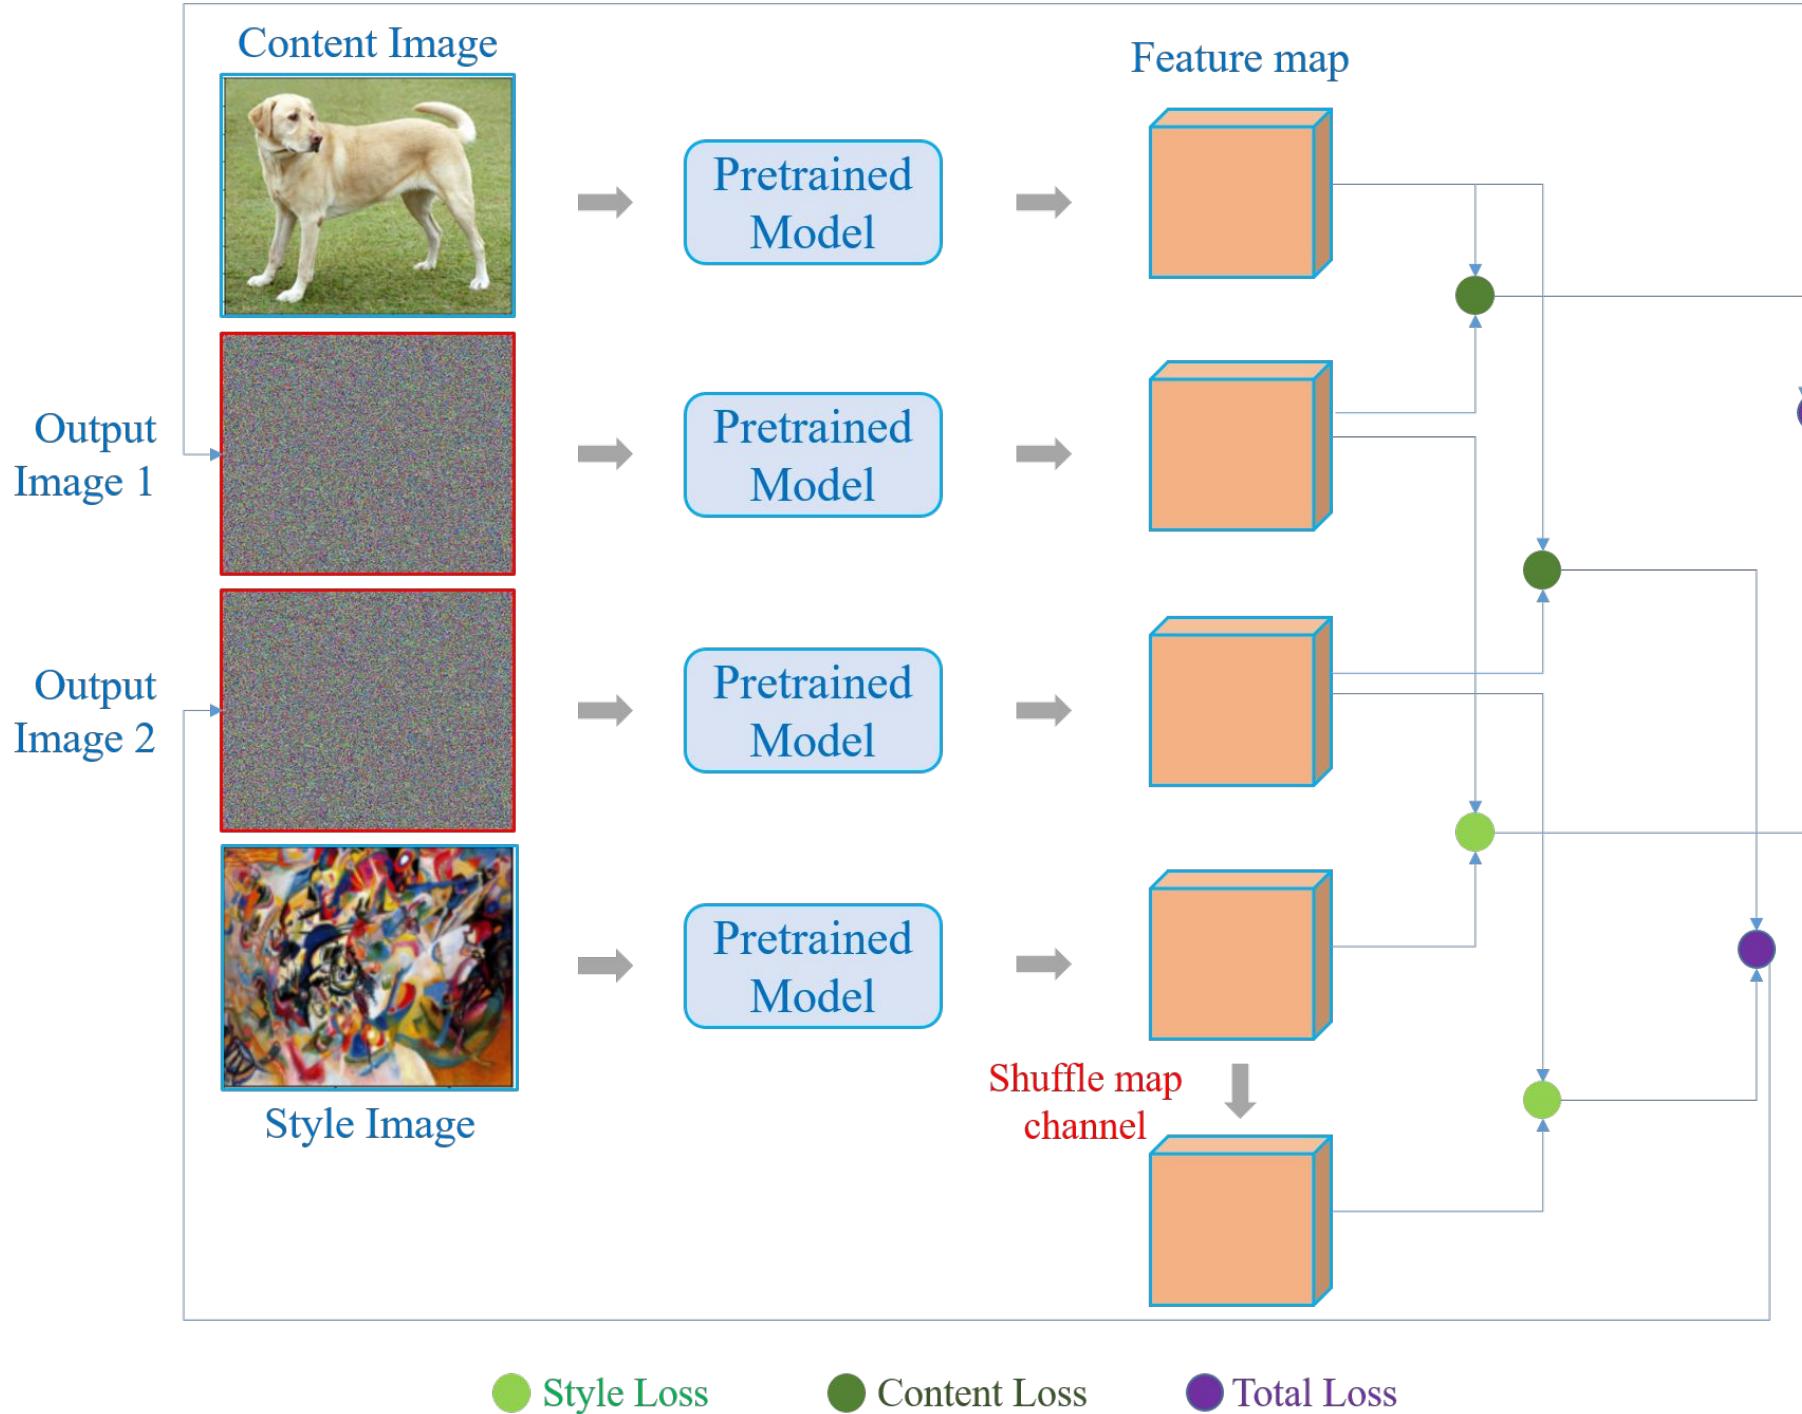


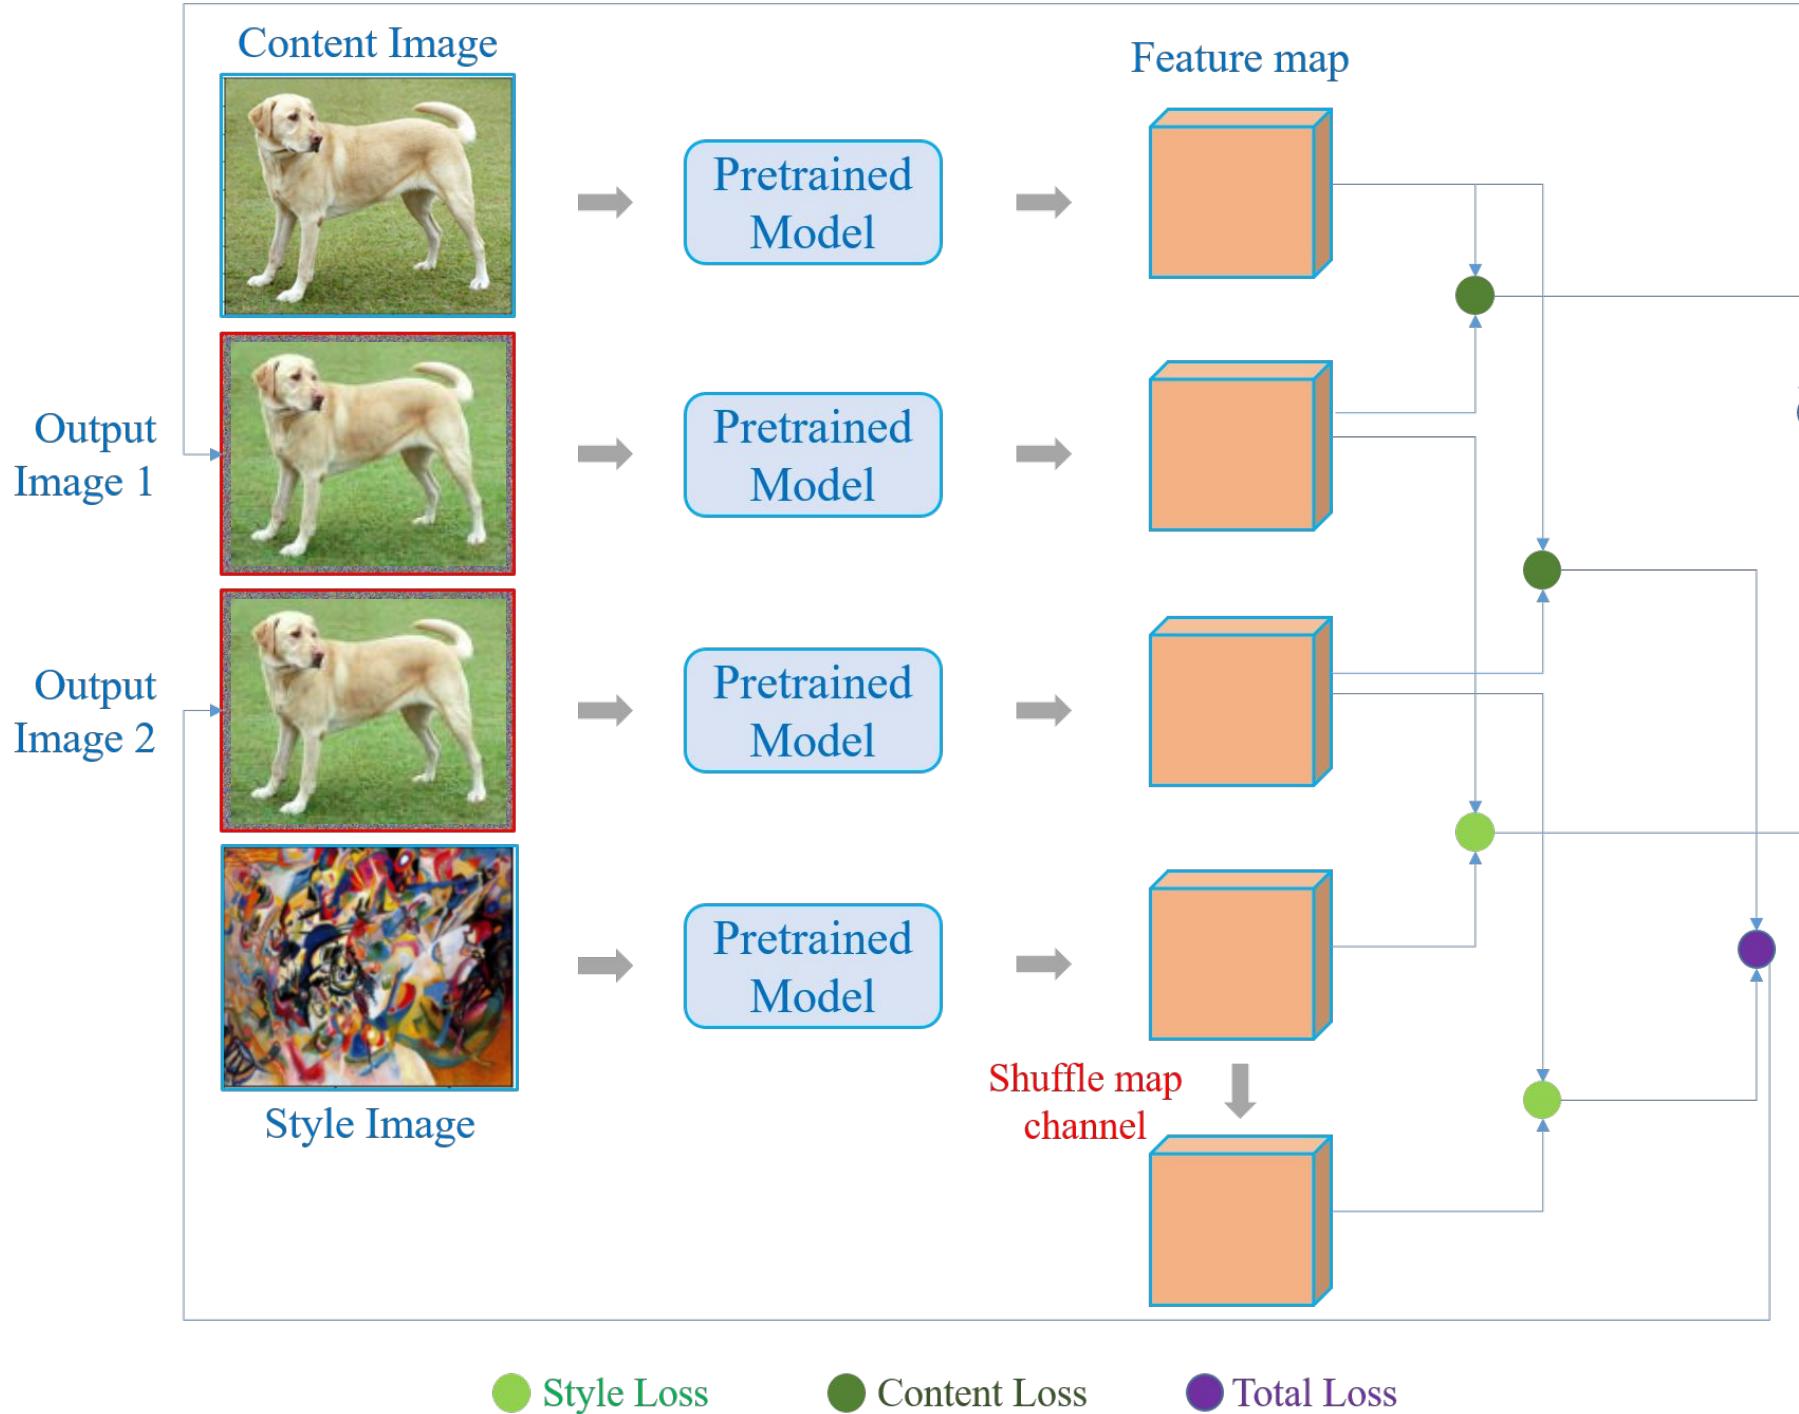
(a) Ảnh đầu vào của bài tập 2, ảnh trái là ảnh style 1, và ảnh phải là ảnh content

Style Transfer with Rotation + DE



(b) Ảnh output của bài tập 2, ảnh trái là ảnh khi transfer theo style 1 và ảnh phải là ảnh khi transfer theo style 1 đã thực hiện các biến đổi rotate 90, 180 và Differential Evolution

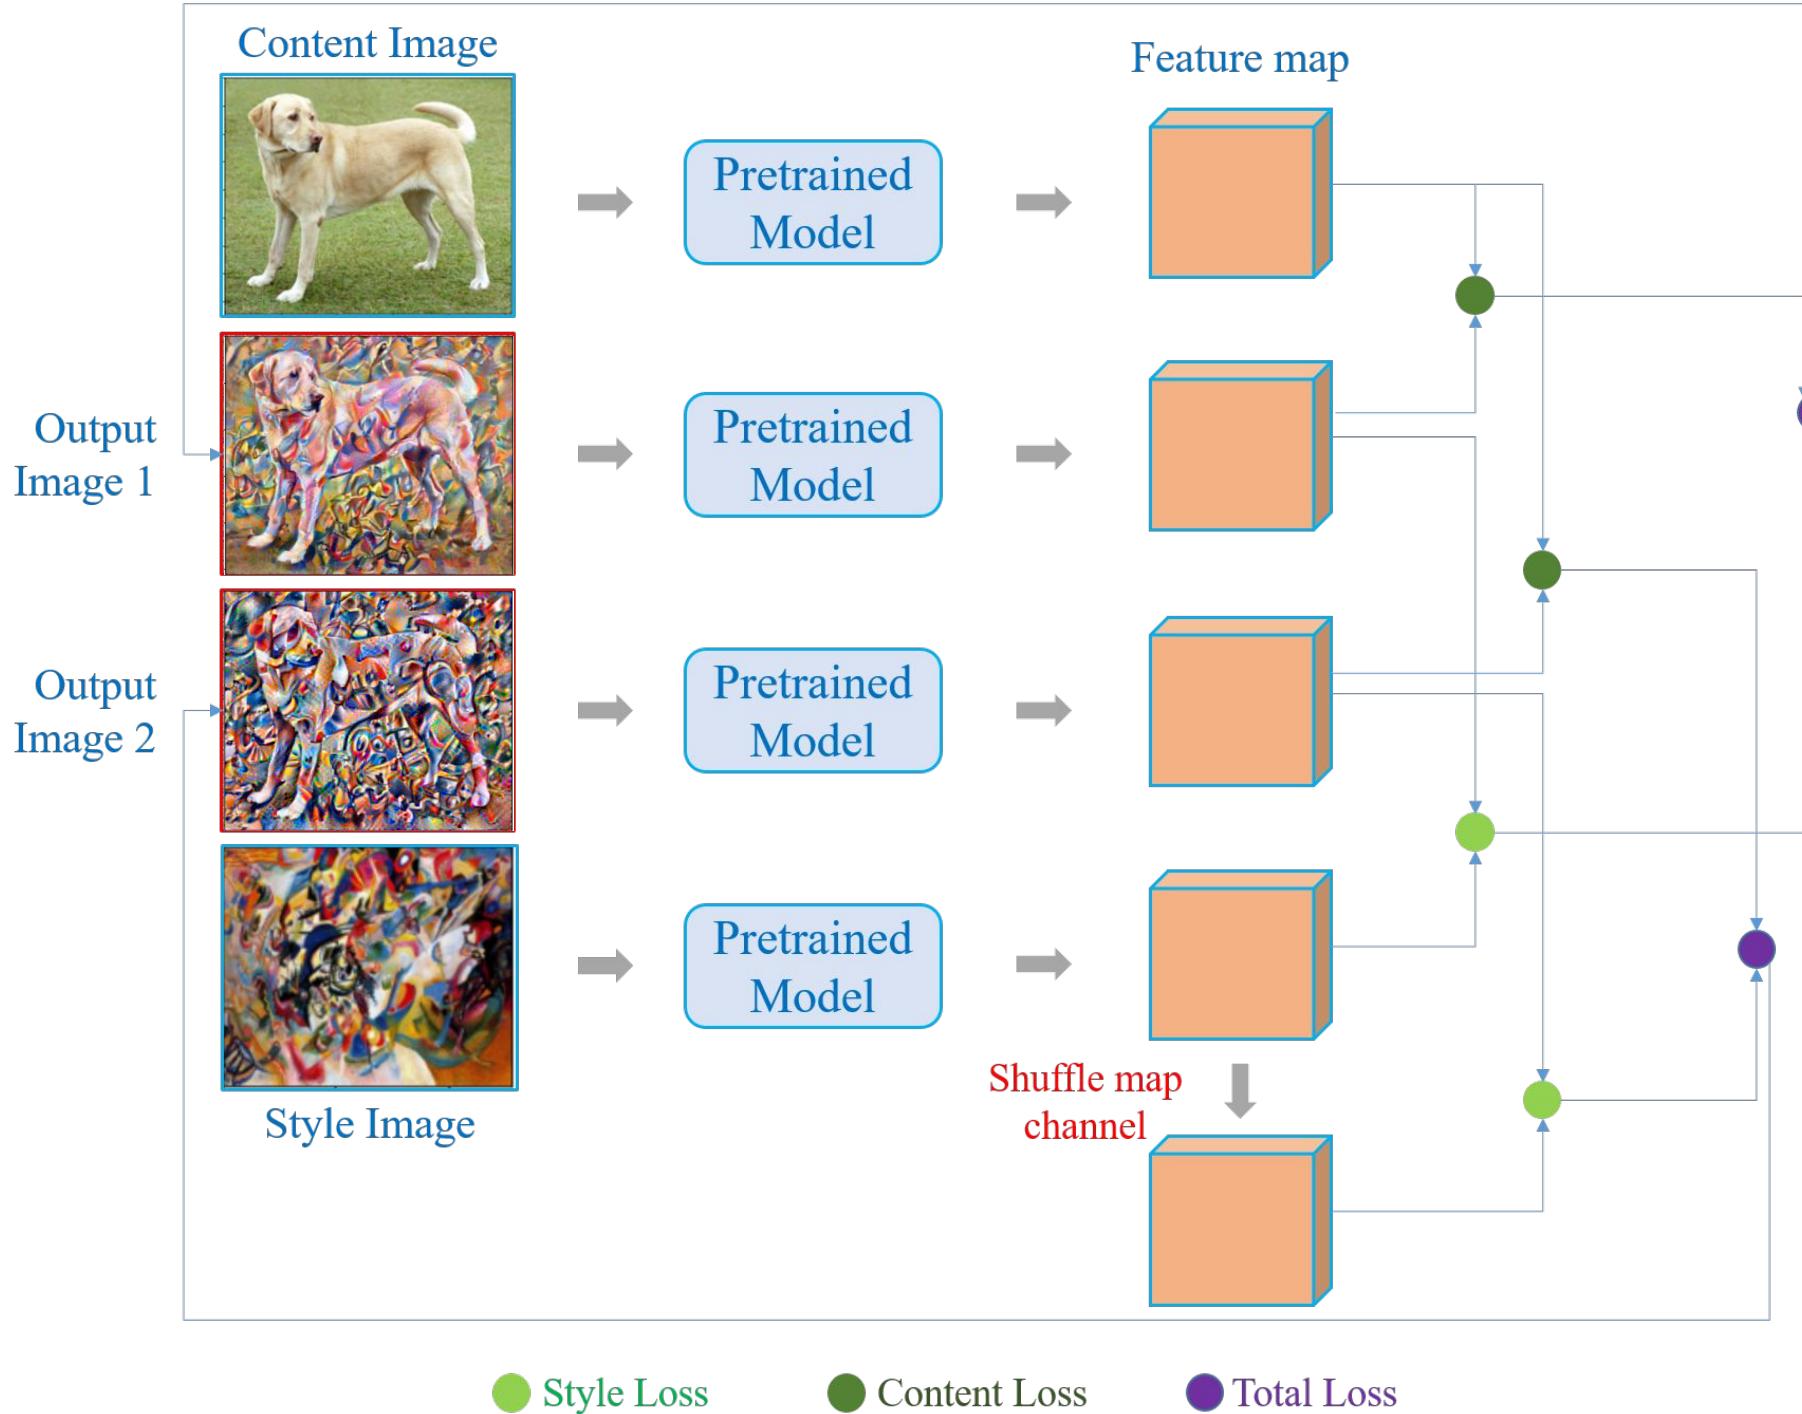




● Style Loss

● Content Loss

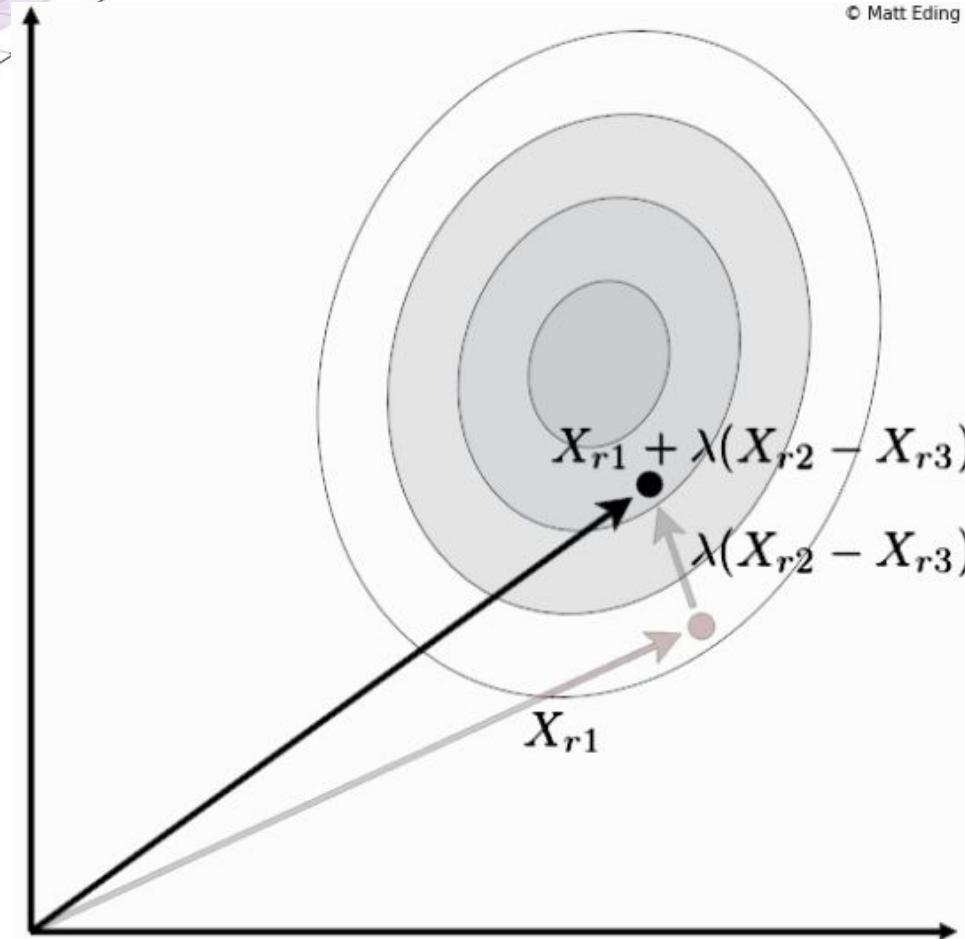
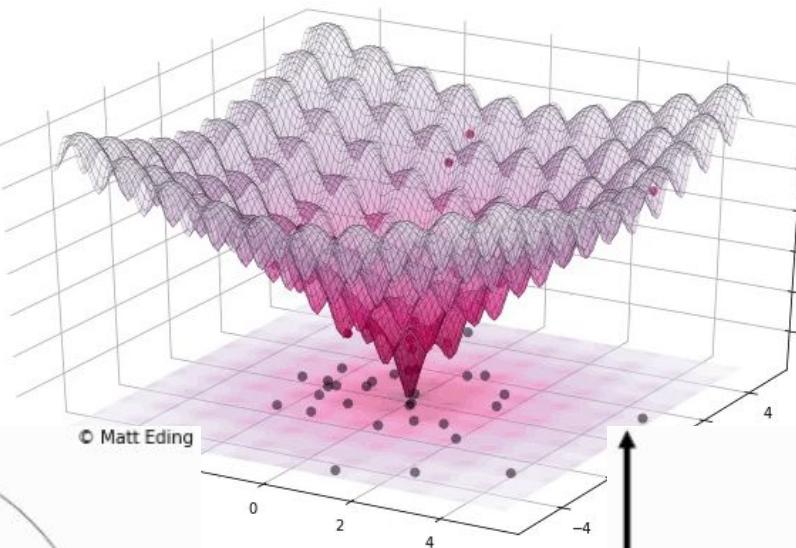
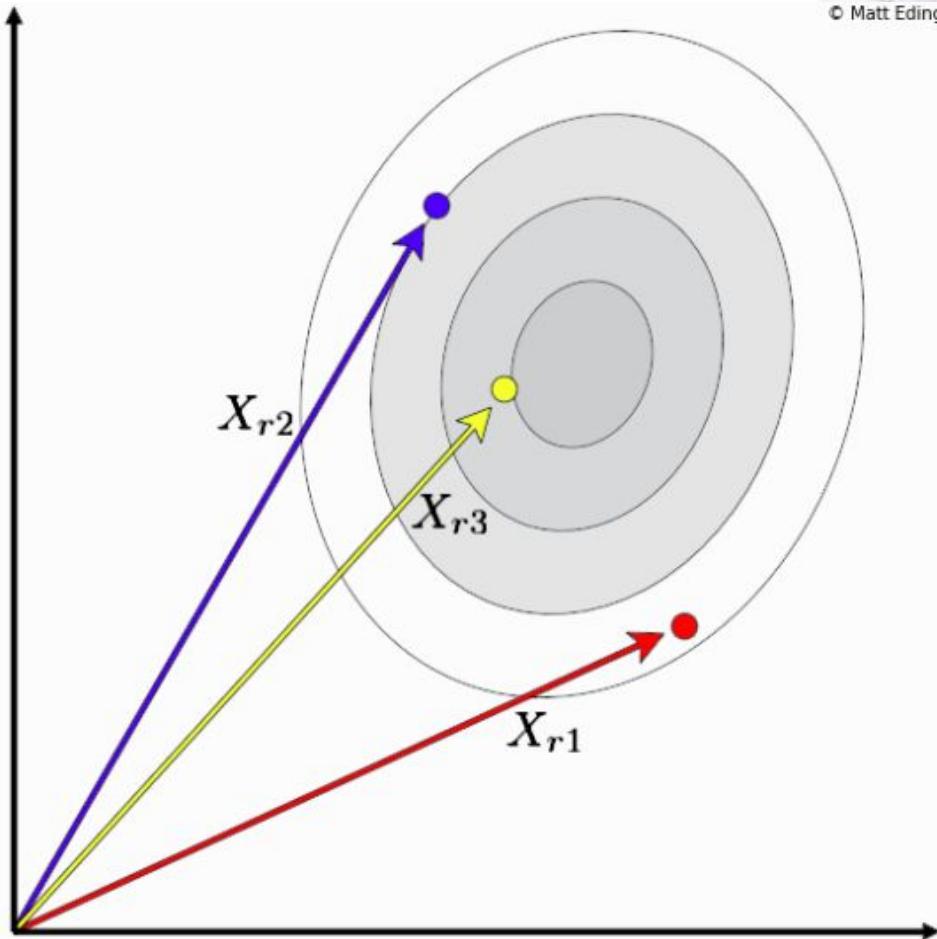
● Total Loss



Style Tra

on + DE

Differential Evolution



Style Transfer with Rotation + DE

0	0	0	0
0	0	0	0
0	0	0	0
1	1	1	1

(a) 0

0	0	0	1
0	0	0	1
0	0	0	1
0	0	0	1

(b) 90

1	1	1	1
0	0	0	0
0	0	0	0
0	0	0	0

(c) 180

$$x_{new} = x + (x_{90} - x_{180})$$

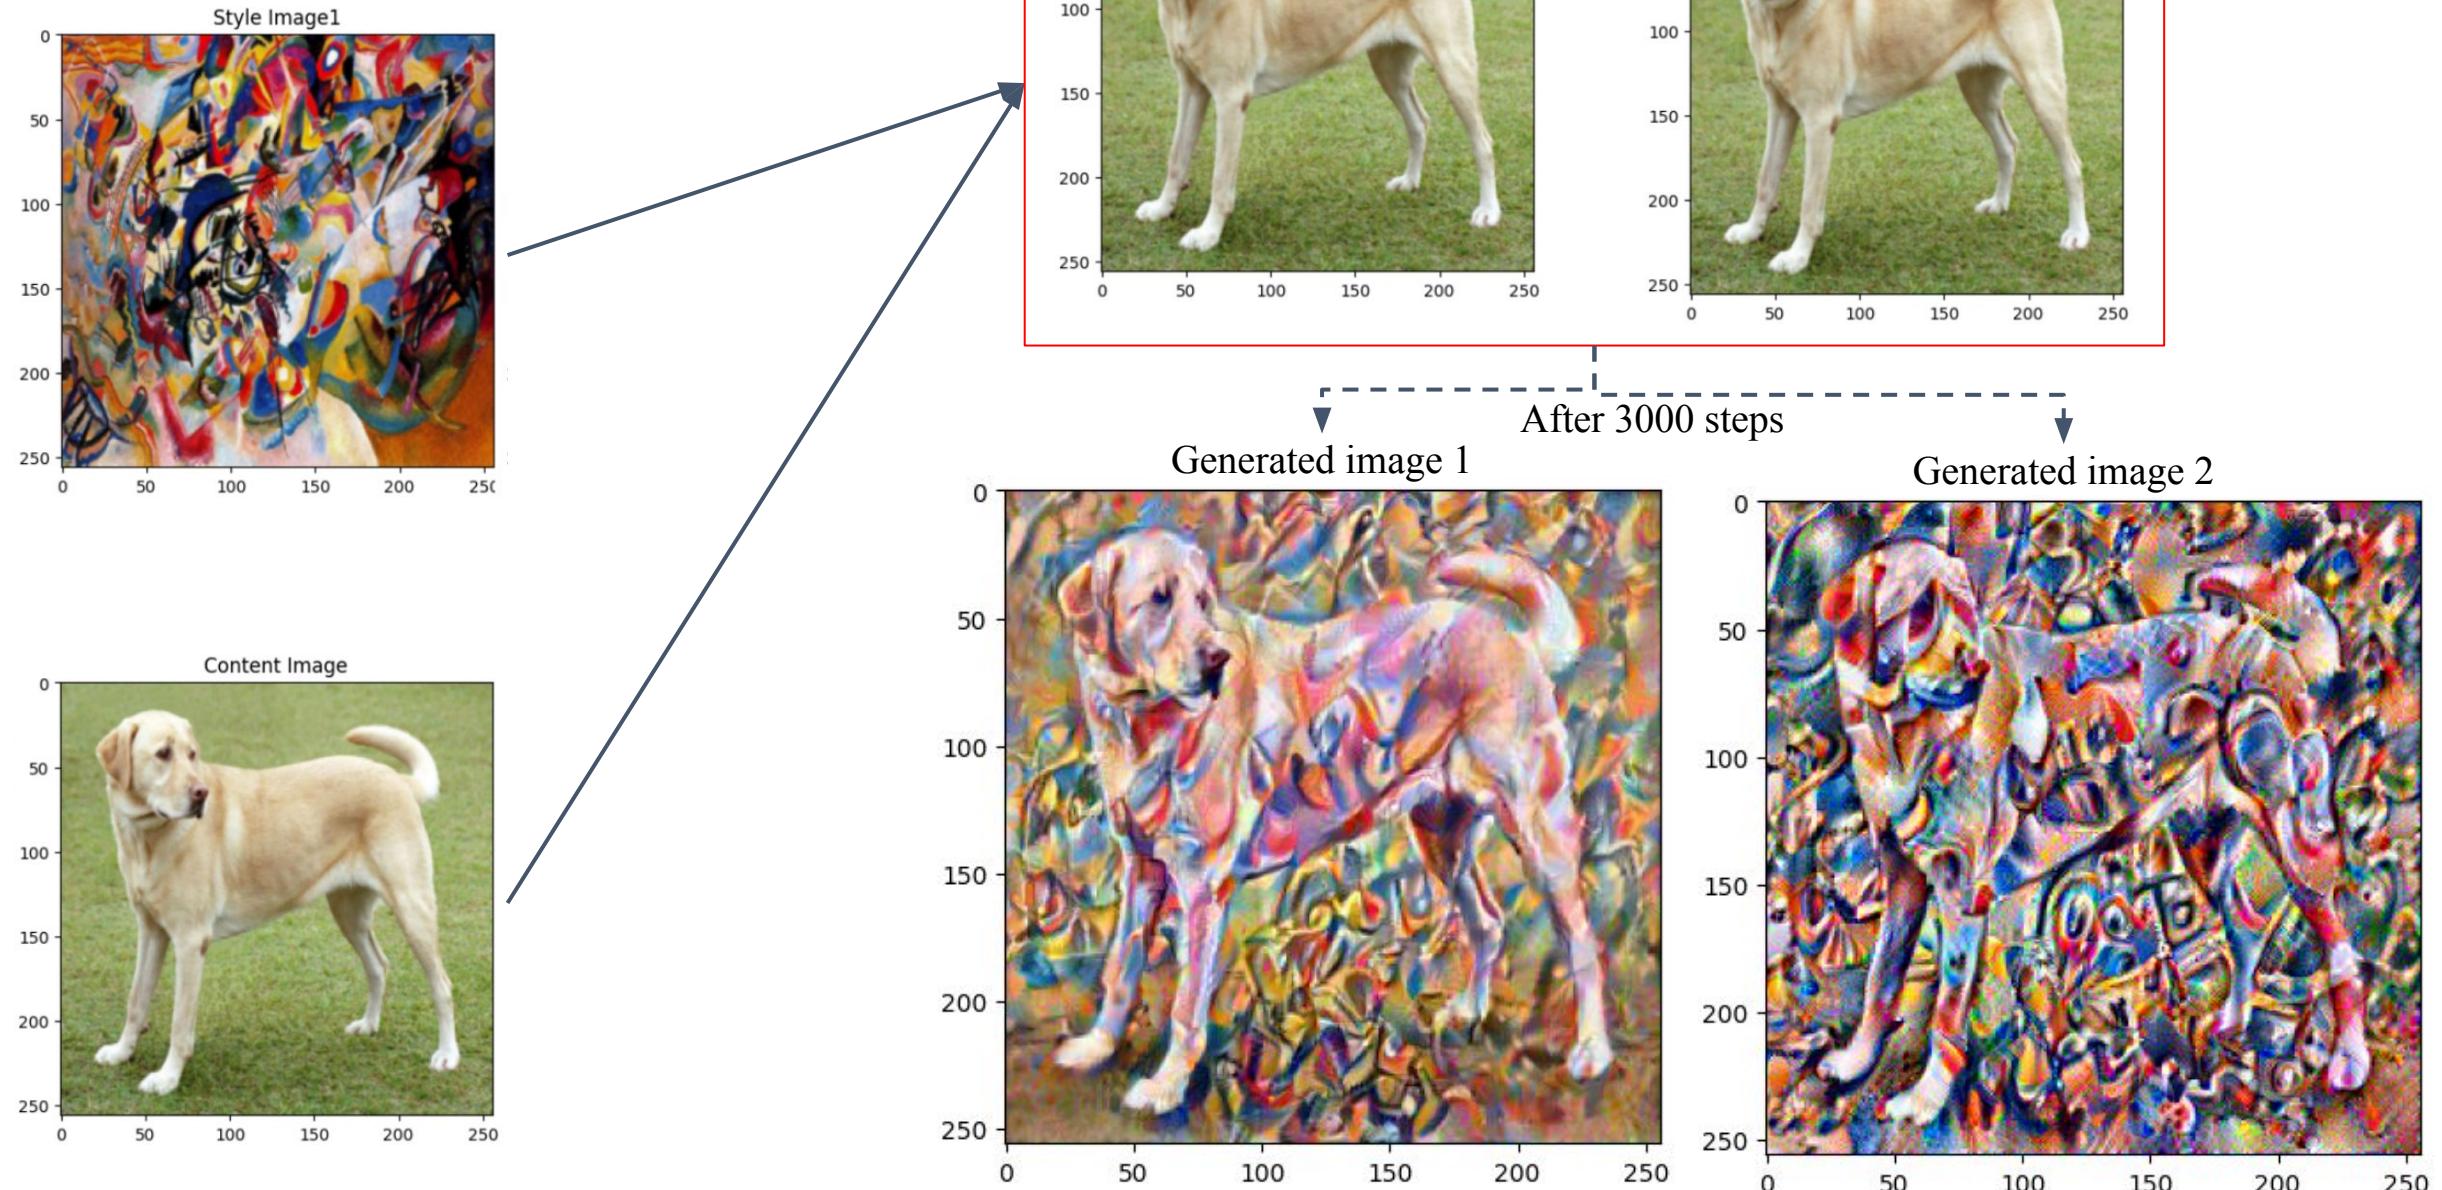
```
1 def rot_style_features(style_features, style_layers):
2     final_rot_style_features = {}
3     for layer in style_layers:
4         sf = style_features[layer].clone()
5         rot90 = torch.rot90(sf.clone(), 1, (2, 3))
6         rot180 = torch.rot90(rot90.clone(), 1, (2, 3))
7         final_rot = sf + (rot90 - rot180)
8         final_rot_style_features[layer] = final_rot
9     return final_rot_style_features
```

```
1 content_features = get_features(VGG19_pretrained, content_img)
2 style_features1 = get_features(VGG19_pretrained, style_img)
3 final_rot_style_features = rot_style_features(style_features1, style_layers)
```

```
1 import torch.optim as optim
2
3 target_img1 = content_img.clone().requires_grad_(True).to(device)
4 target_img2 = content_img.clone().requires_grad_(True).to(device)
5 optimizer1 = optim.Adam([target_img1], lr=0.02)
6 optimizer2 = optim.Adam([target_img2], lr=0.02)
```

```
1 STEPS = 3000
2
3 for step in range(STEPS):
4
5     total_loss1, content_loss1, style_loss1 = style_tranfer_(VGG19_pretrained, optimizer1, target_img1,
6                                         content_features, style_features1,
7                                         style_layers, content_weight, style_weight)
8
9     total_loss2, content_loss2, style_loss2 = style_tranfer_(VGG19_pretrained, optimizer2, target_img2,
10                                         content_features, final_rot_style_features,
11                                         style_layers, content_weight, style_weight)
12
13     if step % 100 == 99:
14         print(f"Epoch [{step+1}/{STEPS}] Total loss1: {total_loss1.item():.6f} - \
15               Content loss1: {content_loss1.item():.6f} - Style loss1: {style_loss1.item():.6f}")
16         print(f"Epoch [{step+1}/{STEPS}] Total loss2: {total_loss2.item():.6f} - \
17               Content loss2: {content_loss2.item():.6f} - Style loss2: {style_loss2.item():.6f}")
18
19     with torch.no_grad():
20         target_img1.clamp_(0, 1)
21         target_img2.clamp_(0, 1)
```

Style Transfer



Summary

- ✓ Reviewed Style Transfer
- ✓ Presented Style Transfer Examples
- ✓ Executed Style Transfer Using Two Style References
- ✓ Executed Style Transfer with Rotation + DE
- ✓ Reviewed fundamental necessary understanding
- ✓ Reviewed Milestone Papers

