


Extra Class

Advanced CNN Architecture

Nguyen Quoc Thai



CONTENT

(1) – ResNet


(2) – Feature Extraction

(3) – Image Similarity

1 – ResNet



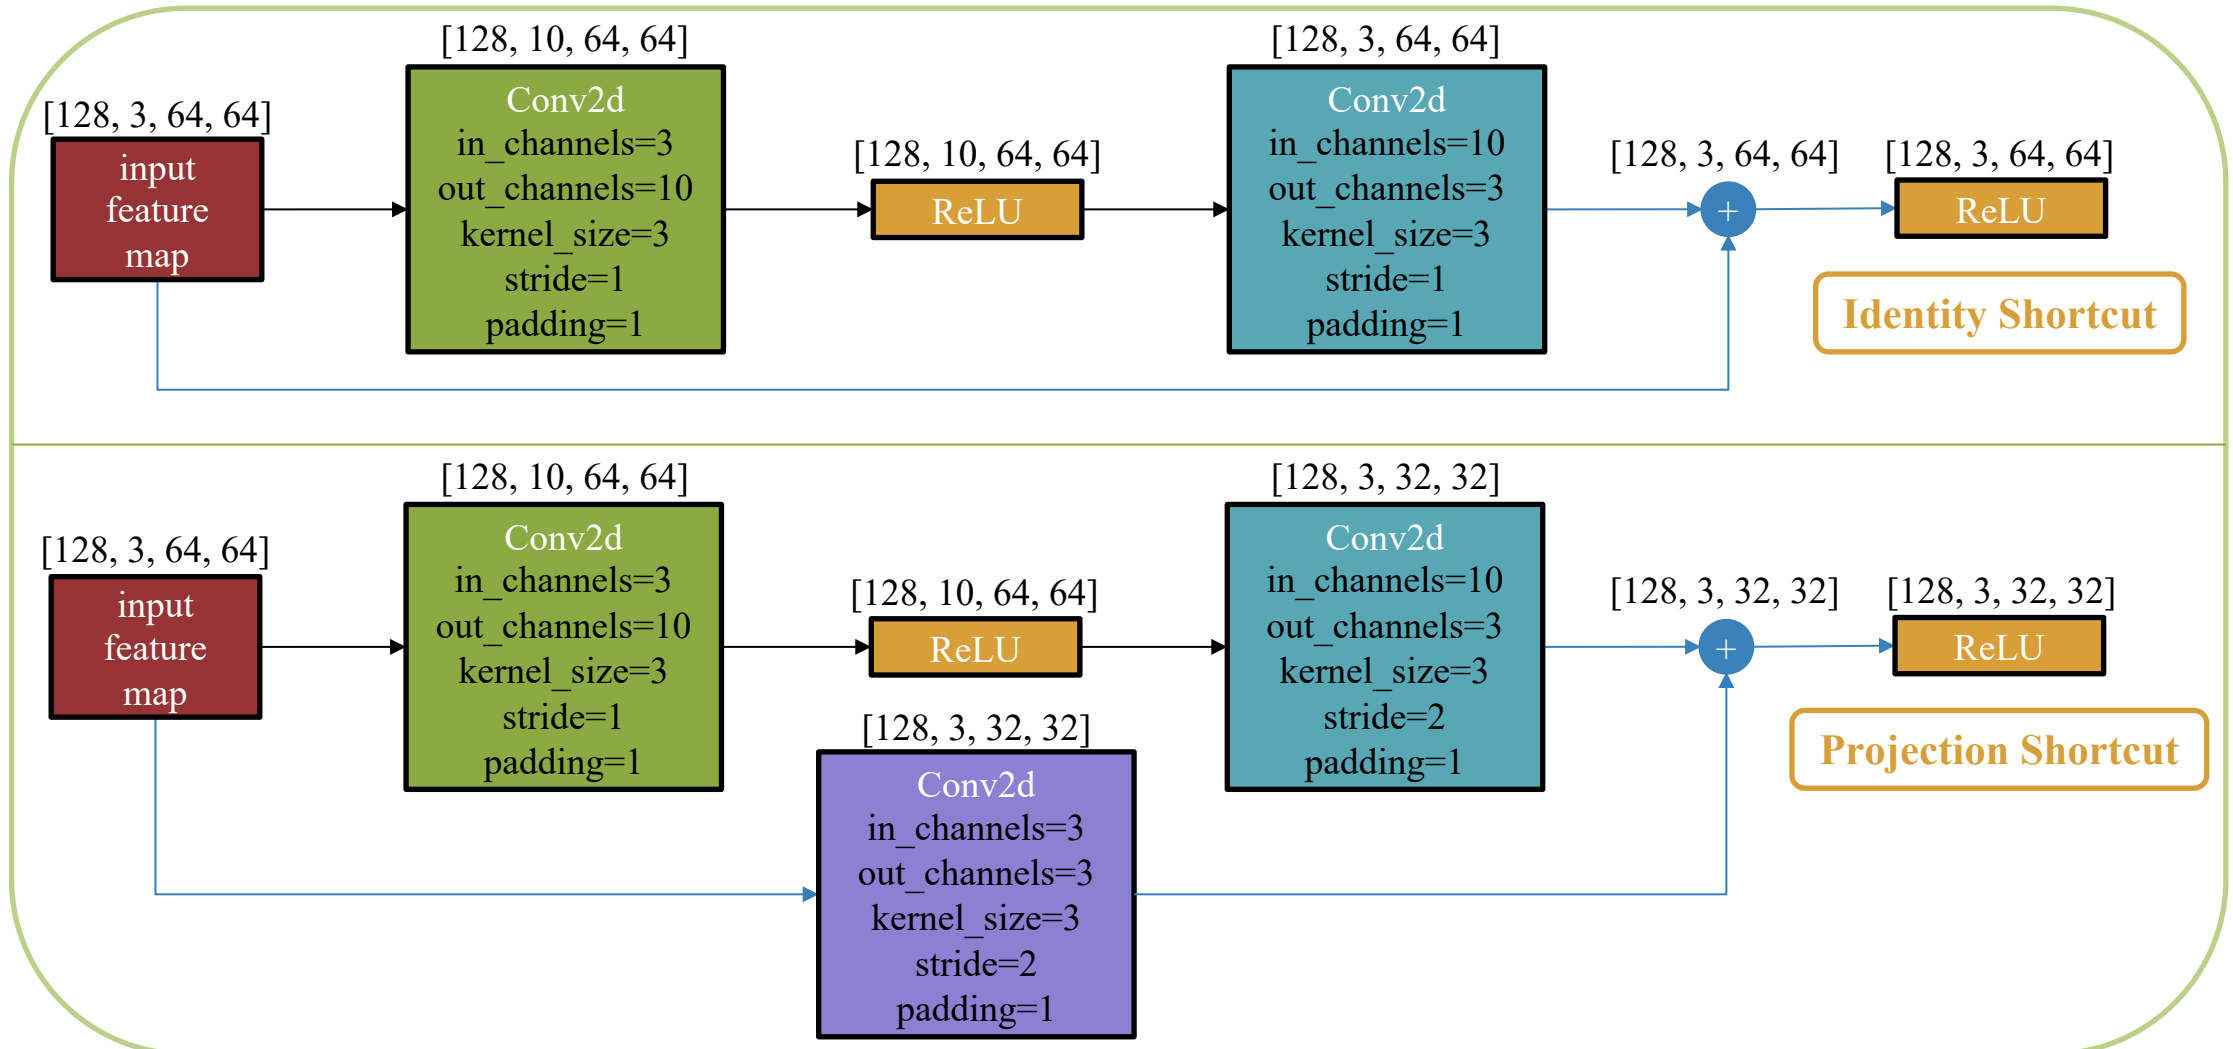
Review: Skip Connection



1 – ResNet




Review: Skip Connection




1 – ResNet



Review: Skip Connection



Short Skip Connection
(ResNet, ...)

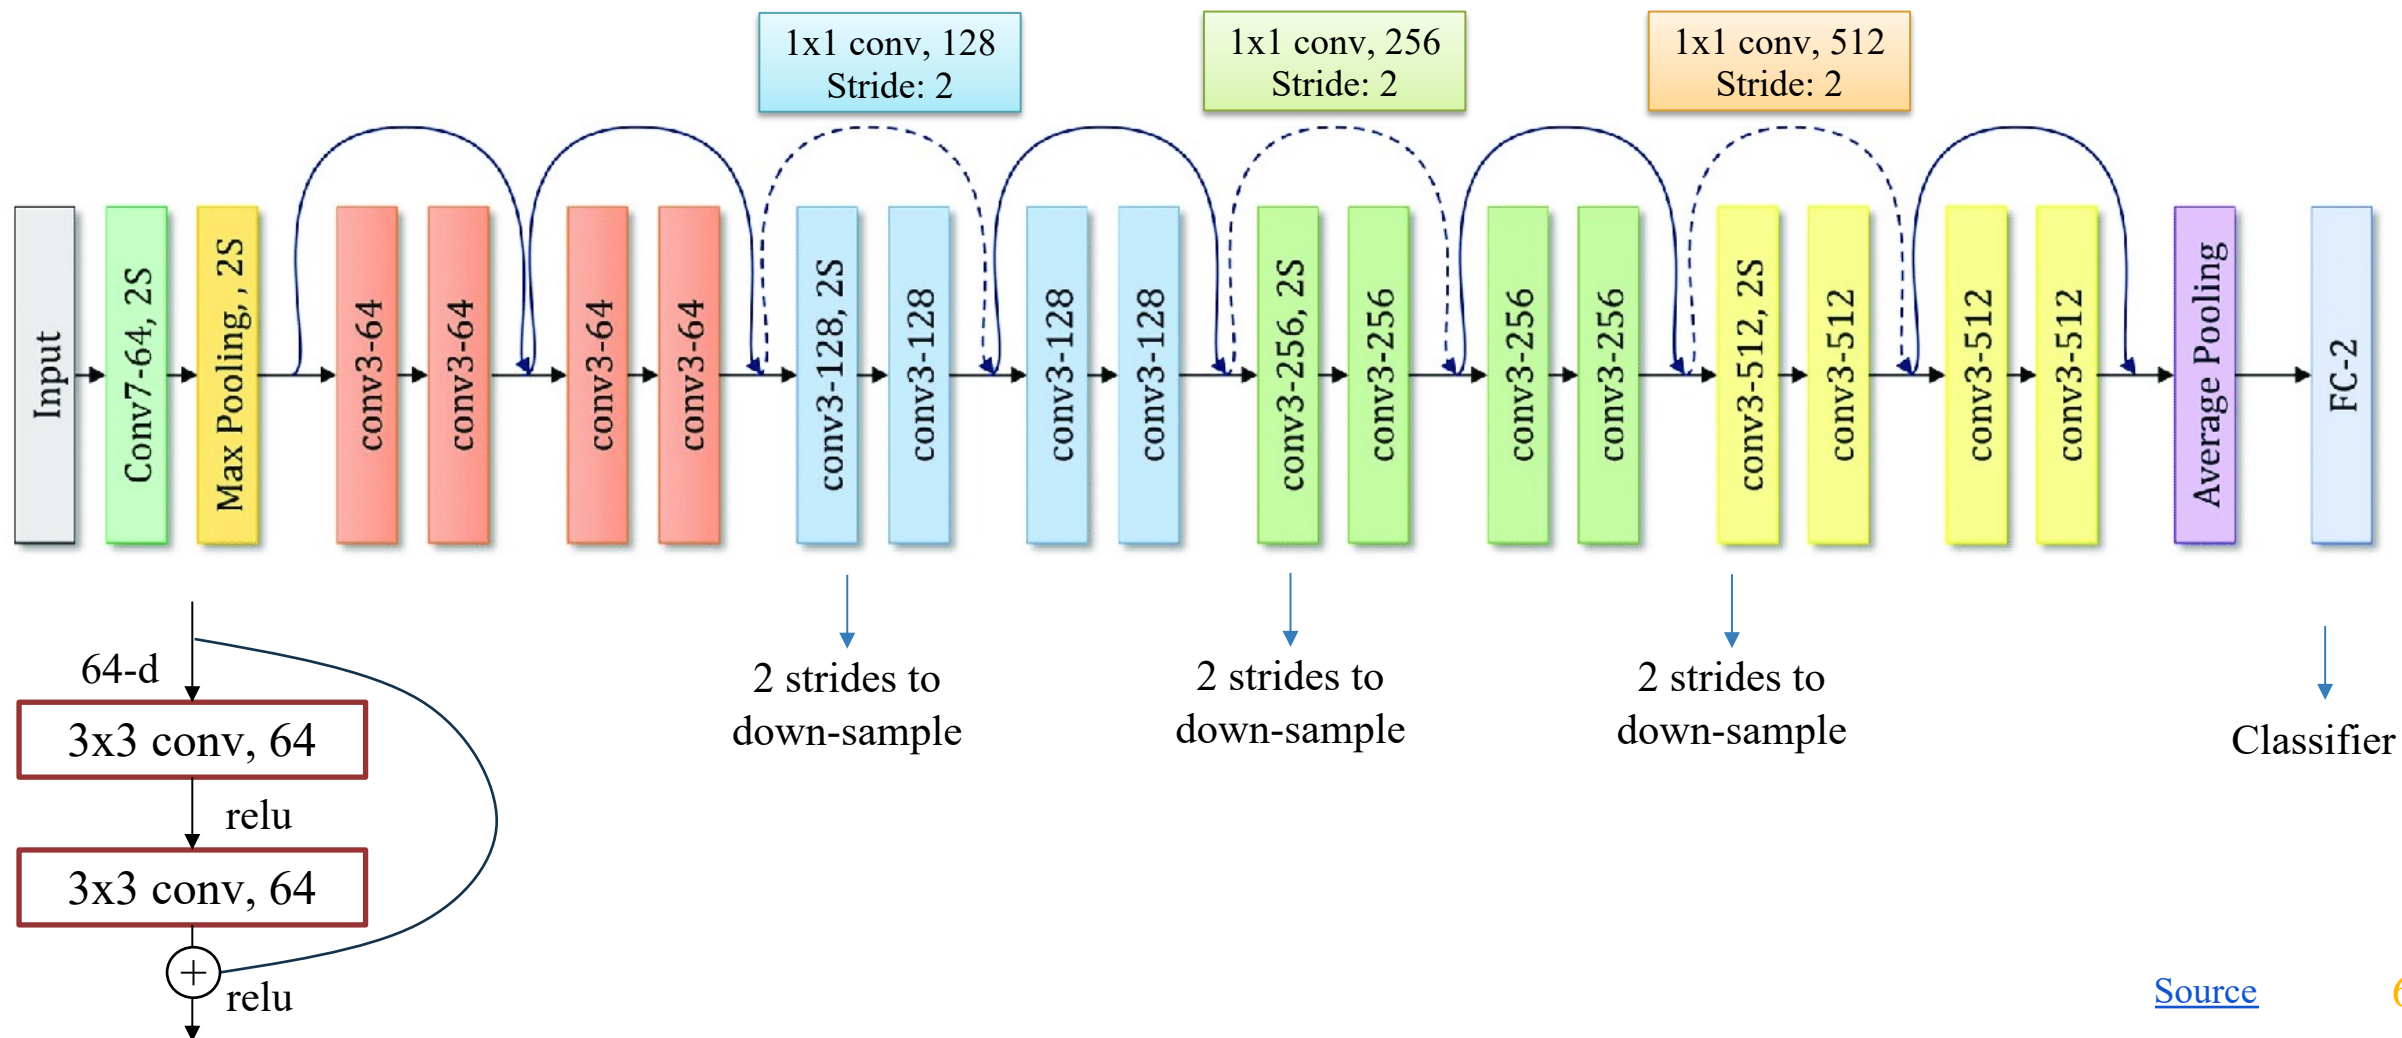


Long Skip Connection
(UNet, ...)

1 – ResNet

!

ResNet18



1 – ResNet



ResNet18

```
import torch

from torchvision import models
from torchsummary import summary

resnet_model = models.resnet18(weights=None)
```

```
=====
Total params: 11,689,512
Trainable params: 11,689,512
Non-trainable params: 0
-----
Input size (MB): 0.57
Forward/backward pass size (MB): 62.79
Params size (MB): 44.59
Estimated Total Size (MB): 107.96
-----
```

1 – ResNet



ResNet (18 – 34 – 50 – 101 – 152)

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	$7 \times 7, 64, \text{stride } 2$				
conv2_x	56×56	$3 \times 3 \text{ max pool, stride } 2$				
		$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				

1 – ResNet



CIFAR10 Classification using ResNet18

- ❖ Training: 50,000 images. Testing: 10,000 images
- ❖ Classes: 10

airplane



automobile



bird



cat



deer




dog




frog



horse



ship



truck



1 – ResNet



CIFAR10 Classification using ResNet18 – Demo

❖ Load Dataset

```
ROOT = './data'

train_data = datasets.CIFAR10(
    root=ROOT,
    train=True,
    download=True,
    transform=transforms.ToTensor()
)

test_data = datasets.CIFAR10(
    root=ROOT,
    train=False,
    download=True,
    transform=transforms.ToTensor()
)
```

1 – ResNet



CIFAR10 Classification using ResNet18 – Demo

❖ Preprocessing

```
train_data_stack = torch.stack([img for img, _ in train_data], dim=3)
```

```
train_data_stack.shape
```

```
torch.Size([3, 32, 32, 50000])
```

```
mean= train_data_stack.view(3,-1).mean(dim=1)
```

```
std= train_data_stack.view(3,-1).std(dim=1)
```

```
mean, std
```

```
(tensor([0.4914, 0.4822, 0.4465]), tensor([0.2470, 0.2435, 0.2616]))
```

```
data_transforms = transforms.Compose([  
    transforms.ToTensor(),  
    transforms.Normalize(mean, std)  
])
```

```
train_data.transform = data_transforms  
test_data.transform = data_transforms
```

```
BATCH_SIZE = 512
```

```
train_dataloader = data.DataLoader(  
    train_data,  
    shuffle=True,  
    batch_size=BATCH_SIZE  
)
```

```
test_dataloader = data.DataLoader(  
    test_data,  
    batch_size=BATCH_SIZE  
)
```

1 – ResNet



CIFAR10 Classification using ResNet18 – Demo

❖ Model

```
model = models.resnet18(weights=None)
in_features = model.fc.in_features
model.fc = nn.Linear(in_features, len(train_data.classes))
model.to(device)
```

```
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=1000, bias=True)
```



```
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=10, bias=True)
```

1 – ResNet



CIFAR10 Classification using ResNet18 – Demo

❖ Training

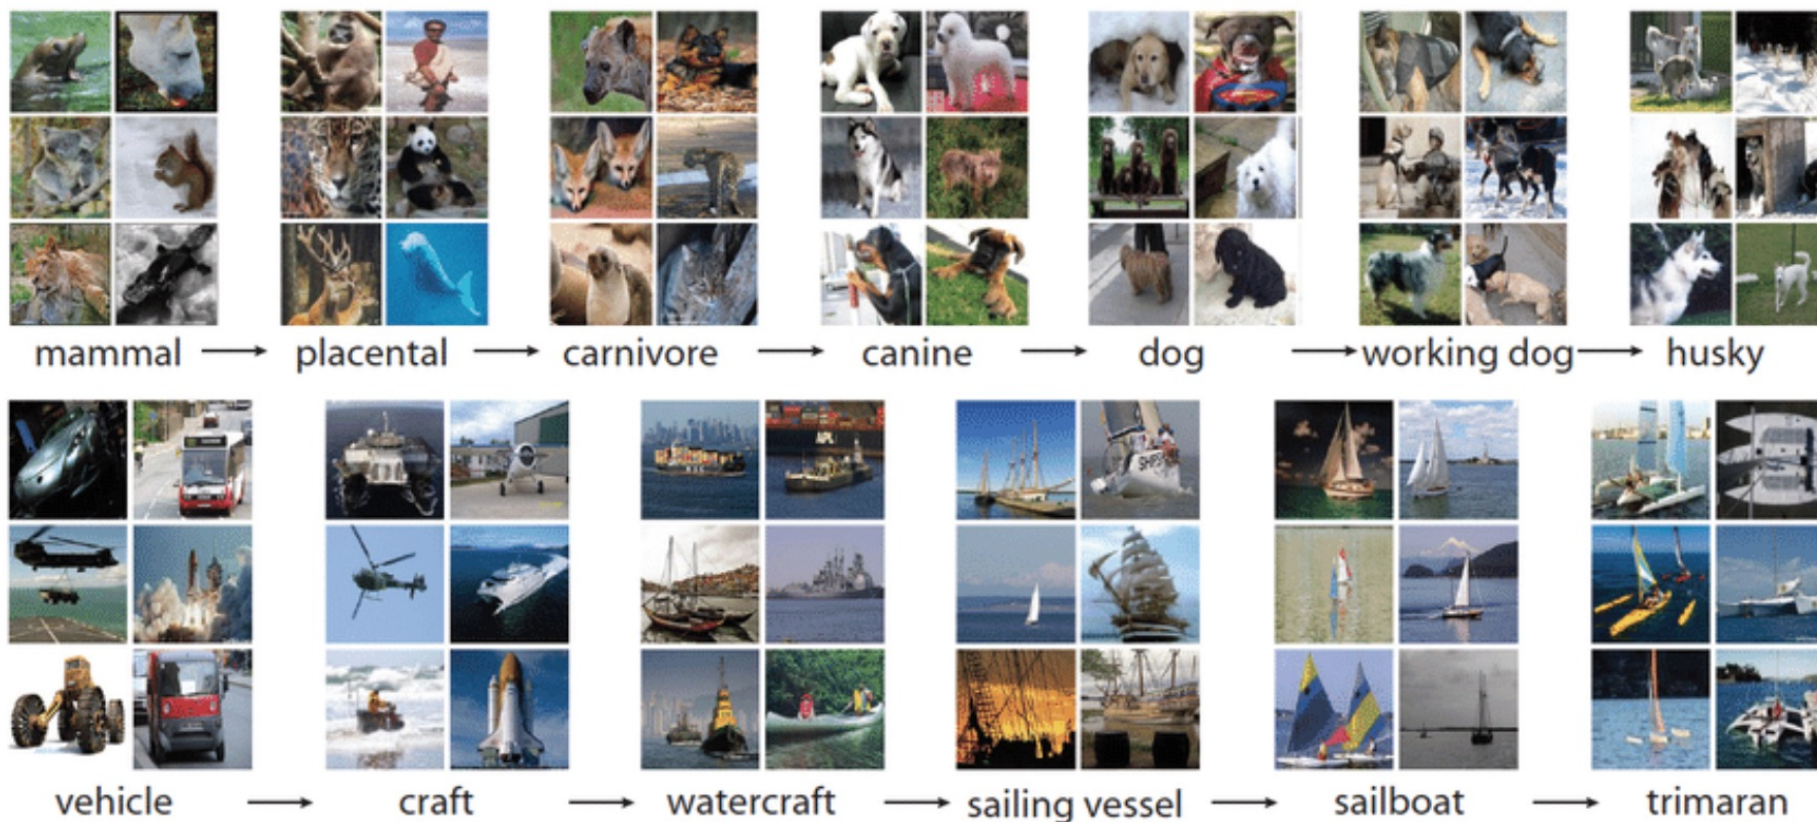


2 – Pretrained Models



ImageNet

- ❖ Training: 1,281,167 images. Validation: 50,000 images. Testing: 100,000 images
- ❖ Object classes: 1,000




2 – Pretrained Models



ImageNet

- ❖ Training: 1,281,167 images. Validation: 50,000 images. Testing: 100,000 images
- ❖ Object classes: 1,000



Large Dataset
(ImageNet,..)



MODEL
(ResNet, VGG)

Trained



Pretrained Model

2 – Pretrained Models



ImageNet

- ❖ Training: 1,281,167 images. Validation: 50,000 images. Testing: 100,000 images
- ❖ Object classes: 1,000




2 – Pretrained Models



ImageNet

- ❖ Training: 1,281,167 images. Validation: 50,000 images. Testing: 100,000 images
- ❖ Object classes: 1,000




2 – Pretrained Models




Transfer Learning

❖ Traditional Learning: Isolated, single task learning



❖ Transfer Learning: Learning of a new task relies on the previous learned tasks

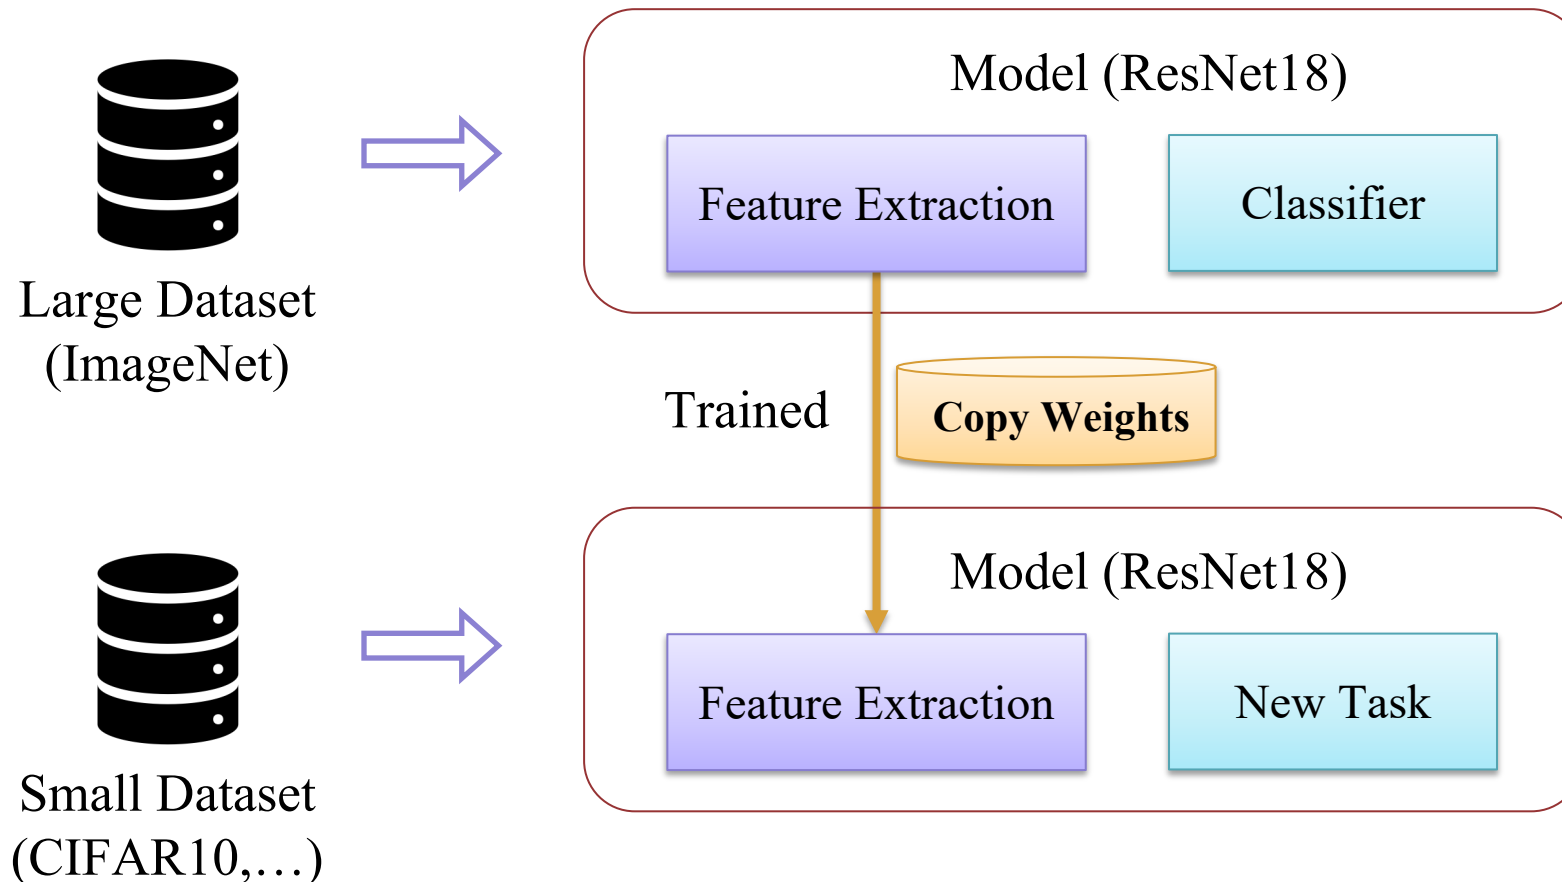


2 – Pretrained Models



Transfer Learning



❖ Transfer Learning: Feature Extractor



2 – Pretrained Models

!

Feature Extraction using Pretrained Models



2 – Pretrained Models



Feature Extraction using Pretrained Models - Demo

```
pretrained_model = models.resnet18(  
    weights=models.ResNet18_Weights.IMAGENET1K_V1  
)
```

Downloading: "[https://download.pytorch.org/models/re](https://download.pytorch.org/models/resnet18-5c136277.pth)
100%|██████████| 44.7M/44.7M [00:00<00:00, 135MB/s]

```
# remove classifier layer (block)  
resnet_feature_extractor = nn.Sequential(  
    *list(pretrained_model.children())[:-1]  
)  
  
resnet_feature_extractor.eval()
```


2 – Pretrained Models



Feature Extraction using Pretrained Models - Demo

```
plt.figure(figsize=(2, 2))  
plt.imshow(train_data.data[4])
```

<matplotlib.image.AxesImage at 0x7ddf7dbf86d0>



```
train_data.data[4].shape
```

(32, 32, 3)

```
processed_image = data_transforms(train_data.data[4])
```

```
processed_image.shape
```

torch.Size([3, 32, 32])

```
prediction = resnet_feature_extractor(processed_image.unsqueeze(0))
```

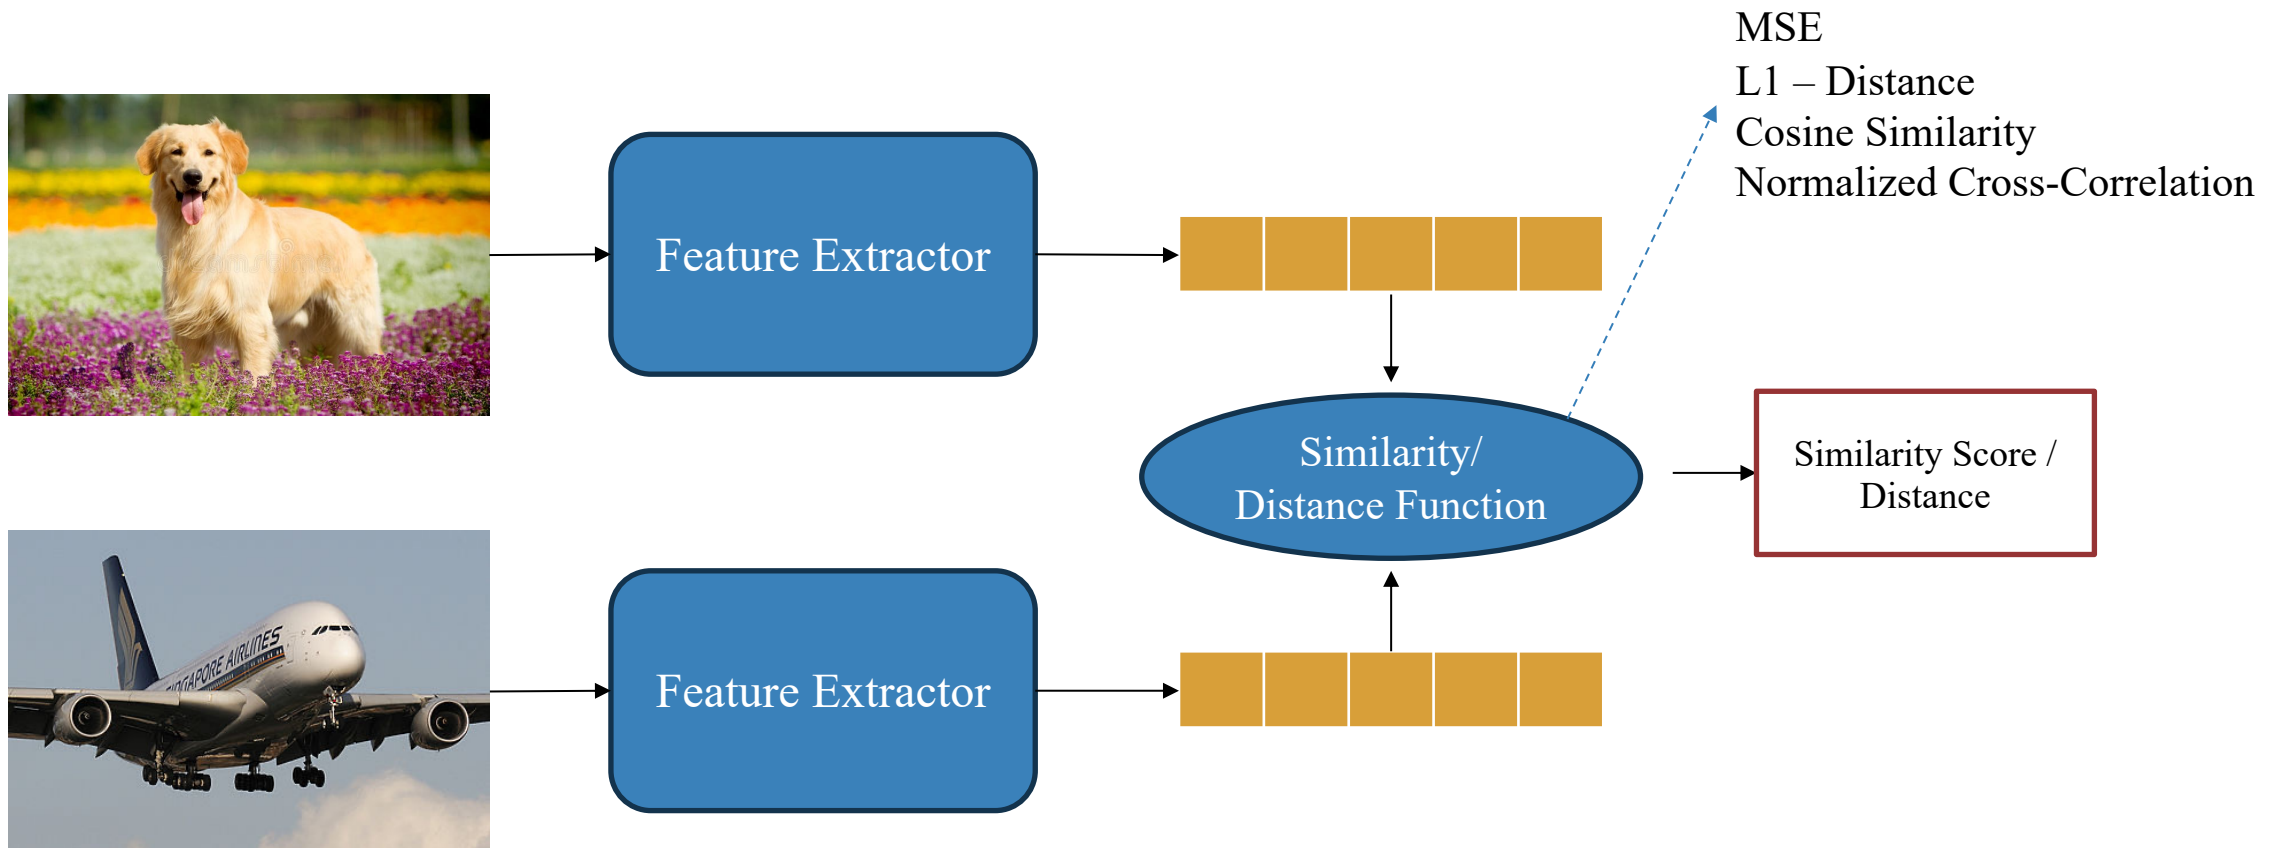
```
prediction.squeeze().shape
```

torch.Size([512])

3 – Image Similarity

1

Similarity between two images



3 – Image Similarity

1


Similarity between two images - MSE

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

$$Y = [y_1, y_2, y_3, \dots, y_n]$$

$$\hat{Y} = [\hat{y}_1, \hat{y}_2, \hat{y}_3, \dots, \hat{y}_n]$$

```
1 def mse(tensor1, tensor2):  
2     return torch.mean((tensor1 - tensor2) ** 2)
```

3 – Image Similarity

1 Similarity between two images - MSE

1.0	2.0	4.0	5.0
-----	-----	-----	-----

-

1.2	2.4	3.4	5.6
-----	-----	-----	-----

=

(

-0.2	-0.4	0.6	-0.6
------	------	-----	------

)² →

0.04	0.16	0.36	0.36
------	------	------	------

mean →

0.23

100	105	98	110
-----	-----	----	-----

-

102	107	97	108
-----	-----	----	-----

=

(

-2	-2	1	?
----	----	---	---

)² →

4	4	1	?
---	---	---	---

mean →


?

3 – Image Similarity

5 Similarity between two images – on raw images



MSE: 0.4063



MSE: 0.3265

3 – Image Similarity

3 Similarity between two images –L1 Distance

$$L1 = \sum_{i=1}^N |p_i - q_i| \quad \begin{array}{l} Y = [y_1, y_2, y_3, \dots, y_n] \\ \hat{Y} = [\hat{y}_1, \hat{y}_2, \hat{y}_3, \dots, \hat{y}_n] \end{array}$$

```
1 def l1_distance(tensor1, tensor2):  
2     return torch.sum(torch.abs(tensor1 - tensor2))
```

3 – Image Similarity

3 Similarity between two images – L1 Distance

100	105	98	110
-----	-----	----	-----

-

=

-2	-2	1	2
----	----	---	---

→

2	2	1	2
---	---	---	---

sum

→

7

102	107	97	108
-----	-----	----	-----

5	8	4	2
---	---	---	---

-

=

2	1	-6	?
---	---	----	---

→

2	1	6	?
---	---	---	---

sum

→

13

3	7	10	6
---	---	----	---


3 – Image Similarity

5

Similarity between two images – on raw images



L1: 51128.43




L1: 40339.5

3 – Image Similarity

5

Similarity between two images – on raw images



L1: 46423.3
MSE: 0.3477



L1: 40339.5
MSE: 0.3265



Problem?

3 – Image Similarity

2

Similarity between two images – Cosine Similarity

$$\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$

$$A = [a_1, a_2, a_3, \dots, a_n]$$
$$B^T = [b_1, b_2, b_3, \dots, b_n]$$

$$A \cdot B = a_1 b_1 + a_2 b_2 + \dots + a_n b_n$$

$$\|A\| = \sqrt{a_1^2 + a_2^2 + \dots + a_n^2}$$

$$\|B\| = \sqrt{b_1^2 + b_2^2 + \dots + b_n^2}$$

```
1 def cosine_similarity(tensor1, tensor2):  
2     return torch.nn.functional.cosine_similarity(tensor1, tensor2, dim=0)
```

```
1 def cosine_similarity_v2(tensor1, tensor2):  
2     tensor1_norm = torch.norm(tensor1)  
3     tensor2_norm = torch.norm(tensor2)  
4     return torch.dot(tensor1, tensor2) / (tensor1_norm * tensor2_norm)
```

3 – Image Similarity

2

Similarity between two images – Cosine Similarity

$$\begin{array}{|c|c|c|c|} \hline 1.0 & 2.0 & 3.0 & 4.0 \\ \hline \end{array} * \begin{array}{|c|} \hline 5.0 \\ \hline 6.0 \\ \hline 7.0 \\ \hline 8.0 \\ \hline \end{array} = \frac{1 * 5 + 2 * 6 + 3 * 7 + 4 * 8}{\sqrt{1^2 + 2^2 + 3^2 + 4^2} \sqrt{5^2 + 6^2 + 7^2 + 8^2}} = 0.97$$


$$\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$$

$$\begin{array}{|c|c|c|c|} \hline 3 & 2 & 4 & 6 \\ \hline \end{array} * \begin{array}{|c|} \hline 9 \\ \hline 3 \\ \hline 2 \\ \hline 6 \\ \hline \end{array} = \frac{?}{\sqrt{?} * \sqrt{?}} = ?$$


3 – Image Similarity

5

Similarity between two images – on raw images



Cosine sim: 1.0



Cosine sim: 0.83

3 – Image Similarity

4

Similarity between two images – Normalized Cross-Correlation

Formula

$$NCC = \frac{1}{IJ} \sum_{j=1}^J \sum_{i=1}^I \frac{(A_{i,j} - a)(B_{i,j} - b)}{\sigma_A \sigma_B}$$

$$a = \frac{1}{IJ} \sum_{j=1}^J \sum_{i=1}^I A_{i,j} \quad b = \frac{1}{IJ} \sum_{j=1}^J \sum_{i=1}^I B_{i,j}$$

$$\sigma_A = \sqrt{\frac{1}{IJ} \sum_{j=1}^J \sum_{i=1}^I (A_{i,j} - a)^2} \quad \sigma_B = \sqrt{\frac{1}{IJ} \sum_{j=1}^J \sum_{i=1}^I (B_{i,j} - b)^2}$$

3 – Image Similarity

4

Similarity between two images – Normalized Cross-Correlation

```
1 def normalized_cross_correlation(tensor1, tensor2):
2     mean1 = torch.mean(tensor1)
3     mean2 = torch.mean(tensor2)
4
5     # Calculate the centered tensors
6     centered1 = tensor1 - mean1
7     centered2 = tensor2 - mean2
8
9     # Calculate the NCC
10    return torch.sum(centered1 * centered2) / (torch.sqrt(torch.sum(centered1 ** 2)) * torch.sqrt(torch.sum(centered2 ** 2)))
```

3 – Image Similarity

4

Similarity between two images – Normalized Cross-Correlation

$$\begin{array}{|c|c|c|c|} \hline 1.0 & 2.0 & 4.0 & 9.0 \\ \hline \end{array} - \begin{array}{|c|} \hline 4 \\ \hline \end{array} = \left(\begin{array}{|c|c|c|c|} \hline -3 & -2 & 0 & 5 \\ \hline \end{array} \right)^2 \longrightarrow \begin{array}{|c|c|c|c|} \hline 9 & 4 & 0 & 25 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|c|} \hline 4.0 & 3.0 & 1.0 & 2.0 \\ \hline \end{array} - \begin{array}{|c|} \hline 2.5 \\ \hline \end{array} = \left(\begin{array}{|c|c|c|c|} \hline 1.5 & 0.5 & -1.5 & -0.5 \\ \hline \end{array} \right)^2 \longrightarrow \begin{array}{|c|c|c|c|} \hline 2.25 & ? & 2.25 & ? \\ \hline \end{array}$$

3 – Image Similarity

4

Similarity between two images – Normalized Cross-Correlation

-3	-2	0	5
----	----	---	---

*

=


-4.5	?	0	-2.5
------	---	---	------

1.5	0.5	-1.5	-0.5
-----	-----	------	------

$$\frac{-8}{\sqrt{38} * \sqrt{5}} = ?$$

9	4	0	25
---	---	---	----


Sum



38

2.25	0.25	2.25	0.25
------	------	------	------

Sum




5

3 – Image Similarity

5

Similarity between two images – on raw images



NCC: 0.99



NCC: 0.3937

3 – Image Similarity

5

Similarity between two images – on raw images




Method	Similarity
L1 Distance	51128.43
MSE	0.4063
Cosine similarity	0.7620
Normalized Cross Correlation	-0.3564

3 – Image Similarity

5

Similarity between two images – on raw images




Method	Similarity
L1 Distance	40339.5
MSE	0.3265
Cosine similarity	0.8294
Normalized Cross Correlation	0.3937

3 – Image Similarity

5

Similarity between two images – on raw images



Method	Similarity
L1 Distance	46423.3
MSE	0.3477
Cosine similarity	1.0000
Normalized Cross Correlation	0.9999

3 – Image Similarity

5 Similarity between two images – on raw images



Method	Similarity
L1 Distance	30073.23
MSE	0.2542
Cosine similarity	0.8691
Normalized Cross Correlation	0.3812



Problem?

3 – Image Similarity

6

Similarity between two images – using resnet18



Feature Extractor



Feature Extractor



3 – Image Similarity

6

Similarity between two images – using resnet18

Load the model →

```
1 import torchvision
2 from torchvision import models
3
4 model = torchvision.models.resnet18(weights=models.ResNet18_Weights.IMAGENET1K_V1)
```

Downloading: "<https://download.pytorch.org/models/resnet18-f37072fd.pth>" to /root/.cache/torch/hub/44.7M/44.7M [00:00<00:00, 127MB/s]

Compute feature →

```
1 tensor_dog1.shape
torch.Size([3, 224, 224])

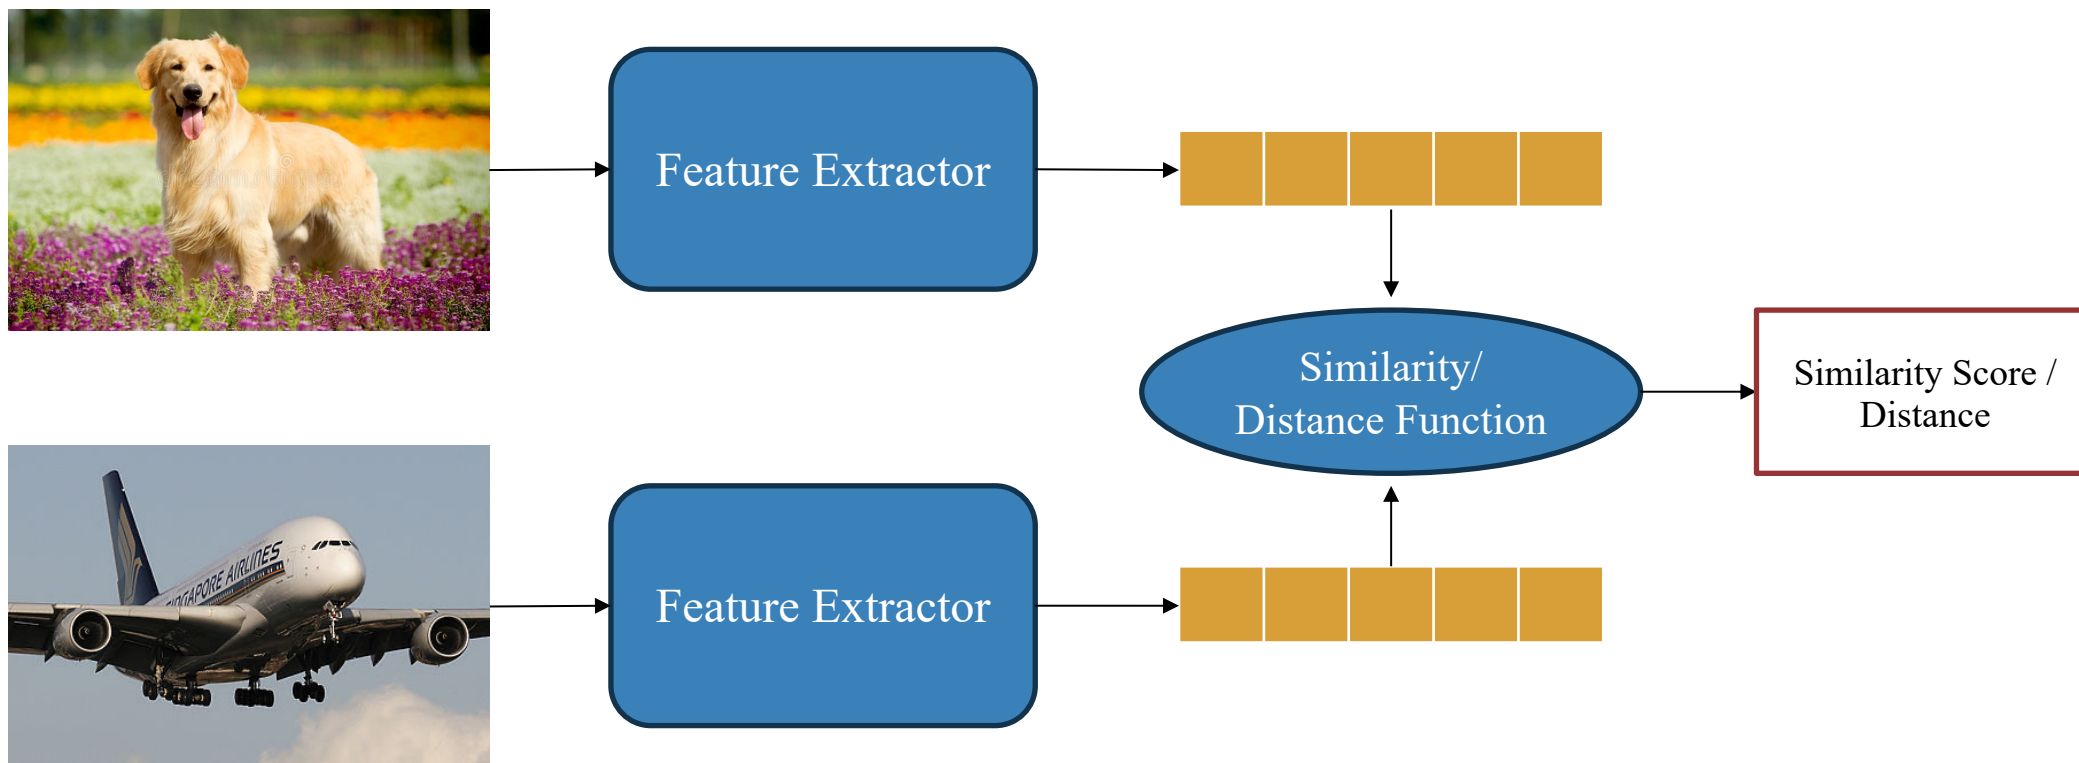
1 model.fc = nn.Identity()
```

```
1 dog1_feature = model(tensor_dog1.unsqueeze(0))
2 dog2_feature = model(tensor_dog2.unsqueeze(0))
3 dog3_feature = model(tensor_dog3.unsqueeze(0))
4 airplane_feature = model(tensor_airplane.unsqueeze(0))
```

3 – Image Similarity

6

Similarity between two images – using resnet18



3 – Image Similarity

6

Similarity between two images – using resnet18




Method	Similarity
L1 Distance	53.34
MSE	0.1329
Cosine similarity	0.8999
Normalized Cross Correlation	0.4541

3 – Image Similarity

6

Similarity between two images – using resnet18




Method	Similarity
L1 Distance	49.47
MSE	0.123
Cosine similarity	0.912
Normalized Cross Correlation	0.595

3 – Image Similarity

6

Similarity between two images – using resnet18



Method	Similarity
L1 Distance	4.344
MSE	0.011
Cosine similarity	0.9993
Normalized Cross Correlation	0.9961

3 – Image Similarity

5

Similarity between two images – on raw images



Method	Similarity
L1 Distance	22.1067
MSE	0.0549
Cosine similarity	0.9983
Normalized Cross Correlation	0.7959



AI VIET NAM

@aivietnam.edu.vn

Thanks!

Any questions?