

Introduction Reinforcement to Learning and Q-Learning

Extra Class: RL

Outline

- Introduction
- Reinforcement Learning
- Bellman Equation
- Monte Carlo and TD Learning
- Q-Learning
- Demo
- Question


Introduction

Introduction


❖ Getting Started

machine learning


unsupervised
learning



supervised
learning




reinforcement
learning




Introduction

❖ Unsupervised vs Supervised vs Reinforcement Learning


Unsupervised Learning is used to train machines using unlabeled data.



Supervised Learning uses unlabeled data to train machines.



Reinforcement Learning uses an agent and an environment to produce actions and rewards.



Introduction

❖ Algorithms

Unsupervised Learning

K-Mean Clustering

PCA

DBSCAN

Supervised Learning

Linear Regression

Logistic Regression

SVM

KNN

Reinforcement Learning

Monte Carlo

Q-Learning


Deep Q-Learning

SARSA

Introduction


❖ Training data

Unsupervised Learning



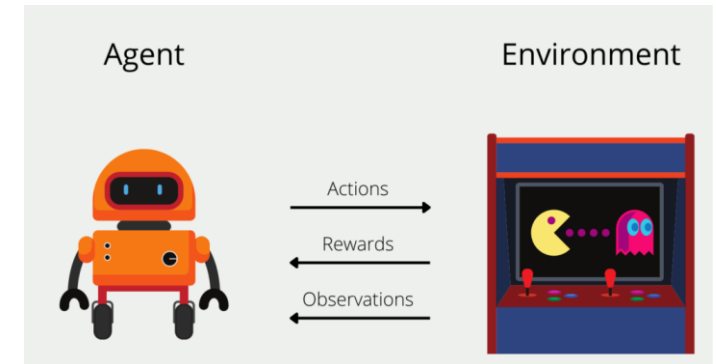
Input data is not labeled

Supervised Learning



Input data is labeled

Reinforcement Learning



No need input data

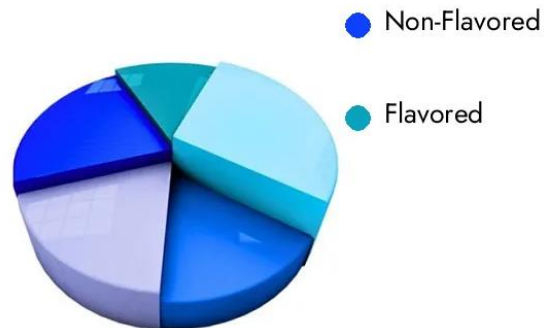
Introduction

❖ Applications

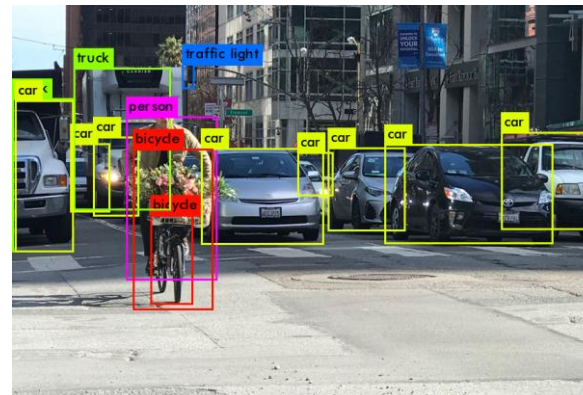

Unsupervised Learning



Soy Chunks Market Analysis By Type



Supervised Learning



Reinforcement Learning




Reinforcement Learning

Reinforcement Learning

❖ Getting Started

- Reinforcement learning is a framework for solving control tasks (also called decision problems)
- By building agents that learn from the environment by interacting with it through trial and error.
- And receiving rewards (positive or negative) as unique feedback.

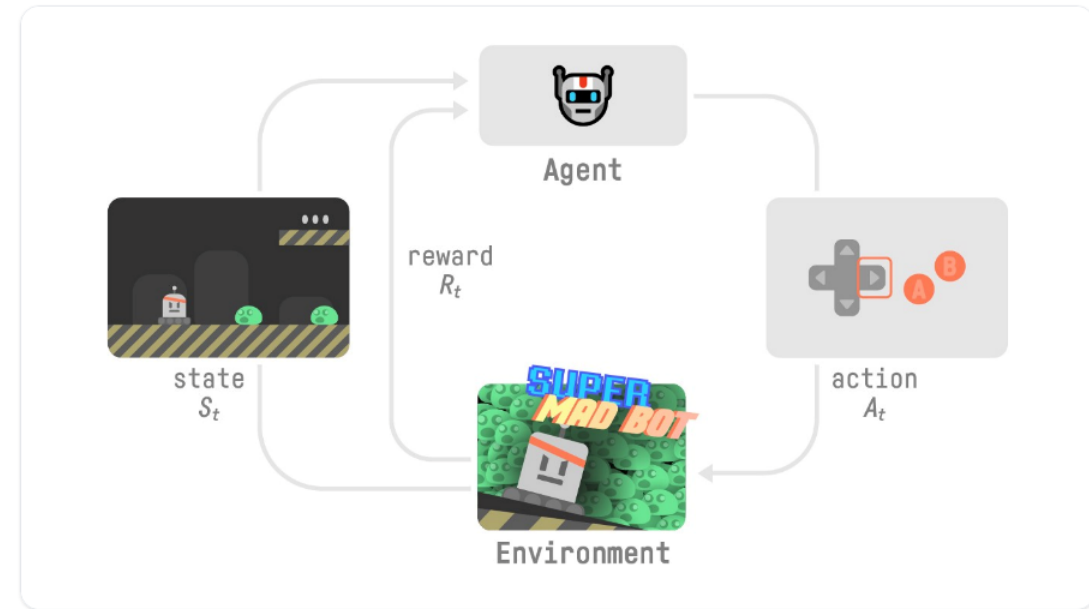


Learning from interactions with the environment comes from our natural experiences.

Reinforcement Learning

❖ Getting Started


1. Agent receives **state** S_0 from environment.
2. Base on S_0 , Agent take a **action** A_0
3. Environment goes to next **state** S_1
4. Environment give some **rewards** R_1 to the Agent



Reinforcement Learning


❖ Environment

The environment for testing RL algorithms is often games




Reinforcement Learning

❖ Observation / State Space



State s : is a complete description of the state of the world.



Observation o : is a partial description of the state.


Reinforcement Learning

❖ Reward

The **goal** is maximizing the **expected return** (expected cumulative reward).

$$R = r_t + r_{t+1} + r_{t+2} + r_{t+3} + \dots$$

$$R = \sum_{k=0}^n r_{t+k+1}$$



Reinforcement Learning

❖ Discounting

It's also reasonable to prefer rewards now to rewards later

$$R = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 r_{t+3} + \dots$$

$$R = \sum_{k=0}^n \gamma^k r_{t+k+1}$$



1

Worth Now





γ

Worth Next Step




γ^2

Worth In Two Steps

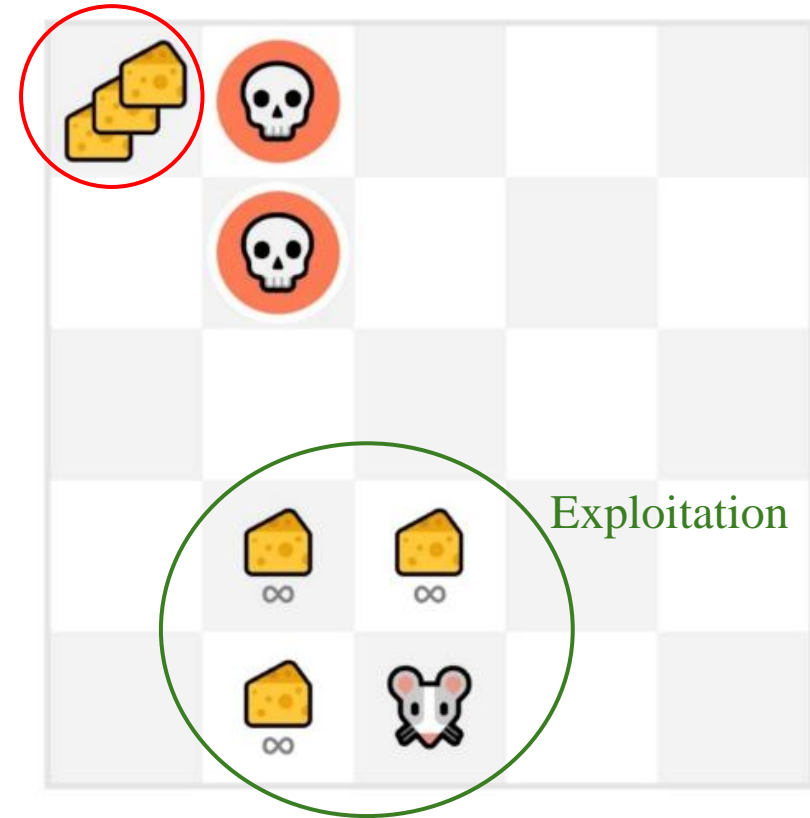


Reinforcement Learning

❖ Exploration / Exploitation trade-off




Exploration



Reinforcement Learning

❖ Approachs for solving RL problem

The **Policy** π is the **brain** of our Agent, it's the function that tells us **what action** to take given the state we are in. So it defines the **agent's behavior** at a given time.

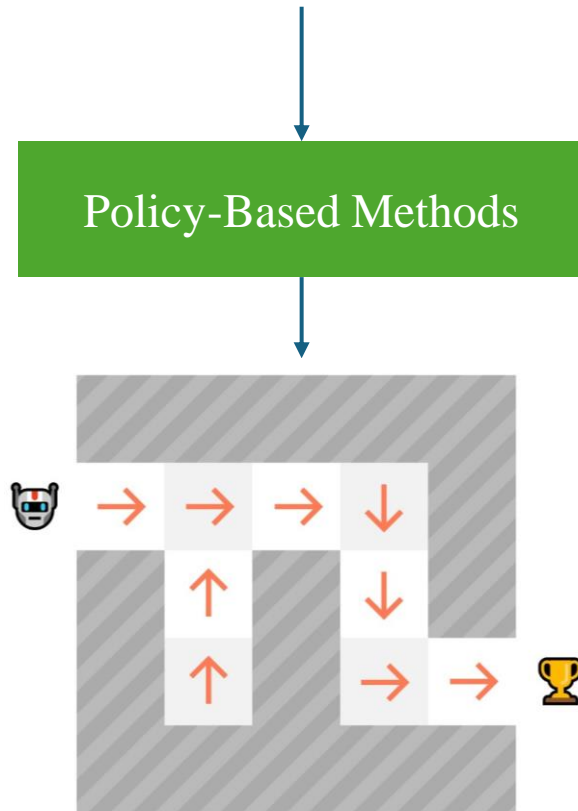


This Policy is the function we want to learn, our goal is to find the **optimal policy** π^* , the policy that **maximizes expected return** when the agent **acts according to it**. We find this π^* through training.

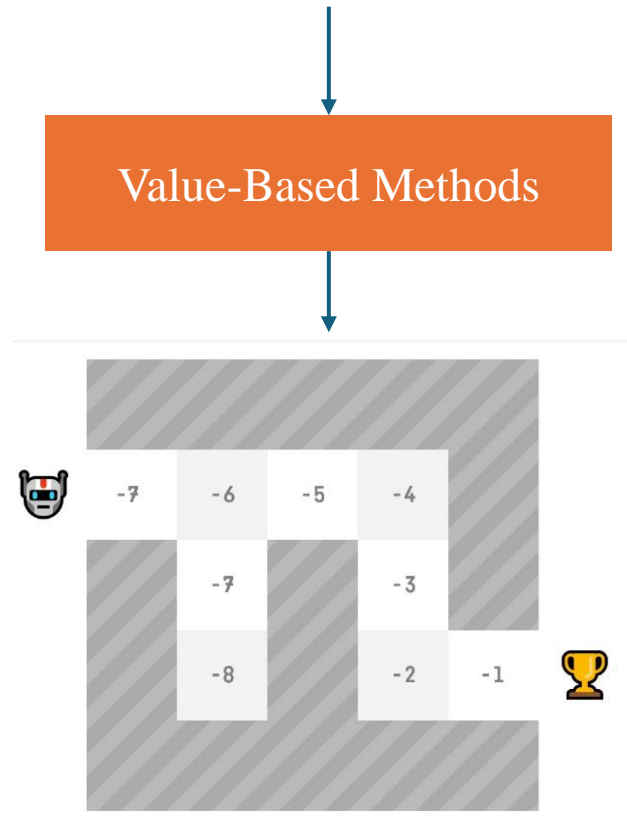
Reinforcement Learning

❖ How to find optimal policy π^*

Directly, by teaching the agent to learn which **action to take**, given the **current state**.




Indirectly, teach the agent to learn which state is **more valuable** and then take the action that **leads to the more valuable states**



Reinforcement Learning

❖ Policy-based methods

Deterministic: action $a_0 = \pi(s_0)$



State $s_0 \rightarrow \pi(s_0) \rightarrow a_0 = \text{Right}$

Stochastic: $\pi(A|s_0) = P[A|s_0]$



State $s_0 \rightarrow \pi(A|s_0) \rightarrow [\text{Left: } 0.1, \text{Right: } 0.7, \text{Jump: } 0.2]$

Policy π is a Neural Network

Reinforcement Learning

❖ Value-based methods

We learn a **value function** that **maps** a **state** to the **expected value** of being at that state.

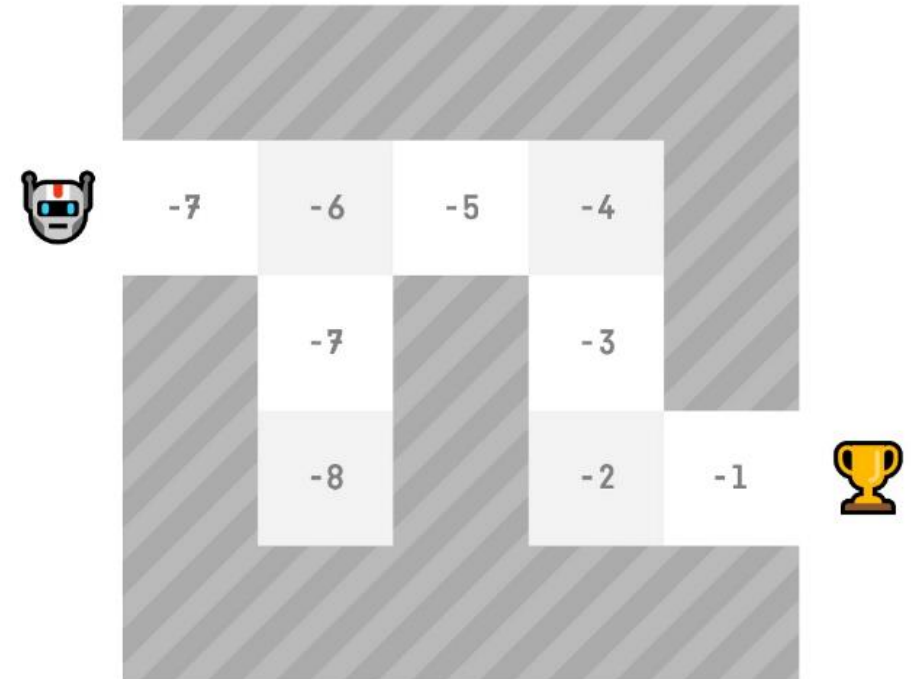
$$\underline{v_\pi(s)} = \underline{\mathbb{E}_\pi[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \mid S_t = s]} \quad \underline{\hspace{10em}}$$

Value
function

Expected discounted return

Starting
at state s

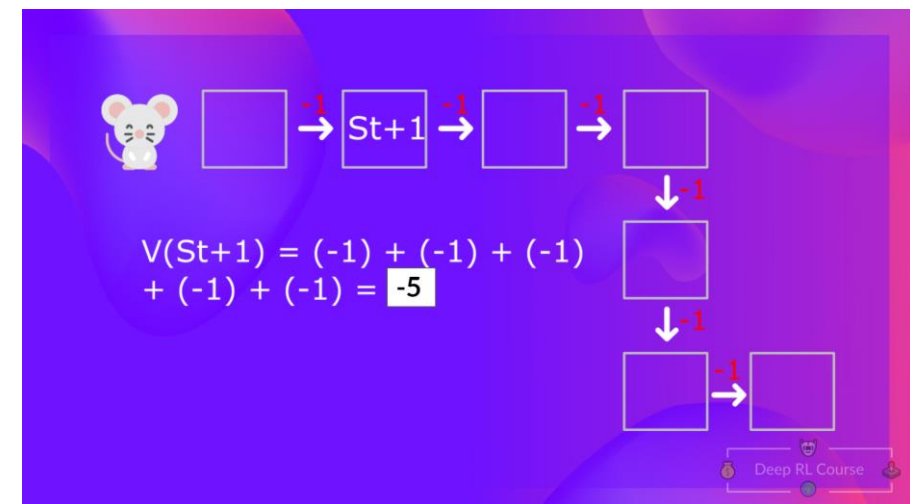
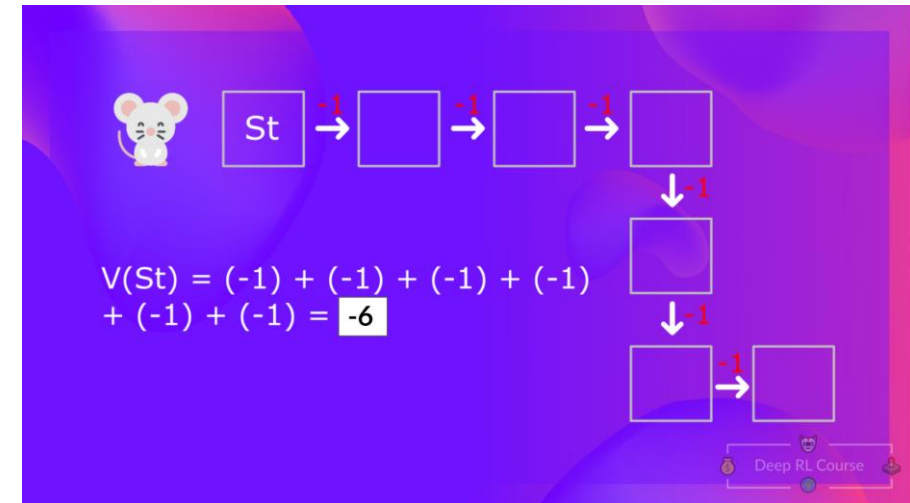

$V_\pi(s)$: value of a state is the **expected discounted return** the agent can get if it **starts in that state**, and then **acts according to our policy**.



Bellman Equation

Bellman Equation

❖ Getting Started



Bellman Equation

❖ Getting Started

The Bellman Equation

$$V_{\pi}(s) = \mathbf{E}_{\pi}[R_{t+1} + \gamma * V_{\pi}(S_{t+1}) | S_t = s]$$


Value of
state s

Expected value of
immediate reward

+ the discounted value of
next_state

If the agent
starts at state s

And uses the policy to
choose its actions for
all time steps



Monte Carlo

Monte Carlo

❖ Getting Started

- **Monte Carlo** is a strategy to train our **value function**.
- It uses **experience** to solve the RL problem.

Monte Carlo Approach:

Monte Carlo: waits until the end of the episode, then calculates G_t (return) and uses it as a target for its value or policy.

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \alpha [G_t - \underline{V(S_t)}]$$

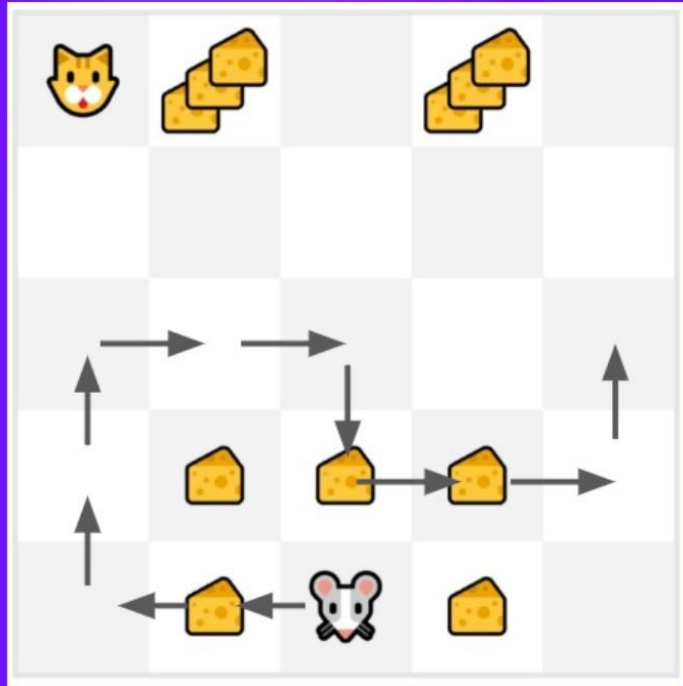
New value of state t	Former estimation of value of state t (= Expected return starting at that state)	Learning Rate	Return at timestep t	Former estimation of value of state t (= Expected return starting at that state)
----------------------	---	---------------	----------------------	---

- Monte Carlo uses an **entire episode** of experience **before learning**.
- So it requires a **complete episode** of interaction before updating our value function.

Monte Carlo

❖ Training

- Initialize value function to zero for each state
- Learning rate (lr) is 0.1 and our discount rate is 1 (= no discount)
- The mouse explores the environment and takes random actions



- Calculate the return G_t .

$$G_t = R_{t+1} + R_{t+2} + R_{t+3} \dots$$

$$G_t = 1 + 0 + 0 + 0 + 0 + 0 + 1 + 1 + 0 + 0$$

$$G_t = 3$$

- We can now update $V(S_0)$.

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

$$\text{New } V(S_0) = V(S_0) + lr * [G_t - V(S_0)]$$

$$\text{New } V(S_0) = 0 + 0.1 * [3 - 0]$$

$$\text{New } V(S_0) = 0.3$$

Temporal Difference

Temporal Difference

❖ Getting Started

- The idea with TD is to update at each step.
- Estimate G_t by R_{t+1} and discounted value of next state.

TD Learning Approach:

Temporal Difference Learning: learning at each time step.

$$\underline{V(S_t)} \leftarrow \underline{V(S_t)} + \underline{\alpha} [\underline{R_{t+1}} + \underline{\gamma V(S_{t+1})} - \underline{V(S_t)}]$$

New value
of state t

Former
estimation of
value of state
 t

Learning
Rate

Reward


Discounted value of next
state

TD Target

Temporal Difference

❖ Training

- Initialize value function to zero for each state
- Learning rate (lr) is 0.1 and our discount rate is 1 (= no discount)
- The mouse explores the environment and takes random actions (**going to the left**)



At the end of one step (State, Action, Reward, Next State):

- We have R_{t+1} and S_{t+1}
 - We update $V(S_t)$:
 - We estimate G_t by adding R_{t+1} and the discounted value of next state.
- TD target : $R_{t+1} + \gamma V(S_{t+1})$

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Now we continue to interact with this environment with our updated value function. By running more and more steps, the agent will learn to play better and better.

Monte Carlo vs TD Learning

❖ Summary

Monte Carlo: $V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$

TD Learning: $V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$

Q-Learning

Q-Learning

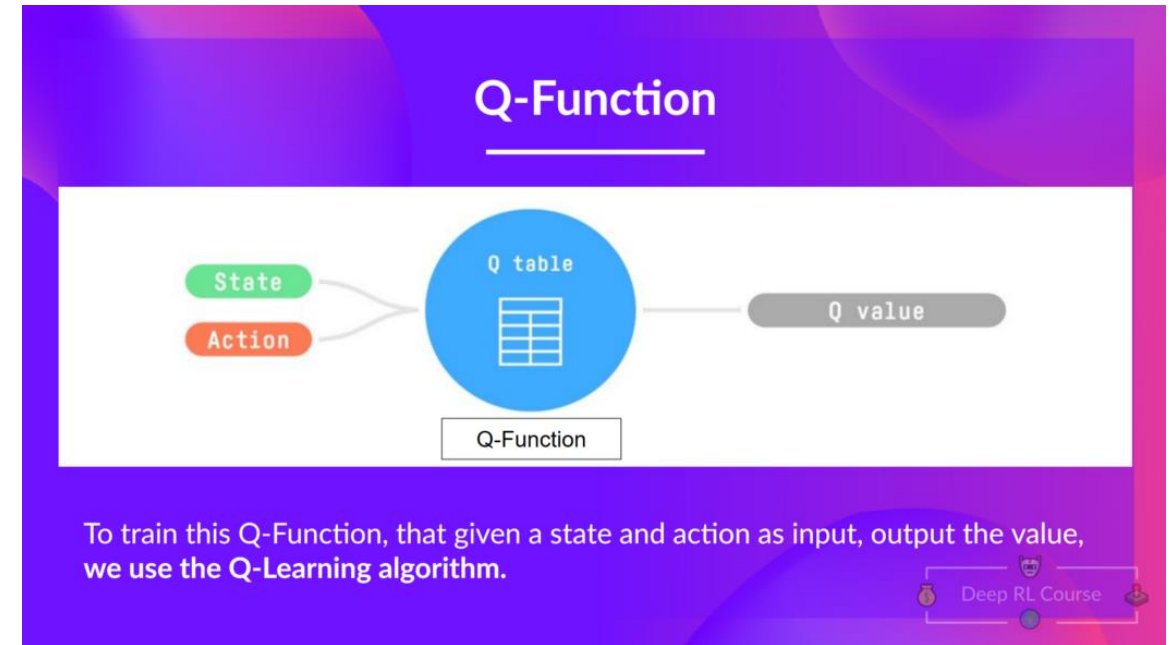
❖ Getting Started

Q-learning is a simple Reinforcement Learning algorithm

Q-learning is a value-based methods.

Q-learning uses a TD approach to train its action-value function.

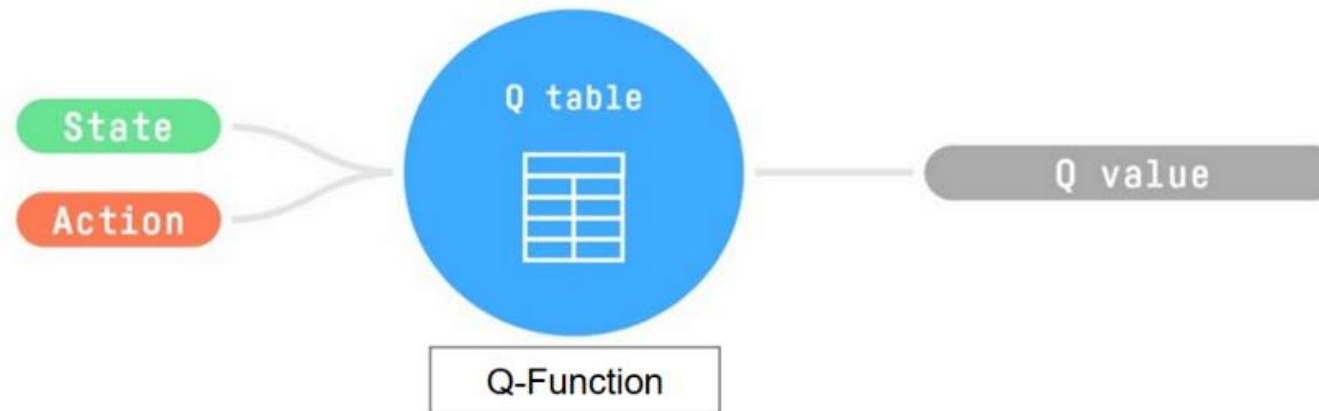
Q-Learning is the algorithm we use to train our **Q-function**, an **action-value function** that determines the **value** of being at a particular **state** and **taking a specific action** at that state.



Q-Learning


❖ Getting Started

- The Q-function uses a Q-table that has the value of each state-action pair.
- Given a state and action, our Q-function will search inside its Q-table to output the value.
- $Q(s, a)$ is the the value by starting from s and take action a .













Q-Learning

❖ Getting Started



States

Actions

				
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

Q-Learning

❖ Training

- Initialize $Q(s, a) = 0$ for each s, a pair
- Select action and observe an experience (s, a, r, s') .
- Update $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

Q-Learning, Step 1

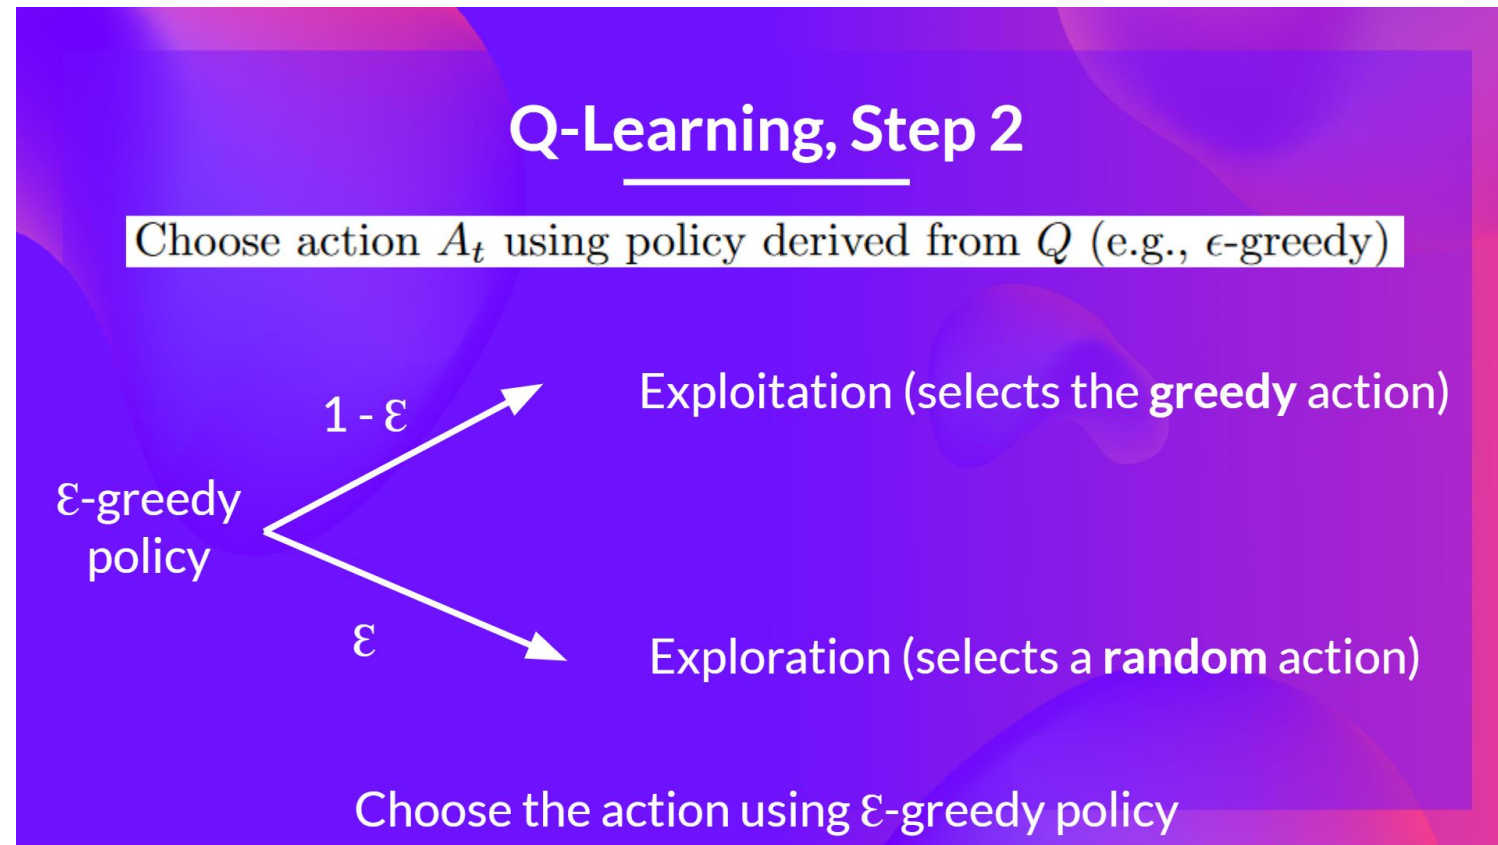
Initialize Q arbitrarily (e.g., $Q(s, a) = 0$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$, and $Q(\text{terminal-state}, \cdot) = 0$)

				
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

We initialize the Q-Table

Q-Learning

❖ Training step 2



Q-Learning

❖ Training step 3

Q-Learning, Step 3

Take action A_t and observe R_{t+1}, S_{t+1}

Q-Learning

❖ Training step 4

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

New
Q-value
estimation

Former
Q-value
estimation

Learning
Rate

Immediate
Reward

Discounted Estimate
optimal Q-value
of next state

Former
Q-value
estimation

TD Target

TD Error

Q-Learning

❖ Rules

- You're a mouse in this tiny maze. You always start at the same starting point.
 - The goal is to eat the big pile of cheese at the bottom right-hand corner and avoid the poison. After all, who doesn't like cheese?
 - The episode ends if we eat the poison, eat the big pile of cheese, or if we take more than five steps.
 - The learning rate is 0.1
 - The discount rate (gamma) is 0.99
-
- +0: Going to a state with no cheese in it.
 - +1: Going to a state with a small cheese in it.
 - +10: Going to the state with the big pile of cheese.
 - -10: Going to the state with the poison and thus dying.
 - +0 If we take more than five steps.











Q-Learning

❖ Training

$$lr = 0.1$$

$$\gamma = 0.99$$



				
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$











Q-Learning

❖ Training

$$\alpha = 0.1$$

$$\gamma = 0.99$$



				
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$


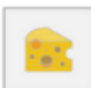




Q-Learning

❖ Training

$$\alpha = 0.1$$

$$\gamma = 0.99$$



	←	→	↑	↓
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0
	0	0	0	0

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$$

Q-Learning

❖ Training

At the end of the training, we'll get an estimate of the optimal Q-function.

	←	→	↑	↓
🐹	0	10.8	0	0
🏠	0	9.9	0	-10
☐	0	0	0	10
■	0	-10	0	0
🍌	0	0	0	0
🧀	0	0	0	0

The link between Value and Policy:

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

Finding an optimal value function leads to having an optimal policy.

Demo

Question

