

# Point Cloud Techniques and Applications

---

Tuan Dang  
Instructor  
Email: [tuan.dang@uta.edu](mailto:tuan.dang@uta.edu)  
Homepage: [www.tuandang.info](http://www.tuandang.info)  
University of Texas at Arlington, USA

Phuc Pham  
Teaching Assistant  
Email: [phuc.phamhuythien@hcmut.edu.vn](mailto:phuc.phamhuythien@hcmut.edu.vn)  
Ho Chi Minh City University of Technology

# Agenda

- Section 1 (24/4): Introduction to Point Cloud and basic techniques to process
- Section 2 (26/4): Machine Learning with Point Cloud

# Requirement

- Linear Algebra
- Basic (numpy)
- Advanced (pytorch)

# Hints to effectively learn this course

- Understanding is the key
- Concepts are connected in a chain: each concept is a link in a chain
- If you don't understand some concepts at certain links, please stop or go back to review them
- Encourage ask questions even sometimes interrupt the instructor

# **Section 1**

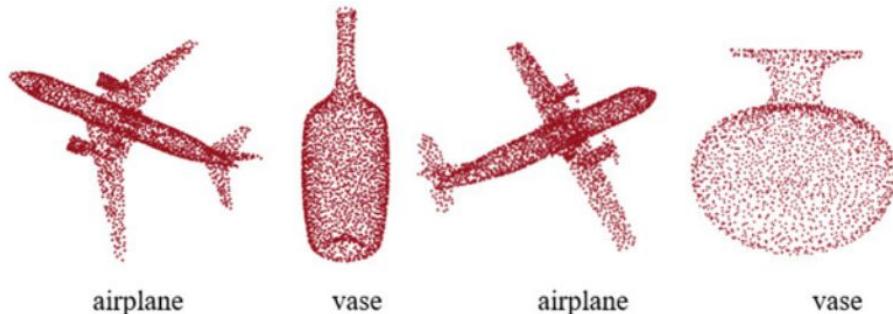
## **Introduction to Point Cloud and basic techniques to process**

# Content

- Concepts of Point Cloud
- Data Acquisition, Sensors
- Camera Pinhole Model
- Noise Removal
- Sampling Techniques
- Surface Normal
- Surface Curvature
- Search on Point Cloud
- Data Structure for Efficient Search
- Registration
- Point Cloud Example Codes

# What is a Point Cloud?

- A list of point 3D points:
  - $[p_1, p_2, p_3, \dots]$
- Each point can be represented
  - $(X,Y,Z[,feature1,feature2\dots])$
  - $X,Y,Z$
  - $X,Y,Z,I$  (intensity)
  - $X,Y,Z,R,G,B$
  - ...
- The number of points may vary from one point cloud to another



# Point Cloud vs Image

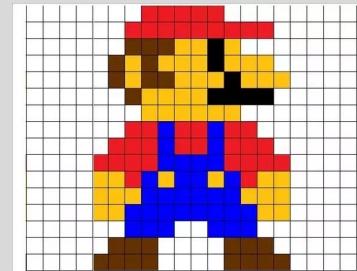
Memory Space

Unordered- irregular structure



$[p_1, p_2, \dots] = [p_2, p_1, \dots]$

Well-defined grid structure



WxHxC

Storage Space

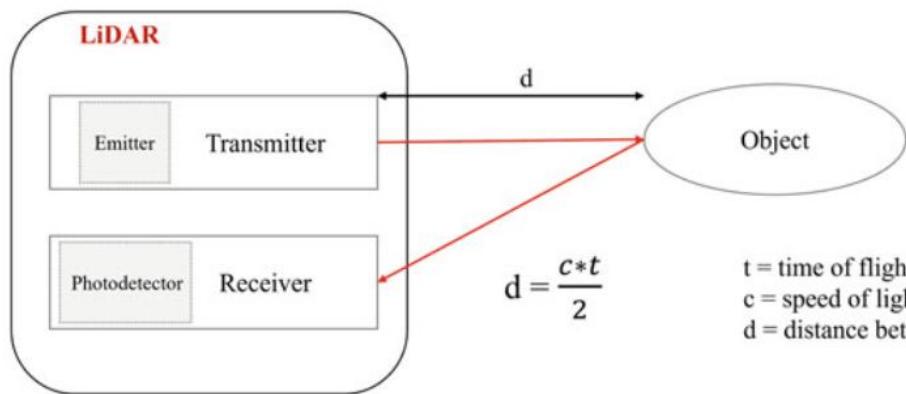
txt, ply, bin

JPEG, PNG, BITMAP

# How to obtain pointcloud?

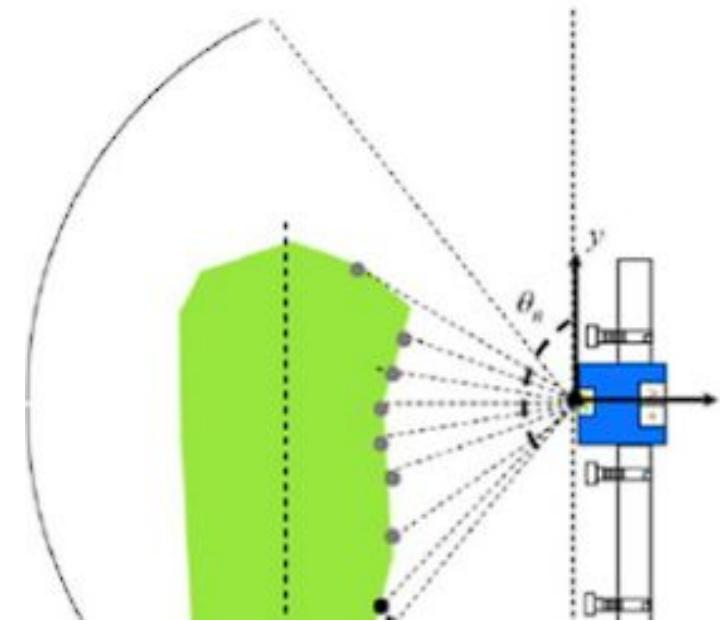
- LiDAR
- Depth images

# LiDAR (outdoor)



Sing Beam

$t$  = time of flight  
 $c$  = speed of light in air  
 $d$  = distance between LiDAR and object



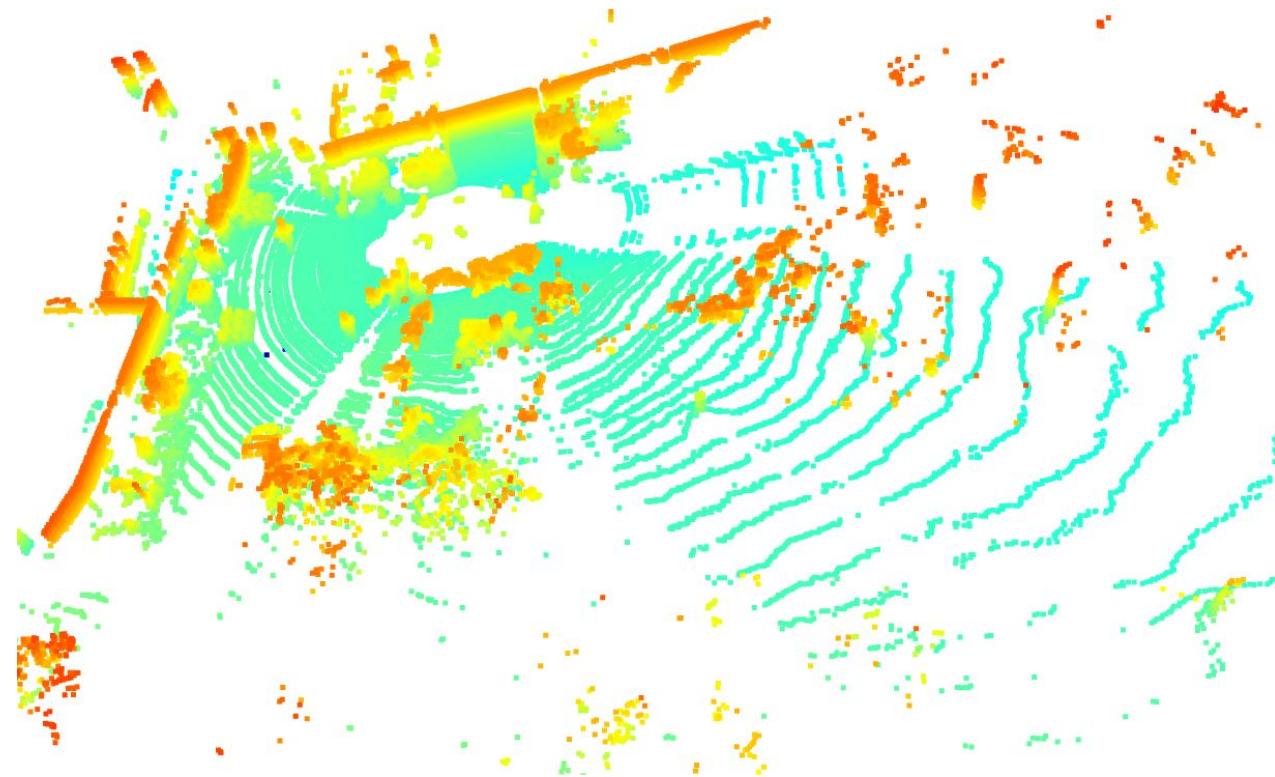
For each point: (X,Y,Z, I)

Multiple beams

# Industry LiDAR

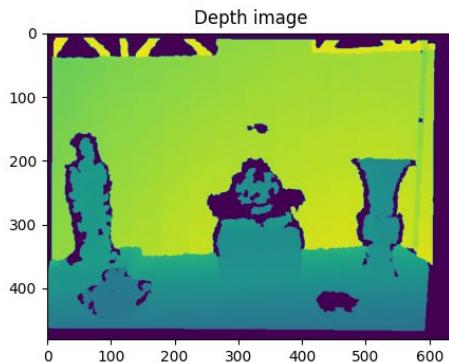
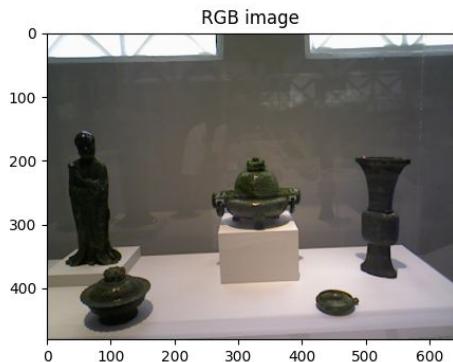


# LiDAR Data



Generated from: [https://github.com/tuantdang/pointcloud\\_lessons/blob/main/session1/load\\_point\\_clouds.py](https://github.com/tuantdang/pointcloud_lessons/blob/main/session1/load_point_clouds.py)

# RGB-D camera (most popular, indoor)



# Commercial RGB-D cameras



Microsoft  
Kinect



Intel  
Realsense D435



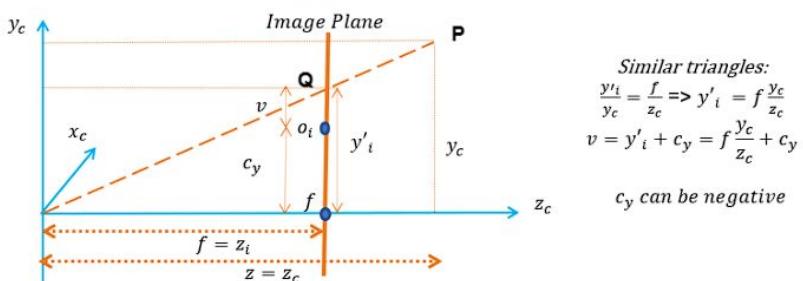
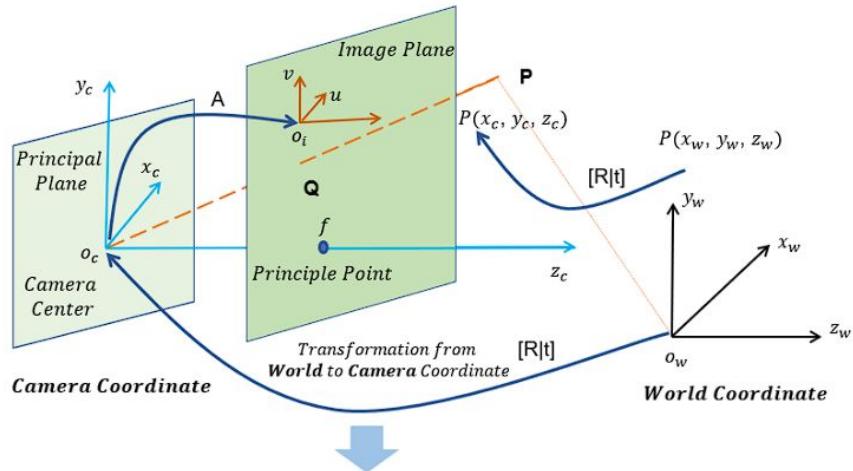
Intel  
Realsense L515

# Depth Measurement

- Stereo Vision
- Structure Light
- LiDAR with MEMS

## From world to camera

# Camera model (Pinhole)



$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \frac{1}{z} \begin{bmatrix} f_x & 0 & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_c \\ y_c \\ 1 \end{bmatrix}$$

## From Camera to World

$$\begin{bmatrix} x_c \\ y_c \\ z_c \end{bmatrix} = R \begin{bmatrix} x_w \\ y_w \\ z_w \end{bmatrix} + t = \begin{bmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{bmatrix} \begin{bmatrix} x_w \\ y_w \\ z_w \end{bmatrix} + \begin{bmatrix} t_x \\ t_y \\ t_z \end{bmatrix}$$

or can be written as:

$$\begin{bmatrix} x_c \\ y_c \\ z_c \\ 1 \end{bmatrix} = \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_w \\ y_w \\ z_w \\ 1 \end{bmatrix}$$

# Point cloud processing

- Filter noises
  - Downsampling
  - Noise Removal
- Search
- Registration

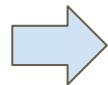
# Voxel Grid Sampling

- Build voxel grid
- Select point strategies
  - Random
  - Center point
  - Centroid point

# Build Voxel Grid Coordinate

$x_{\min} = \min(x_1, x_2, \dots, x_n),$   
 $x_{\max} = \max(x_1, x_2, \dots, x_n),$   
 $y_{\min} = \min(y_1, y_2, \dots, y_n),$   
 $\vdots$   
 $z_{\max} = \max(z_1, z_2, \dots, z_n).$

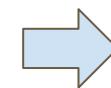
Point Cloud Range



$N_x = (x_{\max} - x_{\min})/r,$   
 $N_y = (y_{\max} - y_{\min})/r,$   
 $N_z = (z_{\max} - z_{\min})/r,$

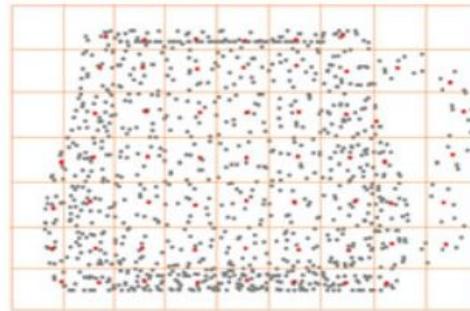
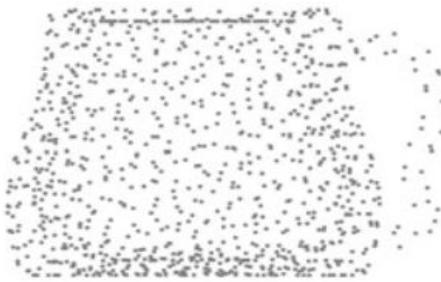
Grid Dimensions

r : voxel size



$i_x = \lfloor (x - x_{\min})/r \rfloor,$   
 $i_y = \lfloor (y - y_{\min})/r \rfloor,$   
 $i_z = \lfloor (z - z_{\min})/r \rfloor,$   
 $i = i_x + i_y * N_x + i_z * N_x * N_z.$

Voxel indexing  
0,1,..., NxNyNz-1



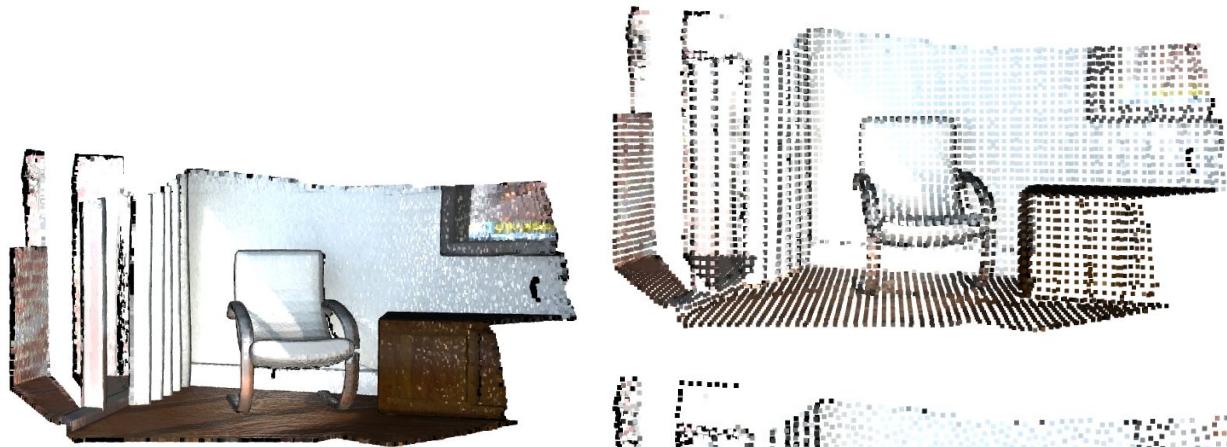
### Center Point Selection



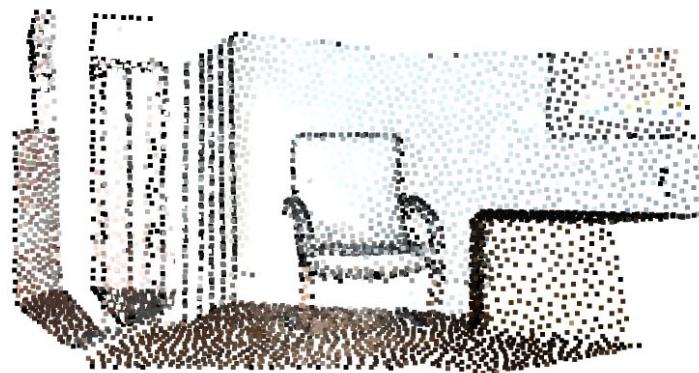
# Furthest Point Sampling (FPS)

- This heuristic algorithm
- Algorithm
  - Step 1: Random select a point
  - Step 2: Calculate distances between current selected point to all remaining points
  - Step 3: Select a point which is the furthest to the previous point
  - Step 4: Loop step 2 & 3 until m points are chosen
- Advantages
  - Evenly across original point cloud
  - Capture Esen





Voxel  
Sampling



FPS

# Surface Normal

- A vector that is **perpendicular** to the plane
- A plane can be determined by
  - Perpendicular vector  $(a,b,c)$
  - A point in the plane  $(x,y,z)$
- Plane equation from perpendicular vector and inner-point:  $\mathbf{ax+by+cz+d = 0}$

$$\mathbf{n} = [n_x, n_y, n_z]^T = \frac{[a, b, c]^T}{\|[a, b, c]^T\|}.$$

# Normal vector in Practise

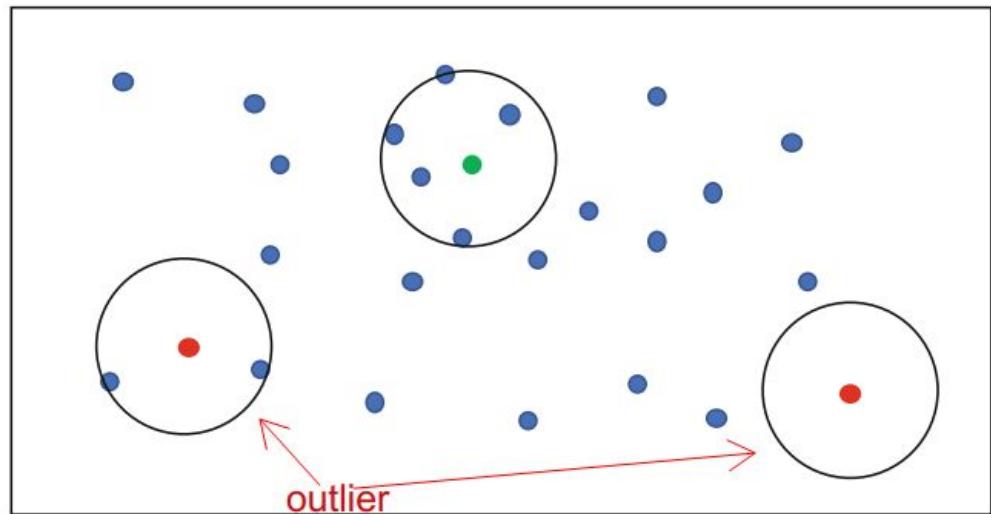
- Not points are perfect in a plane => eigenvector
- Collect a set of point
- Calculate eigenvectors and eigenvalues ( $\lambda_1, \lambda_2, \lambda_3$ )
- The least driven-value is normal vector:  $\mathbf{n} = \lambda_3$

# Surface variance = Curvature

$$c = \frac{\lambda_3}{\lambda_1 + \lambda_2 + \lambda_3}$$

# Noise Removal: Radius Outlier Removal

- Search around query point  $\mathbf{p}$  with radius  $r$ , call  $\text{ball}(\mathbf{p}, r)$
- If number of points in the ball  $< k_{\min}$ ,  $\mathbf{p}$  is outlier.



Radius outlier removal:  $k_{\min} = 4, r = 1$

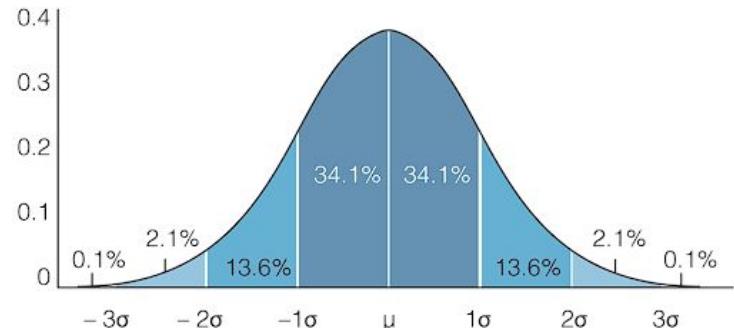
# Statistical Outlier Removal

Model inlier as Gaussian  $\mathbf{N}(\mu, \sigma)$  distribution

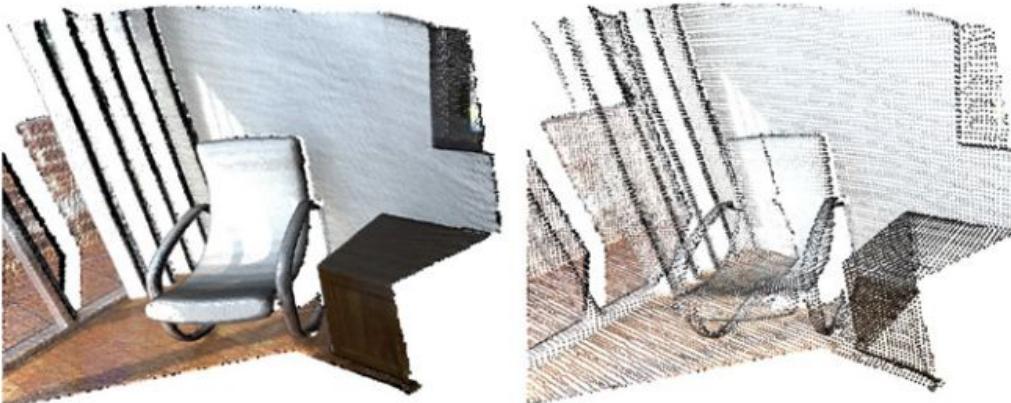
- $n$  : number of points,  $k$  neighbors, **c : variance steps**
- $d_{ij}$  : distance from point  $i$  to point  $j$

$$\mu = \frac{1}{nk} \sum_{i=1}^n \sum_{j=1}^k d_{ij}, \quad \sigma = \sqrt{\frac{1}{nk} \sum_{i=1}^n \sum_{j=1}^k (d_{ij} - \mu)^2}.$$

$$\frac{1}{k} \sum_{j=1}^k d_{ij} > \mu + c\sigma \quad \text{or} \quad \frac{1}{k} \sum_{j=1}^k d_{ij} < \mu - c\sigma, \quad c \in \mathbb{R}^+.$$

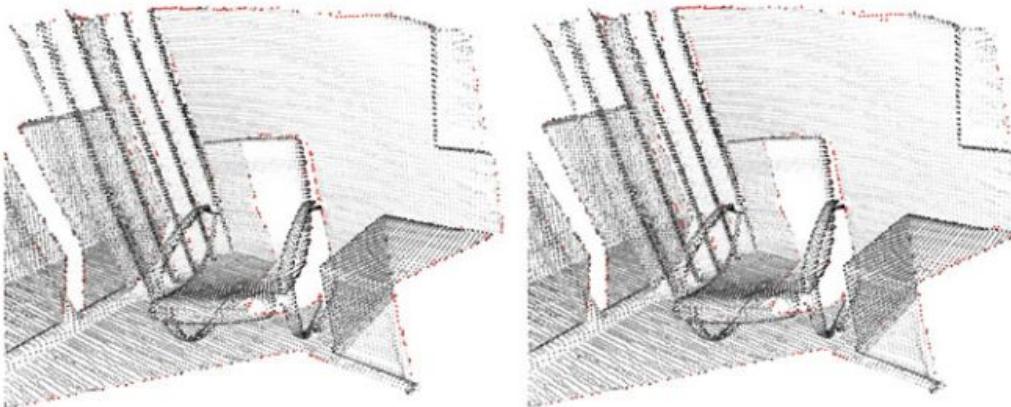


## Examples



Original (a)

(b) Voxel size = 0.02



(c)

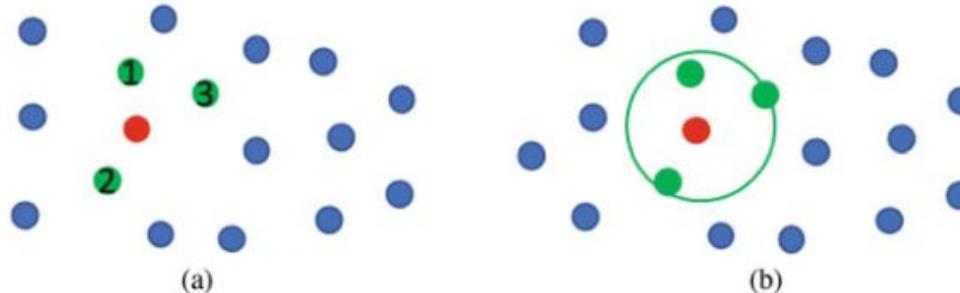
(d)

$k_{min}=16, r=0.05$

Statistical Outlier removal  $k=20, c=2$

# Search Nearest Neighbors (-NN)

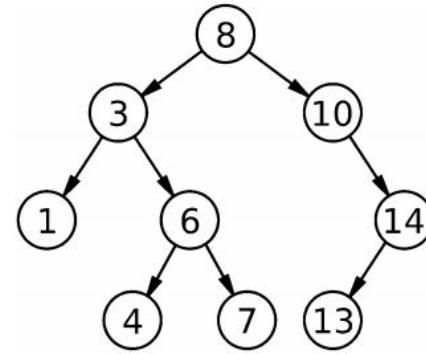
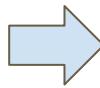
- Why searching neighbors?
- K-NN :  $O(N \log N)$  or  $O(N \log K)$  since using sort algorithm
- Radius-NN :  $O(N)$  compute and compare



Comparison of  $K$ -NN and Radius-NN. (a)  $K$ -NN. (b) Radius-NN

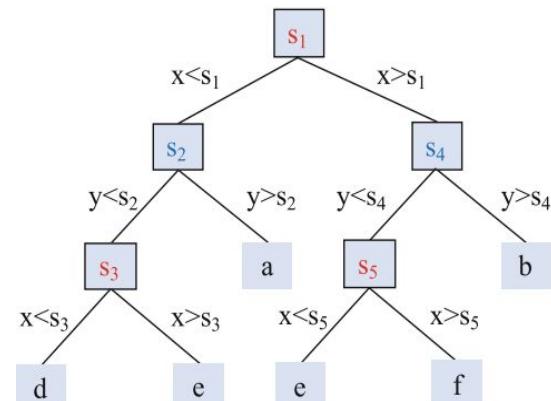
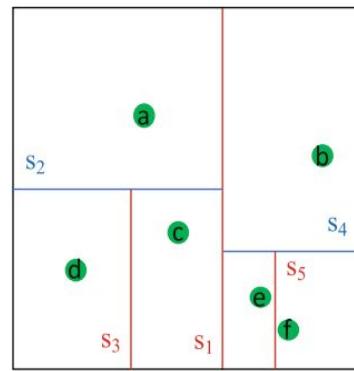
# 1D-BST = 1D Binary Search Tree

8,3,10,1,6,14,4,7,14



# k-dimensional tree = k-d tree

- K-d tree is a BST where each node is k-dimensional point
- Example:
  - Input 2D points: a,b,c,d,e,f
  - Draw a line: perpendicular x-axis to separate two subsets
  - For each subset, draw lines:
    - perpendicular y-axis
    - Separate two subsets
  - Repeat until no more data
- Note: Using round-robin
  - 2D: perpendicular lines:  $x>y>x>y>\dots$
  - 3D: perpendicular planes:  $x>y>z>x>y>z>\dots$

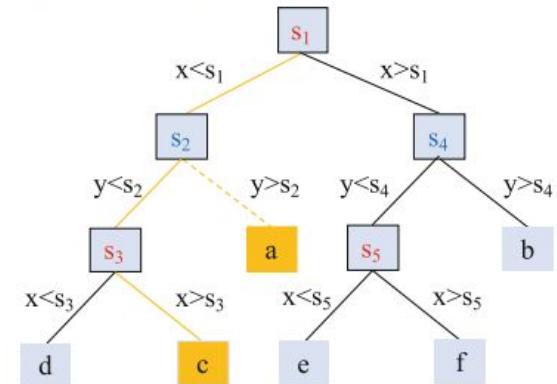
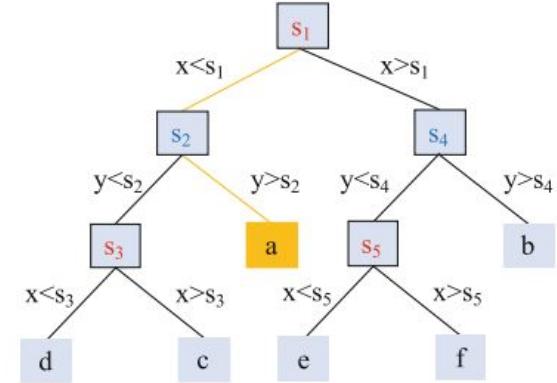
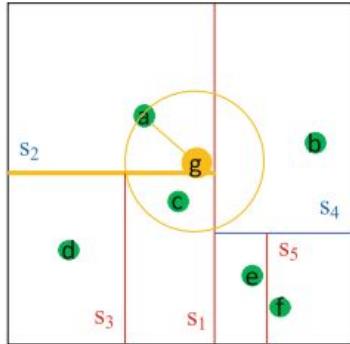


Construct 2-d tree from 2D data

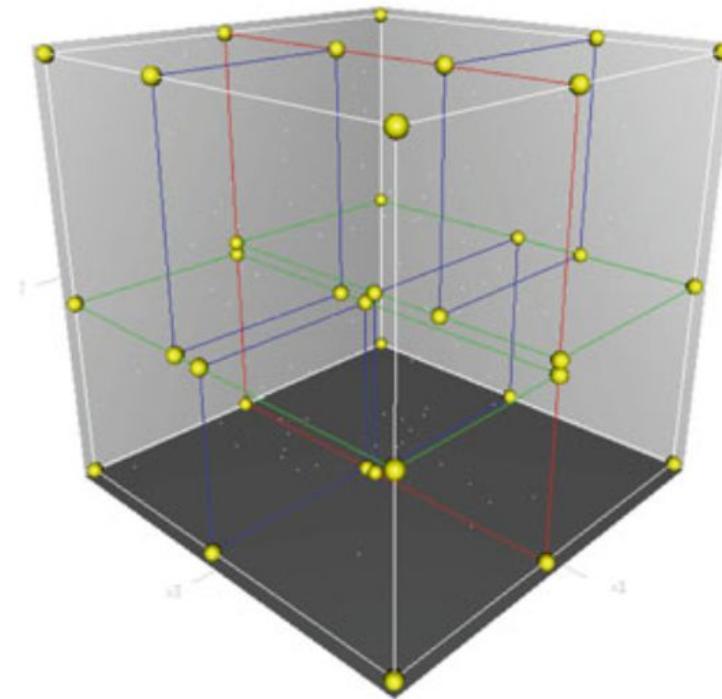
# K-NN search on k-d tree

Example with K=2 and query **g**, with worst distance d = infinity

- Search from root s1
  - $gx < s_1$  and  $gy > s_2$  where  $||g-a|| < d$  store **a**.
  - Update  $d = ||g - a||$
  - Since  $|gx-s_1| < d$  and  $|gy-s_2| < d$ , **c** is stored
  - Update  $d = ||g - c||$
  - ....
- Complexity  $O(\log N)$ : best case and  $O(N)$  worst case

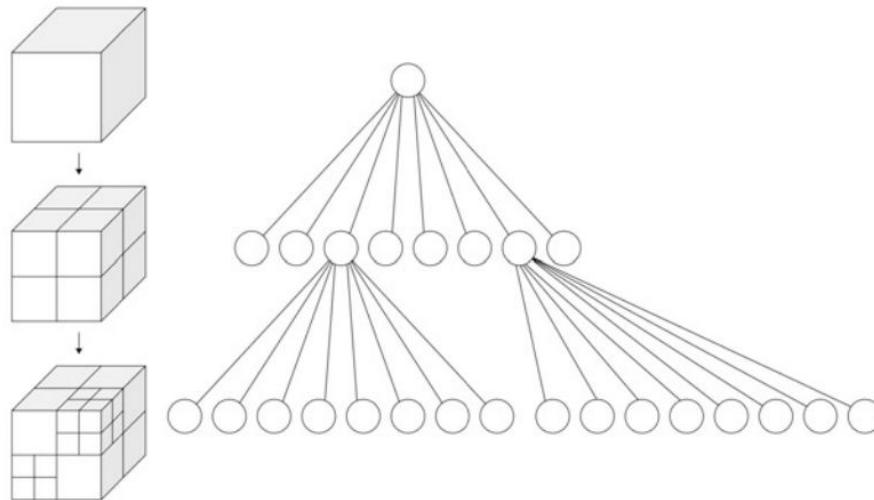


# k-d tree with k=3



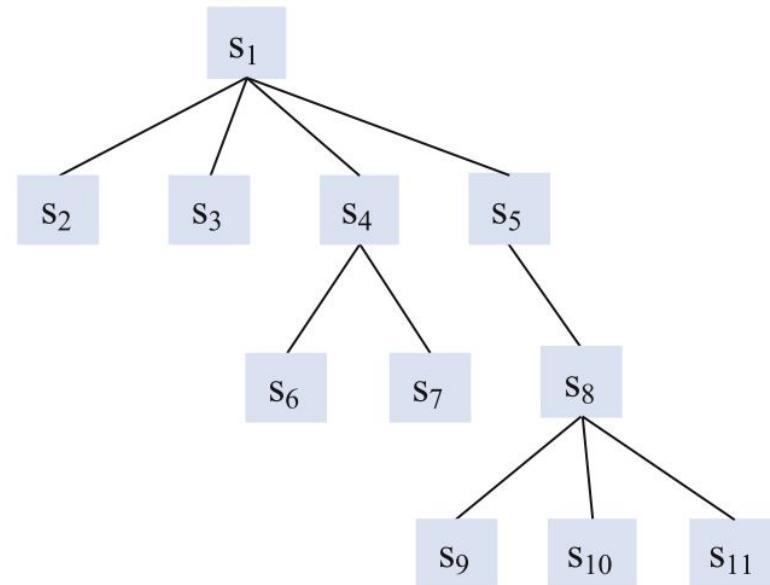
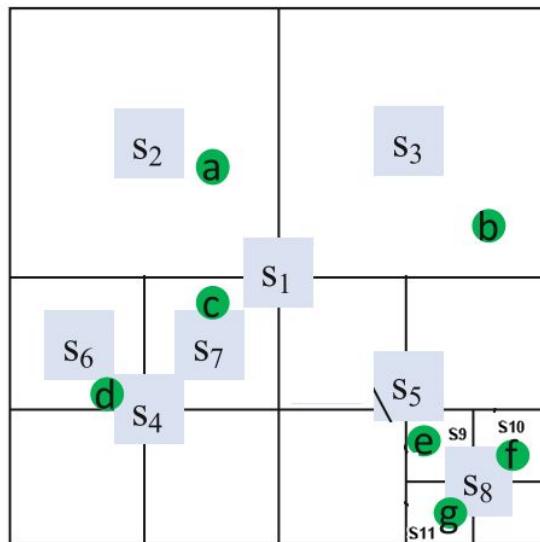
# Octree

- k-d tree split along **a dimension** while octree split along **a point**.
- The center represents for the regions and has 8 children (3D) or 4 children (2D)



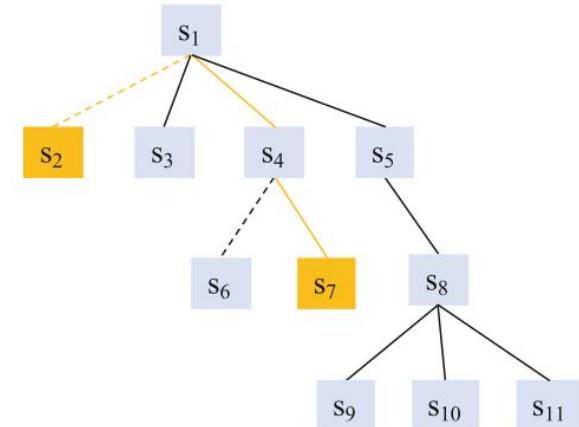
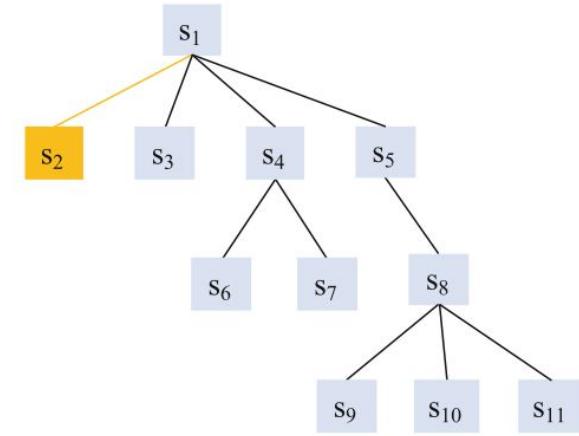
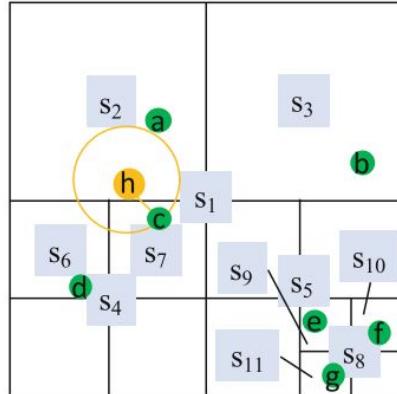
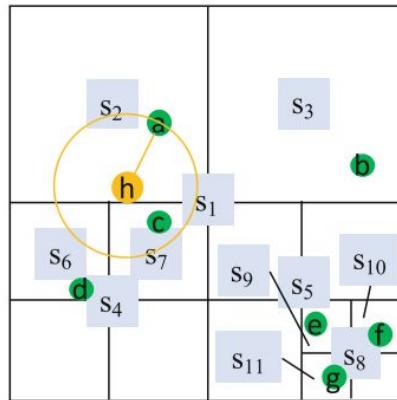
# Octree construction

2D example a(4, 6), b(7,5), c(4,4), d(2,2), e(6, 3), f(8,2), g(7,1)

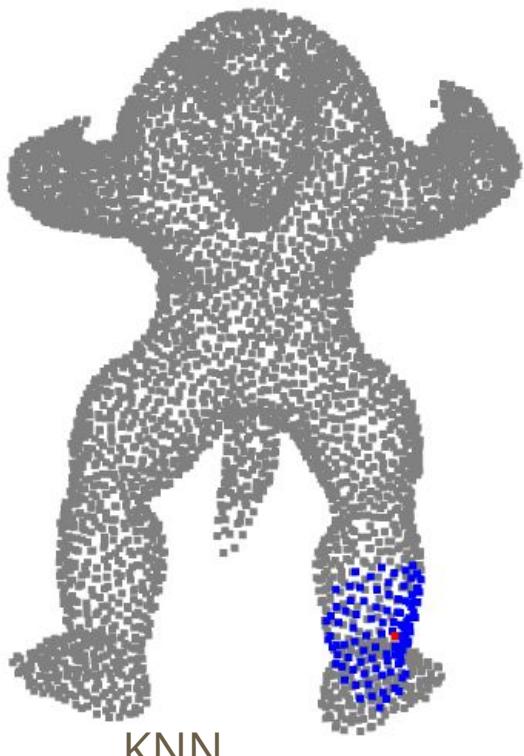


# Octree search with K-NN

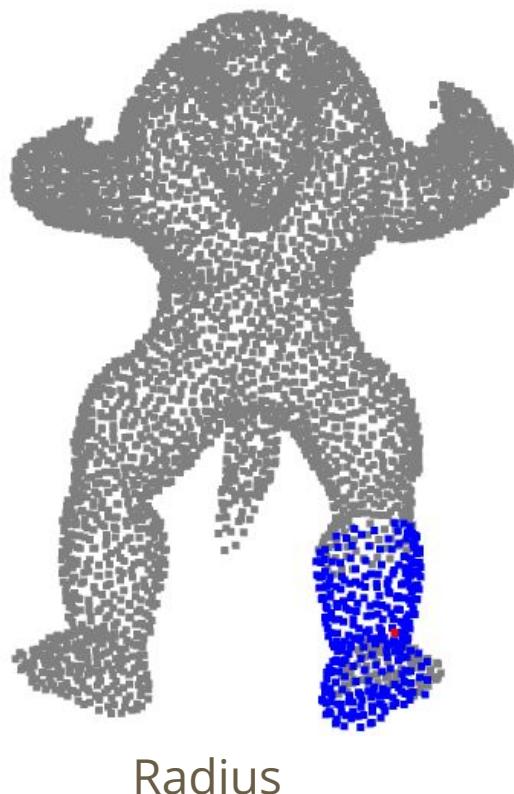
- Use Depth-First Search (DFS)
- Example :
  - Query **h**
  - k=2
  - Initialize distance **d =  $\infty$**
  - Start from **s1** find **h** in **s2** among children > store **a**
  - Update  $d = ||h-a||$
  - Ball circle(**h**,  $d$ ) intersect with **s4**
  - Search children **s4**: **s6,s7**
  - Store **c** where  $d < |h-c|$



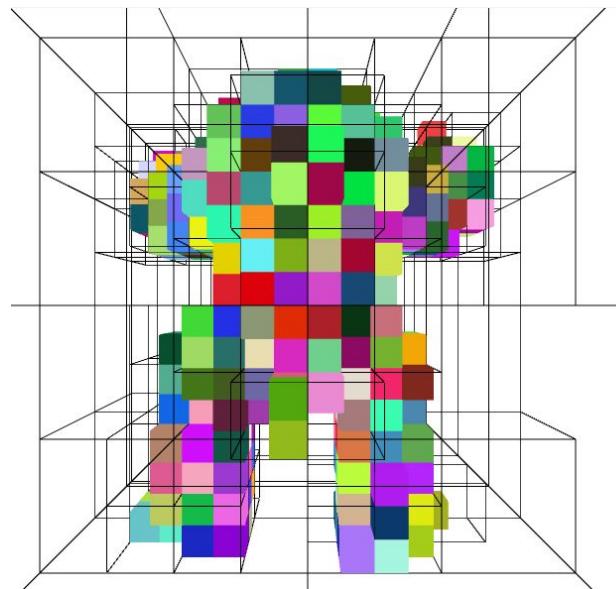
# Example



KNN



Radius



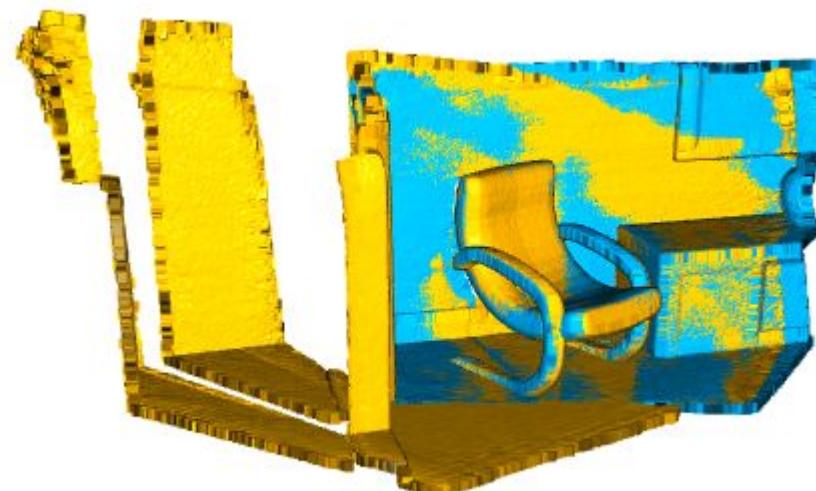
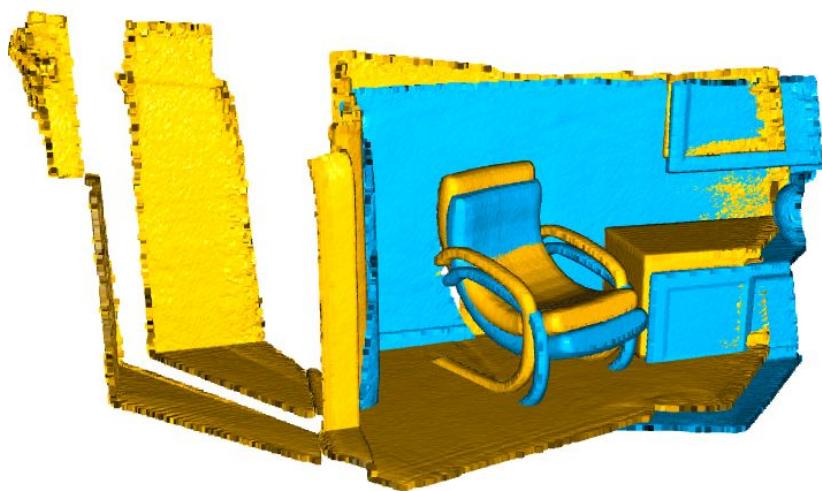
Octree

# Registration

- Align two or more point clouds into **a common coordinate**
- Iterative Closest Point (ICP) is the most common method
  - Find the correspondence between two point sets
  - Find the transformation that minimizes the Euclidean distance between matching points
- Mathematic formation:
  - Let  $A = \{a_i\}$ ,  $B = \{b_j\}$  are two point sets
  - Let  $T$  is transformation
  - Point  $m_i$  in  $A$  which corresponding with  $Tb_j$
  - Find

$$T = \arg \min_T \left\{ \sum_i \|T \cdot b_i - m_i\|^2 \right\}$$

# Example



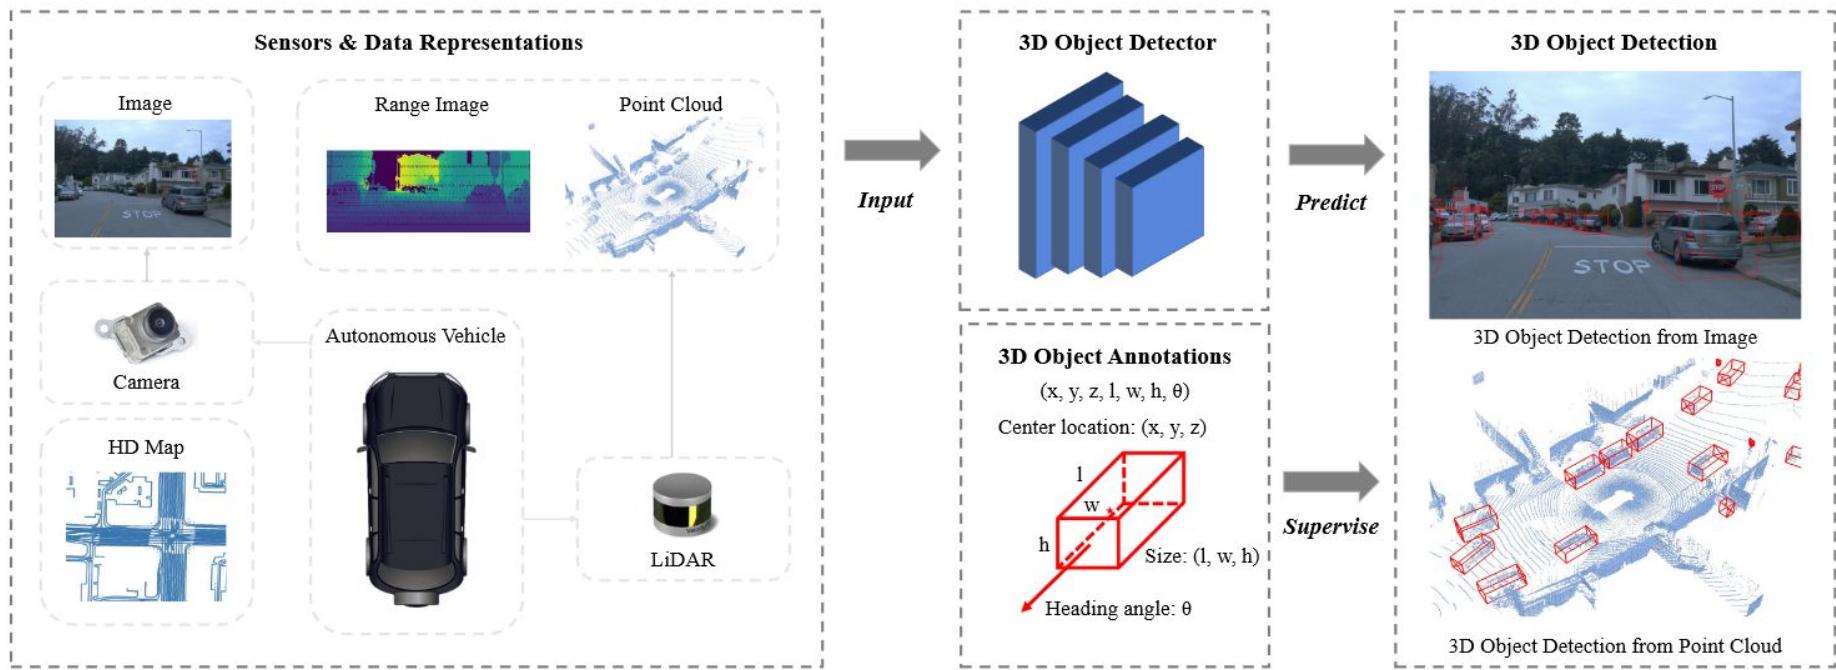
# Coding Examples

- Introduction to **Open3D installation (v0.1)**
- Create folder to store all files below:
- Load a point cloud
  - Bin (KITTI): load\_bin.py
  - Ply (Shapenet): load\_ply.py
  - Txt (ModelNet): load\_txt.py (convert from numpy PCL)
- Downsampling : **downsampling.py**
- Noise Removal : **noise\_removal.py**
- Search neighbors: k-d-tree / octoc.tree (color search results)
- Registration: registration.py
- Reconstruct point cloud from RGB images and depth images:

## Section 2

# Deep Learning on Point Cloud

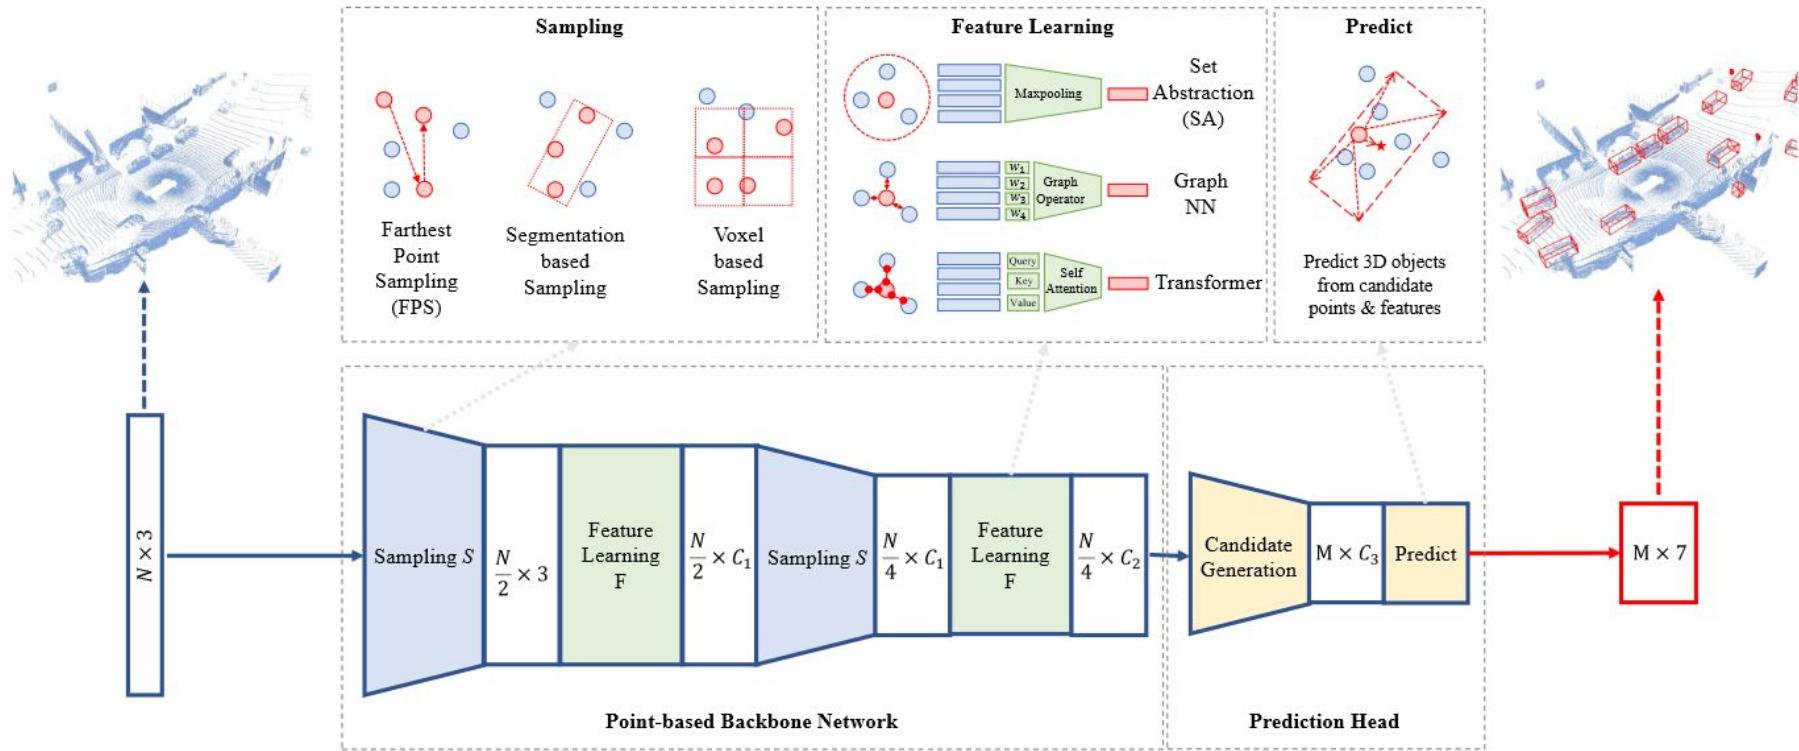
# Overview



# Datasets

Dataset	Year	Size (hr.)	Real-world	LiDAR scans	Images	3D annotations	Classes	night/rain	Locations	Other data
KITTI [80, 81]	2012	1.5	Yes	15k	15k	200k	8	No/No	Germany	-
KAIST [50]	2018	-	Yes	8.9k	8.9k	Yes	3	Yes/No	Korea	thermal images
ApolloScape [104, 166]	2019	100	Yes	20k	144k	475k	6	-/-	China	-
H3D [198]	2019	0.77	Yes	27k	83k	1.1M	8	No/No	USA	-
Lyft L5 [107]	2019	2.5	Yes	46k	323k	1.3M	9	No/No	USA	maps
Argoverse [29]	2019	0.6	Yes	44k	490k	993k	15	Yes/Yes	USA	maps
AIODrive [293]	2020	6.9	No	250k	250k	26M	-	Yes/Yes	-	long-range data
A*3D [202]	2020	55	Yes	39k	39k	230k	7	Yes/Yes	SG	-
A2D2 [82]	2020	-	Yes	12.5k	41.3k	-	14	-/-	Germany	-
Cityscapes 3D [77]	2020	-	Yes	0	5k	-	8	No/No	Germany	-
nuScenes [15]	2020	5.5	Yes	400k	1.4M	1.4M	23	Yes/Yes	SG, USA	maps, radar data
Waymo Open [250]	2020	6.4	Yes	230k	1M	12M	4	Yes/Yes	USA	maps
Cirrus [288]	2021	-	Yes	6.2k	6.2k	-	8	-/-	USA	long-range data
PandaSet [301]	2021	0.22	Yes	8.2k	49k	1.3M	28	Yes/Yes	USA	-
KITTI-360 [142]	2021	-	Yes	80k	300k	68k	37	-/-	Germany	-
Argoverse v2 [295]	2021	-	Yes	-	-	-	30	Yes/Yes	USA	maps
ONCE [172]	2021	144	Yes	1M	7M	417K	5	Yes/Yes	China	-

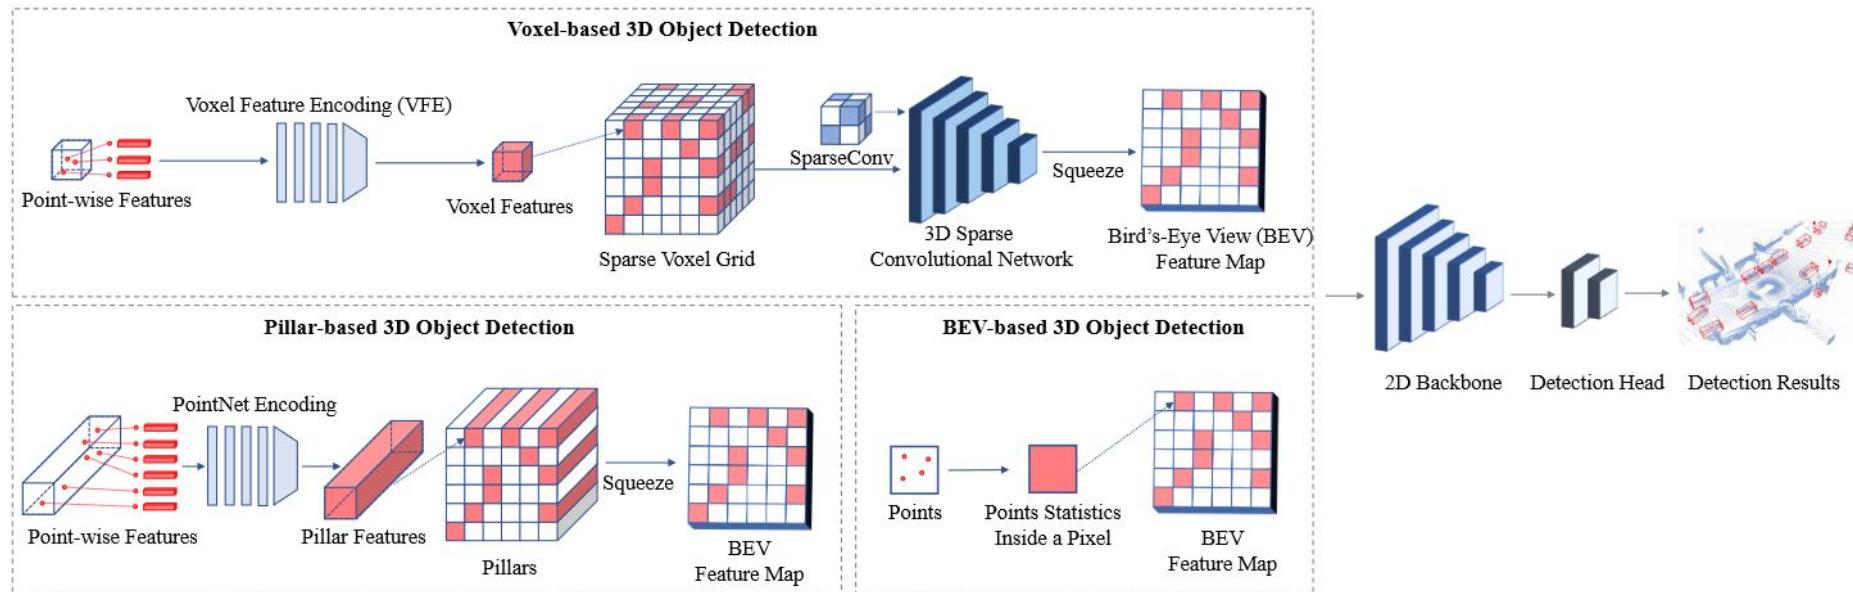
# Point-based 3D methods



# Examples

Method	Context $\Omega$	Sampling $S$	Feature $F$
PointRCNN [234]	Ball Query	FPS	Set Abstraction
IPOD [318]	Ball Query	Seg.	Set Abstraction
STD [319]	Ball Query	FPS	Set Abstraction
3DSSD [321]	Ball Query	Fusion-FPS	Set Abstraction
Point-GNN [238]	Ball Query	Voxel	Graph
StarNet [189]	Ball Query	Targeted-FPS	Graph
Pointformer [193]	Ball Query	FPS + Refine	Transformer

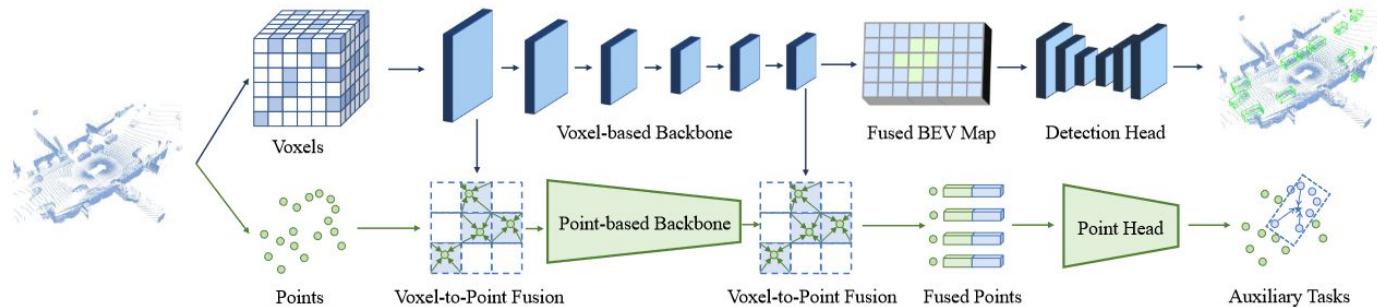
# Voxel (3D grid)-based methods



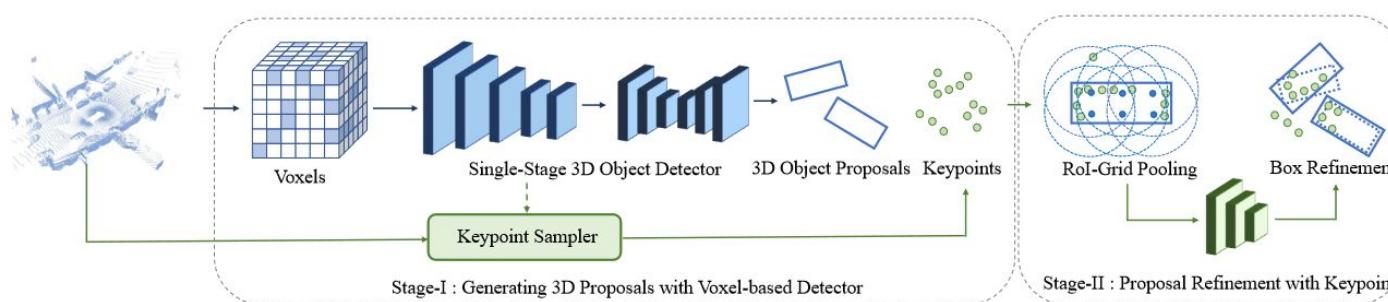
# Example

Method	Representation	Feature Encoding	Backbone Model	Refinement Head
Vote3D [264]	voxel	voxelization	3D sparse CNN w/ voting	-
Vote3Deep [66]	voxel	voxelization	3D sparse CNN w/ voting	-
3D-FCN [118]	voxel	voxelization	3D FCN	-
VeloFCN [119]	BEV map	projection	2D FCN	-
BirdNet [10]	BEV map	projection	2D RPN	-
PIXOR [314]	BEV map	projection	2D CNN	-
HDNet [313]	BEV map, HD map	projection	2D CNN	-
VoxelNet [359]	voxel, BEV map	voxelization, VFE	3D sparse CNN, 2D RPN	-
SECOND [312]	voxel, BEV map	voxelization, VFE	3D sparse CNN, 2D RPN	-
MVF [360]	voxel, BEV+PV map	voxelization, VFE	2D CNN	-
PointPillars [117]	pillar, BEV map	PointNet encoding	2D CNN	-
Pillar-OD [283]	pillar, BEV+CYV map	PointNet encoding	2D CNN	-
Part-A <sup>2</sup> Net [236]	voxel, BEV map	voxelization, VFE	3D sparse CNN, 2D RPN	voxel head
Voxel R-CNN [57]	voxel, BEV map	voxelization, VFE	3D sparse CNN, 2D RPN	voxel head
CenterPoint [329]	voxel, BEV map	voxelization, VFE	3D sparse CNN, 2D RPN	-
Voxel Transformer [173]	voxel, BEV map	voxelization, VFE	3D Transformer, 2D RPN	-

# Hybrid: Point-Voxel based methods



Single-Stage Point-Voxel Detection Framework



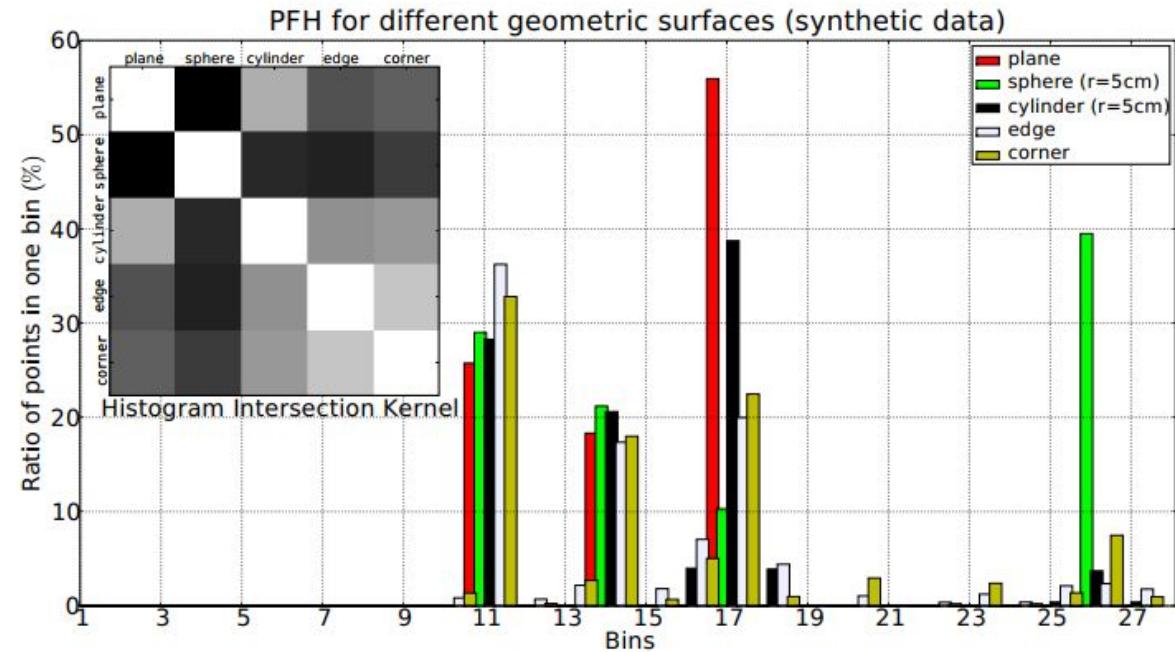
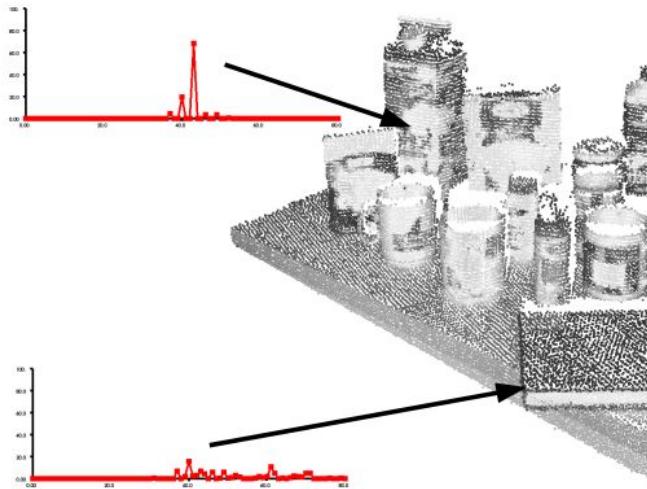
Two-Stage Point-Voxel Detection Framework

# Examples

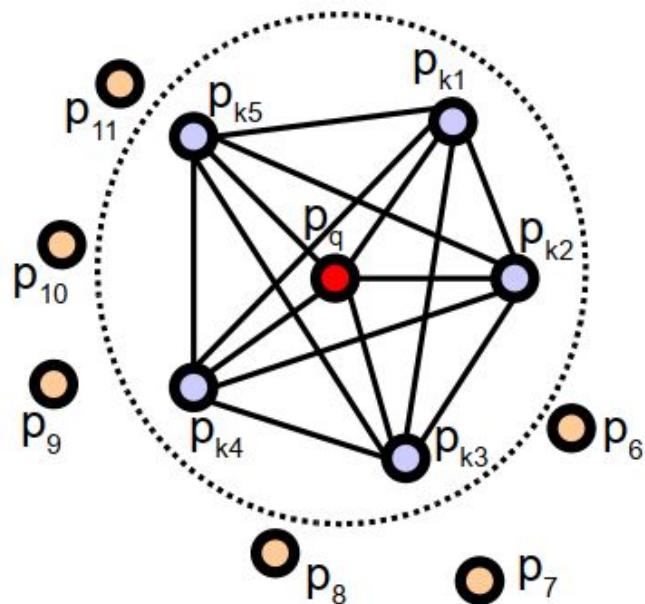
Method	Contribution
Single-Stage Detection Framework	
PVCNN [152]	Point-Voxel Convolution
SPVNAS [254]	Sparse Point-Voxel Convolution
SA-SSD [92]	Auxiliary Point Network
PVGNet [181]	Point-Voxel-Grid Fusion
Two-Stage Detection Framework	
Fast Point R-CNN [43]	RefinerNet
PV-RCNN [235]	RoI-grid Pooling
PV-RCNN++ [237]	VectorPool
Pyramid R-CNN [171]	RoI-grid Attention
LiDAR R-CNN [134]	Scale-aware Pooling
CT3D [233]	Channel-wise Transformer

# Intrinsic Geometry Features (Hand-crafted): Point Feature Histogram (PFH)

- Surface normal and curvature are too simple to represent point features



# Algorithms

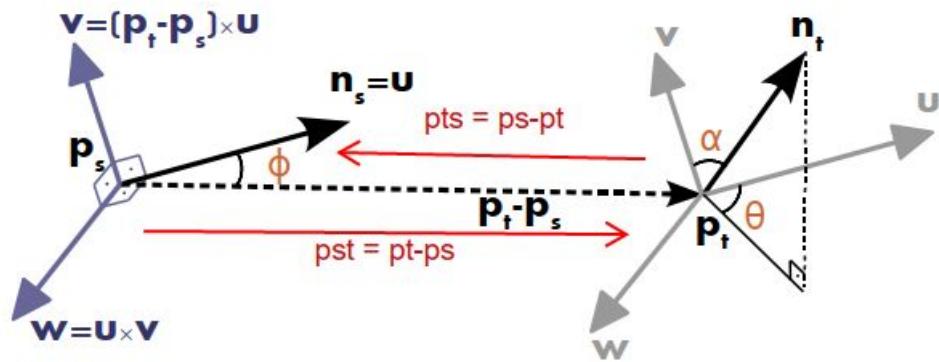


$(\exists) r_1, r_2 \in \mathbb{R}, r_1 < r_2$ , such that  $\begin{cases} r_1 \Rightarrow \mathcal{P}^{k_1} \\ r_2 \Rightarrow \mathcal{P}^{k_2} \end{cases}$ , with  $0 < k_1 < k_2$

## Dual rings:

- First ring: estimate surface normal (see section 1)
- Second ring: estimate PFH

# Darboux frame



$$\begin{cases} u = n_s \\ v = u \times \frac{(p_t - p_s)}{\|p_t - p_s\|_2} \\ w = u \times v \end{cases}$$

$$\alpha = v \cdot n_t$$

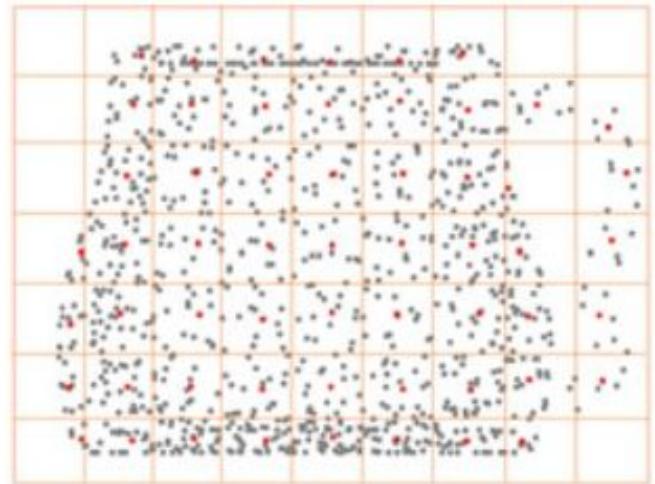
$$\phi = u \cdot \frac{(p_t - p_s)}{d}$$

$$\theta = \arctan(w \cdot n_t, u \cdot n_t)$$

- Binning quadruplets  $(\alpha, \phi, \theta, d)$  and indexing to 1-dimension

# Convolutional Operation on Point Cloud

- Cannot be applied directly as on images
  - Voxelize point cloud into grid coordinate
  - Define a fixed grid coordinate for all pointcloud
  - 3D convolutional operation > memory & computing cost

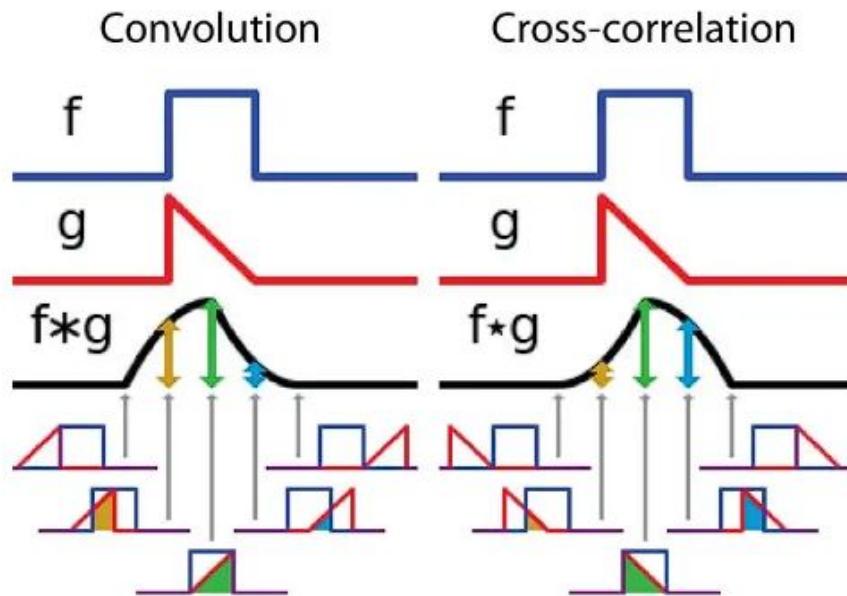


# Reviews Convolution Operation

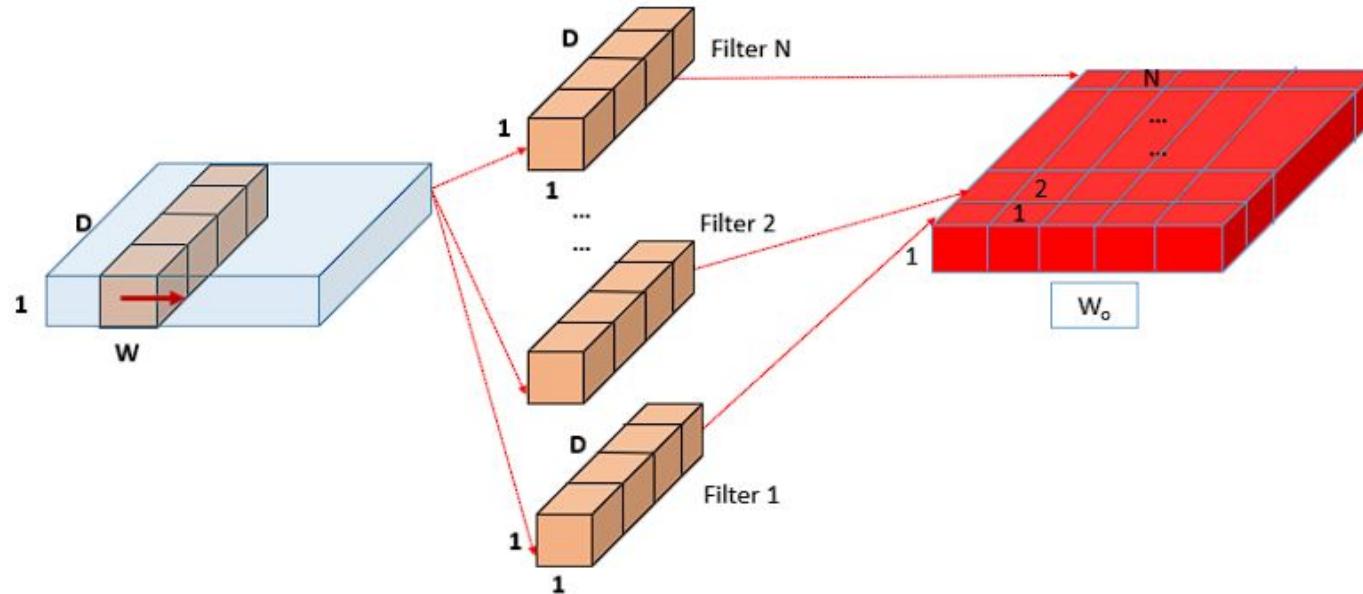
Continuous  $(f * g)(t) = \int_{-\infty}^{\infty} f(\tau) \cdot g(t - \tau) d(\tau)$

Discrete  $(f * g)(t_k) = \sum_{i=0}^m f(\tau_i) \cdot g(t_k - \tau_i)$

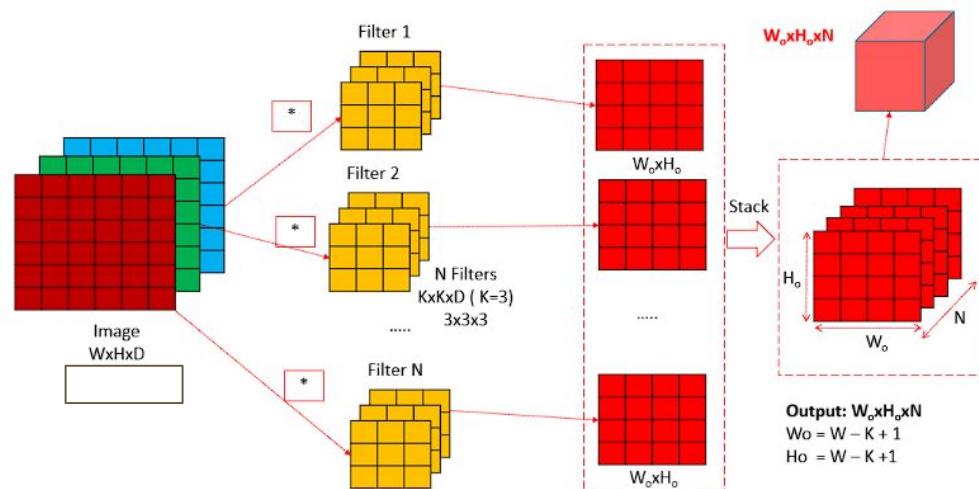
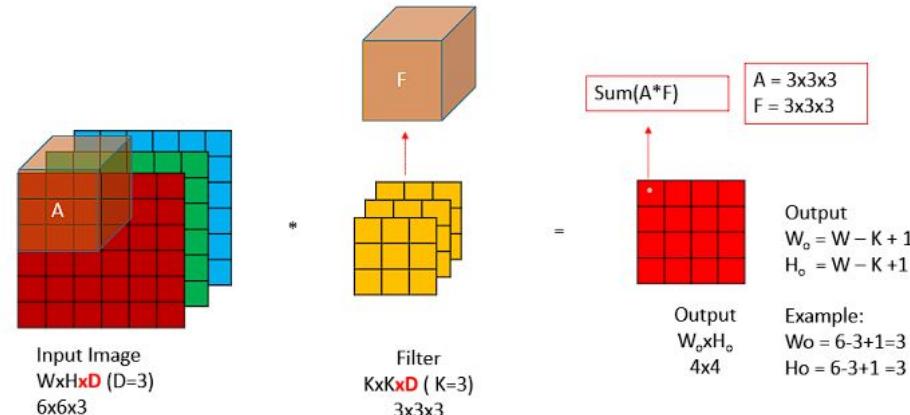
$$\begin{bmatrix} (f * g)(t_1) \\ (f * g)(t_2) \\ \dots \\ (f * g)(t_n) \end{bmatrix} = \begin{bmatrix} \sum_{i=0}^m f(\tau_i) \cdot g(t_1 - \tau_i) \\ \sum_{i=0}^m f(\tau_i) \cdot g(t_2 - \tau_i) \\ \dots \\ \sum_{i=0}^m f(\tau_i) \cdot g(t_n - \tau_i) \end{bmatrix}$$



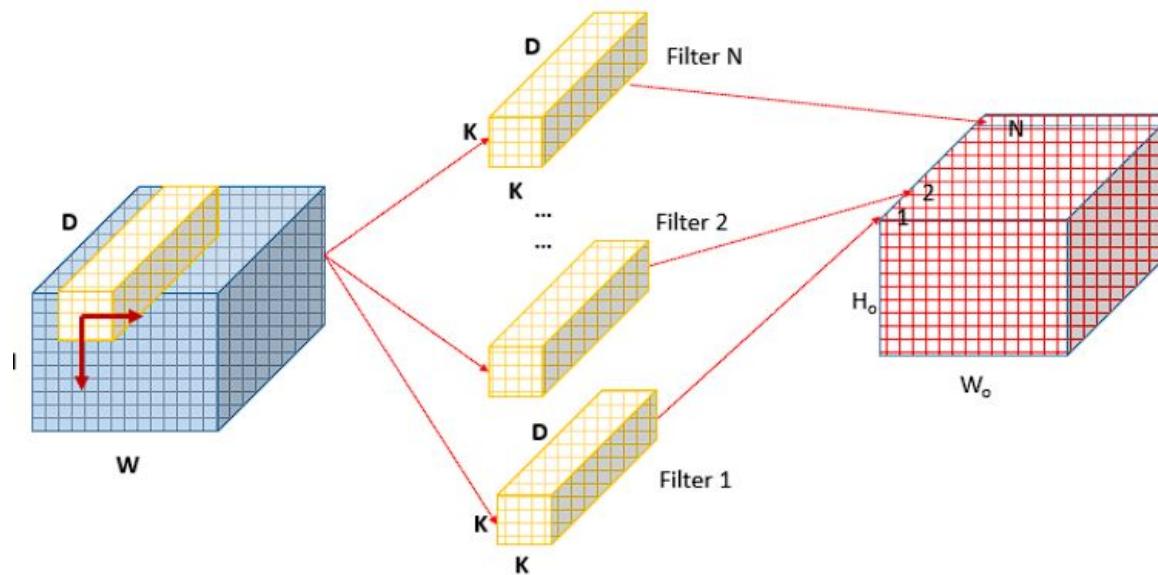
# 1D Convolutions



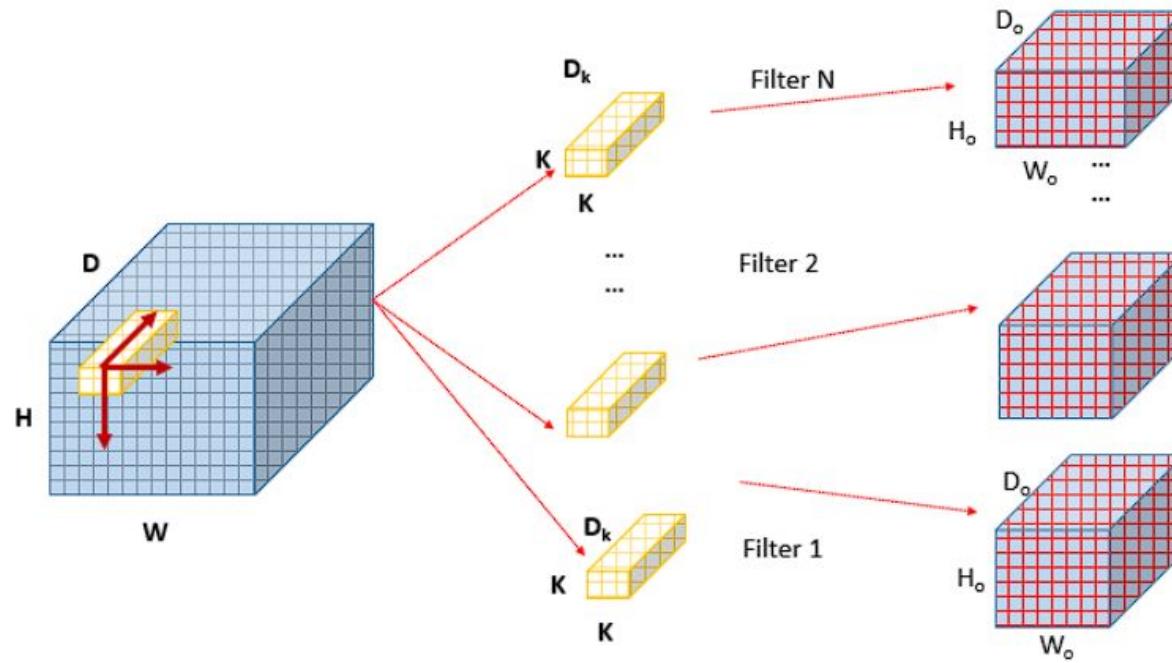
# 2D Convolution



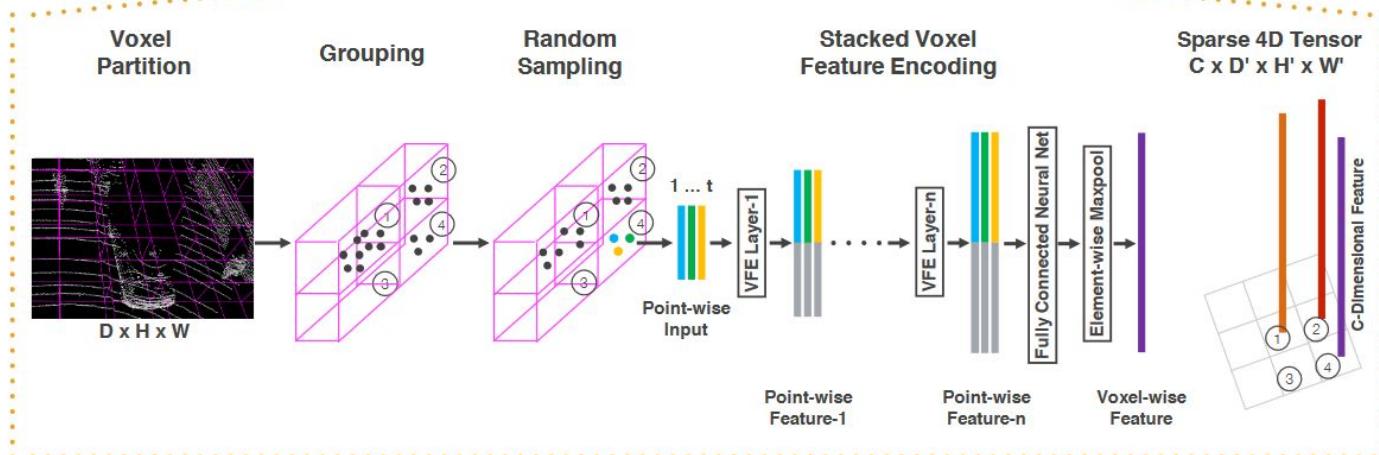
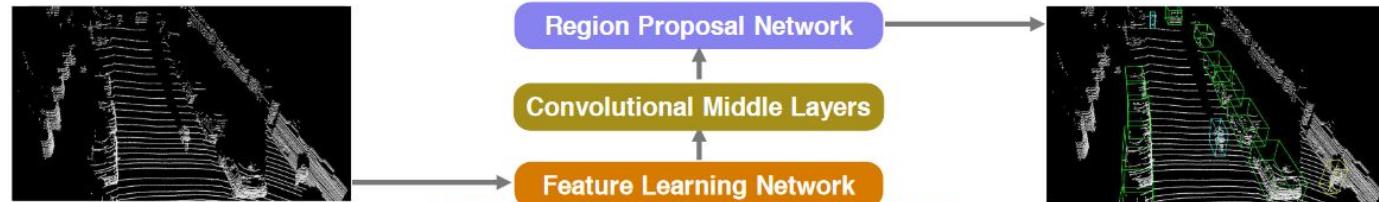
# 2D Convolution (cont')



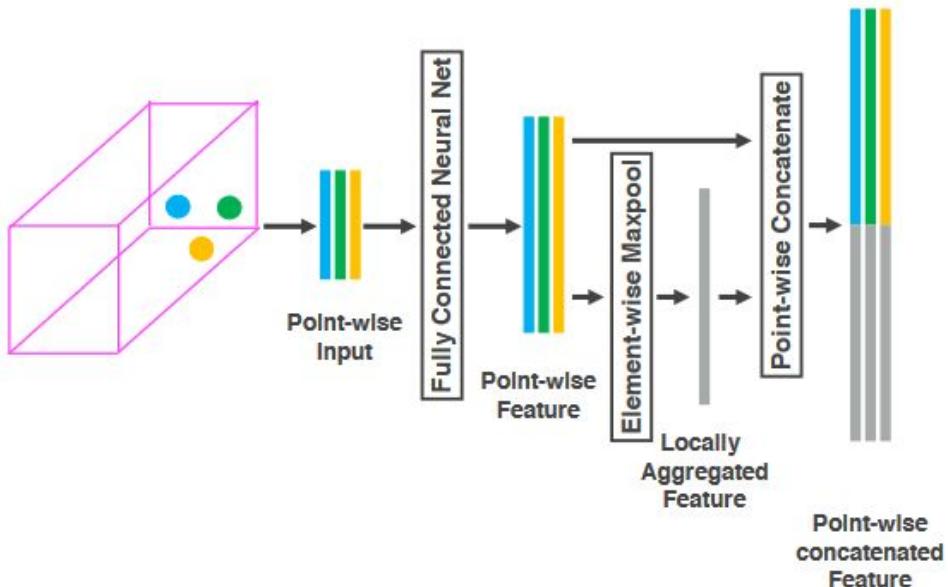
# 3D Convolution



# 3D convolutional operator on 3D voxel grid



# Voxel Feature Encoding (VFE)



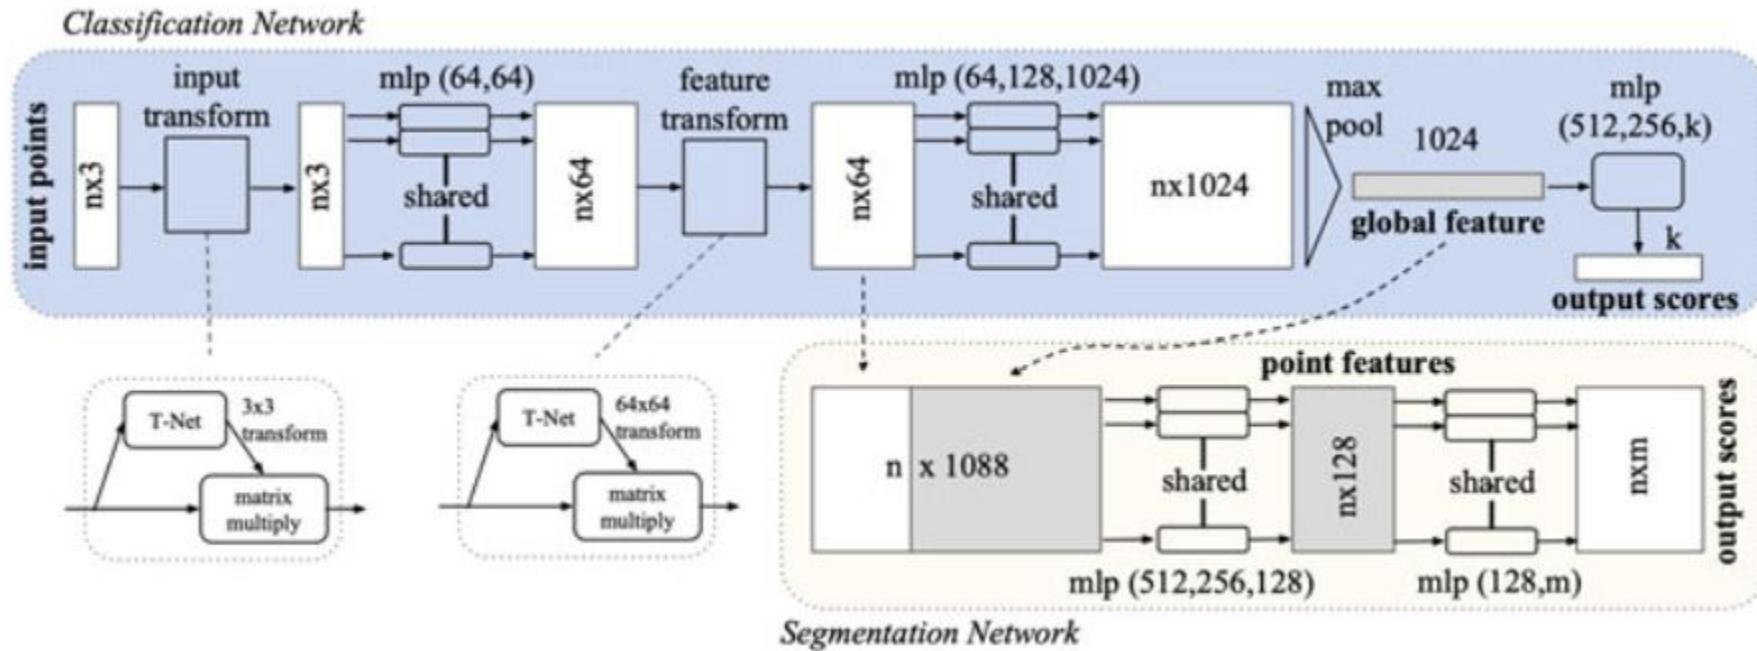
# PointNet

- Consume direct point cloud with **n** points
- Use symmetric function like max-pooling

$$f(x_1, x_2, \dots, x_n) = \gamma \left( \underset{i=1, \dots, n}{\text{MAX}} \{h(x_i)\} \right)$$

where  $\gamma$  and  $h$  are usually multi-layer perceptron (MLP) networks

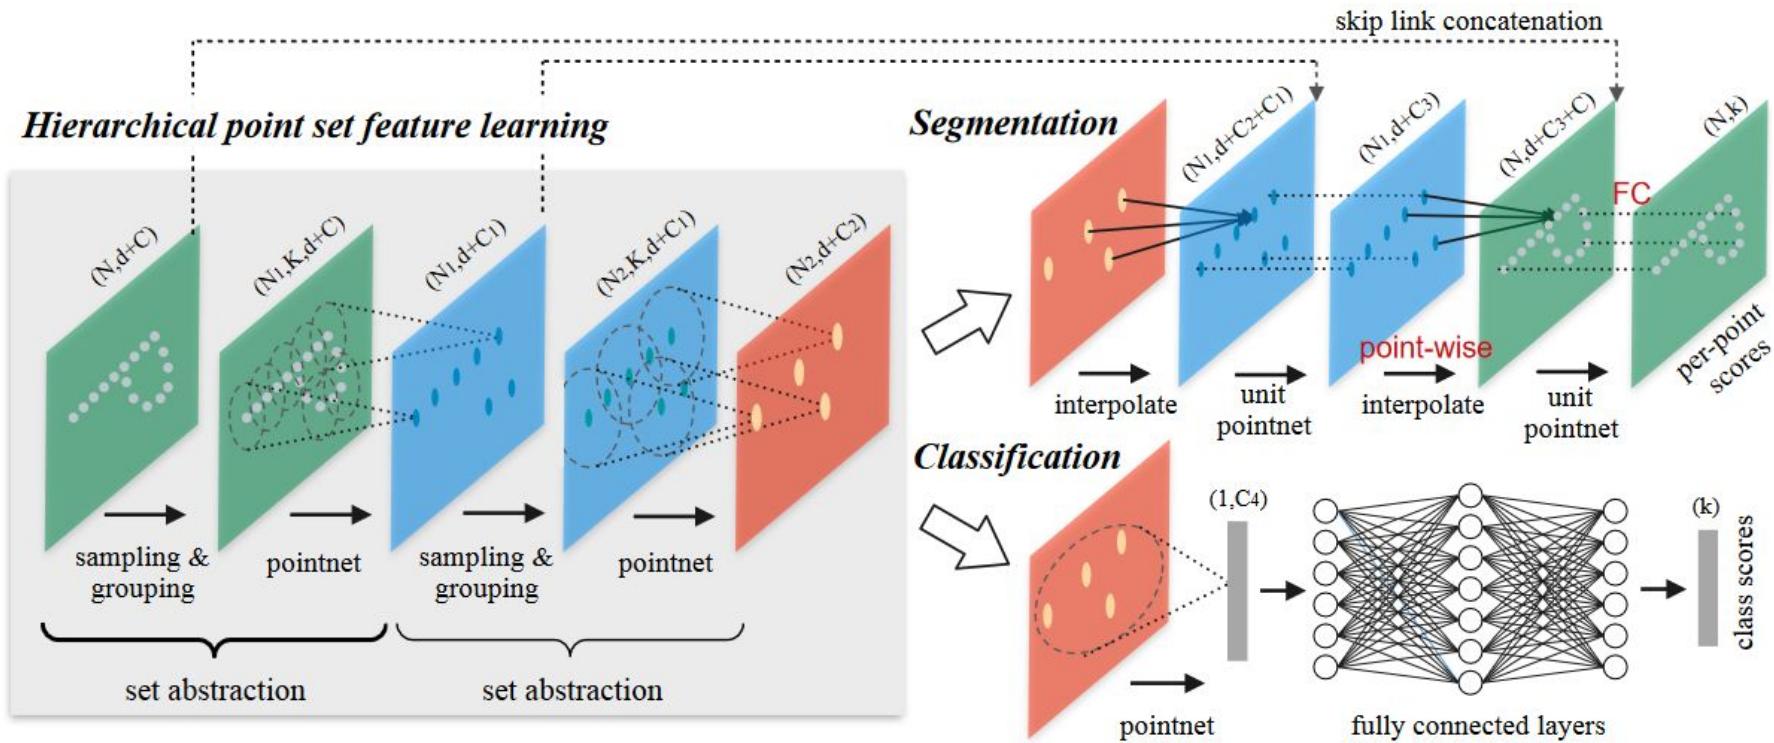
# Architecture



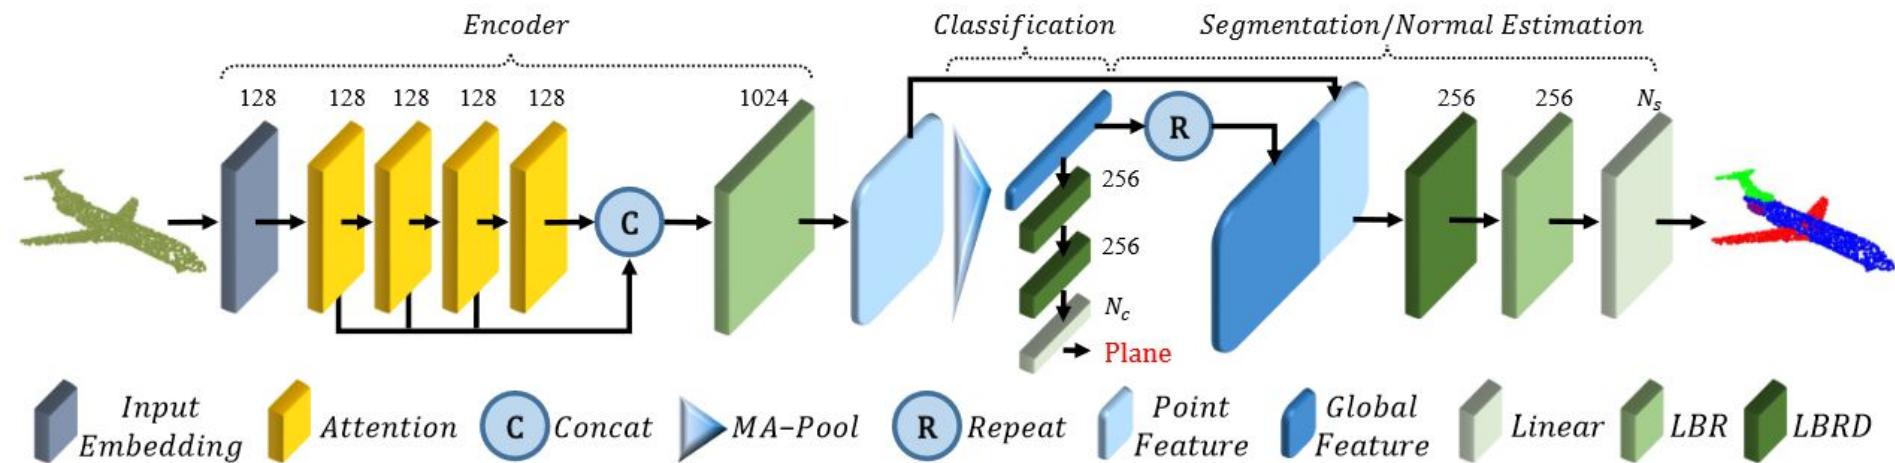
# PointNet++

- Hierarchical Point Set Feature Learning

- Multi-scale grouping (MSG)
- Multi-resolution grouping (MRG)

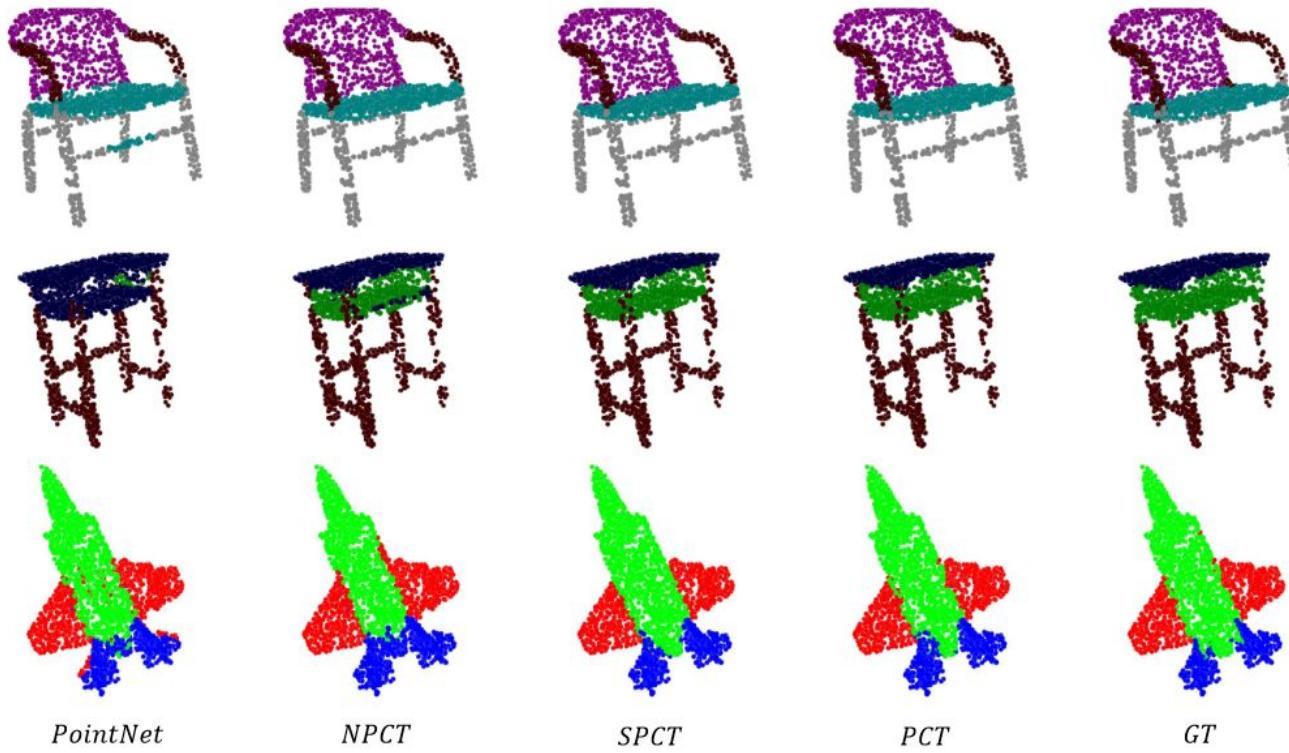


# Point Transformer

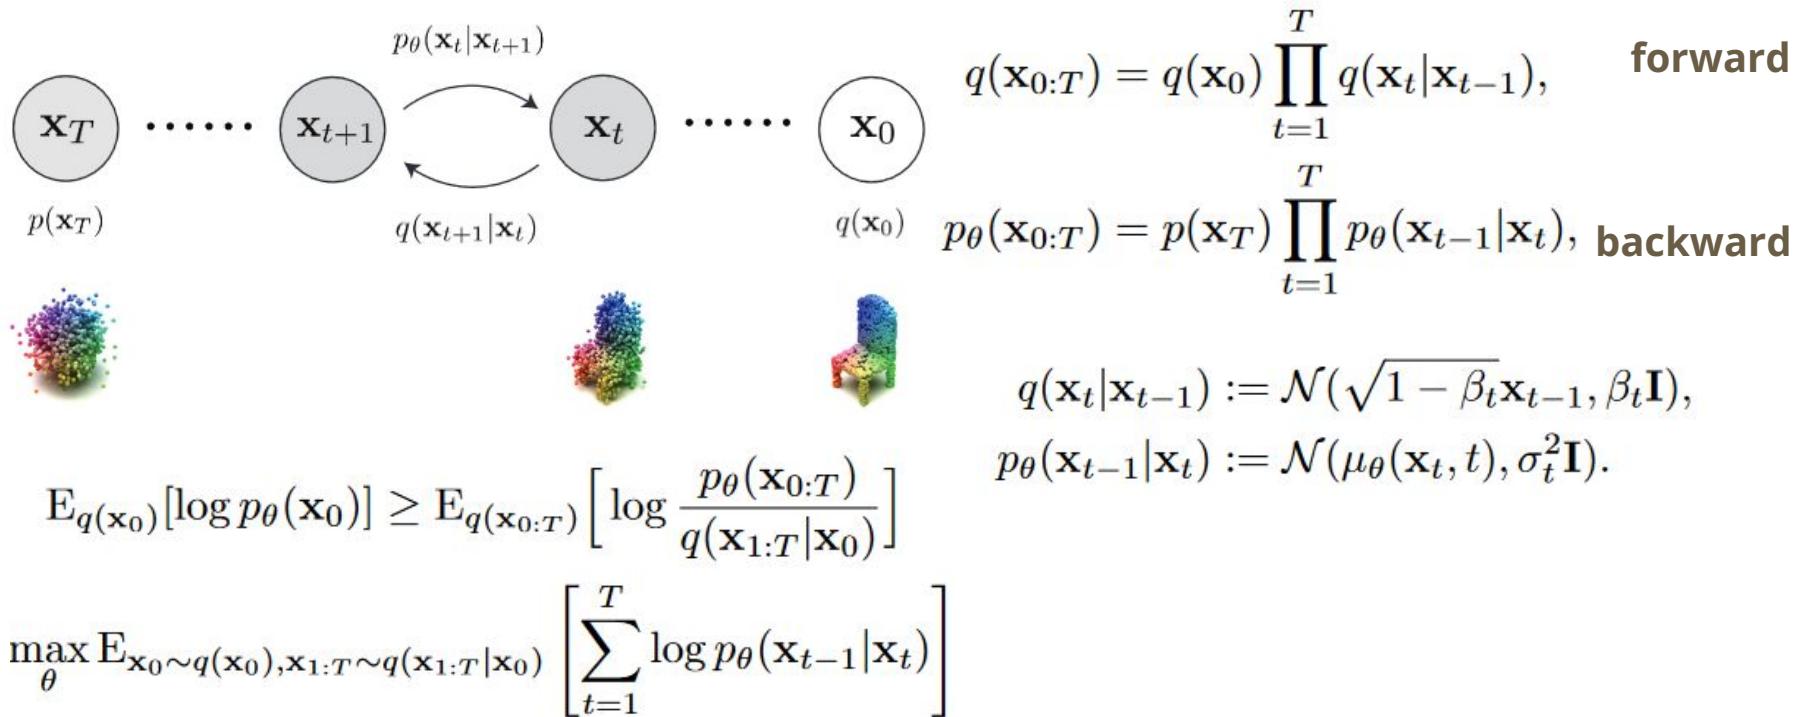


LBR: Linear, Batch, ReLu  
LBRD: Linear, Batch, ReLu, DropOut

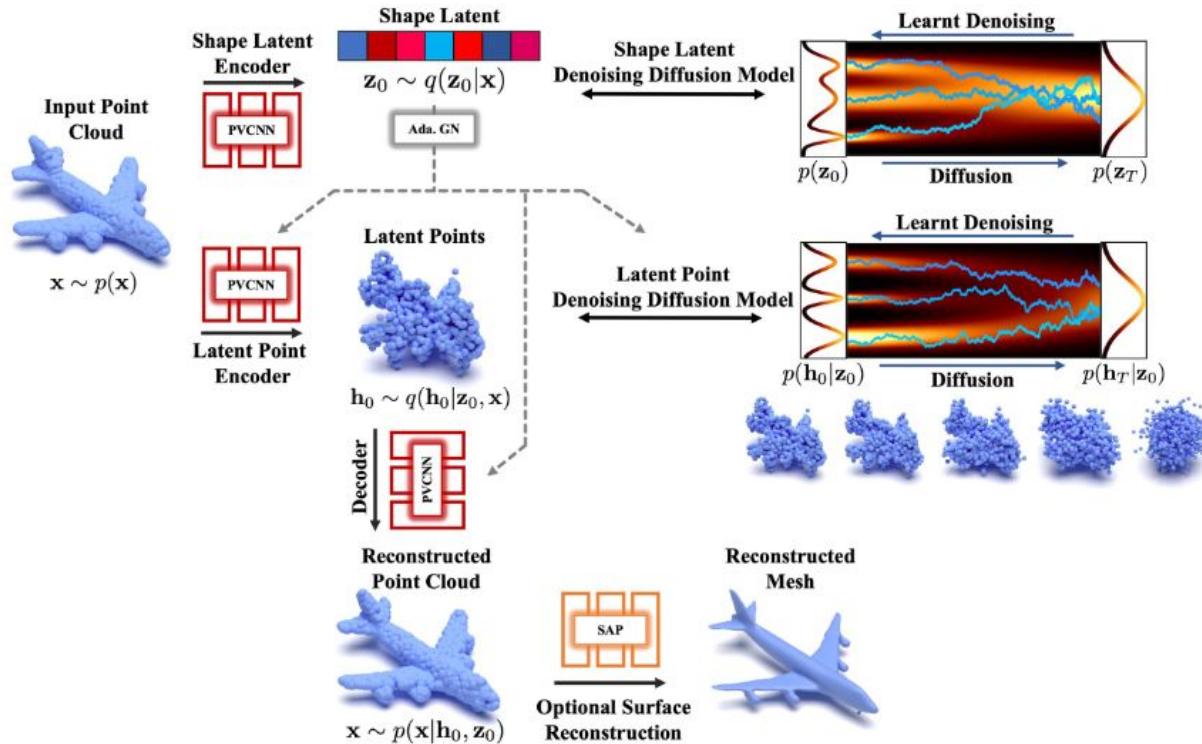
# Some results



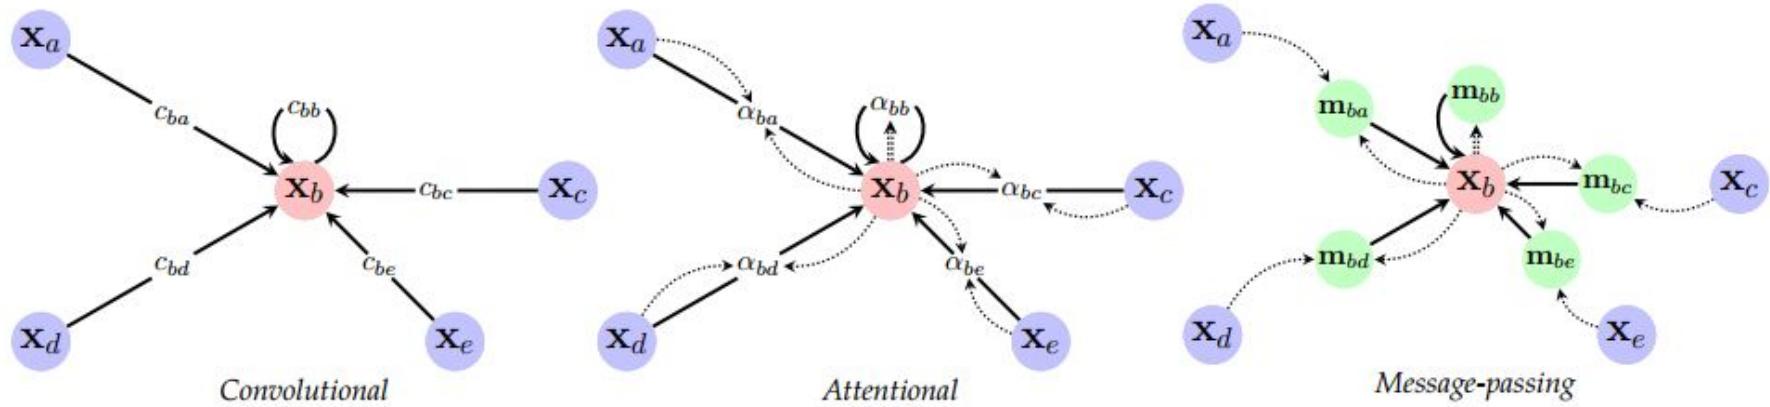
# Generative Point Cloud (Diffusion Approach)



# Latent Point Diffusion



# Advanced: Graph Neural Network on Point Cloud



# Formulation

Graph Convolutional

$$\mathbf{h}_u = \phi \left( \mathbf{x}_u, \bigoplus_{v \in \mathcal{N}_u} c_{uv} \psi(\mathbf{x}_v) \right)$$

Attention

$$\mathbf{h}_u = \phi \left( \mathbf{x}_u, \bigoplus_{v \in \mathcal{N}_u} a(\mathbf{x}_u, \mathbf{x}_v) \psi(\mathbf{x}_v) \right)$$

Message-passing

$$\mathbf{h}_u = \phi \left( \mathbf{x}_u, \bigoplus_{v \in \mathcal{N}_u} \boxed{\psi(\mathbf{x}_u, \mathbf{x}_v)} \right)$$

$\oplus$ : Aggregation, permutation-invariant function

$\Psi$ : Point Feature Transform

$\phi$ : Point Feature Propagation (Diffusion)

$$\begin{aligned}\psi(\mathbf{x}) &= \mathbf{Wx} + \mathbf{b}; \\ \phi(\mathbf{x}, \mathbf{z}) &= \sigma (\mathbf{Wx} + \mathbf{Uz} + \mathbf{b}),\end{aligned}$$

# Geometric Pytorch



latest

Search docs

- INSTALL PYG
  - Installation
- GET STARTED
  - Introduction by Example
  - Colab Notebooks and Video Tutorials
- TUTORIALS
  - Design of Graph Neural Networks
  - Working with Graph Datasets
  - Use-Cases & Applications
  - Distributed Training
- ADVANCED CONCEPTS
  - Advanced Mini-Batching
  - Memory-Efficient Aggregations
  - Hierarchical Neighborhood Sampling
  - Compiled Graph Neural Networks
  - TorchScript Support

PyG Documentation



This documentation is for an **unreleased development version**. Click [here](#) to access the documentation of the current stable release.

## PyG Documentation

 **PyG** (*PyTorch Geometric*) is a library built upon  **PyTorch** to easily write and train Graph Neural Networks (GNNs) for a wide range of applications related to structured data.

It consists of various methods for deep learning on graphs and other irregular structures, also known as **geometric deep learning**, from a variety of published papers. In addition, it consists of easy-to-use mini-batch loaders for operating on many small and single giant graphs, multi GPU-support, `torch.compile` support, `DataPipe` support, a large number of common benchmark datasets (based on simple interfaces to create your own), the `GraphGym` experiment manager, and helpful transforms, both for learning on arbitrary graphs as well as on 3D meshes or point clouds.

 Join our Slack community!

### Install PyG

- [Installation](#)

### Get Started

- [Introduction by Example](#)
- [Colab Notebooks and Video Tutorials](#)

## Point Cloud Processing

This tutorial explains how to leverage Graph Neural Networks (GNNs) for operating and training on point cloud data. Although point clouds do not come with a graph structure by default, we can utilize  PyG transformations to make them applicable for the full suite of GNNs available in  PyG. The key idea is to create a synthetic graph from point clouds, from which we can learn meaningful local geometric structures via a GNN's message passing scheme. These point representations can then be used to, e.g., perform point cloud classification or segmentation.

### 3D Point Cloud Datasets

 PyG provides several point cloud datasets, such as the `PCPNetDataset`, `S3DIS` and `ShapeNet` datasets. To get started, we also provide the `GeometricShapes` dataset, which is a toy dataset that contains various geometric shapes such cubes, spheres or pyramids. Notably, the `GeometricShapes` dataset contains meshes instead of point clouds by default, represented via `pos` and `face` attributes, which hold the information of vertices and their triangular connectivity, respectively:

```
from torch_geometric.datasets import GeometricShapes  
  
dataset = GeometricShapes(root='data/GeometricShapes')  
print(dataset)  
>>> GeometricShapes(40)  
  
data = dataset[0]  
print(data)  
>>> Data(pos=[32, 3], face=[3, 30], y=[1])
```

# Hand-on project

- PointNet
  - Classification
- Visualization:
  - Inference classification