

Module

Self-supervised and Semi-supervised Learning

Ph.D. Ngo Ba Hung

Email: ngohung@jnu.ac.kr

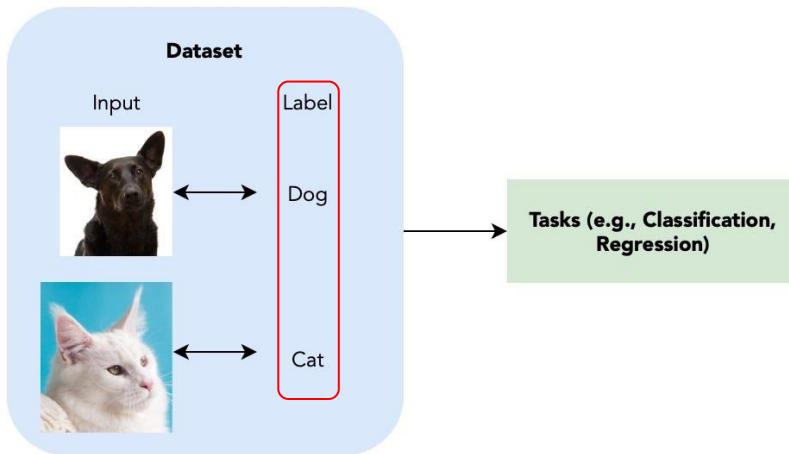
Facebook: <https://www.facebook.com/hung.ngo.7121>

*Graduate School of Data Science
Chonnam National University*

Topic 1: Self-supervised Learning

Overview

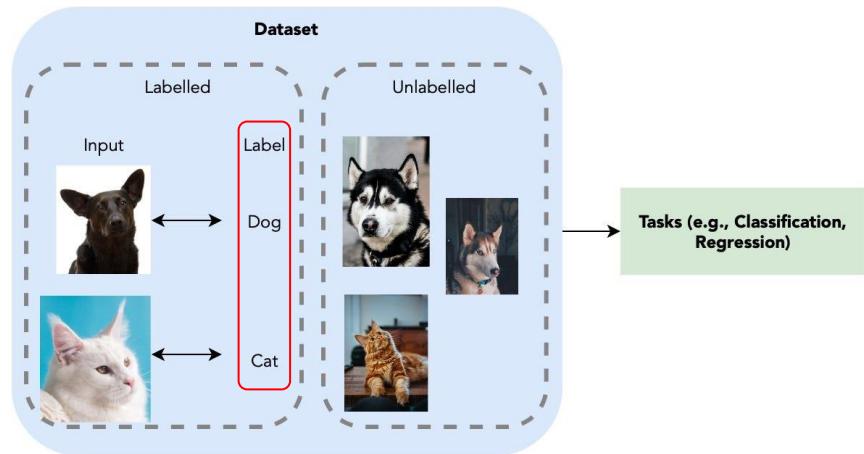
Supervised Learning



supervised learning provides a set of input-output pairs.

The entire dataset to be trained on samples has the corresponding labels.

Semi-supervised Learning



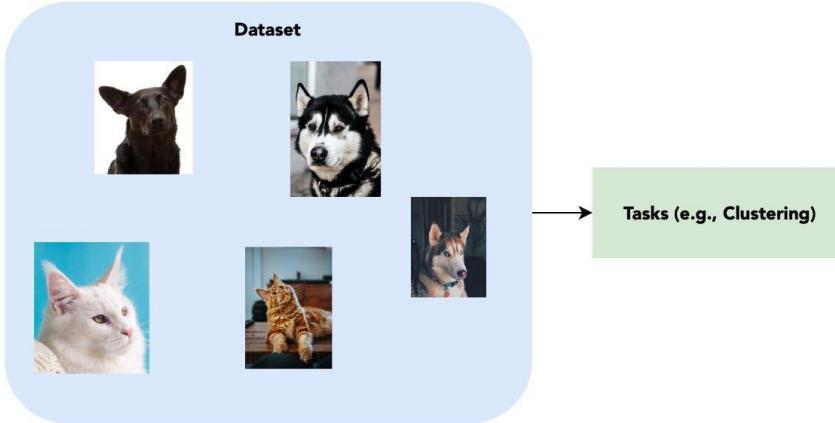
Semi-supervised learning provides a small set of input-output pairs.

The model **trained on small labeled data** is **tested on the large unlabeled data** (the unlabeled data is also used during training).

Topic 1: Self-supervised Learning

Overview

Unsupervised Learning

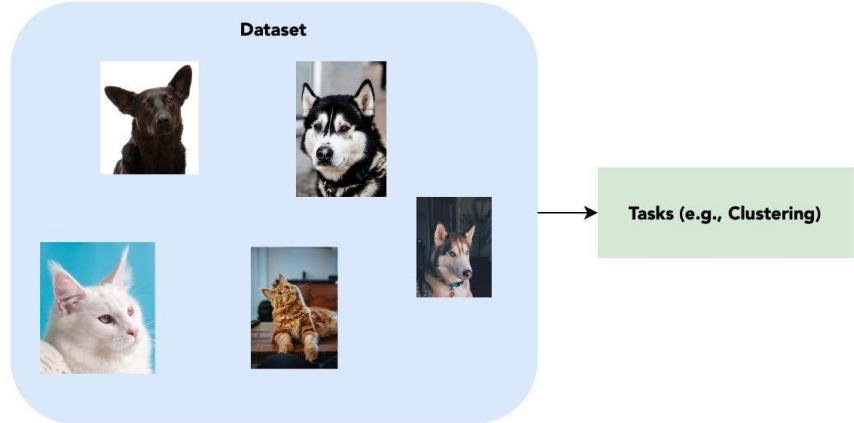


Unsupervised learning only input data has no corresponding labels.

Using method for exploiting extracted representations (samples having the similar features are grouped (clustered) together, but samples having the different features keep far from each other.

(K-NN, K-Mean, K-Medoids methods)

Self-supervised Learning



Self-supervised learning is categorized into to **unsupervised learning** because **no labels were given**.

Topic 1: Self-supervised Learning

Overview

□ Families of SSL

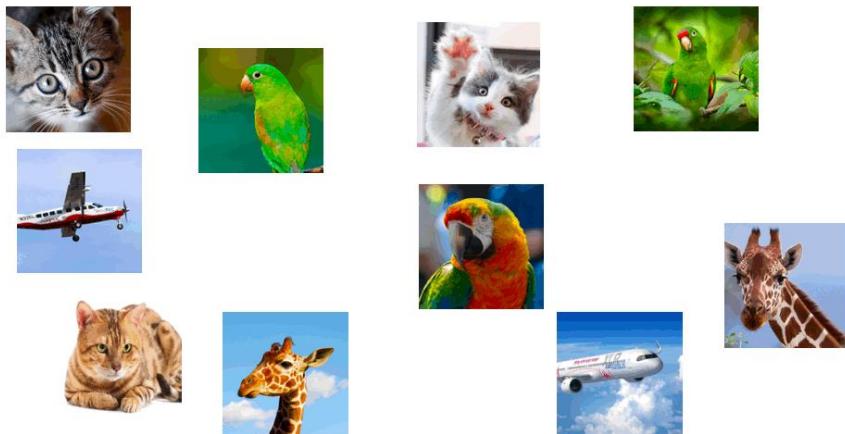
- Deep Metric Learning
 - SimCLR
 - NNCLR
 - SCL
- Self-distillation
 - SimSiam
 - BYOL
 - DINO, DINO v2
- Canonical Correlation Analysis
 - VICReg
 - BarlowTwins
 - SWAV
 - EMP-SSL

Topic 1: Self-supervised Learning

Deep Metric Learning

□ SimCLR*

- Concept of contrastive learning



Mathematic explanation and its formulation:

$$\begin{aligned}\mathcal{L}_{\text{cont}} &= -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)} \\ &= \log \frac{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)} \\ &= \log \left(\frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau) + \sum_{k=1}^{2N} \mathbb{1}_{[k \neq i, j]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)} \right) \\ &= \log \left(1 + \frac{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i, j]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)} \right) \\ \downarrow \mathcal{L}_{\text{cont}} &= \log \left(1 + \frac{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i, j]} \exp(\downarrow \text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}{\exp(\uparrow \text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)} \right)\end{aligned}$$

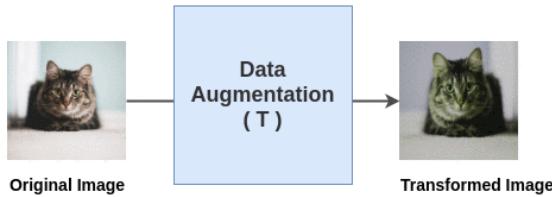
Topic 1: Self-supervised Learning

Deep Metric Learning

□ SimCLR*

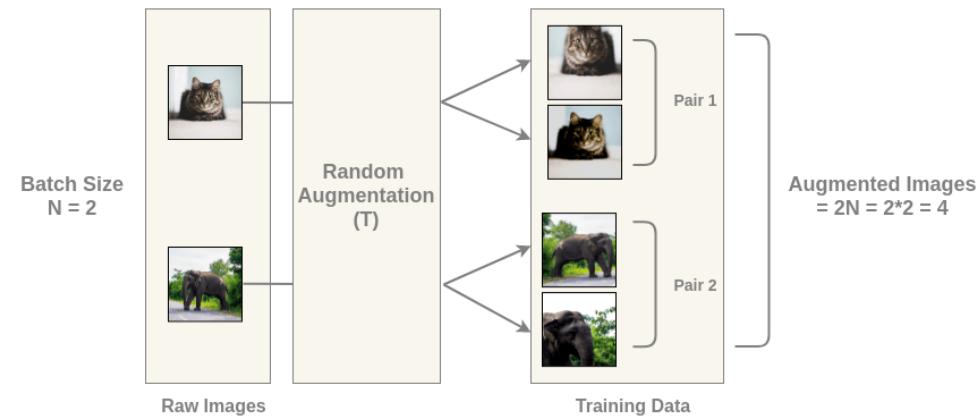
- Image transformation

Random Transformation

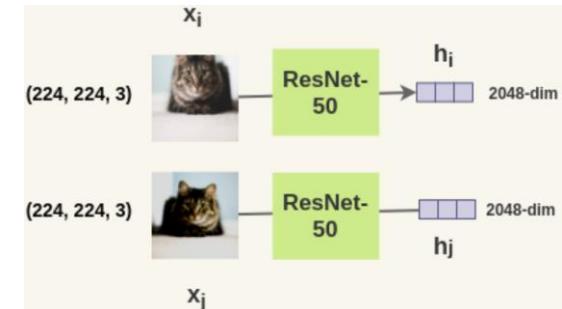
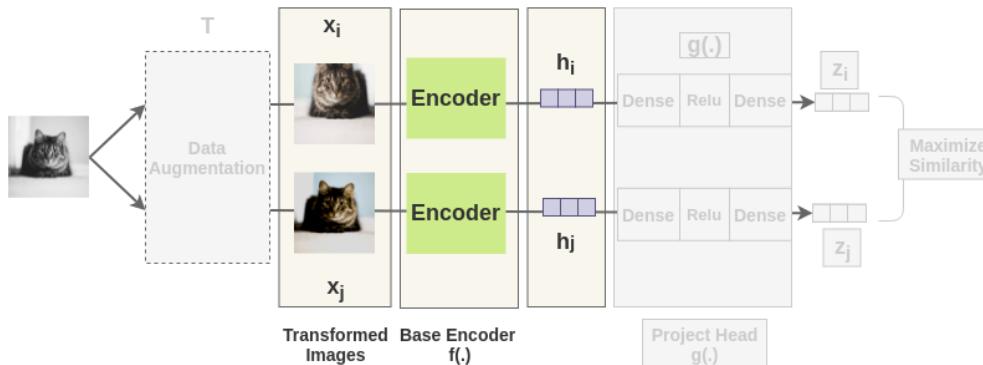


- Getting representation

Preparing similar pairs in a batch



Encoder Component of Framework



Topic 1: Self-supervised Learning

Deep Metric Learning

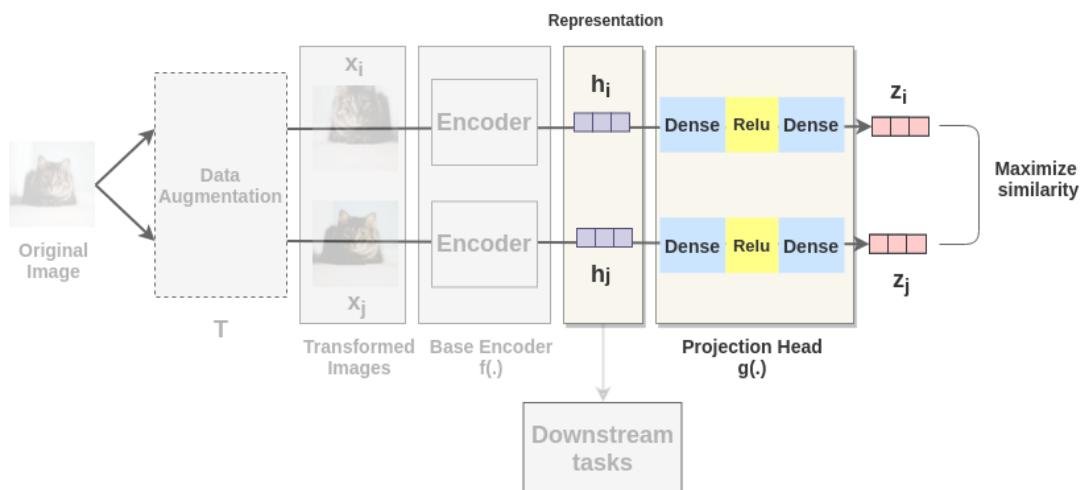
❑ SimCLR*

- Projection head

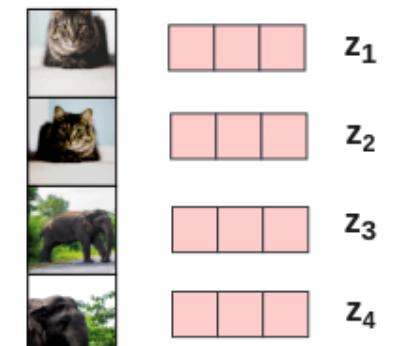
- h_i and h_j are passed through a series of non-linear layers

Dense → Relu → Dense

Projection Head Component



Calculated Embeddings



Topic 1: Self-supervised Learning

Deep Metric Learning

❑ SimCLR*

- Calculation of cosine similarity

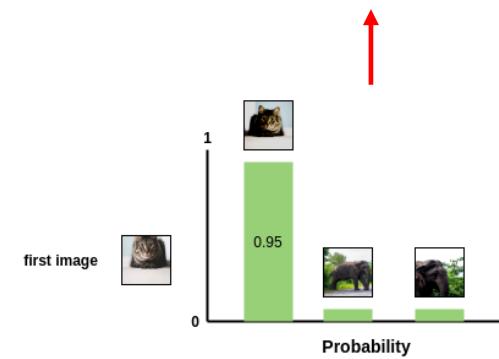
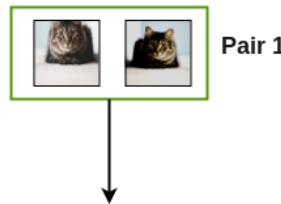
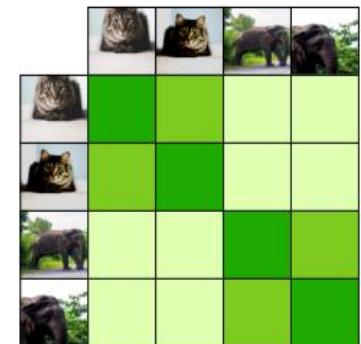
Similarity Calculation of Augmented Images

$$\text{similarity}(\underset{x_i}{\text{cat}}, \underset{x_j}{\text{cat}}) = \text{cosine similarity}(\underset{z_i}{\text{pink}}, \underset{z_j}{\text{pink}})$$

- Calculate the probability of two similarity images

Softmax =
$$\frac{e^{\text{similarity}(\underset{x_i}{\text{cat}}, \underset{x_j}{\text{cat}})}}{e^{\text{similarity}(\underset{x_i}{\text{cat}}, \underset{x_j}{\text{elephant}})} + e^{\text{similarity}(\underset{x_i}{\text{cat}}, \underset{x_j}{\text{elephant}})} + e^{\text{similarity}(\underset{x_i}{\text{elephant}}, \underset{x_j}{\text{elephant}})}}$$

Pairwise cosine similarity



Topic 1: Self-supervised Learning

Self-distillation

□ Simple Siamese network (SimSiam)*

- Architecture
- Pseudo code of SimSiam

Algorithm 1 SimSiam Pseudocode, PyTorch-like

```
# f: backbone + projection mlp
# h: prediction mlp

for x in loader: # load a minibatch x with n samples
    x1, x2 = aug(x), aug(x) # random augmentation
    z1, z2 = f(x1), f(x2) # projections, n-by-d
    p1, p2 = h(z1), h(z2) # predictions, n-by-d

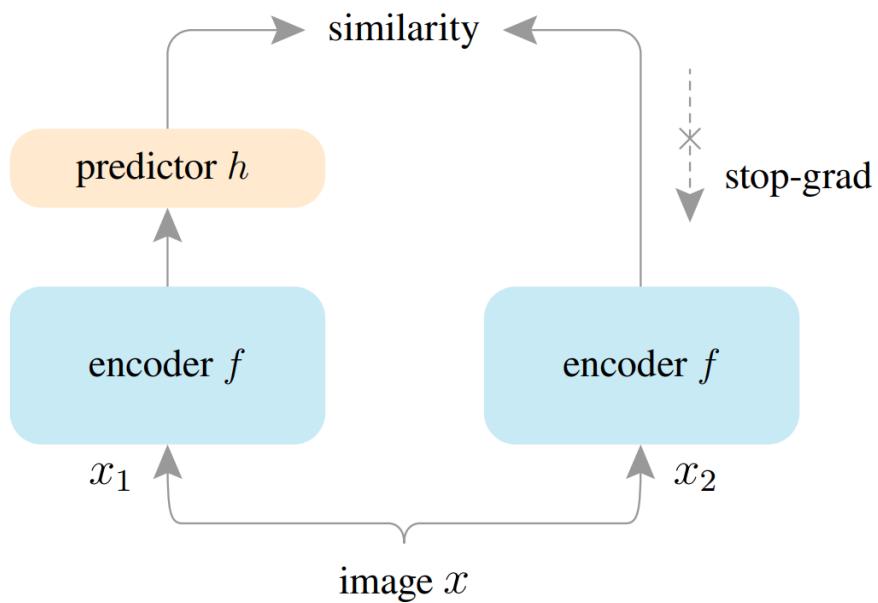
    L = D(p1, z2)/2 + D(p2, z1)/2 # loss

    L.backward() # back-propagate
    update(f, h) # SGD update

def D(p, z): # negative cosine similarity
    z = z.detach() # stop gradient

    p = normalize(p, dim=1) # l2-normalize
    z = normalize(z, dim=1) # l2-normalize
    return -(p*z).sum(dim=1).mean()
```

Encoder uses **ResNet** backbone network
Projection uses **multi-layer perceptron (MLP)**



x_1 and x_2 augmented images of an input image x

Topic 1: Self-supervised Learning

Self-distillation

□ Simple Siamese network (SimSiam)*

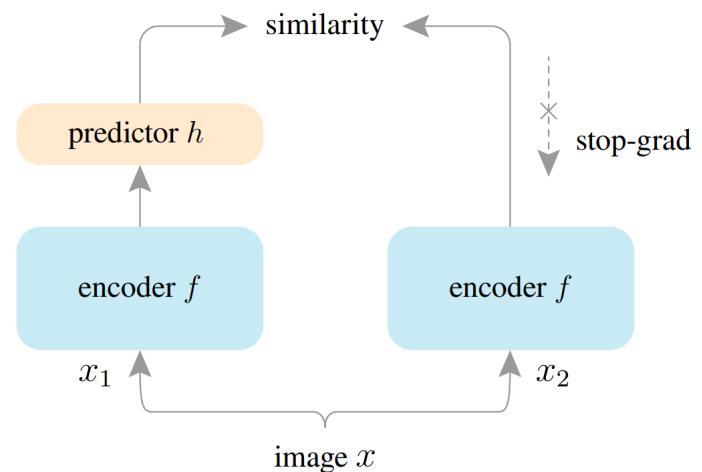
- Output of the encoder

$$z_1 \triangleq f(x_1) \quad z_2 \triangleq f(x_2)$$

- Output of the predictor

$$p_1 \triangleq h(f(x_1)) \quad p_2 \triangleq h(f(x_2))$$

- Minimize the negative cosine similarity



$$\mathcal{D}(p_1, z_2) = -\frac{p_1}{\|p_1\|_2} \cdot \frac{z_2}{\|z_2\|_2}$$

$$\mathcal{L} = \frac{1}{2}\mathcal{D}(p_1, z_2) + \frac{1}{2}\mathcal{D}(p_2, z_1).$$

Topic 1: Self-supervised Learning

Self-distillation

□ Simple Siamese network (SimSiam)*

- Output of the encoder

$$z_1 \triangleq f(x_1) \quad z_2 \triangleq f(x_2)$$

- Output of the predictor

$$p_1 \triangleq h(f(x_1)) \quad p_2 \triangleq h(f(x_2))$$

- Minimize the negative cosine similarity

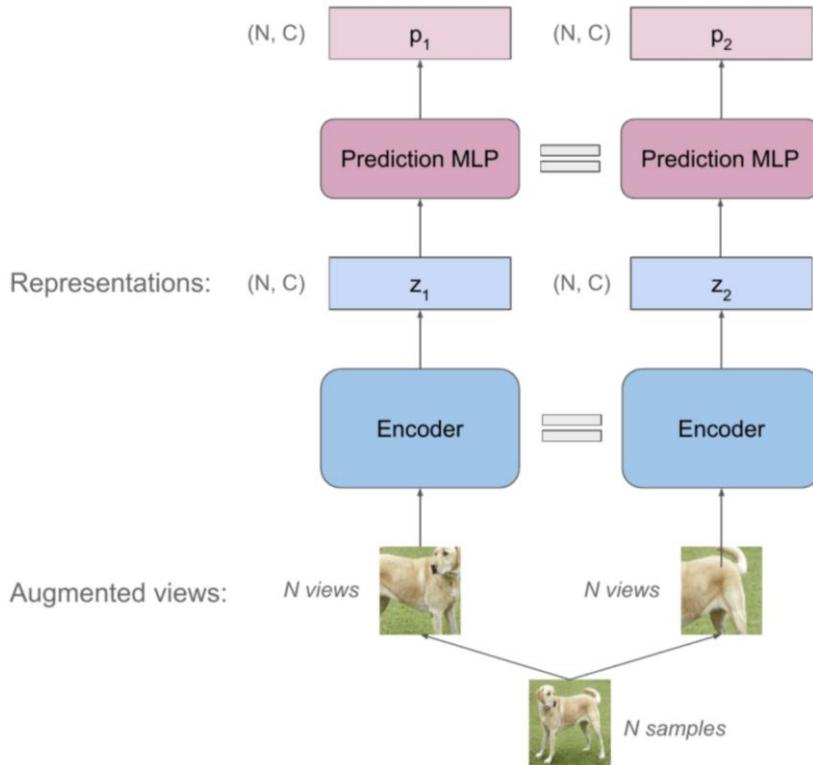
$$\begin{array}{ccc} \mathcal{D}(p_1, z_2) & \xrightarrow{\hspace{1cm}} & \mathcal{L} = \frac{1}{2}\mathcal{D}(p_1, \text{stopgrad}(z_2)) + \frac{1}{2}\mathcal{D}(p_2, \text{stopgrad}(z_1)). \\ \downarrow & & \\ \mathcal{D}(p_1, \text{stopgrad}(z_2)). & & \end{array}$$

Topic 1: Self-supervised Learning

Self-distillation

□ SimSiam*

- Summary



Algorithm 1 SimSiam Pseudocode, PyTorch-like

```

# f: backbone + projection mlp
# h: prediction mlp

for x in loader: # load a minibatch x with n samples
    x1, x2 = aug(x), aug(x) # random augmentation
    z1, z2 = f(x1), f(x2) # projections, n-by-d
    p1, p2 = h(z1), h(z2) # predictions, n-by-d

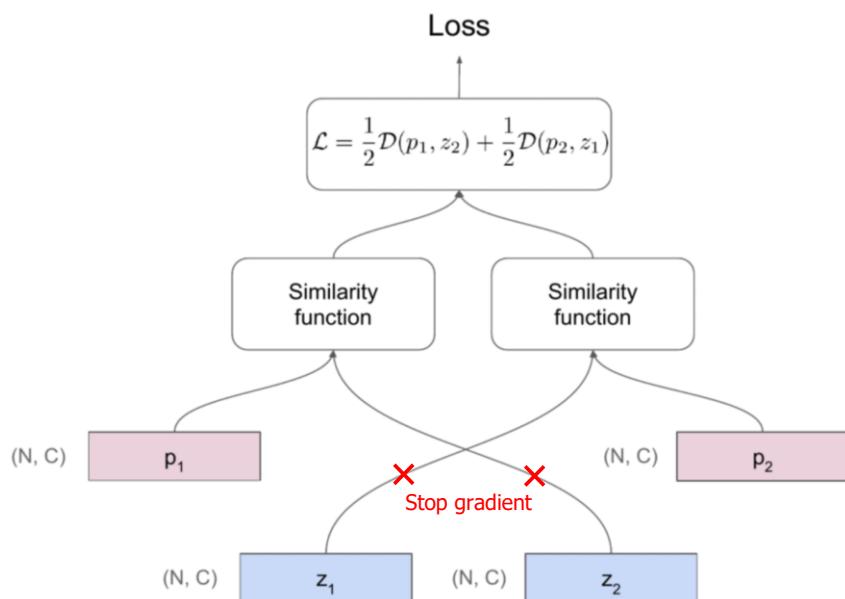
    L = D(p1, z2)/2 + D(p2, z1)/2 # loss

    L.backward() # back-propagate
    update(f, h) # SGD update

def D(p, z): # negative cosine similarity
    z = z.detach() # stop gradient

    p = normalize(p, dim=1) # l2-normalize
    z = normalize(z, dim=1) # l2-normalize
    return -(p*z).sum(dim=1).mean()

```



Topic 1: Self-supervised Learning

Implementation

□ Implementation

● Dataset

- CIFAR-10
 - 4,000 images (labeled)
 - 56,000 image (unlabeled)

● Code

```
parser = argparse.ArgumentParser(description='Train SimSiam')  
parser.add_argument('--feature_dim', default=2048, type=int, help='Feature dim for out vector')  
parser.add_argument('--k', default=200, type=int, help='Top k most similar images used to predict the label')  
parser.add_argument('--batch_size', default=512, type=int, help='Number of images in each mini-batch')  
parser.add_argument('--epochs', default=800, type=int, help='Number of sweeps over the dataset to train')
```

● Data prepare

```
# data prepare  
train_data = CIFAR10Pair(root='data', train=True, transform=train_transform, download=True)  
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True, num_workers=16, pin_memory=True)  
memory_data = CIFAR10Pair(root='data', train=True, transform=test_transform, download=True)  
memory_loader = DataLoader(memory_data, batch_size=batch_size, shuffle=False, num_workers=16, pin_memory=True)  
test_data = CIFAR10Pair(root='data', train=False, transform=test_transform, download=True)  
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False, num_workers=16, pin_memory=True)
```

Topic 1: Self-supervised Learning

Implementation

- Implementation
 - Optimizer

```
# model setup and optimizer config
model = Model(feature_dim).cuda()

optimizer = SGD(model.parameters(), lr=0.03, momentum=0.9, weight_decay=5e-4)
lr_scheduler = LambdaLR(optimizer, lr_lambda=lambda i: 0.5 * (math.cos(i * math.pi / epochs) + 1))
c = len(memory_data.classes)

results = {'train_loss': [], 'test_acc@1': [], 'test_acc@5': []}

# training loop
for epoch in range(1, epochs + 1):
    train_loss = train(model, train_loader, optimizer)

    results['train_loss'].append(train_loss)
    lr_scheduler.step()
    test_acc_1, test_acc_5 = test(model, memory_loader, test_loader)

    results['test_acc@1'].append(test_acc_1)
    results['test_acc@5'].append(test_acc_5)
    # save statistics
    data_frame = pd.DataFrame(data=results, index=range(1, epoch + 1))
    data_frame.to_csv('results/{}_statistics.csv'.format(save_name_pre), index_label='epoch')
    if test_acc_1 > best_acc:
        best_acc = test_acc_1
        torch.save(model.state_dict(), 'results/{}_model.pth'.format(save_name_pre))
```

Algorithm 1 SimSiam Pseudocode, PyTorch-like

```
# f: backbone + projection mlp
# h: prediction mlp

for x in loader: # load a minibatch x with n samples
    x1, x2 = aug(x), aug(x) # random augmentation
    z1, z2 = f(x1), f(x2) # projections, n-by-d
    p1, p2 = h(z1), h(z2) # predictions, n-by-d

    L = D(p1, z2)/2 + D(p2, z1)/2 # loss

    L.backward() # back-propagate
    update(f, h) # SGD update

def D(p, z): # negative cosine similarity
    z = z.detach() # stop gradient

    p = normalize(p, dim=1) # l2-normalize
    z = normalize(z, dim=1) # l2-normalize
    return -(p*z).sum(dim=1).mean()
```

Topic 1: Self-supervised Learning

Implementation

- Implementation
 - Training

```
def train(net, data_loader, train_optimizer):
    net.train()
    total_loss, total_num, train_bar = 0.0, 0, tqdm(data_loader, dynamic_ncols=True)
    for pos_1, pos_2, _ in train_bar:

        pos_1, pos_2 = pos_1.cuda(non_blocking=True), pos_2.cuda(non_blocking=True)

        feature_1, proj_1 = net(pos_1)
        feature_2, proj_2 = net(pos_2)

        # compute loss
        sim_1 = -(F.normalize(proj_1, dim=-1) * F.normalize(feature_2.detach(), dim=-1)).sum(dim=-1).mean()
        sim_2 = -(F.normalize(proj_2, dim=-1) * F.normalize(feature_1.detach(), dim=-1)).sum(dim=-1).mean()

        loss = 0.5 * sim_1 + 0.5 * sim_2

        train_optimizer.zero_grad()
        loss.backward()
        train_optimizer.step()
```

Topic 1: Self-supervised Learning

Implementation

□ Implementation

● Results

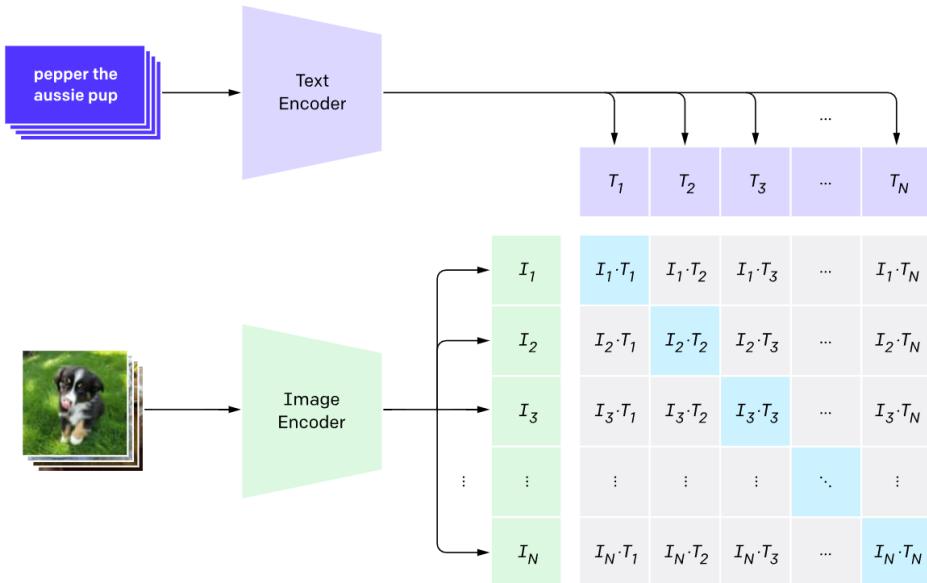
```
Feature extracting: 100%|██████████| 98/98 [00:46<00:00,  2.09it/s]
 5%|█         | 1/20 [00:35<11:16, 35.63s/it]Test Epoch: [470/800] Acc@1:89.06% Acc@5:99.02%
Test Epoch: [470/800] Acc@1:88.28% Acc@5:99.22%
15%|█       | 3/20 [00:35<02:37,  9.27s/it]Test Epoch: [470/800] Acc@1:87.63% Acc@5:99.41%
Test Epoch: [470/800] Acc@1:87.06% Acc@5:99.32%
Test Epoch: [470/800] Acc@1:86.88% Acc@5:99.30%
Test Epoch: [470/800] Acc@1:86.98% Acc@5:99.22%
30%|██      | 6/20 [00:35<00:50,  3.61s/it]Test Epoch: [470/800] Acc@1:86.75% Acc@5:99.16%
Test Epoch: [470/800] Acc@1:86.72% Acc@5:99.15%
45%|███     | 9/20 [00:35<00:21,  1.94s/it]Test Epoch: [470/800] Acc@1:87.00% Acc@5:99.13%
Test Epoch: [470/800] Acc@1:86.64% Acc@5:99.16%
Test Epoch: [470/800] Acc@1:86.70% Acc@5:99.20%
60%|███     | 12/20 [00:36<00:09,  1.18s/it]Test Epoch: [470/800] Acc@1:86.70% Acc@5:99.17%
Test Epoch: [470/800] Acc@1:86.79% Acc@5:99.19%
Test Epoch: [470/800] Acc@1:86.73% Acc@5:99.15%
75%|██████  | 15/20 [00:36<00:03,  1.31it/s]Test Epoch: [470/800] Acc@1:86.73% Acc@5:99.17%
Test Epoch: [470/800] Acc@1:86.74% Acc@5:99.18%
Test Epoch: [470/800] Acc@1:86.70% Acc@5:99.18%
90%|███████ | 18/20 [00:36<00:01,  1.95it/s]Test Epoch: [470/800] Acc@1:86.87% Acc@5:99.16%
Test Epoch: [470/800] Acc@1:86.82% Acc@5:99.18%
Test Epoch: [470/800] Acc@1:86.78% Acc@5:99.20%
100%|████████| 20/20 [00:36<00:00,  1.85s/it]
```

Topic 1: Self-supervised Learning

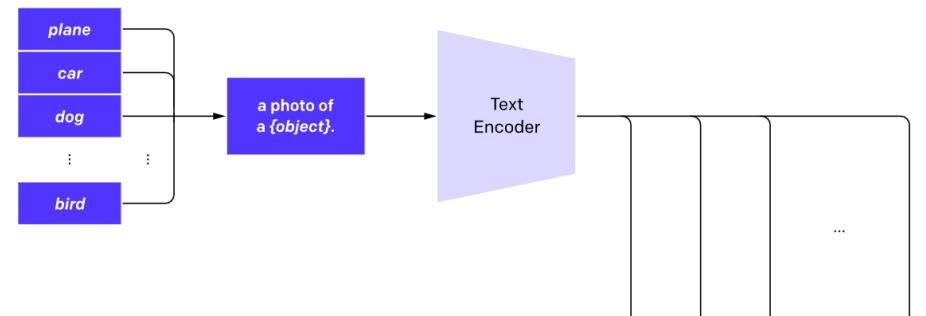
Extension

□ CLIP - Contrastive Language-Image Pre-training

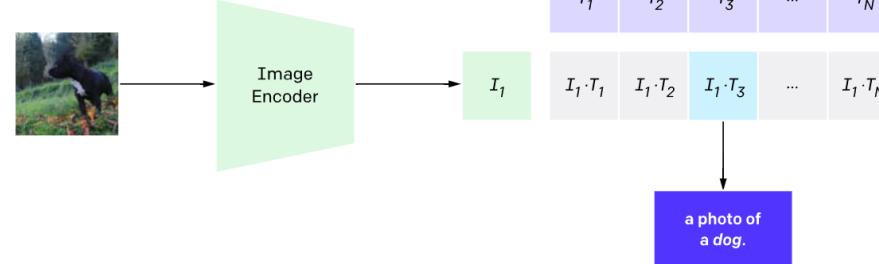
1. Contrastive pre-training



2. Create dataset classifier from label text



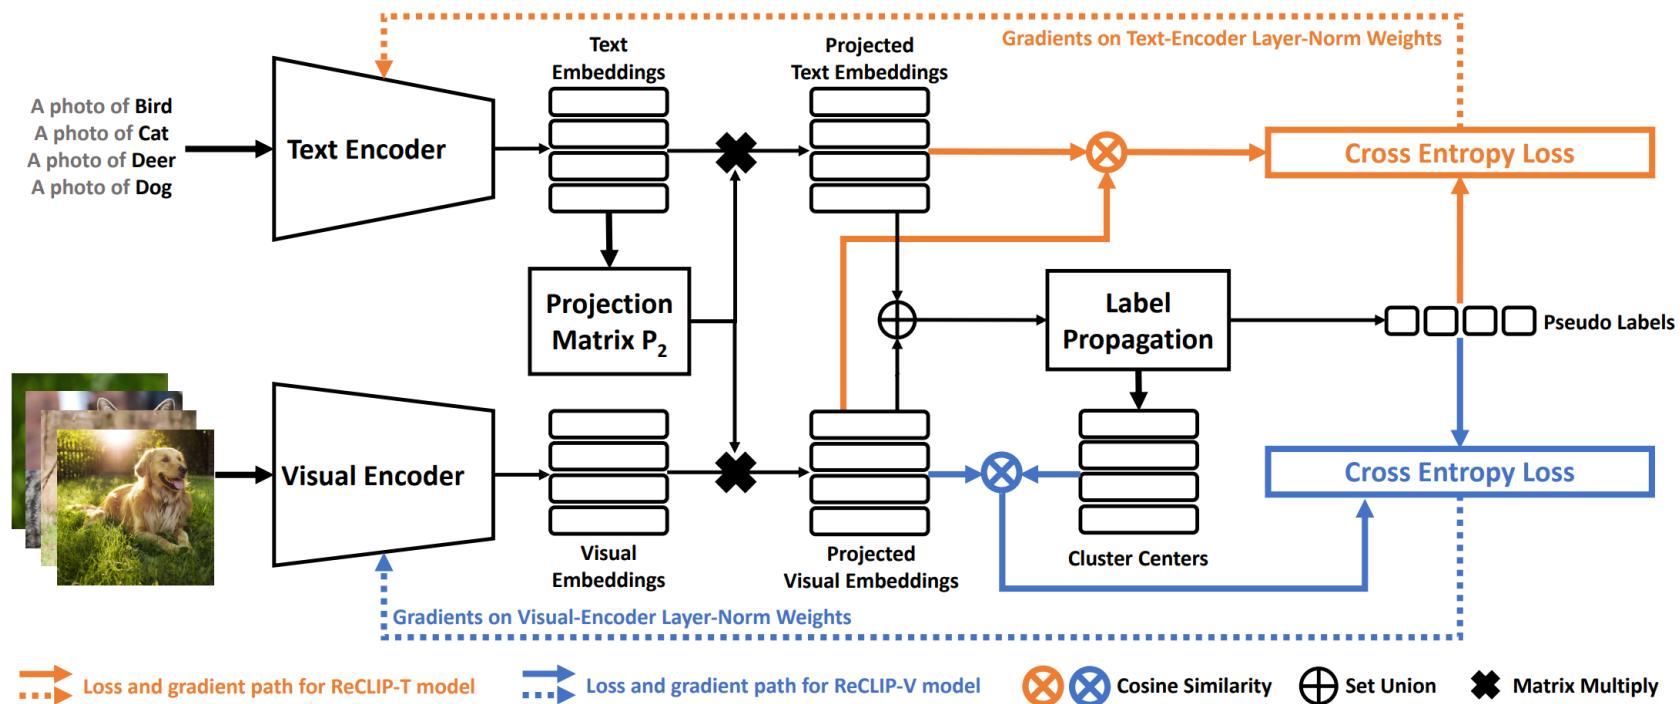
3. Use for zero-shot prediction



Topic 1: Self-supervised Learning

Extension

□ ReCLIP

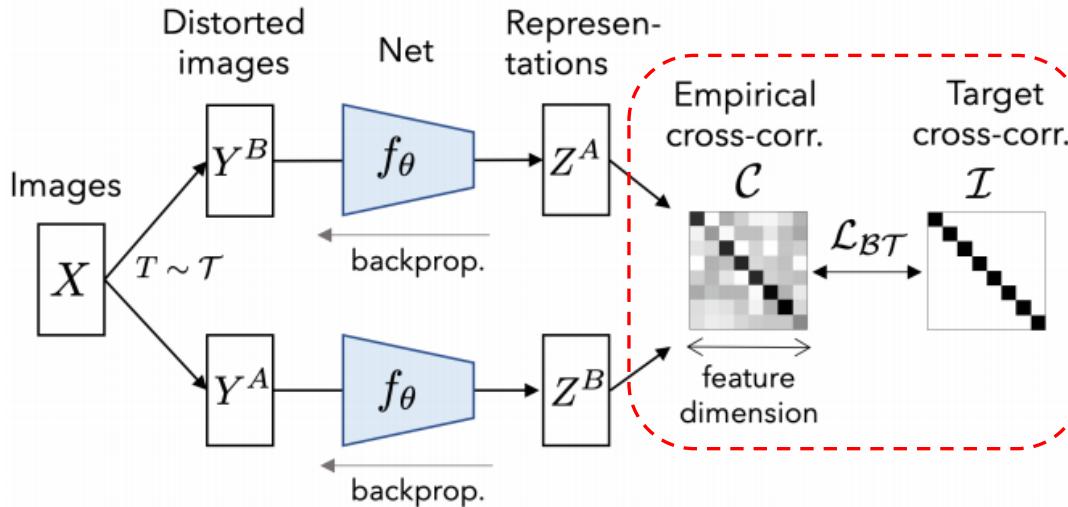


Topic 1: Self-supervised Learning

Method

Canonical Correlation Analysis

- Barlow Twins*
- Self-Supervised Learning via Redundancy Reduction
 - Architecture



Algorithm 1 PyTorch-style pseudocode for Barlow Twins.

```
# f: encoder network
# lambda: weight on the off-diagonal terms
# N: batch size
# D: dimensionality of the embeddings
#
# mm: matrix-matrix multiplication
# off_diagonal: off-diagonal elements of a matrix
# eye: identity matrix

for x in loader: # load a batch with N samples
    # two randomly augmented versions of x
    y_a, y_b = augment(x)

    # compute embeddings
    z_a = f(y_a) # NxD
    z_b = f(y_b) # NxD

    # normalize repr. along the batch dimension
    z_a_norm = (z_a - z_a.mean(0)) / z_a.std(0) # NxD
    z_b_norm = (z_b - z_b.mean(0)) / z_b.std(0) # NxD

    # cross-correlation matrix
    c = mm(z_a_norm.T, z_b_norm) / N # DxD

    # loss
    c_diff = (c - eye(D)).pow(2) # DxD
    # multiply off-diagonal elems of c_diff by lambda
    off_diagonal(c_diff).mul_(lambda)
    loss = c_diff.sum()

    # optimization step
    loss.backward()
    optimizer.step()
```

Topic 1: Self-supervised Learning

Method

Canonical Correlation Analysis

□ Barlow Twins*

● Self-Supervised Learning via Redundancy Reduction

▪ Operation

– Invariance term

» Make the embedding from different image views

(augmentation versions) of the same image close each other

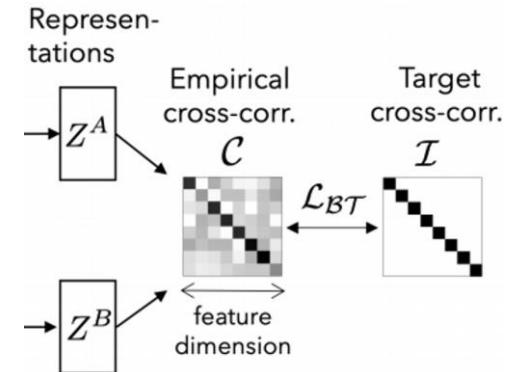
– Redundancy reduction term

$$\mathcal{C}_{ij} \triangleq \frac{\sum_b z_{b,i}^A z_{b,j}^B}{\sqrt{\sum_b (z_{b,i}^A)^2} \sqrt{\sum_b (z_{b,j}^B)^2}}$$

$$\mathcal{L}_{BT} \triangleq \underbrace{\sum_i (1 - \mathcal{C}_{ii})^2}_{\text{invariance term}} + \lambda \underbrace{\sum_i \sum_{j \neq i} \mathcal{C}_{ij}^2}_{\text{redundancy reduction term}}$$

Deal with the on-diagonal terms
As close as possible to one

Deal with the off-diagonal terms
As close as possible to zero

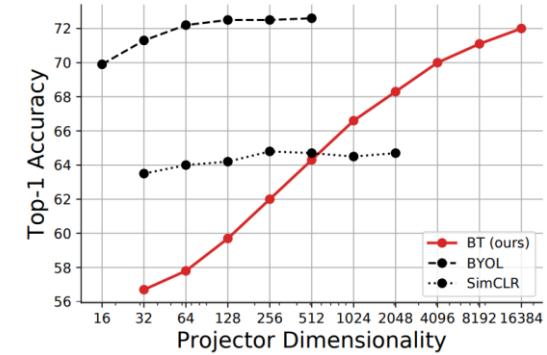
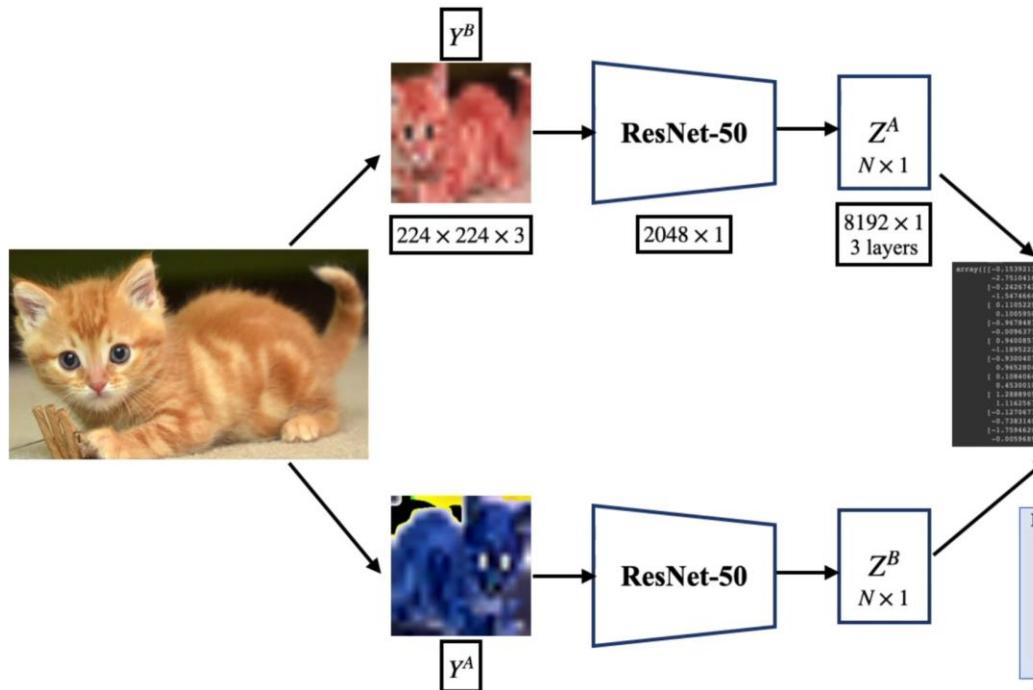


Topic 1: Self-supervised Learning

Method

Canonical Correlation Analysis

- Barlow Twins*
 - Self-Supervised Learning via Redundancy Reduction



```
Encoder = nn.Sequential(  
    resnet50(), # without the final fc layer  
    nn.Linear(2048, 8192), nn.BatchNorm1d(8192), nn.ReLU(),  
    nn.Linear(8192, 8192), nn.BatchNorm1d(8192), nn.ReLU(),  
    nn.Linear(8192, 8192))
```

January 8, 2024

Module

Semi-supervise Learning (SSL)

Ph.D. Ba Hung Ngo

Department of Multimedia Engineering

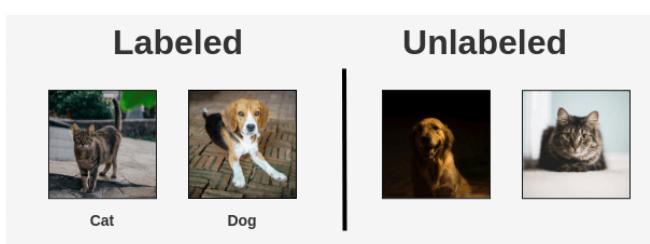
Dongguk University

Topic 2: Semi-supervised Learning

Method

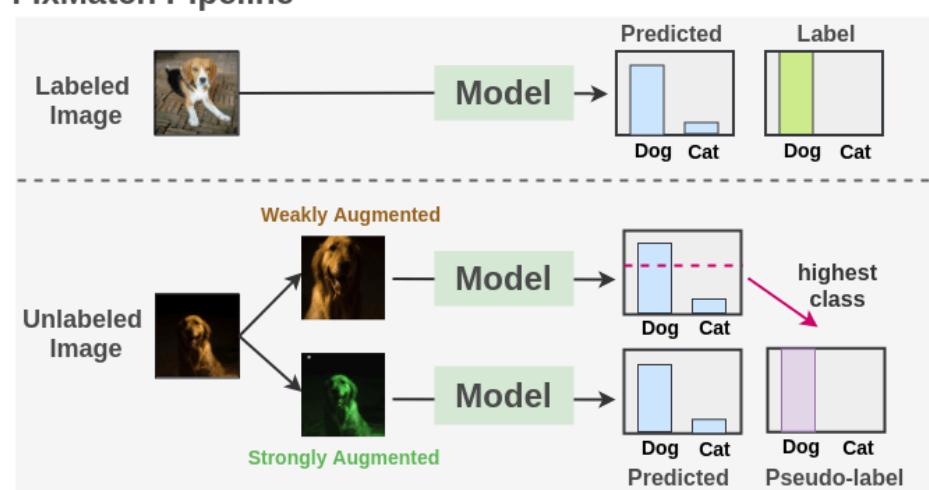
FixMatch*

Data Setting:



Training Procedure:

FixMatch Pipeline

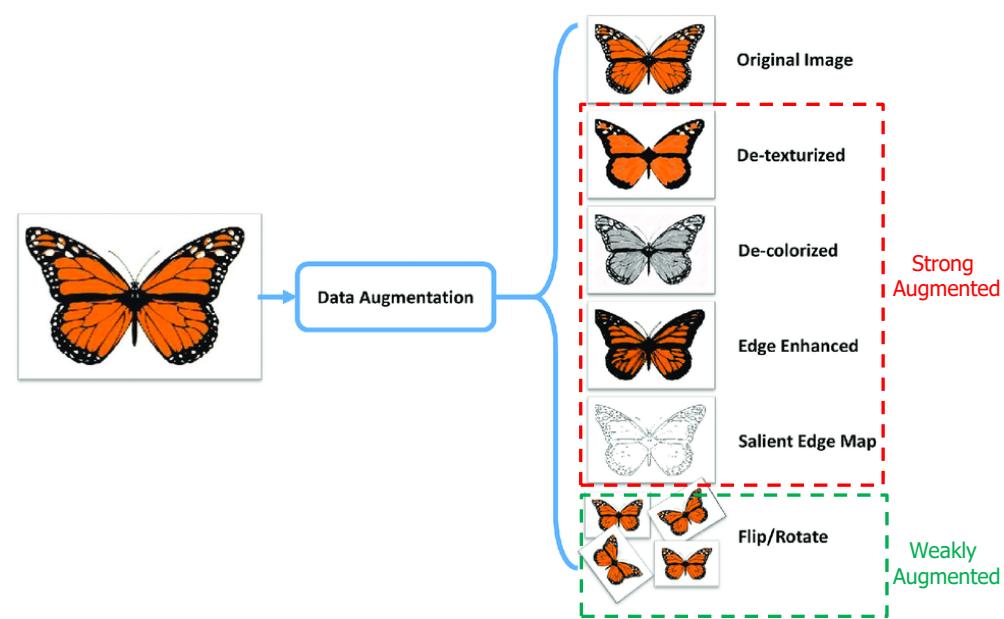
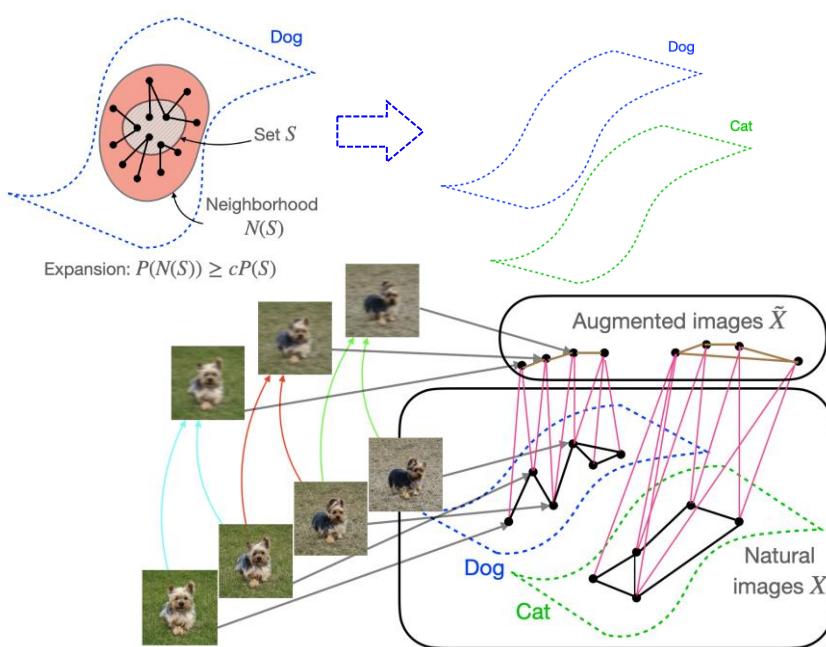


Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Data Augmentation



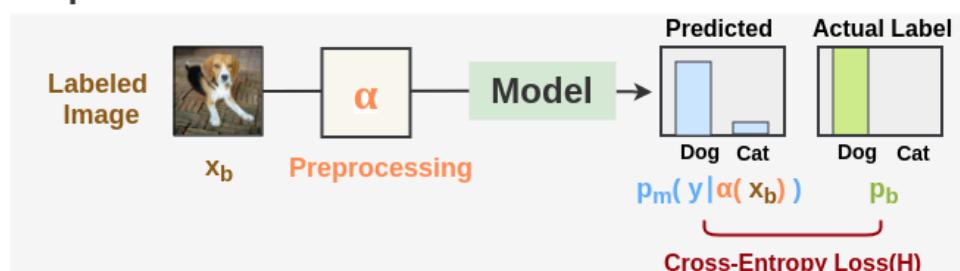
Topic 2: Semi-supervised Learning

Method

FixMatch*

- Model Training and Loss Functions
 - Supervised Learning

Supervised Part of FixMatch



$$l_s = \frac{1}{B} \sum_{b=1}^B H(p_b, p_m(y | \alpha(x_b)))$$

Topic 2: Semi-supervised Learning

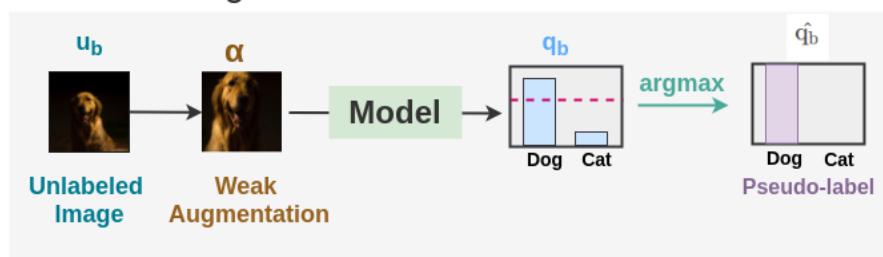
Method

FixMatch*

□ Model Training and Loss Functions

- Unsupervised Learning
 - Pseudo labeling

Pseudo-label generation

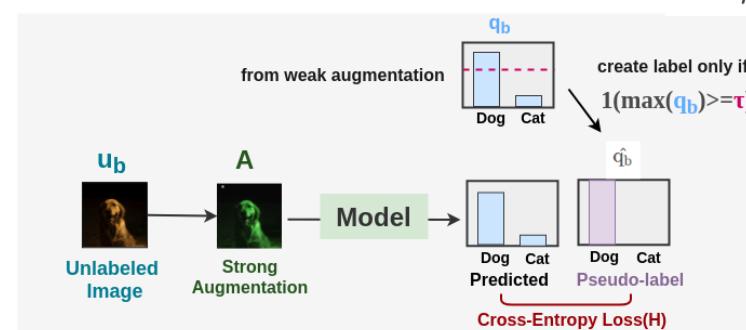


$$q_b = p_m(y|\alpha(u_b))$$

$$\hat{q}_b = \text{argmax}(q_b)$$

- Consistency regularization

Consistency Regularization



$$l_u = \frac{1}{\mu B} \sum_{b=1}^{\mu B} 1(\max(q_b) \geq \tau) H(\hat{q}_b, p_m(y|A(u_b)))$$

Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Implementation

● Dataset

- CIFAR-10
 - 60,000 color images
 - 10 classes
 - 50,000 training images (labeled)
 - 10,000 training images (unlabeled)

airplane



automobile



bird



cat



deer



dog



frog



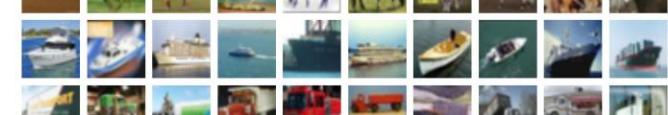
horse



ship



truck



Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Implementation

● Dataset

- CIFAR-10
 - 4,000 images (labeled)
 - 56,000 image (unlabeled)

● Code

```
if args.dataset == 'cifar10':  
    args.num_classes = 10  
    if args.arch == 'wideresnet':  
        args.model_depth = 28  
        args.model_width = 2  
    elif args.arch == 'resnext':  
        args.model_cardinality = 4  
        args.model_depth = 28  
        args.model_width = 4
```

```
elif args.dataset == 'cifar100':  
    args.num_classes = 100  
    if args.arch == 'wideresnet':  
        args.model_depth = 28  
        args.model_width = 8  
    elif args.arch == 'resnext':  
        args.model_cardinality = 8  
        args.model_depth = 29  
        args.model_width = 64
```

Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Implementation

- Data preparation

```
DATASET_GETTERS = {'cifar10': get_cifar10, 'cifar100': get_cifar100}
```

- Data Loader

```
def get_cifar10(args, root):  
  
    transform_labeled = transforms.Compose([  
        transforms.RandomHorizontalFlip(),  
        transforms.RandomCrop(size=32, padding=int(32 * 0.125), padding_mode='reflect'),  
        transforms.ToTensor(),  
        transforms.Normalize(mean=cifar10_mean, std=cifar10_std)  
    ])  
  
    transform_val = transforms.Compose([  
        transforms.ToTensor(),  
        transforms.Normalize(mean=cifar10_mean, std=cifar10_std)  
    ])  
    base_dataset = datasets.CIFAR10(root, train=True, download=True)  
  
    train_labeled_ids, train_unlabeled_ids = x_u_split(args, base_dataset.targets)  
  
    train_labeled_dataset = CIFAR10SSL(root, train_labeled_ids, train=True, transform=transform_labeled)  
  
    train_unlabeled_dataset = CIFAR10SSL(root, train_unlabeled_ids, train=True,  
                                         transform=TransformFixMatch(mean=cifar10_mean, std=cifar10_std))  
  
    test_dataset = datasets.CIFAR10(root, train=False, transform=transform_val, download=False)  
  
    return train_labeled_dataset, train_unlabeled_dataset, test_dataset
```

Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Implementation

- Data transformation

```
class TransformFixMatch(object):
    def __init__(self, mean, std):
        self.weak = transforms.Compose([
            transforms.RandomHorizontalFlip(),
            transforms.RandomCrop(size=32, padding=int(32 * 0.125), padding_mode='reflect')])

        self.strong = transforms.Compose([
            transforms.RandomHorizontalFlip(),
            transforms.RandomCrop(size=32, padding=int(32 * 0.125), padding_mode='reflect'),
            RandAugmentMC(n=2, m=10)])

        self.normalize = transforms.Compose([
            transforms.ToTensor(),
            transforms.Normalize(mean=mean, std=std)])

    def __call__(self, x):
        weak = self.weak(x)
        strong = self.strong(x)
        return self.normalize(weak), self.normalize(strong)
```

Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Implementation

- Model selection

```
def main():
    parser = argparse.ArgumentParser(description='PyTorch FixMatch Training')
    parser.add_argument('--gpu-id', default='0', type=int, help='id(s) for CUDA_VISIBLE_DEVICES')
    parser.add_argument('--num-workers', type=int, default=4, help='number of workers')
    parser.add_argument('--dataset', default='cifar10', type=str, choices=['cifar10', 'cifar100'], help='dataset name')
    parser.add_argument('--num-labeled', type=int, default=4000, help='number of labeled data')
    parser.add_argument("--expand-labels", action="store_true", help="expand labels to fit eval steps")
    parser.add_argument('--arch', default='wideresnet', type=str, choices=['wideresnet', 'resnext'], help='dataset name')
```

```
def create_model(args):
    if args.arch == 'wideresnet':
        import models.wideresnet as models
        model = models.build_wideresnet(depth=args.model_depth, widen_factor=args.model_width, dropout=0, num_classes=args.num_classes)
    elif args.arch == 'resnext':
        import models.resnext as models
        model = models.build_resnext(cardinality=args.model_cardinality, depth=args.model_depth, width=args.model_width, num_classes=args.num_classes)
```

Topic 2: Semi-supervised Learning

FixMatch*

- Implementation
 - Model

```
def forward(self, x):  
    x = self.conv_1_3x3.forward(x)  
    x = self.act(self.bn_1.forward(x))  
    x = self.stage_1.forward(x)  
    x = self.stage_2.forward(x)  
    x = self.stage_3.forward(x)  
    x = F.adaptive_avg_pool2d(x, 1)  
    x = x.view(-1, self.stages[3])  
    return self.classifier(x)
```

```
class CifarResNeXt(nn.Module):  
    """  
    ResNext optimized for the Cifar dataset, as specified in  
    https://arxiv.org/pdf/1611.05431.pdf  
    """  
  
    def __init__(self, cardinality, depth, num_classes,  
                 base_width, widen_factor=4):  
        """ Constructor  
        Args:  
            cardinality: number of convolution groups.  
            depth: number of layers.  
            nlabels: number of classes  
            base_width: base number of channels in each group.  
            widen_factor: factor to adjust the channel dimensionality  
        """  
        super().__init__()  
        self.cardinality = cardinality  
        self.depth = depth  
        self.block_depth = (self.depth - 2) // 9  
        self.base_width = base_width  
        self.widen_factor = widen_factor  
        self.nlabels = num_classes  
        self.output_size = 64  
        self.stages = [64, 64 * self.widen_factor, 128 *  
                      self.widen_factor, 256 * self.widen_factor]  
  
        self.conv_1_3x3 = nn.Conv2d(3, 64, 3, 1, 1, bias=False)  
        self.bn_1 = nn.BatchNorm2d(64, momentum=0.001)  
        self.act = mish  
        self.stage_1 = self.block('stage_1', self.stages[0], self.stages[1], 1)  
        self.stage_2 = self.block('stage_2', self.stages[1], self.stages[2], 2)  
        self.stage_3 = self.block('stage_3', self.stages[2], self.stages[3], 2)  
        self.classifier = nn.Linear(self.stages[3], num_classes)
```

Method

Topic 2: Semi-supervised Learning

Method

FixMatch*

□ Implementation

- Load data

```
labeled_dataset, unlabeled_dataset, test_dataset = DATASET_GETTERS[args.dataset](args, './data')
```

```
labeled_trainloader = DataLoader(  
    labeled_dataset,  
    sampler=train_sampler(labeled_dataset),  
    batch_size=args.batch_size,  
    num_workers=args.num_workers,  
    drop_last=True)  
  
unlabeled_trainloader = DataLoader(  
    unlabeled_dataset,  
    sampler=train_sampler(unlabeled_dataset),  
    batch_size=args.batch_size * args.mu,  
    num_workers=args.num_workers,  
    drop_last=True)  
  
test_loader = DataLoader(  
    test_dataset,  
    sampler=SequentialSampler(test_dataset),  
    batch_size=args.batch_size,  
    num_workers=args.num_workers)
```

Topic 2: Semi-supervised Learning

Method

FixMatch*

- Implementation
 - Training

```
train(args, labeled_trainloader, unlabeled_trainloader, test_loader, model, optimizer, ema_model, scheduler)
```

```
2 usages (1 dynamic)
def train(args, labeled_trainloader, unlabeled_trainloader, test_loader, model, optimizer, ema_model, scheduler):
    if args.amp:
        from apex import amp
    global best_acc
    test_accs = []
    end = time.time()

    if args.world_size > 1:
        labeled_epoch = 0
        unlabeled_epoch = 0
        labeled_trainloader.sampler.set_epoch(labeled_epoch)
        unlabeled_trainloader.sampler.set_epoch(unlabeled_epoch)

    labeled_iter = iter(labeled_trainloader)
    unlabeled_iter = iter(unlabeled_trainloader)
```

```
for batch_idx in range(args.eval_step):
    # read labeled data
    try:
        inputs_x, targets_x = labeled_iter.next()
    except:
        if args.world_size > 1:
            labeled_epoch += 1
            labeled_trainloader.sampler.set_epoch(labeled_epoch)
        labeled_iter = iter(labeled_trainloader)
        inputs_x, targets_x = labeled_iter.next()

    # read unlabeled data
    try:
        (inputs_u_w, inputs_u_s), _ = unlabeled_iter.next()
    except:
        if args.world_size > 1:
            unlabeled_epoch += 1
            unlabeled_trainloader.sampler.set_epoch(unlabeled_epoch)
        unlabeled_iter = iter(unlabeled_trainloader)
        (inputs_u_w, inputs_u_s), _ = unlabeled_iter.next()
```

Topic 2: Semi-supervised Learning

Method

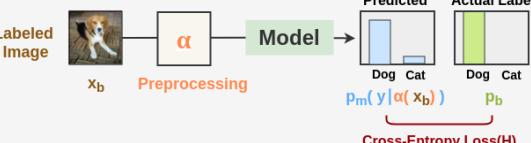
FixMatch*

- Implementation
 - Training
 - On labeled data

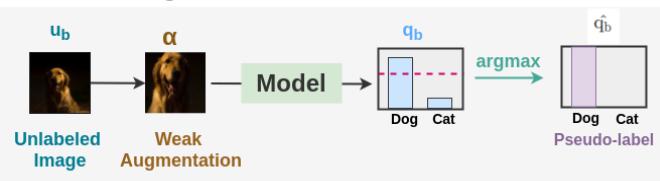
```
representation_labeled_data, _ = model(inputs_x)
representation_unlabeled_data_weak, _ = model(inputs_u_w)
representation_unlabeled_data_strong, _ = model(inputs_u_s)

# extract representation of unlabeled data
logits_labeled_data = model.classifier(representation_labeled_data)
logits_unlabeled_data_weak = model.classifier(representation_unlabeled_data_weak)
logits_unlabeled_data_strong = model.classifier(representation_unlabeled_data_strong)
```

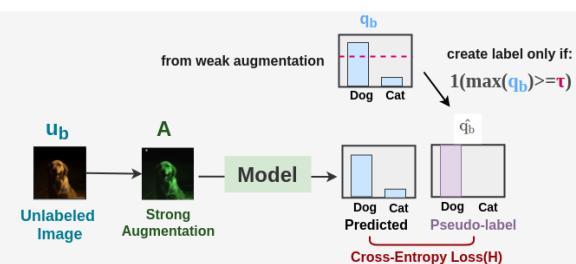
Supervised Part of FixMatch



Pseudo-label generation



Consistency Regularization



```
# Cross-entropy loss on labeled data
Lx = F.cross_entropy(logits_labeled_data, targets_x, reduction='mean')

# Semi-supervised loss on unlabeled data
pseudo_label = torch.softmax(logits_unlabeled_data_weak.detach() / args.T, dim=-1)
max_probs, targets_u = torch.max(pseudo_label, dim=-1)
mask = max_probs.ge(args.threshold).float()

Lu = (F.cross_entropy(logits_unlabeled_data_strong, targets_u, reduction='none') * mask).mean()

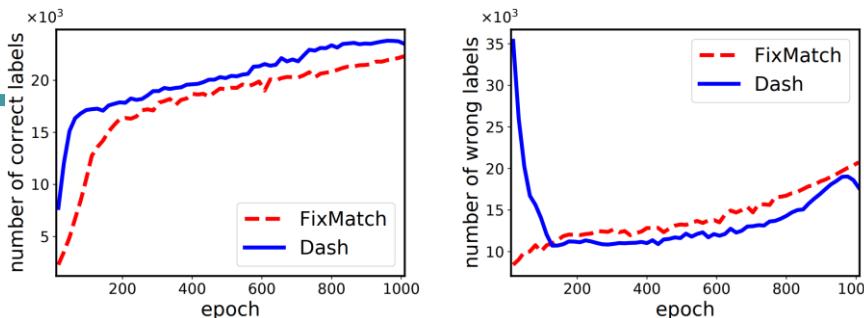
loss = Lx + args.lambda_u * Lu
```

Method

Topic 2: Semi-supervised Learning

DASH*

□ Semi-Supervised Learning with Dynamic Thresholding



Algorithm 1 Dash: Semi-Supervised Learning with Dynamic Thresholding

Input: learning rate η_0 and mini-batch size m_0 for stage one, learning rate η and parameter m of mini-batch size for stage two, two parameters $C > 1$ and $\gamma > 1$ for computing threshold, and violation probability δ .

// Warm-up Stage: run SGD in T_0 iterations.

Initialization: $\mathbf{u}_0 = \mathbf{w}_0$

for $t = 0, 1, \dots, T_0 - 1$ **do**

 Sample m_0 examples $\xi_{t,i}$ ($i = 1, \dots, m_0$) from \mathbf{D}_l ,

$\mathbf{u}_{t+1} = \mathbf{u}_t - \eta_0 \tilde{\mathbf{g}}_t$ where $\tilde{\mathbf{g}}_t = \frac{1}{m_0} \sum_{i=1}^{m_0} \nabla f_s(\mathbf{u}_t; \xi_{t,i})$

end for

// Selection Stage: run SGD in T iterations.

Initialization: $\mathbf{w}_1 = \mathbf{u}_{T_0}$.

Compute the value of $\hat{\rho}$ as in (16). // In practice, $\hat{\rho}$ can be obtained as in (17).

for $t = 1, \dots, T$ **do**

 1) Sample $n_t = m\gamma^{t-1}$ examples from \mathbf{D}_u , where the pseudo labels in \mathbf{D}_u are generated by FixMatch

 2) Set the threshold $\rho_t = C\gamma^{-(t-1)}\hat{\rho}$.

 3) Compute truncated stochastic gradient \mathbf{g}_t as (18).

 4) Update solution by SGD using stochastic gradient \mathbf{g}_t and learning rate η : $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \mathbf{g}_t$.

end for

Output: \mathbf{w}_{T+1}

$$\hat{\rho} \approx \frac{1}{|\mathbf{D}_l|} \sum_{\xi_i \in \mathbf{D}_l} f(\mathbf{w}_1; \xi_i), \quad (17)$$

THANK YOU