

AI VIETNAM
All-in-One Course
(TA Session)

Softmax Regression

Exercise



AI VIET NAM
[@aivietnam.edu.vn](http://aivietnam.edu.vn)

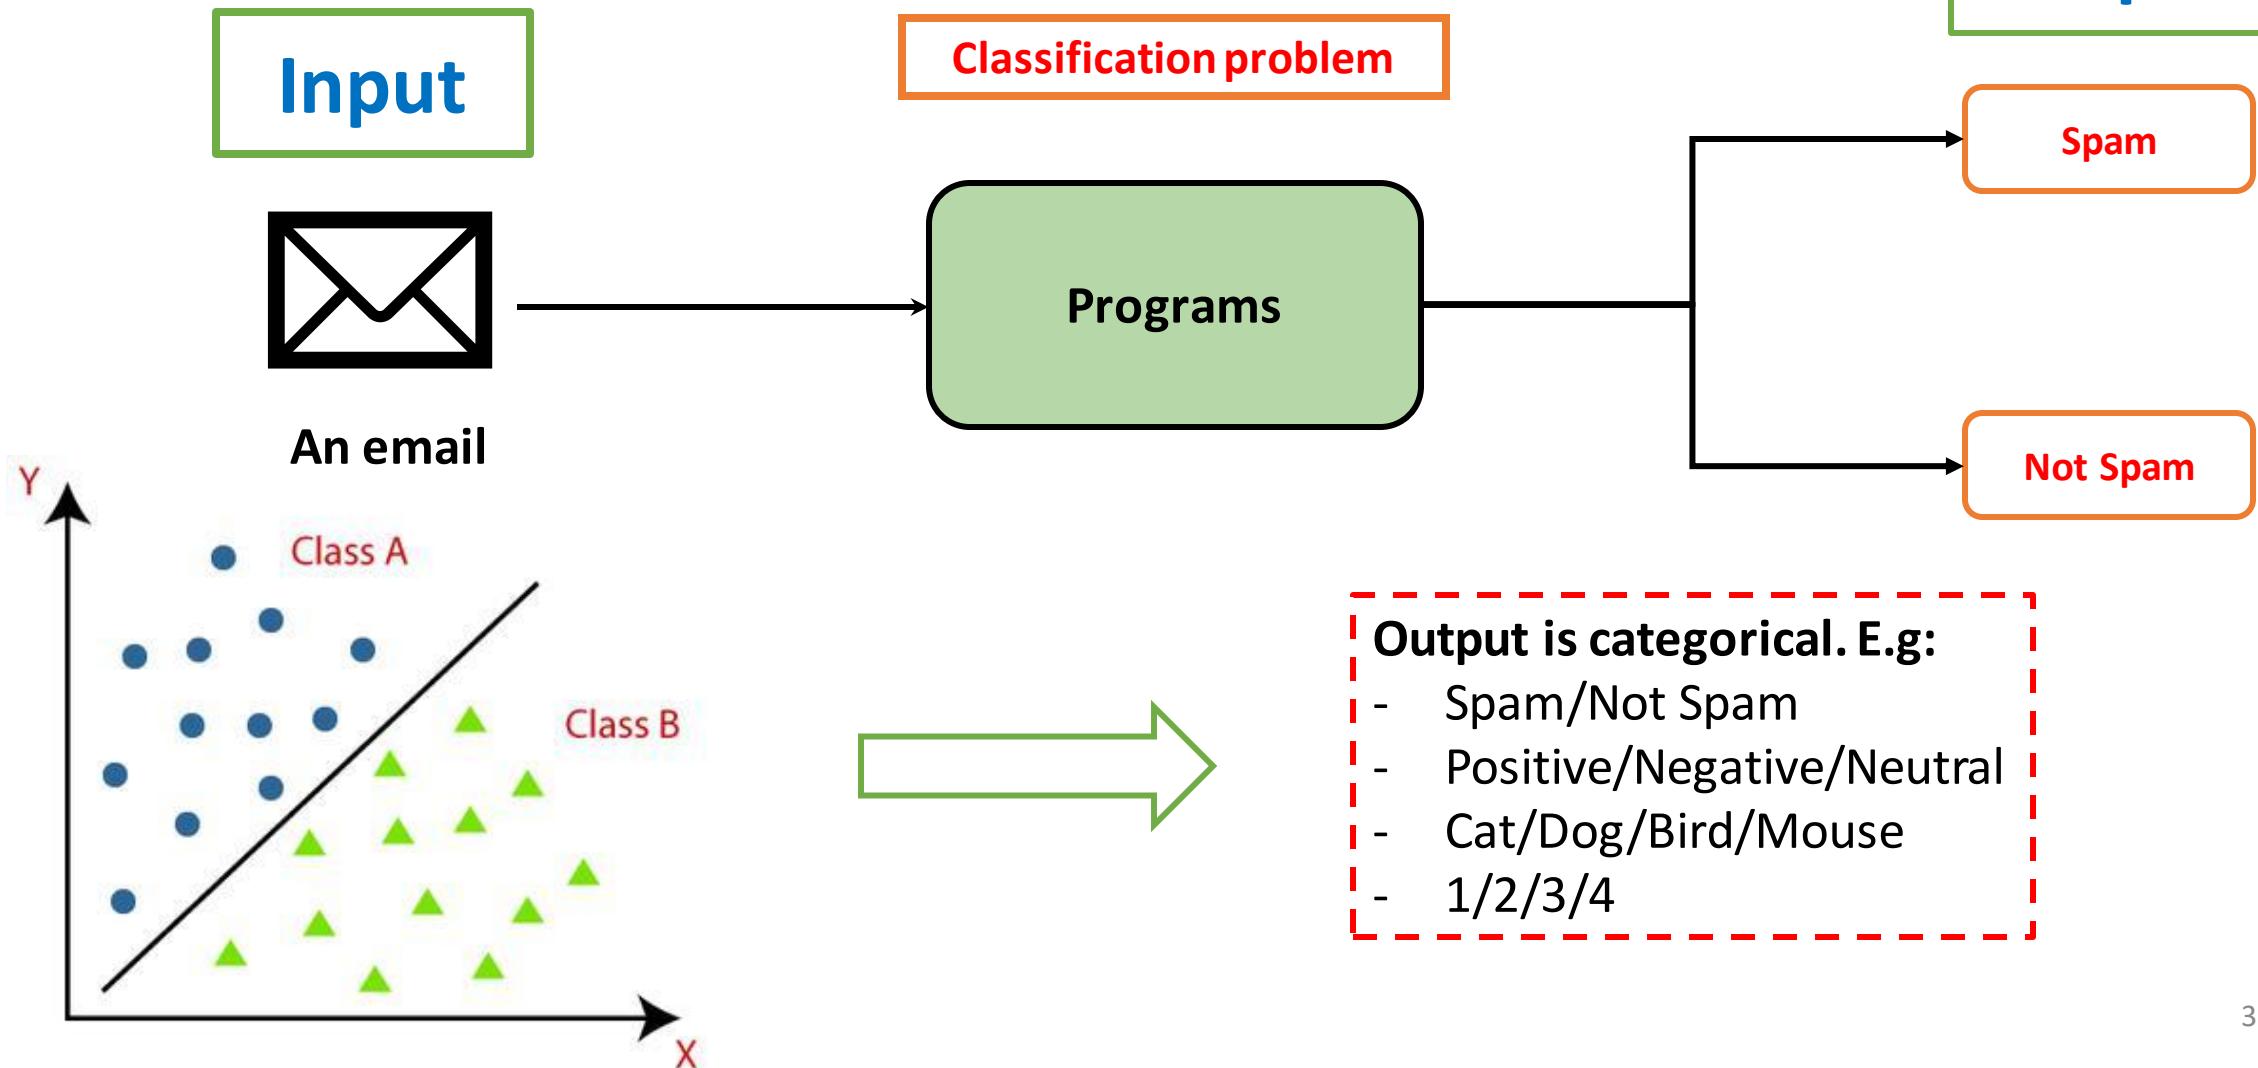
Dinh-Thang Duong – TA

Outline

- Review
- Card Fraud Detection
- Sentiment Analysis
- Question

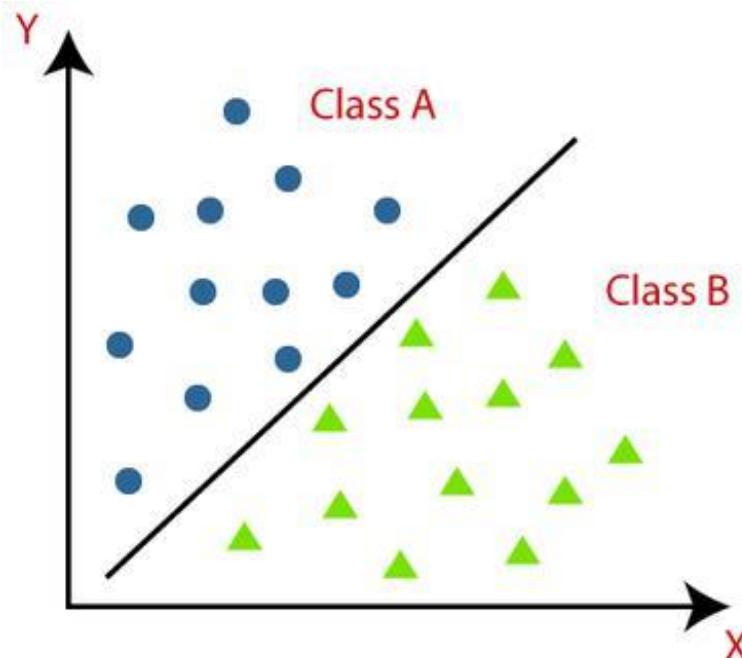
Review

❖ Getting Started

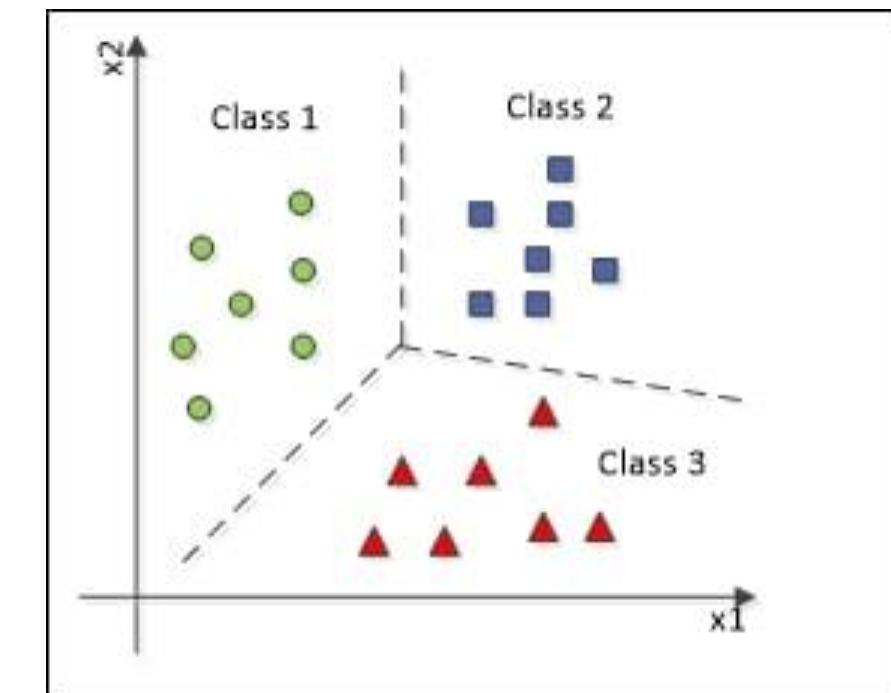


Review

❖ Getting Started



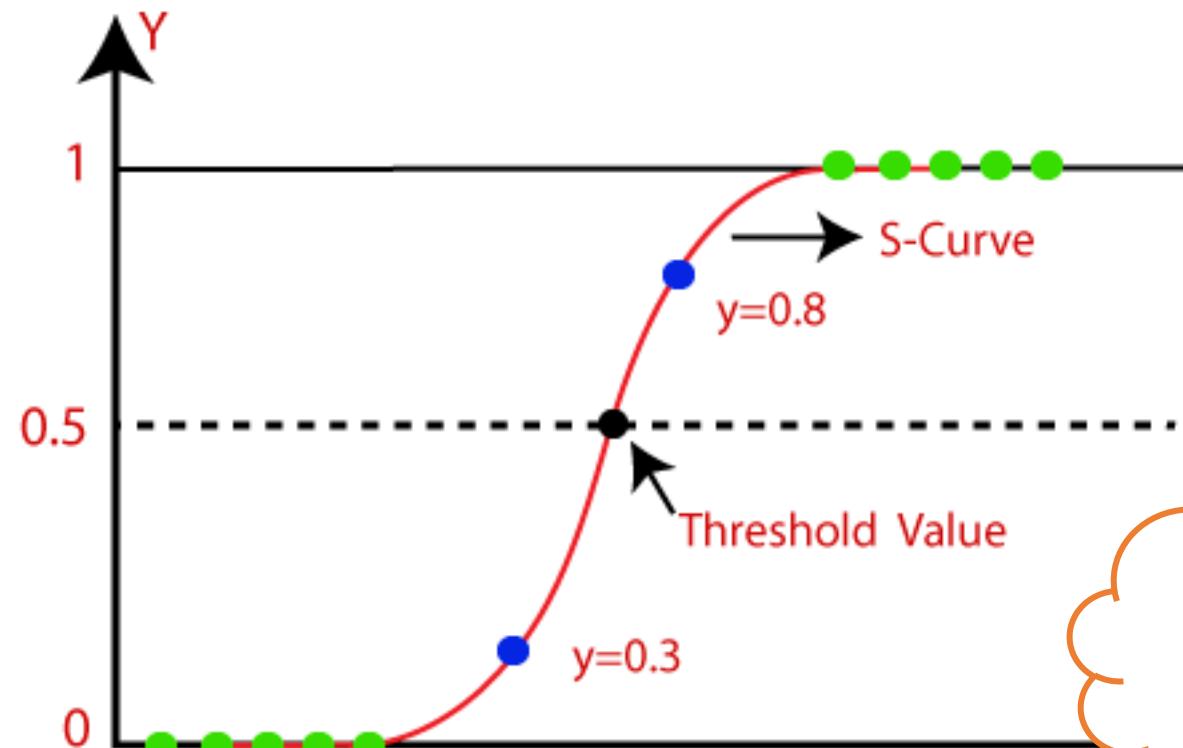
Binary Classification (n_classes=2)



Multi-class Classification (n_classes >= 2)

Review

❖ Logistic Regression for multiclass?

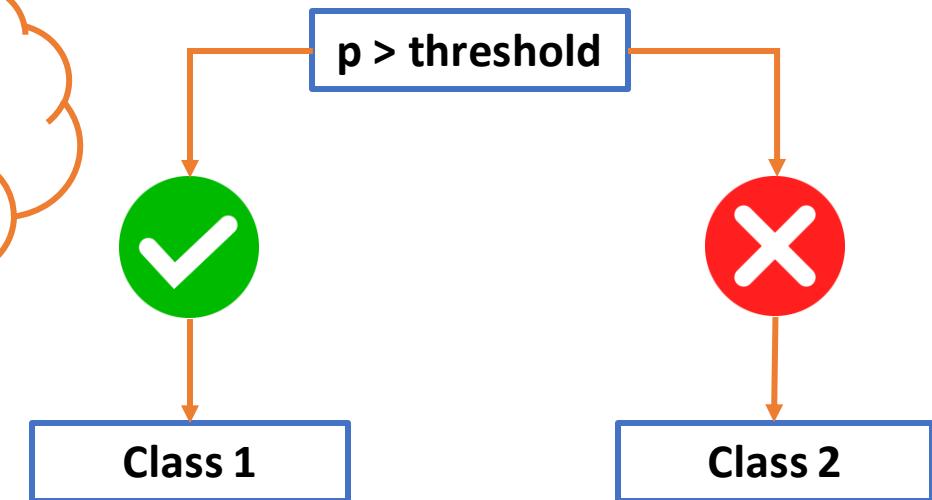


Sigmoid function

Range value: (0, 1)

$$z = b_0 + \sum_{i=1}^{n_features} b_i x_i$$
$$\text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$$

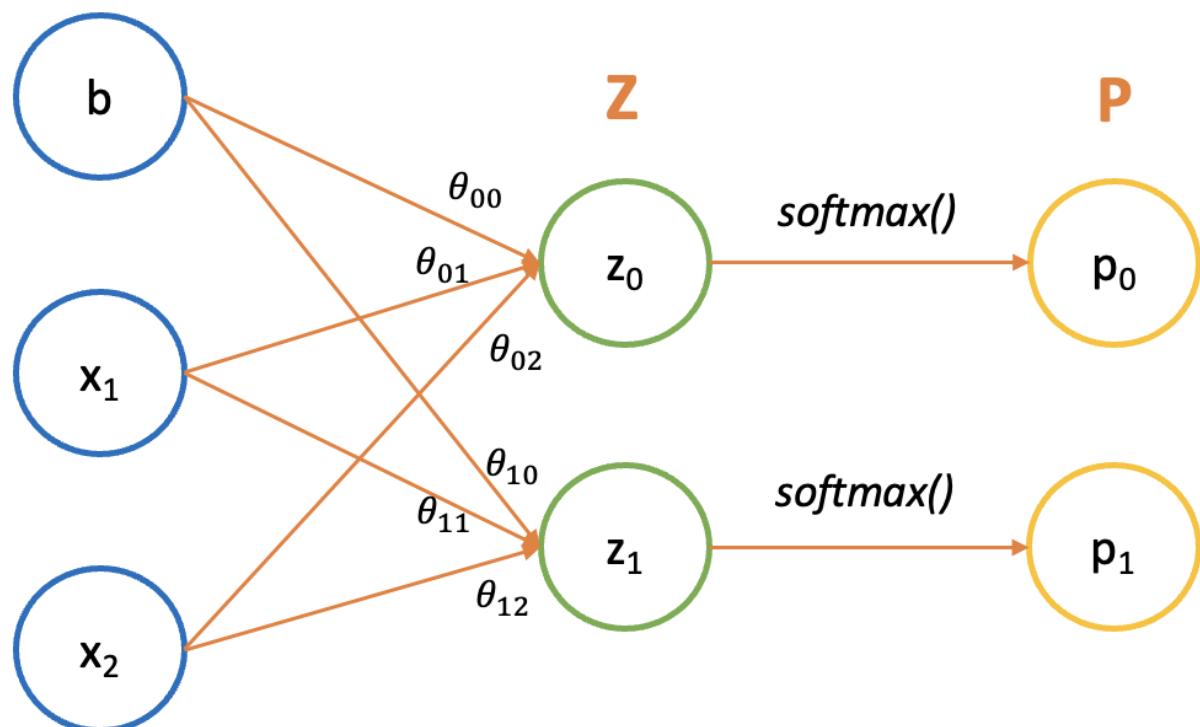
How about
multiclass?



Review

❖ Introduction

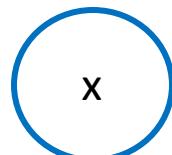
Input

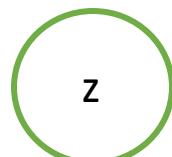


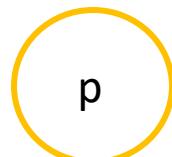
Softmax Regression (Multinomial Logistic Regression): A supervised ML algorithm used for classification tasks. It is a generalization of Logistic Regression to the case where we want to handle multiple classes.

Review

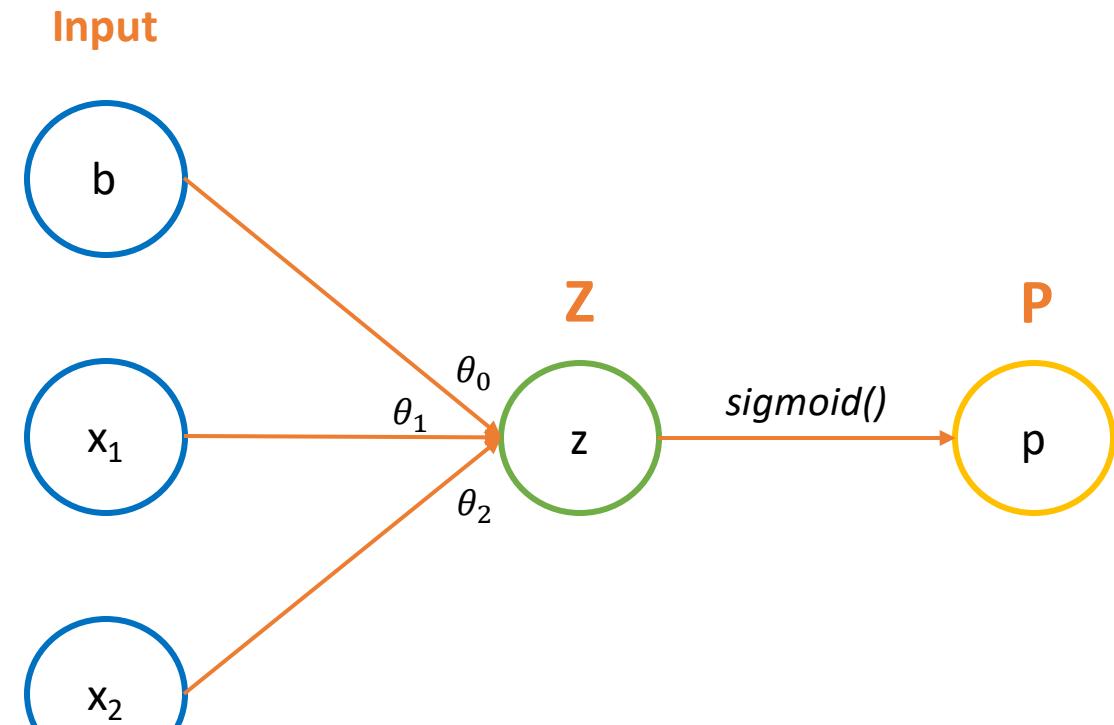
❖ Logistic Regression (Graph)

 : Input data x_i

 : Linear value (z)

 : Activation value (p)

 : Direction

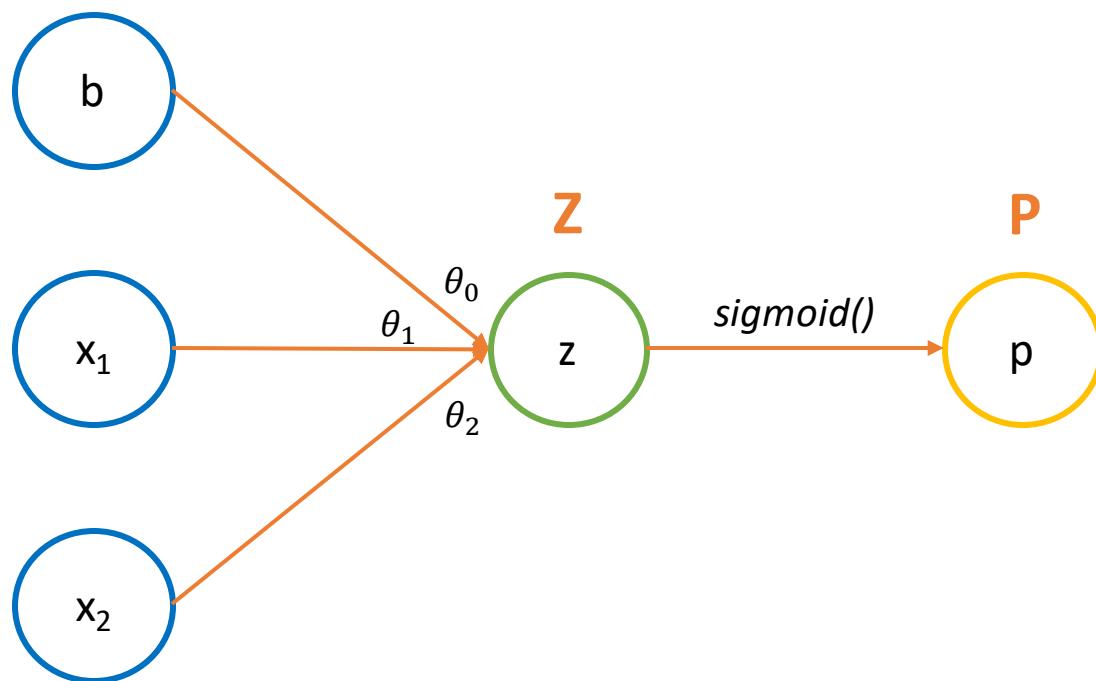


Logistic Regression Computation

Review

❖ Logistic Regression (Graph)

Input

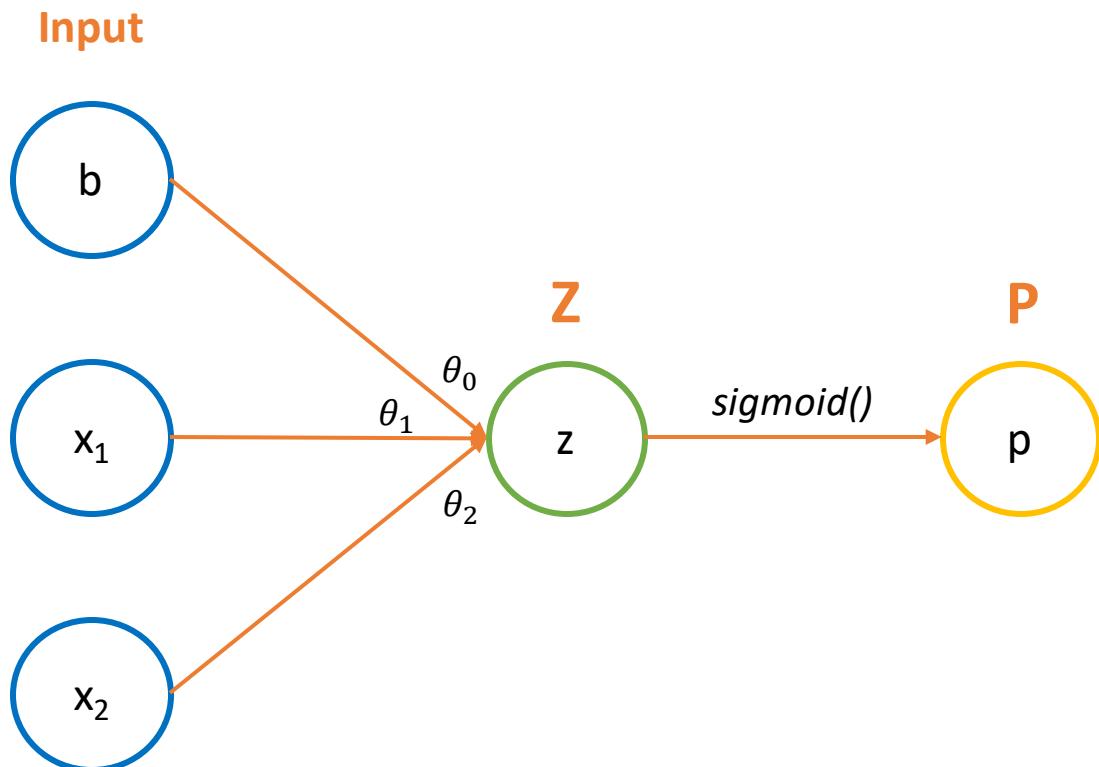


θ : Direction

Each direction is associated with a unique weight theta.

Review

❖ Logistic Regression (Graph)



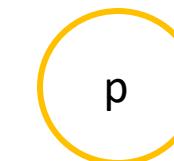
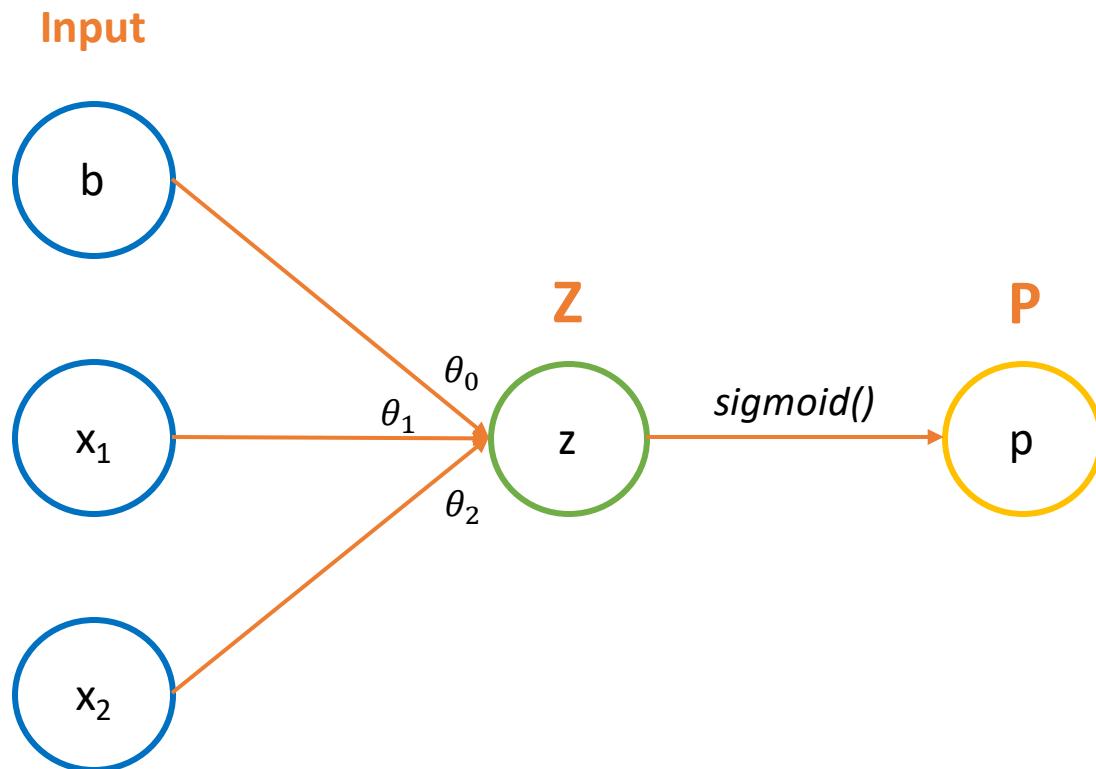
: Linear value

Green node represents the dot product of input features and their weights (including bias):

$$z = X \cdot \theta = \sum_{i=1}^n X_i \theta_i$$

Review

❖ Logistic Regression (Graph)



: Activation value (p)

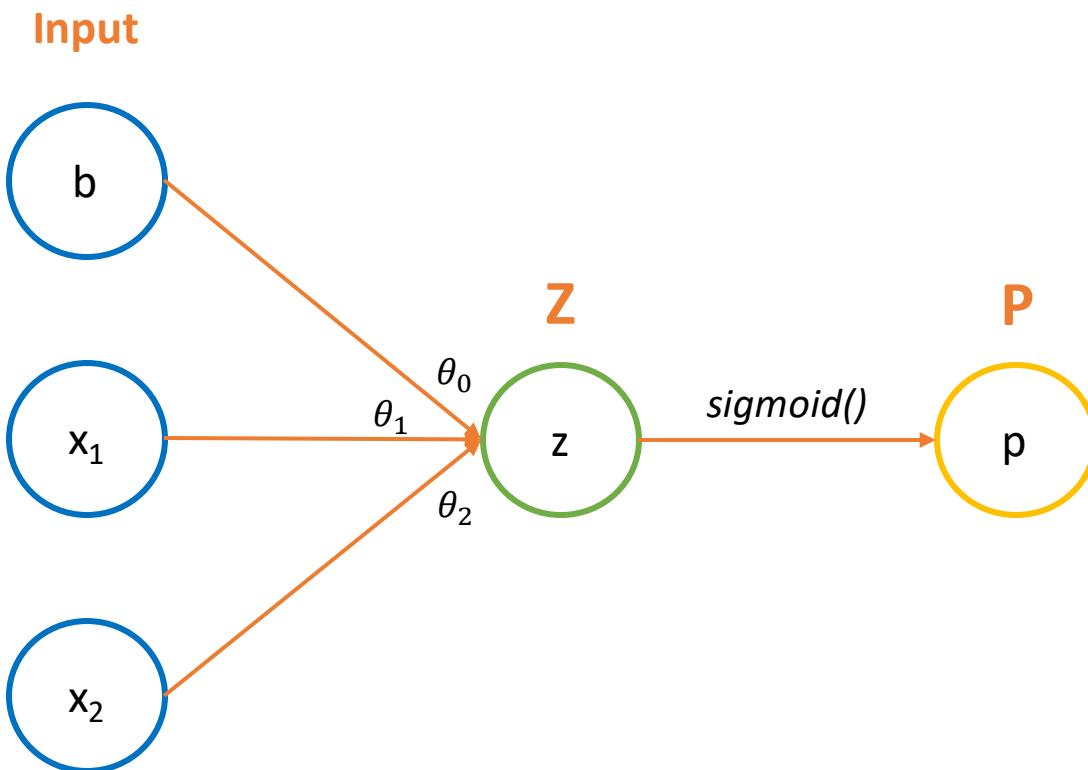
Yellow node represents the output of sigmoid function (activation function) with the input of z:

$$a(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

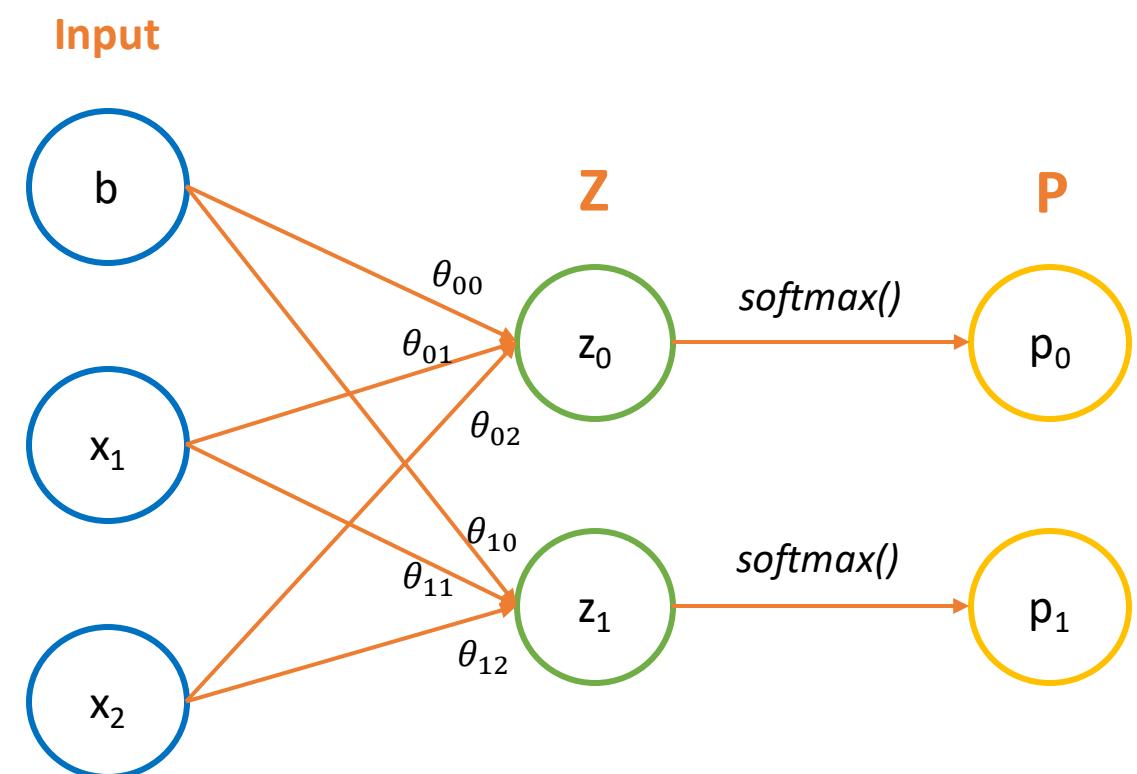
Review

❖ Softmax Regression

n_features = 2, n_classes = 2



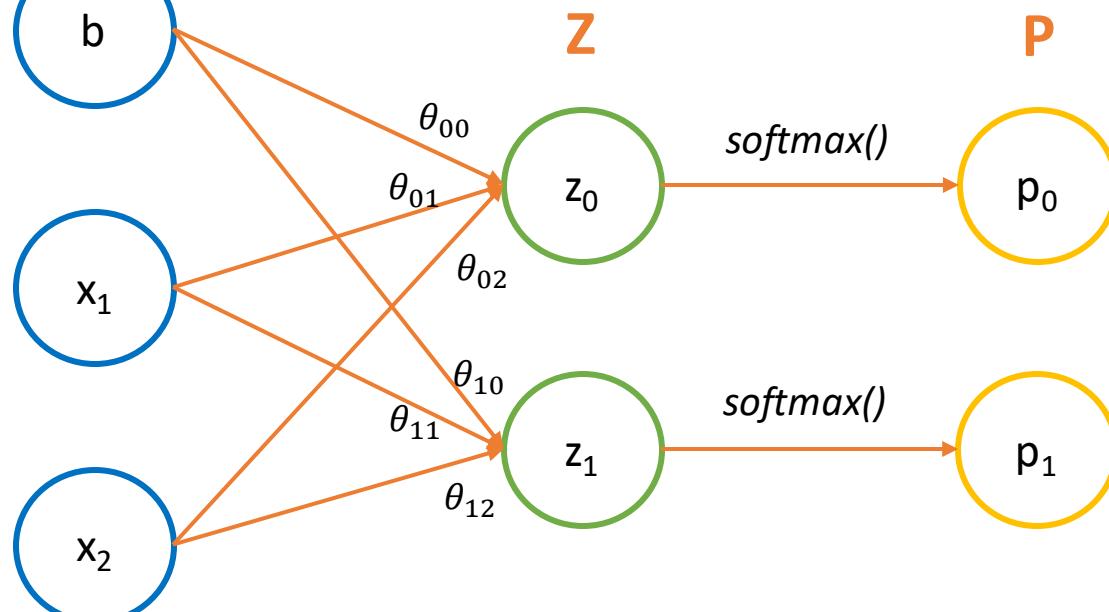
Logistic Regression



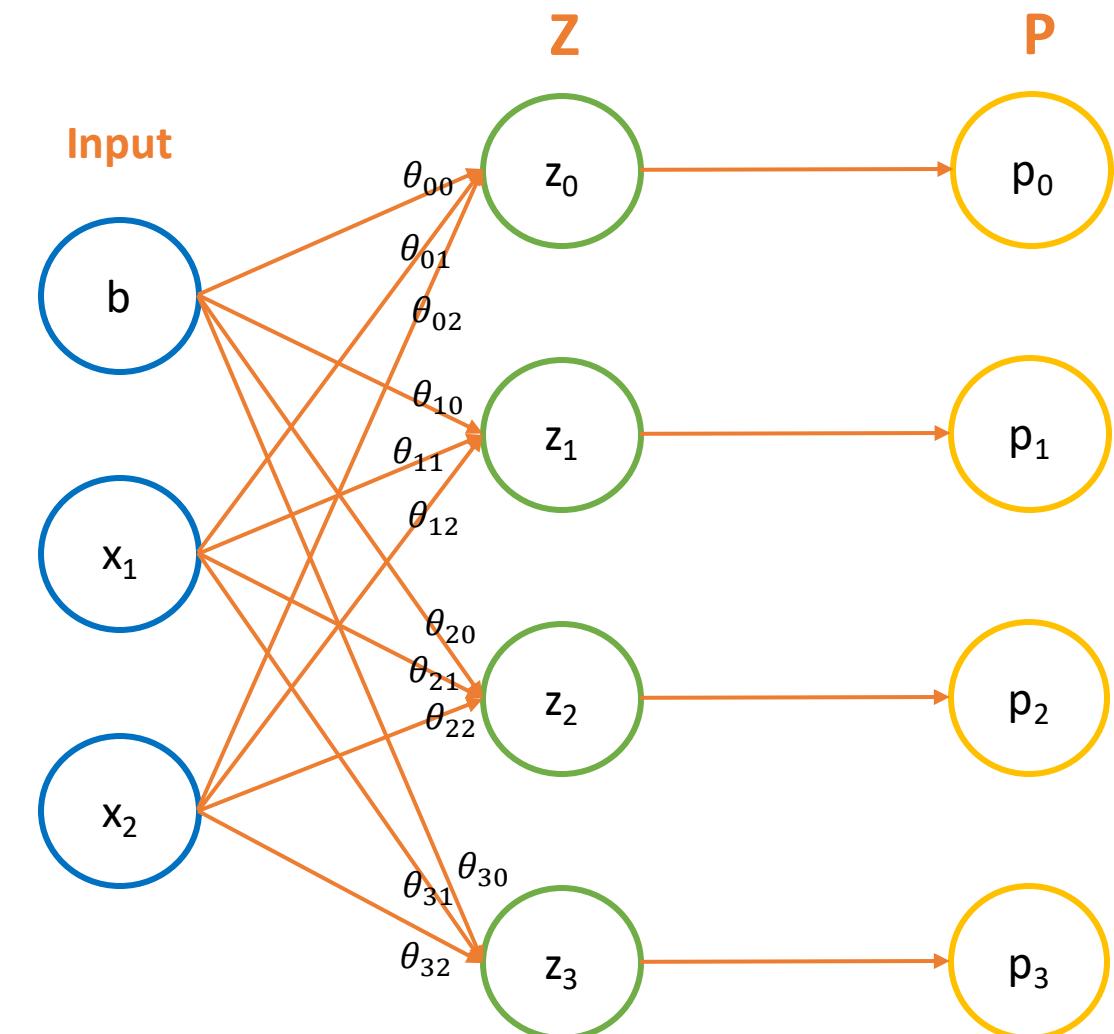
Softmax Regression

Review

❖ Softmax Regression



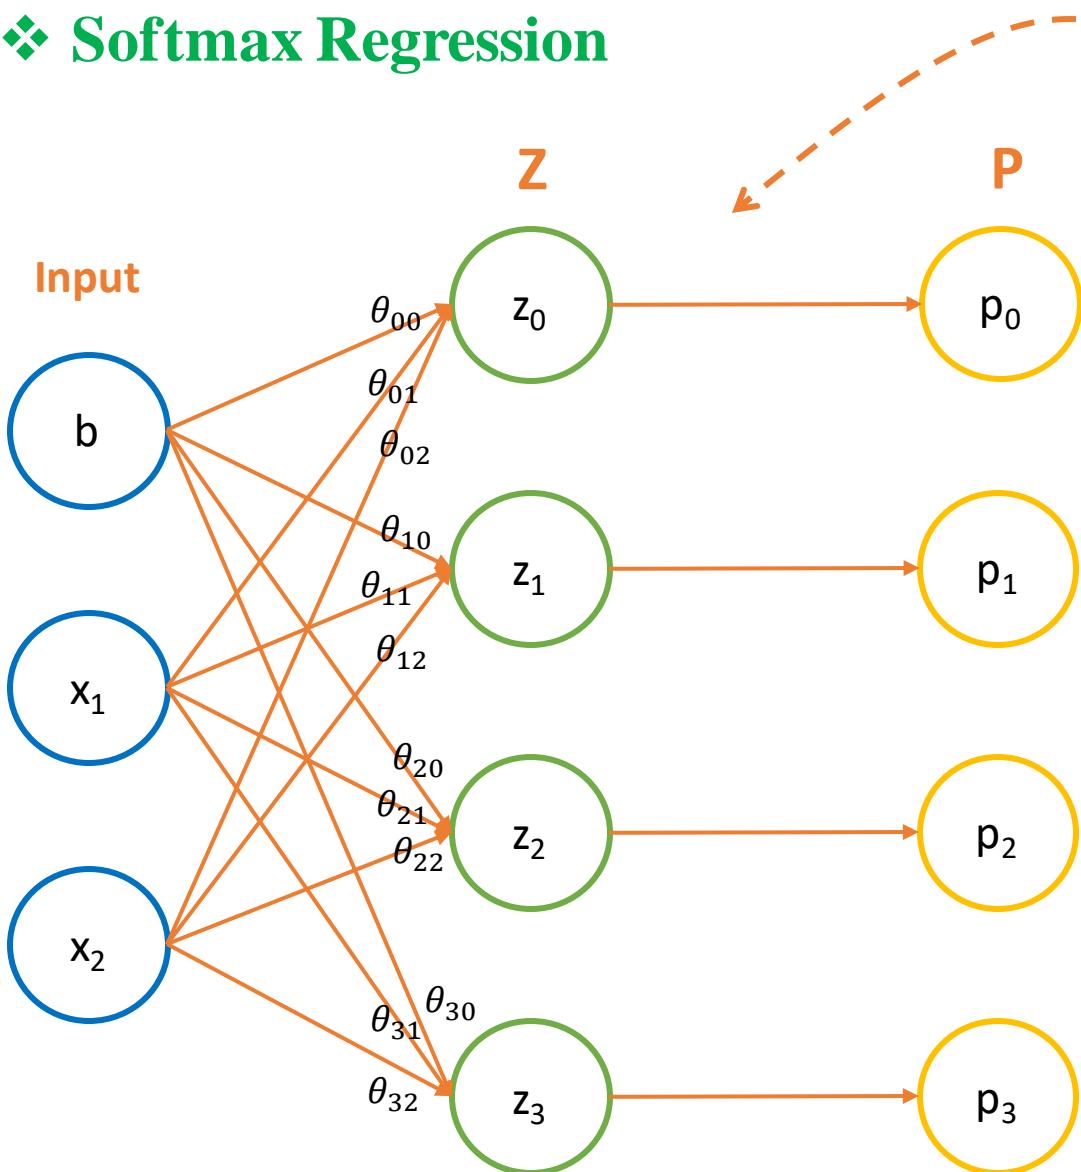
Softmax Regression (2 class)



Softmax Regression (4 class)

Review

❖ Softmax Regression



n_features = 2, n_classes = 4

❖ Calculate z of class i 'th:

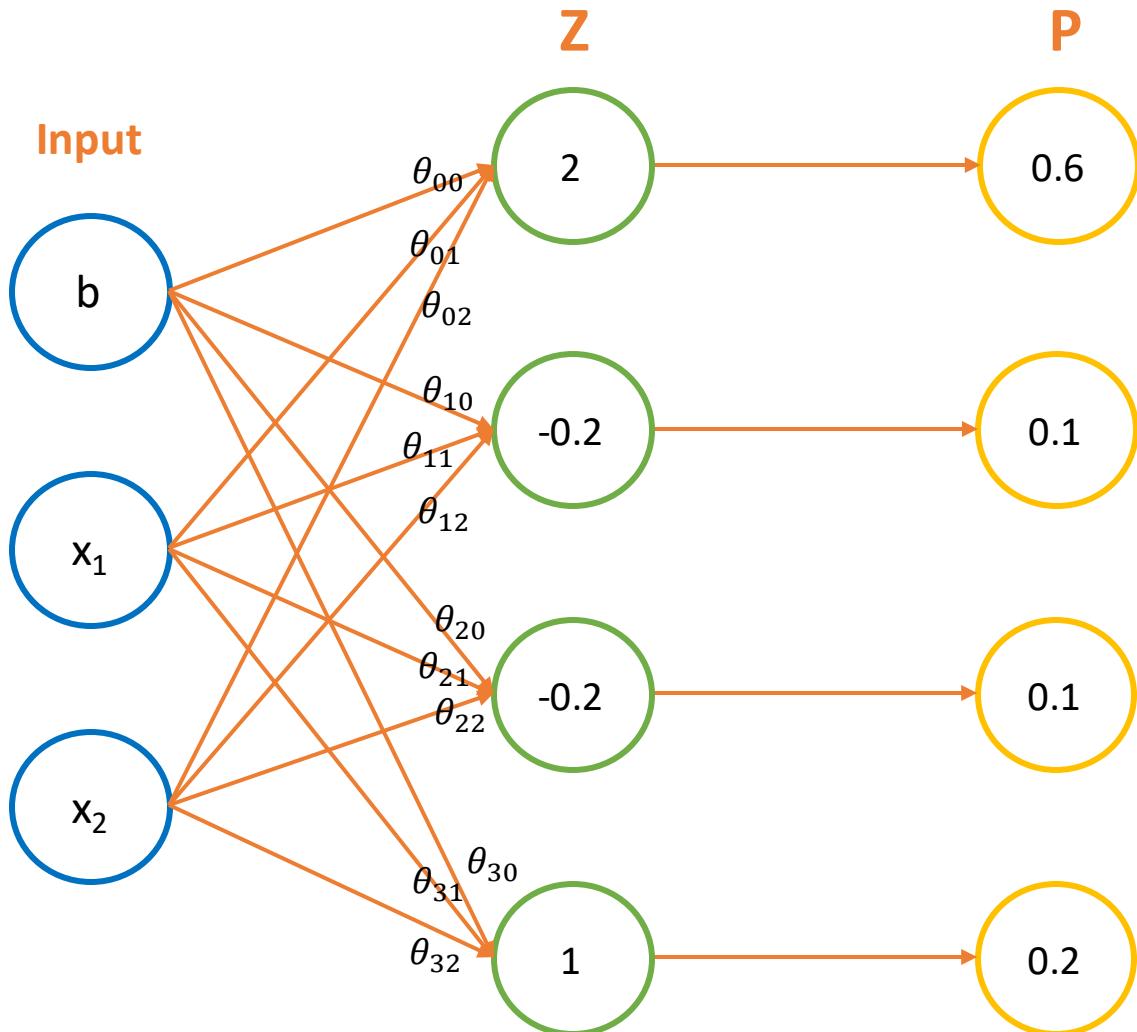
$$z_i = \theta_i^T X = b\theta_{i0} + \sum_{j=1}^{n_features} \theta_{ij} X_j$$

❖ Calculate p of class i 'th:

$$p_i = softmax(Z) = \frac{e^{z_i}}{\sum_{j=1}^{n_classes} e^{z_j}}$$

Review

❖ Properties of Softmax Output

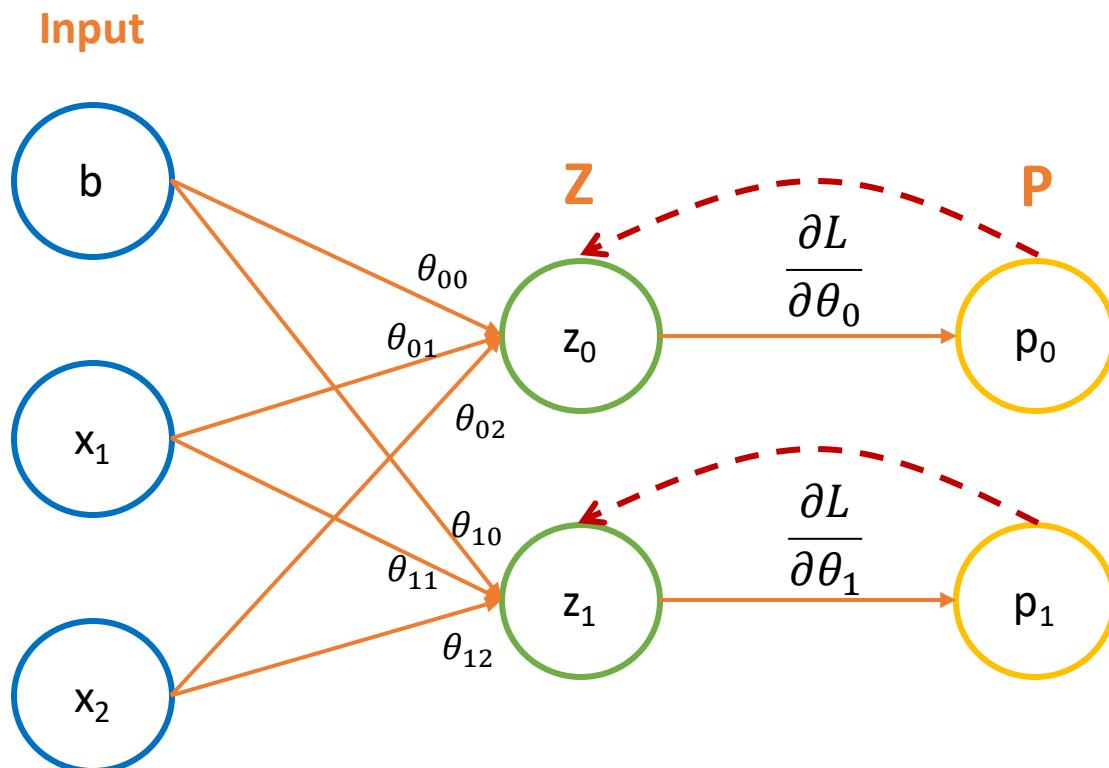


$n_features = 2, n_classes = 4$

- 1) $0 < p_i < 1$
- 2) $\sum_i p_i = 1$
- 3) $z_i < 0, p_i > 0$
- 4) $z_i = z_j, p_i = p_j$
- 5) $\max(z) = \max(p)$

Review

❖ Loss function and derivative



❖ Cross-entropy loss:

$$L(\hat{y}, y) = -\log(y^T \cdot \hat{y}) = -y^T \cdot \log(\hat{y})$$

❖ Derivative:

$$\nabla_{\theta} L = \frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z} \frac{\partial z}{\partial \theta} = X \cdot (\hat{y} - y)^T$$

Review

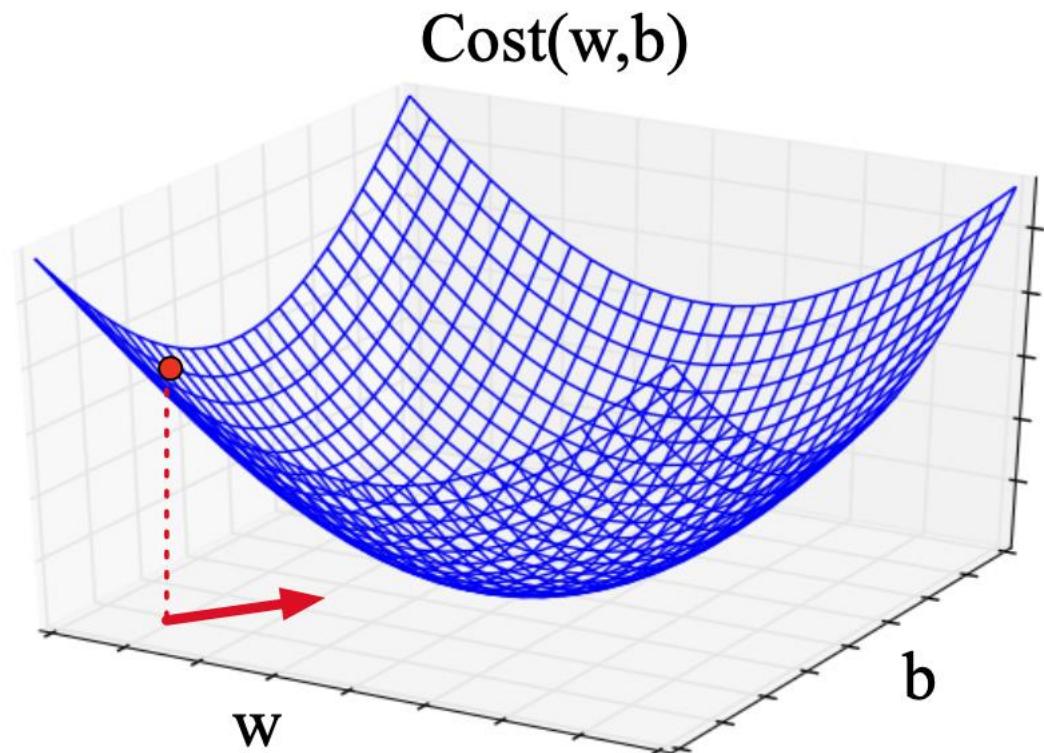
❖ Softmax Regression objective

Our objective is to find the optimal weights and bias (theta) that minimize the loss function:

$$\hat{\theta} = \operatorname{argmin}_{\theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}^i, y^i)$$

For each step, compute the gradient and update the current theta:

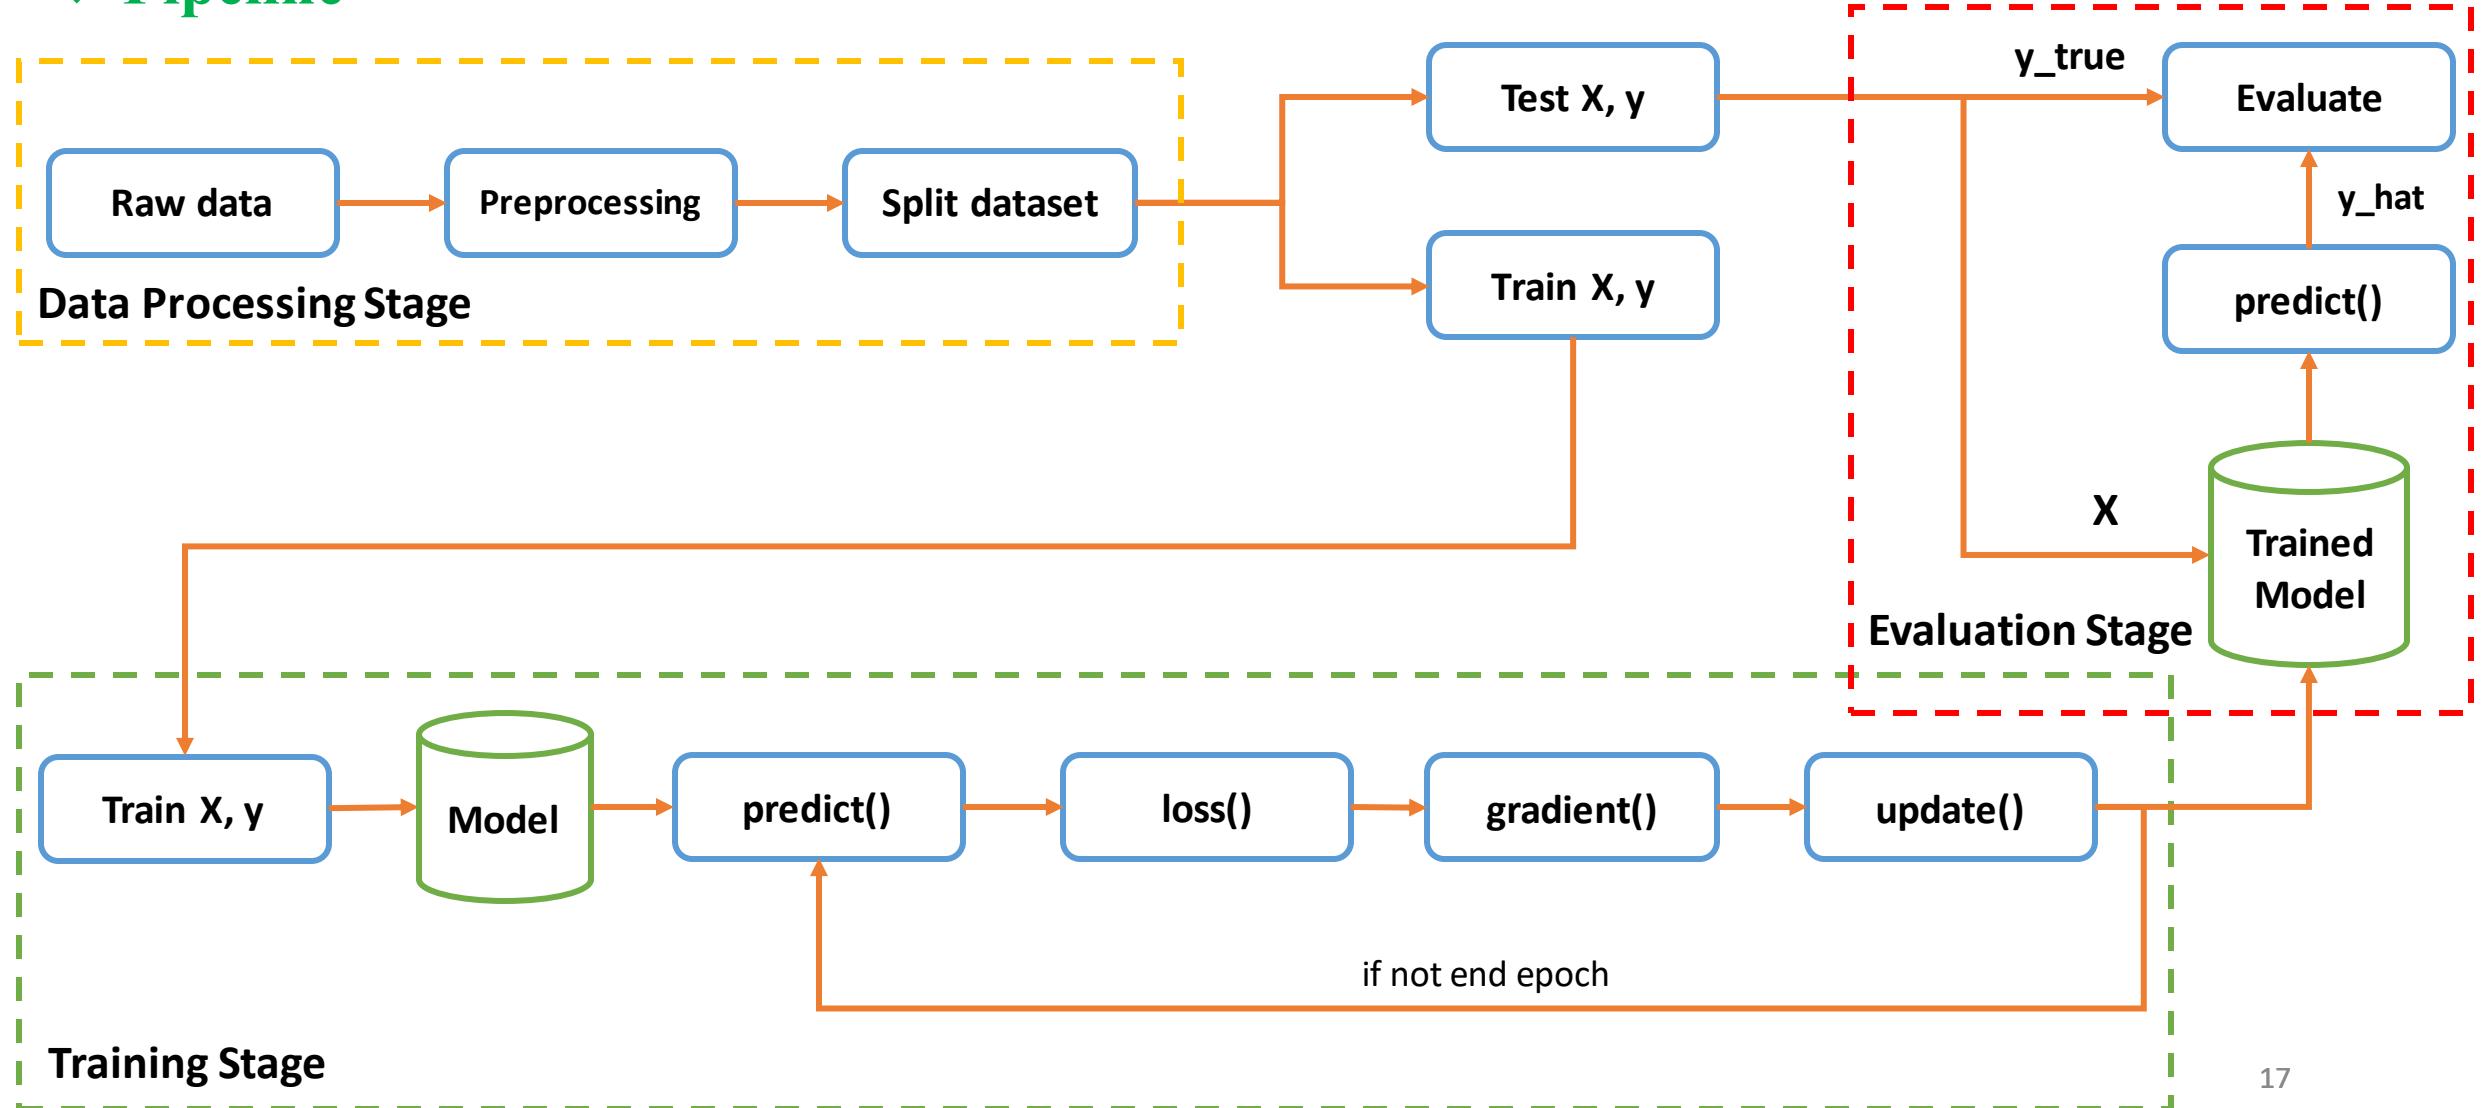
- $\nabla_{\theta_s} L = \frac{1}{N} X^T (\hat{y}^i - y^i)$
- $\theta_{s+1} = \theta_s - \eta \nabla_{\theta_s} L$



Gradient Descent

Review

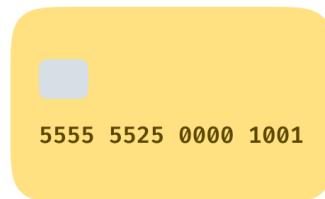
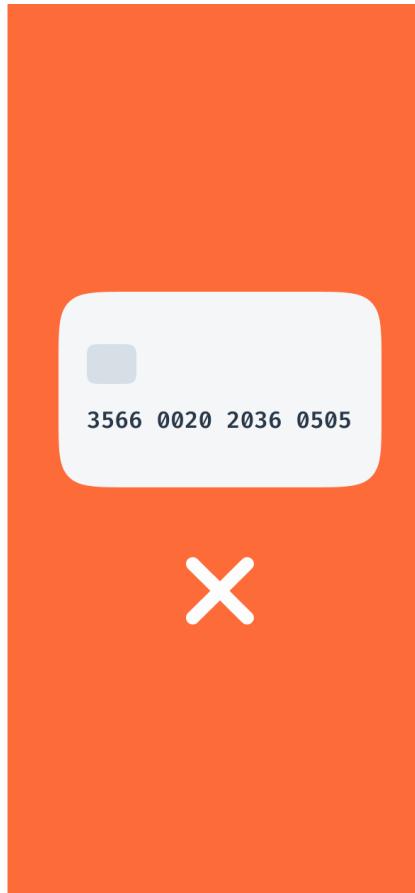
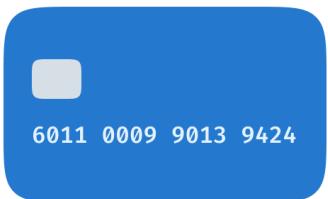
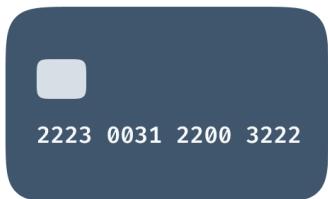
❖ Pipeline



Card Fraud Detection

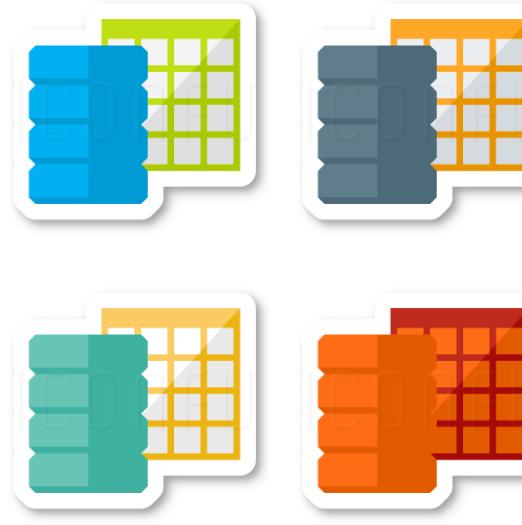
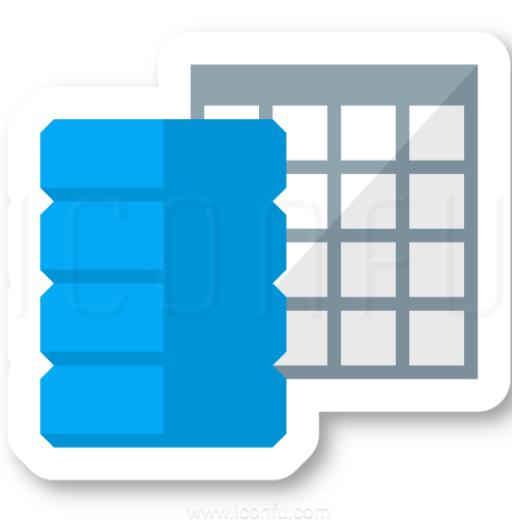
❖ Introduction

Description: Given [card_fraud_detection.csv](#) dataset, build a Card Fraud Detection model using Softmax Regression.



Card Fraud Detection

❖ Step 1: Import libraries and read dataset



	Time	v1	v2	v3	v4	v5	v6	v7	v8	v9
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 from sklearn.model_selection import train_test_split
6 from sklearn.preprocessing import StandardScaler
```

```
1 dataset_path = 'creditcard.csv'
2 df = pd.read_csv(
3     dataset_path
4 )
5 df
```

Card Fraud Detection

❖ Step 2: Get dataset information

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   Time      284807 non-null    float64
 1   V1        284807 non-null    float64
 2   V2        284807 non-null    float64
 3   V3        284807 non-null    float64
 4   V4        284807 non-null    float64
 5   V5        284807 non-null    float64
```

```
1 df.describe()
```

	Time	V1	V2	V3	V4	V5
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01

8 rows × 31 columns

Card Fraud Detection

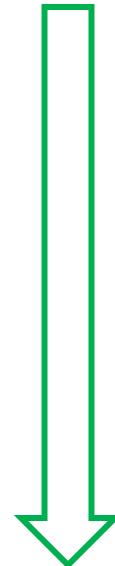
❖ Step 3: Split X, y



Independent Variable (Cause)



Dependent Variable (Result)



Index	Column 1	Column 2	Column 3	...	Column n
...
...
...
...

In our exercise's dataset:

- ❖ Independent variables (features): from column 1 to $n - 1$.
- ❖ Dependent variables (labels): last column.

Card Fraud Detection

❖ Step 3: Split X, y

Time	v1	v2	v3	v4	v5	v6	v7	v8	v9	...	v21	v22	v23	v24	v25	v26	v27	v28	Amount	Class
0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0
0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0
1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0
2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0
...	
!786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1.014480	-0.509348	1.436807	0.250034	0.943651	0.823731	0.77	0
!787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0.012463	-1.016226	-0.606624	-0.395255	0.068472	-0.053527	24.79	0
!788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0.037501	0.640134	0.265745	-0.087371	0.004455	-0.026561	67.88	0
!788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0.163298	0.123205	-0.569159	0.546668	0.108821	0.104533	10.00	0
!792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0.376777	0.008797	-0.473649	-0.818267	-0.002415	0.013649	217.00	0

```
1 dataset_arr = df.to_numpy()
2 X, y = dataset_arr[:, :-1].astype(np.float64), dataset_arr[:, -1].astype(np.uint8)
```

Card Fraud Detection

❖ Step 4: Add the bias term

❖ Calculate z of class i'th:

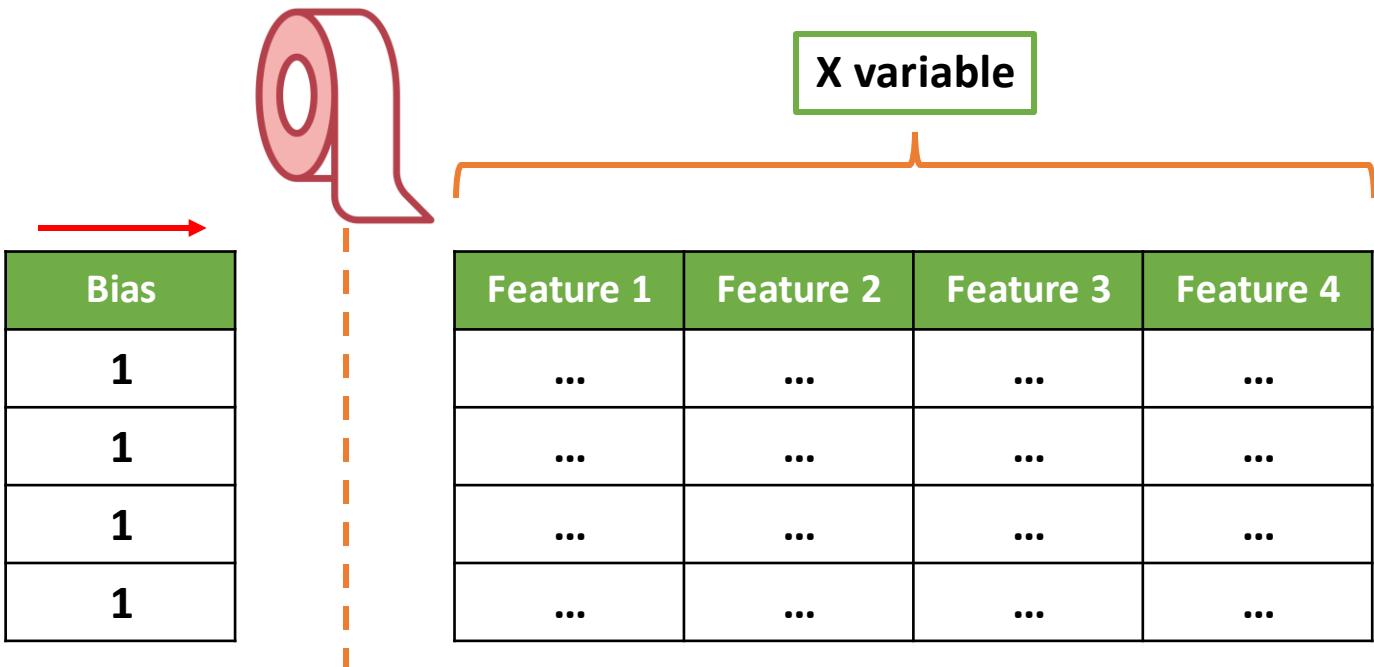
$$z_i = \theta_i^T X = b\theta_{i0} + \sum_{j=1}^{n_features} \theta_{ij}X_j$$

1. Ones Vector (n_elements = n_samples)

idx	0	1	...	n_samples - 1
-----	---	---	-----	---------------

value	1	1	1	1
-------	---	---	---	---

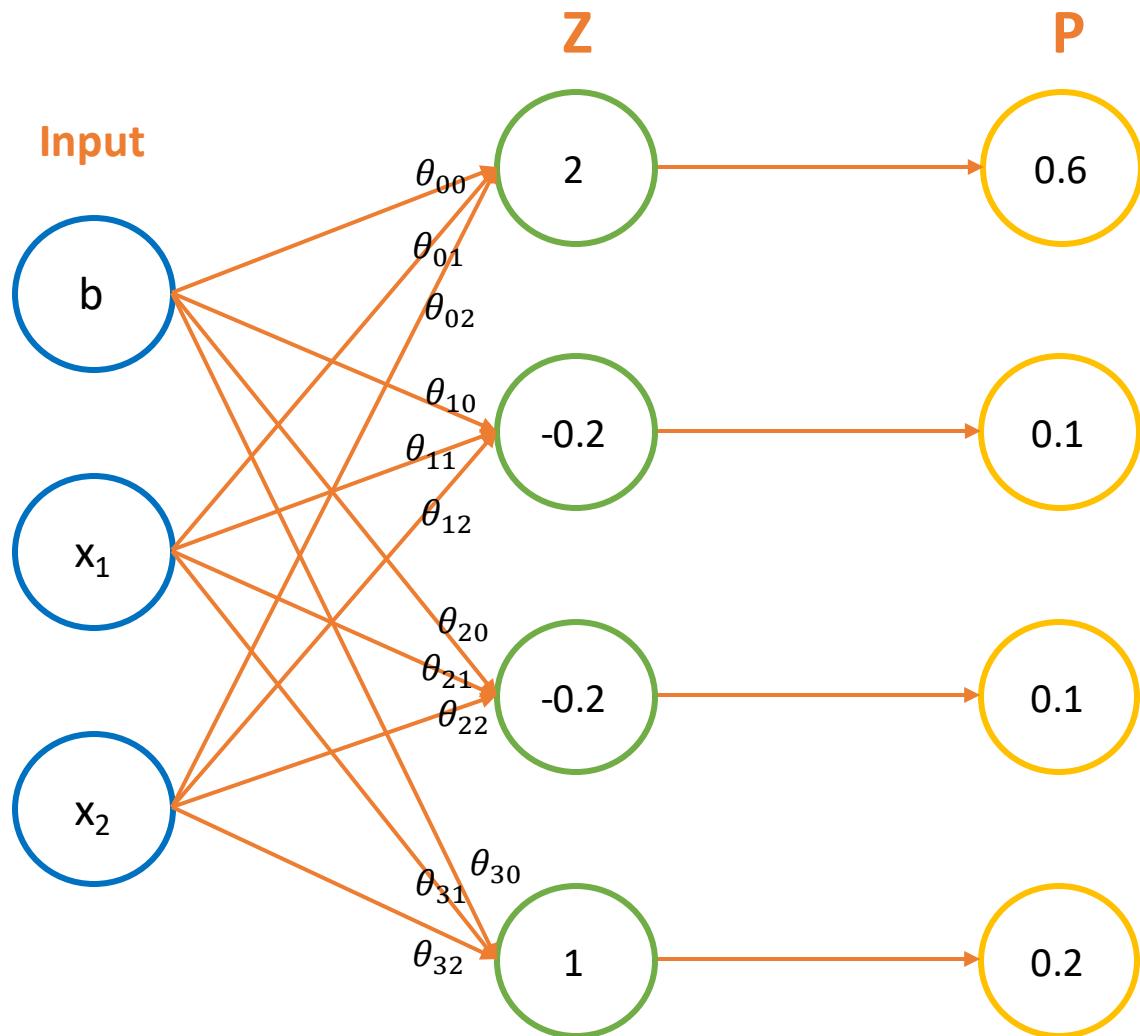
2. Concat by column to independent variable



```
1 intercept = np.ones((  
2 |     x.shape[0], 1)  
3 )  
4 x_b = np.concatenate(  
5 |     (intercept, x),  
6 |     axis=1  
7 )
```

Card Fraud Detection

❖ Step 5: One-hot encoding label



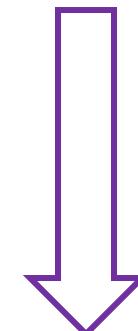
❖ Recall: Predict on sample i'th output an n elements vector

$$\hat{y}_i = [0.6, 0.1, 0.1, 0.2]$$

$$y_i = 0$$

`1 y`
array([0., 0., 0., ..., 0., 0., 0.])

How to calculate loss
between y vs \hat{y} ?



One hot encoding
label

Card Fraud Detection

❖ Step 5: One-hot encoding label

Definition: Convert categorial data to one-hot vector

1. Zeros Vector ($n_elements = n_classes$)

$n_features = 2$, $n_classes = 4$

idx	0	1	2	3
value	0	0	0	0

2. E.g: $y = 2$

$\text{zeros_vector}[y] = 1$

idx	0	1	2	3
value	0	0	1	0

Label encoding

Label

0
0
1
3
2
1
2

One-hot encoding

Label

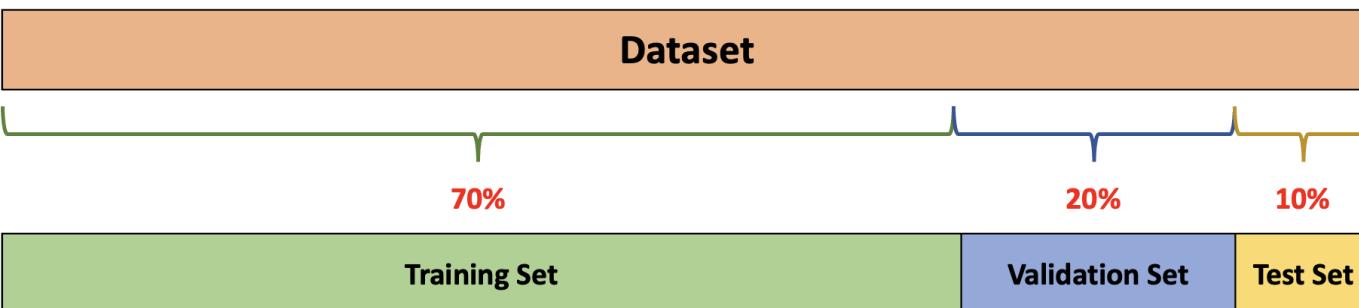
1	0	0	0
1	0	0	0
0	1	0	0
0	0	0	1
0	0	1	0
0	1	0	0
0	0	1	0



```
# One hot encoding label
y = np.array([np.zeros(n_classes) for _ in range(y_idx.shape[0])])
y[np.arange(len(y_idx)), y_idx] = 1
```

Card Fraud Detection

❖ Step 6: Split train, val, test set



```
1 val_size = 0.3
2 test_size = 0.1
3 random_state = 2
4 is_shuffle = True
5
6 X_train, X_val, y_train, y_val = train_test_split(
7     X_b, y_encoded,
8     test_size=val_size,
9     random_state=random_state,
10    shuffle=is_shuffle
11 )
12
13 X_val, X_test, y_val, y_test = train_test_split(
14     X_val, y_val,
15     test_size=test_size,
16     random_state=random_state,
17     shuffle=is_shuffle
18 )
```

```
1 print(f'Number of training samples: {X_train.shape[0]}')
2 print(f'Number of val samples: {X_val.shape[0]}')
3 print(f'Number of test samples: {X_test.shape[0]}')
```

Number of training samples: 199364
Number of val samples: 76898
Number of test samples: 8545

Card Fraud Detection

❖ Step 7: Normalization

Using `sklearn.preprocessing.StandardScaler()` to scale all values in dataset.

$$z = \frac{x_i - \mu}{\sigma}$$

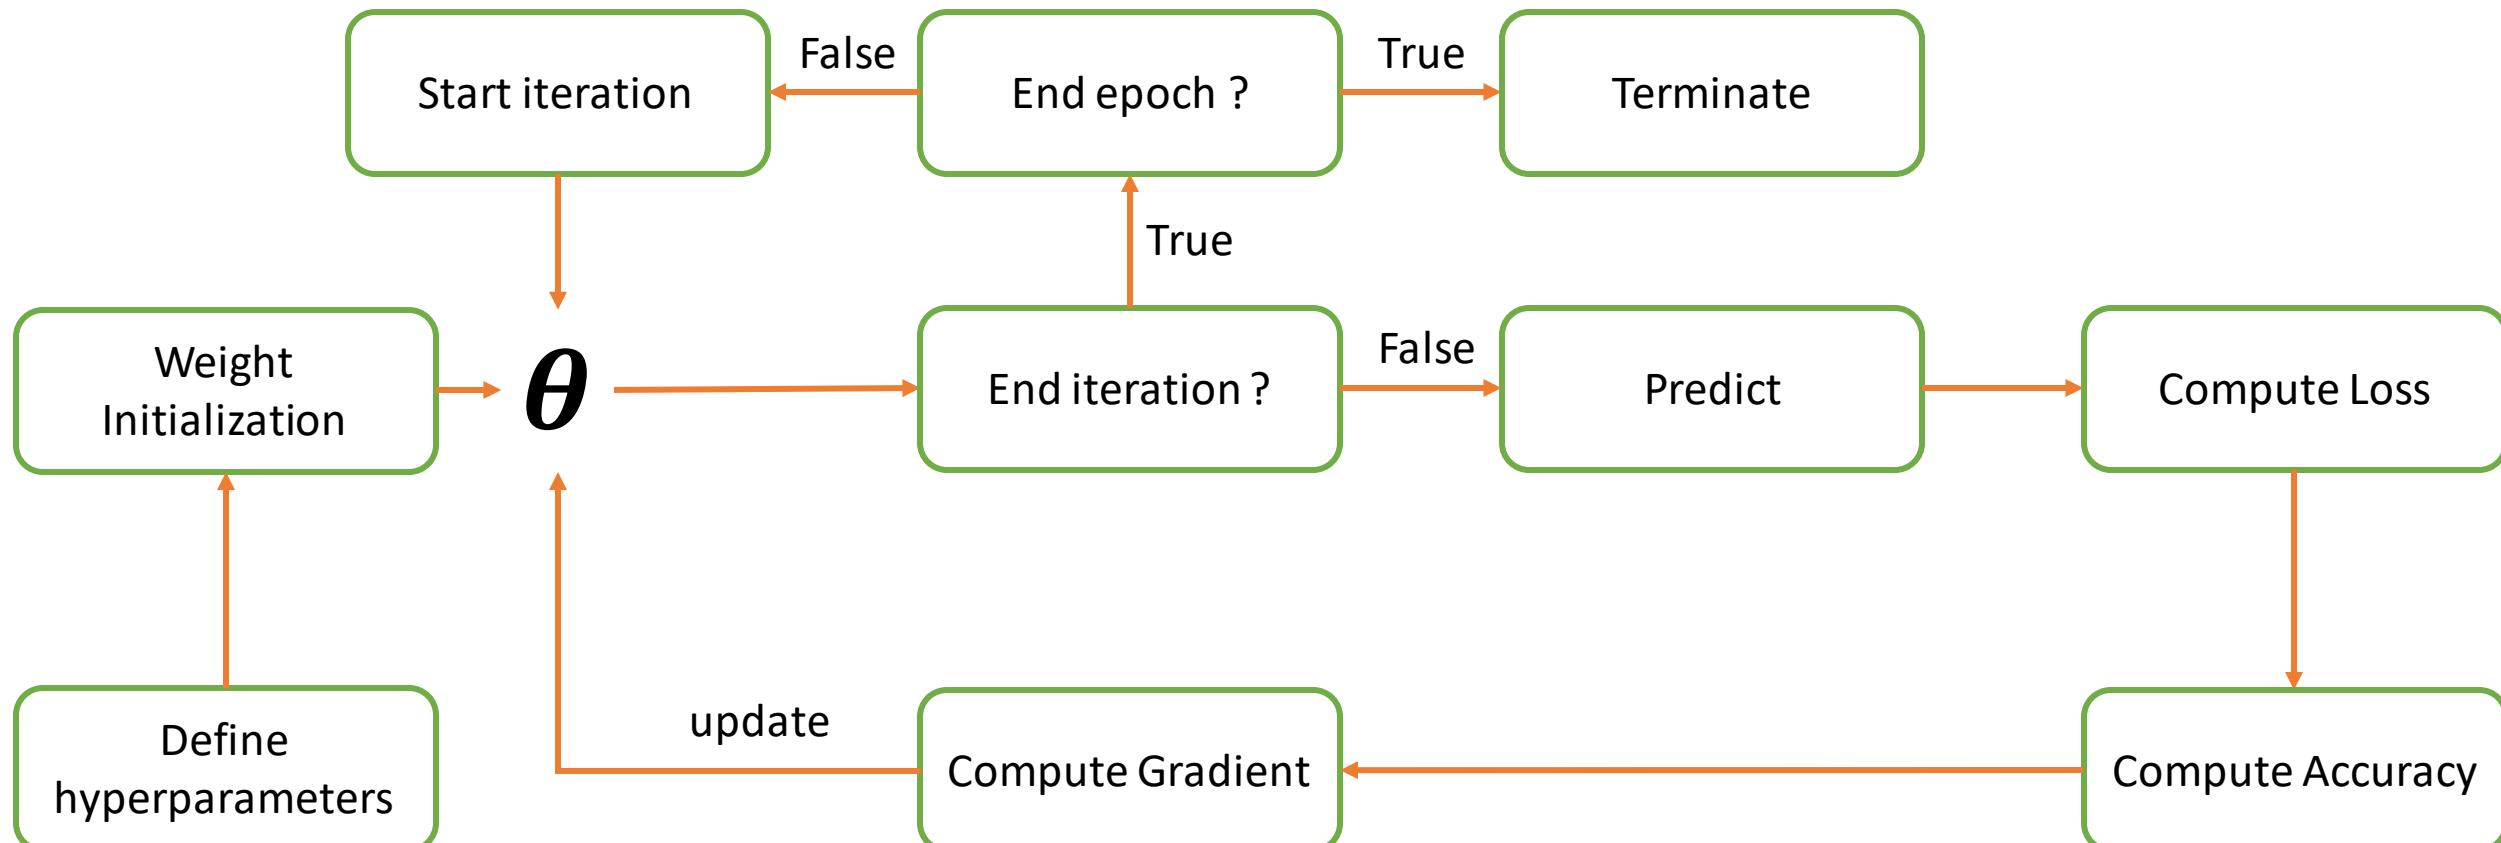
```
1 normalizer = StandardScaler()  
2 X_train[:, 1:] = normalizer.fit_transform(X_train[:, 1:])  
3 X_val[:, 1:] = normalizer.transform(X_val[:, 1:])  
4 X_test[:, 1:] = normalizer.transform(X_test[:, 1:])
```

```
1 X_train  
  
array([[ 1.          , -1.95311891, -0.76256074, ...,  0.79907219,  
       -0.37273829, -0.32754448],  
       [ 1.          ,  1.49473018,  0.96434558, ...,  0.119096 ,  
       -0.11391761,  0.13119296],  
       [ 1.          , -0.4733483 , -1.33737659, ...,  1.16380164,  
       -0.46691221, -0.34433048],  
       ...,  
       [ 1.          , -0.61621632, -0.40746973, ..., -0.05352446,  
       0.15422491, -0.04989603],  
       [ 1.          ,  0.84222097,  1.03085144, ..., -0.02330856,  
       -0.18601998, -0.30104658],  
       [ 1.          , -0.57043285,  0.66268869, ...,  0.19155844,  
       0.08433652, -0.34748785]])
```

Note: We only use the train set to fit the scaler.

Card Fraud Detection

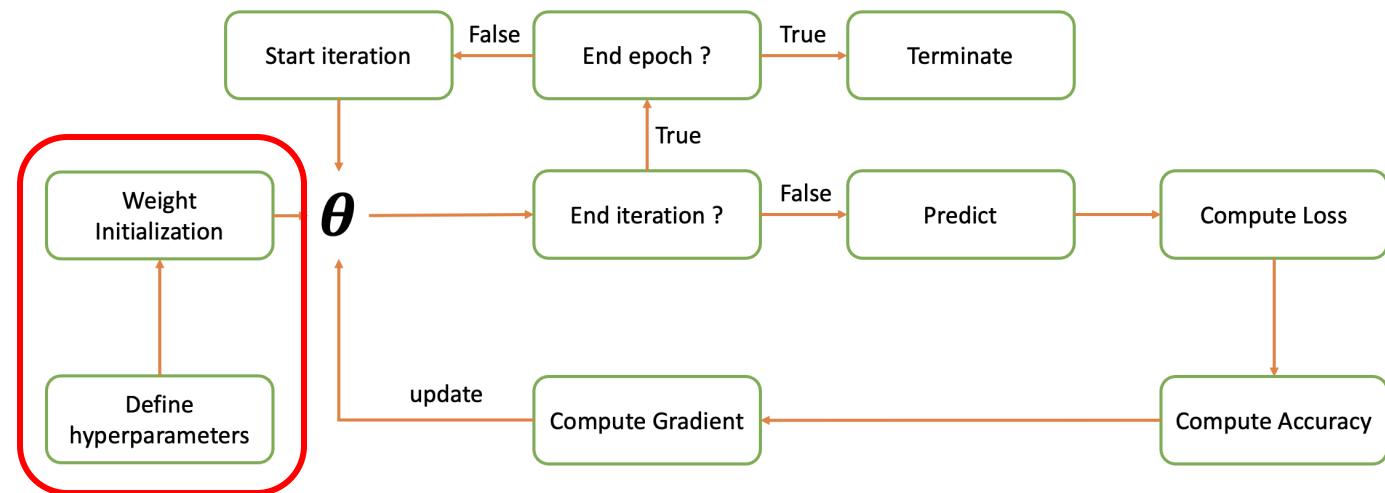
❖ Step 8: Training



Card Fraud Detection

❖ Step 8: Training

```
1 lr = 0.01
2 epochs = 30
3 batch_size = 1024
4 n_features = X_train.shape[1]
5
6 np.random.seed(random_state)
7 theta = np.random.uniform(
8     size=(n_features, n_classes))
9 )
```



Define essential hyperparameters and initialize weights

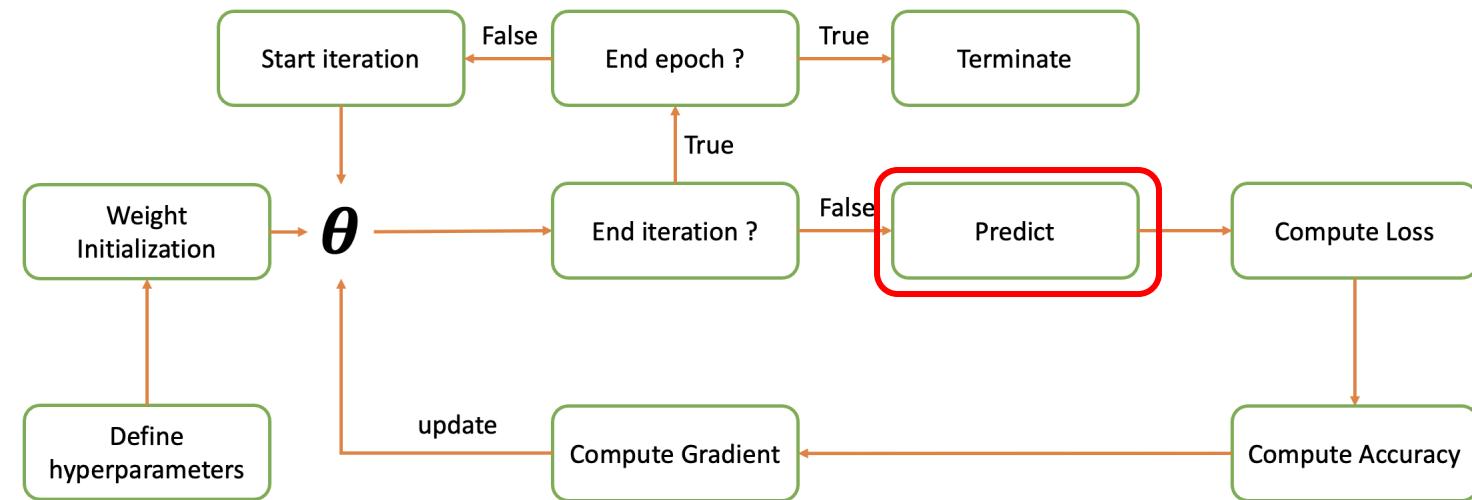
Card Fraud Detection

❖ Step 8: Training

Hypothesis function

- $z = b_0 + \sum_{i=1}^{n_features} b_i x_i$
- $y_{hat} = softmax(z) = \frac{e^{z_i}}{\sum_{j=1}^{n_classes} e^{z_j}}$

```
1 def softmax(z):  
2     exp_z = np.exp(z)  
3  
4     return exp_z / exp_z.sum(axis=1)[:, None]  
5  
6  
7 def predict(x, theta):  
8     z = np.dot(x, theta)  
9     y_hat = softmax(z)  
10  
11    return y_hat
```



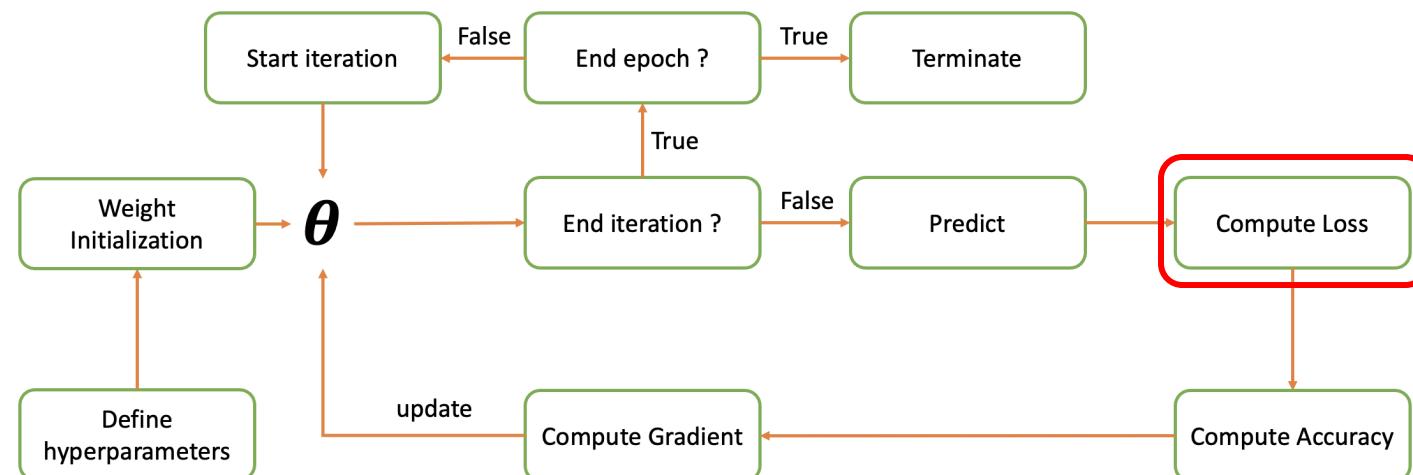
Card Fraud Detection

❖ Step 8: Training

Softmax Regression Loss (Cross-entropy)

$$L(\hat{y}, y) = -\frac{1}{N} \sum_{i=1}^N y_i \log \hat{y}_i$$

```
1 def compute_loss(y_hat, y):  
2     n = y.size  
3  
4     return (-1 / n) * np.sum(y * np.log(y_hat))
```



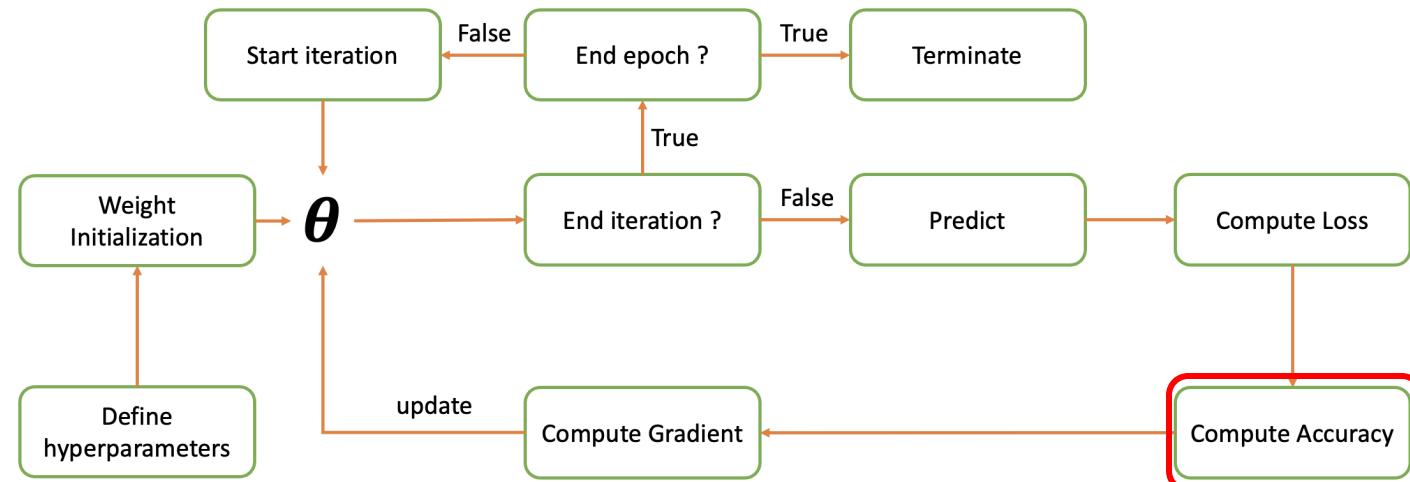
Card Fraud Detection

❖ Step 8: Training

Accuracy Formula

$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(X, y, theta):  
2     y_hat = predict(X, theta)  
3     max_y_hat = np.argmax(y_hat, axis=1)  
4     max_y = np.argmax(y, axis=1)  
5     acc = (max_y_hat == max_y).mean()  
6  
7     return acc
```



Card Fraud Detection

❖ Step 8: Training

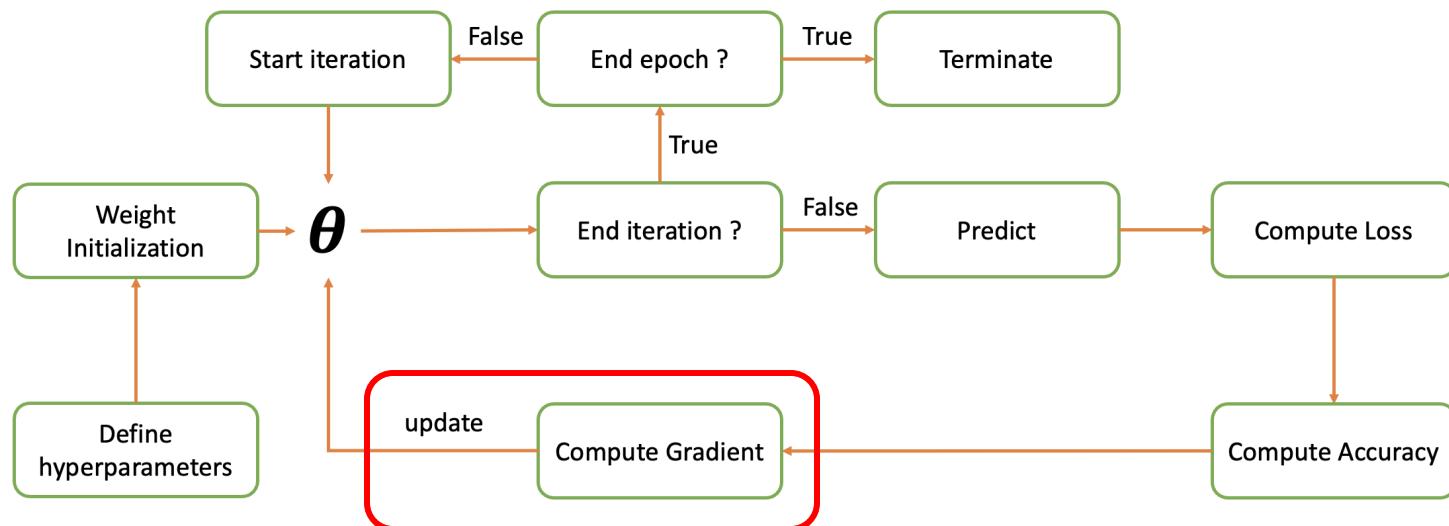
- Gradient computation formula:

$$\nabla_{\theta} L = \frac{1}{N} X^T (\hat{y} - y)$$

- Weights update formula:

$$\theta = \theta - \eta \nabla_{\theta} L$$

```
1 def compute_gradient(X, y, y_hat):  
2     n = y.size  
3  
4     return np.dot(X.T, (y_hat - y)) / n  
5  
6 def update_theta(theta, gradient, lr):  
7     return theta - lr * gradient
```



Card Fraud Detection

❖ Step 8: Training

Line 6: Loop over number of epochs.

Line 7, 8, 9, 10: Declare empty lists to store batch accuracies, losses of train and val sets.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

Card Fraud Detection

❖ Step 8: Training

Line 12: Loop over number of batches (based on batch size).

Line 13, 14: Get X and y data of current batch.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []
11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]
15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)
21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)
24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)
28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)
31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

Card Fraud Detection

❖ Step 8: Training

Line 16: Do prediction on X_i and theta.

Line 17: Compute the loss of $y_{\hat{}}$ and y_i .

Line 18: Compute the gradient.

Line 19: Update theta using the computed gradient.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16    y_hat = predict(X_i, theta)
17    train_loss = compute_loss(y_hat, y_i)
18    gradient = compute_gradient(X_i, y_i, y_hat)
19    theta = update_theta(theta, gradient, lr)
20    train_batch_losses.append(train_loss)

21
22    train_acc = compute_accuracy(X_train, y_train, theta)
23    train_batch_accs.append(train_acc)

24
25    y_val_hat = predict(X_val, theta)
26    val_loss = compute_loss(y_val_hat, y_val)
27    val_batch_losses.append(val_loss)

28
29    val_acc = compute_accuracy(X_val, y_val, theta)
30    val_batch_accs.append(val_acc)

31
32    train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33    val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34    train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35    val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

Card Fraud Detection

❖ Step 8: Training

From line 20 to line 30: Compute and store accuracies and losses on train and val sets.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

Card Fraud Detection

❖ Step 8: Training

From line 32 to line 35: Compute and store batch accuracies and losses of train and val sets.

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []

11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]

15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)

21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)

24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)

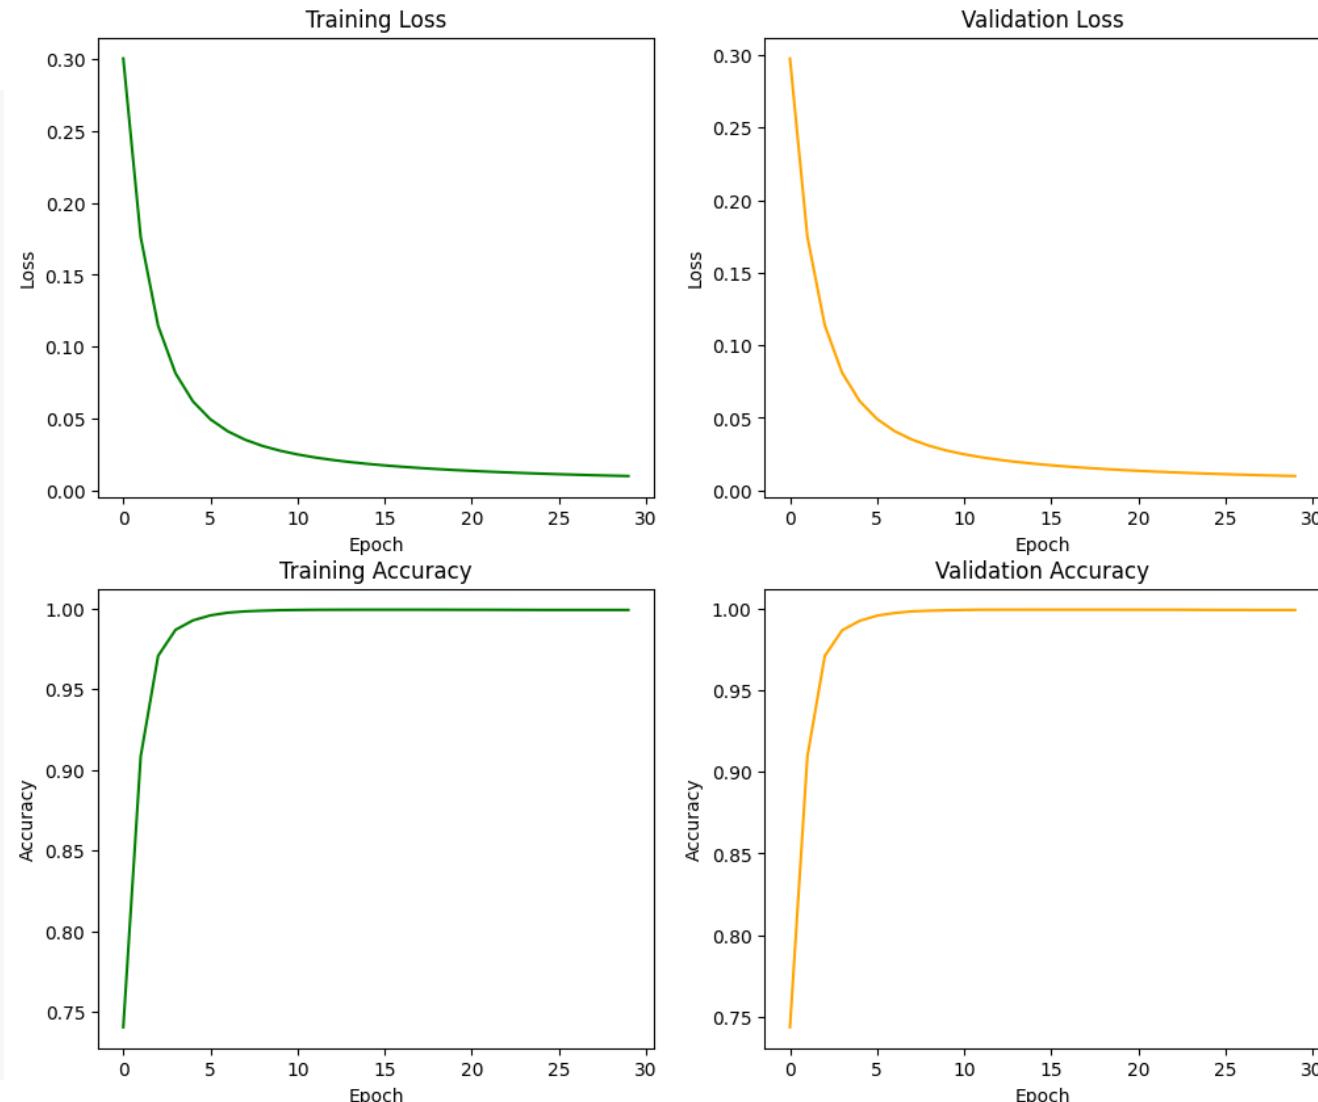
28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)

31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

Card Fraud Detection

❖ Step 8: Training (Visualization)

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses, color='green')
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, color='orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs, color='green')
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, color='orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



Card Fraud Detection

❖ Step 9: Evaluation

$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(X, y, theta):  
2     y_hat = predict(X, theta)  
3     max_y_hat = np.argmax(y_hat, axis=1)  
4     max_y = np.argmax(y, axis=1)  
5     acc = (max_y_hat == max_y).mean()  
6  
7     return acc
```

```
1 # Val set  
2 val_set_acc = compute_accuracy(X_val, y_val, theta)  
3 print('Evaluation on validation set: ')  
4 print(f'Accuracy: {val_set_acc}')
```

Evaluation on validation set:
Accuracy: 0.9992197456370777

```
1 # Test set  
2 test_set_acc = compute_accuracy(X_test, y_test, theta)  
3 print('Evaluation on test set: ')  
4 print(f'Accuracy: {test_set_acc}')
```

Evaluation on test set:
Accuracy: 0.9994148624926857

Sentiment Analysis

❖ Introduction

Description: Given [Twitter Sentiment Analysis](#) dataset, build a Logistic Regression model to determine whether a tweet (text) has a positive, neutral or negative sentiment.



Positive



Neutral



Negative

Sentiment Analysis

❖ Introduction

In this exercise, we will implement Softmax Regression with two approaches: **NumPy** and **PyTorch**.



Approach 1: Build with NumPy



Approach 2: Build with PyTorch

Sentiment Analysis

❖ Step 1: Import libraries

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import re
5 import nltk
6 nltk.download('stopwords')
7
8 from sklearn.model_selection import train_test_split
9 from sklearn.feature_extraction.text import TfidfVectorizer
10 from nltk.corpus import stopwords
11 from nltk.stem import SnowballStemmer
```



Sentiment Analysis

❖ Step 2: Read dataset

```
1 dataset_path = 'Twitter_Data.csv'  
2 df = pd.read_csv(  
3     dataset_path  
4 )  
5 df
```

	clean_text	category
0	when modi promised “minimum government maximum...	-1.0
1	talk all the nonsense and continue all the dra...	0.0
2	what did just say vote for modi welcome bjp t...	1.0
3	asking his supporters prefix chowkidar their n...	1.0
4	answer who among these the most powerful world...	1.0
...
162975	why these 456 crores paid neerav modi not reco...	-1.0
162976	dear rss terrorist payal gawar what about modi...	-1.0
162977	did you cover her interaction forum where she ...	0.0
162978	there big project came into india modi dream p...	0.0
162979	have you ever listen about like gurukul where ...	1.0

162980 rows × 2 columns

Sentiment Analysis

❖ Step 3: Get dataset information

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 162980 entries, 0 to 162979
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   clean_text  162976 non-null   object  
 1   category    162973 non-null   float64
dtypes: float64(1), object(1)
memory usage: 2.5+ MB
```

category
count 162973.000000
mean 0.225436
std 0.781279
min -1.000000
25% 0.000000
50% 0.000000
75% 1.000000
max 1.000000

Sentiment Analysis

❖ Step 4: Data preprocessing

	clean_text	category
0	when modi promised "minimum government maximum...	-1.0
1	talk all the nonsense and continue all the dra...	0.0
2	what did just say vote for modi welcome bjp t...	1.0
3	asking his supporters prefix chowkidar their n...	1.0
4	answer who among these the most powerful world...	1.0
...
162975	why these 456 crores paid neerav modi not reco...	-1.0
162976	dear rss terrorist payal gawar what about modi...	-1.0
162977	did you cover her interaction forum where she ...	0.0
162978	there big project came into india modi dream p...	0.0
162979	have you ever listen about like gurukul where ...	1.0

162980 rows × 2 columns

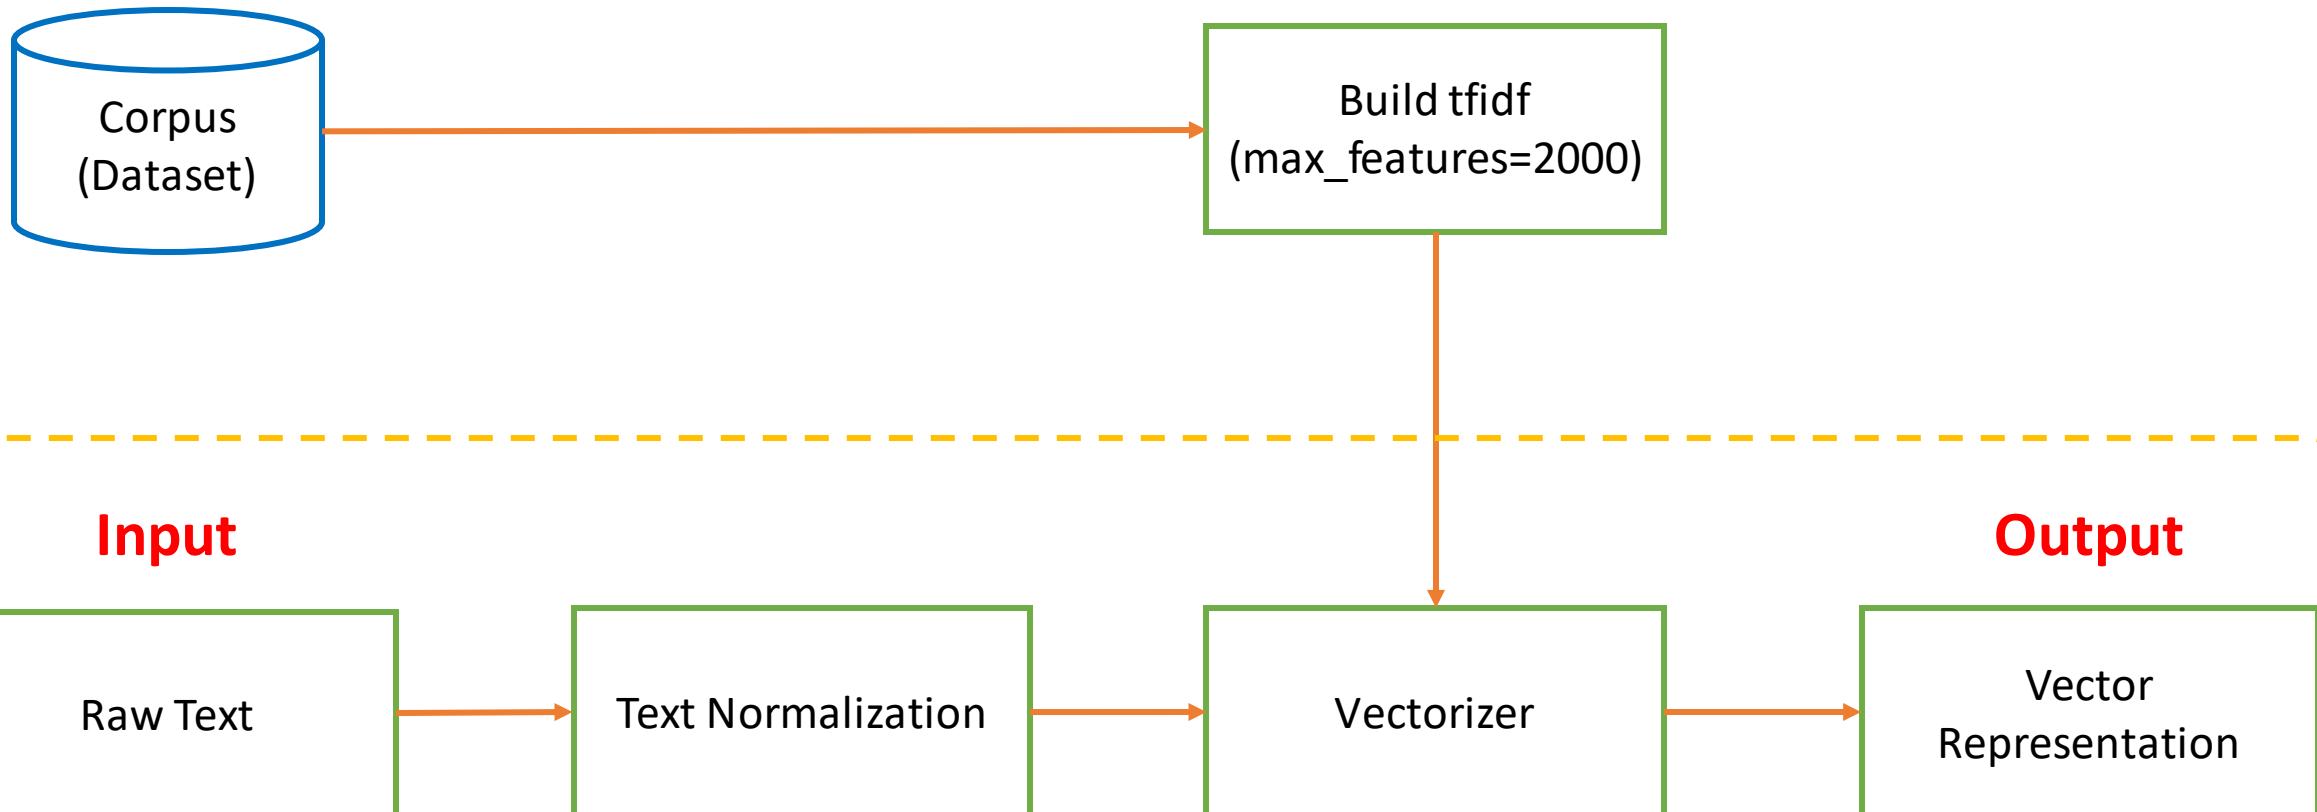
Raw Text: We love this! Would you go? #talk
#makememories #unplug #relax #iphone #smartphone
#wifi #connect... http://fb.me/6N3LsUpCu

Data
Preprocessing

Vector Representation: array([1.0, 4768.0, 1425.0])

Sentiment Analysis

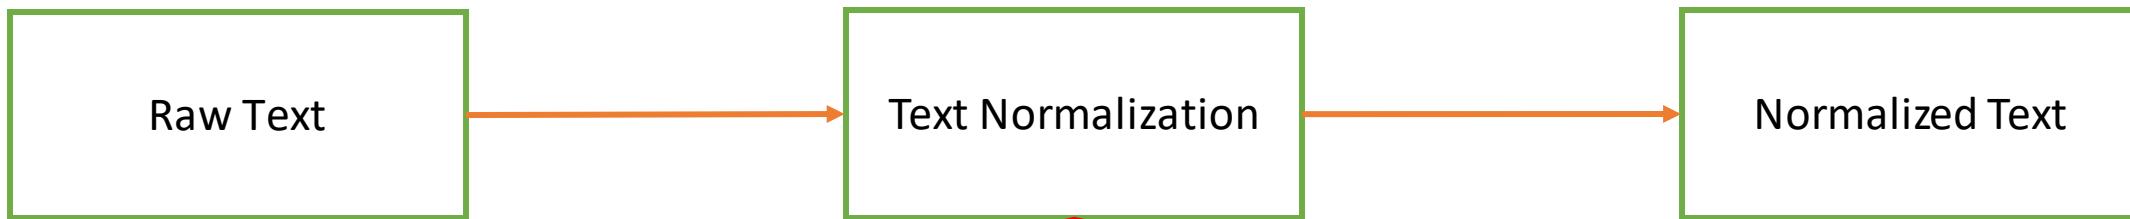
❖ Step 4: Data preprocessing



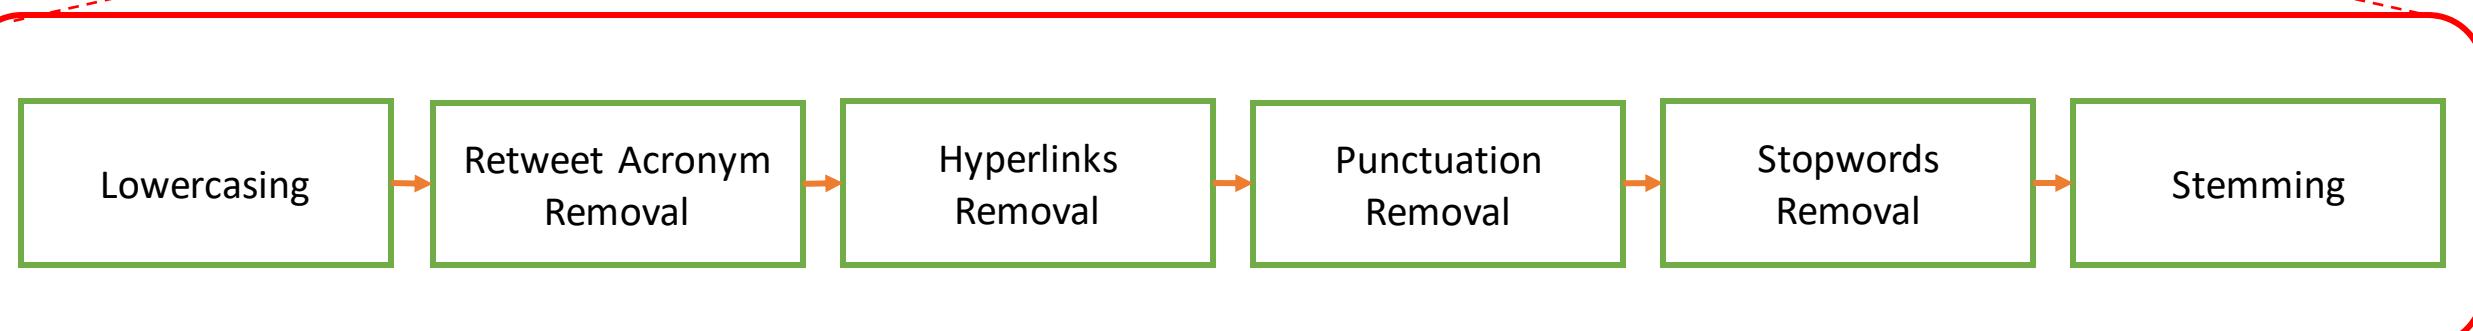
Sentiment Analysis

❖ Step 4: Data preprocessing

Input



Output



Sentiment Analysis

❖ Step 4: Data preprocessing

Lowercasing: Convert text into lowercase.

```
1 text = 'Hello WORLD'  
2 text = text.lower()  
3 print(text)
```

hello world

Retweet Acronym Removal: Remove "RT" at the start of the string

```
1 text = "RT Hello World"  
2 text = re.sub(r'^RT[\s]+', '', text)  
3 print(text)
```

Hello World

Punctuation Removal: Remove all punctuation symbols "#".

```
1 text = "#xinchoa, #hello, #konnichiwa"  
2 text = re.sub(r'[^w\s]', '', text)  
3 print(text)
```

xinchoa hello konnichiwa

Hyperlinks Removal: Remove all hyperlinks in the string.

```
1 text = "http://fb.me/6N3LsUpCu"  
2 text = re.sub(  
3     r'https?:\/\/.*[\r\n]*',  
4     '',  
5     text  
6 )  
7 print(text)
```

Sentiment Analysis

❖ Step 4: Data preprocessing

```
> stopwords("english")
[1] "i"          "me"         "my"        "myself"      "we"
[6] "our"        "ours"       "ourselves" "you"        "your"
[11] "yours"      "yourself"    "yourselves" "he"         "him"
[16] "his"        "himself"    "she"       "her"        "hers"
[21] "herself"    "it"         "its"       "itself"     "they"
[26] "them"        "their"      "theirs"     "themselves" "what"
[31] "which"       "who"        "whom"      "this"       "that"
[36] "these"       "those"      "am"         "is"         "are"
[41] "was"         "were"      "be"         "been"      "being"
[46] "have"        "has"        "had"       "having"     "do"
```

```
1 text = 'Hello world we are one'
2 stop_words = set(stopwords.words('english'))
3 words = text.split()
4 words = [
5     word for word in words \
6     if word not in stop_words
7 ]
8 text = ' '.join(words)
9 print(text)
```

Hello world one

Stopwords Removal: Remove all “stopwords” which are not necessary to represent in vocabulary.

Sentiment Analysis

❖ Step 4: Data preprocessing

Snowball Stemmer Rules

Ends with	Reduced to	
ILY	→ ILI	easily → easili
LY	→	fairly → fair
SS	→ SS	caress → caress
S	→	bats → bat
ING	→	singing → sing
ED	→ E	shared → share

```
1 text = 'I am eating breakfast'
2 stemmer = SnowballStemmer('english')
3 words = text.split()
4 words = [stemmer.stem(word) for word in words]
5 text = ' '.join(words)
6 print(text)
```

i am eat breakfast

Stemming: Convert all the words into its root form
(only use in English or similar languages).

Sentiment Analysis

❖ Step 4: Data preprocessing

```
1 def text_normalize(text):
2     # Lowercasing
3     text = text.lower()
4
5     # Retweet old acronym "RT" removal
6     text = re.sub(r'^rt[\s]+', '', text)
7
8     # Hyperlinks removal
9     text = re.sub(r'https?://.*[\r\n]*', '', text)
10
11    # Punctuation removal
12    text = re.sub(r'[^w\s]', '', text)
13
14    # Remove stopwords
15    stop_words = set(stopwords.words('english'))
16    words = text.split()
17    words = [word for word in words if word not in stop_words]
18    text = ' '.join(words)
19
20    # Stemming
21    stemmer = SnowballStemmer('english')
22    words = text.split()
23    words = [stemmer.stem(word) for word in words]
24    text = ' '.join(words)
25
26
return text
```

```
1 text = """We love this! Would you go?
2 #talk #makememories #unplug
3 #relax #iphone #smartphone #wifi #connect...
4 http://fb.me/6N3LsUpCu
5 """
6 text = text_normalize(text)
7 text
```

'love would go talk makememori unplug relax iphon smartphon wifi connect'

Sentiment Analysis

❖ Step 4: Data preprocessing

$$w_{x,y} = tf_{x,y} \times \log \left(\frac{N}{df_x} \right)$$

TF-IDF

Term x within document y

$tf_{x,y}$ = frequency of x in y

df_x = number of documents containing x

N = total number of documents

To represent normalized text as vector, we will use **Tfidf vectorizer** from sklearn.

```
1 vectorizer = TfidfVectorizer(  
2     max_features=2000  
3 )  
4 X = vectorizer.fit_transform(  
5     df['clean_text']  
6 ).toarray()
```

Sentiment Analysis

❖ Step 4: Data preprocessing

```
1 intercept = np.ones( (  
2 |     X.shape[0], 1)  
3 )  
4 X_b = np.concatenate(  
5 |     (intercept, X),  
6 |     axis=1  
7 )
```

We also add the bias term into each TfIdf vector.

1 intercept	1 x	1 x_b
array([[1.], [1.], [1.], [1.], [1.], [1.], [1.]])	array([[3., 0., 1., [1., 1., 1., [3., 1., 1., [1., 1., 1., [1., 1., 1., [3., 1., 1., [1., 0., 0., [3., 0., 0.]])	array([[1., 3., 1., [1., 1., 1.]])

Sentiment Analysis

❖ Step 5: One-hot encode label

Label encoding

Label
0
0
1
3
2
1
2

One-hot encoding

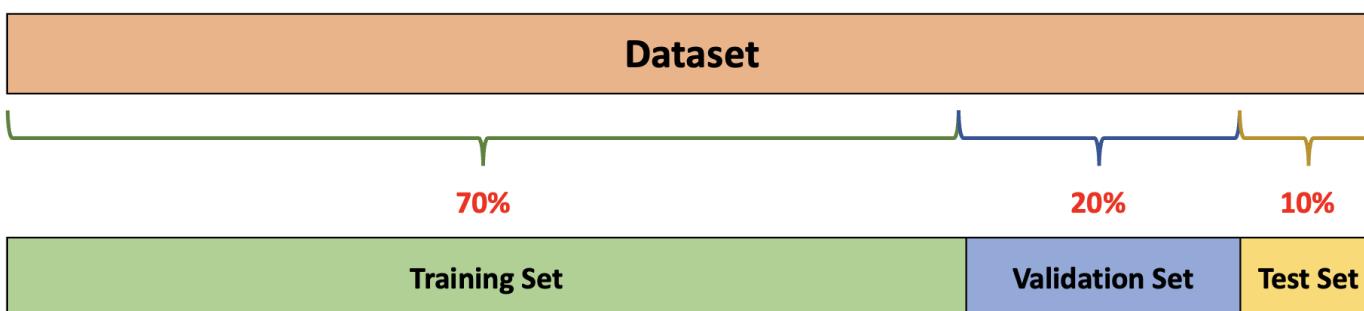
Label	1	0	0	0
0	1	0	0	0
0	1	0	0	0
1	0	1	0	0
3	0	0	0	1
2	0	0	1	0
1	0	1	0	0
2	0	0	1	0



```
1 n_classes = df['category'].nunique()
2 n_samples = df['category'].size
3
4 y = df['category'].to_numpy() + 1
5 y = y.astype(np.uint8)
6 y_encoded = np.array(
7 | [np.zeros(n_classes) for _ in range(n_samples)]
8 )
9 y_encoded[np.arange(n_samples), y] = 1
```

Sentiment Analysis

❖ Step 6: Split train, val, test set



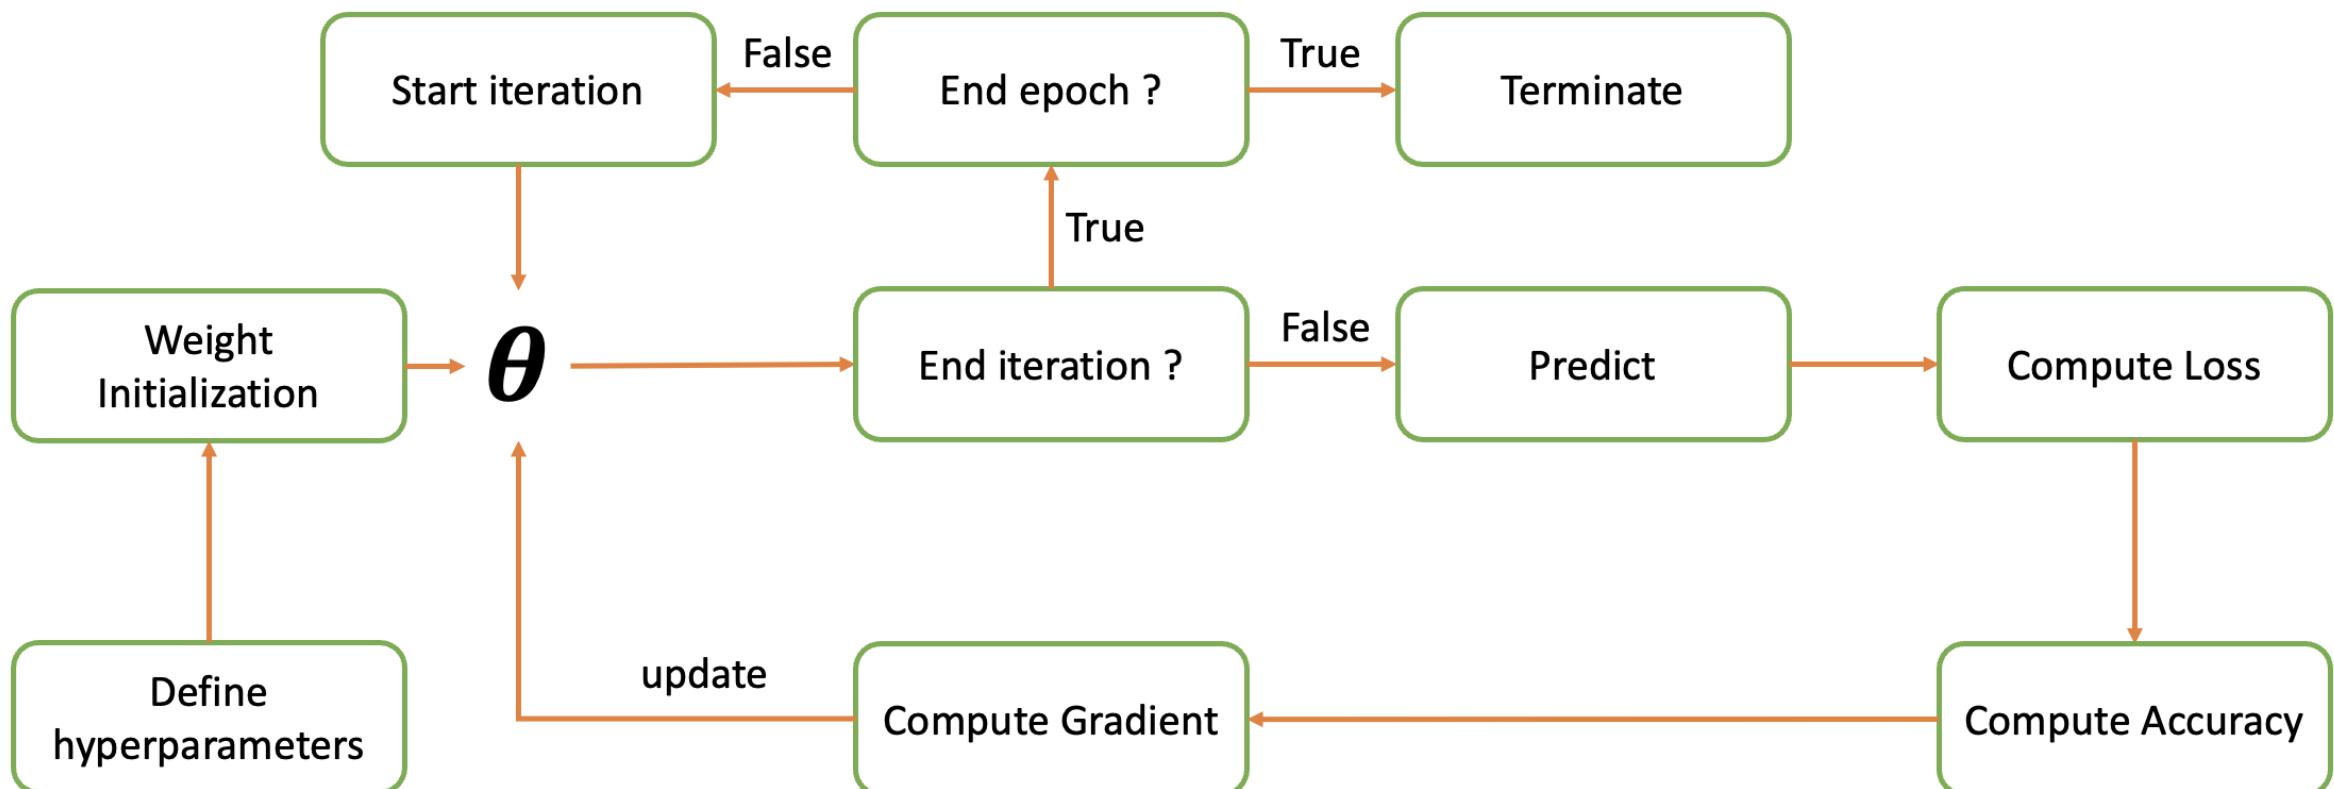
```
1 val_size = 0.3
2 test_size = 0.1
3 random_state = 2
4 is_shuffle = True
5
6 x_train, x_val, y_train, y_val = train_test_split(
7     X_b, y_encoded,
8     test_size=val_size,
9     random_state=random_state,
10    shuffle=is_shuffle
11 )
12
13 x_val, x_test, y_val, y_test = train_test_split(
14     x_val, y_val,
15     test_size=test_size,
16     random_state=random_state,
17     shuffle=is_shuffle
18 )
```

```
1 print(f'Number of training samples: {x_train.shape[0]}')
2 print(f'Number of val samples: {x_val.shape[0]}')
3 print(f'Number of test samples: {x_test.shape[0]}')
```

Number of training samples: 114078
Number of val samples: 44001
Number of test samples: 4890

Sentiment Analysis

❖ Step 7: Training



Sentiment Analysis

❖ Step 7: Training

```
1 lr = 0.1
2 epochs = 200
3 batch_size = X_train.shape[0]
4 n_features = X_train.shape[1]
5
6 np.random.seed(random_state)
7 theta = np.random.uniform(
8     size=(n_features, n_classes))
9 )
```

Define essential functions and hyperparameters

```
1 def softmax(z):
2     exp_z = np.exp(z)
3
4     return exp_z / exp_z.sum(axis=1)[:, None]
5
6 def compute_loss(y_hat, y):
7     n = y.size
8
9     return (-1 / n) * np.sum(y * np.log(y_hat))
10
11 def predict(X, theta):
12     z = np.dot(X, theta)
13     y_hat = softmax(z)
14
15     return y_hat
16
17 def compute_gradient(X, y, y_hat):
18     n = y.size
19
20     return np.dot(X.T, (y_hat - y)) / n
21
22 def update_theta(theta, gradient, lr):
23     return theta - lr * gradient
24
25 def compute_accuracy(X, y, theta):
26     y_hat = predict(X, theta)
27     acc = (np.argmax(y_hat, axis=1) == np.argmax(y, axis=1)).mean()
28
29     return acc
```

Sentiment Analysis

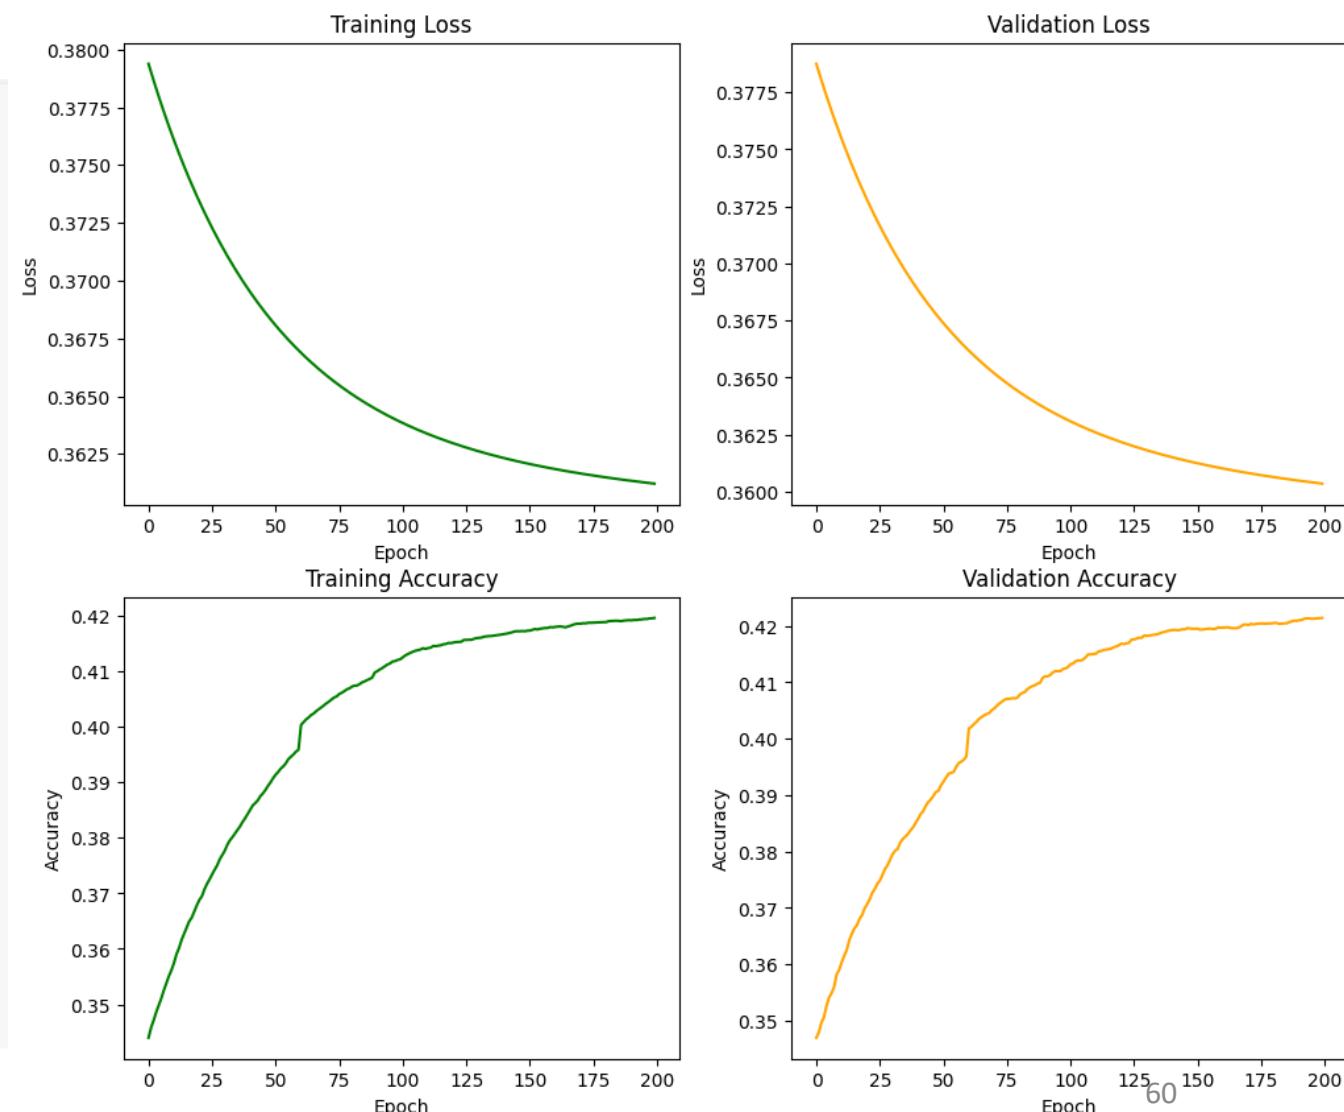
❖ Step 7: Training

```
6 for epoch in range(epochs):
7     train_batch_losses = []
8     train_batch_accs = []
9     val_batch_losses = []
10    val_batch_accs = []
11
12    for i in range(0, X_train.shape[0], batch_size):
13        X_i = X_train[i:i+batch_size]
14        y_i = y_train[i:i+batch_size]
15
16        y_hat = predict(X_i, theta)
17        train_loss = compute_loss(y_hat, y_i)
18        gradient = compute_gradient(X_i, y_i, y_hat)
19        theta = update_theta(theta, gradient, lr)
20        train_batch_losses.append(train_loss)
21
22        train_acc = compute_accuracy(X_train, y_train, theta)
23        train_batch_accs.append(train_acc)
24
25        y_val_hat = predict(X_val, theta)
26        val_loss = compute_loss(y_val_hat, y_val)
27        val_batch_losses.append(val_loss)
28
29        val_acc = compute_accuracy(X_val, y_val, theta)
30        val_batch_accs.append(val_acc)
31
32        train_batch_loss = sum(train_batch_losses) / len(train_batch_losses)
33        val_batch_loss = sum(val_batch_losses) / len(val_batch_losses)
34        train_batch_acc = sum(train_batch_accs) / len(train_batch_accs)
35        val_batch_acc = sum(val_batch_accs) / len(val_batch_accs)
```

Sentiment Analysis

❖ Step 7: Training (Visualization)

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses)
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, 'orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs)
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, 'orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



Sentiment Analysis

❖ Step 8: Evaluation

$$\text{accuracy} = \frac{\text{true_predictions}}{\text{n_samples}}$$

```
1 def compute_accuracy(X, y, theta):  
2     y_hat = predict(X, theta).round()  
3     acc = (y_hat == y).mean()  
4  
5     return acc
```

```
1 # Val set  
2 val_set_acc = compute_accuracy(X_val, y_val, theta)  
3 print('Evaluation on validation set: ')  
4 print(f'Accuracy: {val_set_acc}')
```

Evaluation on validation set:
Accuracy: 0.4214222404036272

```
1 # Test set  
2 test_set_acc = compute_accuracy(X_test, y_test, theta)  
3 print('Evaluation on test set: ')  
4 print(f'Accuracy: {test_set_acc}')
```

Evaluation on test set:
Accuracy: 0.416359918200409

Sentiment Analysis

❖ Implement with PyTorch: Step 1

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import re
5 import torch
6 import torch.nn as nn
7 import torch.optim as optim
8 import nltk
9 nltk.download('stopwords')
10
11 from sklearn.model_selection import train_test_split
12 from sklearn.feature_extraction.text import TfidfVectorizer
13 from nltk.corpus import stopwords
14 from nltk.stem import SnowballStemmer
```



Sentiment Analysis

❖ Implement with PyTorch: Step 2

```
1 dataset_path = 'Twitter_Data.csv'  
2 df = pd.read_csv(  
3     dataset_path  
4 )  
5 df
```

	clean_text	category
0	when modi promised “minimum government maximum...	-1.0
1	talk all the nonsense and continue all the dra...	0.0
2	what did just say vote for modi welcome bjp t...	1.0
3	asking his supporters prefix chowkidar their n...	1.0
4	answer who among these the most powerful world...	1.0
...
162975	why these 456 crores paid neerav modi not reco...	-1.0
162976	dear rss terrorist payal gawar what about modi...	-1.0
162977	did you cover her interaction forum where she ...	0.0
162978	there big project came into india modi dream p...	0.0
162979	have you ever listen about like gurukul where ...	1.0

162980 rows × 2 columns

Sentiment Analysis

❖ Implement with PyTorch: Step 3

```
1 null_rows = df.isnull().any(axis=1)  
2 df[null_rows]
```

		clean_text	category
148		NaN	0.0
130448	the foundation stone northeast gas grid inaugu...	NaN	
155642	dear terrorists you can run but you cant hide ...	NaN	
155698	offense the best defence with mission shakti m...	NaN	
155770	have always heard politicians backing out thei...	NaN	
158693	modi government plans felicitate the faceless ...	NaN	
158694		NaN	-1.0
159442	chidambaram gives praises modinomics	NaN	
159443		NaN	0.0
160559	the reason why modi contested from seats 2014 ...	NaN	
160560		NaN	1.0

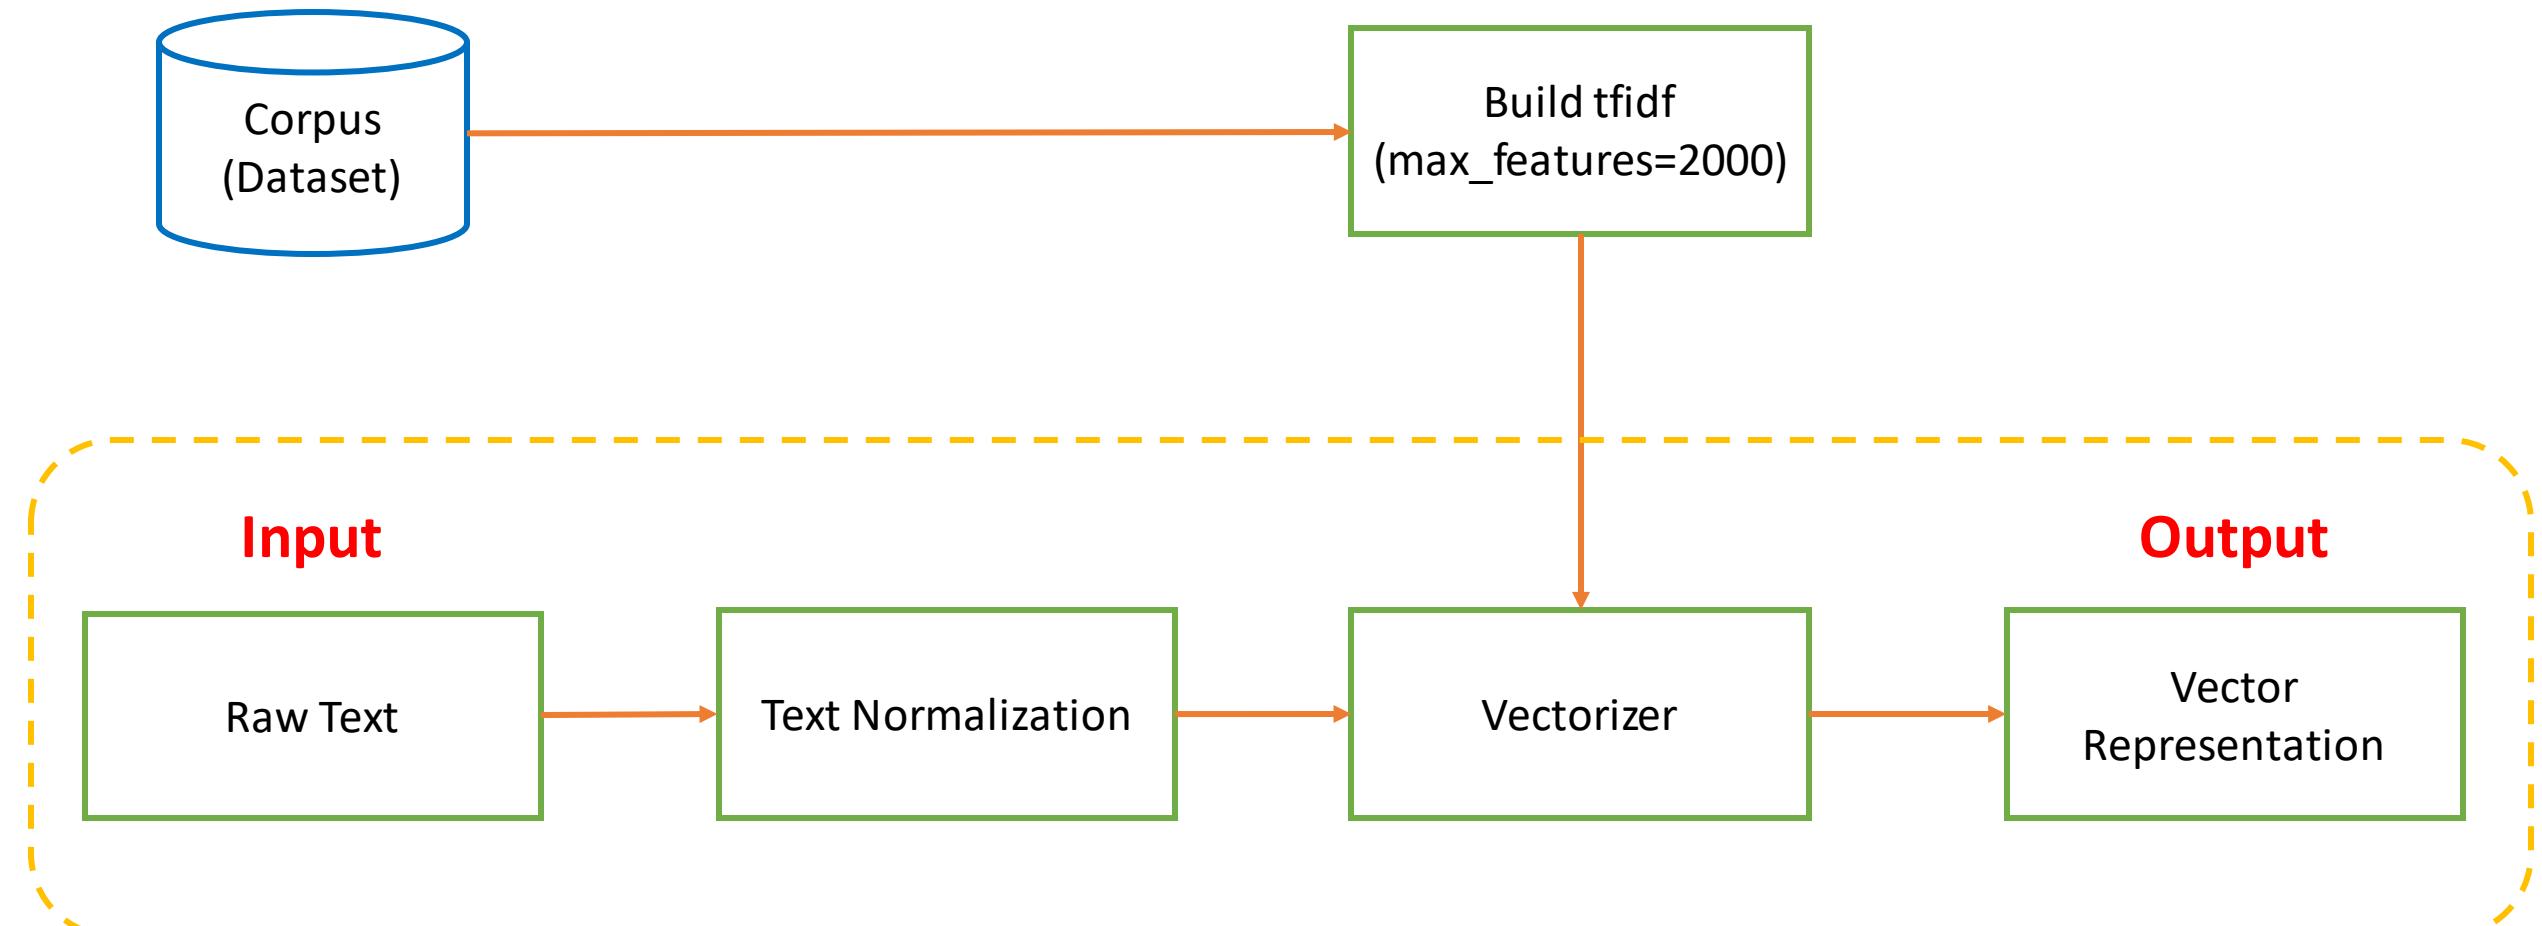
```
1 df = df.dropna()
```

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 162969 entries, 0 to 162979  
Data columns (total 2 columns):  
 #   Column           Non-Null Count  Dtype     
---  --     
 0   clean_text        162969 non-null   object    
 1   category         162969 non-null   float64  
dtypes: float64(1), object(1)  
memory usage: 3.7+ MB
```

Sentiment Analysis

❖ Implement with PyTorch: Step 4



Sentiment Analysis

❖ Implement with PyTorch: Step 4

```
1 def text_normalize(text):
2     # Lowercasing
3     text = text.lower()
4
5     # Retweet old acronym "RT" removal
6     text = re.sub(r'^rt[\s]+', '', text)
7
8     # Hyperlinks removal
9     text = re.sub(r'https?:\/\/.*[\r\n]*', '', text)
10
11    # Punctuation removal
12    text = re.sub(r'[^w\s]', '', text)
13
14    # Remove stopwords
15    stop_words = set(stopwords.words('english'))
16    words = text.split()
17    words = [word for word in words if word not in stop_words]
18    text = ' '.join(words)
19
20    # Stemming
21    stemmer = SnowballStemmer('english')
22    words = text.split()
23    words = [stemmer.stem(word) for word in words]
24    text = ' '.join(words)
25
26
return text
```

```
1 df['clean_text'] = df['clean_text'].apply(
2 |     lambda x: text_normalize(x)
3 )
```

```
1 vectorizer = TfidfVectorizer(
2 |     max_features=2000
3 )
4 X = vectorizer.fit_transform(
5 |     df['clean_text']
6 ).toarray()
```

Sentiment Analysis

❖ Implement with PyTorch: Step 5

Label encoding

Label
0
0
1
3
2
1
2

One-hot encoding

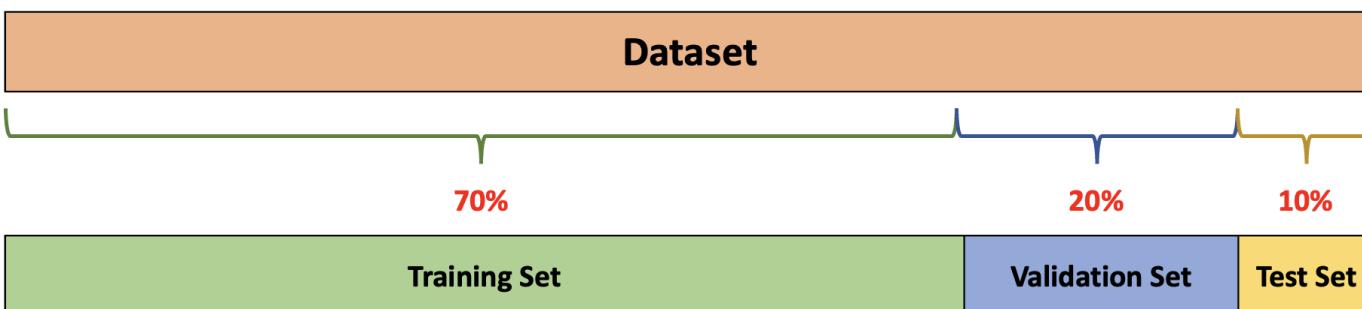
Label	1	0	0	0
0	1	0	0	0
0	1	0	0	0
1	0	1	0	0
3	0	0	0	1
2	0	0	1	0
1	0	1	0	0
2	0	0	1	0



```
1 n_classes = df['category'].nunique()
2 n_samples = df['category'].size
3
4 y = df['category'].to_numpy() + 1
5 y = y.astype(np.uint8)
6 y_encoded = np.array(
7 | [np.zeros(n_classes) for _ in range(n_samples)]
8 )
9 y_encoded[np.arange(n_samples), y] = 1
```

Sentiment Analysis

❖ Implement with PyTorch: Step 6



```
1 X = torch.tensor(X, dtype=torch.float32)
2 y = torch.tensor(y_encoded, dtype=torch.float32)
```

```
1 val_size = 0.3
2 test_size = 0.1
3 random_state = 2
4 is_shuffle = True
5
6 X_train, X_val, y_train, y_val = train_test_split(
7     X_b, y_encoded,
8     test_size=val_size,
9     random_state=random_state,
10    shuffle=is_shuffle
11 )
12
13 X_val, X_test, y_val, y_test = train_test_split(
14     X_val, y_val,
15     test_size=test_size,
16     random_state=random_state,
17     shuffle=is_shuffle
18 )
```

```
1 print(f'Number of training samples: {X_train.shape[0]}')
2 print(f'Number of val samples: {X_val.shape[0]}')
3 print(f'Number of test samples: {X_test.shape[0]}')
```

Number of training samples: 114078
Number of val samples: 44001
Number of test samples: 4890

Sentiment Analysis

❖ Implement with PyTorch: Step 7

Create a SoftmaxRegression class that inherits from nn.Module:

```
1 class SoftmaxRegression(nn.Module):
2     def __init__(self, input_dim, output_dim):
3         super(SoftmaxRegression, self).__init__()
4         self.linear = nn.Linear(
5             input_dim,
6             output_dim,
7             bias=True
8         )
9
10    def forward(self, x):
11        return self.linear(x)
```

Applies a linear transformation to the incoming data: $y = xA^T + b$

This module supports `TensorFloat32`.

On certain ROCm devices, when using float16 inputs this module will use `different precision` for backward.

Parameters

- `in_features` (`int`) – size of each input sample
- `out_features` (`int`) – size of each output sample
- `bias` (`bool`) – If set to `False`, the layer will not learn an additive bias. Default: `True`

We define a linear transformation, which is just the dot product (remember to set `bias=True` to include the bias term).

Sentiment Analysis

❖ Implement with PyTorch: Step 8

```
1 lr = 0.1
2 epochs = 500
3 torch.manual_seed(random_state)
4 if torch.cuda.is_available():
5     torch.cuda.manual_seed(random_state)
6
7 input_dim = X_train.shape[1]
8 output_dim = y_train.shape[1]
9
10 model = SoftmaxRegression(
11     input_dim, output_dim
12 )
13 criterion = nn.CrossEntropyLoss()
14 optimizer = optim.SGD(
15     model.parameters(), lr=lr
16 )
```

We define the loss function (called criterion in PyTorch) using Cross-entropy and optimizer algorithm (in this case is SGD).

Sentiment Analysis

❖ Implement with PyTorch: Step 8

```
1 train_losses = []
2 train_accs = []
3 val_losses = []
4 val_accs = []

5
6 for epoch in range(epochs):
7     model.train()

8
9     # Zero the gradients
10    optimizer.zero_grad()

11
12    # Forward pass
13    y_hat = model(X_train)

14
15    # Compute loss
16    train_loss = criterion(y_hat, y_train)

17
18    train_losses.append(train_loss.item())

19
20    train_acc = compute_accuracy(y_hat, y_train)
21    train_accs.append(train_acc)

22
23    # Backward pass and optimization
24    train_loss.backward()
25    optimizer.step()

26
27    model.eval()
28    # Forward pass for validation data
29    with torch.no_grad():
30        y_val_hat = model(X_val)

31
32        # Compute validation loss
33        val_loss = criterion(y_val_hat, y_val)
34        val_losses.append(val_loss.item())

35
36        val_acc = compute_accuracy(y_val_hat, y_val)
37        val_accs.append(val_acc)

38
39        print(f'\nEPOCH {epoch + 1}:\tTraining loss: {train_loss:.3f}\tValidation loss: {val_loss:.3f}')
```

Line 1, 2, 3, 4: Define empty lists for storing losses and accuracies of train and val sets in each epoch.

Sentiment Analysis

❖ Implement with PyTorch: Step 8

```
1 train_losses = []
2 train_accs = []
3 val_losses = []
4 val_accs = []
5
6 for epoch in range(epochs):
7     model.train()
8
9     # Zero the gradients
10    optimizer.zero_grad()
11
12    # Forward pass
13    y_hat = model(X_train)
14
15    # Compute loss
16    train_loss = criterion(y_hat, y_train)
17
18    train_losses.append(train_loss.item())
19
20    train_acc = compute_accuracy(y_hat, y_train)
21    train_accs.append(train_acc)
22
23    # Backward pass and optimization
24    train_loss.backward()
25    optimizer.step()
26
27    model.eval()
28    # Forward pass for validation data
29    with torch.no_grad():
30        y_val_hat = model(X_val)
31
32        # Compute validation loss
33        val_loss = criterion(y_val_hat, y_val)
34        val_losses.append(val_loss.item())
35
36        val_acc = compute_accuracy(y_val_hat, y_val)
37        val_accs.append(val_acc)
38
39        print(f'\nEPOCH {epoch + 1}:\tTraining loss: {train_loss:.3f}\tValidation loss: {val_loss:.3f}')
```

Line 9, 10: Zero out the gradients of previous epoch before any computations.

Sentiment Analysis

❖ Implement with PyTorch: Step 8

```
1 train_losses = []
2 train_accs = []
3 val_losses = []
4 val_accs = []
5
6 for epoch in range(epochs):
7     model.train()
8
9     # Zero the gradients
10    optimizer.zero_grad()
11
12    # Forward pass
13    y_hat = model(X_train)
14
15    # Compute loss
16    train_loss = criterion(y_hat, y_train)
17
18    train_losses.append(train_loss.item())
19
20    train_acc = compute_accuracy(y_hat, y_train)
21    train_accs.append(train_acc)
22
23    # Backward pass and optimization
24    train_loss.backward()
25    optimizer.step()
26
27    model.eval()
28    # Forward pass for validation data
29    with torch.no_grad():
30        y_val_hat = model(X_val)
31
32        # Compute validation loss
33        val_loss = criterion(y_val_hat, y_val)
34        val_losses.append(val_loss.item())
35
36        val_acc = compute_accuracy(y_val_hat, y_val)
37        val_accs.append(val_acc)
38
39        print(f'\nEPOCH {epoch + 1}:\tTraining loss: {train_loss:.3f}\tValidation loss: {val_loss:.3f}')
```

Line 12, 13: Do prediction on all data points of current epoch.

Sentiment Analysis

❖ Implement with PyTorch: Step 8

Line 15, 16: Calculating the training loss and append the result to train_losses list.
Line 20, 21: Calculate the accuracy on the train set.

```
1 train_losses = []
2 train_accs = []
3 val_losses = []
4 val_accs = []
5
6 for epoch in range(epochs):
7     model.train()
8
9     # Zero the gradients
10    optimizer.zero_grad()
11
12    # Forward pass
13    y_hat = model(X_train)
14
15    # Compute loss
16    train_loss = criterion(y_hat, y_train)
17
18    train_losses.append(train_loss.item())
19
20    train_acc = compute_accuracy(y_hat, y_train)
21    train_accs.append(train_acc)
22
23    # Backward pass and optimization
24    train_loss.backward()
25    optimizer.step()
26
27    model.eval()
28    # Forward pass for validation data
29    with torch.no_grad():
30        y_val_hat = model(X_val)
31
32        # Compute validation loss
33        val_loss = criterion(y_val_hat, y_val)
34        val_losses.append(val_loss.item())
35
36        val_acc = compute_accuracy(y_val_hat, y_val)
37        val_accs.append(val_acc)
38
39        print(f'\nEPOCH {epoch + 1}:\tTraining loss: {train_loss:.3f}\tValidation loss: {val_loss:.3f}')
```

Sentiment Analysis

❖ Implement with PyTorch: Step 8

Line 23, 24, 25: Compute the gradient and update the weights.

```
1 train_losses = []
2 train_accs = []
3 val_losses = []
4 val_accs = []
5
6 for epoch in range(epochs):
7     model.train()
8
9     # Zero the gradients
10    optimizer.zero_grad()
11
12    # Forward pass
13    y_hat = model(X_train)
14
15    # Compute loss
16    train_loss = criterion(y_hat, y_train)
17
18    train_losses.append(train_loss.item())
19
20    train_acc = compute_accuracy(y_hat, y_train)
21    train_accs.append(train_acc)
22
23    # Backward pass and optimization
24    train_loss.backward()
25    optimizer.step()
26
27    model.eval()
28    # Forward pass for validation data
29    with torch.no_grad():
30        y_val_hat = model(X_val)
31
32        # Compute validation loss
33        val_loss = criterion(y_val_hat, y_val)
34        val_losses.append(val_loss.item())
35
36        val_acc = compute_accuracy(y_val_hat, y_val)
37        val_accs.append(val_acc)
38
39        print(f'\nEPOCH {epoch + 1}:\tTraining loss: {train_loss:.3f}\tValidation loss: {val_loss:.3f}')
```

Sentiment Analysis

❖ Implement with PyTorch: Step 8

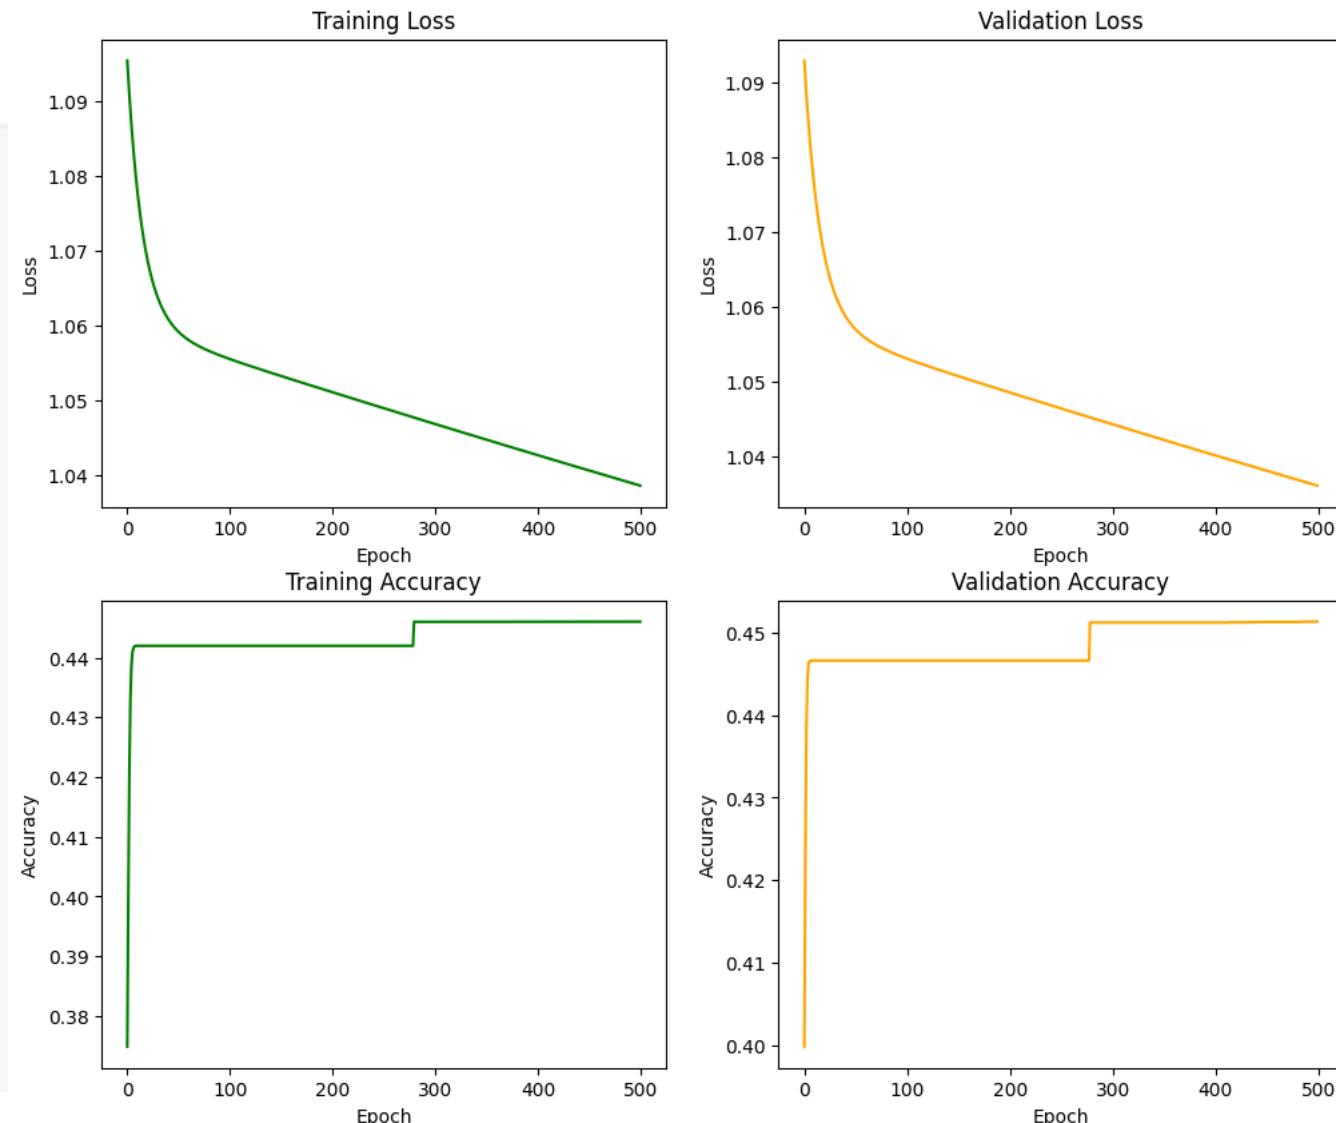
From line 27 to line 37: Change the model to evaluation mode and calculate the loss and accuracy on val set. Then append those scores into according lists.

```
1 train_losses = []
2 train_accs = []
3 val_losses = []
4 val_accs = []
5
6 for epoch in range(epochs):
7     model.train()
8
9     # Zero the gradients
10    optimizer.zero_grad()
11
12    # Forward pass
13    y_hat = model(X_train)
14
15    # Compute loss
16    train_loss = criterion(y_hat, y_train)
17
18    train_losses.append(train_loss.item())
19
20    train_acc = compute_accuracy(y_hat, y_train)
21    train_accs.append(train_acc)
22
23    # Backward pass and optimization
24    train_loss.backward()
25    optimizer.step()
26
27    model.eval()
28    # Forward pass for validation data
29    with torch.no_grad():
30        y_val_hat = model(X_val)
31
32        # Compute validation loss
33        val_loss = criterion(y_val_hat, y_val)
34        val_losses.append(val_loss.item())
35
36        val_acc = compute_accuracy(y_val_hat, y_val)
37        val_accs.append(val_acc)
38
39        print(f'\nEPOCH {epoch + 1}:\tTraining loss: {train_loss:.3f}\tValidation loss: {val_loss:.3f}')
```

Sentiment Analysis

❖ Implement with PyTorch: Step 9

```
1 fig, ax = plt.subplots(2, 2, figsize=(12, 10))
2 ax[0, 0].plot(train_losses)
3 ax[0, 0].set(xlabel='Epoch', ylabel='Loss')
4 ax[0, 0].set_title('Training Loss')
5
6 ax[0, 1].plot(val_losses, 'orange')
7 ax[0, 1].set(xlabel='Epoch', ylabel='Loss')
8 ax[0, 1].set_title('Validation Loss')
9
10 ax[1, 0].plot(train_accs)
11 ax[1, 0].set(xlabel='Epoch', ylabel='Accuracy')
12 ax[1, 0].set_title('Training Accuracy')
13
14 ax[1, 1].plot(val_accs, 'orange')
15 ax[1, 1].set(xlabel='Epoch', ylabel='Accuracy')
16 ax[1, 1].set_title('Validation Accuracy')
17
18 plt.show()
```



Sentiment Analysis

❖ Implement with PyTorch: Step 10

```
1 def compute_accuracy(y_hat, y_true):
2     _, y_hat = torch.max(y_hat, dim=1)
3     _, y_true = torch.max(y_true, dim=1)
4
5     correct = (y_hat == y_true).sum().item()
6
7     accuracy = (correct / len(y_true))
8
9     return accuracy
```

```
1 # Val set
2 model.eval()
3 with torch.no_grad():
4     y_hat = model(X_val)
5     val_set_acc = compute_accuracy(y_hat, y_val)
6     print('Evaluation on validation set:')
7     print(f'Accuracy: {val_set_acc}')
```

Evaluation on validation set:
Accuracy: 0.4512851980636804

```
1 # Test set
2 model.eval()
3 with torch.no_grad():
4     y_hat = model(X_test)
5     test_set_acc = compute_accuracy(y_hat, y_test)
6     print('Evaluation on test set:')
7     print(f'Accuracy: {test_set_acc}')
```

Evaluation on test set:
Accuracy: 0.4494887525562372

Question

