

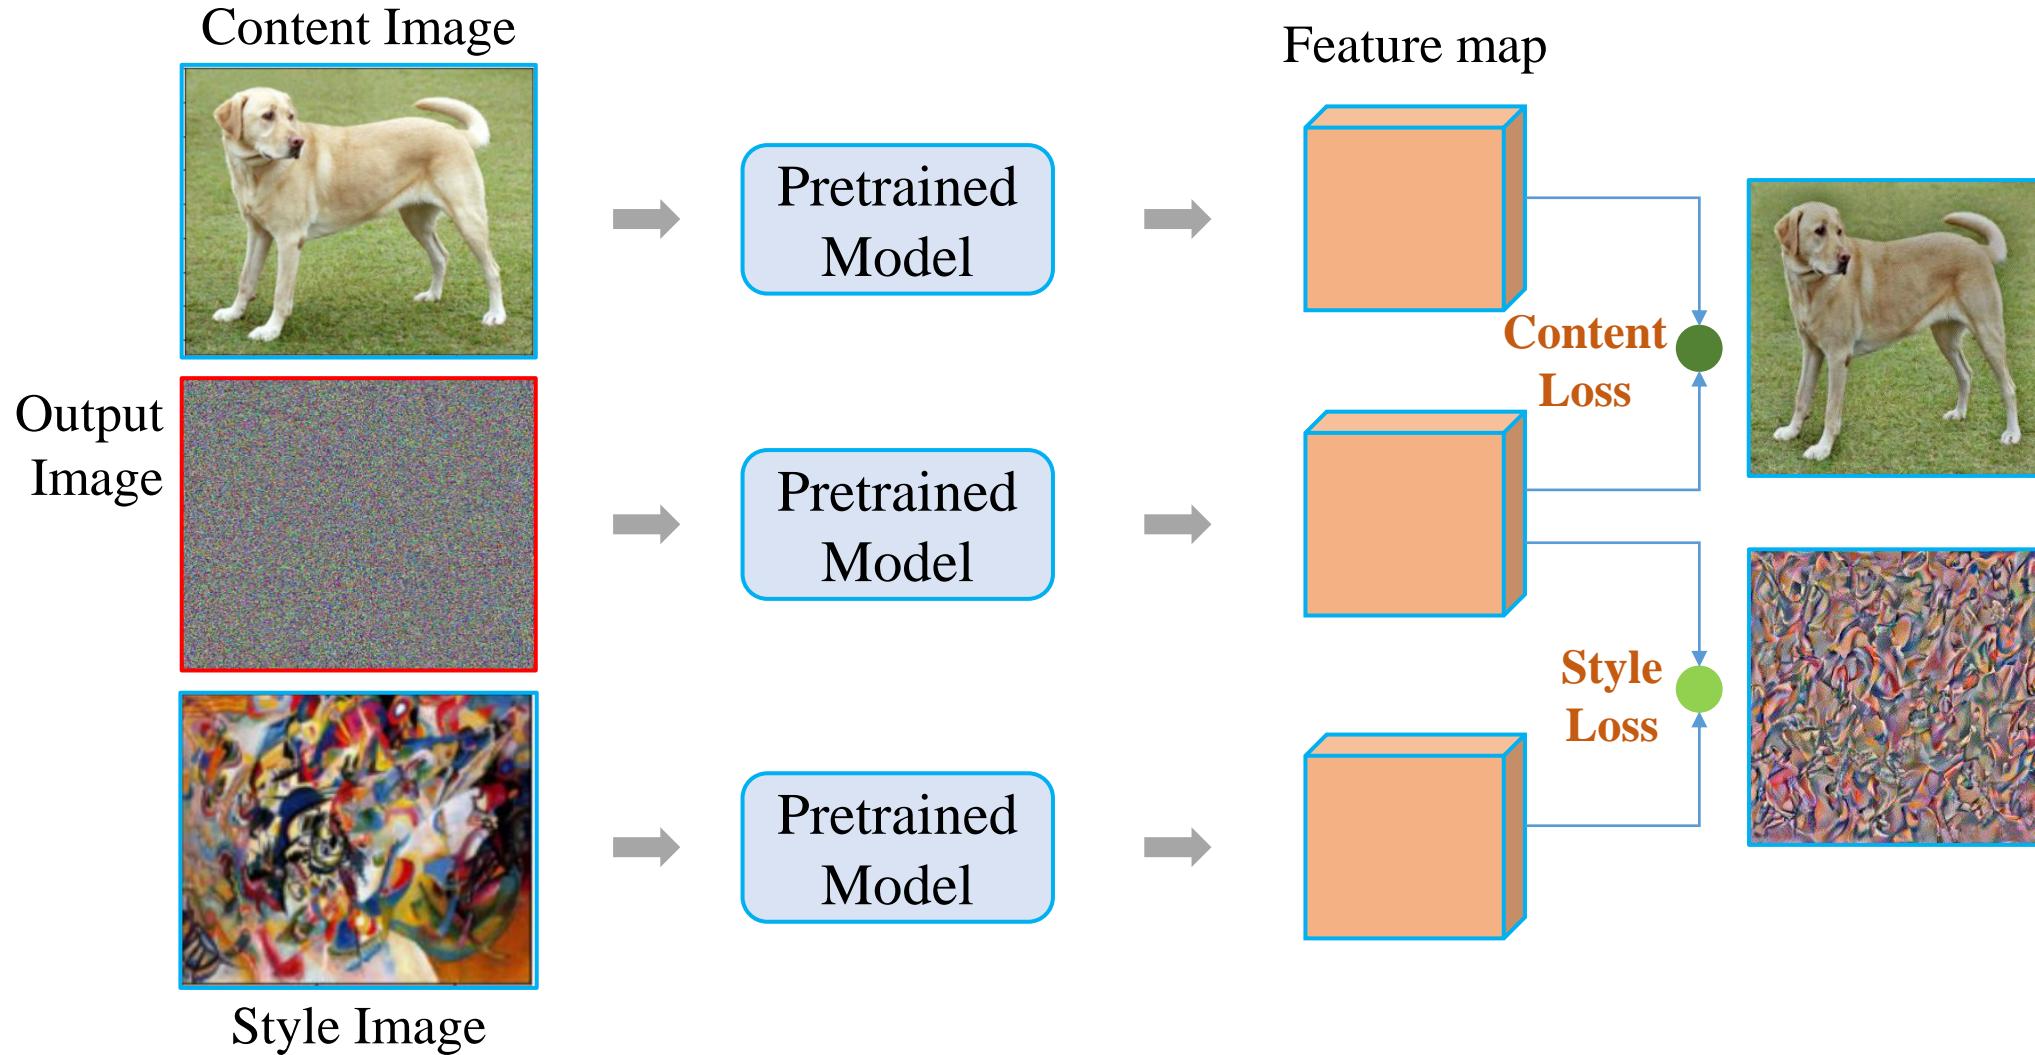
Style Transfer

Understanding Content Loss and Style Loss

Quang-Vinh Dinh
Ph.D. in Computer Science

Objectives

- ✓ Understand deeply about content loss
- ✓ Understand deeply about style loss

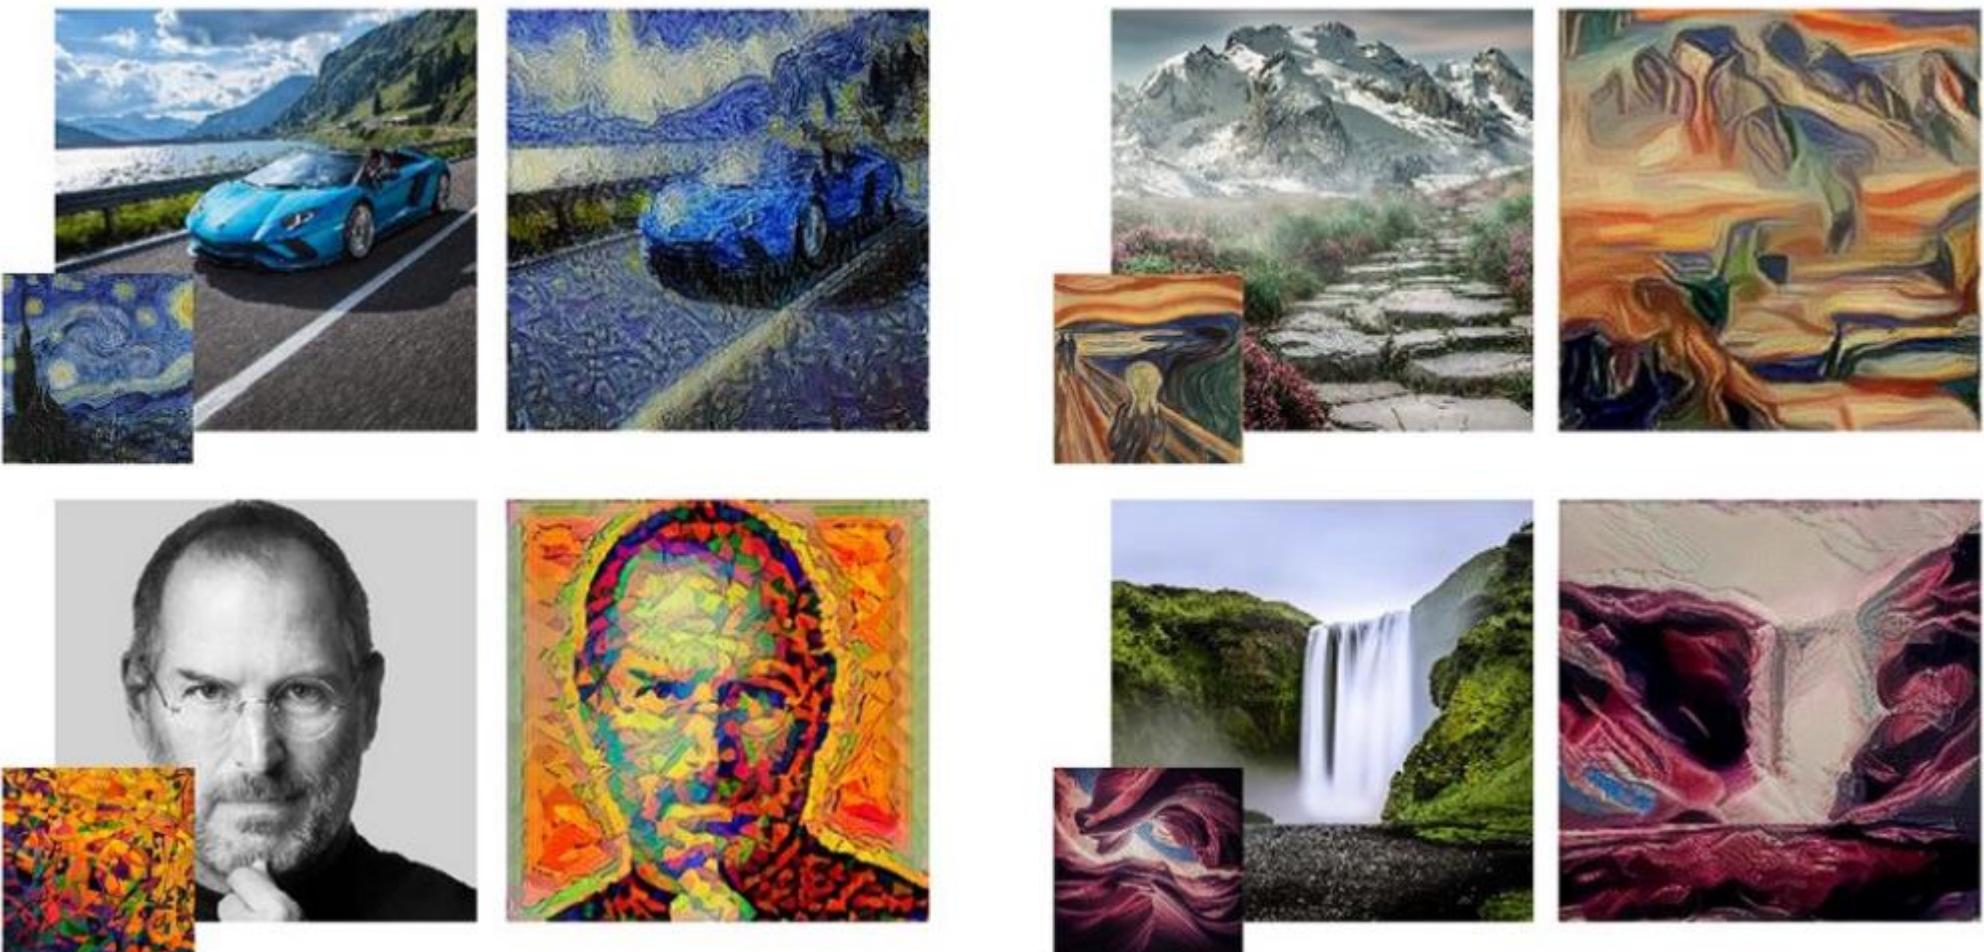


Outline

- Introduction
- Inputs as Variables
- Content Loss
- Style Loss

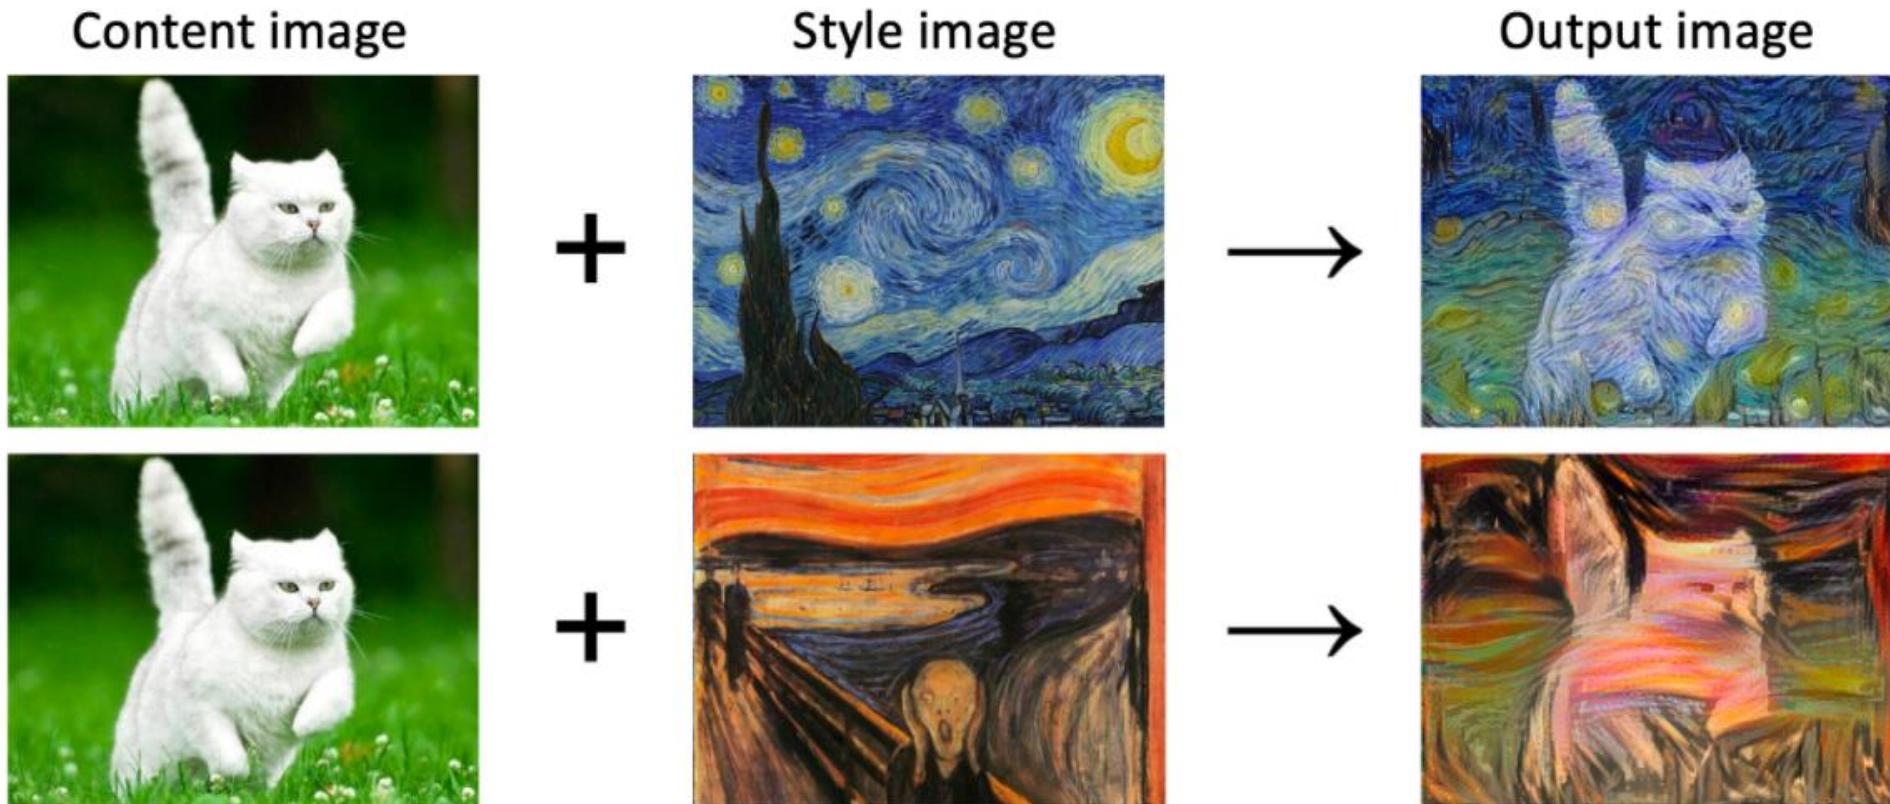
Style Transfer

❖ From different viewpoints



Style Transfer

❖ From different viewpoints



Style Transfer

❖ From different viewpoints

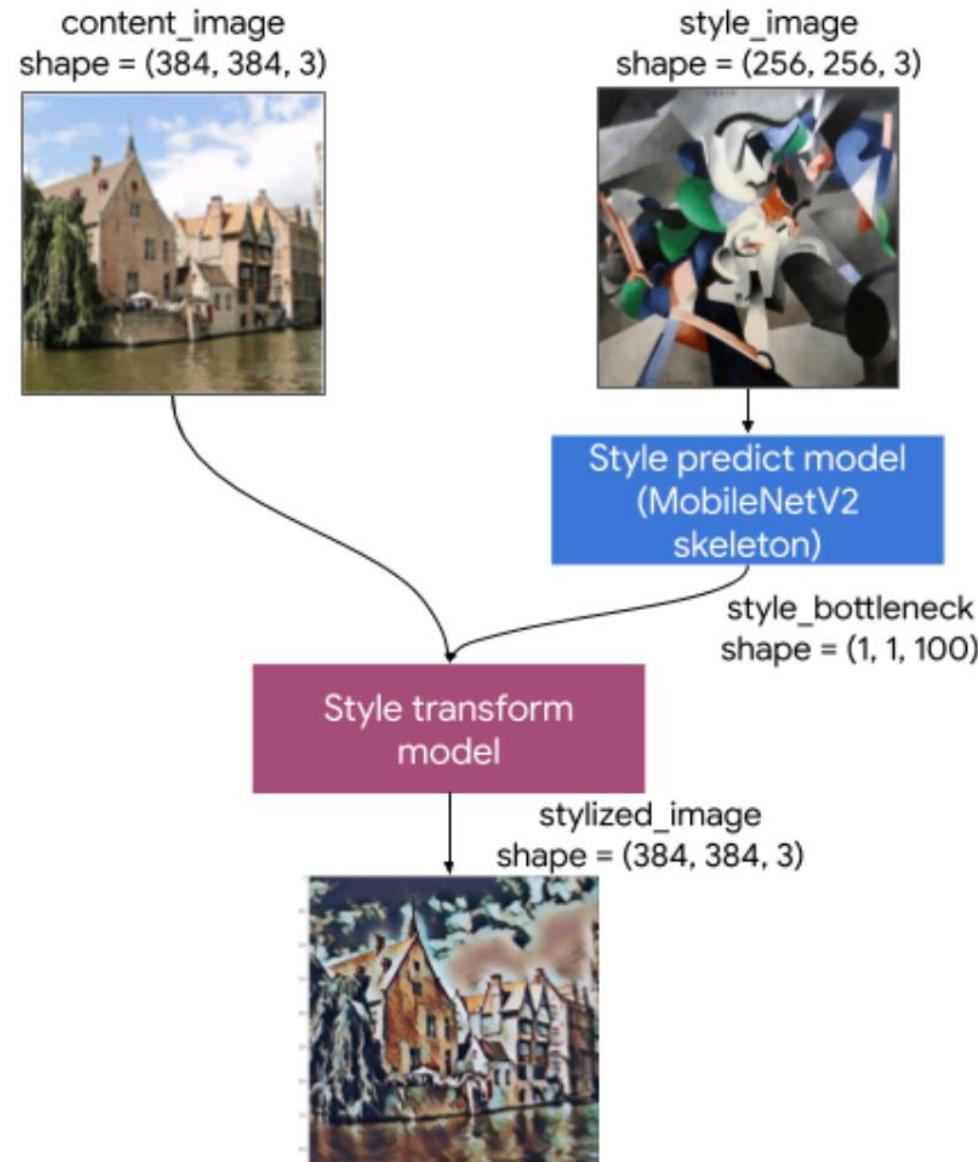


https://www.tensorflow.org/lite/examples/style_transfer/overview

Style Transfer

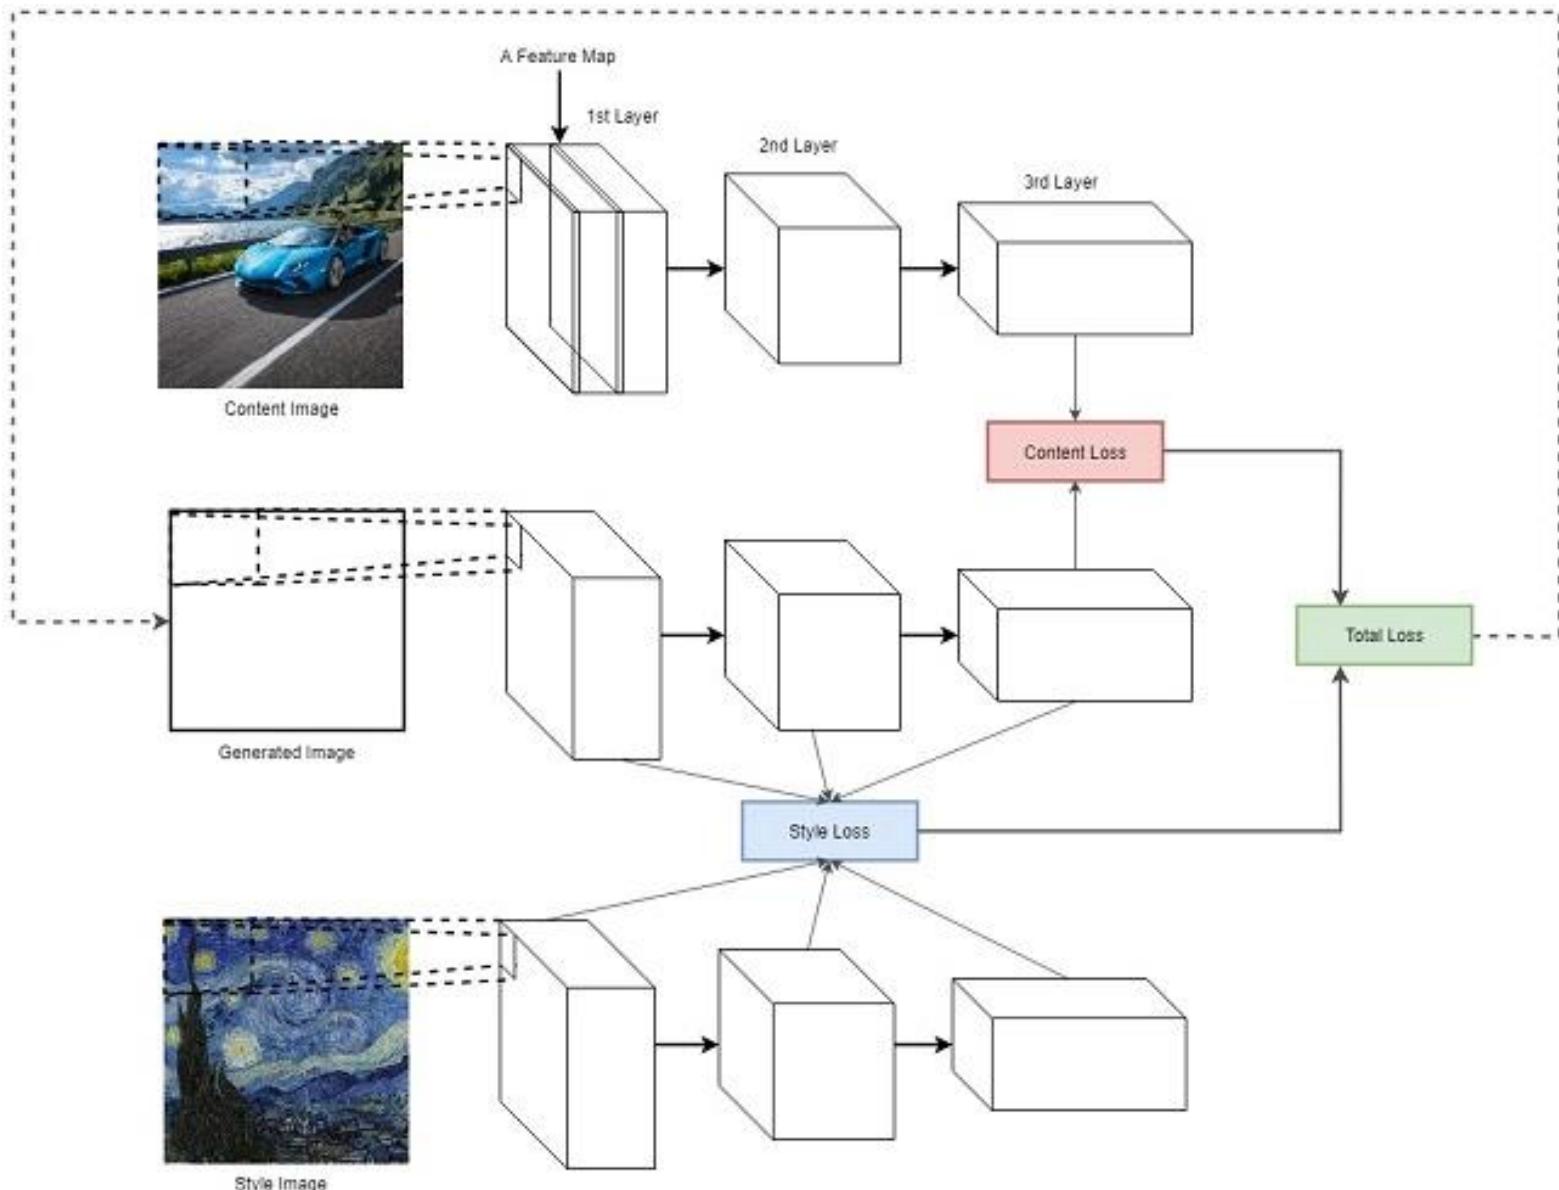
❖ From different viewpoints

https://www.tensorflow.org/lite/examples/style_transfer/overview



Style Transfer

From different
viewpoints

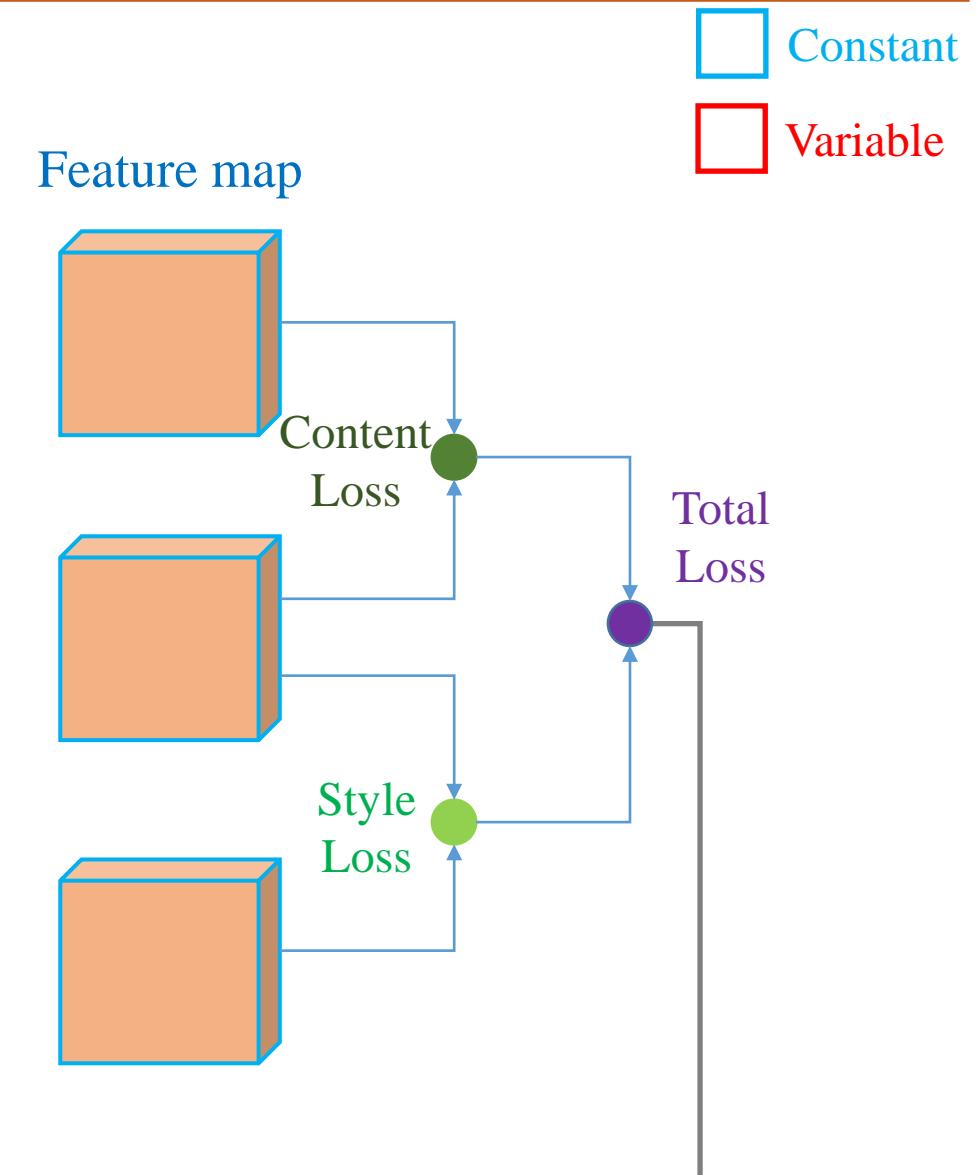
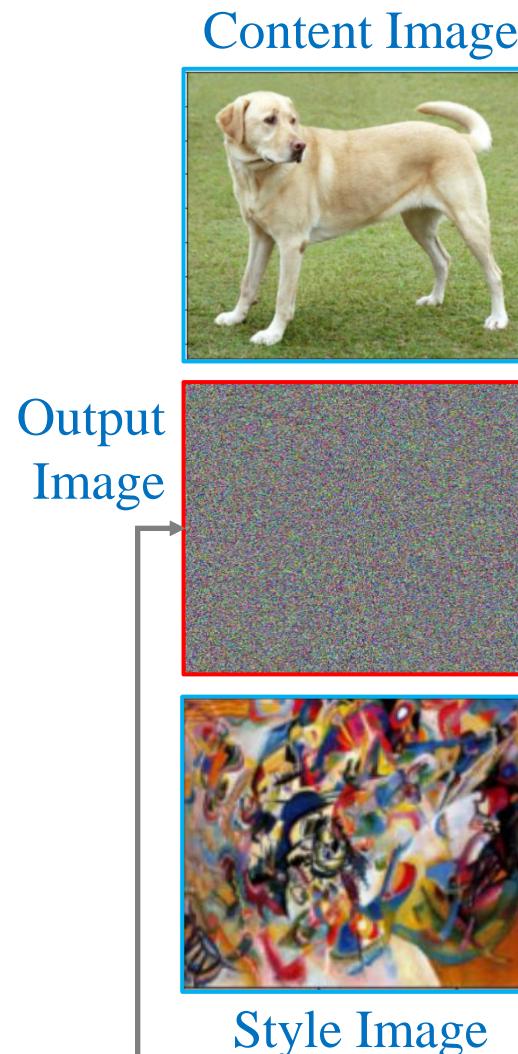


Outline

- Introduction
- Inputs as Variables
- Content Loss
- Style Loss

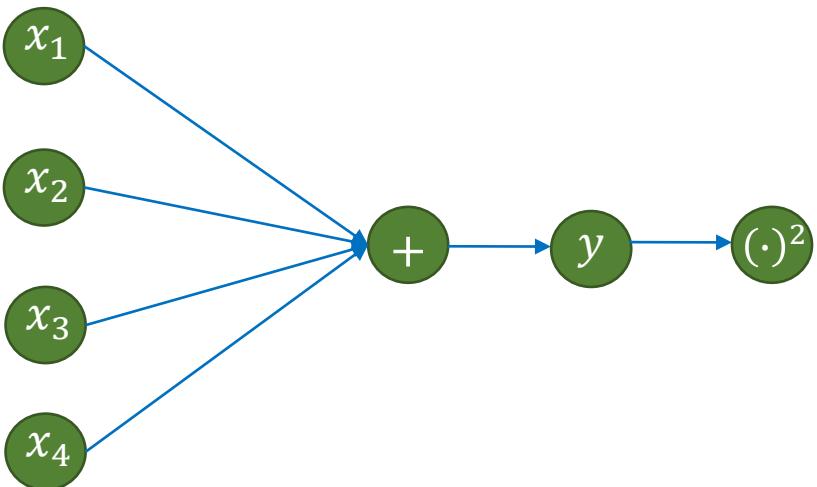
Style Transfer

❖ How to use inputs as variables



Inputs as Variables

- ❖ **Variable:** Represents a tensor whose value can be changed
- ❖ Gradient computation

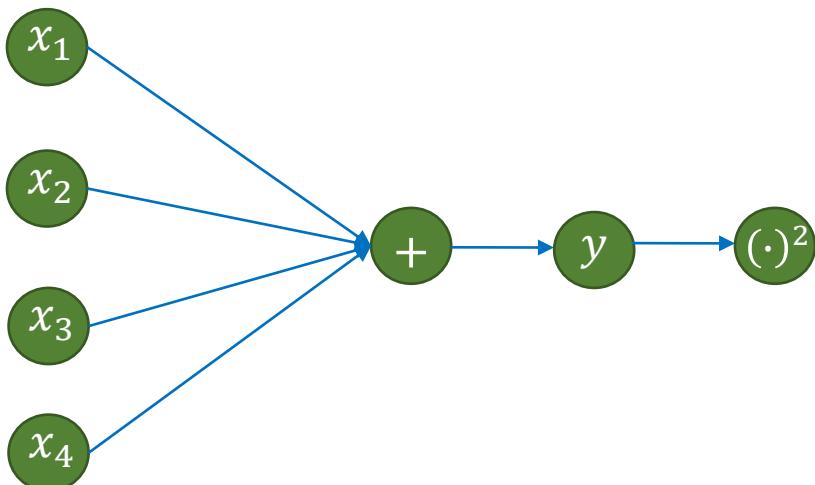


```
tensor([[1., 1.],  
        [1., 1.]], requires_grad=True)  
tensor([[8., 8.],  
        [8., 8.]])
```

```
import torch  
from torch.autograd import Variable  
  
x = torch.Tensor([[1.0, 1.0],  
                  [1.0, 1.0]])  
  
x = Variable(x, requires_grad=True)  
y = x.sum()  
z = y*y  
  
# compute gradient  
z.backward()  
  
# print out  
print(x)  
print(x.grad)
```

Inputs as Variables

- ❖ Variable
- ❖ Gradient computation

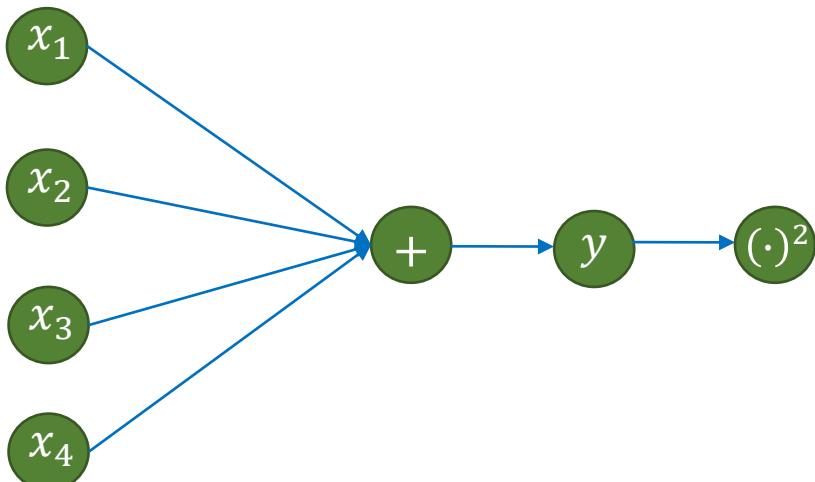


```
tensor([[0.2000, 0.2000],  
        [0.2000, 0.2000]], requires_grad=True)  
tensor([[8., 8.],  
        [8., 8.]])
```

```
import torch  
from torch.autograd import Variable  
  
x = torch.Tensor([[1.0, 1.0],  
                  [1.0, 1.0]])  
  
x = Variable(x, requires_grad=True)  
y = x.sum()  
z = y*y  
  
# compute gradient  
optimizer = optim.SGD([x], lr=0.1)  
optimizer.zero_grad()  
z.backward()  
optimizer.step()  
  
# print out  
print(x)  
print(x.grad)
```

Inputs as Variables

- ❖ Change to Variable
- ❖ Gradient computation



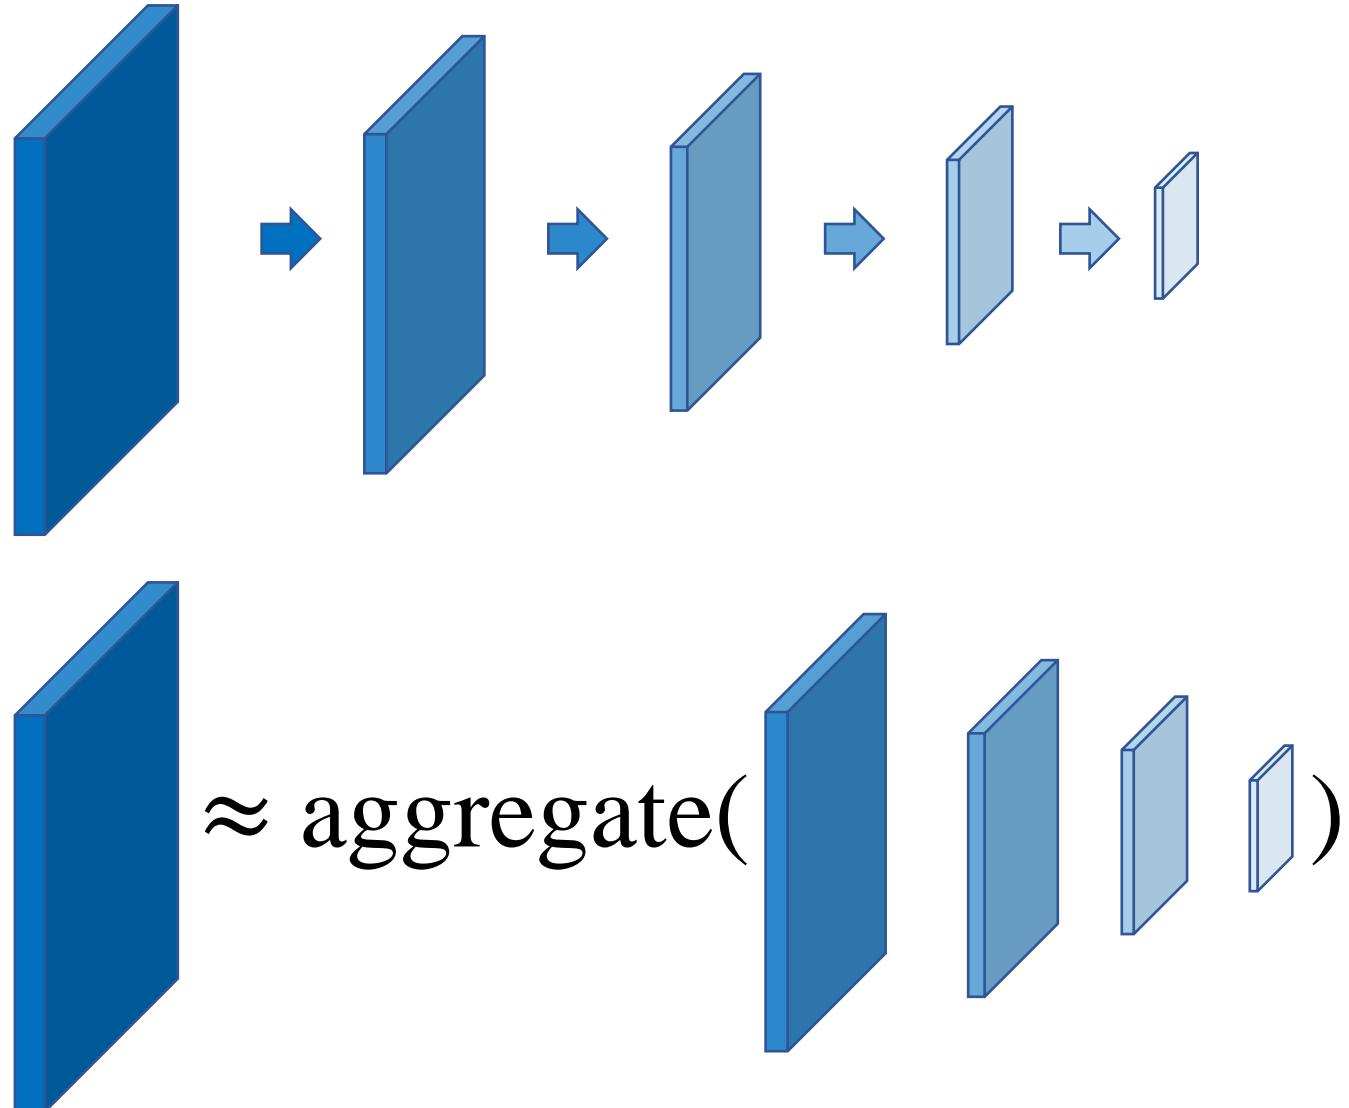
```
tensor([[0.2000, 0.2000],  
        [0.2000, 0.2000]], requires_grad=True)  
tensor([[8., 8.],  
        [8., 8.]])
```

```
import torch  
import torch.optim as optim  
  
x = torch.Tensor([[1.0, 1.0],  
                  [1.0, 1.0]])  
x = x.requires_grad_(True)  
  
y = x.sum()  
z = y*y  
  
# compute gradient  
optimizer = optim.SGD([x], lr=0.1)  
optimizer.zero_grad()  
z.backward()  
optimizer.step()  
  
# print out  
print(x)  
print(x.grad)
```

Outline

- Introduction
- Inputs as Variables
- Content Loss
- Style Loss

Local and Global Similarity

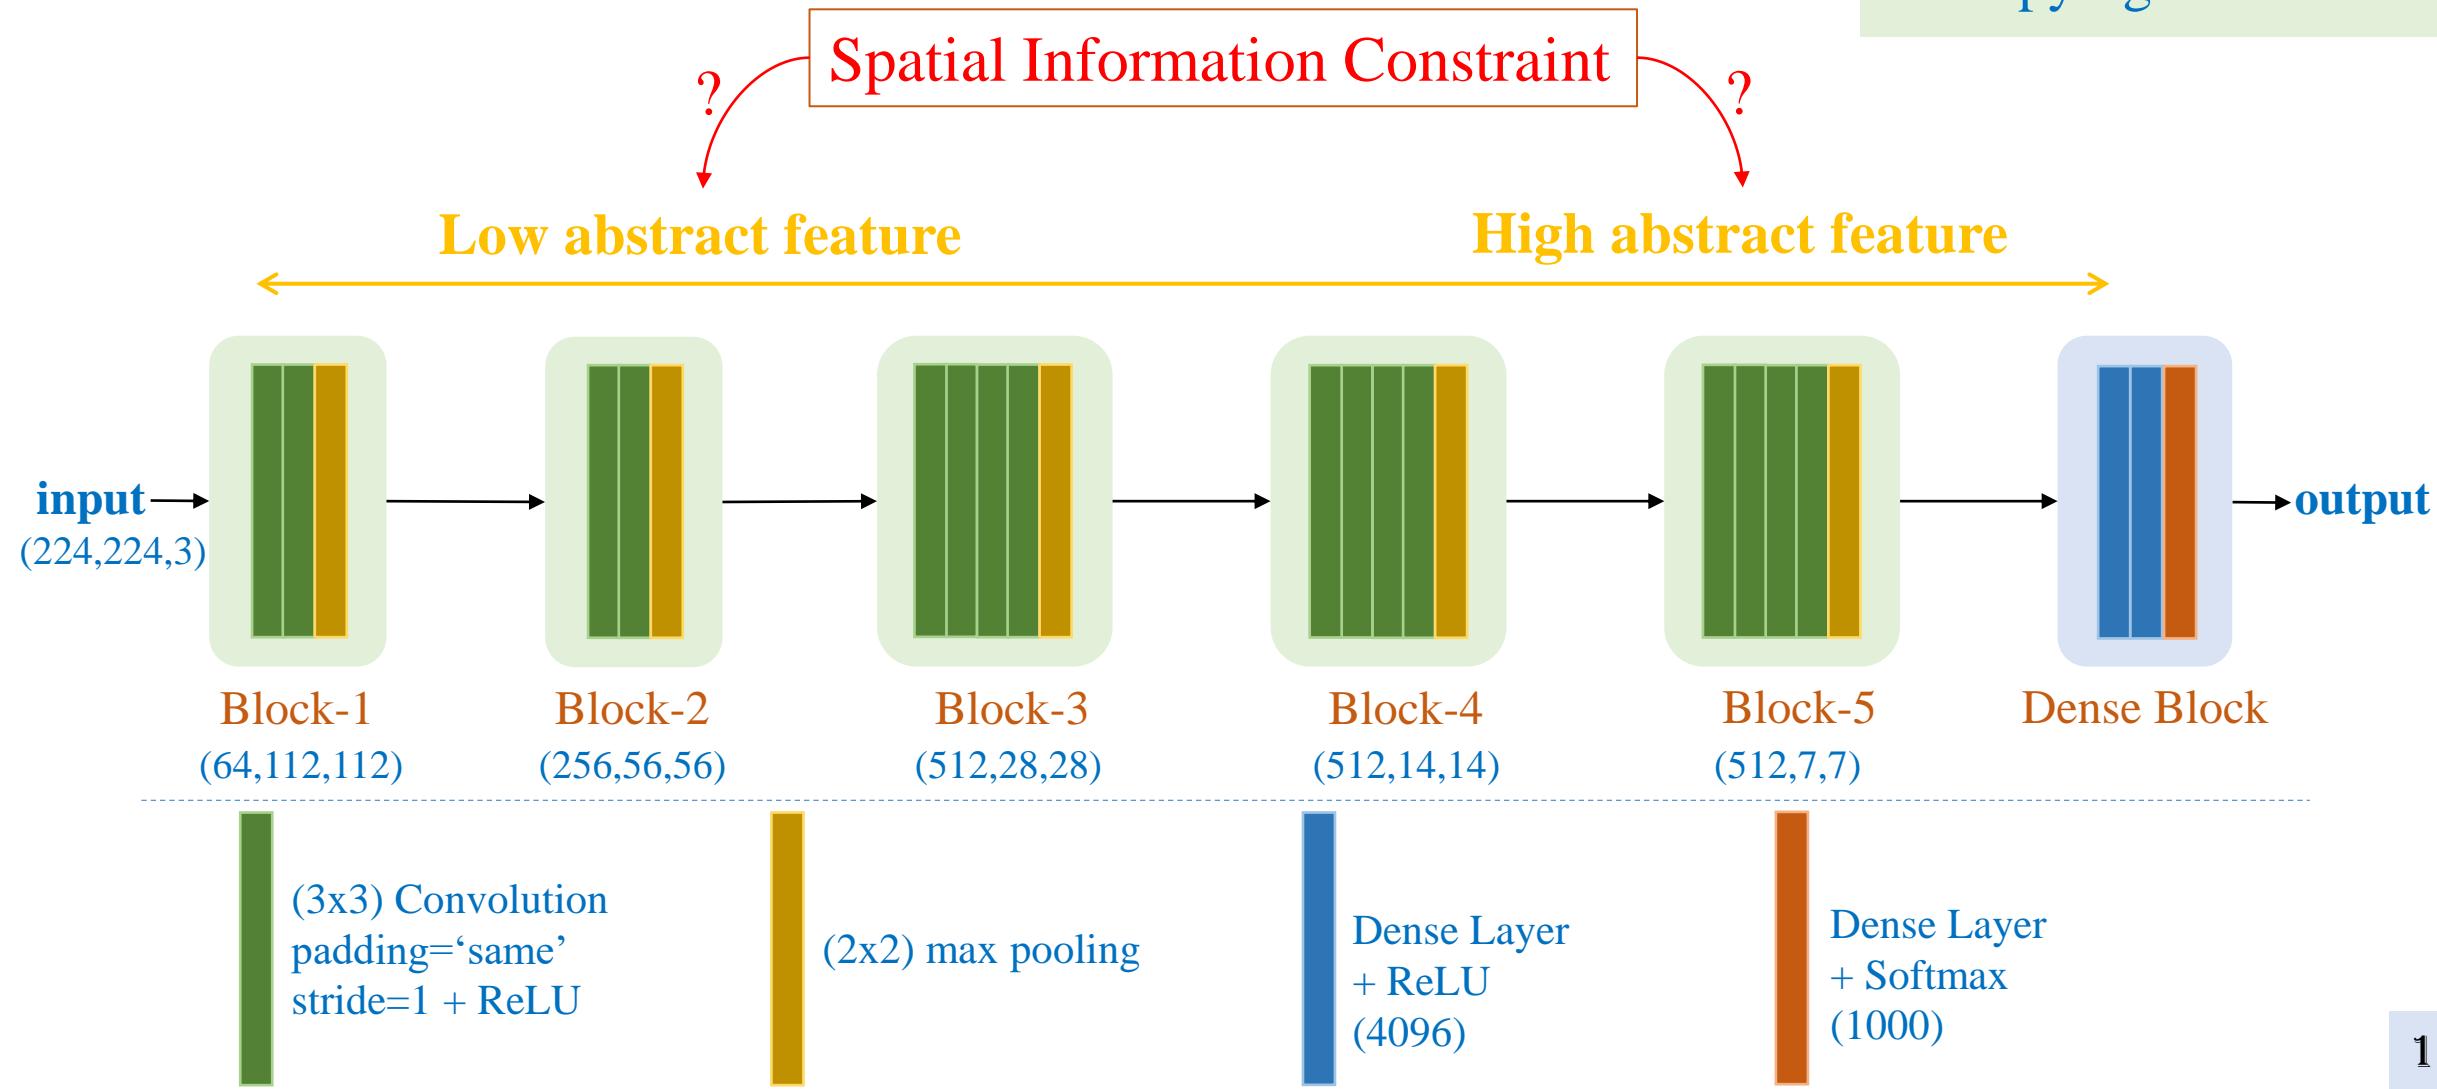


Style Transfer



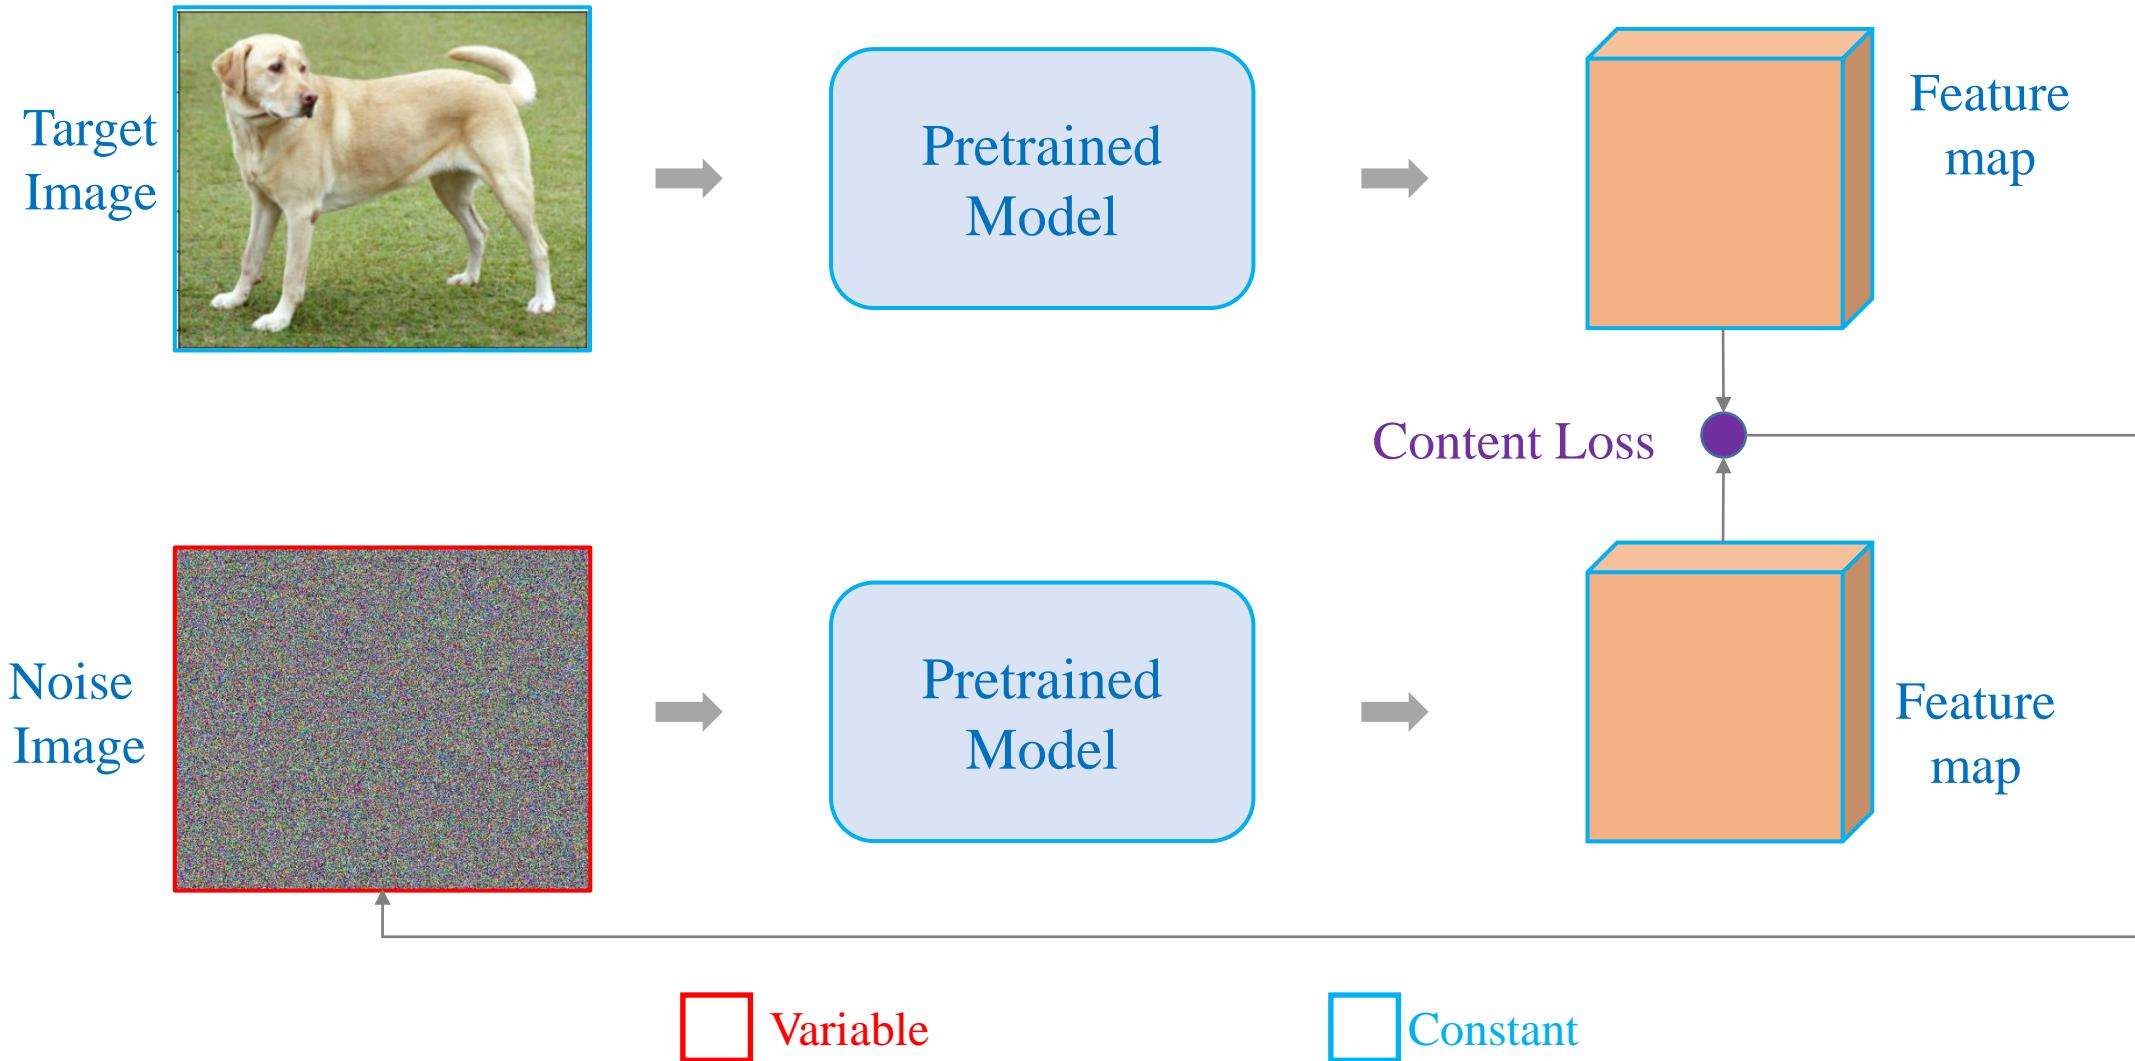
Local and Global Detection Problem

Copying Problem



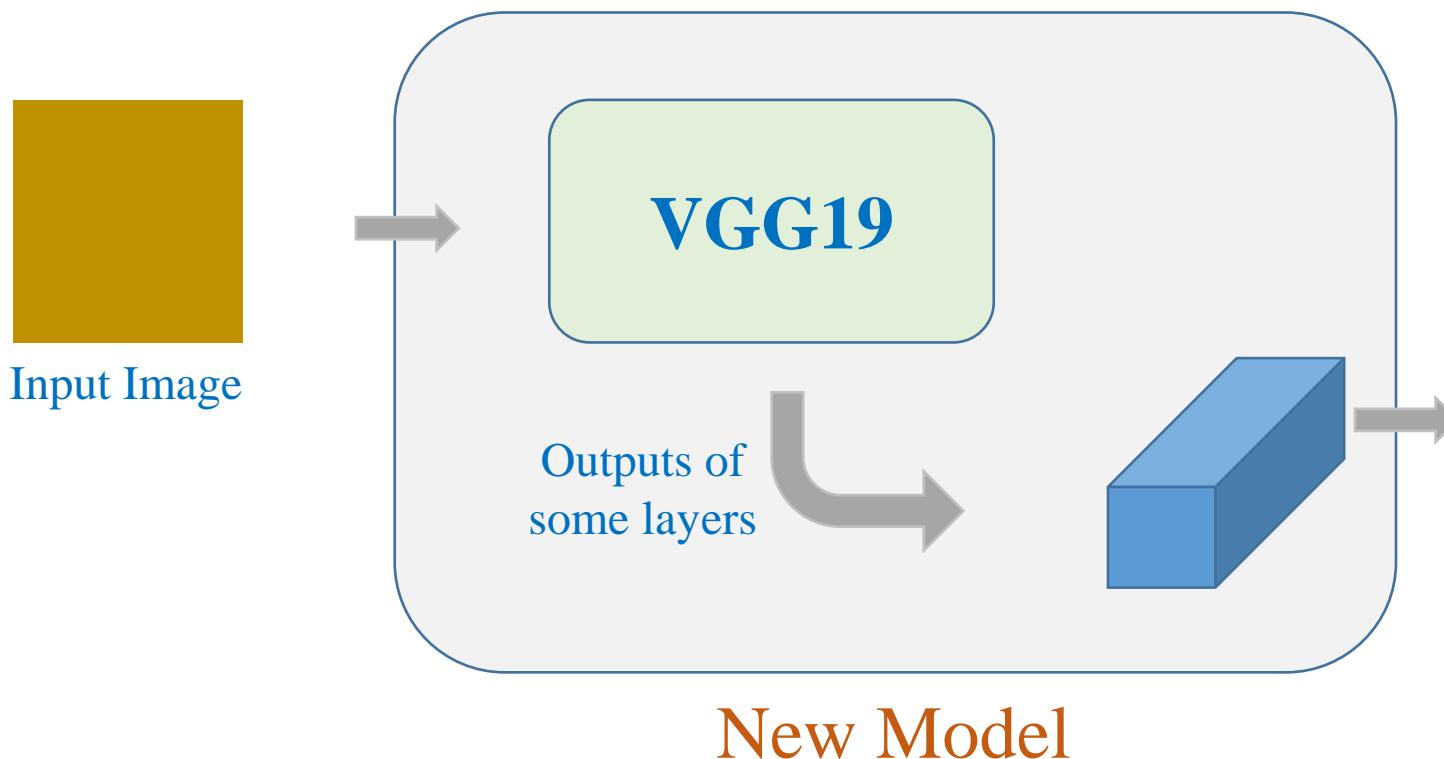
Style Transfer

❖ Content Loss



Style Transfer

- ❖ Content Loss
 - ❖ Create a model from some specific layers



nn.Sequential

```
● ● ●  
1 import torch.nn as nn  
2 model = nn.Sequential(  
3     nn.Conv2d(1, 20, 5),  
4     nn.ReLU(),  
5     nn.Conv2d(20, 64, 5),  
6     nn.ReLU()  
7 )  
8  
9 >> Sequential(  
10    (0): Conv2d(1, 20, kernel_size=(5, 5), stride=(1, 1))  
11    (1): ReLU()  
12    (2): Conv2d(20, 64, kernel_size=(5, 5), stride=(1, 1))  
13    (3): ReLU()  
14 >> )
```

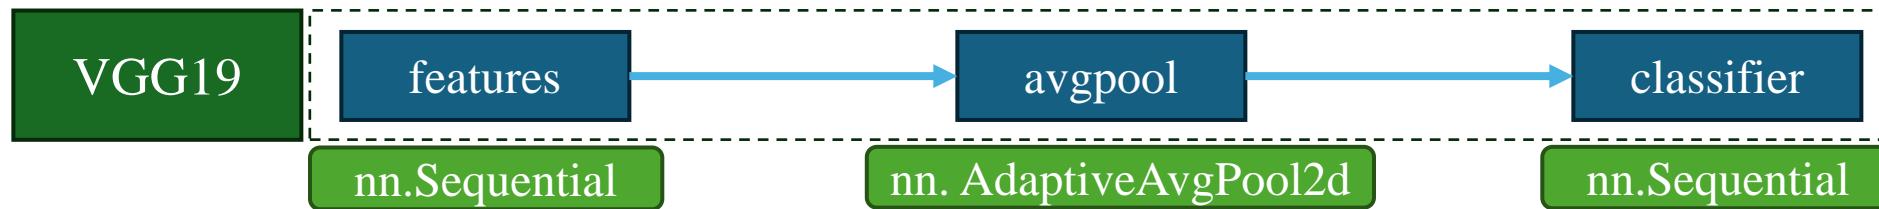


```
1 import torch.nn as nn  
2 from typing import OrderedDict  
3  
4 model = nn.Sequential(OrderedDict([  
5     ('conv1', nn.Conv2d(1,20,5)),  
6     ('relu1', nn.ReLU()),  
7     ('conv2', nn.Conv2d(20,64,5)),  
8     ('relu2', nn.ReLU())  
9 ]))  
10  
11 >> Sequential(  
12    (conv1): Conv2d(1, 20, kernel_size=(5, 5), stride=(1, 1))  
13    (relu1): ReLU()  
14    (conv2): Conv2d(20, 64, kernel_size=(5, 5), stride=(1, 1))  
15    (relu2): ReLU()  
16 >> )
```

```
● ● ●  
1 for name, layer in model._modules.items():  
2     print(f"Layer name : {name} - Module: {layer}")  
3  
4 >> Layer name : 0 - Module: Conv2d(1, 20, kernel_size=(5, 5), stride=(1, 1))  
5 >> Layer name : 1 - Module: ReLU()  
6 >> Layer name : 2 - Module: Conv2d(20, 64, kernel_size=(5, 5), stride=(1, 1))  
7 >> Layer name : 3 - Module: ReLU()
```

```
● ● ●  
1 for name, layer in model._modules.items():  
2     print(f"Layer name : {name} - Module: {layer}")  
3  
4 >> Layer name : conv1 - Module: Conv2d(1, 20, kernel_size=(5, 5), stride=(1, 1))  
5 >> Layer name : relu1 - Module: ReLU()  
6 >> Layer name : conv2 - Module: Conv2d(20, 64, kernel_size=(5, 5), stride=(1, 1))  
7 >> Layer name : relu2 - Module: ReLU()
```

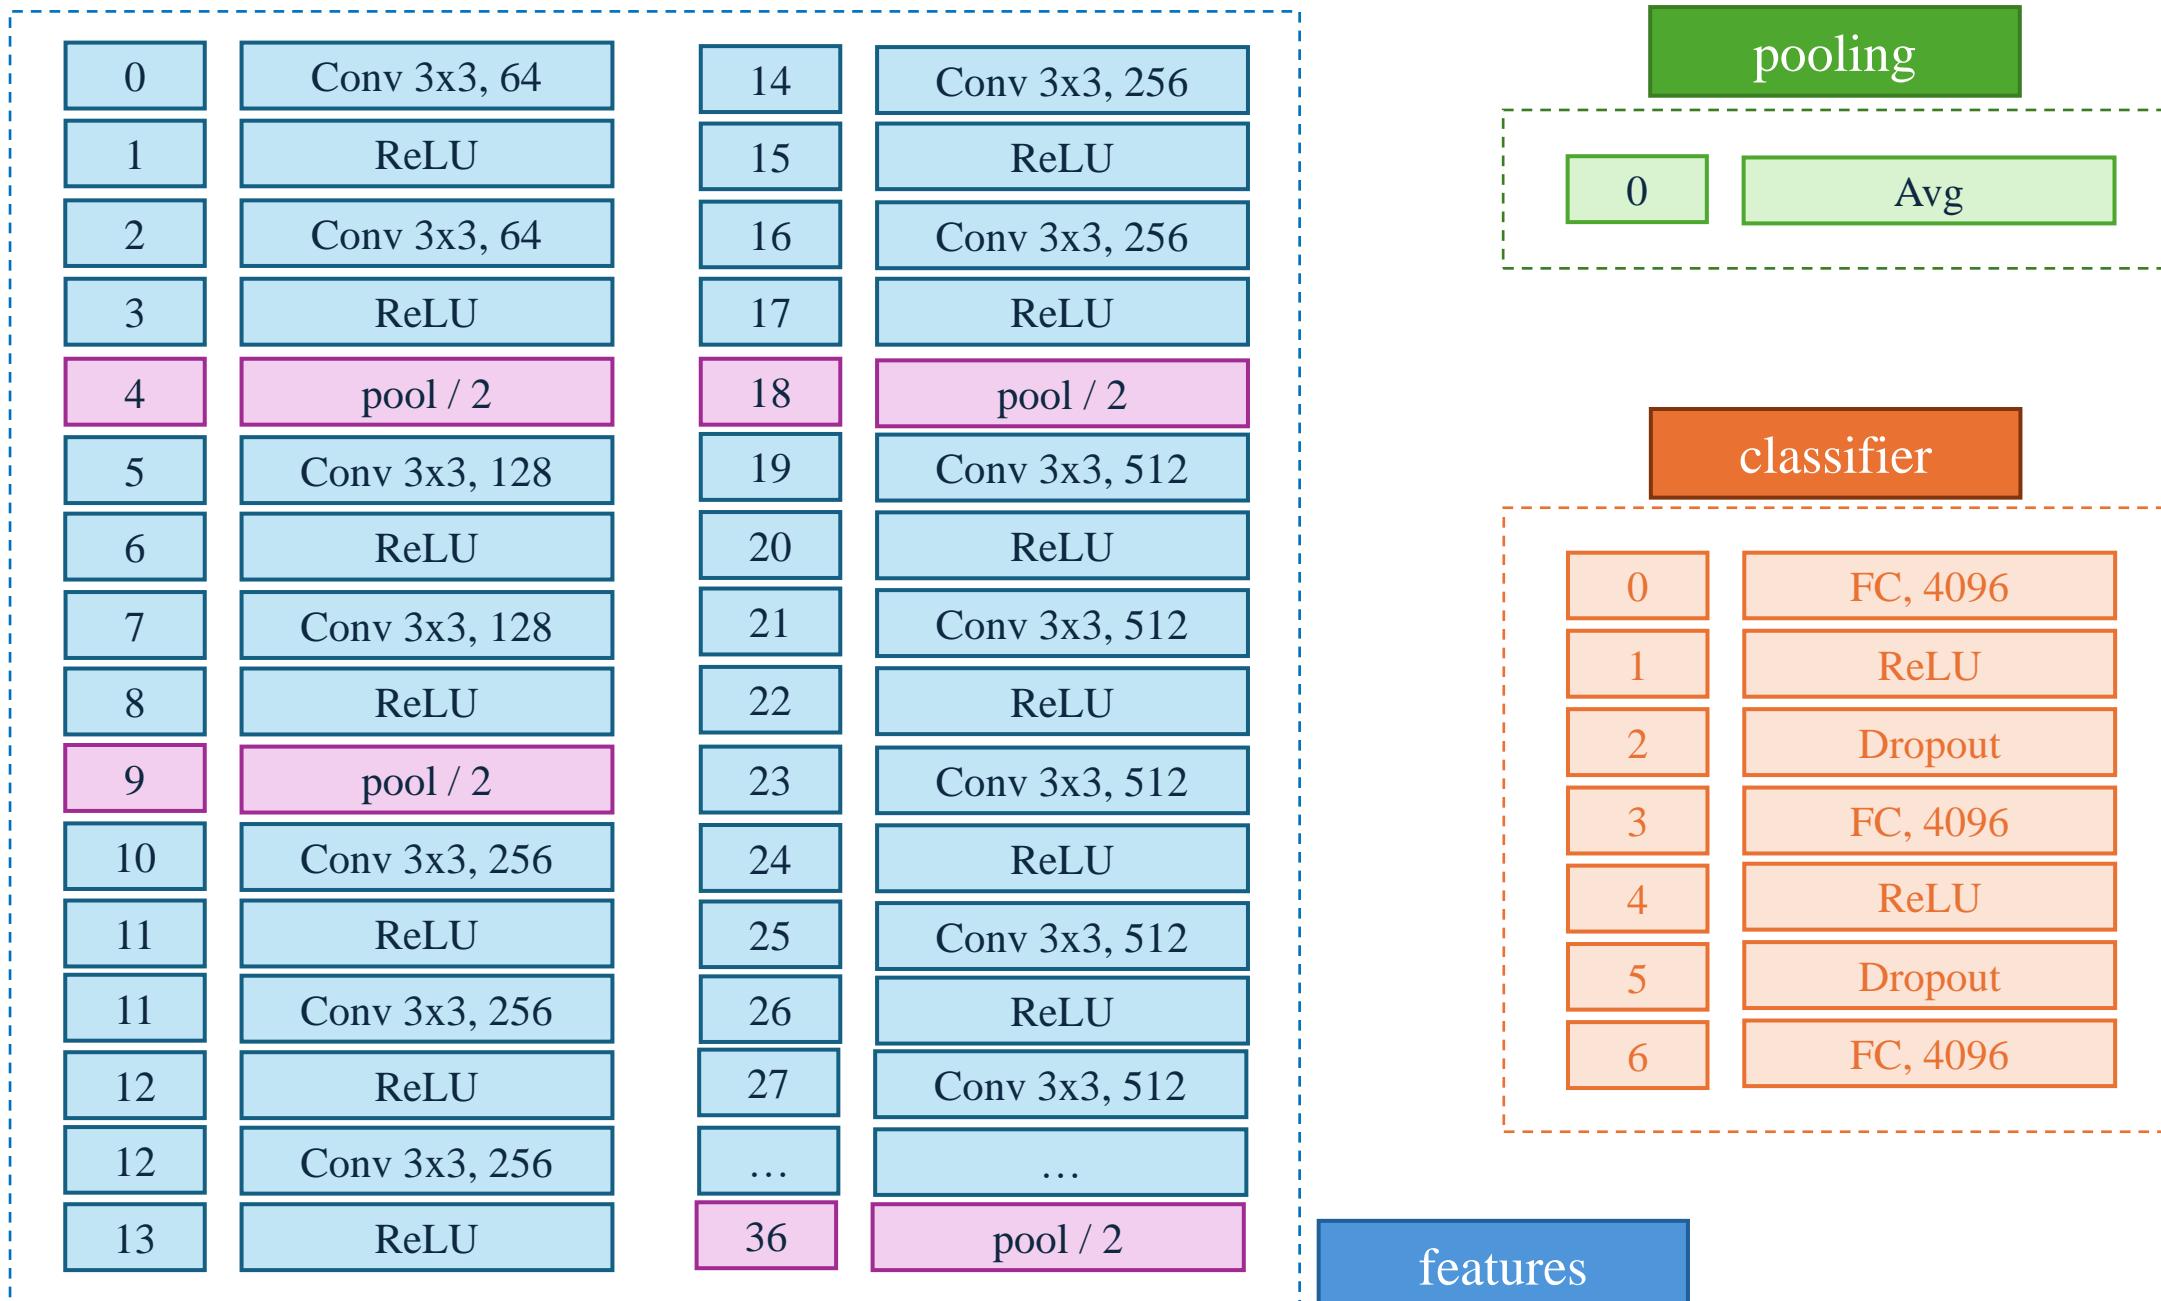
```
from torchvision.models import vgg19
```



```
● ● ●  
1 (features): Sequential(  
2   (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
3   (1): ReLU(inplace=True)  
4   (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
5   (3): ReLU(inplace=True)  
6   (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
7   (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
8   (6): ReLU(inplace=True)  
9   (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
10  (8): ReLU(inplace=True)  
11  (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
12  (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
13  (11): ReLU(inplace=True)  
14  (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
15  (13): ReLU(inplace=True)  
16  (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
17  (15): ReLU(inplace=True)  
18  (16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
19  (17): ReLU(inplace=True)  
20  (18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
21  (19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
22  (20): ReLU(inplace=True)  
23  (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
24  (22): ReLU(inplace=True)  
25  (23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
26  (24): ReLU(inplace=True)  
27  (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
28  (26): ReLU(inplace=True)  
29  (27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
30  (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
31  (29): ReLU(inplace=True)  
32  (30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
33  (31): ReLU(inplace=True)  
34  (32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
35  (33): ReLU(inplace=True)  
36  (34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
37  (35): ReLU(inplace=True)  
38  (36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```
● ● ●  
1 (classifier): Sequential(  
2   (0): Linear(in_features=25088, out_features=4096, bias=True)  
3   (1): ReLU(inplace=True)  
4   (2): Dropout(p=0.5, inplace=False)  
5   (3): Linear(in_features=4096, out_features=4096, bias=True)  
6   (4): ReLU(inplace=True)  
7   (5): Dropout(p=0.5, inplace=False)  
8   (6): Linear(in_features=4096, out_features=1000, bias=True)  
9 )
```

```
● ● ●  
1 import torchvision  
2  
3 vgg19_model = torchvision.models.vgg19()  
4 vgg19_features = torchvision.models.vgg19().features  
5 vgg19_features
```



Style Transfer

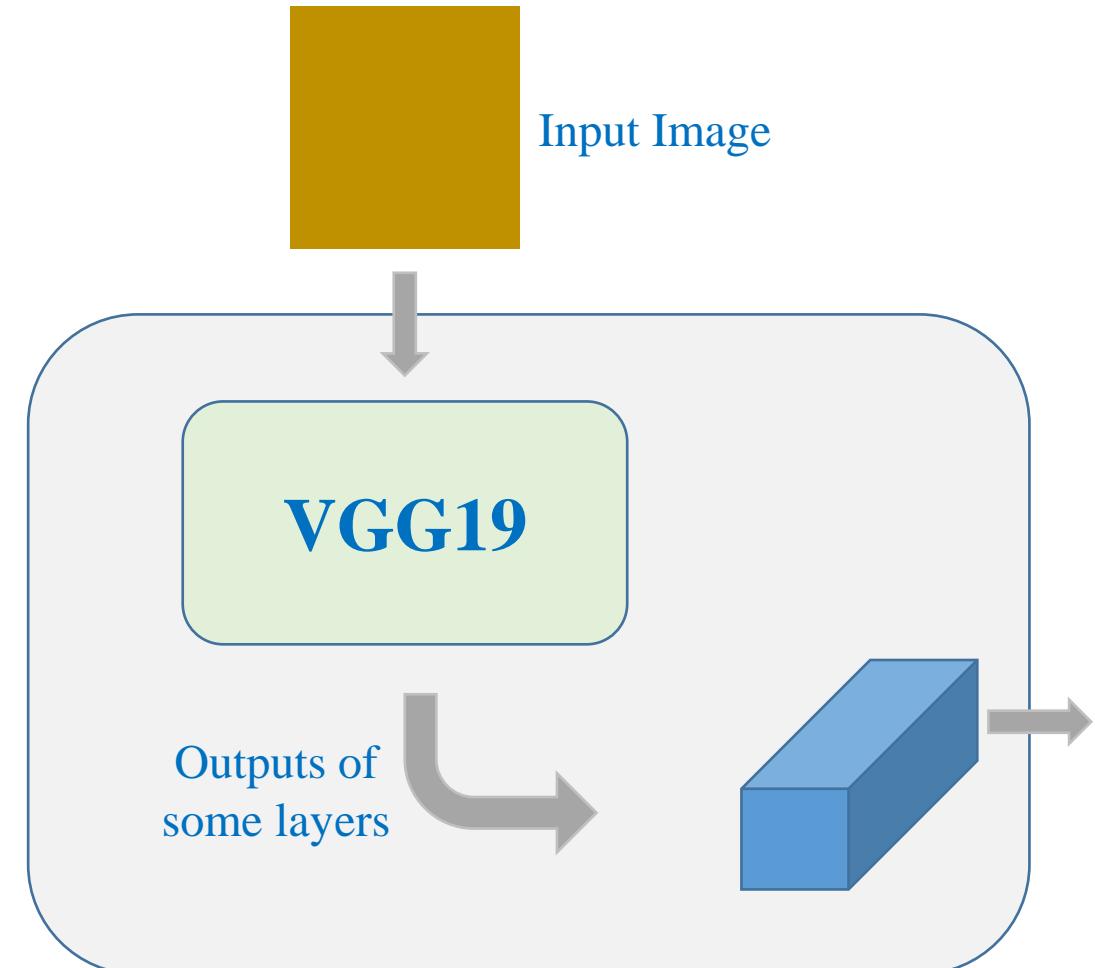
- ❖ Content Loss: Demo
 - ❖ Some utility functions

```
from PIL import Image
import torchvision.transforms as transforms

imsize = 256
img_transforms = transforms.Compose([
    transforms.Resize((imsize, imszie)),
    transforms.ToTensor()
])

def image_loader(image_name):
    image = Image.open(image_name)
    image = img_transforms(image).unsqueeze(0)
    return image.to(device, torch.float)

content_img = image_loader("content_img.jpg")
```

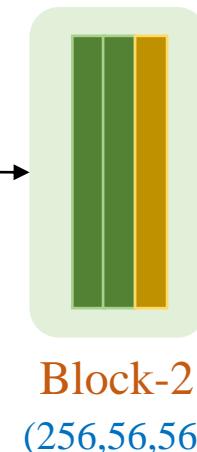
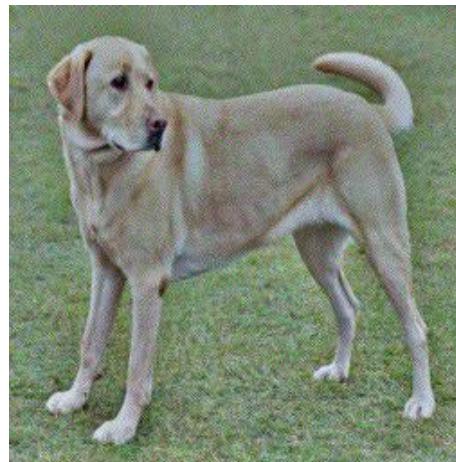


❖ Content Loss

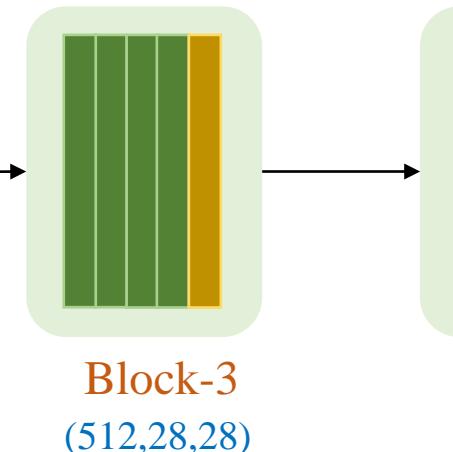
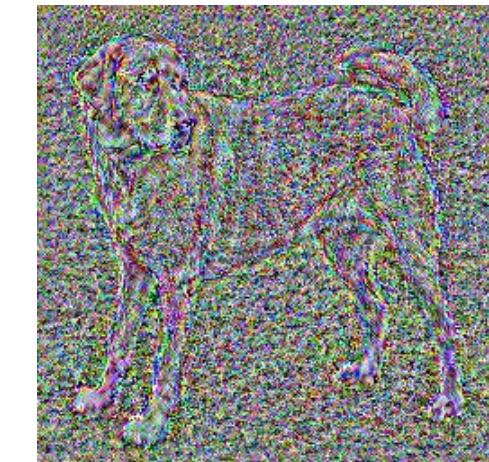
❖ Using different features



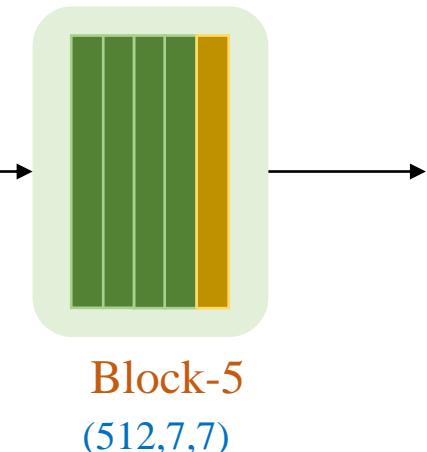
input →
(224,224,3)



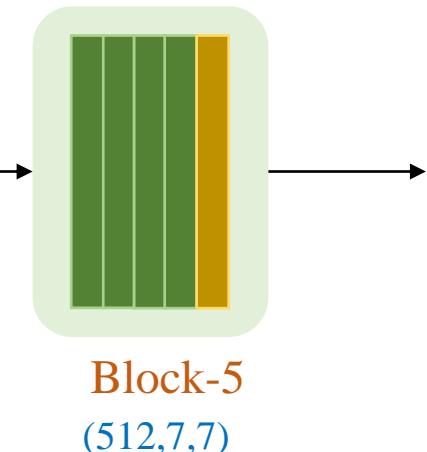
Block-1
(64,112,112)



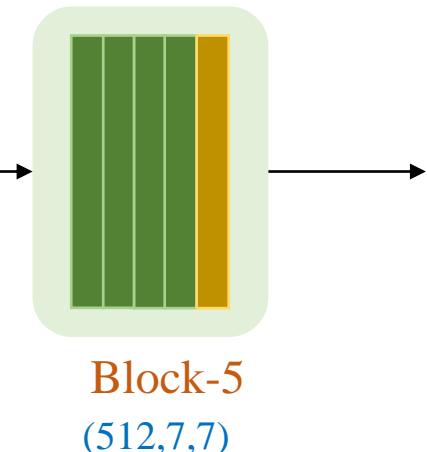
Block-2
(256,56,56)



Block-3
(512,28,28)



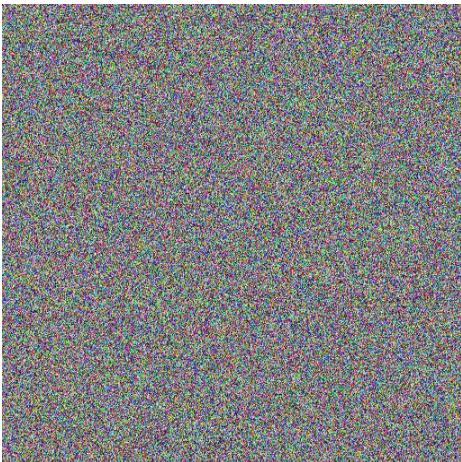
Block-4
(512,14,14)



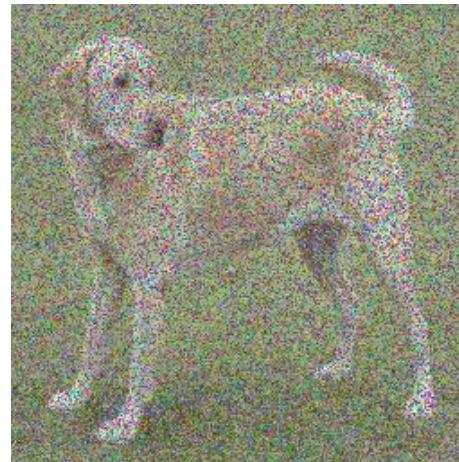
Block-5
(512,7,7)

Style Transfer

❖ Content Loss: Using the Content Image



Initial Image



Epoch 11



Epoch 21



Epoch 31



Epoch 41



Epoch 71



Epoch 91



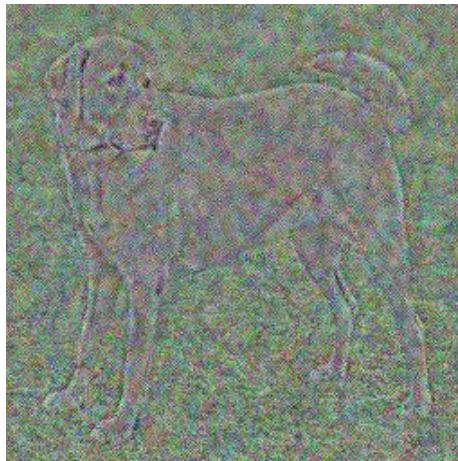
Target Image

Style Transfer

❖ Content Loss: Using Conv1



Initial Image



Epoch 11



Epoch 21



Epoch 31



Epoch 41



Epoch 71

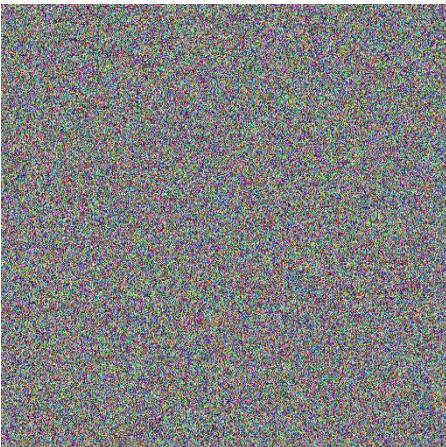


Epoch 91

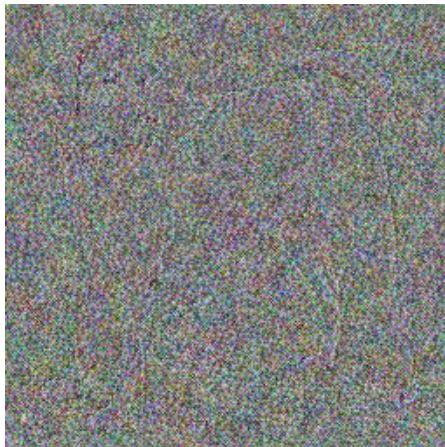


Target Image

Initial Image



Epoch 11



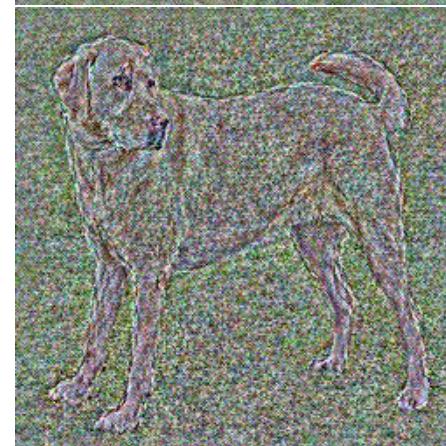
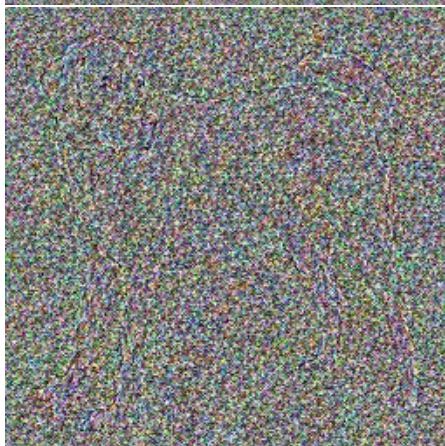
Epoch 41



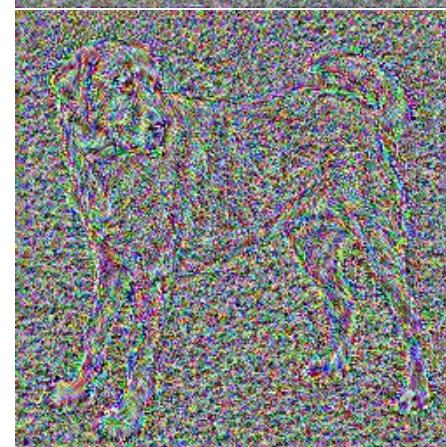
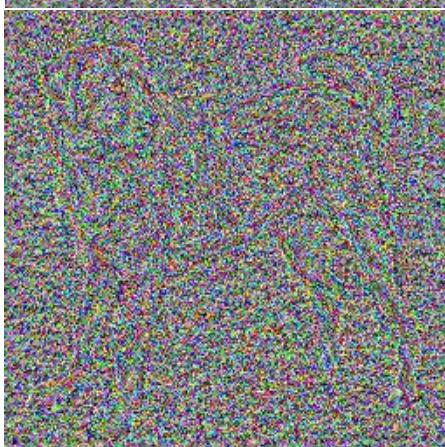
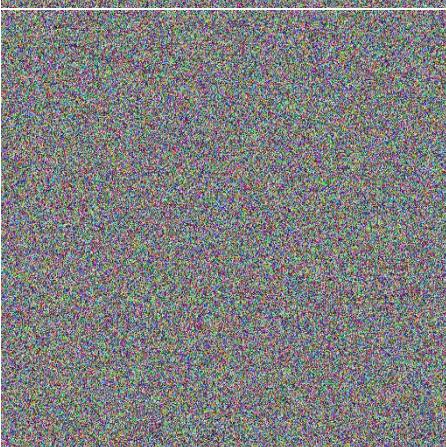
Epoch 91



Using Conv2



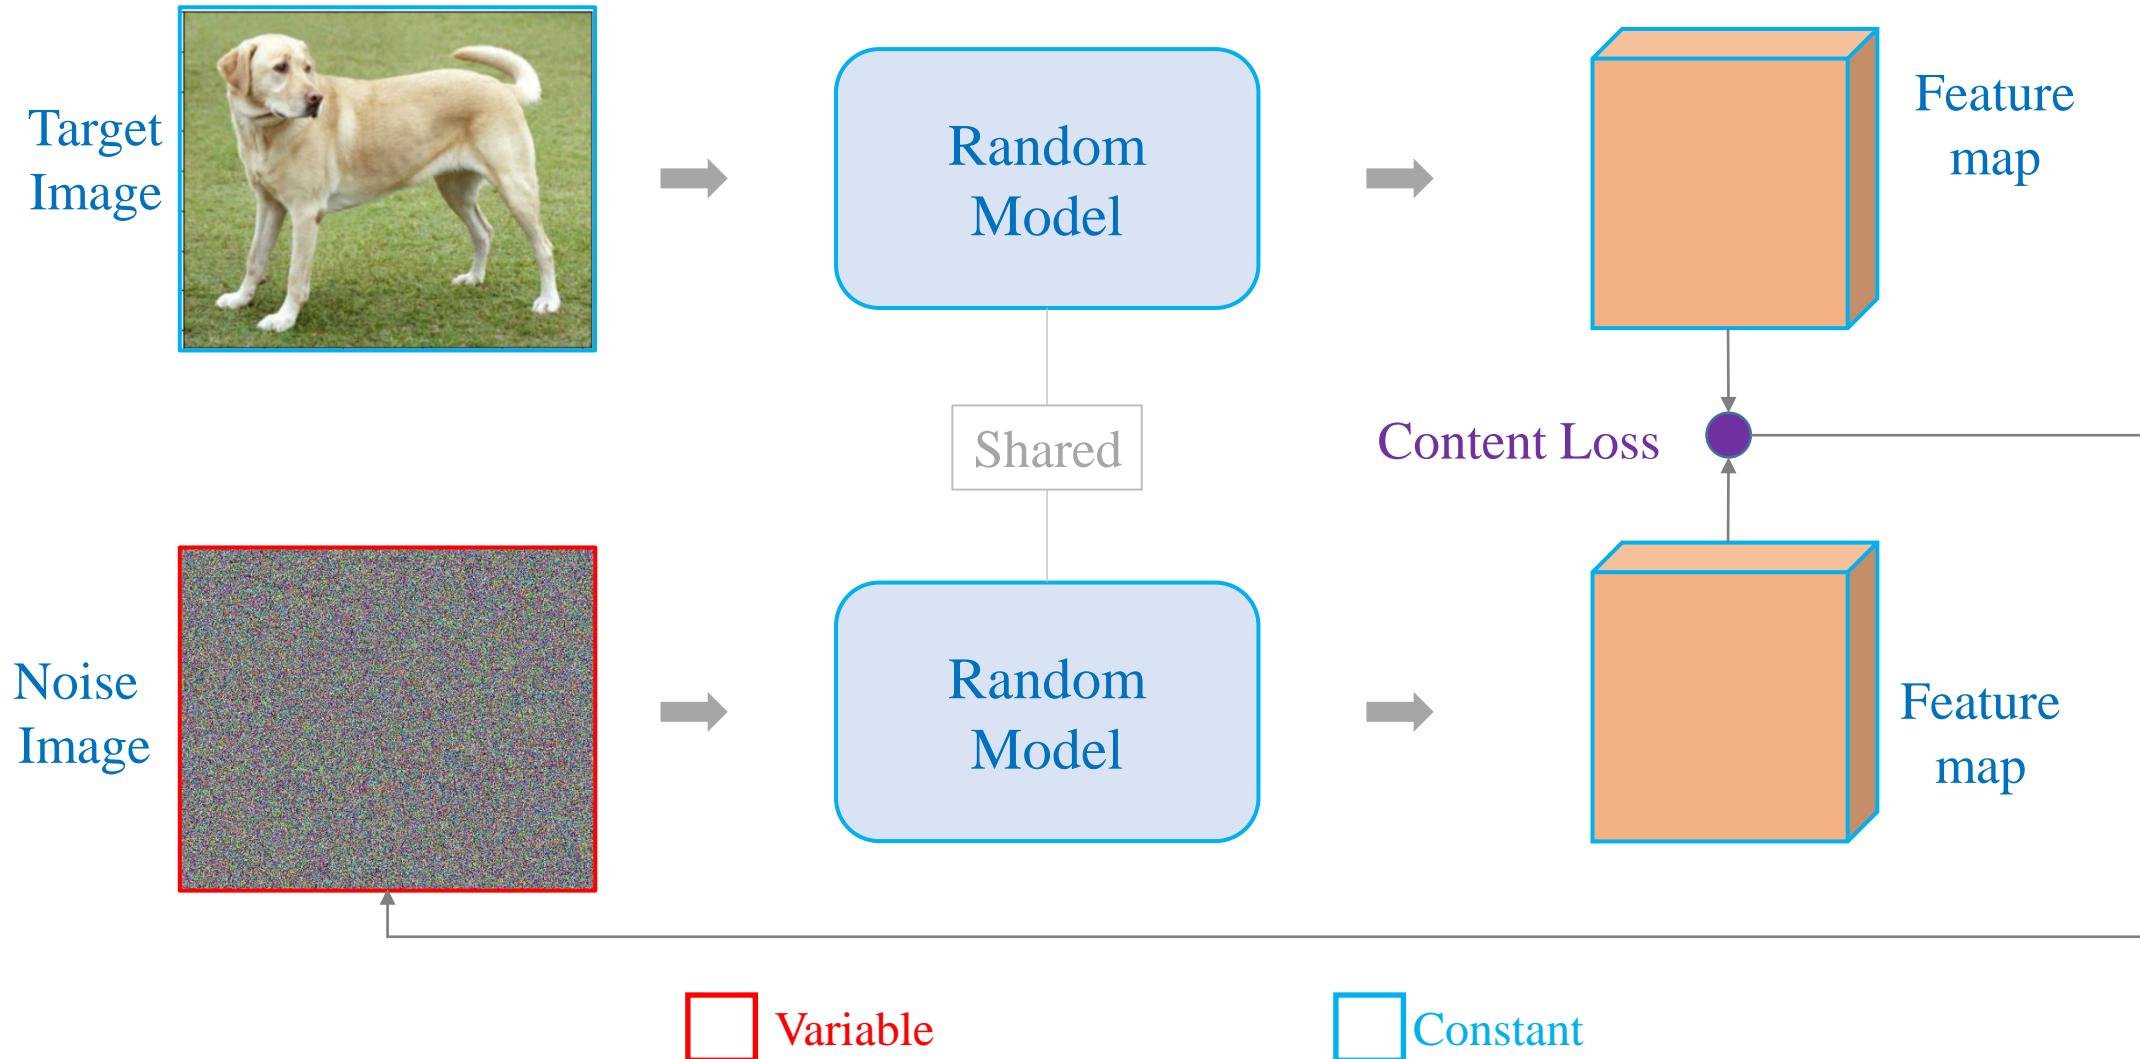
Using Conv3



Using Conv4

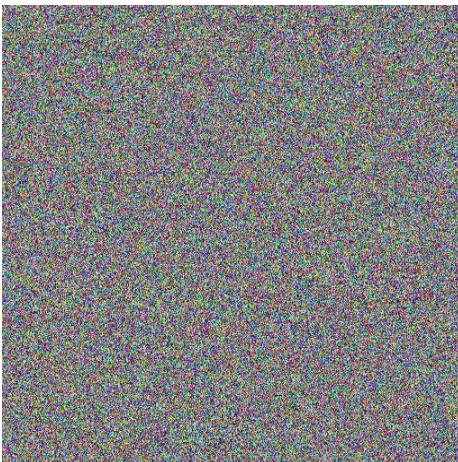
Style Transfer

❖ Content Loss: Using RandomNet

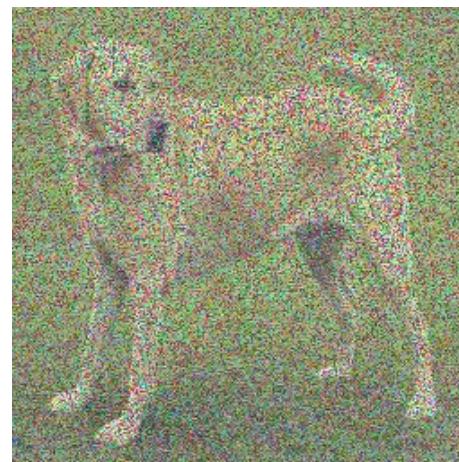


Style Transfer

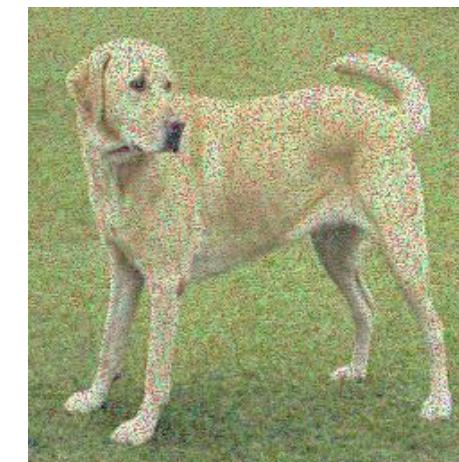
❖ Content Loss: Using RandomNet and the conv1 features (Why?)



Initial Image



Epoch 11



Epoch 21



Epoch 31



Epoch 41



Epoch 71



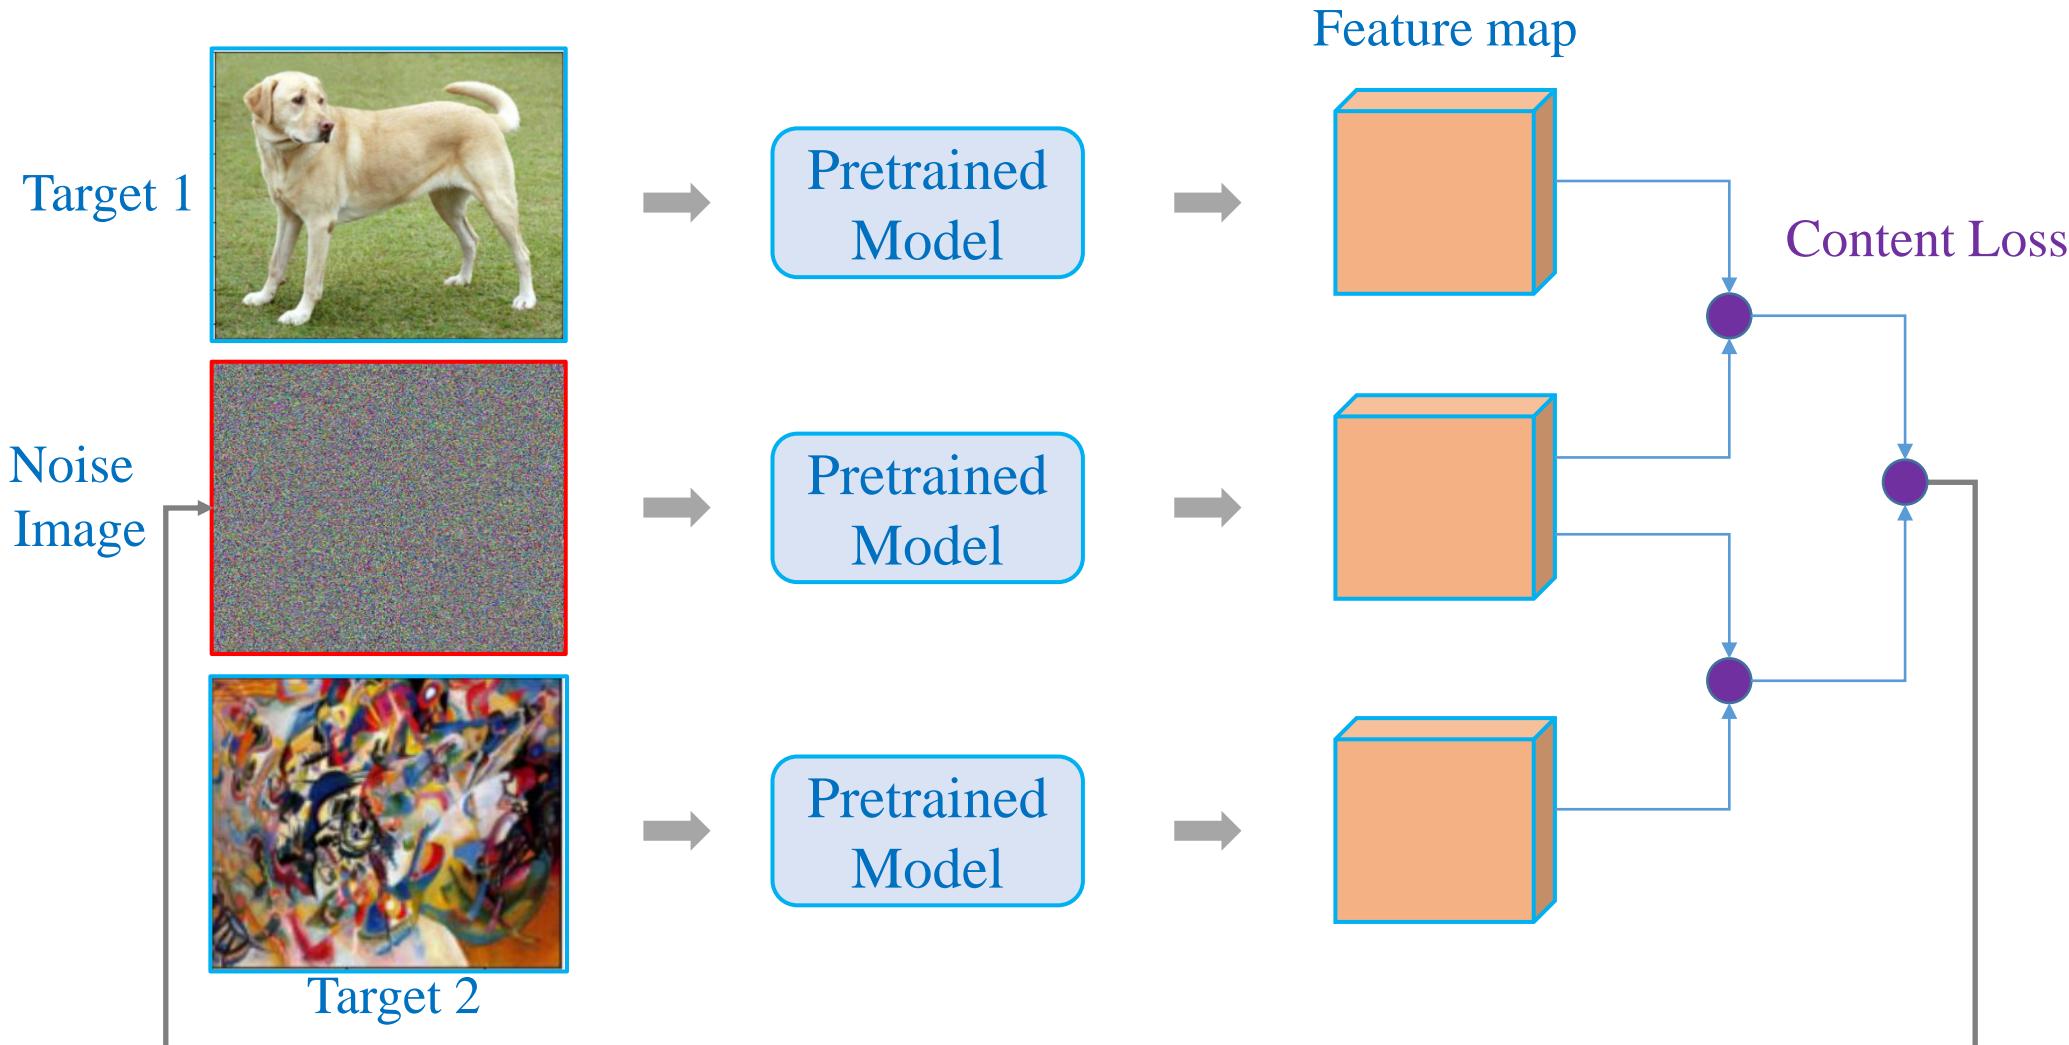
Epoch 91



Target Image

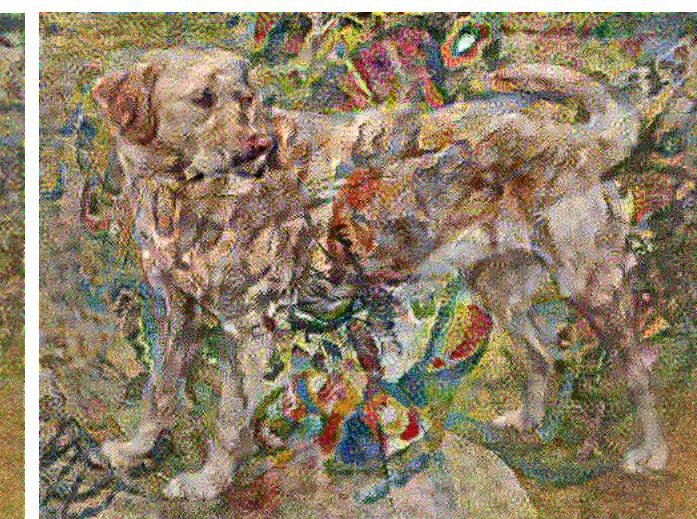
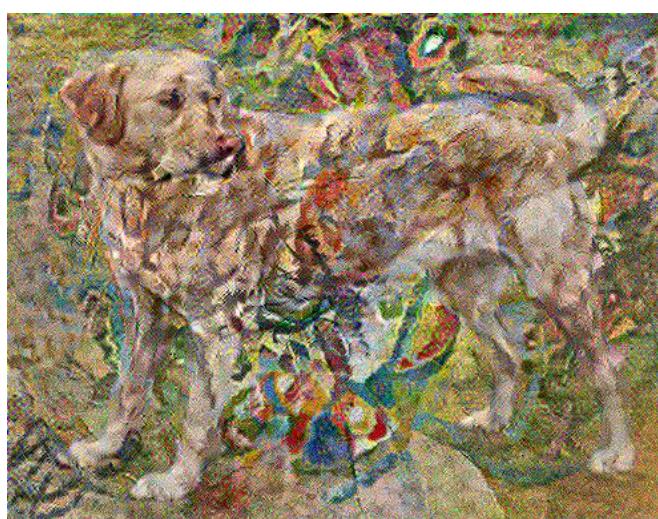
Style Transfer

❖ Content Loss: Using Two Target



Style Transfer

❖ Content Loss: Using Two Target



Outline

- Introduction
- Inputs as Variables
- Content Loss
- Style Loss

Vector Operations

Hadamard product

$$\vec{v} = \begin{bmatrix} v_1 \\ \dots \\ v_n \end{bmatrix} \quad \vec{u} = \begin{bmatrix} u_1 \\ \dots \\ u_n \end{bmatrix}$$

$$\vec{v} \odot \vec{u} = \begin{bmatrix} v_1 \\ \dots \\ v_n \end{bmatrix} \odot \begin{bmatrix} u_1 \\ \dots \\ u_n \end{bmatrix} = \begin{bmatrix} v_1 \times u_1 \\ \dots \\ v_n \times u_n \end{bmatrix}$$

```
1 def Hadamard_product(vector1, vector2):  
2     """  
3         Compute Hadamard product between two vectors  
4         Output is a vector  
5     """  
6     return [v1*v2 for v1, v2 in zip(vector1, vector2)]  
7  
8 # test case  
9 vector1 = [1, 2]  
10 vector2 = [3, 4]  
11  
12 output = Hadamard_product(vector1, vector2)  
13 print(output)
```

[3, 8]

Vector Operations

Dot product

$$\vec{v} = \begin{bmatrix} v_1 \\ \dots \\ v_n \end{bmatrix} \quad \vec{u} = \begin{bmatrix} u_1 \\ \dots \\ u_n \end{bmatrix}$$

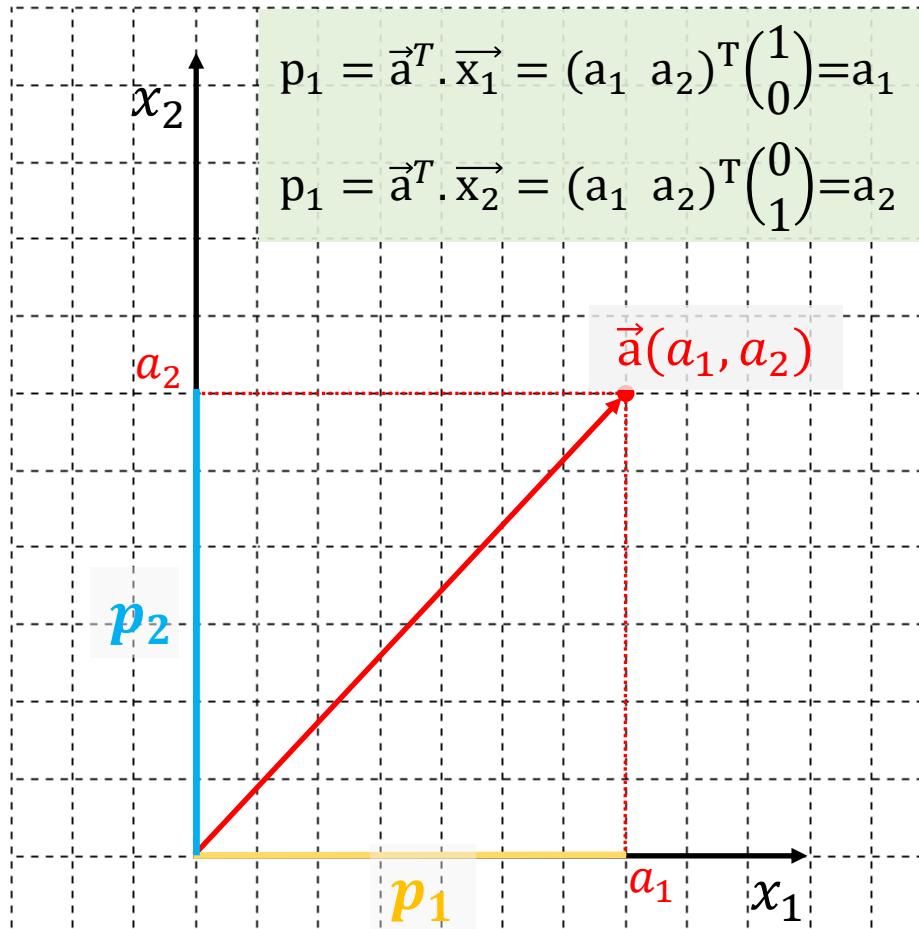
$$\vec{v} \cdot \vec{u} = v_1 \times u_1 + \dots + v_n \times u_n$$

```
1 def dot_product(vector1, vector2):
2     """
3         Compute dot product between two vectors
4         Output is a floating-point number
5     """
6
7     return sum([v1*v2 for v1, v2 in zip(vector1, vector2)])
8
9 # test case
10 vector1 = [1, 2, 3]
11 vector2 = [2, 3, 4]
12
13 output = dot_product(vector1, vector2)
14 print(output)
```

20

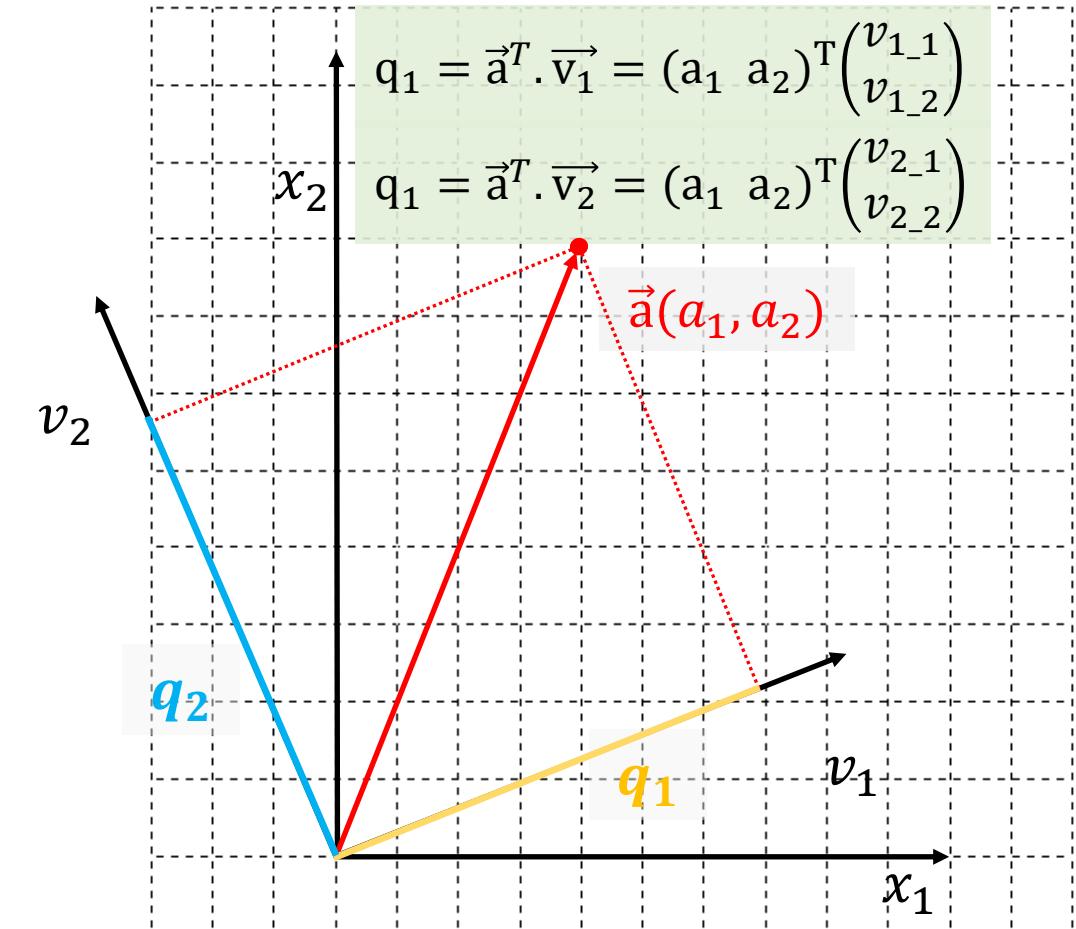
Dot Product

$$\vec{x}_1 = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \quad \vec{x}_2 = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$



Tìm độ dài hình chiếu của \vec{a} lên \vec{x}_1 và \vec{x}_2

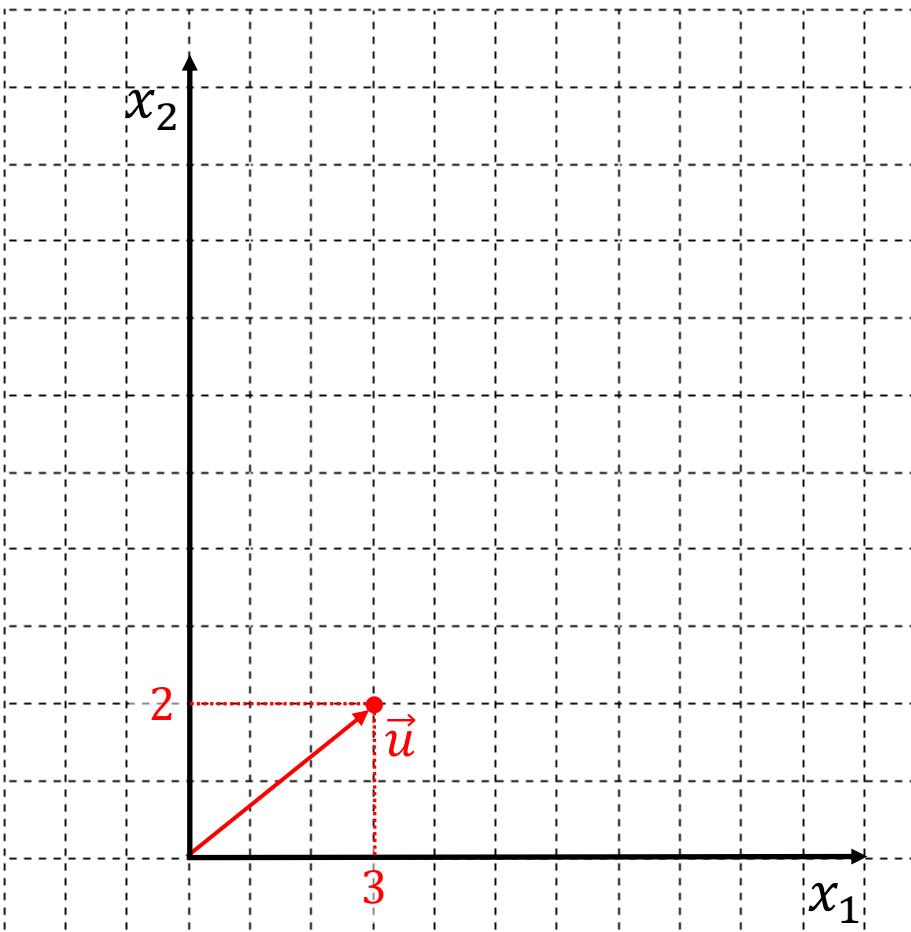
$$\vec{v}_1 = \begin{pmatrix} v_{1_1} \\ v_{1_2} \end{pmatrix} \quad \vec{v}_2 = \begin{pmatrix} v_{2_1} \\ v_{2_2} \end{pmatrix}$$



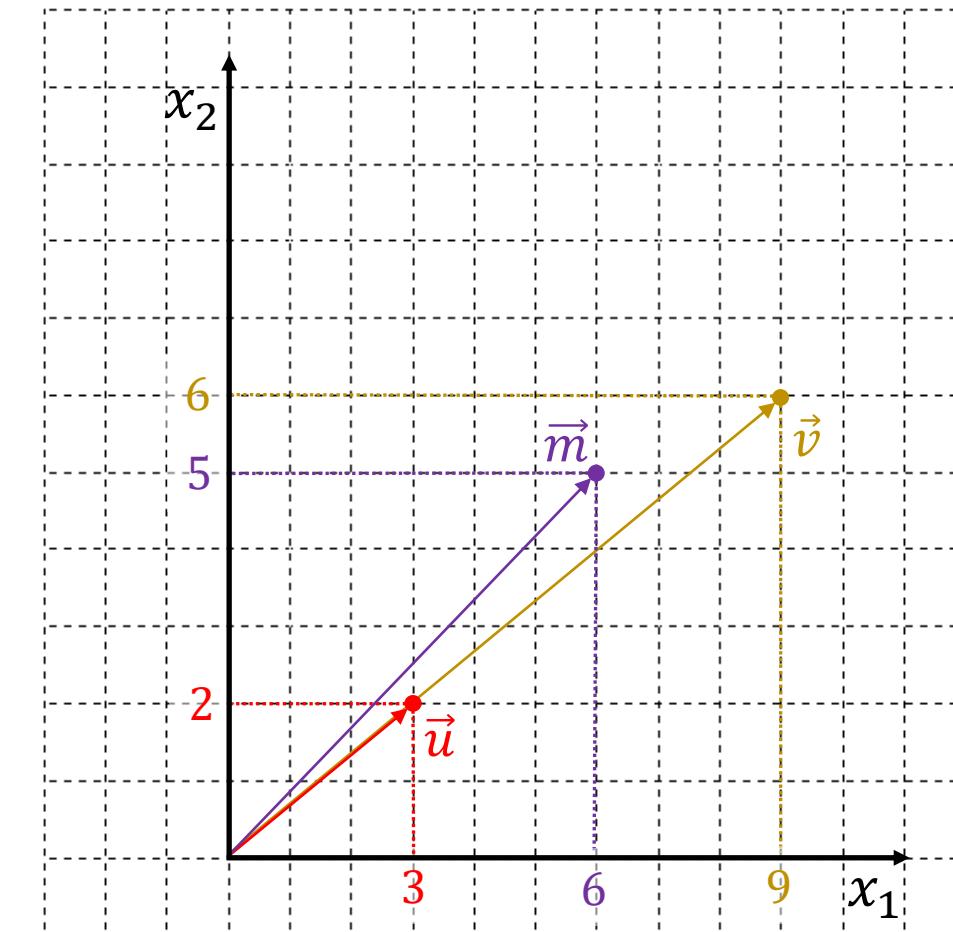
Tìm độ dài hình chiếu của \vec{a} lên \vec{v}_1 và \vec{v}_2

Dot Product

$$\vec{u} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \begin{pmatrix} 3 \\ 2 \end{pmatrix}$$

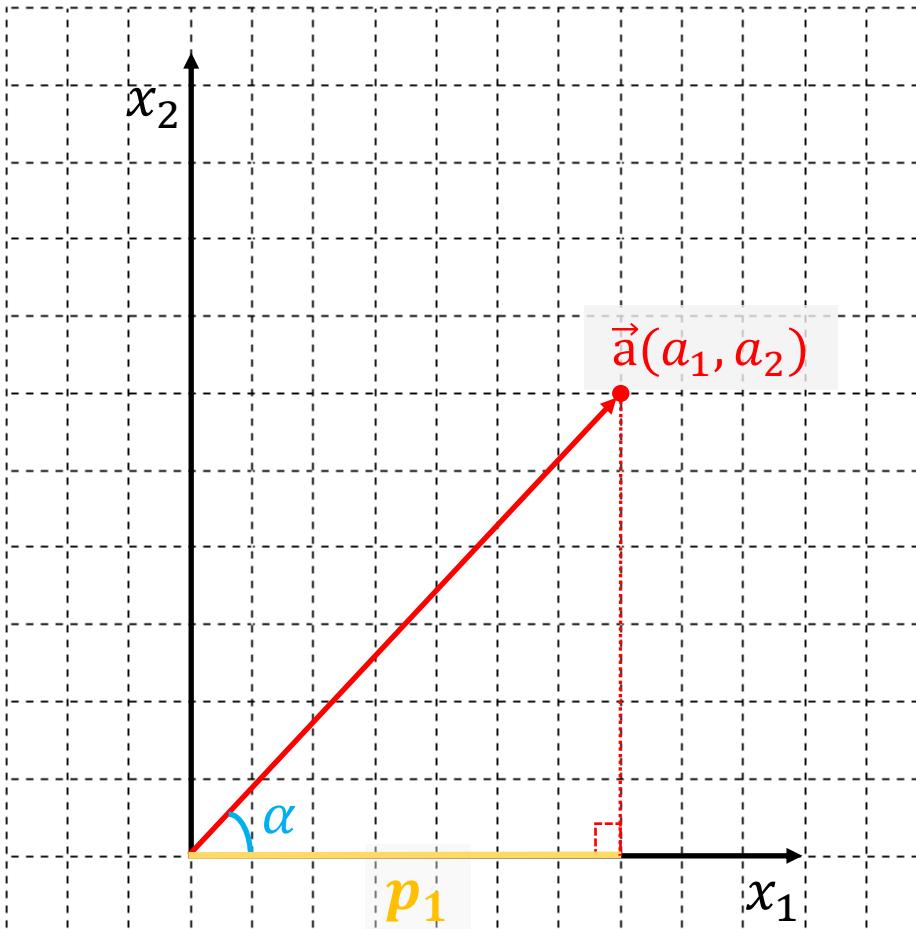


$$\vec{u} = \begin{pmatrix} 3 \\ 2 \end{pmatrix} \quad \vec{v} = 3 * \vec{u} = \begin{pmatrix} 9 \\ 6 \end{pmatrix} \quad \vec{m} = 3 + \vec{u} = \begin{pmatrix} 6 \\ 5 \end{pmatrix}$$

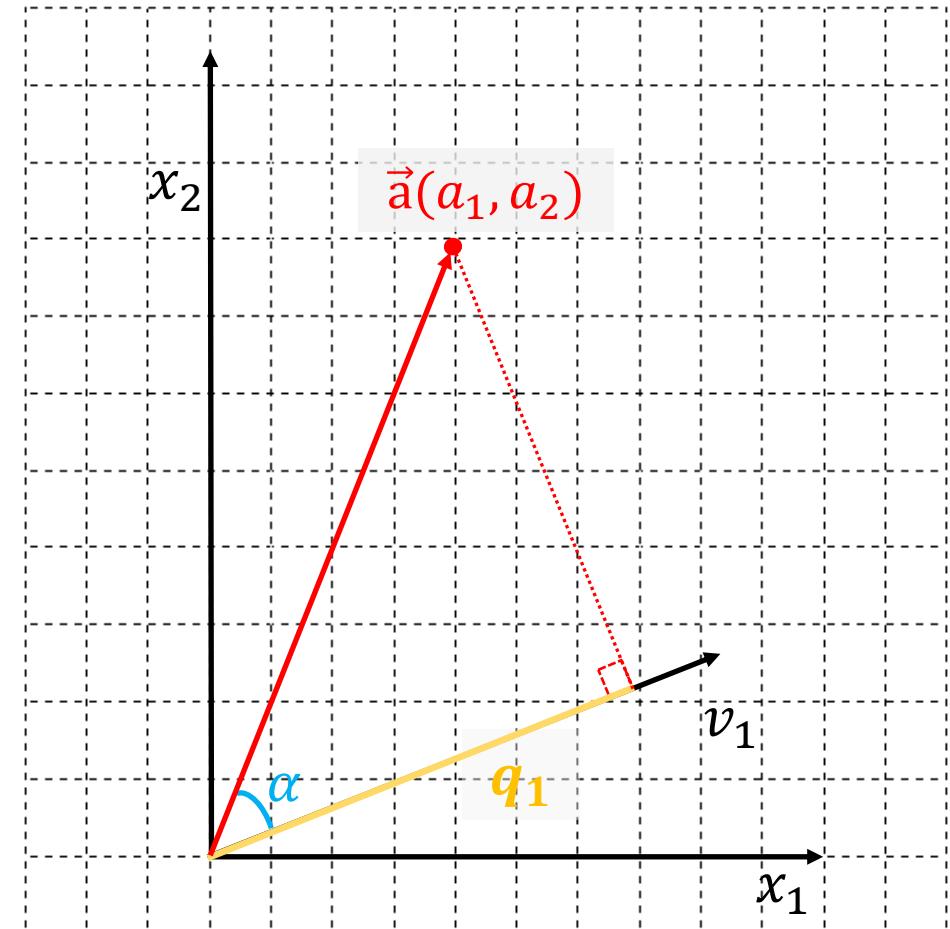


Dot Product

$$\vec{x}_1 = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \quad \vec{x}_2 = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$



$$\vec{v}_1 = \begin{pmatrix} v_{1_1} \\ v_{1_2} \end{pmatrix} \quad \vec{v}_2 = \begin{pmatrix} v_{2_1} \\ v_{2_2} \end{pmatrix}$$



Dot Product

❖ Definition

Algebra

$$\vec{x} \cdot \vec{y} = \sum_{i=1}^n x_i y_i$$

$$\begin{bmatrix} 4 & 0 & 3 & 1 \end{bmatrix} \cdot \begin{bmatrix} 4 \\ 0 \\ 3 \\ 1 \end{bmatrix} = 26$$

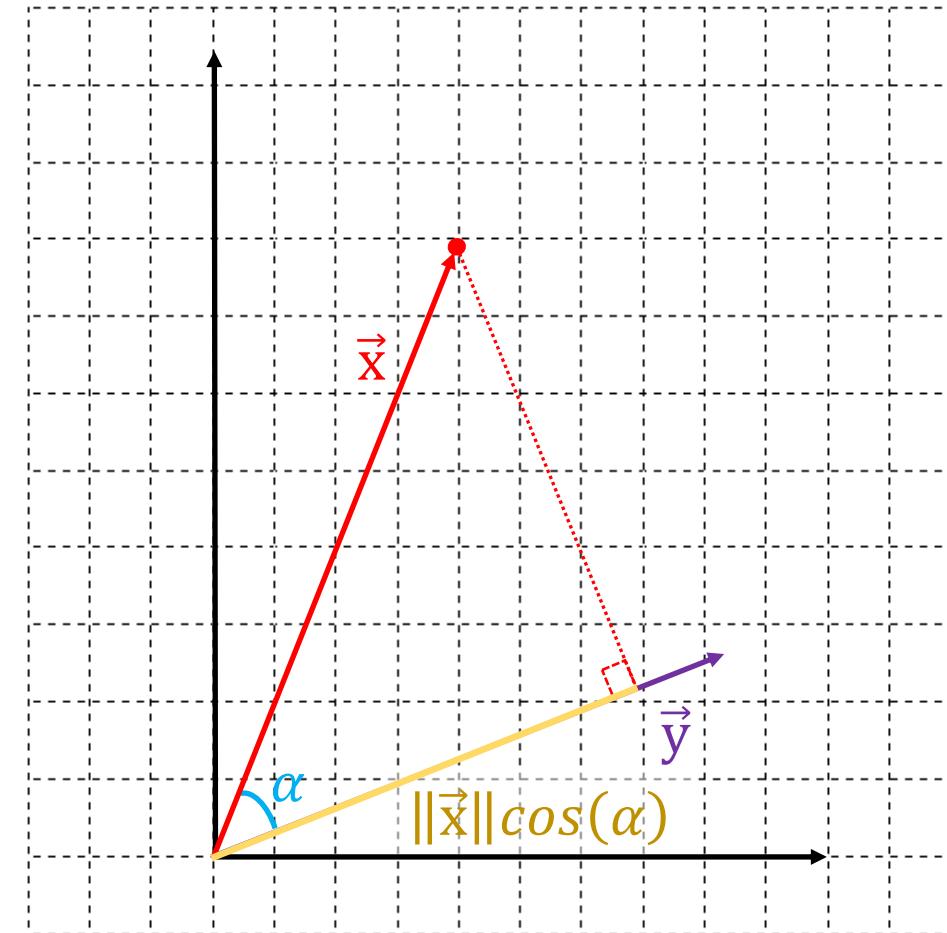
Position-invariant measure

Geometry

$$\vec{x} \cdot \vec{y} = \|\vec{x}\| \|\vec{y}\| \cos(\alpha)$$

$$\begin{bmatrix} 3 & 0 & 4 & 1 \end{bmatrix} \cdot \begin{bmatrix} 3 \\ 0 \\ 4 \\ 1 \end{bmatrix} = 26$$

$$\|\vec{x}\| \cos(\alpha)$$



Dot Product

❖ Definition

Algebra

$$\vec{x} \cdot \vec{y} = \sum_1^n x_i y_i$$

Geometry

$$\vec{x} \cdot \vec{y} = \|\vec{x}\| \|\vec{y}\| \cos(\alpha)$$

```
1 import numpy as np
2
3 u = np.array([0, 5])
4 v = np.array([3, 3])
5
6 dproduct_1 = np.dot(u, v)
7 dproduct_2 = np.dot(v, u)
8
9 print(dproduct_1)
10 print(dproduct_2)
```

```
1 # geometry
2
3 norm_u = np.linalg.norm(u)
4 norm_v = np.linalg.norm(v)
5
6 alpha = 3.14/4 # 45 degree
7 dproduct_3 = np.cos() * norm_u * norm_v
8 print(dproduct_3)
```

Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

light green	light green	light green	light green
light blue	light blue	light blue	light blue
light yellow	light yellow	light yellow	light yellow

V

light green	light blue	light yellow
light green	light blue	light yellow
light green	light blue	light yellow

V^T

light orange	light blue	light orange
light orange	light blue	light orange
light orange	light blue	light orange

Gram matrix

```
import torch

x = torch.Tensor([[4, 0, 3, 1],
                  [3, 0, 1, 2],
                  [4, 5, 0, 3]])
G = torch.mm(x, x.t())
print(G)

tensor([[26., 17., 19.],
        [17., 14., 18.],
        [19., 18., 50.]])
```

4	0	3	1
3	0	1	2
4	5	0	3

V

4	3	4
0	0	5
3	1	0
1	2	3

V^T

26	17	19
17	14	18
19	18	50

Gram matrix

Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

light green	light green	light green	light green
light blue	light blue	light blue	light blue
light yellow	light yellow	light yellow	light yellow

V

light green	light blue	light yellow
light green	light blue	light yellow
light green	light blue	light yellow

•

V^T

light orange	light blue	light green
light orange	light blue	light green
light orange	light blue	light green

Gram matrix

```
import torch

x = torch.Tensor([[4, 0, 3, 1],
                  [3, 0, 1, 2],
                  [4, 5, 0, 3]])
G = torch.einsum('ij,jk->ik', x, x.t())
print(G)

tensor([[26., 17., 19.],
        [17., 14., 18.],
        [19., 18., 50.]])
```

4	0	3	1
3	0	1	2
4	5	0	3

V

4	3	4
0	0	5
3	1	0
1	2	3

•

V^T

26	17	19
17	14	18
19	18	50

Gram matrix

Detailed discussion on the Einsum scheduled for Friday

Gram Matrix

$$G_{ij} = \sum_k V_{ik} V_{jk}$$

4	0	3	1
3	0	1	2
4	5	0	3

V

4	3	4
0	0	5
3	1	0

V^T

26	17	19
17	14	18
19	18	50

Gram matrix

Inner Product

Algebra

$$\vec{x} \cdot \vec{y} = \sum_1^n x_i y_i$$

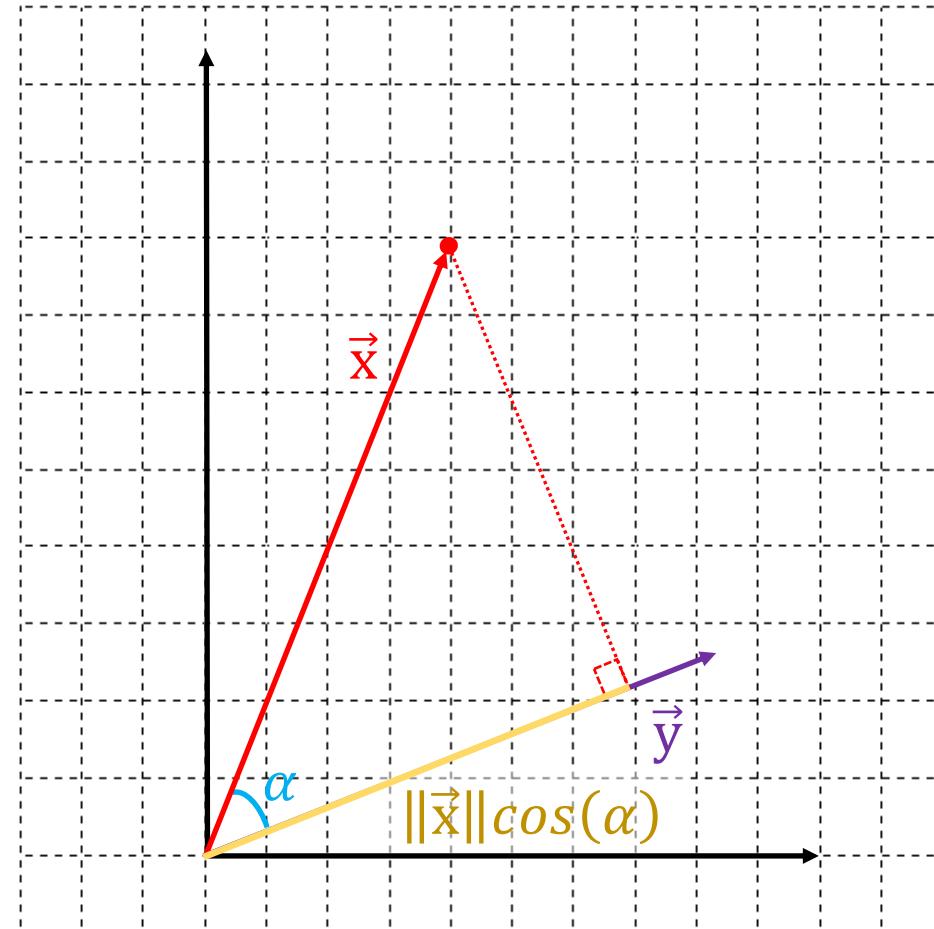
$$\begin{bmatrix} 4 & 0 & 3 & 1 \end{bmatrix} \cdot \begin{bmatrix} 4 \\ 0 \\ 3 \\ 1 \end{bmatrix} = 26$$

Geometry

$$\vec{x} \cdot \vec{y} = \|\vec{x}\| \|\vec{y}\| \cos(\alpha)$$

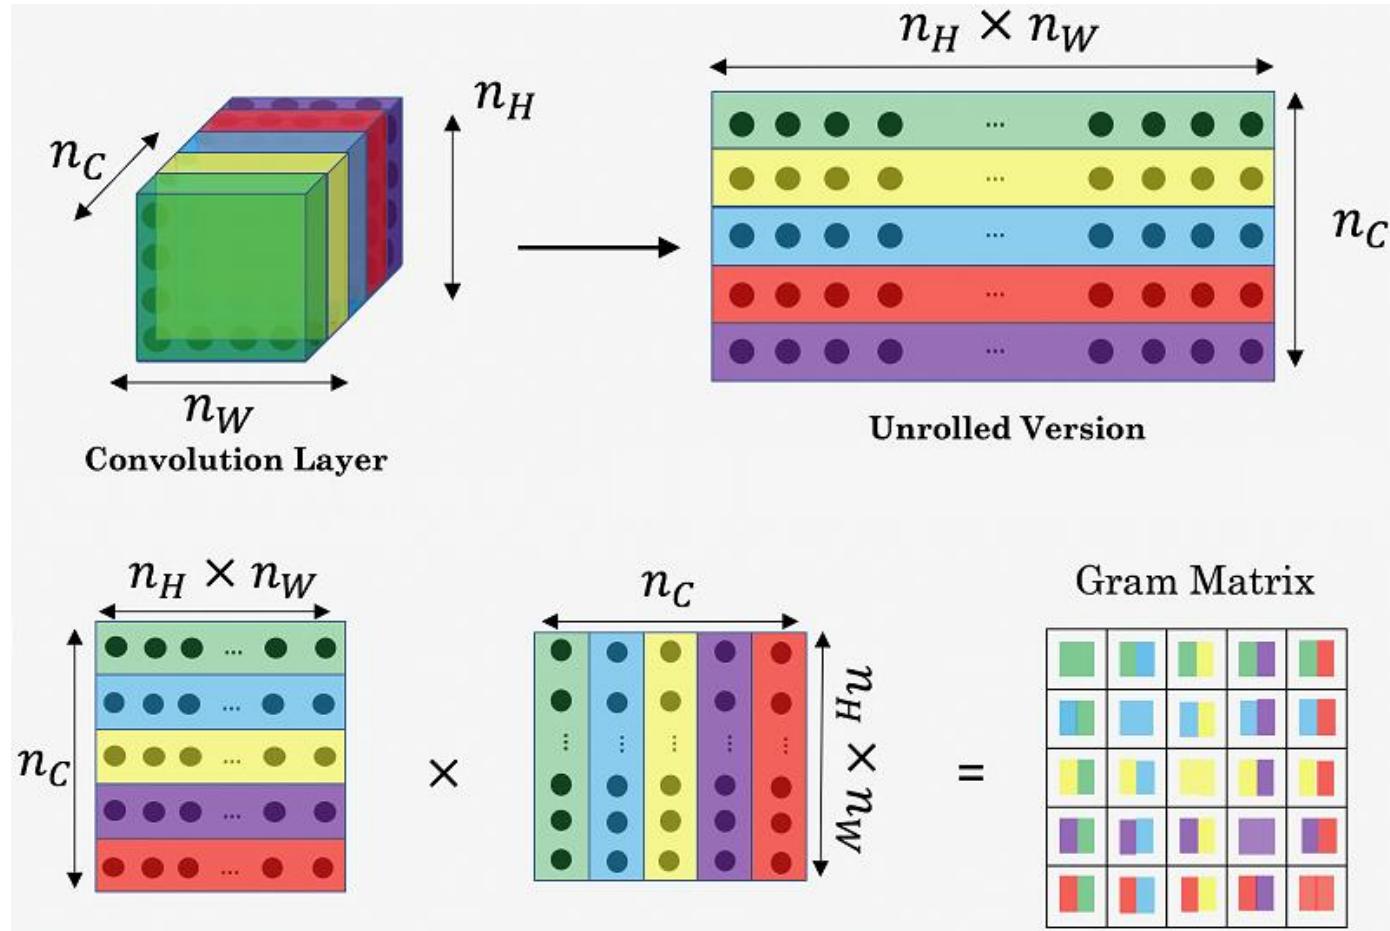
$$\begin{bmatrix} 3 & 0 & 4 & 1 \end{bmatrix} \cdot \begin{bmatrix} 3 \\ 0 \\ 4 \\ 1 \end{bmatrix} = 26$$

Position-invariant measure

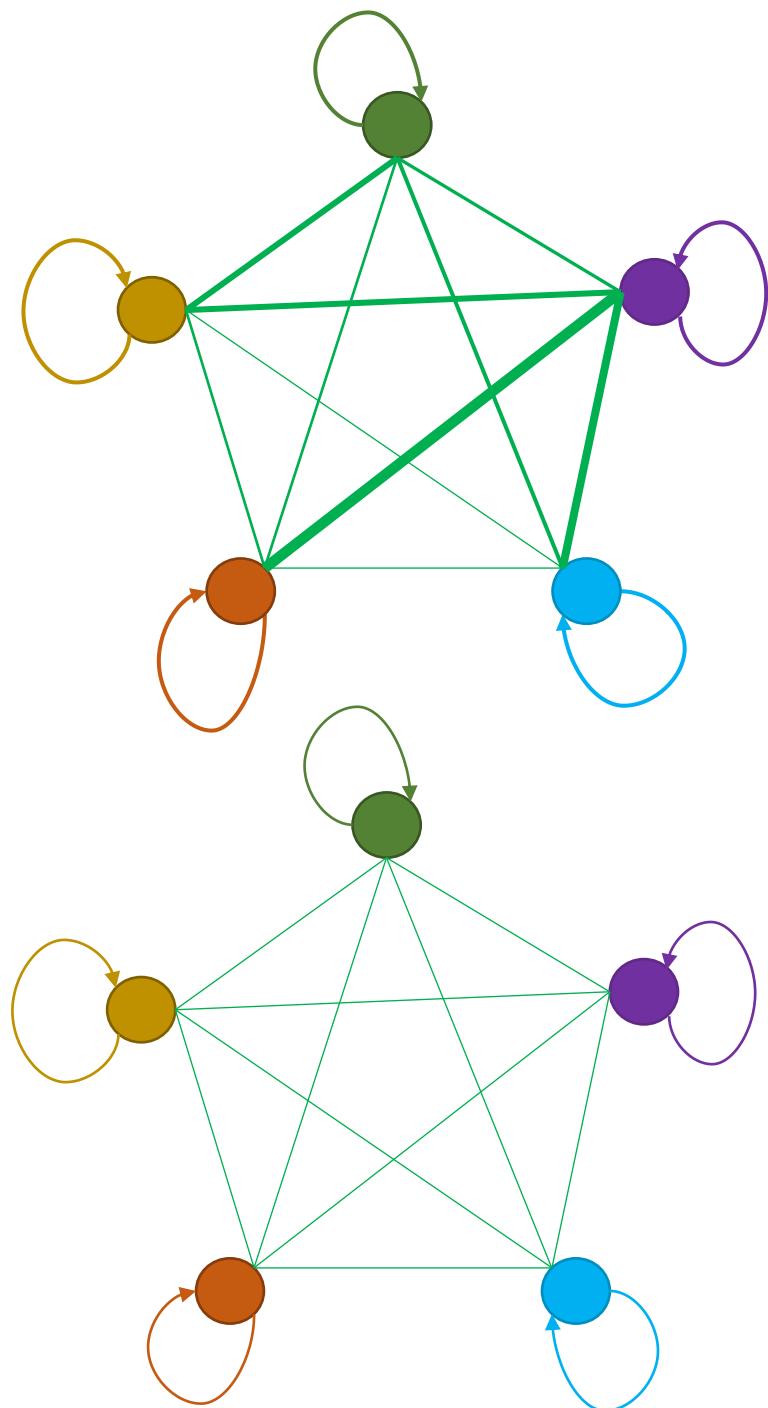


Style Transfer

❖ Gram matrix

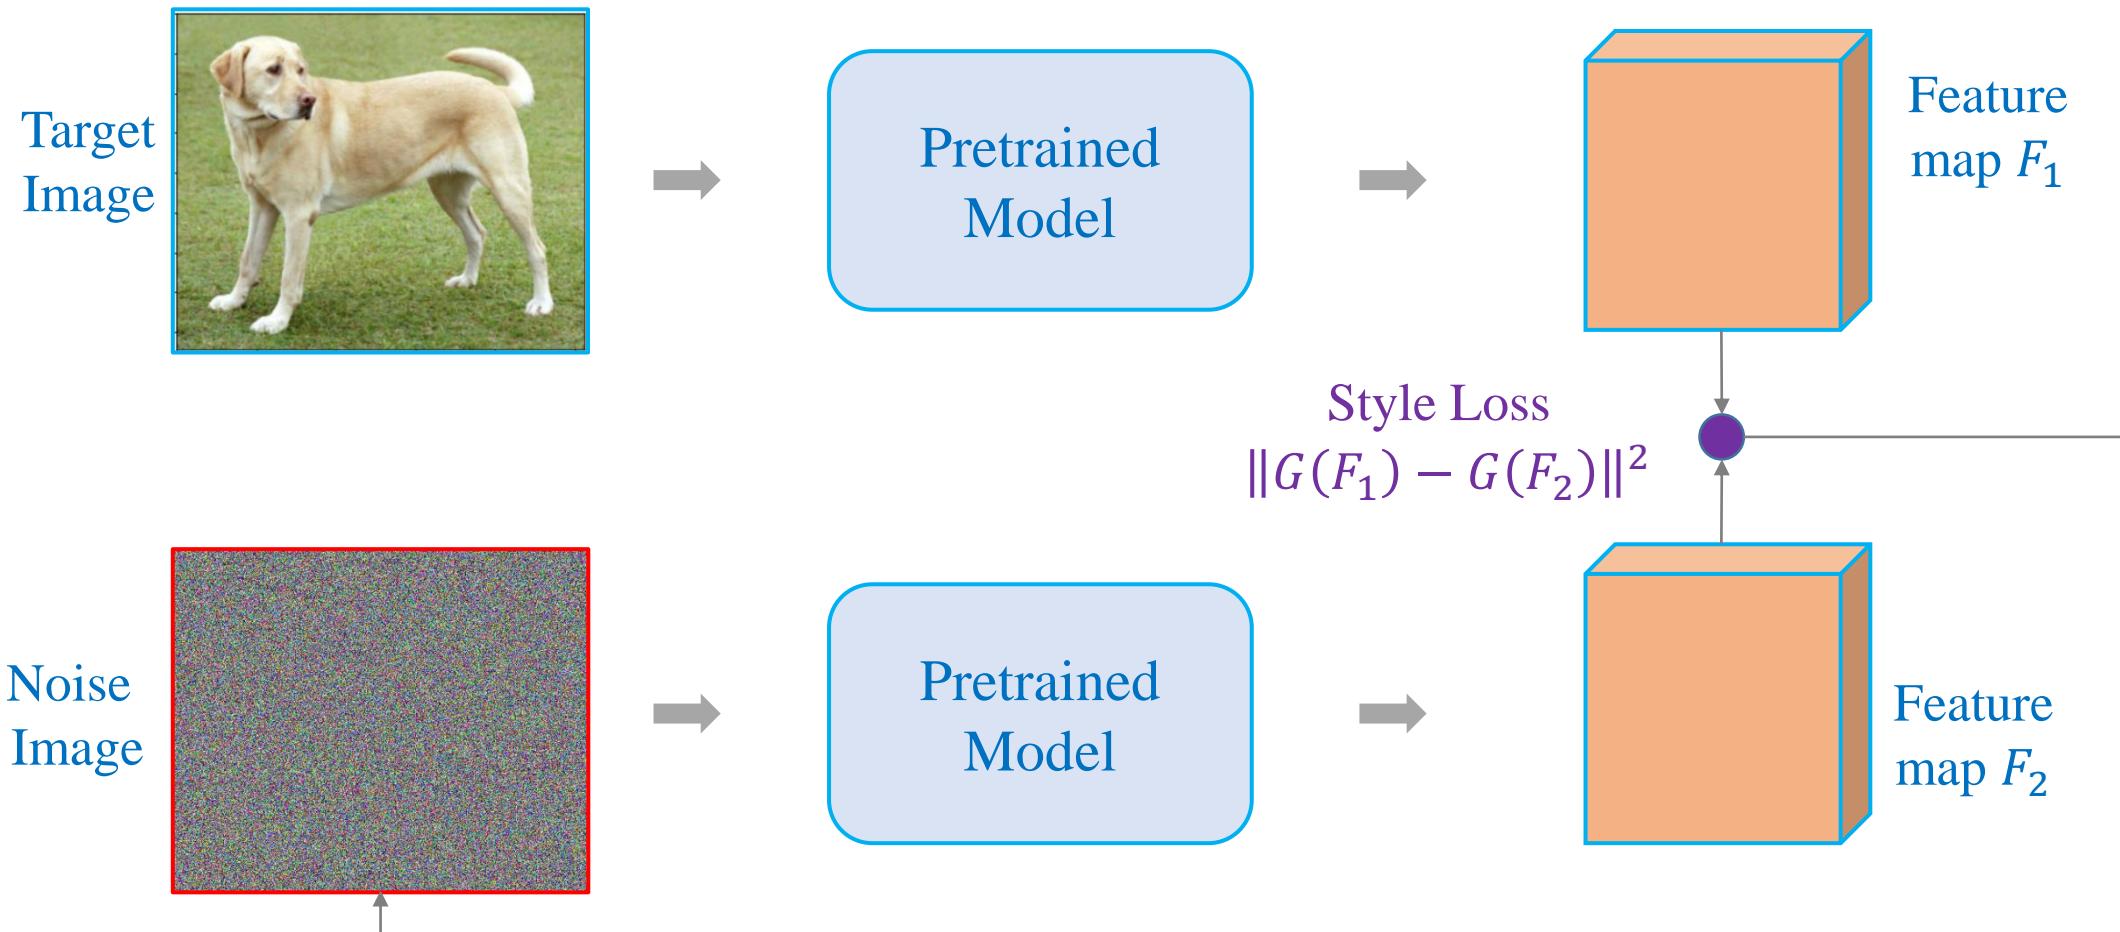


<https://pytorchtaipei.github.io/articles/PyTorchTP-Style-Transfer/>



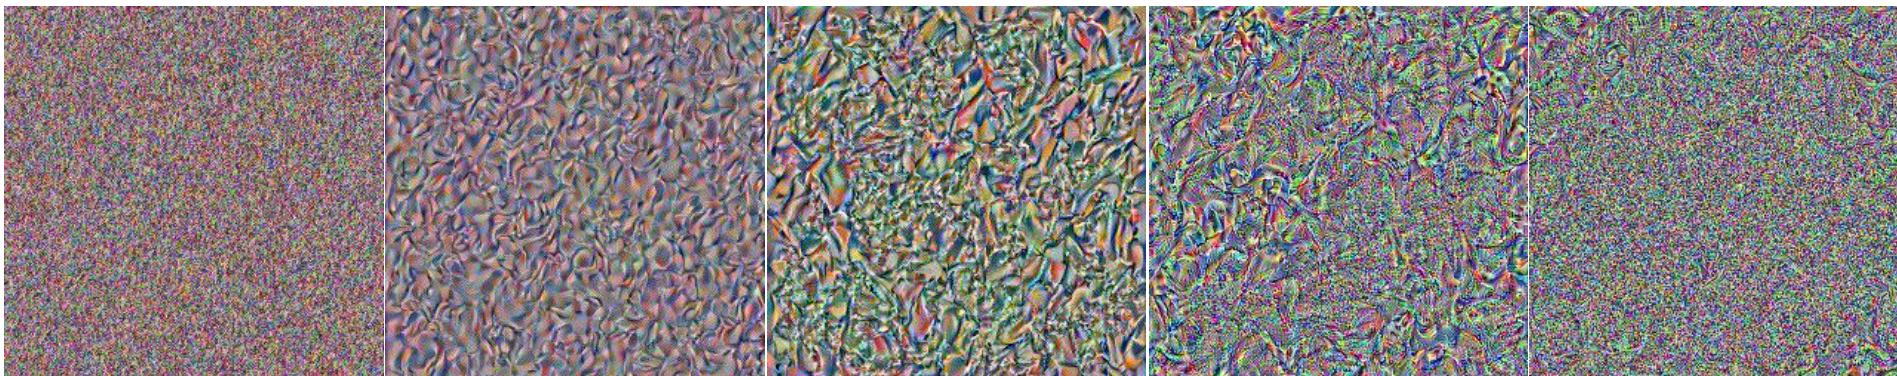
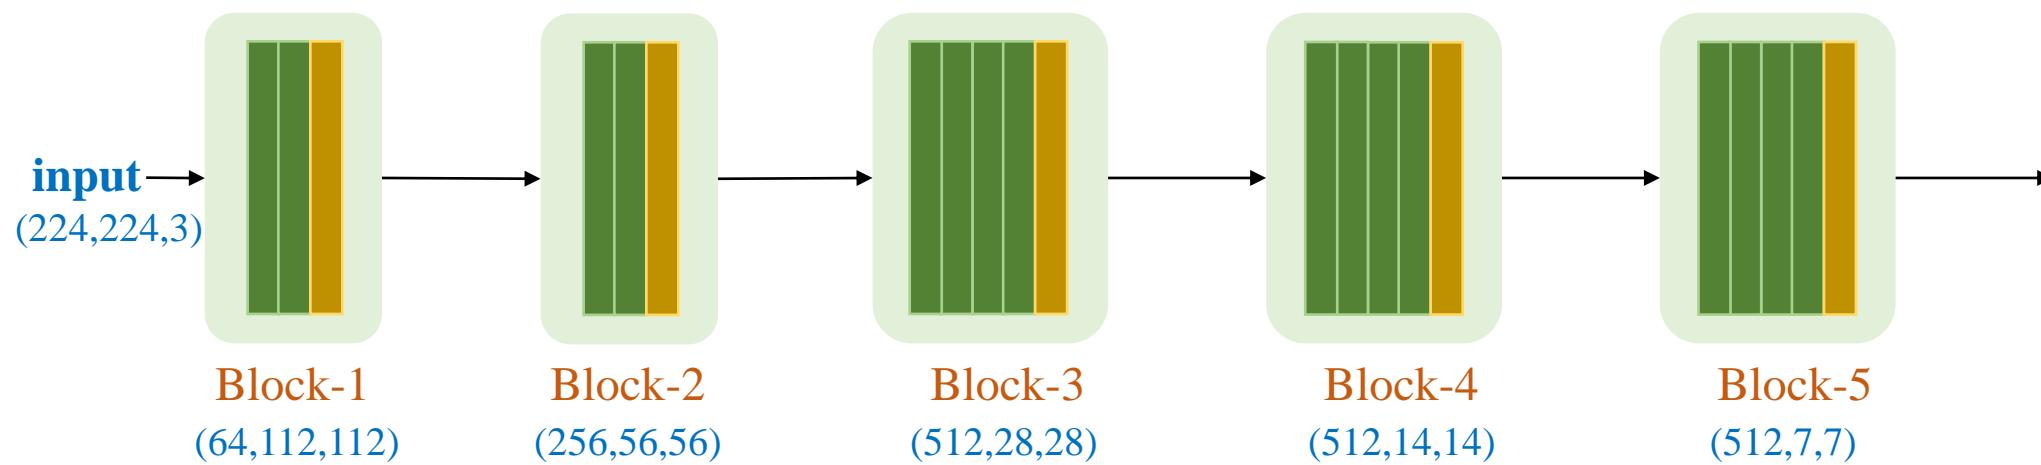
Style Transfer

❖ Style Loss



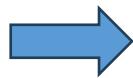
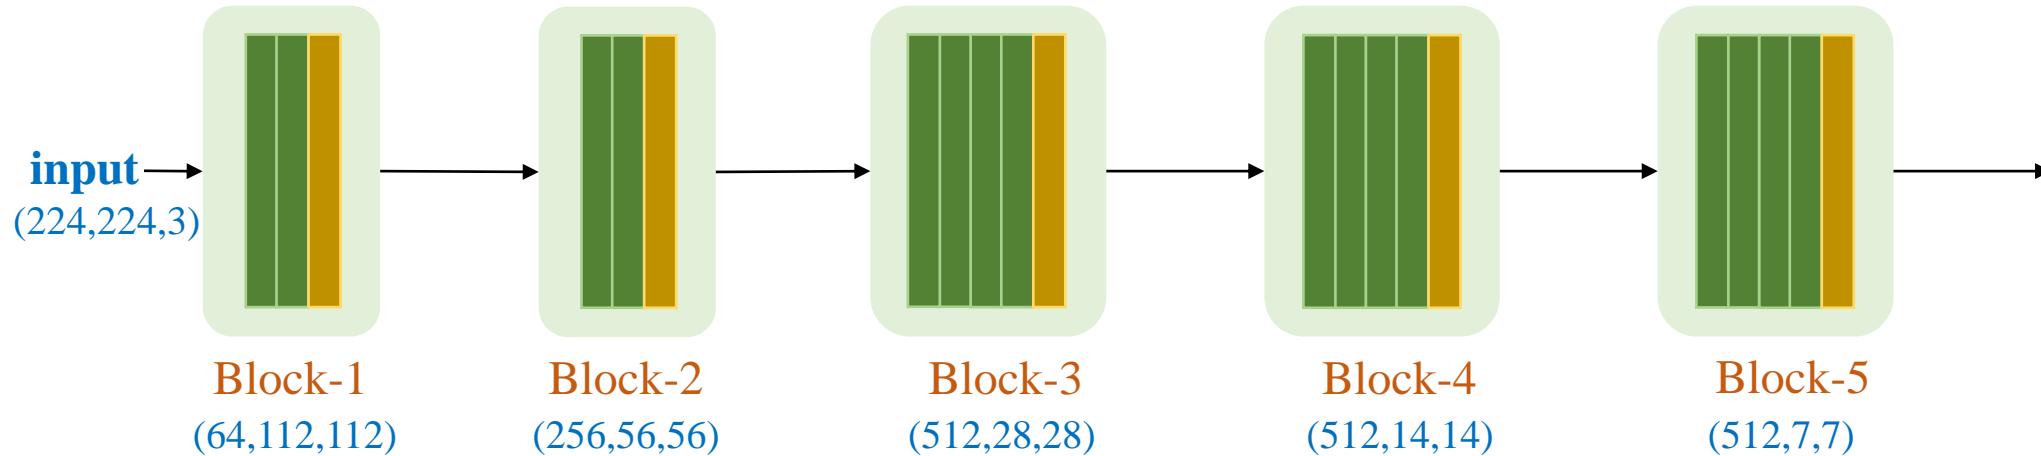
Style Transfer

❖ Style Loss: Use which features?



Style Transfer

❖ Style Loss: Use which features?





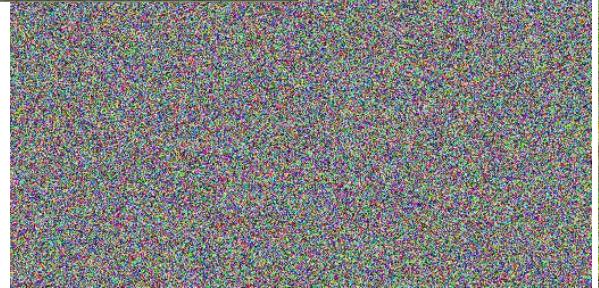
Initial Image



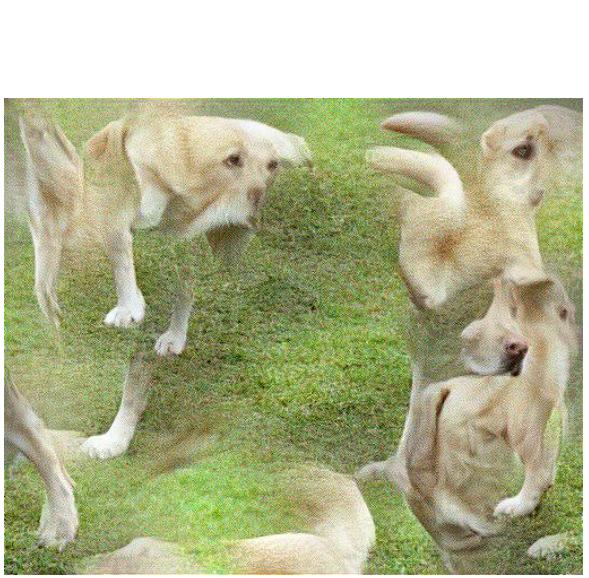
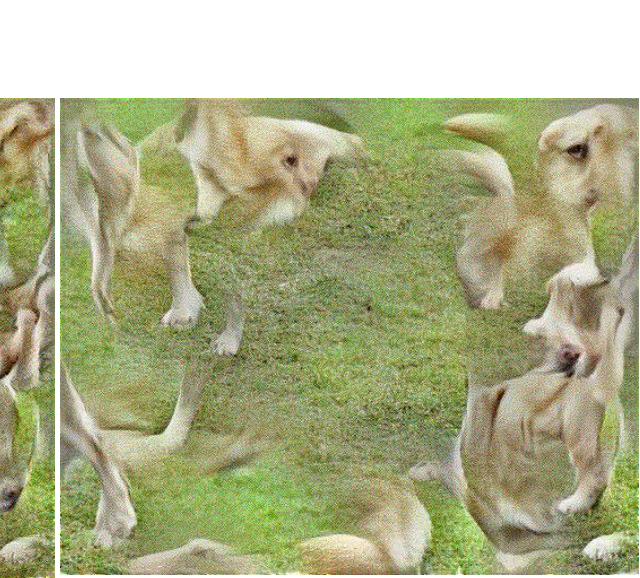
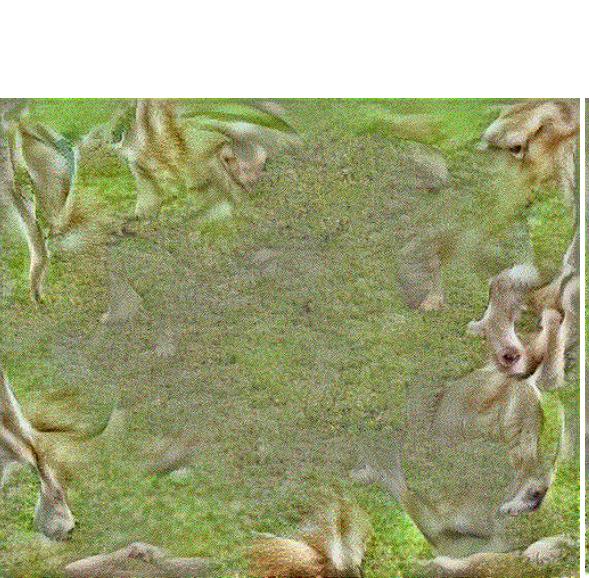
Epoch 10

Epoch 40

Epoch 90



Initial Image



Epoch 10

Epoch 40

Epoch 90

QUIZ TIME

Outline

- Introduction
- Inputs as Variables
- Content Loss
- Style Loss

Summary

- ✓ Studied about content loss
- ✓ Studied about style loss

- ✓ Studying about the total loss
- ✓ Extend to some variants

