

# Time-series Forecasting Project

*(MLP, RNN, LSTM, Bi-LSTM, XGBoost)*

Vinh Dinh Nguyen  
PhD in Computer Science

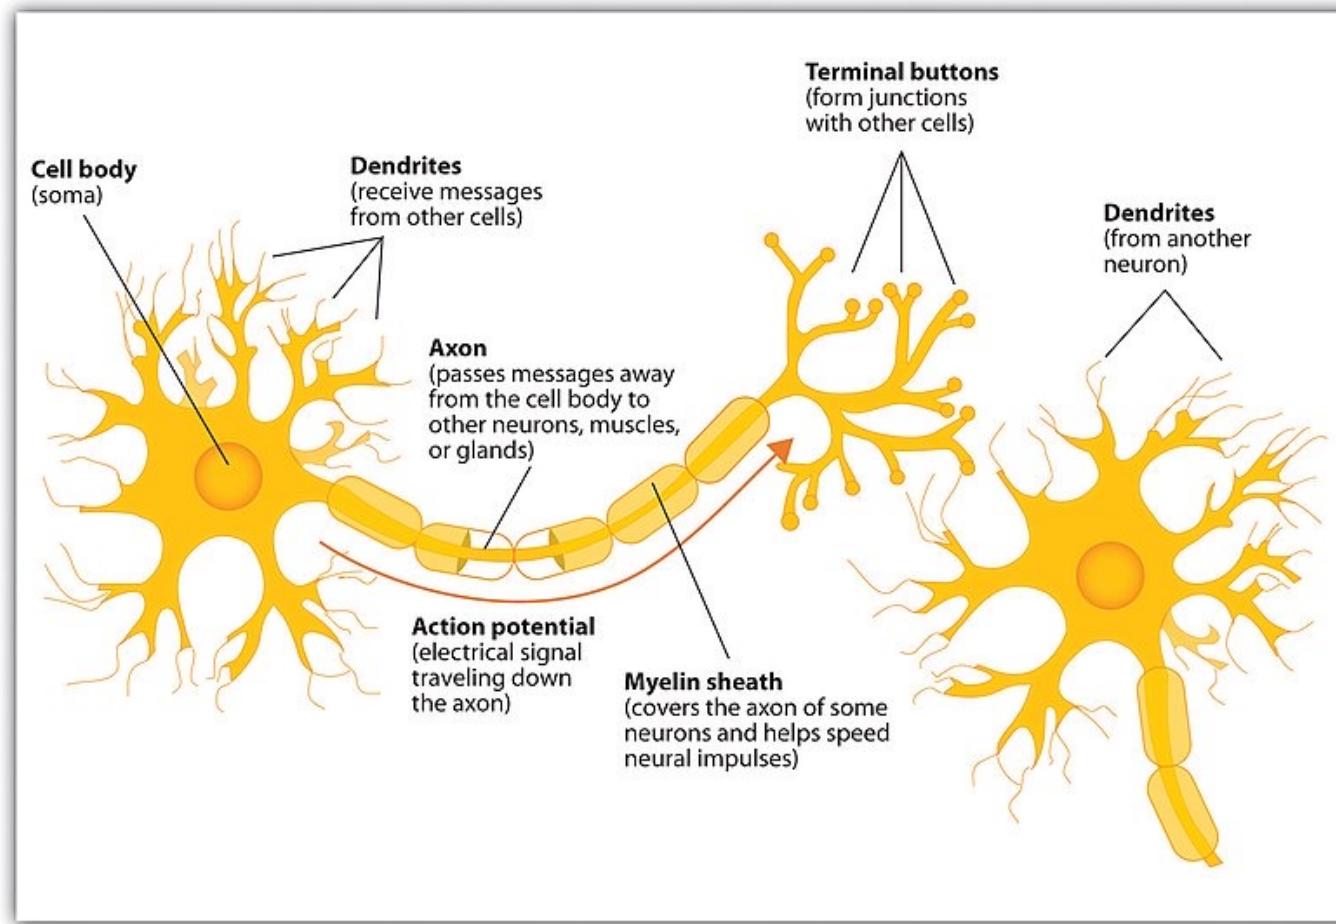
# Outline

- Multilayer Perceptron and Time-series Data Forecasting
- RNN and Time-series Data Forecasting
- LSTM and Time-series Data Forecasting
- Bi-LSTM and Time-series Data Forecasting
- XGBoost and Time-series Data Forecasting

# Outline

- **Multilayer Perceptron and Time-series Data Forecasting**
- **RNN and Time-series Data Forecasting**
- **LSTM and Time-series Data Forecasting**
- **Bi-LSTM and Time-series Data Forecasting**
- **XGBoost and Time-series Data Forecasting**

# Neuron: Simple Explanation



Neuron and it's different components

In the early 1940's [Warren McCulloch](#), a neurophysiologist, teamed up with logician [Walter Pitts](#) to create a model of how brains work. It was a simple linear model that produced a positive or negative output, given a set of inputs and weights.

$$f(x, w) = \underbrace{x_1 w_1 + \cdots + x_n w_n}_{\text{inputs}} / \text{weights}$$

# Neuron: Simple Explanation

The first application of the neuron replicated a logic gate, where you have one or two binary inputs, and a boolean function that only gets activated given the right inputs and weights

AND FUNCTION

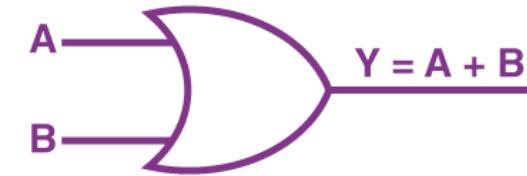


A	B	Y
0	0	0
0	1	0
1	0	0
1	1	1

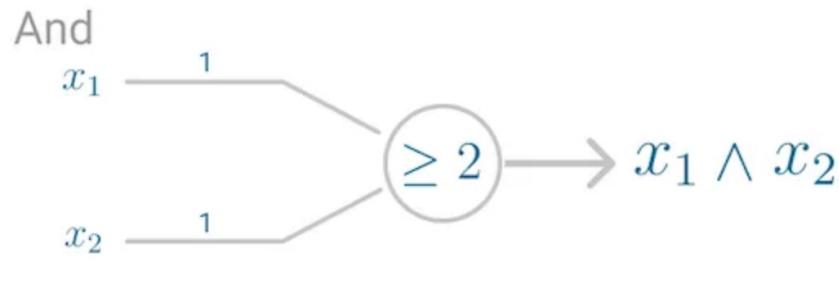


It couldn't *learn* like the brain.

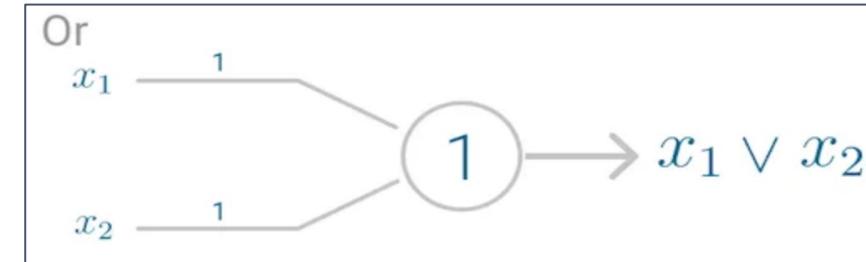
OR FUNCTION



A	B	Y
0	0	0
0	1	1
1	0	1
1	1	1

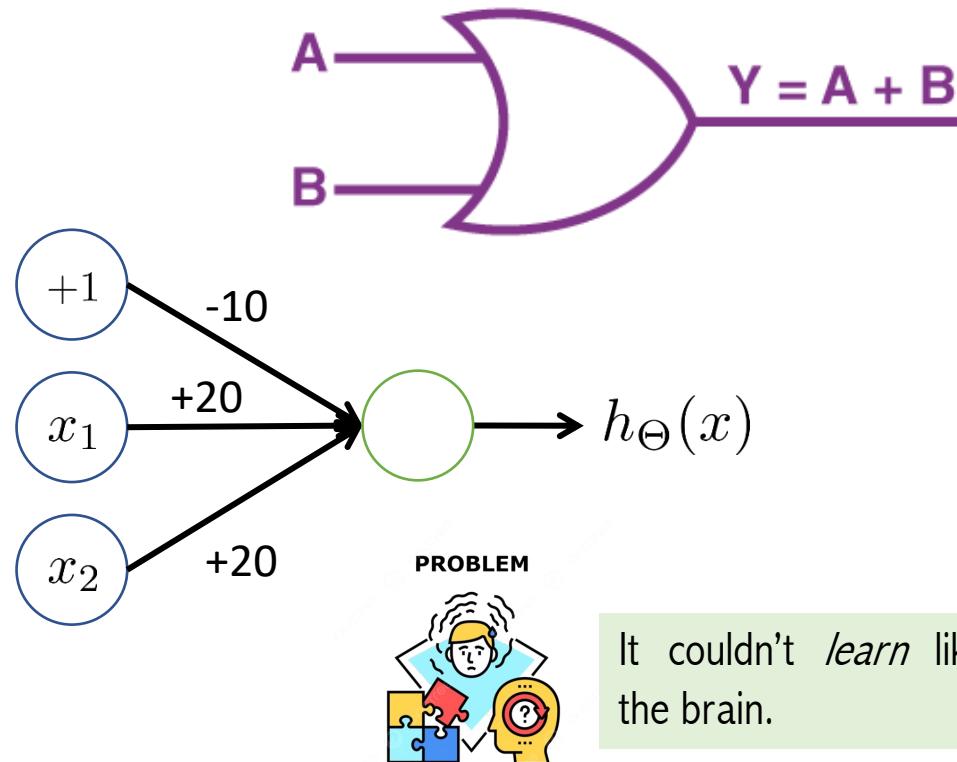


Frank Rosenblatt  
extended his model



# Neuron: Simple Explanation

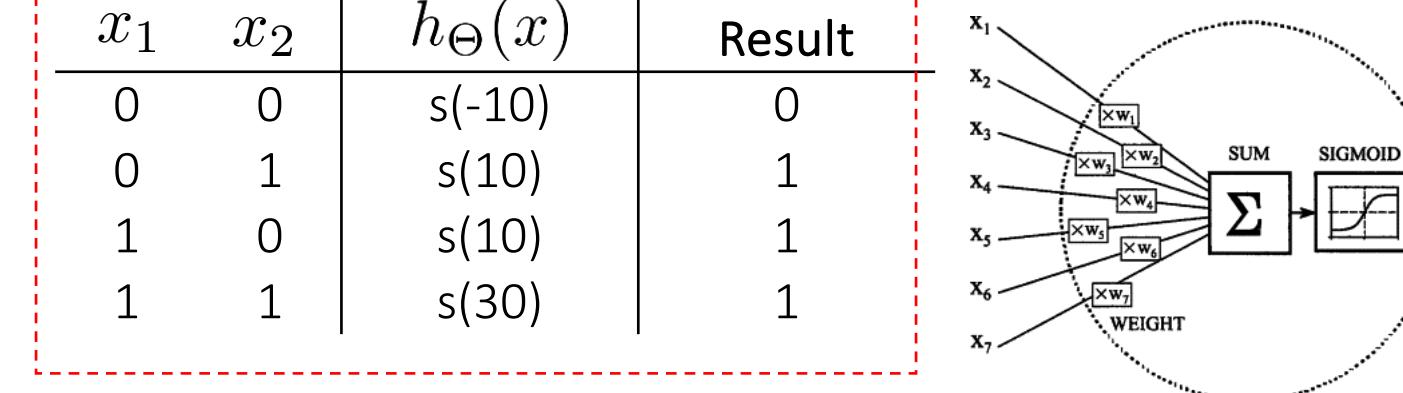
OR FUNCTION



A	B	Y
0	0	0
0	1	1
1	0	1
1	1	1

$x_1$	$x_2$	$h_{\Theta}(x)$	Result
0	0	$s(-10)$	0
0	1	$s(10)$	1
1	0	$s(10)$	1
1	1	$s(30)$	1

$$S(x) = \frac{1}{1 + e^{-x}}$$

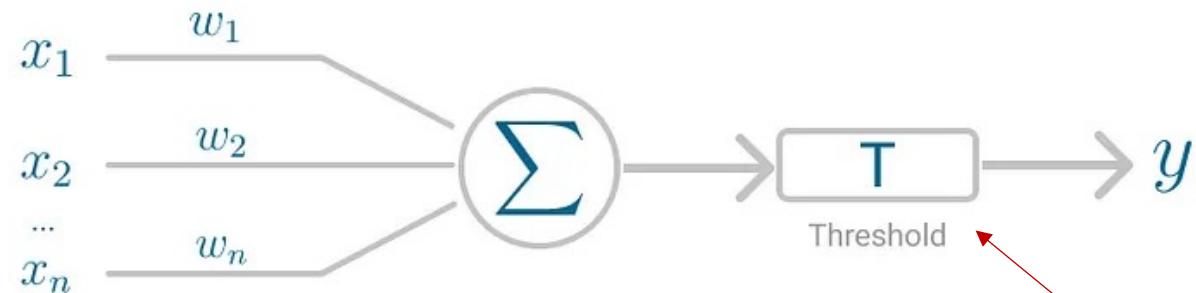


It was only a decade later that [Frank Rosenblatt](#) extended this model, and created an algorithm that could *learn* the weights in order to generate an output.

# Perceptron: Simple Explanation

[Frank Rosenblatt](#) extended neuron model, and created an algorithm that could learn the weights in order to generate an output.  
It gets its name from performing the human-like function of perception, seeing and recognizing images.

Perceptrons neuron model (left) and threshold logic (right).



$$y = \begin{cases} 1, & \text{if } \frac{\text{weighted sum}}{\sum_i w_i x_i - T} > 0 \\ 0, & \text{otherwise} \end{cases}$$

threshold

# Perception: Activation Functions

ACTIVATION FUNCTION	PLOT	EQUATION	DERIVATIVE	RANGE
Linear		$f(x) = x$	$f'(x) = 1$	$(-\infty, \infty)$
Binary Step		$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{if } x \neq 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$	$\{0, 1\}$
Sigmoid		$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$	$(0, 1)$
Hyperbolic Tangent(tanh)		$f(x) = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$	$f'(x) = 1 - f(x)^2$	$(-1, 1)$
Rectified Linear Unit(ReLU)		$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ \text{undefined} & \text{if } x = 0 \end{cases}$	$[0, \infty)$
Softplus		$f(x) = \ln(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$	$(0, \infty)$
Leaky ReLU		$f(x) = \begin{cases} 0.01x & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0.01 & \text{if } x < 0 \\ 1 & \text{if } x \geq 0 \end{cases}$	$(-1, 1)$
Exponential Linear Unit(ELU)		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{if } x \leq 0 \\ x & \text{if } x > 0 \end{cases}$	$f'(x) = \begin{cases} \alpha e^x & \text{if } x < 0 \\ 1 & \text{if } x > 0 \\ 1 & \text{if } x = 0 \text{ and } \alpha = 1 \end{cases}$	$[0, \infty)$

But you might be wondering, Doesn't Perceptron actually learn the weights?

It does! Perceptron uses Stochastic Gradient Descent to find!

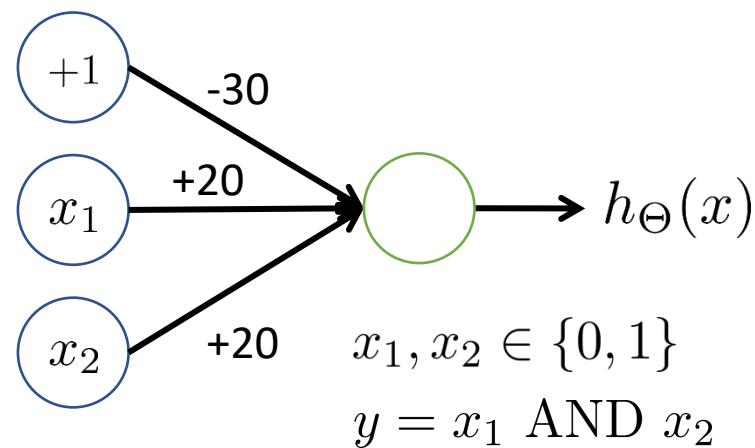
# Perceptron: Simple Explanation

The first application of the neuron replicated a logic gate, where you have one or two binary inputs, and a boolean function that only gets activated given the right inputs and weights

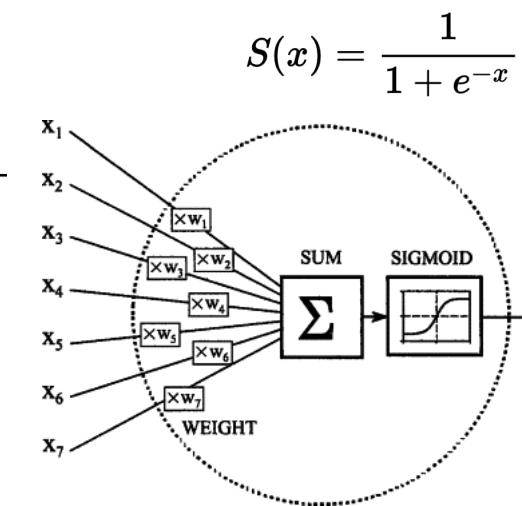
## AND FUNCTION



A	B	Y
0	0	0
0	1	0
1	0	0
1	1	1

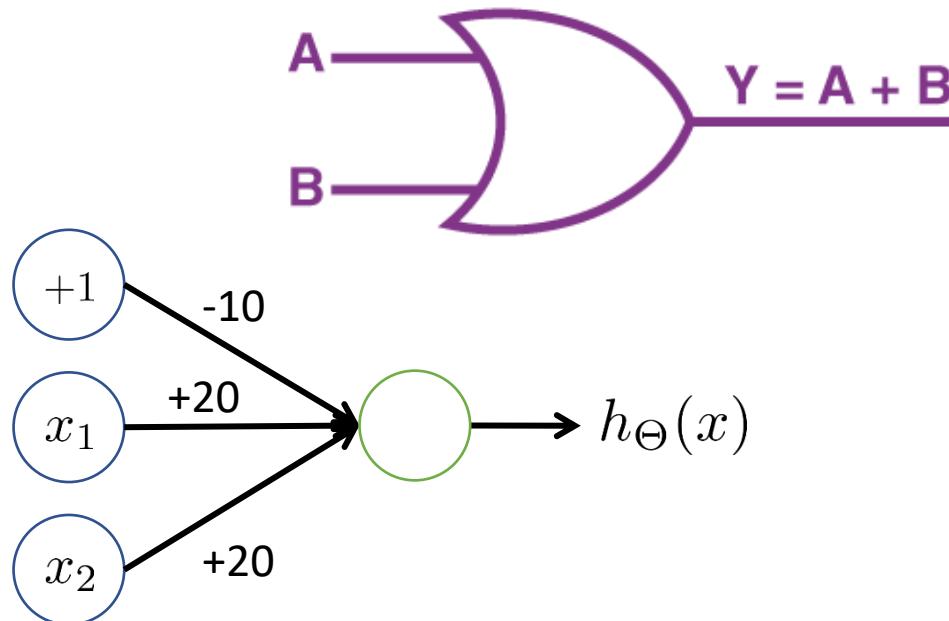


$x_1$	$x_2$	$h_{\Theta}(x)$	Result
0	0	$s(-30)$	0
0	1	$s(-10)$	0
1	0	$s(-10)$	0
1	1	$s(10)$	1



# Perceptron: Simple Explanation

OR FUNCTION



PROBLEM

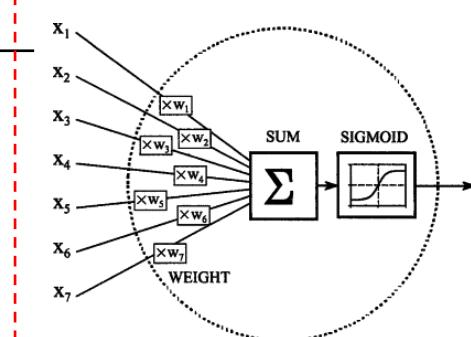


it couldn't represent the [XOR gate](#), *exclusive OR*, where the gate only returns 1 if the inputs are different

A	B	Y
0	0	0
0	1	1
1	0	1
1	1	1

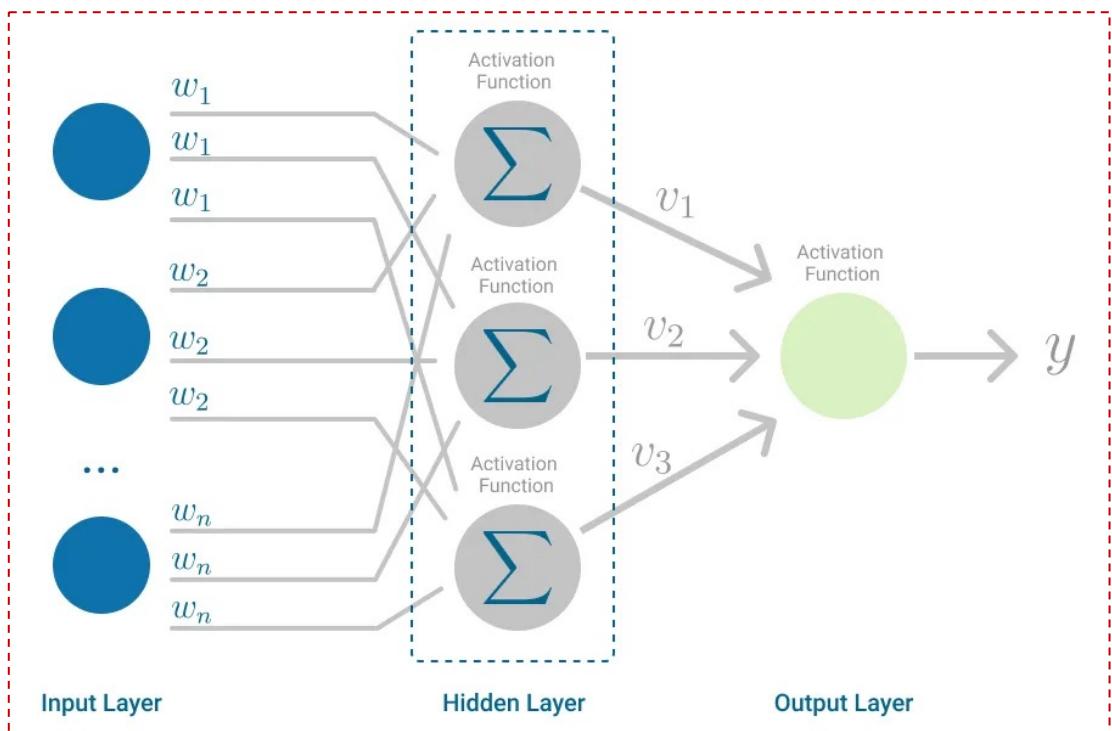
x <sub>1</sub>	x <sub>2</sub>	h <sub>Θ</sub> (x)	Result
0	0	s(-10)	0
0	1	s(10)	1
1	0	s(10)	1
1	1	s(30)	1

$$S(x) = \frac{1}{1 + e^{-x}}$$

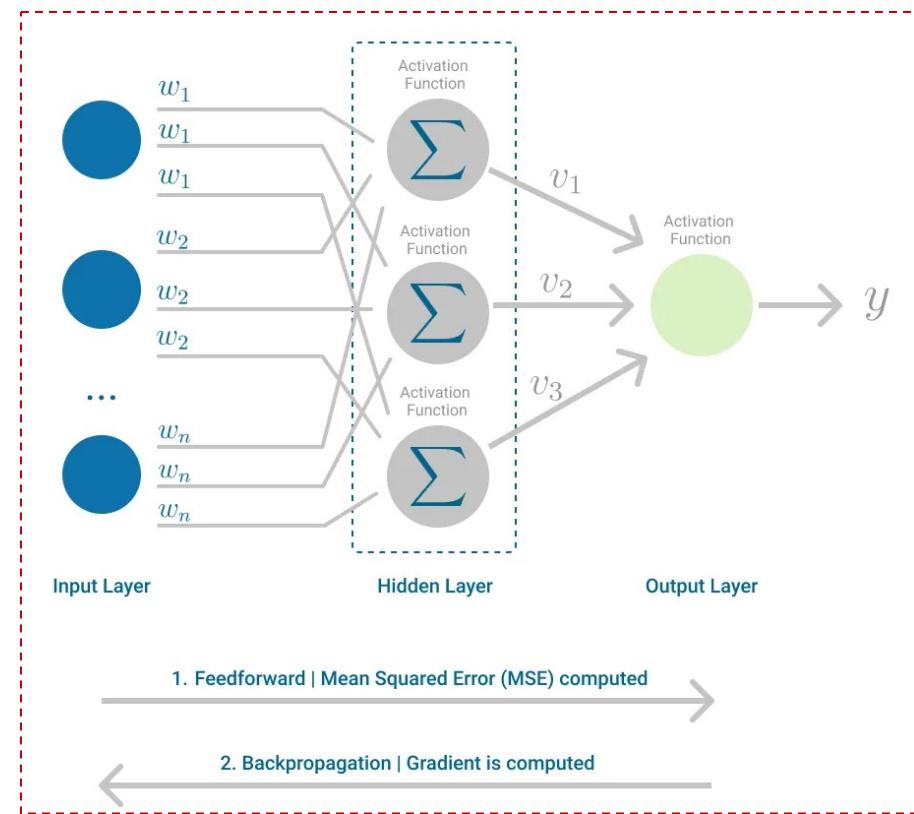


# Multilayer Perceptron: Simple Explanation

The **Multilayer Perceptron** was developed to tackle this limitation. It is a neural network where the mapping between inputs and output is non-linear.



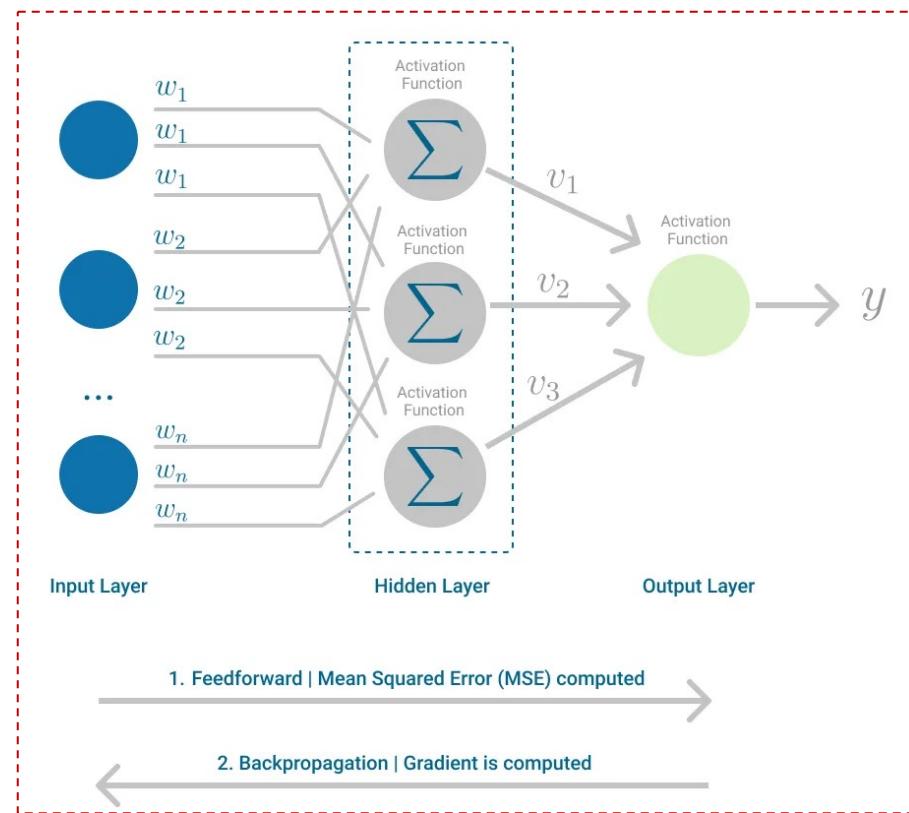
*Multilayer Perceptron*



*the Feedforward and Backpropagation steps*

# Multilayer Perceptron: Simple Explanation

The **Multilayer Perceptron** was developed to tackle this limitation. It is a neural network where the mapping between inputs and output is non-linear.

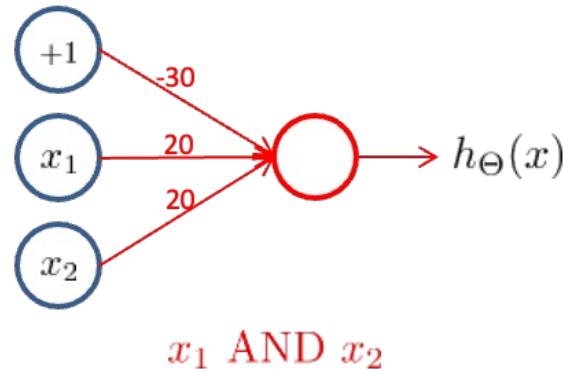


$$\Delta_w(t) = -\varepsilon \frac{\frac{\partial E}{\partial w(t)}}{\frac{\partial E}{\partial b}} + \alpha \Delta_w(t-1)$$

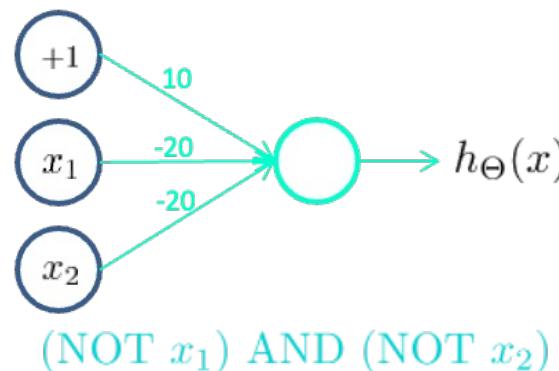
## *One iteration of Gradient Descent*

### *the Feedforward and Backpropagation steps*

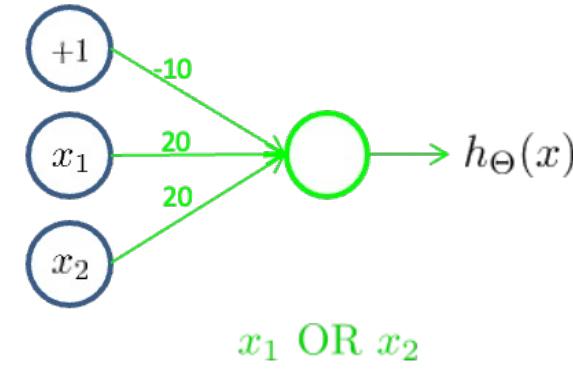
# Multilayer Perceptron: Simple Explanation



$x_1$  AND  $x_2$

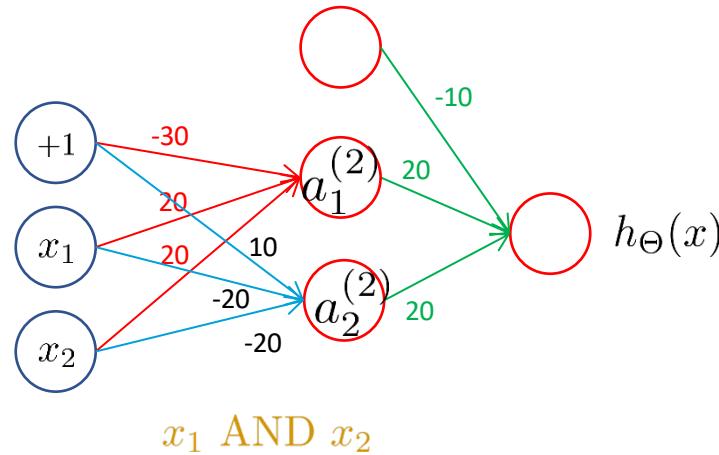


(NOT  $x_1$ ) AND (NOT  $x_2$ )



$x_1$  OR  $x_2$

$x_1$  XNOR  $x_2$



$x_1$  AND  $x_2$

$x_1$	$x_2$	$a_1^{(2)}$	$a_2^{(2)}$	$h_{\Theta}(x)$
0	0	0	1	1
0	1	0	0	0
1	0	0	0	0
1	1	1	0	1

# Multilayer Perceptron: Time-series Data Project

## Electricity Transformer Dataset (ETDataset)

	date	HUFL	HULL	MUFL	MULL	LUFL	ULLL	OT
0	2016-07-01 00:00:00	5.827	2.009	1.599	0.462	4.203	1.340	30.531000
1	2016-07-01 00:15:00	5.760	2.076	1.492	0.426	4.264	1.401	30.459999
2	2016-07-01 00:30:00	5.760	1.942	1.492	0.391	4.234	1.310	30.038000
3	2016-07-01 00:45:00	5.760	1.942	1.492	0.426	4.234	1.310	27.013000
4	2016-07-01 01:00:00	5.693	2.076	1.492	0.426	4.142	1.371	27.787001

Field	date	HUFL	HULL	MUFL	MULL	LUFL	ULLL	OT
Description	The recorded date	High Useful Load	High Useless Load	Middle Useful Load	Middle Useless Load	Low Useful Load	Low Useless Load	Oil Temperature (target)

## Electricity Transformer Dataset (ETDataset)

### Data Windowing Techniques

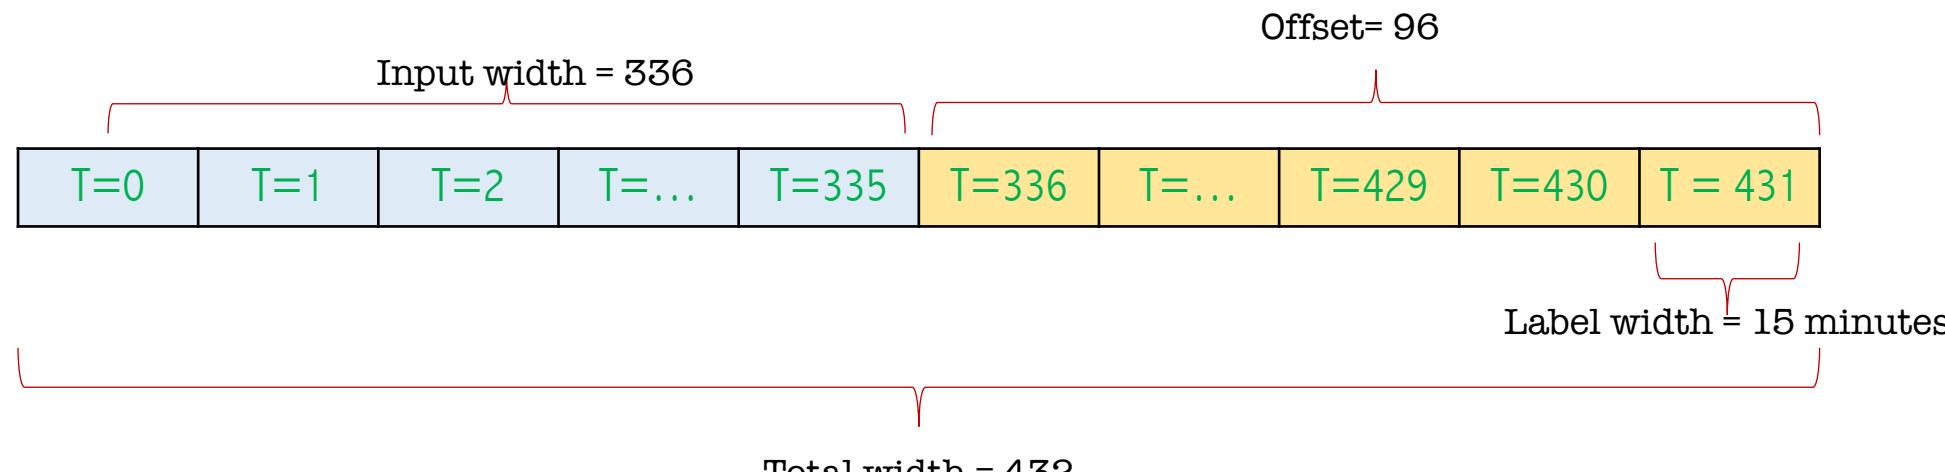
Case 1 : Multivariate to Multivariate Prediction

(Input, Output) = (336, 96)

(Input shape)= (15932, 336, 7)

(Output shape)= (15932, 96, 7)

	date	HUFL	HULL	MUFL	MULL	LUFL	ULLL	OT
0	2016-07-01 00:00:00	5.827	2.009	1.599	0.462	4.203	1.340	30.531000
1	2016-07-01 00:15:00	5.760	2.076	1.492	0.426	4.264	1.401	30.459999
2	2016-07-01 00:30:00	5.760	1.942	1.492	0.391	4.234	1.310	30.038000
3	2016-07-01 00:45:00	5.760	1.942	1.492	0.426	4.234	1.310	27.013000
4	2016-07-01 01:00:00	5.693	2.076	1.492	0.426	4.142	1.371	27.787001



## Electricity Transformer Dataset (ETDataset)

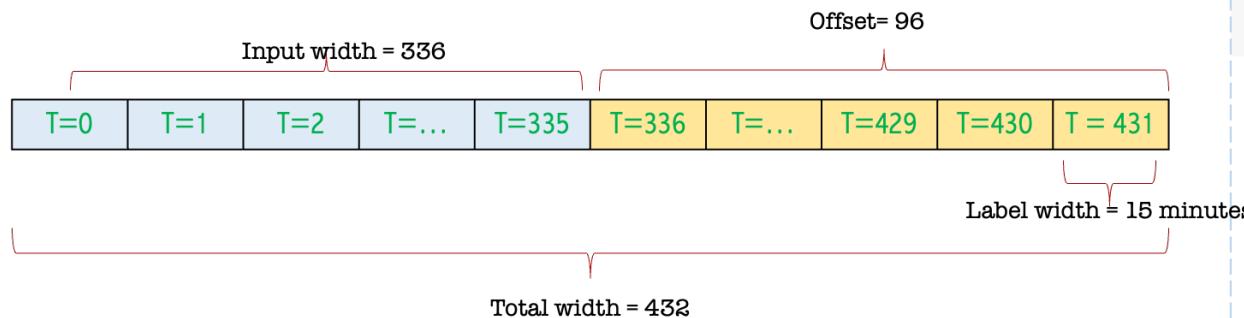
### Data Windowing Techniques

Case 1 : Multivariate to Multivariate Prediction

(Input, Output) = (336, 96)

(Input shape)= (15932, 336, 7)

(Output shape)= (15932, 96, 7)



### Multivariate-to-Multivartiate

▶ features\_type='M'  
sub\_dir = 'multi2multi'  
os.makedirs(os.path.join(weight\_dir, sub\_dir), exist\_ok=True)  
multi2multi\_loader = TimeSeriesDataLoader(file\_path,  
input\_size=input\_size,  
label\_size=label\_size,  
offset=offset,  
train\_size=train\_size,  
val\_size=val\_size,  
target\_name=target\_name,  
features\_type=features\_type,  
date\_column=date\_column)

⌚ Offset will be change from 1 to 96  
self.X\_train.shape = (11762, 336, 7)  
self.y\_train.shape = (11762, 96, 7)  
self.X\_val.shape = (1309, 336, 7)  
self.y\_val.shape = (1309, 96, 7)  
self.X\_test.shape = (2861, 336, 7)  
self.y\_test.shape = (2861, 96, 7)

## Electricity Transformer Dataset (ETDataset)

### Data Windowing Techniques

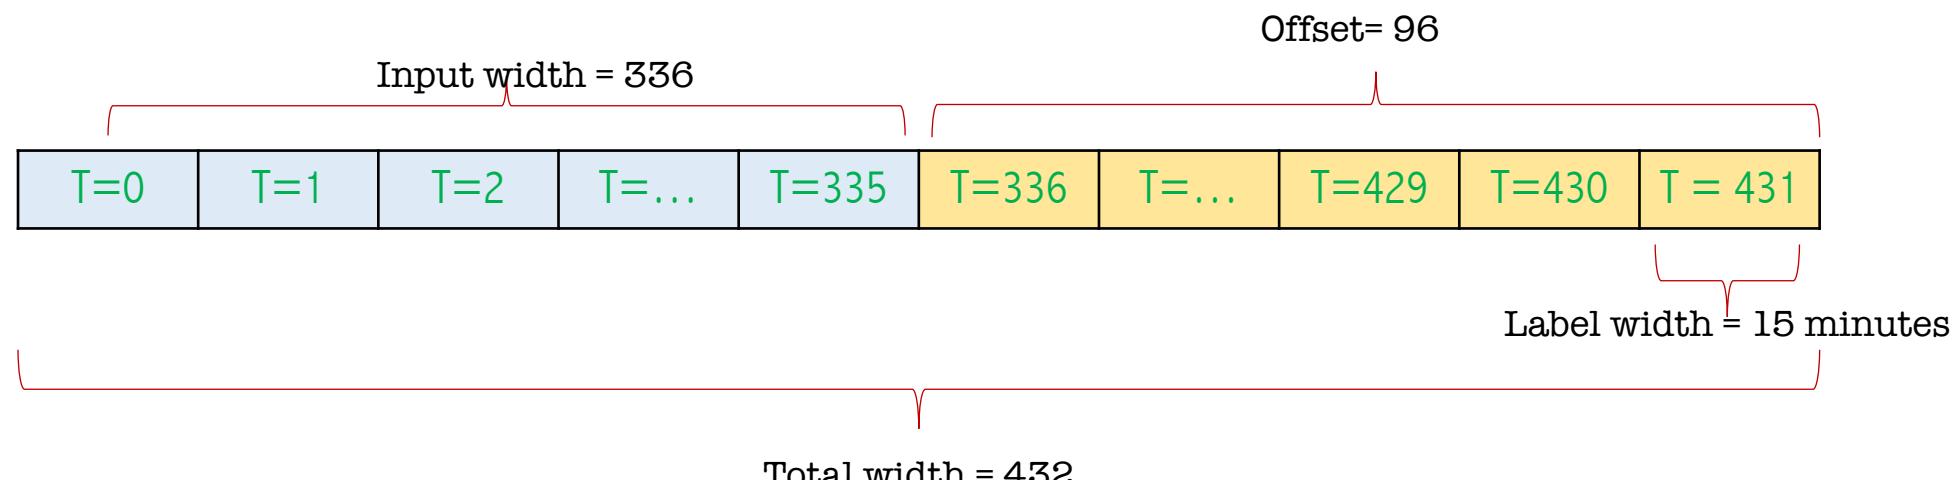
Case 2: Multivariate to Univariate Prediction

(Input, Output) = (336, 96)

(Input shape)= (15932, 336, 7)

(Output shape)= (15932, 96, 1)

	date	HUFL	HULL	MUFL	MULL	LUFL	ULLL	OT
0	2016-07-01 00:00:00	5.827	2.009	1.599	0.462	4.203	1.340	30.531000
1	2016-07-01 00:15:00	5.760	2.076	1.492	0.426	4.264	1.401	30.459999
2	2016-07-01 00:30:00	5.760	1.942	1.492	0.391	4.234	1.310	30.038000
3	2016-07-01 00:45:00	5.760	1.942	1.492	0.426	4.234	1.310	27.013000
4	2016-07-01 01:00:00	5.693	2.076	1.492	0.426	4.142	1.371	27.787001



## Electricity Transformer Dataset (ETDataset)

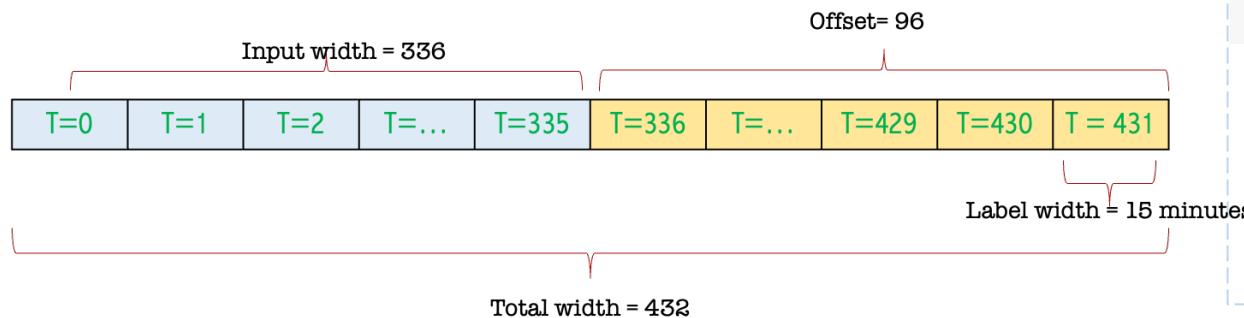
### Data Windowing Techniques

#### Case 2: Multivariate to Univariate Prediction

(Input, Output) = (336, 96)

(Input shape)= (15932, 336, 7)

(Output shape)= (15932, 96, 1)



#### ▼ Multivariate-to-Univariante

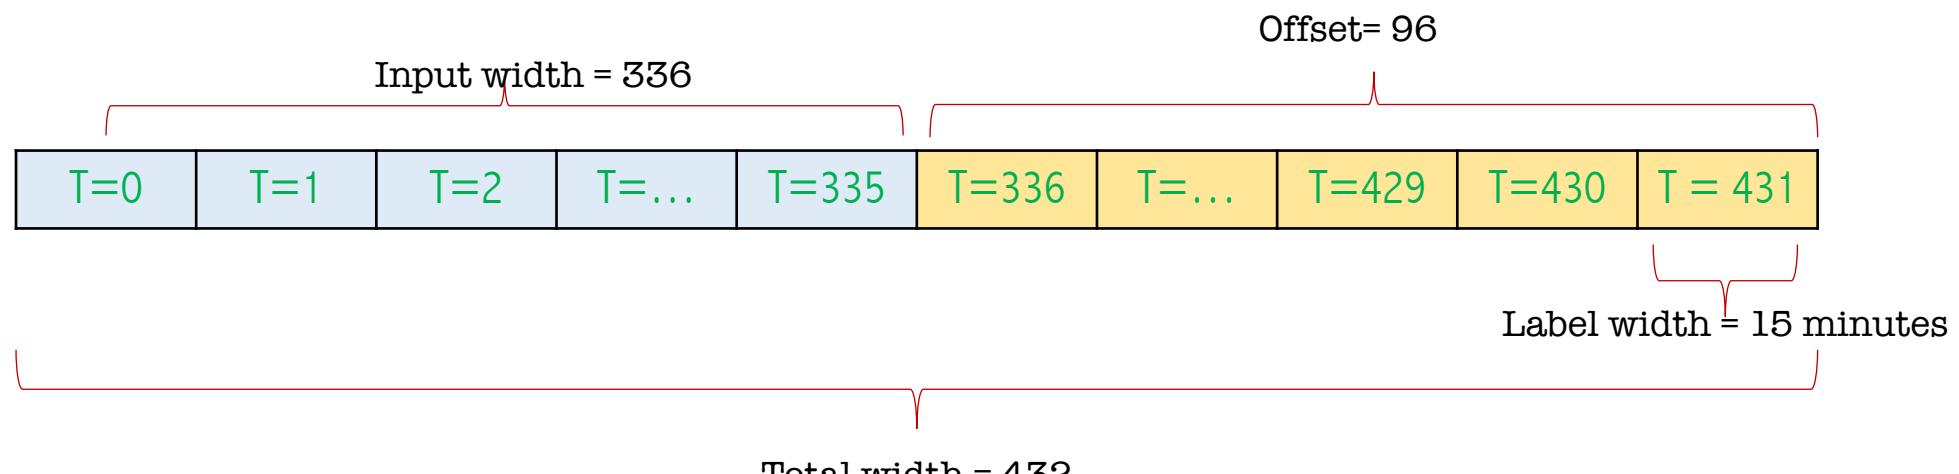
```
▶ features_type='MS'  
    sub_dir = 'multi2uni'  
    os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
    multi2uni_loader = TimeSeriesDataLoader(file_path,  
                                              input_size=input_size,  
                                              label_size=label_size,  
                                              offset=offset,  
                                              train_size=train_size,  
                                              val_size=val_size,  
                                              target_name=target_name,  
                                              features_type=features_type,  
                                              date_column=date_column)
```

👤 Offset will be change from 1 to 96  
self.X\_train.shape = (11762, 336, 7)  
self.y\_train.shape = (11762, 96, 1)  
self.X\_val.shape = (1309, 336, 7)  
self.y\_val.shape = (1309, 96, 1)  
self.X\_test.shape = (2861, 336, 7)  
self.y\_test.shape = (2861, 96, 1)

## Electricity Transformer Dataset (ETDataset)

Data Windowing Techniques	
Case 3: Univariate to Univariate Prediction	
(Input, Output) = (336, 96)	
(Input shape)= (15932, 336, 1)	
(Output shape)= (15932, 96, 1)	

	date	HUFL	HULL	MUFL	MULL	LUFL	ULLL	OT
0	2016-07-01 00:00:00	5.827	2.009	1.599	0.462	4.203	1.340	30.531000
1	2016-07-01 00:15:00	5.760	2.076	1.492	0.426	4.264	1.401	30.459999
2	2016-07-01 00:30:00	5.760	1.942	1.492	0.391	4.234	1.310	30.038000
3	2016-07-01 00:45:00	5.760	1.942	1.492	0.426	4.234	1.310	27.013000
4	2016-07-01 01:00:00	5.693	2.076	1.492	0.426	4.142	1.371	27.787001



## Electricity Transformer Dataset (ETDataset)

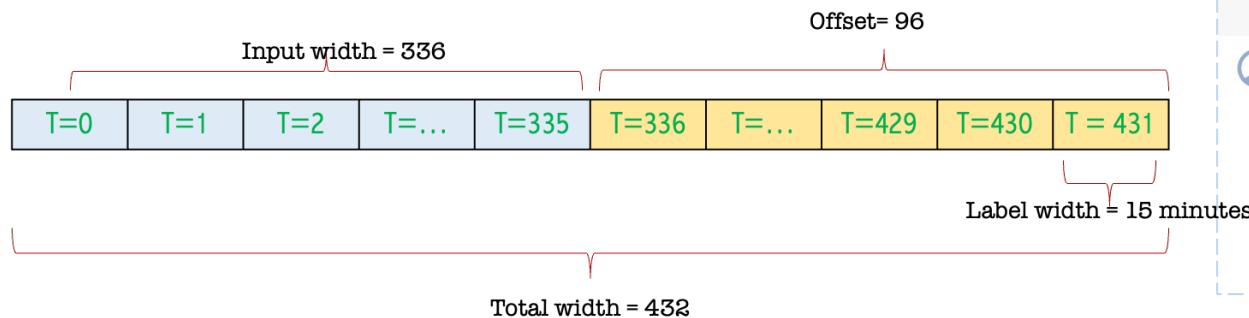
### Data Windowing Techniques

Case 3: **Univariate to Univariate** Prediction

(Input, Output) = (336, 96)

(Input shape)= (15932, 336, 1)

(Output shape)= (15932, 96, 1)



### Univariate-to-Univariate

```
▶ features_type='S'  
sub_dir = 'uni2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
uni2uni_loader = TimeSeriesDataLoader(file_path,  
input_size=input_size,  
label_size=label_size,  
offset=offset,  
train_size=train_size,  
val_size=val_size,  
target_name=target_name,  
features_type=features_type,  
date_column=date_column)
```

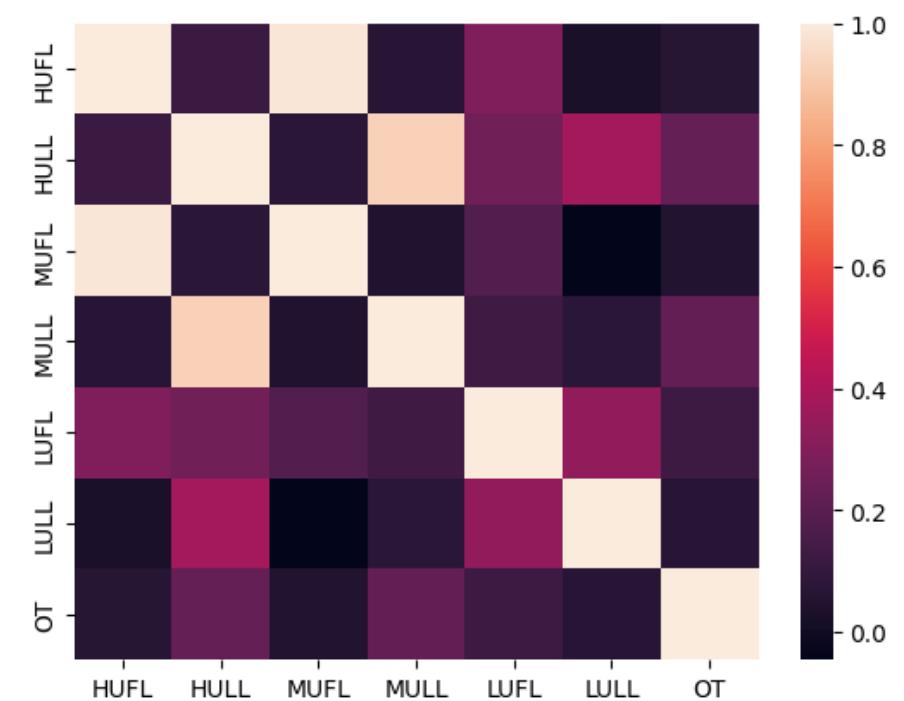
Offset will be change from 1 to 96  
self.X\_train.shape = (11762, 336, 1)  
self.y\_train.shape = (11762, 96, 1)  
self.X\_val.shape = (1309, 336, 1)  
self.y\_val.shape = (1309, 96, 1)  
self.X\_test.shape = (2861, 336, 1)  
self.y\_test.shape = (2861, 96, 1)

# Time-series Data Project: Data Exploration

```
[ ] train_df = pl.read_csv(file_path)

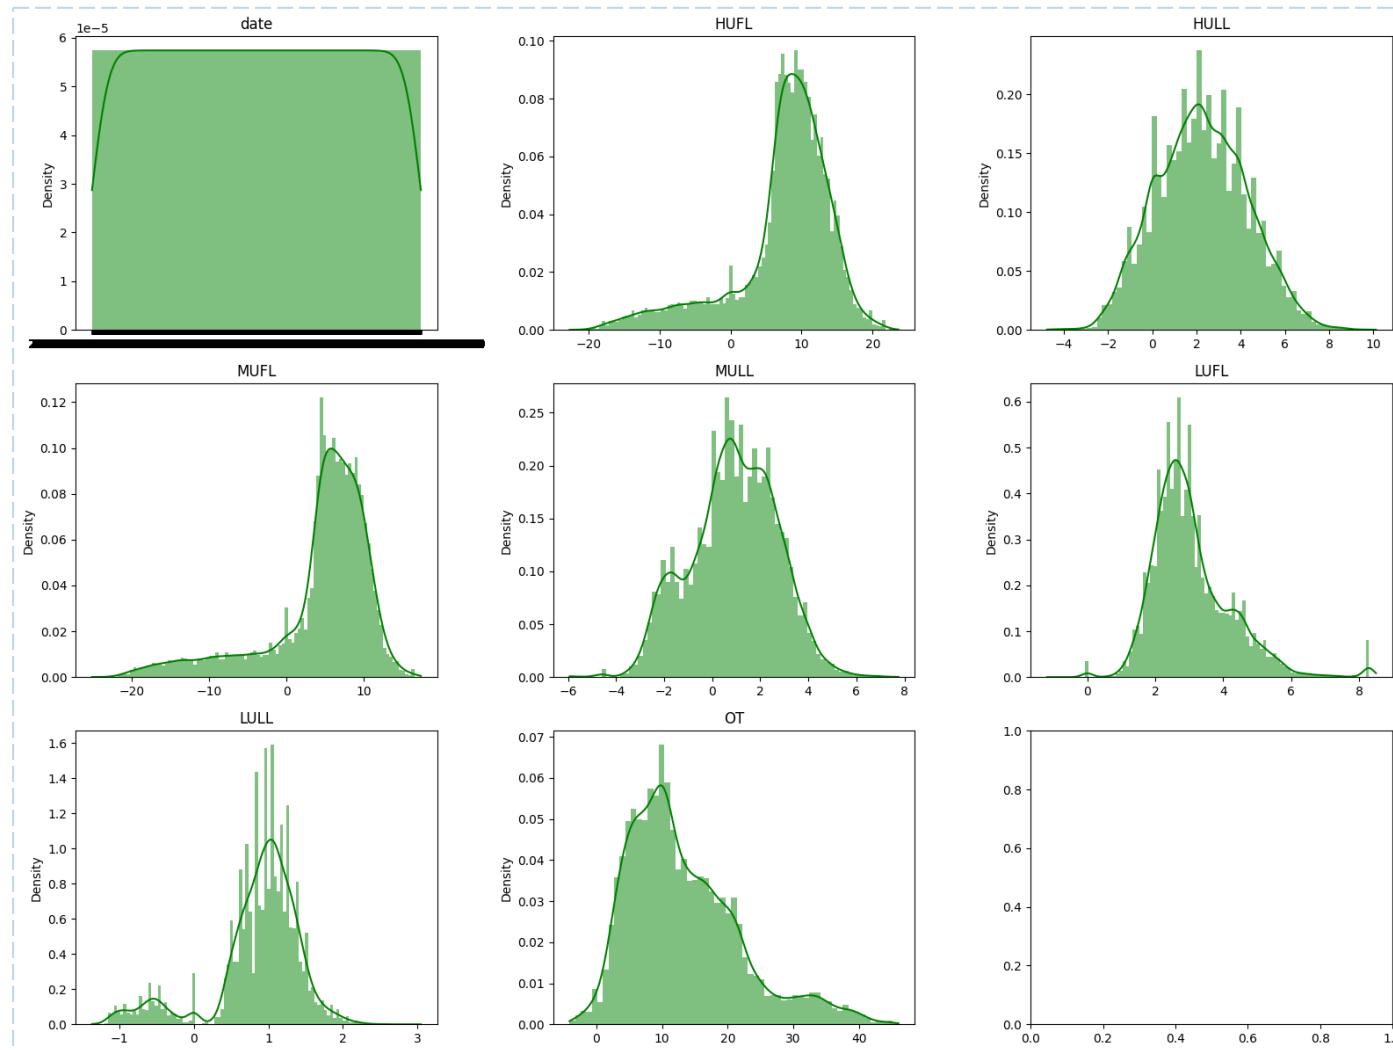
▶ train_df.describe()

shape: (7, 9)
describe      date      HUFL      HULL      MUFL      MULL      LUFL      LULL      OT
str           str       f64       f64       f64       f64       f64       f64       f64
"count"      "17420"   17420.0  17420.0  17420.0  17420.0  17420.0  17420.0  17420.0
"null_count" "0"        0.0       0.0       0.0       0.0       0.0       0.0       0.0
"mean"        null     7.375141  2.242242  4.300239  0.881568  3.066062  0.856932  13.324672
"std"         null     7.067744  2.042342  6.826978  1.809293  1.164506  0.599552  8.566946
"min"        "2016-07-01 00:00:00" -22.705999 -4.756    -25.087999 -5.934    -1.188    -1.371    -4.08
"max"        "2018-06-26 19:59:59" 23.643999 10.114    17.341    7.747    8.498    3.046    46.007
"median"      null     8.774    2.21      5.97      0.959    2.833    0.975    11.396
```



```
train = train_df.to_pandas()
pearson_corr = train.corr(method="pearson")
sns.heatmap(pearson_corr)
```

# Time-series Data Project: Data Exploration



```
columns_to_plot = train_df.columns

num_columns = 3
num_rows = int(np.ceil(len(columns_to_plot) / num_columns))
grid_layout = (num_rows, num_columns)

fig, axes = plt.subplots(*grid_layout, figsize=(16, 12))

axes = axes.flatten()

for i, column in enumerate(columns_to_plot):
    sns.histplot(train_df[column], kde=True, ax=axes[i],
                 color='green', stat="density", linewidth=0)
    axes[i].set_title(column)

plt.tight_layout()
plt.show()
```

## Define MLP Structure in Pytorch

```
▶ class MLP(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, ahead):
        super(MLP, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, output_size * ahead) # Adjust for output sequence length
        self.ahead = ahead
        self.output_size = output_size

    def forward(self, x):
        # Flatten the input
        x = x.view(x.size(0), -1) # Reshape from [batch, lag, features] to [batch, lag * features]
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        return x.view(-1, self.ahead, self.output_size) # Reshape to [batch, ahead, features]
```

```
# Determine in_variable and out_variable based on features_type
if features_type == 'S':
    self.in_variable = 1
    self.out_variable = 1
elif features_type == 'M':
    self.in_variable = len(self.df.columns)
    self.out_variable = len(self.df.columns)
elif features_type == 'MS':
    self.in_variable = len(self.df.columns)
    self.out_variable = 1
else:
    raise ValueError("Invalid features_type. Choose from 'S' for Univariate-to-Univariate,
```

## Configuration

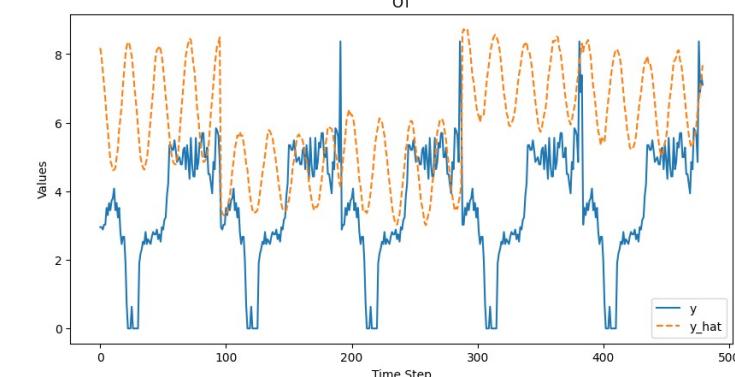
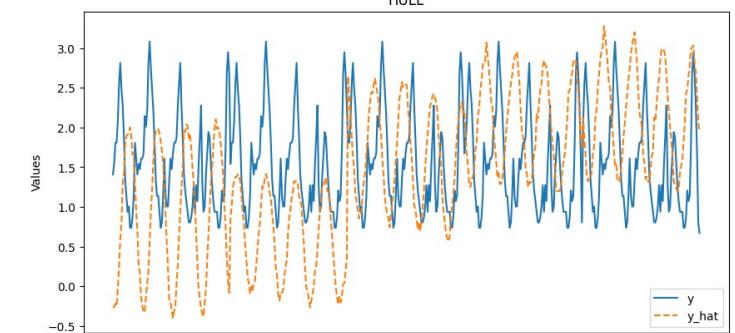
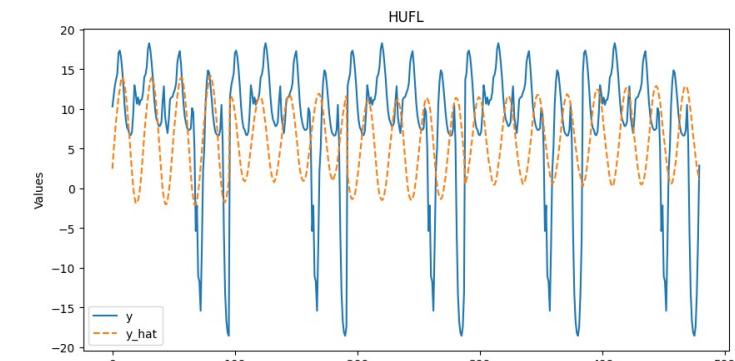
```
▶ input_size = 336
label_size = 96
offset = 1
train_size = 0.7
val_size = 0.1
num_epochs = 1_000_000
patience = 20
learning_rate = 0.001
hidden_size = 64
num_layers = 2
ele = 5
target_name = 'OT'
date_column = 'date'
file_path = 'ETTh1.csv'
plot_dir = 'plots'
weight_dir = 'weights'
results = []
```

## Define MLP Structure in Pytorch

```
[ ] MLP_multi2multi = MLP(input_size=multi2multi_loader.in_variable*input_size,
                           hidden_size=hidden_size,
                           output_size=multi2multi_loader.out_variable,
                           ahead=label_size)
MLP_multi2multi_manager = ModelManager(model=MLP_multi2multi,
                                         train_loader=multi2multi_loader.train_loader,
                                         val_loader=multi2multi_loader.val_loader,
                                         lr=learning_rate,
                                         patience=patience)
MLP_multi2multi_manager.train(num_epochs=num_epochs,
                               save_dir=os.path.join(weight_dir, sub_dir))
results.append({
    "Name": MLP_multi2multi_manager.model.__class__.__name__,
    "Type": sub_dir,
    "MAE": MLP_multi2multi_manager.evaluate(loader=multi2multi_loader.test_loader)
})
results[-1]

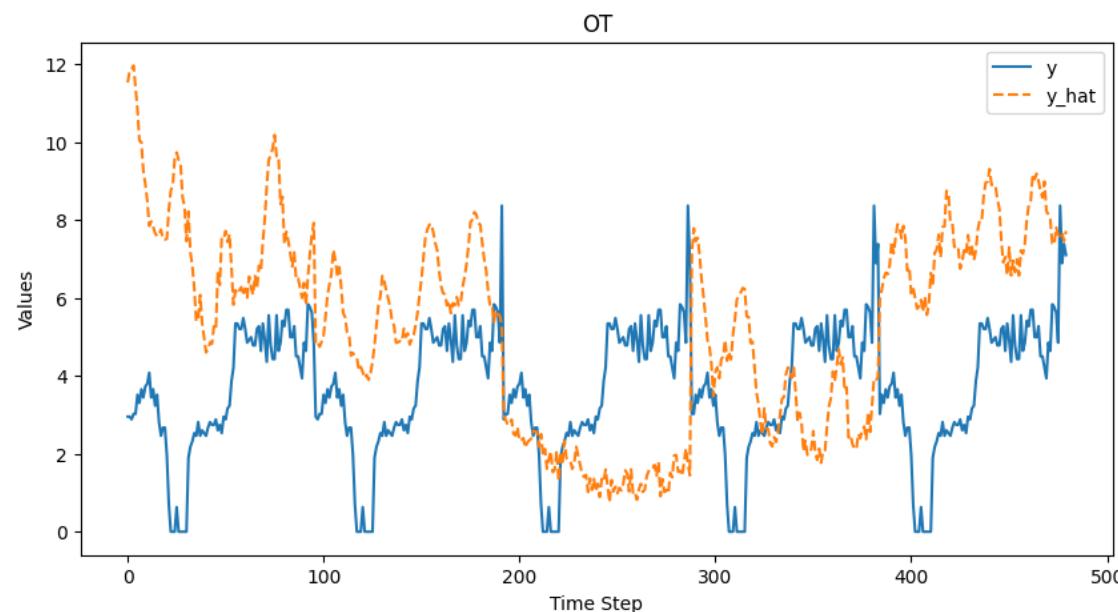
Model saved to weights\multi2multi\best-MLP.pth
Epoch [1/1000000], time: 0s, loss: 3.0666, val_loss: 3.6309
Model saved to weights\multi2multi\best-MLP.pth
Epoch [2/1000000], time: 0s, loss: 2.3158, val_loss: 3.4847
```

## Case 1: Multivariate to Multivariate Prediction



## Define MLP Structure in Pytorch

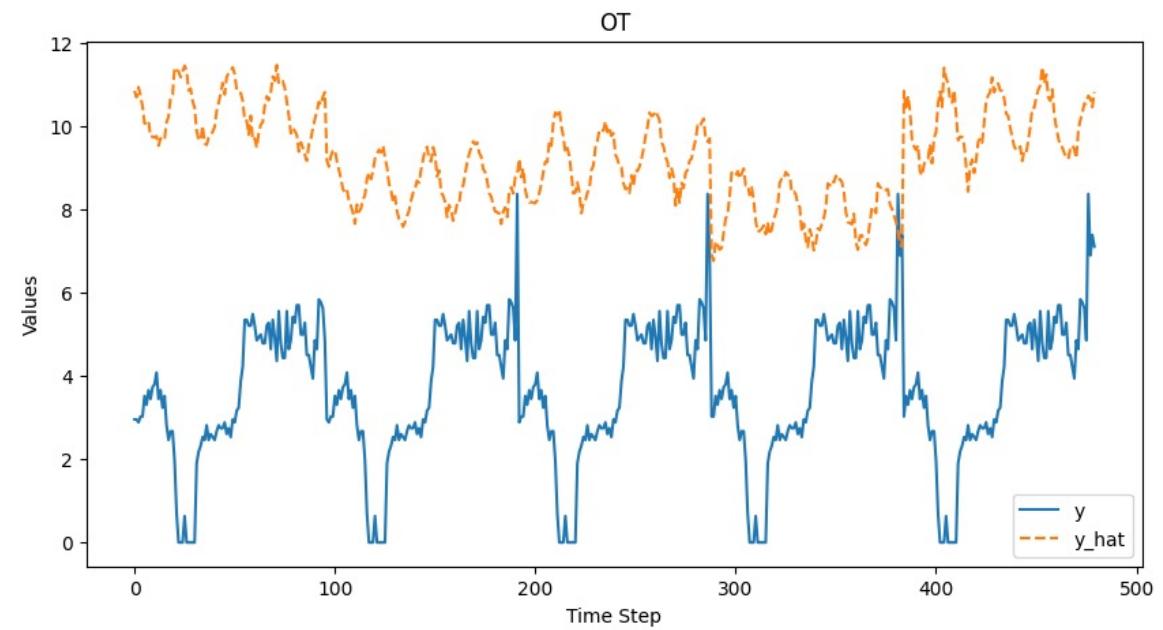
```
▶ features_type='MS'  
sub_dir = 'multi2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
multi2uni_loader = TimeSeriesDataLoader(file_path,  
                                         input_size=input_size,  
                                         label_size=label_size,  
                                         offset=offset,  
                                         train_size=train_size,  
                                         val_size=val_size,  
                                         target_name=target_name,  
                                         features_type=features_type,  
                                         date_column=date_column)  
  
👤 Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 7)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 7)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 7)  
self.y_test.shape = (2861, 96, 1)
```



## Case 2: Multivariate to Univariate Prediction

## Define MLP Structure in Pytorch

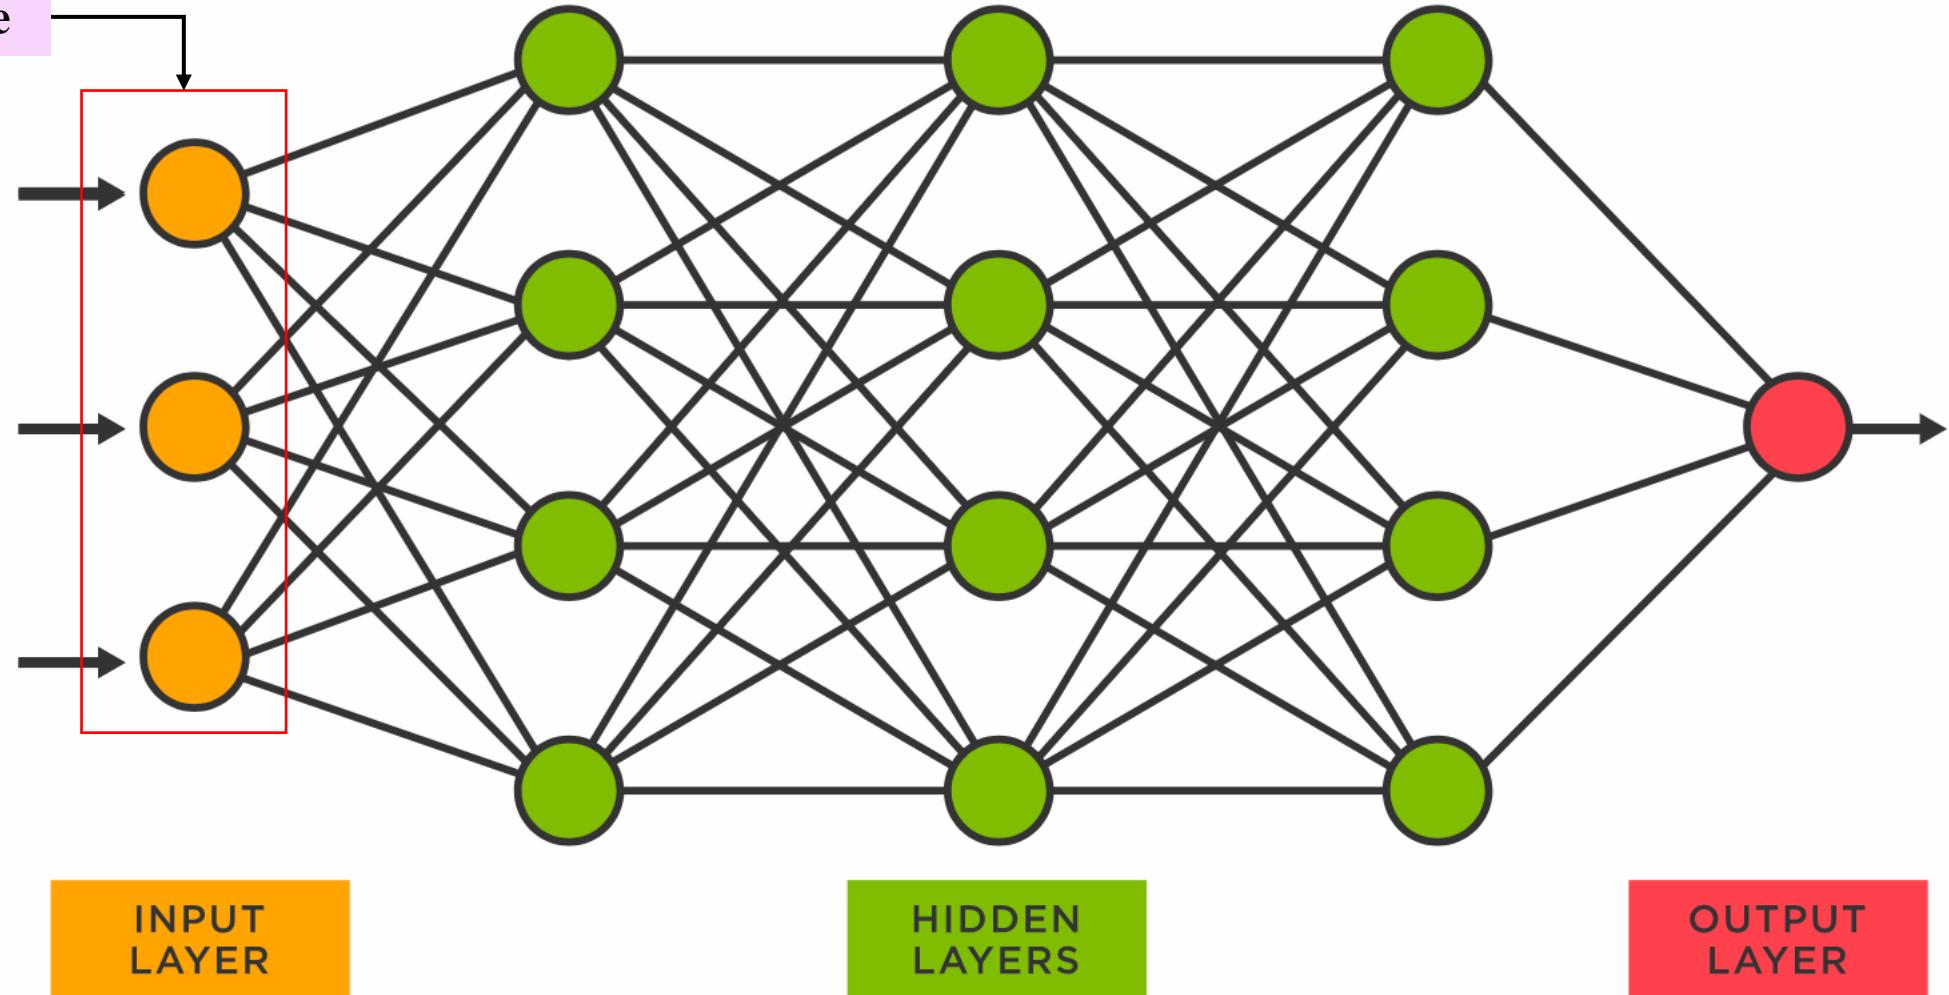
```
▶ features_type='S'  
sub_dir = 'uni2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
uni2uni_loader = TimeSeriesDataLoader(file_path,  
                                       input_size=input_size,  
                                       label_size=label_size,  
                                       offset=offset,  
                                       train_size=train_size,  
                                       val_size=val_size,  
                                       target_name=target_name,  
                                       features_type=features_type,  
                                       date_column=date_column)  
  
👤 Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 1)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 1)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 1)  
self.y_test.shape = (2861, 96, 1)
```



## Case 3: Univariate to Univariate Prediction

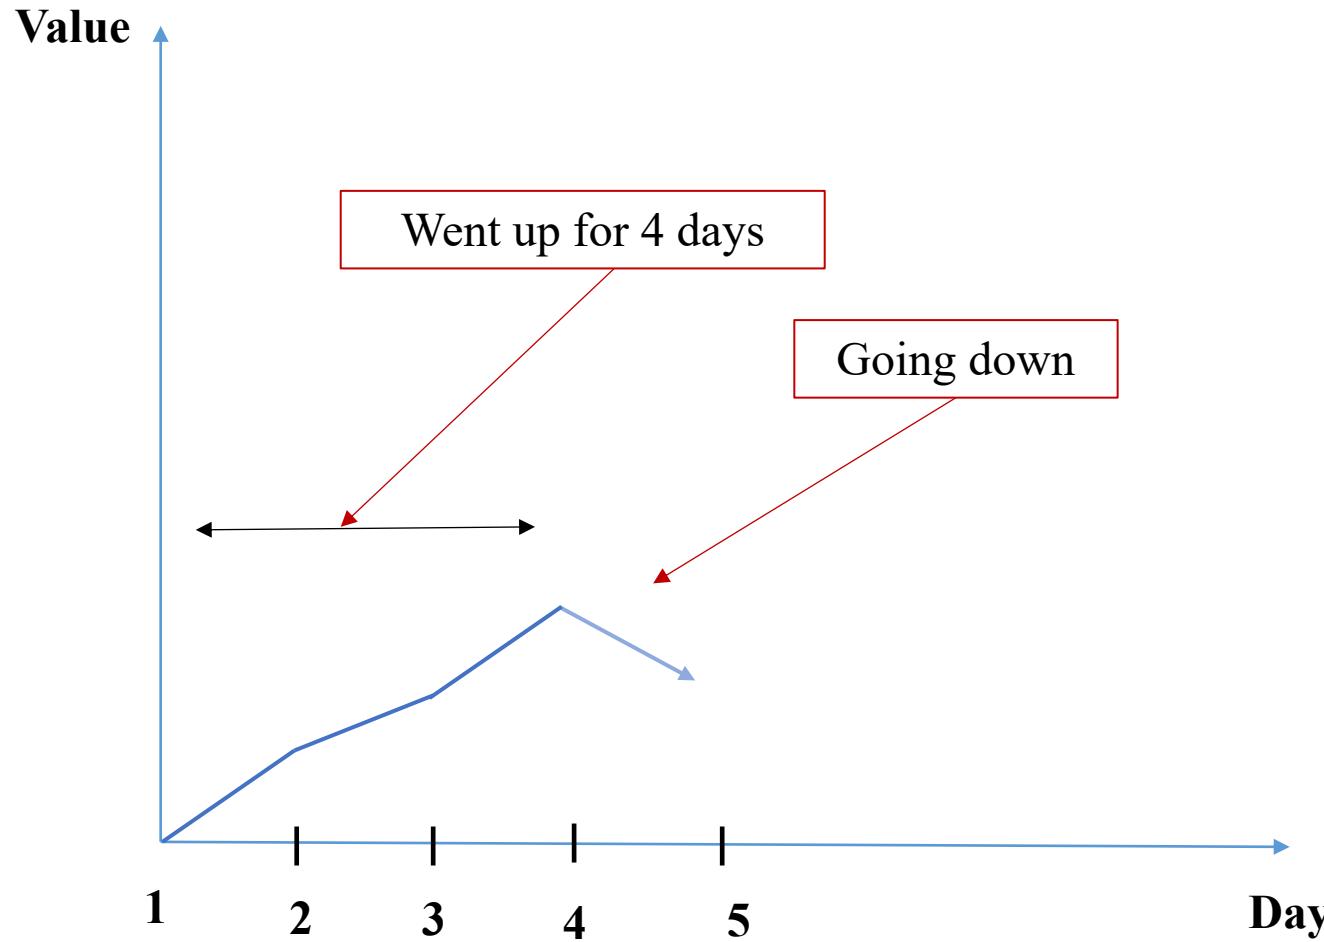
# Neural Network : Limitations

Fixed input size



# Outline

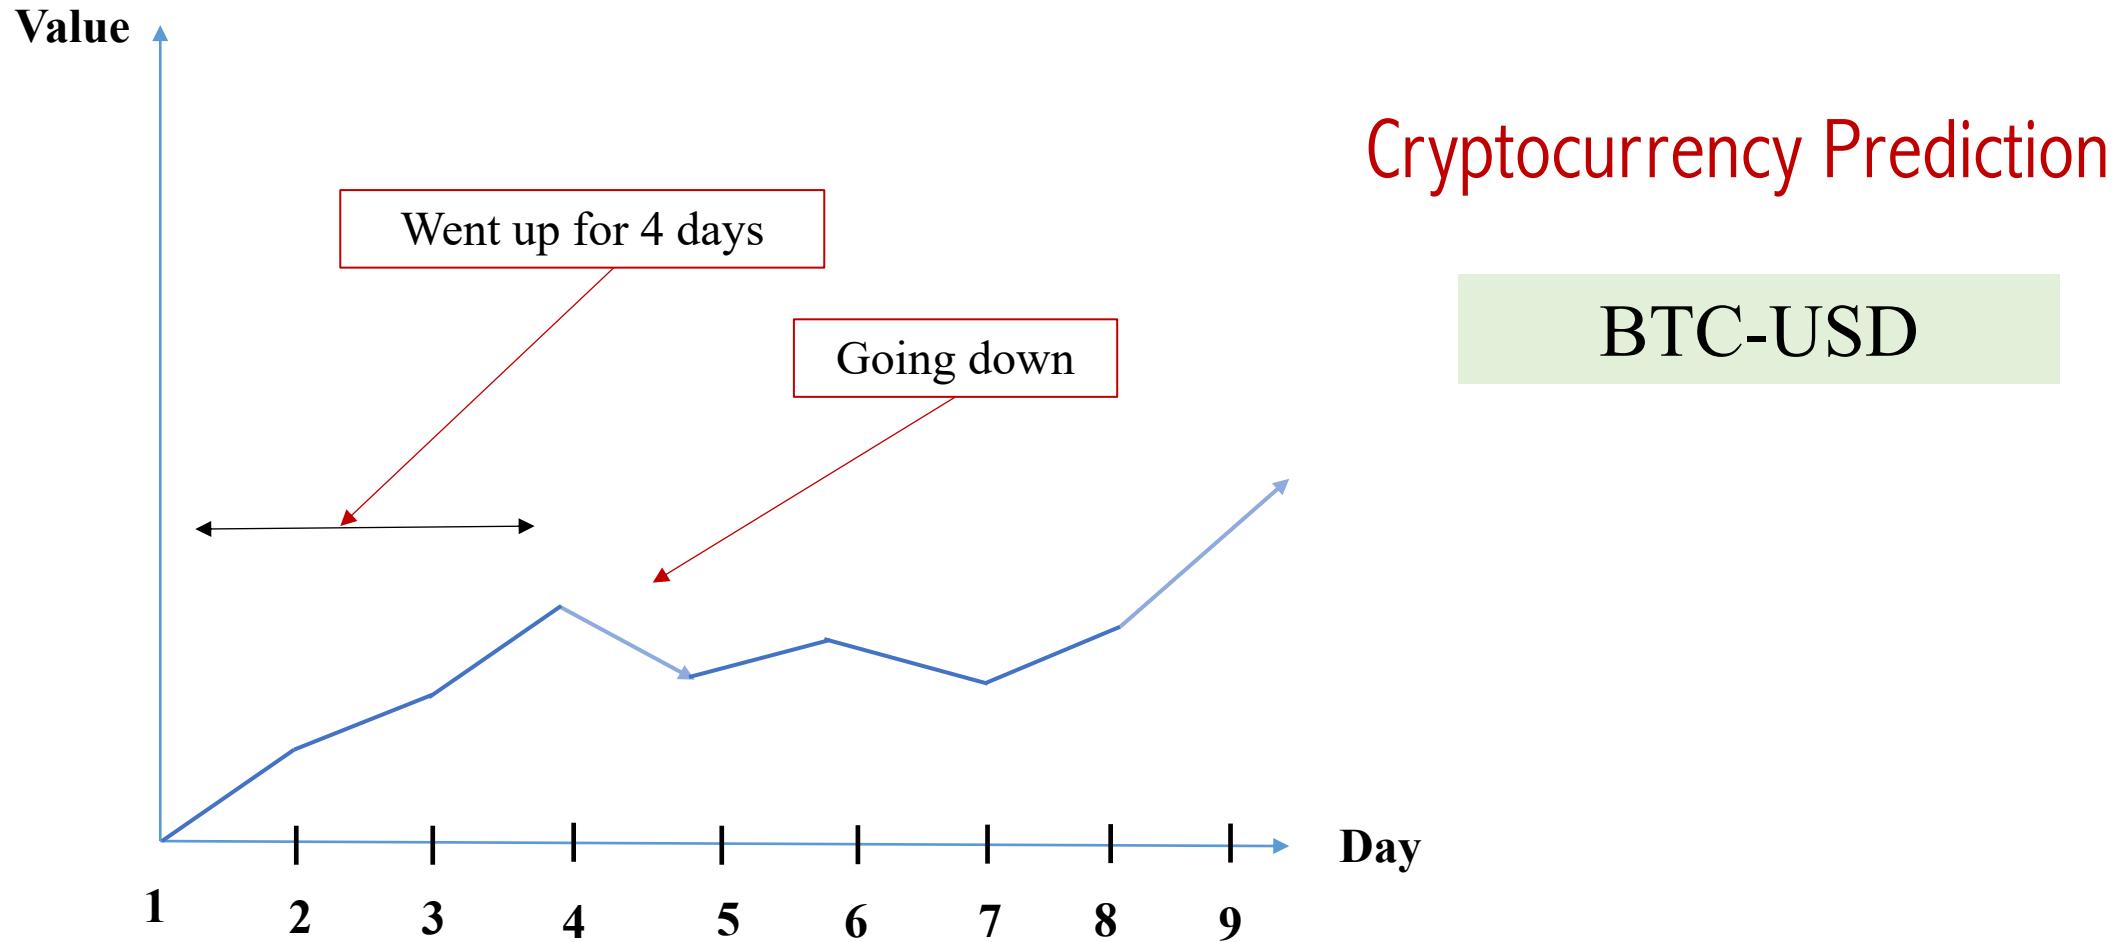
- Multilayer Perceptron and Time-series Data Forecasting
- RNN and Time-series Data Forecasting
- LSTM and Time-series Data Forecasting
- Bi-LSTM and Time-series Data Forecasting
- XGBoost and Time-series Data Forecasting



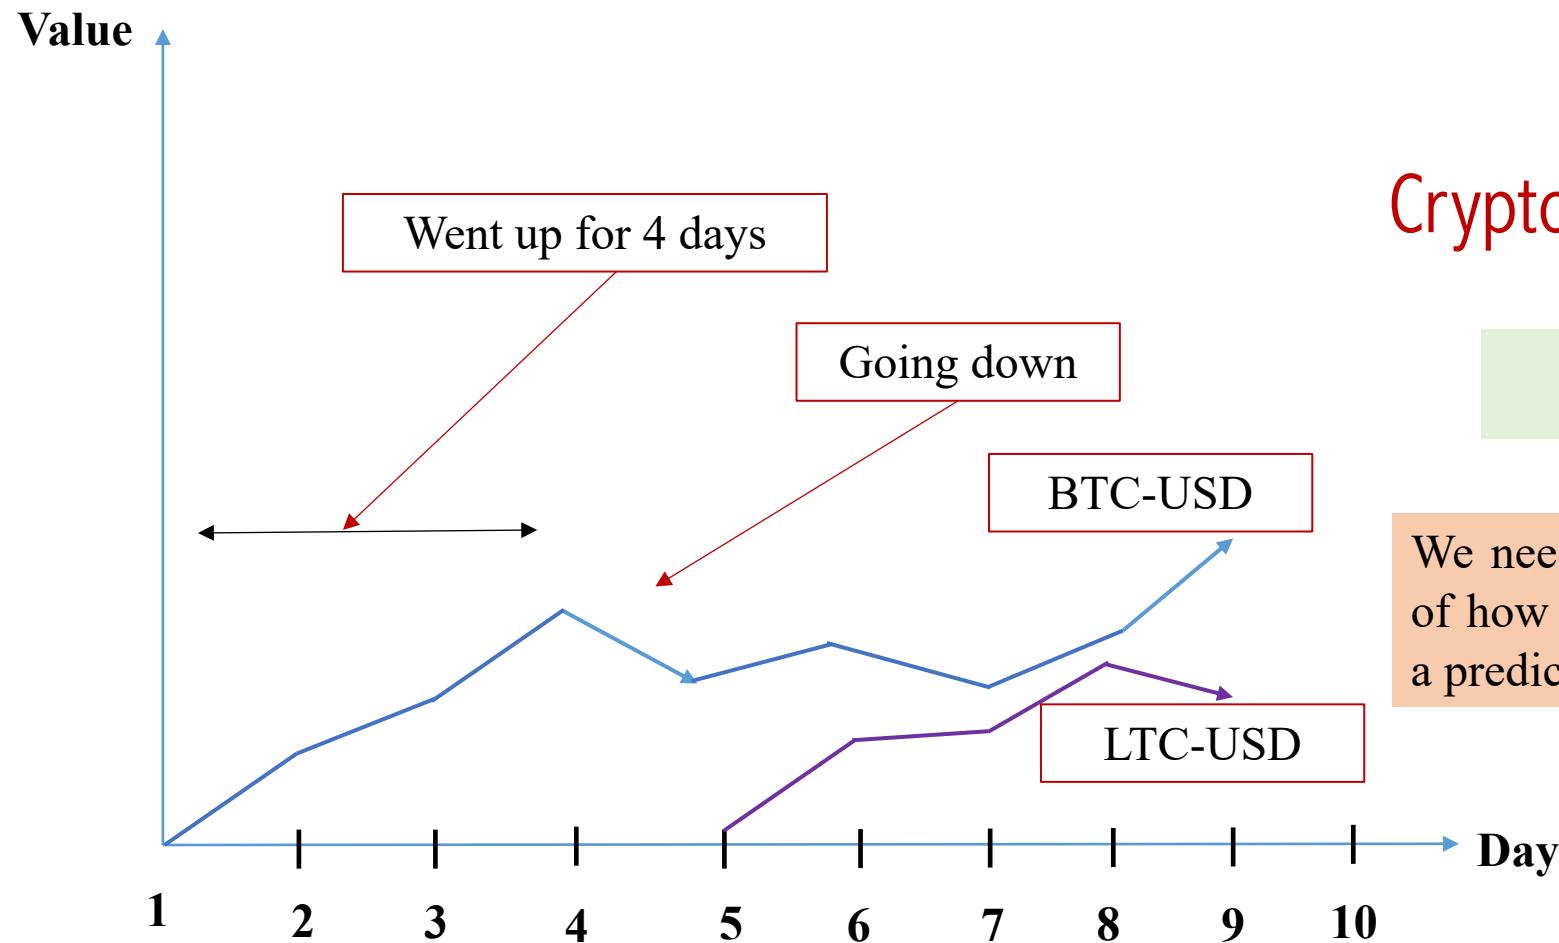
Cryptocurrency Prediction

BTC-USD

# Recurrent Neural Network : Simple Explanation



# Recurrent Neural Network : Simple Explanation



Cryptocurrency Prediction

BTC-USD

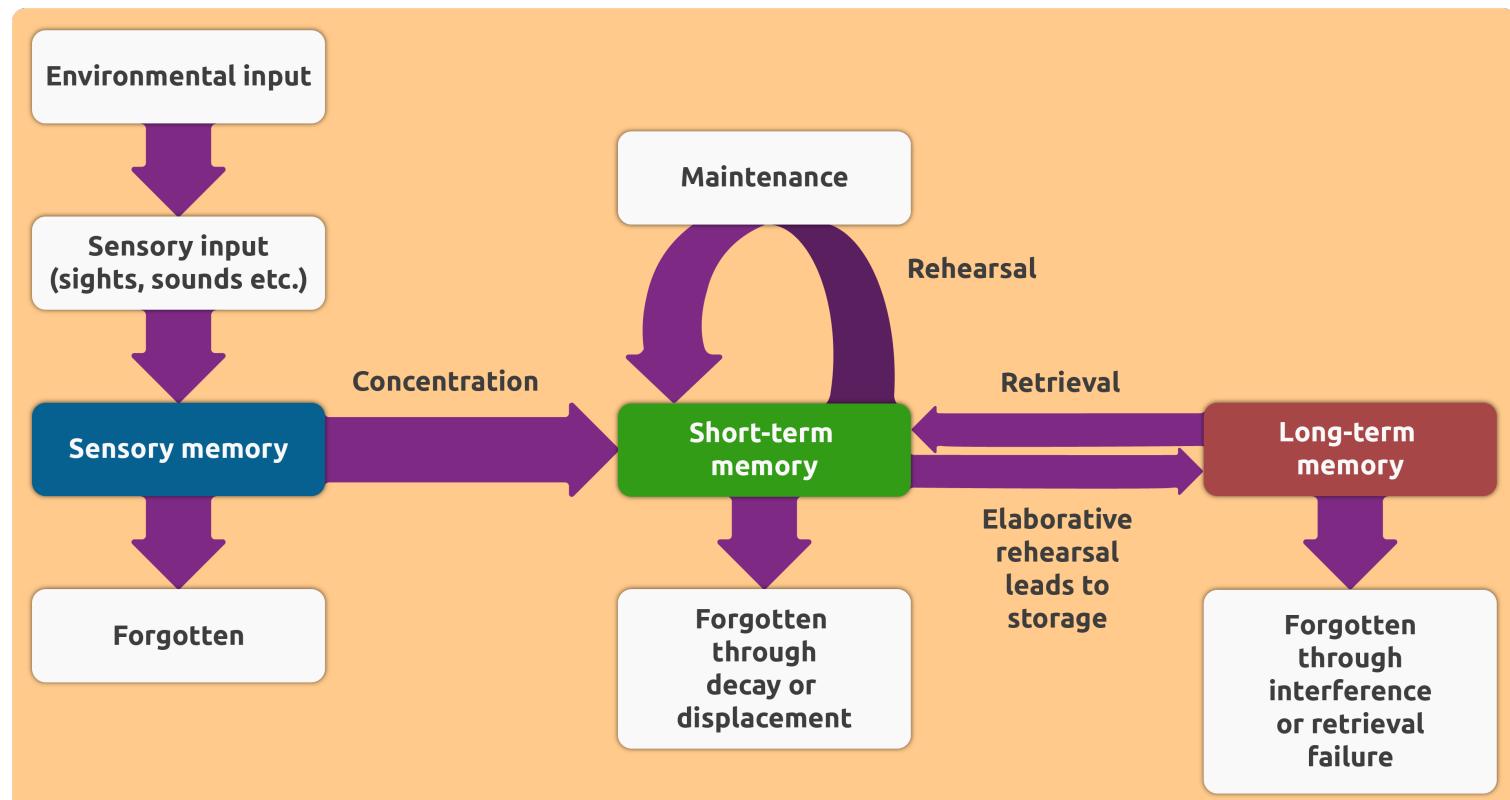
We need the NN to be more flexible in term of how much sequential data we use to make a prediction

*Predict LTC's price, we need data for 5 days (5,6,7,8,9)*

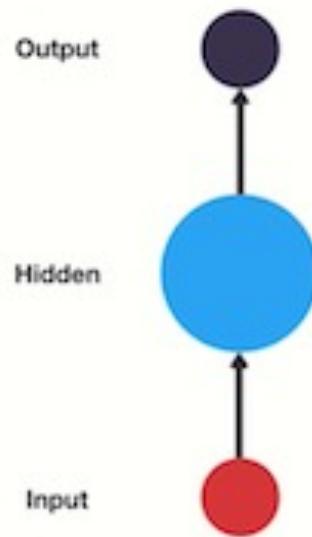
*Predict BTC's price, we need data for 9 days (1,2,3,4,5,6,7,8,9)*

# Recurrent Neural Network : Simple Explanation

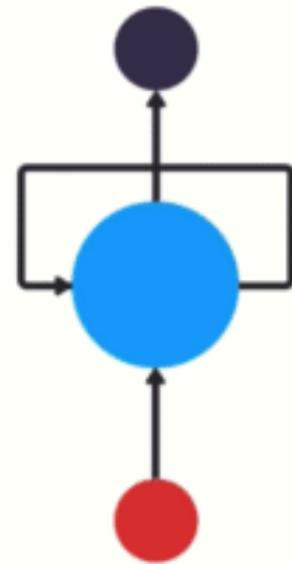
Humans tend to retrieve information from **memory**, short or long, use **current information** with it and derive logic to take next action (or impulse/habit, again based on previous experiences).



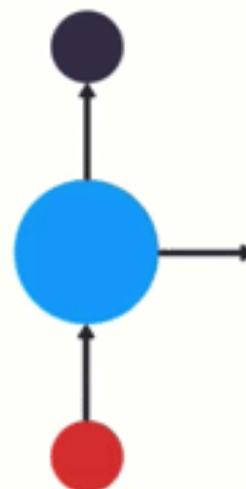
# Recurrent Neural Network : Simple Explanation



NN

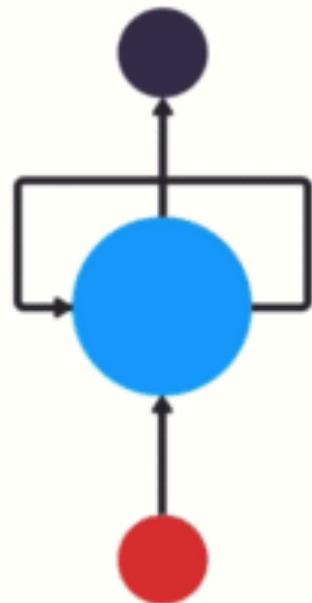


RNN



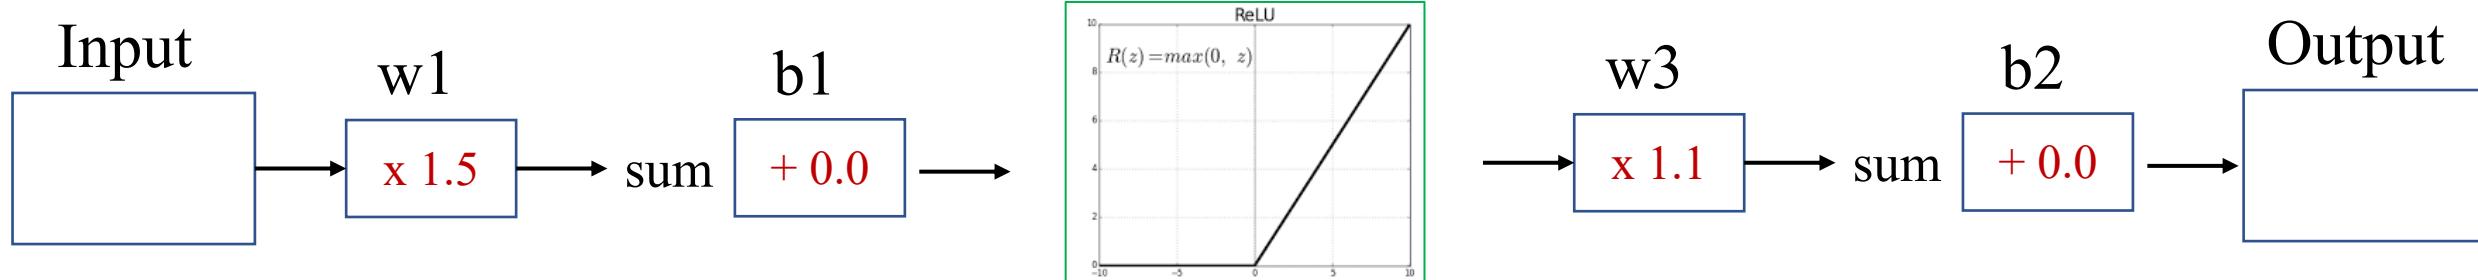
RNN

# Recurrent Neural Network : Simple Explanation

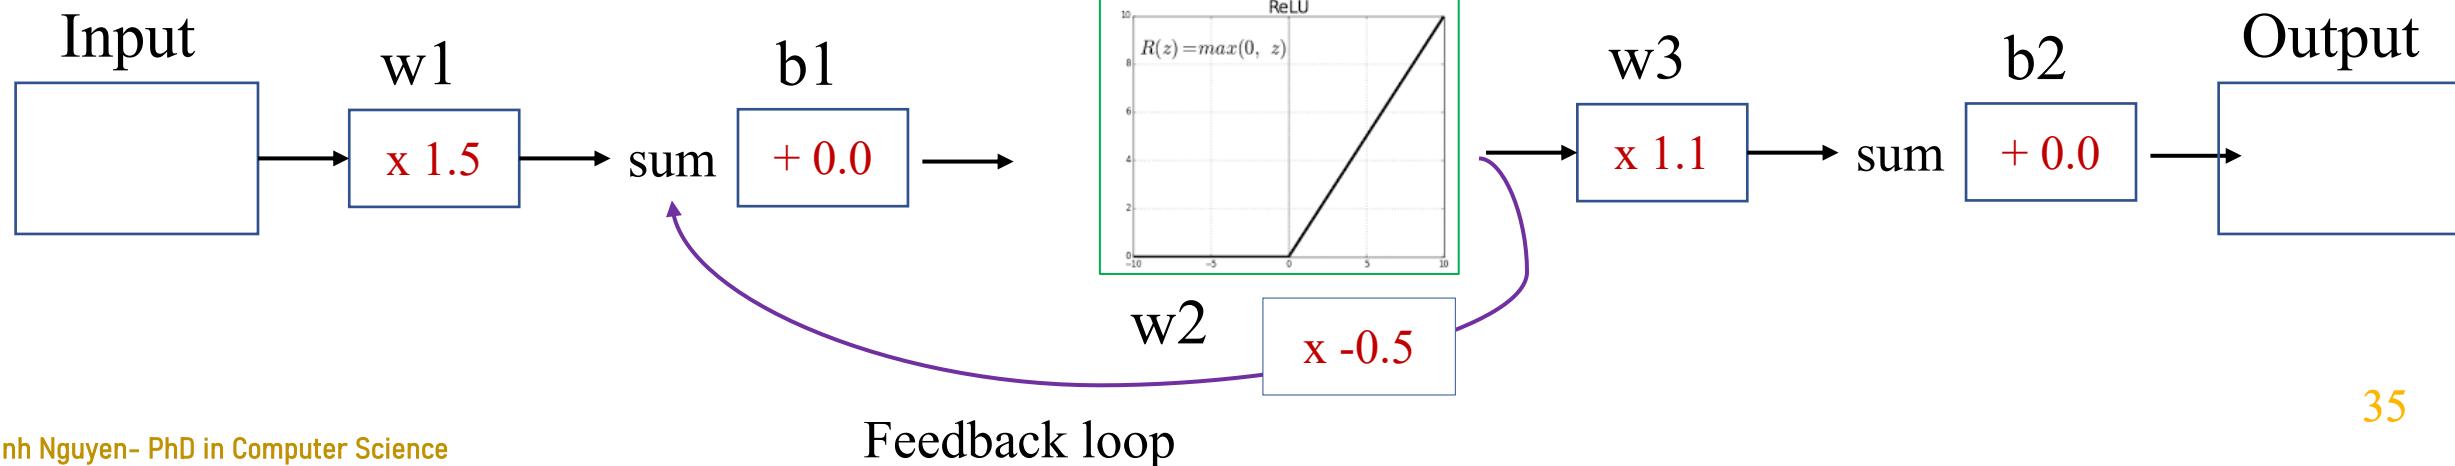


# Recurrent Neural Network : Simple Explanation

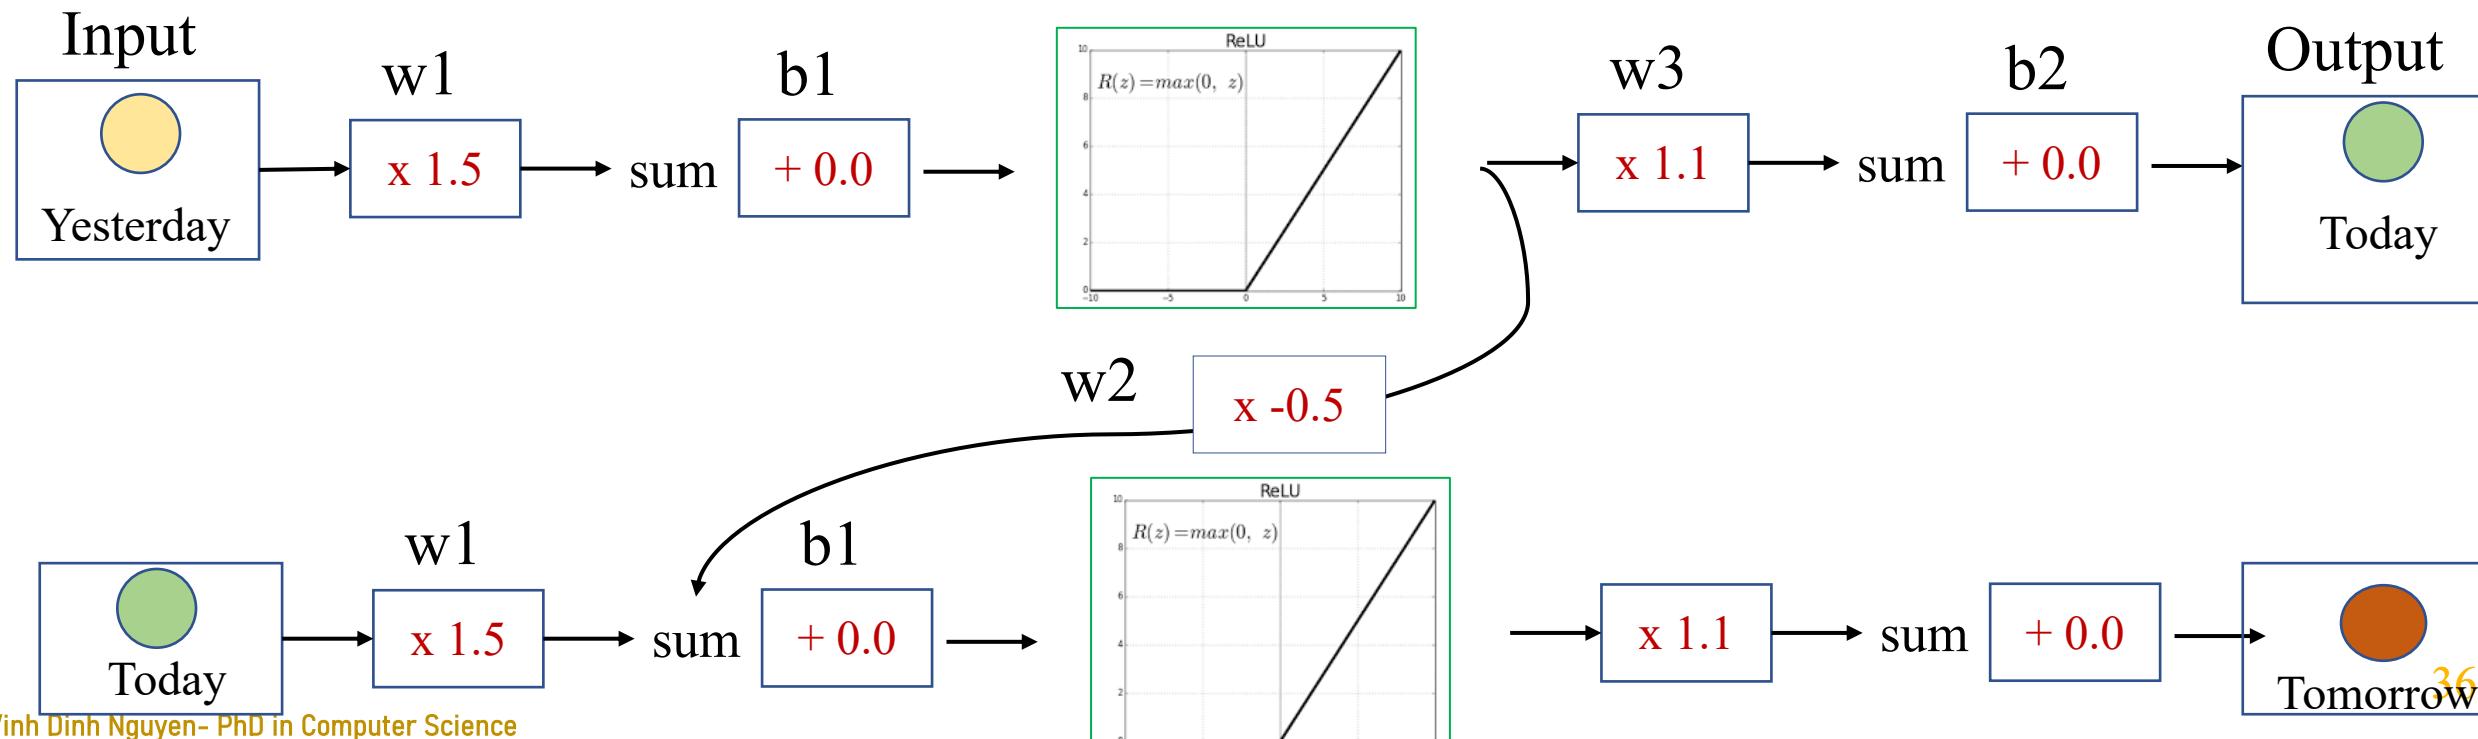
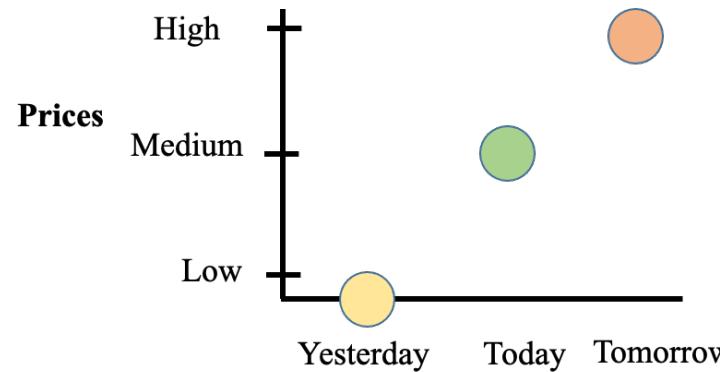
Traditional Neural Network



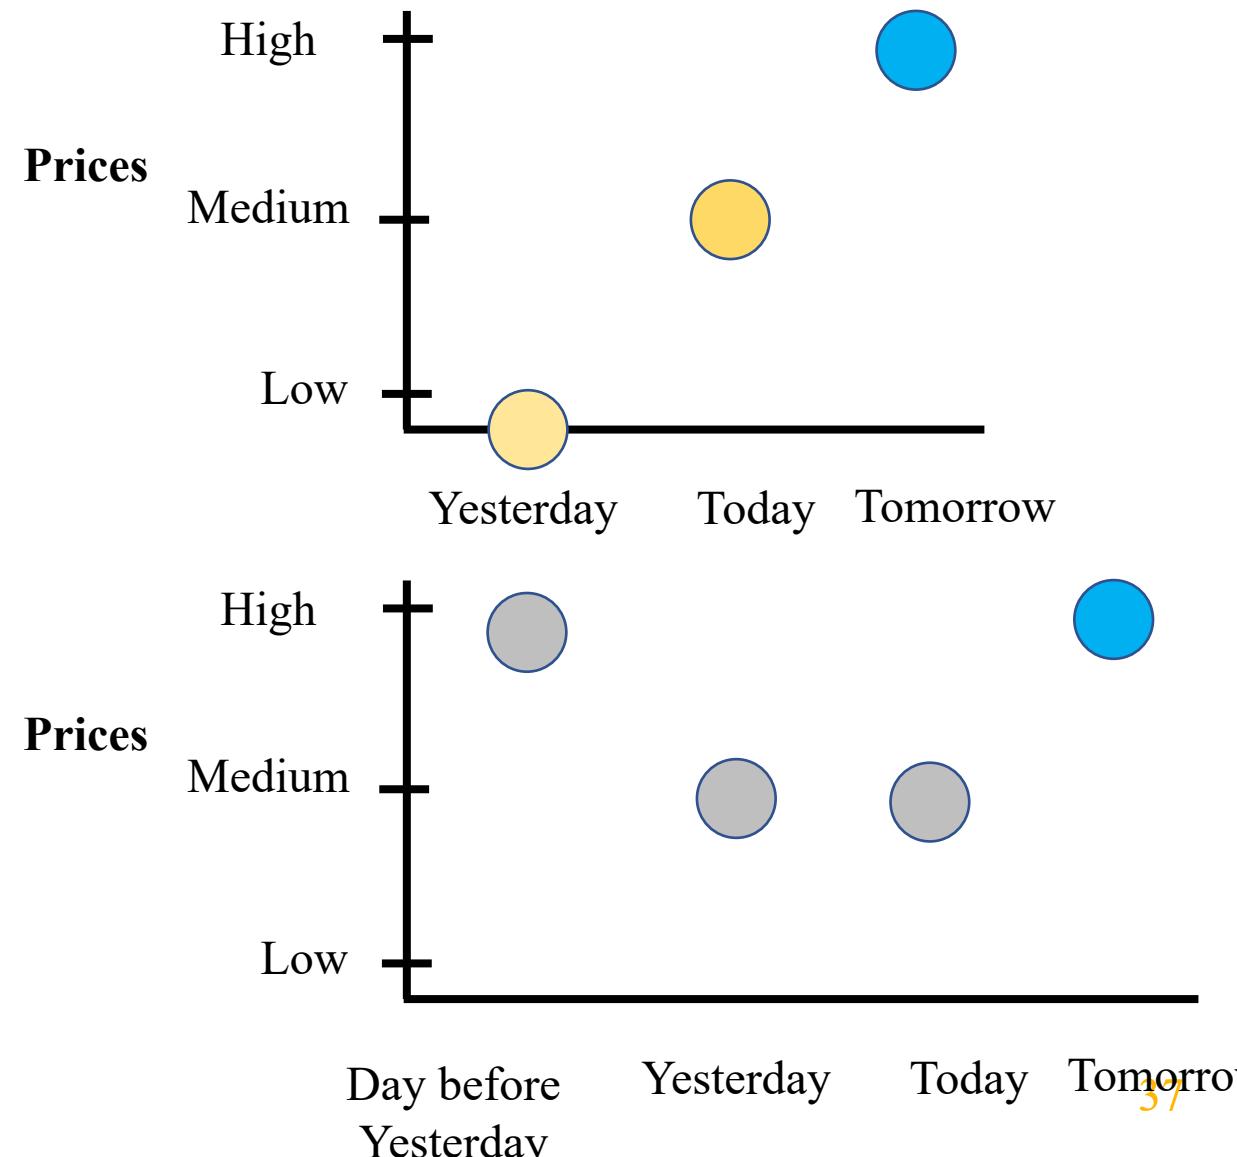
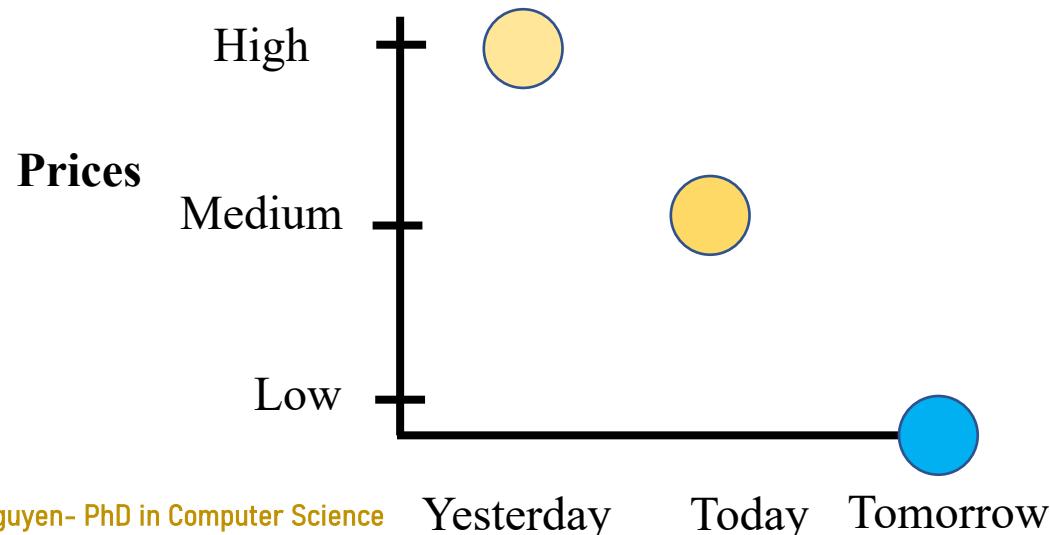
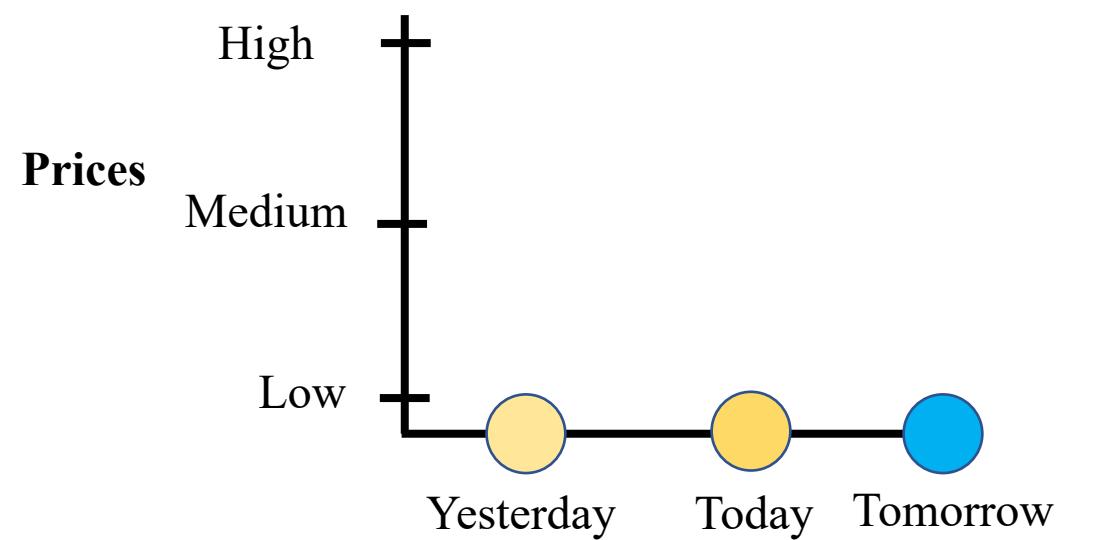
Recurrent Neural Network for sequential input values



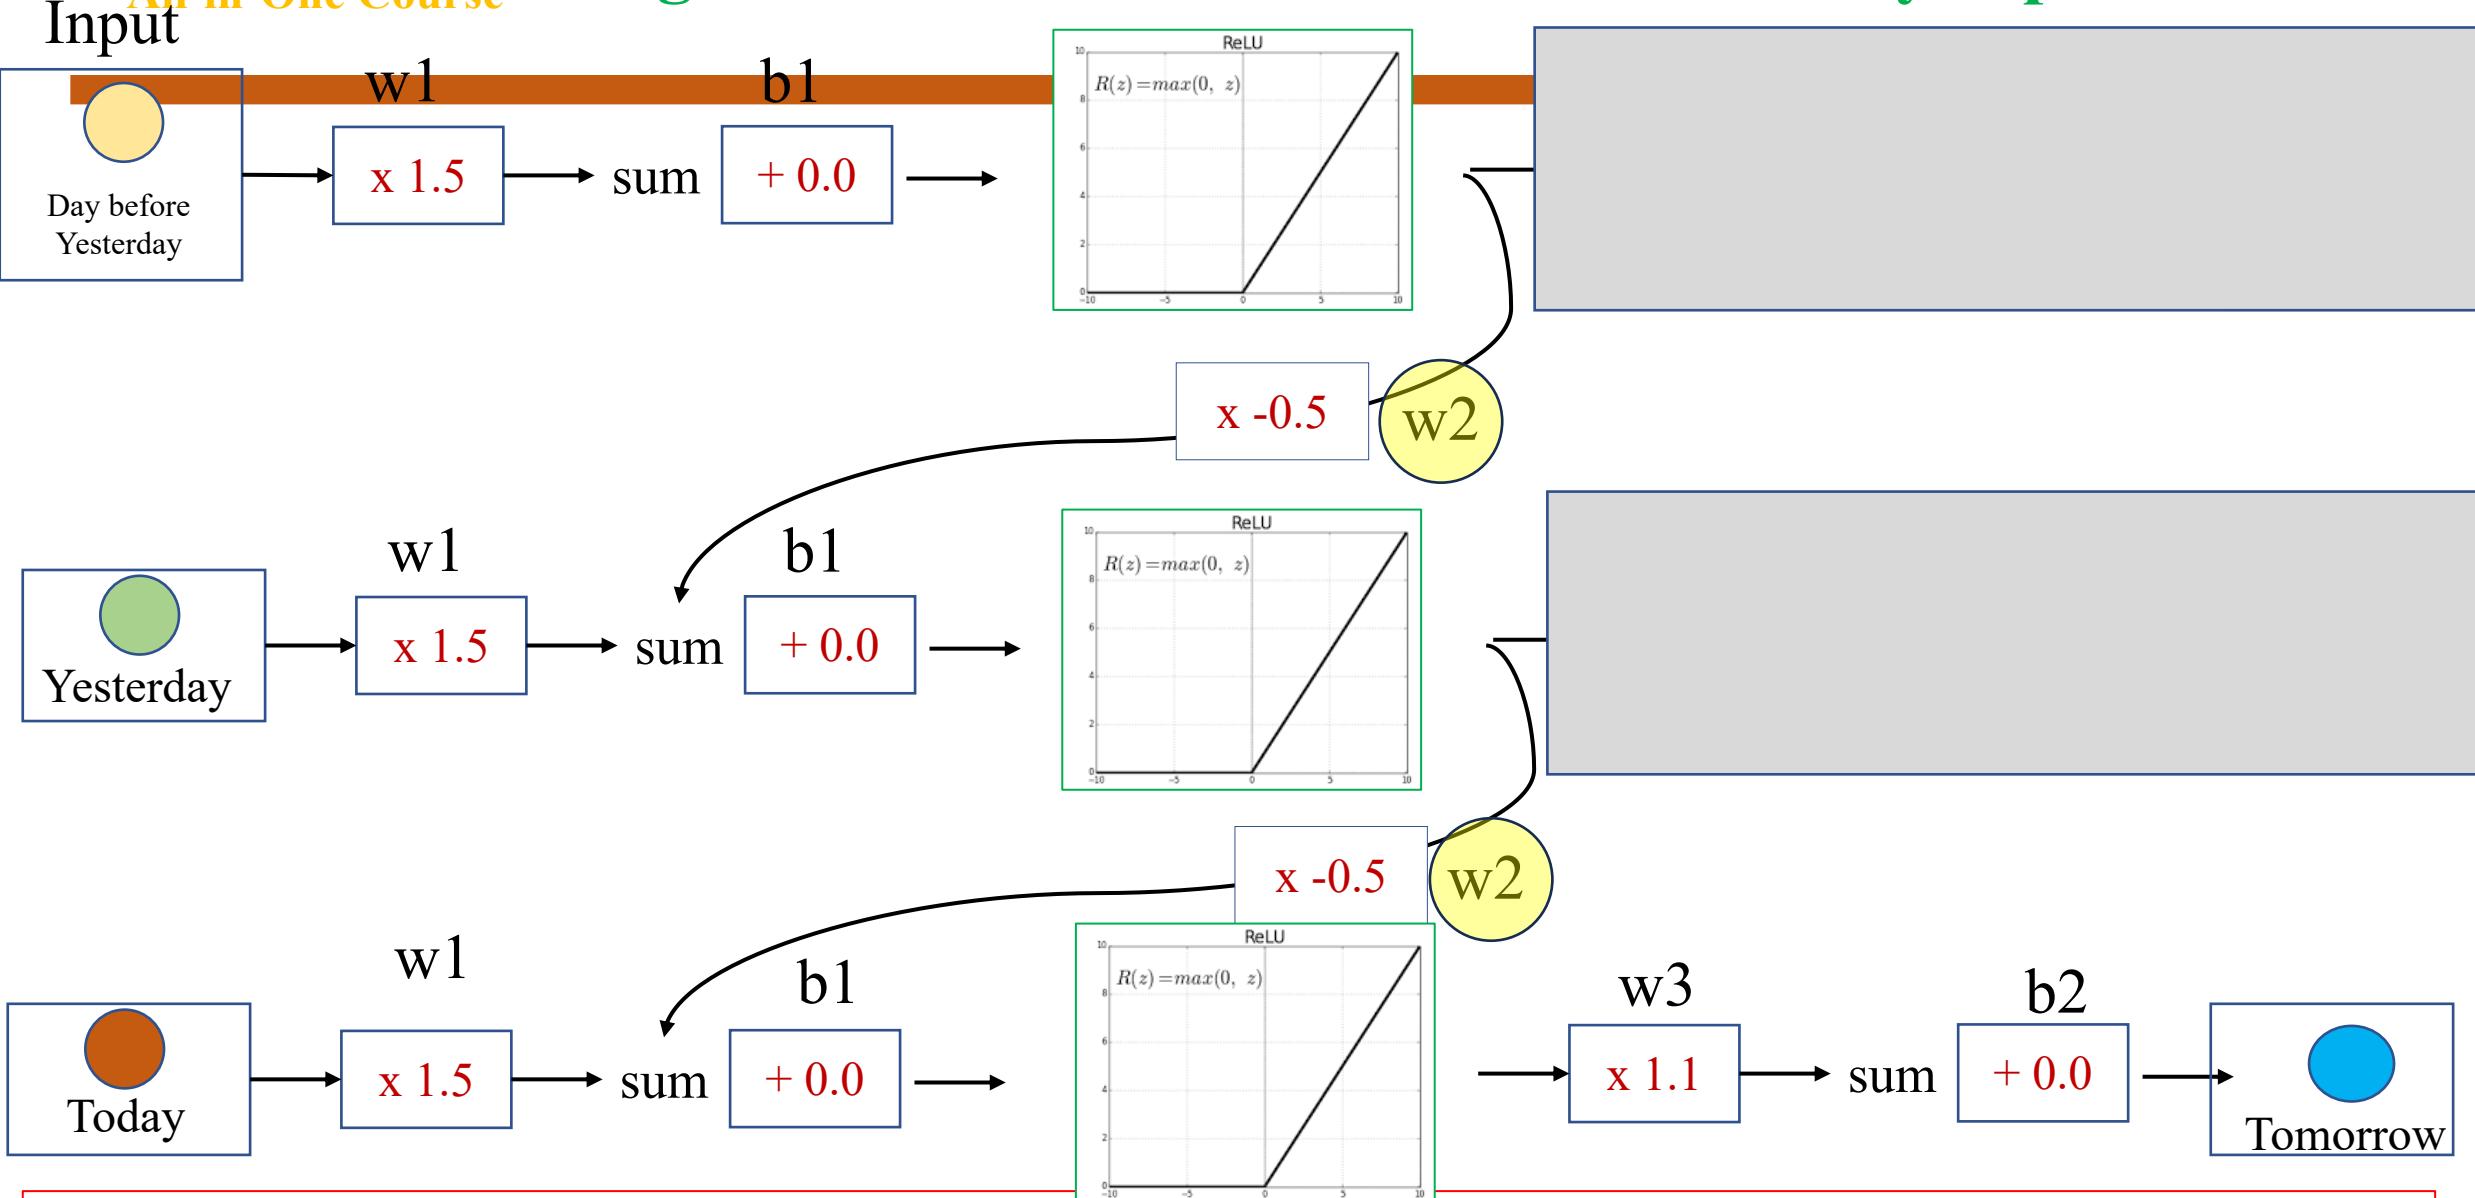
# Recurrent Neural Network : Simple Explanation



# Recurrent Neural Network : Simple Explanation



# Weight and Biases Are Shared Across Every Input



No matter how many times we unroll a recurrent neural network, we never increase number of weights and bias to train

## Define RNN Structure in Pytorch

```
▶ class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, ahead):
        super(RNN, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.ahead = ahead
        self.output_size = output_size

        # RNN Layer - can be replaced with nn.LSTM or nn.GRU
        self.rnn = nn.RNN(input_size, hidden_size, num_layers, batch_first=True)

        # Output layer
        self.fc = nn.Linear(hidden_size, output_size * ahead)

    def forward(self, x):
        # Initialize hidden state
        h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)

        # Forward propagate RNN
        out, _ = self.rnn(x, h0)

        # Get the last time step's output for each sequence
        out = out[:, -1, :]

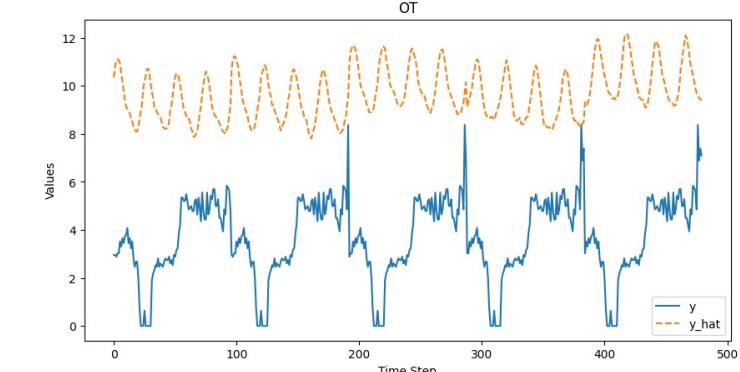
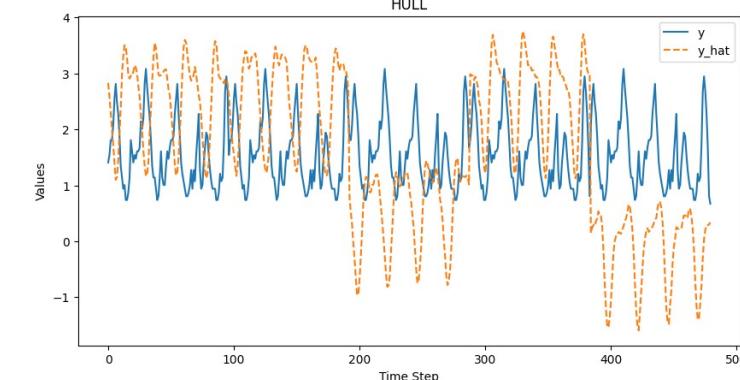
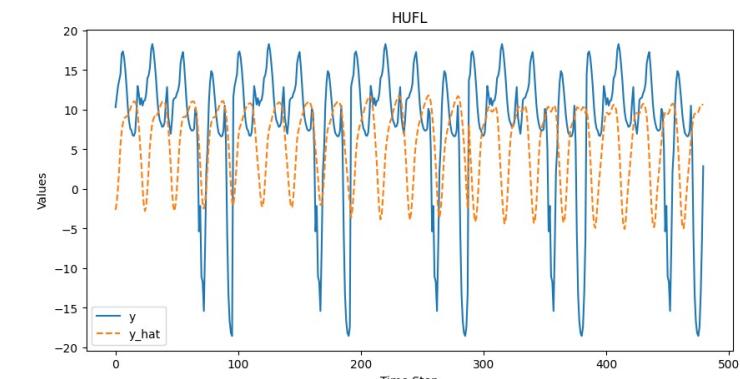
        # Pass through the linear layer and reshape
        out = self.fc(out).view(-1, self.ahead, self.output_size)
        return out
```

## Configuration

```
▶ input_size = 336
label_size = 96
offset = 1
train_size = 0.7
val_size = 0.1
num_epochs = 1_000_000
patience = 20
learning_rate = 0.001
hidden_size = 64
num_layers = 2
ele = 5
target_name = 'OT'
date_column = 'date'
file_path = 'ETTh1.csv'
plot_dir = 'plots'
weight_dir = 'weights'
results = []
```

## Define RNN Structure in Pytorch

```
▶ RNN_multi2multi = RNN(input_size=multi2multi_loader.in_variable,  
                         hidden_size=hidden_size,  
                         output_size=multi2multi_loader.out_variable,  
                         ahead=label_size,  
                         num_layers=num_layers)  
RNN_multi2multi_manager = ModelManager(model=RNN_multi2multi,  
                                         train_loader=multi2multi_loader.train_loader,  
                                         val_loader=multi2multi_loader.val_loader,  
                                         lr=learning_rate,  
                                         patience=patience)  
RNN_multi2multi_manager.train(num_epochs=num_epochs, Loading...  
                               save_dir=os.path.join(weight_dir, sub_dir))  
results.append({  
    "Name": RNN_multi2multi_manager.model.__class__.__name__,  
    "Type": sub_dir,  
    "MAE": RNN_multi2multi_manager.evaluate(loader=multi2multi_loader.test_loader)  
})  
results[-1]  
  
👤 Epoch [29/1000000], time: 13s, loss: 1.7758, val_loss: 2.7365  
Model saved to weights\multi2multi\best-RNN.pth  
Epoch [30/1000000], time: 13s, loss: 1.7012, val_loss: 2.5789  
Epoch [31/1000000], time: 13s, loss: 1.6243, val_loss: 2.6399  
Epoch [32/1000000], time: 13s, loss: 1.6731, val_loss: 2.6522  
Model saved to weights\multi2multi\best-RNN.pth
```

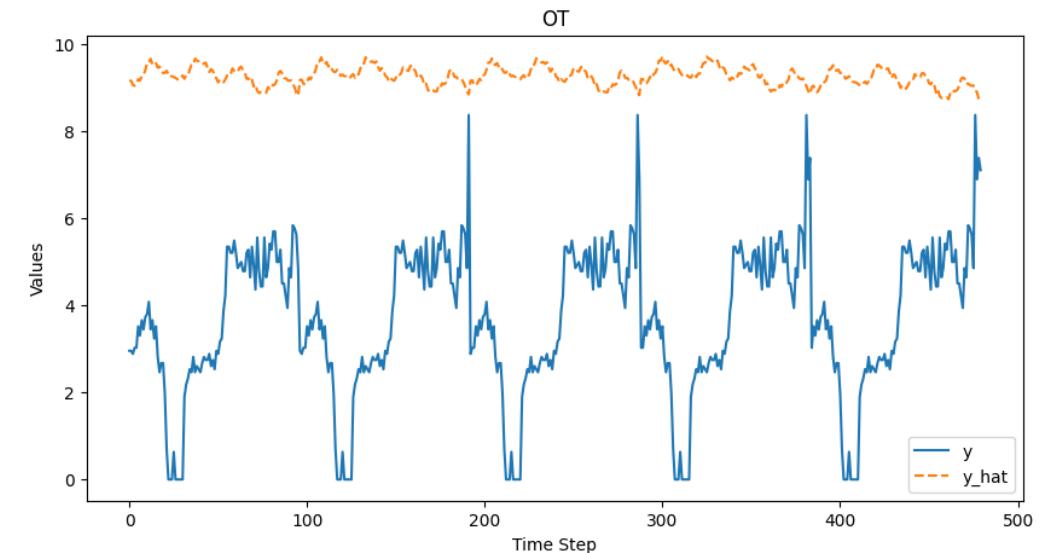


## Case 1: Multivariate to Multivariate Prediction

*Experiments were performed by TA Khoa Nguyen*

## Define RNN Structure in Pytorch

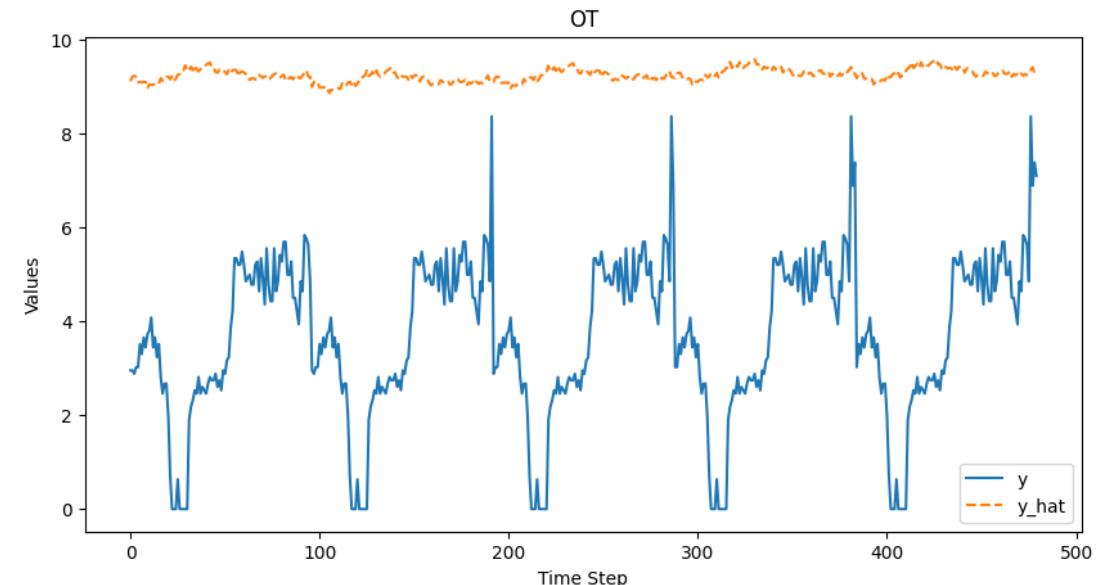
```
▶ features_type='MS'  
sub_dir = 'multi2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
multi2uni_loader = TimeSeriesDataLoader(file_path,  
                                         input_size=input_size,  
                                         label_size=label_size,  
                                         offset=offset,  
                                         train_size=train_size,  
                                         val_size=val_size,  
                                         target_name=target_name,  
                                         features_type=features_type,  
                                         date_column=date_column)  
  
👤 Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 7)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 7)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 7)  
self.y_test.shape = (2861, 96, 1)
```



## Case 2: Multivariate to Univariate Prediction

## Define RNN Structure in Pytorch

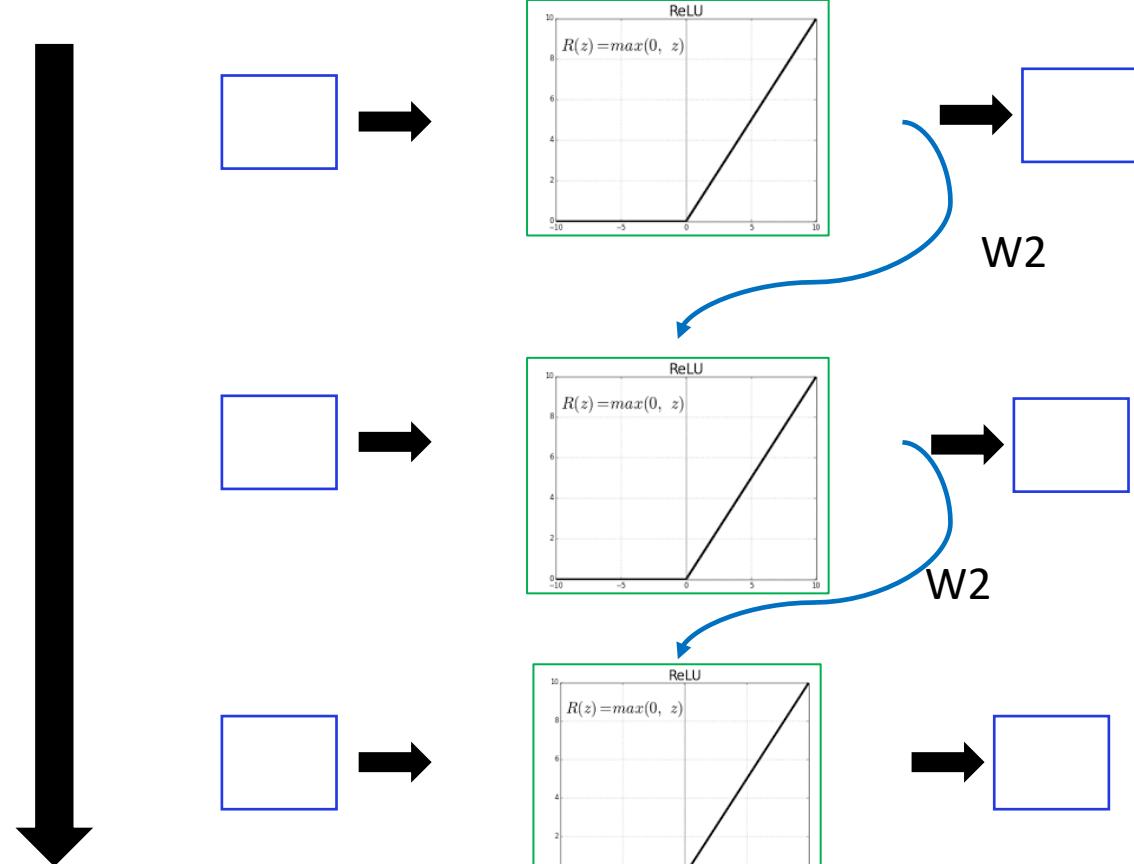
```
▶ features_type='S'  
sub_dir = 'uni2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
uni2uni_loader = TimeSeriesDataLoader(file_path,  
                                       input_size=input_size,  
                                       label_size=label_size,  
                                       offset=offset,  
                                       train_size=train_size,  
                                       val_size=val_size,  
                                       target_name=target_name,  
                                       features_type=features_type,  
                                       date_column=date_column)  
  
⌚ Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 1)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 1)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 1)  
self.y_test.shape = (2861, 96, 1)
```



## Case 3: Univariate to Univariate Prediction

# Why RNN is not Use Often

The more we unroll the recurrent neural network, the harder it is to train



Vanishing / Exploding Gradient Problem

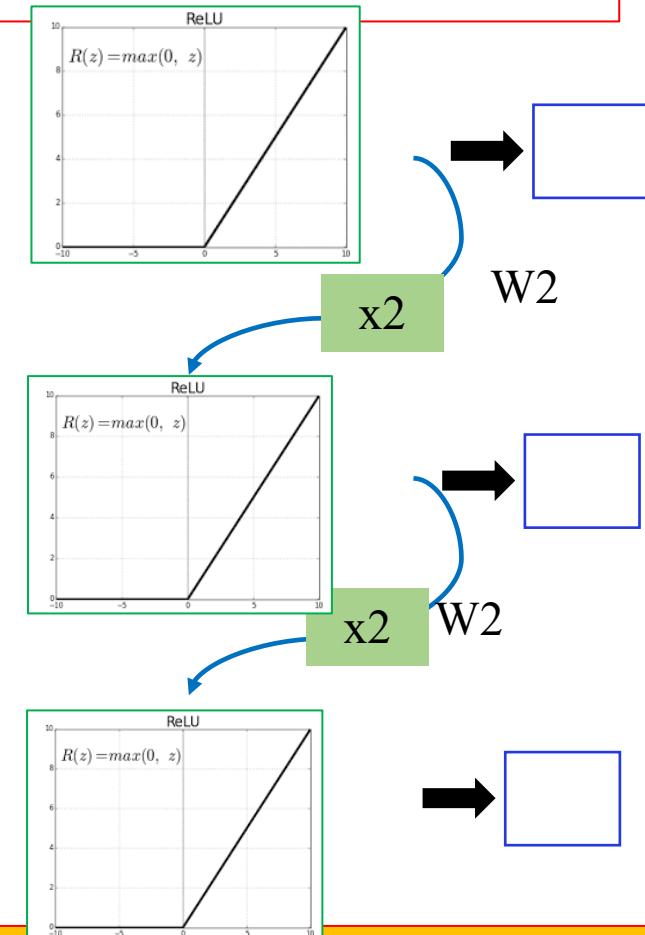
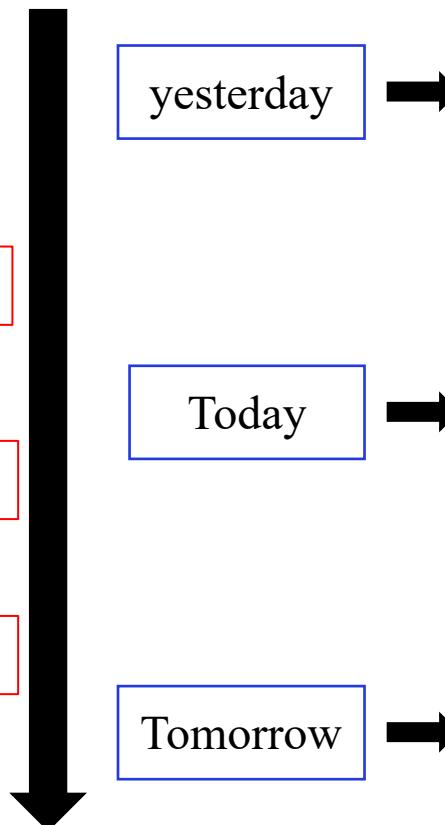
# Why RRN is not use Often

If we set  $W2 \in [1, +\infty]$ , Let's start with  $W2 = 2$

$$\text{yesterday} \times 2 \times 2 = \text{yesterday} \times 2^2 = \text{yesterday} \times W_2^2$$

$$\text{Yesterday} \times 2^{100} = \text{yesterday} \times 2^{100} = \text{yesterday} \times W_2^{100}$$

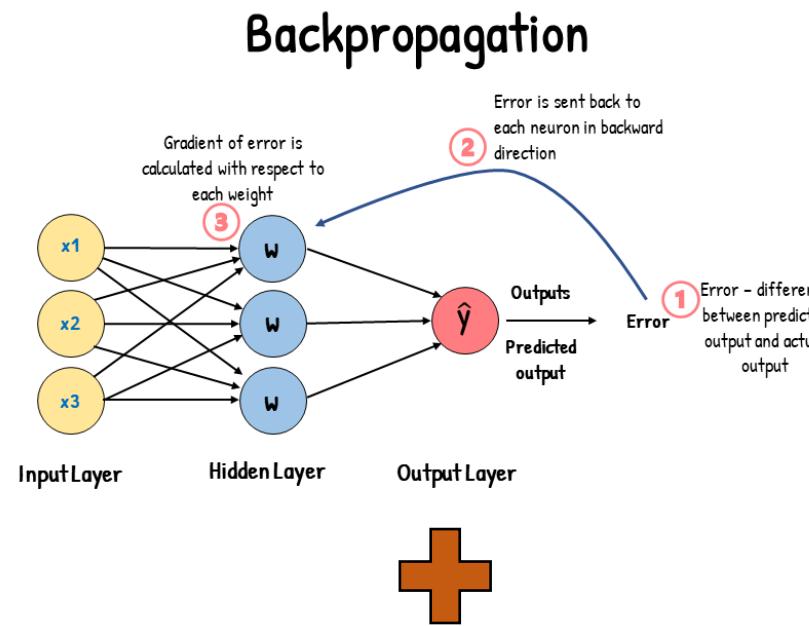
Yesterday  $\times$  Huge Number



Exploding Gradient Problem

# Why RRN is not use Often

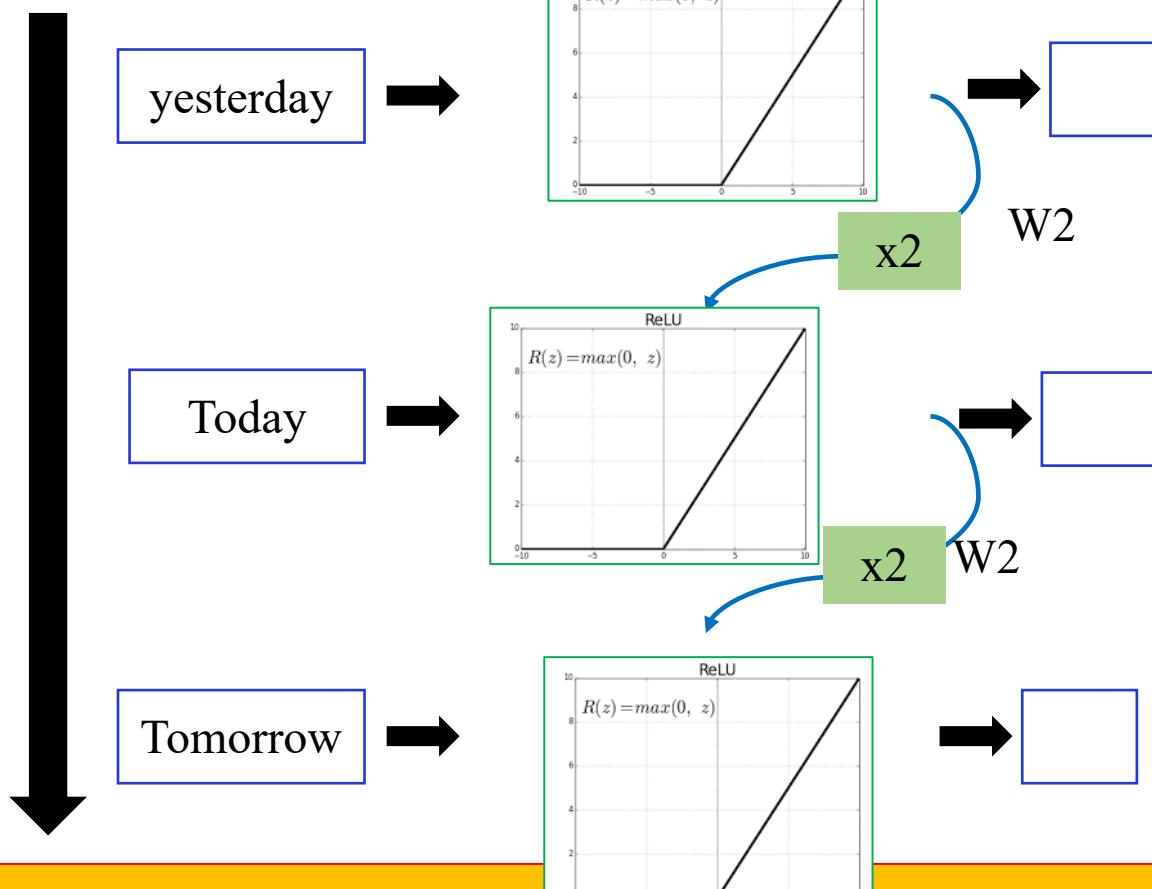
If we set  $W2 \in [1, +\infty]$ , Let's start with  $W2 = 2$



**W is a huge numbers**

**Overshoot the optimal solution**

**Exploding Gradient Problem**

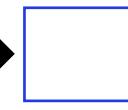
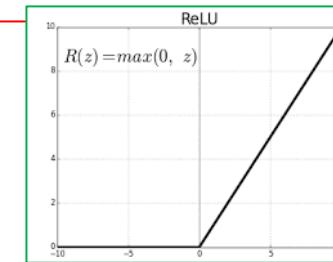


# Why RRN is not use Often

If we set  $W2 \in [1, +\infty]$ , Let's start with  $W2 = 2$

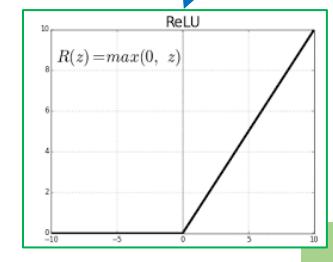


yesterday →



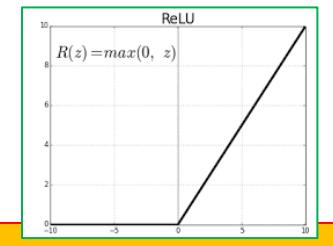
$W2$

Today →



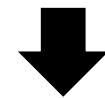
$x2$   
 $W2$

Tomorrow →



$x2$   
 $W2$

$$\text{yesterday} \times 0.5 \times 0.5 = \text{yesterday} \times 0.5^2 = \text{yesterday} \times W_2^2$$



$$\text{Yesterday} \times 0.5^{100} = \text{yesterday} \times 0.5^{100} = \text{yesterday} \times W_2^{100}$$



Yesterday × Really Small Number

Vanishing Gradient Problem

# How to avoid Exploding and Vanishing Gradient Problem in RNN

I grew up in France... I speak fluent *French*

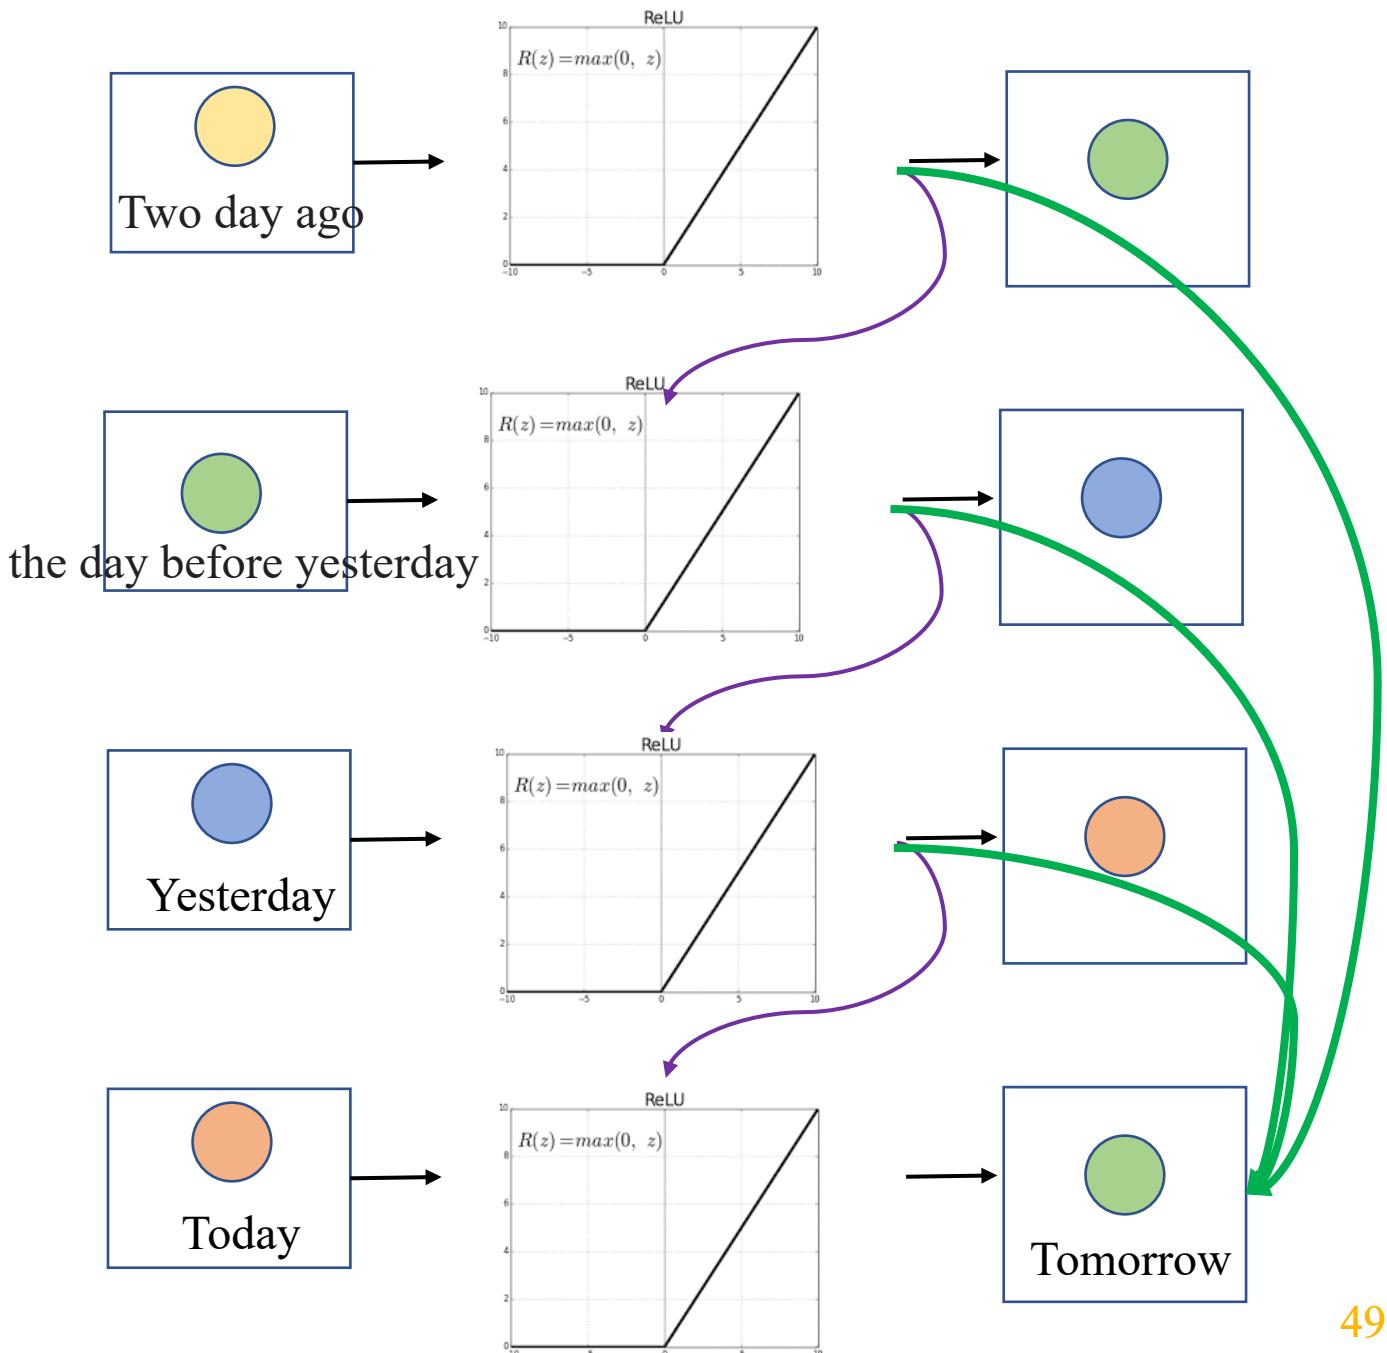
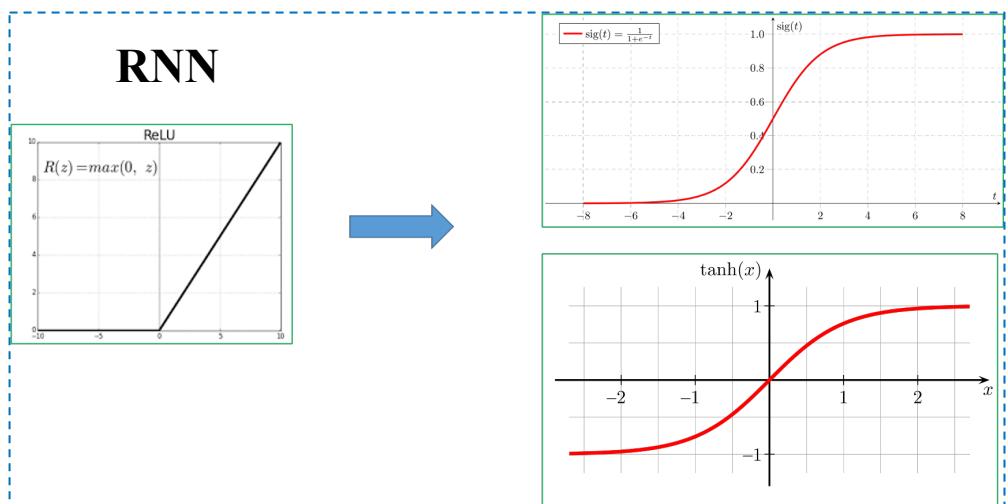
RNN fail to handle

# Outline

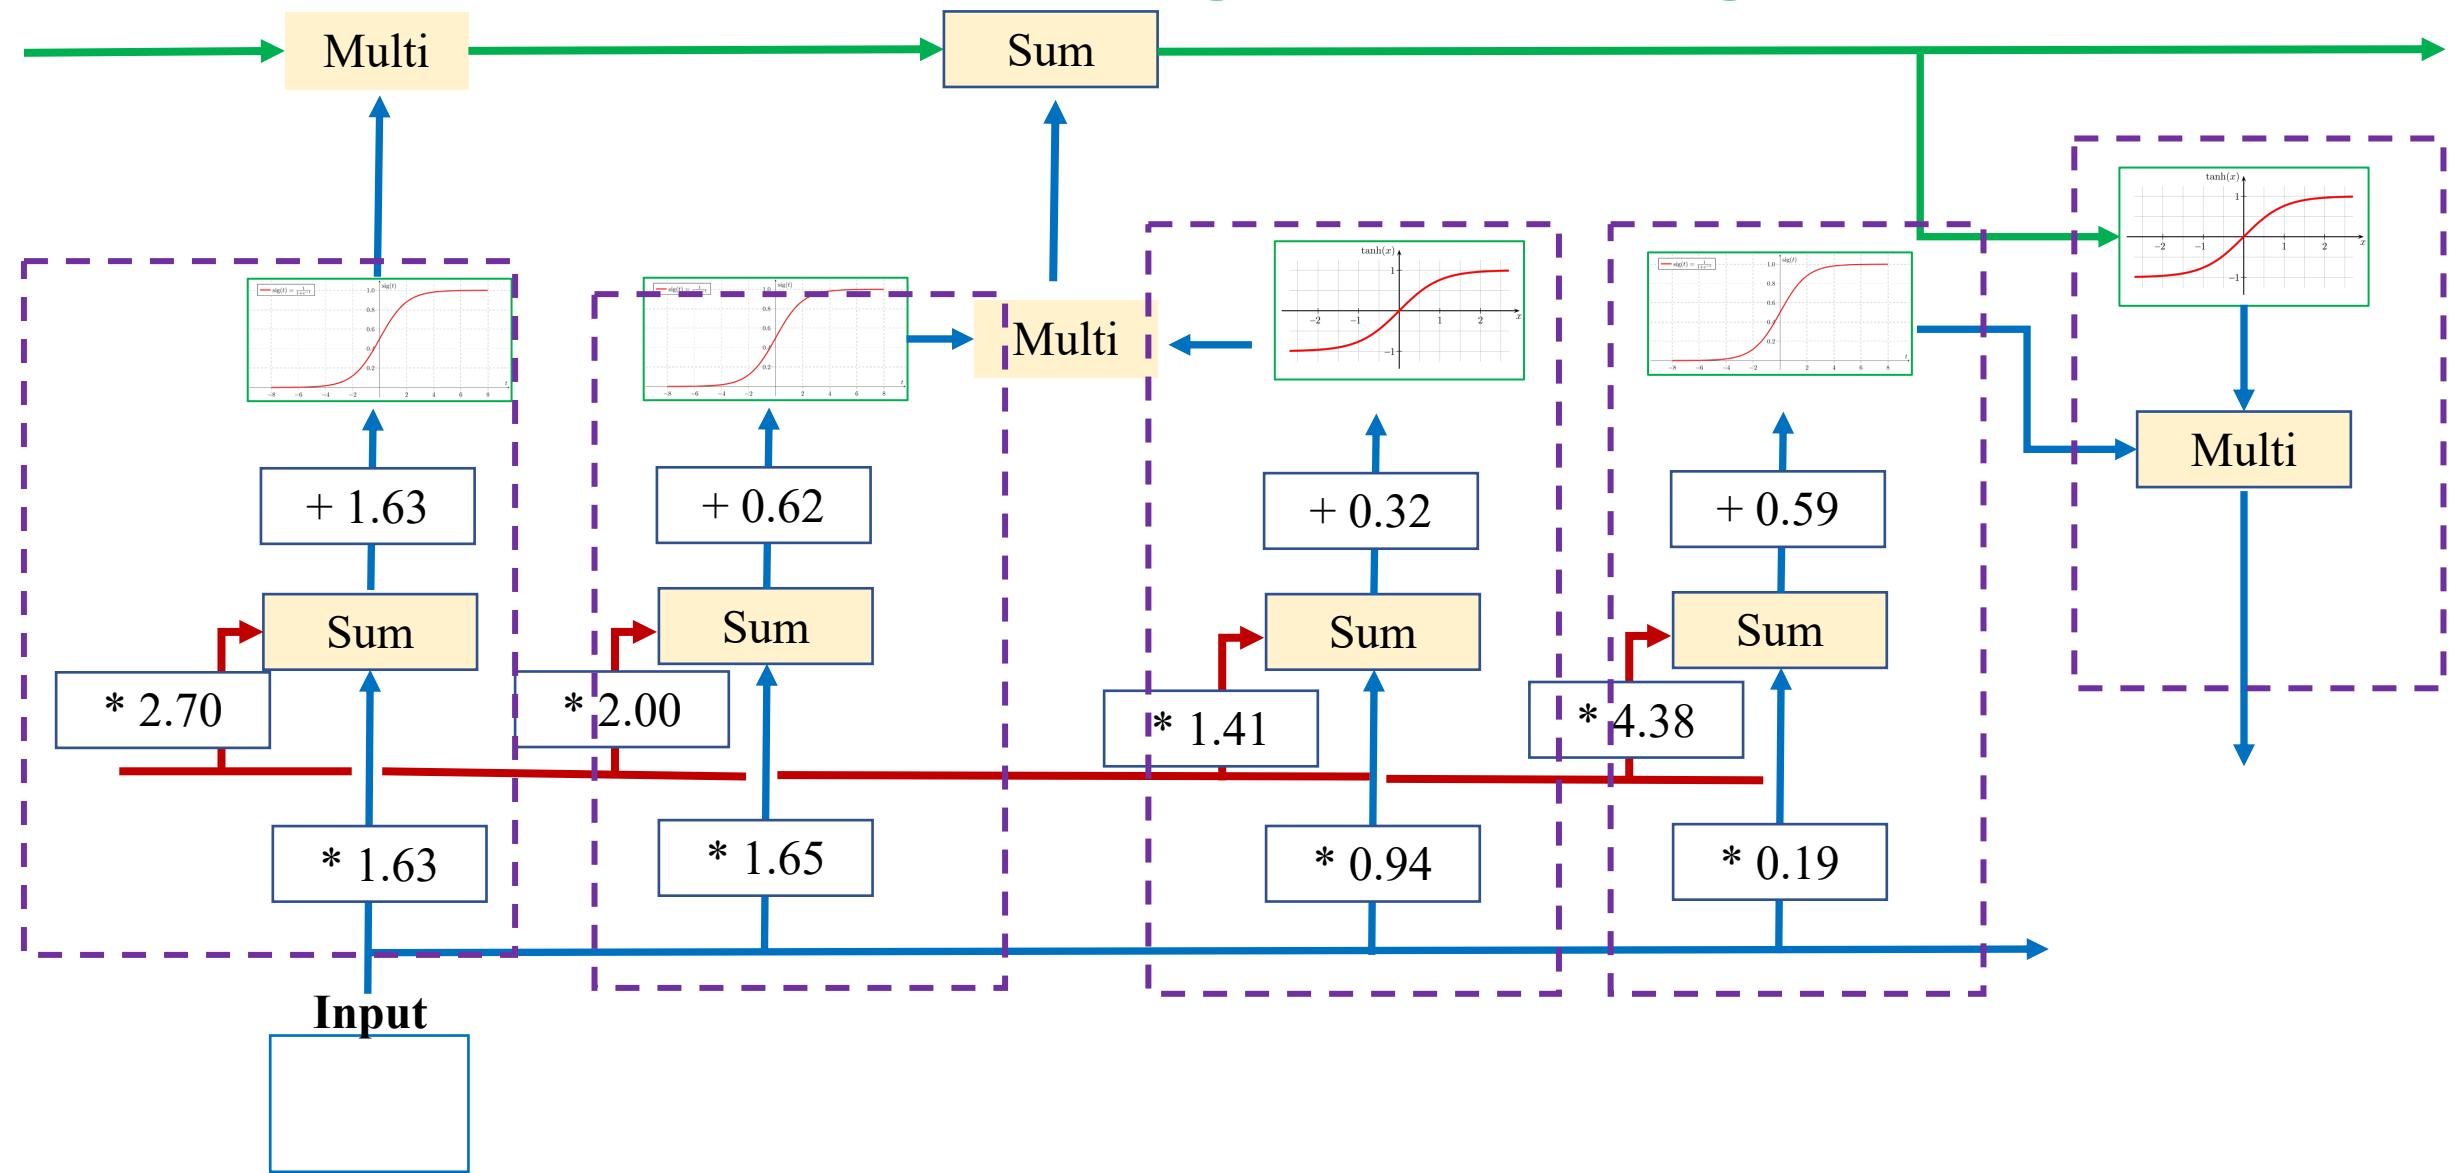
- Multilayer Perceptron and Time-series Data Forecasting
- RNN and Time-series Data Forecasting
- LSTM and Time-series Data Forecasting
- Bi-LSTM and Time-series Data Forecasting
- XGBoost and Time-series Data Forecasting

# Long short-term Memory: Simple Explanation

Short memory  
→  
Long memory



# Three Stages in a Single LSTM Unit

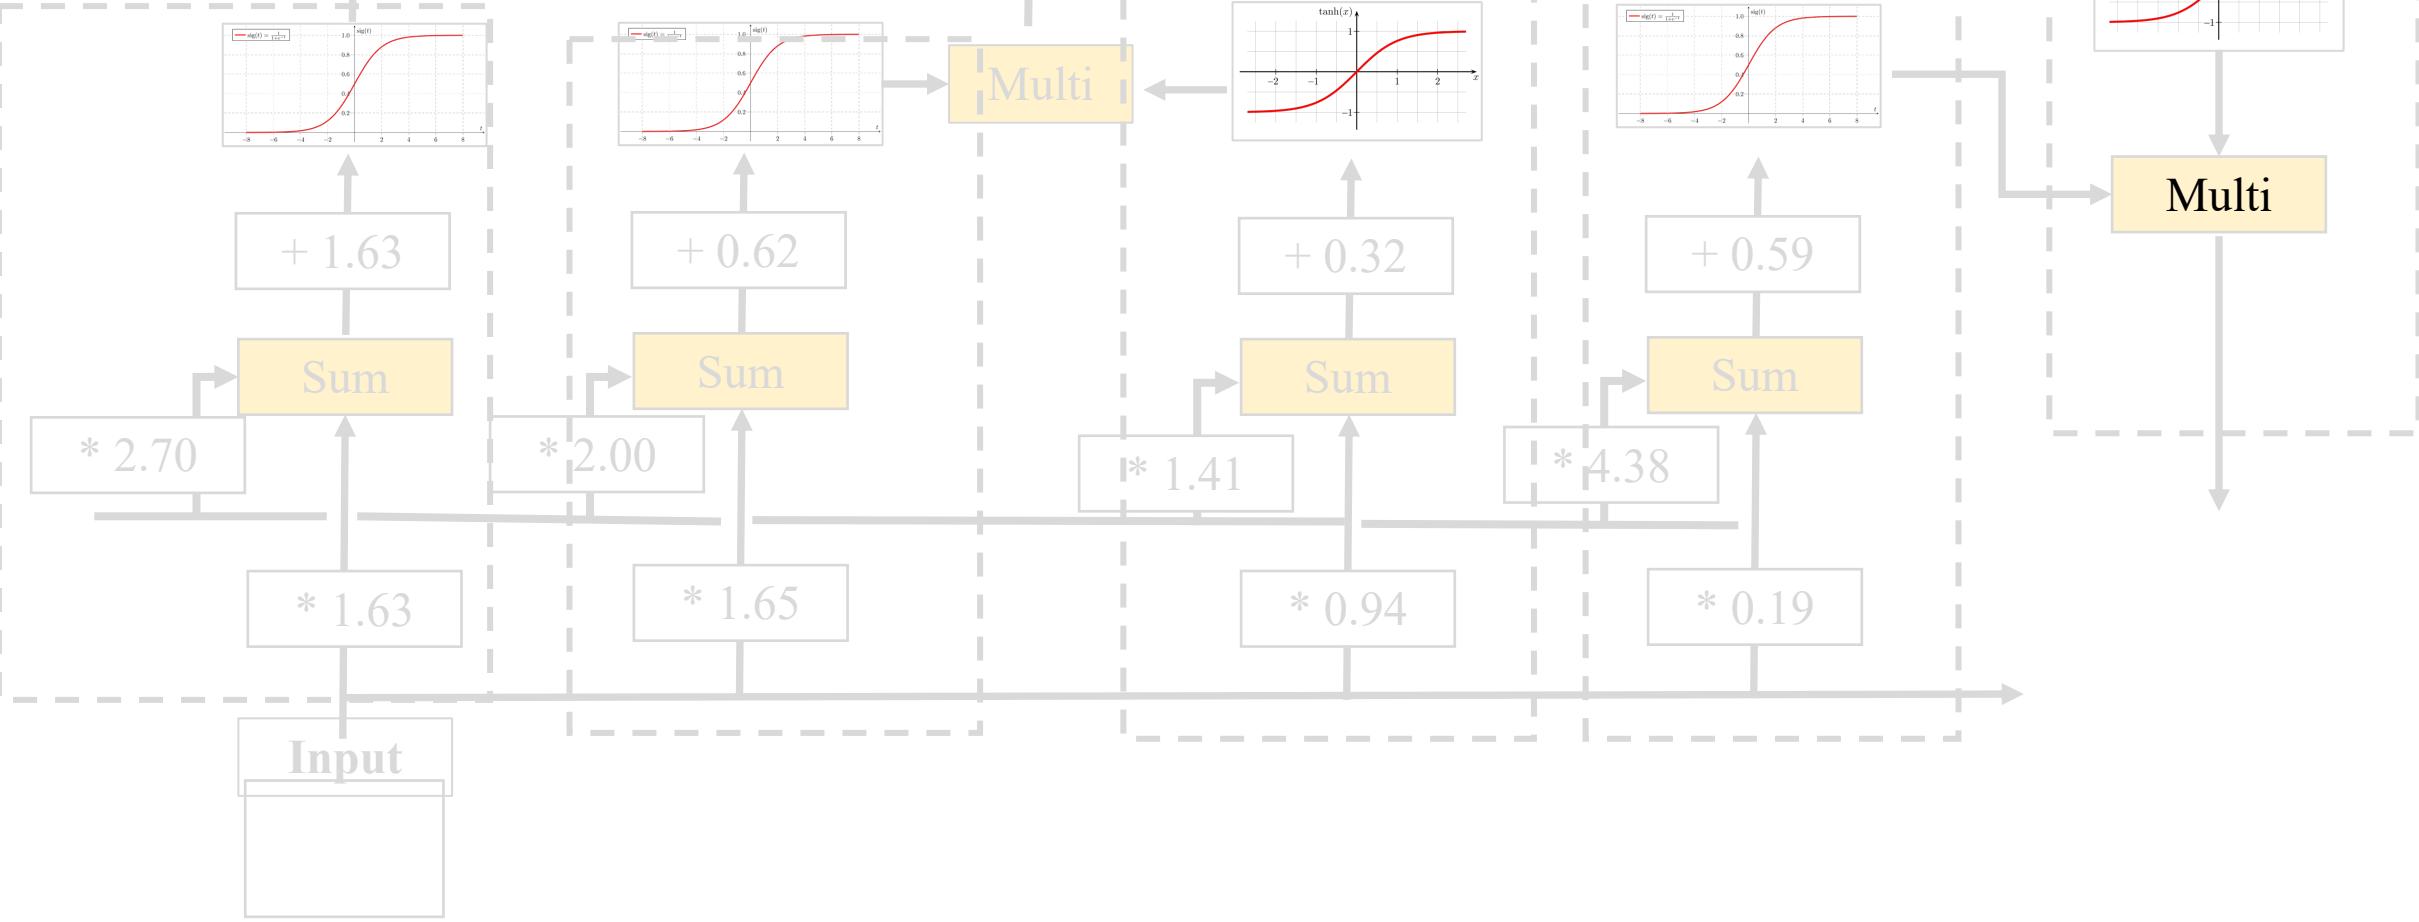
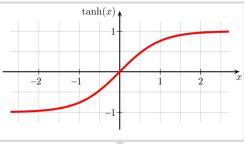


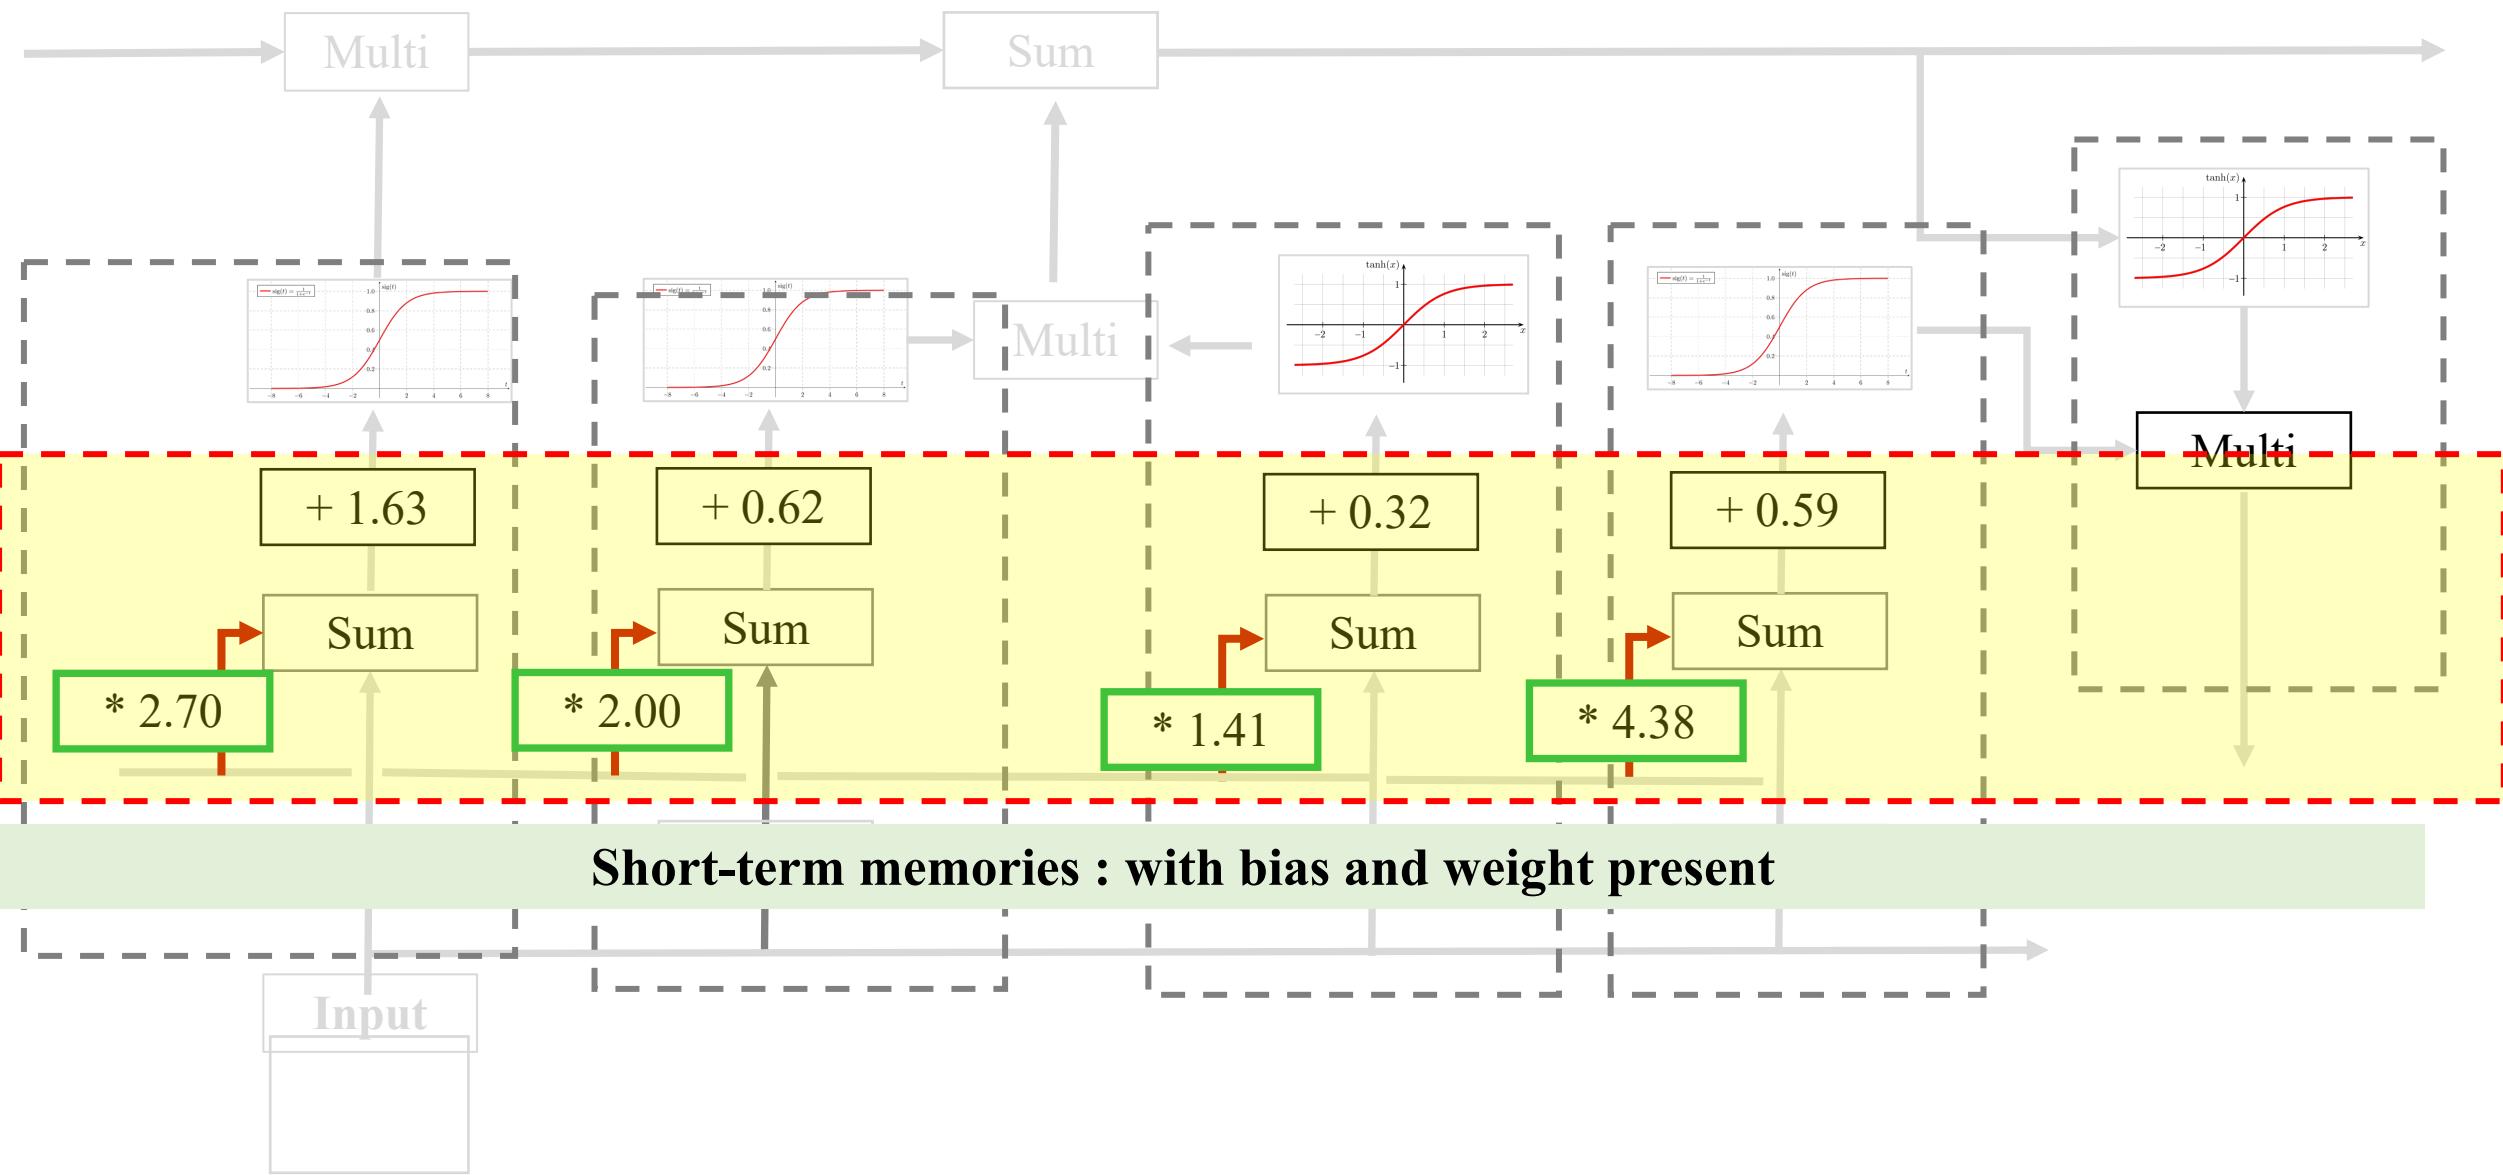
Long-term memories : no bias and weight present => avoid causing exploding and vanishing gradient

All-in-One Course

Multi

Sum





Long-term memories : no bias and weight present => avoid causing exploding and vanishing gradient

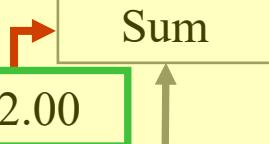
All-in-One Course

Multi

Sum



How do long-term memory and short-term memories work?



Short-term memories : with bias and weight present

Reduce the long-term memory

AI in One Course

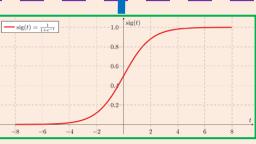
2

Multi

1.99

Sum

0.977



+ 1.62

Sum

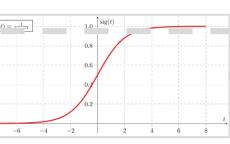
1

\* 2.70

\* 1.63

Input

1

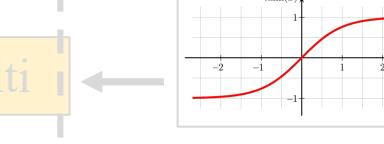


+ 0.62

Sum

\* 2.00

\* 1.65

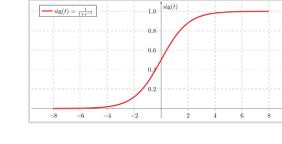


+ 0.32

Sum

\* 1.41

\* 0.94

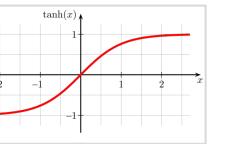


+ 0.59

Sum

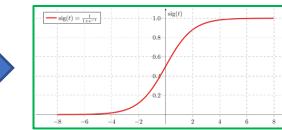
\* 4.38

\* 0.19

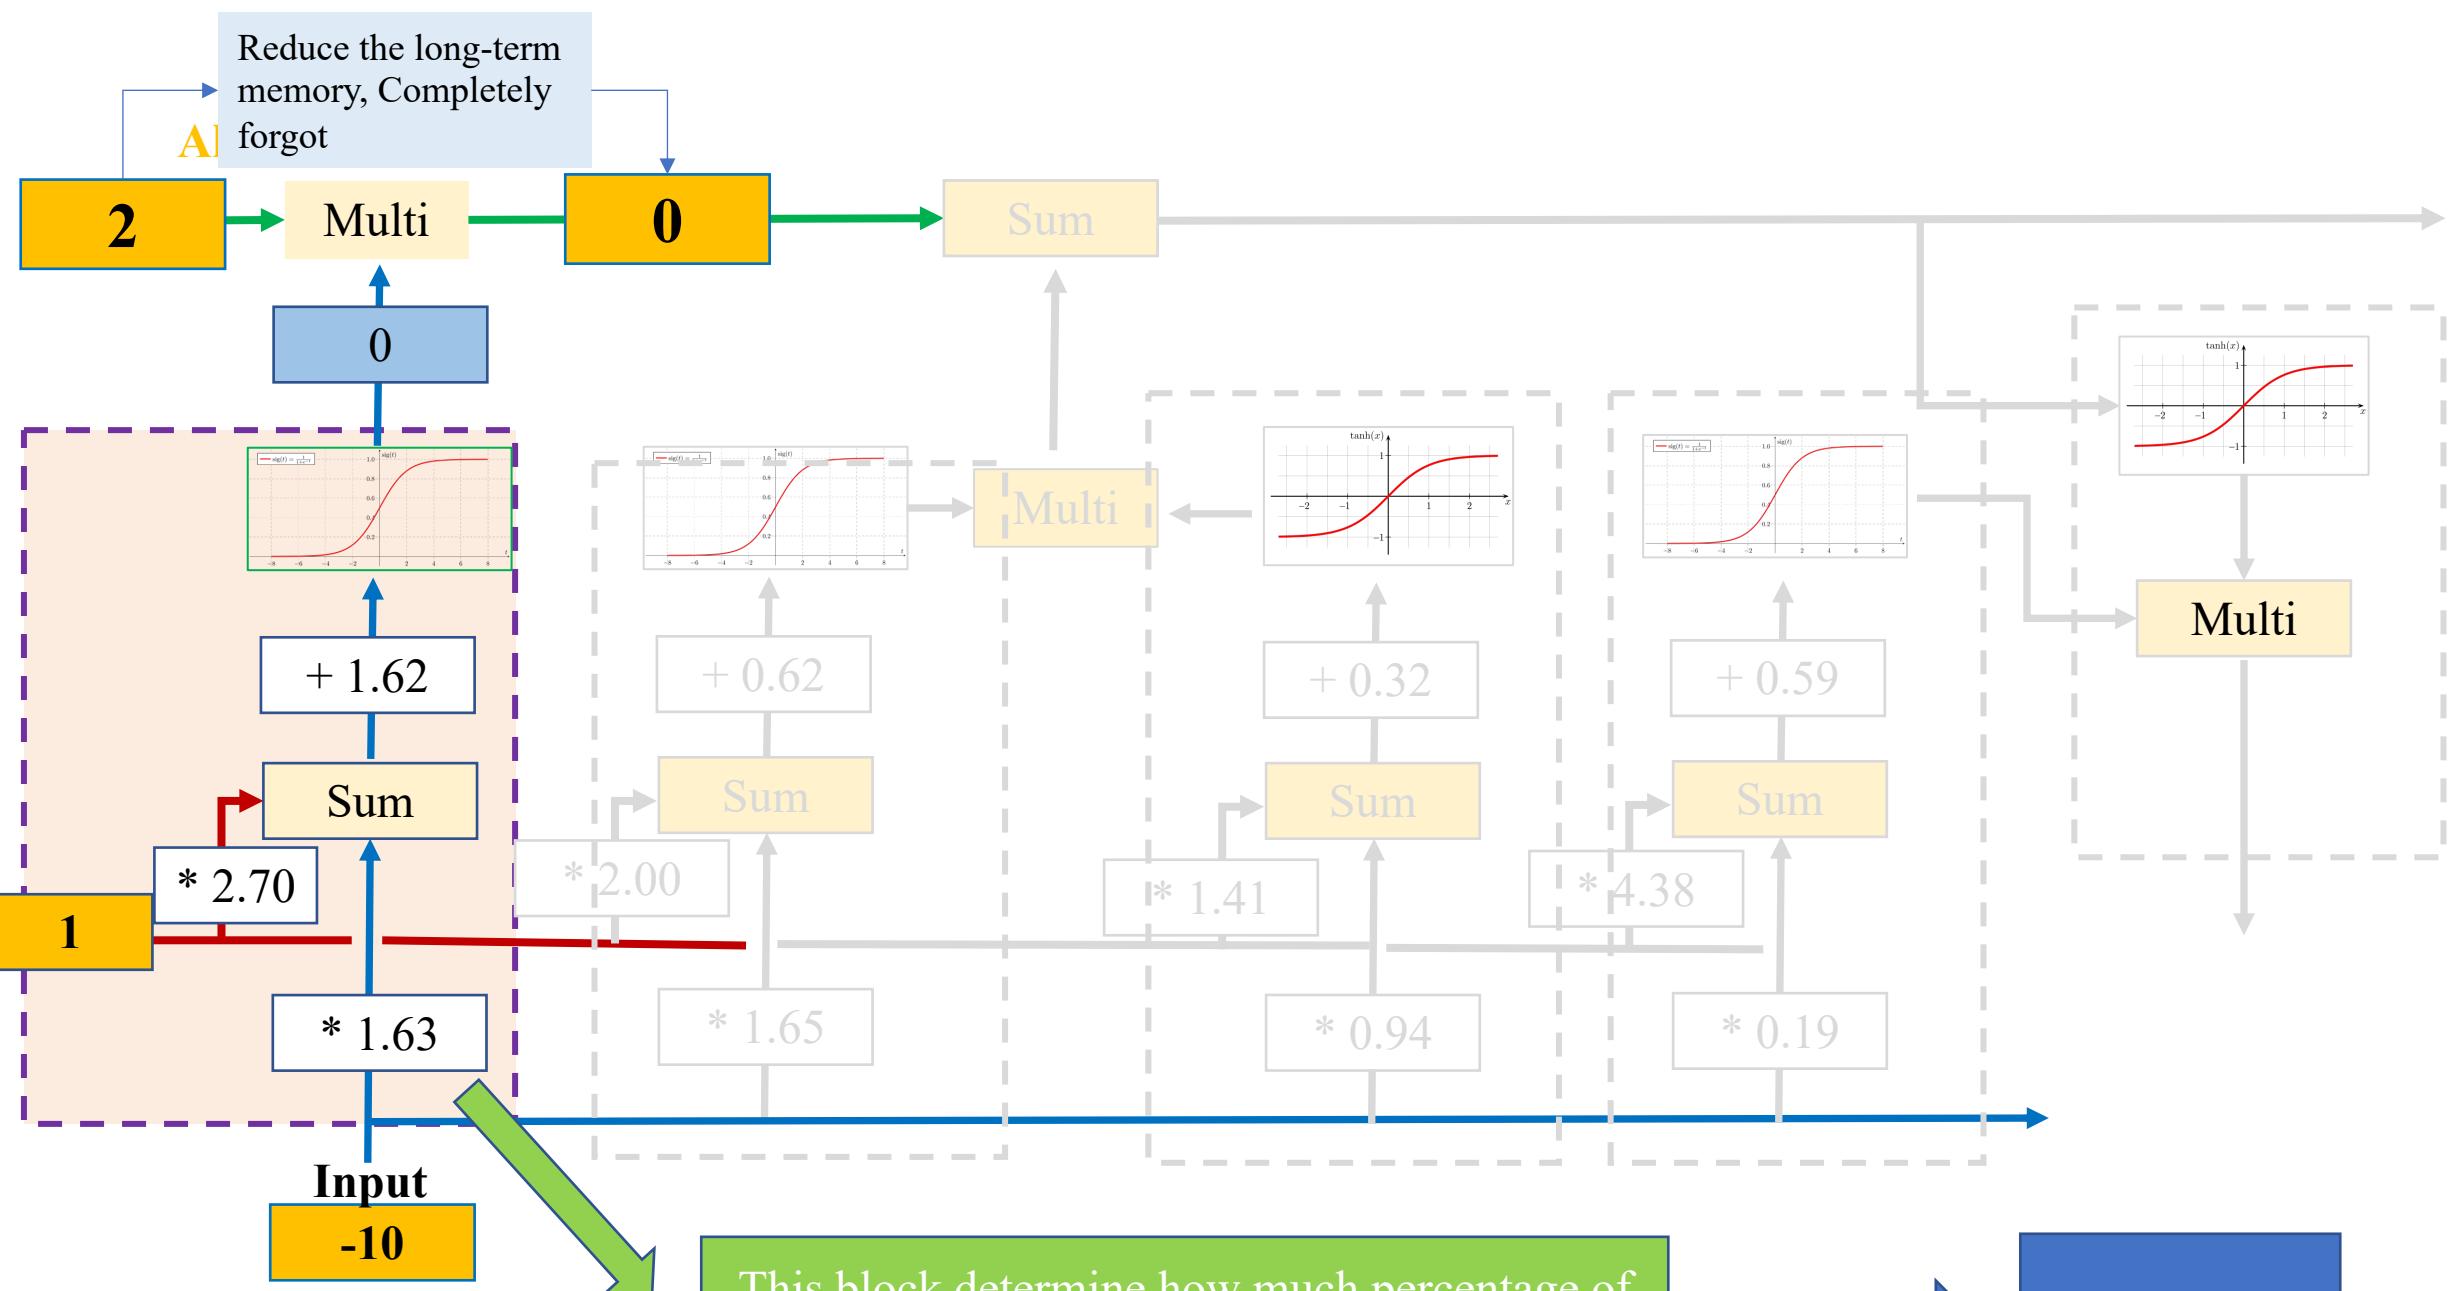


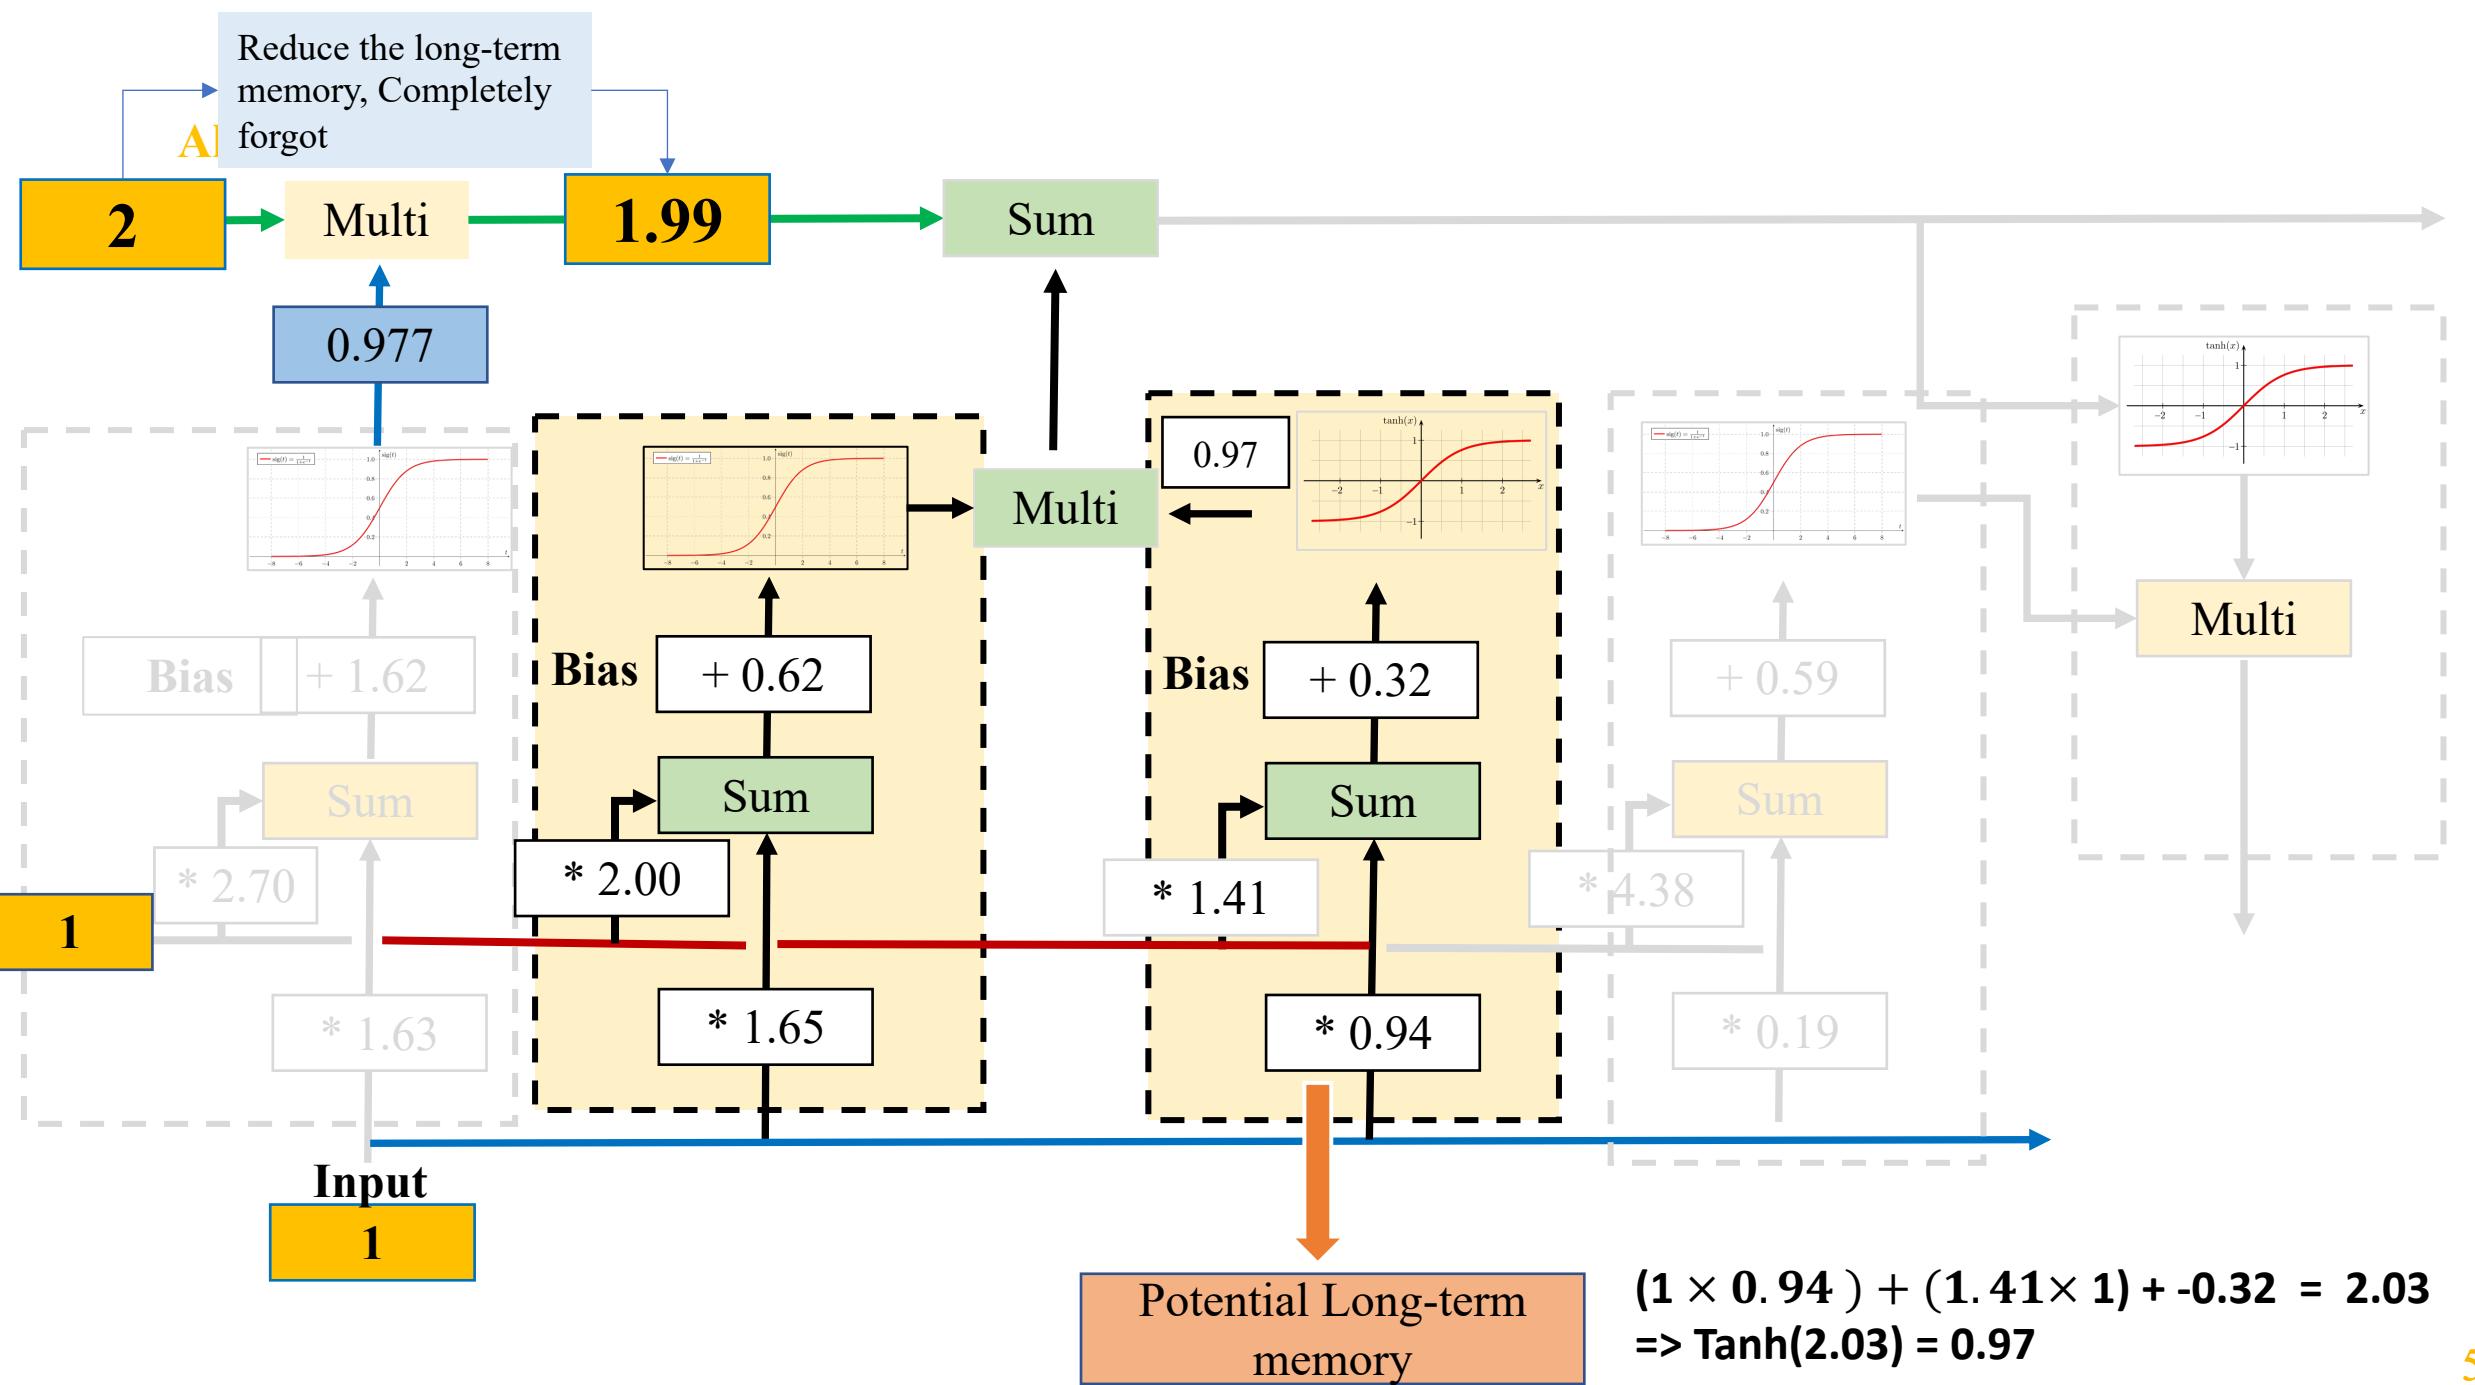
Multi

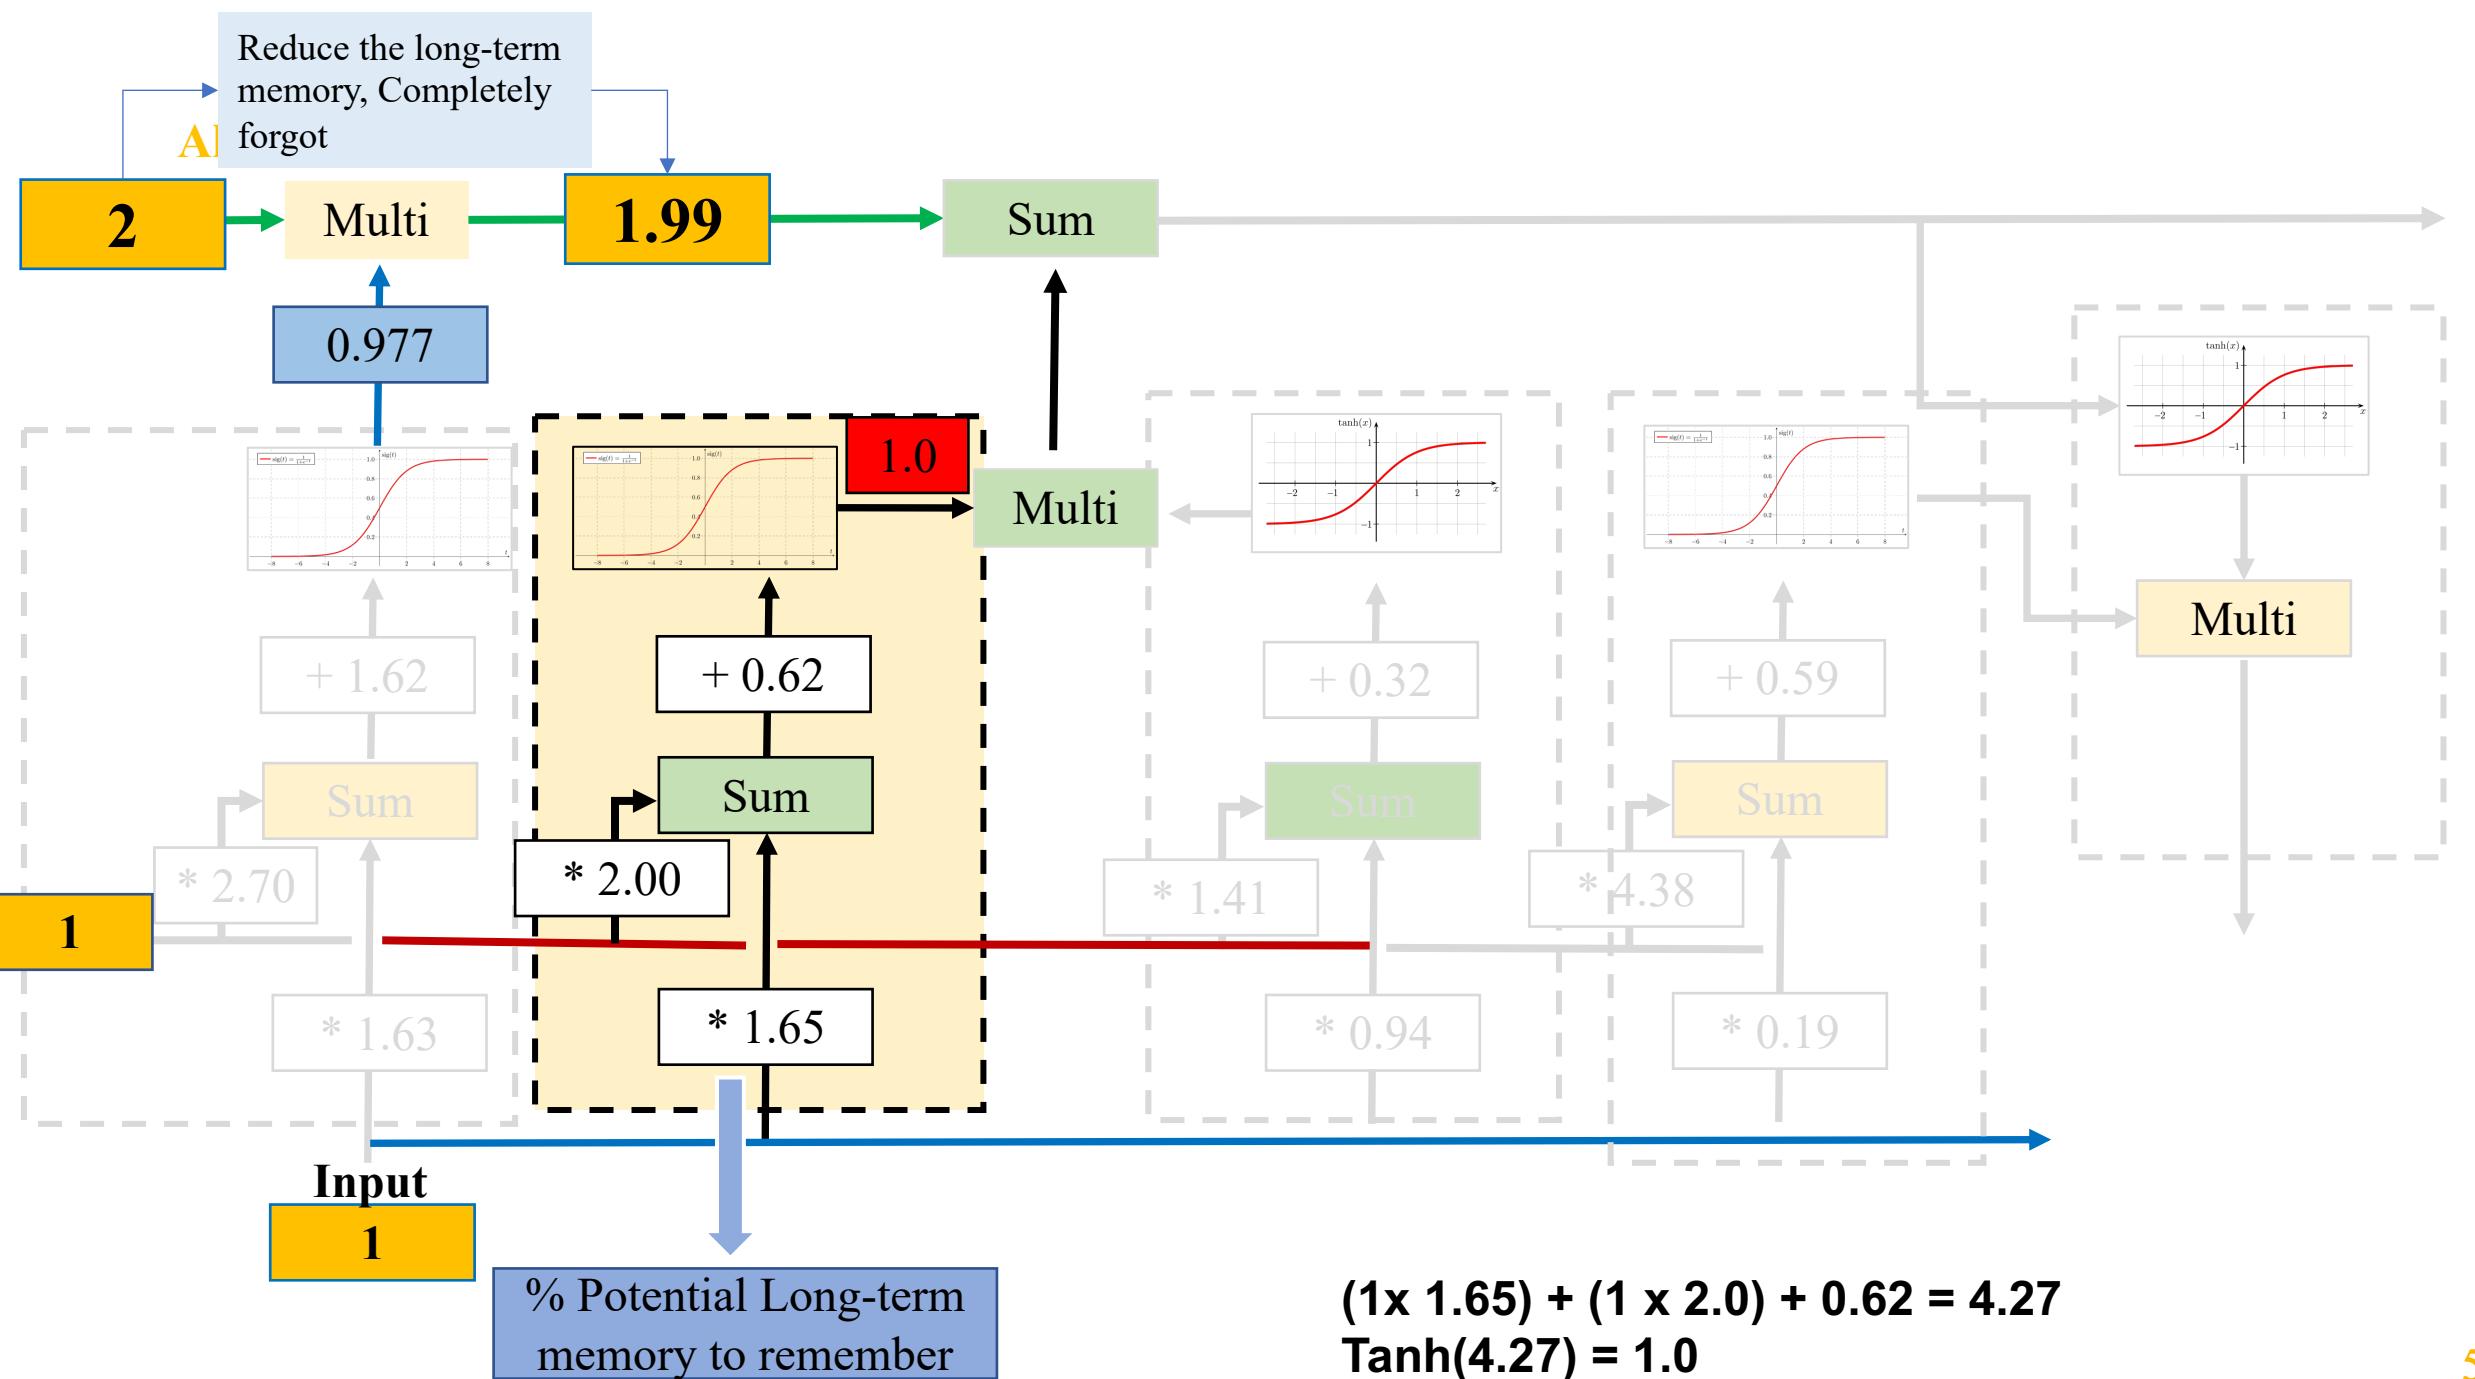
$$(1 \times 1.63) + (1 \times 2.70) + 1.62 = 5.95$$

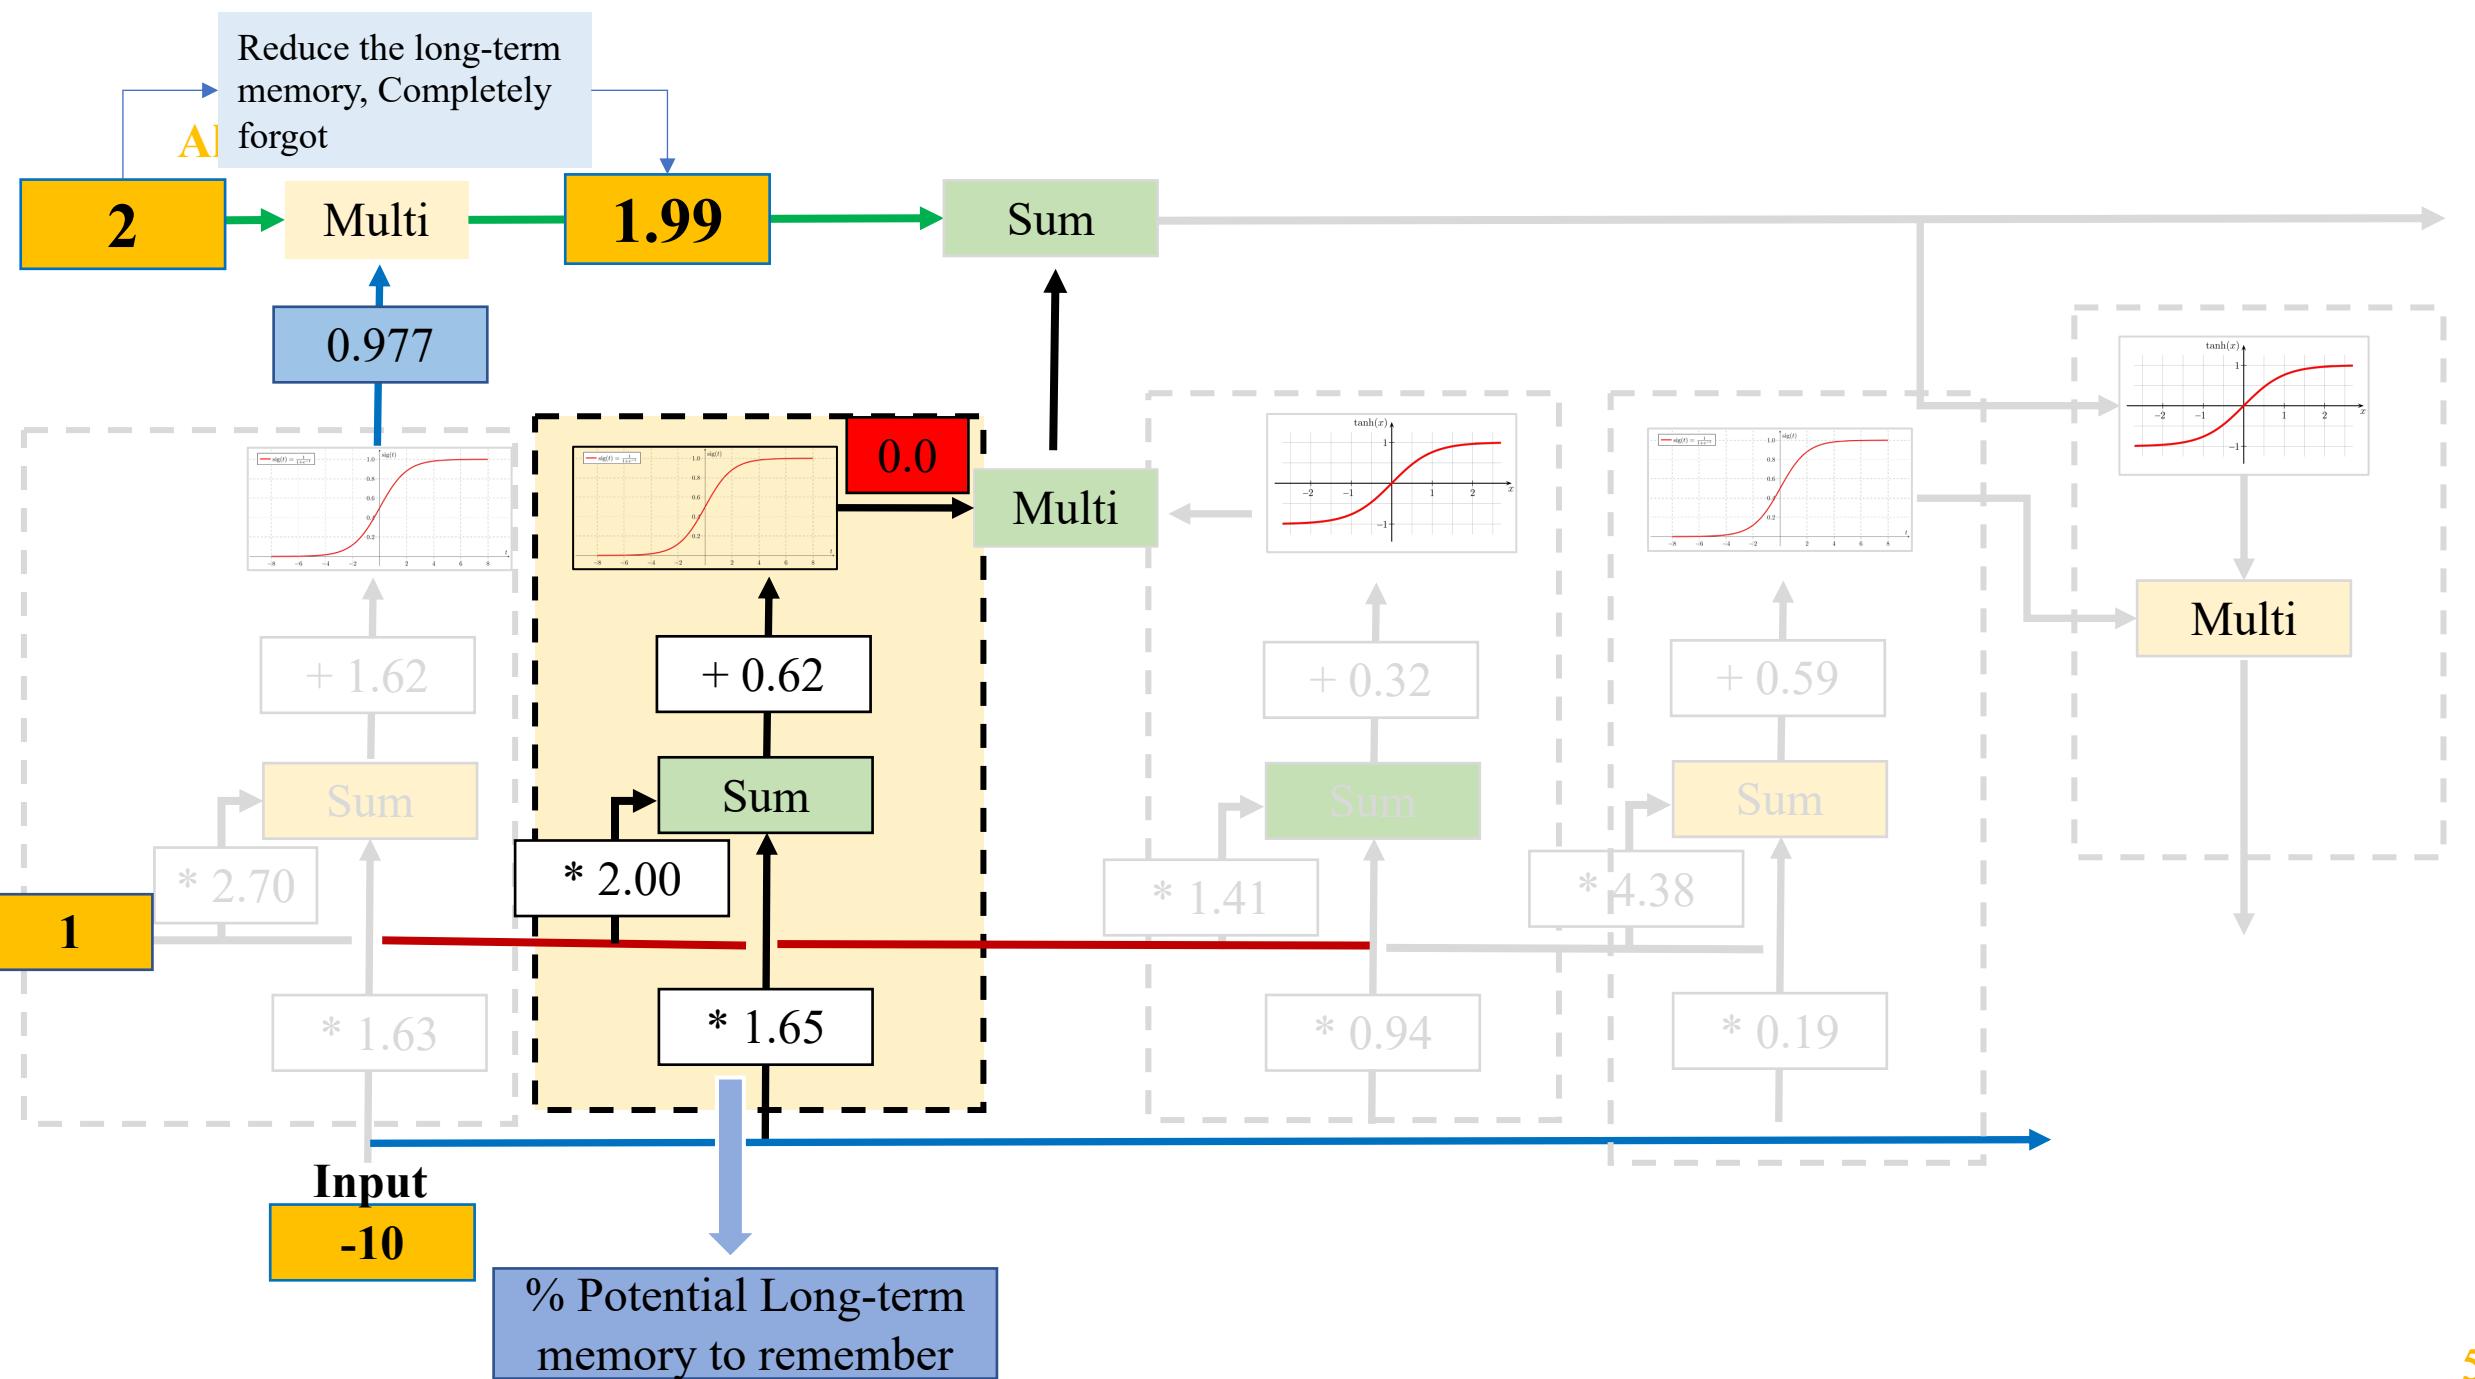


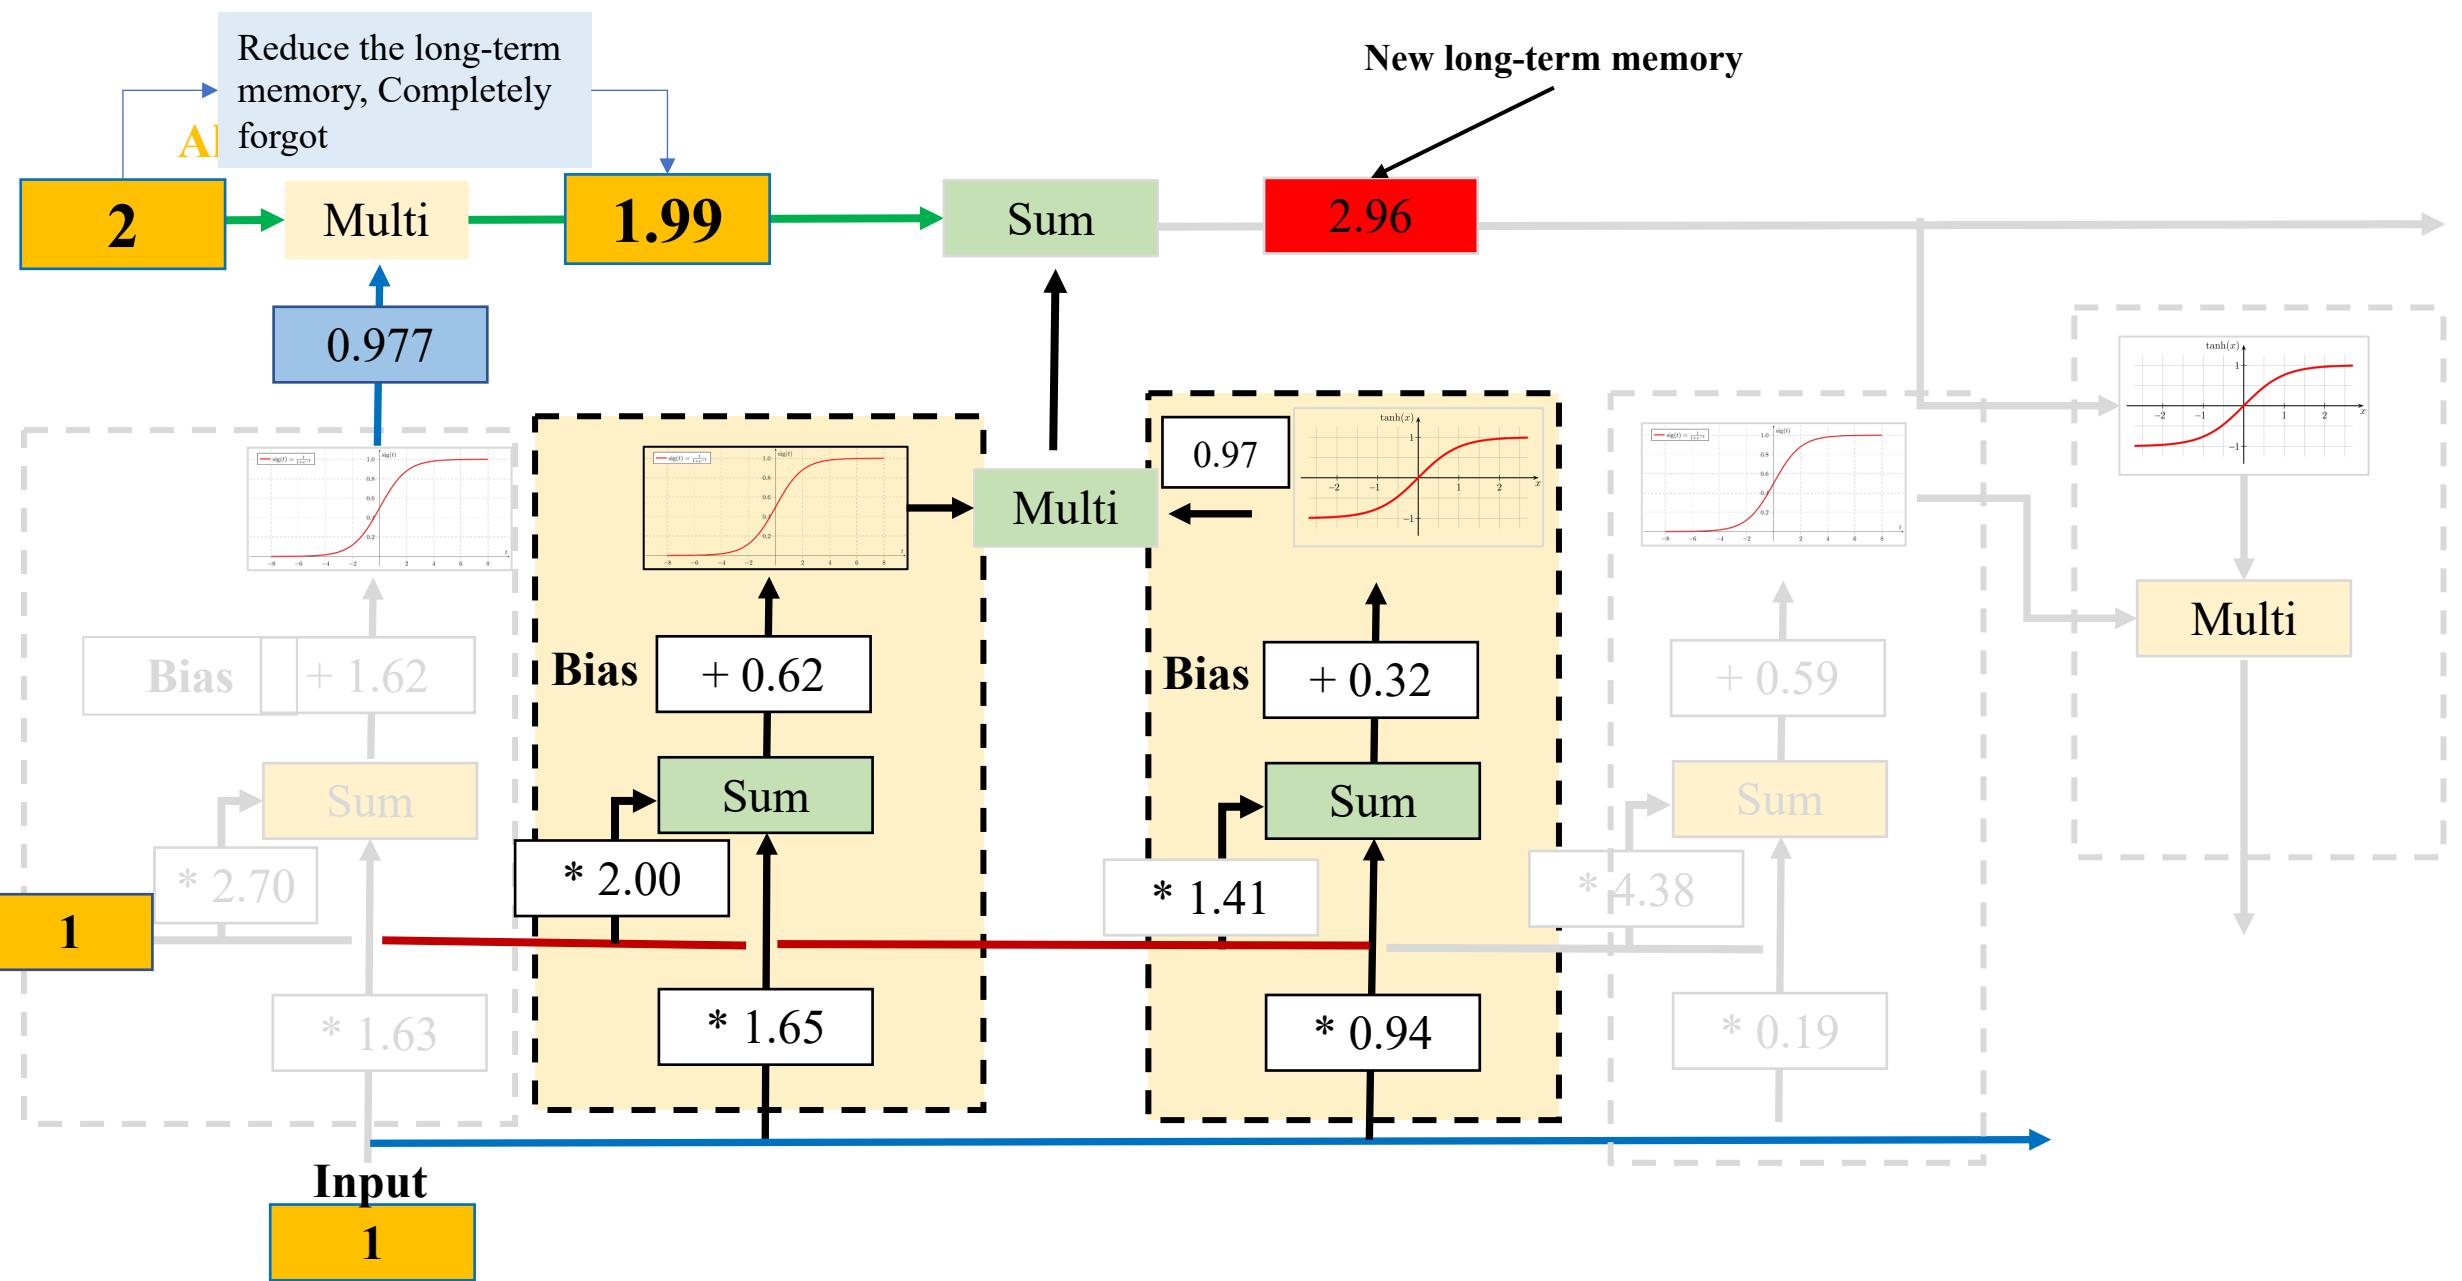
$$0.997 \times 2 = 1.99$$

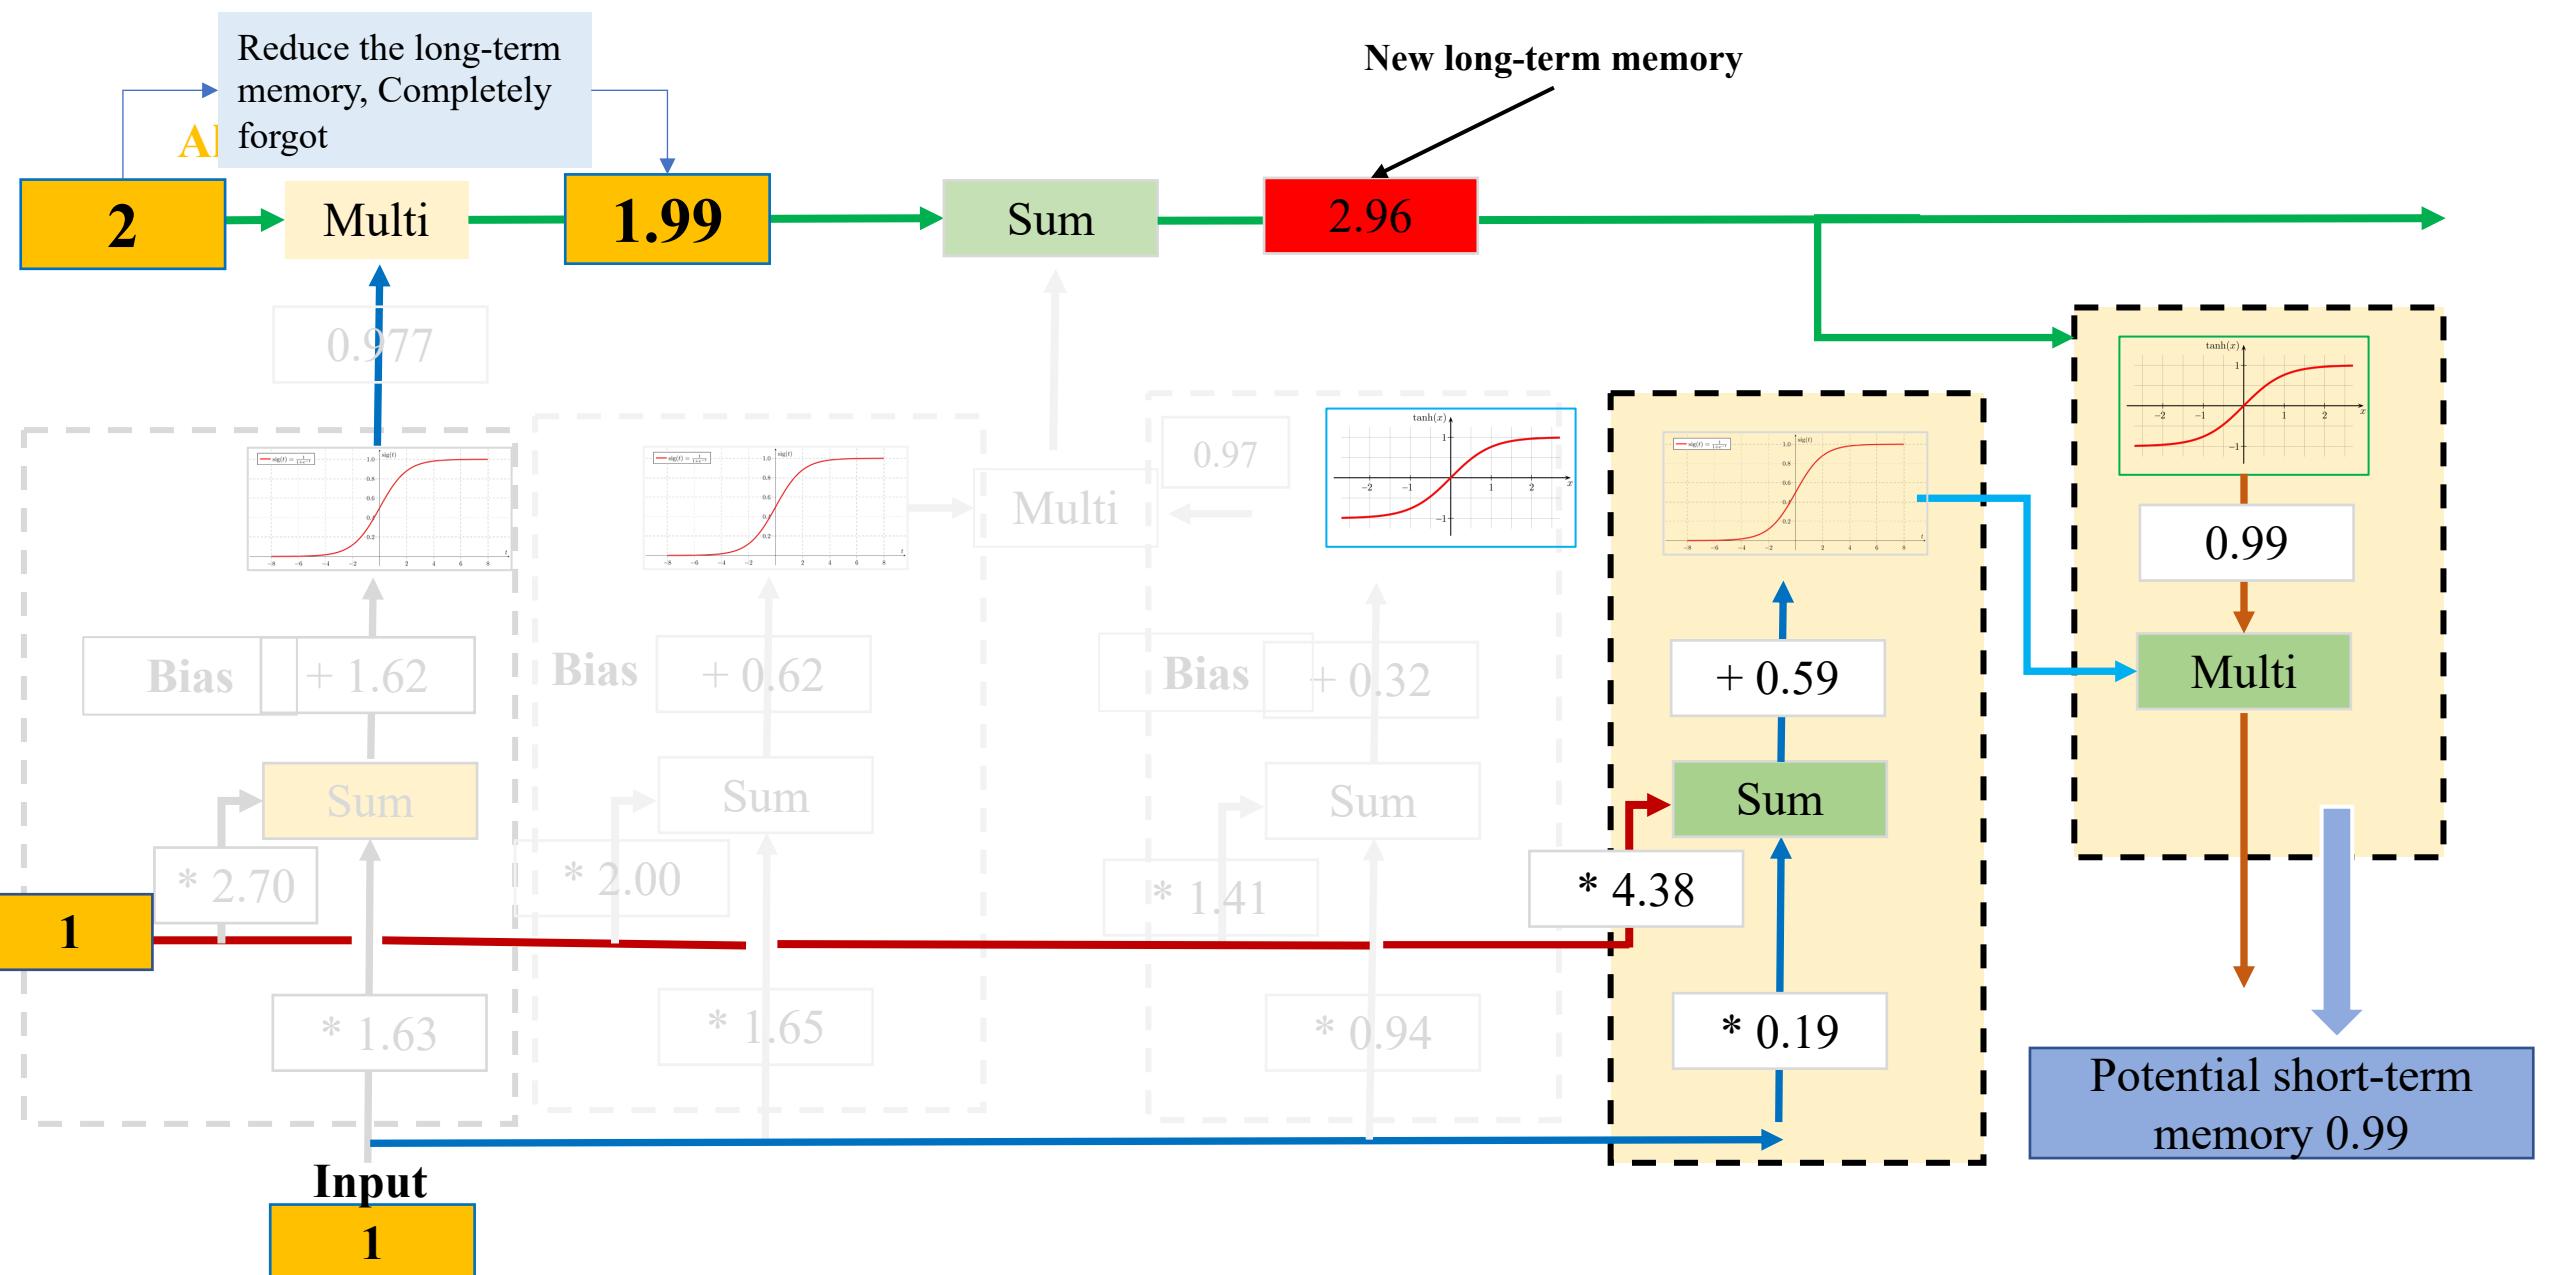


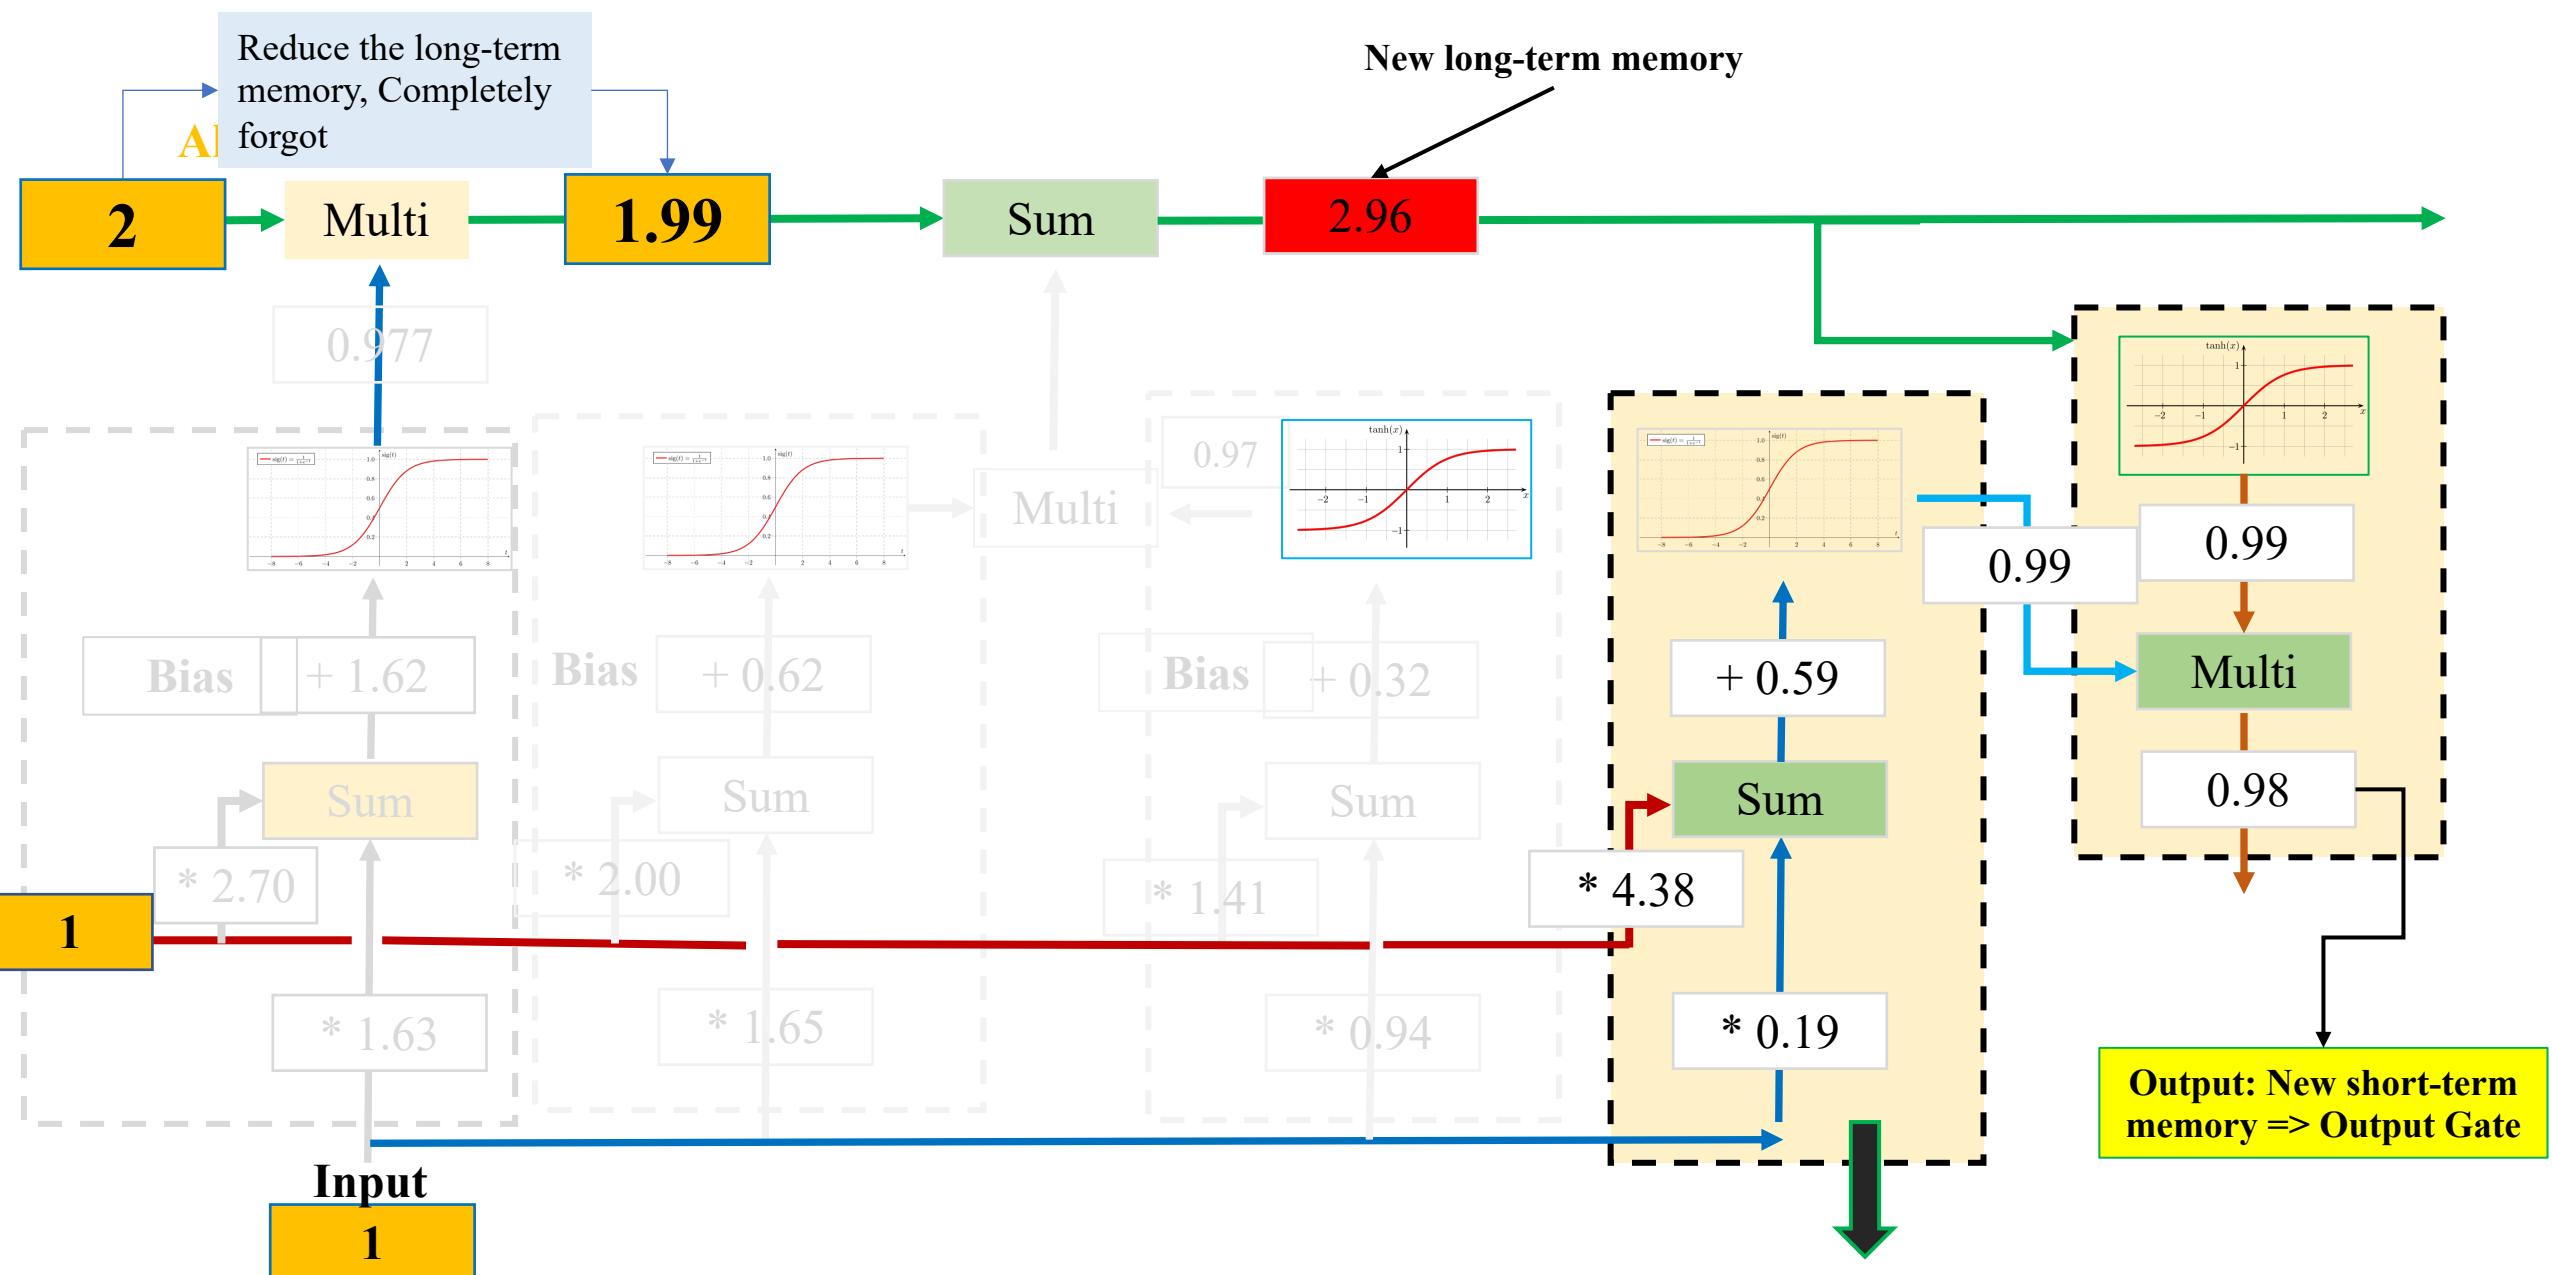




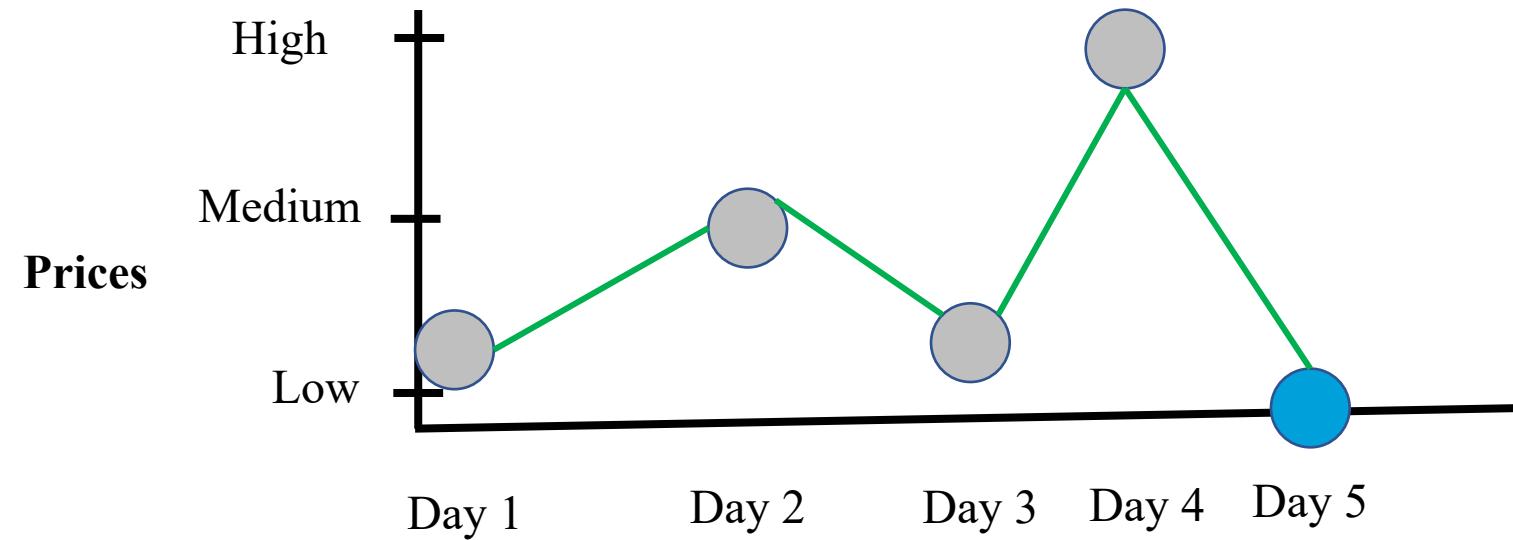








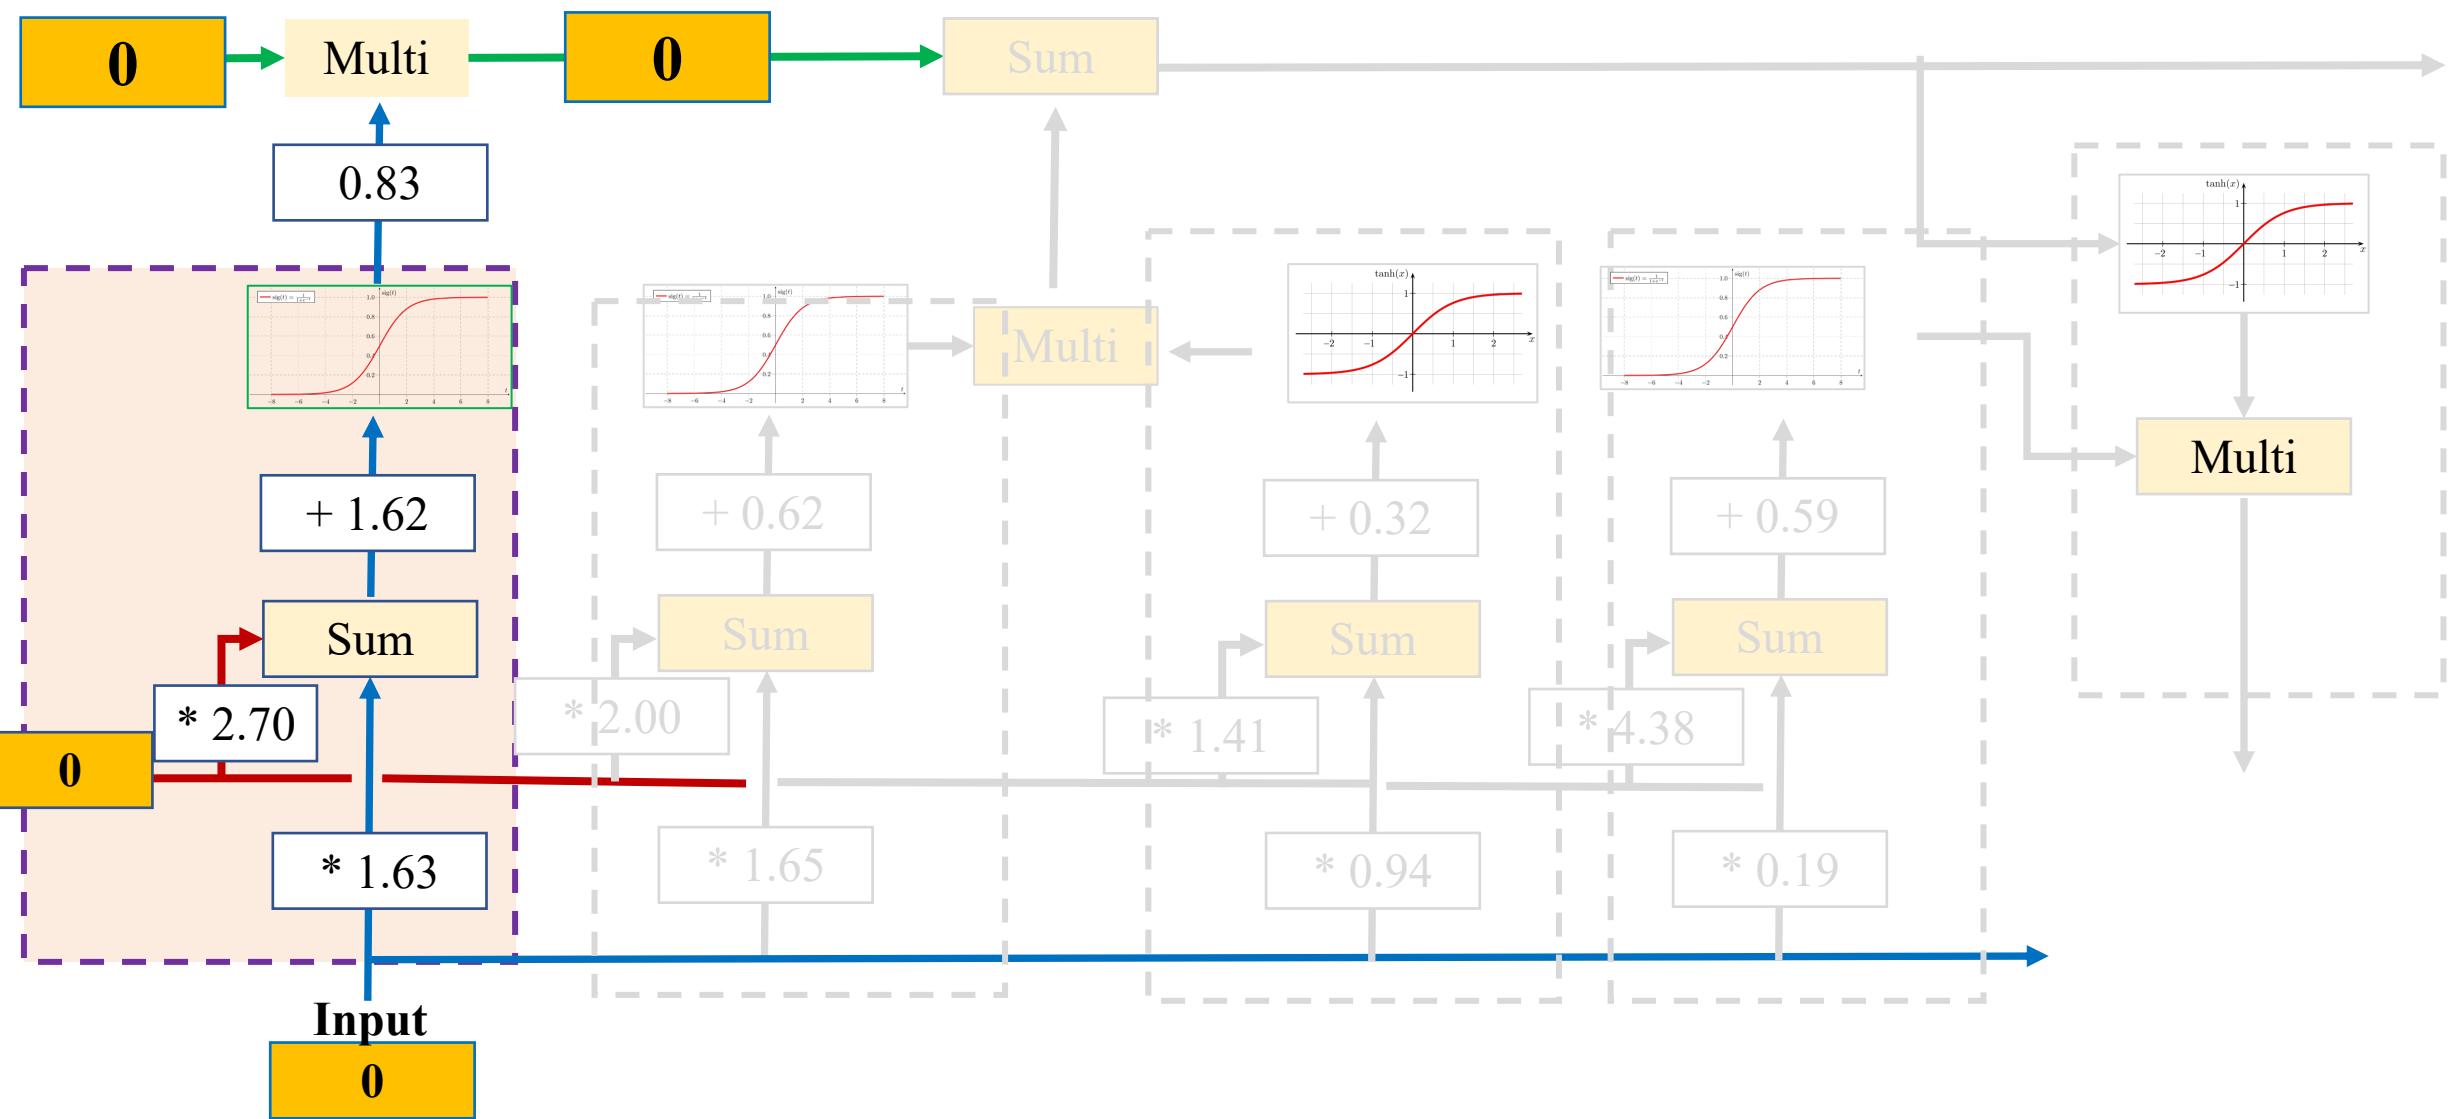
# LSTM Unit: Example

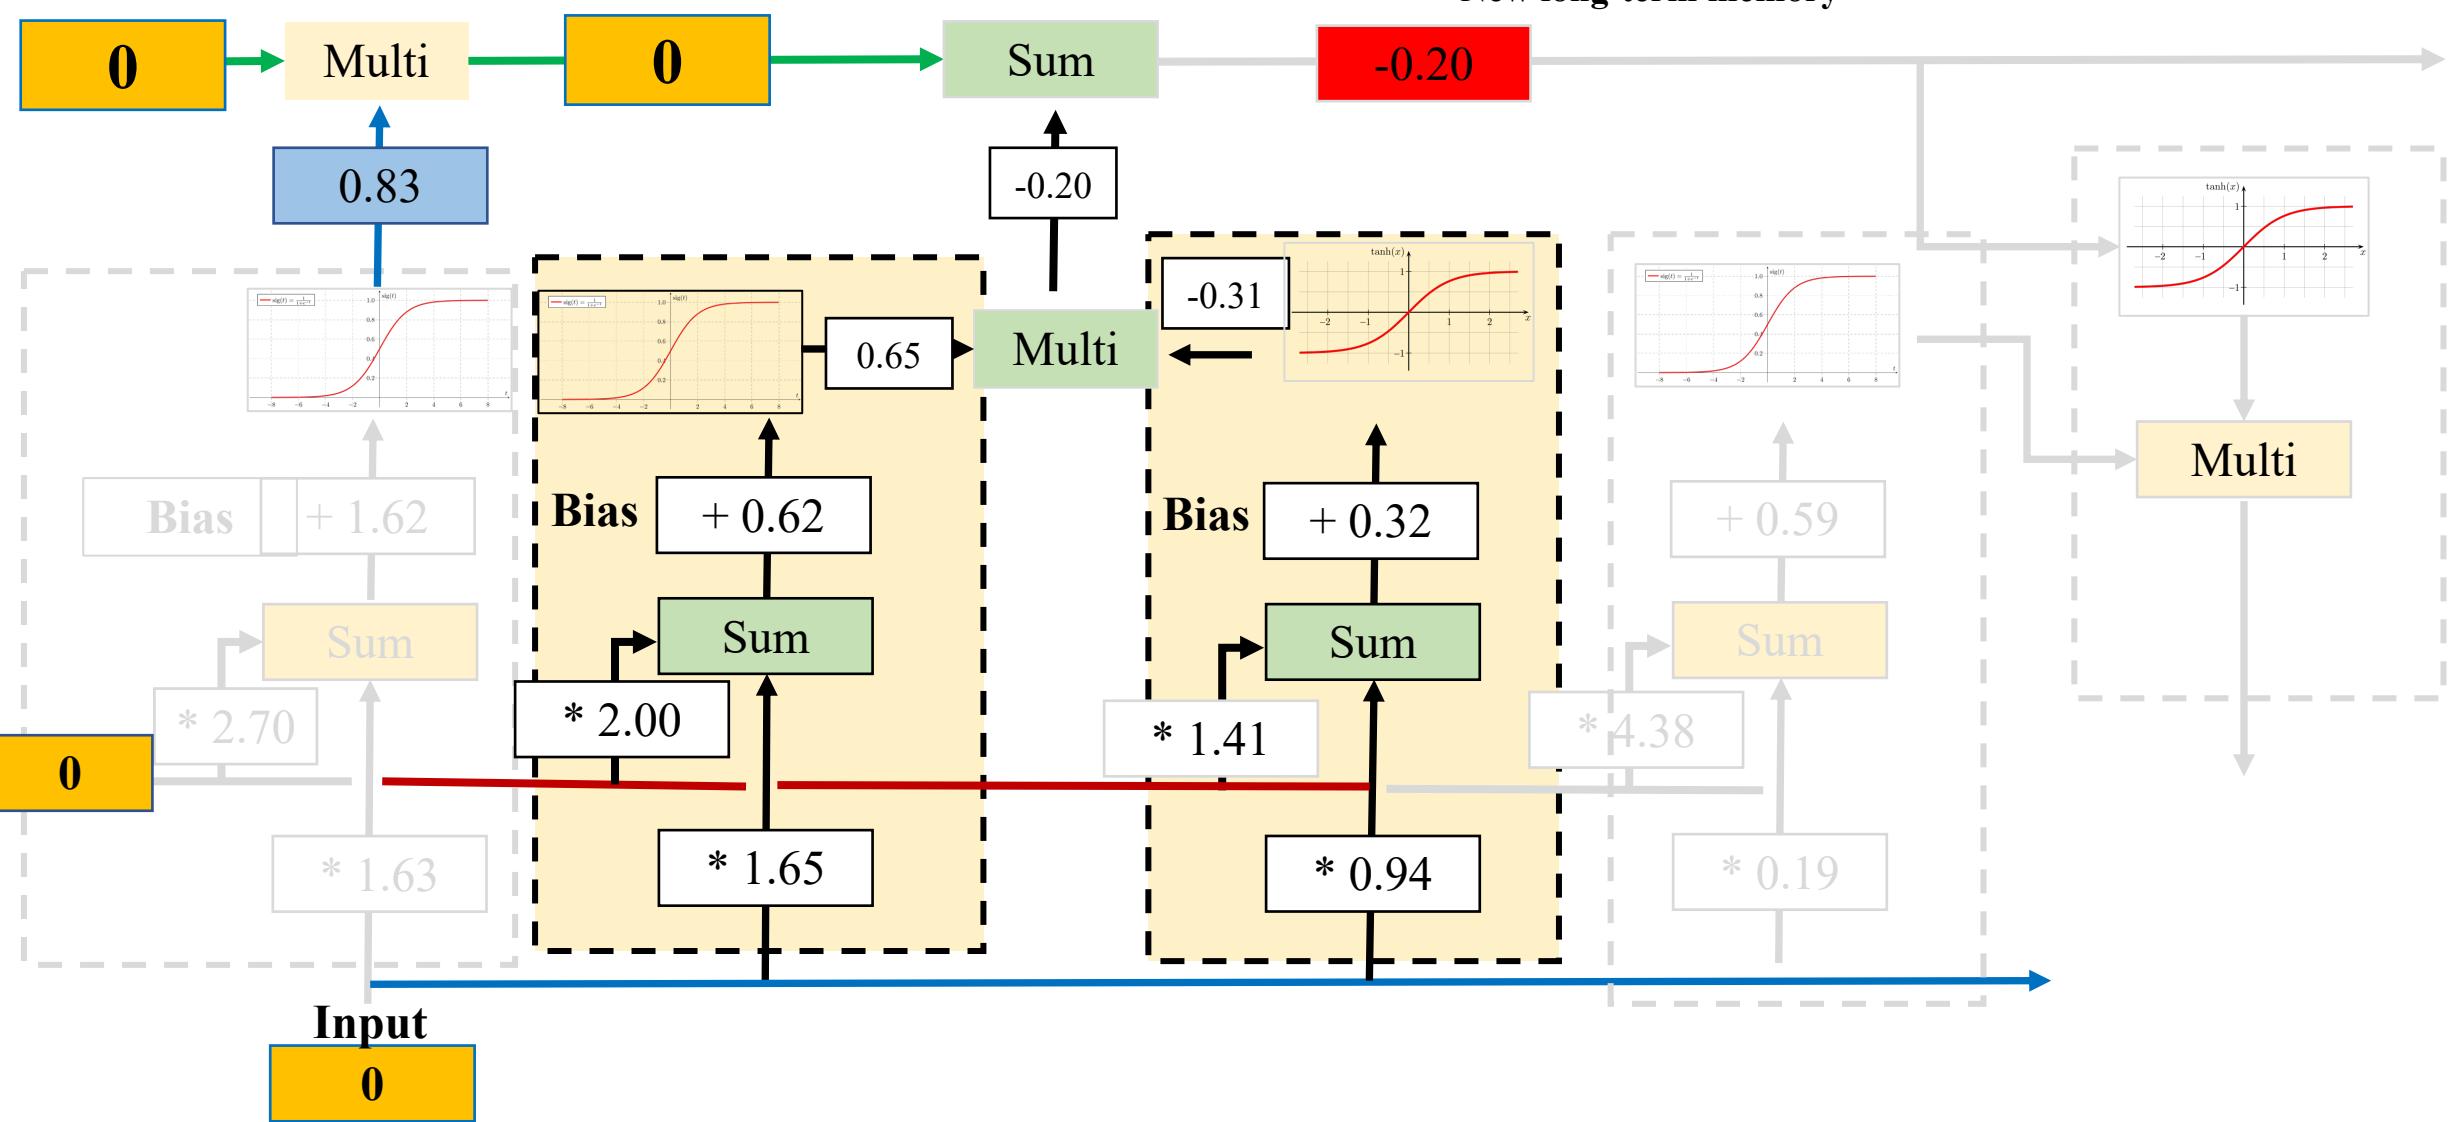


Input: day 1, day 2, day 3, day 4

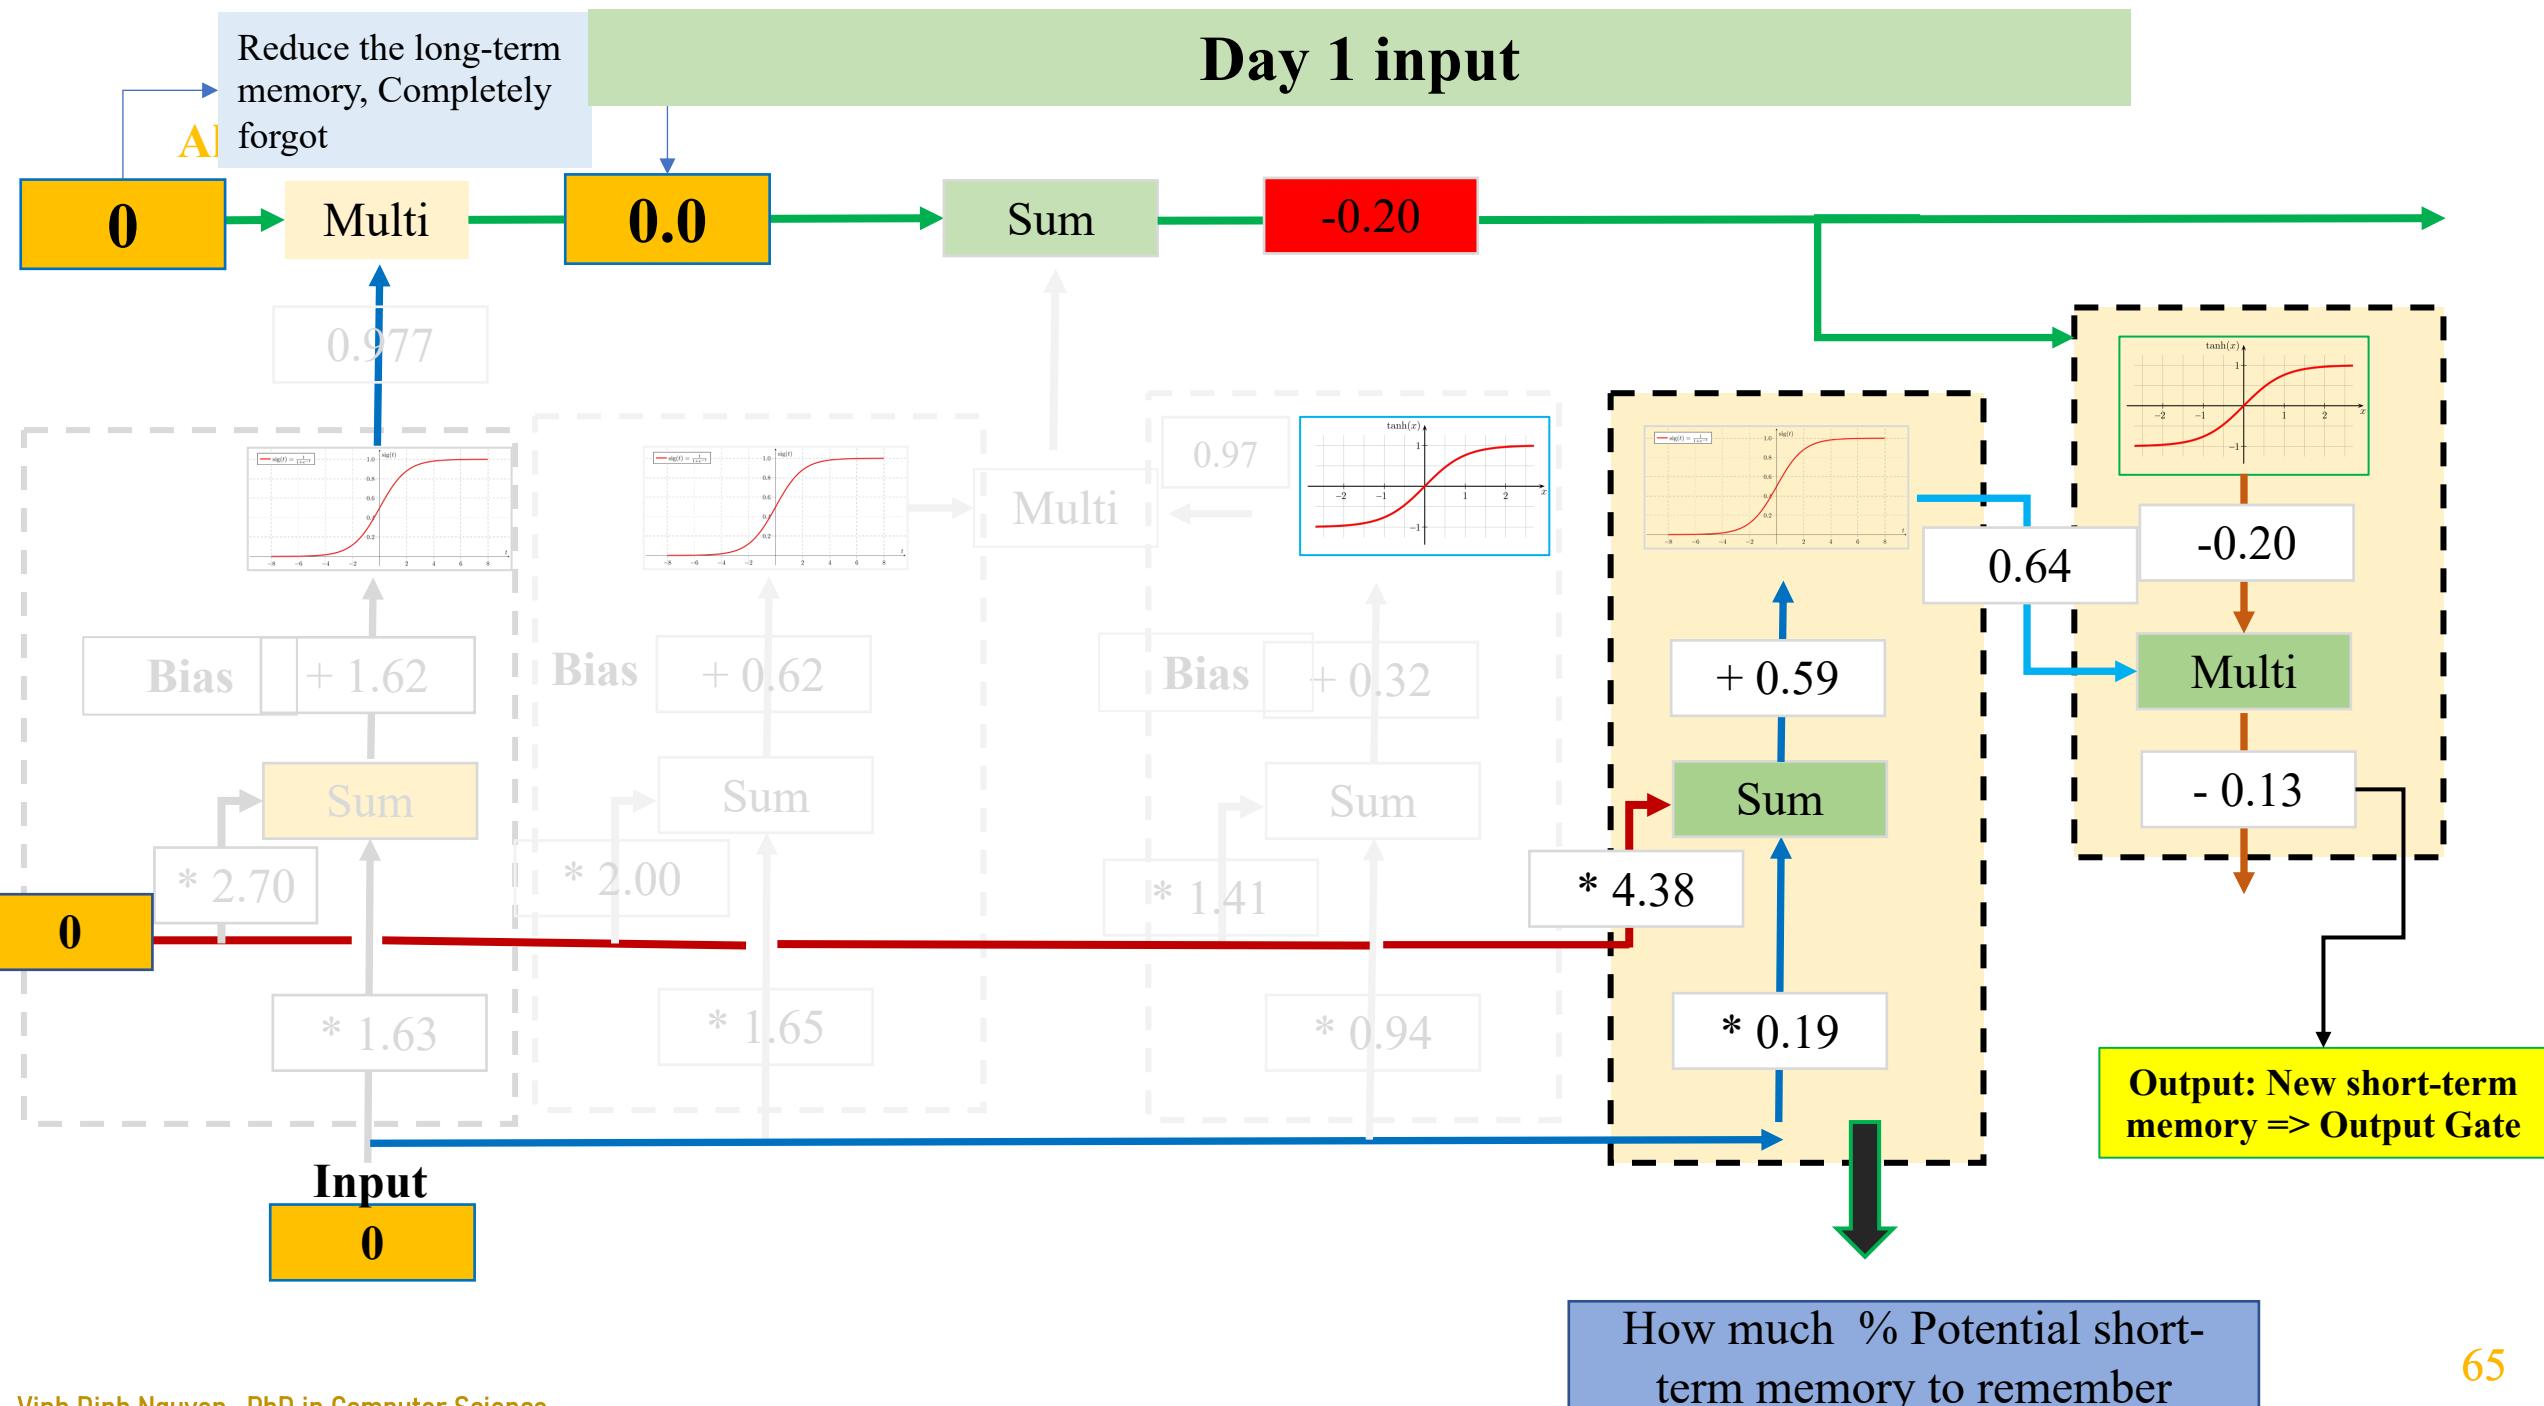
Output: day 5

Algorithm: Long-short term memory

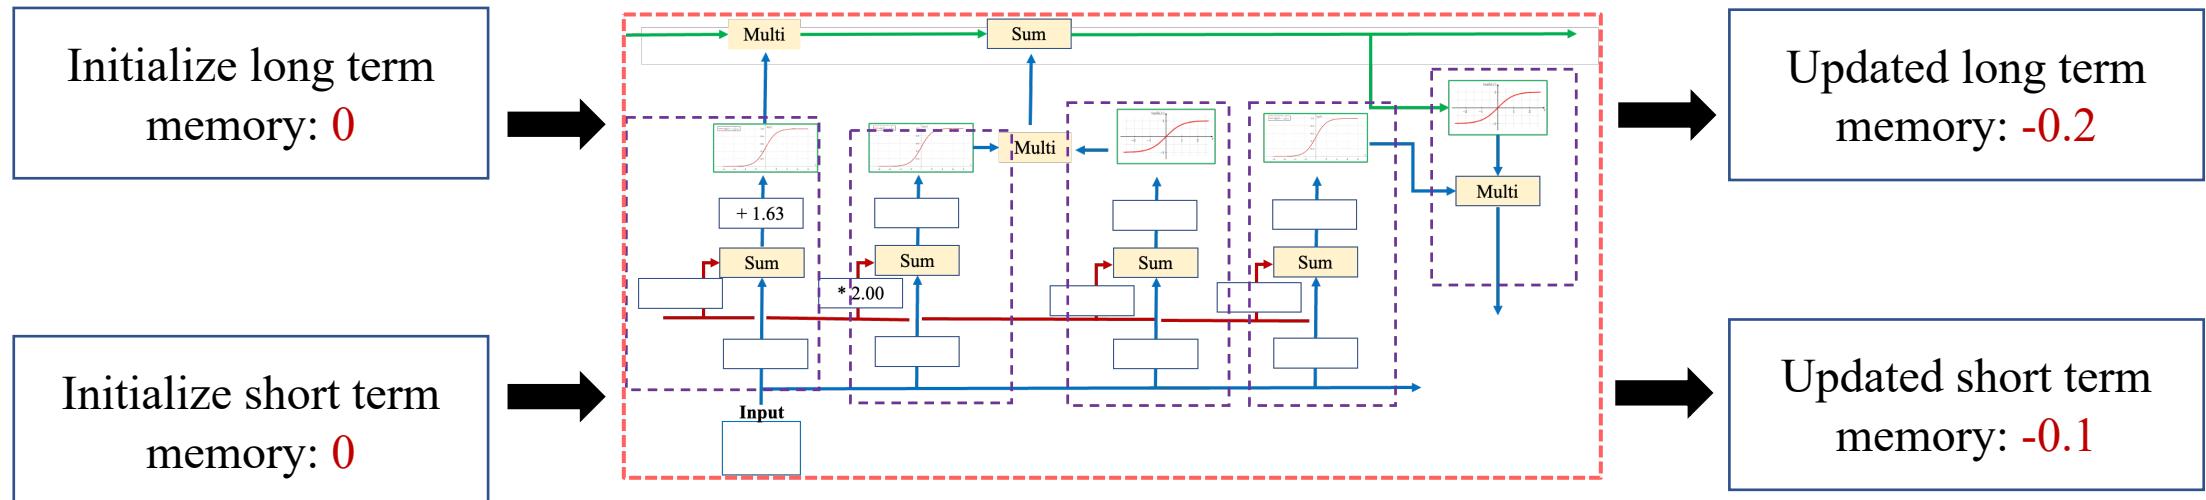




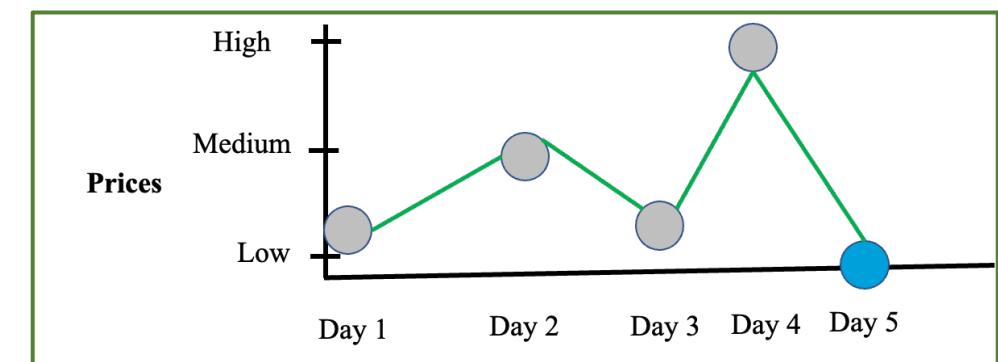
# Day 1 input

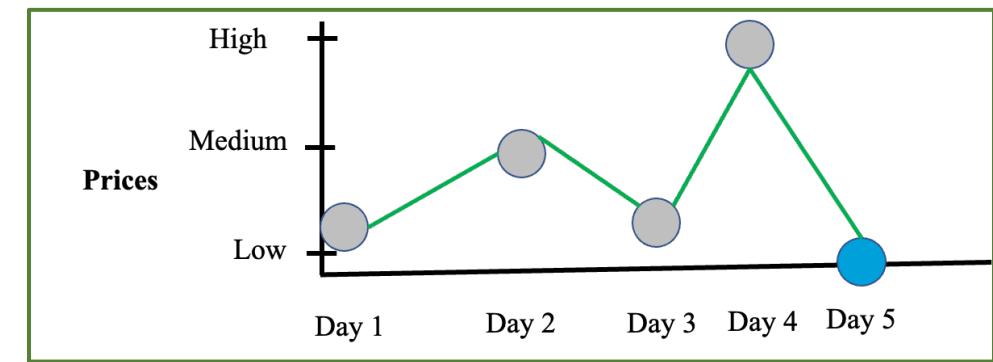
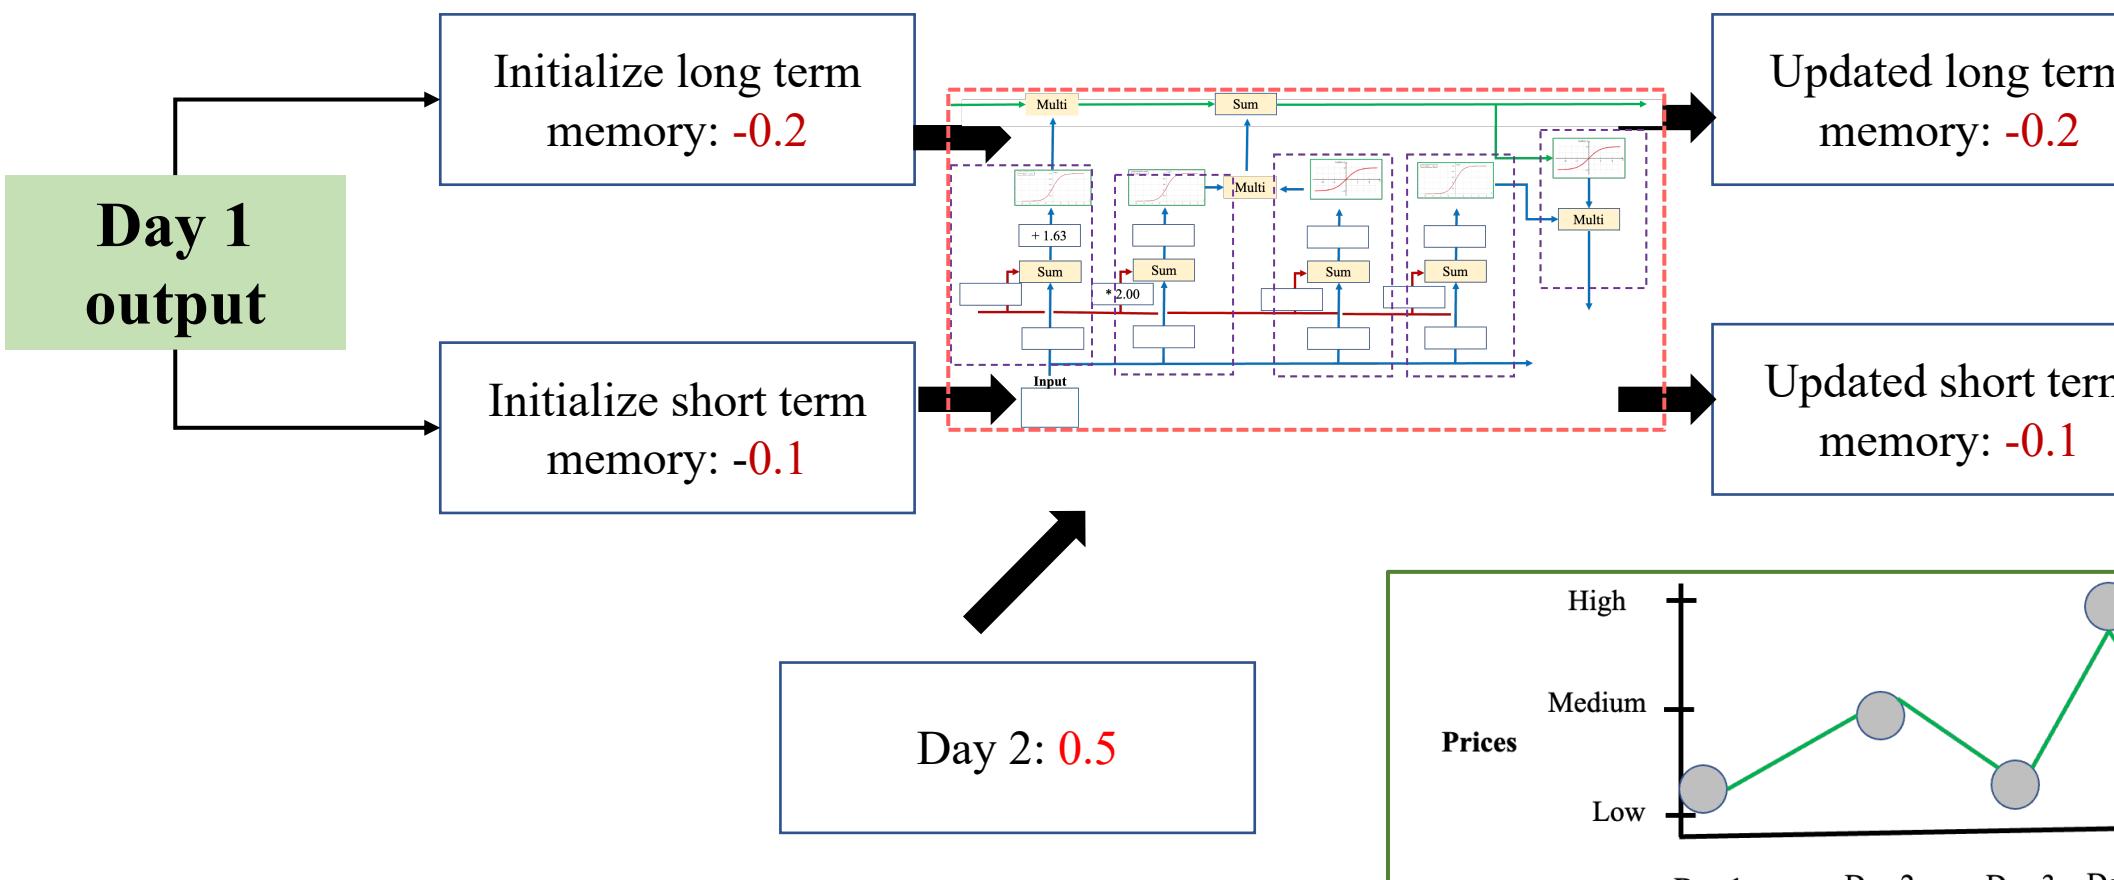


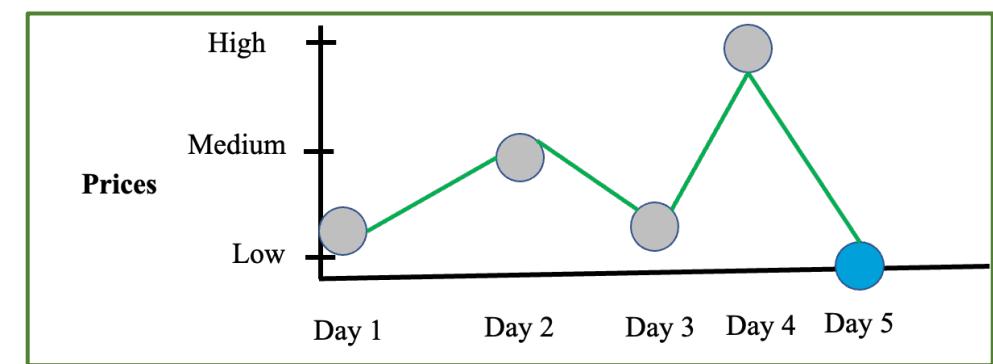
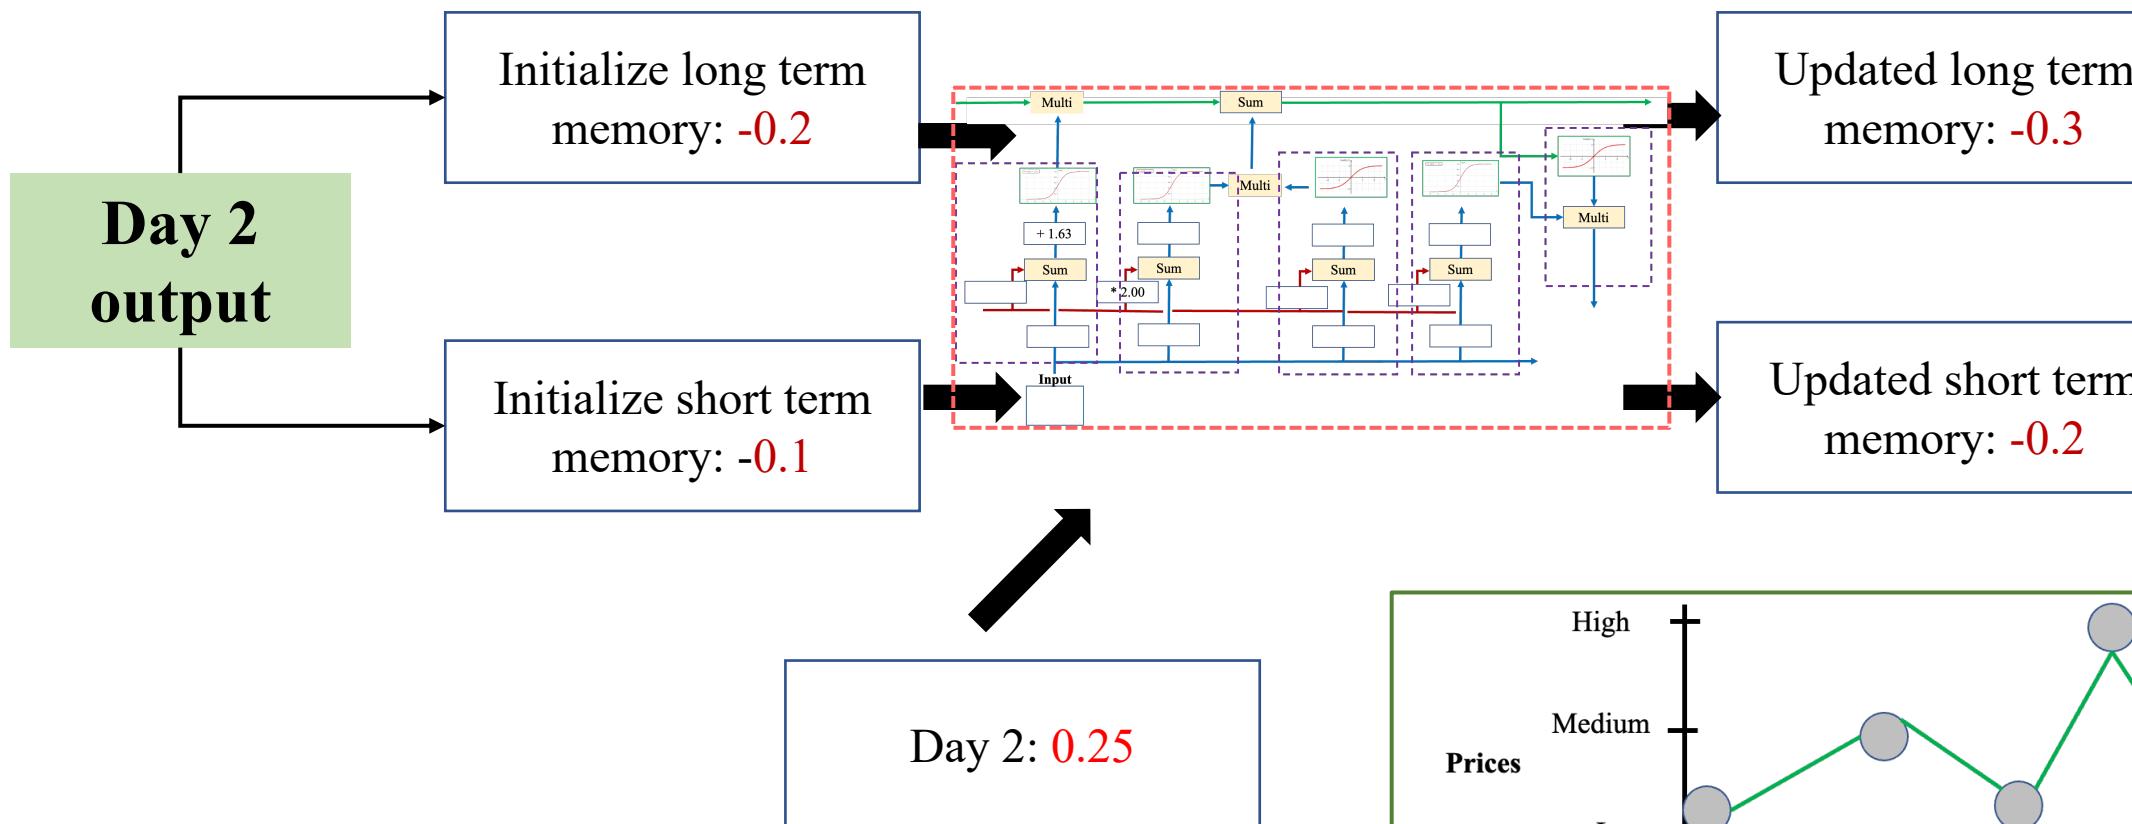
## Day 1 input

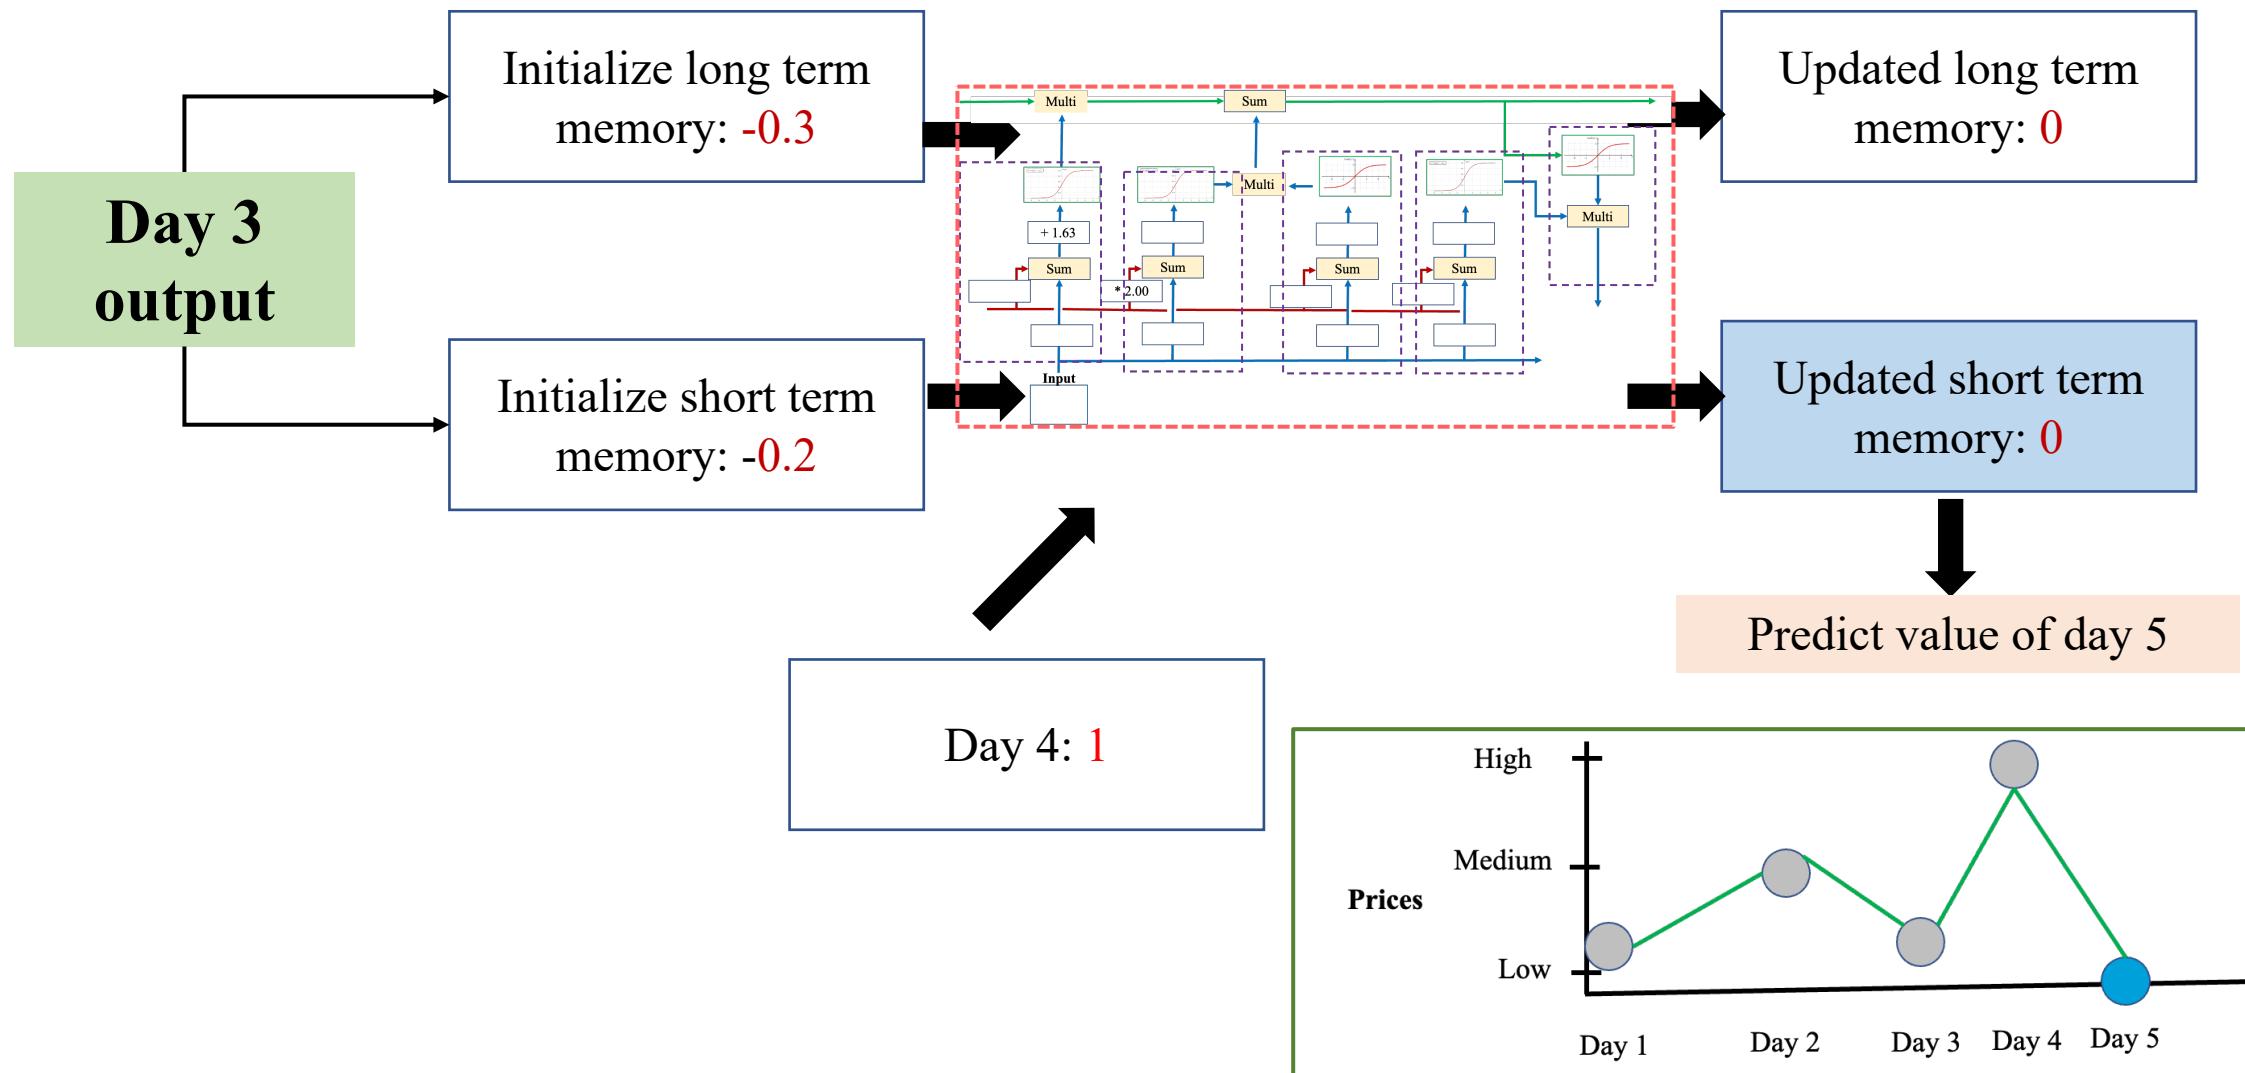


Day 1: 0









## Define LSTM Structure in Pytorch

```
class LSTM(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, ahead):
        super(LSTM, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.ahead = ahead
        self.output_size = output_size

        # LSTM Layer
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True)

        # Output layer
        self.fc = nn.Linear(hidden_size, output_size * ahead)

    def forward(self, x):
        # Initialize hidden and cell states
        h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)
        c0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size).to(x.device)

        # Forward propagate LSTM
        out, _ = self.lstm(x, (h0, c0))

        # Get the last time step's output for each sequence
        out = out[:, -1, :]

        # Pass through the linear layer and reshape
        out = self.fc(out).view(-1, self.ahead, self.output_size)
        return out
```

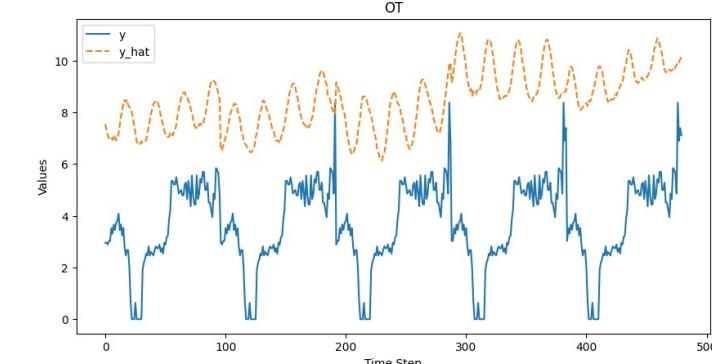
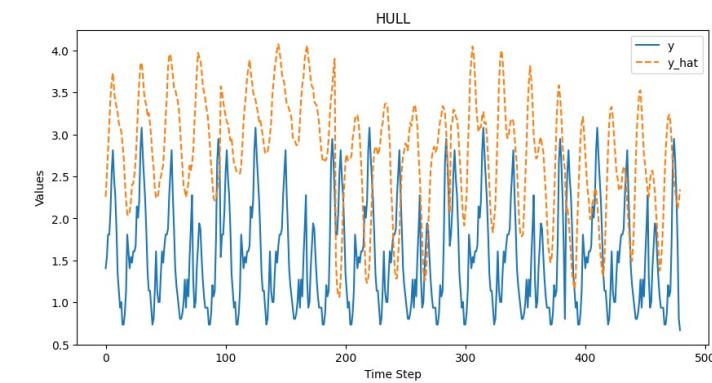
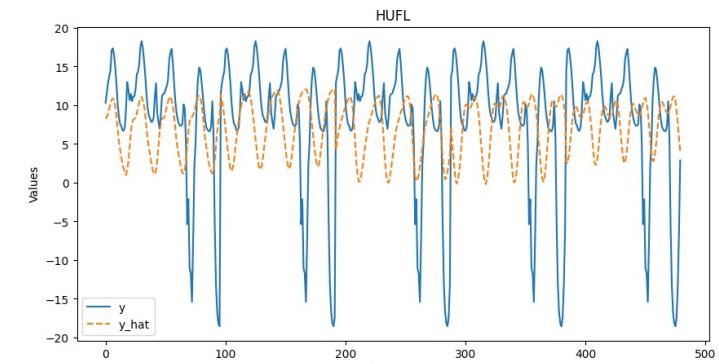
## Configuration

▶ input\_size = 336  
label\_size = 96  
offset = 1  
train\_size = 0.7  
val\_size = 0.1  
num\_epochs = 1\_000\_000  
patience = 20  
learning\_rate = 0.001  
hidden\_size = 64  
num\_layers = 2  
ele = 5  
target\_name = 'OT'  
date\_column = 'date'  
file\_path = 'ETTh1.csv'  
plot\_dir = 'plots'  
weight\_dir = 'weights'  
results = []

## Define LSTM Structure in Pytorch

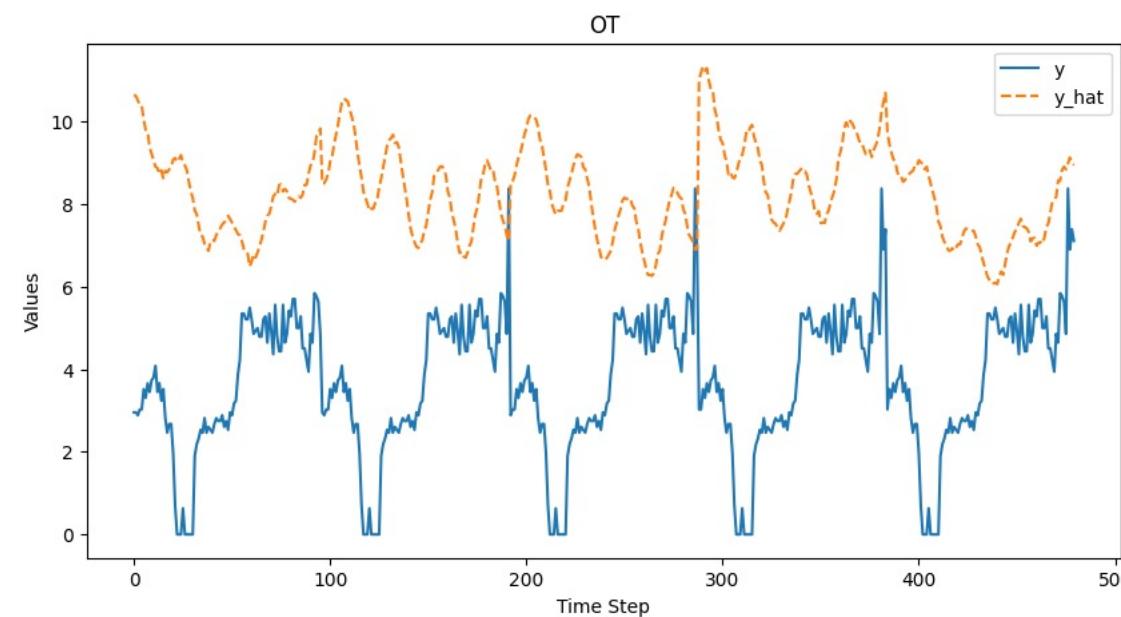
```
LSTM_multi2multi = LSTM(input_size=multi2multi_loader.in_variable,  
                        hidden_size=hidden_size,  
                        output_size=multi2multi_loader.out_variable,  
                        ahead_label_size=  
                        num_layers=num_layers)  
LSTM_multi2multi_manager = ModelManager(model=LSTM_multi2multi,  
                                         train_loader=multi2multi_loader.train_loader,  
                                         val_loader=multi2multi_loader.val_loader,  
                                         lr=learning_rate,  
                                         patience=patience)  
LSTM_multi2multi_manager.train(num_epochs=num_epochs,  
                               save_dir=os.path.join(weight_dir, sub_dir))  
results.append({  
    "Name": LSTM_multi2multi_manager.model.__class__.__name__,  
    "Type": sub_dir,  
    "MAE": LSTM_multi2multi_manager.evaluate(loader=multi2multi_loader.test_loader)  
})  
results[-1]  
  
Model saved to weights\multi2multi\best-LSTM.pth  
Epoch [1/1000000], time: 50s, loss: 3.6222, val_loss: 3.2154  
Epoch [2/1000000], time: 49s, loss: 2.7683, val_loss: 3.6213  
Epoch [3/1000000], time: 49s, loss: 2.7491, val_loss: 3.6428  
Epoch [4/1000000], time: 49s, loss: 2.5998, val_loss: 3.5903  
Epoch [5/1000000], time: 49s, loss: 2.3548, val_loss: 3.5180  
Epoch [6/1000000], time: 49s, loss: 2.2306, val_loss: 3.2928  
Model saved to weights\multi2multi\best-LSTM.pth  
Epoch [7/1000000], time: 49s, loss: 2.1372, val_loss: 3.0611  
Model saved to weights\multi2multi\best-LSTM.pth
```

## Case 1: Multivariate to Multivariate Prediction



## Define LSTM Structure in Pytorch

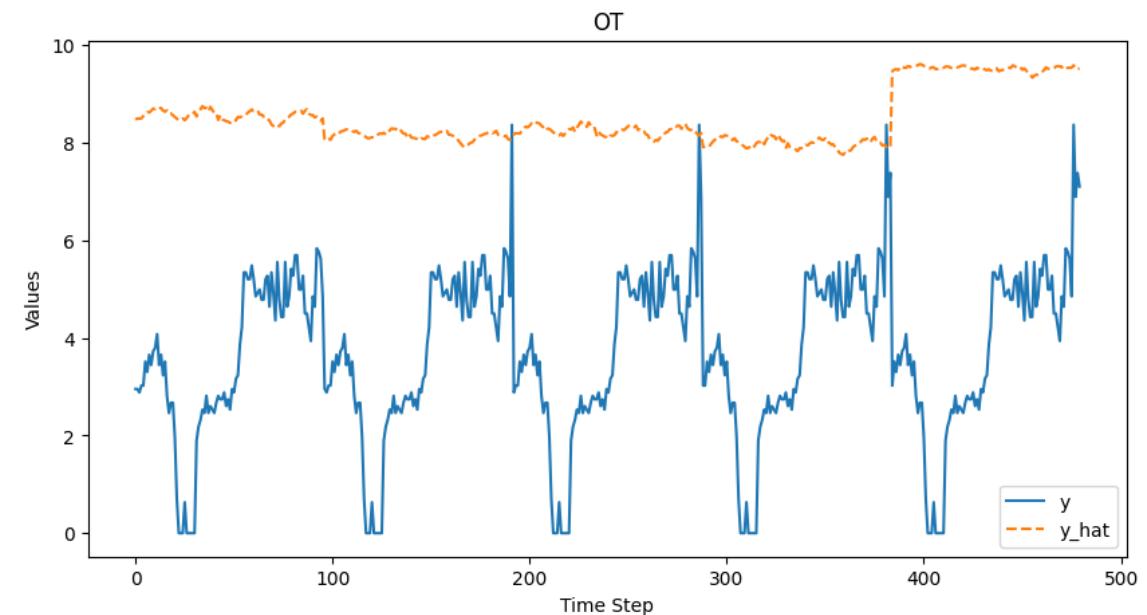
```
▶ features_type='MS'  
sub_dir = 'multi2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
multi2uni_loader = TimeSeriesDataLoader(file_path,  
                                         input_size=input_size,  
                                         label_size=label_size,  
                                         offset=offset,  
                                         train_size=train_size,  
                                         val_size=val_size,  
                                         target_name=target_name,  
                                         features_type=features_type,  
                                         date_column=date_column)  
  
👤 Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 7)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 7)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 7)  
self.y_test.shape = (2861, 96, 1)
```



## Case 2: Multivariate to Univariate Prediction

## Define LSTM Structure in Pytorch

```
▶ features_type='S'  
sub_dir = 'uni2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
uni2uni_loader = TimeSeriesDataLoader(file_path,  
                                       input_size=input_size,  
                                       label_size=label_size,  
                                       offset=offset,  
                                       train_size=train_size,  
                                       val_size=val_size,  
                                       target_name=target_name,  
                                       features_type=features_type,  
                                       date_column=date_column)  
  
👤 Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 1)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 1)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 1)  
self.y_test.shape = (2861, 96, 1)
```

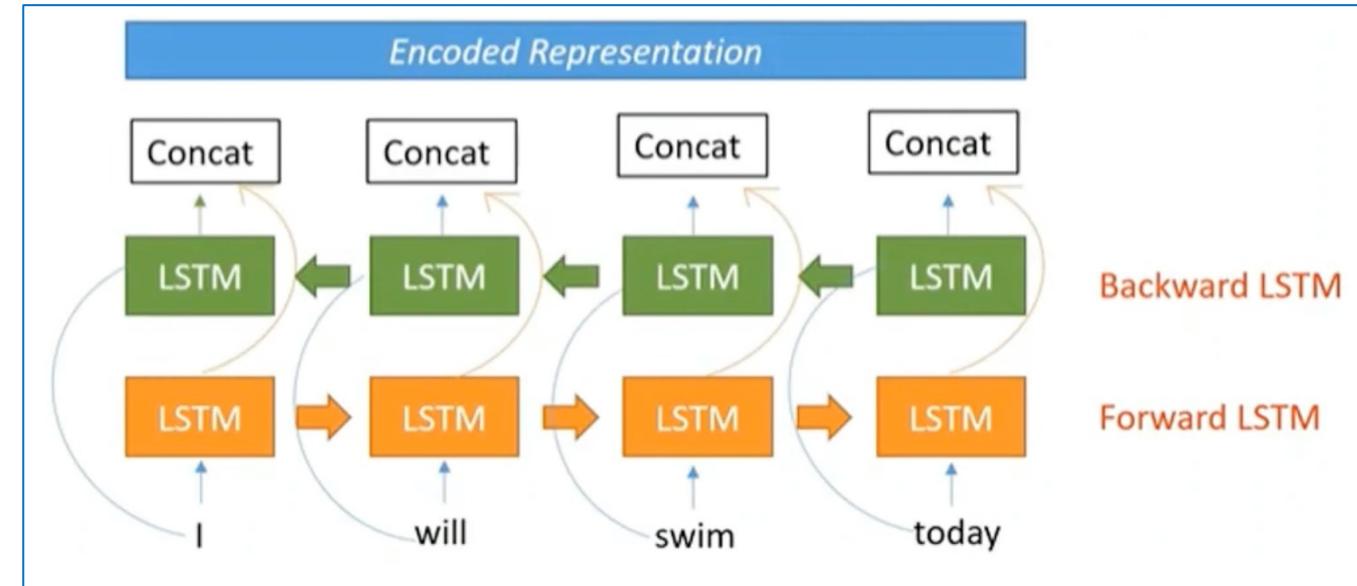
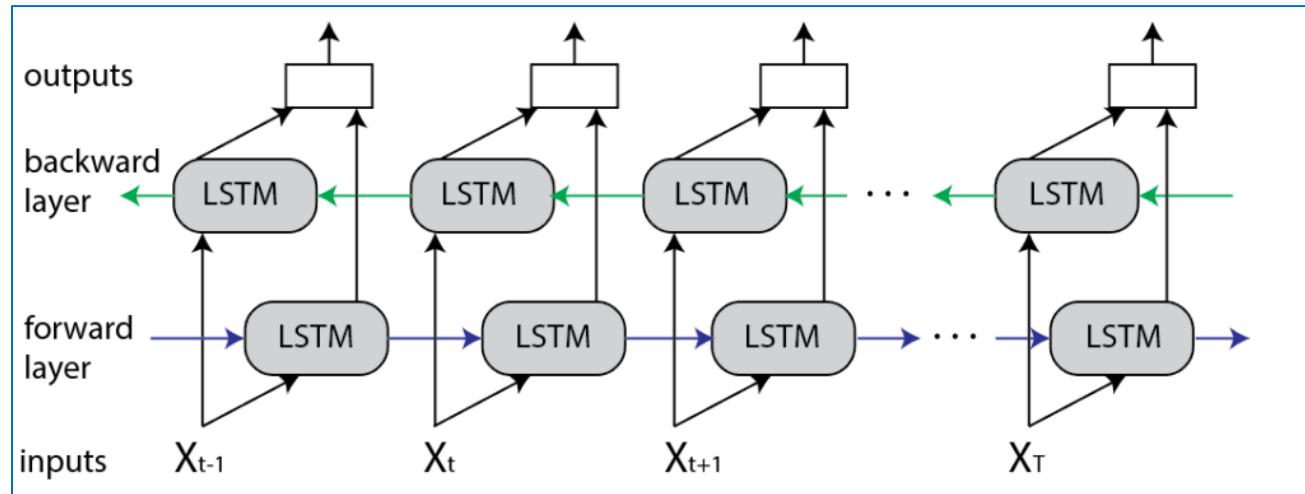


## Case 3: Univariate to Univariate Prediction

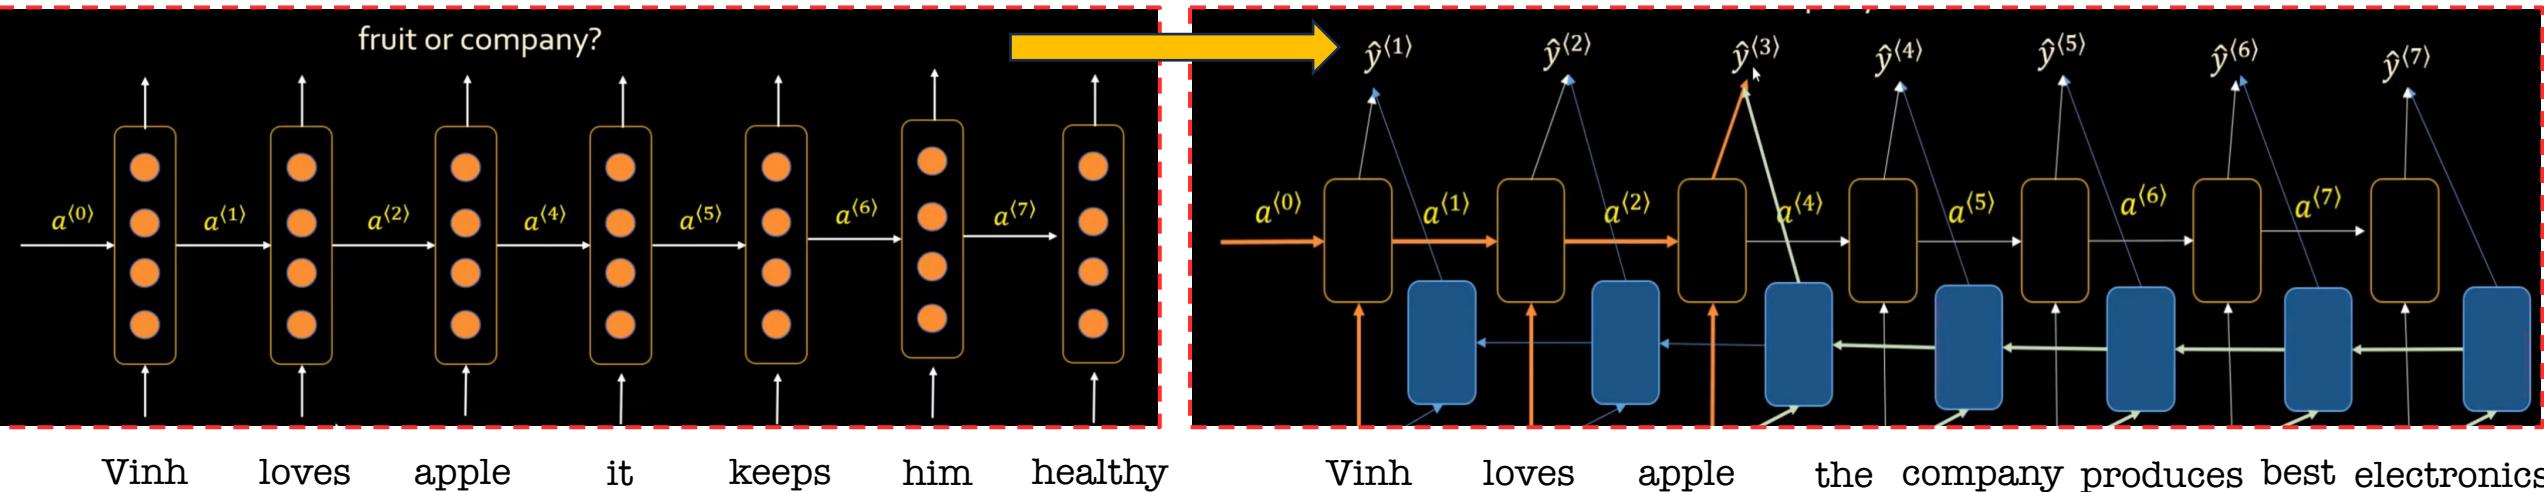
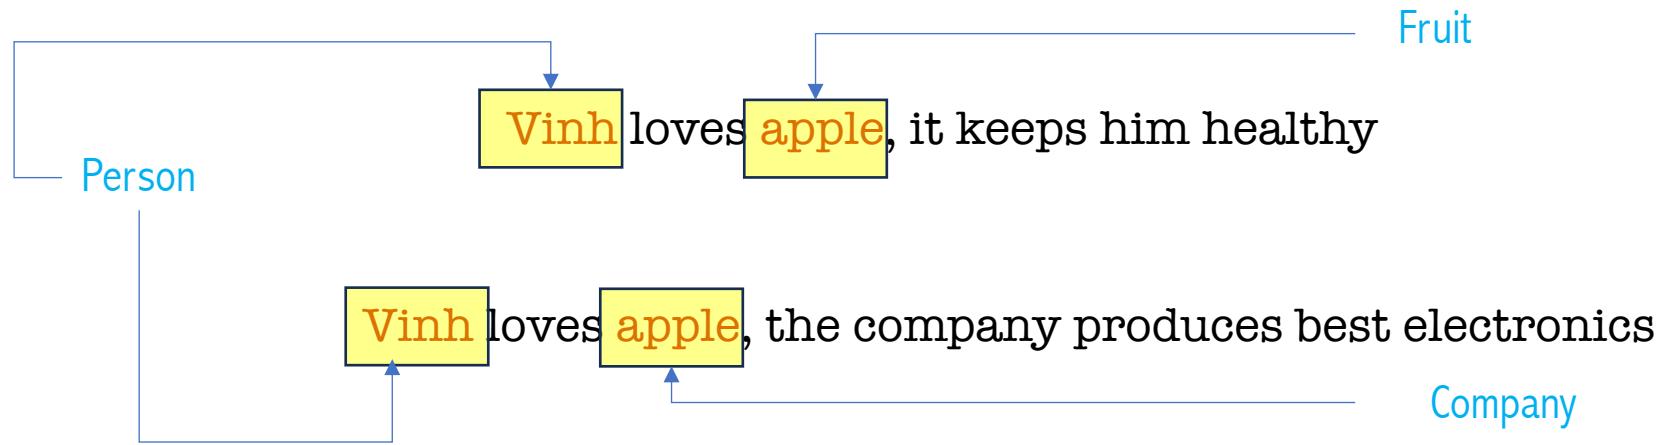
# Outline

- Multilayer Perceptron and Time-series Data Forecasting
- RNN and Time-series Data Forecasting
- LSTM and Time-series Data Forecasting
- Bi-LSTM and Time-series Data Forecasting
- XGBoost and Time-series Data Forecasting

# Bidirectional-LSTM: Simple Explanation



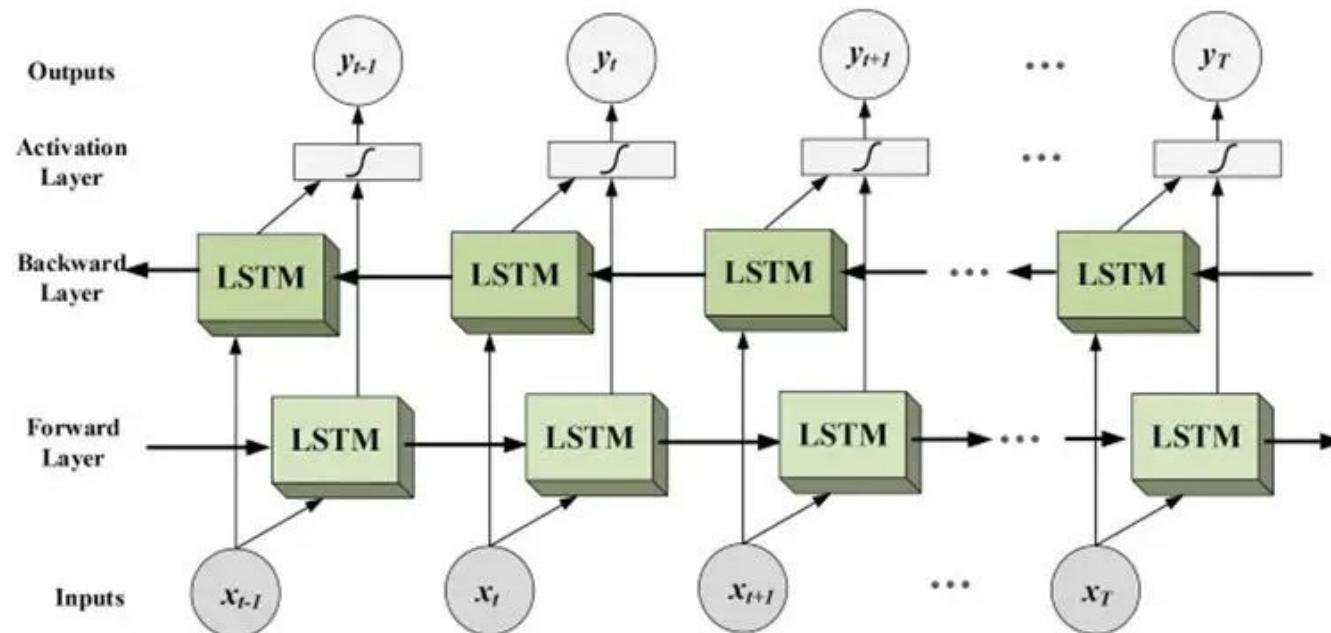
# Bidirectional-LSTM: Simple Explanation



# Bidirectional-LSTM: Simple Explanation

In the sentence “boys go to .....” we can not fill the blank space.

We have a future sentence “boys come out of school”



## Define Bi-LSTM Structure in Pytorch

```
▶ class BiLSTM(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers, ahead):
        super(BiLSTM, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers
        self.ahead = ahead
        self.output_size = output_size

        # BiLSTM Layer
        self.lstm = nn.LSTM(input_size, hidden_size, num_layers, batch_first=True, bidirectional=True)

        # Output layer
        # Output size is doubled because BiLSTM has two hidden states for each layer
        self.fc = nn.Linear(hidden_size * 2, output_size * ahead)

    def forward(self, x):
        # Initialize hidden and cell states
        h0 = torch.zeros(self.num_layers * 2, x.size(0), self.hidden_size).to(x.device) # 2 for bidirectional
        c0 = torch.zeros(self.num_layers * 2, x.size(0), self.hidden_size).to(x.device)

        # Forward propagate BiLSTM
        out, _ = self.lstm(x, (h0, c0))

        # Get the last time step's output for each sequence
        out = out[:, -1, :]

        # Pass through the linear layer and reshape
        out = self.fc(out).view(-1, self.ahead, self.output_size)
        return out
```

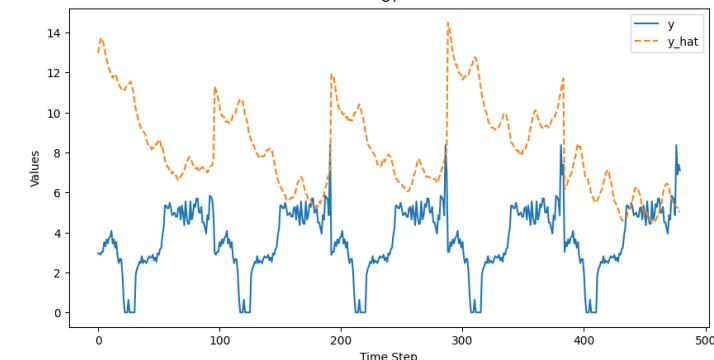
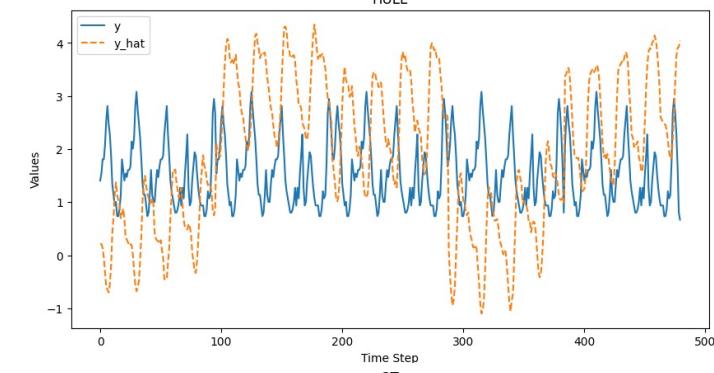
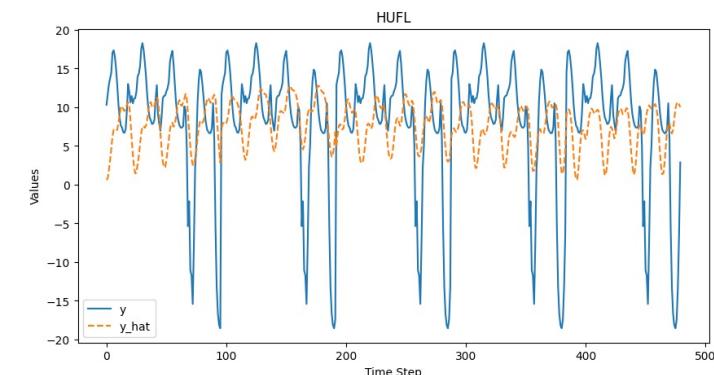
## Configuration

```
▶ input_size = 336
label_size = 96
offset = 1
train_size = 0.7
val_size = 0.1
num_epochs = 1_000_000
patience = 20
learning_rate = 0.001
hidden_size = 64
num_layers = 2
ele = 5
target_name = 'OT'
date_column = 'date'
file_path = 'ETTh1.csv'
plot_dir = 'plots'
weight_dir = 'weights'
results = []
```

## Define Bi-LSTM Structure in Pytorch

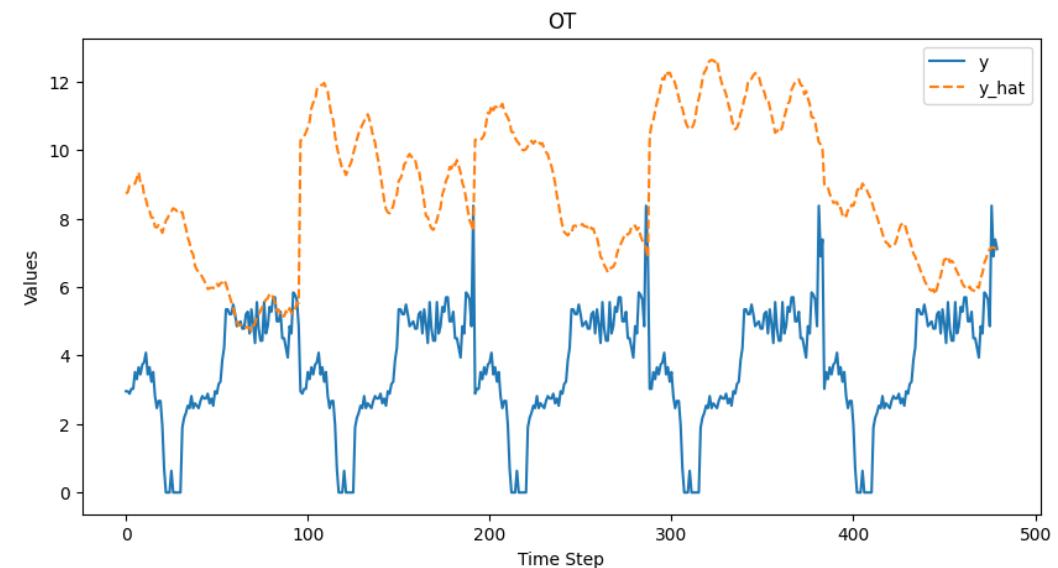
```
▶ BiLSTM_multi2multi = BiLSTM(input_size=multi2multi_loader.in_variable,  
                               hidden_size=hidden_size,  
                               output_size=multi2multi_loader.out_variable,  
                               ahead=label_size,  
                               num_layers=num_layers)  
BiLSTM_multi2multi_manager = ModelManager(model=BiLSTM_multi2multi,  
                                            train_loader=multi2multi_loader.train_loader,  
                                            val_loader=multi2multi_loader.val_loader,  
                                            lr=learning_rate,  
                                            patience=patience)  
BiLSTM_multi2multi_manager.train(num_epochs=num_epochs,  
                                  save_dir=os.path.join(weight_dir, sub_dir))  
results.append({  
    "Name": BiLSTM_multi2multi_manager.model.__class__.__name__,  
    "Type": sub_dir,  
    "MAE": BiLSTM_multi2multi_manager.evaluate(loader=multi2multi_loader.test_loader)  
})  
results[-1]  
  
🕒 Model saved to weights\multi2multi\best-BiLSTM.pth  
Epoch [1/1000000], time: 101s, loss: 3.3011, val_loss: 3.5839  
Model saved to weights\multi2multi\best-BiLSTM.pth  
Epoch [2/1000000], time: 101s, loss: 2.5628, val_loss: 3.5039  
Model saved to weights\multi2multi\best-BiLSTM.pth  
🕒 [2/1000000], time: 101s, loss: 2.5628, val_loss: 3.5039
```

## Case 1: Multivariate to Multivariate Prediction



## Define Bi-LSTM Structure in Pytorch

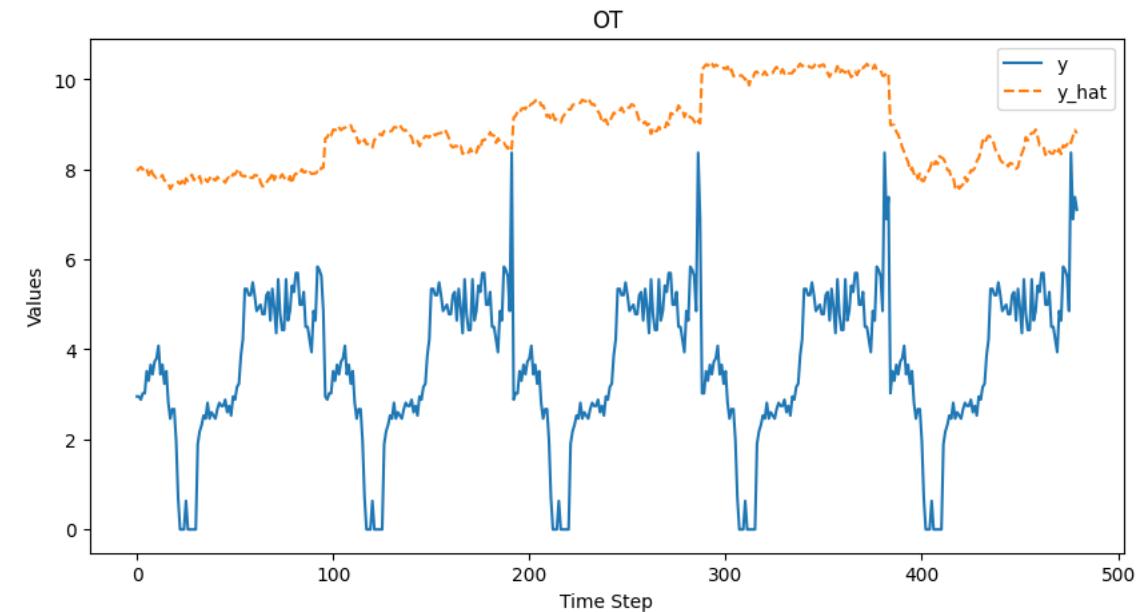
```
▶ features_type='MS'  
sub_dir = 'multi2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
multi2uni_loader = TimeSeriesDataLoader(file_path,  
                                         input_size=input_size,  
                                         label_size=label_size,  
                                         offset=offset,  
                                         train_size=train_size,  
                                         val_size=val_size,  
                                         target_name=target_name,  
                                         features_type=features_type,  
                                         date_column=date_column)  
  
👤 Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 7)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 7)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 7)  
self.y_test.shape = (2861, 96, 1)
```



## Case 2: Multivariate to Univariate Prediction

## Define Bi-LSTM Structure in Pytorch

```
▶ features_type='S'  
sub_dir = 'uni2uni'  
os.makedirs(os.path.join(weight_dir, sub_dir), exist_ok=True)  
uni2uni_loader = TimeSeriesDataLoader(file_path,  
                                       input_size=input_size,  
                                       label_size=label_size,  
                                       offset=offset,  
                                       train_size=train_size,  
                                       val_size=val_size,  
                                       target_name=target_name,  
                                       features_type=features_type,  
                                       date_column=date_column)  
  
⌚ Offset will be change from 1 to 96  
self.X_train.shape = (11762, 336, 1)  
self.y_train.shape = (11762, 96, 1)  
self.X_val.shape = (1309, 336, 1)  
self.y_val.shape = (1309, 96, 1)  
self.X_test.shape = (2861, 336, 1)  
self.y_test.shape = (2861, 96, 1)
```

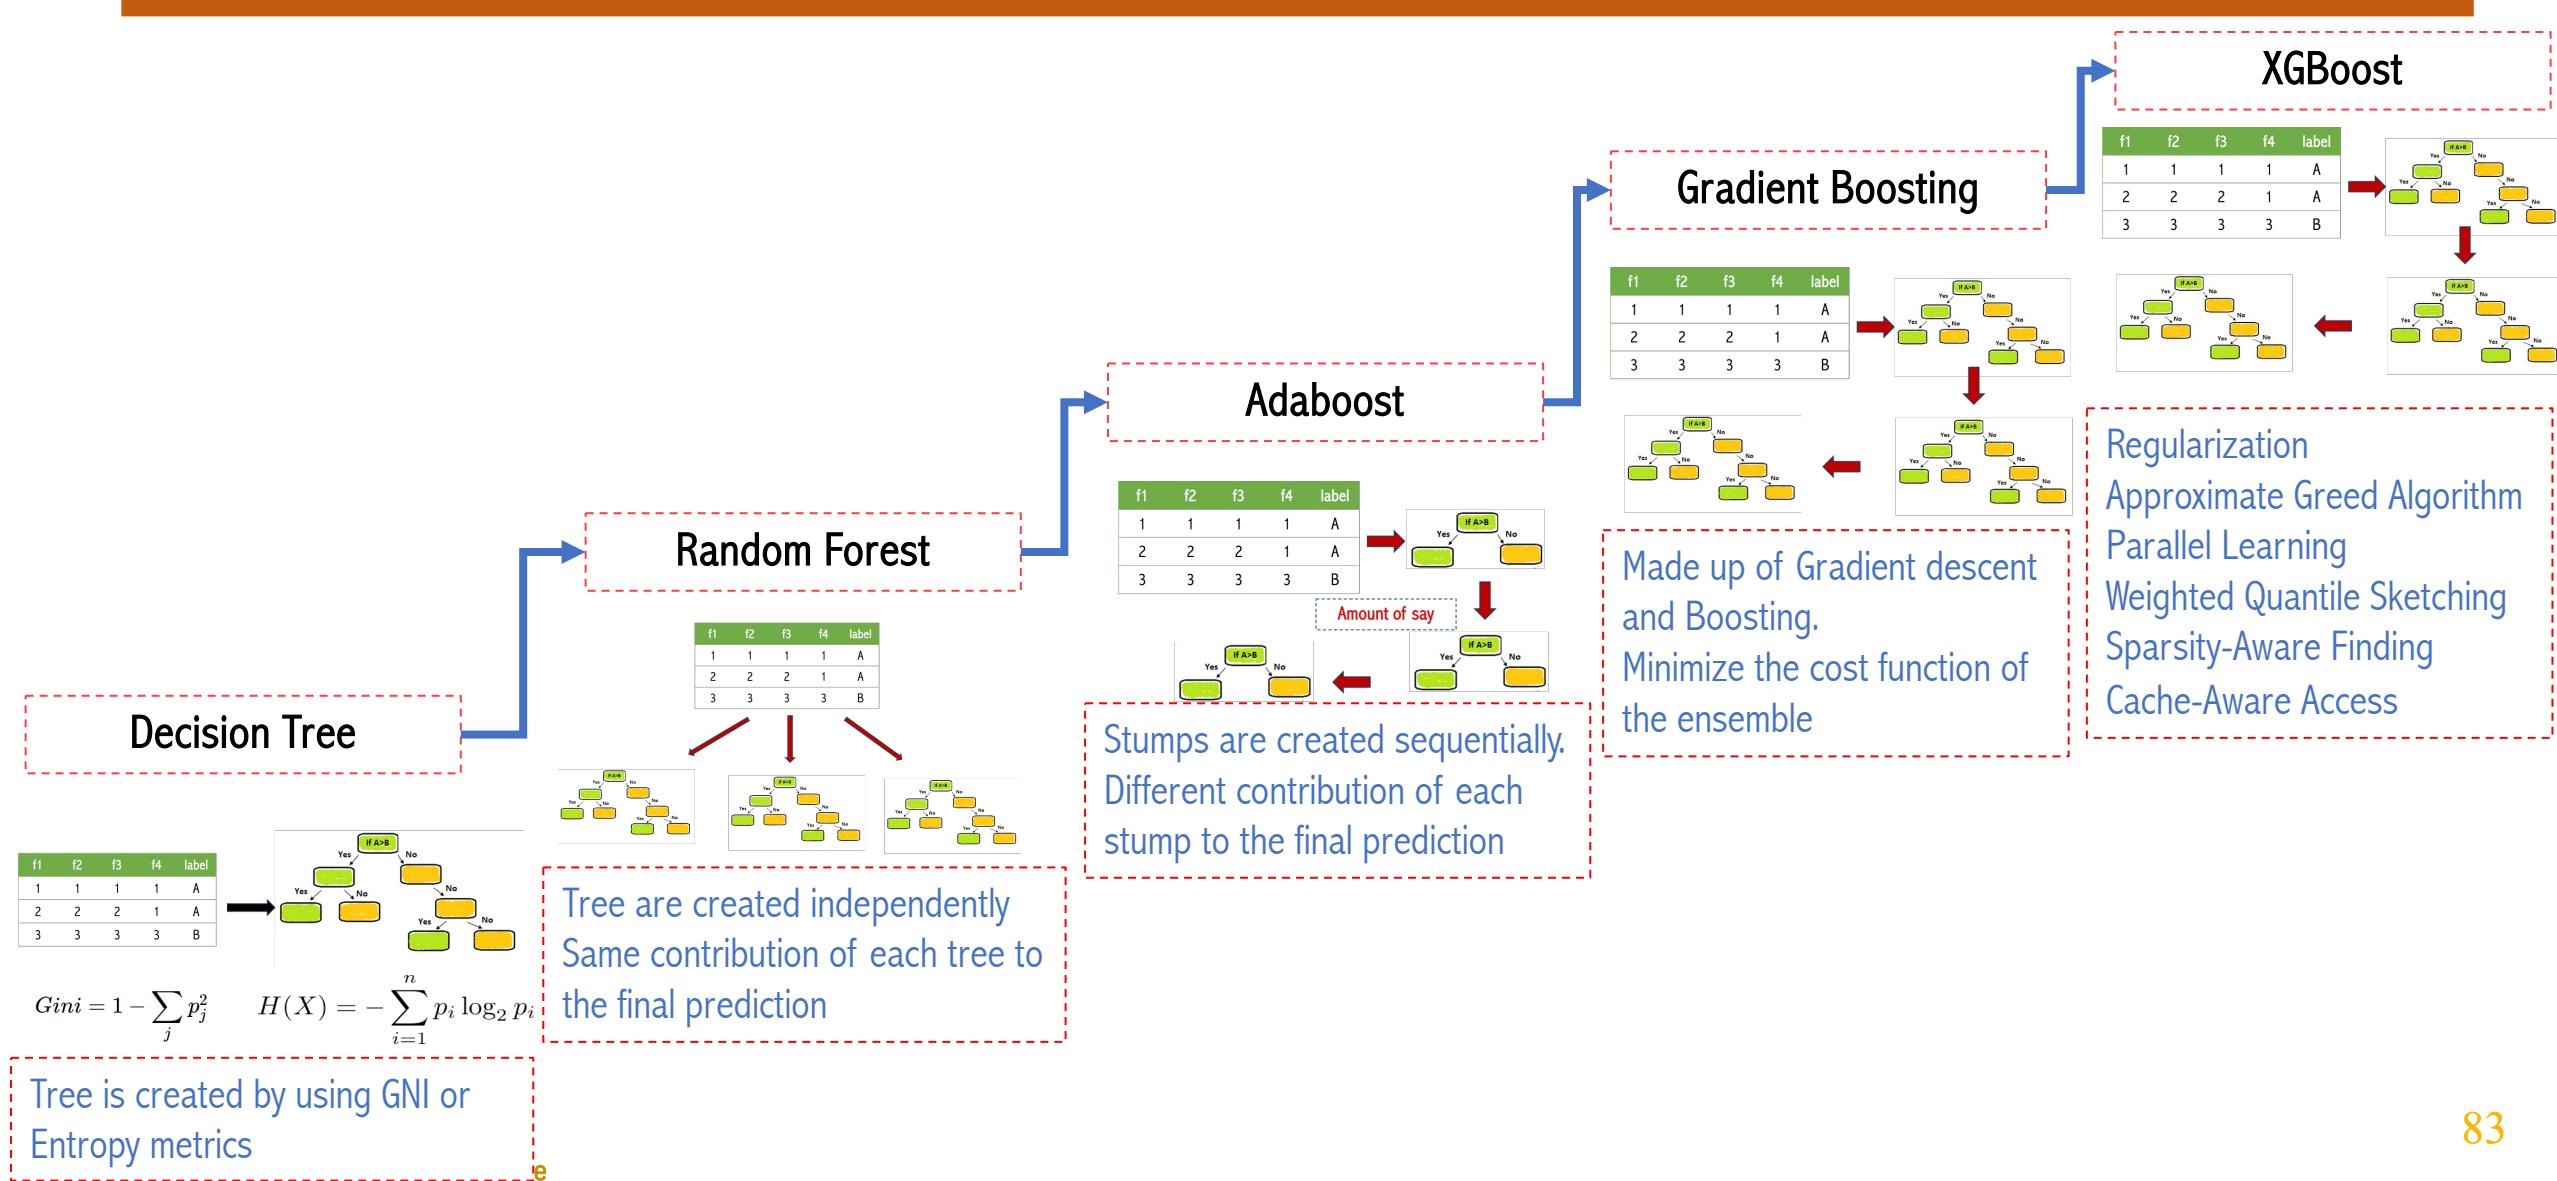


## Case 3: Univariate to Univariate Prediction

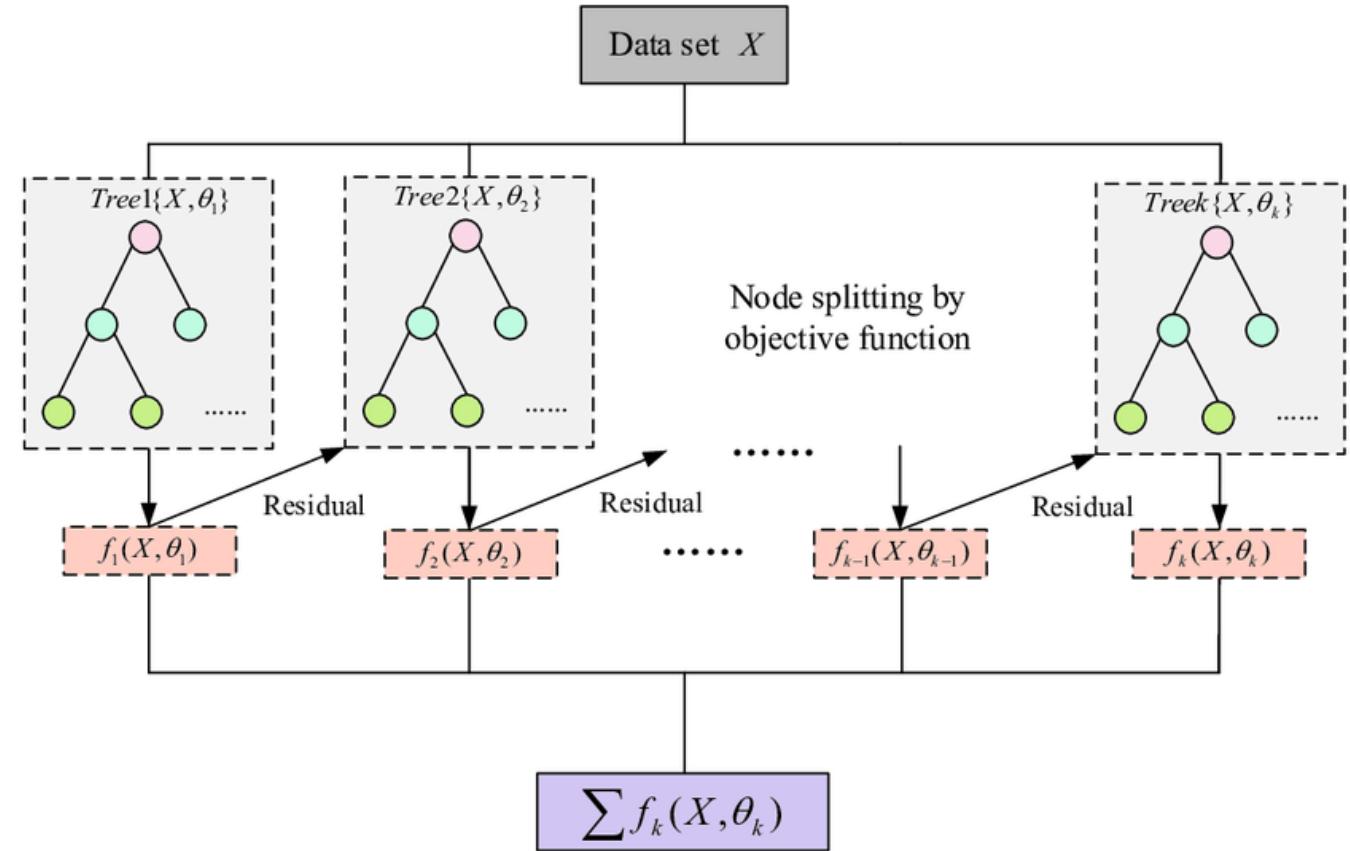
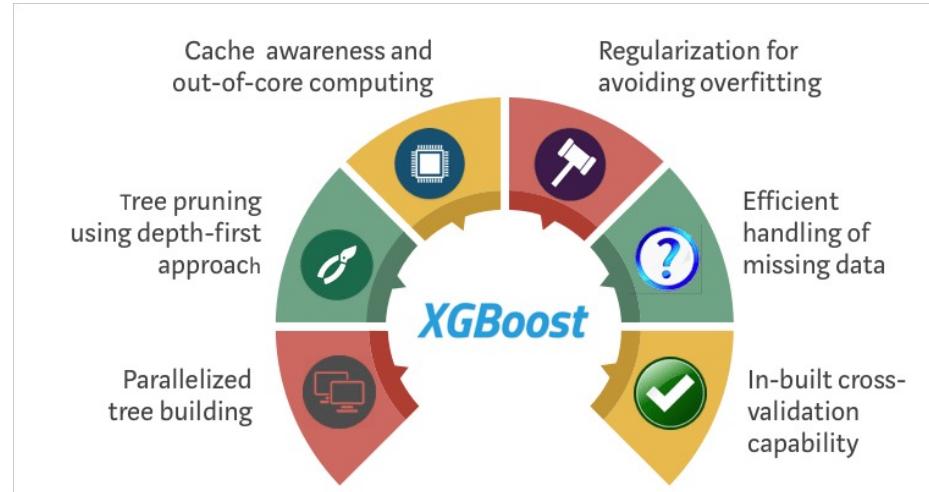
# Outline

- Multilayer Perceptron and Time-series Data Forecasting
- RNN and Time-series Data Forecasting
- LSTM and Time-series Data Forecasting
- Bi-LSTM and Time-series Data Forecasting
- XGBoost and Time-series Data Forecasting

# Evolution of Tree and Its Variant



# XGBoost: Simple Explanation



# XGBoost: Simple Explanation

Gradient Boost

Regularization

Approximate Greedy Algorithm

Parallel Learning

Weighted Quantile Sketch

Sparsity-Aware Split Finding

Cache-Aware Access

## Time Series Analysis



# XGBoost : Time-series Data Project

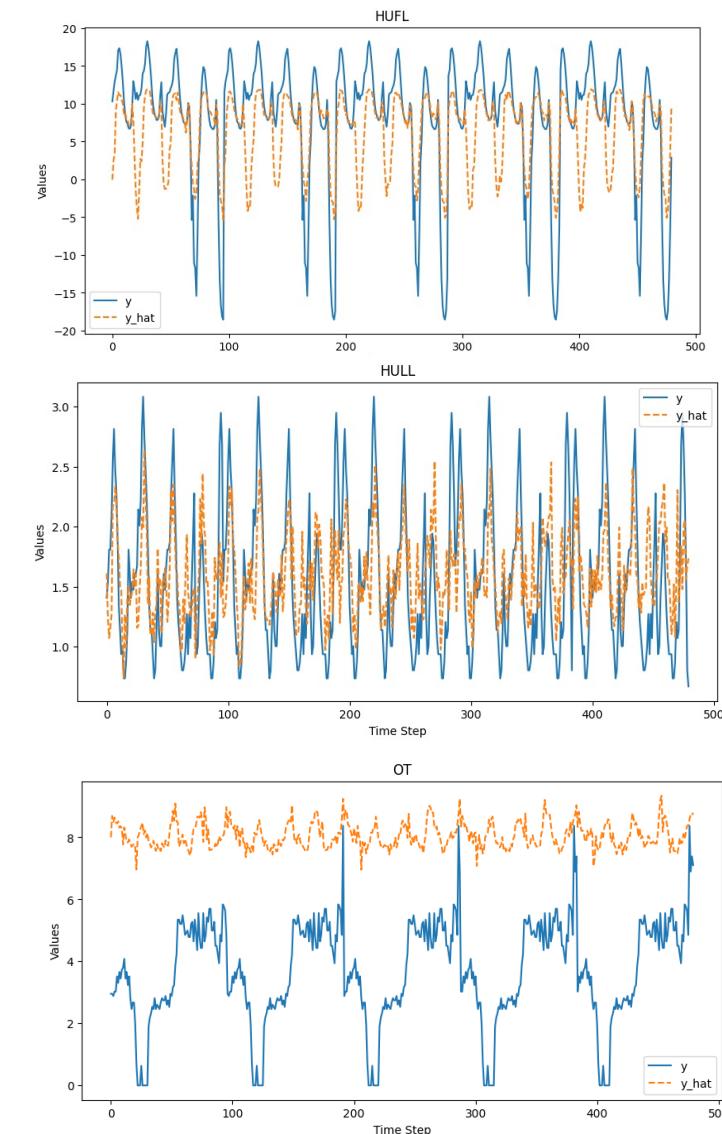
## Define XGBoost Structure in Pytorch

```
XGBoost_multi2multi = XGBRegressor(**xgboost_config)
XGBoost_multi2multi_manager = MachineLearningModelManager(model=XGBoost_multi2multi,
                                                          xtrain=multi2multi_loader.X_train,
                                                          ytrain=multi2multi_loader.y_train,
                                                          xval=multi2multi_loader.X_val,
                                                          yval=multi2multi_loader.y_val)

XGBoost_multi2multi_manager.train(save_dir=os.path.join(weight_dir, sub_dir))
results.append({
    "Name": XGBoost_multi2multi_manager.model.__class__.__name__,
    "Type": sub_dir,
    "MAE": XGBoost_multi2multi_manager.evaluate(x=multi2multi_loader.X_test, y=multi2multi_loader.y_test
})
results[-1]

[0] validation_0-mae:3.90950
[1] validation_0-mae:3.61946
[2] validation_0-mae:3.38305
[3] validation_0-mae:3.19688
[4] validation_0-mae:3.03072
```

## Multivariate to Multivariate Prediction

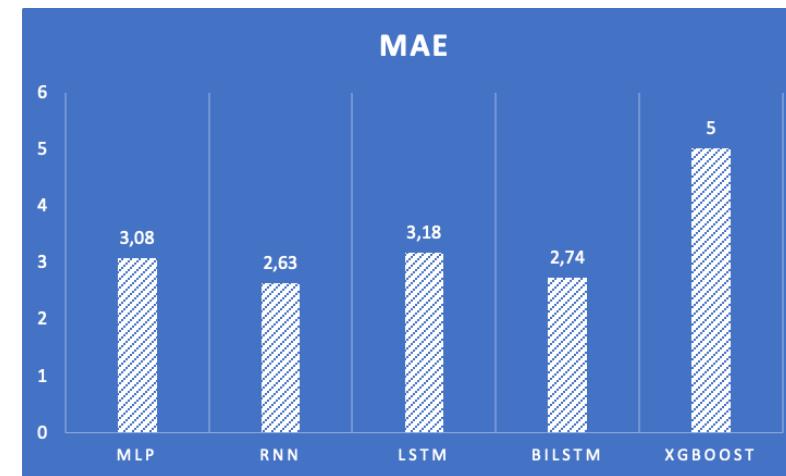


*Experiments were performed by TA Khoa Nguyen*

# Performance Evaluation



Multivariate to Multivariate



Multivariate to Univariate



Multivariate to Univariate

# Further Reading

Publisher	Year	Paper	Model
AAAI	2023	Are Transformers Effective for Time Series Forecasting?	NLinearDLinear
ICLR	2023	A TIME SERIES IS WORTH 64 WORDS: LONG-TERM FORECASTING WITH TRANSFORMERS	PatchTST
	2023	iTransformer: Inverted Transformers Are Effective for Time Series Forecasting	iTransformer
ICLR	2023	TimesNet: Temporal 2D-Variation Modeling for General Time Series Analysis	TimesNet
	2023	An adaptive standardisation model for Day-Ahead electricity price forecasting	None
ICLR	2023	MICN: Multi-scale Local and Global Context Modeling for Long-term Series Forecasting	MICN
PeerJ Computer Science	2023	Multi-horizon short-term load forecasting using hybrid of LSTM and modified split convolution	LSTM-SC
	2023	PatchMixer: A Patch-Mixing Architecture for Long-Term Time Series Forecasting	PatchMixer
	2023	SegRNN: Segment Recurrent Neural Network for Long-Term Time Series Forecasting	SegRNN
TMLR	2023	TSMixer: An All-MLP Architecture for Time Series Forecasting	TSMixer
NeurIPS	2023	One Fits All: Power General Time Series Analysis by Pretrained LM	None
	2023	Distilling Universal and Joint Knowledge for Cross-Domain Model Compression on Time Series Data	UNI-KD
	2023	TSMixer: An All-MLP Architecture for Time Series Forecasting	TSMixer
IEEE	2023	Prophet: Prompting Large Language Models with Complementary Answer Heuristics for Knowledge-based Visual Question Answering	Prophet
	2023	Alternative Telescopic Displacement: An Efficient Multimodal Alignment Method	None
	2023	Long-term Forecasting with TiDE: Time-series Dense Encoder	TiDE
NeurIPS	2022	Non-stationary Transformers: Exploring the Stationarity in Time Series Forecasting	Stationary
ICML	2022	FEDformer: Frequency Enhanced Decomposed Transformer for Long-term Series Forecasting	FEDformer
NeurIPS	2022	SCINet: Time Series Modeling and Forecasting with Sample Convolution and Interaction	SCINet
ICLR	2022	CROSSFORMER: A VERSATILE VISION TRANSFORMER HINGING ON CROSS-SCALE ATTENTION	Crossformer
	2022	Less Is More: Fast Multivariate Time Series Forecasting with Light Sampling-oriented MLP Structures	LightTS
	2022	N-HiTS: Neural Hierarchical Interpolation for Time Series Forecasting	N-HiTS
	2022	Self-Supervised Transformer for Sparse and Irregularly Sampled Multivariate Clinical Time-Series	STraTS
NeurIPS	2022	WaveBound: Dynamic Error Bounds for Stable Time Series Forecasting	None
	2022	ETSformer: Exponential Smoothing Transformers for Time-series Forecasting	ETSformer
NeurIPS	2021	Autoformer: Decomposition Transformers with Auto-Correlation for Long-Term Series Forecasting	Autoformer
AAAI	2021	Informer: Beyond Efficient Transformer for Long Sequence Time-Series Forecasting	Informer
CIKM	2021	AdaRNN: Adaptive Learning and Forecasting of Time Series	AdaRNN
	2021	Prediction of the Position of External Markers Using a Recurrent Neural Network Trained With Unbiased Online Recurrent Optimization for Safe Lung Cancer Radiotherapy	UORO
	2021	Long-term series forecasting with Query Selector -- efficient model of sparse attention	QuerySelector
	2021	A hybrid method of exponential smoothing and recurrent neural networks for time series forecasting	ES-RNN
	2021	Autoregressive Denoising Diffusion Models for Multivariate Probabilistic Time Series Forecasting	TimeGrad
	2021	Systematic Generalization in Neural Networks-based Multivariate Time Series Forecasting Models	None
	2021	ES-dRNN: A Hybrid Exponential Smoothing and Dilated Recurrent Neural Network Model for Short-Term Load Forecasting	ES-dRNN
ICLR	2021	Multivariate Probabilistic Time Series Forecasting via Conditioned Normalizing Flows	LSTM-MAF
	2020	Temporal Fusion Transformers for Interpretable Multi-horizon Time Series Forecasting	TFT
NeurIPS	2020	Normalizing Kalman Filters for Multivariate Time Series Analysis	NKF

<https://docs.google.com/spreadsheets/d/17J6ldQogp186AZDZv8ZVIBF5s2EaeZwQ/edit?usp=sharing&ouid=118245386452054287147&rtpof=true&sd=true>

