

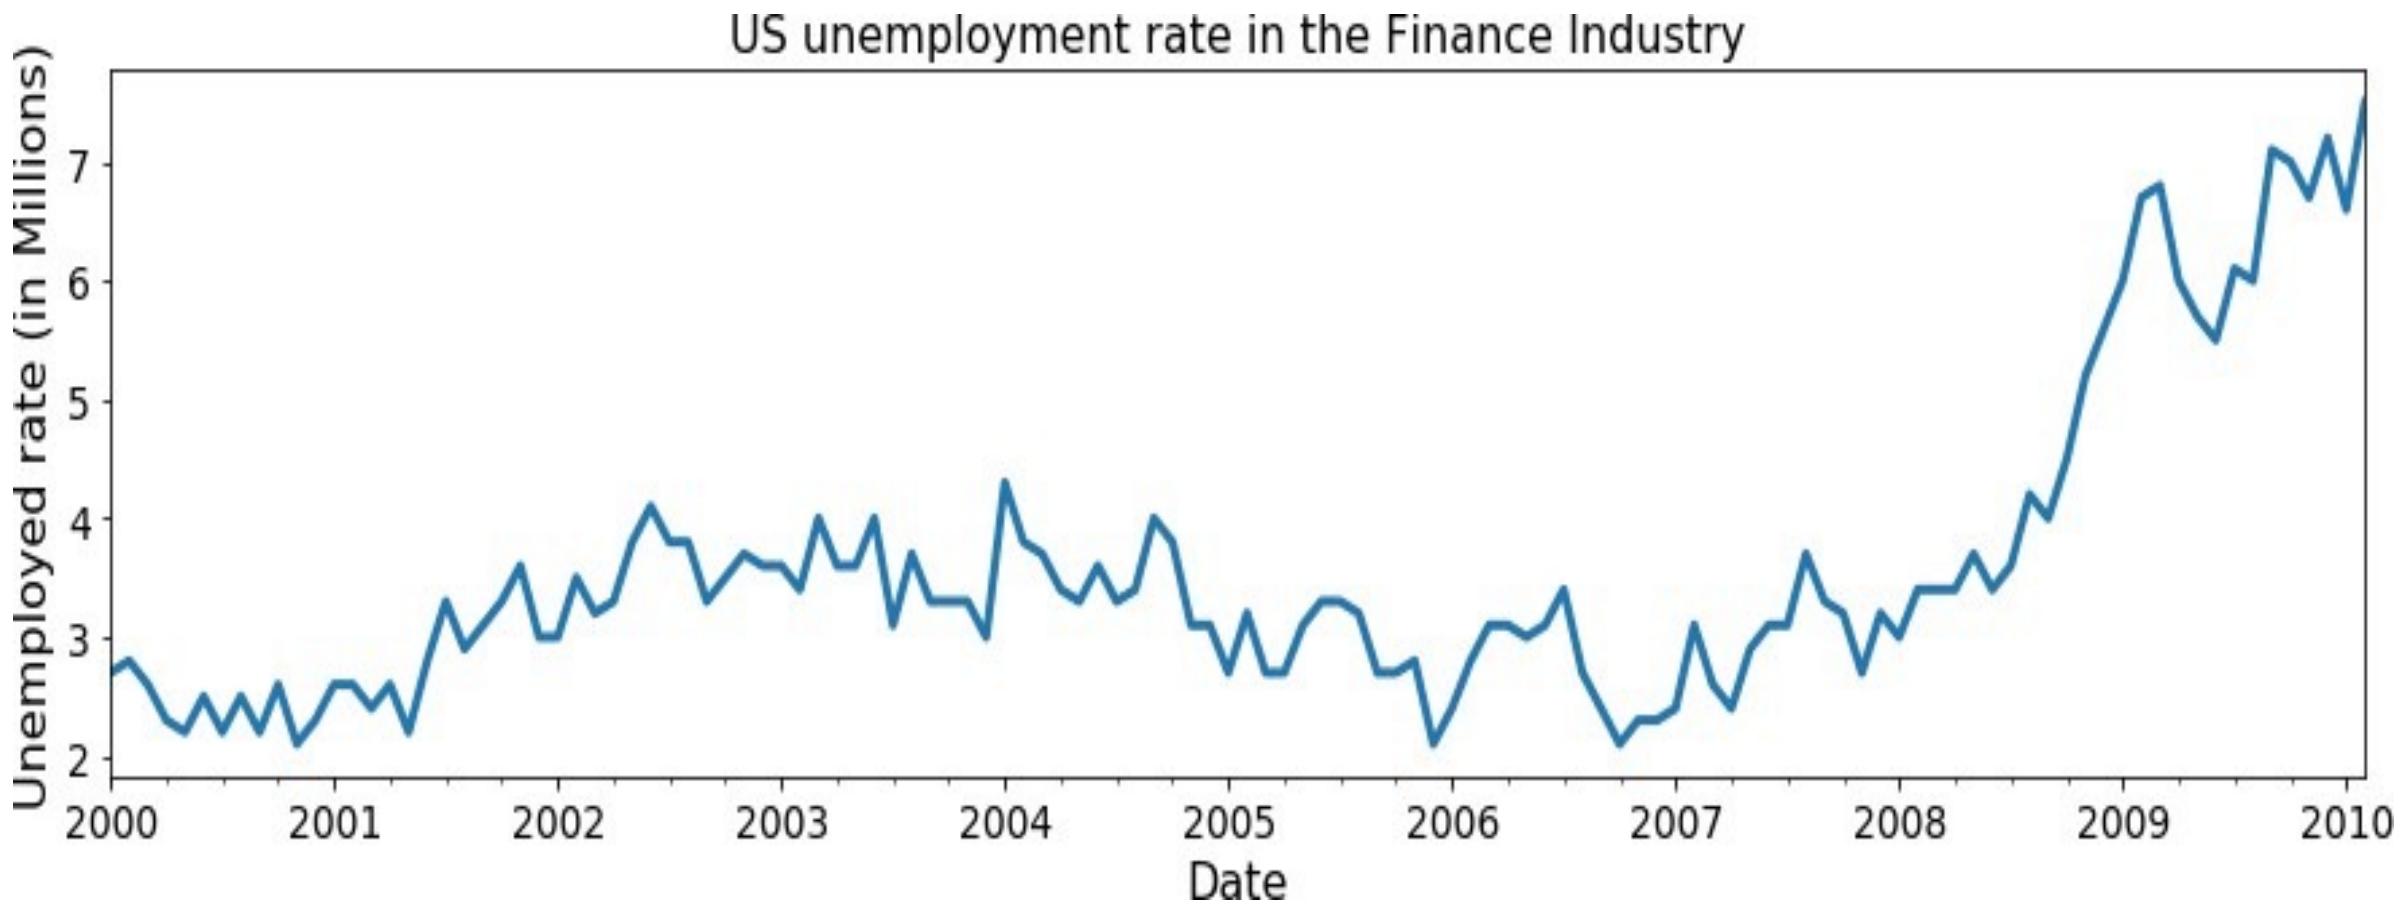
# Visualize Time Series Data in Python

Trung-Nghia Cao  
Minh-Duc Bui

# Outline

- Line Plots
- Summary Statistics and Diagnostics
- Seasonality, Trend and Noise
- Work with Multiple Time Series
- Case Study: Unemployment Rate

# Time Series in the Field of Data Science



# Time Series in the Field of Data Science

## Univariate Time Series

S&P 500 Stock Prices

Currency in USD

Date	Open
Oct 30, 2020	3,293.59
Oct 29, 2020	3,277.17
Oct 28, 2020	3,342.48
Oct 27, 2020	3,403.15
Oct 26, 2020	3,441.42
Oct 23, 2020	3,464.90
Oct 22, 2020	3,438.50
Oct 21, 2020	3,439.91
Oct 20, 2020	3,439.38
Oct 19, 2020	3,493.66
Oct 16, 2020	3,493.50
Oct 15, 2020	3,453.72

Source: [www.finance.yahoo.com](http://www.finance.yahoo.com)

## Multivariate Time Series

S&P 500 Stock Prices

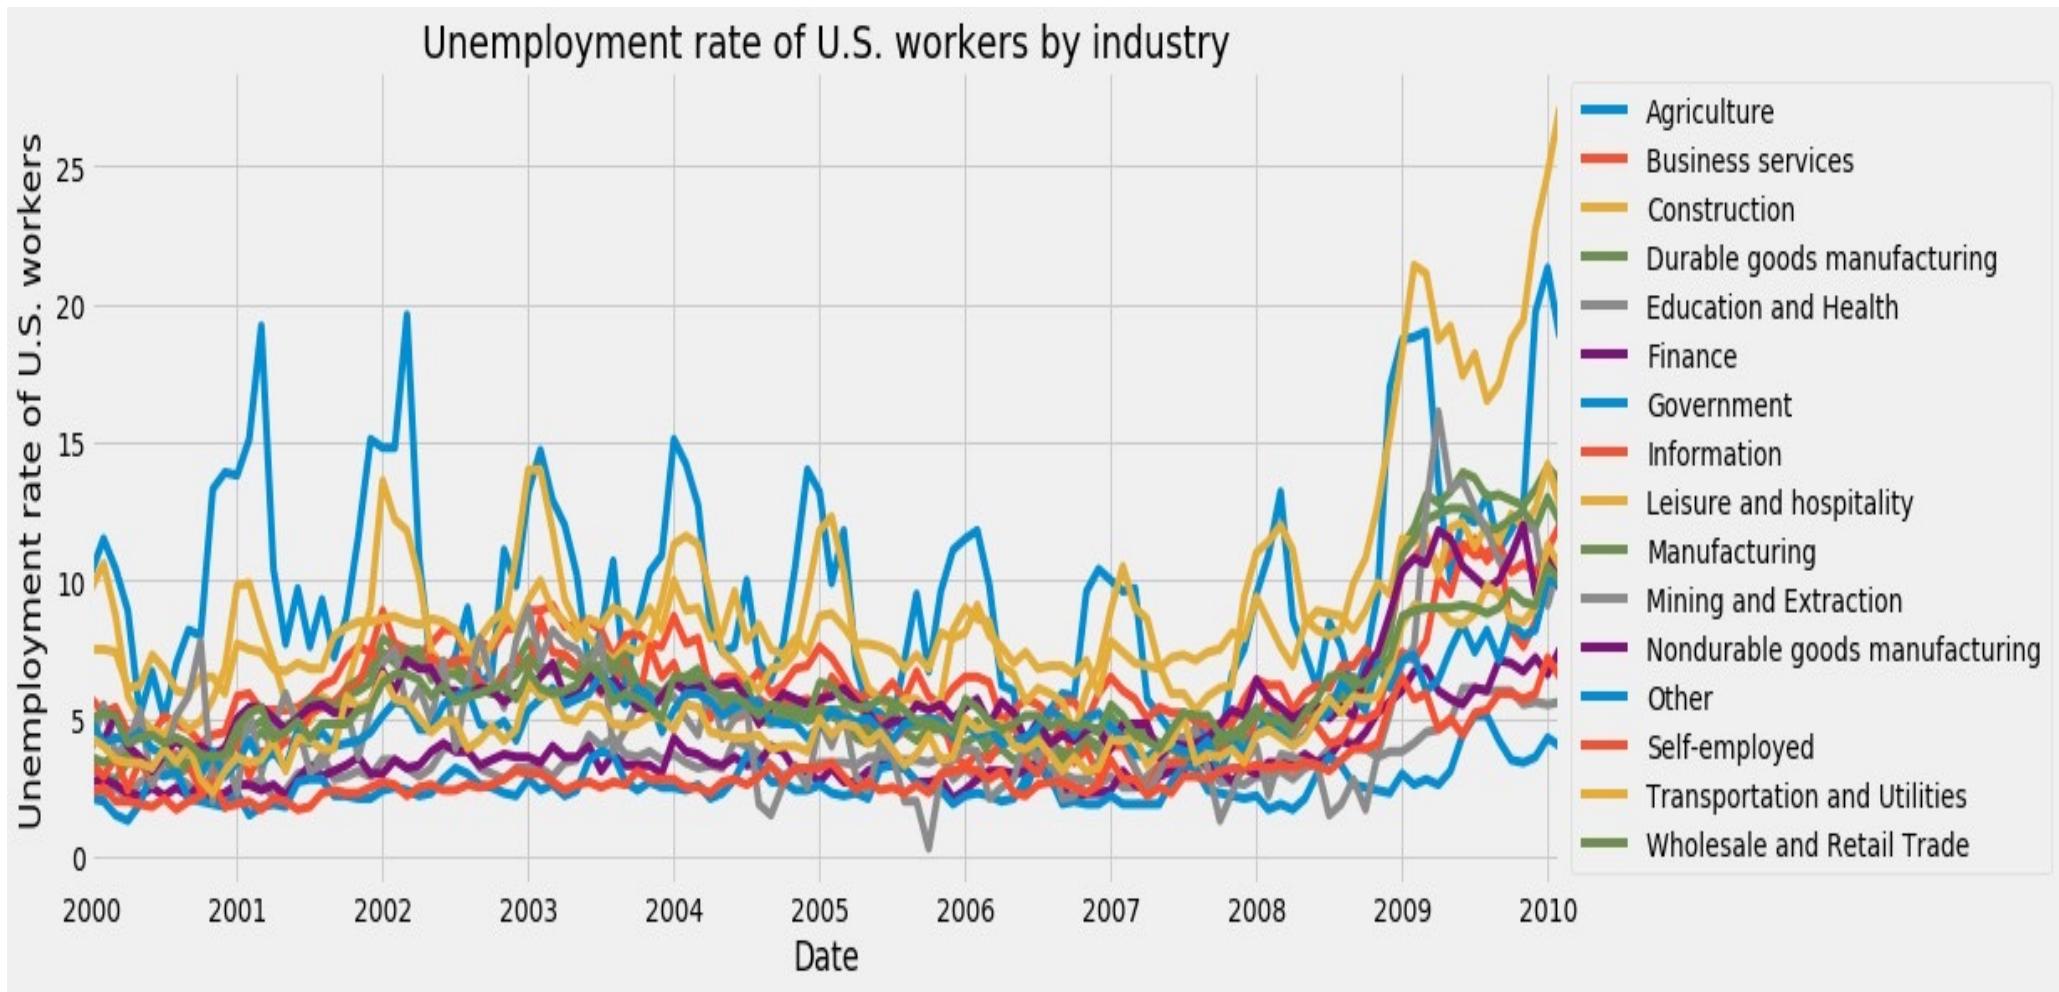
Currency in USD

Date	Open	High	Low	Close*	Adj Close**	Volume
Oct 30, 2020	3,293.59	3,304.93	3,233.94	3,269.96	3,269.96	4,840,450,000
Oct 29, 2020	3,277.17	3,341.05	3,259.82	3,310.11	3,310.11	4,903,070,000
Oct 28, 2020	3,342.48	3,342.48	3,268.89	3,271.03	3,271.03	5,129,860,000
Oct 27, 2020	3,403.15	3,409.51	3,388.71	3,390.68	3,390.68	3,946,990,000
Oct 26, 2020	3,441.42	3,441.42	3,364.86	3,400.97	3,400.97	3,988,080,000
Oct 23, 2020	3,464.90	3,466.46	3,440.45	3,465.39	3,465.39	3,646,570,000
Oct 22, 2020	3,438.50	3,460.53	3,415.34	3,453.49	3,453.49	4,163,630,000
Oct 21, 2020	3,439.91	3,464.86	3,433.06	3,435.56	3,435.56	4,097,750,000
Oct 20, 2020	3,439.38	3,476.93	3,435.65	3,443.12	3,443.12	3,901,260,000
Oct 19, 2020	3,493.66	3,502.42	3,419.93	3,426.92	3,426.92	4,086,200,000
Oct 16, 2020	3,493.50	3,515.76	3,480.45	3,483.81	3,483.81	4,675,890,000
Oct 15, 2020	3,453.72	3,489.08	3,440.89	3,483.34	3,483.34	3,717,640,000

Source: [www.finance.yahoo.com](http://www.finance.yahoo.com)

# Time Series in the Field of Data Science

## Multiple Time Series



# Outline

- Line Plots
- Summary Statistics and Diagnostics
- Seasonality, Trend and Noise
- Work with Multiple Time Series
- Case Study: Unemployment Rate

# Reading data with Pandas

```
import pandas as pd  
  
df = pd.read_csv('ch2_co2_levels.csv')  
  
print(df)
```

```
      datestamp    co2  
0    1958-03-29  316.1  
1    1958-04-05  317.3  
2    1958-04-12  317.6  
...  
...  
...  
2281  2001-12-15  371.2  
2282  2001-12-22  371.3  
2283  2001-12-29  371.5
```

# Preview data with Pandas

```
print(df.head(n=5))
```

```
print(df.tail(n=5))
```

```
datestamp      co2
0   1958-03-29  316.1
1   1958-04-05  317.3
2   1958-04-12  317.6
3   1958-04-19  317.5
4   1958-04-26  316.4
```

```
datestamp      co2
2279  2001-12-01  370.3
2280  2001-12-08  370.8
2281  2001-12-15  371.2
2282  2001-12-22  371.3
2283  2001-12-29  371.5
```

# Check data types with Pandas

```
print(df.dtypes)
```

```
datestamp          object
co2                float64
dtype:            object
```

# Working with dates

To work with time series data in **pandas**, your date columns needs to be of the **datetime64** type.

```
pd.to_datetime(['2009/07/31', 'test'])
```

```
ValueError: Unknown string format)
```

```
pd.to_datetime(['2009/07/31', 'test'], errors='coerce')
```

```
DatetimeIndex(['2009-07-31', 'NaT'],
                dtype='datetime64[ns]', freq=None)
```

# The Matplotlib library

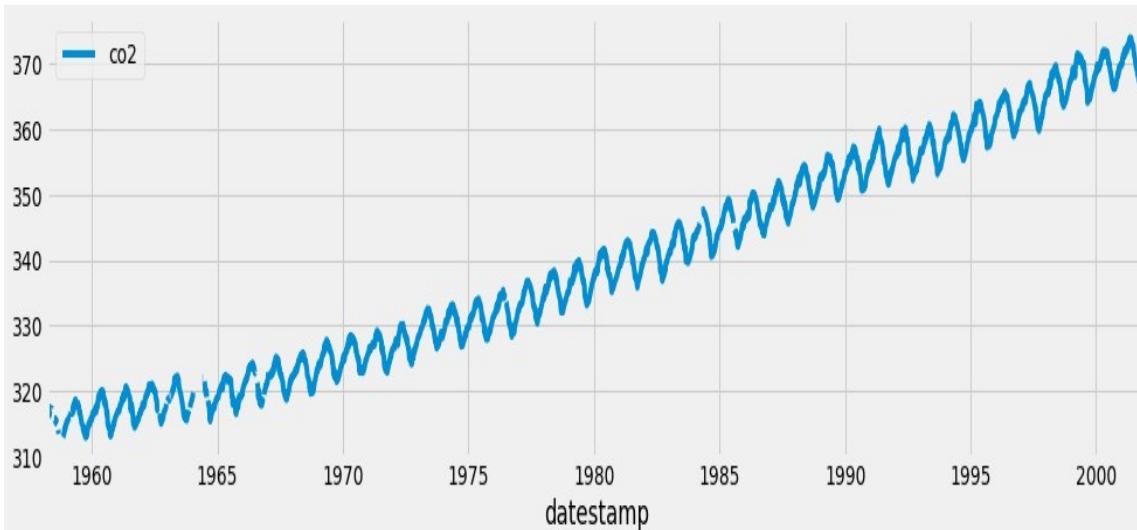
---

In Python, matplotlib is an extensive package used to plot data  
The pyplot submodule of matplotlib is traditionally imported using  
the **plt** alias

```
import matplotlib.pyplot as plt
```

# Plotting time series data

```
import matplotlib.pyplot as plt  
  
import pandas as pd  
  
plt.style.use('fivethirtyeight')  
  
df = df.set_index('date_column')  
df.plot()  
plt.show()
```



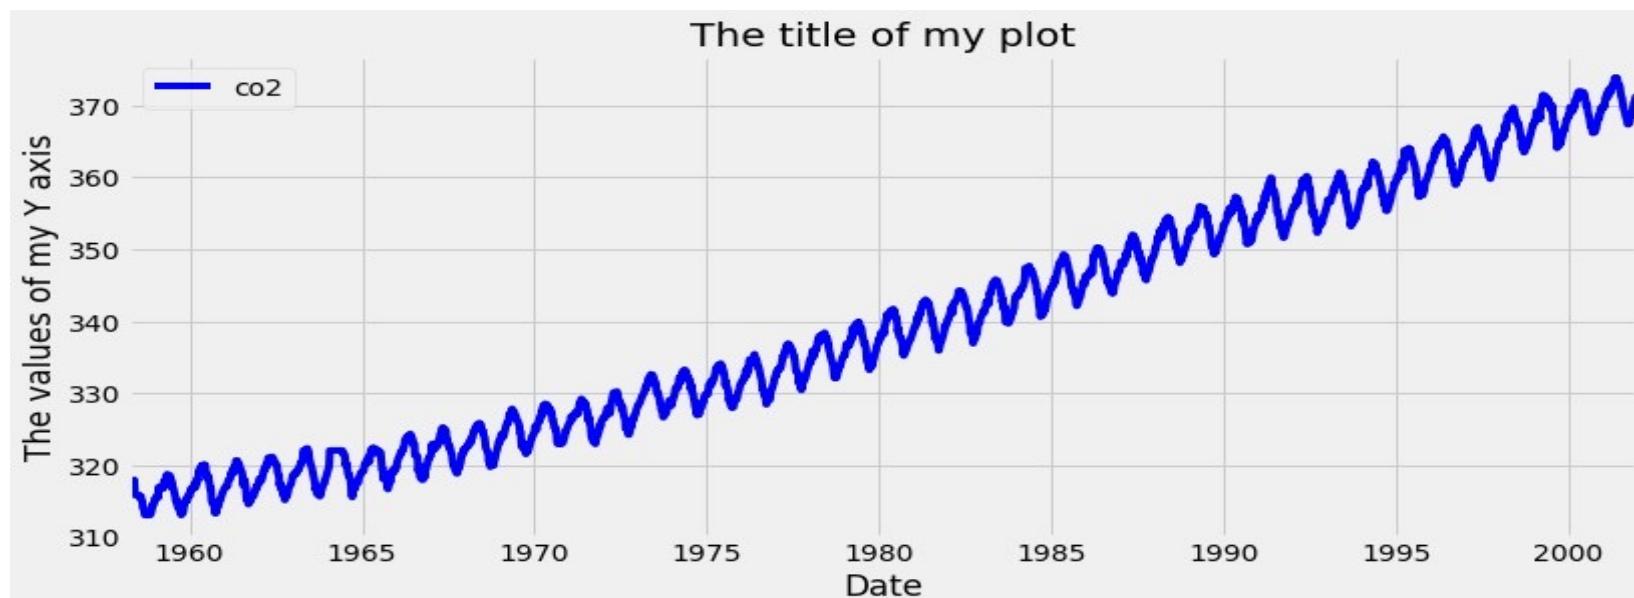
```
print(plt.style.available)
```

```
['seaborn-dark-palette', 'seaborn-darkgrid',  
'seaborn-dark', 'seaborn-notebook', 'seaborn-  
pastel', 'seaborn-white', 'classic', 'ggplot',  
'grayscale', 'dark_background', 'seaborn-  
poster', 'seaborn-muted', 'seaborn',  
'bmh', 'seaborn-paper', 'seaborn-whitegrid',  
'seaborn-bright', 'seaborn-talk',  
'fivethirtyeight', 'seaborn-colorblind',  
'seaborn-deep', 'seaborn-ticks']
```

# Describing your graphs with labels

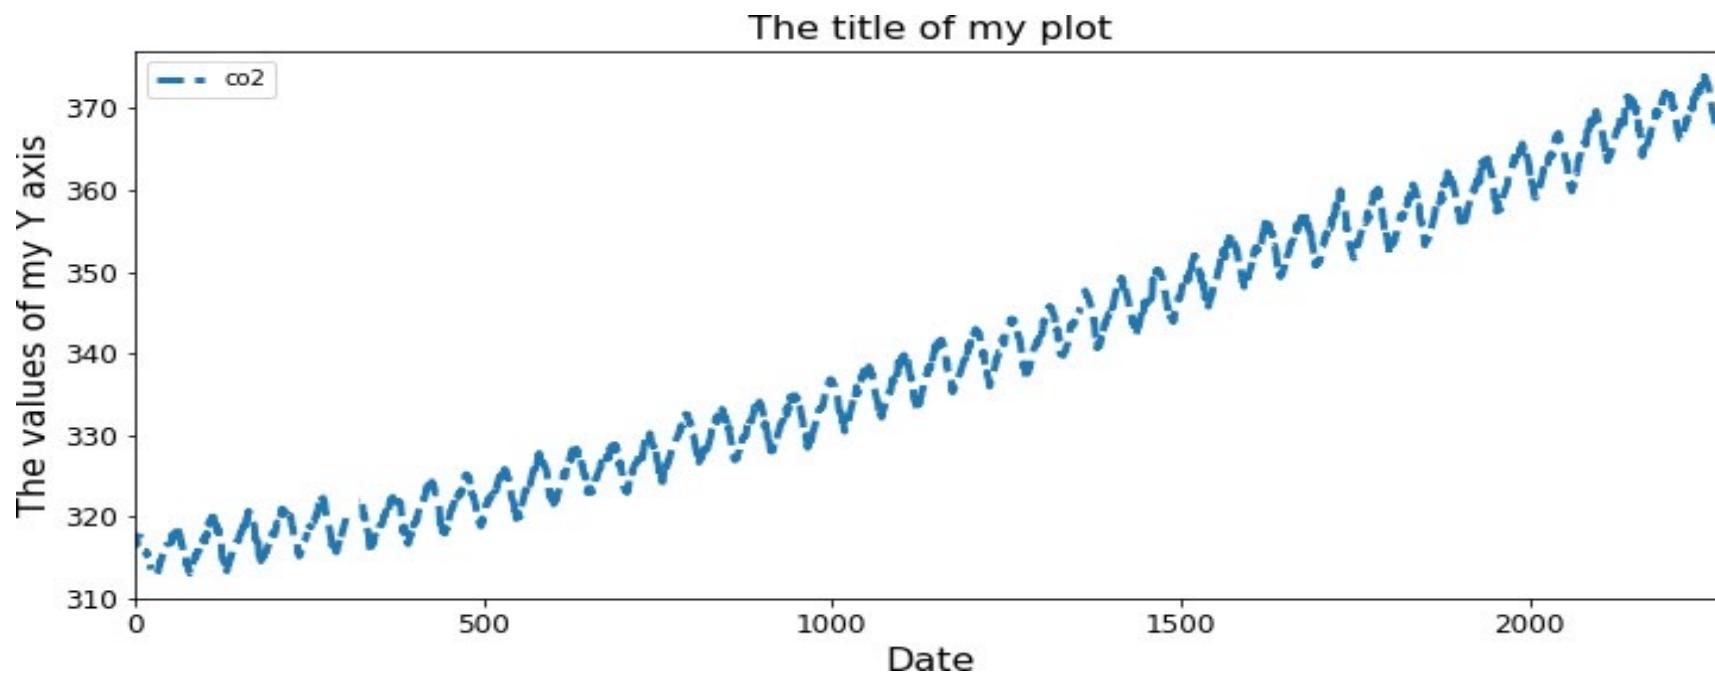
```
ax = df.plot(color='blue')

ax.set_xlabel('Date')
ax.set_ylabel('The values of my Y axis')
ax.set_title('The title of my plot')
plt.show()
```



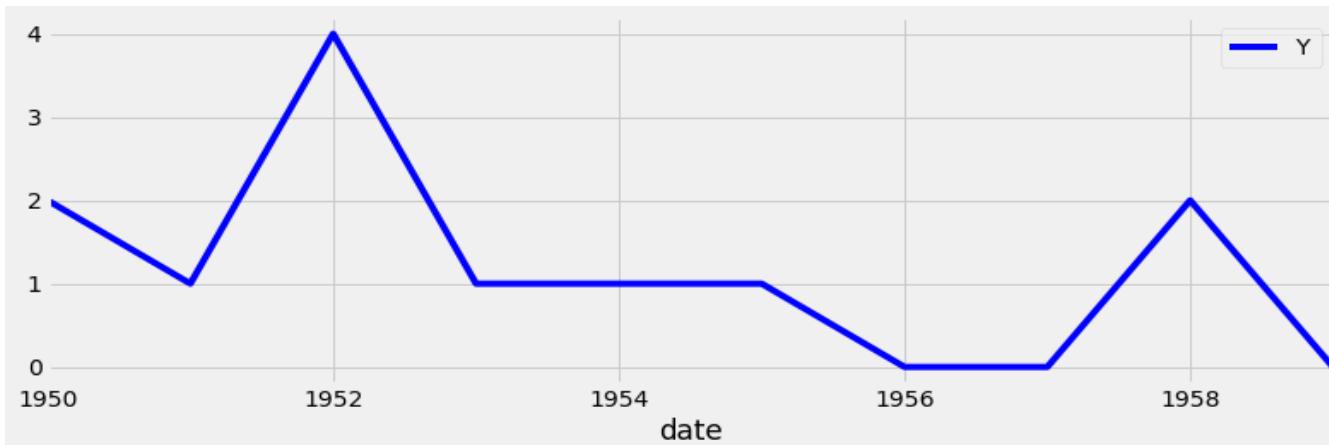
# Figure size, linewidth, linestyle and fontsize

```
ax = df.plot(figsize=(12, 5), fontsize=12,  
             linewidth=3, linestyle='--')  
  
ax.set_xlabel('Date', fontsize=16)  
ax.set_ylabel('The values of my Y axis', fontsize=16)  
ax.set_title('The title of my plot', fontsize=16)  
plt.show()
```



# Plotting subset of your time series data

```
import matplotlib.pyplot as plt  
  
plt.style.use('fivethirtyeight')  
  
df_subset = discoveries['1950':'1960']  
  
ax = df_subset.plot(color='blue', fontsize=14)  
plt.show()
```



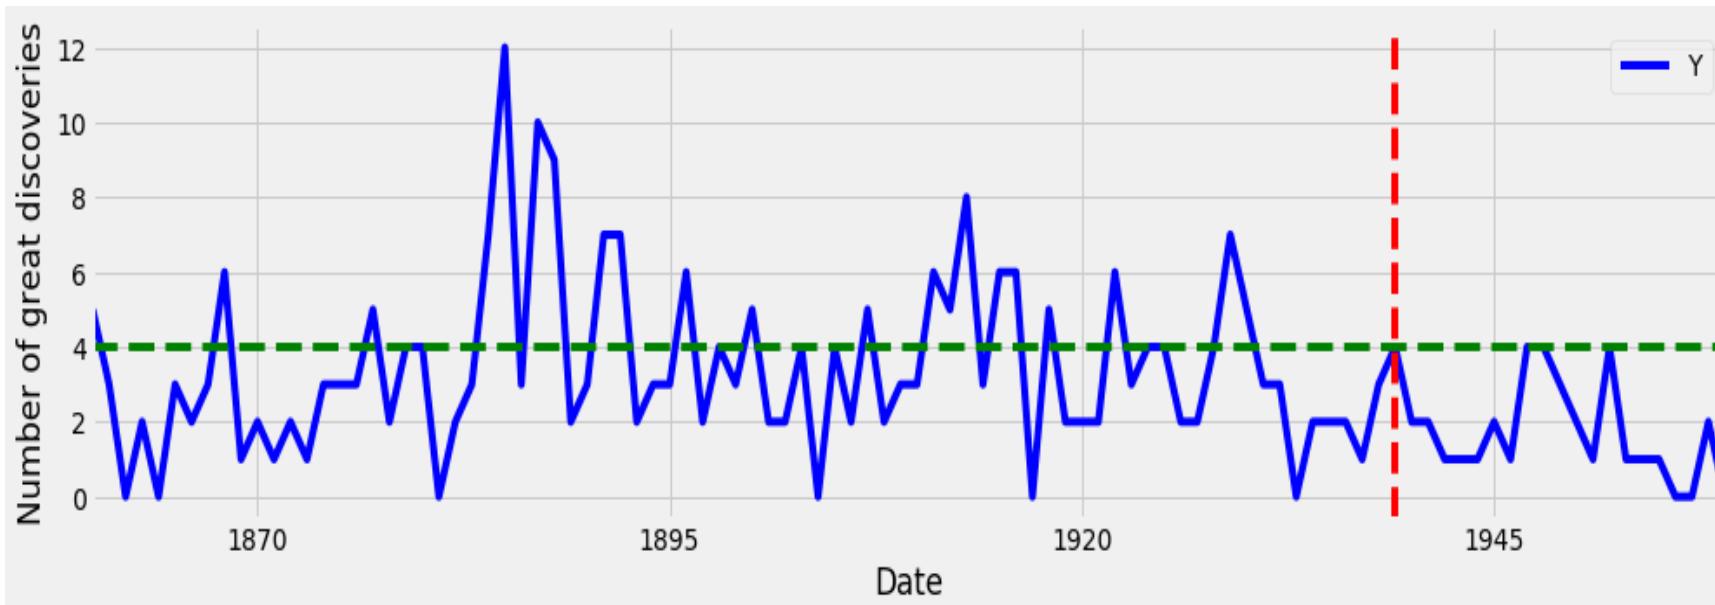
```
discoveries['1960':'1970']
```

```
discoveries['1950-01':'1950-12']
```

```
discoveries['1960-01-01':'1960-01-15']
```

# Using markers

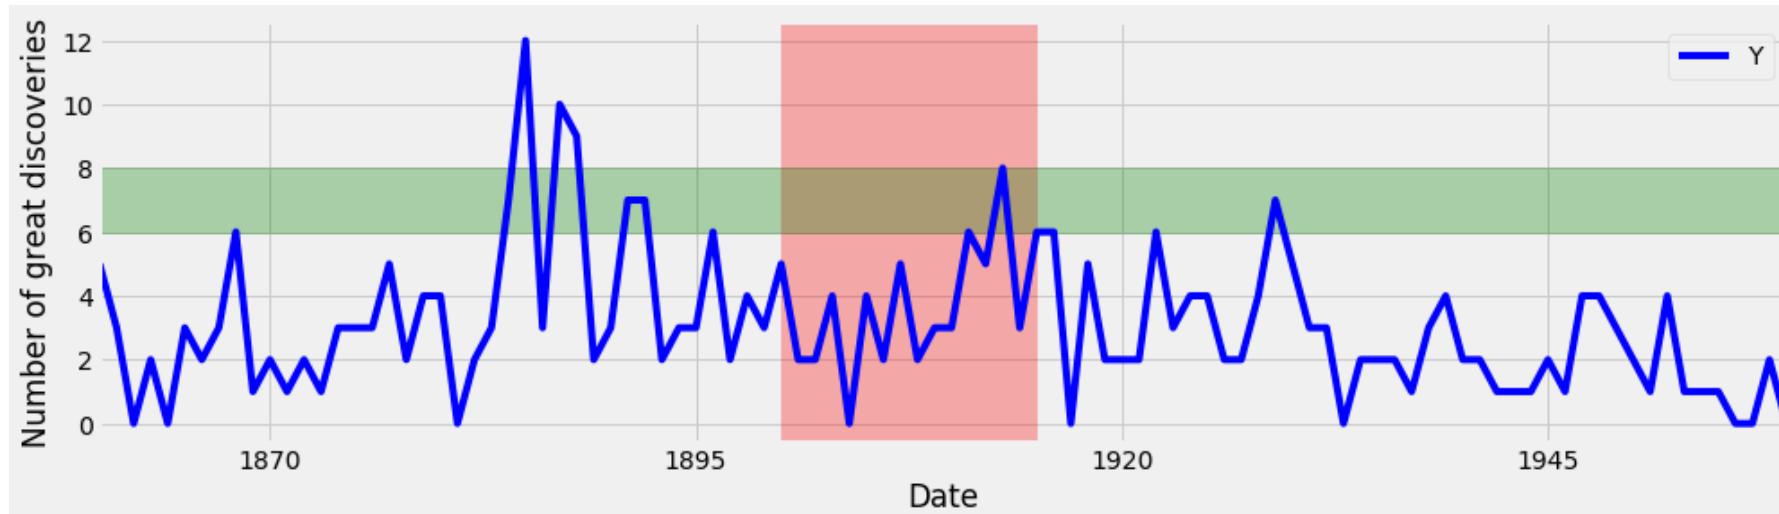
```
ax = discoveries.plot(color='blue')
ax.set_xlabel('Date')
ax.set_ylabel('Number of great discoveries')
ax.axvline('1940-01-01', color='red', linestyle='--')
ax.axhline(4, color='green', linestyle='--')
```



# Highlighting regions of interest

```
ax = discoveries.plot(color='blue')  
ax.set_xlabel('Date')  
ax.set_ylabel('Number of great discoveries')
```

```
ax.axvspan('1900-01-01', '1915-01-01', color='red', alpha=0.3)  
ax.axhspan(8, 6, color='green', alpha=0.3)
```



# Outline

- Line Plots
- **Summary Statistics and Diagnostics**
- Seasonality, Trend and Noise
- Work with Multiple Time Series
- Case Study: Unemployment Rate
- Feature Engineering in General
- Feature Engineering for Time Series Data

# The CO<sub>2</sub> level time series

A snippet of the weekly measurements of CO<sub>2</sub> levels at the Mauna Loa Observatory, Hawaii.

datastamp	co2
1958-03-29	316.1
1958-04-05	317.3
1958-04-12	317.6
...	
...	
2001-12-15	371.2
2001-12-22	371.3
2001-12-29	371.5

# Finding missing values in a DataFrame

```
print(df.isnull())
```

```
datestamp      co2
1958-03-29    False
1958-04-05    False
1958-04-12    False
```

```
print(df.notnull())
```

```
datestamp      co2
1958-03-29    True
1958-04-05    True
1958-04-12    True
...
...
```

# Counting missing values in a DataFrame

```
print(df.isnull().sum())
```

```
datestamp      0
co2           59
dtype: int64
```

# Replacing missing values in a DataFrame

```
print(df)
```

```
...
5 1958-05-03 316.9
6 1958-05-10     NaN
7 1958-05-17 317.5
...
```

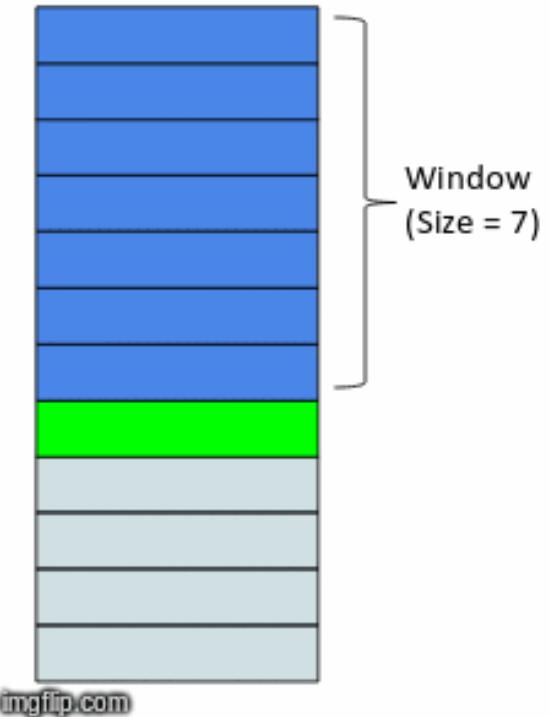
```
df = df.fillna(method='bfill')
print(df)
```

```
...
5 1958-05-03 316.9
6 1958-05-10 317.5
7 1958-05-17 317.5
...
```

# Moving averages

In the field of time series analysis, a moving average can be used for many different purposes:

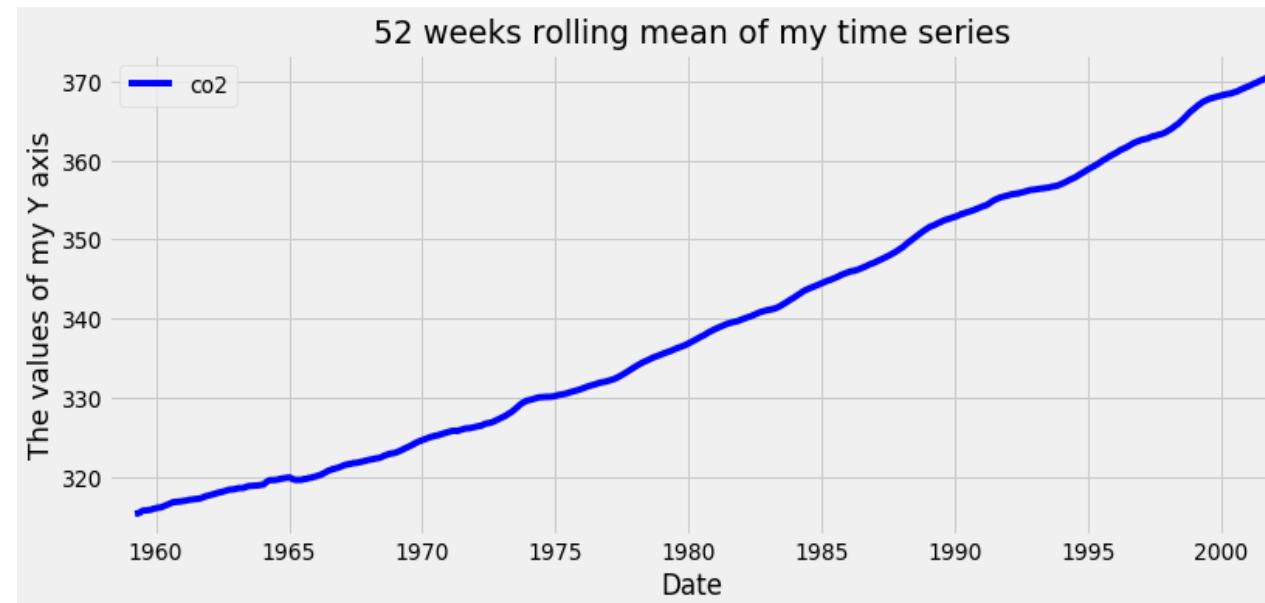
- smoothing out short-term fluctuations
- removing outliers
- highlighting long-term trends or cycles



# The moving average model

```
co2_levels_mean = co2_levels.rolling(window=52).mean()

ax = co2_levels_mean.plot()
ax.set_xlabel("Date")
ax.set_ylabel("The values of my Y axis")
ax.set_title("52 weeks rolling mean of my time series")
plt.show()
```



# Computing aggregate values

```
co2_levels.index
```

```
DatetimeIndex(['1958-03-29', '1958-04-05', ...],  
               dtype='datetime64[ns]', name='datestamp',  
               length=2284, freq=None)
```

```
print(co2_levels.index.month)
```

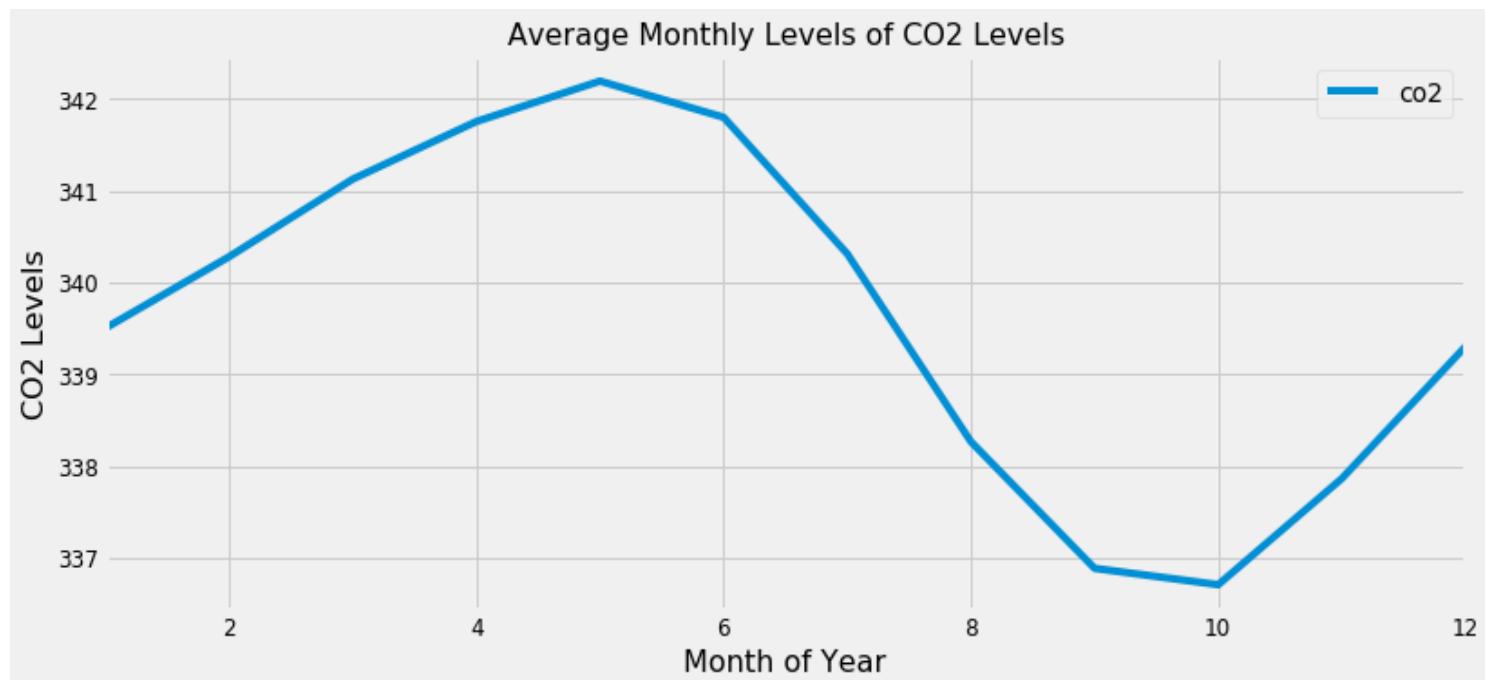
```
array([ 3,  4,  4, ..., 12, 12, 12], dtype=int32)
```

```
print(co2_levels.index.year)
```

```
array([1958, 1958, 1958, ..., 2001,  
      2001, 2001], dtype=int32)
```

# Plotting aggregate values

```
index_month = co2_levels.index.month  
co2_levels_by_month = co2_levels.groupby(index_month).mean()  
co2_levels_by_month.plot()  
  
plt.show()
```



# Obtaining numerical summaries

---

What is the **average value** of this data?

What is the **maximum value** observed in this time series?



# Obtaining numerical summaries

The `.describe()` method automatically computes key statistics of all numeric columns in your DataFrame

```
print(df.describe())
```

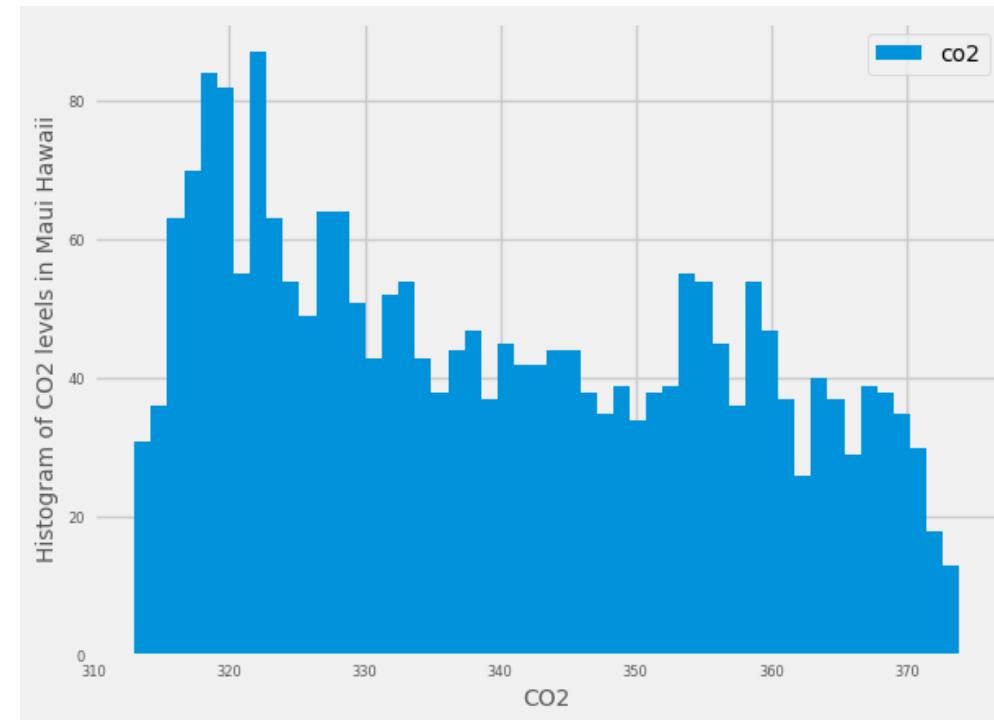
```
          co2
count    2284.000000
mean     339.657750
std      17.100899
min     313.000000
25%     323.975000
50%     337.700000
75%     354.500000
max     373.900000
```

# Histogram plot

```
# Generate a histogram
ax = co2_levels.plot(kind='hist', bins=50)

# Annotate labels
ax.set_xlabel('CO2', fontsize=10)
ax.set_ylabel('Histogram plot of CO2 levels in Maui Hawaii', fontsize=10)

plt.show()
```

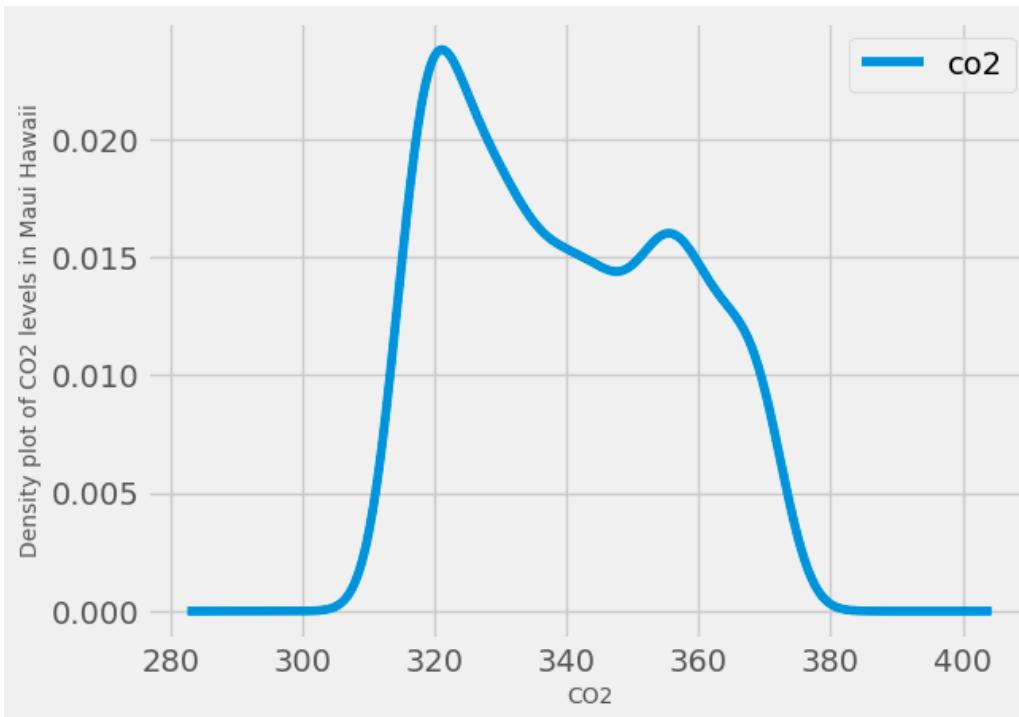


# Density plot

```
# Display density plot of CO2 levels values
ax = co2_levels.plot(kind='density', linewidth=4)

# Annotate labels
ax.set_xlabel('CO2', fontsize=10)
ax.set_ylabel('Density plot of CO2 levels in Maui Hawaii', fontsize=10)

plt.show()
```



The continuous and smoothed version of the Histogram

# Histogram & Density plots

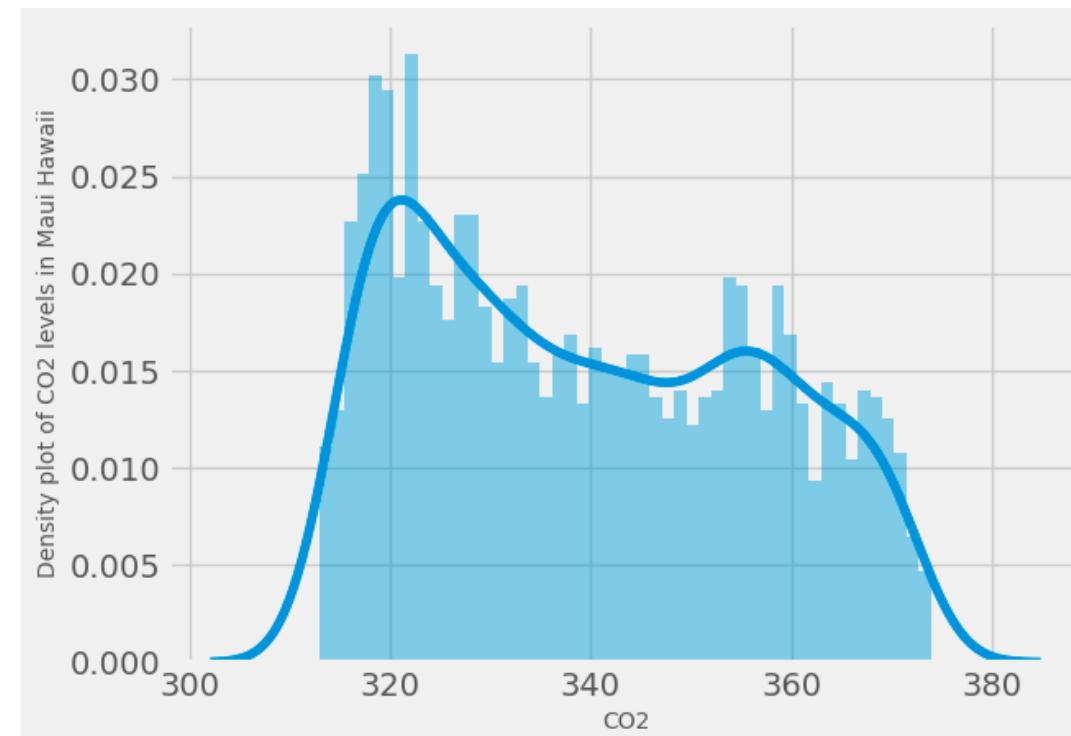
```
import seaborn as sns

ax = sns.distplot(co2_levels, hist=True, kde=True, bins=50)

# Annotate labels
ax.set_xlabel('CO2', fontsize=10)
ax.set_ylabel('Density plot of CO2 levels in Maui Hawaii', fontsize=10)

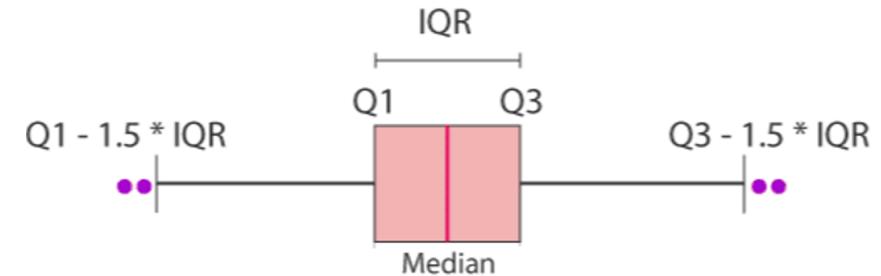
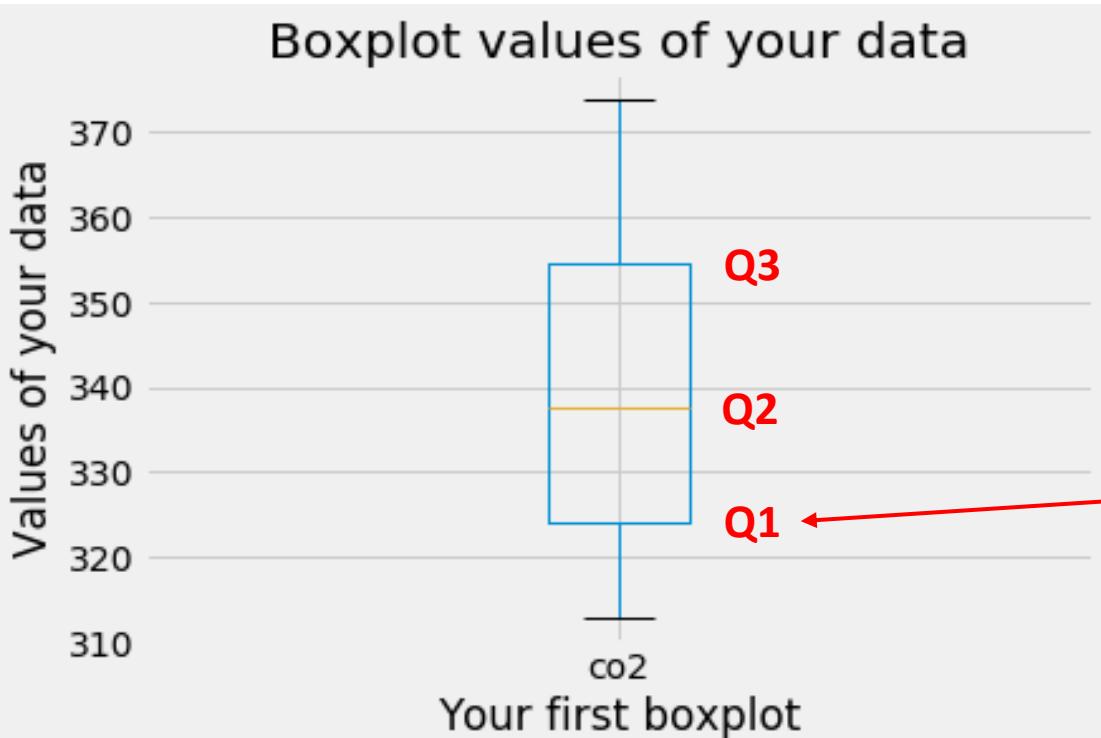
plt.show()
```

kernel density estimate



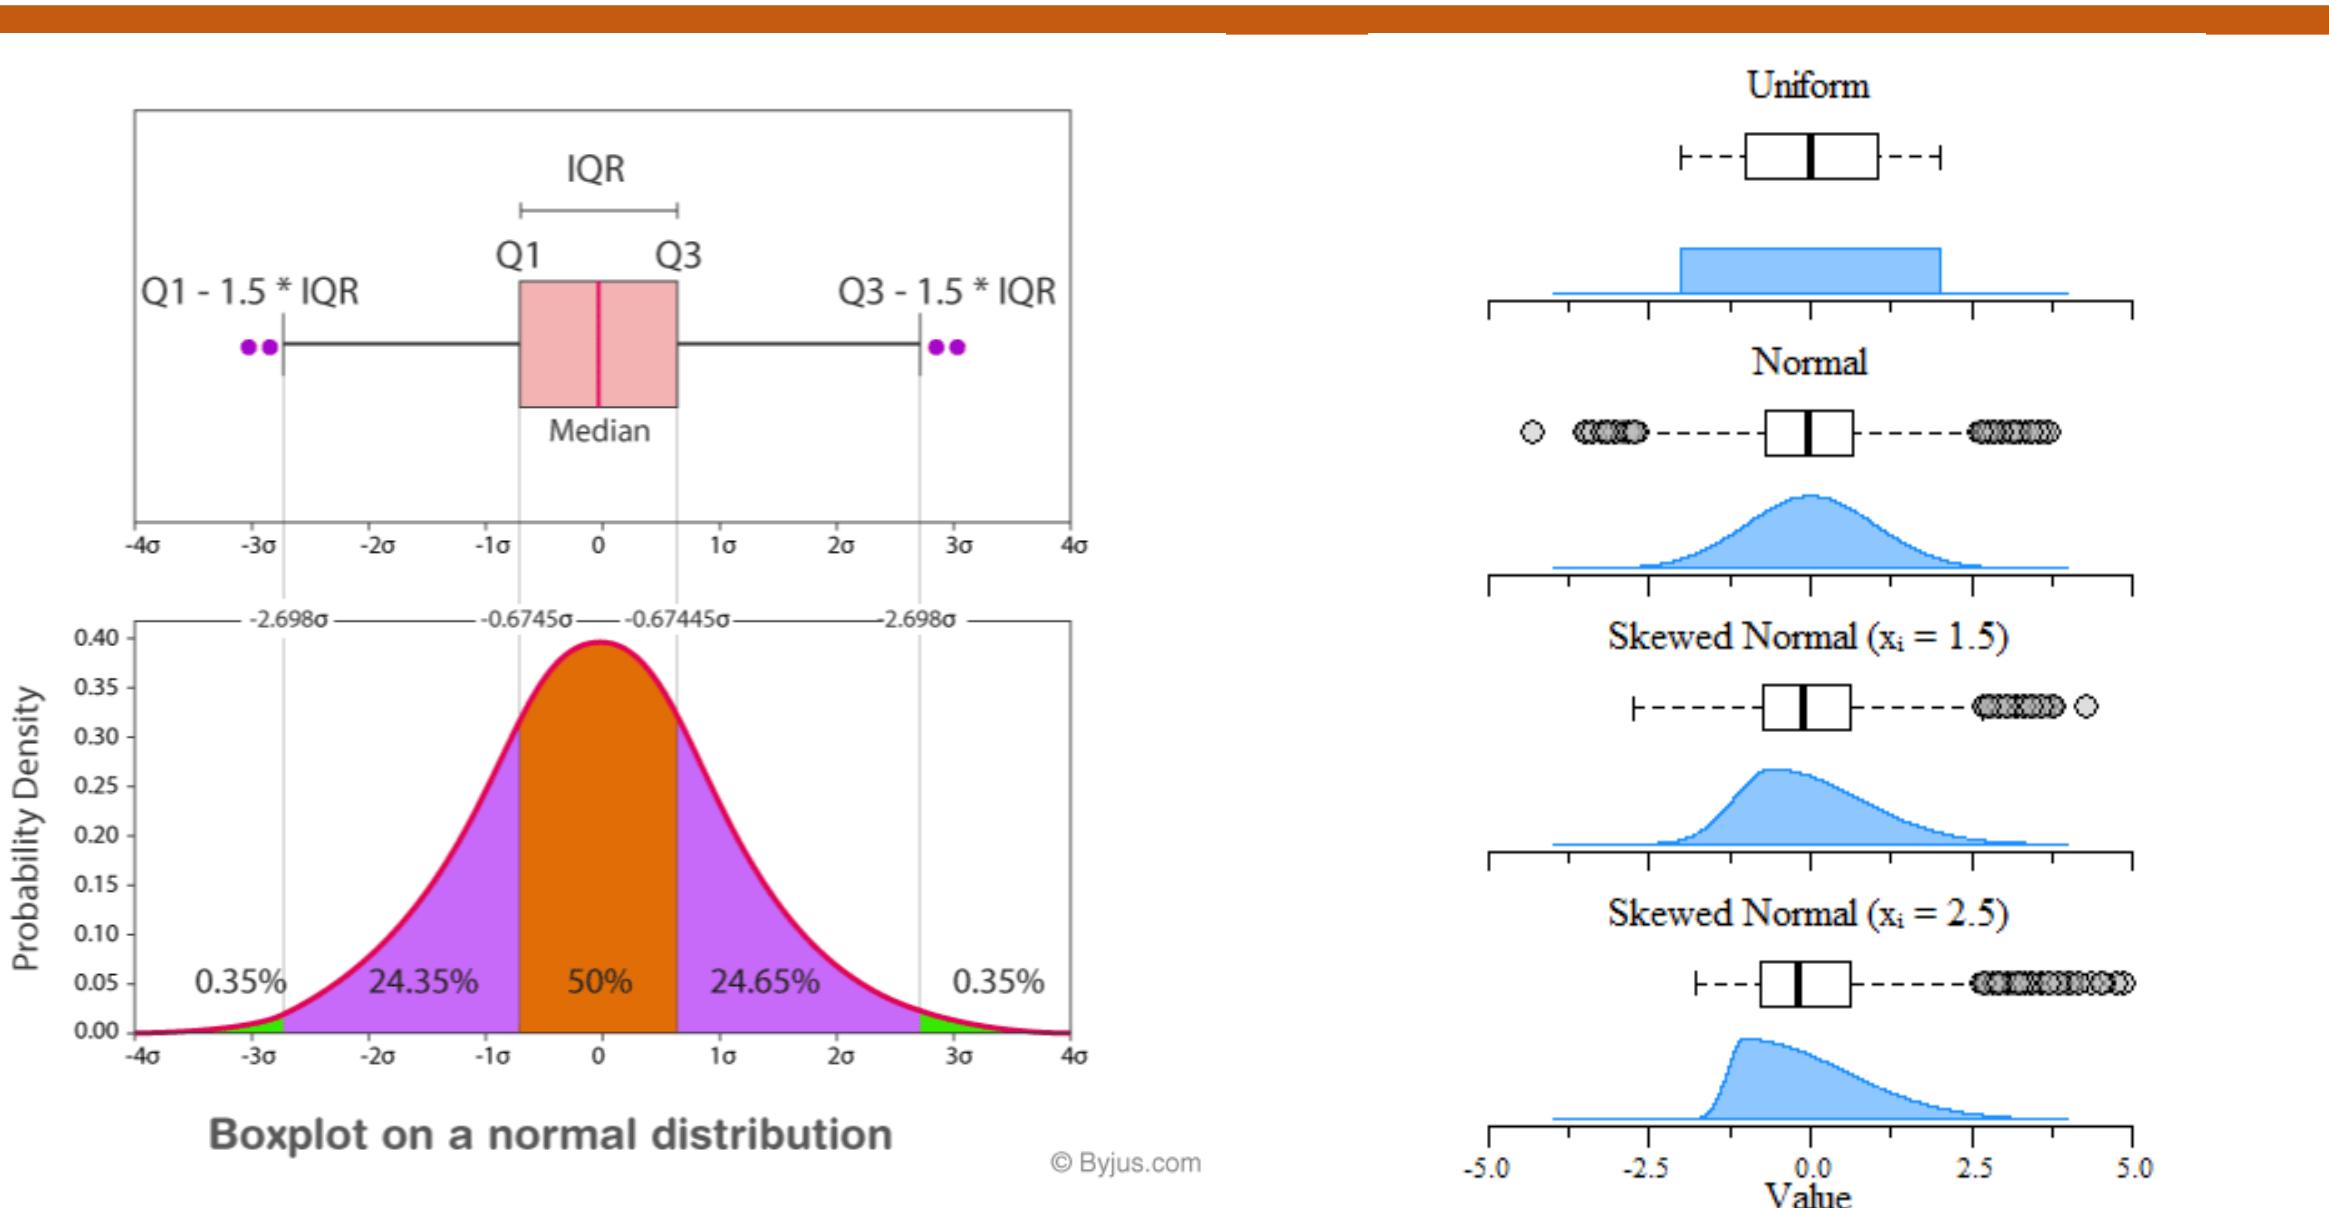
# Box Plot

```
ax1 = df.boxplot()  
ax1.set_xlabel('Your first boxplot')  
ax1.set_ylabel('Values of your data')  
ax1.set_title('Boxplot values of your data')  
  
plt.show()
```



	CO2
count	2284.000000
mean	339.657750
std	17.100899
min	313.000000
25% (Q1)	323.975000
50% (Q2)	337.700000
75% (Q3)	354.500000
max	373.900000

# Box Plot



# Outline

- Line Plots
- Summary Statistics and Diagnostics
- **Seasonality, Trend and Noise**
- Work with Multiple Time Series
- Case Study: Unemployment Rate

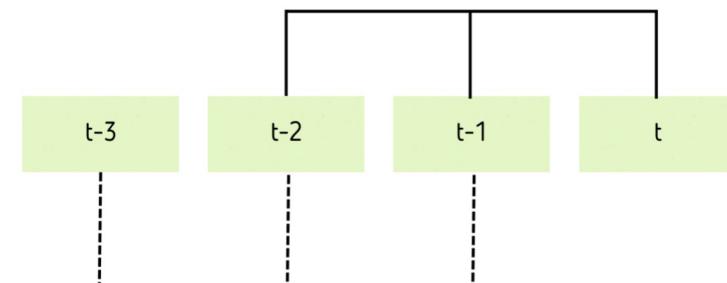
# Statsmodels

**Statsmodels** is a Python module that provides classes and functions for the estimation of many different statistical models, as well as for conducting **statistical tests**, and **statistical data exploration**.

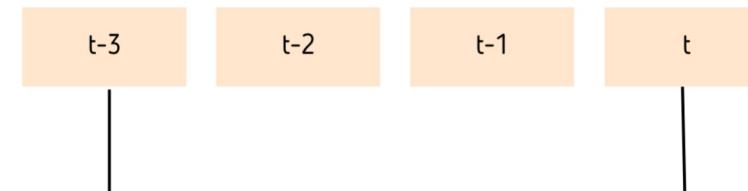


# Autocorrelation in time series data

- **Autocorrelation** is measured as the correlation between a time series and a delayed copy of itself.
- For example, an autocorrelation of order 3 returns the correlation between a time series at points ( $t_{-1}$ ,  $t_0$ ,  $t_1$ , ...) and its own values lagged by 3 time points, i.e. ( $t_{-4}$ ,  $t_{-3}$ ,  $t_{-2}$ , ...)
- It is used to find **repetitive patterns** or **periodic signal** in time series



Indirect



Direct

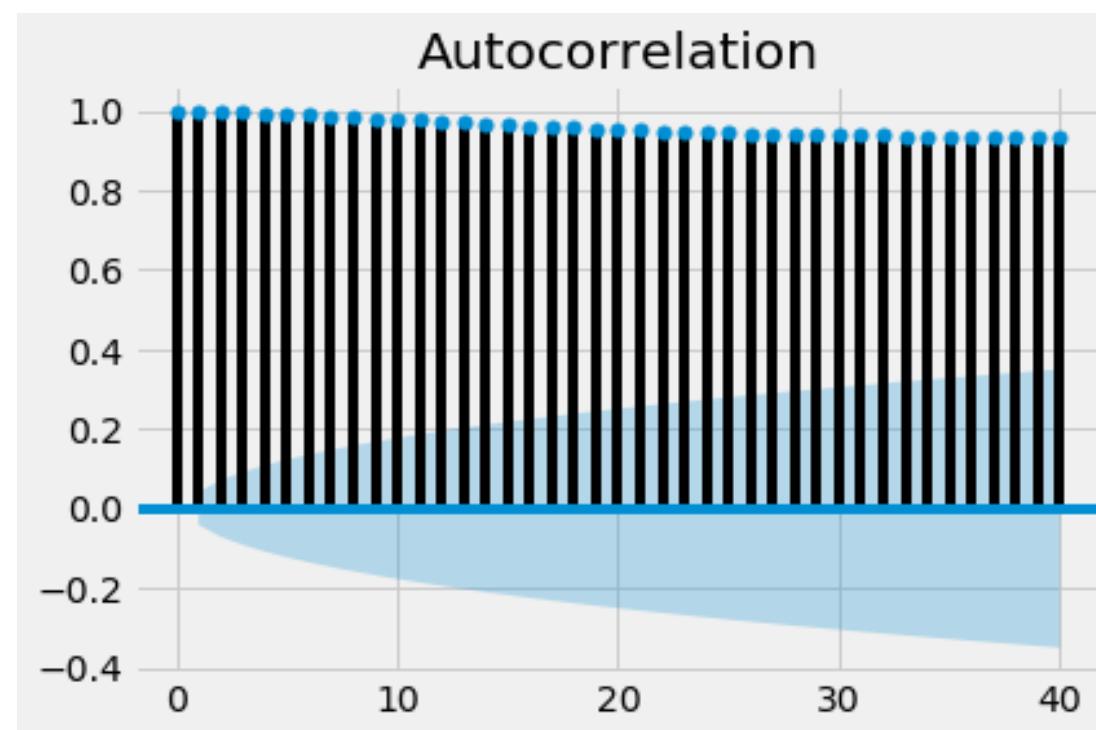
# Plotting Autocorrelations

```
import matplotlib.pyplot as plt

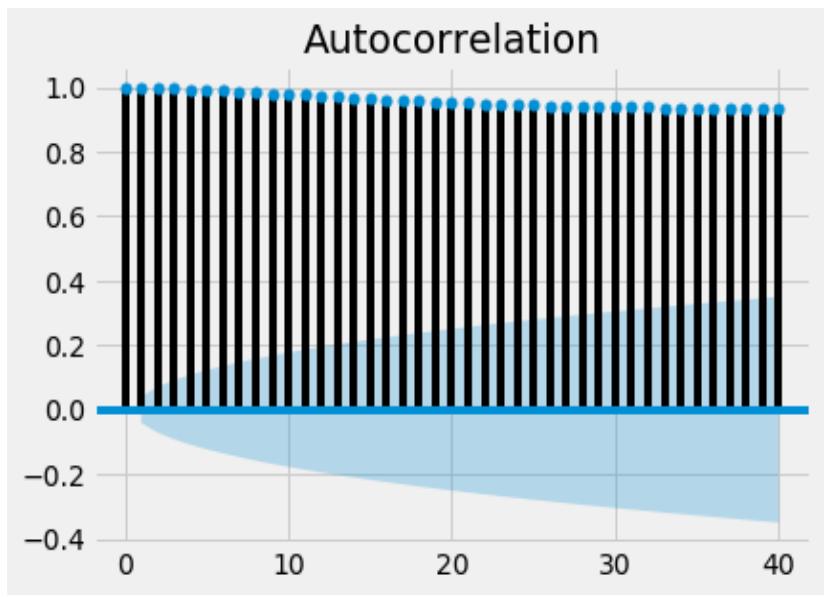
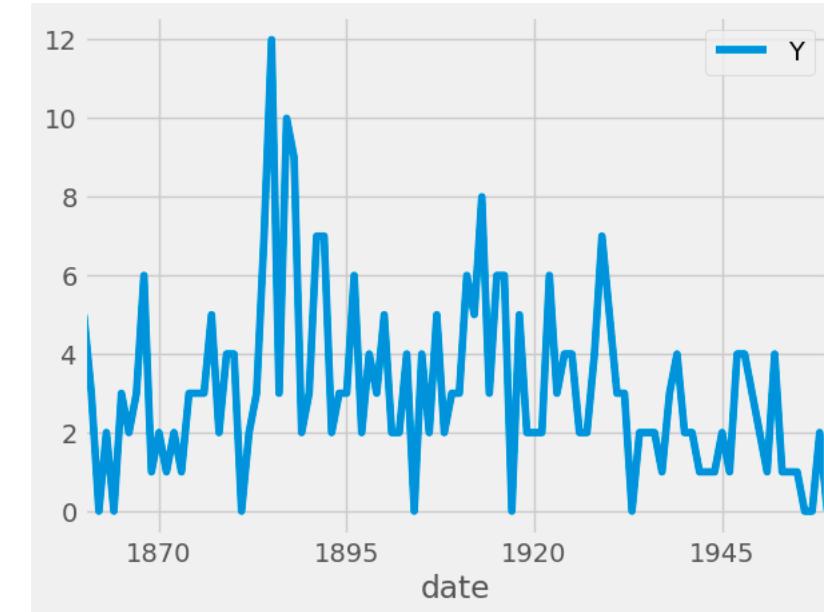
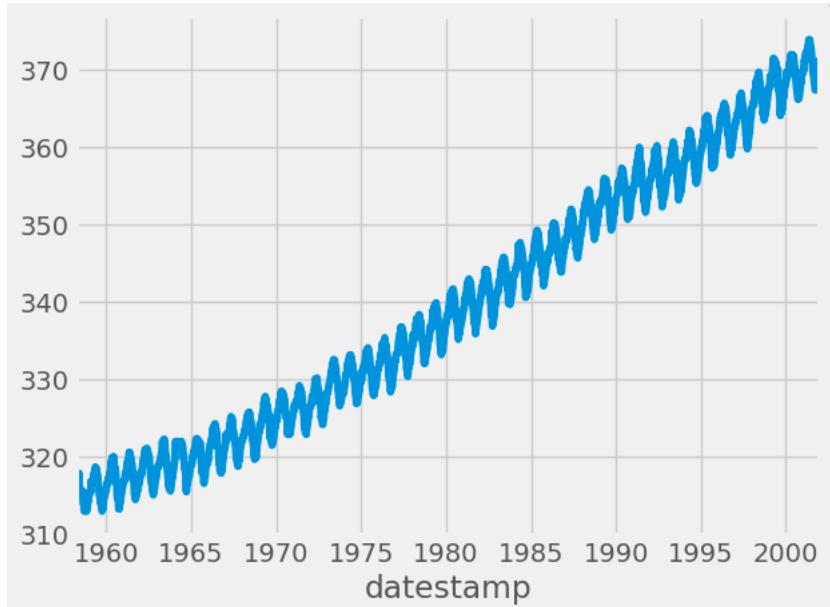
from statsmodels.graphics import tsaplots

fig = tsaplots.plot_acf(co2_levels['co2'], lags=40)

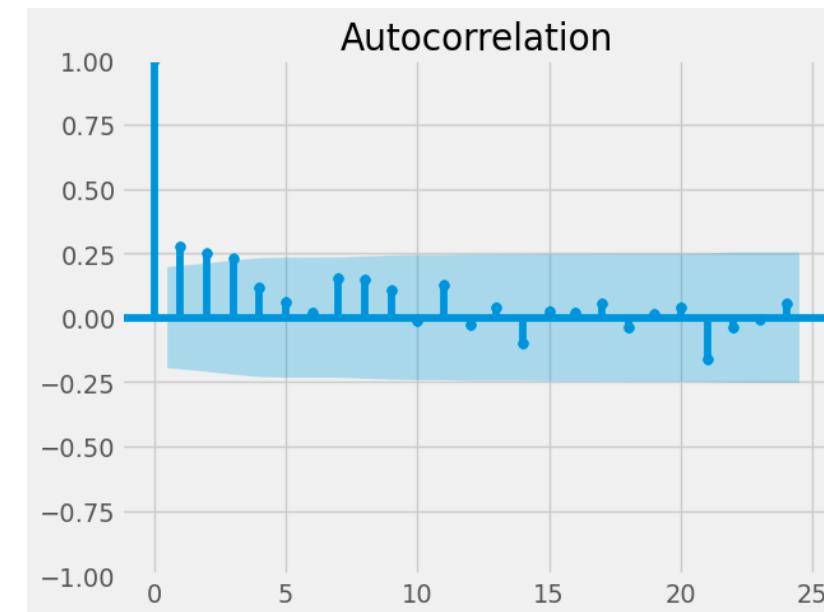
plt.show()
```



ACF Plot  
(Autocorrelations Function)



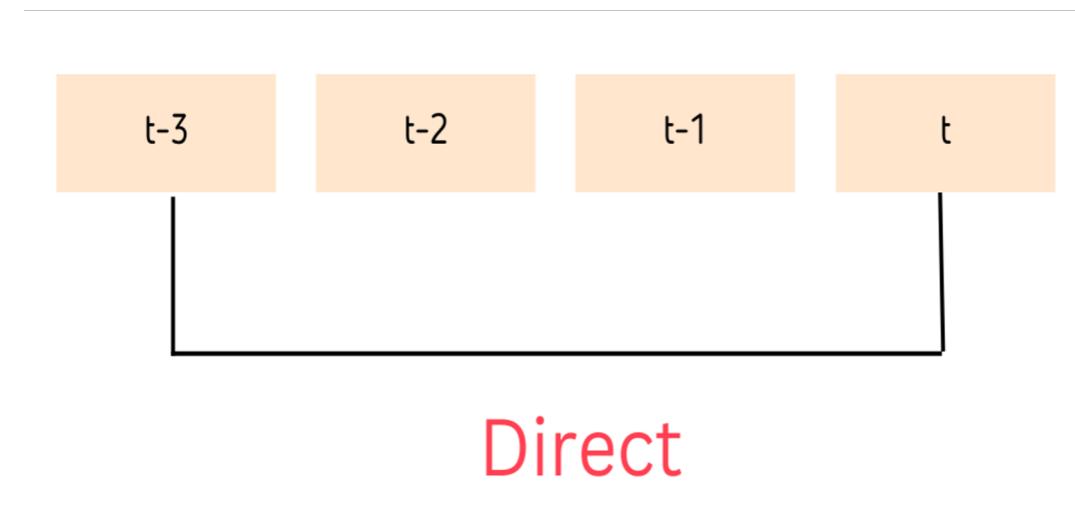
ch2\_co2\_levels.csv



ch1\_discoveries.csv

# Partial autocorrelation in time series data

- Contrary to autocorrelation, **partial autocorrelation** removes the effect of previous time points.
- For example, a partial autocorrelation function of **order 3** returns the correlation between our time series ( $t_1, t_2, t_3, \dots$ ) and lagged values of itself by 3 time points ( $t_4, t_5, t_6, \dots$ ), but only after removing all effects attributable to lags 1 and 2.



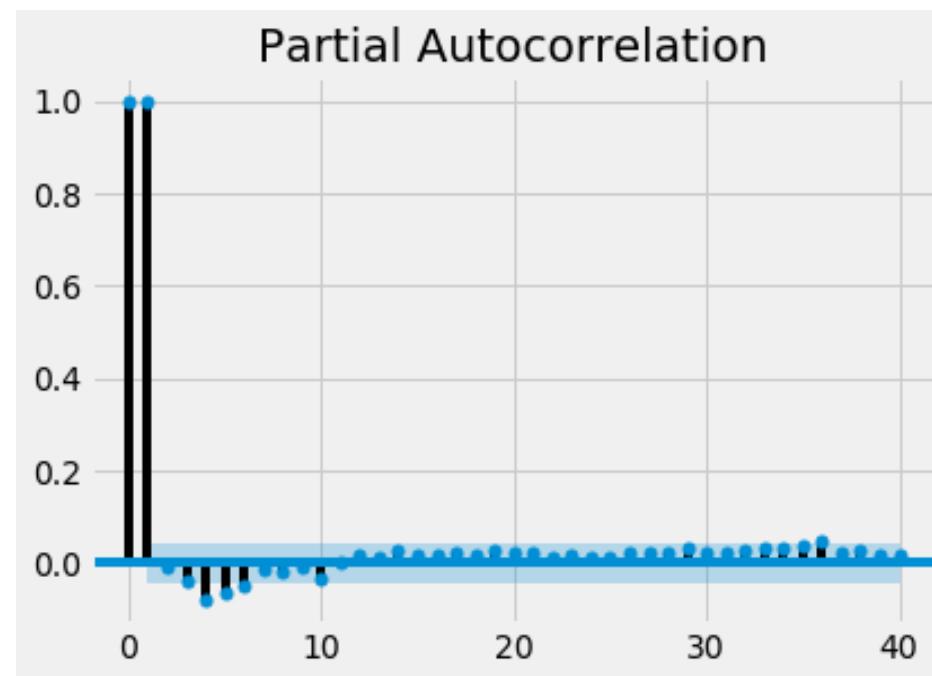
# Plotting Partial Autocorrelations

```
import matplotlib.pyplot as plt

from statsmodels.graphics import tsaplots

fig = tsaplots.plot_pacf(co2_levels['co2'], lags=40)

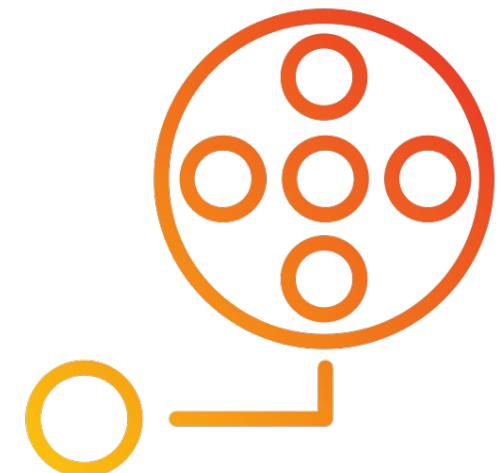
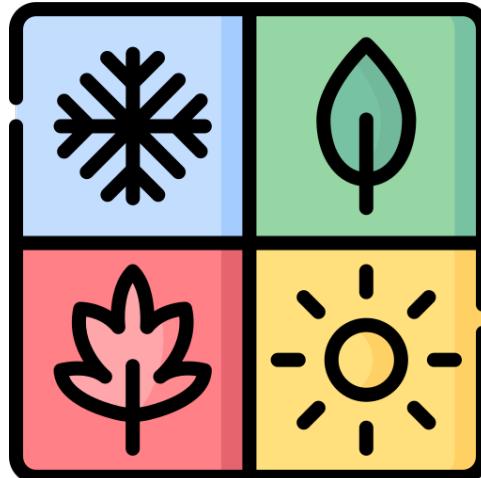
plt.show()
```



PACF Plot  
(Partial Autocorrelations Function)

# Properties of time series

- **Seasonality**: does the data display a clear **periodic pattern**?
- **Trend**: does the data follow a **consistent** upwards or downwards slope?
- **Noise**: are there any **outlier** points or **missing** values that are not consistent with the rest of the data?

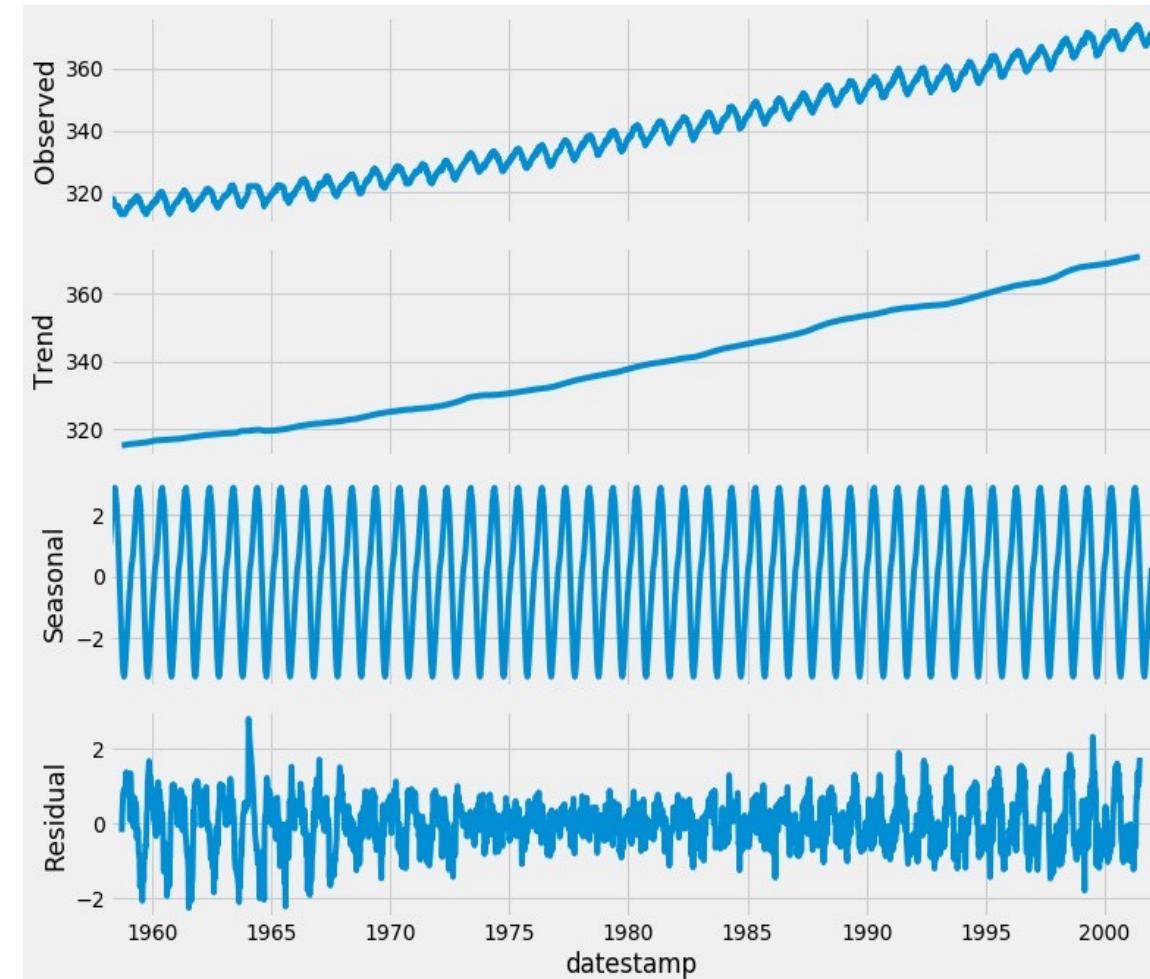


# Time Series Decomposition

```
import statsmodels.api as sm
import matplotlib.pyplot as plt
from pylab import rcParams

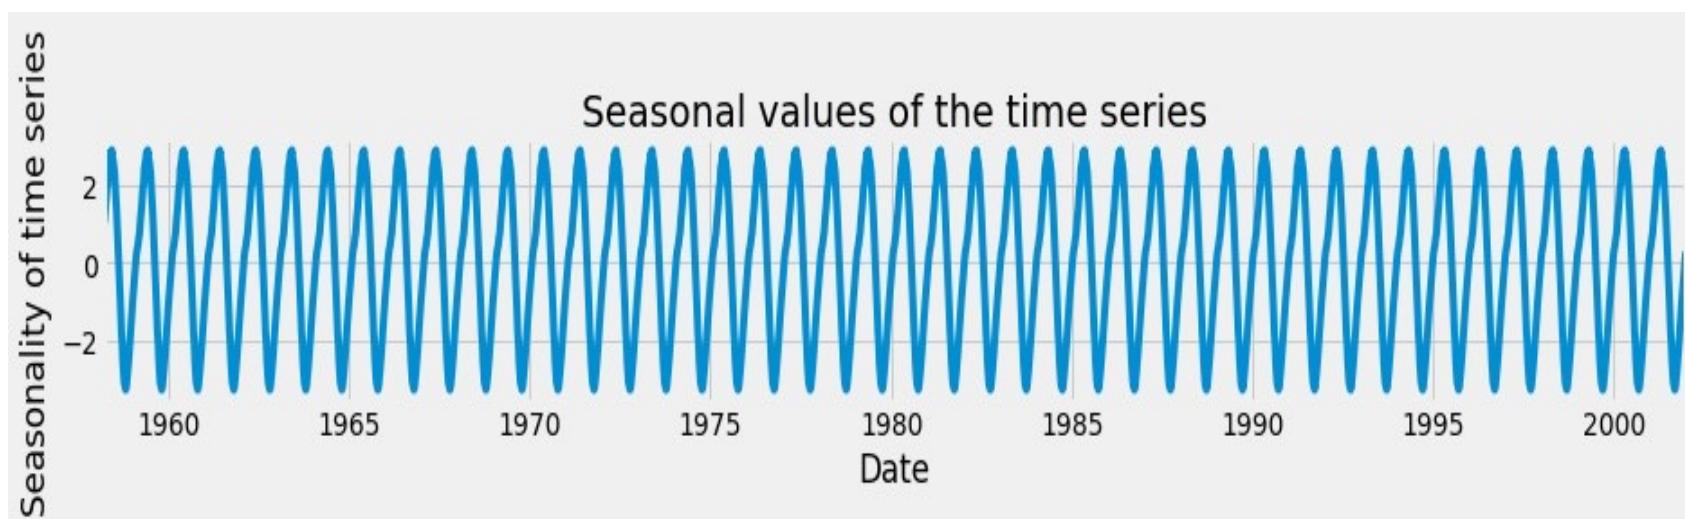
rcParams['figure.figsize'] = 11, 9
decomposition = sm.tsa.seasonal_decompose(
    co2_levels['co2'])
fig = decomposition.plot()

plt.show()
```



# Seasonality Component

```
decomposition = sm.tsa.seasonal_decompose(  
    co2_levels['co2'])  
  
decomp_seasonal = decomposition.seasonal  
  
ax = decomp_seasonal.plot(figsize=(14, 2))  
ax.set_xlabel('Date')  
ax.set_ylabel('Seasonality of time series')  
ax.set_title('Seasonal values of the time series')  
  
plt.show()
```



# Outline

- Line Plots
- Summary Statistics and Diagnostics
- Seasonality, Trend and Noise
- **Work with Multiple Time Series**
- Case Study: Unemployment Rate

# Working with multiple time series

## An isolated time series

date	ts1
1949-01	112
1949-02	118
1949-03	132

## A file with multiple time series

date	ts1	ts2	ts3	ts4	ts5	ts6	ts7
2012-01-01	21138	10.4	1987.0	12.1	3091.8	43.2	476.7
2012-02-01	2009.0	9.8	1882.9	12.3	2954.0	38.8	466.8
2012-03-01	2159.8	10.0	1987.9	14.3	3043.7	40.1	502.1

# The Meat Production Dataset

```
import pandas as pd

meat = pd.read_csv("meat.csv")

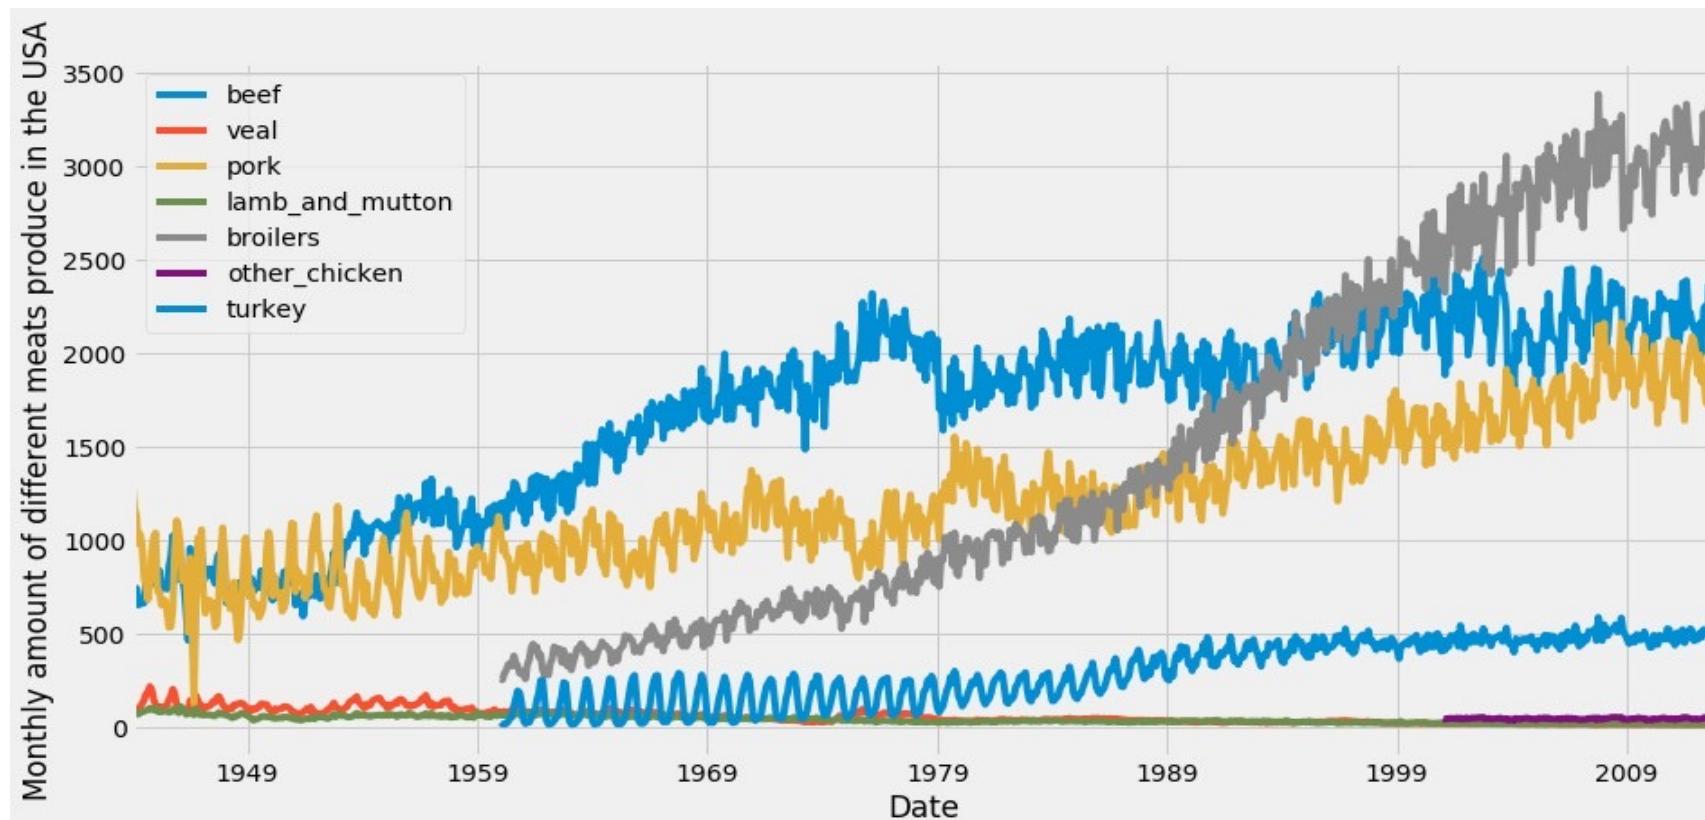
print(meat.head(10))
```

	beef	veal	pork	lamb_and_mutton	broilers	other_chicken	turkey
date							
1944-01-01	751.0	85.0	1280.0		89.0	NaN	NaN
1944-02-01	713.0	77.0	1169.0		72.0	NaN	NaN
1944-03-01	741.0	90.0	1128.0		75.0	NaN	NaN
1944-04-01	650.0	89.0	978.0		66.0	NaN	NaN
1944-05-01	681.0	106.0	1029.0		78.0	NaN	NaN
1944-06-01	658.0	125.0	962.0		79.0	NaN	NaN
1944-07-01	662.0	142.0	796.0		82.0	NaN	NaN
1944-08-01	787.0	175.0	748.0		87.0	NaN	NaN
1944-09-01	774.0	182.0	678.0		91.0	NaN	NaN
1944-10-01	834.0	215.0	777.0		100.0	NaN	NaN

# Summarizing and Plotting Multiple Time Series

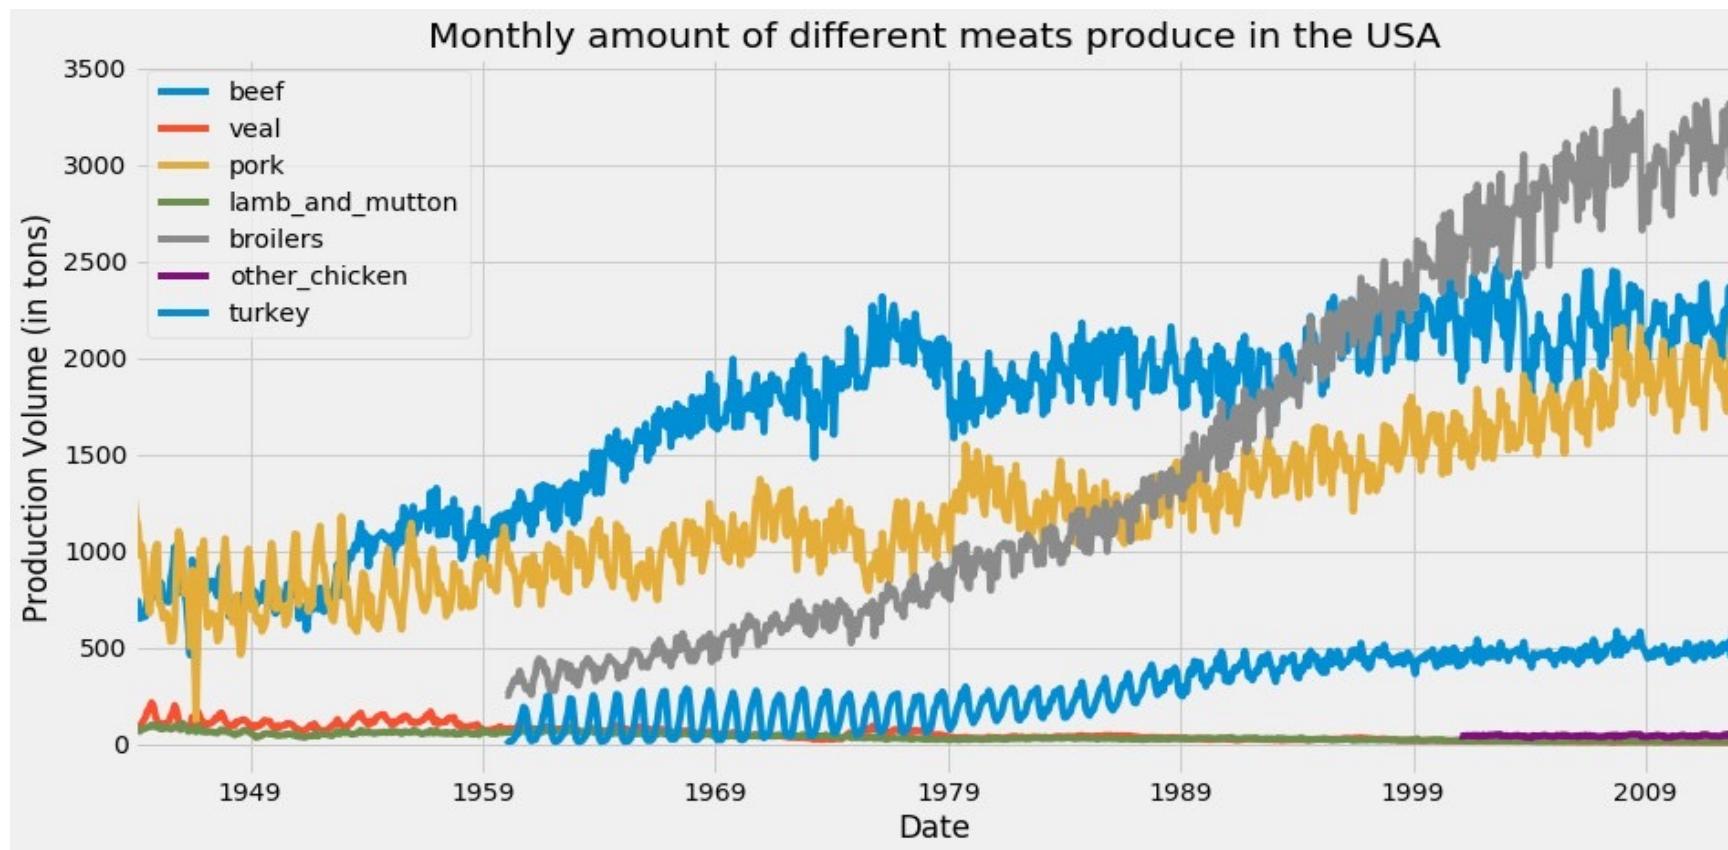
```
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
ax = df.plot(figsize=(12, 4), fontsize=14)

plt.show()
```



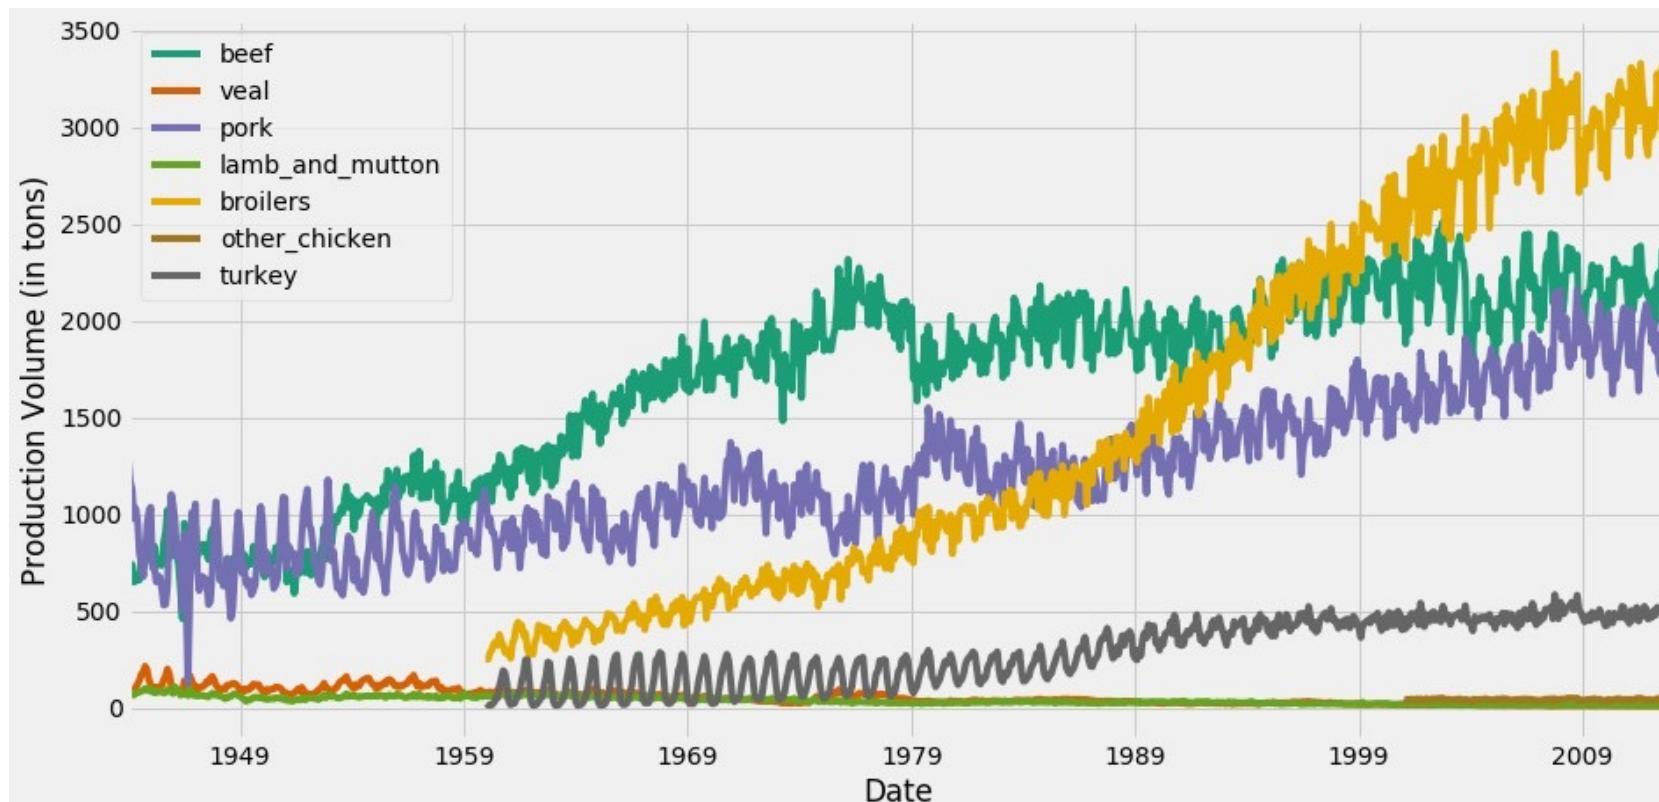
# Clarity is key

In this plot, the default **matplotlib** color scheme assigns the same color to the **beef** and **turkey** time series.



# The colormap argument

```
ax = df.plot(colormap='Dark2', figsize=(14, 7))  
  
ax.set_xlabel('Date')  
  
ax.set_ylabel('Production Volume (in tons)')  
  
plt.show()
```



# Enhancing your plot with information

```
ax = df.plot(colormap='Dark2', figsize=(14, 7))

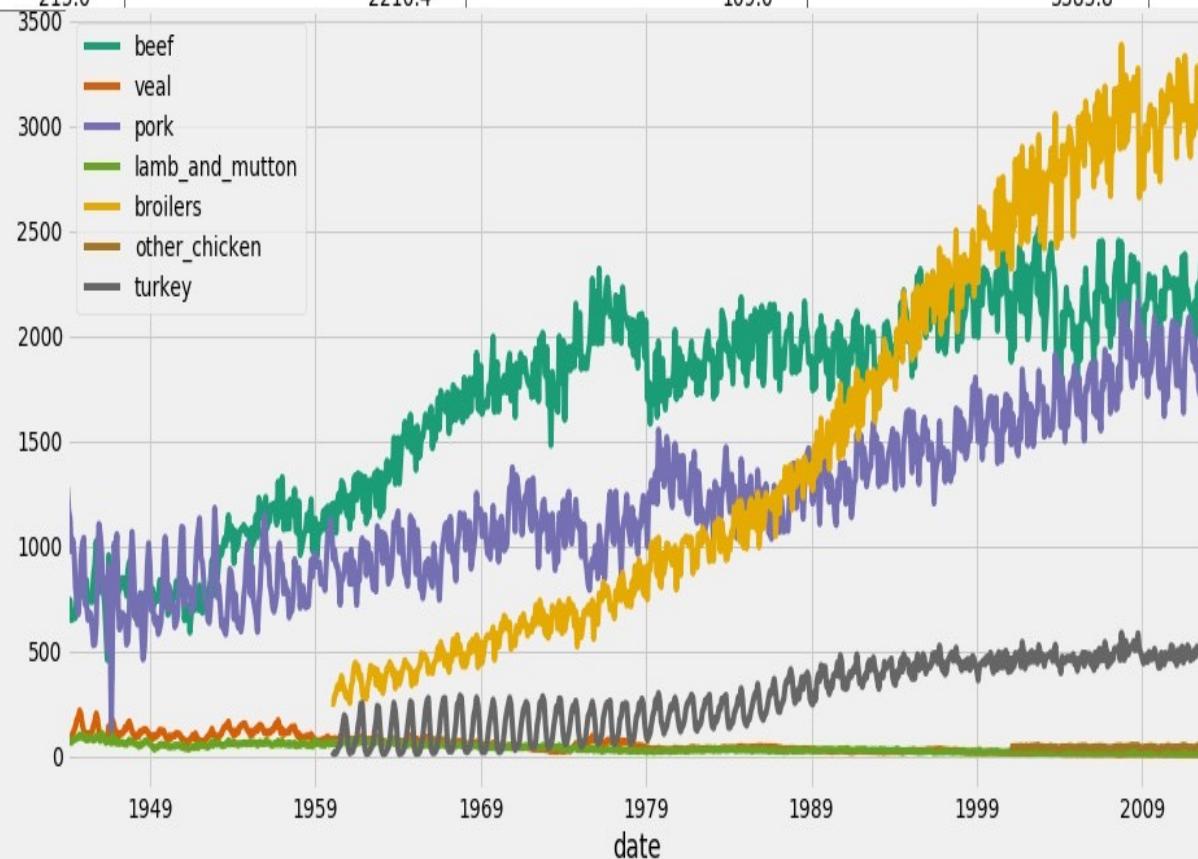
df_summary = df.describe()

# Specify values of cells in the table
ax.table(cellText=df_summary.values,
          # Specify width of the table
          colWidths=[0.3]*len(df.columns),
          # Specify row labels
          rowLabels=df_summary.index,
          # Specify column labels
          colLabels=df_summary.columns,
          # Specify location of the table
          loc='top')

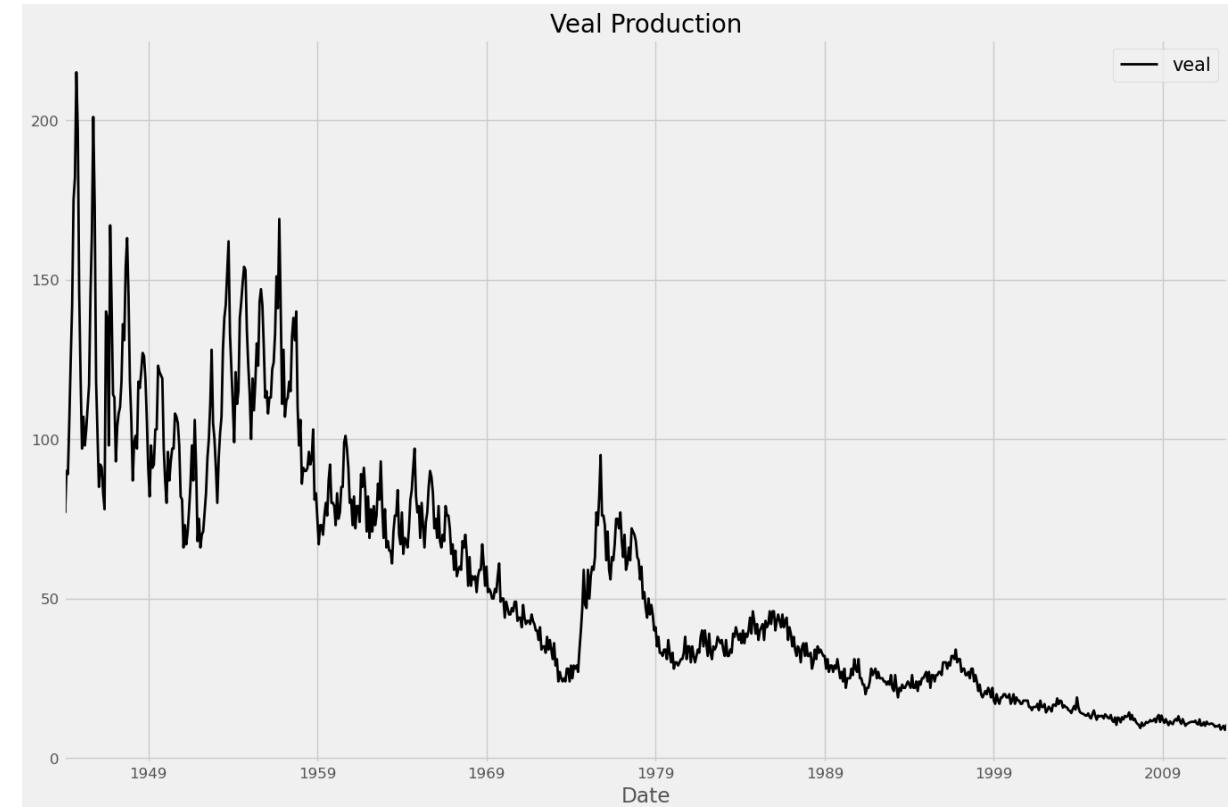
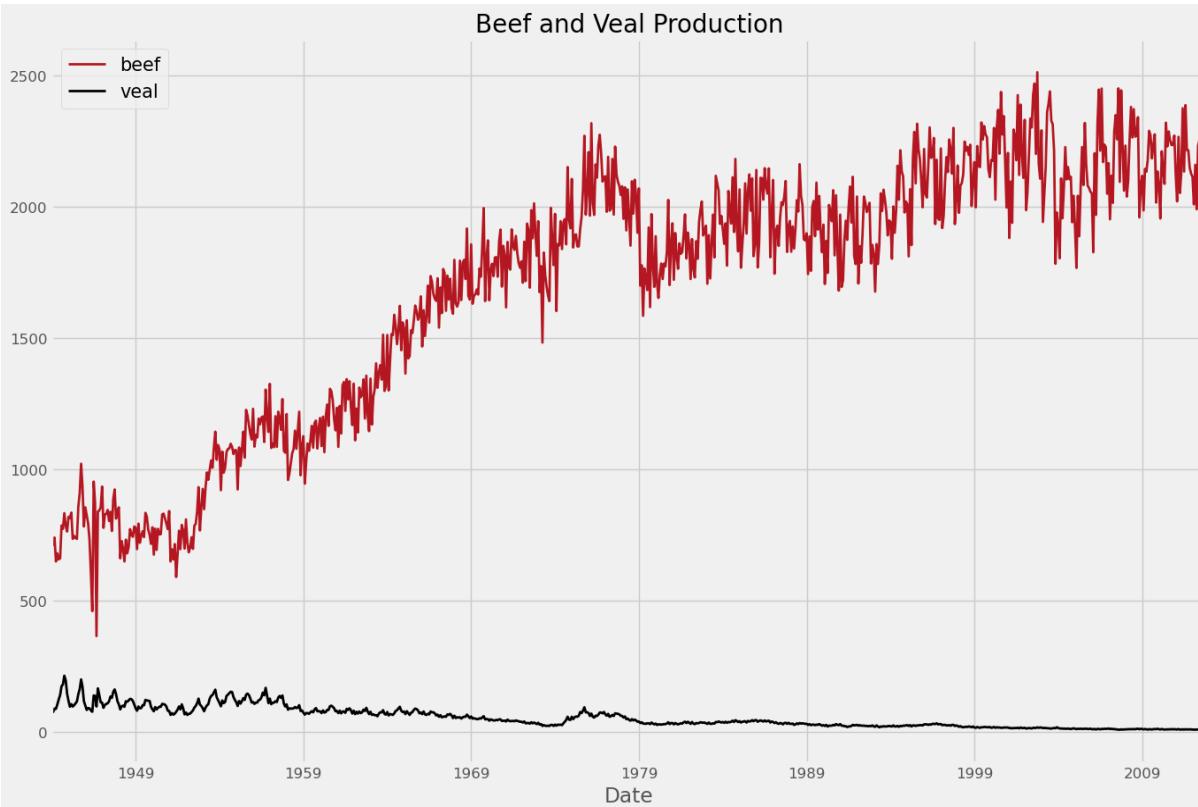
plt.show()
```

# Adding Statistical summaries to your plots

	beef	veal	pork	lamb and mutton	broilers	other chicken	turkey
count	827.0	827.0	827.0	827.0	635.0	143.0	635.0
mean	1683.46336155	54.1985489722	1211.68379686	38.3607013301	1516.58251969	43.0335664336	292.814645669
std	501.698479574	39.0628037539	371.311802273	19.6243398061	963.012101054	3.86714117313	162.482637819
min	366.0	8.8	124.0	10.9	250.9	32.3	12.4
25%	1231.5	24.0	934.5	23.0	636.35	40.2	154.15
50%	1853.0	40.0	1156.0	31.0	1211.3	43.4	278.3
75%	2070.0	79.0	1466.0	55.0	2426.65	45.65	449.15
max	2512.0	215.0	2210.4	109.0	3383.8	51.1	585.1

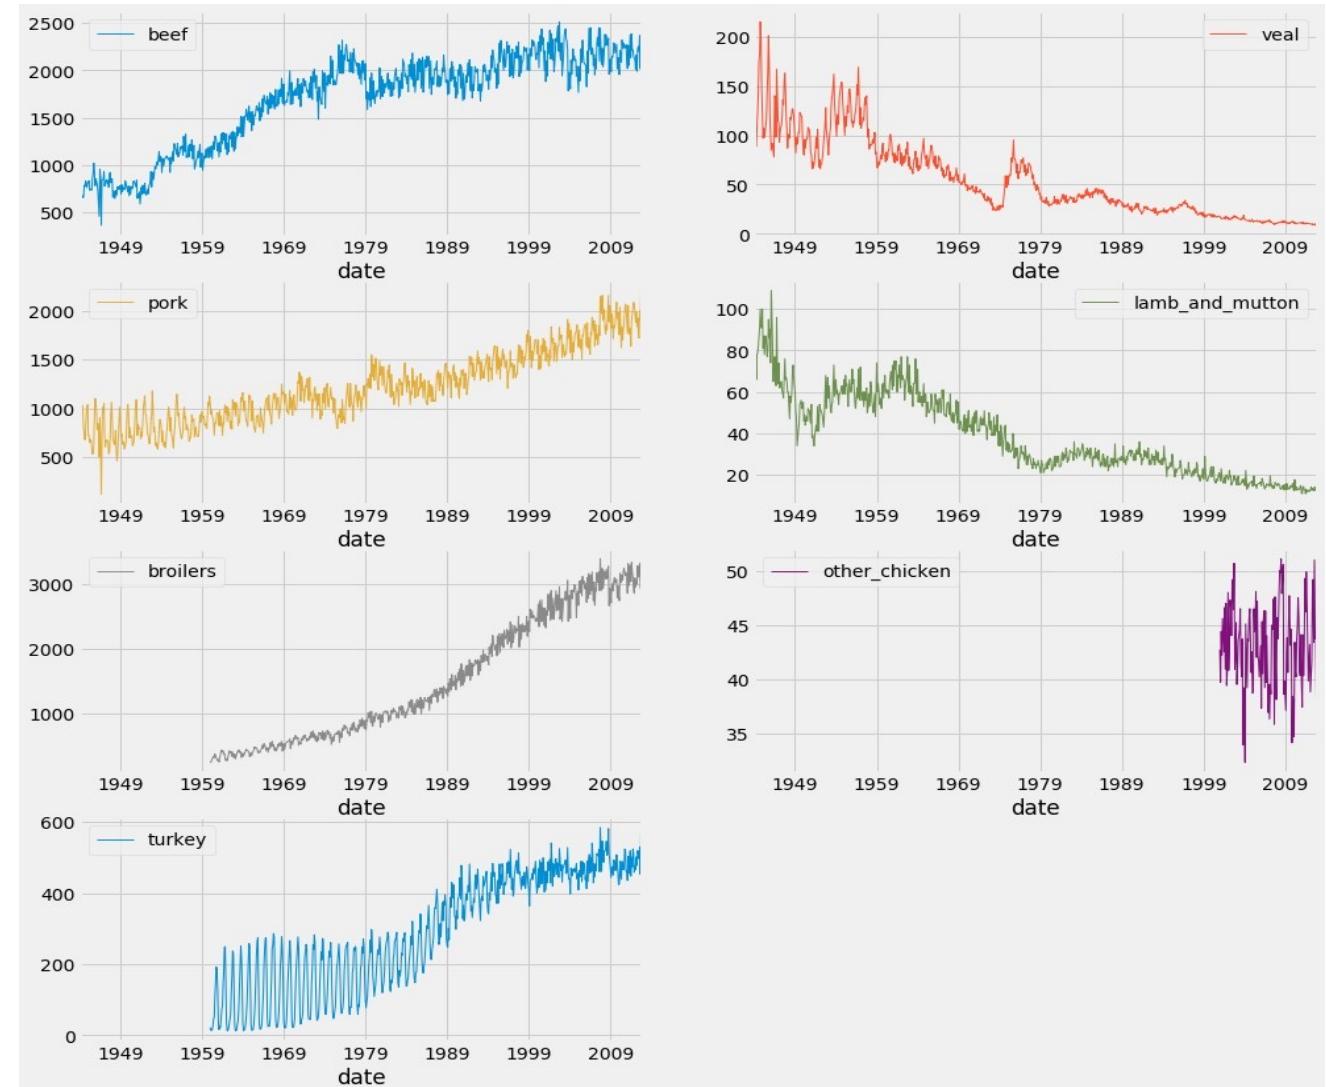


# Dealing with different scales



# Facet plots

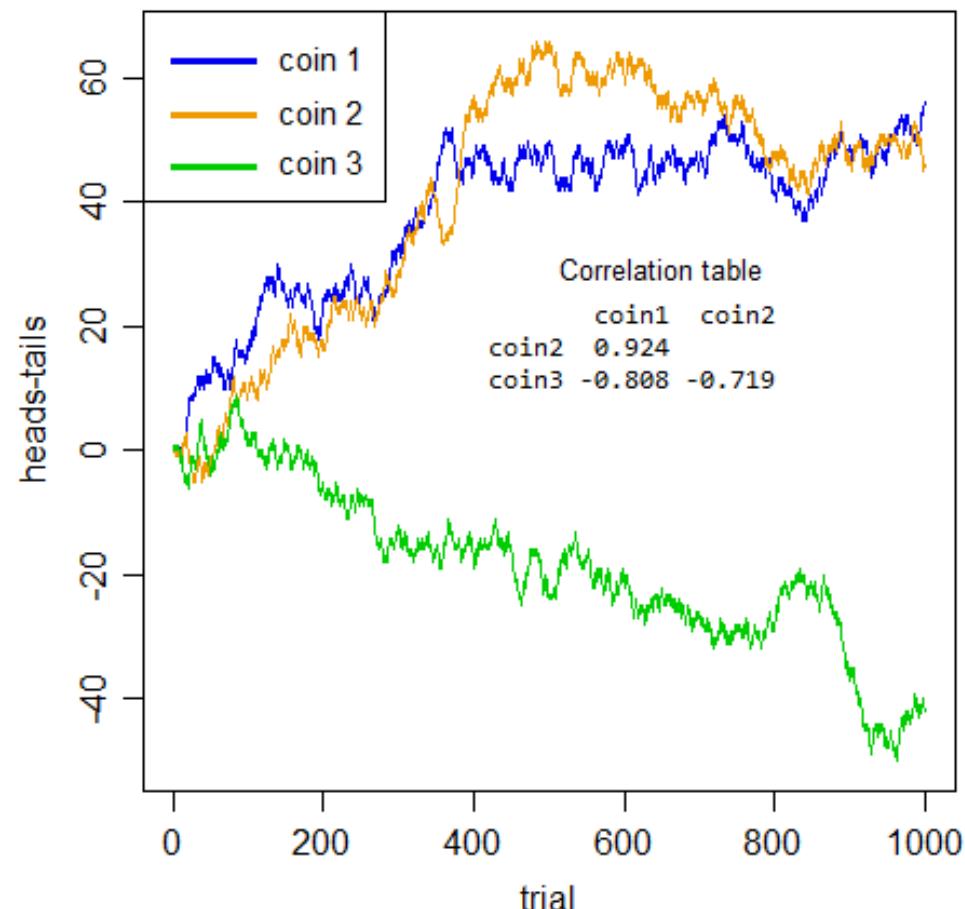
```
df.plot(subplots=True,  
        linewidth=0.5,  
        layout=(2, 4),  
        figsize=(16, 10),  
        sharex=False,  
        sharey=False)  
  
plt.show()
```



# Correlations between two variables

The **correlation coefficient** is a measure used to determine the strength or lack of relationship between two variables:

- **Pearson's coefficient** can be used to compute the correlation coefficient between variables for which the relationship is thought to be **linear**.
- **Kendall Tau** or **Spearman rank** can be used to compute the correlation coefficient between variables for which the relationship is thought to be **non-linear**.



# Compute Correlations

```
from scipy.stats.stats import pearsonr
from scipy.stats.stats import spearmanr
from scipy.stats.stats import kendalltau

x = [1, 2, 4, 7]
y = [1, 3, 4, 8]
pearsonr(x, y)
```

```
SpearmanResult(correlation=0.9843, pvalue=0.01569)
```

```
spearmanr(x, y)
```

```
SpearmanResult(correlation=1.0, pvalue=0.0)
```

```
kendalltau(x, y)
```

```
KendalltauResult(correlation=1.0, pvalue=0.0415)
```

# What is a correlation matrix?

---

When computing the correlation coefficient between more than two variables, you obtain a **correlation matrix**:

- Range: [-1, 1]
- 0: no relationship
- 1: strong positive relationship
- -1: strong negative relationship

# What is a correlation matrix?

- A correlation matrix is always "**symmetric**"
- The **diagonal values** will always be equal to **1**

	x	y	z
x	1.00	-0.46	0.49
y	-0.46	1.00	-0.61
z	0.49	-0.61	1.00

# Computing Correlation Matrices with Pandas

```
corr_p = meat[ ['beef', 'veal', 'turkey']].corr(method='pearson')
print(corr_p)
```

	beef	veal	turkey
beef	1.000	-0.829	0.738
veal	-0.829	1.000	-0.768
turkey	0.738	-0.768	1.000

```
corr_p = meat[ ['beef', 'veal', 'turkey']].corr(method='spearman')
print(corr_p)
```

	beef	veal	turkey
beef	1.000	-0.812	0.778
veal	-0.812	1.000	-0.829
turkey	0.778	-0.829	1.000

# Computing Correlation Matrices with Pandas

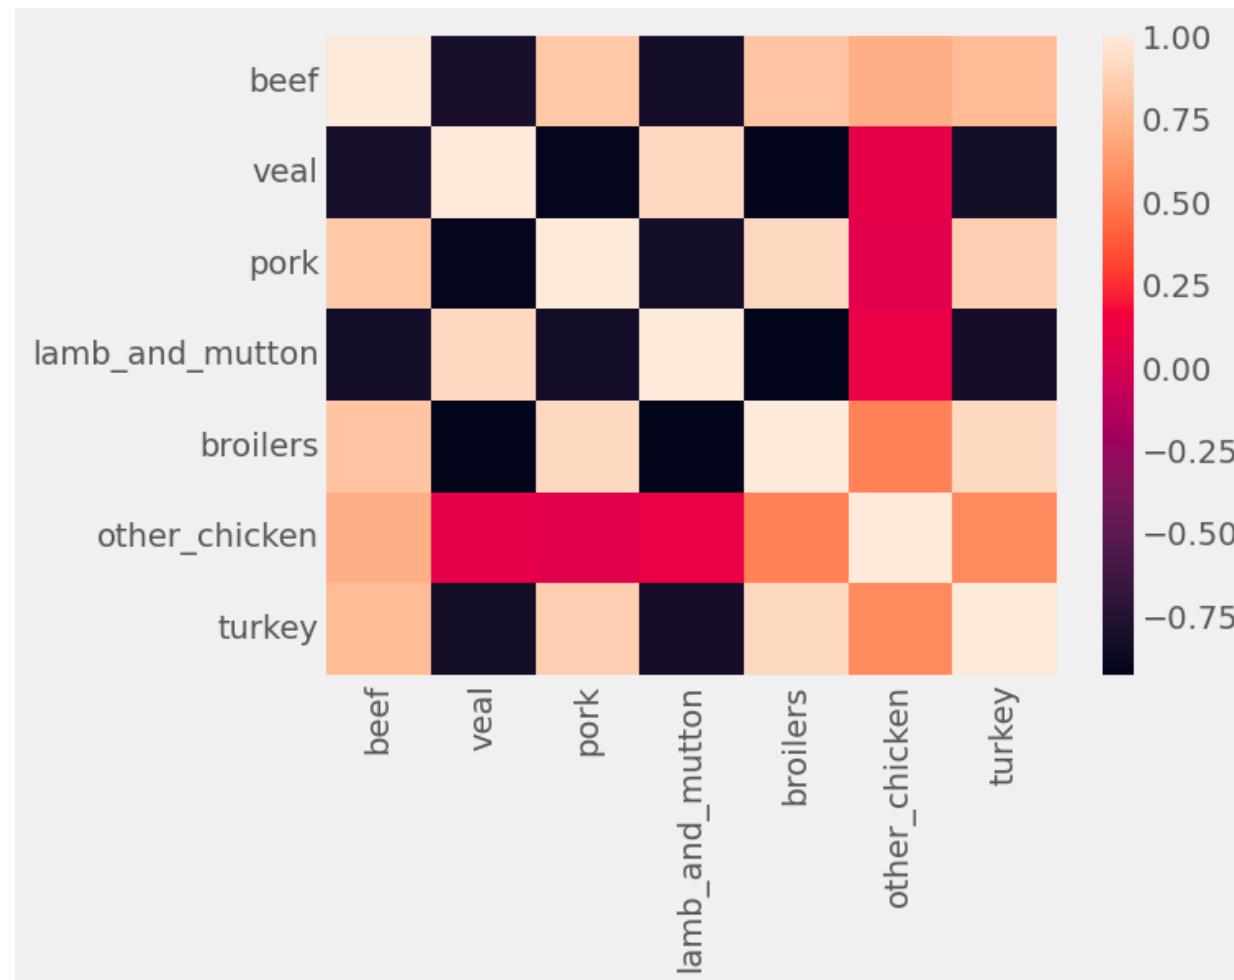
If you want to compute the correlation between all time series in your DataFrame, simply remove the references to the columns.

```
corr_mat = meat.corr(method='pearson')
```

	beef	veal	pork	lamb_and_mutton	broilers	other_chicken	turkey
date							
1944-01-01	751.0	85.0	1280.0		89.0	NaN	NaN
1944-02-01	713.0	77.0	1169.0		72.0	NaN	NaN
1944-03-01	741.0	90.0	1128.0		75.0	NaN	NaN
1944-04-01	650.0	89.0	978.0		66.0	NaN	NaN
1944-05-01	681.0	106.0	1029.0		78.0	NaN	NaN
1944-06-01	658.0	125.0	962.0		79.0	NaN	NaN
1944-07-01	662.0	142.0	796.0		82.0	NaN	NaN
1944-08-01	787.0	175.0	748.0		87.0	NaN	NaN
1944-09-01	774.0	182.0	678.0		91.0	NaN	NaN
1944-10-01	834.0	215.0	777.0		100.0	NaN	NaN

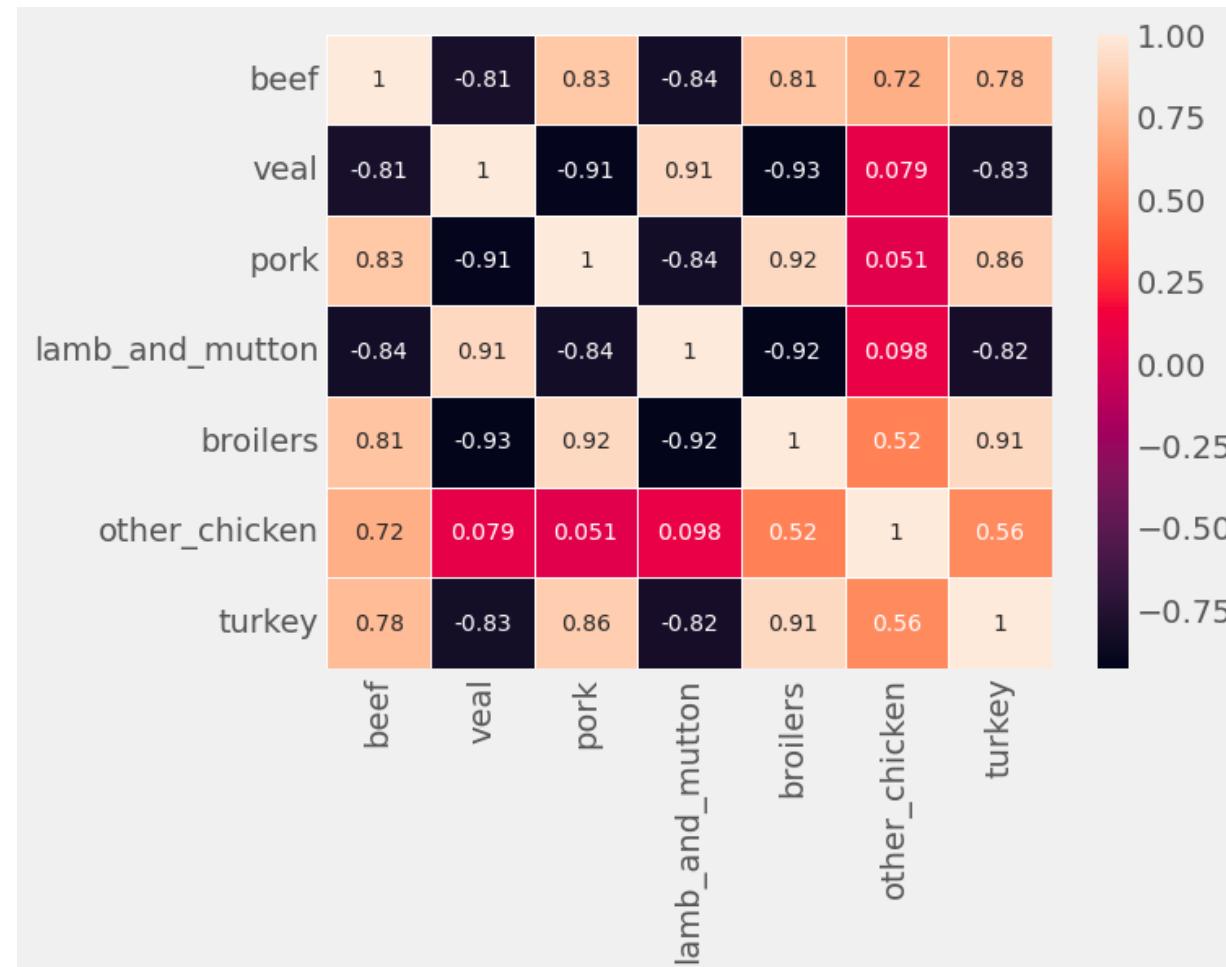
# Heatmap

```
import seaborn as sns  
  
sns.heatmap(corr_mat)
```



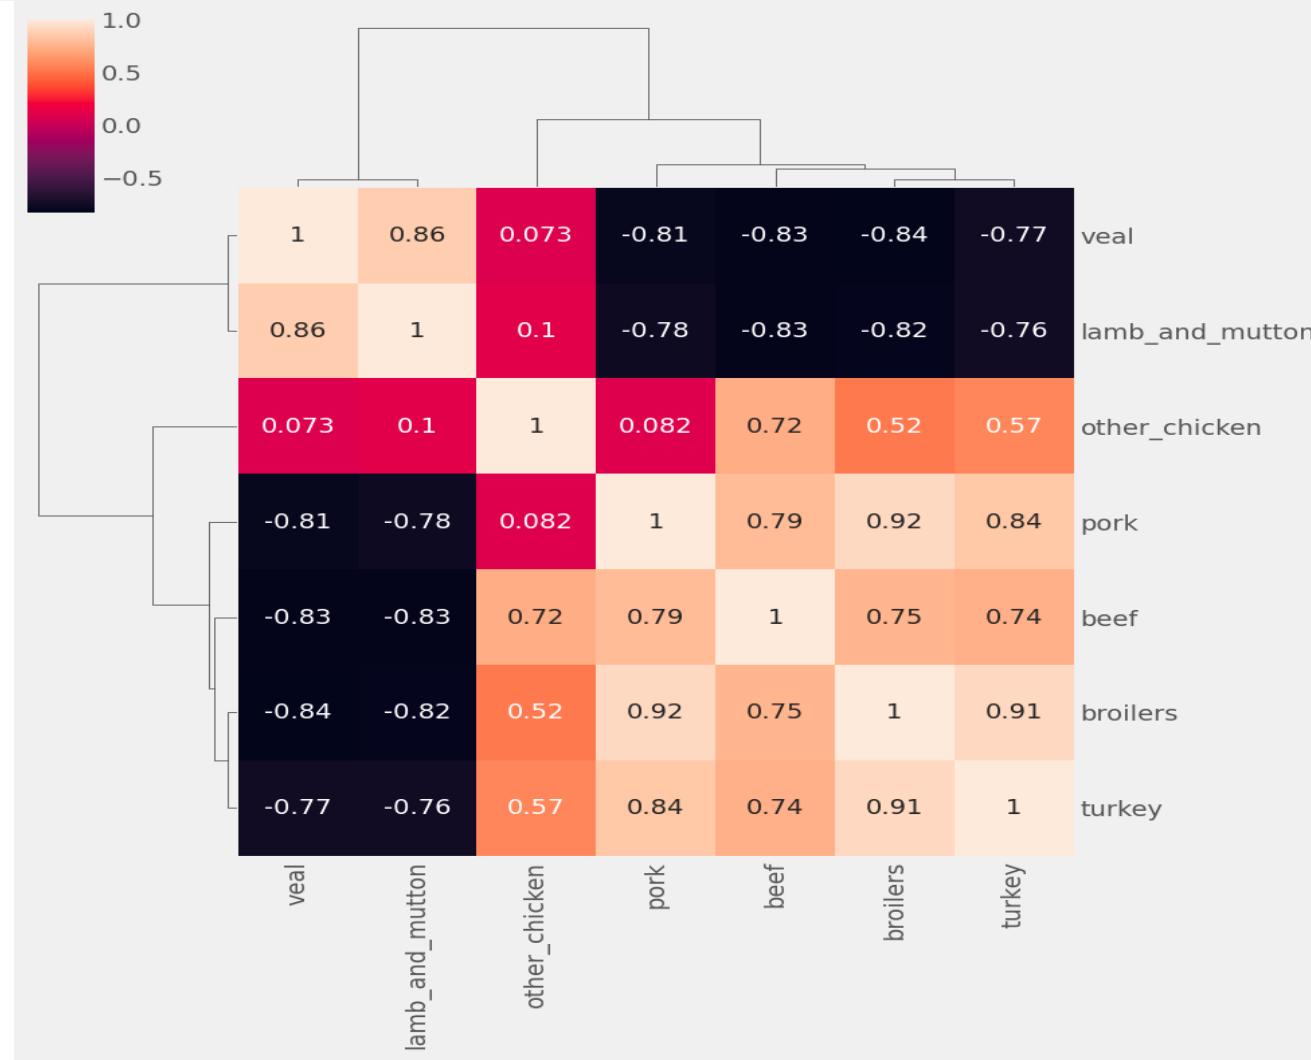
# Heatmap

```
import seaborn as sns  
  
sns.heatmap(corr_mat, annot=True)
```



# Clustermap

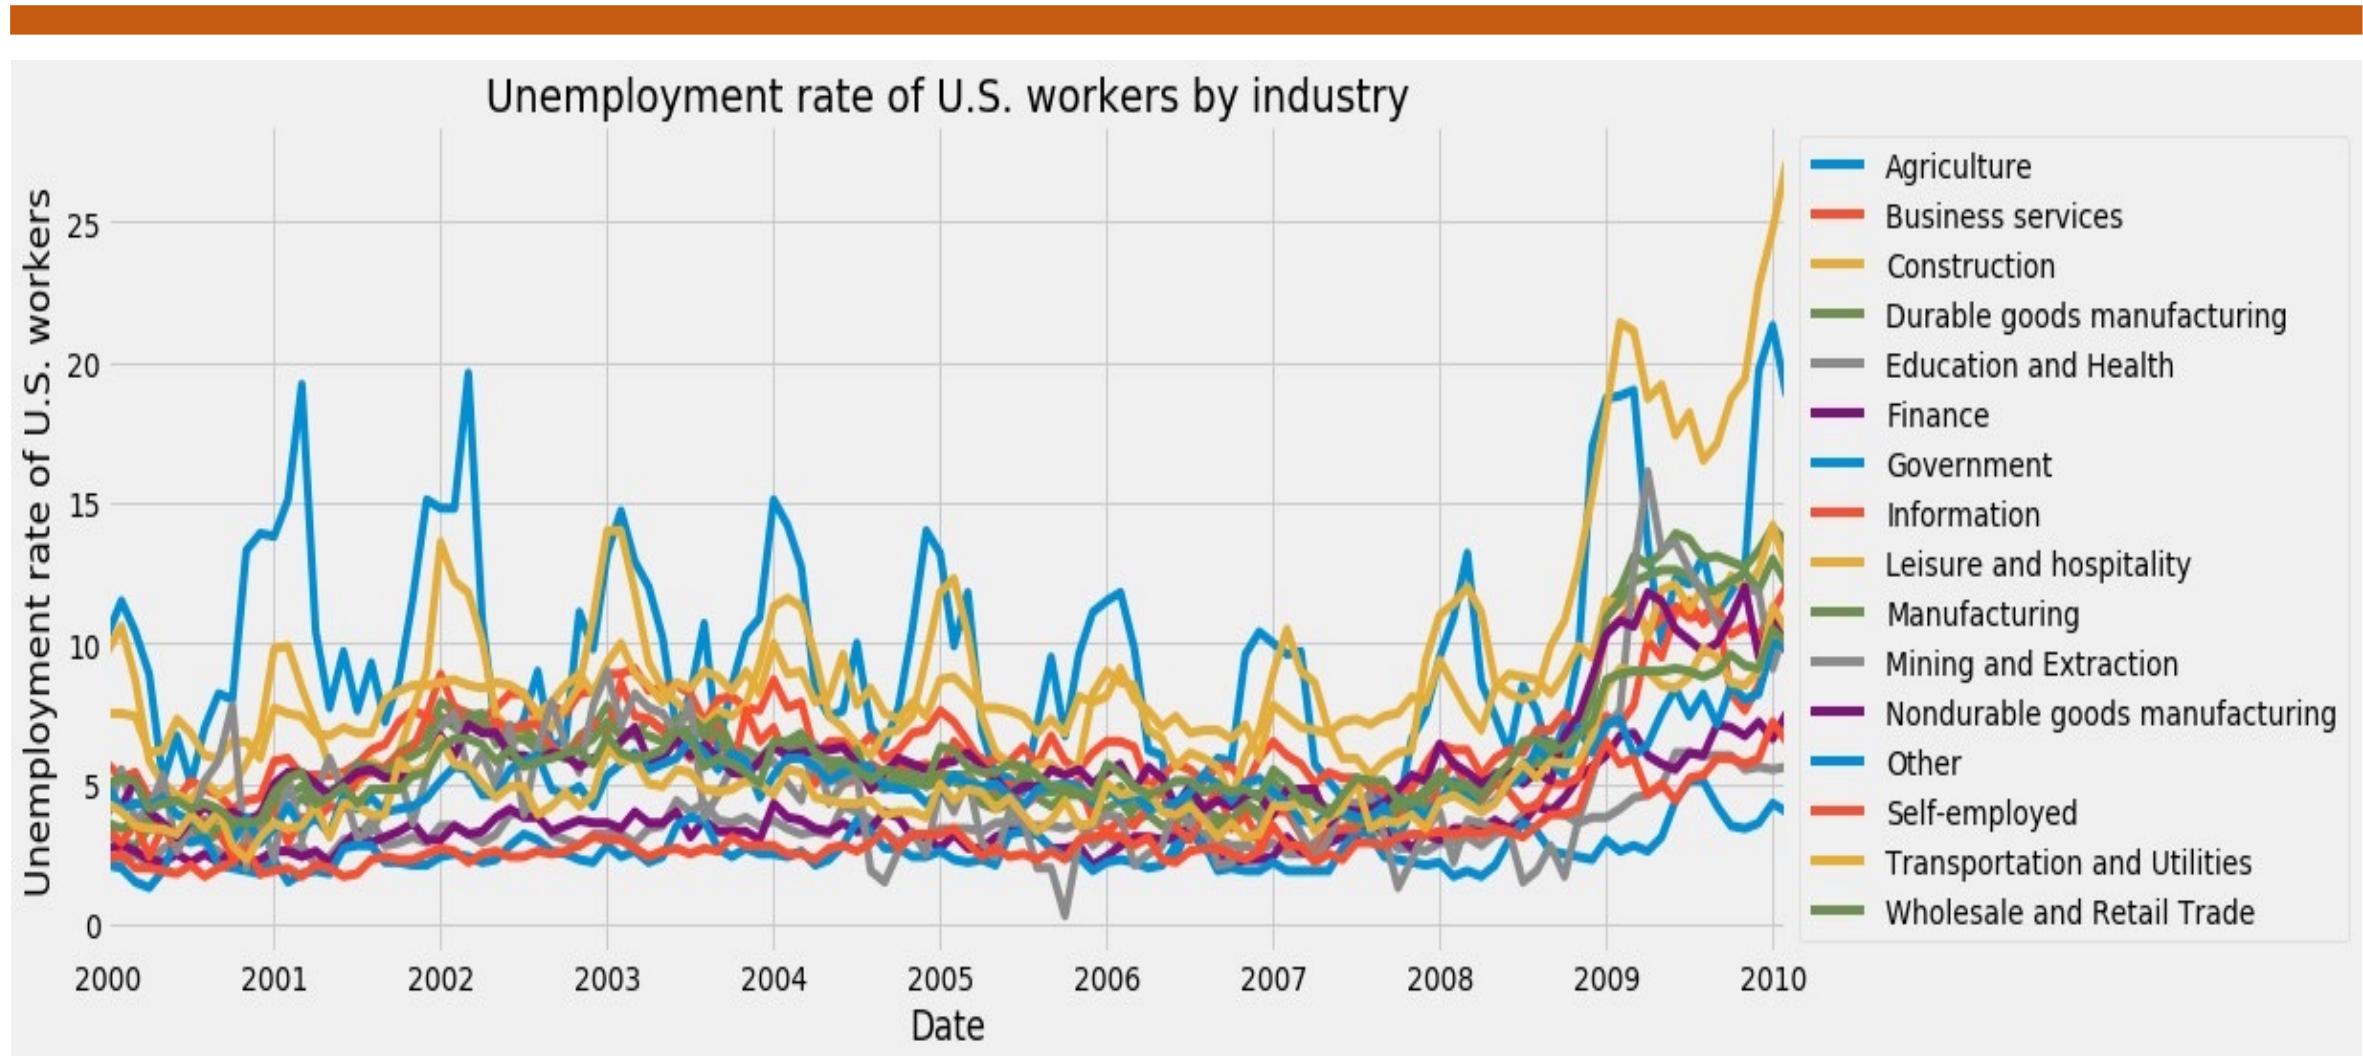
```
sns.clustermap(corr_mat)
```



# Outline

- Line Plots
- Summary Statistics and Diagnostics
- Seasonality, Trend and Noise
- Work with Multiple Time Series
- **Case Study: Unemployment Rate**

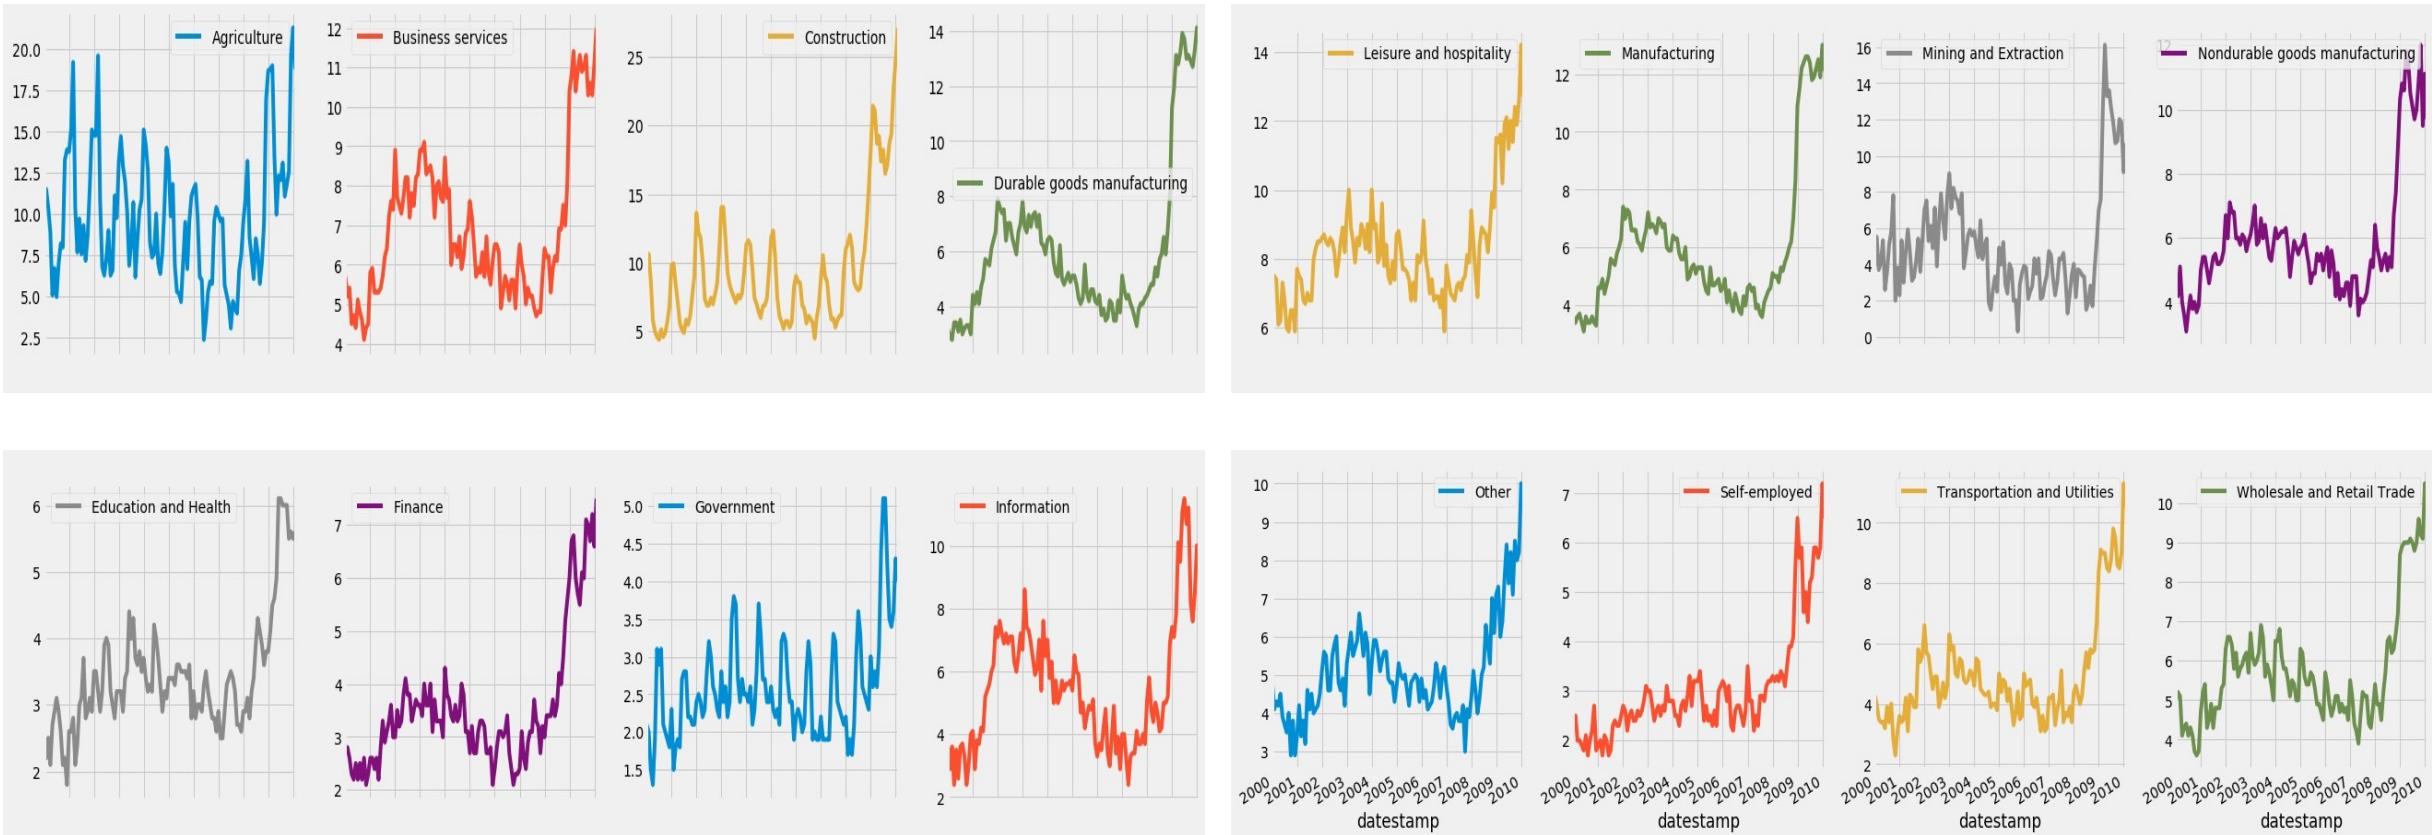
# The Jobs dataset



# Facet plots of the jobs dataset

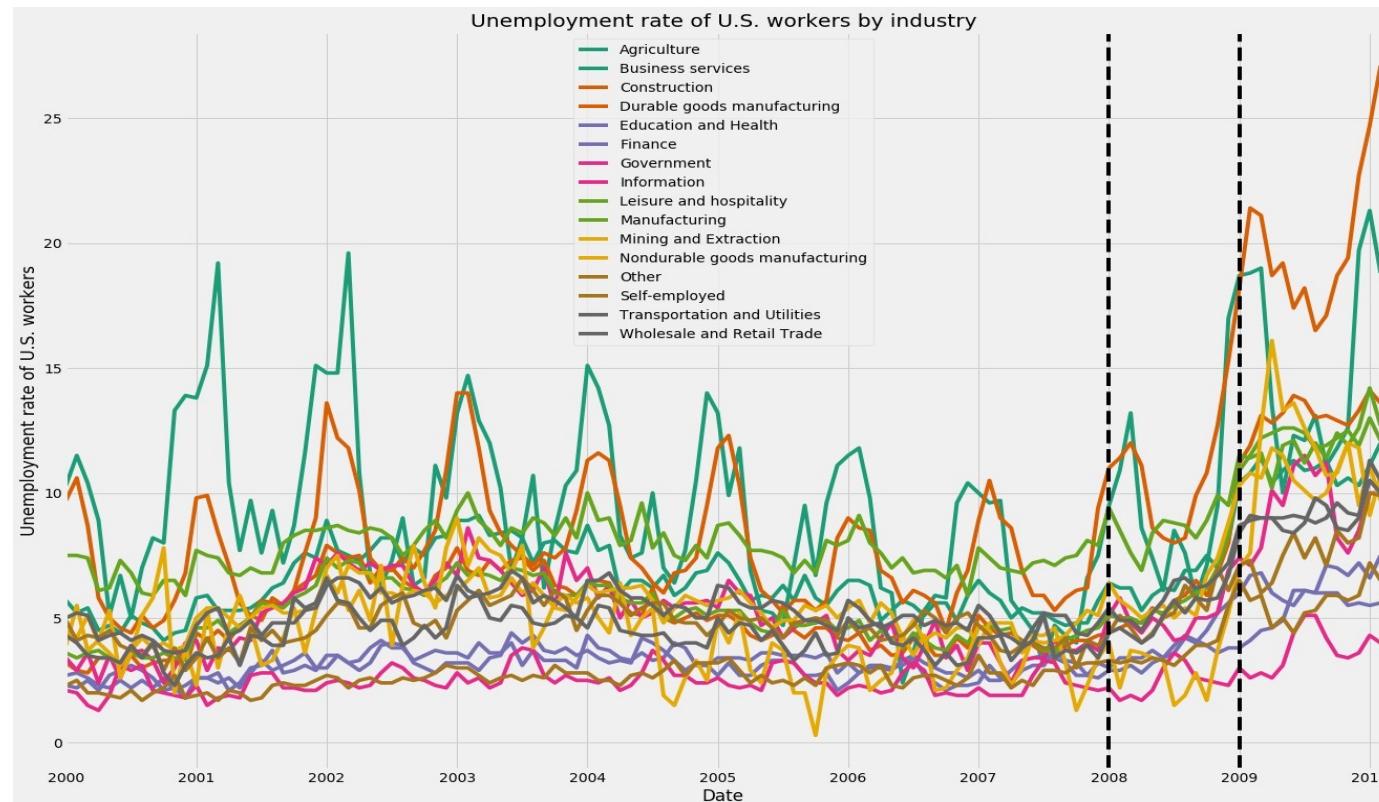
```
jobs.plot(subplots=True,  
          layout=(4, 4),  
          figsize=(20, 16),  
          sharex=True,  
          sharey=False)  
  
plt.show()
```

# Facet plots



# Annotating events in the jobs dataset

```
ax = jobs.plot(figsize=(20, 14), colormap='Dark2')  
  
ax.axvline('2008-01-01', color='black',  
           linestyle='--')  
  
ax.axvline('2009-01-01', color='black',  
           linestyle='--')
```



# Taking seasonal average in the jobs dataset

```
print(jobs.index)
```

```
DatetimeIndex(['2000-01-01', '2000-02-01', '2000-03-01',
 '2000-04-01', '2009-09-01', '2009-10-01',
 '2009-11-01', '2009-12-01', '2010-01-01', '2010-02-01'],
 dtype='datetime64[ns]', name='datestamp',
 length=122, freq=None)
```

```
index_month = jobs.index.month
jobs_by_month = jobs.groupby(index_month).mean()
print(jobs_by_month)
```

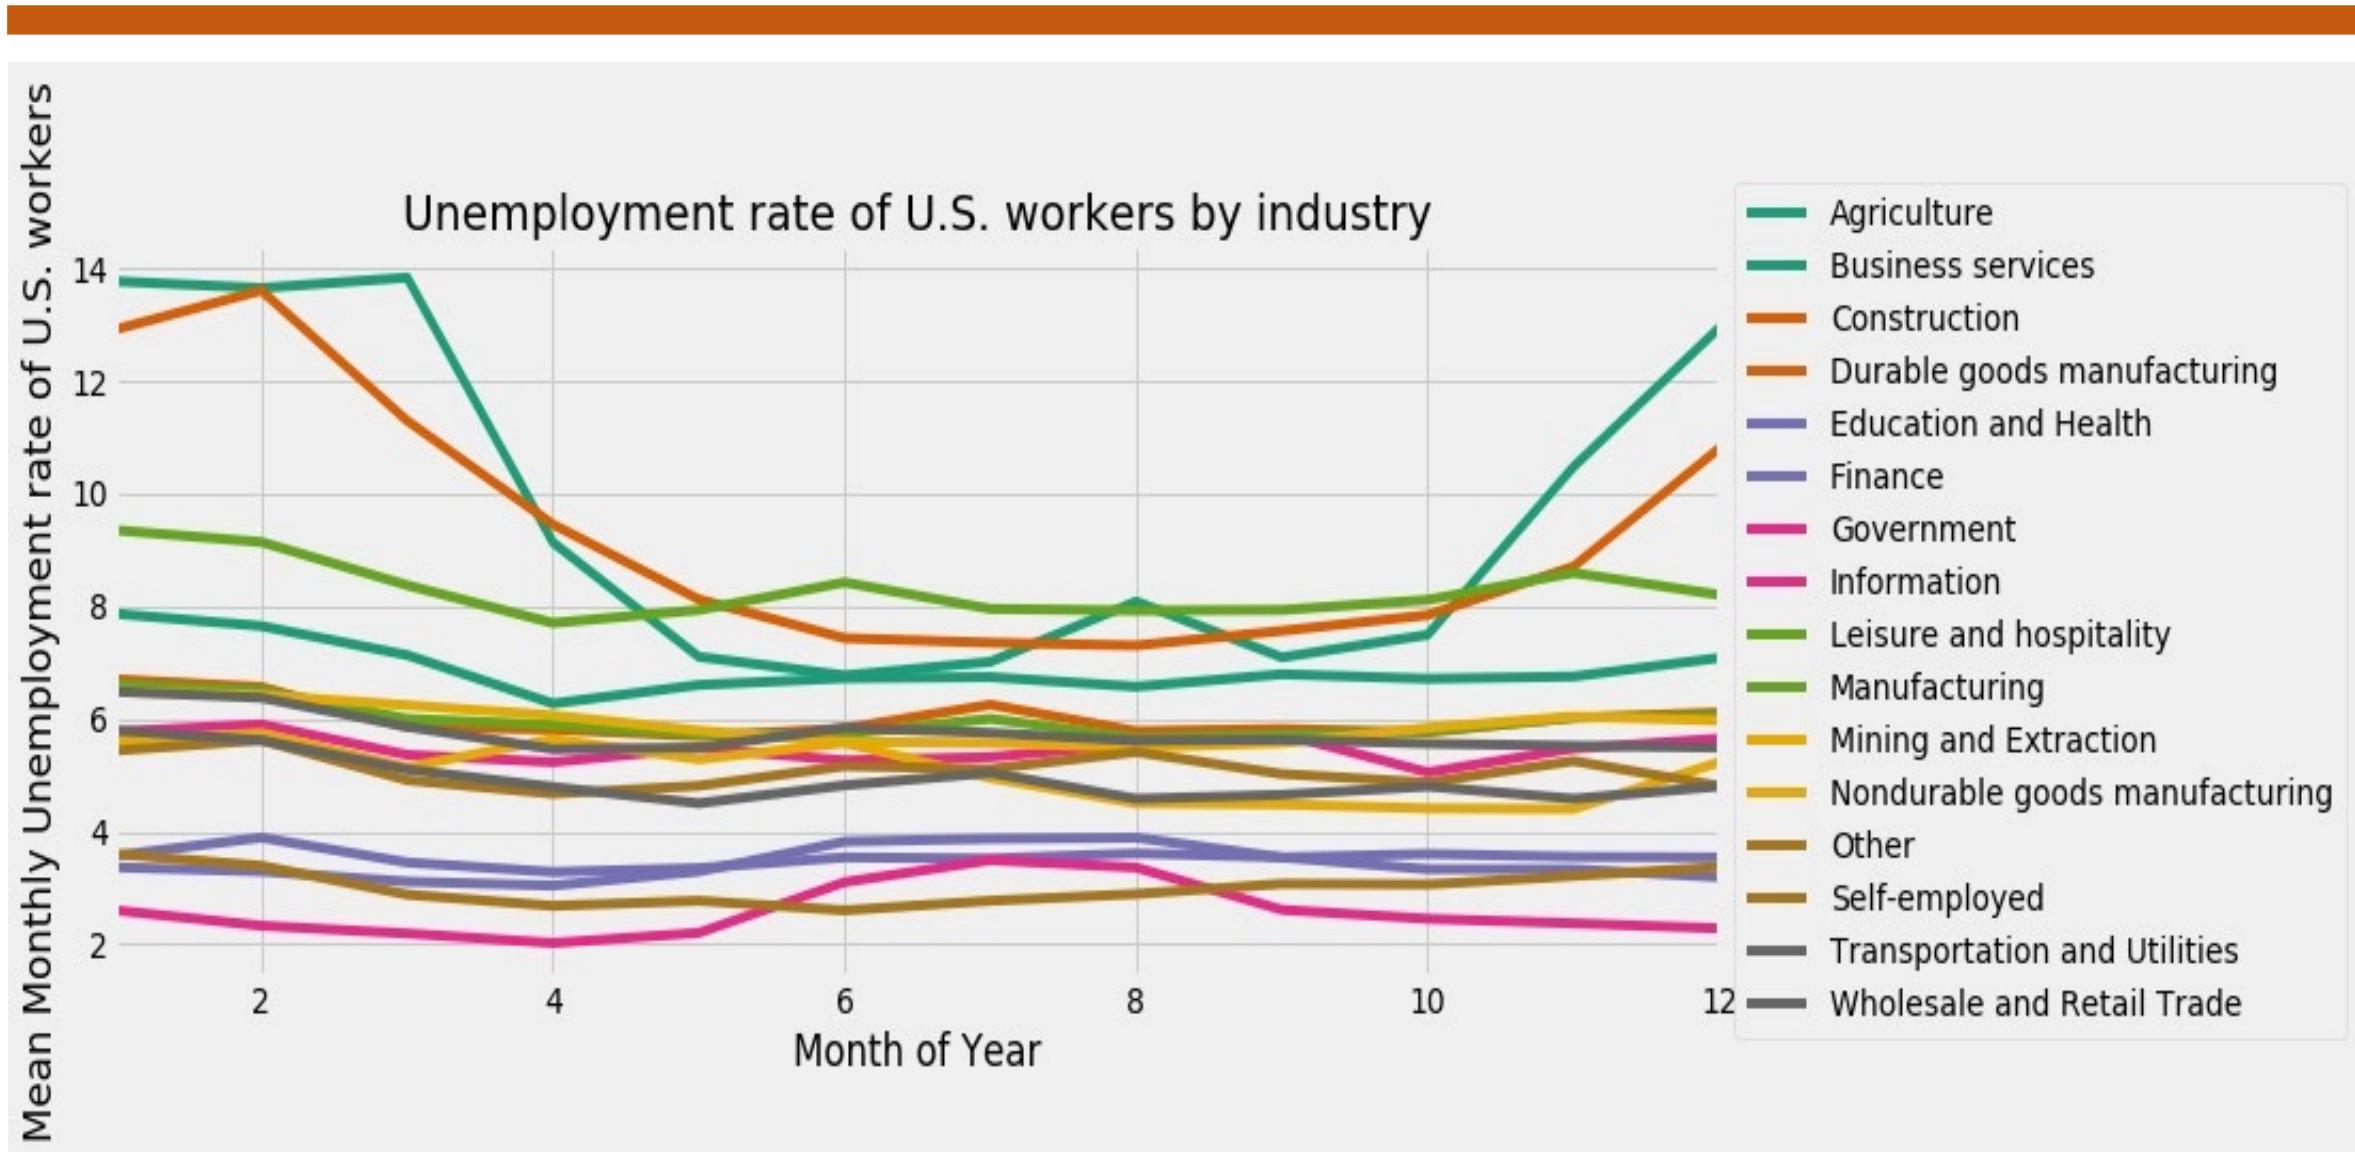
datestamp	Agriculture	Business services	Construction
1	13.763636	7.863636	12.909091
2	13.645455	7.645455	13.600000
3	13.830000	7.130000	11.290000
4	9.130000	6.270000	9.450000
5	7.100000	6.600000	8.120000
...			

# Monthly averages in the jobs dataset

```
ax = jobs_by_month.plot(figsize=(12, 5),  
colormap='Dark2')
```

```
ax.legend(bbox_to_anchor=(1.0, 0.5),  
loc='center left')
```

# Monthly averages in the jobs dataset



# Python dictionaries

```
# Initialize a Python dictionary
my_dict = {}

# Add a key and value to your dictionary
my_dict['your_key'] = 'your_value'

# Add a second key and value to your dictionary
my_dict['your_second_key'] = 'your_second_value'

# Print out your dictionary
print(my_dict)

{'your_key': 'your_value',
 'your_second_key': 'your_second_value'}
```

# Decomposing multiple time series with Python dictionaries

```
# Import the statsmodel library
import statsmodels.api as sm
# Initialize a dictionary
my_dict = {}
# Extract the names of the time series
ts_names = df.columns
print(ts_names)
```

```
['ts1', 'ts2', 'ts3']
```

```
# Run time series decomposition
for ts in ts_names:
    ts_decomposition = sm.tsa.seasonal_decompose(jobs[ts])
    my_dict[ts] = ts_decomposition
```

# Extract decomposition components of multiple time series

```
# Initialize a new dictionary
my_dict_trend = {}

# Extract the trend component
for ts in ts_names:
    my_dict_trend[ts] = my_dict[ts].trend

# Convert to a DataFrame
trend_df = pd.DataFrame.from_dict(my_dict_trend)
print(trend_df)
```

	ts1	ts2	ts3
datestamp			
2000-01-01	2.2	1.3	3.6
2000-02-01	3.4	2.1	4.7
...			

# Trends in Jobs data

```
print(trend_df)
```

datestamp	Agriculture	Business services	Construction
-----------	-------------	-------------------	--------------

2000-01-01	NaN	NaN	NaN
------------	-----	-----	-----

2000-02-01	NaN	NaN	NaN
------------	-----	-----	-----

2000-03-01	NaN	NaN	NaN
------------	-----	-----	-----

2000-04-01	NaN	NaN	NaN
------------	-----	-----	-----

2000-05-01	NaN	NaN	NaN
------------	-----	-----	-----

2000-06-01	NaN	NaN	NaN
------------	-----	-----	-----

2000-07-01	9.170833	4.787500	6.329167
------------	----------	----------	----------

2000-08-01	9.466667	4.820833	6.304167
------------	----------	----------	----------

...

# Plotting a clustermap of the jobs correlation matrix

```
# Get correlation matrix of the seasonality_df DataFrame
trend_corr = trend_df.corr(method='spearmann')

# Customize the clustermap of the seasonality_corr
correlation matrix
fig = sns.clustermap(trend_corr, annot=True, linewidth=0.4)

plt.setp(fig.ax_heatmap.yaxis.get_majorticklabels(),
rotation=0)

plt.setp(fig.ax_heatmap.xaxis.get_majorticklabels(),
rotation=90)
```

# The jobs correlation matrix

